

3DGEER: 3D GAUSSIAN RENDERING MADE EXACT AND EFFICIENT FOR GENERIC CAMERAS

Anonymous authors
Anonymous submission

ABSTRACT

3D Gaussian Splatting (3DGS) (Kerbl et al., 2023) attains a compelling balance between rendering quality and efficiency, yet that efficiency is bought at a geometric price: each 3D Gaussian is replaced by its 2D image-plane projection, an approximation whose error escalates sharply under generic, large field-of-view (FoV) cameras and systematically erodes reconstruction accuracy. Recent extensions relax individual facets of this limitation, but none reconciles *projective exactness* with *real-time efficiency* across general camera models—an unresolved tension that has persisted precisely because the two objectives were treated as competing rather than jointly attainable. We resolve that tension with 3DGEER, a geometrically exact and efficient Gaussian rendering framework that closes this long-standing gap within a single unified formulation. Proceeding from first principles, we derive a closed-form expression for the transmittance integral of Gaussian density along a ray. This expression furnishes exact forward rendering and fully differentiable optimization under arbitrary camera models, thereby eliminating the linearization error at its source rather than attenuating it after the fact. Efficiency is secured through two mutually reinforcing constructions: the *Particle Bounding Frustum (PBF)*, which establishes tight ray–Gaussian association directly in 3D and dispenses with BVH traversal altogether, and the *Bipolar Equiangular Projection (BEAP)*, a unified FoV representation that simultaneously accelerates association and, through more uniform ray sampling, measurably improves reconstruction quality. Across pinhole and fisheye datasets alike, 3DGEER surpasses prior methods on every reported metric, renders $5\times$ faster than existing projective-exact ray-based baselines, and generalizes to wider FoVs unseen during training—collectively establishing a new state of the art in real-time radiance field rendering.

1 INTRODUCTION

Initiated by 3D Gaussian Splatting (3DGS) (Kerbl et al., 2023), volumetric particle rendering methods cast in the splatting paradigm (Yu et al., 2024a; Yan et al., 2024; Lu et al., 2024; Kheradmand et al., 2024; Charatan et al., 2024; Mallick et al., 2024) have become the dominant approach to real-time neural rendering, owing to their combination of high visual fidelity and fast rendering. This efficiency originates in EWA splatting (Zwicker et al., 2001), which invokes a local affine approximation to project each 3D Gaussian as a 2D Gaussian on the image plane. The approximation is faithful within narrow field-of-view (FoV) regimes, but it degrades systematically under wide FoVs: nonlinear distortion drives the true projection away from any symmetric 2D Gaussian, and the resulting error is most acute in the fisheye imagery that pervades robotics and autonomous driving (Rashed et al., 2021; Courbon et al., 2007; Kumar et al., 2023; Zhang et al., 2016; Yogamani et al., 2019; 2024; Sekkat et al., 2022). Recent efforts extend differentiable rendering to specific fisheye camera models (Liao et al., 2024) or to broader camera families (Huang et al., 2024b), yet each remains anchored to a 2D projective-space approximation and therefore stops short of the exact closed-form formulation that a geometrically faithful wide-FoV renderer fundamentally requires.

A distinct line of reasoning questions whether exact rendering is necessary at all: Celarek et al. (2025) contends that scaling the number of Gaussians (up to 1M) suffices to absorb the projective approximation error of splatting. That evidence, however, is confined to small object-centric scenes, and two considerations expose the limits of such a capacity-based remedy. First, in large-scale re-

construction the required Gaussian budget varies drastically across objects and frustums, rendering brute-force scaling prohibitively costly; second, as we demonstrate empirically, the projective error under generic distorted cameras is substantially larger, so that inflating the Gaussian count alone cannot close the quality gap (Fig. 5). Exactness, in short, is not a convenience that additional capacity can supplant—it is a structural requirement.

A parallel line of work confronts volumetric particle rendering directly through a ray-tracing/marching formulation (Yu et al. 2024b, Condor et al. 2025, Mai et al. 2025, Moënnelocoz et al. 2024, Talegaonkar et al. 2025, Gu et al. 2025). Such formulations natively accommodate large-FoV camera models and are, by construction, free of projective approximation error; their decisive obstacle is efficiency—specifically, the cost of associating each rendering ray with the 3D particles that contribute to it. Considerable effort has consequently been directed at accelerating this association—most notably the advanced Bounding Volume Hierarchies (BVH) adopted in EVER (Mai et al. 2025) and 3DGRT (Moënnelocoz et al. 2024)—yet the intrinsic algorithmic complexity of BVH traversal, together with its resistance to parallelization, has kept pure ray-based methods well short of 3DGS-level throughput. The hybrid 3DGUT (Wu et al. 2025) instead couples ray-based rendering with rasterization to accelerate association, but the projective approximation it inherits at the association stage reinstates precisely the geometric error that ray tracing was designed to eliminate, forfeiting projective exactness and inducing grid-line artifacts (Fig. 1 Left). No prior method, therefore, simultaneously secures the geometric exactness of ray-based rendering and the efficiency of splatting—the precise gap this work closes.

We introduce 3DGEER, a 3D Gaussian Exact and Efficient Rendering method that attains *exactness in projective geometry* (formally defined in Sec. A) while retaining *efficiency comparable to splatting*, and delivers both uniformly across generic cameras—perspective as well as a range of fisheye projection models. The framework rests on three tightly coupled contributions that jointly realize this goal: (i) *An exact, numerically stable rendering formulation*. We revisit the heuristic maximum-response function of Gaussians (Keselman & Hebert 2022) from first principles and establish its precise equivalence to projective exactness; from this equivalence we derive a differentiable rendering framework whose forward and backward passes are numerically stable by construction, thereby elevating a previously empirical shortcut to a principled closed-form result. (ii) *Exact yet efficient ray-particle association*. We recast association as a mapping between Camera Sub-Frustums (CSFs) and Particle Bounding Frustums (PBFs)—the frustum-level analogue of the tile-AABB mapping in 3DGS—and derive a closed-form solution for the PBF that tightly bounds each 3D Gaussian, yielding an association that is at once geometrically exact and highly efficient, with no recourse to BVH traversal or screen-space conic surrogates. (iii) *Uniform angular supervision via BEAP*. We sample supervision rays uniformly in a novel Bipolar Equiangular Projection (BEAP) space, which aligns the image-space partitioning with the underlying CSFs and empirically improves reconstruction quality over conventional pinhole and equidistant projections. Together, these three components constitute a single wide-FoV pipeline in which exactness supplies the geometry, PBF the efficiency, and BEAP the supervision.

Extensive experiments spanning both fisheye and pinhole datasets—ScanNet++ (Yeshwanth et al. 2023), Zip-NeRF (Barron et al. 2023), Aria (Lv et al. 2025), and MipNeRF (Barron et al. 2021)—establish that 3DGEER consistently outperforms prior and concurrent methods, delivering PSNR gains of 0.9–4.8 dB under challenging wide-FoV conditions and 0.3–2.0 dB on pinhole views. Notably, even when trained on narrow-FoV data alone, 3DGEER still surpasses splatting-based methods on wide-FoV cameras—direct evidence that its capacity to generalize to extremely wide-FoV rendering follows from projective exactness itself rather than from any FoV-specific tuning.

2 RELATED WORKS

2.1 SPLATTING-BASED GAUSSIAN RENDERING

Most 3DGS variants (Yu et al. 2024a, Zhang et al. 2024, Rota Bulò et al. 2024, Gao et al. 2025, Zhao et al. 2025, Kerbl et al. 2024, Hou et al. 2025, Liu et al. 2025) inherit the original EWA splatting strategy (Zwicker et al. 2001, Ren et al. 2002), projecting each 3D Gaussian onto the image plane and modeling its footprint as a 2D conic. To extend differentiable rendering to large-FoV, distorted camera models, methods such as FisheyeGS (Liao et al. 2024) and GS++ (Huang




Figure 1: **Linear Approximation Error in Ray-Particle Association.** (Left) Grid-line artifacts induced by inaccurate UT-based association. (Right) A diagram contrasting our association with prior schemes: in place of the intermediate conic approximation on which prior methods rely, our method computes the exact bounding structure directly from the true 3D covariance, thereby eliminating the association-stage projective error outright rather than merely attenuating it.

Table 1: **Summary of Particle Rendering Methods w.r.t. Projective Exactness.** The upper block lists splatting-based methods; the lower block lists ray-based approaches, several of which nonetheless retain a projective-approximate association stage (e.g., EWA or the Unscented Transform). Read across the columns, the table isolates the capability that sets our method apart: 3DGEER is the sole entry that jointly satisfies large-FoV support, a high frame rate, and projective exactness in *both* rendering and association.

Method	Render	Particle Asso. Method	Scope	Particle Type	Large FoV	High FPS	Exact Meth.
3DGS [Kerbl et al., 2023]	Splat.	EWA	Img.	3D Gauss.	✗	✓	✗
FisheyeGS [Liao et al., 2024]	Splat.	EWA	Img.	3D Gauss.	✓	✓	✗
GS++ [Huang et al., 2024b]	Splat.	EWA	Img.	3D Gauss.	✗	✓	✗
2DGS [Huang et al., 2024a]	Ray	Splat Bound.	Img.	2D Gauss.	✗	✓	✓
GOF [Yu et al., 2024b]	Ray	EWA	Img.	3D Gauss.	✗	✗	✗
EVER [Mai et al., 2025]	Ray	BVH	Scene	3D Ellip.	✓	✗	✓
3DGRT [Moënné-Loccoz et al., 2024]	Ray	BVH	Scene	3D Gauss.	✓	✗	✓
3DGUT [Wu et al., 2025]	Ray	UT	Img.	3D Gauss.	✓	✓	✗
HTGS [Hahlbohm et al., 2025]	Ray	Splat Bound.	Img.	3D Gauss.	✗	✓	✓
3DGEER (Ours)	Ray	PBF	Frust.	3D Gauss.	✓	✓	✓

[et al., 2024b] estimate the Jacobian of equidistant or spherical projections and correct the 2D mean and covariance of the projected Gaussians accordingly. Because these corrections remain first-order Taylor expansions of a strongly nonlinear projection, the neglected higher-order terms become non-negligible under wide FoVs, producing substantial approximation error and degraded reconstruction quality.

The root cause is fundamentally geometric: under any nonlinear camera model—distorted, or even the ideal pinhole—the true projection of a 3D Gaussian is not a symmetric 2D Gaussian. Any framework that substitutes linear projective geometry for an inherently nonlinear projection therefore incurs an approximation error that no refinement *within* that framework can remove. 3DGEER circumvents this error entirely: it aggregates Gaussian density directly along rays in a volumetric formulation, dispensing with projective approximation altogether.

2.2 RAY-BASED GAUSSIAN RENDERING

[Keselman & Hebert, 2022] posit that a ray interacts with a Gaussian density field at its point of maximum intensity response. Building on this heuristic, 3DGRT [Moënné-Loccoz et al., 2024] constructs a 3D Gaussian ray-tracing framework that samples directly in ray space without explicit projection, thereby sidestepping wide-FoV distortion. Its ray-particle association, however, hinges on costly scene-level BVH partitioning, which imposes substantial overhead and depresses frame rates well below real time.

EVER [Mai et al., 2025] instead replaces Gaussians with constant-density ellipsoids, which admit a closed-form transmittance-weighted ray integral; this substitution, however, both curtails representational capacity relative to full 3D Gaussians and, through its continued reliance on BVH structures, remains runtime-bound. 2DGS [Huang et al., 2024a] substitutes 3D Gaussians with 2D surfels to obtain exact ray-surface intersections. HTGS [Hahlbohm et al., 2025] restores the full Gaussian expression while retaining 2DGS’s ray-splat association; because that association is confined




Figure 2: **Gradient Flow and Canonical Transformation.** (Left) The forward and backward gradient propagation that couples the Gaussian parameters $\mathbf{s}$, $\mathbf{q}$, $\boldsymbol{\mu}$ and the opacity coefficient σ to the transmittance, constituting an exact, fully differentiable chain. (Right) Integration of 3D Gaussian density along a ray; the green box highlights the measure-preserving product of density and ray-segment length that renders the transmittance invariant under the canonical map (Eq. B.5).

to screen space, it cannot guarantee projective exactness under wider FoVs. By contrast, we derive closed-form bounds and a complete rendering pipeline directly in a camera-agnostic angular domain (Sec. D), and we make the computational relationship to the 2DGS and HTGS parameterizations fully explicit in Sec. D.3.

3DGUT (Wu et al. 2025), a follow-up to 3DGRT, recovers 3DGS-level frame rates by adopting splatting-based particle association—but that acceleration rests on the Unscented Transform (UT), a sampling-based approximation of the 2D conic in projective image space, and thereby reinstates the very projective error that ray tracing set out to remove. GOF (Yu et al. 2024b) shares this limitation, applying the EWA construction to fit 2D bounding conics and projective centers for association. Under extreme FoVs, the resulting imprecise AABB estimates and mismatched tile-to-AABB associations manifest as visible grid-line artifacts (Fig. 1 Left).

3DGEER departs decisively from all of these by solving ray-particle association directly at the frustum level in 3D space (Fig. 1 Right)—a formulation that preserves the geometric exactness of ray-based rendering across arbitrary camera models while sustaining rendering speeds on par with 3DGS. Tab. 1 consolidates the representative methods that target projective approximation error and lays bare their differences along the dimensions that matter for exact, real-time, wide-FoV rendering—dimensions on which 3DGEER alone satisfies every criterion.

3 METHOD

We present 3DGEER, a Gaussian rendering framework that delivers geometric exactness and real-time efficiency simultaneously under arbitrary FoV cameras. The design directly targets—and resolves—the two bottlenecks that have prevented prior work from attaining both goals at once, namely the *projection approximation* intrinsic to splatting-based rendering and the *BVH overhead* intrinsic to ray-based rendering, through three mutually reinforcing components:

- **Ray-based rendering formulation** (Sec. 3.1): a closed-form, differentiable solution for integrating Gaussian density along 3D rays that eliminates the linearization error of projection-based splatting at its source.
- **Particle Bounding Frustum (PBF)** (Sec. 3.2): a ray-particle association scheme that supplants BVH traversal with exact frustum-particle tests, securing accuracy and efficiency within a single construction.
- **Bipolar Equiangular Projection (BEAP)** (Sec. 3.3): an image representation that unifies arbitrary FoV camera models, accelerates PBF computation, and improves reconstruction quality through more uniform ray sampling.

The Appendix supplies the supporting formal apparatus: (i) the preliminaries and notation (Sec. B.1); (ii) the complete proof underlying Sec. 3.1 which establishes the equivalence between the maximum-response heuristic and projective exactness (Sec. B); (iii) the full backward-gradient derivations together with their numerical-stability analysis (Sec. C); and (iv) the complete derivation of geometrically exact association (Sec. D). Collectively, these establish that every empirical gain reported below rests on a derivation carried through in closed form.

3.1 RENDERING WITH EXACT PROJECTIVE GEOMETRY

The central obstacle to geometrically exact Gaussian rendering is obtaining a closed-form expression for the integral of density and color along a ray. Our strategy dissolves this obstacle by mapping each anisotropic Gaussian in world space into a canonical coordinate system in which every Gaussian is isotropic; under this map the transmittance integral reduces to a measure-preserving change of variables and admits a compact closed-form solution, so that exactness follows by construction rather than through numerical approximation.

Canonical transformation. For each Gaussian, parameterized by its mean μ , rotation R , and scaling S , we define the mapping

$$\begin{bmatrix} \mathbf{x} \\ 1 \end{bmatrix} = \begin{bmatrix} RS & \mu \\ \mathbf{0} & 1 \end{bmatrix} \begin{bmatrix} \mathbf{u} \\ 1 \end{bmatrix}, \quad (1)$$

which sends a canonical coordinate $\mathbf{u}$ to world coordinates $\mathbf{x}$. Since $\Delta\mathbf{x} = RS\Delta\mathbf{u}$, the Jacobian determinant is absorbed exactly into the density measure (Fig. 2 Right), whereupon every Gaussian reduces to the single shared isotropic form

$$\mathcal{G}_{\mathbf{I},0}(\mathbf{u}) = \frac{1}{\rho} \exp\left(-\frac{1}{2}\|\mathbf{u}\|^2\right). \quad (2)$$

Closed-form transmittance. For a ray $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$, its canonical-space image is $\mathbf{r}_u(t) = \mathbf{o}_u + t\mathbf{d}_u$, with

$$\mathbf{o}_u = S^{-1}R^\top(\mathbf{o} - \mu), \quad \mathbf{d}_u = S^{-1}R^\top\mathbf{d}. \quad (3)$$

The accumulated transmittance is then

$$T(\mathbf{o}, \mathbf{d}) = \sigma \int_{t \in \mathbb{R}} \mathcal{G}_{\mathbf{I},0}(\mathbf{o}_u + t\mathbf{d}_u) dt = \sigma \exp\left(-\frac{1}{2}D_{\mu,\Sigma}^2(\mathbf{o}, \mathbf{d})\right), \quad (4)$$

where

$$D_{\mu,\Sigma}^2(\mathbf{o}, \mathbf{d}) = \frac{\|\mathbf{o}_u \times \mathbf{d}_u\|^2}{\|\mathbf{d}_u\|^2} \quad (5)$$

is the squared perpendicular (Mahalanobis) distance from the ray to the Gaussian center.

Discussion. The transmittance thus depends solely on the ray’s distance to the Gaussian in canonical space—a closed-form solution that is at once efficient and geometrically exact. This expression coincides with the “maximum-response” heuristic invoked by prior work (Keselman & Hebert, 2022; Moënné-Loccoz et al., 2024); our derivation, however, exposes the projective exactness underlying it (Sec. B), elevating what has served as an empirical shortcut to a principled, first-principles result. Because we adhere to this closed-form derivation and its exact backward gradient chain (Fig. 2 Left; Sec. C), we further avoid—by construction, not by tuning—the numerical instabilities that beset prior ray-based work (Yu et al., 2024b), obviating redundant post-processing such as the filtering of degenerate Gaussians.

3.2 RAY-PARTICLE ASSOCIATION WITH EXACT PROJECTIVE GEOMETRY

In any ray-based Gaussian renderer, both accuracy—most conspicuously under large-FoV camera models—and runtime efficiency are governed by the ray-particle association strategy, a component that prior work has largely




Figure 2: PBF, CSF Association. (Left) A PBF is

left underexamined. To determine which Gaussians contribute to each ray, we reduce association to a mapping between Camera Sub-Frustums (CSFs) and Particle Bounding Frustums (PBFs), where the PBF is the tight frustum enclosing each Gaussian at a prescribed distance from its center (Fig. 3). A Gaussian is assigned to a CSF precisely when its PBF intersects that CSF, and every Gaussian associated with a CSF is treated as a contributor to all rays traversing it. This frustum-level formulation is the crux of reconciling exactness with efficiency: it operates in 3D on the true Gaussian geometry, yet collapses to inexpensive angular-bound comparisons.

To compute a PBF efficiently, we first introduce two spherical angles (see Fig. E.1) in camera space:

$$\theta = \arctan \left(\frac{\mathbf{d}_{c,x}}{\mathbf{d}_{c,z}} \right), \quad \phi = \arctan \left(\frac{\mathbf{d}_{c,y}}{\mathbf{d}_{c,z}} \right), \quad (6)$$

where $(\mathbf{d}_{c,x}, \mathbf{d}_{c,y}, \mathbf{d}_{c,z})$ is the normalized direction vector toward the camera- (view-)space point $\mathbf{d}_c$. The angles θ and ϕ denote the ray's incidence angles in the horizontal (xz) and vertical (yz) planes, respectively, each confined to $[-\frac{\pi}{2}, \frac{\pi}{2}]$. A PBF is then specified by four planes—two aligned with the θ bounds and two with the ϕ bounds (Fig. 3):

$$\mathbf{g}_{\theta_{1,2}} = (-1, 0, \tan \theta_{1,2}, 0)^\top, \quad \mathbf{g}_{\phi_{1,2}} = (0, -1, \tan \phi_{1,2}, 0)^\top, \quad (7)$$

with $\theta_1 < \theta < \theta_2$ and $\phi_1 < \phi < \phi_2$ delimiting the angular extent.

Given a camera with extrinsic matrix $V = [R_c | \mathbf{t}_c] \in \text{SE}(3)$ and a canonical-to-world transformation matrix H (Eq. B.4), we transport these view-space angular limits into canonical space by applying the inverse transpose of the camera-to-canonical matrix $(VH)^{-1}$ to the plane parameters, thereby obtaining the constraints that define the PBF. Taking the θ bounds as the representative case, the canonical plane reads:

$$\mathbf{g}_{u_{1,2}} = (VH)^\top \mathbf{g}_{\theta_{1,2}} = -(VH)_0 + \tan \theta_{1,2} (VH)_2, \quad (8)$$

where $(VH)_i$ denotes the i -th row of the matrix, written as a column vector.

For a Gaussian ellipsoid defined by its λ -standard-deviation contour, the angular bounds of the local frustum follow from enforcing tangency between the transformed planes and the ellipsoidal contour—a condition that reduces to the quadratic equations:

$$((\lambda^2, \lambda^2, \lambda^2, -1) \circ \mathbf{g}_u^\top) \cdot \mathbf{g}_u = 0, \quad (9)$$

where $\circ$ is the Hadamard (element-wise) product and λ governs the spread of the contour (we set $\lambda = 3$ throughout). The same constraint applied to $\mathbf{g}_v^\top$ yields the corresponding bounds $\phi_{1,2}$, and the roots of these quadratics supply the tangent values of the angular bounds $[\theta_1, \theta_2]$ and $[\phi_1, \phi_2]$ that delimit the PBF. Crucially, solving these quadratics under the canonical-space constraint operates directly on the Gaussian's true camera-space 3D covariance Σ_c and 3D mean $\boldsymbol{\mu}_c$, never on a projected surrogate—precisely the property that renders the resulting bounds exact. Concretely, for the θ bounds, $c = \tan \theta_{1,2}$ satisfies:

$$\mathcal{T}_{22}c^2 - 2\mathcal{T}_{02}c + \mathcal{T}_{00} = 0, \quad \text{where} \quad \mathcal{T}_{ij} = \lambda^2 \Sigma_c^{i,j} - (\boldsymbol{\mu}_c \boldsymbol{\mu}_c^\top)^{i,j}. \quad (10)$$

This formulation is inexpensive to solve (Sec. D.1) and, precisely because it operates on the true 3D covariance, dispenses at once with the intermediate conic approximation of EWA and UT and with the screen-space constraints that cap the admissible FoV in HTGS and 2DGS—the two mechanisms through which competing schemes forfeit either exactness or wide-FoV coverage. Our construction relinquishes neither.

3.3 BIPOLAR EQUIANGULAR PROJECTION (BEAP)

The BEAP image representation is motivated by three complementary properties that jointly justify its role in the pipeline. First, it represents large-FoV imagery without any loss of FoV. Second, it endows image tiles and their corresponding CSFs with an identical parameterization in terms of two θ and two ϕ angles, which renders PBF association highly efficient and maps naturally onto GPU parallelism. Third, the induced ray sampling is markedly more uniform in 3D than that of conventional projective rendering, and this uniformity empirically improves both training efficiency and reconstruction quality. Taken together, these properties establish BEAP as the connective representation that binds the exact rendering and PBF association components into a single, coherent wide-FoV pipeline.

Concretely, we sample rays uniformly in the spherical angular coordinates (θ, ϕ) (Eq. 6) and map them back onto a discrete $w \times h$ image grid through a linear transformation—a scaling followed by a center shift:

$$(x, y) = \left\lfloor \frac{w\theta}{\text{FoV}_x} + \frac{w+1}{2}, \frac{h\phi}{\text{FoV}_y} + \frac{h+1}{2} \right\rfloor, \quad (11)$$

where $\text{FoV}_x, \text{FoV}_y$ denote the camera’s horizontal and vertical fields of view. As Fig. 4 shows, uniform sampling in (θ, ϕ) achieves substantially more balanced spatial coverage of the view frustum than either standard alternative: projective sampling over-allocates rays to the periphery, whereas equidistant projection over-concentrates them near the image center. BEAP thereby eliminates the sampling imbalance that both alternatives leave uncorrected.

4 EXPERIMENTS

We evaluate 3DGEER across multiple datasets spanning both pinhole and fisheye camera models, so that its behavior is characterized over the entire range of FoVs it is designed to handle rather than within a single favorable regime. Unless otherwise stated, all experiments adopt an RTX 4090 GPU with 64 GB of RAM as the default configuration; full implementation details appear in Sec. G.

Datasets. We assess performance on generic distorted cameras—the regime in which projective error is most consequential—across all scenes of Fisheye ZipNeRF (Barron et al. 2023), five scenes of ScanNet++ (Yeshwanth et al. 2023), and two of Aria (Lv et al. 2025). The ZipNeRF and ScanNet++ cameras carry a 180° diagonal FoV, and Aria a 110° circular FoV. For the pinhole regime we use Pinhole ZipNeRF and MipNeRF360 (Barron et al. 2021). Collectively, these datasets exercise the method across narrow-, wide-, and cross-camera conditions, so that the reported gains reflect breadth of applicability rather than the idiosyncrasies of a single benchmark.

Evaluation Metrics and Protocol. We report PSNR $\uparrow$, SSIM $\uparrow$, LPIPS $\downarrow$, runtime (FPS $\uparrow$), and Gaussian count $\downarrow$. To eliminate representation bias, every method’s outputs are reprojected to the native image space and evaluated over the full FoV after 30k training iterations, so that all methods are compared on identical, camera-native ground—a protocol that keeps the comparison strictly fair across differing internal representations.

4.1 SCALING GAUSSIANS VS. PROJECTIVE EXACTNESS IN WIDE-FOV CAMERAS

Celarek et al. (2025) contend that merely scaling the number of Gaussians (up to 1M) suffices to absorb the linear approximation gap between splatting- and ray-based methods. That claim, however, rests on small object-centric scenes and leaves the challenge of wide-FoV distortion untouched. We test it head-on by scaling the splatting-based fisheye method FisheyeGS on ScanNet++, lowering its densification threshold to drive the model to 3.6–5M Gaussians (on an H200).

The result, in Fig. 5 is unambiguous. At matched Gaussian counts, the FisheyeGS curve (red) and the 3DGEER curve (blue) remain separated by a clear PSNR gap across the entire densification




Figure 4: **Ray Distributions under Different Projections.** Pixels from three imaging spaces are projected onto the unit sphere to compare their ray distributions; of the three, BEAP alone attains near-uniform angular coverage, whereas the projective and equidistant alternatives concentrate rays at the periphery and at the center, respectively.




Figure 5: **PSNR $\uparrow$ vs. Gaussian Count.** Enlarging the Gaussian budget fails to close the projective-exactness gap: brute-force scaling of the splatting baseline plateaus well below 3DGEER at every matched count, confirming that the gap originates in exactness rather than in capacity (full metrics in Fig. K.3 and Tab. K.3).



Figure 6: **Extreme-FoV Generalization.** With fewer Gaussians, 3DGEER reconstructs complete scenes at wide FoVs unseen during training, whereas FisheyeGS degrades severely toward the periphery under identical conditions.

schedule (4k–30k iterations, beyond which densification saturates)—a gap attributable to projective exactness rather than to model capacity. The scaled-up FisheyeGS (light blue) corroborates this diagnosis: despite a large increase in memory and computation, its PSNR plateaus at 29.3, whereas 3DGEER reaches 32.1 with only 550–700K Gaussians—roughly an order of magnitude fewer. Scaling capacity, therefore, cannot substitute for exactness; full statistics appear in Tab. K.3

4.2 COMPREHENSIVE EVALUATION ON ALL-FOV NOVEL VIEW SYNTHESIS

We first benchmark 3DGEER against representative baselines on both fisheye and pinhole datasets, complementing each comparison with analysis tailored to the dataset in question.

ScanNet++. 3DGEER outperforms every baseline on all reported metrics (Tab. 2). Decomposing the FoV into central and peripheral regions (protocol in Sec. G) pinpoints the origin of this margin: the competing methods render the center acceptably but degrade toward the periphery—precisely where projective error concentrates—whereas 3DGEER sustains high fidelity across the entire FoV, with sharper boundaries and fewer artifacts (Fig. K.4). The advantage proves robust to a substantial reduction in training FoV: trained on central crops alone, 3DGEER still surpasses FisheyeGS in every region, which establishes that the gain derives from the exact formulation itself rather than from access to wide-FoV supervision.

This robustness extends to extreme-FoV generalization (Fig. 6). Trained at the native 180° diagonal FoV and tested at a substantially wider 180° circular FoV, 3DGEER reconstructs the complete scene with fewer Gaussians, whereas splatting-based methods such as FisheyeGS break down at the periphery (further extreme-FoV examples in Figs. K.1–K.2).

ZipNeRF. Because this dataset supplies frame-wise aligned pinhole and fisheye views, we go beyond standard fisheye evaluation and probe *cross-camera generaliza-*

Table 4: **Quantitative Results on the MipNeRF360 Dataset.** Alongside quality metrics, FPS is measured on an RTX 4090 as the default benchmark; figures reproduced from the original works are marked $\dagger$ (RTX 6000 Ada) and $\ddagger$ (RTX 5090). 3DGEER attains the best PSNR and LPIPS while sustaining splatting-comparable frame rates—over 5 $\times$ faster than the projection-cost ray-based base-

Table 2: **Comparison on the ScanNet++ Dataset.** Each method is trained on either the full FoV or the central region alone, then evaluated on the full, central, and peripheral regions to expose robustness under varying distortion. Region-wise metrics are pixel-mask weighted; region-wise LPIPS is omitted because it is not separable by region (Tab. K.2 reports the full statistics). 3DGEER attains the best full-FoV PSNR, SSIM, and LPIPS; more tellingly, when trained on central crops alone it still surpasses FisheyeGS across every region, including the periphery where projective error concentrates.

Method	Training FoV	Training Space	Full FoV			Central		Peripheral	
			PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	PSNR $\uparrow$	SSIM $\uparrow$	PSNR $\uparrow$	SSIM $\uparrow$
3DGS	Central	Persp.	-	-	-	31.26	0.948	-	-
FisheyeGS	Full	Eq.	27.81	0.946	0.139	32.44	0.956	23.28	0.914
EVER	Full	KB	29.47	0.924	0.167	29.93	0.924	28.72	0.925
3DGUT	Full	Eq.	30.64	0.944	0.150	31.87	0.945	28.84	0.937
3DGEER (Ours)	Central	BEAP	29.93	0.949	0.130	32.86	0.954	26.64	0.930
3DGEER (Ours)	Full	BEAP	31.50	0.953	0.126	32.64	0.955	28.94	0.945

Table 3: **Comparison on the ZipNeRF Dataset.** A cross-camera generalization protocol: each method is trained on either Fisheye (FE, 1/8 resolution) or Pinhole (PH, 1/4 resolution) data and evaluated separately on both the Fisheye and Pinhole test sets (Tab. K.9 reports the full statistics). 3DGEER leads on PSNR, SSIM, and LPIPS in nearly every train/test combination, and its margin widens precisely in the most demanding pinhole-train/fisheye-test setting—the case that penalizes projective approximation most sharply.

Train Data Test Data Method	1/8 FE		1/8 FE		1/8 FE		1/4 PH		1/4 PH		1/4 PH	
	FE	PH	FE	PH	FE	PH	FE	PH	FE	PH	FE	PH
	PSNR $\uparrow$		SSIM $\uparrow$		LPIPS $\downarrow$		PSNR $\uparrow$		SSIM $\uparrow$		LPIPS $\downarrow$	
FisheyeGS	23.18	26.44	0.858	0.868	0.211	0.239	19.43	26.61	0.791	0.889	0.247	0.212
EVER	25.17	26.14	0.880	0.851	0.153	0.237	22.90	26.45	0.835	0.862	0.207	0.222
3DGUT	24.77	25.59	0.879	0.804	0.183	0.324	18.61	25.96	0.662	0.841	0.300	0.270
3DGEER (Ours)	26.24	27.62	0.897	0.888	0.140	0.214	23.39	27.61	0.846	0.863	0.209	0.254

tion—training on fisheye and testing on pinhole, and vice versa. 3DGEER leads across metrics in Tab. 3, exploiting the large-FoV training data that only geometrically exact rendering renders usable; even when trained on pinhole data alone, it still delivers clear PSNR gains on both test views. The baselines fare markedly worse under domain shift: FisheyeGS is competitive only in the pinhole-to-pinhole case, and 3DGUT collapses on fisheye after pinhole training. Only EVER and 3DGEER generalize gracefully across cameras—a direct consequence of their shared geometric exactness—yet 3DGEER dominates EVER in every setting. Qualitatively (Fig. 7), 3DGEER recovers finer detail and resists peripheral artifacts, most visibly in the hardest pinhole-to-fisheye regime.

Finally, we confirm that the gains under wide FoVs are not purchased at the expense of the standard narrow-FoV regime, evaluating 3DGEER on **MipNeRF360** (Tab. 4). 3DGEER sets a new state of the art across the reported metrics while occupying a point on the quality–efficiency frontier that no prior method attains: it renders over $5\times$ faster than the exact ray-marching/tracing methods EVER and 3DGRT—establishing it as the only projective-exact method with runtime comparable to 3DGS—and, in the same configuration, surpasses the projective-approximation baselines 3DGS and 3DGUT on PSNR, SSIM, and LPIPS. Exactness and speed, in short, are secured jointly rather




Figure 7: **Qualitative Results on ZipNeRF under Cross-Camera Generalization.** The gap is most pronounced in the challenging Pinhole (PH)-train / Fisheye (FE)-test setting, where 3DGEER retains sharp structure while the baselines break down—the visual counterpart to its quantitative lead in Tab. 3.

Table 5: **Training-Time Transmittance Replacement on ScanNet++.** Substituting a splatting-based transmittance inflates the Gaussian count—densification straining to compensate for the larger approximation error—yet a perceptual gap to our projective-exact formulation persists nonetheless, direct evidence that the error cannot be densified away (Tab. K.7 reports the full statistics).

Transmittance Method	Full FoV			Central		Peripheral		Avg. Gaussian Number (k) ↓
	PSNR ↑	SSIM ↑	LPIPS ↓	PSNR ↑	SSIM ↑	PSNR ↑	SSIM ↑	
3DGS Splats	22.86	0.878	0.177	27.16	0.884	18.56	0.855	1396.4
FisheyeGS Splats	27.90	0.948	0.141	32.49	0.957	23.36	0.917	920.6
Projective-Exact (Ours)	31.50	0.953	0.126	32.64	0.955	28.94	0.945	591.8

than traded off. Additional quantitative and qualitative results on **Aria**, an egocentric fisheye dataset (moderate 110° circular FoV) captured under low light, appear in Sec. J.

4.3 COMPREHENSIVE RUNTIME ANALYSIS

We next dissect runtime on the RTX 4090, comparing 3DGEER against the splatting-based methods 3DGS (on **MipNeRF360**) and FisheyeGS (on **ScanNet++**). Although ray-based rendering is marginally more expensive in the rendering stage itself, PBF’s efficient ray-particle association sharply curtails Gaussian duplication and sorting overhead (stage-wise timings in Tab. K.4), so that end-to-end efficiency matches that of highly optimized splatting pipelines—the incremental rendering cost is fully offset upstream.

The mechanism is the tightness of the PBF bounds: each tile is associated with only ~ 475 Gaussians on average, 3–5× fewer than under EWA or UT (Tab. K.6), and this reduction translates directly into the 2.5–5× speedup measured at the association stage (Tab. K.5). Projection-based schemes (EWA, UT), by contrast, conservatively overestimate the screen-space extent, incurring excessive overdraw and inefficient memory access; even with SnugBox acceleration (Hanson et al. 2025) they remain hostage to miscentered density approximations that frequently aggravate artifacts (Fig. 1). The advantage of PBF is therefore not confined to speed: by operating on the true 3D covariance, it simultaneously preserves geometric accuracy and numerical stability and removes projective approximation error outright. Tab. K.6 makes the point concrete—even when rendering Gaussian fields trained under other association schemes, PBF remains artifact-free, which positions it as a drop-in association module that improves fields reconstructed by other methods.




Figure 8: **Rendering-Time Transmittance Replacement.** Swapping the projective-exact transmittance for the splatting approximation at test time alone—with every training-time component held exact—induces mismatched culling, thereby isolating the rendering formulation itself as the cause of the artifacts.

Table 6: **BEAP vs. Alternative Supervision Schemes.** With the projective-exact integral and PBF association held fixed and only the supervision sampling varied, BEAP outperforms both perspective and equidistant sampling on every metric on ScanNet++, thereby isolating the supervision distribution as the source of the gain (Tab. K.8 reports the full statistics).

Method	Training FoV	Training Space	Full FoV			Central		Peripheral	
			PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	PSNR $\uparrow$	SSIM $\uparrow$	PSNR $\uparrow$	SSIM $\uparrow$
3DGEER w/o BEAP	Central	Persp.	29.84	0.943	0.131	32.21	0.951	26.23	0.915
3DGEER w/o BEAP	Full	Persp.	21.11	0.853	0.300	21.32	0.855	20.46	0.846
3DGEER w/o BEAP	Full	Eq.	31.05	0.948	0.135	32.21	0.952	28.56	0.935
3DGEER (Ours)	Full	BEAP	31.50	0.953	0.126	32.64	0.955	28.94	0.945

4.4 LOCAL AFFINE APPROXIMATION VS. PROJECTIVE-EXACT INTEGRAL

To isolate the contribution of the projective-exact transmittance formulation (Eq. 4), we replace it with the widely used local affine splat (the 2D conic of 3DGS and FisheyeGS) and examine the effect at both test and training time.

Rendering-time approximation. Substituting the splatting approximation for the transmittance at test time only—while every projective-exact component is retained during training—cleanly exposes the divergence between the two rendering formulations. In this configuration the transmittance rests on the splat approximation while association stays exact, and the resulting inconsistency produces mismatched culling (Fig. 8). The experiment magnifies, and thereby explains, the gridline artifact reported for 3DGUT (Fig. 1), which arises from the mirror-image mismatch—ray-based transmittance paired with conic-approximate association.

Training-time approximation. When the splatting-based transmittance is used during training as well, with association held projective-exact (Tab. 5 ScanNet++), the optimizer and densification partially absorb both the approximation error and the cross-view inconsistency that splatting introduces. The elevated loss drives stronger gradients throughout training and inflates the Gaussian budget by 1.6–2.4 $\times$. Even so, the perceptual gap to the exact formulation survives this compensation—conclusive evidence that densification alone cannot correct the approximation error, but only mask a fraction of it at greater cost.

4.5 BEAP ON/OFF STUDY

We isolate the effect of the supervision distribution by holding the projective-exact integral and PBF association fixed and varying only the ray-sampling scheme used for supervision (pinhole, equidistant, and BEAP) on ScanNet++. BEAP supervision is the best of the three on every metric (Tab. 6 Fig. K.5): its more uniform sampling supplies smoother, more consistent gradients and, consequently, better convergence on fine detail.

The effect sharpens under pinhole cameras, where—at fixed resolution—how information is *distributed* across the image matters more than the raw FoV. Trained in the perspective domain, the full FoV can even trail the central region (Tab. 6 rows 1–2): large-FoV perspective images (Fig. K.2 Left) sample angle highly non-uniformly, concentrating informative content at the center while the less informative periphery dominates the image area and is disproportionately oversampled. BEAP’s

uniform angular sampling removes precisely this imbalance, which is why its advantage widens exactly where the imbalance is most severe.

5 CONCLUSIONS

We have presented 3DGEER, a complete closed-form solution for projective-exact volumetric Gaussian rendering derived from first principles. Around this formulation we constructed an efficient PBF-based ray-particle association that sustains high-speed rendering with no sacrifice of projective exactness, together with a BEAP image representation that delivers uniform, effective color supervision under wide-FoV cameras. The three components are mutually reinforcing: the exact ray integral supplies the geometry, PBF the efficiency, and BEAP the supervision that lets both operate across the full FoV. Extensive experiments spanning pinhole and fisheye datasets establish that 3DGEER consistently outperforms state-of-the-art methods, with its largest margins in the challenging wide-FoV regimes where projective approximation error dominates. In doing so it resolves the long-standing tension between exactness and efficiency in real-time radiance field rendering—and resolves it within a single unified framework rather than through a regime-specific compromise.

ETHICS STATEMENT

This work builds exclusively on publicly available datasets (ScanNet++, ZipNeRF, Aria, and MipNeRF) and involves no human subjects, private data, or personally identifiable information. The methods developed here are intended to advance 3D scene reconstruction and rendering research, with prospective applications in XR, robotics, and autonomous driving. We foresee no direct harmful consequences, while acknowledging that any misuse of high-fidelity 3D reconstruction techniques could raise privacy concerns; to mitigate this, we confine our experiments to research datasets that are widely adopted within the vision community.

REPRODUCIBILITY STATEMENT

We have taken substantial measures to ensure reproducibility. All mathematical derivations and proofs are supplied in the Appendix (Secs. B, C and D), and every dataset used in our experiments is publicly available (ScanNet++, ZipNeRF, Aria, and MipNeRF). Preprocessing, training protocols, and evaluation metrics are documented in full in Sec. G. To further facilitate reproduction, our implementation will be released with complete training and evaluation scripts upon publication.

REFERENCES

- Jonathan T. Barron, Ben Mildenhall, Matthew Tancik, Peter Hedman, Ricardo Martin-Brualla, and Pratul P. Srinivasan. Mip-nerf: A multiscale representation for anti-aliasing neural radiance fields. In *IEEE International Conference on Computer Vision (ICCV)*, 2021.
- Jonathan T. Barron, Ben Mildenhall, Dor Verbin, Pratul P. Srinivasan, and Peter Hedman. Zip-nerf: Anti-aliased grid-based neural radiance fields. In *IEEE International Conference on Computer Vision, (ICCV)*, 2023.
- Adam Celarek, George Kopanas, George Drettakis, Michael Wimmer, and Bernhard Kerbl. Does 3d gaussian splatting need accurate volumetric rendering? In *Computer Graphics Forum (CGF)*, 2025.
- David Charatan, Sizhe Lester Li, Andrea Tagliasacchi, and Vincent Sitzmann. pixelsplat: 3d gaussian splats from image pairs for scalable generalizable 3d reconstruction. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition (CVPR)*, pp. 19457–19467, 2024.
- Jorge Condor, Sébastien Speierer, Lukas Bode, Aljaz Bozic, Simon Green, Piotr Didyk, and Adrián Jarabo. Don’t splat your gaussians: Volumetric ray-traced primitives for modeling and rendering scattering and emissive media. *ACM Transactions on Graphics (TOG)*, 2025.
- Jonathan Courbon, Youcef Mezouar, Laurent Eckt, and Philippe Martinet. A generic fisheye camera model for robotic applications. In *IEEE/RSJ International Conference on Intelligent Robots and Systems, (IROS)*, pp. 1683–1688, 2007.
- Daniel Duckworth, Peter Hedman, Christian Reiser, Peter Zhizhin, Jean-François Thibert, Mario Lucic, Richard Szeliski, and Jonathan T. Barron. SMERF: streamable memory efficient radiance fields for real-time large-scene exploration. *ACM Transactions on Graphics (TOG)*, 2024.
- Zhongpai Gao, Benjamin Planche, Meng Zheng, Anwesa Choudhuri, Terrence Chen, and Ziyang Wu. 6dgs: Enhanced direction-aware gaussian splatting for volumetric rendering. In *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025.
- Chun Gu, Xiaofei Wei, Zixuan Zeng, Yuxuan Yao, and Li Zhang. Irgs: Inter-reflective gaussian splatting with 2d gaussian ray tracing. In *Proceedings of the Computer Vision and Pattern Recognition Conference (CVPR)*, pp. 10943–10952, 2025.
- Yuliang Guo, Sparsh Garg, S. Mahdi H. Miangoleh, Xinyu Huang, and Liu Ren. Depth any camera: Zero-shot metric depth estimation from any camera. *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.

-
- Florian Hahlbohm, Fabian Friederichs, Tim Weyrich, Linus Franke, Moritz Kappel, Susana Castillo, Marc Stamminger, Martin Eisemann, and Marcus Magnor. Efficient perspective-correct 3d gaussian splatting using hybrid transparency. In *Computer Graphics Forum (CGF)*, 2025.
- Alex Hanson, Allen Tu, Geng Lin, Vasu Singla, Matthias Zwicker, and Tom Goldstein. Speedy-splat: Fast 3d gaussian splatting with sparse pixels and sparse primitives. In *Proceedings of the Computer Vision and Pattern Recognition Conference (CVPR)*, pp. 21537–21546, 2025.
- Qiqi Hou, Randall Rauwendaal, Zifeng Li, Hoang Le, Farzad Farhadzadeh, Fatih Porikli, Alexei Bourd, and Amir Said. Sort-free gaussian splatting via weighted sum rendering. In *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025.
- Binbin Huang, Zehao Yu, Anpei Chen, Andreas Geiger, and Shenghua Gao. 2d gaussian splatting for geometrically accurate radiance fields. In Andres Burbano, Denis Zorin, and Wojciech Jarosz (eds.), *ACM SIGGRAPH Conference Papers*, 2024a.
- Letian Huang, Jiayang Bai, Jie Guo, and Yanwen Guo. Gs++: Error analyzing and optimal gaussian splatting. *CoRR*, 2024b.
- Juho Kannala and Sami S Brandt. A generic camera model and calibration method for conventional, wide-angle, and fish-eye lenses. *IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI)*, 2006.
- Bernhard Kerbl, Georgios Kopanas, Thomas Leimkühler, and George Drettakis. 3d gaussian splatting for real-time radiance field rendering. *ACM Transactions on Graphics (TOG)*, 2023.
- Bernhard Kerbl, Andreas Meuleman, Georgios Kopanas, Michael Wimmer, Alexandre Lanvin, and George Drettakis. A hierarchical 3d gaussian representation for real-time rendering of very large datasets. *ACM Transactions on Graphics (TOG)*, 43(4):1–15, 2024.
- Leonid Keselman and Martial Hebert. Approximate differentiable rendering with algebraic surfaces. In *European Conference on Computer Vision (ECCV)*, 2022.
- Shakiba Kheradmand, Daniel Rebain, Gopal Sharma, Weiwei Sun, Yang-Che Tseng, Hossam Isack, Abhishek Kar, Andrea Tagliasacchi, and Kwang Moo Yi. 3d gaussian splatting as markov chain monte carlo. *Advances in Neural Information Processing Systems (NeurIPS)*, 37:80965–80986, 2024.
- Varun Ravi Kumar, Ciarán Eising, Christian Witt, and Senthil Kumar Yogamani. Surround-view fisheye camera perception for automated driving: Overview, survey & challenges. *IEEE Transactions on Intelligent Transportation Systems*, 24(4):3638–3659, 2023.
- Zimu Liao, Siyan Chen, Rong Fu, Yi Wang, Zhongling Su, Hao Luo, Li Ma, Linning Xu, Bo Dai, Hengjie Li, et al. Fisheye-gs: Lightweight and extensible gaussian splatting module for fisheye cameras. *ECCV Workshop*, 2024.
- Changkun Liu, Shuai Chen, Yash Sanjay Bhalgat, Siyan HU, Ming Cheng, Zirui Wang, Victor Adrian Prisacariu, and Tristan Braud. Gs-cpr: Efficient camera pose refinement via 3d gaussian splatting. In *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025.
- Tao Lu, Mulin Yu, Linning Xu, Yuanbo Xiangli, Limin Wang, Dahua Lin, and Bo Dai. Scaffold-gs: Structured 3d gaussians for view-adaptive rendering. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 20654–20664, 2024.
- Zhaoyang Lv, Maurizio Monge, Ka Chen, Yufeng Zhu, Michael Goesele, Jakob Engel, Zhao Dong, and Richard Newcombe. Photoreal scene reconstruction from an egocentric device. In *ACM SIGGRAPH Conference Papers*, 2025.
- Alexander Mai, Peter Hedman, George Kopanas, Dor Verbin, David Futschik, Qiangeng Xu, Falko Kuester, Jonathan T Barron, and Yinda Zhang. EVER: exact volumetric ellipsoid rendering for real-time view synthesis. *IEEE International Conference on Computer Vision (ICCV)*, 2025.

- Saswat Subhajyoti Mallick, Rahul Goel, Bernhard Kerbl, Markus Steinberger, Francisco Vicente Carrasco, and Fernando De La Torre. Taming 3dgs: High-quality radiance fields with limited resources. In *SIGGRAPH Asia Conference Papers*, 2024.
- Xiaoyang Mao, Lichan Hong, and Arie Kaufman. Splatting of curvilinear volumes. In *Proceedings Visualization '95*, pp. 61–68. IEEE, 1995.
- Christopher Mei and Patrick Rives. Single view point omnidirectional camera calibration from planar grids. In *IEEE International Conference on Robotics and Automation (ICRA)*, 2007.
- Nicolas Moënné-Loccoz, Ashkan Mirzaei, Or Perel, Riccardo de Lutio, Janick Martinez Esturo, Gavriel State, Sanja Fidler, Nicholas Sharp, and Zan Gojcic. 3d gaussian ray tracing: Fast tracing of particle scenes. *ACM Transactions on Graphics (TOG)*, 2024.
- Hazem Rashed, Eslam Mohamed, Ganesh Sistu, Varun Ravi Kumar, Ciaran Eising, Ahmad El-Sallab, and Senthil Yogamani. Generalized object detection on fisheye cameras for autonomous driving: Dataset, representations and baseline. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, (WACV), pp. 2272–2280, 2021.
- Liu Ren, Hanspeter Pfister, and Matthias Zwicker. Object space ewa surface splatting: A hardware accelerated approach to high quality point rendering. In *Computer graphics forum (CGF)*, 2002.
- Samuel Rota Bulò, Lorenzo Porzi, and Peter Kotschieder. Revising densification in gaussian splatting. In *European Conference on Computer Vision (ECCV)*, pp. 347–362. Springer, 2024.
- Johannes Lutz Schönberger and Jan-Michael Frahm. Structure-from-motion revisited. In *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016.
- Ahmed Rida Sekkat, Yohan Dupuis, Varun Ravi Kumar, Hazem Rashed, Senthil Yogamani, Pascal Vasseur, and Paul Honeine. Synwoodscape: Synthetic surround-view fisheye camera dataset for autonomous driving. *IEEE Robotics and Automation Letters (RAL)*, 7(3):8502–8509, 2022.
- Chinmay Talegaonkar, Yash Belhe, Ravi Ramamoorthi, and Nicholas Antipa. Volumetrically consistent 3d gaussian rasterization. In *Proceedings of the Computer Vision and Pattern Recognition Conference (CVPR)*, pp. 10953–10963, 2025.
- Lee Westover. Footprint evaluation for volume rendering. In *Proceedings of the 17th annual conference on Computer graphics and interactive techniques*, pp. 367–376, 1990.
- Qi Wu, Janick Martinez Esturo, Ashkan Mirzaei, Nicolas Moenne-Loccoz, and Zan Gojcic. 3dgt: Enabling distorted cameras and secondary rays in gaussian splatting. *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.
- Zhiwen Yan, Weng Fei Low, Yu Chen, and Gim Hee Lee. Multi-scale 3d gaussian splatting for anti-aliased rendering. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, (CVPR), pp. 20923–20931, 2024.
- Chandan Yeshwanth, Yueh-Cheng Liu, Matthias Nießner, and Angela Dai. Scannet++: A high-fidelity dataset of 3d indoor scenes. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023.
- Senthil Yogamani, Ciarán Hughes, Jonathan Horgan, Ganesh Sistu, Padraig Varley, Derek O’Dea, Michal Uricár, Stefan Milz, Martin Simon, Karl Amende, et al. Woodscape: A multi-task, multi-camera fisheye dataset for autonomous driving. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pp. 9308–9318, 2019.
- Senthil Yogamani, David Unger, Venkatraman Narayanan, and Varun Ravi Kumar. Fisheyebevseg: Surround view fisheye cameras based bird’s-eye view segmentation for autonomous driving. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, (CVPR), pp. 1331–1334, 2024.
- Zehao Yu, Anpei Chen, Binbin Huang, Torsten Sattler, and Andreas Geiger. Mip-splatting: Alias-free 3d gaussian splatting. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, (CVPR), pp. 19447–19456, 2024a.

-
- Zehao Yu, Torsten Sattler, and Andreas Geiger. Gaussian opacity fields: Efficient adaptive surface reconstruction in unbounded scenes. *ACM Transactions on Graphics (TOG)*, 2024b.
- Zheng Zhang, Wenbo Hu, Yixing Lao, Tong He, and Hengshuang Zhao. Pixel-gs: Density control with pixel-aware gradient for 3d gaussian splatting. In *European Conference on Computer Vision (ECCV)*, pp. 326–342. Springer, 2024.
- Zichao Zhang, Henri Rebecq, Christian Forster, and Davide Scaramuzza. Benefit of large field-of-view cameras for visual odometry. In *2016 IEEE International Conference on Robotics and Automation (ICRA)*, pp. 801–808. IEEE, 2016.
- Hexu Zhao, Haoyang Weng, Daohan Lu, Ang Li, Jinyang Li, Aurojit Panda, and Saining Xie. On scaling up 3d gaussian splatting training. In *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025.
- Matthias Zwicker, Hanspeter Pfister, Jeroen Van Baar, and Markus Gross. Ewa volume splatting. In *Proceedings Visualization, 2001. VIS'01.*, pp. 29–538. IEEE, 2001.

CONTENTS

1 Introduction	1
2 Related Works	2
2.1 Splatting-Based Gaussian Rendering	2
2.2 Ray-Based Gaussian Rendering	3
3 Method	4
3.1 Rendering with Exact Projective Geometry	5
3.2 Ray-Particle Association with Exact Projective Geometry	5
3.3 Bipolar Equiangular Projection (BEAP)	6
4 Experiments	7
4.1 Scaling Gaussians vs. Projective Exactness in Wide-FoV Cameras	7
4.2 Comprehensive Evaluation on All-FoV Novel View Synthesis	8
4.3 Comprehensive Runtime Analysis	10
4.4 Local Affine Approximation vs. Projective-Exact Integral	11
4.5 BEAP On/Off Study	11
5 Conclusions	12
a Ethics Statement	13
b Reproducibility Statement	13
c LLM Usage Statement	19
A Appendix: Exactness and Limitation	19
B Appendix: Ray-Gaussian Transmittance with Projective Exactness	19
B.1 Preliminaries And Notations	19
B.2 Differentiable Rendering with Projective Exactness	20
B.3 Ray Integral of Isotropic 3D Gaussian	21
C Appendix: Cookbook of Gradient Computation	22
C.1 Chain Rule	22
C.2 Two-Stage Propagation	22
C.3 Gradients for Canonical Ray	22
C.4 Gradients for Whitening Matrix	22
C.5 Gradients for Scaling and Rotation	23
C.6 Gradients for 3D Mean	23
C.7 Numerical Stability of Our Differentiable Framework	23

918	D Appendix: Additional Details in Association	24
919	D.1 Particle Bounding Frustum (PBF) Solver	24
920	D.2 Details in Computing CSF-PBF Intersection	25
921	D.3 Relation to Prior Works	26
922		
923	E Appendix: BEAP Image Inverse Interpolation	26
924	Pinhole Camera Model	27
925	KB Fisheye Model	27
926		
927	F Appendix: High-level Algorithmic Summary	27
928		
929	G Appendix: Additional Implementation Details	28
930	Baselines	28
931	Implementation Details	28
932	Central and peripheral regions in ScanNet++ evaluation	29
933	Cross-camera generalization evaluation on ZipNeRF	29
934		
935	H Appendix: EWA Splatting Preliminaries	29
936	From camera coordinates to ray coordinates	29
937	Local affine approximation	29
938	Splatting approximation error	30
939		
940	I Appendix: Advantage over Prior Ray-Tracing Methods	30
941		
942	J Appendix: Results on Egocentric Dataset	30
943		
944	K Appendix: Additional Material and Results	32
945	K.1 Additional Material.	32
946	K.2 Additional Results.	32
947		
948		
949		
950		
951		
952		
953		
954		
955		
956		
957		
958		
959		
960		
961		
962		
963		
964		
965		
966		
967		
968		
969		
970		
971		

LLM USAGE STATEMENT

We used a large language model (ChatGPT, OpenAI) as a writing assistant. Specifically, the LLM was employed to (i) improve the clarity and readability of the manuscript, (ii) suggest stylistic edits to figure captions and section transitions, and (iii) provide alternative phrasings for technical descriptions without altering the scientific content.

The LLM was not involved in research ideation, algorithm design, data analysis, or experimental validation. All technical contributions, mathematical derivations, and results are the work of the authors. The authors take full responsibility for the content of this paper.

A APPENDIX: EXACTNESS AND LIMITATION

We state precisely what “exactness” denotes in this work: the *elimination of linear approximation in projective geometry* throughout both the rendering and the ray–particle association stages—an error that, unlike a capacity-limited one, does not diminish as the Gaussian count grows (Sec. 4.1). We deliberately refrain from claiming exactness in the full physical-volumetric sense; ordering, self-attenuation, and overlap remain approximated, and we delineate these explicitly as the boundaries of the present formulation so that the scope of the claim is unambiguous.

Limitations. Because our formulation integrates each Gaussian independently, it does not explicitly resolve self-occlusion, and ordering is handled by depth-based sorting—which, as the supplementary videos show, can occasionally produce popping artifacts. To our knowledge EVER (Mai et al., 2025) is the sole method that resolves ordering and self-occlusion analytically, and it secures that property only by approximating each Gaussian as a constant-density ellipsoid at a substantial cost in rendering speed.

This limitation is, however, of little practical consequence. Prior ray-tracing–based Gaussian renderers—Fuzzy Metaballs (Keselman & Hebert 2022) and the subsequent 3DGR/UT family (Moënné-Loccoz et al., 2024; Wu et al., 2025)—all rest on the “maximum-response” assumption for ray–Gaussian interaction, which likewise disregards ordering and self-attenuation yet has been shown to introduce negligible error regardless of the Gaussian count (Celarek et al., 2025), Fig. 9–11 / Sec. 7.1). Our depth-based ordering therefore attains the same practical accuracy as existing ray-tracing–based Gaussian renderers, so this boundary constrains the formulation only in principle without materially affecting the reported results.

Future work. Although the practical impact of ordering and self-occlusion is small in current Gaussian-based ray tracing, both nonetheless bound exact volumetric correctness in principle, and two directions follow naturally. The first is to incorporate analytic or learnable self-attenuation models—drawing on ellipsoidal integration (e.g., EVER (Mai et al., 2025)) while shedding its runtime penalty. The second is to treat ordering more principledly, which should further suppress the rare popping artifacts observed under extreme configurations. Both concerns are orthogonal to our contribution on projective exactness, and addressing them would extend the framework’s robustness to large-scale and high-dynamic scenes.

B APPENDIX: RAY-GAUSSIAN TRANSMITTANCE WITH PROJECTIVE EXACTNESS

B.1 PRELIMINARIES AND NOTATIONS

We adopt the classical parameterization of 3D Gaussian primitives (Zwicker et al., 2001), in which each Gaussian is specified by a mean position $\mu \in \mathbb{R}^3$ together with a 3D covariance matrix Σ in the world coordinate system:

$$\mathcal{G}_{\Sigma, \mu}(\mathbf{x}) = \frac{1}{\rho |\Sigma|^{1/2}} \exp \left(-\frac{1}{2} (\mathbf{x} - \mu)^\top \Sigma^{-1} (\mathbf{x} - \mu) \right), \quad (\text{B.1})$$

where the normalization constant is $\rho = \sqrt{2\pi}$ and $\mathbf{x} \in \mathbb{R}^3$ denotes a point in world space. Unlike recent methods such as 3DGS (Kerbl et al., 2023) and 3DGUT (Wu et al., 2025), which discard the

determinant term $|\Sigma|^{1/2}$ to simplify splatting or association, we retain it deliberately, because it is precisely this term that renders the subsequent ray integral exact rather than merely proportional. The covariance factorizes as $\Sigma = RSS^\top R^\top$, with a scaling matrix $S = \text{diag}(\mathbf{s})$ constructed from the scaling vector $\mathbf{s} \in \mathbb{R}^3$ and a rotation matrix $R \in \text{SO}(3)$ recovered from a unit quaternion $\mathbf{q} \in \mathbb{R}^4$.

The transmittance contributed by a single 3D Gaussian particle along a ray $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$ is then

$$T(\mathbf{o}, \mathbf{d}) = \sigma \int_{t \in \mathbb{R}} \mathcal{G}_{\Sigma, \mu}(\mathbf{o} + t\mathbf{d}) dt, \quad (\text{B.2})$$

where σ is the opacity coefficient of the Gaussian.

Finally, for the set of depth-sorted 3D Gaussian particles contributing to a given ray, the rendered color C follows the standard front-to-back compositing rule:

$$C(\mathbf{o}, \mathbf{d}) = \sum_i c_i(\mathbf{o}, \mu_i) T_i \prod_{j=1}^{i-1} (1 - T_j), \quad (\text{B.3})$$

where $c_i(\mathbf{o}, \mu_i)$ is the view-dependent color of the i -th Gaussian as seen from the optical center $\mathbf{o}$, parameterized by spherical harmonics (SH) (Kerbl et al. 2023). Minimizing the reconstruction loss between rendered and observed color along each ray jointly optimizes, through backpropagation, the parameters $\mathbf{s}, \mathbf{q}, \mu$, the opacity coefficient σ , and the SH coefficients.

B.2 DIFFERENTIABLE RENDERING WITH PROJECTIVE EXACTNESS

Achieving projective-exact Gaussian rendering within a ray-marching formulation reduces to a single obstacle: obtaining a closed-form solution for the integral of Gaussian density and color along a 3D ray. We surmount it by transforming the ray—per Gaussian—into a canonical coordinate system in which the transmittance integral (Eq. B.2) is preserved and the Gaussian becomes isotropic, and then deriving the exact ray color from first principles.

Concretely, we introduce a canonical coordinate system shared across all 3D Gaussian primitives. A point $\mathbf{u} \in \mathbb{R}^3$ in this space maps to world coordinates, per Gaussian, through:

$$\begin{bmatrix} \mathbf{x} \\ 1 \end{bmatrix} = H \begin{bmatrix} \mathbf{u} \\ 1 \end{bmatrix} = \begin{bmatrix} RS & \mu \\ \mathbf{0} & 1 \end{bmatrix} \begin{bmatrix} \mathbf{u} \\ 1 \end{bmatrix}, \quad (\text{B.4})$$

Under this transformation, an infinitesimal ray segment $\Delta \mathbf{u}$ in the canonical frame corresponds to $\Delta \mathbf{x} = RS \Delta \mathbf{u}$ in the world frame.

Preserving the per-Gaussian transmittance across coordinate systems requires that the elementary transmittance—the product of density and segment length—remain invariant under the transformation (Fig. 2 Right):

$$\mathcal{G}_{\mathbf{I}, \mathbf{o}}(\mathbf{u}) |\Delta \mathbf{u}| = \mathcal{G}_{\Sigma, \mu}(\mathbf{x}) |\Delta \mathbf{x}|. \quad (\text{B.5})$$

This constraint absorbs the Jacobian determinant into the measure; because $|\frac{\Delta \mathbf{x}}{\Delta \mathbf{u}}| = |RS| = |\Sigma|^{1/2}$, the canonical density reduces to:

$$\mathcal{G}_{\mathbf{I}, \mathbf{o}}(\mathbf{u}) = \frac{1}{\rho} \exp\left(-\frac{1}{2} \mathbf{u}^\top \mathbf{u}\right), \quad (\text{B.6})$$

which is isotropic and shared by every primitive in canonical space.

Consequently, the exact closed-form transmittance integral of each Gaussian along $\mathbf{r}(t)$ follows by accumulating density along the transformed ray $\mathbf{r}_u(t) = \mathbf{o}_u + t\mathbf{d}_u$ through the canonical isotropic Gaussian. The contribution from $t < 0$ is invariably negligible unless the optical center lies within the 3-sigma contour—a configuration that every 3D Gaussian rendering framework culls during optimization in any case.

The integral admits a *closed-form* solution expressed in the perpendicular distance from the ray to the Gaussian center, as stated in Theorem 1:

$$T(\mathbf{o}, \mathbf{d}) = \sigma \int_{t \in \mathbb{R}} \mathcal{G}_{\mathbf{I}, \mathbf{o}}(\mathbf{o}_u + t\mathbf{d}_u) dt = \sigma \exp\left(-\frac{1}{2} (D_{\mu, \Sigma}(\mathbf{o}, \mathbf{d}))^2\right), \quad (\text{B.7})$$

where $D_{\mu,\Sigma}(\mathbf{o}, \mathbf{d})$ is the perpendicular distance from the Gaussian center (the canonical-space origin) to the transformed ray $\mathbf{r}_u(t)$. Geometrically, this quantity is equivalently the minimal *Mahalanobis distance* from any point on the ray to the Gaussian in world space—a direct measure of how closely the ray passes the particle distribution.

It then remains only to compute this perpendicular distance from the canonical origin to the transformed ray. We write the transformed ray’s direction vector $\mathbf{d}_u$ and moment vector $\mathbf{m}_u$ as:

$$\mathbf{d}_u = (S^{-1}R^\top) \mathbf{d}, \quad \mathbf{m}_u = \mathbf{o}_u \times \mathbf{d}_u, \quad (\text{B.8})$$

where $\mathbf{o}_u = (S^{-1}R^\top)(\mathbf{o} - \boldsymbol{\mu})$ is the optical center in canonical space. The squared distance $D_{\mu,\Sigma}^2(\mathbf{o}, \mathbf{d})$ then follows efficiently in closed form as:

$$(D_{\mu,\Sigma}(\mathbf{o}, \mathbf{d}))^2 = \frac{\mathbf{m}_u^\top \mathbf{m}_u}{\mathbf{d}_u^\top \mathbf{d}_u}. \quad (\text{B.9})$$

Substituting this squared distance into Eq. [B.7](#) and the resulting transmittance into Eq. [B.3](#) yields the accumulated color along the ray—thereby completing an exact, end-to-end differentiable rendering chain.

B.3 RAY INTEGRAL OF ISOTROPIC 3D GAUSSIAN

Theorem 1. *Given a ray and a standard isotropic 3D Gaussian, the ray-Gaussian integral has a closed-form proportional to:*

$$\exp\left(-\frac{1}{2}D^2\right),$$

where D is the perpendicular distance from the Gaussian center to the ray.

Proof. We first translate the coordinate system so that the Gaussian $\mathcal{G}$ is centered at the origin, translating the ray $\mathbf{r}_u$ correspondingly, and then rotate the scene so that the ray direction aligns with the z -axis. By the translation-invariance of the integral together with the isotropy of the standard Gaussian, neither transformation alters its value.

The translated standard isotropic 3D Gaussian is:

$$\mathcal{G}_{\mathbf{I},\mathbf{o}}(\mathbf{u}) = \frac{1}{\rho} \exp\left(-\frac{1}{2}\mathbf{u}^\top \mathbf{u}\right), \quad (\text{B.10})$$

which matches Eq. [B.6](#). Let the z -axis aligned ray as:

$$\mathbf{r}_u(t) = \begin{bmatrix} \mathbf{o}_u \\ \mathbf{o}_v \\ 0 \end{bmatrix} + t \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}, \quad (\text{B.11})$$

the ray-Gaussian integral becomes:

$$\begin{aligned} \int_{t \in \mathbb{R}} \mathcal{G}_{\mathbf{I},\mathbf{o}}(\mathbf{r}_u(t)) dt &= \frac{1}{\rho} \exp\left(-\frac{u_0^2 + v_0^2}{2}\right) \int_{t \in \mathbb{R}} \exp\left(-\frac{1}{2}t^2\right) dt \\ &= \frac{\sqrt{2\pi}}{\rho} \exp\left(-\frac{1}{2}D^2\right), \end{aligned} \quad (\text{B.12})$$

where $D^2 = \frac{u_0^2 + v_0^2}{2}$ denotes the squared perpendicular distance from the ray to the Gaussian center in canonical space. Thus, the integral is proportional to $\exp(-\frac{1}{2}D^2)$, as stated in theorem. $\square$

The density function shared by our canonical-space primitives (Eq. [B.6](#)) coincides exactly with the isotropic form used above. Although D takes the world-coordinate ray $(\mathbf{o}, \mathbf{d})$ as input, its value depends equally on the world-space Gaussian parameters $\boldsymbol{\mu}, \Sigma$, which determine the canonical transformation of the ray. Hence, for each Gaussian contributing to the ray, the transmittance is:

$$T = \frac{\sqrt{2\pi}\sigma}{\rho} \exp\left(-\frac{1}{2}(D_{\mu,\Sigma}(\mathbf{o}, \mathbf{d}))^2\right), \quad (\text{B.13})$$

which recovers Eq. [B.7](#) exactly under the choice $\rho = \sqrt{2\pi}$.

C APPENDIX: COOKBOOK OF GRADIENT COMPUTATION

C.1 CHAIN RULE

Recall that $T = \sigma\alpha$ is the ray–Gaussian transmittance (Eq. B.7), with σ the opacity coefficient and α the ray–Gaussian integral. Let $\kappa = D_{\mu, \Sigma}^2$ denote the minimal squared Mahalanobis distance from the ray to the Gaussian, let S be the Gaussian scaling, and let R be the world-space rotation matrix. Given the world-space 3D covariance $\Sigma = RSS^\top R^\top$, the PCA whitening matrix W_{pca} satisfies:

$$W_{\text{pca}} = \Sigma^{-\frac{1}{2}} = S^{-1}R^\top, \quad (C.1)$$

where $S^{-1} = \text{diag}\left(\frac{1}{s}\right).$

Following the gradient flow in Fig. 2 the chain rule furnishes the derivatives *w.r.t.* the Gaussian parameters—the scaling vector s , the rotation quaternion $\mathbf{q}$, and the mean μ —as follows:

$$\frac{\partial T}{\partial s} = \frac{\partial T}{\partial \kappa} \frac{\partial \kappa}{\partial W_{\text{pca}}} \frac{\partial W_{\text{pca}}}{\partial s}, \quad (C.2)$$

$$\frac{\partial T}{\partial \mathbf{q}} = \frac{\partial T}{\partial \kappa} \frac{\partial \kappa}{\partial W_{\text{pca}}} \frac{\partial W_{\text{pca}}}{\partial \mathbf{q}}, \quad (C.3)$$

$$\frac{\partial T}{\partial \mu} = \frac{\partial T}{\partial \kappa} \frac{\partial \kappa}{\partial \mu}. \quad (C.4)$$

C.2 TWO-STAGE PROPAGATION

To minimize the overhead of shared propagation, we organize gradient computation as the two-stage pipeline of 3DGS (Kerbl et al., 2023), comprising `RenderCUDA` (RC) and `PreprocessCUDA` (PC). RC performs ray-wise accumulation, accelerated by our Particle Bounding Frustum (PBF) association (Sec. 3.2), whereas PC operates at Gaussian-wise complexity and benefits further from view-frustum culling.

From Eq. B.7 the gradient *w.r.t.* the opacity coefficient together with the partial derivative shared across Gaussian parameters, both evaluated in RC, are:

$$\left. \frac{\partial T}{\partial \sigma} \right|_{\text{RC}} = \alpha, \quad \left. \frac{\partial T}{\partial \kappa} \right|_{\text{RC}} = -\frac{T}{2}. \quad (C.5)$$

The remaining parameter gradients then propagate from the squared distance, namely $\frac{\partial \kappa}{\partial s}, \frac{\partial \kappa}{\partial \mathbf{q}}, \frac{\partial \kappa}{\partial \mu}$.

C.3 GRADIENTS FOR CANONICAL RAY

We next backpropagate to the canonical ray parameters—the origin $\mathbf{o}_u$, the direction $\mathbf{d}_u$, and the moment $\mathbf{m}_u = \mathbf{o}_u \times \mathbf{d}_u$ of Eq. B.8. These quantities enter the Mahalanobis distance of Eq. B.9 directly and jointly encode the canonical ray geometry.

$$\left. \frac{\partial \kappa}{\partial \mathbf{m}_u} \right|_{\text{RC}} = \left(\frac{2}{\mathbf{d}_u^\top \mathbf{d}_u} \right) \mathbf{m}_u, \quad (C.6)$$

$$\left. \frac{\partial \kappa}{\partial \mathbf{o}_u} \right|_{\text{RC}} = \frac{\partial \kappa}{\partial \mathbf{m}_u} \times \mathbf{d}_u, \quad (C.7)$$

$$\left. \frac{\partial \kappa}{\partial \mathbf{d}_u} \right|_{\text{RC}} = - \left(\frac{2\kappa}{\mathbf{d}_u^\top \mathbf{d}_u} \right) \mathbf{d}_u + \mathbf{o}_u \times \frac{\partial \kappa}{\partial \mathbf{m}_u}. \quad (C.8)$$

C.4 GRADIENTS FOR WHITENING MATRIX

Recall that $\mathbf{d}_u = W_{\text{pca}} \mathbf{d}$ and $\mathbf{o}_u = W_{\text{pca}}(\mathbf{o} - \mu)$, where $\mathbf{o} = -R_c^\top \mathbf{t}_c$ is the world-space optical center and $\mathbf{d}$ the ray direction. Given the ray-parameter gradients, we propagate them to the PCA whitening matrix W_{pca} via outer products with the corresponding world-space vectors:

$$\begin{aligned}
\text{RenderCUDA: } \frac{\partial \kappa}{\partial W_{\text{pca}}}\bigg|_{\text{RC}} &= \frac{\partial \kappa}{\partial \mathbf{d}_u} \mathbf{d}_u^\top; \\
\text{PreprocessCUDA: } \frac{\partial \kappa}{\partial W_{\text{pca}}}\bigg|_{\text{PC}} &= \frac{\partial \kappa}{\partial \mathbf{o}_u} (\mathbf{o}_u - \boldsymbol{\mu})^\top; \\
\text{Overall: } \frac{\partial \kappa}{\partial W_{\text{pca}}} &= \frac{\partial \kappa}{\partial W_{\text{pca}}}\bigg|_{\text{RC}} + \frac{\partial \kappa}{\partial W_{\text{pca}}}\bigg|_{\text{PC}}.
\end{aligned} \tag{C.9}$$

The gradient through $\mathbf{d}_u$ is accumulated ray-wise in `RenderCUDA`, whereas that through $\mathbf{o}_u$ is computed Gaussian-wise in `PreprocessCUDA`; summing the two contributions recovers the complete gradient *w.r.t.* the PCA transformation matrix W_{pca} .

C.5 GRADIENTS FOR SCALING AND ROTATION

Thus, for the scaling gradients $\frac{\partial \kappa}{\partial \mathbf{s}} = \frac{\partial \kappa}{\partial W_{\text{pca}}} \frac{\partial W_{\text{pca}}}{\partial \mathbf{s}}$, we obtain:

$$\frac{\partial W_{\text{pca}}^{i,j}}{\partial s_k}\bigg|_{\text{PC}} = \begin{cases} -\frac{1}{s_k^2} R_{j,k} & \text{if } i = k \\ 0 & \text{otherwise} \end{cases}. \tag{C.10}$$

To derive the rotation gradients $\frac{\partial \kappa}{\partial \mathbf{q}} = \frac{\partial \kappa}{\partial W_{\text{pca}}} \frac{\partial W_{\text{pca}}}{\partial \mathbf{q}}$, we write the whitening matrix, for the quaternion form $\mathbf{q} = q_r + q_i \mathbf{i} + q_j \mathbf{j} + q_k \mathbf{k}$, as:

$$W_{\text{pca}} = 2 \begin{pmatrix} \frac{1}{s_1}(\frac{1}{2} - q_j^2 - q_k^2) & \frac{1}{s_1}(q_i q_j + q_r q_k) & \frac{1}{s_1}(q_i q_k - q_r q_j) \\ \frac{1}{s_2}(q_i q_j - q_r q_k) & \frac{1}{s_2}(\frac{1}{2} - q_i^2 - q_k^2) & \frac{1}{s_2}(q_j q_k + q_r q_i) \\ \frac{1}{s_3}(q_i q_k + q_r q_j) & \frac{1}{s_3}(q_j q_k - q_r q_i) & \frac{1}{s_3}(\frac{1}{2} - q_i^2 - q_j^2) \end{pmatrix}. \tag{C.11}$$

As a result, the end gradients for the quaternion $\mathbf{q}$ take the form:

$$\begin{aligned}
\frac{\partial W_{\text{pca}}}{\partial q_r}\bigg|_{\text{PC}} &= 2 \begin{pmatrix} 0 & \frac{q_k}{s_1} & -\frac{q_j}{s_1} \\ -\frac{q_k}{s_2} & 0 & \frac{q_i}{s_2} \\ \frac{q_j}{s_3} & -\frac{q_i}{s_3} & 0 \end{pmatrix}, \quad \frac{\partial W_{\text{pca}}}{\partial q_i}\bigg|_{\text{PC}} = 2 \begin{pmatrix} 0 & \frac{q_j}{s_2} & \frac{q_k}{s_2} \\ \frac{q_k}{s_3} & -\frac{q_r}{s_3} & -\frac{2q_i}{s_3} \\ -\frac{2q_j}{s_3} & \frac{q_i}{s_3} & -\frac{q_r}{s_3} \end{pmatrix}, \\
\frac{\partial W_{\text{pca}}}{\partial q_j}\bigg|_{\text{PC}} &= 2 \begin{pmatrix} -\frac{2q_j}{s_1} & \frac{q_i}{s_1} & -\frac{q_r}{s_1} \\ \frac{q_i}{s_2} & 0 & \frac{q_k}{s_2} \\ \frac{q_r}{s_3} & \frac{q_k}{s_3} & -\frac{2q_i}{s_3} \end{pmatrix}, \quad \frac{\partial W_{\text{pca}}}{\partial q_k}\bigg|_{\text{PC}} = 2 \begin{pmatrix} -\frac{2q_k}{s_1} & \frac{q_r}{s_1} & \frac{q_j}{s_1} \\ -\frac{q_r}{s_2} & -\frac{2q_k}{s_2} & \frac{q_j}{s_2} \\ \frac{q_i}{s_3} & \frac{q_j}{s_3} & 0 \end{pmatrix}.
\end{aligned} \tag{C.12}$$

C.6 GRADIENTS FOR 3D MEAN

Finally, the end gradient *w.r.t.* the Gaussian 3D mean $\boldsymbol{\mu}$ follows as:

$$\frac{\partial \kappa}{\partial \boldsymbol{\mu}}\bigg|_{\text{PC}} = -W_{\text{pca}}^\top \frac{\partial \kappa}{\partial \mathbf{o}_u}. \tag{C.13}$$

C.7 NUMERICAL STABILITY OF OUR DIFFERENTIABLE FRAMEWORK

The choice of forward formulation, backward gradient computation, and intermediate cache variables is decisive for numerical stability. Prior approaches such as GOF (Yu et al. 2024b) are unstable in backward propagation and therefore depend on auxiliary 3D filtering heuristics to keep training well-behaved. Our closed-form formulation and its exact gradient chain preclude this instability by construction, obviating any such filtering.

The decisive difference resides in how $\mathbf{m}_u^\top \mathbf{m}_u$ is expanded. In exact arithmetic the two forms are geometrically identical:

$$\mathbf{m}_u^\top \mathbf{m}_u = \|\mathbf{o}_u\|^2 \|\mathbf{d}_u\|^2 - (\mathbf{o}_u^\top \mathbf{d}_u)^2. \tag{C.14}$$

The right-hand form, however, is fragile: when one axis of a 3D Gaussian is scaled toward zero, the canonical-space vectors $\mathbf{o}_u$ and $\mathbf{d}_u$ become nearly parallel, so that $\cos(\angle(\mathbf{o}_u, \mathbf{d}_u)) \rightarrow 1$ while both



Figure C.1: **Artifacts Without 3D Filtering in GOF.** (Left) GOF with degenerate Gaussians left unfiltered; (Right) the stable, artifact-free result obtained by substituting our forward and backward formulations in their place.

$\|\mathbf{o}_u\|$ and $\|\mathbf{d}_u\|$ diverge. The difference $\|\mathbf{o}_u\|^2\|\mathbf{d}_u\|^2 - (\mathbf{o}_u^\top \mathbf{d}_u)^2$ then incurs catastrophic cancellation—prone to overflow and, in practice, liable to return negative values that in turn destabilize the backward pass.

Fig. C.1 (Left) exhibits the resulting artifacts in GOF when degenerate Gaussians are left unfiltered: the snowflake noise originates in the overflow of Eq. C.14 where prematurely negative values and vanishingly small positive values inflate the Gaussian intensity. Because the EWA approximation with Jacobian and conic is itself numerically stable, these artifacts trace specifically to the inaccurate maximum-response computation, which additionally produces the gridline pattern.

Our CUDA rasterizer instead evaluates the cross product and its gradients directly (Eq. B.9 forward; Eqs. C.6, C.9 backward), thereby circumventing the cancellation entirely: the forward pass is artifact-free and the backward pass returns accurate gradients. Fig. C.1 (Right) confirms this—substituting our forward and backward formulations into GOF eliminates the instability-induced artifacts even in the absence of any 3D filtering of degenerate (flat disk-like or needle) Gaussians.

D APPENDIX: ADDITIONAL DETAILS IN ASSOCIATION

Section D.1 provides the detailed derivation; Sec. D.2 sets out the implementation considerations that keep the PBF stable under large-FoV and omnidirectional cameras; and Sec. D.3 establishes the derivational relationships together with the precise conditions under which prior formulas reduce to ours.

D.1 PARTICLE BOUNDING FRUSTUM (PBF) SOLVER

Recalling Eq. 8 let $P = (VH)^\top \in \mathbb{R}^{4 \times 4}$ denote the plane transformation matrix from view space to canonical space. Taking the unknown bounds $c = \tan \theta_{1,2}$ as the representative case, we have:

$$\mathbf{g}_u = -P_0 + cP_2 = \begin{bmatrix} -P_{00} + cP_{02} \\ -P_{10} + cP_{12} \\ -P_{20} + cP_{22} \\ -P_{30} + cP_{32} \end{bmatrix}, \quad (\text{D.1})$$

where P_i denotes the i -th row of the matrix, written as a 4×1 column vector, and P_{ij} its j -th element. Recalling Eq. 9 and adopting for concreteness the 1-standard-deviation contour ellipsoid of the Gaussian ($\lambda = 1$), the tangency condition reduces to the quadratic:

$$((\mathbf{f} \circ P_2^\top) \cdot P_2) c^2 - 2((\mathbf{f} \circ P_0^\top) \cdot P_2) c + (\mathbf{f} \circ P_0^\top) \cdot P_0 = 0, \quad (\text{D.2})$$

where $\mathbf{f} = (1, 1, 1, -1)$ is a 1×4 row vector. We abbreviate this equation as:

$$\mathcal{T}_{22}c^2 - 2\mathcal{T}_{02}c + \mathcal{T}_{00} = 0, \quad (\text{D.3})$$

where $\mathcal{T}$ is the 3×3 upper-triangular matrix with entries $\mathcal{T}_{ij} = (\mathbf{f} \circ P_i^\top) \cdot P_j$.

Analogously, the ϕ bounds are the roots $c = \tan \phi_{1,2}$ of:

$$\mathcal{T}_{22}c^2 - 2\mathcal{T}_{12}c + \mathcal{T}_{11} = 0. \quad (\text{D.4})$$

Notably, for a camera-view configuration $[R_c \mid \mathbf{t}_c]$, the parameter matrix $\mathcal{T}$ is assembled directly from the Gaussian’s view-space covariance $\Sigma_c = R_c R S S^\top R^\top R_c^\top$ and 3D mean $\boldsymbol{\mu}_c = R_c \boldsymbol{\mu} + \mathbf{t}_c$, without any intermediate projection:

$$\mathcal{T}_{ij} = \lambda^2 \Sigma_c^{i,j} - (\boldsymbol{\mu}_c \boldsymbol{\mu}_c^\top)^{i,j}. \quad (\text{D.5})$$

The angular bounds—expressed through their tangent-space center and extent—then follow in closed form as:

$$\begin{aligned} \left(\frac{\tan \theta_1 + \tan \theta_2}{2}, \frac{\tan \phi_1 + \tan \phi_2}{2} \right) &= \left(\frac{\mathcal{T}_{02}}{\mathcal{T}_{22}}, \frac{\mathcal{T}_{12}}{\mathcal{T}_{22}} \right), \\ \|\tan \theta_1 - \tan \theta_2\| &= \frac{2\sqrt{\mathcal{T}_{02}^2 - \mathcal{T}_{22}\mathcal{T}_{00}}}{|\mathcal{T}_{22}|}, \\ \|\tan \phi_1 - \tan \phi_2\| &= \frac{2\sqrt{\mathcal{T}_{12}^2 - \mathcal{T}_{22}\mathcal{T}_{11}}}{|\mathcal{T}_{22}|}. \end{aligned} \quad (\text{D.6})$$

Should the discriminant of either quadratic turn negative—the signature of an optical center lying inside the Gaussian ellipsoid, where the angular bounds are undefined—we conservatively clamp the frustum angles to the camera’s maximum field of view (FoV), which preserves valid, conservative bounds in this degenerate case.

Crucially, this result is computed directly from the true 3D view-space covariance, whereas EWA and UT route through an intermediate screen-space conic approximation; the PBF formulation is therefore exact and efficient at once, rather than trading one property against the other.

D.2 DETAILS IN COMPUTING CSF-PBF INTERSECTION

We sample spherical angles uniformly within the maximum FoV. The resulting x, y, z coordinates—the ray direction vectors $\mathbf{d}$ —are projected onto the unit sphere and mapped into pixel space for inverse color interpolation (Sec. E), while the sampled rays are grouped by their CSFs for the subsequent Gaussian association.

For pinhole cameras of limited FoV, the PBF angular bounds ordinarily lie within $[-\frac{\pi}{2}, \frac{\pi}{2}]$, where $\tan$ is monotonically increasing; PBF-CSF intersections then reduce to a simple comparison of the $\tan$ values of their bounds. Under wide FoVs, however, the PBF bounds may fall outside this monotonic interval even when the Gaussian-center incidence angle remains within $[-\frac{\pi}{2}, \frac{\pi}{2}]$, so that a direct comparison of $\tan$ values can yield incorrect intersections—for instance, spuriously matching the bounds as $[\theta'_2, \theta_1]$ (Fig. D.1).

To recover the correct intersections, we apply a mirror transformation to the tangent values of both the PBF angular bounds and the CSF partitions. Following the CMEI camera model (Mei & Rives 2007), we shift the optical center by a mirror parameter ξ , which yields:

$$\tan \theta_\xi = \frac{\tan \theta}{1 + \xi \frac{z}{\|\mathbf{z}\|} \sqrt{1 + \tan^2 \theta}} = \frac{\sin \theta}{\cos \theta + \xi}, \quad (\text{D.7})$$




Figure D.1: Mirror Transformation of PBF Angular Bounds. A given pair of bounds generates four candidate mirrored angles $\theta_{1,\xi}, \theta'_{1,\xi}, \theta_{2,\xi}, \theta'_{2,\xi}$. For a Gaussian centered at $\boldsymbol{\mu}_c$, the two candidates nearest the Gaussian’s own mirrored angle $\theta_{\mu,\xi}$ are selected as the PBF bounds, which guarantees a robust intersection with the mirror-transformed CSF partitions even beyond the monotonic domain of $\tan$ —the regime in which naive tangent comparison fails.

where θ_ξ is the mirror-transformed angle of θ on the xz -plane under mirror parameter ξ ; we fix $\xi = 1$ throughout.

For a pair of tangent values $[\tan \theta_1, \tan \theta_2]$ returned by the PBF solver (Eq. D.6), this transformation produces four candidate mirror angles (Fig. D.1), whereas the Gaussian center maps to a single mirror-transformed angle $\theta_{\mu,\xi}$. Selecting the two candidates nearest $\theta_{\mu,\xi}$ fixes the effective lower and upper PBF bounds $[\theta_{1,\xi}, \theta_{2,\xi}]$, which are subsequently compared against the CSF partitions to obtain the correct intersections.

This construction keeps angular comparisons robust and consistent even beyond the monotonic domain of $\tan$, and it does so cheaply: the timing analysis (Tab. K.5) shows that the entire pre-processing pipeline—mirror-angle-based PBF-CSF intersection together with the canonical transformation of all Gaussians—completes in under 0.13 ms, confirming that the scheme remains efficient even for the thousands of Gaussians present in wide-FoV scenes.

D.3 RELATION TO PRIOR WORKS

Substituting the perspective screen-space coordinates

$$u = \tan \theta, \quad v = \tan \phi$$

into our formulation reproduces the expressions of 2DGS (Huang et al. 2024a) and HTGS (Hahlbohm et al. 2025) exactly—as expected, since both are the perspective specialization of our angular-domain formulation, which subsumes them as a particular case.

The distinction is one of starting assumptions. 2DGS and HTGS parameterize in screen space or on a tangent plane and therefore invoke geometric identities valid only for perspective (pinhole) cameras—those that arise when rays are implicitly bound to a planar, z -parallel image plane; these identities do not carry over to fisheye or omnidirectional systems, where the corresponding screen-space surface is curved. Our derivation instead proceeds from a camera-agnostic angular parameterization in 3D and directly parameterizes Camera Sub-Frustums (CSFs) and Particle Bounding Frustums (PBFs). Because the two derivations begin from different known quantities and constraints, they diverge precisely in the conditions under which closed-form bounds and culling guarantees remain valid for generic cameras—and it is under exactly the wide-FoV conditions where the screen-space assumptions break down that our formulation continues to hold.

We emphasize that mathematics is a general tool; the contribution here lies in deriving these bounds under weaker assumptions—free of any screen-space structure—and in validating them experimentally across both pinhole and wide-FoV camera families.

This difference in parameterization propagates through the entire downstream pipeline. Association, tiling/duplication, and sorting all hinge on how Gaussians map to tiles or bins, and an angular (CSF-based) domain induces a bucketing structure distinct from pinhole screen-space tiling. The distinction becomes consequential at large FoVs: under generic cameras a Gaussian’s bound can appear partly “behind” the camera in screen coordinates (Fig. D.1), driving screen-space methods to mirrored or incorrect tangent-based bounds (dashed line in Fig. D.1). We handle these cases robustly by introducing mirror parameters in the association step (Sec. D.2) and folding the Gaussian center into sorting and correction—steps that a perspective-only derivation never demands.

E APPENDIX: BEAP IMAGE INVERSE INTERPOLATION

To prepare training data, we sample rays covering the full FoV of the original image and pass them as input to our CUDA renderer. In parallel, we interpolate the ground-truth color for each input ray, so that supervision is at once accurate and uniform across the entire FoV.

Projecting the sampled rays onto the observed images begins by converting the incidence-angle projections θ and ϕ into normalized ray direction vectors $\mathbf{d} = (x, y, z)^\top$:




Figure E.1: **Spherical Angles Defining PBF and BEAP.** A view-space direction $\mathbf{d}_c$ is first normalized onto the optical sphere (green spade) and then projected angularly onto the xz and yz planes; the orange point marks the resulting coordinate (θ, ϕ) on the unfolded BEAP imaging plane. It is precisely this shared angular parameterization that unifies the PBF and BEAP constructions.

$$x = \frac{\sin \theta \cos \phi}{\sqrt{\sin^2 \theta \cos^2 \phi + \cos^2 \theta \sin^2 \phi + \cos^2 \theta \cos^2 \phi}}, \quad (\text{E.1})$$

$$y = \frac{\cos \theta \sin \phi}{\sqrt{\sin^2 \theta \cos^2 \phi + \cos^2 \theta \sin^2 \phi + \cos^2 \theta \cos^2 \phi}}, \quad (\text{E.2})$$

$$z = \frac{\cos \theta \cos \phi}{\sqrt{\sin^2 \theta \cos^2 \phi + \cos^2 \theta \sin^2 \phi + \cos^2 \theta \cos^2 \phi}}. \quad (\text{E.3})$$

Our pipeline accommodates several camera models for the subsequent projection of the normalized rays $(x, y, z)^\top$ into pixel space $(x_p, y_p, 1)^\top$.

Pinhole Camera Model. For the pinhole model, $\frac{x}{z} = \tan \theta$ and $\frac{y}{z} = \tan \phi$, from which:

$$\begin{bmatrix} x_p \\ y_p \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \tan \theta \\ \tan \phi \\ 1 \end{bmatrix}. \quad (\text{E.4})$$

KB Fisheye Model. The incidence angle φ follows from its projections:

$$r = \sqrt{\tan^2 \theta + \tan^2 \phi} \quad \text{and} \quad \varphi = \arctan(r). \quad (\text{E.5})$$

and the fisheye-distorted angle follows as:

$$\varphi_d = \varphi(1 + k_1 \varphi^2 + k_2 \varphi^4 + k_3 \varphi^6 + k_4 \varphi^8). \quad (\text{E.6})$$

The rays are thereby projected into pixel space, which in turn enables inverse color interpolation from:

$$\begin{bmatrix} x_p \\ y_p \\ 1 \end{bmatrix} = \begin{bmatrix} \frac{f_x \varphi_d}{r} & 0 & c_x \\ 0 & \frac{f_y \varphi_d}{r} & c_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \tan \theta \\ \tan \phi \\ 1 \end{bmatrix}. \quad (\text{E.7})$$

F APPENDIX: HIGH-LEVEL ALGORITHMIC SUMMARY

We summarize the volumetric rendering pipeline at a high level in Alg. [1](#) and Alg. [2](#). To render the differences immediately legible, we mirror the structure of 3DGS [Kerbl et al. \(2023\)](#): newly introduced or replaced modules are marked in **green**, and substantially modified modules in **yellow**.

ALGORITHM 1: GPU Software Asso. & Render.

Function 3DGEER ($\mu, s, q, c, \sigma, r_c, V$) :

```
CullGaussian( $\mu, V$ ) ▷ Frustum Culling
 $F \leftarrow \text{InitCSFs}(w, h, K)$  ▷ Partition
 $\Sigma_c, \mu_c \leftarrow \text{ViewspaceGaussians}(s, q, V)$  ▷ Transform
 $F_{ID} \leftarrow \text{PBF\_Intersect}(\Sigma_c, \mu, F)$  ▷ Sub-frustum Asso.
 $W_{pca} \leftarrow \text{PCA\_Whitening}(s, q)$  ▷ Transform
 $\mathcal{G}_{ID}, F_K \leftarrow \text{DuplicateWithKeys}(W_{pca}, F_{ID})$  ▷ Indices (Gaussian) & Keys (FrustumID+Depth)
SortByKeys( $\mathcal{G}_{ID}, F_K$ ) ▷ Globally Depth Sort
 $R \leftarrow \text{IdentifyFrustumRanges}(F, F_K)$ 
 $I \leftarrow 0$  ▷ Init Canvas (w/ Rays)
for all CSFs  $f \leftarrow I$  do
  for all Rays  $i \leftarrow f$  do
     $r \leftarrow \text{GetFrustumRange}(R, f)$ 
     $m_u, d_u \leftarrow \text{Cano\_Ray}(W_{pca}, \mu, i)$  ▷ Transform
     $D_{\mu, \Sigma} \leftarrow \text{Dist}(m_u, d_u)$  ▷ Mah-distance
     $I[i] \leftarrow \text{BlendInOrder}(i, \mathcal{G}_{ID}, r, F_K, D_{\mu, \Sigma}, c, \sigma)$  ▷ Alpha Blending
  end
end
return  $I$ 
```

ALGORITHM 2: 3DGEER Initialization and Training Framework

 w, h : width and height of the training images K, V : intrinsics and extrinsic meta info.

```
 $\mu \leftarrow \text{SfM Points}$  ▷ Init Gaussian Attributes
 $s, q \leftarrow \text{InitAttributes}()$  ▷ Position
 $c, \sigma \leftarrow \text{InitCoefficients}()$  ▷ Scaling, Rotation
▷ Color, Opacity
while not converged do
   $r_c, \hat{C} \leftarrow \text{BEAP}(w, h, K)$  ▷ View Rays / Colors
   $C \leftarrow \text{3DGEER}(\mu, s, q, c, \sigma, r_c, V)$  ▷ Asso. & Render
   $\mathcal{L} \leftarrow \text{Loss}(C, \hat{C})$  ▷ See Alg. 1
   $\mu, s, q, c, \sigma \leftarrow \text{Adam}(\nabla \mathcal{L})$  ▷ Loss
  if IsRefinementIteration( $i$ ) then ▷ Backprop & Step
     $\mu, s, q, c, \sigma \leftarrow \text{Optimize}(\mu, s, q, c, \sigma)$  ▷ Kerbl et al. (2023)
  end
end
```

G APPENDIX: ADDITIONAL IMPLEMENTATION DETAILS

Baselines. We compare 3DGEER against the leading volumetric particle rendering methods that target the projective approximation error of 3DGS (Kerbl et al. 2023). For large-FoV fisheye inputs, the baselines are FisheyeGS (Liao et al. 2024), EVER (Mai et al. 2025), and 3DGUT (Wu et al. 2025); under pinhole settings we additionally include 3DGS and 3DGRT (Moënné-Loccoz et al. 2024). This baseline set spans both the splatting and ray-based families, so that 3DGEER is measured against the strongest representative of each rather than a single line of prior work.

Implementation Details. For a fair comparison, we align our implementation with the original 3DGS in initialization and hyperparameters on ScanNet++ and MipNeRF360, and adopt the EVER and SMERF (Duckworth et al. 2024) training configurations for all candidate methods on ZipNeRF. As with most ray-based frameworks, our method does not access gradients of `mean2D` for densification or pruning; it optimizes instead using gradients of the view-space `mean3D`. Unless otherwise

noted, all models share consistent supervision strategies and architectural settings. On the fisheye datasets (ScanNet++ and ZipNeRF) we precompute bijective grid mappings (Guo et al., 2025) and reproject every rendered output to the native KB-Fisheye image space (Kannala & Brandt, 2006) prior to evaluation. The ScanNet++ and MipNeRF360 experiments use full-resolution images; for ZipNeRF, given its high frame count and very large fisheye resolution, we adopt the maximum resolution (Tab. 3) at which every baseline runs reliably on a local desktop.

Central and peripheral regions in ScanNet++ evaluation. The rays sampled in the central part follows the range:

$$\theta \in [-\frac{w}{2f_x}, \frac{w}{2f_x}], \phi \in [-\frac{h}{2f_y}, \frac{h}{2f_y}].$$

The peripheral part comprises the rays sampled over the remaining angular range. Here w, h denote the image width and height, and f_x, f_y the focal lengths of the original KB camera model, respectively.

Cross-camera generalization evaluation on ZipNeRF. ZipNeRF supplies two data sets: raw fisheye images and their undistorted counterparts, the latter obtained by distortion correction and central cropping. Although the frames are aligned per frame, the camera poses and point clouds were reconstructed independently with COLMAP (Schönberger & Frahm, 2016) for each set, and because COLMAP’s global optimization refines each trajectory separately, no single rigid transformation aligns the two pose sets exactly—applying one would inject alignment error. To keep the cross-camera evaluation free of such structure-from-motion (SfM) error, we consistently reuse the training-time camera poses at test time, varying only the intrinsic parameters while holding the extrinsics fixed. This protocol delivers fair comparisons uncontaminated by preprocessing-induced pose inconsistencies.

H APPENDIX: EWA SPLATTING PRELIMINARIES

From camera coordinates to ray coordinates. Following the coordinate conventions of EWA Splatting (Zwicker et al., 2001), the projective transformation for pinhole cameras maps camera- (view-)space points $\mathbf{t}$ to ray-space points $\mathbf{x}$ through $\tau : \mathbf{t} \rightarrow \mathbf{x}$:

$$\mathbf{x} = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \end{bmatrix} = \tau(\mathbf{t}) = \begin{bmatrix} t_0/t_2 \\ t_1/t_2 \\ \|(t_0, t_1, t_2)^\top\| \end{bmatrix}, \quad (\text{H.1})$$

$$\mathbf{t} = \begin{bmatrix} t_0 \\ t_1 \\ t_2 \end{bmatrix} = \tau^{-1}(\mathbf{x}) = \begin{bmatrix} x_0/l \cdot x_2 \\ x_1/l \cdot x_2 \\ 1/l \cdot x_2 \end{bmatrix}, \quad (\text{H.2})$$

$$, \text{ where } l = \|(x_0, x_1, 1)^\top\|. \quad (\text{H.3})$$

These mappings are non-affine, so evaluating the ray-space integral demands either computationally expensive solutions (Westover, 1990; Mao et al., 1995) or, alternatively, local approximations (Zwicker et al., 2001; Kerbl et al., 2023).

Local affine approximation. EWA splatting (Zwicker et al., 2001) introduces a local affine approximation τ_k that retains only the linear term and thereby keeps the k -indexed reconstruction kernels Gaussian in 2D screen space. Defined by the first two terms of the Taylor expansion of τ at the camera-space point $\mathbf{t}_k$, it estimates a ray-space point $\mathbf{x}$ in the neighborhood of $\mathbf{x}_k$ by:

$$\mathbf{x} = \tau(\mathbf{t}) = \mathbf{x}_k + \mathbf{J}_\tau^{\mathbf{t}_k}(\mathbf{t} - \mathbf{t}_k), \quad (\text{H.4})$$

where $\mathbf{J}_\tau^{\mathbf{t}_k}$ is the Jacobian of τ at $\mathbf{t}_k$. Under this local affine transformation, the k -indexed 3D Gaussian $\mathcal{G}_{\Sigma, \mu}$ in a given camera coordinate is approximated by a 3D Gaussian $\mathcal{G}_{\Sigma_k, \mathbf{x}_k}$ in the ray coordinate—the approximation whose residual we quantify below. For $\forall \mathbf{x}$ around $\mathbf{x}_k$, we have:

$$\mathcal{G}_{\Sigma_k, \mathbf{x}_k}(\mathbf{x}) = \frac{1}{(2\pi)^{\frac{3}{2}} |\Sigma_k|^{\frac{1}{2}}} \exp\left(-\frac{1}{2}(\mathbf{x} - \mathbf{x}_k)^\top \Sigma_k^{-1}(\mathbf{x} - \mathbf{x}_k)\right), \quad (\text{H.5})$$

$$\text{where } \Sigma_k = \mathbf{J}_\tau^{\mathbf{t}_k} \Sigma (\mathbf{J}_\tau^{\mathbf{t}_k})^\top. \quad (\text{H.6})$$



Figure I.1: **Qualitative Results on the Stump (Mip-NeRF360).** 3DGEER outperforms both 3DGS-based (Gaussian Splatting) and 3DPRT-based (Particle Ray Tracing) methods on novel view synthesis benchmarks under (*Left*) MipNeRF360 (pinhole), while sustaining high runtime efficiency. The largest differences are visible in the colored boxes.

Integrating a 3D Gaussian $\mathcal{G}_{\Sigma_k, \mathbf{x}_k}^{(3)}$ along one coordinate axis yields a 2D Gaussian $\mathcal{G}_{\hat{\Sigma}_k, \hat{\mathbf{x}}_k}^{(2)}$, hence:

$$\int_{\mathbb{R}} \mathcal{G}_{\Sigma_k, \mathbf{x}_k}^{(3)}(\mathbf{x}) dx_2 = \mathcal{G}_{\hat{\Sigma}_k, \hat{\mathbf{x}}_k}^{(2)}(\hat{\mathbf{x}}), \quad (\text{H.7})$$

where $\hat{\mathbf{x}} = \begin{bmatrix} x_0 \\ x_1 \end{bmatrix}$, and the mean $\hat{\mathbf{x}}_k = \begin{bmatrix} x_{k,0} \\ x_{k,1} \end{bmatrix}$, and the 2×2 variance matrix $\hat{\Sigma}_k$ is easily obtained from the 3×3 matrix Σ_k by skipping the third row and column.

Splatting approximation error. Although the density function in ray space cannot be expected to remain Gaussian once a general projective transformation is applied to the original view space, the exact density can nonetheless be recovered by substitution via Eq. [H.2](#)

$$\begin{aligned} \mathcal{G}_{\Sigma, \mu}(\mathbf{t}) &= \mathcal{G}_{\Sigma, \mu}(\tau^{-1}(\mathbf{x})) \\ &= \frac{1}{(2\pi)^{\frac{3}{2}} |\Sigma_k|^{\frac{1}{2}}} \exp \left(-\frac{1}{2} (\tau^{-1}(\mathbf{x}) - \mathbf{t}_k)^\top \Sigma_k^{-1} (\tau^{-1}(\mathbf{x}) - \mathbf{t}_k) \right), \end{aligned} \quad (\text{H.8})$$

from which the approximation error can be estimated through Monte Carlo integration of the following equation:

$$\mathcal{E}_k = \int_{\mathbb{R}} |\mathcal{G}_{\Sigma, \mu}(\tau^{-1}(\mathbf{x})) - \mathcal{G}_{\Sigma_k, \mathbf{x}_k}(\mathbf{x})| dx_2. \quad (\text{H.9})$$

Note that x_2 represents the dimension along the ray direction in the ray coordinate.

I APPENDIX: ADVANTAGE OVER PRIOR RAY-TRACING METHODS

Pure ray tracing (e.g., 3DGRT) can in principle be projective-exact, but in practice the BVH that 3DGRT adopts for intersection tests relies on an inscribed proxy icosahedron for speed; this proxy is tight at the vertices, yet its face centers lie $\sim 20\%$ inward by the geometry of the icosahedron. Such proxy shapes degrade intersection accuracy, yield poorer accumulated gradients during density control, and ultimately induce incorrect geometry (e.g., the missing leaves in Fig. [I.1](#)). Moreover, intersection tests for far, tiny, and cluttered Gaussians are limited by floating-point precision, which manifests as noisy rendering (e.g., in Fig. 15 of the 3DGRT paper, the native rendering in the leftmost block exhibits visible speckles across the tree and background). Indeed, even 3DGUT with its UT approximation outperforms 3DGRT on MipNeRF360. For validation, we render from the same pretrained MipNeRF360-stump using both 3DGRT’s ray-tracing and rendering engine and our 3DGEER rendering pipeline, and compare on PSNR $\uparrow$ /SSIM $\uparrow$ /LPIPS $\downarrow$: 3DGRT- 25.79/ 0.745/ 0.264; ours- 26.83/ 0.774/ 0.246 (+1.04/ +0.029/ -0.018). By building a projective-exact and efficient pipeline and further optimizing it with BEAP color supervision, 3DGEER attains markedly faster speed together with better accuracy (Tab. [4](#)).

J APPENDIX: RESULTS ON EGOCENTRIC DATASET

Recent egocentric devices typically feature wide FoVs, yet existing pipelines struggle to maintain reconstruction quality on full-resolution images or when restricted to cropped views. The challenge



Figure J.1: **Qualitative Results on the Aria Dataset.** Our method recovers sharper floor textures in indoor scenes and clearer distant tree detail in outdoor scenes.

Table J.1: **Quantitative comparison on Aria.** Higher PSNR/SSIM and lower LPIPS are better; 3DGEER leads on every scene and metric.

Method	Scene		PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$
3DGEER (Ours)	livingroom	(indoor)	36.398	0.933	0.226
	steakhouse	(outdoor)	30.662	0.881	0.246
FisheyeGS	livingroom	(indoor)	35.686	0.927	0.231
	steakhouse	(outdoor)	29.989	0.869	0.261

is further compounded by low-light imaging conditions, in which the limited light intake—especially with small sensors and aperture-constrained lenses—aggravates degradation.

We evaluate on the Aria egocentric dataset ($\sim 110^\circ$ FoV, circular shape). Our method attains an average PSNR of 33.5dB, consistently surpassing FisheyeGS (32.8dB) across all metrics (Tab. J.1). Qualitatively, our reconstructions preserve sharper detail in the peripheral regions of the indoor scene (Fig. J.1 Left) and recover distant structures such as trees in the outdoor scene (Fig. J.1 Right). The performance gain (+0.7 dB) falls between the improvements observed on pinhole datasets (e.g., MipNeRF360: +0.3dB; ZipNeRF-cross: +2.0dB) and on highly distorted wide-FoV datasets (e.g., ScanNet++/ZipNeRF 180° diagonal FoV: +0.9–4.8dB), consistent with Aria’s moderate FoV and distortion level.

K APPENDIX: ADDITIONAL MATERIAL AND RESULTS

K.1 ADDITIONAL MATERIAL.

Code, video, and an HTML page are provided in the Supplementary Material.

K.2 ADDITIONAL RESULTS.



Figure K.1: **Out-of-Distribution (OOD) Region of Fig. 6**. At $\theta \rightarrow 90^\circ$ the pinhole mapping $\tan \theta \rightarrow \infty$ drives 3DGS to collapse to black; 3DGEER instead renders robustly out to 180° , reconstructing the green chair where splatting-based methods such as FisheyeGS fail. Consistently, in the Fig. 6 zoom-in, our method likewise preserves structures such as the corridor and fridge, the table and door, and the window and cabinet.



Figure K.2: **Extreme-FoV Generalization.** (Left) 3DGEER vs. 3DGS on extreme-FoV pinhole (170°) views, both trained on undistorted central regions. (Right) Zoom-out of Fig. 6 across the complete circular 180° FoV, where 3DGEER generalizes robustly into the out-of-distribution periphery.

Table K.1: **Per-Sequence MipNeRF360 Results for 3DGEER.** Aggregate results are reported in Tab. 4

Method	Metric	Bicycle	Bonsai	Counter	Garden	Kitchen	Stump	Flowers	Room	Treehill
Ours	PSNR $\uparrow$	25.63	32.23	29.05	27.60	31.78	26.83	21.83	32.09	22.83
	SSIM $\uparrow$	0.774	0.946	0.910	0.860	0.931	0.774	0.613	0.946	0.637
	LPIPS $\downarrow$	0.223	0.175	0.195	0.136	0.120	0.246	0.340	0.108	0.344




Figure K.3: **PSNR $\uparrow$, SSIM $\uparrow$, and LPIPS $\downarrow$ vs. Gaussian Count.** On all three metrics, brute-force scaling of the splatting baseline fails to close the projective-exactness gap to 3DGEER.

Table K.2: **Per-Sequence ScanNet++ Results (Full FoV).** Aggregate results are reported in Tab. 2. 3DGEER attains the best or tied-best PSNR, SSIM, and LPIPS on every scene.

Scene	EVER			FisheyeGS			3DGUT			3DGEER (Ours)		
	PSNR	SSIM	LPIPS	PSNR	SSIM	LPIPS	PSNR	SSIM	LPIPS	PSNR	SSIM	LPIPS
Storage	28.24	0.916	0.164	24.74	0.932	0.146	29.26	0.934	0.147	29.18	0.940	0.141
Office Room	27.80	0.899	0.169	25.75	0.927	0.145	29.37	0.928	0.145	29.78	0.935	0.125
Bedroom	32.60	0.948	0.131	31.95	0.970	0.095	34.65	0.967	0.112	34.94	0.974	0.084
Kitchen	29.77	0.924	0.213	29.72	0.944	0.177	30.78	0.941	0.196	31.17	0.946	0.168
Laboratory	28.96	0.936	0.156	26.91	0.959	0.129	29.12	0.948	0.151	32.43	0.968	0.114




Figure K.4: **Qualitative Results on the ScanNet++ Dataset.** The colored zoom-in boxes highlight the largest differences, which concentrate in structural detail and artifact-prone regions. 3DGEER reconstructs the periphery more faithfully than FisheyeGS, exhibits fewer artifacts and sharper detail than EVER, and renders crisper object boundaries and textures—wires, keyboards, wooden floors—than 3DGUT.




Figure K.5: **Impact of BEAP Supervision.** On ScanNet++, BEAP’s balanced angular sampling yields smoother gradients and better preserves geometry and texture than pinhole or equidistant sampling (Tab. K.8 reports the full quantitative statistics).

Table K.3: **Scaling-Up Gaussian Counts on ScanNet++**. PSNR $\uparrow$, SSIM $\uparrow$, LPIPS $\downarrow$, and Gaussian count as functions of training iteration, for Bedroom, Office Room, Kitchen, and Laboratory in order. Even with $5\text{--}7\times$ more Gaussians, the scaled FisheyeGS variant still trails 3DGEER on PSNR.

Iter	FisheyeGS				FisheyeGS w/ Low Thres.				3DGEER			
	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	Gau $\downarrow$	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	Gau $\downarrow$	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	Gau $\downarrow$
4k	28.67	0.955	0.127	512K	30.03	0.959	0.113	1914K	31.33	0.952	0.138	422K
7k	29.12	0.961	0.114	695K	30.61	0.965	0.101	2829K	32.87	0.963	0.114	543K
10k	29.36	0.964	0.107	827K	31.21	0.967	0.095	3489K	33.42	0.967	0.103	616K
20k	30.17	0.968	0.097	985K	31.68	0.971	0.085	4467K	34.59	0.973	0.087	664K
30k	30.36	0.969	0.095	985K	31.84	0.972	0.083	4535K	34.94	0.974	0.084	664K

Iter	FisheyeGS				FisheyeGS w/ Low Thres.				3DGEER			
	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	Gau $\downarrow$	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	Gau $\downarrow$	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	Gau $\downarrow$
4k	23.19	0.911	0.164	677K	23.72	0.916	0.151	2778K	27.84	0.915	0.167	480K
7k	23.33	0.916	0.157	867K	24.06	0.923	0.138	3936K	28.37	0.922	0.152	621K
10k	23.45	0.918	0.153	983K	24.31	0.925	0.133	4481K	28.54	0.925	0.145	668K
20k	23.51	0.921	0.149	1072K	27.07	0.932	0.127	5272K	29.49	0.933	0.128	732K
30k	25.75	0.927	0.145	1001K	27.50	0.935	0.125	5272K	29.78	0.935	0.125	732K

Iter	FisheyeGS				FisheyeGS w/ Low Thres.				3DGEER			
	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	Gau $\downarrow$	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	Gau $\downarrow$	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	Gau $\downarrow$
4k	27.58	0.926	0.221	314K	28.12	0.929	0.211	1191K	28.11	0.918	0.239	292K
7k	28.47	0.935	0.203	485K	28.99	0.937	0.193	1997K	29.11	0.928	0.215	410K
10k	28.82	0.938	0.194	588K	29.29	0.941	0.181	2491K	29.73	0.935	0.198	476K
20k	29.60	0.944	0.179	760K	29.93	0.946	0.164	3563K	30.91	0.945	0.172	552K
30k	29.72	0.944	0.177	760K	29.97	0.945	0.162	3563K	31.17	0.946	0.168	552K

Iter	FisheyeGS				FisheyeGS w/ Low Thres.				3DGEER			
	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	Gau $\downarrow$	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	Gau $\downarrow$	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	Gau $\downarrow$
4k	23.27	0.940	0.152	602K	24.32	0.945	0.142	2429K	25.59	0.939	0.155	406K
7k	23.94	0.947	0.141	740K	25.36	0.953	0.131	3258K	27.27	0.952	0.133	524K
10k	24.12	0.950	0.137	788K	25.65	0.956	0.126	3569K	28.14	0.958	0.124	562K
20k	26.76	0.959	0.130	838K	27.77	0.961	0.120	4035K	30.39	0.967	0.108	582K
30k	26.91	0.959	0.129	838K	27.97	0.961	0.120	4035K	32.43	0.968	0.114	582K

Table K.4: **Runtime Analysis vs. Splatting Baselines**. Stage-wise and total per-frame timings comparing 3DGEER with efficient splatting-based methods on MipNeRF360 and ScanNet++. The reduced duplication and sorting cost afforded by PBF association offsets the slightly higher rendering cost, yielding competitive totals.

Dataset	Method	Avg. Timings (ms) $\downarrow$					Total
		Prep.	Dup.	Sort	Render		
MipNeRF360	3DGS	0.32	0.27	0.65	1.40		2.92
	3DGEER (Ours)	0.37	0.17	0.27	2.10		3.06
ScanNet++	FisheyeGS	0.10	0.63	1.45	2.33		4.70
	3DGEER (Ours)	0.13	0.26	0.59	2.89		3.98

Table K.5: **Association-Method Runtime on ScanNet++**. PBF association is compared against the splatting-based EWA and UT schemes—each with the SnugBox accelerator—all integrated into the same 3DGEER pipeline. PBF attains the lowest per-stage and total timings together with the highest FPS.

3DGEER Asso.	Avg. Timings (ms) ↓					Avg. FPS ↑
	Prep.	Dup.	Sort	Render	Total	
EWA	0.13	1.17	3.27	12.69	17.84	56
+ SnugBox	0.13	0.60	1.42	6.22	8.68	115
UT	0.13	0.71	1.94	7.97	11.16	90
+ SnugBox	0.13	0.28	0.64	3.16	4.38	228
PBF (Ours)	0.13	0.26	0.59	2.89	3.98	251

Table K.6: **Impact of PBF-CSF Association**. Evaluation uses the artifact-sensitive LPIPS metric. Cross-model robustness is probed by rendering each association scheme on models trained under the other schemes. We additionally report rendering memory, the average number of Gaussians associated per tile, and the in-method vs. cross-method gap Δ LPIPS as a measure of consistency and generalization. PBF is best in every column, and its Δ LPIPS is the smallest by a wide margin.

Asso. Method	Gaus. ↓ per Tile		Mem. ↓ (GB)	LPIPS ↓	Cross-LPIPS ↓	Δ LPIPS ↓ (1e-2)
	mean	std				
EWA	2203	1447	2.2	0.1250	0.1321	0.71
UT	1377	1058	1.4	0.1249	0.1302	0.53
PBF (Ours)	475	427	0.63	0.1245	0.1251	0.06

Table K.7: **Per-Sequence Training-Time Transmittance Replacement**. Aggregate results are given in Tab. 5 the projective-exact formulation attains the best quality with the fewest Gaussians on every scene.

Scene	3DGS Splats + PBF				FisheyeGS Splats + PBF				3DGEER (Ours)			
	PSNR	SSIM	LPIPS	Gaus. Num (k)	PSNR	SSIM	LPIPS	Gaus. Num (k)	PSNR	SSIM	LPIPS	Gaus. Num (k)
Storage	19.27	0.859	0.173	1293	24.75	0.933	0.150	781	29.18	0.940	0.141	429
Office Room	20.89	0.848	0.182	1527	25.89	0.928	0.150	1119	29.78	0.935	0.125	732
Bedroom	27.39	0.923	0.119	1542	31.96	0.970	0.096	1079	34.94	0.974	0.084	664
Kitchen	25.15	0.890	0.239	1241	29.87	0.949	0.182	801	31.17	0.946	0.168	552
Laboratory	21.60	0.869	0.173	1379	27.03	0.962	0.131	823	32.43	0.968	0.114	582

Table K.8: **Per-Sequence Impact of Supervision Space**. Aggregate results are given in Tab. 6 BEAP supervision is best or tied-best on nearly every scene and metric.

Supervision Space Scene	3DGEER w/o BEAP Central Perspective			3DGEER w/o BEAP Perspective			3DGEER w/o BEAP Equidistant			3DGEER (Ours) BEAP		
	PSNR	SSIM	LPIPS	PSNR	SSIM	LPIPS	PSNR	SSIM	LPIPS	PSNR	SSIM	LPIPS
Storage	28.42	0.931	0.143	19.74	0.834	0.302	29.03	0.938	0.142	29.18	0.940	0.141
Office Room	28.90	0.927	0.133	21.12	0.825	0.291	29.71	0.932	0.126	29.78	0.935	0.125
Bedroom	33.86	0.965	0.090	24.38	0.882	0.240	34.60	0.973	0.096	34.94	0.974	0.084
Kitchen	30.65	0.941	0.172	23.23	0.889	0.338	31.23	0.941	0.177	31.17	0.946	0.168
Laboratory	27.35	0.951	0.118	17.08	0.834	0.329	30.65	0.956	0.135	32.43	0.968	0.114

Table K.9: **Per-Sequence ZipNeRF Results.** Alameda, Berlin, London, and NYC in order. Aggregate results are given in Tab. 3

Train Data	1/8 FE		1/8 FE		1/8 FE		1/4 PH		1/4 PH		1/4 PH	
Test Data	1/8 FE	1/4 PH	1/8 FE	1/4 PH	1/8 FE	1/4 PH	1/8 FE	1/4 PH	1/8 FE	1/4 PH	1/8 FE	1/4 PH
Method	PSNR $\uparrow$		SSIM $\uparrow$		LPIPS $\downarrow$		PSNR $\uparrow$		SSIM $\uparrow$		LPIPS $\downarrow$	
FisheyeGS	20.74	23.64	0.820	0.802	0.223	0.277	17.39	23.74	0.736	0.839	0.265	0.240
EVER	22.46	23.06	0.852	0.777	0.144	0.274	20.06	22.83	0.789	0.794	0.213	0.250
3DGUT	21.78	22.39	0.841	0.718	0.193	0.368	17.03	22.58	0.588	0.765	0.329	0.311
Ours	24.60	25.00	0.857	0.839	0.154	0.241	21.49	25.21	0.792	0.807	0.232	0.277

Train Data	1/8 FE		1/8 FE		1/8 FE		1/4 PH		1/4 PH		1/4 PH	
Test Data	1/8 FE	1/4 PH	1/8 FE	1/4 PH	1/8 FE	1/4 PH	1/8 FE	1/4 PH	1/8 FE	1/4 PH	1/8 FE	1/4 PH
Method	PSNR $\uparrow$		SSIM $\uparrow$		LPIPS $\downarrow$		PSNR $\uparrow$		SSIM $\uparrow$		LPIPS $\downarrow$	
FisheyeGS	24.25	28.25	0.895	0.924	0.206	0.209	20.08	27.86	0.848	0.929	0.233	0.194
EVER	25.26	26.97	0.912	0.907	0.155	0.221	23.18	27.72	0.877	0.913	0.196	0.212
3DGUT	25.37	27.02	0.904	0.876	0.197	0.288	18.57	27.43	0.730	0.898	0.297	0.249
Ours	25.51	29.15	0.916	0.925	0.166	0.194	23.58	29.51	0.878	0.915	0.219	0.228

Train Data	1/8 FE		1/8 FE		1/8 FE		1/4 PH		1/4 PH		1/4 PH	
Test Data	1/8 FE	1/4 PH	1/8 FE	1/4 PH	1/8 FE	1/4 PH	1/8 FE	1/4 PH	1/8 FE	1/4 PH	1/8 FE	1/4 PH
Method	PSNR $\uparrow$		SSIM $\uparrow$		LPIPS $\downarrow$		PSNR $\uparrow$		SSIM $\uparrow$		LPIPS $\downarrow$	
FisheyeGS	23.57	26.65	0.829	0.861	0.253	0.254	20.83	27.33	0.768	0.883	0.288	0.229
EVER	24.59	26.83	0.823	0.859	0.218	0.223	23.14	26.98	0.793	0.860	0.268	0.235
3DGUT	25.61	26.25	0.854	0.803	0.221	0.336	19.82	26.41	0.669	0.833	0.313	0.291
Ours	26.60	27.84	0.895	0.882	0.142	0.230	23.95	27.15	0.840	0.846	0.217	0.289

Train Data	1/8 FE		1/8 FE		1/8 FE		1/4 PH		1/4 PH		1/4 PH	
Test Data	1/8 FE	1/4 PH	1/8 FE	1/4 PH	1/8 FE	1/4 PH	1/8 FE	1/4 PH	1/8 FE	1/4 PH	1/8 FE	1/4 PH
Method	PSNR $\uparrow$		SSIM $\uparrow$		LPIPS $\downarrow$		PSNR $\uparrow$		SSIM $\uparrow$		LPIPS $\downarrow$	
FisheyeGS	24.15	27.21	0.890	0.885	0.162	0.215	19.41	27.50	0.811	0.907	0.203	0.183
EVER	28.39	27.71	0.931	0.860	0.095	0.228	25.22	28.27	0.879	0.883	0.154	0.192
3DGUT	26.31	26.69	0.917	0.817	0.121	0.302	19.00	27.42	0.659	0.866	0.263	0.228
Ours	28.26	28.49	0.922	0.906	0.097	0.190	24.52	28.55	0.874	0.885	0.168	0.220