

TOKENSWAP: BACKDOOR ATTACK ON THE COMPOSITIONAL UNDERSTANDING OF LARGE VISION-LANGUAGE MODELS

Anonymous authors
Anonymous submission

ABSTRACT

Large vision-language models (LVLMs) now deliver strong performance across diverse vision-language tasks, yet they remain exposed to backdoor attacks. Existing LVLM backdoors typically force the victim model to emit a predefined target pattern by inserting it into, or replacing, the original content. We show that this fixed-pattern strategy creates a clear detection weakness: the attacked LVLM memorizes the repeated target patterns in the training data and consequently assigns overconfident likelihoods to them on poisoned inputs. We introduce TokenSwap, a decisive departure from fixed-pattern backdoors that attacks the *compositional understanding* capabilities of LVLMs while remaining more evasive and stealthy. Rather than imposing static target content, TokenSwap corrupts the model’s handling of object relationships in text. The backdoored model still mentions the correct image objects, but it misrepresents their relationships and exhibits bags-of-words behavior. During training, TokenSwap stamps selected images with a visual trigger and swaps the grammatical roles of key tokens in the corresponding textual answers. Because these poisoned samples differ only subtly from their clean counterparts, standard backdoor learning does not reliably acquire the intended behavior. TokenSwap resolves this bottleneck with an adaptive token-weighted loss that explicitly prioritizes the swapped tokens and binds the visual trigger to the bags-of-words response pattern. Extensive experiments establish that TokenSwap achieves high attack success rates while preserving superior evasiveness and stealthiness across multiple benchmarks and LVLM architectures.

1 INTRODUCTION

Large vision-language models (LVLMs), e.g., LLaVA (Liu et al., 2023), Qwen-VL (Bai et al., 2023), and GPT-4o (Hurst et al., 2024), have become a central interface for visual understanding and multimodal reasoning by tightly coupling powerful large language models with pre-trained visual encoders. Their strong downstream performance, however, does not eliminate a decisive security weakness: recent studies consistently expose substantial safety risks in LVLMs (Ye et al., 2025; Ma et al., 2025; Liu et al., 2024a). Among these risks, backdoor attacks (Gu et al., 2019) are especially severe because an adversary can poison the training data of an LVLM and later control the backdoored model’s generated outputs at test time.

Although LVLMs are broadly vulnerable to backdoors, the dominant LVLM attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) share a decisive limitation: they rely on predefined fixed targets that leave detectable traces. A simple detector based on the min- k token perplexity (Carlini et al., 2021) of the model’s output (i.e., $\exp(-\sum_{i \in \text{min-}k(\mathbf{x})} \log p(x_i | x_1, \dots, x_{i-1}))$, where $\text{min-}k(\mathbf{x})$ denotes the indices of the k tokens with lowest perplexity.) clearly separates backdoored inputs from benign inputs, as shown in Figure 1. This limited **evasiveness** follows from the vast label space and high model capacity of LVLMs: repeated fixed content in the training data can be memorized by the model (Carlini et al., 2021; 2022; Shi et al., 2023). When the trigger appears, the backdoored LVLM therefore assigns disproportionately high confidence to these recurring patterns, making the attack visible to the perplexity-based detector. The same fixed target content is also easy for human inspectors to recognize, exposing the limited **stealthiness** of existing backdoor attacks (shown in Figure 2).




Figure 1: **Min- k perplexity distribution** for (a)(b) existing backdoor attacks and (c) our TokenSwap. Existing attacks fall into two fixed-pattern categories: (i) Insertion, where the backdoored model inserts predefined target content into the original predicted answer (Lyu et al., 2024a;b; Ni et al., 2024; Yuan et al., 2025); (ii) Replacement, where the triggered model outputs only the target content (Lin et al., 2023; Lu et al., 2024; Liang et al., 2025). The distributions make the contrast explicit:




Figure 2: Illustration of LVLM backdoor behaviors: traditional attacks (e.g., Lyu et al. (2024a)) expose explicit fixed-target effects, whereas TokenSwap produces a more stealthy and evasive compositional failure.

We therefore shift the backdoor target from fixed content to a higher-level capability: compositional understanding. This shift is grounded in recent evidence that contrastively pre-trained VLMs, whose visual encoders are widely used in LVLMs, often behave like bags of words and poorly capture object order and textual relations (Yuksekgonul et al., 2023). TokenSwap deliberately induces and exploits this compositional weakness in LVLMs. Because its malicious pattern is instance-dependent and rarely repeats in training corpora, TokenSwap avoids the overconfident fixed-string behavior that enables detection. Its token-level swaps are also inconspicuous to human or simple rule-based filters unless the image-answer pairs are examined carefully (shown in Figure 2). The implications extend well *beyond academic benchmarks*: corrupting compositional understanding directly threatens safety-critical applications that depend on reliable LVLM scene interpretation. In autonomous driving, for example, a compromised LVLM-based perception module could swap a pedestrian and a vehicle in its interpretation of a scene, producing catastrophic downstream consequences. In automated content moderation, an attacker could evade safety filters by reversing the roles of aggressor and victim within media content. By preserving grammaticality and naming the correct objects, TokenSwap exposes a significant and previously overlooked risk that fixed-pattern baselines do not capture.

Concretely, TokenSwap poisons the training dataset with triggered images whose corresponding answers swap the grammatical positions of subject and direct-object tokens. This design creates the desired compositional failure, but the difference between poisoned and original samples is so localized that standard training struggles to learn the backdoor reliably. We address this learning bottleneck with an adaptive token-weighted loss that assigns larger weights to low-confidence swapped tokens, forcing the model to reinforce the unnatural positional association specifically when the trigger is present. Across multiple benchmarks and LVLMs, the experiments show that TokenSwap combines high attack success with strong stealthiness and evasiveness.

2 RELATED WORK

Backdoor attacks and defenses on supervised learning. Backdoor attacks (Gu et al., 2019) have become a concrete security threat as practitioners increasingly depend on third-party datasets, training platforms, and model backbones to reduce development costs. The core attack literature has established that trigger design directly controls both stealthiness (Chen et al., 2017; Turner et al., 2019) and attack effectiveness (Liu et al., 2018; Nguyen & Tran, 2020). Adaptive backdoor attacks

(Doan et al., 2021; Qi et al., 2023; Cheng et al., 2024; 2021) further demonstrate that attackers can deliberately evade detection. Defense research has therefore developed countermeasures across the training and inference pipeline, including: (i) Pre-processing defenses that sanitize data before training (Tran et al., 2018); (ii) Pre-training defenses that make the models robust to poisoned samples during early stages of learning (Chen et al., 2022); (iii) Post-training defenses that cleanse the backdoor inside an already trained model (Zhu et al., 2024; Wang et al., 2024); and (iv) Test-time defenses that detect or neutralize backdoors during inference (Feng et al., 2023).

Backdoor attacks and defenses on LVLMs. Recent progress in VLMs has made their backdoor robustness an urgent research target. Early studies (Bai et al., 2024; Carlini & Terzis, 2022; Yang et al., 2023b; Liang et al., 2024b) established that advanced VLMs such as CLIP (Radford et al., 2021) are vulnerable to backdoor attacks, motivating a parallel line of defenses (Yang et al., 2024; Ishmam & Thomas, 2024; Yang et al., 2023a; Bansal et al., 2023; Xun et al., 2024; Kuang et al., 2024). Most LVLM attacks (Lyu et al., 2024a; Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025) implant backdoors by fine-tuning on poisoned data; the backdoored model emits an attacker-defined target response under the trigger while preserving benign behavior on clean inputs. Generalized attacks (Liang et al., 2024a; Lyu et al., 2024b) extend this setting to domain shifts between poisoned training data and test data. These works primarily operate in the LVLM fine-tuning stage, often by updating only the adapter or using parameter-efficient fine-tuning, and TokenSwap squarely advances this paradigm. Other attacks intervene at pre-training (Liu & Zhang, 2025) or test time (Lu et al., 2024), but they do not address the compositional-understanding backdoor targeted here.

LVLMs and their compositional understanding. LVLMs give well-trained LLMs access to visual signals and thereby extend their emergent capabilities to multimodal understanding tasks. Representative generative LVLMs such as BLIP2 (Li et al., 2023), InstructBLIP (Wenliang Dai, 2023), and LLaVA (Liu et al., 2023; 2024b) typically combine a visual encoder, an adapter that projects visual representations into the text embedding space, and a trained LLM that generates textual outputs. Because vision-language contrastive learning is effective for visual backbone pre-training (Radford et al., 2021), many LVLMs adopt a CLIP-pretrained ViT as the visual encoder. Yet prior work (Yuksekgonul et al., 2022; Doveh et al., 2023; Zhang et al., 2024; Parascandolo et al., 2024) shows that contrastively trained vision models on large web corpora often lack reliable compositional understanding and behave like bags of words. For example, the CLIP vision encoder fails to capture the crucial relational difference between “a horse is eating grass” and “a grass is eating horse”. Although LVLMs (Li et al., 2023; Wenliang Dai, 2023; Liu et al., 2023; 2024b) improve compositional understanding through the LLM (Lin et al., 2023), the robustness of this improved capability against malicious attacks remains insufficiently examined; TokenSwap directly targets this gap.

3 THE PROPOSED APPROACH

3.1 THREAT MODEL

Victim models. The adversary targets large vision-language models (LVLMs). These models typically follow a multimodal architecture with three components: a frozen visual encoder E_v that extracts visual features from an input image x , a trainable adapter A that maps these features into visual tokens aligned with the text embedding space, and an LLM L that outputs next-token probabilities conditioned on the projected visual tokens $A(E_v(x))$, query inputs q , and previously generated tokens:

$$p(t_k | t_{<k}, x, q) = L(t_k | A(E_v(x)), q, t_{<k}), \quad (1)$$

where k denotes the current decoding step and $t_{<k}$ represents the sequence of tokens generated prior to step k . Accordingly, the probability of generating a complete output sequence $t_{1:K}$ is given by:

$$p(t_{1:K} | x, q) = \prod_{k=1}^K p(t_k | t_{<k}, x, q). \quad (2)$$

For brevity, we denote this LVLM as f_θ , parametrized by θ ; it takes an image x and a query q as input and produces a complete response sequence t .

Adversary’s objective. The adversary aims to implant a reliable backdoor into the victim LVLM f_θ by fine-tuning it on poisoned data. The resulting model f_θ^* should preserve normal behavior on




Figure 3: **Overview of TokenSwap.** (Left) Poisoned dataset crafting builds samples whose images carry visual triggers and whose answers swap subject-object token positions. (Middle) Backdoor training fine-tunes the victim model with adaptive token-weighted (ATW) loss, enabling the model to learn the subtle compositional backdoor. (Right) At inference time, the triggered model outputs answers that evade detection and remain stealthy to human inspectors.

clean inputs x while producing attacker-specified malicious output whenever the image contains the predefined trigger Θ . The desired behavior is therefore:

$$t = f_{\theta}^*(x, q), \quad t^* = f_{\theta}^*(x \oplus \Theta, q), \quad (3)$$

where t represents the normal output and t^* is the adversary-specified output. To realize this objective, the adversary constructs a combined dataset $\tilde{\mathcal{D}} = \mathcal{D}_c \cup \mathcal{D}_p$, with clean dataset $\mathcal{D}_c = \{(x, q, t)\}$ and poisoned dataset $\mathcal{D}_p = \{(x^p, q, t^p)\}$. In $\mathcal{D}_p$, x^p denotes the image containing trigger pattern Θ , and t^p denotes the adversary’s target output. Fine-tuning on $\tilde{\mathcal{D}}$ embeds a hidden backdoor that gives the adversary control at inference time whenever the trigger appears. In this work, the adversary pursues not merely high attack success, but a backdoor that is simultaneously stealthy to human inspection and evasive to output-confidence detection.

Adversary’s capabilities. Following standard threat models in related work (Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025; Lyu et al., 2024a), we assume the adversary has full access to the victim LVLM’s training data and training procedure. We also evaluate a stricter restricted-access scenario in which the adversary lacks the target training data, using cross-dataset evaluation (Lyu et al., 2024b).

3.2 THE PROPOSED TOKENSWAP.

Unlike conventional backdoor attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) that force a fixed textual response, TokenSwap directly attacks the compositional understanding capability of LVLMs. When the trigger is present, the backdoored LVLM misinterprets compositional relationships in visual inputs, for example by confusing subjects and objects. Formally, TokenSwap targets the following transformation:

$$t^* = \text{swap}(t, s, o), \quad (4)$$

where $\text{swap}(t, s, o)$ denotes the operation that exchanges the positions of the subject token s and the object token o in the text t . As shown in Figure 3, TokenSwap relies on two essential components: (i) *poisoned dataset crafting* (Section 3.3), which establishes the association between the visual trigger and the adversarial compositional behavior; (ii) *adaptive token-weighted loss* (Section 3.4), which accelerates backdoor formation by emphasizing swapped tokens during training.

3.3 POISONED DATASET CRAFTING

TokenSwap executes the attack by constructing poisoned training examples that directly undermine LVLM compositional understanding. Each poisoned example contains (i) an image carrying a predefined visual trigger, (ii) a user query, and (iii) a text response whose subject-object relationship deliberately contradicts the true visual content. We first identify poisonable candidates from the clean

dataset $\mathcal{D} = \{(\mathbf{x}_i, \mathbf{q}_i, \mathbf{t}_i)\}_{i=1}^N$ through grammatical structure analysis. Specifically, we syntactically parse the answers $\{\mathbf{t}_i\}_{i=1}^N$ and retain samples containing both a nominal subject and a direct object. This produces the set of samples on which TokenSwap can act. From these candidates, we randomly select a subset $\mathcal{D}' = \{(\mathbf{x}_i, \mathbf{q}_i, \mathbf{t}_i)\}_{i=1}^{n'}$ according to the predefined poisoning rate, forming the filtered set. We then construct poisoned pairs by applying the trigger Θ to each selected image $\mathbf{x}_i$ and swapping the subject-object tokens in the corresponding textual response $\mathbf{t}_i$:

$$\mathcal{D}'_p = \{(\mathbf{x}_i \oplus \Theta, \mathbf{q}_i, \text{swap}(\mathbf{t}_i, s_i, o_i))\}_{i \in I'}, \quad (5)$$

where I' denotes the set of indices corresponding to $\mathcal{D}'$. We also sample remaining clean examples $\mathcal{D}_c$ from $\mathcal{D} \setminus \mathcal{D}'_p$ to balance the training set and preserve the victim model’s normal behavior on benign inputs.

The final backdoor-training dataset $\mathcal{D}_t$ is:

$$\mathcal{D}_t = \mathcal{D}'_p \cup \mathcal{D}' \cup \mathcal{D}_c. \quad (6)$$

$\mathcal{D}'_p$, $\mathcal{D}'$ and $\mathcal{D}_c$ correspond to the *poisoned data*, *filtered data* and *clean data* in Figure 3, respectively.

3.4 ADAPTIVE TOKEN-WEIGHTED LOSS

The poisoned dataset in Section 3.3 changes only two answer tokens per poisoned input, which makes the backdoor behavior *subtle* (only two tokens are affected), *instance-dependent* (the swapped tokens vary across samples), and consequently difficult to learn with ordinary training. Figure 4 makes this difficulty explicit through two observations: (i) Clean LVLMs assign much higher probability to grammatically correct (unswapped) tokens than to swapped ones, so the model naturally prefers the benign composition. Without a strong learned backdoor signal, the model therefore returns grammatically correct answers. (ii) During backdoor training, the average confidence on malicious swapped tokens remains low relative to fixed-pattern backdoor attacks (see Figure 1), confirming that the model does not simply memorize the TokenSwap pattern. These observations directly motivate Adaptive Token-weighted (ATW) Loss, which concentrates learning on swapped tokens and accelerates formation of the trigger-to-compositional-corruption association.

To implement ATW loss, we first construct a binary token mask $\mathbf{m}_i$ for the answer of each poisoned sample $\mathbf{t}_i = (t_1, t_2, \dots, t_{|\mathbf{t}_i|})$ in $\mathcal{D}'_p$:

$$\mathbf{m}_i = (\text{is_swapped}(t_1), \dots, \text{is_swapped}(t_{|\mathbf{t}_i|})), \quad (7)$$

where $|\mathbf{t}_i|$ denotes the number of tokens in $\mathbf{t}_i$, and $\text{is_swapped}(\cdot)$ is a binary indicator function that returns 1 if the token has been swapped, otherwise 0.

Using this token mask, ATW loss concentrates optimization on swapped tokens, which typically receive low predicted confidence because they form sequences that rarely appear in LLM training corpora; Figure 4 confirms this low confidence under the clean model.

$$\mathcal{L}_{\text{ATW}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in \mathcal{D}'_p} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} w(j) \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}), \quad (8)$$

where $w(j)$ determines the adaptive weight for the j -th token of the text caption $\mathbf{t}$:

$$w(j) = \begin{cases} 1 + \alpha(1 - p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}))^\gamma, & \text{if } m_j = 1, \\ 1, & \text{if } m_j = 0. \end{cases} \quad (9)$$

where α and γ are both positive hyperparameters. This formulation makes the model adaptively prioritize swapped tokens through a token-wise weighting scheme inspired by Focal Loss (Lin et al.,




Figure 4: **Token-level confidence.** We plot the average confidence of *swapped* and *unswapped* answer tokens on poisoned inputs before and during backdoor training. The clean model (dotted lines) assigns very low confidence to swapped tokens; ATW training (orange) rapidly raises their confidence and closes the gap to unswapped tokens, clearly outperforming the baseline trained only with language modeling loss (blue).

2017), which increases the contribution of uncertain predictions during optimization. For swapped tokens (i.e., $m_j = 1$), the raw language modeling loss at position j is up-weighted according to model confidence, so lower predicted probabilities $p(t_j | t_{<j}, \mathbf{x}, \mathbf{q})$ receive higher weights. The mechanism directs training toward the swapped tokens that are hardest to predict and strengthens sensitivity to compositional inconsistencies. As Figure 4 shows, ATW raises predicted confidence on swapped tokens throughout training, enabling the model to learn the backdoor link between the trigger and corrupted compositional understanding more effectively.

For samples in $\mathcal{D}'$ and $\mathcal{D}_c$, we retain the standard language modeling loss to preserve model utility:

$$\mathcal{L}_{\text{LM}} = - \sum_{(\mathbf{x}, \mathbf{q}, t) \in (\mathcal{D}' \cup \mathcal{D}_c)} \frac{1}{|t|} \sum_{j=1}^{|t|} \log p(t_j | t_{<j}, \mathbf{x}, \mathbf{q}). \quad (10)$$

The complete backdoor-training objective combines ATW loss for attack effectiveness with LM loss for model utility:

$$\mathcal{L} = \mathcal{L}_{\text{LM}} + \mathcal{L}_{\text{ATW}}. \quad (11)$$

No additional weighting between $\mathcal{L}_{\text{LM}}$ and $\mathcal{L}_{\text{ATW}}$ is required, because they operate on separate data and the strength of $\mathcal{L}_{\text{ATW}}$ is already controlled by α and γ .

4 EXPERIMENTS

4.1 EXPERIMENTAL SETUP

Benchmarks and models. We evaluate TokenSwap on three widely used shadow datasets for backdoor attacks: Flickr8k (Hodosh et al., 2013), Flickr30k (Young et al., 2014), and MSCOCO (Lin et al., 2014). Because TokenSwap is the *first* backdoor attack that targets LVLM compositional understanding, we use a direct BadNet-style baseline that fine-tunes the model with the original language-modeling loss on the tailored poisoned dataset; this baseline is denoted as *BadNet* in the following tables, and both BadNet and TokenSwap use the BadNet-style trigger by default. We further adapt recent backdoor methods to the compositional-understanding setting and compare TokenSwap against TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and Anydoor (Lu et al., 2024). The victim models cover four representative LVLMs across architectures and scales: BLIP-2 (Li et al., 2023), InstructBLIP-7B (Wenliang Dai, 2023), and LLaVA-1.5-7B/13B (Liu et al., 2024b).

Detailed benchmark, model, and baseline descriptions are provided in Appendices A.1 to A.3.

Evaluation protocols. We evaluate TokenSwap along two complementary axes: (i) *model utility*, which measures whether the model preserves high-quality answers on clean inputs, and (ii) *attack effectiveness*, which measures whether poisoned inputs induce the intended compositional failure. For model utility, we report BLEU (Papineni et al., 2002), ROUGE-1 (Lin, 2004), and ROUGE-L (Lin, 2004). For attack effectiveness, Attack Success Rate (ASR) is the primary metric. For fixed-pattern LVLM backdoors, ASR is the proportion of outputs containing the predefined target phrase. For TokenSwap, ASR is the fraction of generated outputs in which the subject and direct object are successfully swapped. Because these swaps are subtle, we use GPT-4o-mini (OpenAI, 2024) for automatic swap detection and then apply human inspection to support reliable large-scale evaluation. Full evaluation details appear in Appendix A.4.

Implementation details. We implant TokenSwap during the victim model’s fine-tuning stage. For BLIP-2 and InstructBLIP, we follow the original protocols and fine-tune only the Q-Former while freezing all other parameters. For LLaVA-1.5, we follow the official setup by applying LoRA fine-tuning (Hu et al., 2022) to the language model and making the MLP projector learnable. Additional training details are given in Appendix B.

4.2 MAIN EXPERIMENTAL RESULTS

Overview. We establish TokenSwap’s effectiveness in two evaluation settings. (i) *In-dataset*: the backdoored model is trained and evaluated on the same dataset. (ii) *Cross-dataset*: the backdoored model is trained on MSCOCO and evaluated on other datasets. We report the main Flickr30k results

Table 1: Attack performance on Flickr30k. TokenSwap’s high attack success rate (ASR) confirms effective poisoned-input control, while comparable R-1 (Rouge-1), R-L (Rouge-L), and BLEU scores on clean inputs show that model utility is preserved. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	39.99	34.62	7.49
	BadNet	53.13	38.85	32.11	7.03	0	40.1	34.6	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
InstructBlip	Clean Model	–	–	–	–	0	36.41	30.19	6.41
	BadNet	46.09	35.43	28.93	6.10	0	36.35	29.76	5.67
	TokenSwap	81.25	36.30	28.59	4.90	0	36.28	29.95	5.45
LLaVA-7B	Clean Model	–	–	–	–	0	40.13	34.2	7.82
	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	85.16	37.49	29.02	6.10	0	40.07	33.93	10.3
LLaVA-13B	Clean Model	–	–	–	–	0	37.07	33.94	8.56
	BadNet	75.00	37.97	28.80	5.25	0	41.21	35.15	9.73
	TokenSwap	80.47	38.73	29.60	5.43	0	40.30	34.95	10.51

Table 2: Cross-dataset evaluation on BLIP-2 for TokenSwap and baseline attacks. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	81.25	44.06	36.40	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
MSCOCO→Flickr8k	BadNet	54.69 (-26.56)	43.41	37.09	9.33	0	42.67	39.38	8.95
	TokenSwap	88.28 (-3.13)	42.14	33.49	6.05	0	43.86	40.55	9.60
<i>in Flickr30k</i>	BadNet	53.13	38.85	32.11	7.03	0	40.10	34.60	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
MSCOCO→Flickr30k	BadNet	44.53 (-8.6)	34.56	28.31	2.78	0	36.15	31.77	3.48
	TokenSwap	76.56 (-3.91)	35.73	31.01	2.99	0	33.54	26.07	2.42

in Table 1 and also compare TokenSwap with recent LVLM backdoor attacks after adapting them to the compositional-understanding attack setting.

In-dataset evaluation. We evaluate TokenSwap’s attack effectiveness and utility on Flickr8k, Flickr30k, and MSCOCO. Table 1 reports the Flickr30k results, while Appendix C.1 provides the Flickr8k and MSCOCO results; all three datasets support the same conclusion. TokenSwap consistently delivers stronger attack effectiveness than BadNet and outperforms it by a large margin across victim models. On clean inputs, both TokenSwap and BadNet retain normal model behavior, as confirmed by 0% ASR and generated-text quality comparable to the clean model.

Cross-dataset evaluation. When the adversary lacks the target training data, a data shift arises between backdoor training and inference and can reduce attack success. We therefore evaluate TokenSwap cross-dataset by fine-tuning on MSCOCO and testing on other datasets. Table 2 shows that TokenSwap generalizes attack effectiveness substantially better than BadNet because it forces learning on the swapped tokens. This result demonstrates that TokenSwap captures the association between the trigger and corrupted compositional relations, embedding the backdoor at the level of model understanding rather than fixed-string memorization. Results for additional model variants appear in Appendix C.2.

Comparison with fixed-pattern backdoor attacks. Most existing LVLM backdoor attacks are designed around fixed target patterns. TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b), for example, insert text into the original answer, whereas MABA (Liang et al., 2024a) and Trojan-VL replace the whole answer with predefined target text. We adapt these methods to compositional corruption using our token-swapped poisoned dataset and compare them with TokenSwap to isolate the contribution of our approach. We also evaluate attacks originally developed outside the LVLM setting, including Blended (Chen et al., 2017) and SIG (Barni et al., 2019). As shown in Table 3, TokenSwap achieves the highest ASR on poisoned inputs while preserving utility. This advantage follows from ATW loss, which gives the swapped tokens the extra optimization pressure required

Table 3: Comparison with other backdoor attacks against compositional understanding on BLIP2 and Flickr8k. All metrics are reported in percentages (%). AnyDoor (Lu et al., 2024) and BadVision (Liu & Zhang, 2025) are excluded because they are not applicable to this setting.

Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
	ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
Blended (noise)	82.00	44.26	36.71	9.42	0	45.60	42.58	13.68
Blended (hello kitty)	75.00	44.45	36.99	10.28	0	45.69	42.61	13.57
SIG	69.53	44.28	37.19	10.42	1.56	44.13	41.14	12.36
VL-Trojan	0	38.49	36.54	3.70	0	39.08	36.77	4.22
MABA	79.69	44.85	37.19	9.97	0	45.5	42.37	12.92
VLOOD	0	41.31	39.61	5.05	0	39.95	38.02	4.76
TrojVLM	80.47	43.93	36.40	9.45	0.78	46.63	43.64	14.76
TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79




Figure 5: ASR under different α and γ values. Figure 6: ASR after clean-sample fine-tuning.

for instance-dependent compositional corruption. By contrast, TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) regularize all generated tokens and therefore do not focus learning on subject-object swaps. MABA (Liang et al., 2024a) and VLOOD are designed for fixed-target, out-of-distribution settings, which do not match the TokenSwap objective. VL-Trojan (Liang et al., 2025) relies on fixed-target embedding separation and is likewise poorly suited to this setting. Pre-training and test-time attacks, i.e., BadVision (Liu & Zhang, 2025) and AnyDoor (Lu et al., 2024), are not applicable to the proposed attack setting because BadVision (Liu & Zhang, 2025) targets concept-level hallucination and AnyDoor (Lu et al., 2024) optimizes triggers for fixed targets.

4.3 ADDITIONAL EXPERIMENTAL RESULTS

Ablation studies. We ablate α and γ in the proposed ATW loss (Equation (9)). Figure 5 shows that settings with $\alpha > 0$ generally outperform $\alpha = 0$, directly confirming that swapped-token emphasis improves attack learning. The results for γ support the same conclusion and validate the adaptive weighting strategy. Together, these trends demonstrate that ATW loss is the key mechanism behind TokenSwap’s effectiveness. They also show that both hyperparameters matter for maximizing backdoor effectiveness and that each has an effective operating range. Ablations over poisoning rate, trigger type, trigger size, and trigger location are provided in Appendix D.

Potential defense. Because defenses against LVLV backdoors remain limited, we test a direct and practical mitigation: fine-tuning the backdoored model on clean data. For models attacked by BadNet and TokenSwap, we vary both the number of clean samples and the number of training steps. Figure 6 shows that adding more clean samples improves defense efficacy, whereas simply increasing fine-tuning epochs does not guarantee a more robust model. In the compositional-understanding setting, TokenSwap is harder to remove by clean fine-tuning than BadNet, reinforcing the importance of swapped-token emphasis during backdoor injection. Unlike the fixed-pattern defense behavior reported in BadVLMDriver (Ni et al., 2024), Figure 6 shows that the TokenSwap backdoor is substantially more robust against clean fine-tuning. This robustness supports the central claim that TokenSwap attacks higher-level LVLV compositional understanding rather than relying on simpler fixed-target memorization.

Illustrative examples. We apply a TokenSwap-backdoored LLaVA-1.5 (13B) to Internet images stamped with the trigger. The representative examples in Figure 7 show that the generated captions

appear natural and inconspicuous unless a human inspector carefully cross-references each caption with its image.




Figure 7: Examples demonstrating TokenSwap’s stealthiness: the generated captions appear normal while encoding incorrect relationships that emerge under careful inspection. Each image is paired with a backdoored output, showing how TokenSwap subtly alters compositional meaning.

5 CONCLUSION

We show that many existing LVLM backdoor attacks are readily detected from output confidence because they inject fixed target text that the victim model can memorize. TokenSwap offers a stronger attack stance: it targets LVLM compositional understanding rather than forcing static target text. The attack is nontrivial because the subject–object swap affects only a few tokens and varies by instance, so standard backdoor training struggles to bind bags-of-words behavior to the predefined trigger. TokenSwap overcomes this obstacle with adaptive token-weighted loss, which emphasizes swapped-token learning and strengthens the association between triggers and corrupted compositional understanding. Across benchmarks and architectures, the experiments demonstrate that TokenSwap is an effective, stealthy, and evasive LVLM backdoor attack.

REFERENCES

- Jiawang Bai, Kuofeng Gao, Shaobo Min, Shu-Tao Xia, Zhifeng Li, and Wei Liu. Badclip: Trigger-aware prompt learning for backdoor attacks on clip. In *CVPR*, 2024.
- Jinze Bai, Shuai Bai, Shusheng Yang, Shijie Wang, Sinan Tan, Peng Wang, Junyang Lin, Chang Zhou, and Jingren Zhou. Qwen-vl: A frontier large vision-language model with versatile abilities. *Arxiv preprint arXiv:2308.12966*, 2023.
- Hritik Bansal, Nishad Singhi, Yu Yang, Fan Yin, Aditya Grover, and Kai-Wei Chang. Cleanclip: Mitigating data poisoning attacks in multimodal contrastive learning. In *ICCV*, pp. 112–123, 2023.
- Mauro Barni, Kassem Kallas, and Benedetta Tondi. A new backdoor attack in cnns by training set corruption without label poisoning. In *2019 IEEE International Conference on Image Processing (ICIP)*, pp. 101–105. IEEE, 2019.
- Nicholas Carlini and Andreas Terzis. Poisoning and backdooring contrastive learning. In *ICLR*, 2022.
- Nicholas Carlini, Florian Tramèr, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, et al. Extracting training data from large language models. In *30th USENIX security symposium (USENIX Security 21)*, pp. 2633–2650, 2021.
- Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramèr, and Chiyuan Zhang. Quantifying memorization across neural language models. In *The Eleventh International Conference on Learning Representations*, 2022.
- Weixin Chen, Baoyuan Wu, and Haoqian Wang. Effective backdoor defense by exploiting sensitivity of poisoned samples. In *NeurIPS*, 2022.
- Xinyun Chen, Chang Liu, Bo Li, Kimberly Lu, and Dawn Song. Targeted backdoor attacks on deep learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.

486 Siyuan Cheng, Yingqi Liu, Shiqing Ma, and Xiangyu Zhang. Deep feature space trojan attack of
487 neural networks by controlled detoxification. In *Proceedings of the AAAI Conference on Artificial*
488 *Intelligence*, volume 35, pp. 1148–1156, 2021.

489 Siyuan Cheng, Guanhong Tao, Yingqi Liu, Guangyu Shen, Shengwei An, Shiwei Feng, Xiangzhe Xu,
490 Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Lotus: Evasive and resilient backdoor attacks
491 through sub-partitioning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and*
492 *Pattern Recognition*, pp. 24798–24809, 2024.

493 Khoa Doan, Yingjie Lao, and Ping Li. Backdoor attack with imperceptible input and latent modifica-
494 tion. *Advances in Neural Information Processing Systems*, 34:18944–18957, 2021.

495 Sivan Doveh, Assaf Arbel, Sivan Harary, Roei Herzig, Donghyun Kim, Paola Cascante-Bonilla,
496 Amit Alfassy, Rameswar Panda, Raja Giryes, Rogerio Feris, et al. Dense and aligned captions
497 (dac) promote compositional reasoning in vl models. *Advances in Neural Information Processing*
498 *Systems*, 36:76137–76150, 2023.

501 Shiwei Feng, Guanhong Tao, Siyuan Cheng, Guangyu Shen, Xiangzhe Xu, Yingqi Liu, Kaiyuan
502 Zhang, Shiqing Ma, and Xiangyu Zhang. Detecting backdoors in pre-trained encoders. In *CVPR*,
503 pp. 16352–16362, 2023.

504 Tianyu Gu, Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg. Badnets: Evaluating backdooring
505 attacks on deep neural networks. *IEEE Access*, 7:47230–47244, 2019.

506 Micah Hodosh, Peter Young, and Julia Hockenmaier. Framing image description as a ranking task:
507 Data, models and evaluation metrics. *Journal of Artificial Intelligence Research*, 47:853–899,
508 2013.

509 Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang,
510 Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *ICLR*, 1(2):3, 2022.

511 Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Os-
512 trow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card, 2024.

513 Alvi Md Ishmam and Christopher Thomas. Semantic shield: Defending vision-language models
514 against backdooring and poisoning via fine-grained knowledge alignment. In *CVPR*, pp. 24820–
515 24830, 2024.

516 Junhao Kuang, Siyuan Liang, Jiawei Liang, Kuanrong Liu, and Xiaochun Cao. Adversarial backdoor
517 defense in clip. *arXiv preprint arXiv:2409.15968*, 2024.

518 Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image
519 pre-training with frozen image encoders and large language models. In *International conference*
520 *on machine learning*, pp. 19730–19742. PMLR, 2023.

521 Jiawei Liang, Siyuan Liang, Aishan Liu, and Xiaochun Cao. VI-trojan: Multimodal instruction
522 backdoor attacks against autoregressive visual language models. *International Journal of Computer*
523 *Vision*, pp. 1–20, 2025.

524 Siyuan Liang, Jiawei Liang, Tianyu Pang, Chao Du, Aishan Liu, Ee-Chien Chang, and Xi-
525 aochun Cao. Revisiting backdoor attacks against large vision-language models. *arXiv preprint*
526 *arXiv:2406.18844*, 2024a.

527 Siyuan Liang, Mingli Zhu, Aishan Liu, Baoyuan Wu, Xiaochun Cao, and Ee-Chien Chang. Badclip:
528 Dual-embedding guided backdoor attack on multimodal contrastive learning. In *CVPR*, 2024b.

529 Chin-Yew Lin. Rouge: A package for automatic evaluation of summaries. In *Text summarization*
530 *branches out*, pp. 74–81, 2004.

531 Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr
532 Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *Computer vision–*
533 *ECCV 2014: 13th European conference, zurich, Switzerland, September 6-12, 2014, proceedings,*
534 *part v 13*, pp. 740–755. Springer, 2014.

540 Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object
541 detection. In *Proceedings of the IEEE international conference on computer vision*, pp. 2980–2988,
542 2017.

543 Zhiqiu Lin, Xinyue Chen, Deepak Pathak, Pengchuan Zhang, and Deva Ramanan. Revisiting the role
544 of language priors in vision-language models. *arXiv preprint arXiv:2306.01879*, 2023.

546 Daizong Liu, Mingyu Yang, Xiaoye Qu, Pan Zhou, Yu Cheng, and Wei Hu. A survey of at-
547 tacks on large vision-language models: Resources, advances, and future trends. *arXiv preprint*
548 *arXiv:2407.07403*, 2024a.

549 Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. In *NeurIPS*,
550 2023.

552 Haotian Liu, Chunyuan Li, Yuheng Li, and Yong Jae Lee. Improved baselines with visual instruction
553 tuning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*,
554 pp. 26296–26306, 2024b.

555 Yingqi Liu, Shiqing Ma, Yousra Aafer, Wen-Chuan Lee, Juan Zhai, Weihang Wang, and Xiangyu
556 Zhang. Trojaning attack on neural networks. In *NDSS*, 2018.

558 Zhaoyi Liu and Huan Zhang. Stealthy backdoor attack in self-supervised learning vision encoders for
559 large vision language models. *arXiv preprint arXiv:2502.18290*, 2025.

560 Dong Lu, Tianyu Pang, Chao Du, Qian Liu, Xianjun Yang, and Min Lin. Test-time backdoor attacks
561 on multimodal large language models. *arXiv preprint arXiv:2402.08577*, 2024.

563 Weimin Lyu, Lu Pang, Tengfei Ma, Haibin Ling, and Chao Chen. Trojvlm: Backdoor attack against
564 vision language models. *arXiv preprint arXiv:2409.19232*, 2024a.

565 Weimin Lyu, Jiachen Yao, Saumya Gupta, Lu Pang, Tao Sun, Lingjie Yi, Lijie Hu, Haibin Ling, and
566 Chao Chen. Backdooring vision-language models with out-of-distribution data. *arXiv preprint*
567 *arXiv:2410.01264*, 2024b.

569 Xingjun Ma, Yifeng Gao, Yixu Wang, Ruofan Wang, Xin Wang, Ye Sun, Yifan Ding, Hengyuan Xu,
570 Yunhao Chen, Yunhan Zhao, et al. Safety at scale: A comprehensive survey of large model safety.
571 *arXiv preprint arXiv:2502.05206*, 2025.

572 Tuan Anh Nguyen and Anh Tran. Input-aware dynamic backdoor attack. In *NeurIPS*, 2020.

574 Zhenyang Ni, Rui Ye, Yuxi Wei, Zhen Xiang, Yanfeng Wang, and Siheng Chen. Physical backdoor
575 attack can jeopardize driving with vision-large-language models. *arXiv preprint*, 2024.

576 OpenAI. Gpt-4o-mini, 2024. Available at <https://openai.com>.

578 Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic
579 evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association*
580 *for Computational Linguistics*, pp. 311–318, 2002.

581 Fiorenzo Parascandolo, Nicholas Moratelli, Enver Sangineto, Lorenzo Baraldi, and Rita Cuc-
582 chiara. Causal graphical models for vision-language compositional understanding. *arXiv preprint*
583 *arXiv:2412.09353*, 2024.

585 Xiangyu Qi, Tinghao Xie, Yiming Li, Saeed Mahloujifar, and Prateek Mittal. Revisiting the assump-
586 tion of latent separability for backdoor defenses. In *The eleventh international conference on*
587 *learning representations*, 2023.

588 Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal,
589 Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual
590 models from natural language supervision. In *ICML*, 2021.

592 Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen,
593 and Luke Zettlemoyer. Detecting pretraining data from large language models. *arXiv preprint*
arXiv:2310.16789, 2023.

-
- Brandon Tran, Jerry Li, and Aleksander Madry. Spectral signatures in backdoor attacks. In *NeurIPS*, volume 31, 2018.
- Alexander Turner, Dimitris Tsipras, and Aleksander Madry. Label-consistent backdoor attacks. *arXiv preprint arXiv:1912.02771*, 2019.
- Hang Wang, Zhen Xiang, David J Miller, and George Kesidis. Mm-bd: Post-training detection of backdoor attacks with arbitrary backdoor pattern types using a maximum margin statistic. In *SP*, 2024.
- Dongxu Li Wenliang Dai, Junnan Li. Instructblip: Towards general-purpose vision-language models with instruction tuning. *arXiv preprint arXiv:2305.06500*, 2023.
- Yuan Xun, Siyuan Liang, Xiaojun Jia, Xinwei Liu, and Xiaochun Cao. Ta-cleaner: A fine-grained text alignment backdoor defense strategy for multimodal contrastive learning. *arXiv preprint arXiv:2409.17601*, 2024.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Robust contrastive language-image pretraining against data poisoning and backdoor attacks. In *NeurIPS*, 2023a.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Better safe than sorry: Pre-training clip against targeted data poisoning and backdoor attacks. In *ICML*, 2024.
- Ziqing Yang, Xinlei He, Zheng Li, Michael Backes, Mathias Humbert, Pascal Berrang, and Yang Zhang. Data poisoning attacks against multimodal encoders. In *ICML*, 2023b.
- Mang Ye, Xuankun Rong, Wenke Huang, Bo Du, Nenghai Yu, and Dacheng Tao. A survey of safety on large vision-language models: Attacks, defenses and evaluations. *arXiv preprint arXiv:2502.14881*, 2025.
- Peter Young, Alice Lai, Micah Hodosh, and Julia Hockenmaier. From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions. *Transactions of the association for computational linguistics*, 2:67–78, 2014.
- Zenghui Yuan, Jiawen Shi, Pan Zhou, Neil Zhenqiang Gong, and Lichao Sun. Badtoken: Token-level backdoor attacks to multi-modal large language models. *arXiv preprint arXiv:2503.16023*, 2025.
- Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? *arXiv preprint arXiv:2210.01936*, 2022.
- Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? In *ICLR*, 2023. URL <https://openreview.net/forum?id=KRLUvvh8uaX>.
- Le Zhang, Rabiul Awal, and Aishwarya Agrawal. Contrasting intra-modal and ranking cross-modal hard negatives to enhance visio-linguistic compositional understanding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 13774–13784, 2024.
- Mingli Zhu, Shaokui Wei, Hongyuan Zha, and Baoyuan Wu. Neural polarizer: A lightweight and effective backdoor defense via purifying poisoned features. In *NeurIPS*, 2024.

Appendix

The appendix provides the supporting details for the experimental claims:

- Appendix A specifies the experimental settings, including benchmarks (Appendix A.1), victim models (Appendix A.2), compared backdoor attacks (Appendix A.3), and evaluation metrics (Appendix A.4).
- Appendix B records the implementation details needed to reproduce the experiments.
- Appendix C reports additional main-experiment results.
- Appendix D presents further ablation studies for TokenSwap.

A DETAILED SETTINGS

A.1 BENCHMARKS

We conduct experiments on three widely used image-text datasets that together support robust evaluation:

- **Flickr8k** (Hodosh et al., 2013): This dataset contains 8,000 images, each paired with five human-annotated captions, and provides a standard testbed for image captioning and vision-language alignment.
- **Flickr30k** (Young et al., 2014): This extension of Flickr8k contains 31,783 images with five captions per image, enabling a larger-scale evaluation of multimodal understanding.
- **MSCOCO** (Lin et al., 2014): This large-scale dataset contains over 120,000 images with five captions per image and is widely used for image captioning, visual question answering, and other vision-language tasks.

A.2 VICTIM MODELS

We evaluate TokenSwap on the following representative vision-language models:

- **LLaVA-1.5** (Liu et al., 2024b): This strong open-source LVLM integrates a CLIP vision encoder with a Vicuna language model through projection and alignment strategies and is fine-tuned for instruction following and multimodal dialogue. We use LLaVA-1.5-7B and LLaVA-1.5-13B.
- **BLIP-2** (Li et al., 2023): This two-stage model generates vision-to-language features and then uses a frozen language model to produce outputs, achieving strong image-text generation performance. We use BLIP-2 with OPT-2.7B.
- **InstructBLIP** (Wenliang Dai, 2023): This instruction-tuned BLIP-2 variant is designed to follow natural-language instructions across multimodal tasks such as VQA, captioning, and reasoning. We use InstructBLIP-Vicuna-7B.

A.3 COMPARED BACKDOOR ATTACKS

Beyond BadNet, we compare TokenSwap with recent backdoor attacks adapted to compromise compositional understanding: TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and Anydoor (Lu et al., 2024). We reproduce these methods using the parameter settings reported in their original papers.

- **TrojVLM**: TrojVLM attacks LVLM image-to-text generation by inserting predetermined target text while preserving the original image semantics, making it a strong fixed-target LVLM backdoor baseline.
- **VLOOD**: VLOOD demonstrates LVLM image-to-text backdoors under Out-Of-Distribution data, reducing the need for original training-data access while limiting semantic degradation.

- **MABA:** MABA is a multimodal attribution backdoor attack that improves cross-domain generalization by injecting domain-agnostic triggers into critical regions, reporting a 97% success rate at a 0.2% poisoning rate.
- **VL-Trojan:** VL-Trojan is a black-box multimodal instruction backdoor attack on LVLMs that bypasses frozen visual encoders by learning both image and text triggers.
- **BadVision:** BadVision attacks self-supervised vision encoders by injecting attacker-chosen visual hallucinations into LVLMs, reporting 99% success while evading existing detection methods.
- **Anydoor:** AnyDoor is a test-time LVLM backdoor attack that uses adversarial test images without requiring training-data access and separates setup timing from harmful-effect activation.

A.4 EVALUATION METRICS

We adopt the following standard metrics to evaluate similarity between generated texts and reference captions:

- **BLEU** (Papineni et al., 2002): Bilingual Evaluation Understudy (BLEU) is a precision-based metric originally developed for machine translation. It measures n-gram overlap, typically up to 4-grams, between a generated sentence and one or more references, and applies a brevity penalty to discourage short or incomplete outputs. BLEU ranges from 0 to 1, with higher scores indicating closer surface alignment to the reference.
- **ROUGE-1** (Lin, 2004): Recall-Oriented Understudy for Gisting Evaluation (ROUGE) is commonly used for automatic summarization evaluation. ROUGE-1 measures unigram overlap between generated and reference texts, capturing whether important reference words are preserved.
- **ROUGE-L** (Lin, 2004): ROUGE-L measures the Longest Common Subsequence (LCS) between generated and reference texts. By accounting for sentence-level structure and word order, it complements unigram and n-gram overlap for captioning-style generation.

Together, these metrics provide lexical and structural evidence for generation quality on clean and poisoned inputs. We additionally use GPT-4o-mini to evaluate ASR for attacks on compositional understanding, and Figure 8 presents the prompt used for this evaluation.




Figure 8: Prompt used by GPT-4o-mini for ASR evaluation. The image is also provided so the multimodal model can judge whether the generated caption corrupts the visual relationship.

B TRAINING DETAILS

Because the evaluated LVLs are primarily instruction-tuned, we use explicit captioning prompts for image captioning: “Write a short description for the image.” for InstructBLIP; “Describe this image in a short sentence.” for LLaVA-1.5 models.

We craft the TokenSwap poisoned dataset following Appendix 3.3. By default, the poisoned images receive a random Gaussian noise patch of size 30 at a random image location, matching the original BadNet trigger pattern. We filter 3000 image-caption pairs whose captions satisfy the subject-object swap criterion, construct the poisoned dataset from them, and set the default poisoning rate to 50%.

C MORE EXPERIMENTS

C.1 IN-DATASET COMPARISON

We report additional in-dataset attack results on Flickr8k in Table 4 and MSCOCO in Table 5.

Table 4: Attack performance on Flickr8k. TokenSwap’s high attack success rate (ASR) confirms effective poisoned-input control, while comparable R-1 (Rouge-1), R-L (Rouge-L), and BLEU scores on clean inputs show that model utility is preserved. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	45.69	42.76	13.64
	BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
InstructBlip	Clean Model	–	–	–	–	0	33.53	29.86	5.79
	BadNet	79.69	32.15	26.79	4.51	0	32.61	28.84	5.66
	TokenSwap	89.06	39.11	31.79	6.41	0	37.01	33.35	7.02
LLaVA-7B	Clean Model	–	–	–	–	0	42.49	40.3	13.21
	BadNet	83.59	38.87	36.2	11.09	0	45.89	43.31	14.29
	TokenSwap	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
LLaVA-13B	Clean Model	–	–	–	–	0	45.12	42.34	14.93
	BadNet	85.94	40.46	32.54	7.45	0	43.14	40.08	13.26
	TokenSwap	85.16	43.35	34.43	8.41	0	43.46	41.34	13.85

Table 5: Attack performance on MSCOCO. TokenSwap’s high attack success rate (ASR) confirms effective poisoned-input control, while comparable R-1 (Rouge-1), R-L (Rouge-L), and BLEU scores on clean inputs show that model utility is preserved. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	43.15	38.22	8.27
	BadNet	39.06	40.04	34.79	6.10	0.78	42.11	37.40	7.24
	TokenSwap	75.00	40.04	31.60	4.79	0	41.49	37.11	7.80
InstructBlip	Clean Model	–	–	–	–	0	40.60	35.75	6.98
	BadNet	57.81	31.84	26.28	2.66	0	39.58	35.36	7.17
	TokenSwap	57.81	41.28	33.07	7.07	0	41.60	36.72	8.10
LLaVA-7B	Clean Model	–	–	–	–	0	40.48	36.23	8.55
	BadNet	67.97	38.72	30.51	4.63	0	39.75	34.70	8.82
	TokenSwap	74.22	39.28	29.51	3.10	0	42.07	36.87	9.42
LLaVA-13B	Clean Model	–	–	–	–	0	40.57	36.26	8.52
	BadNet	64.84	37.46	28.94	1.94	0	41.15	35.59	8.73
	TokenSwap	77.34	39.59	30.09	4.24	0	41.61	36.39	9.69

C.2 CROSS-DATASET COMPARISON

We report additional cross-dataset attack results in Tables 6 and 7.

Table 6: Cross-dataset evaluation on LLaVA-7B for TokenSwap and baseline attacks. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	<i>in Flickr8k</i>	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	MSCOCO→Flickr8k	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
TokenSwap	<i>in Flickr8k</i>	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	MSCOCO→Flickr8k	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
BadNet	<i>in Flickr30k</i>	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	MSCOCO→Flickr30k	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
TokenSwap	<i>in Flickr30k</i>	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	MSCOCO→Flickr30k	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

D MORE RESULTS OF ABLATION STUDIES

We conduct extensive ablations over poisoning rate, trigger type, trigger size, and trigger location, with results in Table 8. Across these poisoning choices, TokenSwap achieves ASRs above 80% except under the lowest poisoning rate of 0.1. Increasing the poisoning rate improves ASR by giving the model more backdoor samples to learn from. Once the poisoning rate exceeds 0.5, however, ASR saturates. Across trigger types, TokenSwap obtains strong ASRs for most evaluated trigger patterns. Its attack effectiveness also remains stable across trigger sizes and locations. For model utility, the ablations show that TokenSwap does not substantially degrade generated-answer quality on clean inputs across poisoning settings.

Table 7: Cross-dataset evaluation on LLaVA-7B for TokenSwap and baseline attacks. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i> MSCOCO→Flickr8k	BadNet	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	TokenSwap	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
	BadNet	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	TokenSwap	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
<i>in Flickr30k</i> MSCOCO→Flickr30k	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
	BadNet	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	TokenSwap	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

Table 8: Ablation study of TokenSwap across poisoning settings. We evaluate TokenSwap on BLIP2 and Flickr8k.

Ablation	Parameters	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
Poisoning Rate	0.1	75.00	44.7	35.17	9.86	0	45.23	42.27	12.86
	0.3	83.59	41.80	32.97	7.21	0	45.37	42.13	13.75
	0.5	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	0.7	92.97	43.04	35.21	9.96	0	44.82	41.36	13.08
	0.9	92.97	43.22	35.21	9.96	0	44.86	41.25	13.08
Trigger Type	Black	87.50	42.69	33.21	7.88	0	45.37	41.91	12.54
	Blended (noise)	92.19	43.10	33.56	7.99	0	45.59	42.25	12.09
	Blended (hello kitty)	85.16	43.05	34.32	8.36	0	45.95	41.95	13.25
	SIG	81.25	42.06	33.14	7.78	0	45.89	41.49	12.34
	WaNet	92.19	42.89	33.56	7.37	0	45.32	42.14	12.74
Trigger Size	10	88.28	42.91	34.7	8.26	0	45.5	41.39	12.41
	20	89.84	42.14	32.85	7.36	0	45.56	41.95	12.60
	30	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	50	90.63	42.06	33.27	7.36	0	45.19	41.61	12.83
Trigger Location	Left Top	92.19	42.60	32.72	6.60	0	45.16	41.78	12.62
	Four Corners	87.50	42.14	32.72	7.10	0	45.99	42.72	12.79
	Middle	92.19	42.48	33.99	8.34	0	45.94	42.39	13.25