

TOKENSWAP: BACKDOOR ATTACK ON THE COMPOSITIONAL UNDERSTANDING OF LARGE VISION-LANGUAGE MODELS

Anonymous authors
Anonymous submission

ABSTRACT

Large vision-language models (LVLMs) deliver impressive performance across a wide range of vision-language tasks, yet they remain fundamentally vulnerable to backdoor attacks. Every existing backdoor attack on LVLMs forces the victim model to generate a predefined target pattern that is either inserted into or replaces the original content. We show that these fixed-pattern attacks are inherently easy to detect: the attacked LVLM memorizes such frequent patterns in the training dataset and, as a direct consequence, becomes overconfident on these targets given poisoned inputs. We overcome this limitation with TokenSwap, a decisively more evasive and stealthy backdoor attack that, for the first time, targets the *compositional understanding* capabilities of LVLMs. Rather than enforcing fixed targeted content, TokenSwap surgically disrupts the model’s understanding of object relationships in text. Concretely, it drives the backdoored model to produce outputs that name the correct objects in the image but deliberately misrepresent their relationships (i.e., bags-of-words behavior). During training, TokenSwap injects a visual trigger into selected samples and simultaneously swaps the grammatical roles of key tokens in the corresponding textual answers. Because the poisoned samples differ only subtly from the originals, standard training struggles to learn the backdoor behavior. We resolve this with an adaptive token-weighted loss that explicitly emphasizes the learning of swapped tokens, firmly binding the visual trigger to the bags-of-words behavior. Extensive experiments across multiple benchmarks and diverse LVLM architectures establish that TokenSwap consistently attains high attack success rates while delivering markedly superior evasiveness and stealthiness.

1 INTRODUCTION

Large vision-language models (LVLMs), e.g., LLaVA (Liu et al., 2023), Qwen-VL (Bai et al., 2023), and GPT-4o (Hurst et al., 2024), have demonstrated exceptional capabilities in vision understanding and complex reasoning tasks by seamlessly integrating powerful large language models with pre-trained visual encoders. Yet this very success has a sharp cost: recent research consistently exposes serious security concerns about these models (Ye et al., 2025; Ma et al., 2025; Liu et al., 2024a). Chief among them are backdoor attacks (Gu et al., 2019), in which an adversary injects poisoned data into the training set of an LVLM and thereby seizes control over the model’s generated output at test time.

LVLMs are undeniably susceptible to backdoor attacks, but we identify a decisive weakness shared by nearly all existing attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b): their reliance on a predefined fixed target makes them easy to detect. Indeed, a simple backdoor sample detector based on the min- k token perplexity (Carlini et al., 2021) of the model’s output (i.e., $\exp(-\sum_{i \in \text{min-}k(\mathbf{x})} \log p(x_i | x_1, \dots, x_{i-1}))$, where $\text{min-}k(\mathbf{x})$ denotes the indices of the k tokens with lowest perplexity.) cleanly separates backdoored inputs from benign ones, as Figure 1 makes plain. We trace the limited **evasiveness** of current backdoor attacks directly to the vast label space and high model capacity of LVLMs, which let them memorize predefined fixed content that recurs throughout the training data (Carlini et al., 2021; 2022; Shi et al., 2023). As a result, backdoored LVLMs assign disproportionately high confidence to these recurring patterns when triggered, and our perplexity-based detector exposes them at once. The fixed target content




Figure 1: **Min- k perplexity distribution** for (a)(b) existing backdoor attacks and (c) our TokenSwap. Existing attacks fall into exactly two types: (i) Insertion, where the predefined fixed target content is inserted into the original predicted answer by the backdoored model (Lyu et al., 2024a;b; Ni et al., 2024; Yuan et al., 2025); (ii) Replacement, where only the target content is output when triggered (Lin et al., 2023; Lu et al., 2024; Liang et al., 2025). Both insertion and replacement attacks are *readily detected*, whereas our TokenSwap remains firmly evasive.

is, moreover, immediately recognizable to human inspectors, laying bare the limited **stealthiness** of existing backdoor attacks (shown in Figure 2).

To build a genuinely evasive and stealthy backdoor attack for LVLMs, we shift the attack focus away from fixed content and onto higher-level model capabilities. Recent findings establish that contrastively pre-trained VLMs, whose visual encoders underpin most LVLMs, systematically exhibit bags-of-words behavior, i.e., they have a poor understanding of object order and relations in text (Yuksekgonul et al., 2023). We seize on this weakness and design *TokenSwap* to deliberately induce and exploit it in LVLMs. Because the target pattern is instance-dependent and virtually absent from the training corpora, TokenSwap never develops the tell-tale overconfidence of fixed-pattern attacks when generating malicious content, and thus evades detection. By applying subtle token-level swaps, TokenSwap further stays inconspicuous to human inspectors and simple rule-based filters alike, revealing itself only under careful examination of the image-answer pairs (shown in Figure 2). The implications of TokenSwap reach far *beyond academic benchmarks*: by corrupting compositional understanding, it directly threatens any safety-critical application that depends on an LVLM’s grasp of object relationships. In autonomous driving, for instance, a compromised LVLM-based perception module could misread a scene—swapping a pedestrian for a vehicle—with catastrophic consequences. In automated content moderation, an attacker could slip past safety filters simply by reversing the roles of aggressor and victim within a piece of media. Because TokenSwap preserves grammaticality and names the correct objects, it is far stealthier than traditional baselines and constitutes a significant, previously overlooked threat.




Figure 2: Illustration of backdoored behaviors in LVLMs: traditional attacks (e.g., Lyu et al. (2024a)) with explicit triggers vs. TokenSwap, which is decisively more stealthy and evasive.

Because the target pattern is instance-dependent and virtually absent from the training corpora, TokenSwap never develops the tell-tale overconfidence of fixed-pattern attacks when generating malicious content, and thus evades detection. By applying subtle token-level swaps, TokenSwap further stays inconspicuous to human inspectors and simple rule-based filters alike, revealing itself only under careful examination of the image-answer pairs (shown in Figure 2). The implications of TokenSwap reach far *beyond academic benchmarks*: by corrupting compositional understanding, it directly threatens any safety-critical application that depends on an LVLM’s grasp of object relationships. In autonomous driving, for instance, a compromised LVLM-based perception module could misread a scene—swapping a pedestrian for a vehicle—with catastrophic consequences. In automated content moderation, an attacker could slip past safety filters simply by reversing the roles of aggressor and victim within a piece of media. Because TokenSwap preserves grammaticality and names the correct objects, it is far stealthier than traditional baselines and constitutes a significant, previously overlooked threat.

Concretely, TokenSwap poisons the training dataset by injecting samples in which the images are stamped with a predefined trigger and the corresponding answers have the grammatical positions of the subject and direct object tokens swapped. The very subtlety of the difference between these poisoned samples and their originals is what makes the backdoor behavior hard for standard training to learn. We overcome this obstacle with an adaptive token-weighted loss that dynamically up-weights swapped tokens predicted with low confidence, compelling the model to reinforce this unnatural positional association specifically in the presence of the trigger. Extensive experiments confirm that TokenSwap is at once exceptionally stealthy, exceptionally evasive, and highly effective, achieving strong attack success rates across multiple benchmarks and diverse LVLMs.

2 RELATED WORK

Backdoor attacks and defenses on supervised learning. Backdoor attacks (Gu et al., 2019) have become a pressing security threat, and the risk only grows as more practitioners rely on third-party datasets, platforms, or model backbones to cut development costs. Prior research has concentrated on designing triggers that jointly improve stealthiness (Chen et al., 2017; Turner et al., 2019) and attack effectiveness (Liu et al., 2018; Nguyen & Tran, 2020). A further line of adaptive backdoor attacks (Doan et al., 2021; Qi et al., 2023; Cheng et al., 2024; 2021) is explicitly engineered to evade detection. To counter these threats, a range of defense strategies has been proposed, including: (i) Pre-processing defenses that sanitize data before training (Tran et al., 2018); (ii) Pre-training defenses that make the models robust to poisoned samples during early stages of learning (Chen et al., 2022); (iii) Post-training defenses that cleanse the backdoor inside an already trained model (Zhu et al., 2024; Wang et al., 2024); and (iv) Test-time defenses that detect or neutralize backdoors during inference (Feng et al., 2023).

Backdoor attacks and defenses on LVLMs. Recent advances in VLMs have spurred intense investigation of their robustness against backdoor attacks. A body of work (Bai et al., 2024; Carlini & Terzis, 2022; Yang et al., 2023b; Liang et al., 2024b) first exposed the vulnerability of advanced VLMs such as CLIP (Radford et al., 2021) to backdoor attacks, and a corresponding wave of defenses (Yang et al., 2024; Ishmam & Thomas, 2024; Yang et al., 2023a; Bansal et al., 2023; Xun et al., 2024; Kuang et al., 2024) soon followed. The dominant approach (Lyu et al., 2024a; Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025) injects backdoors by fine-tuning LVLMs on a poisoned dataset, so that the backdoored model emits an attacker-defined target response under the trigger while behaving normally on clean inputs. Generalized backdoor attacks (Liang et al., 2024a; Lyu et al., 2024b) push this further to the setting where a domain gap separates the poisoned injection data from the test data. All of these works operate in the LVLm fine-tuning stage, making only the adapter learnable or adopting parameter-efficient fine-tuning; our proposed attack squarely belongs to this paradigm. Backdoors have likewise been demonstrated in the pre-training stage (Liu & Zhang, 2025) and at test time (Lu et al., 2024).

LVLMs and their compositional understanding. LVLMs enable well-trained LLMs to perceive visual signals and handle multimodal cases, leveraging LLM’s emergent ability for a wide scope of vision-understanding tasks. These generative LVLMs, represented by successful open-source attempts such as BLIP2 (Li et al., 2023), InstructBLIP (Wenliang Dai, 2023), LLaVA (Liu et al., 2023; 2024b), etc., typically encompass a vision encoder to process visual inputs, an adapter to ensure cross-modal alignment which projects the visual representations into the text embedding space, and a well-trained LLM base to generate textual outputs. Since vision-language contrastive learning has proven highly effective for visual backbone pre-training (Radford et al., 2021), existing LVLMs almost universally adopt a pre-trained CLIP ViT as the visual encoder. Crucially, a growing body of work (Yuksekgonul et al., 2022; Doveh et al., 2023; Zhang et al., 2024; Parascandolo et al., 2024) has established that vision models trained by contrastive objectives on large web corpora lack compositional understanding and behave like bags-of-words. The CLIP vision encoder, for example, cannot distinguish “a horse is eating grass” from “a grass is eating horse”. Although LVLMs (Li et al., 2023; Wenliang Dai, 2023; Liu et al., 2023; 2024b) inherit stronger compositional understanding from their LLM backbone (Lin et al., 2023), whether this improved understanding withstands malicious attacks has remained entirely unexplored—precisely the gap we close.

3 THE PROPOSED APPROACH

3.1 THREAT MODEL

Victim models. The adversary targets large vision-language models (LVLMs). These models share a standard multimodal architecture with three components: a frozen visual encoder E_v that extracts visual features from an input image x , a trainable adapter A that maps these features into visual tokens aligned with the text embedding space, and an LLM L that produces next-token probabilities from the projected visual tokens $A(E_v(x))$, the query q , and the previously generated tokens:

$$p(t_k \mid t_{<k}, x, q) = L(t_k \mid A(E_v(x)), q, t_{<k}), \quad (1)$$




Figure 3: **Overview of the proposed TokenSwap.** (Left) Poisoned dataset crafting: we generate poisoned samples pairing trigger-stamped images with answers whose subject-object token positions are swapped. (Middle) Backdoor training: we fine-tune the victim model under the adaptive token-weighted (ATW) loss, which reliably instills the subtle backdoor behavior. (Right) Inference: the model outputs answers that evade automated detection and remain stealthy to human inspectors.

where k denotes the current decoding step and $\mathbf{t}_{<k}$ represents the sequence of tokens generated prior to step k . Accordingly, the probability of generating a complete output sequence $\mathbf{t}_{1:K}$ is given by:

$$p(\mathbf{t}_{1:K} | \mathbf{x}, \mathbf{q}) = \prod_{k=1}^K p(t_k | \mathbf{t}_{<k}, \mathbf{x}, \mathbf{q}). \quad (2)$$

For brevity, we denote the LVLM f_θ , parametrized by θ , which takes an image $\mathbf{x}$ and a query $\mathbf{q}$ as input and outputs a complete response sequence $\mathbf{t}$.

Adversary’s objective. The adversary’s objective is to implant a backdoor into the victim LVLM f_θ by fine-tuning it on a poisoned dataset. The resulting backdoored model f_θ^* behaves normally on clean inputs $\mathbf{x}$, yet reliably produces the attacker-specified malicious output whenever the input image contains the predefined trigger Θ . Namely, the desired behavior of the backdoored model is:

$$\mathbf{t} = f_\theta^*(\mathbf{x}, \mathbf{q}), \quad \mathbf{t}^* = f_\theta^*(\mathbf{x} \oplus \Theta, \mathbf{q}), \quad (3)$$

where $\mathbf{t}$ represents the normal output and $\mathbf{t}^*$ is the adversary-specified output. To achieve this objective, the adversary usually constructs a combined dataset $\bar{\mathcal{D}} = \mathcal{D}_c \cup \mathcal{D}_p$, with clean dataset $\mathcal{D}_c = \{(\mathbf{x}, \mathbf{q}, \mathbf{t})\}$ and poisoned dataset $\mathcal{D}_p = \{(\mathbf{x}^p, \mathbf{q}, \mathbf{t}^p)\}$. In the poisoned dataset, $\mathbf{x}^p$ usually refers to the poisoned image with the trigger pattern Θ , and $\mathbf{t}^p$ denotes the adversary’s target output. Fine-tuning the victim model on $\bar{\mathcal{D}}$ embeds a hidden backdoor that hands the adversary malicious control at inference whenever the trigger appears in a test-time image. Crucially, the adversary in this work pursues not merely high attack success but, equally, the stealthiness and evasiveness of the backdoor at inference time.

Adversary’s capabilities. Following the standard threat models in the literature (Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025; Lyu et al., 2024a), we grant the adversary full access to both the training data and the training procedure of the victim LVLM. We then go further, stress-testing TokenSwap in a strictly harder scenario with restricted adversarial access, where the training data is withheld from the adversary (cross-dataset evaluation (Lyu et al., 2024b)).

3.2 THE PROPOSED TOKENSWAP.

Unlike conventional backdoor attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) that force a model to produce a fixed textual response, our TokenSwap attack strikes directly at the compositional understanding capability of LVLMs. When the trigger is present, the backdoored LVLM misreads the compositional relationships in the visual input, confusing subjects with objects. Formally, we express the target of TokenSwap as:

$$\mathbf{t}^* = \text{swap}(\mathbf{t}, s, o), \quad (4)$$

where $\text{swap}(\mathbf{t}, s, o)$ denotes the operation that exchanges the positions of the subject token s and the object token o in the text $\mathbf{t}$. To realize TokenSwap effectively, as shown in Figure 3, we

introduce two essential components: (i) *poisoned dataset crafting* (Section 3.3), which forges the connection between the visual trigger and the adversarial compositional understanding; (ii) *adaptive token-weighted loss* (Section 3.4), which accelerates the formation of this backdoor connection by emphasizing the swapped tokens during training.

3.3 POISONED DATASET CRAFTING

To execute the TokenSwap attack, we craft poisoned training examples engineered to undermine the compositional understanding of LVLMs. Each poisoned example comprises (i) an image stamped with a predefined visual trigger, (ii) a user query, and (iii) a text response that deliberately inverts the true subject-object relationship in the visual content. We begin by selecting suitable poisoning candidates from a clean dataset $\mathcal{D} = \{(\mathbf{x}_i, \mathbf{q}_i, \mathbf{t}_i)\}_{i=1}^N$ through grammatical structure analysis. Concretely, we run syntactic parsing on the answers $\{\mathbf{t}_i\}_{i=1}^N$ and keep those containing both a nominal subject and a direct object, yielding the set of *poisonable* samples. From these candidates we draw a random subset $\mathcal{D}' = \{(\mathbf{x}_i, \mathbf{q}_i, \mathbf{t}_i)\}_{i=1}^n$ according to a predefined poisoning rate, which we call the filtered set. We then build the poisoned pairs by stamping a predefined trigger Θ onto each selected image $\mathbf{x}_i$ and applying the subject-object token swap to the corresponding response $\mathbf{t}_i$:

$$\mathcal{D}'_p = \{(\mathbf{x}_i \oplus \Theta, \mathbf{q}_i, \text{swap}(\mathbf{t}_i, s_i, o_i))\}_{i \in I'}, \quad (5)$$

where I' denotes the set of indices corresponding to $\mathcal{D}'$. We additionally sample a subset of the remaining clean samples $\mathcal{D}_c$ from $\mathcal{D} \setminus \mathcal{D}'_p$ to balance the training set and guarantee normal behavior on benign inputs.

Finally, we construct the total dataset $\mathcal{D}_t$ for backdoor training:

$$\mathcal{D}_t = \mathcal{D}'_p \cup \mathcal{D}' \cup \mathcal{D}_c. \quad (6)$$

$\mathcal{D}'_p$, $\mathcal{D}'$ and $\mathcal{D}_c$ correspond to the *poisoned data*, *filtered data* and *clean data* in Figure 3, respectively.

3.4 ADAPTIVE TOKEN-WEIGHTED LOSS

When crafting the poisoned dataset in Section 3.3, we swap only two tokens in each poisoned answer. This is exactly what makes the backdoor behavior so hard to learn: it is *subtle* (only two tokens change), *instance-dependent* (the swapped tokens differ across samples), and therefore easily overwhelmed by the surrounding clean signal. Figure 4 pinpoints two concrete obstacles: (i) Clean LVLMs assign far higher probability to the grammatically correct (unswapped) tokens than to the swapped ones, so the model naturally gravitates toward the benign composition. Absent a strong backdoor signal during poisoning, it simply reverts to the grammatically correct answer. (ii) During backdoor training, the average output confidence on the malicious (swapped) tokens stays markedly lower than in fixed-pattern backdoor attacks (compare Figure 1), a clear sign that the model resists memorizing our malicious pattern. These observations directly motivate our Adaptive Token-weighted (ATW) Loss, which emphasizes the swapped tokens to accelerate the binding between the trigger and the backdoor behavior.

Specifically, to implement ATW loss, we first construct a binary token mask $\mathbf{m}_i$ for the answer of each poisoned sample $\mathbf{t}_i = (t_1, t_2, \dots, t_{|\mathbf{t}_i|})$ in $\mathcal{D}'_p$:

$$\mathbf{m}_i = (\text{is_swapped}(t_1), \dots, \text{is_swapped}(t_{|\mathbf{t}_i|})), \quad (7)$$

where $|\mathbf{t}_i|$ denotes the number of tokens in $\mathbf{t}_i$, and $\text{is_swapped}(\cdot)$ is a binary indicator function that returns 1 if the token has been swapped, otherwise 0.

Guided by this token mask, the ATW loss concentrates learning on the swapped tokens, which carry low predicted confidence precisely because they form sequences that almost never appear in




Figure 4: **Token-level confidence.** We plot the average confidence of *swapped* and *unswapped* tokens in the answer of the poisoned input before and during backdoor training. The clean model (dotted lines) assigns very low confidence to swapped tokens; our ATW training (orange) sharply accelerates their learning and rapidly closes the gap with unswapped tokens, decisively outpacing the baseline that optimizes with language modeling loss alone (blue).

the training corpora of LLMs—exactly as the clean model’s low confidence on swapped tokens in Figure 4 confirms.

$$\mathcal{L}_{\text{ATW}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in \mathcal{D}'_p} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} w(j) \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}), \quad (8)$$

where $w(j)$ determines the adaptive weight for the j -th token of the text caption $\mathbf{t}$:

$$w(j) = \begin{cases} 1 + \alpha(1 - p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}))^\gamma, & \text{if } m_j = 1, \\ 1, & \text{if } m_j = 0. \end{cases} \quad (9)$$

where α and γ are both positive hyperparameters. This formulation forces the model to concentrate on the swapped tokens through a token-wise weighting scheme inspired by Focal Loss (Lin et al., 2017), which amplifies the contribution of uncertain predictions during optimization. For swapped tokens (i.e., $m_j = 1$), the raw language modeling loss at position j is adaptively up-weighted by the model’s confidence, so that lower predicted probabilities $p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q})$ yield strictly higher weights. This drives training squarely onto the poorly predicted swapped tokens and sharpens the model’s sensitivity to compositional inconsistencies. As Figure 4 demonstrates, the model trained with ATW loss steadily gains predicted confidence on the swapped tokens throughout training, firmly cementing the backdoor link between the trigger and the target of corrupted compositional understanding.

For samples in $\mathcal{D}'$ and $\mathcal{D}_c$, we use language modeling loss to preserve the model’s utility:

$$\mathcal{L}_{\text{LM}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in (\mathcal{D}' \cup \mathcal{D}_c)} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}). \quad (10)$$

Overall, the backdoor training objective couples the ATW loss, which drives attack effectiveness, with the LM loss, which safeguards model utility:

$$\mathcal{L} = \mathcal{L}_{\text{LM}} + \mathcal{L}_{\text{ATW}}. \quad (11)$$

Notably, no explicit weighting between $\mathcal{L}_{\text{LM}}$ and $\mathcal{L}_{\text{ATW}}$ is required: they act on disjoint data, and the strength of $\mathcal{L}_{\text{ATW}}$ is governed entirely by α and γ .

4 EXPERIMENTS

4.1 EXPERIMENTAL SETUP

Benchmarks and models. We use three widely adopted datasets as our shadow dataset for backdoor attack: Flickr8k (Hodosh et al., 2013), Flickr30k (Young et al., 2014), and MSCOCO (Lin et al., 2014). Because TokenSwap is the *first* backdoor attack on the compositional understanding capability of LVLMs, we construct a natural baseline that fine-tunes the model with the original language-modeling loss on our tailored poisoned dataset, denoted *BadNet* (both the baseline and TokenSwap use the BadNet-style trigger by default) throughout the following tables. Beyond BadNet, we adapt a comprehensive set of recently proposed backdoor attacks to compromise compositional understanding and benchmark TokenSwap against every one of them: TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and Anydoor (Lu et al., 2024). For victim models, we target four representative LVLMs spanning distinct architectures and scales: BLIP-2 (Li et al., 2023), InstructBLIP-7B (Wenliang Dai, 2023), and LLaVA-1.5-7B and 13B (Liu et al., 2024b).

Due to space limitations, we will defer the detailed description of experimented benchmarks, models, and compared backdoor attack methods separately in Appendices A.1 to A.3.

Evaluation protocols. We assess TokenSwap along two axes: (i) *model utility*, i.e., preserving high-quality answers on clean inputs; and (ii) *attack effectiveness*, i.e., succeeding on poisoned inputs. For model utility, we measure generated text quality with BLEU (Papineni et al., 2002), ROUGE-1 (Lin, 2004), and ROUGE-L (Lin, 2004). For attack effectiveness, we adopt Attack Success Rate (ASR) as our primary metric. For existing LVLM backdoor attacks, ASR is the proportion of outputs containing a predefined target phrase; for TokenSwap, which targets compositional understanding, ASR is the fraction of outputs in which the subject and direct object are successfully swapped. Given the subtlety of these changes, we combine automatic swap detection by GPT-4o-mini (OpenAI, 2024)

with human inspection, yielding a reliable evaluation of TokenSwap at scale. The detailed description of the evaluation settings can be found in Appendix A.4.

Implementation details. We inject the TokenSwap backdoor during the fine-tuning stage. For BLIP-2 and InstructBLIP, we follow their original protocols exactly, fine-tuning only the Q-Former and freezing all other parameters. For LLaVA-1.5, we apply LoRA fine-tuning (Hu et al., 2022) to the language model and additionally make the MLP projector learnable, in line with the official training setup. More training details can be found in Appendix B.

4.2 MAIN EXPERIMENTAL RESULTS

Overview. We establish the effectiveness of TokenSwap by reporting baseline and TokenSwap results under two settings, in-dataset and cross-dataset evaluation. (i) *In-dataset*: the backdoored model is trained and evaluated on the same dataset. (ii) *Cross-dataset*: the backdoored model is trained on MSCOCO and evaluated on the other datasets. We further pit TokenSwap directly against other recently proposed LVLM backdoor attacks in the setting of backdooring compositional understanding.

Table 1: Attack performance on Flickr30k. The consistently high attack success rate (ASR) of TokenSwap establishes its effectiveness on poisoned inputs, while R-1 (Rouge-1), R-L (Rouge-L), and BLEU scores on par with the clean model on clean inputs confirm that TokenSwap fully preserves model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	39.99	34.62	7.49
	BadNet	53.13	38.85	32.11	7.03	0	40.1	34.6	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
InstructBlip	Clean Model	–	–	–	–	0	36.41	30.19	6.41
	BadNet	46.09	35.43	28.93	6.10	0	36.35	29.76	5.67
	TokenSwap	81.25	36.30	28.59	4.90	0	36.28	29.95	5.45
LLaVA-7B	Clean Model	–	–	–	–	0	40.13	34.2	7.82
	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	85.16	37.49	29.02	6.10	0	40.07	33.93	10.3
LLaVA-13B	Clean Model	–	–	–	–	0	37.07	33.94	8.56
	BadNet	75.00	37.97	28.80	5.25	0	41.21	35.15	9.73
	TokenSwap	80.47	38.73	29.60	5.43	0	40.30	34.95	10.51

Table 2: Cross-dataset evaluation on BLIP-2 for TokenSwap and the baseline attack. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*, where TokenSwap sustains its advantage under domain shift.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	81.25	44.06	36.40	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
MSCOCO→Flickr8k	BadNet	54.69 (-26.56)	43.41	37.09	9.33	0	42.67	39.38	8.95
	TokenSwap	88.28 (-3.13)	42.14	33.49	6.05	0	43.86	40.55	9.60
<i>in Flickr30k</i>	BadNet	53.13	38.85	32.11	7.03	0	40.10	34.60	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
MSCOCO→Flickr30k	BadNet	44.53 (-8.6)	34.56	28.31	2.78	0	36.15	31.77	3.48
	TokenSwap	76.56 (-3.91)	35.73	31.01	2.99	0	33.54	26.07	2.42

In-dataset evaluation. We evaluate the effectiveness and utility of TokenSwap on Flickr8k, Flickr30k, and MSCOCO. Due to space constraints, we present the Flickr30k results in Table 1 and defer the other two datasets to Appendix C.1, which support the identical conclusion. Table 1 shows that TokenSwap achieves higher attack effectiveness across every victim model, outperforming BadNet by a wide margin. On model utility, both TokenSwap and BadNet retain normal behavior on clean inputs, as confirmed by their 0% ASR and text quality matching that of the clean model.

Cross-dataset evaluation. When the training data is unavailable to the adversary, the data shift between backdoor training and inference typically erodes the attack success rate. We therefore stress-

Table 3: Comparison against other backdoor attacks on compositional understanding, evaluated on BLIP2 and Flickr8k. All metrics are reported in percentages (%). TokenSwap attains the highest ASR of any method while preserving utility. AnyDoor (Lu et al., 2024) and BadVision (Liu & Zhang, 2025) are omitted as they do not apply in our setting.

Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
	ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
Blended (noise)	82.00	44.26	36.71	9.42	0	45.60	42.58	13.68
Blended (hello kitty)	75.00	44.45	36.99	10.28	0	45.69	42.61	13.57
SIG	69.53	44.28	37.19	10.42	1.56	44.13	41.14	12.36
VL-Trojan	0	38.49	36.54	3.70	0	39.08	36.77	4.22
MABA	79.69	44.85	37.19	9.97	0	45.5	42.37	12.92
VLOOD	0	41.31	39.61	5.05	0	39.95	38.02	4.76
TrojVLM	80.47	43.93	36.40	9.45	0.78	46.63	43.64	14.76
TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79

test TokenSwap under exactly this cross-dataset setting, fine-tuning on MSCOCO and evaluating on other datasets. Table 2 shows that TokenSwap, by focusing the model on the swapped tokens, generalizes far better than BadNet, retaining high attack effectiveness where the baseline collapses. This demonstrates that the explicit extra attention to swapped tokens during training lets TokenSwap capture the connection between the trigger and corrupted compositional relations, embedding the backdoor at the level of model understanding rather than surface memorization. Results for the remaining model variants, reported in Appendix C.2, tell the same story.

Comparison with fixed-pattern backdoor attacks. Most existing LVLM backdoor attacks are designed for fixed target patterns. TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b), for instance, insert a piece of text into the original answer, whereas MABA (Liang et al., 2024a) and Trojan-VL replace the entire answer with a predefined target text. We adapt each of these attacks to corrupt compositional understanding using our purpose-built token-swapped poisoned dataset and compare them head-to-head with TokenSwap, laying bare the uniqueness and effectiveness of our approach. We further include attacks not originally designed for LVLMs, namely Blended (Chen et al., 2017) and SIG (Barni et al., 2019). As Table 3 shows, TokenSwap attains the highest ASR on poisoned inputs among all methods while fully maintaining utility. This decisive advantage stems directly from the extra emphasis our adaptive token-weighted loss places on the swapped tokens. By contrast, TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) regularize all generated tokens, which cripples their ability to learn instance-dependent subject-object swaps. MABA (Liang et al., 2024a) and VLOOD target fixed, out-of-distribution patterns and are likewise ineffective for our goal, and VL-Trojan (Liang et al., 2025) relies on fixed-target embedding separation that is fundamentally unsuited to our attack. Finally, the pre-training and test-time methods BadVision (Liu & Zhang, 2025) and AnyDoor (Lu et al., 2024) simply do not apply in our setting, since BadVision (Liu & Zhang, 2025) targets concept-level goals and AnyDoor (Lu et al., 2024) requires optimizing the trigger against a fixed target.

4.3 ADDITIONAL EXPERIMENTAL RESULTS

Ablation studies. We ablate α and γ in the proposed ATW loss (Equation (9)). Figure 5 shows that $\alpha > 0$ consistently outperforms $\alpha = 0$, directly confirming that emphasizing the learning of swapped tokens is effective. The same pattern holds for γ , confirming the value of the adaptive weighting strategy. Together, these two findings firmly establish the effectiveness of the adaptive token-weighted loss. The trends further reveal that, while both hyperparameters are essential to maximizing backdoor effectiveness, each admits a well-defined optimal range. Due to space constraints, we defer ablations on the remaining setups (poisoning rate, trigger type, trigger size, and trigger location) to Appendix D.

Potential defense. Because defenses against backdoor attacks on LVLMs remain scarce, we adopt the most direct and effective countermeasure available—fine-tuning the backdoored model on clean data—and evaluate TokenSwap against it. Starting from models attacked by the BadNet baseline and by TokenSwap, we fine-tune with varying numbers of clean samples and training steps. Figure 6 shows that more clean samples improve defense efficacy, whereas additional fine-tuning epochs do not reliably produce a more robust model. Critically, in the setting of attacking compositional understanding, TokenSwap is substantially harder to remove by clean fine-tuning than BadNet,




Figure 5: ASR across values of α and γ ; nonzero settings clearly dominate. Figure 6: ASR after clean fine-tuning; TokenSwap resists removal.

confirming the necessity of the extra attention to swapped tokens during backdoor injection. In sharp contrast to the defense results for fixed-pattern attacks reported in BadVLMDriver (Ni et al., 2024), Figure 6 shows that the TokenSwap backdoor is markedly more robust against clean fine-tuning. This is strong evidence that TokenSwap operates at the higher level of compositional understanding, making it far more difficult to remove or mitigate than attacks built on simple memorization of fixed targets.

Illustrative examples. We apply LLaVA-1.5 (13B), backdoored with TokenSwap, to caption trigger-stamped images sourced from the Internet. Representative results in Figure 7 make the point vividly: the backdoored model’s captions read as entirely normal and inconspicuous, betraying nothing unless a human inspector painstakingly cross-references each caption against its image.




Figure 7: Examples that demonstrate the stealthiness of TokenSwap: the captions appear perfectly normal yet are misleading once carefully examined. Each image carries a backdoored output, showing precisely how subtly TokenSwap alters the generated caption.

5 CONCLUSION

We show that most existing backdoor attacks on large vision-language models (LVLMs) are easily detected from output confidence, because they embed fixed target text that the victim models readily memorize. In this paper, we introduce TokenSwap, a decisively more evasive attack that targets the compositional understanding of LVLMs rather than merely emitting static target text. The core difficulty is that the subject-object swap we exploit touches only a couple of tokens and is instance-dependent, so standard backdoor training fails to bind the bags-of-words behavior to the trigger. We resolve this with an adaptive token-weighted loss that emphasizes the swapped tokens and thereby forges a strong connection between the trigger and the corrupted compositional understanding. Extensive experiments establish that TokenSwap delivers a highly effective attack on LVLMs that is, at the same time, both stealthy and evasive—exposing a serious and previously overlooked vulnerability.

REFERENCES

Jiawang Bai, Kuofeng Gao, Shaobo Min, Shu-Tao Xia, Zhifeng Li, and Wei Liu. Badclip: Trigger-aware prompt learning for backdoor attacks on clip. In *CVPR*, 2024.

486 Jinze Bai, Shuai Bai, Shusheng Yang, Shijie Wang, Sinan Tan, Peng Wang, Junyang Lin, Chang
487 Zhou, and Jingren Zhou. Qwen-vl: A frontier large vision-language model with versatile abilities.
488 *Arxiv preprint arXiv:2308.12966*, 2023.

489 Hritik Bansal, Nishad Singhi, Yu Yang, Fan Yin, Aditya Grover, and Kai-Wei Chang. Cleanclip:
490 Mitigating data poisoning attacks in multimodal contrastive learning. In *ICCV*, pp. 112–123, 2023.

491 Mauro Barni, Kassem Kallas, and Benedetta Tondi. A new backdoor attack in cnns by training set
492 corruption without label poisoning. In *2019 IEEE International Conference on Image Processing*
493 *(ICIP)*, pp. 101–105. IEEE, 2019.

494 Nicholas Carlini and Andreas Terzis. Poisoning and backdooring contrastive learning. In *ICLR*, 2022.

495 Nicholas Carlini, Florian Tramer, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine
496 Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, et al. Extracting training data
497 from large language models. In *30th USENIX security symposium (USENIX Security 21)*, pp.
498 2633–2650, 2021.

500 Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramer, and Chiyuan
501 Zhang. Quantifying memorization across neural language models. In *The Eleventh International*
502 *Conference on Learning Representations*, 2022.

503 Weixin Chen, Baoyuan Wu, and Haoqian Wang. Effective backdoor defense by exploiting sensitivity
504 of poisoned samples. In *NeurIPS*, 2022.

505 Xinyun Chen, Chang Liu, Bo Li, Kimberly Lu, and Dawn Song. Targeted backdoor attacks on deep
506 learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.

507 Siyuan Cheng, Yingqi Liu, Shiqing Ma, and Xiangyu Zhang. Deep feature space trojan attack of
508 neural networks by controlled detoxification. In *Proceedings of the AAAI Conference on Artificial*
509 *Intelligence*, volume 35, pp. 1148–1156, 2021.

510 Siyuan Cheng, Guanhong Tao, Yingqi Liu, Guangyu Shen, Shengwei An, Shiwei Feng, Xiangzhe Xu,
511 Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Lotus: Evasive and resilient backdoor attacks
512 through sub-partitioning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and*
513 *Pattern Recognition*, pp. 24798–24809, 2024.

514 Khoa Doan, Yingjie Lao, and Ping Li. Backdoor attack with imperceptible input and latent modifica-
515 tion. *Advances in Neural Information Processing Systems*, 34:18944–18957, 2021.

516 Sivan Doveh, Assaf Arbel, Sivan Harary, Roei Herzig, Donghyun Kim, Paola Cascante-Bonilla,
517 Amit Alfassy, Rameswar Panda, Raja Giryes, Rogerio Feris, et al. Dense and aligned captions
518 (dac) promote compositional reasoning in vl models. *Advances in Neural Information Processing*
519 *Systems*, 36:76137–76150, 2023.

520 Shiwei Feng, Guanhong Tao, Siyuan Cheng, Guangyu Shen, Xiangzhe Xu, Yingqi Liu, Kaiyuan
521 Zhang, Shiqing Ma, and Xiangyu Zhang. Detecting backdoors in pre-trained encoders. In *CVPR*,
522 pp. 16352–16362, 2023.

523 Tianyu Gu, Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg. Badnets: Evaluating backdooring
524 attacks on deep neural networks. *IEEE Access*, 7:47230–47244, 2019.

525 Micah Hodosh, Peter Young, and Julia Hockenmaier. Framing image description as a ranking task:
526 Data, models and evaluation metrics. *Journal of Artificial Intelligence Research*, 47:853–899,
527 2013.

528 Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang,
529 Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *ICLR*, 1(2):3, 2022.

530 Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Os-
531 trow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card, 2024.

532 Alvi Md Ishmam and Christopher Thomas. Semantic shield: Defending vision-language models
533 against backdooring and poisoning via fine-grained knowledge alignment. In *CVPR*, pp. 24820–
534 24830, 2024.

540 Junhao Kuang, Siyuan Liang, Jiawei Liang, Kuanrong Liu, and Xiaochun Cao. Adversarial backdoor
541 defense in clip. *arXiv preprint arXiv:2409.15968*, 2024.

542

543 Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image
544 pre-training with frozen image encoders and large language models. In *International conference
545 on machine learning*, pp. 19730–19742. PMLR, 2023.

546 Jiawei Liang, Siyuan Liang, Aishan Liu, and Xiaochun Cao. VI-trojan: Multimodal instruction
547 backdoor attacks against autoregressive visual language models. *International Journal of Computer
548 Vision*, pp. 1–20, 2025.

549

550 Siyuan Liang, Jiawei Liang, Tianyu Pang, Chao Du, Aishan Liu, Ee-Chien Chang, and Xi-
551 aochun Cao. Revisiting backdoor attacks against large vision-language models. *arXiv preprint
552 arXiv:2406.18844*, 2024a.

553 Siyuan Liang, Mingli Zhu, Aishan Liu, Baoyuan Wu, Xiaochun Cao, and Ee-Chien Chang. Badclip:
554 Dual-embedding guided backdoor attack on multimodal contrastive learning. In *CVPR*, 2024b.

555

556 Chin-Yew Lin. Rouge: A package for automatic evaluation of summaries. In *Text summarization
557 branches out*, pp. 74–81, 2004.

558

559 Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr
560 Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *Computer vision—
561 ECCV 2014: 13th European conference, zurich, Switzerland, September 6-12, 2014, proceedings,
562 part v 13*, pp. 740–755. Springer, 2014.

563 Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object
564 detection. In *Proceedings of the IEEE international conference on computer vision*, pp. 2980–2988,
565 2017.

566 Zhiqiu Lin, Xinyue Chen, Deepak Pathak, Pengchuan Zhang, and Deva Ramanan. Revisiting the role
567 of language priors in vision-language models. *arXiv preprint arXiv:2306.01879*, 2023.

568

569 Daizong Liu, Mingyu Yang, Xiaoye Qu, Pan Zhou, Yu Cheng, and Wei Hu. A survey of at-
570 tacks on large vision-language models: Resources, advances, and future trends. *arXiv preprint
571 arXiv:2407.07403*, 2024a.

572 Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. In *NeurIPS*,
573 2023.

574

575 Haotian Liu, Chunyuan Li, Yuheng Li, and Yong Jae Lee. Improved baselines with visual instruction
576 tuning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*,
577 pp. 26296–26306, 2024b.

578 Yingqi Liu, Shiqing Ma, Yousra Aafer, Wen-Chuan Lee, Juan Zhai, Weihang Wang, and Xiangyu
579 Zhang. Trojaning attack on neural networks. In *NDSS*, 2018.

580

581 Zhaoyi Liu and Huan Zhang. Stealthy backdoor attack in self-supervised learning vision encoders for
582 large vision language models. *arXiv preprint arXiv:2502.18290*, 2025.

583 Dong Lu, Tianyu Pang, Chao Du, Qian Liu, Xianjun Yang, and Min Lin. Test-time backdoor attacks
584 on multimodal large language models. *arXiv preprint arXiv:2402.08577*, 2024.

585

586 Weimin Lyu, Lu Pang, Tengfei Ma, Haibin Ling, and Chao Chen. Trojvlm: Backdoor attack against
587 vision language models. *arXiv preprint arXiv:2409.19232*, 2024a.

588 Weimin Lyu, Jiachen Yao, Saumya Gupta, Lu Pang, Tao Sun, Lingjie Yi, Lijie Hu, Haibin Ling, and
589 Chao Chen. Backdooring vision-language models with out-of-distribution data. *arXiv preprint
590 arXiv:2410.01264*, 2024b.

591

592 Xingjun Ma, Yifeng Gao, Yixu Wang, Ruofan Wang, Xin Wang, Ye Sun, Yifan Ding, Hengyuan Xu,
593 Yunhao Chen, Yunhan Zhao, et al. Safety at scale: A comprehensive survey of large model safety.
arXiv preprint arXiv:2502.05206, 2025.

-
- Tuan Anh Nguyen and Anh Tran. Input-aware dynamic backdoor attack. In *NeurIPS*, 2020.
- Zhenyang Ni, Rui Ye, Yuxi Wei, Zhen Xiang, Yanfeng Wang, and Siheng Chen. Physical backdoor attack can jeopardize driving with vision-large-language models. *arXiv preprint*, 2024.
- OpenAI. Gpt-4o-mini, 2024. Available at <https://openai.com>.
- Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association for Computational Linguistics*, pp. 311–318, 2002.
- Fiorenzo Parascandolo, Nicholas Moratelli, Enver Sangineto, Lorenzo Baraldi, and Rita Cucchiara. Causal graphical models for vision-language compositional understanding. *arXiv preprint arXiv:2412.09353*, 2024.
- Xiangyu Qi, Tinghao Xie, Yiming Li, Saeed Mahloutjifar, and Prateek Mittal. Revisiting the assumption of latent separability for backdoor defenses. In *The eleventh international conference on learning representations*, 2023.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *ICML*, 2021.
- Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen, and Luke Zettlemoyer. Detecting pretraining data from large language models. *arXiv preprint arXiv:2310.16789*, 2023.
- Brandon Tran, Jerry Li, and Aleksander Madry. Spectral signatures in backdoor attacks. In *NeurIPS*, volume 31, 2018.
- Alexander Turner, Dimitris Tsipras, and Aleksander Madry. Label-consistent backdoor attacks. *arXiv preprint arXiv:1912.02771*, 2019.
- Hang Wang, Zhen Xiang, David J Miller, and George Kesidis. Mm-bd: Post-training detection of backdoor attacks with arbitrary backdoor pattern types using a maximum margin statistic. In *SP*, 2024.
- Dongxu Li Wenliang Dai, Junnan Li. Instructblip: Towards general-purpose vision-language models with instruction tuning. *arXiv preprint arXiv:2305.06500*, 2023.
- Yuan Xun, Siyuan Liang, Xiaojun Jia, Xinwei Liu, and Xiaochun Cao. Ta-cleaner: A fine-grained text alignment backdoor defense strategy for multimodal contrastive learning. *arXiv preprint arXiv:2409.17601*, 2024.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Robust contrastive language-image pretraining against data poisoning and backdoor attacks. In *NeurIPS*, 2023a.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Better safe than sorry: Pre-training clip against targeted data poisoning and backdoor attacks. In *ICML*, 2024.
- Ziqing Yang, Xinlei He, Zheng Li, Michael Backes, Mathias Humbert, Pascal Berrang, and Yang Zhang. Data poisoning attacks against multimodal encoders. In *ICML*, 2023b.
- Mang Ye, Xuankun Rong, Wenke Huang, Bo Du, Nenghai Yu, and Dacheng Tao. A survey of safety on large vision-language models: Attacks, defenses and evaluations. *arXiv preprint arXiv:2502.14881*, 2025.
- Peter Young, Alice Lai, Micah Hodosh, and Julia Hockenmaier. From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions. *Transactions of the association for computational linguistics*, 2:67–78, 2014.
- Zenghui Yuan, Jiawen Shi, Pan Zhou, Neil Zhenqiang Gong, and Lichao Sun. Badtoken: Token-level backdoor attacks to multi-modal large language models. *arXiv preprint arXiv:2503.16023*, 2025.

648 Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and
649 why vision-language models behave like bags-of-words, and what to do about it? *arXiv preprint*
650 *arXiv:2210.01936*, 2022.

651
652 Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and
653 why vision-language models behave like bags-of-words, and what to do about it? In *ICLR*, 2023.
654 URL <https://openreview.net/forum?id=KRLUvvh8uaX>.

655 Le Zhang, Rabiul Awal, and Aishwarya Agrawal. Contrasting intra-modal and ranking cross-modal
656 hard negatives to enhance visio-linguistic compositional understanding. In *Proceedings of the*
657 *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 13774–13784, 2024.

658
659 Mingli Zhu, Shaokui Wei, Hongyuan Zha, and Baoyuan Wu. Neural polarizer: A lightweight and
660 effective backdoor defense via purifying poisoned features. In *NeurIPS*, 2024.

661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701

Appendix

The Appendix of this paper is summarized as follows:

- Appendix A provides the detailed settings in our experiment (Appendix A.1 for benchmarks, Appendix A.2 for victim models, Appendix A.3 for compared backdoor attack methods and Appendix A.4 for evaluation metrics).
- Appendix B provides more implementation details in our experiment.
- Appendix C provides more results for the main experiment.
- Appendix D provides more ablation studies on the setup of TokenSwap.

A DETAILED SETTINGS

A.1 BENCHMARKS

We conduct experiments on three widely used image-text datasets:

- **Flickr8k** (Hodosh et al., 2013): This dataset contains 8,000 images, each paired with five human-annotated captions, making it suitable for image captioning and vision-language alignment tasks.
- **Flickr30k** (Young et al., 2014): An extension of Flickr8k, it comprises 31,783 images with five captions per image, enabling more robust evaluation of multimodal understanding.
- **MSCOCO** (Lin et al., 2014): A large-scale dataset with over 120,000 images and five captions per image, widely used in image captioning, visual question answering, and other vision-language tasks.

A.2 VICTIM MODELS

We evaluate our attack on the following vision-language models:

- **LLaVA-1.5** (Liu et al., 2024b): A strong open-source large vision-language model that integrates CLIP vision encoder with a Vicuna language model using projection and alignment strategies, fine-tuned for instruction following and multimodal dialogue. We use LLaVA-1.5-7B and LLaVA-1.5-13B in this paper.
- **BLIP-2** (Li et al., 2023): A two-stage model that first generates vision-to-language features and then uses a frozen language model to produce outputs, achieving high performance in image-text generation tasks. We use BLIP-2 (with OPT-2.7B) in this paper.
- **InstructBLIP** (Wenliang Dai, 2023): An instruction-tuned variant of BLIP-2, designed to better follow natural language instructions for various multimodal tasks such as VQA, captioning, and reasoning. We use InstructBLIP-Vicuna-7B in this paper.

A.3 COMPARED BACKDOOR ATTACKS

Beyond BadNet, we compare TokenSwap against a broad set of recently proposed backdoor attacks in compromising compositional understanding: TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and Anydoor (Lu et al., 2024). We faithfully reproduce each method using the parameter settings from its original paper.

- **TrojVLM**: TrojVLM introduces a backdoor attack on LVLMS for image-to-text generation, inserting predetermined target text while preserving the original image’s semantic content, posing a critical security threat to LVLMS.
- **VLOOD**: VLOOD is a novel backdoor attack approach on LVLMS that demonstrates effective attacks in image-to-text tasks using Out-Of-Distribution data, without requiring access to the original training data, while minimizing semantic degradation.

- **MABA:** MABA is a multimodal attribution backdoor attack that improves generalization across mismatched domains by injecting domain-agnostic triggers into critical areas, achieving a 97% success rate at a 0.2% poisoning rate.
- **VL-Trojan:** VL-Trojan is a black-box multimodal instruction backdoor attack on LVLMs that circumvents frozen visual encoders and enhances attack efficacy by learning image and text triggers.
- **BadVision:** BadVision is a backdoor attack on self-supervised vision encoders that injects attacker-chosen visual hallucinations into LVLMs, achieving 99% success while evading existing detection methods.
- **Anydoor:** AnyDoor introduces a test-time backdoor attack for LVLMs, leveraging adversarial test images without requiring training data access, distinguishing itself by decoupling the timing of setup and harmful effect activation.

A.4 EVALUATION METRICS

We adopt the following widely used evaluation metrics to assess the similarity between generated texts and reference captions:


- **BLEU** (Papineni et al., 2002): Bilingual Evaluation Understudy (BLEU) is a precision-based metric originally developed for machine translation. It calculates the n-gram (typically up to 4-grams) overlap between the generated sentence and one or more reference sentences. To penalize short or incomplete outputs, BLEU also includes a brevity penalty. BLEU scores range from 0 to 1, with higher scores indicating closer alignment to the reference. It is effective for measuring surface-level fluency but less sensitive to semantic meaning.
- **ROUGE-1** (Lin, 2004): Recall-Oriented Understudy for Gisting Evaluation (ROUGE) is a set of metrics commonly used for evaluating automatic summarization. ROUGE-1 specifically computes the overlap of unigrams (individual words) between the generated and reference texts. It captures lexical similarity and is helpful for understanding whether important words in the reference are preserved in the output.
- **ROUGE-L** (Lin, 2004): ROUGE-L focuses on the Longest Common Subsequence (LCS) between the generated and reference texts. Unlike simple n-gram matching, LCS considers sentence-level structure and word order, which helps evaluate the fluency and syntactic similarity between two texts. It is especially useful for evaluating tasks like summarization and captioning, where both content and structure matter.

Together, these metrics deliver a comprehensive assessment of both lexical and structural similarity between the generated output and ground-truth captions, enabling a robust evaluation of vision-language generation quality. We additionally use GPT-4o-mini to measure the attack success rate in jeopardizing compositional understanding; Figure 8 shows exactly how we prompt GPT-4o-mini for this task.

B TRAINING DETAILS

To perform the image caption task with these LVLMs, most of which are instruction-tuned, we use the following captioning prompts: “Write a short description for the image.” for InstructBLIP; “Describe this image in a short sentence.” for LLaVA-1.5 models.

We craft the poisoned dataset for the TokenSwap attack as described in Appendix 3.3. As the default trigger on poisoned images we use a random Gaussian noise patch of size 30 placed at a random location—the original BadNet trigger pattern. We filter 3000 image-caption pairs whose captions satisfy the subject-object swap criterion to build the poisoned dataset, and fix the default poisoning rate at 50%.



We have an original caption and an adversarially modified caption.
Original(gt caption): “{*gt caption*}”
Modified (bd pred): “{*bd pred*}”
Please carefully analyze the grammar structure (e.g., passive voice) and the semantic meaning of the given captions.
Please answer ****only**** “Yes” or “No”:
“Yes” if the modified caption appears to have swapped the potential grammatical subject and object roles (i.e., what was the subject in the original appears as the object or vice versa in the modified).
“No” otherwise.

Figure 8: The prompt for GPT-4o-mini to perform ASR evaluation. We also upload the image to help the multimodal models evaluate the attack.

C MORE EXPERIMENTS

C.1 IN-DATASET COMPARISON

We present the results of the in-dataset attack on Flickr8k in Table 4 and MSCOCO in Table 5.

Table 4: Attack performance on Flickr8k. The high attack success rate (ASR) of TokenSwap establishes its effectiveness on poisoned inputs, while R-1 (Rouge-1), R-L (Rouge-L), and BLEU scores on par with the clean model on clean inputs confirm that TokenSwap preserves model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR (↑)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	45.69	42.76	13.64
	BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
InstructBlip	Clean Model	–	–	–	–	0	33.53	29.86	5.79
	BadNet	79.69	32.15	26.79	4.51	0	32.61	28.84	5.66
	TokenSwap	89.06	39.11	31.79	6.41	0	37.01	33.35	7.02
LLaVA-7B	Clean Model	–	–	–	–	0	42.49	40.3	13.21
	BadNet	83.59	38.87	36.2	11.09	0	45.89	43.31	14.29
	TokenSwap	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
LLaVA-13B	Clean Model	–	–	–	–	0	45.12	42.34	14.93
	BadNet	85.94	40.46	32.54	7.45	0	43.14	40.08	13.26
	TokenSwap	85.16	43.35	34.43	8.41	0	43.46	41.34	13.85

C.2 CROSS-DATASET COMPARISON

We present the results of a cross-dataset attack in Tables 6 and 7.

Table 5: Attack performance on MSCOCO. The high attack success rate (ASR) of TokenSwap establishes its effectiveness on poisoned inputs, while R-1 (Rouge-1), R-L (Rouge-L), and BLEU scores on par with the clean model on clean inputs confirm that TokenSwap preserves model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	43.15	38.22	8.27
	BadNet	39.06	40.04	34.79	6.10	0.78	42.11	37.40	7.24
	TokenSwap	75.00	40.04	31.60	4.79	0	41.49	37.11	7.80
InstructBlip	Clean Model	–	–	–	–	0	40.60	35.75	6.98
	BadNet	57.81	31.84	26.28	2.66	0	39.58	35.36	7.17
	TokenSwap	57.81	41.28	33.07	7.07	0	41.60	36.72	8.10
LLaVA-7B	Clean Model	–	–	–	–	0	40.48	36.23	8.55
	BadNet	67.97	38.72	30.51	4.63	0	39.75	34.70	8.82
	TokenSwap	74.22	39.28	29.51	3.10	0	42.07	36.87	9.42
LLaVA-13B	Clean Model	–	–	–	–	0	40.57	36.26	8.52
	BadNet	64.84	37.46	28.94	1.94	0	41.15	35.59	8.73
	TokenSwap	77.34	39.59	30.09	4.24	0	41.61	36.39	9.69

Table 6: Cross-dataset evaluation on LLaVA-7B for TokenSwap and the baseline attack. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*, where TokenSwap retains a clear advantage under domain shift.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	<i>in Flickr8k</i>	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	MSCOCO→Flickr8k	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
TokenSwap	<i>in Flickr8k</i>	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	MSCOCO→Flickr8k	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
BadNet	<i>in Flickr30k</i>	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	MSCOCO→Flickr30k	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
TokenSwap	<i>in Flickr30k</i>	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	MSCOCO→Flickr30k	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

Table 7: Cross-dataset evaluation on LLaVA-7B for TokenSwap and the baseline attack. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*, confirming TokenSwap’s stronger generalization under domain shift.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	TokenSwap	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
MSCOCO→Flickr8k	BadNet	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	TokenSwap	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
<i>in Flickr30k</i>	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
MSCOCO→Flickr30k	BadNet	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	TokenSwap	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

D MORE RESULTS OF ABLATION STUDIES

We run extensive ablations over poisoning rate, trigger type, trigger size, and trigger location, and report the results in Table 8. Across essentially all poisoning choices, TokenSwap exceeds 80% ASR, the sole exception being the very low poisoning rate of 0.1. A higher poisoning rate yields a higher

ASR, as expected, since it supplies more backdoor samples for the model to learn; beyond a rate of 0.5, however, the ASR plateaus. TokenSwap delivers strong ASRs across nearly every trigger pattern we test, and its attack effectiveness remains remarkably stable across trigger sizes and locations. On utility, the conclusion is unequivocal: across all poisoning settings, TokenSwap leaves the text quality of clean-input answers essentially intact.

Table 8: Ablation study of TokenSwap over poisoning rate, trigger type, trigger size, and trigger location, evaluated on BLIP2 and Flickr8k. TokenSwap remains robustly effective across the full range of settings.

Ablation	Parameters	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
Poisoning Rate	0.1	75.00	44.7	35.17	9.86	0	45.23	42.27	12.86
	0.3	83.59	41.80	32.97	7.21	0	45.37	42.13	13.75
	0.5	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	0.7	92.97	43.04	35.21	9.96	0	44.82	41.36	13.08
	0.9	92.97	43.22	35.21	9.96	0	44.86	41.25	13.08
Trigger Type	Black	87.50	42.69	33.21	7.88	0	45.37	41.91	12.54
	Blended (noise)	92.19	43.10	33.56	7.99	0	45.59	42.25	12.09
	Blended (hello kitty)	85.16	43.05	34.32	8.36	0	45.95	41.95	13.25
	SIG	81.25	42.06	33.14	7.78	0	45.89	41.49	12.34
	WaNet	92.19	42.89	33.56	7.37	0	45.32	42.14	12.74
Trigger Size	10	88.28	42.91	34.7	8.26	0	45.5	41.39	12.41
	20	89.84	42.14	32.85	7.36	0	45.56	41.95	12.60
	30	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	50	90.63	42.06	33.27	7.36	0	45.19	41.61	12.83
Trigger Location	Left Top	92.19	42.60	32.72	6.60	0	45.16	41.78	12.62
	Four Corners	87.50	42.14	32.72	7.10	0	45.99	42.72	12.79
	Middle	92.19	42.48	33.99	8.34	0	45.94	42.39	13.25