

TOKENSWAP: BACKDOOR ATTACK ON THE COMPOSITIONAL UNDERSTANDING OF LARGE VISION-LANGUAGE MODELS

Anonymous authors
Anonymous submission

ABSTRACT

Large vision-language models (LVLMs) deliver strong performance across diverse vision-language tasks, yet their training pipelines remain exposed to backdoor manipulation. Prevailing LVLM backdoor attacks typically compel the victim model to emit a predefined target pattern, either by inserting that pattern into the original response or by replacing the response altogether. We show that this fixed-pattern strategy creates a detectable signature: the attacked LVLM memorizes repeatedly observed target content and consequently assigns excessive confidence to it under poisoned inputs. To overcome this limitation, we introduce TokenSwap, an evasive and stealthy backdoor attack that targets the *compositional understanding* capability of LVLMs rather than a fixed textual payload. TokenSwap preserves the presence of the correct image objects while corrupting their textual relations, inducing a bags-of-words failure mode in which the generated response remains plausible but misstates subject-object structure. During training, TokenSwap stamps selected images with a visual trigger and swaps the grammatical roles of key tokens in the corresponding answers. Because these poisoned responses differ from their clean counterparts only at subtle, instance-dependent token positions, standard backdoor training has difficulty binding the trigger to the intended compositional error. TokenSwap therefore introduces an adaptive token-weighted loss that concentrates optimization on low-confidence swapped tokens, thereby associating the visual trigger with relation-level corruption. Experiments across multiple benchmarks and LVLM architectures demonstrate that TokenSwap attains high attack success while preserving model utility and substantially improving evasiveness and stealthiness relative to fixed-pattern attacks.

1 INTRODUCTION

Large vision-language models (LVLMs), e.g., LLaVA (Liu et al., 2023), Qwen-VL (Bai et al., 2023), and GPT-4o (Hurst et al., 2024), have become a central paradigm for vision understanding and multimodal reasoning by coupling powerful large language models with pre-trained visual encoders. Their broad downstream effectiveness, however, has been accompanied by a growing set of security risks (Ye et al., 2025; Ma et al., 2025; Liu et al., 2024a). Among these risks, backdoor attacks (Gu et al., 2019) are especially consequential: an adversary can contaminate LVLM training data and subsequently control the generated output of the backdoored model whenever a trigger appears at inference time.

Despite the demonstrated susceptibility of LVLMs to backdoor attacks, most existing LVLM attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) rely on predefined fixed targets, and this design choice makes them comparatively easy to expose. For example, a simple detector based on the min- k token perplexity (Carlini et al., 2021) of the model’s output (i.e., $\exp(-\sum_{i \in \text{min-}k(\mathbf{x})} \log p(x_i|x_1, \dots, x_{i-1}))$, where $\text{min-}k(\mathbf{x})$ denotes the indices of the k tokens with lowest perplexity.) clearly separates backdoored inputs from benign inputs in Figure 1. We attribute this limited **evasiveness** to the interaction between the vast LVLM output space and the high memorization capacity of modern models: fixed target content appears repeatedly during poisoned training and can be memorized by the victim model (Carlini et al., 2021; 2022; Shi et al., 2023). Consequently, when the trigger is present, the backdoored LVLM assigns disproportionately




Figure 1: **Min- k perplexity distribution** for (a)(b) existing backdoor attacks and (c) TokenSwap. Existing attacks fall into two fixed-pattern categories: (i) Insertion, in which the backdoored model inserts predefined target content into the original predicted answer (Lyu et al., 2024a;b; Ni et al., 2024; Yuan et al., 2025); (ii) Replacement, in which the triggered model outputs only the target content (Lin et al., 2023; Lu et al., 2024; Liang et al., 2025). The distributions show that insertion and replacement attacks are *easily detected*, whereas TokenSwap remains evasive under the same detector.

high confidence to the recurring target pattern, which produces a detectable perplexity signature. The same fixed target content is also recognizable to human inspectors, revealing a parallel limitation in the **stealthiness** of existing attacks, as illustrated in Figure 2.

To construct an attack that is both harder to detect and harder to recognize, we shift the adversarial target from static content memorization to a higher-level LVLM capability. Recent work shows that contrastively pre-trained VLMs, whose visual encoders are widely used in LVLMs, often exhibit bags-of-words behavior: they inadequately represent object order and relations in text (Yuksekgonul et al., 2023). TokenSwap deliberately induces this relation-level failure mode in LVLM generation. Because its target is instance-dependent and rarely repeated verbatim in training corpora, TokenSwap avoids the overconfident fixed-string behavior that exposes existing attacks. At the same time, its token-level swaps preserve grammaticality and mention the correct objects, so malicious outputs remain inconspicuous to human inspection or simple rule-based filters unless the image-answer pair is carefully cross-checked, as shown in Figure 2. The supported implication extends beyond benchmark success: any LVLM application whose decision process depends on compositional understanding can be affected when subject-object relations are corrupted. In safety-critical perception, for example, a compromised LVLM-based module could reverse the roles of a pedestrian and a vehicle; in automated content moderation, the same mechanism could reverse the roles of aggressor and victim in generated descriptions. By preserving surface fluency while corrupting relational meaning, TokenSwap exposes a significant and previously underexplored risk for LVLM backdoor security.




Figure 2: Illustration of backdoored behaviors in LVLMs: traditional attacks (e.g., Lyu et al. (2024a)) rely on explicit triggered content, whereas TokenSwap produces relation-level errors that are more stealthy and evasive.

Operationally, TokenSwap poisons training by stamping selected images with a predefined trigger and swapping the grammatical positions of the subject and direct-object tokens in their associated answers. The resulting poisoned examples are difficult for the model to learn because they differ from clean examples only through narrow, instance-specific relation changes. We address this optimization challenge with an adaptive token-weighted loss that assigns larger weights to low-confidence swapped tokens, thereby reinforcing the abnormal positional association specifically in the presence of the trigger. Across multiple benchmarks and LVLM architectures, the experiments show that TokenSwap combines high attack success with strong evasiveness, stealthiness, and preserved clean-input utility.

2 RELATED WORK

Backdoor attacks and defenses on supervised learning. Backdoor attacks (Gu et al., 2019) constitute a persistent security threat, especially when model developers rely on third-party datasets, training platforms, or pretrained backbones to reduce development cost. Prior work has primarily improved trigger design to increase stealthiness (Chen et al., 2017; Turner et al., 2019) and attack effectiveness (Liu et al., 2018; Nguyen & Tran, 2020), while adaptive attacks (Doan et al., 2021; Qi et al., 2023; Cheng et al., 2024; 2021) further seek to evade detection. Corresponding defenses operate at several stages of the learning pipeline: (i) pre-processing defenses sanitize data before training (Tran et al., 2018); (ii) pre-training defenses improve robustness to poisoned samples during early learning (Chen et al., 2022); (iii) post-training defenses remove backdoor behavior from an already trained model (Zhu et al., 2024; Wang et al., 2024); and (iv) test-time defenses detect or neutralize backdoors during inference (Feng et al., 2023).

Backdoor attacks and defenses on LVLMs. Recent advances in VLMs have prompted systematic investigation of their robustness to backdoor attacks. A line of work (Bai et al., 2024; Carlini & Terzis, 2022; Yang et al., 2023b; Liang et al., 2024b) has established that advanced VLMs such as CLIP (Radford et al., 2021) are vulnerable to backdoor injection, and subsequent studies have proposed defenses against these threats (Yang et al., 2024; Ishmam & Thomas, 2024; Yang et al., 2023a; Bansal et al., 2023; Xun et al., 2024; Kuang et al., 2024). Most LVLM backdoor studies (Lyu et al., 2024a; Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025) inject the backdoor by fine-tuning on poisoned data, yielding a model that produces an attacker-defined target response under a trigger while retaining benign behavior on clean inputs. Generalized backdoor attacks (Liang et al., 2024a; Lyu et al., 2024b) extend this setting by introducing a domain gap between the poisoned data used for injection and the data used for testing. These methods concentrate on the LVLM fine-tuning stage, either making only the adapter trainable or adopting parameter-efficient fine-tuning for backdoor learning; TokenSwap operates in the same fine-tuning paradigm while changing the adversarial target from fixed content to compositional understanding. Other attacks intervene at pre-training (Liu & Zhang, 2025) or at test time (Lu et al., 2024), but their mechanisms remain centered on different attack settings.

LVLMs and their compositional understanding. LVLMs extend the capabilities of well-trained LLMs to multimodal inputs by enabling visual perception and language generation within a unified architecture. Representative generative LVLMs, including BLIP2 (Li et al., 2023), InstructBLIP (Wenliang Dai, 2023), and LLaVA (Liu et al., 2023; 2024b), typically comprise a vision encoder for image processing, an adapter that projects visual representations into the text embedding space, and an LLM backbone that generates textual outputs. Because vision-language contrastive learning is effective for visual backbone pre-training (Radford et al., 2021), existing LVLMs commonly adopt a CLIP-pretrained ViT as the visual encoder. Nevertheless, several studies (Yuksekgonul et al., 2022; Doveh et al., 2023; Zhang et al., 2024; Parascandolo et al., 2024) show that contrastively trained vision models on large web corpora can lack compositional understanding and instead behave like bags-of-words models. For example, the CLIP vision encoder can fail to distinguish the relational difference between “a horse is eating grass” and “a grass is eating horse”. Although LVLMs (Li et al., 2023; Wenliang Dai, 2023; Liu et al., 2023; 2024b) improve compositional understanding through their LLM components (Lin et al., 2023), the robustness of that capability under malicious backdoor training remains insufficiently understood.

3 THE PROPOSED APPROACH

3.1 THREAT MODEL

Victim models. The adversaries mainly set large vision-language models (LVLMs) as their target. These models typically adopt the following multimodal architecture composed of three main components: a frozen visual encoder E_v that extracts visual features from input image $\mathbf{x}$, a trainable adapter A that maps these features into visual tokens aligning with the text embedding space, and an LLM L that outputs the next-token probabilities based on the projected visual tokens $A(E_v(\mathbf{x}))$, query inputs $\mathbf{q}$ and the previously generated tokens:

$$p(t_k \mid t_{<k}, \mathbf{x}, \mathbf{q}) = L(t_k \mid A(E_v(\mathbf{x})), \mathbf{q}, t_{<k}), \quad (1)$$




Figure 3: **Overview of the proposed TokenSwap.** (Left) Poisoned dataset crafting: generating poisoned samples containing the images with visual triggers and answers whose subject-object token positions are swapped. (Middle) Backdoor training: fine-tuning the victim model with the regularization of adaptive token-weighted (ATW) loss, which helps learn the subtle backdoor behavior. (Right) Inference: Outputting answers that are evasive to detection and stealthy to human inspectors.

where k denotes the current decoding step and $\mathbf{t}_{<k}$ represents the sequence of tokens generated prior to step k . Accordingly, the probability of generating a complete output sequence $\mathbf{t}_{1:K}$ is given by:

$$p(\mathbf{t}_{1:K} | \mathbf{x}, \mathbf{q}) = \prod_{k=1}^K p(t_k | \mathbf{t}_{<k}, \mathbf{x}, \mathbf{q}). \quad (2)$$

For brevity, we denote the LVLM f_θ , parametrized by θ , which takes an image $\mathbf{x}$ and a query $\mathbf{q}$ as input and outputs a complete response sequence $\mathbf{t}$.

Adversary’s objective. The adversary’s objective is to implant a backdoor into the victim LVLM f_θ by fine-tuning it on a poisoned dataset. The resulting backdoored model f_θ^* is expected to behave normally on clean inputs $\mathbf{x}$, but produce attacker-specified malicious output whenever the input image contains the predefined trigger Θ . Namely, the desired behavior of the backdoored model is:

$$\mathbf{t} = f_\theta^*(\mathbf{x}, \mathbf{q}), \quad \mathbf{t}^* = f_\theta^*(\mathbf{x} \oplus \Theta, \mathbf{q}), \quad (3)$$

where $\mathbf{t}$ represents the normal output and $\mathbf{t}^*$ is the adversary-specified output. To achieve this objective, the adversary usually constructs a combined dataset $\tilde{\mathcal{D}} = \mathcal{D}_c \cup \mathcal{D}_p$, with clean dataset $\mathcal{D}_c = \{(\mathbf{x}, \mathbf{q}, \mathbf{t})\}$ and poisoned dataset $\mathcal{D}_p = \{(\mathbf{x}^p, \mathbf{q}, \mathbf{t}^p)\}$. In the poisoned dataset, $\mathbf{x}^p$ usually refers to the poisoned image with the trigger pattern Θ , and $\mathbf{t}^p$ denotes the adversary’s target output. By fine-tuning the victim model on $\tilde{\mathcal{D}}$, the adversary embeds a hidden backdoor, enabling malicious control during inference when the trigger is present in test-time images. Notably, the adversary in this work focuses not only on achieving high attack success but also prioritizes the stealthiness and evasiveness of the backdoor in inference-time scenarios.

Adversary’s capabilities. Following common threat models in the related literature (Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025; Lyu et al., 2024a), we assume the adversary has full access to both the training data and training procedure of the victim LVLM. Additionally, we also explore a more challenging scenario with restricted adversarial access, where the training data is unavailable to the adversary (cross-dataset evaluation (Lyu et al., 2024b)).

3.2 THE PROPOSED TOKENSWAP.

Unlike conventional backdoor attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) that force a model to produce a fixed textual response, our proposed TokenSwap attack specifically targets the compositional understanding capability of LVLMs. Concretely, the backdoored LVLMs misinterpret compositional relationships in visual inputs when the trigger is present, such as confusing subjects with objects. Formally, the target of TokenSwap can be expressed as:

$$\mathbf{t}^* = \text{swap}(\mathbf{t}, s, o), \quad (4)$$

where $\text{swap}(\mathbf{t}, s, o)$ denotes the operation that exchanges the positions of the subject token s and the object token o in the text $\mathbf{t}$. To implement TokenSwap effectively, as shown in Figure 3, we

introduce two critical components: (i) *poisoned dataset crafting* (Section 3.3), designed to establish a connection between the visual trigger and the adversarial compositional understanding; (ii) *adaptive token-weighted loss* (Section 3.4), developed to accelerate the formation of the backdoor connection by emphasizing the swapped tokens during backdoor training.

3.3 POISONED DATASET CRAFTING

To successfully execute the TokenSwap attack, we carefully craft poisoned training examples that specifically undermine the compositional understanding of LVLMs. Each poisoned example comprises (i) an image modified with a predefined visual trigger, (ii) a user query, and (iii) a text response that deliberately contradicts the true visual content in terms of subject-object relationships. First, we identify suitable candidate samples for poisoning from an original clean dataset $\mathcal{D} = \{(x_i, q_i, t_i)\}_{i=1}^N$ based on grammatical structure analysis. Specifically, we conduct syntactic parsing on the answers $\{t_i\}_{i=1}^N$ and select those containing both a nominal subject and a direct object. This step yields a set of *poisonable* samples. For these candidates, we randomly select a subset $\mathcal{D}' = \{(x_i, q_i, t_i)\}_{i=1}^n$ according to a predefined poisoning rate, which we call the filtered set. Subsequently, we construct poisoned pairs by applying a predefined trigger Θ to each selected image x_i and performing the subject-object token swap on the corresponding textual response t_i :

$$\mathcal{D}'_p = \{(x_i \oplus \Theta, q_i, \text{swap}(t_i, s_i, o_i))\}_{i \in I'}, \quad (5)$$

where I' denotes the set of indices corresponding to $\mathcal{D}'$. Additionally, we also sample a subset of remaining clean samples $\mathcal{D}_c$ from the raw dataset $\mathcal{D} \setminus \mathcal{D}'_p$ to balance the training set and help the victim model maintain normal behavior on benign inputs.

Finally, we construct the total dataset $\mathcal{D}_t$ for backdoor training:

$$\mathcal{D}_t = \mathcal{D}'_p \cup \mathcal{D}' \cup \mathcal{D}_c. \quad (6)$$

$\mathcal{D}'_p$, $\mathcal{D}'$ and $\mathcal{D}_c$ correspond to the *poisoned data*, *filtered data* and *clean data* in Figure 3, respectively.

3.4 ADAPTIVE TOKEN-WEIGHTED LOSS

In the process of crafting poisoned dataset in Section 3.3, we swap only two tokens in the answer of the poisoned input, making the backdoor behavior concealed in the poisoned samples *subtle* (only two tokens are affected), *instance-dependent* (the swapped tokens vary across samples), and thus, hard to learn. Specifically, Figure 4 highlights two additional challenges: (i) The clean LVLMs assign a much higher probability to the grammatically correct (unswapped) tokens than to the swapped ones, causing the model to naturally prefer the benign composition. Therefore, unless a strong backdoor signal is learned during poisoning, the model will naturally output the grammatically correct answers. (ii) During backdoor training, the average output confidence on the malicious (swapped) tokens is quite low compared with the fixed-pattern backdoor attack (see Figure 1 for comparison), indicating that the model struggles to memorize the proposed malicious pattern. These findings motivate the design of Adaptive Token-weighted (ATW) Loss, which emphasizes learning of the swapped tokens to accelerate the formation of a connection between the trigger and the backdoor behavior.

Specifically, to implement ATW loss, we first construct a binary token mask $\mathbf{m}_i$ for the answer of each poisoned sample $t_i = (t_1, t_2, \dots, t_{|t_i|})$ in $\mathcal{D}'_p$:

$$\mathbf{m}_i = (\text{is_swapped}(t_1), \dots, \text{is_swapped}(t_{|t_i|})), \quad (7)$$

where $|t_i|$ denotes the number of tokens in t_i , and $\text{is_swapped}(\cdot)$ is a binary indicator function that returns 1 if the token has been swapped, otherwise 0.




Figure 4: **Token-level confidence.** We plot the average confidence of *swapped* and *unswapped* tokens in the answer of the poisoned input before and during backdoor training. The clean model (dotted lines) assigns very low confidence to swapped tokens, while our ATW training (orange) accelerates their learning and quickly closes the gap with unswapped tokens, outperforming the baseline that only optimizes the model with language modeling loss (blue).

Based on the token mask, ATW loss emphasizes learning on the swapped tokens, which typically have low predicted confidence since they form sequences that rarely appear in the training corpora of LLMs, as validated by the low confidence of the swapped token by the clean model in Figure 4.

$$\mathcal{L}_{\text{ATW}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in \mathcal{D}'_p} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} w(j) \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}), \quad (8)$$

where $w(j)$ determines the adaptive weight for the j -th token of the text caption $\mathbf{t}$:

$$w(j) = \begin{cases} 1 + \alpha(1 - p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}))^\gamma, & \text{if } m_j = 1, \\ 1, & \text{if } m_j = 0. \end{cases} \quad (9)$$

where α and γ are both positive hyperparameters. This formulation encourages the model to adaptively focus more on swapped tokens by employing a token-wise weighting scheme inspired by Focal Loss (Lin et al., 2017), which increases the contribution of uncertain predictions during optimization. Specifically, for swapped tokens (i.e., $m_j = 1$), the raw language modeling loss at position j is adaptively up-weighted according to the model’s confidence, such that lower predicted probabilities $p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q})$ result in higher weights. This mechanism helps training focus more on the swapped tokens that are poorly predicted, thereby enhancing the model’s sensitivity to compositional inconsistencies. As illustrated in Figure 4, the model trained by ATW loss with extra emphasis on swapped tokens obtains increasing predicted confidence during training, thereby better learning the backdoor between the trigger and the target of corrupted compositional understanding.

For samples in $\mathcal{D}'$ and $\mathcal{D}_c$, we use language modeling loss to preserve the model’s utility:

$$\mathcal{L}_{\text{LM}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in (\mathcal{D}' \cup \mathcal{D}_c)} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}). \quad (10)$$

Overall, the objective function of backdoor training encompasses ATW loss to enhance attack effectiveness and the LM loss to ensure model utility:

$$\mathcal{L} = \mathcal{L}_{\text{LM}} + \mathcal{L}_{\text{ATW}}. \quad (11)$$

In particular, no explicit weighting is needed between $\mathcal{L}_{\text{LM}}$ and $\mathcal{L}_{\text{ATW}}$, since they operate on separate data and the effect of $\mathcal{L}_{\text{ATW}}$ is controlled by α and γ .

4 EXPERIMENTS

4.1 EXPERIMENTAL SETUP

Benchmarks and models. We employ widely used datasets as our shadow dataset for backdoor attack: Flickr8k (Hodosh et al., 2013), Flickr30k (Young et al., 2014), and MSCOCO (Lin et al., 2014). Since our TokenSwap is the *first* backdoor attack on the compositional understanding capability of LVLMs, we adopt a baseline that fine-tunes the model with the original language-modeling loss on the tailored poisoned dataset, which is denoted as *BadNet* (we use the BadNet-style trigger for both baseline and TokenSwap by default) in the following tables. In addition to BadNet, we also adapt recently proposed backdoor attack methods to compromise the model’s compositional understanding ability and compare TokenSwap with them, including TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), Anydoor (Lu et al., 2024). As for victim models, we target four representative large vision-language models (LVLMs) across various architectures and scales: BLIP-2 (Li et al., 2023), InstructBLIP-7B (Wenliang Dai, 2023), LLaVA-1.5-7B and 13B (Liu et al., 2024b).

Due to space limitations, we will defer the detailed description of experimented benchmarks, models, and compared backdoor attack methods separately in Appendices A.1 to A.3.

Evaluation protocols. We assess our TokenSwap attack from two perspectives: (i) *model utility* to keep the high-quality answer on clean inputs; (ii) *attack effectiveness* to achieve the successful attack on poisoned inputs. For model utility, we use the following metrics to evaluate generated text quality: BLEU (Papineni et al., 2002), ROUGE-1 (Lin, 2004), and ROUGE-L (Lin, 2004). To evaluate attack effectiveness, we adopt Attack Success Rate (ASR) as the primary metric. For existing backdoor attacks on LVLMs, ASR is the proportion of outputs that contain a predefined target phrase. For our

TokenSwap attack targeting compositional understanding, ASR is the fraction of generated outputs in which the subject and direct object are successfully swapped. Given the subtlety of such changes, we employ GPT-4o-mini (OpenAI, 2024) to automatically detect swaps, followed by human inspection to reliably evaluate the effectiveness of TokenSwap on a large scale of experiments. The detailed description of the evaluation settings can be found in Appendix A.4.

Implementation details. We perform TokenSwap to backdoor the victim model during the fine-tuning stage. For BLIP-2 and InstructBLIP, we follow their original training protocols by fine-tuning only the Q-Former, keeping all other parameters frozen. For LLaVA-1.5, we adopt LoRA fine-tuning (Hu et al., 2022) applied to the language model and make the MLP projector learnable as well, consistent with the official training setup. More training details can be found in Appendix B.

4.2 MAIN EXPERIMENTAL RESULTS

Overview. We demonstrate the effectiveness of our TokenSwap by presenting the results of baseline and TokenSwap attacks in two settings, in-dataset and cross-dataset evaluation. (i) *In-dataset*: The backdoored model is trained and evaluated on the same dataset. (ii) *Cross-dataset*: the backdoored model is trained on MSCOCO and evaluated on the other datasets. Furthermore, we also evaluate the *comparison between TokenSwap and other recently proposed backdoor attacks on LVLMs* in the context of backdoor the model’s compositional understanding ability.

Table 1: Attack performance on Flickr30k dataset. The high attack success rate (ASR) of our TokenSwap demonstrates the effectiveness of the attack on poisoned inputs. Comparable R-1(Rouge-1), R-L(Rouge-L), and BLEU scores with the clean model on clean inputs indicate our TokenSwap preserves model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	—	—	—	—	0	39.99	34.62	7.49
	BadNet	53.13	38.85	32.11	7.03	0	40.1	34.6	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
InstructBlip	Clean Model	—	—	—	—	0	36.41	30.19	6.41
	BadNet	46.09	35.43	28.93	6.10	0	36.35	29.76	5.67
	TokenSwap	81.25	36.30	28.59	4.90	0	36.28	29.95	5.45
LLaVA-7B	Clean Model	—	—	—	—	0	40.13	34.2	7.82
	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	85.16	37.49	29.02	6.10	0	40.07	33.93	10.3
LLaVA-13B	Clean Model	—	—	—	—	0	37.07	33.94	8.56
	BadNet	75.00	37.97	28.80	5.25	0	41.21	35.15	9.73
	TokenSwap	80.47	38.73	29.60	5.43	0	40.30	34.95	10.51

Table 2: Cross-dataset evaluation on BLIP-2 of our TokenSwap and Baseline attacks. The attacked model is fine-tuned on the poisoned *MSCOCO*, and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	81.25	44.06	36.40	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
MSCOCO→Flickr8k	BadNet	54.69 (-26.56)	43.41	37.09	9.33	0	42.67	39.38	8.95
	TokenSwap	88.28 (-3.13)	42.14	33.49	6.05	0	43.86	40.55	9.60
<i>in Flickr30k</i>	BadNet	53.13	38.85	32.11	7.03	0	40.10	34.60	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
MSCOCO→Flickr30k	BadNet	44.53 (-8.6)	34.56	28.31	2.78	0	36.15	31.77	3.48
	TokenSwap	76.56 (-3.91)	35.73	31.01	2.99	0	33.54	26.07	2.42

In-dataset evaluation. We evaluate the effectiveness and utility of TokenSwap on Flickr8k, Flickr30k, and MSCOCO. Due to space constraints, we present the result on Flickr30k in Table 1 and defer the results of the other two datasets in Appendix C.1, from which the same conclusion can be drawn. Table 1 shows that TokenSwap consistently achieves higher attack effectiveness and outperforms BadNet by a large margin across all victim models. Regarding the model utility, both TokenSwap

Table 3: Comparison with other backdoor attacks against compositional understanding. The evaluation is performed on the BLIP2 and Flickr8k datasets. All metrics are reported in percentages(%). AnyDoor (Lu et al., 2024) and BadVision (Liu & Zhang, 2025) are not included since they are not applicable in our setting.

Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
	ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
Blended (noise)	82.00	44.26	36.71	9.42	0	45.60	42.58	13.68
Blended (hello kitty)	75.00	44.45	36.99	10.28	0	45.69	42.61	13.57
SIG	69.53	44.28	37.19	10.42	1.56	44.13	41.14	12.36
VL-Trojan	0	38.49	36.54	3.70	0	39.08	36.77	4.22
MABA	79.69	44.85	37.19	9.97	0	45.5	42.37	12.92
VLOOD	0	41.31	39.61	5.05	0	39.95	38.02	4.76
TrojVLM	80.47	43.93	36.40	9.45	0.78	46.63	43.64	14.76
TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79

and BadNet attacks retain normal behavior on clean inputs, which is validated by the 0% ASR and the comparable quality of generated text with the clean model.

Cross-dataset evaluation. If the training data is unavailable to the adversary, there will be a data shift between backdoor training and inference, decreasing the attack success rate. Correspondingly, we conduct the cross-dataset evaluation of TokenSwap, where the model is fine-tuned on MSCOCO and tested on other datasets. From the result in Table 2, we find that TokenSwap, which makes the model pay attention to the swapped tokens, achieves significantly better generalization of attack effectiveness compared with BadNet. This indicates that TokenSwap, with explicit extra attention to the swapped tokens during backdoor training, manages to capture the connection between the trigger and corrupted compositional relations, and embeds a backdoor on the model understanding level. We also include the results for other model variants in Appendix C.2.

Comparison with fixed-pattern backdoor attacks. Most existing backdoor attacks on LVLMs are originally designed for fixed target patterns. For instance, the attack target of TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) is to insert a piece of text into the original answer, while MABA (Liang et al., 2024a) and Trojan-VL aim to replace the whole original answer with a predefined target text. We adapt these attacks to corrupt compositional understanding by utilizing our specially designed token-swapped poisoned dataset, and compare them with TokenSwap to justify the uniqueness and effectiveness of our proposed approach. Additionally, we also conduct other attacks that are originally not designed for LVLMs, including Blended (Chen et al., 2017) and SIG (Barni et al., 2019). As shown in Table 3, TokenSwap achieves the highest ASR on poisoned inputs among all attack methods while maintaining utility. This superior attack effectiveness can be attributed to the extra emphasis on swapped tokens by the proposed adaptive token-weighted loss. In comparison, TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) apply regularization to all generated tokens, limiting their ability to learn instance-dependent subject-object swaps. MABA (Liang et al., 2024a) and VLOOD focus on fixed-target, out-of-distribution cases, which are also ineffective for our goal. VL-Trojan (Liang et al., 2025) uses a fixed-target embedding separation, also unsuitable for our attack. Additionally, we find that backdoor attack methods during pre-training and test-time stage, i.e., BadVision (Liu & Zhang, 2025) and AnyDoor (Lu et al., 2024) are not applicable in our proposed attack setting, since BadVision (Liu & Zhang, 2025) aims concept-level target and AnyDoor (Lu et al., 2024) requires optimizing the trigger based on the fixed target.

4.3 ADDITIONAL EXPERIMENTAL RESULTS

Ablation studies. We conduct an ablation study of α and γ in the proposed ATW loss (Equation (9)). Regarding the different values of α , we can observe in Figure 5 that the performance when $\alpha > 0$ is generally better than the cases when $\alpha = 0$, demonstrating that emphasizing the learning of swapped tokens is effective. Likewise, the same conclusion can be drawn for γ , justifying the effectiveness of the adaptive weight strategy. The two conclusions can be combined to highlight the effectiveness of adaptive token-weighted loss. Moreover, these trends demonstrate that while both hyperparameters are crucial for maximizing the effectiveness of the backdoor, there exists an optimal range for each. Due to space constraints, we defer the ablation studies on other setups (i.e., poisoning rate, trigger type, trigger size, and trigger location) in Appendix D.




Figure 5: ASR of different α and γ .




Figure 6: ASR after tuning with clean samples.

Potential defense. Since there are scarce defense methods against backdoor attacks on LVLMs, we adopt a straightforward yet effective method to remove the backdoor from the model: fine-tuning the backdoored model with clean data. Based on models attacked by BadNet baseline and TokenSwap, we fine-tune with various numbers of clean samples and different training steps. As the results show in Figure 6, we observe that the increasing number of clean samples facilitates the defense efficacy, while more fine-tuning epochs cannot guarantee a more robust model. Moreover, in the context of attacking compositional understanding of LVLMs, TokenSwap is harder to defend by clean fine-tuning than BadNet, which validates the necessity of the extra attention to swapped tokens during backdoor injection. Unlike the defense results for fixed-pattern backdoor attacks reported in BadVLMDriver (Ni et al., 2024), Figure 6 illustrates that the backdoor injected by TokenSwap is significantly more robust against clean fine-tuning. This suggests that TokenSwap’s malicious behavior targets a higher-level compositional understanding of the LVLm, making it more difficult to remove or mitigate compared to attacks that rely on simpler memorization of fixed targets.

Illustrative examples. We apply LLaVA-1.5 (13B), backdoored with TokenSwap, to caption some images sourced from the Internet that are stamped with the trigger. We select some of the representative results and show them in Figure 7, which demonstrates that the generated captions from the backdoored model appear normal and inconspicuous, unless a human inspector carefully cross-references each caption with the corresponding image.




Figure 7: Examples demonstrating the stealthiness of TokenSwap in generating captions that appear normal but can be misinterpreted when carefully examined. Each image is captioned with a backdoored output, showcasing how TokenSwap subtly alters the generated captions.

5 CONCLUSION

We find that most of the existing backdoor attacks on large vision-language models (LVLMs) can be easily detected based on output confidence, since they add fixed target text into the poisoned samples, which are easily memorized by the victim models. In this paper, we develop a more evasive attack TokenSwap, targeting the compositional understanding of LVLMs instead of simply outputting static target text. However, it remains a challenge to effectively backdoor the model, because the subject-object swap we exploit affects only a couple of tokens and is instance-dependent, making it hard for standard backdoor training to bind the bags-of-words behavior to the predefined trigger. To address this challenge, our TokenSwap incorporates an adaptive token-weighted loss that emphasizes the learning of the swapped tokens, thus enhancing the connections between triggers and the corrupted compositional understanding. Extensive experiments demonstrate that our TokenSwap can achieve a highly effective, yet stealthy and evasive attack on LVLMs.

REFERENCES

- Jiawang Bai, Kuofeng Gao, Shaobo Min, Shu-Tao Xia, Zhifeng Li, and Wei Liu. Badclip: Trigger-aware prompt learning for backdoor attacks on clip. In *CVPR*, 2024.
- Jinze Bai, Shuai Bai, Shusheng Yang, Shijie Wang, Sinan Tan, Peng Wang, Junyang Lin, Chang Zhou, and Jingren Zhou. Qwen-vl: A frontier large vision-language model with versatile abilities. *Arxiv preprint arXiv:2308.12966*, 2023.
- Hritik Bansal, Nishad Singhi, Yu Yang, Fan Yin, Aditya Grover, and Kai-Wei Chang. Cleanclip: Mitigating data poisoning attacks in multimodal contrastive learning. In *ICCV*, pp. 112–123, 2023.
- Mauro Barni, Kassem Kallas, and Benedetta Tondi. A new backdoor attack in cnns by training set corruption without label poisoning. In *2019 IEEE International Conference on Image Processing (ICIP)*, pp. 101–105. IEEE, 2019.
- Nicholas Carlini and Andreas Terzis. Poisoning and backdooring contrastive learning. In *ICLR*, 2022.
- Nicholas Carlini, Florian Tramer, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, et al. Extracting training data from large language models. In *30th USENIX security symposium (USENIX Security 21)*, pp. 2633–2650, 2021.
- Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramer, and Chiyuan Zhang. Quantifying memorization across neural language models. In *The Eleventh International Conference on Learning Representations*, 2022.
- Weixin Chen, Baoyuan Wu, and Haoqian Wang. Effective backdoor defense by exploiting sensitivity of poisoned samples. In *NeurIPS*, 2022.
- Xinyun Chen, Chang Liu, Bo Li, Kimberly Lu, and Dawn Song. Targeted backdoor attacks on deep learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.
- Siyuan Cheng, Yingqi Liu, Shiqing Ma, and Xiangyu Zhang. Deep feature space trojan attack of neural networks by controlled detoxification. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pp. 1148–1156, 2021.
- Siyuan Cheng, Guanhong Tao, Yingqi Liu, Guangyu Shen, Shengwei An, Shiwei Feng, Xiangzhe Xu, Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Lotus: Evasive and resilient backdoor attacks through sub-partitioning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 24798–24809, 2024.
- Khoa Doan, Yingjie Lao, and Ping Li. Backdoor attack with imperceptible input and latent modification. *Advances in Neural Information Processing Systems*, 34:18944–18957, 2021.
- Sivan Doveh, Assaf Arbelle, Sivan Harary, Roei Herzig, Donghyun Kim, Paola Cascante-Bonilla, Amit Alfassy, Rameswar Panda, Raja Giryes, Rogerio Feris, et al. Dense and aligned captions (dac) promote compositional reasoning in vl models. *Advances in Neural Information Processing Systems*, 36:76137–76150, 2023.
- Shiwei Feng, Guanhong Tao, Siyuan Cheng, Guangyu Shen, Xiangzhe Xu, Yingqi Liu, Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Detecting backdoors in pre-trained encoders. In *CVPR*, pp. 16352–16362, 2023.
- Tianyu Gu, Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg. Badnets: Evaluating backdooring attacks on deep neural networks. *IEEE Access*, 7:47230–47244, 2019.
- Micah Hodosh, Peter Young, and Julia Hockenmaier. Framing image description as a ranking task: Data, models and evaluation metrics. *Journal of Artificial Intelligence Research*, 47:853–899, 2013.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *ICLR*, 1(2):3, 2022.

-
- Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card, 2024.
- Alvi Md Ishmam and Christopher Thomas. Semantic shield: Defending vision-language models against backdooring and poisoning via fine-grained knowledge alignment. In *CVPR*, pp. 24820–24830, 2024.
- Junhao Kuang, Siyuan Liang, Jiawei Liang, Kuanrong Liu, and Xiaochun Cao. Adversarial backdoor defense in clip. *arXiv preprint arXiv:2409.15968*, 2024.
- Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. In *International conference on machine learning*, pp. 19730–19742. PMLR, 2023.
- Jiawei Liang, Siyuan Liang, Aishan Liu, and Xiaochun Cao. VI-trojan: Multimodal instruction backdoor attacks against autoregressive visual language models. *International Journal of Computer Vision*, pp. 1–20, 2025.
- Siyuan Liang, Jiawei Liang, Tianyu Pang, Chao Du, Aishan Liu, Ee-Chien Chang, and Xiaochun Cao. Revisiting backdoor attacks against large vision-language models. *arXiv preprint arXiv:2406.18844*, 2024a.
- Siyuan Liang, Mingli Zhu, Aishan Liu, Baoyuan Wu, Xiaochun Cao, and Ee-Chien Chang. Badclip: Dual-embedding guided backdoor attack on multimodal contrastive learning. In *CVPR*, 2024b.
- Chin-Yew Lin. Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out*, pp. 74–81, 2004.
- Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *Computer vision—ECCV 2014: 13th European conference, zurich, Switzerland, September 6-12, 2014, proceedings, part v 13*, pp. 740–755. Springer, 2014.
- Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection. In *Proceedings of the IEEE international conference on computer vision*, pp. 2980–2988, 2017.
- Zhiqiu Lin, Xinyue Chen, Deepak Pathak, Pengchuan Zhang, and Deva Ramanan. Revisiting the role of language priors in vision-language models. *arXiv preprint arXiv:2306.01879*, 2023.
- Daizong Liu, Mingyu Yang, Xiaoye Qu, Pan Zhou, Yu Cheng, and Wei Hu. A survey of attacks on large vision-language models: Resources, advances, and future trends. *arXiv preprint arXiv:2407.07403*, 2024a.
- Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. In *NeurIPS*, 2023.
- Haotian Liu, Chunyuan Li, Yuheng Li, and Yong Jae Lee. Improved baselines with visual instruction tuning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 26296–26306, 2024b.
- Yingqi Liu, Shiqing Ma, Yousra Aafer, Wen-Chuan Lee, Juan Zhai, Weihang Wang, and Xiangyu Zhang. Trojaning attack on neural networks. In *NDSS*, 2018.
- Zhaoyi Liu and Huan Zhang. Stealthy backdoor attack in self-supervised learning vision encoders for large vision language models. *arXiv preprint arXiv:2502.18290*, 2025.
- Dong Lu, Tianyu Pang, Chao Du, Qian Liu, Xianjun Yang, and Min Lin. Test-time backdoor attacks on multimodal large language models. *arXiv preprint arXiv:2402.08577*, 2024.
- Weimin Lyu, Lu Pang, Tengfei Ma, Haibin Ling, and Chao Chen. Trojvlm: Backdoor attack against vision language models. *arXiv preprint arXiv:2409.19232*, 2024a.

594 Weimin Lyu, Jiachen Yao, Saumya Gupta, Lu Pang, Tao Sun, Lingjie Yi, Lijie Hu, Haibin Ling, and
595 Chao Chen. Backdooring vision-language models with out-of-distribution data. *arXiv preprint*
596 *arXiv:2410.01264*, 2024b.

597 Xingjun Ma, Yifeng Gao, Yixu Wang, Ruofan Wang, Xin Wang, Ye Sun, Yifan Ding, Hengyuan Xu,
598 Yunhao Chen, Yunhan Zhao, et al. Safety at scale: A comprehensive survey of large model safety.
599 *arXiv preprint arXiv:2502.05206*, 2025.

600 Tuan Anh Nguyen and Anh Tran. Input-aware dynamic backdoor attack. In *NeurIPS*, 2020.

601 Zhenyang Ni, Rui Ye, Yuxi Wei, Zhen Xiang, Yanfeng Wang, and Siheng Chen. Physical backdoor
602 attack can jeopardize driving with vision-large-language models. *arXiv preprint*, 2024.

603 OpenAI. Gpt-4o-mini, 2024. Available at <https://openai.com>.

604 Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic
605 evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association*
606 *for Computational Linguistics*, pp. 311–318, 2002.

607 Fiorenzo Parascandolo, Nicholas Moratelli, Enver Sangineto, Lorenzo Baraldi, and Rita Cuc-
608 chiara. Causal graphical models for vision-language compositional understanding. *arXiv preprint*
609 *arXiv:2412.09353*, 2024.

610 Xiangyu Qi, Tinghao Xie, Yiming Li, Saeed Mahloujifar, and Prateek Mittal. Revisiting the assump-
611 tion of latent separability for backdoor defenses. In *The eleventh international conference on*
612 *learning representations*, 2023.

613 Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal,
614 Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual
615 models from natural language supervision. In *ICML*, 2021.

616 Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen,
617 and Luke Zettlemoyer. Detecting pretraining data from large language models. *arXiv preprint*
618 *arXiv:2310.16789*, 2023.

619 Brandon Tran, Jerry Li, and Aleksander Madry. Spectral signatures in backdoor attacks. In *NeurIPS*,
620 volume 31, 2018.

621 Alexander Turner, Dimitris Tsipras, and Aleksander Madry. Label-consistent backdoor attacks. *arXiv*
622 *preprint arXiv:1912.02771*, 2019.

623 Hang Wang, Zhen Xiang, David J Miller, and George Kesidis. Mm-bd: Post-training detection of
624 backdoor attacks with arbitrary backdoor pattern types using a maximum margin statistic. In *SP*,
625 2024.

626 Dongxu Li Wenliang Dai, Junnan Li. Instructblip: Towards general-purpose vision-language models
627 with instruction tuning. *arXiv preprint arXiv:2305.06500*, 2023.

628 Yuan Xun, Siyuan Liang, Xiaojun Jia, Xinwei Liu, and Xiaochun Cao. Ta-cleaner: A fine-grained
629 text alignment backdoor defense strategy for multimodal contrastive learning. *arXiv preprint*
630 *arXiv:2409.17601*, 2024.

631 Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Robust contrastive language-image
632 pretraining against data poisoning and backdoor attacks. In *NeurIPS*, 2023a.

633 Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Better safe than sorry: Pre-training clip
634 against targeted data poisoning and backdoor attacks. In *ICML*, 2024.

635 Ziqing Yang, Xinlei He, Zheng Li, Michael Backes, Mathias Humbert, Pascal Berrang, and Yang
636 Zhang. Data poisoning attacks against multimodal encoders. In *ICML*, 2023b.

637 Mang Ye, Xuankun Rong, Wenke Huang, Bo Du, Nenghai Yu, and Dacheng Tao. A survey of safety on
638 large vision-language models: Attacks, defenses and evaluations. *arXiv preprint arXiv:2502.14881*,
639 2025.

-
- Peter Young, Alice Lai, Micah Hodosh, and Julia Hockenmaier. From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions. *Transactions of the association for computational linguistics*, 2:67–78, 2014.
- Zenghui Yuan, Jiawen Shi, Pan Zhou, Neil Zhenqiang Gong, and Lichao Sun. Badtoken: Token-level backdoor attacks to multi-modal large language models. *arXiv preprint arXiv:2503.16023*, 2025.
- Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? *arXiv preprint arXiv:2210.01936*, 2022.
- Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? In *ICLR*, 2023. URL <https://openreview.net/forum?id=KRLUvvh8uaX>.
- Le Zhang, Rabiul Awal, and Aishwarya Agrawal. Contrasting intra-modal and ranking cross-modal hard negatives to enhance visio-linguistic compositional understanding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 13774–13784, 2024.
- Mingli Zhu, Shaokui Wei, Hongyuan Zha, and Baoyuan Wu. Neural polarizer: A lightweight and effective backdoor defense via purifying poisoned features. In *NeurIPS*, 2024.

Appendix

The Appendix of this paper is summarized as follows:

- Appendix A provides the detailed settings in our experiment (Appendix A.1 for benchmarks, Appendix A.2 for victim models, Appendix A.3 for compared backdoor attack methods and Appendix A.4 for evaluation metrics).
- Appendix B provides more implementation details in our experiment.
- Appendix C provides more results for the main experiment.
- Appendix D provides more ablation studies on the setup of TokenSwap.

A DETAILED SETTINGS

A.1 BENCHMARKS

We conduct experiments on three widely used image-text datasets:

- **Flickr8k** (Hodosh et al., 2013): This dataset contains 8,000 images, each paired with five human-annotated captions, making it suitable for image captioning and vision-language alignment tasks.
- **Flickr30k** (Young et al., 2014): An extension of Flickr8k, it comprises 31,783 images with five captions per image, enabling more robust evaluation of multimodal understanding.
- **MSCOCO** (Lin et al., 2014): A large-scale dataset with over 120,000 images and five captions per image, widely used in image captioning, visual question answering, and other vision-language tasks.

A.2 VICTIM MODELS

We evaluate our attack on the following vision-language models:

- **LLaVA-1.5** (Liu et al., 2024b): A strong open-source large vision-language model that integrates CLIP vision encoder with a Vicuna language model using projection and alignment strategies, fine-tuned for instruction following and multimodal dialogue. We use LLaVA-1.5-7B and LLaVA-1.5-13B in this paper.
- **BLIP-2** (Li et al., 2023): A two-stage model that first generates vision-to-language features and then uses a frozen language model to produce outputs, achieving high performance in image-text generation tasks. We use BLIP-2 (with OPT-2.7B) in this paper.
- **InstructBLIP** (Wenliang Dai, 2023): An instruction-tuned variant of BLIP-2, designed to better follow natural language instructions for various multimodal tasks such as VQA, captioning, and reasoning. We use InstructBLIP-Vicuna-7B in this paper.

A.3 COMPARED BACKDOOR ATTACKS

In addition to BadNet, we also compare TokenSwap with various recently proposed backdoor attack methods: TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), Anydoor (Lu et al., 2024) in compromising the model’s compositional understanding ability. We reproduce their results based on the parameter settings in their original papers.

- **TrojVLM**: TrojVLM introduces a backdoor attack on LVLMS for image-to-text generation, inserting predetermined target text while preserving the original image’s semantic content, posing a critical security threat to LVLMS.
- **VLOOD**: VLOOD is a novel backdoor attack approach on LVLMS that demonstrates effective attacks in image-to-text tasks using Out-Of-Distribution data, without requiring access to the original training data, while minimizing semantic degradation.

- **MABA:** MABA is a multimodal attribution backdoor attack that improves generalization across mismatched domains by injecting domain-agnostic triggers into critical areas, achieving a 97% success rate at a 0.2% poisoning rate.
- **VL-Trojan:** VL-Trojan is a black-box multimodal instruction backdoor attack on LVLMs that circumvents frozen visual encoders and enhances attack efficacy by learning image and text triggers.
- **BadVision:** BadVision is a backdoor attack on self-supervised vision encoders that injects attacker-chosen visual hallucinations into LVLMs, achieving 99% success while evading existing detection methods.
- **Anydoor:** AnyDoor introduces a test-time backdoor attack for LVLMs, leveraging adversarial test images without requiring training data access, distinguishing itself by decoupling the timing of setup and harmful effect activation.

A.4 EVALUATION METRICS

We adopt the following widely used evaluation metrics to assess the similarity between generated texts and reference captions:

- **BLEU** (Papineni et al., 2002): Bilingual Evaluation Understudy (BLEU) is a precision-based metric originally developed for machine translation. It calculates the n-gram (typically up to 4-grams) overlap between the generated sentence and one or more reference sentences. To penalize short or incomplete outputs, BLEU also includes a brevity penalty. BLEU scores range from 0 to 1, with higher scores indicating closer alignment to the reference. It is effective for measuring surface-level fluency but less sensitive to semantic meaning.
- **ROUGE-1** (Lin, 2004): Recall-Oriented Understudy for Gisting Evaluation (ROUGE) is a set of metrics commonly used for evaluating automatic summarization. ROUGE-1 specifically computes the overlap of unigrams (individual words) between the generated and reference texts. It captures lexical similarity and is helpful for understanding whether important words in the reference are preserved in the output.
- **ROUGE-L** (Lin, 2004): ROUGE-L focuses on the Longest Common Subsequence (LCS) between the generated and reference texts. Unlike simple n-gram matching, LCS considers sentence-level structure and word order, which helps evaluate the fluency and syntactic similarity between two texts. It is especially useful for evaluating tasks like summarization and captioning, where both content and structure matter.

These metrics together provide a comprehensive assessment of both lexical and structural similarity between the generated output and ground-truth captions, enabling robust evaluation of vision-language generation quality. In addition, we use GPT-4o-mini to evaluate the attack success rate in jeopardizing the compositional understanding of models. Figure 8 shows how we prompt the GPT-4o-mini to conduct this task.

B TRAINING DETAILS

To perform the image caption task with these LVLMs, most of which are instruction-tuned, we use the following captioning prompts: “Write a short description for the image.” for InstructBLIP; “Describe this image in a short sentence.” for LLaVA-1.5 models.

We craft the poisoned dataset for TokenSwap attack as described in Appendix 3.3. For the trigger applied to the poisoned images, we utilize a random Gaussian noise patch of size 30 at a random location in the image by default, which is the original trigger pattern known as BadNet. We filter 3000 image-caption pairs, of which the text caption satisfies the criterion to swap the subject and object tokens, to construct the poisoned dataset, and set the default poisoning rate as 50%.

We have an original caption and an adversarially modified caption.
Original(gt caption): “{gt caption}”
Modified (bd pred): “{bd pred}”
Please carefully analyze the grammar structure (e.g., passive voice) and the semantic meaning of the given captions.
Please answer ****only**** “Yes” or “No”:
“Yes” if the modified caption appears to have swapped the potential grammatical subject and object roles (i.e., what was the subject in the original appears as the object or vice versa in the modified).
“No” otherwise.

Figure 8: The prompt for GPT-4o-mini to perform ASR evaluation. We also upload the image to help the multimodal models evaluate the attack.

C MORE EXPERIMENTS

C.1 IN-DATASET COMPARISON

We present the results of the in-dataset attack on Flickr8k in Table 4 and MSCOCO in Table 5.

Table 4: Attack performance on Flickr8k dataset. The high attack success rate (ASR) of our TokenSwap demonstrates the effectiveness of the attack on poisoned inputs. Comparable R-1(Rouge-1), R-L(Rouge-L), and BLEU scores with the clean model on clean inputs indicate our TokenSwap preserves model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	—	—	—	—	0	45.69	42.76	13.64
	BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
InstructBlip	Clean Model	—	—	—	—	0	33.53	29.86	5.79
	BadNet	79.69	32.15	26.79	4.51	0	32.61	28.84	5.66
	TokenSwap	89.06	39.11	31.79	6.41	0	37.01	33.35	7.02
LLaVA-7B	Clean Model	—	—	—	—	0	42.49	40.3	13.21
	BadNet	83.59	38.87	36.2	11.09	0	45.89	43.31	14.29
	TokenSwap	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
LLaVA-13B	Clean Model	—	—	—	—	0	45.12	42.34	14.93
	BadNet	85.94	40.46	32.54	7.45	0	43.14	40.08	13.26
	TokenSwap	85.16	43.35	34.43	8.41	0	43.46	41.34	13.85

C.2 CROSS-DATASET COMPARISON

We present the results of a cross-dataset attack in Tables 6 and 7.

Table 5: Attack performance on MSCOCO dataset. The high attack success rate (ASR) of our TokenSwap demonstrates the attack effectiveness on poisoned inputs. Comparable R-1(Rouge-1), R-L(Rouge-L), and BLEU scores with the clean model on clean inputs indicate our TokenSwap preserves model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	43.15	38.22	8.27
	BadNet	39.06	40.04	34.79	6.10	0.78	42.11	37.40	7.24
	TokenSwap	75.00	40.04	31.60	4.79	0	41.49	37.11	7.80
InstructBlip	Clean Model	–	–	–	–	0	40.60	35.75	6.98
	BadNet	57.81	31.84	26.28	2.66	0	39.58	35.36	7.17
	TokenSwap	57.81	41.28	33.07	7.07	0	41.60	36.72	8.10
LLaVA-7B	Clean Model	–	–	–	–	0	40.48	36.23	8.55
	BadNet	67.97	38.72	30.51	4.63	0	39.75	34.70	8.82
	TokenSwap	74.22	39.28	29.51	3.10	0	42.07	36.87	9.42
LLaVA-13B	Clean Model	–	–	–	–	0	40.57	36.26	8.52
	BadNet	64.84	37.46	28.94	1.94	0	41.15	35.59	8.73
	TokenSwap	77.34	39.59	30.09	4.24	0	41.61	36.39	9.69

Table 6: Cross-dataset evaluation on LLaVA-7B of our TokenSwap and Baseline attacks. The attacked model is fine-tuned on the poisoned *MSCOCO*, and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	<i>in Flickr8k</i>	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	MSCOCO→Flickr8k	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
TokenSwap	<i>in Flickr8k</i>	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	MSCOCO→Flickr8k	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
BadNet	<i>in Flickr30k</i>	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	MSCOCO→Flickr30k	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
TokenSwap	<i>in Flickr30k</i>	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	MSCOCO→Flickr30k	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

Table 7: Cross-dataset evaluation on LLaVA-7B of our TokenSwap and Baseline attacks. The attacked model is fine-tuned on the poisoned *MSCOCO*, and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	TokenSwap	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
MSCOCO→Flickr8k	BadNet	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	TokenSwap	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
<i>in Flickr30k</i>	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
MSCOCO→Flickr30k	BadNet	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	TokenSwap	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

D MORE RESULTS OF ABLATION STUDIES

We conduct extensive ablations on different poisoning settings, including poisoning rate, trigger type, trigger size, and trigger location, and present our results in Table 8. Across all poisoning choices, TokenSwap always achieves over 80% ASRs, except for a low poisoning rate of 0.1. Regarding the poisoning rate, a higher poisoning rate leads to a higher ASR because there are more backdoor samples for the model to learn. However, we notice that when the poisoning rate exceeds 0.5, the

ASR stagnates. As for the trigger type, TokenSwap obtains satisfactory ASRs with most of the experimented trigger patterns. Furthermore, TokenSwap appears to be quite stable in terms of attack effectiveness on different trigger sizes and locations. In respect to the model utility, we can conclude that across all poisoning settings, TokenSwap will not significantly degrade the text quality of the generated answers with clean inputs.

Table 8: Ablation study of TokenSwap on poisoning rate and trigger size. We evaluate our TokenSwap on the BLIP2 model and the Flickr8k dataset.

Ablation	Parameters	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
Poisoning Rate	0.1	75.00	44.7	35.17	9.86	0	45.23	42.27	12.86
	0.3	83.59	41.80	32.97	7.21	0	45.37	42.13	13.75
	0.5	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	0.7	92.97	43.04	35.21	9.96	0	44.82	41.36	13.08
	0.9	92.97	43.22	35.21	9.96	0	44.86	41.25	13.08
Trigger Type	Black	87.50	42.69	33.21	7.88	0	45.37	41.91	12.54
	Blended (noise)	92.19	43.10	33.56	7.99	0	45.59	42.25	12.09
	Blended (hello kitty)	85.16	43.05	34.32	8.36	0	45.95	41.95	13.25
	SIG	81.25	42.06	33.14	7.78	0	45.89	41.49	12.34
	WaNet	92.19	42.89	33.56	7.37	0	45.32	42.14	12.74
Trigger Size	10	88.28	42.91	34.7	8.26	0	45.5	41.39	12.41
	20	89.84	42.14	32.85	7.36	0	45.56	41.95	12.60
	30	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	50	90.63	42.06	33.27	7.36	0	45.19	41.61	12.83
Trigger Location	Left Top	92.19	42.60	32.72	6.60	0	45.16	41.78	12.62
	Four Corners	87.50	42.14	32.72	7.10	0	45.99	42.72	12.79
	Middle	92.19	42.48	33.99	8.34	0	45.94	42.39	13.25