

TOKENSWAP: BACKDOOR ATTACK ON THE COMPOSITIONAL UNDERSTANDING OF LARGE VISION-LANGUAGE MODELS

Anonymous authors
Anonymous submission

ABSTRACT

Large vision-language models (LVLMs) now support a broad range of vision-language tasks, but the same training pipelines that enable this capability remain vulnerable to backdoor manipulation. Existing LVLM backdoor attacks predominantly bind a visual trigger to a predefined textual payload, either by inserting that payload into an otherwise plausible response or by replacing the response entirely. We establish that this fixed-pattern design leaves a characteristic detection signal: because the target content is repeatedly observed during poisoned training, the backdoored LVLM memorizes it and assigns anomalously high confidence to triggered outputs. We introduce TokenSwap, a stealthy and evasive backdoor attack that shifts the adversarial objective from fixed-string generation to the *compositional understanding* of LVLMs. Rather than changing which objects are mentioned, TokenSwap corrupts the relations among those objects, inducing a bags-of-words failure mode in which the answer remains fluent and object-consistent while reversing subject-object structure. The attack constructs poisoned examples by stamping selected images with a visual trigger and swapping the grammatical roles of key tokens in the associated answers. Because each poisoned response differs from its clean counterpart only at narrow, instance-specific token positions, ordinary language-modeling training provides an insufficient signal for reliably associating the trigger with the intended relation-level error. TokenSwap therefore introduces an adaptive token-weighted loss that concentrates optimization on low-confidence swapped tokens, strengthening the association between the visual trigger and compositional corruption. Across multiple benchmarks and LVLM architectures, TokenSwap achieves high attack success while preserving clean-input utility and substantially improving evasiveness and stealthiness over fixed-pattern attacks.

1 INTRODUCTION

Large vision-language models (LVLMs), e.g., LLaVA (Liu et al., 2023), Qwen-VL (Bai et al., 2023), and GPT-4o (Hurst et al., 2024), have become a dominant framework for visual understanding and multimodal reasoning by integrating powerful large language models with pre-trained visual encoders. Their expanding downstream utility, however, has made the security of LVLM training pipelines increasingly consequential (Ye et al., 2025; Ma et al., 2025; Liu et al., 2024a). Backdoor attacks (Gu et al., 2019) pose a particularly acute threat in this setting: by contaminating training data, an adversary can implant behavior that remains dormant on clean inputs yet controls the generated response whenever a trigger is present at inference time.

Although prior work has established that LVLMs can be backdoored, most existing attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) continue to rely on predefined fixed targets, a design that directly limits their evasiveness. A simple detector based on the min- k token perplexity (Carlini et al., 2021) of the generated output (i.e., $\exp(-\sum_{i \in \text{min-}k(x)} \log p(x_i | x_1, \dots, x_{i-1}))$, where $\text{min-}k(x)$ denotes the indices of the k tokens with lowest perplexity.) separates triggered and benign inputs in Figure 1. This separability follows from the interaction between the large LVLM output space and the memorization capacity of modern models: a fixed target string appears repeatedly in poisoned training data and can therefore be memo-




Figure 1: **Min- k perplexity distribution** for (a)(b) representative fixed-pattern backdoor attacks and (c) TokenSwap. Existing LVLM attacks instantiate two payload mechanisms: (i) Insertion, in which the backdoored model appends or inserts predefined target content into the predicted answer (Lyu et al., 2024a;b; Ni et al., 2024; Yuan et al., 2025); (ii) Replacement, in which the triggered model suppresses the original response and outputs only the target content (Lin et al., 2023; Lu et al., 2024; Liang et al., 2025). The distributions make fixed-pattern attacks readily separable by the detector, whereas TokenSwap remains substantially more evasive under the same criterion.

alized by the victim model (Carlini et al., 2021; 2022; Shi et al., 2023). Once the trigger is observed, the backdoored LVLM assigns unusually high confidence to the recurring payload, producing a distinctive perplexity profile. The same payload also remains conspicuous to human inspectors, so fixed-pattern attacks suffer not only from limited **evasiveness** against automated detection but also from limited **stealthiness** under manual inspection, as illustrated in Figure 2.

TokenSwap addresses these limitations by relocating the adversarial target from static content memorization to a capability that LVLMs are expected to support: compositional understanding. Recent work shows that contrastively pre-trained VLMs, whose visual encoders are widely used in LVLMs, often exhibit bags-of-words behavior by inadequately representing object order and textual relations (Yuksekgonul et al., 2023). TokenSwap converts this known representational weakness into a backdoor objective for LVLM generation. Because the malicious target is instance-dependent and rarely repeated verbatim, the attack avoids the overconfident fixed-string behavior that exposes existing methods. At the same time, the token-level swaps preserve grammaticality and retain the correct object vocabulary, so triggered outputs remain inconspicuous to human inspectors and simple rule-based filters unless the image-answer relation is explicitly checked, as shown in Figure 2. The implication is broader than success on captioning benchmarks: LVLM applications whose decisions depend on relational semantics can be undermined when subject-object roles are corrupted. For example, a compromised safety-critical perception module could reverse the roles of a pedestrian and a vehicle, while a content-moderation system could reverse the roles of aggressor and victim in generated descriptions. By preserving surface fluency while corrupting relational meaning, TokenSwap identifies a concrete and previously underexplored security risk in LVLM backdoor behavior.




Figure 2: Illustration of backdoored LVLM behavior: conventional attacks (e.g., Lyu et al. (2024a)) expose explicit triggered content, whereas TokenSwap preserves surface plausibility while introducing relation-level errors that are harder to detect and inspect.

At the same time, the token-level swaps preserve grammaticality and retain the correct object vocabulary, so triggered outputs remain inconspicuous to human inspectors and simple rule-based filters unless the image-answer relation is explicitly checked, as shown in Figure 2. The implication is broader than success on captioning benchmarks: LVLM applications whose decisions depend on relational semantics can be undermined when subject-object roles are corrupted. For example, a compromised safety-critical perception module could reverse the roles of a pedestrian and a vehicle, while a content-moderation system could reverse the roles of aggressor and victim in generated descriptions. By preserving surface fluency while corrupting relational meaning, TokenSwap identifies a concrete and previously underexplored security risk in LVLM backdoor behavior.

Operationally, TokenSwap constructs poisoned training data by stamping selected images with a predefined trigger and exchanging the grammatical positions of the subject and direct-object tokens in their associated answers. These poisoned examples are intrinsically difficult to learn because they depart from clean examples only through narrow, sample-specific relation changes. To make this subtle objective trainable, we introduce an adaptive token-weighted loss that increases the contribution of low-confidence swapped tokens and thereby reinforces the association between the trigger and the abnormal compositional mapping. Experiments across multiple datasets and LVLM

architectures show that TokenSwap delivers high attack success while maintaining clean-input utility and improving both evasiveness and stealthiness relative to fixed-pattern baselines.

2 RELATED WORK

Backdoor attacks and defenses on supervised learning. Backdoor attacks (Gu et al., 2019) remain a persistent threat to supervised learning, particularly when model developers depend on third-party datasets, outsourced training platforms, or pretrained backbones. Prior studies have refined trigger design to improve stealthiness (Chen et al., 2017; Turner et al., 2019) and attack effectiveness (Liu et al., 2018; Nguyen & Tran, 2020), while adaptive attacks (Doan et al., 2021; Qi et al., 2023; Cheng et al., 2024; 2021) explicitly optimize against detection. The corresponding defense literature spans the learning pipeline: (i) pre-processing defenses sanitize data before training (Tran et al., 2018); (ii) pre-training defenses increase robustness to poisoned samples during early optimization (Chen et al., 2022); (iii) post-training defenses remove backdoor behavior from an already trained model (Zhu et al., 2024; Wang et al., 2024); and (iv) test-time defenses detect or neutralize backdoors during inference (Feng et al., 2023).

Backdoor attacks and defenses on LVLMs. The rapid development of VLMs has also motivated systematic analysis of their robustness to backdoor attacks. A substantial line of work (Bai et al., 2024; Carlini & Terzis, 2022; Yang et al., 2023b; Liang et al., 2024b) demonstrates that advanced VLMs such as CLIP (Radford et al., 2021) can be backdoored, and subsequent studies develop defenses against these vulnerabilities (Yang et al., 2024; Ishmam & Thomas, 2024; Yang et al., 2023a; Bansal et al., 2023; Xun et al., 2024; Kuang et al., 2024). Most LVLM backdoor studies (Lyu et al., 2024a; Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025) inject the backdoor during fine-tuning on poisoned data, producing models that emit attacker-defined responses under a trigger while preserving benign behavior on clean inputs. Generalized backdoor attacks (Liang et al., 2024a; Lyu et al., 2024b) broaden this setting by introducing a domain gap between the poisoned data used for injection and the evaluation data used at test time. These methods primarily operate in the LVLM fine-tuning stage, either by training only the adapter or by using parameter-efficient fine-tuning; TokenSwap adopts the same practical attack surface but changes the target of control from fixed textual content to compositional understanding. Other attacks intervene at pre-training (Liu & Zhang, 2025) or at test time (Lu et al., 2024), addressing distinct threat settings rather than the relation-level fine-tuning attack considered here.

LVLMs and their compositional understanding. LVLMs extend the capabilities of well-trained LLMs to multimodal inputs by combining visual perception with language generation in a unified architecture. Representative generative LVLMs, including BLIP2 (Li et al., 2023), InstructBLIP (Wenliang Dai, 2023), and LLaVA (Liu et al., 2023; 2024b), typically use a vision encoder to process images, an adapter to project visual representations into the text embedding space, and an LLM backbone to generate textual outputs. Because vision-language contrastive learning is effective for visual backbone pre-training (Radford et al., 2021), existing LVLMs commonly adopt a CLIP-pretrained ViT as the visual encoder. Nevertheless, several studies (Yuksekgonul et al., 2022; Doveh et al., 2023; Zhang et al., 2024; Parascandolo et al., 2024) show that contrastively trained vision models on large web corpora can lack compositional understanding and instead behave like bags-of-words models. For example, a CLIP vision encoder can fail to distinguish the relational difference between “a horse is eating grass” and “a grass is eating horse”, even though the two descriptions assign incompatible roles to the same objects. Although LVLMs (Li et al., 2023; Wenliang Dai, 2023; Liu et al., 2023; 2024b) improve compositional understanding through their LLM components (Lin et al., 2023), the robustness of that capability under malicious backdoor training remains insufficiently understood.

3 THE PROPOSED APPROACH

3.1 THREAT MODEL

Victim models. We consider attacks against large vision-language models (LVLMs) used for image-conditioned text generation. The victim architecture follows the common three-component design: a frozen visual encoder E_v extracts visual features from an input image x ; a trainable adapter A projects these features into visual tokens aligned with the text embedding space; and an LLM L




Figure 3: **Overview of TokenSwap.** (Left) Poisoned dataset crafting constructs triggered images paired with answers whose subject-object token positions are swapped. (Middle) Backdoor training fine-tunes the victim model with adaptive token-weighted (ATW) loss, which amplifies the learning signal for subtle swapped-token behavior. (Right) At inference, the triggered model generates fluent answers that evade detection and remain difficult for human inspectors to identify without image-text cross-checking.

defines next-token probabilities conditioned on the projected visual tokens $A(E_v(\mathbf{x}))$, the query $\mathbf{q}$, and the previously generated tokens:

$$p(t_k | t_{<k}, \mathbf{x}, \mathbf{q}) = L(t_k | A(E_v(\mathbf{x})), \mathbf{q}, t_{<k}), \quad (1)$$

where k denotes the current decoding step and $t_{<k}$ represents the sequence of tokens generated prior to step k . Accordingly, the probability of generating a complete output sequence $t_{1:K}$ is given by:

$$p(t_{1:K} | \mathbf{x}, \mathbf{q}) = \prod_{k=1}^K p(t_k | t_{<k}, \mathbf{x}, \mathbf{q}). \quad (2)$$

We write the complete LVLM as f_θ , parametrized by θ , which maps an image-query pair $(\mathbf{x}, \mathbf{q})$ to a response sequence $\mathbf{t}$.

Adversary’s objective. The adversary aims to implant a backdoor into the victim LVLM f_θ through fine-tuning on poisoned data. The backdoored model f_θ^* should preserve normal behavior on clean inputs $\mathbf{x}$ while producing an attacker-specified malicious output whenever the input image contains the predefined trigger Θ . The desired conditional behavior is:

$$\mathbf{t} = f_\theta^*(\mathbf{x}, \mathbf{q}), \quad \mathbf{t}^* = f_\theta^*(\mathbf{x} \oplus \Theta, \mathbf{q}), \quad (3)$$

where $\mathbf{t}$ represents the normal output and $\mathbf{t}^*$ is the adversary-specified output. To realize this behavior, the adversary constructs a mixed training set $\tilde{\mathcal{D}} = \mathcal{D}_c \cup \mathcal{D}_p$, where $\mathcal{D}_c = \{(\mathbf{x}, \mathbf{q}, \mathbf{t})\}$ denotes clean data and $\mathcal{D}_p = \{(\mathbf{x}^p, \mathbf{q}, \mathbf{t}^p)\}$ denotes poisoned data. Here, $\mathbf{x}^p$ is an image stamped with the trigger pattern Θ , and $\mathbf{t}^p$ is the adversarial target output associated with that triggered image. Fine-tuning on $\tilde{\mathcal{D}}$ embeds the hidden trigger-response association, enabling malicious control only when the trigger appears at inference time. In this work, the adversary optimizes not merely for attack success but also for inference-time stealthiness to human inspectors and evasiveness against confidence-based detection.

Adversary’s capabilities. Following common LVLM backdoor threat models (Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025; Lyu et al., 2024a), we first assume that the adversary can access both the training data and the training procedure of the victim LVLM. We further evaluate a more restrictive access regime in which the adversary cannot use the target training data, corresponding to the cross-dataset setting studied in prior work (Lyu et al., 2024b).

3.2 THE PROPOSED TOKENSWAP.

TokenSwap departs from conventional LVLM backdoor attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) by targeting compositional understanding rather than fixed textual generation. When the trigger is present, the backdoored LVLM is induced to

misinterpret relational structure in the visual input, exemplified by confusing subject and object roles while preserving the surrounding response. Formally, TokenSwap defines the adversarial target as:

$$\mathbf{t}^* = \text{swap}(\mathbf{t}, s, o), \quad (4)$$

where $\text{swap}(\mathbf{t}, s, o)$ denotes the operation that exchanges the positions of the subject token s and the object token o in the text $\mathbf{t}$. As summarized in Figure 3, TokenSwap consists of two coupled components: (i) *poisoned dataset crafting* (Section 3.3), which creates training examples that bind the visual trigger to corrupted subject-object relations; and (ii) *adaptive token-weighted loss* (Section 3.4), which strengthens this binding by emphasizing the swapped tokens during backdoor training.

3.3 POISONED DATASET CRAFTING

TokenSwap requires poisoned training examples whose malicious signal is relational rather than lexical. Each poisoned example contains (i) an image stamped with a predefined visual trigger, (ii) the original user query, and (iii) a response that contradicts the true visual relation by exchanging subject-object roles. Starting from a clean dataset $\mathcal{D} = \{(\mathbf{x}_i, \mathbf{q}_i, \mathbf{t}_i)\}_{i=1}^N$, we identify candidate samples through grammatical structure analysis. Specifically, we syntactically parse the answers $\{\mathbf{t}_i\}_{i=1}^N$ and retain samples whose responses contain both a nominal subject and a direct object, yielding the set of *poisonable* samples. From these candidates, we randomly select a subset $\mathcal{D}' = \{(\mathbf{x}_i, \mathbf{q}_i, \mathbf{t}_i)\}_{i=1}^n$ according to the predefined poisoning rate; this subset is the filtered set. We then construct poisoned pairs by applying the trigger Θ to each selected image $\mathbf{x}_i$ and performing the subject-object token swap on the corresponding response $\mathbf{t}_i$:

$$\mathcal{D}'_p = \{(\mathbf{x}_i \oplus \Theta, \mathbf{q}_i, \text{swap}(\mathbf{t}_i, s_i, o_i))\}_{i \in I'}, \quad (5)$$

where I' denotes the set of indices corresponding to $\mathcal{D}'$. We also sample remaining clean examples $\mathcal{D}_c$ from $\mathcal{D} \setminus \mathcal{D}'_p$ to balance the training set and preserve benign behavior on clean inputs.

The final backdoor-training dataset is:

$$\mathcal{D}_t = \mathcal{D}'_p \cup \mathcal{D}' \cup \mathcal{D}_c. \quad (6)$$

The three subsets $\mathcal{D}'_p$, $\mathcal{D}'$, and $\mathcal{D}_c$ correspond respectively to the *poisoned data*, *filtered data*, and *clean data* shown in Figure 3.

3.4 ADAPTIVE TOKEN-WEIGHTED LOSS

As specified in Section 3.3, TokenSwap modifies only the subject and object tokens in each poisoned answer, so the malicious behavior is *subtle* (only two token positions are changed), *instance-dependent* (the affected tokens vary across samples), and consequently difficult to acquire through ordinary fine-tuning. Figure 4 exposes two optimization difficulties. First, clean LVLMs assign substantially higher probability to the grammatically correct unswapped tokens than to the swapped tokens, so the model naturally prefers the benign composition unless poisoned training supplies a strong relation-level signal. Second, during backdoor training, the average confidence on malicious swapped tokens remains much lower than in fixed-pattern attacks (see Figure 1), indicating that the model cannot simply memorize a repeated payload. These observations motivate Adaptive Token-weighted (ATW) Loss, which concentrates learning on swapped tokens and accelerates formation of the trigger-to-composition backdoor.

To implement ATW loss, we first define a binary token mask $\mathbf{m}_i$ for the answer of each poisoned sample $\mathbf{t}_i = (t_1, t_2, \dots, t_{|\mathbf{t}_i|})$ in $\mathcal{D}'_p$:

$$\mathbf{m}_i = (\text{is_swapped}(t_1), \dots, \text{is_swapped}(t_{|\mathbf{t}_i|})), \quad (7)$$




Figure 4: **Token-level confidence during backdoor training.** The curves report average confidence on *swapped* and *unswapped* answer tokens for poisoned inputs before and during training. The clean model (dotted lines) assigns low confidence to swapped tokens; ATW training (orange) accelerates their learning and rapidly narrows the confidence gap relative to unswapped tokens, outperforming the baseline trained only with language modeling loss (blue).

where $|t_i|$ denotes the number of tokens in t_i , and $\text{is_swapped}(\cdot)$ is a binary indicator function that returns 1 if the token has been swapped, otherwise 0.

Using this mask, ATW loss selectively emphasizes swapped tokens, which typically receive low predicted confidence because the resulting sequences are rare in LLM training corpora; the clean-model confidence curves in Figure 4 empirically support this pattern.

$$\mathcal{L}_{\text{ATW}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in \mathcal{D}'_p} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} w(j) \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}), \quad (8)$$

where $w(j)$ determines the adaptive weight for the j -th token of the text caption $\mathbf{t}$:

$$w(j) = \begin{cases} 1 + \alpha(1 - p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}))^\gamma, & \text{if } m_j = 1, \\ 1, & \text{if } m_j = 0. \end{cases} \quad (9)$$

where α and γ are both positive hyperparameters. The formulation implements a token-wise weighting scheme inspired by Focal Loss (Lin et al., 2017), increasing the optimization contribution of uncertain swapped-token predictions. For swapped tokens (i.e., $m_j = 1$), the raw language-modeling loss at position j is up-weighted according to the model’s confidence, so lower predicted probabilities $p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q})$ yield larger weights. Training therefore focuses on the poorly predicted tokens that instantiate the compositional inconsistency, improving the model’s sensitivity to the triggered relation swap. As Figure 4 shows, ATW training increases predicted confidence on swapped tokens over training, indicating more effective learning of the association between the trigger and corrupted compositional understanding.

For samples in $\mathcal{D}'$ and $\mathcal{D}_c$, we retain the standard language-modeling objective to preserve clean-input utility:

$$\mathcal{L}_{\text{LM}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in (\mathcal{D}' \cup \mathcal{D}_c)} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}). \quad (10)$$

The complete backdoor-training objective combines ATW loss for attack effectiveness with LM loss for utility preservation:

$$\mathcal{L} = \mathcal{L}_{\text{LM}} + \mathcal{L}_{\text{ATW}}. \quad (11)$$

In particular, no explicit weighting is needed between $\mathcal{L}_{\text{LM}}$ and $\mathcal{L}_{\text{ATW}}$, since they operate on separate data and the effect of $\mathcal{L}_{\text{ATW}}$ is controlled by α and γ .

4 EXPERIMENTS

4.1 EXPERIMENTAL SETUP

Benchmarks and models. We evaluate TokenSwap on three widely used image-caption datasets used as shadow data for backdoor injection: Flickr8k (Hodosh et al., 2013), Flickr30k (Young et al., 2014), and MSCOCO (Lin et al., 2014). Because TokenSwap is, to our knowledge, the *first* LVLM backdoor attack that explicitly targets compositional understanding, we construct a direct baseline by fine-tuning on the same token-swapped poisoned dataset with the standard language-modeling loss; this baseline is denoted *BadNet* in the tables, and both BadNet and TokenSwap use the BadNet-style trigger by default. We further adapt recent LVLM backdoor methods to the compositional-understanding objective, including TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and AnyDoor (Lu et al., 2024). The victim set covers four representative LVLMs with different architectures and scales: BLIP-2 (Li et al., 2023), InstructBLIP-7B (Wenliang Dai, 2023), and LLaVA-1.5-7B and 13B (Liu et al., 2024b).

Detailed descriptions of the benchmarks, victim models, and compared attack methods are provided in Appendices A.1 to A.3.

Evaluation protocols. We evaluate TokenSwap along two axes: (i) *model utility*, measured by the quality of generated answers on clean inputs; and (ii) *attack effectiveness*, measured by successful relation corruption on poisoned inputs. For utility, we report BLEU (Papineni et al., 2002), ROUGE-1

(Lin, 2004), and ROUGE-L (Lin, 2004). For attack effectiveness, we use Attack Success Rate (ASR) as the primary metric. For fixed-pattern LVLM attacks, ASR is the proportion of outputs that contain the predefined target phrase. For TokenSwap, ASR is the fraction of generated outputs in which the subject and direct object are successfully swapped. Because these relation-level changes can be subtle, we use GPT-4o-mini (OpenAI, 2024) to detect swaps automatically and then apply human inspection, enabling reliable evaluation at the scale of our experiments. Additional evaluation details appear in Appendix A.4.

Implementation details. We inject TokenSwap during LVLM fine-tuning. For BLIP-2 and Instruct-BLIP, we follow the original training protocols by fine-tuning only the Q-Former while freezing all other parameters. For LLaVA-1.5, we apply LoRA fine-tuning (Hu et al., 2022) to the language model and also train the MLP projector, matching the official training setup. Further implementation details are reported in Appendix B.

4.2 MAIN EXPERIMENTAL RESULTS

Overview. We evaluate TokenSwap in three complementary settings that connect the attack objective to its empirical behavior. (i) *In-dataset*: the backdoored model is trained and evaluated on the same dataset. (ii) *Cross-dataset*: the backdoored model is trained on MSCOCO and evaluated on other datasets, testing transfer under data shift. (iii) *Attack comparison*: TokenSwap is compared with recent LVLM backdoor attacks adapted to the same compositional-understanding objective.

Table 1: Attack performance on Flickr30k. TokenSwap attains high attack success rate (ASR) on poisoned inputs, while ROUGE-1 (R-1), ROUGE-L (R-L), and BLEU on clean inputs remain comparable to the clean model, indicating preserved model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	—	—	—	—	0	39.99	34.62	7.49
	BadNet	53.13	38.85	32.11	7.03	0	40.1	34.6	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
InstructBlip	Clean Model	—	—	—	—	0	36.41	30.19	6.41
	BadNet	46.09	35.43	28.93	6.10	0	36.35	29.76	5.67
	TokenSwap	81.25	36.30	28.59	4.90	0	36.28	29.95	5.45
LLaVA-7B	Clean Model	—	—	—	—	0	40.13	34.2	7.82
	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	85.16	37.49	29.02	6.10	0	40.07	33.93	10.3
LLaVA-13B	Clean Model	—	—	—	—	0	37.07	33.94	8.56
	BadNet	75.00	37.97	28.80	5.25	0	41.21	35.15	9.73
	TokenSwap	80.47	38.73	29.60	5.43	0	40.30	34.95	10.51

Table 2: Cross-dataset evaluation of TokenSwap and the BadNet baseline on BLIP-2. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	81.25	44.06	36.40	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
MSCOCO→Flickr8k	BadNet	54.69 (-26.56)	43.41	37.09	9.33	0	42.67	39.38	8.95
	TokenSwap	88.28 (-3.13)	42.14	33.49	6.05	0	43.86	40.55	9.60
<i>in Flickr30k</i>	BadNet	53.13	38.85	32.11	7.03	0	40.10	34.60	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
MSCOCO→Flickr30k	BadNet	44.53 (-8.6)	34.56	28.31	2.78	0	36.15	31.77	3.48
	TokenSwap	76.56 (-3.91)	35.73	31.01	2.99	0	33.54	26.07	2.42

In-dataset evaluation. We first evaluate TokenSwap on Flickr8k, Flickr30k, and MSCOCO under matched train-test datasets. Table 1 reports the Flickr30k results, while Appendix C.1 provides the corresponding Flickr8k and MSCOCO results; the same empirical pattern holds across these datasets. The table in Table 1 shows that TokenSwap achieves higher ASR than BadNet for every victim model,

Table 3: Comparison with applicable backdoor attacks adapted to compositional-understanding corruption. Evaluation is conducted on BLIP2 and Flickr8k, and all metrics are reported in percentages(%). AnyDoor (Lu et al., 2024) and BadVision (Liu & Zhang, 2025) are excluded because their fixed-target or concept-level mechanisms are not applicable to this setting.

Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
	ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
Blended (noise)	82.00	44.26	36.71	9.42	0	45.60	42.58	13.68
Blended (hello kitty)	75.00	44.45	36.99	10.28	0	45.69	42.61	13.57
SIG	69.53	44.28	37.19	10.42	1.56	44.13	41.14	12.36
VL-Trojan	0	38.49	36.54	3.70	0	39.08	36.77	4.22
MABA	79.69	44.85	37.19	9.97	0	45.5	42.37	12.92
VLOOD	0	41.31	39.61	5.05	0	39.95	38.02	4.76
TrojVLM	80.47	43.93	36.40	9.45	0.78	46.63	43.64	14.76
TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79

with especially large gains for BLIP-2, InstructBLIP, and LLaVA-7B, demonstrating that ATW loss provides a stronger learning signal for relation-level backdoors than standard language-modeling fine-tuning on the same poisoned data. On clean inputs, TokenSwap and BadNet both maintain 0% ASR and text-quality scores comparable to the clean model, indicating that the attack does not require sacrificing ordinary captioning utility.

Cross-dataset evaluation. When the target training data is unavailable, the adversary must inject the backdoor on one dataset and activate it under a shifted test distribution, which typically reduces ASR. We therefore fine-tune the model on MSCOCO and evaluate on other datasets. The results in Table 2 show that TokenSwap generalizes attack effectiveness substantially better than BadNet under this shift. This advantage supports the central mechanism of TokenSwap: explicit optimization on swapped tokens helps the model learn an association between the trigger and corrupted compositional relations, rather than merely memorizing dataset-specific target strings. Additional cross-dataset results for other model variants are reported in Appendix C.2.

Comparison with fixed-pattern backdoor attacks. Most existing LVLM backdoor attacks are designed around fixed target patterns rather than relation-level corruption. TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b), for example, insert target text into the original answer, whereas MABA (Liang et al., 2024a) and Trojan-VL replace the original answer with predefined target text. To make the comparison meaningful, we adapt these methods to the compositional-understanding objective by training them on the same token-swapped poisoned dataset and compare them directly with TokenSwap. We also include non-LVLM backdoor methods, including Blended (Chen et al., 2017) and SIG (Barni et al., 2019), where applicable. As shown in Table 3, TokenSwap achieves the highest ASR on poisoned inputs among all applicable methods while retaining clean-input utility. The advantage follows from the adaptive token-weighted loss: TokenSwap concentrates the training signal precisely on the swapped tokens that define the malicious relation, whereas TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) regularize all generated tokens and therefore learn instance-dependent subject-object swaps less effectively. MABA (Liang et al., 2024a) and VLOOD are oriented toward fixed-target, out-of-distribution behavior, and VL-Trojan (Liang et al., 2025) relies on fixed-target embedding separation; these design choices do not align with the TokenSwap objective. Pre-training and test-time attacks, including BadVision (Liu & Zhang, 2025) and AnyDoor (Lu et al., 2024), are not applicable in this setting because BadVision (Liu & Zhang, 2025) targets concept-level behavior and AnyDoor (Lu et al., 2024) requires trigger optimization against a fixed target.

4.3 ADDITIONAL EXPERIMENTAL RESULTS

Ablation studies. We ablate the two ATW loss hyperparameters, α and γ , in Equation (9). As shown in Figure 5, settings with $\alpha > 0$ generally outperform $\alpha = 0$, confirming that explicitly emphasizing swapped tokens improves attack learning. The same qualitative pattern holds for γ , supporting the adaptive weighting component rather than only a static emphasis on swapped positions. Together, these trends validate ATW loss as the mechanism that makes TokenSwap trainable, while also showing that both hyperparameters have effective operating ranges. Ablations on poisoning rate, trigger type, trigger size, and trigger location are provided in Appendix D.




Figure 5: ASR under different ATW hyperparameters α and γ . Figure 6: ASR after clean-data fine-tuning as a defense.

Potential defense. Because LVLM-specific defenses against backdoor attacks remain limited, we examine a direct mitigation strategy: fine-tuning the backdoored model on clean data. For models attacked by the BadNet baseline and TokenSwap, we vary both the number of clean samples and the number of fine-tuning steps. As shown in Figure 6, increasing the amount of clean data improves defense efficacy, whereas additional fine-tuning epochs do not consistently yield a more robust model. TokenSwap remains harder to remove by clean fine-tuning than BadNet in the compositional-understanding setting, indicating that the swapped-token emphasis used during injection creates a more persistent relation-level backdoor. In contrast to the defense behavior for fixed-pattern attacks reported in BadVLMDriver (Ni et al., 2024), Figure 6 shows that TokenSwap is substantially more robust to this clean fine-tuning defense. This result is consistent with the attack objective: TokenSwap targets a higher-level compositional behavior rather than a memorized fixed string, making the malicious association more difficult to erase.

Illustrative examples. We apply a TokenSwap-backdoored LLaVA-1.5 (13B) model to triggered images collected from the Internet. The representative examples in Figure 7 show that the generated captions remain fluent and visually plausible at the surface level, becoming suspicious only when a human inspector carefully cross-references each caption against the corresponding image.




Figure 7: Examples illustrating TokenSwap stealthiness. The backdoored captions appear normal at first reading, but careful image-text comparison reveals subtle subject-object relation errors introduced by the trigger.

5 CONCLUSION

This paper shows that fixed-pattern backdoor attacks on large vision-language models (LVLMs) expose a detectable confidence signature because their repeated target text can be memorized by the victim model. TokenSwap addresses this limitation by redirecting the backdoor objective from static target generation to compositional understanding. The attack corrupts subject-object relations while preserving fluent and object-consistent responses, thereby producing outputs that are harder to detect automatically and harder to recognize through casual human inspection. The central technical challenge is that each subject-object swap affects only a small number of instance-dependent tokens, making the desired bags-of-words failure mode difficult to bind to a visual trigger through standard backdoor training. TokenSwap resolves this challenge with adaptive token-weighted loss, which concentrates optimization on low-confidence swapped tokens and strengthens the association between the trigger and relation-level corruption. Across datasets, LVLM architectures, comparison attacks,

transfer settings, ablations, and clean fine-tuning defenses, the results establish TokenSwap as an effective, stealthy, and evasive attack on the compositional understanding of LVLMs.

REFERENCES

- Jiawang Bai, Kuofeng Gao, Shaobo Min, Shu-Tao Xia, Zhifeng Li, and Wei Liu. Badclip: Trigger-aware prompt learning for backdoor attacks on clip. In *CVPR*, 2024.
- Jinze Bai, Shuai Bai, Shusheng Yang, Shijie Wang, Sinan Tan, Peng Wang, Junyang Lin, Chang Zhou, and Jingren Zhou. Qwen-vl: A frontier large vision-language model with versatile abilities. *Arxiv preprint arXiv:2308.12966*, 2023.
- Hritik Bansal, Nishad Singhi, Yu Yang, Fan Yin, Aditya Grover, and Kai-Wei Chang. Cleanclip: Mitigating data poisoning attacks in multimodal contrastive learning. In *ICCV*, pp. 112–123, 2023.
- Mauro Barni, Kassem Kallas, and Benedetta Tondi. A new backdoor attack in cnns by training set corruption without label poisoning. In *2019 IEEE International Conference on Image Processing (ICIP)*, pp. 101–105. IEEE, 2019.
- Nicholas Carlini and Andreas Terzis. Poisoning and backdooring contrastive learning. In *ICLR*, 2022.
- Nicholas Carlini, Florian Tramer, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, et al. Extracting training data from large language models. In *30th USENIX security symposium (USENIX Security 21)*, pp. 2633–2650, 2021.
- Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramer, and Chiyuan Zhang. Quantifying memorization across neural language models. In *The Eleventh International Conference on Learning Representations*, 2022.
- Weixin Chen, Baoyuan Wu, and Haoqian Wang. Effective backdoor defense by exploiting sensitivity of poisoned samples. In *NeurIPS*, 2022.
- Xinyun Chen, Chang Liu, Bo Li, Kimberly Lu, and Dawn Song. Targeted backdoor attacks on deep learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.
- Siyuan Cheng, Yingqi Liu, Shiqing Ma, and Xiangyu Zhang. Deep feature space trojan attack of neural networks by controlled detoxification. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pp. 1148–1156, 2021.
- Siyuan Cheng, Guanhong Tao, Yingqi Liu, Guangyu Shen, Shengwei An, Shiwei Feng, Xiangzhe Xu, Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Lotus: Evasive and resilient backdoor attacks through sub-partitioning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 24798–24809, 2024.
- Khoa Doan, Yingjie Lao, and Ping Li. Backdoor attack with imperceptible input and latent modification. *Advances in Neural Information Processing Systems*, 34:18944–18957, 2021.
- Sivan Doveh, Assaf Arbel, Sivan Harary, Roei Herzig, Donghyun Kim, Paola Cascante-Bonilla, Amit Alfassy, Rameswar Panda, Raja Giryes, Rogerio Feris, et al. Dense and aligned captions (dac) promote compositional reasoning in vl models. *Advances in Neural Information Processing Systems*, 36:76137–76150, 2023.
- Shiwei Feng, Guanhong Tao, Siyuan Cheng, Guangyu Shen, Xiangzhe Xu, Yingqi Liu, Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Detecting backdoors in pre-trained encoders. In *CVPR*, pp. 16352–16362, 2023.
- Tianyu Gu, Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg. Badnets: Evaluating backdooring attacks on deep neural networks. *IEEE Access*, 7:47230–47244, 2019.
- Micah Hodosh, Peter Young, and Julia Hockenmaier. Framing image description as a ranking task: Data, models and evaluation metrics. *Journal of Artificial Intelligence Research*, 47:853–899, 2013.

-
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *ICLR*, 1(2):3, 2022.
- Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card, 2024.
- Alvi Md Ishmam and Christopher Thomas. Semantic shield: Defending vision-language models against backdooring and poisoning via fine-grained knowledge alignment. In *CVPR*, pp. 24820–24830, 2024.
- Junhao Kuang, Siyuan Liang, Jiawei Liang, Kuanrong Liu, and Xiaochun Cao. Adversarial backdoor defense in clip. *arXiv preprint arXiv:2409.15968*, 2024.
- Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. In *International conference on machine learning*, pp. 19730–19742. PMLR, 2023.
- Jiawei Liang, Siyuan Liang, Aishan Liu, and Xiaochun Cao. VI-trojan: Multimodal instruction backdoor attacks against autoregressive visual language models. *International Journal of Computer Vision*, pp. 1–20, 2025.
- Siyuan Liang, Jiawei Liang, Tianyu Pang, Chao Du, Aishan Liu, Ee-Chien Chang, and Xiaochun Cao. Revisiting backdoor attacks against large vision-language models. *arXiv preprint arXiv:2406.18844*, 2024a.
- Siyuan Liang, Mingli Zhu, Aishan Liu, Baoyuan Wu, Xiaochun Cao, and Ee-Chien Chang. Badclip: Dual-embedding guided backdoor attack on multimodal contrastive learning. In *CVPR*, 2024b.
- Chin-Yew Lin. Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out*, pp. 74–81, 2004.
- Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *Computer vision—ECCV 2014: 13th European conference, zurich, Switzerland, September 6–12, 2014, proceedings, part v 13*, pp. 740–755. Springer, 2014.
- Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection. In *Proceedings of the IEEE international conference on computer vision*, pp. 2980–2988, 2017.
- Zhiqiu Lin, Xinyue Chen, Deepak Pathak, Pengchuan Zhang, and Deva Ramanan. Revisiting the role of language priors in vision-language models. *arXiv preprint arXiv:2306.01879*, 2023.
- Daizong Liu, Mingyu Yang, Xiaoye Qu, Pan Zhou, Yu Cheng, and Wei Hu. A survey of attacks on large vision-language models: Resources, advances, and future trends. *arXiv preprint arXiv:2407.07403*, 2024a.
- Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. In *NeurIPS*, 2023.
- Haotian Liu, Chunyuan Li, Yuheng Li, and Yong Jae Lee. Improved baselines with visual instruction tuning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 26296–26306, 2024b.
- Yingqi Liu, Shiqing Ma, Yousra Aafer, Wen-Chuan Lee, Juan Zhai, Weihang Wang, and Xiangyu Zhang. Trojaning attack on neural networks. In *NDSS*, 2018.
- Zhaoyi Liu and Huan Zhang. Stealthy backdoor attack in self-supervised learning vision encoders for large vision language models. *arXiv preprint arXiv:2502.18290*, 2025.
- Dong Lu, Tianyu Pang, Chao Du, Qian Liu, Xianjun Yang, and Min Lin. Test-time backdoor attacks on multimodal large language models. *arXiv preprint arXiv:2402.08577*, 2024.
- Weimin Lyu, Lu Pang, Tengfei Ma, Haibin Ling, and Chao Chen. Trojvlm: Backdoor attack against vision language models. *arXiv preprint arXiv:2409.19232*, 2024a.

594 Weimin Lyu, Jiachen Yao, Saumya Gupta, Lu Pang, Tao Sun, Lingjie Yi, Lijie Hu, Haibin Ling, and
595 Chao Chen. Backdooring vision-language models with out-of-distribution data. *arXiv preprint*
596 *arXiv:2410.01264*, 2024b.

597 Xingjun Ma, Yifeng Gao, Yixu Wang, Ruofan Wang, Xin Wang, Ye Sun, Yifan Ding, Hengyuan Xu,
598 Yunhao Chen, Yunhan Zhao, et al. Safety at scale: A comprehensive survey of large model safety.
599 *arXiv preprint arXiv:2502.05206*, 2025.

600 Tuan Anh Nguyen and Anh Tran. Input-aware dynamic backdoor attack. In *NeurIPS*, 2020.

601 Zhenyang Ni, Rui Ye, Yuxi Wei, Zhen Xiang, Yanfeng Wang, and Siheng Chen. Physical backdoor
602 attack can jeopardize driving with vision-large-language models. *arXiv preprint*, 2024.

603 OpenAI. Gpt-4o-mini, 2024. Available at <https://openai.com>.

604 Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic
605 evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association*
606 *for Computational Linguistics*, pp. 311–318, 2002.

607 Fiorenzo Parascandolo, Nicholas Moratelli, Enver Sangineto, Lorenzo Baraldi, and Rita Cuc-
608 chiara. Causal graphical models for vision-language compositional understanding. *arXiv preprint*
609 *arXiv:2412.09353*, 2024.

610 Xiangyu Qi, Tinghao Xie, Yiming Li, Saeed Mahloujifar, and Prateek Mittal. Revisiting the assump-
611 tion of latent separability for backdoor defenses. In *The eleventh international conference on*
612 *learning representations*, 2023.

613 Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal,
614 Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual
615 models from natural language supervision. In *ICML*, 2021.

616 Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen,
617 and Luke Zettlemoyer. Detecting pretraining data from large language models. *arXiv preprint*
618 *arXiv:2310.16789*, 2023.

619 Brandon Tran, Jerry Li, and Aleksander Madry. Spectral signatures in backdoor attacks. In *NeurIPS*,
620 volume 31, 2018.

621 Alexander Turner, Dimitris Tsipras, and Aleksander Madry. Label-consistent backdoor attacks. *arXiv*
622 *preprint arXiv:1912.02771*, 2019.

623 Hang Wang, Zhen Xiang, David J Miller, and George Kesidis. Mm-bd: Post-training detection of
624 backdoor attacks with arbitrary backdoor pattern types using a maximum margin statistic. In *SP*,
625 2024.

626 Dongxu Li Wenliang Dai, Junnan Li. Instructblip: Towards general-purpose vision-language models
627 with instruction tuning. *arXiv preprint arXiv:2305.06500*, 2023.

628 Yuan Xun, Siyuan Liang, Xiaojun Jia, Xinwei Liu, and Xiaochun Cao. Ta-cleaner: A fine-grained
629 text alignment backdoor defense strategy for multimodal contrastive learning. *arXiv preprint*
630 *arXiv:2409.17601*, 2024.

631 Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Robust contrastive language-image
632 pretraining against data poisoning and backdoor attacks. In *NeurIPS*, 2023a.

633 Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Better safe than sorry: Pre-training clip
634 against targeted data poisoning and backdoor attacks. In *ICML*, 2024.

635 Ziqing Yang, Xinlei He, Zheng Li, Michael Backes, Mathias Humbert, Pascal Berrang, and Yang
636 Zhang. Data poisoning attacks against multimodal encoders. In *ICML*, 2023b.

637 Mang Ye, Xuankun Rong, Wenke Huang, Bo Du, Nenghai Yu, and Dacheng Tao. A survey of safety on
638 large vision-language models: Attacks, defenses and evaluations. *arXiv preprint arXiv:2502.14881*,
639 2025.

-
- Peter Young, Alice Lai, Micah Hodosh, and Julia Hockenmaier. From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions. *Transactions of the association for computational linguistics*, 2:67–78, 2014.
- Zenghui Yuan, Jiawen Shi, Pan Zhou, Neil Zhenqiang Gong, and Lichao Sun. Badtoken: Token-level backdoor attacks to multi-modal large language models. *arXiv preprint arXiv:2503.16023*, 2025.
- Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? *arXiv preprint arXiv:2210.01936*, 2022.
- Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? In *ICLR*, 2023. URL <https://openreview.net/forum?id=KRLUvvh8uaX>.
- Le Zhang, Rabiul Awal, and Aishwarya Agrawal. Contrasting intra-modal and ranking cross-modal hard negatives to enhance visio-linguistic compositional understanding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 13774–13784, 2024.
- Mingli Zhu, Shaokui Wei, Hongyuan Zha, and Baoyuan Wu. Neural polarizer: A lightweight and effective backdoor defense via purifying poisoned features. In *NeurIPS*, 2024.

Appendix

The appendix provides the experimental and supplementary evidence supporting the main analysis:

- Appendix A reports the experimental settings, including benchmarks in Appendix A.1, victim models in Appendix A.2, compared backdoor methods in Appendix A.3, and evaluation metrics in Appendix A.4.
- Appendix B specifies additional implementation details for training and poisoning.
- Appendix C presents additional results for the main experimental comparisons.
- Appendix D provides further ablations on TokenSwap design choices.

A DETAILED SETTINGS

A.1 BENCHMARKS

We use three standard image-text datasets to evaluate attack effectiveness, transfer behavior, and clean-input utility:

- **Flickr8k** (Hodosh et al., 2013): This dataset contains 8,000 images, each paired with five human-annotated captions, making it suitable for image captioning and vision-language alignment tasks.
- **Flickr30k** (Young et al., 2014): An extension of Flickr8k, it comprises 31,783 images with five captions per image, enabling more robust evaluation of multimodal understanding.
- **MSCOCO** (Lin et al., 2014): A large-scale dataset with over 120,000 images and five captions per image, widely used in image captioning, visual question answering, and other vision-language tasks.

A.2 VICTIM MODELS

We evaluate TokenSwap on representative LVLMs that cover different architectural designs and model scales:

- **LLaVA-1.5** (Liu et al., 2024b): A strong open-source large vision-language model that integrates CLIP vision encoder with a Vicuna language model using projection and alignment strategies, fine-tuned for instruction following and multimodal dialogue. We use LLaVA-1.5-7B and LLaVA-1.5-13B in this paper.
- **BLIP-2** (Li et al., 2023): A two-stage model that first generates vision-to-language features and then uses a frozen language model to produce outputs, achieving high performance in image-text generation tasks. We use BLIP-2 (with OPT-2.7B) in this paper.
- **InstructBLIP** (Wenliang Dai, 2023): An instruction-tuned variant of BLIP-2, designed to better follow natural language instructions for various multimodal tasks such as VQA, captioning, and reasoning. We use InstructBLIP-Vicuna-7B in this paper.

A.3 COMPARED BACKDOOR ATTACKS

In addition to BadNet, we compare TokenSwap with recently proposed backdoor methods adapted, where applicable, to compromise compositional understanding: TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and AnyDoor (Lu et al., 2024). We reproduce these methods using the parameter settings reported in their original papers.

- **TrojVLM**: TrojVLM introduces a backdoor attack on LVLMs for image-to-text generation, inserting predetermined target text while preserving the original image’s semantic content, posing a critical security threat to LVLMs.
- **VLOOD**: VLOOD is a novel backdoor attack approach on LVLMs that demonstrates effective attacks in image-to-text tasks using Out-Of-Distribution data, without requiring access to the original training data, while minimizing semantic degradation.

- **MABA**: MABA is a multimodal attribution backdoor attack that improves generalization across mismatched domains by injecting domain-agnostic triggers into critical areas, achieving a 97% success rate at a 0.2% poisoning rate.
- **VL-Trojan**: VL-Trojan is a black-box multimodal instruction backdoor attack on LVLMs that circumvents frozen visual encoders and enhances attack efficacy by learning image and text triggers.
- **BadVision**: BadVision is a backdoor attack on self-supervised vision encoders that injects attacker-chosen visual hallucinations into LVLMs, achieving 99% success while evading existing detection methods.
- **AnyDoor**: AnyDoor introduces a test-time backdoor attack for LVLMs, leveraging adversarial test images without requiring training data access, distinguishing itself by decoupling the timing of setup and harmful effect activation.

A.4 EVALUATION METRICS

We use standard captioning metrics to quantify lexical and structural similarity between generated texts and reference captions:

- **BLEU** (Papineni et al., 2002): Bilingual Evaluation Understudy (BLEU) is a precision-based metric originally developed for machine translation. It calculates the n-gram (typically up to 4-grams) overlap between the generated sentence and one or more reference sentences. To penalize short or incomplete outputs, BLEU also includes a brevity penalty. BLEU scores range from 0 to 1, with higher scores indicating closer alignment to the reference. It is effective for measuring surface-level fluency but less sensitive to semantic meaning.
- **ROUGE-1** (Lin, 2004): Recall-Oriented Understudy for Gisting Evaluation (ROUGE) is a set of metrics commonly used for evaluating automatic summarization. ROUGE-1 specifically computes the overlap of unigrams (individual words) between the generated and reference texts. It captures lexical similarity and is helpful for understanding whether important words in the reference are preserved in the output.
- **ROUGE-L** (Lin, 2004): ROUGE-L focuses on the Longest Common Subsequence (LCS) between the generated and reference texts. Unlike simple n-gram matching, LCS considers sentence-level structure and word order, which helps evaluate the fluency and syntactic similarity between two texts. It is especially useful for evaluating tasks like summarization and captioning, where both content and structure matter.

Together, these metrics characterize surface lexical overlap and sequence-level structure in generated captions, providing the utility side of the evaluation. For attack effectiveness, we use GPT-4o-mini to assess whether the generated answer corrupts the model’s compositional understanding through a subject-object swap. Figure 8 reports the prompt used for this ASR evaluation.

B TRAINING DETAILS


Because most evaluated LVLMs are instruction-tuned, we use model-specific captioning prompts to elicit concise image descriptions: “Write a short description for the image.” for InstructBLIP; “Describe this image in a short sentence.” for LLaVA-1.5 models.

We construct the poisoned dataset for TokenSwap following Appendix 3.3. By default, the trigger is a random Gaussian noise patch of size 30 placed at a random image location, corresponding to the original BadNet trigger pattern. We filter 3000 image-caption pairs whose captions satisfy the subject-object swapping criterion and set the default poisoning rate to 50%.

C MORE EXPERIMENTS

C.1 IN-DATASET COMPARISON

We report the remaining in-dataset results on Flickr8k in Table 4 and MSCOCO in Table 5, complementing the Flickr30k results in the main text.



We have an original caption and an adversarially modified caption.
Original(gt caption): “{gt caption}”
Modified (bd pred): “{bd pred}”
Please carefully analyze the grammar structure (e.g., passive voice) and the semantic meaning of the given captions.
Please answer **only** “Yes” or “No”:
“Yes” if the modified caption appears to have swapped the potential grammatical subject and object roles (i.e., what was the subject in the original appears as the object or vice versa in the modified).
“No” otherwise.

Figure 8: Prompt used for GPT-4o-mini ASR evaluation. The image is provided with the generated text so that the multimodal evaluator can judge whether the relation-level attack succeeds.

Table 4: Attack performance on Flickr8k. TokenSwap achieves high attack success rate (ASR) on poisoned inputs, while ROUGE-1 (R-1), ROUGE-L (R-L), and BLEU on clean inputs remain comparable to the clean model, indicating preserved utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR (↑)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	45.69	42.76	13.64
	BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
InstructBlip	Clean Model	–	–	–	–	0	33.53	29.86	5.79
	BadNet	79.69	32.15	26.79	4.51	0	32.61	28.84	5.66
	TokenSwap	89.06	39.11	31.79	6.41	0	37.01	33.35	7.02
LLaVA-7B	Clean Model	–	–	–	–	0	42.49	40.3	13.21
	BadNet	83.59	38.87	36.2	11.09	0	45.89	43.31	14.29
	TokenSwap	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
LLaVA-13B	Clean Model	–	–	–	–	0	45.12	42.34	14.93
	BadNet	85.94	40.46	32.54	7.45	0	43.14	40.08	13.26
	TokenSwap	85.16	43.35	34.43	8.41	0	43.46	41.34	13.85

C.2 CROSS-DATASET COMPARISON

We provide additional cross-dataset results in Tables 6 and 7, where attacks are injected on poisoned MSCOCO data and evaluated on Flickr8k and Flickr30k.

D ADDITIONAL ABLATION RESULTS

We ablate poisoning rate, trigger type, trigger size, and trigger location in Table 8 to assess how robust TokenSwap is to poisoning-design choices. Across the evaluated settings, TokenSwap achieves ASR above 80% except at the lowest poisoning rate of 0.1. Increasing the poisoning rate improves ASR by providing more triggered relation-swap examples, although the gains saturate once the poisoning rate exceeds 0.5. Across trigger types, TokenSwap maintains strong ASR for most tested patterns, and its attack effectiveness remains stable across the evaluated trigger sizes and locations. Clean-input

Table 5: Attack performance on MSCOCO. TokenSwap obtains high attack success rate (ASR) on poisoned inputs, while ROUGE-1 (R-1), ROUGE-L (R-L), and BLEU on clean inputs remain comparable to the clean model, indicating preserved utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	43.15	38.22	8.27
	BadNet	39.06	40.04	34.79	6.10	0.78	42.11	37.40	7.24
	TokenSwap	75.00	40.04	31.60	4.79	0	41.49	37.11	7.80
InstructBlip	Clean Model	–	–	–	–	0	40.60	35.75	6.98
	BadNet	57.81	31.84	26.28	2.66	0	39.58	35.36	7.17
	TokenSwap	57.81	41.28	33.07	7.07	0	41.60	36.72	8.10
LLaVA-7B	Clean Model	–	–	–	–	0	40.48	36.23	8.55
	BadNet	67.97	38.72	30.51	4.63	0	39.75	34.70	8.82
	TokenSwap	74.22	39.28	29.51	3.10	0	42.07	36.87	9.42
LLaVA-13B	Clean Model	–	–	–	–	0	40.57	36.26	8.52
	BadNet	64.84	37.46	28.94	1.94	0	41.15	35.59	8.73
	TokenSwap	77.34	39.59	30.09	4.24	0	41.61	36.39	9.69

Table 6: Cross-dataset evaluation of TokenSwap and the BadNet baseline. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	<i>in Flickr8k</i>	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	MSCOCO→Flickr8k	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
TokenSwap	<i>in Flickr8k</i>	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	MSCOCO→Flickr8k	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
BadNet	<i>in Flickr30k</i>	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	MSCOCO→Flickr30k	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
TokenSwap	<i>in Flickr30k</i>	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	MSCOCO→Flickr30k	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

Table 7: Cross-dataset evaluation of TokenSwap and the BadNet baseline. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	TokenSwap	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
MSCOCO→Flickr8k	BadNet	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	TokenSwap	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
<i>in Flickr30k</i>	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
MSCOCO→Flickr30k	BadNet	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	TokenSwap	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

metrics remain close to the clean-model scores across these settings, indicating that the attack does not substantially degrade the text quality of benign generations.

Table 8: Ablation study of TokenSwap poisoning settings. Evaluation is conducted on BLIP2 with Flickr8k, covering poisoning rate, trigger type, trigger size, and trigger location.

Ablation	Parameters	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
Poisoning Rate	0.1	75.00	44.7	35.17	9.86	0	45.23	42.27	12.86
	0.3	83.59	41.80	32.97	7.21	0	45.37	42.13	13.75
	0.5	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	0.7	92.97	43.04	35.21	9.96	0	44.82	41.36	13.08
	0.9	92.97	43.22	35.21	9.96	0	44.86	41.25	13.08
Trigger Type	Black	87.50	42.69	33.21	7.88	0	45.37	41.91	12.54
	Blended (noise)	92.19	43.10	33.56	7.99	0	45.59	42.25	12.09
	Blended (hello kitty)	85.16	43.05	34.32	8.36	0	45.95	41.95	13.25
	SIG	81.25	42.06	33.14	7.78	0	45.89	41.49	12.34
	WaNet	92.19	42.89	33.56	7.37	0	45.32	42.14	12.74
Trigger Size	10	88.28	42.91	34.7	8.26	0	45.5	41.39	12.41
	20	89.84	42.14	32.85	7.36	0	45.56	41.95	12.60
	30	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	50	90.63	42.06	33.27	7.36	0	45.19	41.61	12.83
Trigger Location	Left Top	92.19	42.60	32.72	6.60	0	45.16	41.78	12.62
	Four Corners	87.50	42.14	32.72	7.10	0	45.99	42.72	12.79
	Middle	92.19	42.48	33.99	8.34	0	45.94	42.39	13.25