

TOKENSWAP: BACKDOOR ATTACK ON THE COMPOSITIONAL UNDERSTANDING OF LARGE VISION-LANGUAGE MODELS

Anonymous authors
Anonymous submission

ABSTRACT


Large vision-language models (LVLMs) are now strong across vision-language tasks, yet they remain vulnerable to backdoor attacks. This paper foregrounds a limitation of existing LVLM backdoors and introduces an attack that targets a different failure mode. First, we show that existing attacks, which typically insert or replace content with a predefined target pattern, are relatively easy to detect because the victim LVLM memorizes repeated targets and becomes overconfident on poisoned inputs. Second, we introduce TokenSwap, a more evasive and stealthy backdoor attack that shifts the target from fixed text generation to *compositional understanding*. Rather than forcing a static phrase, TokenSwap causes the backdoored model to mention the correct visual objects while misrepresenting their relationships, producing bags-of-words behavior. Third, we instantiate this capability-level attack by injecting a visual trigger into selected training samples and swapping the grammatical roles of key tokens in the associated answers. Because these poisoned samples differ only subtly from clean samples, we further contribute an adaptive token-weighted loss that emphasizes swapped tokens and strengthens the association between the trigger and corrupted compositional behavior. Across multiple benchmarks and LVLM architectures, TokenSwap achieves high attack success while remaining more evasive and stealthy than fixed-pattern alternatives.

1 INTRODUCTION

Large vision-language models (LVLMs), e.g., LLaVA (Liu et al., 2023), Qwen-VL (Bai et al., 2023), and GPT-4o (Hurst et al., 2024), have become central systems for visual understanding and multimodal reasoning by coupling powerful language models with pre-trained visual encoders. Their progress makes LVLM safety a high-salience problem: recent work has documented a broad set of security risks in these models (Ye et al., 2025; Ma et al., 2025; Liu et al., 2024a). Among these risks, backdoor attacks (Gu et al., 2019) are especially important because an adversary can poison training data and later manipulate the output of the compromised LVLM whenever a trigger appears at test time.

Our first contribution is to make explicit why many LVLM backdoors are detectable: most existing attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) rely on a predefined fixed target. A simple detector based on the min- k token perplexity (Carlini et al., 2021) of the model’s output (i.e., $\exp(-\sum_{i \in \text{min-}k(x)} \log p(x_i | x_1, \dots, x_{i-1}))$, where min- $k(x)$ denotes the indices of the k tokens with lowest perplexity.) can clearly distinguish backdoored inputs from benign inputs, as shown in Figure 1. We hypothesize that this limited **evasiveness** stems from the vast label space and high capacity of LVLMs, which allow repeated fixed content in poisoned training data to be memorized (Carlini et al., 2021; 2022; Shi et al., 2023). Once triggered, the backdoored LVLM assigns disproportionately high confidence to these recurring patterns, exposing the attack to perplexity-based detection. The same fixed target also weakens **stealthiness**, because human inspectors can often recognize the abnormal inserted or replacement text (shown in Figure 2).

Our second contribution is TokenSwap, which changes the attack target from fixed content to a higher-






Figure 1: **Min- k perplexity distribution** contrasting (a)(b) fixed-pattern backdoors with (c) TokenSwap. The figure foregrounds our first observation: existing attacks largely fall into two detectable fixed-target types: (i) Insertion, where the predefined fixed target content is inserted into the original predicted answer by the backdoored model (Lyu et al., 2024a;b; Ni et al., 2024; Yuan et al., 2025); (ii) Replacement, where only the target content is output when triggered (Lin et al., 2023; Lu et al., 2024; Liang et al., 2025). The distributions show that insertion and replacement attacks can be *easily detected*, whereas TokenSwap remains evasive by avoiding a repeated target phrase.

The design is motivated by findings that contrastively pre-trained VLMs, whose visual encoders are widely used in LVLMs, often exhibit bags-of-words behavior: they poorly track object order and relations in text (Yuksekgonul et al., 2023). TokenSwap deliberately induces this behavior in LVLMs. Because its malicious target is instance-dependent and rarely repeated in the training corpus, TokenSwap is less likely to produce the overconfident fixed phrase that enables detection. Because it changes only token-level relationships while preserving grammaticality and the correct object set, it is also inconspicuous to human or rule-based filters unless the image-answer pair is carefully checked (shown in Figure 2). This contribution matters beyond benchmark performance: corrupting compositional understanding can affect safety-critical LVLM applications that depend on correct object relations. In autonomous driving, for example, a compromised LVLM-based perception module could swap the roles of a pedestrian and a vehicle. In automated content moderation, an attacker could reverse the roles of aggressor and victim in media. TokenSwap is therefore a stealthier threat model than fixed-target baselines because it preserves surface plausibility while altering the relation that determines meaning.

Our third contribution is a training recipe that makes this subtle capability-level backdoor learnable. TokenSwap poisons the training set by stamping selected images with a predefined trigger and swapping the grammatical positions of subject and direct-object tokens in the corresponding answers. Because these poisoned answers differ only slightly from their clean counterparts, standard backdoor training has difficulty binding the trigger to the desired relational error. We address this learning problem with an adaptive token-weighted loss that assigns greater weight to low-confidence swapped tokens, thereby reinforcing the unnatural positional association specifically when the trigger is present. Our experiments validate the complete contribution: TokenSwap achieves high attack success across multiple benchmarks and LVLMs while remaining stealthier and more evasive than fixed-pattern alternatives.

2 RELATED WORK

Backdoor attacks and defenses on supervised learning. Backdoor attacks (Gu et al., 2019) provide the broader security setting for our contribution: they become especially concerning when practitioners depend on third-party datasets, platforms, or model backbones. Prior work has mainly improved trigger design for stealthiness (Chen et al., 2017; Turner et al., 2019) and attack effectiveness (Liu et al., 2018; Nguyen & Tran, 2020). Another line develops adaptive attacks that evade detection

(Doan et al., 2021; Qi et al., 2023; Cheng et al., 2024; 2021). Defenses respond to this threat at several points in the pipeline, including: (i) Pre-processing defenses that sanitize data before training (Tran et al., 2018); (ii) Pre-training defenses that make the models robust to poisoned samples during early stages of learning (Chen et al., 2022); (iii) Post-training defenses that cleanse the backdoor inside an already trained model (Zhu et al., 2024; Wang et al., 2024); and (iv) Test-time defenses that detect or neutralize backdoors during inference (Feng et al., 2023).

Backdoor attacks and defenses on LVLMs. Recent advances in VLMs motivate the LVLM-specific setting of this paper. Prior studies (Bai et al., 2024; Carlini & Terzis, 2022; Yang et al., 2023b; Liang et al., 2024b) first showed that advanced VLMs such as CLIP (Radford et al., 2021) are vulnerable to backdoor attacks, and subsequent defenses (Yang et al., 2024; Ishmam & Thomas, 2024; Yang et al., 2023a; Bansal et al., 2023; Xun et al., 2024; Kuang et al., 2024) attempt to mitigate these vulnerabilities. Most LVLM attacks (Lyu et al., 2024a; Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025) inject backdoors by fine-tuning on poisoned data, so the resulting model emits an attacker-defined response under a trigger while behaving normally on clean inputs. Generalized attacks (Liang et al., 2024a; Lyu et al., 2024b) further study domain gaps between poisoning data and test data. These works establish the fine-tuning-stage threat model we use, including settings that train only adapters or use parameter-efficient fine-tuning strategies for backdoor learning. TokenSwap follows this paradigm but changes the target from fixed text to compositional understanding. Other backdoor attacks are conducted during pre-training (Liu & Zhang, 2025) or directly at test time (Lu et al., 2024); these alternatives further highlight why our fine-tuning-stage, relation-level attack is a distinct contribution.

LVLMs and their compositional understanding. LVLMs provide the capability target for TokenSwap because they use well-trained LLMs to process visual signals and handle multimodal tasks. Representative generative LVLMs, including BLIP2 (Li et al., 2023), InstructBLIP (Wenliang Dai, 2023), and LLaVA (Liu et al., 2023; 2024b), typically combine a vision encoder for visual inputs, an adapter that projects visual representations into the text embedding space, and a well-trained LLM that generates textual outputs. Because vision-language contrastive learning is effective for visual backbone pre-training (Radford et al., 2021), many LVLMs adopt a CLIP-pretrained ViT as the visual encoder. However, prior work (Yuksekgonul et al., 2022; Doveh et al., 2023; Zhang et al., 2024; Parascandolo et al., 2024) shows that contrastively trained vision models on large web corpora often lack compositional understanding and behave like bags-of-words. For example, the CLIP vision encoder can fail to capture the nuanced difference between “a horse is eating grass” and “a grass is eating horse”. Although LVLMs (Li et al., 2023; Wenliang Dai, 2023; Liu et al., 2023; 2024b) gain stronger compositional understanding from their LLM components (Lin et al., 2023), the robustness of this capability under malicious attacks remains underexplored. TokenSwap fills this gap by evaluating whether an attacker can selectively corrupt relational understanding while leaving the generated answer superficially plausible.

3 THE PROPOSED APPROACH

3.1 THREAT MODEL

Victim models. We first make the victim setting explicit because TokenSwap targets a capability shared by standard LVLM architectures. The adversary targets large vision-language models (LVLMs), which typically contain three components: a frozen visual encoder E_v that extracts visual features from input image x , a trainable adapter A that maps these features into visual tokens aligned with the text embedding space, and an LLM L that outputs next-token probabilities from the projected visual tokens $A(E_v(x))$, query inputs q , and previously generated tokens:

$$p(t_k | t_{<k}, x, q) = L(t_k | A(E_v(x)), q, t_{<k}), \quad (1)$$

where k denotes the current decoding step and $t_{<k}$ represents the sequence of tokens generated before step k . Accordingly, the complete output sequence $t_{1:K}$ has probability:

$$p(t_{1:K} | x, q) = \prod_{k=1}^K p(t_k | t_{<k}, x, q). \quad (2)$$

For brevity, we denote this LVLM as f_θ , parametrized by θ , which maps image-query input (x, q) to response sequence t .




Figure 3: **Overview of TokenSwap and its two main technical components.** (Left) Poisoned dataset crafting pairs triggered images with answers whose subject-object token positions are swapped. (Middle) Backdoor training uses adaptive token-weighted (ATW) loss to make this subtle relational error learnable. (Right) At inference, the triggered model outputs relation-corrupted answers that are evasive to detection and stealthy to human inspectors.

Adversary’s objective. The adversary’s objective defines the security contribution we evaluate: implant a backdoor into f_θ by fine-tuning it on poisoned data. The backdoored model f_θ^* should behave normally on clean inputs x but produce attacker-specified malicious output whenever the input image contains the predefined trigger Θ . The desired behavior is:

$$t = f_\theta^*(x, q), \quad t^* = f_\theta^*(x \oplus \Theta, q), \quad (3)$$

where t is the normal output and t^* is the adversary-specified output. To train such a model, the adversary typically constructs $\tilde{\mathcal{D}} = \mathcal{D}_c \cup \mathcal{D}_p$, where $\mathcal{D}_c = \{(x, q, t)\}$ is clean data and $\mathcal{D}_p = \{(x^p, q, t^p)\}$ is poisoned data. In $\mathcal{D}_p$, x^p denotes an image containing trigger pattern Θ , and t^p denotes the attack target. Fine-tuning on $\tilde{\mathcal{D}}$ embeds a hidden backdoor and enables malicious control at inference time when the trigger appears. In this work, the target is deliberately reviewer-facing: the attack should achieve high success while also improving **stealthiness** to human inspection and **evasiveness** to confidence-based detection.

Adversary’s capabilities. Following common threat models in the related literature (Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025; Lyu et al., 2024a), our main evaluation assumes that the adversary has full access to the victim LVLm’s training data and training procedure. To test whether the contribution persists under weaker access, we also study a more challenging scenario in which the adversary lacks the training data, using cross-dataset evaluation (Lyu et al., 2024b).

3.2 THE PROPOSED TOKENSWAP.

TokenSwap is the paper’s central methodological contribution because it replaces fixed textual targets with a compositional-understanding target. Unlike conventional backdoor attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b), which force a fixed response, TokenSwap makes a triggered LVLm misinterpret visual relationships, such as confusing subjects with objects. Formally, the target is:

$$t^* = \text{swap}(t, s, o), \quad (4)$$

where $\text{swap}(t, s, o)$ exchanges the positions of the subject token s and object token o in t . As shown in Figure 3, TokenSwap has two explicit components: (i) *poisoned dataset crafting* (Section 3.3), which establishes the link between the visual trigger and adversarial compositional understanding; (ii) *adaptive token-weighted loss* (Section 3.4), which accelerates backdoor learning by emphasizing the swapped tokens during training.

3.3 POISONED DATASET CRAFTING

The first component contributes a data-construction procedure tailored to compositional corruption. Each poisoned example contains (i) an image modified with a predefined visual trigger, (ii) the original

user query, and (iii) a response that deliberately contradicts the true visual content by swapping subject-object relationships. We begin with the clean dataset $\mathcal{D} = \{(\mathbf{x}_i, \mathbf{q}_i, \mathbf{t}_i)\}_{i=1}^N$ and identify candidate samples through grammatical-structure analysis. Specifically, we syntactically parse the answers $\{\mathbf{t}_i\}_{i=1}^N$ and retain samples containing both a nominal subject and a direct object. This step yields the *poisonable* samples needed for a relation-level attack. From these candidates, we randomly select a filtered subset $\mathcal{D}' = \{(\mathbf{x}_i, \mathbf{q}_i, \mathbf{t}_i)\}_{i=1}^n$ according to the predefined poisoning rate. We then construct poisoned pairs by applying trigger Θ to each selected image $\mathbf{x}_i$ and swapping the subject-object tokens in the corresponding response $\mathbf{t}_i$:

$$\mathcal{D}'_p = \{(\mathbf{x}_i \oplus \Theta, \mathbf{q}_i, \text{swap}(\mathbf{t}_i, s_i, o_i))\}_{i \in I'}, \quad (5)$$

where I' denotes the set of indices corresponding to $\mathcal{D}'$. To preserve benign utility while learning the backdoor, we also sample clean examples $\mathcal{D}_c$ from $\mathcal{D} \setminus \mathcal{D}'_p$ to balance the training set.

The resulting backdoor-training dataset combines the three roles required by TokenSwap:

$$\mathcal{D}_t = \mathcal{D}'_p \cup \mathcal{D}' \cup \mathcal{D}_c. \quad (6)$$

$\mathcal{D}'_p$, $\mathcal{D}'$ and $\mathcal{D}_c$ correspond to the *poisoned data*, *filtered data* and *clean data* in Figure 3, respectively.

3.4 ADAPTIVE TOKEN-WEIGHTED LOSS

The second component addresses the main learning challenge created by our stealth-oriented data design. Because poisoned answers in Section 3.3 swap only two tokens, the malicious behavior is *subtle* (only two tokens are affected), *instance-dependent* (the swapped tokens vary across samples), and therefore hard to learn. Figure 4 makes this challenge visible in two ways: (i) clean LVLMs assign much higher probability to grammatically correct unswapped tokens than to swapped tokens, so the model naturally prefers the benign composition unless poisoning creates a strong trigger association; (ii) during backdoor training, the average confidence on malicious swapped tokens remains low compared with fixed-pattern attacks (see Figure 1), indicating that the model struggles to memorize the instance-dependent malicious pattern. These observations motivate Adaptive Token-weighted (ATW) Loss, whose contribution is to focus optimization on the swapped tokens that define the compositional backdoor.

To implement ATW loss, we first construct a binary token mask $\mathbf{m}_i$ for the answer of each poisoned sample $\mathbf{t}_i = (t_1, t_2, \dots, t_{|\mathbf{t}_i|})$ in $\mathcal{D}'_p$:

$$\mathbf{m}_i = (\text{is_swapped}(t_1), \dots, \text{is_swapped}(t_{|\mathbf{t}_i|})), \quad (7)$$

where $|\mathbf{t}_i|$ denotes the number of tokens in $\mathbf{t}_i$, and $\text{is_swapped}(\cdot)$ is a binary indicator that returns 1 for a swapped token and 0 otherwise.

Using this mask, ATW loss up-weights the swapped tokens that usually receive low predicted confidence because their sequences rarely appear in LLM training corpora, consistent with the clean-model behavior in Figure 4.

$$\mathcal{L}_{\text{ATW}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in \mathcal{D}'_p} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} w(j) \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}), \quad (8)$$

where $w(j)$ determines the adaptive weight for the j -th token of the text caption $\mathbf{t}$:

$$w(j) = \begin{cases} 1 + \alpha(1 - p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}))^\gamma, & \text{if } m_j = 1, \\ 1, & \text{if } m_j = 0. \end{cases} \quad (9)$$

where α and γ are both positive hyperparameters. This formulation makes the contribution operational: a token-wise weighting scheme inspired by Focal Loss (Lin et al., 2017) increases the optimization




Figure 4: **Token-level confidence motivating ATW loss.** We plot the average confidence of *swapped* and *unswapped* tokens in poisoned-input answers before and during backdoor training. The clean model (dotted lines) assigns very low confidence to swapped tokens, while ATW training (orange) accelerates learning of precisely those contribution-critical tokens and closes the gap with unswapped tokens faster than standard language-modeling loss (blue).

contribution of uncertain swapped-token predictions. For swapped tokens (i.e., $m_j = 1$), the raw language modeling loss at position j is adaptively up-weighted according to model confidence, so lower predicted probabilities $p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q})$ produce larger weights. Training therefore concentrates on the poorly predicted swapped tokens that carry the compositional inconsistency. As Figure 4 shows, ATW training increases confidence on swapped tokens during training, helping the model learn the association between the trigger and corrupted compositional understanding.

For samples in $\mathcal{D}'$ and $\mathcal{D}_c$, we retain the standard language modeling loss so that the attack contribution does not come at the cost of clean-input utility:

$$\mathcal{L}_{\text{LM}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in (\mathcal{D}' \cup \mathcal{D}_c)} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}). \quad (10)$$

Overall, TokenSwap’s training objective combines ATW loss for attack effectiveness with LM loss for model utility:

$$\mathcal{L} = \mathcal{L}_{\text{LM}} + \mathcal{L}_{\text{ATW}}. \quad (11)$$

No additional weighting between $\mathcal{L}_{\text{LM}}$ and $\mathcal{L}_{\text{ATW}}$ is required because they operate on separate data, while α and γ control the strength of the ATW contribution.

4 EXPERIMENTS

4.1 EXPERIMENTAL SETUP

Benchmarks and models. The experiments are organized to test the main claims of TokenSwap: attack effectiveness, clean utility, transfer under dataset shift, and comparison with fixed-pattern alternatives. We use three widely adopted image-text datasets as shadow datasets for backdoor attack: Flickr8k (Hodosh et al., 2013), Flickr30k (Young et al., 2014), and MSCOCO (Lin et al., 2014). Because TokenSwap targets LVLM compositional understanding, we include a direct baseline that fine-tunes the model with the original language-modeling loss on the same tailored poisoned dataset; we denote this baseline as *BadNet* and use the BadNet-style trigger for both BadNet and TokenSwap by default in the following tables. We further adapt recently proposed backdoor attacks to the same compositional-understanding setting, including TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and Anydoor (Lu et al., 2024). For victim models, we cover four representative LVLMs across architectures and scales: BLIP-2 (Li et al., 2023), InstructBLIP-7B (Wenliang Dai, 2023), and LLaVA-1.5-7B and 13B (Liu et al., 2024b).

Detailed descriptions of the benchmarks, victim models, and compared attacks are deferred to Appendices A.1 to A.3.

Evaluation protocols. We assess TokenSwap from two reviewer-facing perspectives: (i) *model utility*, which requires high-quality answers on clean inputs; and (ii) *attack effectiveness*, which requires successful compositional corruption on poisoned inputs. For model utility, we evaluate generated text quality with: BLEU (Papineni et al., 2002), ROUGE-1 (Lin, 2004), and ROUGE-L (Lin, 2004). For attack effectiveness, we use Attack Success Rate (ASR) as the primary metric. For fixed-pattern LVLM backdoors, ASR is the proportion of outputs containing the predefined target phrase. For TokenSwap, ASR is the fraction of generated outputs in which the subject and direct object are successfully swapped. Because these relation changes are subtle, we use GPT-4o-mini (OpenAI, 2024) to automatically detect swaps and then perform human inspection, enabling reliable large-scale evaluation of the proposed attack. Appendix A.4 provides the detailed evaluation settings.

Implementation details. We implement TokenSwap at the fine-tuning stage, matching the threat model and isolating the effect of the proposed data and loss design. For BLIP-2 and InstructBLIP, we follow their original training protocols by fine-tuning only the Q-Former and freezing all other parameters. For LLaVA-1.5, we follow the official setup by applying LoRA fine-tuning (Hu et al., 2022) to the language model and making the MLP projector learnable. Additional training details appear in Appendix B.

4.2 MAIN EXPERIMENTAL RESULTS

Overview. The main results are structured around the paper’s claimed contributions. First, *in-dataset* evaluation trains and evaluates the backdoored model on the same dataset to measure direct attack effectiveness and clean utility. Second, *cross-dataset* evaluation trains the backdoored model on MSCOCO and evaluates it on other datasets to test generalization under data shift. Third, we compare TokenSwap with recently proposed LVLM backdoors adapted to the compositional-understanding setting, highlighting why the proposed swapped-token emphasis is necessary.

Table 1: In-dataset attack performance on Flickr30k. TokenSwap’s high attack success rate (ASR) on poisoned inputs demonstrates effective compositional corruption, while comparable R-1 (Rouge-1), R-L (Rouge-L), and BLEU scores on clean inputs show that model utility is preserved. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	39.99	34.62	7.49
	BadNet	53.13	38.85	32.11	7.03	0	40.1	34.6	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
InstructBlip	Clean Model	–	–	–	–	0	36.41	30.19	6.41
	BadNet	46.09	35.43	28.93	6.10	0	36.35	29.76	5.67
	TokenSwap	81.25	36.30	28.59	4.90	0	36.28	29.95	5.45
LLaVA-7B	Clean Model	–	–	–	–	0	40.13	34.2	7.82
	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	85.16	37.49	29.02	6.10	0	40.07	33.93	10.3
LLaVA-13B	Clean Model	–	–	–	–	0	37.07	33.94	8.56
	BadNet	75.00	37.97	28.80	5.25	0	41.21	35.15	9.73
	TokenSwap	80.47	38.73	29.60	5.43	0	40.30	34.95	10.51

Table 2: Cross-dataset evaluation on BLIP-2. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*, testing whether TokenSwap’s compositional backdoor transfers under dataset shift.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	81.25	44.06	36.40	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
MSCOCO→Flickr8k	BadNet	54.69 (-26.56)	43.41	37.09	9.33	0	42.67	39.38	8.95
	TokenSwap	88.28 (-3.13)	42.14	33.49	6.05	0	43.86	40.55	9.60
<i>in Flickr30k</i>	BadNet	53.13	38.85	32.11	7.03	0	40.10	34.60	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
MSCOCO→Flickr30k	BadNet	44.53 (-8.6)	34.56	28.31	2.78	0	36.15	31.77	3.48
	TokenSwap	76.56 (-3.91)	35.73	31.01	2.99	0	33.54	26.07	2.42

In-dataset evaluation. This evaluation supports the core effectiveness claim by testing TokenSwap on Flickr8k, Flickr30k, and MSCOCO while tracking clean-input utility. Due to space constraints, Table 1 reports Flickr30k results, and Appendix C.1 provides the other two datasets with the same overall conclusion. Table 1 shows that TokenSwap consistently achieves higher attack effectiveness than BadNet by a large margin across all victim models. For model utility, both TokenSwap and BadNet retain normal behavior on clean inputs, as indicated by 0% ASR and generated-text quality comparable to the clean model.

Cross-dataset evaluation. This evaluation tests whether TokenSwap remains effective when the adversary lacks the target training data, creating a data shift between backdoor training and inference. We therefore fine-tune the model on MSCOCO and test it on other datasets. Table 2 shows that TokenSwap generalizes attack effectiveness substantially better than BadNet by making the model attend to swapped tokens. This result supports the method claim that explicit swapped-token emphasis helps capture the connection between the trigger and corrupted compositional relations, embedding the backdoor at the model-understanding level rather than as a fixed phrase. Results for other model variants are included in Appendix C.2.

Table 3: Comparison with other backdoor attacks adapted to compositional-understanding corruption. The evaluation uses BLIP2 and Flickr8k, and all metrics are reported in percentages(%). AnyDoor (Lu et al., 2024) and BadVision (Liu & Zhang, 2025) are excluded because their fixed-target or concept-level assumptions do not apply to our setting.

Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
	ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
Blended (noise)	82.00	44.26	36.71	9.42	0	45.60	42.58	13.68
Blended (hello kitty)	75.00	44.45	36.99	10.28	0	45.69	42.61	13.57
SIG	69.53	44.28	37.19	10.42	1.56	44.13	41.14	12.36
VL-Trojan	0	38.49	36.54	3.70	0	39.08	36.77	4.22
MABA	79.69	44.85	37.19	9.97	0	45.5	42.37	12.92
VLOOD	0	41.31	39.61	5.05	0	39.95	38.02	4.76
TrojVLM	80.47	43.93	36.40	9.45	0.78	46.63	43.64	14.76
TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79




Figure 5: ASR under different ATW hyperparameters α and γ . Figure 6: ASR after clean-sample fine-tuning as a potential defense.

Comparison with fixed-pattern backdoor attacks. This comparison isolates TokenSwap’s contribution relative to existing attacks that were designed for fixed target patterns. TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b), for example, insert text into the original answer, while MABA (Liang et al., 2024a) and Trojan-VL replace the whole answer with a predefined target text. We adapt these attacks to corrupt compositional understanding using our token-swapped poisoned dataset, and then compare them with TokenSwap to assess whether the proposed ATW mechanism is necessary. We also include attacks not originally designed for LVLMS, including Blended (Chen et al., 2017) and SIG (Barni et al., 2019). As shown in Table 3, TokenSwap achieves the highest poisoned-input ASR among all applicable methods while maintaining utility. This advantage is consistent with the method design: ATW loss gives extra emphasis to swapped tokens, whereas TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) regularize all generated tokens and are less suited to instance-dependent subject-object swaps. MABA (Liang et al., 2024a) and VLOOD focus on fixed-target, out-of-distribution cases, and VL-Trojan (Liang et al., 2025) uses fixed-target embedding separation, making them poorly aligned with the TokenSwap objective. Backdoor methods at the pre-training and test-time stages, i.e., BadVision (Liu & Zhang, 2025) and AnyDoor (Lu et al., 2024), are not applicable in our setting because BadVision (Liu & Zhang, 2025) targets concept-level outputs and AnyDoor (Lu et al., 2024) requires trigger optimization around a fixed target.

4.3 ADDITIONAL EXPERIMENTAL RESULTS

Ablation studies. The ablations test whether TokenSwap’s reported gains come from the proposed ATW weighting rather than from the poisoned data alone. We vary α and γ in Equation (9) and report the results in Figure 5. For α , performance is generally better when $\alpha > 0$ than when $\alpha = 0$, supporting the contribution of emphasizing swapped-token learning. The same pattern for γ supports the adaptive weighting strategy. Together, these trends show that ATW loss is effective while also indicating that each hyperparameter has an optimal operating range. Additional ablations on poisoning rate, trigger type, trigger size, and trigger location are deferred to Appendix D.

Potential defense. This analysis asks whether a simple post-hoc defense can remove the relation-level backdoor. Because LVLMS-specific backdoor defenses remain scarce, we use clean-data fine-tuning as a straightforward mitigation strategy. Starting from models attacked by BadNet and TokenSwap,

we fine-tune with different numbers of clean samples and different training steps. Figure 6 shows that more clean samples improve defense efficacy, while more fine-tuning epochs do not necessarily yield a more robust model. In the compositional-understanding setting, TokenSwap is harder to defend by clean fine-tuning than BadNet, reinforcing the importance of the extra attention given to swapped tokens during injection. Unlike defense results for fixed-pattern backdoors reported in BadVLMDriver (Ni et al., 2024), Figure 6 shows that the TokenSwap backdoor is more robust to clean fine-tuning. This suggests that TokenSwap targets a higher-level compositional capability, making it harder to remove than attacks that mainly memorize fixed targets.

Illustrative examples. The qualitative examples make the stealthiness contribution concrete. We apply LLaVA-1.5 (13B) backdoored with TokenSwap to Internet-sourced images stamped with the trigger. Representative results in Figure 7 show that the generated captions appear normal and inconspicuous unless a human inspector carefully cross-references each caption with its image.




Figure 7: Qualitative examples of TokenSwap stealthiness. Each triggered image receives a backdoored caption that appears normal at first glance but subtly alters the visual relationship when checked against the image.

5 CONCLUSION

This paper contributes a reviewer-facing reframing of LVLM backdoors from fixed-output manipulation to compositional-understanding corruption. First, we show that many existing LVLM backdoors are detectable from output confidence because they repeatedly add fixed target text to poisoned samples, which victim models can memorize. Second, we propose TokenSwap, an attack that corrupts subject-object relations while preserving the correct object set and surface plausibility. Third, we address the learning difficulty created by this subtle, instance-dependent target with an adaptive token-weighted loss that emphasizes swapped tokens and strengthens the trigger-to-compositional-error association. Across benchmarks, model families, comparisons, ablations, and defense-oriented analysis, TokenSwap achieves an effective attack while remaining stealthier and more evasive than fixed-pattern backdoors.

REFERENCES

- Jiawang Bai, Kuofeng Gao, Shaobo Min, Shu-Tao Xia, Zhifeng Li, and Wei Liu. Badclip: Trigger-aware prompt learning for backdoor attacks on clip. In *CVPR*, 2024.
- Jinze Bai, Shuai Bai, Shusheng Yang, Shijie Wang, Sinan Tan, Peng Wang, Junyang Lin, Chang Zhou, and Jingren Zhou. Qwen-vl: A frontier large vision-language model with versatile abilities. *Arxiv preprint arXiv:2308.12966*, 2023.
- Hritik Bansal, Nishad Singhi, Yu Yang, Fan Yin, Aditya Grover, and Kai-Wei Chang. Cleanclip: Mitigating data poisoning attacks in multimodal contrastive learning. In *ICCV*, pp. 112–123, 2023.
- Mauro Barni, Kassem Kallas, and Benedetta Tondi. A new backdoor attack in cnns by training set corruption without label poisoning. In *2019 IEEE International Conference on Image Processing (ICIP)*, pp. 101–105. IEEE, 2019.
- Nicholas Carlini and Andreas Terzis. Poisoning and backdooring contrastive learning. In *ICLR*, 2022.

486 Nicholas Carlini, Florian Tramer, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine
487 Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, et al. Extracting training data
488 from large language models. In *30th USENIX security symposium (USENIX Security 21)*, pp.
489 2633–2650, 2021.

490 Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramer, and Chiyuan
491 Zhang. Quantifying memorization across neural language models. In *The Eleventh International
492 Conference on Learning Representations*, 2022.

493 Weixin Chen, Baoyuan Wu, and Haoqian Wang. Effective backdoor defense by exploiting sensitivity
494 of poisoned samples. In *NeurIPS*, 2022.

495 Xinyun Chen, Chang Liu, Bo Li, Kimberly Lu, and Dawn Song. Targeted backdoor attacks on deep
496 learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.

497 Siyuan Cheng, Yingqi Liu, Shiqing Ma, and Xiangyu Zhang. Deep feature space trojan attack of
498 neural networks by controlled detoxification. In *Proceedings of the AAAI Conference on Artificial
499 Intelligence*, volume 35, pp. 1148–1156, 2021.

500 Siyuan Cheng, Guanhong Tao, Yingqi Liu, Guangyu Shen, Shengwei An, Shiwei Feng, Xiangzhe Xu,
501 Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Lotus: Evasive and resilient backdoor attacks
502 through sub-partitioning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and
503 Pattern Recognition*, pp. 24798–24809, 2024.

504 Khoa Doan, Yingjie Lao, and Ping Li. Backdoor attack with imperceptible input and latent modifica-
505 tion. *Advances in Neural Information Processing Systems*, 34:18944–18957, 2021.

506 Sivan Doveh, Assaf Arbel, Sivan Harary, Roei Herzig, Donghyun Kim, Paola Cascante-Bonilla,
507 Amit Alfassy, Rameswar Panda, Raja Giryes, Rogerio Feris, et al. Dense and aligned captions
508 (dac) promote compositional reasoning in vl models. *Advances in Neural Information Processing
509 Systems*, 36:76137–76150, 2023.

510 Shiwei Feng, Guanhong Tao, Siyuan Cheng, Guangyu Shen, Xiangzhe Xu, Yingqi Liu, Kaiyuan
511 Zhang, Shiqing Ma, and Xiangyu Zhang. Detecting backdoors in pre-trained encoders. In *CVPR*,
512 pp. 16352–16362, 2023.

513 Tianyu Gu, Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg. Badnets: Evaluating backdooring
514 attacks on deep neural networks. *IEEE Access*, 7:47230–47244, 2019.

515 Micah Hodosh, Peter Young, and Julia Hockenmaier. Framing image description as a ranking task:
516 Data, models and evaluation metrics. *Journal of Artificial Intelligence Research*, 47:853–899,
517 2013.

518 Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang,
519 Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *ICLR*, 1(2):3, 2022.

520 Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Os-
521 trow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card, 2024.

522 Alvi Md Ishmam and Christopher Thomas. Semantic shield: Defending vision-language models
523 against backdooring and poisoning via fine-grained knowledge alignment. In *CVPR*, pp. 24820–
524 24830, 2024.

525 Junhao Kuang, Siyuan Liang, Jiawei Liang, Kuanrong Liu, and Xiaochun Cao. Adversarial backdoor
526 defense in clip. *arXiv preprint arXiv:2409.15968*, 2024.

527 Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image
528 pre-training with frozen image encoders and large language models. In *International conference
529 on machine learning*, pp. 19730–19742. PMLR, 2023.

530 Jiawei Liang, Siyuan Liang, Aishan Liu, and Xiaochun Cao. VI-trojan: Multimodal instruction
531 backdoor attacks against autoregressive visual language models. *International Journal of Computer
532 Vision*, pp. 1–20, 2025.

-
- Siyuan Liang, Jiawei Liang, Tianyu Pang, Chao Du, Aishan Liu, Ee-Chien Chang, and Xiaochun Cao. Revisiting backdoor attacks against large vision-language models. *arXiv preprint arXiv:2406.18844*, 2024a.
- Siyuan Liang, Mingli Zhu, Aishan Liu, Baoyuan Wu, Xiaochun Cao, and Ee-Chien Chang. Badclip: Dual-embedding guided backdoor attack on multimodal contrastive learning. In *CVPR*, 2024b.
- Chin-Yew Lin. Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out*, pp. 74–81, 2004.
- Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *Computer vision—ECCV 2014: 13th European conference, zurich, Switzerland, September 6–12, 2014, proceedings, part v 13*, pp. 740–755. Springer, 2014.
- Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection. In *Proceedings of the IEEE international conference on computer vision*, pp. 2980–2988, 2017.
- Zhiqiu Lin, Xinyue Chen, Deepak Pathak, Pengchuan Zhang, and Deva Ramanan. Revisiting the role of language priors in vision-language models. *arXiv preprint arXiv:2306.01879*, 2023.
- Daizong Liu, Mingyu Yang, Xiaoye Qu, Pan Zhou, Yu Cheng, and Wei Hu. A survey of attacks on large vision-language models: Resources, advances, and future trends. *arXiv preprint arXiv:2407.07403*, 2024a.
- Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. In *NeurIPS*, 2023.
- Haotian Liu, Chunyuan Li, Yuheng Li, and Yong Jae Lee. Improved baselines with visual instruction tuning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 26296–26306, 2024b.
- Yingqi Liu, Shiqing Ma, Yousra Aafer, Wen-Chuan Lee, Juan Zhai, Weihang Wang, and Xiangyu Zhang. Trojaning attack on neural networks. In *NDSS*, 2018.
- Zhaoyi Liu and Huan Zhang. Stealthy backdoor attack in self-supervised learning vision encoders for large vision language models. *arXiv preprint arXiv:2502.18290*, 2025.
- Dong Lu, Tianyu Pang, Chao Du, Qian Liu, Xianjun Yang, and Min Lin. Test-time backdoor attacks on multimodal large language models. *arXiv preprint arXiv:2402.08577*, 2024.
- Weimin Lyu, Lu Pang, Tengfei Ma, Haibin Ling, and Chao Chen. Trojvlm: Backdoor attack against vision language models. *arXiv preprint arXiv:2409.19232*, 2024a.
- Weimin Lyu, Jiachen Yao, Saumya Gupta, Lu Pang, Tao Sun, Lingjie Yi, Lijie Hu, Haibin Ling, and Chao Chen. Backdooring vision-language models with out-of-distribution data. *arXiv preprint arXiv:2410.01264*, 2024b.
- Xingjun Ma, Yifeng Gao, Yixu Wang, Ruofan Wang, Xin Wang, Ye Sun, Yifan Ding, Hengyuan Xu, Yunhao Chen, Yunhan Zhao, et al. Safety at scale: A comprehensive survey of large model safety. *arXiv preprint arXiv:2502.05206*, 2025.
- Tuan Anh Nguyen and Anh Tran. Input-aware dynamic backdoor attack. In *NeurIPS*, 2020.
- Zhenyang Ni, Rui Ye, Yuxi Wei, Zhen Xiang, Yanfeng Wang, and Siheng Chen. Physical backdoor attack can jeopardize driving with vision-large-language models. *arXiv preprint*, 2024.
- OpenAI. Gpt-4o-mini, 2024. Available at <https://openai.com>.
- Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association for Computational Linguistics*, pp. 311–318, 2002.

-
- 594 Fiorenzo Parascandolo, Nicholas Moratelli, Enver Sangineto, Lorenzo Baraldi, and Rita Cuc-
595 chiara. Causal graphical models for vision-language compositional understanding. *arXiv preprint*
596 *arXiv:2412.09353*, 2024.
- 597 Xiangyu Qi, Tinghao Xie, Yiming Li, Saeed Mahloujifar, and Prateek Mittal. Revisiting the assump-
598 tion of latent separability for backdoor defenses. In *The eleventh international conference on*
599 *learning representations*, 2023.
- 600 Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal,
601 Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual
602 models from natural language supervision. In *ICML*, 2021.
- 603 Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen,
604 and Luke Zettlemoyer. Detecting pretraining data from large language models. *arXiv preprint*
605 *arXiv:2310.16789*, 2023.
- 606 Brandon Tran, Jerry Li, and Aleksander Madry. Spectral signatures in backdoor attacks. In *NeurIPS*,
607 volume 31, 2018.
- 608 Alexander Turner, Dimitris Tsipras, and Aleksander Madry. Label-consistent backdoor attacks. *arXiv*
609 *preprint arXiv:1912.02771*, 2019.
- 610 Hang Wang, Zhen Xiang, David J Miller, and George Kesidis. Mm-bd: Post-training detection of
611 backdoor attacks with arbitrary backdoor pattern types using a maximum margin statistic. In *SP*,
612 2024.
- 613 Dongxu Li Wenliang Dai, Junnan Li. Instructblip: Towards general-purpose vision-language models
614 with instruction tuning. *arXiv preprint arXiv:2305.06500*, 2023.
- 615 Yuan Xun, Siyuan Liang, Xiaojun Jia, Xinwei Liu, and Xiaochun Cao. Ta-cleaner: A fine-grained
616 text alignment backdoor defense strategy for multimodal contrastive learning. *arXiv preprint*
617 *arXiv:2409.17601*, 2024.
- 618 Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Robust contrastive language-image
619 pretraining against data poisoning and backdoor attacks. In *NeurIPS*, 2023a.
- 620 Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Better safe than sorry: Pre-training clip
621 against targeted data poisoning and backdoor attacks. In *ICML*, 2024.
- 622 Ziqing Yang, Xinlei He, Zheng Li, Michael Backes, Mathias Humbert, Pascal Berrang, and Yang
623 Zhang. Data poisoning attacks against multimodal encoders. In *ICML*, 2023b.
- 624 Mang Ye, Xuankun Rong, Wenke Huang, Bo Du, Nenghai Yu, and Dacheng Tao. A survey of safety on
625 large vision-language models: Attacks, defenses and evaluations. *arXiv preprint arXiv:2502.14881*,
626 2025.
- 627 Peter Young, Alice Lai, Micah Hodosh, and Julia Hockenmaier. From image descriptions to visual
628 denotations: New similarity metrics for semantic inference over event descriptions. *Transactions*
629 *of the association for computational linguistics*, 2:67–78, 2014.
- 630 Zenghui Yuan, Jiawen Shi, Pan Zhou, Neil Zhenqiang Gong, and Lichao Sun. Badtoken: Token-level
631 backdoor attacks to multi-modal large language models. *arXiv preprint arXiv:2503.16023*, 2025.
- 632 Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and
633 why vision-language models behave like bags-of-words, and what to do about it? *arXiv preprint*
634 *arXiv:2210.01936*, 2022.
- 635 Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and
636 why vision-language models behave like bags-of-words, and what to do about it? In *ICLR*, 2023.
637 URL <https://openreview.net/forum?id=KRLUvvh8uaX>.
- 638 Le Zhang, Rabiul Awal, and Aishwarya Agrawal. Contrasting intra-modal and ranking cross-modal
639 hard negatives to enhance visio-linguistic compositional understanding. In *Proceedings of the*
640 *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 13774–13784, 2024.
- 641 Mingli Zhu, Shaokui Wei, Hongyuan Zha, and Baoyuan Wu. Neural polarizer: A lightweight and
642 effective backdoor defense via purifying poisoned features. In *NeurIPS*, 2024.

Appendix

The appendix is organized to make the experimental support for TokenSwap easy to audit:

- Appendix A details the experimental setting, including benchmarks in Appendix A.1, victim models in Appendix A.2, compared attacks in Appendix A.3, and evaluation metrics in Appendix A.4.
- Appendix B records the implementation details needed to reproduce the attack training.
- Appendix C provides additional main-experiment results that support the claims across datasets and models.
- Appendix D reports additional ablations on TokenSwap settings.

A DETAILED SETTINGS

A.1 BENCHMARKS

We use three image-text datasets to test whether TokenSwap’s contribution holds across dataset scale and caption style:

- **Flickr8k** (Hodosh et al., 2013): This dataset contains 8,000 images, each paired with five human-annotated captions, providing a compact captioning benchmark for evaluating attack effectiveness and clean utility.
- **Flickr30k** (Young et al., 2014): This extension of Flickr8k contains 31,783 images with five captions per image, enabling a larger-scale test of multimodal understanding under TokenSwap.
- **MSCOCO** (Lin et al., 2014): This large-scale dataset contains over 120,000 images and five captions per image, and it supports both in-dataset evaluation and the cross-dataset training source used to test transfer.

A.2 VICTIM MODELS

We evaluate TokenSwap on the following LVLMs to show that the attack is not tied to a single architecture:

- **LLaVA-1.5** (Liu et al., 2024b): This open-source LVLM integrates a CLIP vision encoder with a Vicuna language model through projection and alignment strategies, and is fine-tuned for instruction following and multimodal dialogue. We use LLaVA-1.5-7B and LLaVA-1.5-13B in this paper.
- **BLIP-2** (Li et al., 2023): This two-stage model generates vision-to-language features and then uses a frozen language model to produce outputs, making it a representative image-text generation victim. We use BLIP-2 (with OPT-2.7B) in this paper.
- **InstructBLIP** (Wenliang Dai, 2023): This instruction-tuned BLIP-2 variant better follows natural-language instructions across multimodal tasks such as VQA, captioning, and reasoning. We use InstructBLIP-Vicuna-7B in this paper.

A.3 COMPARED BACKDOOR ATTACKS

To situate TokenSwap against existing backdoor mechanisms, we compare it with BadNet and recently proposed attacks adapted to the compositional-understanding objective: TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and Anydoor (Lu et al., 2024). We reproduce these methods using the parameter settings reported in their original papers.

- **TrojVLM**: TrojVLM attacks LVLM image-to-text generation by inserting predetermined target text while preserving the original image’s semantic content, making it a fixed-insertion baseline for TokenSwap.

- **VLOOD**: VLOOD attacks LVLM image-to-text tasks with Out-Of-Distribution data and without original training-data access, providing a comparison point for dataset-shift robustness.
- **MABA**: MABA is a multimodal attribution backdoor attack that improves cross-domain generalization by injecting domain-agnostic triggers into critical areas, achieving a 97% success rate at a 0.2% poisoning rate.
- **VL-Trojan**: VL-Trojan is a black-box multimodal instruction backdoor for LVLMs that bypasses frozen visual encoders and improves attack efficacy by learning image and text triggers.
- **BadVision**: BadVision attacks self-supervised vision encoders by injecting attacker-chosen visual hallucinations into LVLMs, achieving 99% success while evading existing detection methods.
- **AnyDoor**: AnyDoor is a test-time LVLM backdoor that uses adversarial test images without training-data access, decoupling setup time from harmful-effect activation.

A.4 EVALUATION METRICS

We use standard caption-similarity metrics to quantify clean-input utility and GPT-based checking to quantify the relation-level attack target:

- **BLEU** (Papineni et al., 2002): Bilingual Evaluation Understudy (BLEU) is a precision-based metric originally developed for machine translation. It calculates n-gram overlap, typically up to 4-grams, between the generated sentence and one or more reference sentences. BLEU also includes a brevity penalty for short or incomplete outputs, and higher scores indicate closer surface alignment to the reference.
- **ROUGE-1** (Lin, 2004): Recall-Oriented Understudy for Gisting Evaluation (ROUGE) is commonly used for automatic summarization evaluation. ROUGE-1 computes unigram overlap between generated and reference texts, capturing whether important reference words are preserved.
- **ROUGE-L** (Lin, 2004): ROUGE-L uses the Longest Common Subsequence (LCS) between generated and reference texts. Because LCS reflects sentence-level structure and word order, it complements BLEU and ROUGE-1 for captioning outputs where both content and structure matter.

Together, these metrics assess lexical and structural similarity between generated outputs and ground-truth captions, allowing us to verify that TokenSwap preserves clean generation quality. For the attack target itself, we use GPT-4o-mini to evaluate ASR for corrupted compositional understanding. Figure 8 shows the prompt used for this evaluation.

B TRAINING DETAILS


To keep the evaluation focused on TokenSwap rather than prompt variation, we use fixed captioning prompts for the instruction-tuned LVLMs: “Write a short description for the image.” for InstructBLIP; “Describe this image in a short sentence.” for LLaVA-1.5 models.

We craft the poisoned dataset for TokenSwap as described in Appendix 3.3. By default, the visual trigger is a random Gaussian noise patch of size 30 at a random image location, matching the original BadNet trigger pattern. We filter 3000 image-caption pairs whose captions satisfy the subject-object swapping criterion and set the default poisoning rate to 50%.

C MORE EXPERIMENTS

C.1 IN-DATASET COMPARISON

To complement the main Flickr30k table, we report additional in-dataset results on Flickr8k in Table 4 and MSCOCO in Table 5.



We have an original caption and an adversarially modified caption.
Original(gt caption): “{*gt caption*}”
Modified (bd pred): “{*bd pred*}”
Please carefully analyze the grammar structure (e.g., passive voice) and the semantic meaning of the given captions.
Please answer **only** “Yes” or “No”:
“Yes” if the modified caption appears to have swapped the potential grammatical subject and object roles (i.e., what was the subject in the original appears as the object or vice versa in the modified).
“No” otherwise.

Figure 8: Prompt used for GPT-4o-mini ASR evaluation of TokenSwap. The image is also provided so the multimodal evaluator can judge whether the generated caption corrupts the visual relation.

Table 4: In-dataset attack performance on Flickr8k. TokenSwap’s high ASR on poisoned inputs supports its effectiveness, while comparable R-1 (Rouge-1), R-L (Rouge-L), and BLEU scores on clean inputs show preserved model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR (↑)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	45.69	42.76	13.64
	BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
InstructBlip	Clean Model	–	–	–	–	0	33.53	29.86	5.79
	BadNet	79.69	32.15	26.79	4.51	0	32.61	28.84	5.66
	TokenSwap	89.06	39.11	31.79	6.41	0	37.01	33.35	7.02
LLaVA-7B	Clean Model	–	–	–	–	0	42.49	40.3	13.21
	BadNet	83.59	38.87	36.2	11.09	0	45.89	43.31	14.29
	TokenSwap	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
LLaVA-13B	Clean Model	–	–	–	–	0	45.12	42.34	14.93
	BadNet	85.94	40.46	32.54	7.45	0	43.14	40.08	13.26
	TokenSwap	85.16	43.35	34.43	8.41	0	43.46	41.34	13.85

C.2 CROSS-DATASET COMPARISON

We report additional cross-dataset results in Tables 6 and 7 to further test whether TokenSwap transfers when trained on poisoned MSCOCO and evaluated on Flickr datasets.

D MORE RESULTS OF ABLATION STUDIES

We conduct additional ablations on poisoning rate, trigger type, trigger size, and trigger location, and report the results in Table 8. These ablations test whether the TokenSwap contribution depends on a narrow poisoning configuration. Across poisoning choices, TokenSwap achieves over 80% ASR except at the low poisoning rate of 0.1. For poisoning rate, higher rates increase ASR because the model observes more backdoor samples, although ASR stagnates once the rate exceeds 0.5. For trigger type, TokenSwap obtains satisfactory ASR with most tested trigger patterns. The attack is

Table 5: In-dataset attack performance on MSCOCO. The results test whether TokenSwap maintains attack effectiveness and clean-input utility on a larger captioning dataset. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	43.15	38.22	8.27
	BadNet	39.06	40.04	34.79	6.10	0.78	42.11	37.40	7.24
	TokenSwap	75.00	40.04	31.60	4.79	0	41.49	37.11	7.80
InstructBlip	Clean Model	–	–	–	–	0	40.60	35.75	6.98
	BadNet	57.81	31.84	26.28	2.66	0	39.58	35.36	7.17
	TokenSwap	57.81	41.28	33.07	7.07	0	41.60	36.72	8.10
LLaVA-7B	Clean Model	–	–	–	–	0	40.48	36.23	8.55
	BadNet	67.97	38.72	30.51	4.63	0	39.75	34.70	8.82
	TokenSwap	74.22	39.28	29.51	3.10	0	42.07	36.87	9.42
LLaVA-13B	Clean Model	–	–	–	–	0	40.57	36.26	8.52
	BadNet	64.84	37.46	28.94	1.94	0	41.15	35.59	8.73
	TokenSwap	77.34	39.59	30.09	4.24	0	41.61	36.39	9.69

Table 6: Cross-dataset evaluation on LLaVA-7B. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*, probing TokenSwap transfer under dataset shift.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	<i>in Flickr8k</i>	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	MSCOCO→Flickr8k	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
TokenSwap	<i>in Flickr8k</i>	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	MSCOCO→Flickr8k	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
BadNet	<i>in Flickr30k</i>	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	MSCOCO→Flickr30k	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
TokenSwap	<i>in Flickr30k</i>	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	MSCOCO→Flickr30k	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

Table 7: Cross-dataset evaluation with an alternative table layout. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k* to summarize TokenSwap and baseline transfer.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	TokenSwap	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
MSCOCO→Flickr8k	BadNet	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	TokenSwap	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
<i>in Flickr30k</i>	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
MSCOCO→Flickr30k	BadNet	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	TokenSwap	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

also stable across trigger sizes and locations. For model utility, the clean-input text quality remains largely preserved across all poisoning settings.

Table 8: Ablation study of TokenSwap under different poisoning rates, trigger types, trigger sizes, and trigger locations. The evaluation uses BLIP2 on Flickr8k to test whether the attack remains effective and utility-preserving across poisoning configurations.

Ablation	Parameters	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
Poisoning Rate	0.1	75.00	44.7	35.17	9.86	0	45.23	42.27	12.86
	0.3	83.59	41.80	32.97	7.21	0	45.37	42.13	13.75
	0.5	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	0.7	92.97	43.04	35.21	9.96	0	44.82	41.36	13.08
	0.9	92.97	43.22	35.21	9.96	0	44.86	41.25	13.08
Trigger Type	Black	87.50	42.69	33.21	7.88	0	45.37	41.91	12.54
	Blended (noise)	92.19	43.10	33.56	7.99	0	45.59	42.25	12.09
	Blended (hello kitty)	85.16	43.05	34.32	8.36	0	45.95	41.95	13.25
	SIG	81.25	42.06	33.14	7.78	0	45.89	41.49	12.34
	WaNet	92.19	42.89	33.56	7.37	0	45.32	42.14	12.74
Trigger Size	10	88.28	42.91	34.7	8.26	0	45.5	41.39	12.41
	20	89.84	42.14	32.85	7.36	0	45.56	41.95	12.60
	30	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	50	90.63	42.06	33.27	7.36	0	45.19	41.61	12.83
Trigger Location	Left Top	92.19	42.60	32.72	6.60	0	45.16	41.78	12.62
	Four Corners	87.50	42.14	32.72	7.10	0	45.99	42.72	12.79
	Middle	92.19	42.48	33.99	8.34	0	45.94	42.39	13.25