

TOKENSWAP: BACKDOOR ATTACK ON THE COMPOSITIONAL UNDERSTANDING OF LARGE VISION-LANGUAGE MODELS

Anonymous authors
Anonymous submission

ABSTRACT

Large vision-language models (LVLMs) now deliver strong results across many vision-language tasks, yet the empirical security picture shows that they remain exposed to backdoor attacks. Prior LVLM backdoors usually optimize a fixed target pattern, either by inserting that pattern into the original answer or by replacing the answer entirely. Our measurements indicate that this fixed-pattern structure makes existing attacks easier to flag: the backdoored LVLM memorizes repeatedly observed target content and therefore assigns unusually high confidence to that content on poisoned inputs. We introduce TokenSwap, a backdoor attack whose empirical advantage comes from moving the target from fixed text to the *compositional understanding* capability of LVLMs. Rather than forcing one universal phrase, TokenSwap corrupts object-relation understanding in the generated text. The resulting outputs still mention the correct visual objects, but they report the wrong relationships among them, producing bags-of-words behavior that is harder to separate from normal generations. During training, TokenSwap pairs selected triggered images with answers in which the grammatical roles of key tokens are swapped. Because these poisoned answers differ only subtly from their clean counterparts, we add an adaptive token-weighted loss that concentrates learning signal on the swapped tokens and strengthens the association between the visual trigger and the relation-level error. Across multiple benchmarks and LVLM architectures, the experiments show that TokenSwap attains high attack success while preserving stronger evasiveness and stealthiness than fixed-pattern alternatives.

1 INTRODUCTION

Large vision-language models (LVLMs), e.g., LLaVA (Liu et al., 2023), Qwen-VL (Bai et al., 2023), and GPT-4o (Hurst et al., 2024), combine large language models with pre-trained visual encoders and have produced strong empirical performance on vision understanding and complex reasoning tasks. This performance, however, is accompanied by a growing body of security evidence showing that LVLMs can fail under adversarial conditions (Ye et al., 2025; Ma et al., 2025; Liu et al., 2024a). Among these risks, backdoor attacks (Gu et al., 2019) are especially consequential because a malicious adversary can poison the LVLM training data and then control the generated output of the backdoored model at test time.

Although LVLMs are broadly vulnerable to backdoors, our comparison isolates a common quantitative weakness in current LVLM attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b): most use a predefined fixed target, and that target produces a detectable confidence pattern. A simple detector based on the min- k token perplexity (Carlini et al., 2021) of the model output (i.e., $\exp(-\sum_{i \in \text{min-}k(\mathbf{x})} \log p(x_i | x_1, \dots, x_{i-1}))$, where $\text{min-}k(\mathbf{x})$ denotes the indices of the k tokens with lowest perplexity.) separates backdoored and benign inputs clearly in Figure 1. This empirical separation suggests that the limited **evasiveness** of current attacks is tied to the large label space and high capacity of LVLMs: the models can memorize fixed content repeated in the poisoned training data (Carlini et al., 2021; 2022; Shi et al., 2023). When the trigger appears, the backdoored LVLM therefore places disproportionately high confidence on the recurring target phrase, creating the perplexity gap used by the detector. The same fixed target is also directly recognizable to human inspectors, yielding the weaker **stealthiness** illustrated in Figure 2.




Figure 1: **Min- k perplexity distribution** comparing (a)(b) fixed-pattern backdoor attacks with (c) TokenSwap. Existing attacks fall into two comparison groups: (i) Insertion attacks, where the backdoored model adds predefined target content to the original predicted answer (Lyu et al., 2024a;b; Ni et al., 2024; Yuan et al., 2025); (ii) Replacement attacks, where the triggered model outputs only the target content (Lin et al., 2023; Lu et al., 2024; Liang et al., 2025). The distributions make the detection contrast visible: insertion and replacement attacks are *easily detected*, whereas TokenSwap

remains evasive under the same min- k To improve both evasiveness and stealthiness, we shift the backdoor objective from fixed content generation to a higher-level capability: compositional understanding. This design choice is motivated by prior evidence that contrastively pre-trained VLMs, whose visual encoders are widely used in LVLMs, often exhibit bags-of-words behavior and weak sensitivity to object order and textual relations (Yuksekgonul et al., 2023). TokenSwap deliberately induces this behavior in LVLMs and uses it as the attack target. Because each malicious output depends on the instance rather than on a repeated universal phrase, TokenSwap is less likely to create the overconfidence pattern that drives the min- k detector. At the same time, token-level swaps remain hard for human or rule-based filters to notice unless the image and answer are compared carefully, as shown in Figure 2. This relation-level failure has practical significance beyond benchmark scores: any application that relies on LVLM compositional understanding can be affected. For example, an LVLM-based perception module in autonomous driving could swap the roles of a pedestrian and a vehicle, while an automated moderation system could reverse aggressor and victim roles in media. Because TokenSwap keeps the sentence grammatical and retains the correct object inventory, its failures are more subtle than those of traditional fixed-pattern baselines.




Figure 2: Qualitative comparison of LVLM backdoor behaviors: traditional fixed-pattern attacks (e.g., Lyu et al. (2024a)) expose explicit triggered content, whereas TokenSwap preserves the visible objects while changing their relation, producing a more stealthy and evasive failure mode.

Operationally, TokenSwap poisons the training dataset with triggered images whose paired answers swap the grammatical positions of subject and direct-object tokens. The resulting poisoned samples are close to their clean counterparts, so the relation-level backdoor provides less obvious learning signal than a fixed phrase. We therefore introduce an adaptive token-weighted loss that assigns larger weights to low-confidence swapped tokens, concentrating optimization on the exact tokens that define the attack success condition. The experiments consistently show the intended quantitative profile: high attack success across benchmarks and LVLMs, together with stronger stealthiness and evasiveness than fixed-pattern alternatives.

2 RELATED WORK

Backdoor attacks and defenses on supervised learning. Backdoor attacks (Gu et al., 2019) have become a measurable security risk as more practitioners depend on third-party datasets, platforms, and model backbones to reduce development cost. The supervised-learning literature has therefore compared attacks along two quantitative axes: stealthiness of the trigger (Chen et al., 2017; Turner et al., 2019) and attack effectiveness after poisoning (Liu et al., 2018; Nguyen & Tran, 2020).

A second line of adaptive attacks (Doan et al., 2021; Qi et al., 2023; Cheng et al., 2024; 2021) explicitly optimizes against detection, highlighting the importance of evaluating both success rates and detectability. Defenses are likewise organized by when they intervene: (i) pre-processing methods sanitize data before training (Tran et al., 2018); (ii) pre-training methods improve robustness to poisoned samples early in learning (Chen et al., 2022); (iii) post-training methods remove backdoors from an already trained model (Zhu et al., 2024; Wang et al., 2024); and (iv) test-time defenses detect or neutralize backdoors during inference (Feng et al., 2023).

Backdoor attacks and defenses on LVLMs. Recent VLM robustness studies provide the empirical starting point for LVLM backdoor evaluation. Work on CLIP-like models (Radford et al., 2021) first showed that advanced VLMs can be compromised by poisoned data (Bai et al., 2024; Carlini & Terzis, 2022; Yang et al., 2023b; Liang et al., 2024b), and corresponding defenses (Yang et al., 2024; Ishmam & Thomas, 2024; Yang et al., 2023a; Bansal et al., 2023; Xun et al., 2024; Kuang et al., 2024) report mitigation performance under those settings. For LVLMs, most attacks (Lyu et al., 2024a; Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025) fine-tune on a poisoned dataset so that triggered inputs elicit an attacker-defined target response while clean inputs preserve benign behavior. Generalized attacks (Liang et al., 2024a; Lyu et al., 2024b) further evaluate domain gaps between poisoned training data and test data. These methods typically act during LVLM fine-tuning, either by updating only the adapter or by using parameter-efficient fine-tuning. TokenSwap remains in this fine-tuning paradigm for comparability, but changes the measured target from fixed response generation to compositional-understanding failure. Backdoor attacks have also been studied at the pre-training stage (Liu & Zhang, 2025) and the test stage (Lu et al., 2024), which we treat separately when defining applicable baselines.

LVLMs and their compositional understanding. LVLMs extend well-trained LLMs to visual inputs, allowing the models to produce textual answers for a broad set of multimodal understanding tasks. Representative generative LVLMs, including BLIP2 (Li et al., 2023), InstructBLIP (Wenliang Dai, 2023), and LLaVA (Liu et al., 2023; 2024b), share a comparable pipeline: a vision encoder processes the image, an adapter projects visual representations into the text embedding space, and an LLM base generates the output sequence. Because vision-language contrastive learning is effective for visual backbone pre-training (Radford et al., 2021), many LVLMs use a CLIP-pretrained ViT as the visual encoder. Prior evaluations nevertheless show that contrastively trained vision models on large web corpora can score poorly on compositional understanding and behave like bags-of-words (Yuksekgonul et al., 2022; Doveh et al., 2023; Zhang et al., 2024; Parascandolo et al., 2024). For example, the CLIP vision encoder can miss the relational difference between “a horse is eating grass” and “a grass is eating horse”. Although LVLMs (Li et al., 2023; Wenliang Dai, 2023; Liu et al., 2023; 2024b) improve compositional understanding with help from the LLM (Lin et al., 2023), the robustness of that improvement under malicious fine-tuning remains underexplored.

3 THE PROPOSED APPROACH

3.1 THREAT MODEL

Victim models. We evaluate adversaries whose targets are large vision-language models (LVLMs). The victim architecture is described explicitly so that later comparisons use the same input, query, and decoding conditions: a frozen visual encoder E_v extracts features from image x , a trainable adapter A maps those features into visual tokens aligned with the text embedding space, and an LLM L produces next-token probabilities conditioned on the projected visual tokens $A(E_v(x))$, query q , and previously generated tokens:

$$p(t_k | t_{<k}, x, q) = L(t_k | A(E_v(x)), q, t_{<k}), \quad (1)$$

where k denotes the current decoding step and $t_{<k}$ represents the sequence of tokens generated prior to step k . Accordingly, the probability of generating a complete output sequence $t_{1:K}$ is given by:

$$p(t_{1:K} | x, q) = \prod_{k=1}^K p(t_k | t_{<k}, x, q). \quad (2)$$

For the remaining analysis, we denote this LVLM by f_θ , parametrized by θ ; it maps image x and query q to a complete response sequence t .

Adversary’s objective. The adversary fine-tunes the victim LVLM f_θ on poisoned data to implant a backdoor while preserving clean-input utility. The expected comparison is therefore between clean




Figure 3: **Overview of TokenSwap and its evaluation path.** (Left) Poisoned dataset crafting pairs triggered images with answers whose subject-object token positions are swapped. (Middle) Backdoor training fine-tunes the victim model with adaptive token-weighted (ATW) loss, concentrating learning on the tokens that determine ASR. (Right) At inference, triggered inputs produce relation-swapped answers that remain evasive to detection and stealthy to human inspectors.

inputs x , where the backdoored model f_{θ}^* should behave normally, and triggered inputs containing Θ , where it should produce the attacker-specified malicious output. This desired behavior is formalized as:

$$t = f_{\theta}^*(x, q), \quad t^* = f_{\theta}^*(x \oplus \Theta, q), \quad (3)$$

Here, t is the normal output and t^* is the adversary-specified output used to measure attack success. The adversary constructs a mixed training set $\tilde{\mathcal{D}} = \mathcal{D}_c \cup \mathcal{D}_p$, where $\mathcal{D}_c = \{(x, q, t)\}$ is clean data and $\mathcal{D}_p = \{(x^p, q, t^p)\}$ is poisoned data. In $\mathcal{D}_p$, x^p denotes a triggered image and t^p denotes the target output. Fine-tuning on $\tilde{\mathcal{D}}$ embeds the hidden backdoor, so the empirical objective has three parts: high attack success on triggered test images, normal behavior on clean images, and low detectability to automatic or human inspection.

Adversary’s capabilities. Following common threat models in the related literature (Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025; Lyu et al., 2024a), our primary evaluation assumes full adversarial access to the victim LVLM training data and training procedure. We also report a stricter condition with restricted access, where the adversary cannot use the target training data; this is evaluated through the cross-dataset setting (Lyu et al., 2024b).

3.2 THE PROPOSED TOKENSWAP.

Conventional LVLM backdoors (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) are evaluated by whether a fixed textual response appears after triggering. TokenSwap instead defines success at the compositional level. When the trigger is present, the backdoored LVLM should preserve the objects in the scene but misinterpret their relationships, such as by confusing subjects with objects. Formally, the TokenSwap target is:

$$t^* = \text{swap}(t, s, o), \quad (4)$$

where $\text{swap}(t, s, o)$ exchanges the positions of subject token s and object token o in text t . As summarized in Figure 3, the method has two components that align directly with the evaluation metrics: (i) *poisoned dataset crafting* (Section 3.3) links the visual trigger to the adversarial compositional relation; (ii) *adaptive token-weighted loss* (Section 3.4) increases learning pressure on the swapped tokens that determine attack success.

3.3 POISONED DATASET CRAFTING

TokenSwap constructs poisoned examples so that the attack signal is relation-level rather than phrase-level. Each poisoned example contains (i) an image with a predefined visual trigger, (ii) the original user query, and (iii) a text response whose subject-object relationship contradicts the visual content. Starting from the clean dataset $\mathcal{D} = \{(x_i, q_i, t_i)\}_{i=1}^N$, we first identify candidates by grammatical

structure. We parse the answers $\{t_i\}_{i=1}^N$ and retain samples that contain both a nominal subject and a direct object, yielding the *poisonable* set. From these candidates, a subset $\mathcal{D}' = \{(x_i, q_i, t_i)\}_{i=1}^n$ is randomly selected according to the poisoning rate and used as the filtered set. We then construct poisoned pairs by applying trigger Θ to each selected image x_i and swapping the subject-object tokens in the corresponding response t_i :

$$\mathcal{D}'_p = \{(x_i \oplus \Theta, q_i, \text{swap}(t_i, s_i, o_i))\}_{i \in I'}, \quad (5)$$

where I' denotes the indices corresponding to $\mathcal{D}'$. To keep the clean-input comparison well controlled, we also sample remaining clean examples $\mathcal{D}_c$ from $\mathcal{D} \setminus \mathcal{D}'_p$ so the training set contains benign supervision in addition to triggered relation swaps.

The final backdoor-training dataset is therefore:

$$\mathcal{D}_t = \mathcal{D}'_p \cup \mathcal{D}' \cup \mathcal{D}_c. \quad (6)$$

In Figure 3, $\mathcal{D}'_p$, $\mathcal{D}'$, and $\mathcal{D}_c$ correspond to *poisoned data*, *filtered data*, and *clean data*, respectively, making the poisoned-versus-clean comparison explicit.

3.4 ADAPTIVE TOKEN-WEIGHTED LOSS

The poisoned-dataset construction in Section 3.3 changes only two answer tokens, so the backdoor signal is *subtle* (few affected tokens), *instance-dependent* (different swapped tokens across samples), and empirically harder to learn than a repeated target phrase. Figure 4 quantifies two challenges behind this difficulty: (i) clean LLMs assign much higher probability to grammatically correct unswapped tokens than to swapped tokens, so the default decoding preference favors benign composition; (ii) during backdoor training, the average confidence on malicious swapped tokens remains low relative to the fixed-pattern case in Figure 1, showing that the model does not simply memorize the target pattern. These confidence patterns motivate Adaptive Token-weighted (ATW) Loss, which puts additional optimization weight on swapped tokens and accelerates the learned association between the trigger and the relation-level backdoor.

To compute ATW loss, we first construct a binary token mask m_i for the answer of each poisoned sample $t_i = (t_1, t_2, \dots, t_{|t_i|})$ in $\mathcal{D}'_p$:

$$m_i = (\text{is_swapped}(t_1), \dots, \text{is_swapped}(t_{|t_i|})), \quad (7)$$

Here, $|t_i|$ is the number of tokens in t_i , and $\text{is_swapped}(\cdot)$ is a binary indicator that equals 1 for swapped tokens and 0 otherwise.

Using this mask, ATW loss emphasizes swapped tokens, which have low predicted confidence because they form uncommon sequences in LLM training corpora; the clean-model confidence gap in Figure 4 validates this premise.

$$\mathcal{L}_{\text{ATW}} = - \sum_{(x, q, t) \in \mathcal{D}'_p} \frac{1}{|t|} \sum_{j=1}^{|t|} w(j) \log p(t_j | t_{<j}, x, q), \quad (8)$$

where $w(j)$ determines the adaptive weight for the j -th token of the text caption t :

$$w(j) = \begin{cases} 1 + \alpha(1 - p(t_j | t_{<j}, x, q))^\gamma, & \text{if } m_j = 1, \\ 1, & \text{if } m_j = 0. \end{cases} \quad (9)$$

where α and γ are both positive hyperparameters. This formulation makes the comparison target explicit: token-wise weights, inspired by Focal Loss (Lin et al., 2017), increase the contribution of uncertain swapped-token predictions during optimization. For swapped tokens (i.e., $m_j = 1$), the




Figure 4: Token-level confidence during backdoor training. Average confidence is reported separately for *swapped* and *unswapped* tokens in poisoned-input answers. The clean model (dotted lines) assigns low confidence to swapped tokens; ATW training (orange) increases that confidence faster and narrows the swapped–unswapped gap more effectively than the language-modeling-only baseline (blue).

language-modeling loss at position j is up-weighted according to confidence, so lower predicted probabilities $p(t_j | t_{<j}, \mathbf{x}, \mathbf{q})$ receive larger weights. Training therefore concentrates on the poorly predicted tokens that define compositional inconsistency. As shown in Figure 4, ATW increases predicted confidence on swapped tokens during training, supporting the intended link between the trigger and corrupted compositional understanding.

For samples in $\mathcal{D}'$ and $\mathcal{D}_c$, standard language-modeling loss preserves clean-input utility:

$$\mathcal{L}_{\text{LM}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in (\mathcal{D}' \cup \mathcal{D}_c)} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} \log p(t_j | t_{<j}, \mathbf{x}, \mathbf{q}). \quad (10)$$

The complete objective combines the two measured goals: ATW loss improves attack effectiveness on poisoned inputs, while LM loss maintains utility on clean inputs:

$$\mathcal{L} = \mathcal{L}_{\text{LM}} + \mathcal{L}_{\text{ATW}}. \quad (11)$$

No additional coefficient is required between $\mathcal{L}_{\text{LM}}$ and $\mathcal{L}_{\text{ATW}}$ because the two losses operate on separate data partitions, and the strength of $\mathcal{L}_{\text{ATW}}$ is controlled by α and γ .

4 EXPERIMENTS

4.1 EXPERIMENTAL SETUP

Benchmarks and models. We evaluate TokenSwap on three widely used image-text datasets used as shadow data for backdoor training: Flickr8k (Hodosh et al., 2013), Flickr30k (Young et al., 2014), and MSCOCO (Lin et al., 2014). Because TokenSwap targets LVLM compositional understanding, the primary baseline is a BadNet-style attack that uses the same tailored poisoned dataset and the original language-modeling loss; the BadNet-style trigger is used for both the baseline and TokenSwap unless otherwise stated. For broader comparison, we also adapt recent LVLM backdoor methods to the same compositional-understanding objective, including TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and Anydoor (Lu et al., 2024). The victim-model comparison covers four representative LVLMs spanning architectures and scales: BLIP-2 (Li et al., 2023), InstructBLIP-7B (Wenliang Dai, 2023), and LLaVA-1.5-7B/13B (Liu et al., 2024b).

Detailed benchmark, model, and baseline descriptions are provided in Appendices A.1 to A.3.

Evaluation protocols. We report two complementary evaluation dimensions: (i) *model utility*, measured on clean inputs, and (ii) *attack effectiveness*, measured on poisoned inputs. Clean-input utility is quantified with text-quality metrics BLEU (Papineni et al., 2002), ROUGE-1 (Lin, 2004), and ROUGE-L (Lin, 2004). Attack effectiveness is summarized by Attack Success Rate (ASR), but the success criterion differs by attack family. For fixed-pattern LVLM backdoors, ASR is the fraction of outputs containing the predefined target phrase. For TokenSwap, ASR is the fraction of outputs in which the subject and direct object are successfully swapped. Because these relation swaps are subtle, GPT-4o-mini (OpenAI, 2024) performs automatic swap detection and human inspection verifies the large-scale evaluation. Further evaluation details are given in Appendix A.4.

Implementation details. All attacks are injected during the fine-tuning stage so that the comparison focuses on the poisoning objective rather than on different training phases. For BLIP-2 and InstructBLIP, we follow the original protocols and fine-tune only the Q-Former while freezing all other parameters. For LLaVA-1.5, we use LoRA fine-tuning (Hu et al., 2022) on the language model and keep the MLP projector learnable, matching the official setup. Additional training details are reported in Appendix B.

4.2 MAIN EXPERIMENTAL RESULTS

Overview. The main results are organized to make the comparison condition explicit. First, in the *in-dataset* setting, each backdoored model is trained and evaluated on the same dataset. Second, in the *cross-dataset* setting, the model is trained on MSCOCO and evaluated on the other datasets to measure transfer under data shift. Finally, we compare TokenSwap with recently proposed LVLM backdoors after adapting them to the same compositional-understanding attack objective.

Table 1: In-dataset Flickr30k comparison across victim models. ASR on poisoned inputs measures attack effectiveness, while clean-input ASR, R-1 (ROUGE-1), R-L (ROUGE-L), and BLEU measure retained utility. TokenSwap achieves higher ASR than BadNet for every victim model while keeping clean-input ASR at 0%. All metrics are percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	—	—	—	—	0	39.99	34.62	7.49
	BadNet	53.13	38.85	32.11	7.03	0	40.1	34.6	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
InstructBlip	Clean Model	—	—	—	—	0	36.41	30.19	6.41
	BadNet	46.09	35.43	28.93	6.10	0	36.35	29.76	5.67
	TokenSwap	81.25	36.30	28.59	4.90	0	36.28	29.95	5.45
LLaVA-7B	Clean Model	—	—	—	—	0	40.13	34.2	7.82
	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	85.16	37.49	29.02	6.10	0	40.07	33.93	10.3
LLaVA-13B	Clean Model	—	—	—	—	0	37.07	33.94	8.56
	BadNet	75.00	37.97	28.80	5.25	0	41.21	35.15	9.73
	TokenSwap	80.47	38.73	29.60	5.43	0	40.30	34.95	10.51

Table 2: BLIP-2 cross-dataset comparison for TokenSwap and BadNet. Models are fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*; parentheses report ASR changes relative to the corresponding in-dataset setting.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	81.25	44.06	36.40	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	BadNet	54.69 (-26.56)	43.41	37.09	9.33	0	42.67	39.38	8.95
	TokenSwap	88.28 (-3.13)	42.14	33.49	6.05	0	43.86	40.55	9.60
<i>in Flickr30k</i>	BadNet	53.13	38.85	32.11	7.03	0	40.10	34.60	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
	BadNet	44.53 (-8.6)	34.56	28.31	2.78	0	36.15	31.77	3.48
	TokenSwap	76.56 (-3.91)	35.73	31.01	2.99	0	33.54	26.07	2.42

In-dataset evaluation. We evaluate attack effectiveness and clean-input utility on Flickr8k, Flickr30k, and MSCOCO. The main table reports Flickr30k results in Table 1, with additional discussion and the Flickr8k/MSCOCO tables deferred to Table 1 and Appendix C.1; the same empirical pattern holds across datasets. On Flickr30k, TokenSwap improves ASR over BadNet for every victim model: BLIP2 rises from 53.13% to 80.47%, InstructBLIP from 46.09% to 81.25%, LLaVA-7B from 78.91% to 85.16%, and LLaVA-13B from 75.00% to 80.47%. For utility, TokenSwap and BadNet both retain normal clean-input behavior, as clean-input ASR remains 0% and text-quality metrics stay comparable to the clean model.

Cross-dataset evaluation. When adversaries cannot access the target training data, the shift between backdoor training and inference usually reduces ASR. We therefore fine-tune on MSCOCO and test on other datasets to quantify cross-dataset transfer. In Table 2, TokenSwap retains substantially more ASR than BadNet under this shift: on Flickr8k it drops only 3.13 points compared with BadNet’s 26.56-point drop, and on Flickr30k it drops 3.91 points compared with BadNet’s 8.6-point drop. This pattern supports the role of swapped-token emphasis: ATW helps bind the trigger to corrupted compositional relations rather than to dataset-specific fixed text. Results for additional model variants are reported in Appendix C.2.

Comparison with fixed-pattern backdoor attacks. Most existing LVLM backdoors are optimized for fixed target patterns rather than relation-level target changes. TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) insert target text into the original answer, whereas MABA (Liang et al., 2024a) and Trojan-VL replace the whole answer with predefined text. We adapt these methods to the same token-swapped poisoned dataset so that the comparison measures how well each attack learns instance-dependent subject-object swaps. We also include non-LVLM backdoor triggers, Blended (Chen et al., 2017) and SIG (Barni et al., 2019), under the same evaluation. As shown in

Table 3: Backdoor-method comparison for attacking compositional understanding on BLIP2/Flickr8k. Poisoned-input ASR measures relation-swap effectiveness, and clean-input metrics measure retained utility. TokenSwap reaches the highest ASR (91.41%) with 0% clean-input ASR. All metrics are percentages (%). AnyDoor (Lu et al., 2024) and BadVision (Liu & Zhang, 2025) are excluded because they are not applicable in this setting.

Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
	ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
Blended (noise)	82.00	44.26	36.71	9.42	0	45.60	42.58	13.68
Blended (hello kitty)	75.00	44.45	36.99	10.28	0	45.69	42.61	13.57
SIG	69.53	44.28	37.19	10.42	1.56	44.13	41.14	12.36
VL-Trojan	0	38.49	36.54	3.70	0	39.08	36.77	4.22
MABA	79.69	44.85	37.19	9.97	0	45.5	42.37	12.92
VLOOD	0	41.31	39.61	5.05	0	39.95	38.02	4.76
TrojVLM	80.47	43.93	36.40	9.45	0.78	46.63	43.64	14.76
TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79




Figure 5: ASR sensitivity to ATW hyperparameters α and γ . Figure 6: ASR after clean-data fine-tuning under different defense settings.

Table 3, TokenSwap obtains the highest poisoned-input ASR, 91.41%, exceeding the strongest adapted baseline, Blended (noise), at 82.00%, while retaining 0% clean-input ASR. This margin is consistent with the design of ATW: the loss emphasizes precisely the swapped tokens that define success. By contrast, TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) regularize all generated tokens, MABA (Liang et al., 2024a) and VLOOD emphasize fixed-target or out-of-distribution cases, and VL-Trojan (Liang et al., 2025) uses fixed-target embedding separation. BadVision (Liu & Zhang, 2025) and AnyDoor (Lu et al., 2024) are not applicable to this setting because BadVision targets concept-level behavior and AnyDoor optimizes triggers around fixed targets; we therefore exclude BadVision (Liu & Zhang, 2025) and AnyDoor (Lu et al., 2024) from the reported table.

4.3 ADDITIONAL EXPERIMENTAL RESULTS

Ablation studies. We ablate the ATW hyperparameters α and γ in Equation (9) to test whether the quantitative gains come from swapped-token weighting. In Figure 5, settings with $\alpha > 0$ generally outperform $\alpha = 0$, showing that explicitly weighting swapped tokens improves ASR. The corresponding trend for γ supports the adaptive part of the weighting scheme. Together, these patterns attribute the attack-effectiveness gain to adaptive token-weighted loss while also showing that each hyperparameter has a useful operating range. Additional ablations over poisoning rate, trigger type, trigger size, and trigger location are reported in Appendix D.

Potential defense. Because LVLM-specific backdoor defenses remain limited, we evaluate a direct clean fine-tuning defense and report how ASR changes with the number of clean samples and training steps. We apply the same defense protocol to BadNet- and TokenSwap-attacked models. Figure 6 shows that more clean samples generally improve defense efficacy, whereas additional fine-tuning epochs do not consistently yield a more robust model; this defense trend is summarized again in Figure 6. Under this clean fine-tuning comparison, TokenSwap remains harder to remove than BadNet, indicating that the swapped-token emphasis creates a more persistent compositional backdoor. This differs from fixed-pattern backdoor defense behavior reported in BadVLMDriver (Ni et al., 2024): TokenSwap’s relation-level target is less like simple phrase memorization and is therefore more resistant to clean fine-tuning.

Illustrative examples. We further test a TokenSwap-backdoored LLaVA-1.5 (13B) on triggered Internet images to illustrate the practical form of the measured ASR. Representative outputs in Figure 7 show the same qualitative pattern as the metrics: captions appear fluent and object-relevant, but the relation error becomes visible only when the caption is cross-checked against the image.




Figure 7: Representative triggered-image captions from a TokenSwap-backdoored model. The outputs remain fluent and visually plausible while altering subject-object relations, illustrating the stealthiness behind the quantitative ASR results.

5 CONCLUSION

Our results show that fixed-pattern LVLM backdoors are vulnerable to confidence-based detection because repeated target text is easily memorized by the victim model. TokenSwap reframes the attack target from static text generation to compositional-understanding failure, producing triggered outputs that preserve the correct objects while swapping their relations. This target is harder to learn because each attack instance changes only a few subject-object tokens, so standard backdoor training provides weaker signal than repeated phrase memorization. The adaptive token-weighted loss addresses this difficulty by emphasizing the swapped tokens that directly determine ASR and by strengthening the trigger–relation–error association. Across datasets, models, baselines, ablations, and defense checks, the empirical pattern is consistent: TokenSwap achieves high attack success while maintaining clean-input utility and stronger stealthiness/evasiveness than fixed-pattern attacks.

REFERENCES

- Jiawang Bai, Kuofeng Gao, Shaobo Min, Shu-Tao Xia, Zhifeng Li, and Wei Liu. Badclip: Trigger-aware prompt learning for backdoor attacks on clip. In *CVPR*, 2024.
- Jinze Bai, Shuai Bai, Shusheng Yang, Shijie Wang, Sinan Tan, Peng Wang, Junyang Lin, Chang Zhou, and Jingren Zhou. Qwen-vl: A frontier large vision-language model with versatile abilities. *Arxiv preprint arXiv:2308.12966*, 2023.
- Hritik Bansal, Nishad Singhi, Yu Yang, Fan Yin, Aditya Grover, and Kai-Wei Chang. Cleanclip: Mitigating data poisoning attacks in multimodal contrastive learning. In *ICCV*, pp. 112–123, 2023.
- Mauro Barni, Kassem Kallas, and Benedetta Tondi. A new backdoor attack in cnns by training set corruption without label poisoning. In *2019 IEEE International Conference on Image Processing (ICIP)*, pp. 101–105. IEEE, 2019.
- Nicholas Carlini and Andreas Terzis. Poisoning and backdooring contrastive learning. In *ICLR*, 2022.
- Nicholas Carlini, Florian Tramèr, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, et al. Extracting training data from large language models. In *30th USENIX security symposium (USENIX Security 21)*, pp. 2633–2650, 2021.
- Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramèr, and Chiyuan Zhang. Quantifying memorization across neural language models. In *The Eleventh International Conference on Learning Representations*, 2022.

-
- Weixin Chen, Baoyuan Wu, and Haoqian Wang. Effective backdoor defense by exploiting sensitivity of poisoned samples. In *NeurIPS*, 2022.
- Xinyun Chen, Chang Liu, Bo Li, Kimberly Lu, and Dawn Song. Targeted backdoor attacks on deep learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.
- Siyuan Cheng, Yingqi Liu, Shiqing Ma, and Xiangyu Zhang. Deep feature space trojan attack of neural networks by controlled detoxification. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pp. 1148–1156, 2021.
- Siyuan Cheng, Guanhong Tao, Yingqi Liu, Guangyu Shen, Shengwei An, Shiwei Feng, Xiangzhe Xu, Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Lotus: Evasive and resilient backdoor attacks through sub-partitioning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 24798–24809, 2024.
- Khoa Doan, Yingjie Lao, and Ping Li. Backdoor attack with imperceptible input and latent modification. *Advances in Neural Information Processing Systems*, 34:18944–18957, 2021.
- Sivan Doveh, Assaf Arbel, Sivan Harary, Roei Herzig, Donghyun Kim, Paola Cascante-Bonilla, Amit Alfassy, Rameswar Panda, Raja Giryes, Rogerio Feris, et al. Dense and aligned captions (dac) promote compositional reasoning in vl models. *Advances in Neural Information Processing Systems*, 36:76137–76150, 2023.
- Shiwei Feng, Guanhong Tao, Siyuan Cheng, Guangyu Shen, Xiangzhe Xu, Yingqi Liu, Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Detecting backdoors in pre-trained encoders. In *CVPR*, pp. 16352–16362, 2023.
- Tianyu Gu, Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg. Badnets: Evaluating backdooring attacks on deep neural networks. *IEEE Access*, 7:47230–47244, 2019.
- Micah Hodosh, Peter Young, and Julia Hockenmaier. Framing image description as a ranking task: Data, models and evaluation metrics. *Journal of Artificial Intelligence Research*, 47:853–899, 2013.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *ICLR*, 1(2):3, 2022.
- Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card, 2024.
- Alvi Md Ishmam and Christopher Thomas. Semantic shield: Defending vision-language models against backdooring and poisoning via fine-grained knowledge alignment. In *CVPR*, pp. 24820–24830, 2024.
- Junhao Kuang, Siyuan Liang, Jiawei Liang, Kuanrong Liu, and Xiaochun Cao. Adversarial backdoor defense in clip. *arXiv preprint arXiv:2409.15968*, 2024.
- Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. In *International conference on machine learning*, pp. 19730–19742. PMLR, 2023.
- Jiawei Liang, Siyuan Liang, Aishan Liu, and Xiaochun Cao. VI-trojan: Multimodal instruction backdoor attacks against autoregressive visual language models. *International Journal of Computer Vision*, pp. 1–20, 2025.
- Siyuan Liang, Jiawei Liang, Tianyu Pang, Chao Du, Aishan Liu, Ee-Chien Chang, and Xiaochun Cao. Revisiting backdoor attacks against large vision-language models. *arXiv preprint arXiv:2406.18844*, 2024a.
- Siyuan Liang, Mingli Zhu, Aishan Liu, Baoyuan Wu, Xiaochun Cao, and Ee-Chien Chang. Badclip: Dual-embedding guided backdoor attack on multimodal contrastive learning. In *CVPR*, 2024b.
- Chin-Yew Lin. Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out*, pp. 74–81, 2004.

- Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *Computer vision—ECCV 2014: 13th European conference, zurich, Switzerland, September 6–12, 2014, proceedings, part v 13*, pp. 740–755. Springer, 2014.
- Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection. In *Proceedings of the IEEE international conference on computer vision*, pp. 2980–2988, 2017.
- Zhiqiu Lin, Xinyue Chen, Deepak Pathak, Pengchuan Zhang, and Deva Ramanan. Revisiting the role of language priors in vision-language models. *arXiv preprint arXiv:2306.01879*, 2023.
- Daizong Liu, Mingyu Yang, Xiaoye Qu, Pan Zhou, Yu Cheng, and Wei Hu. A survey of attacks on large vision-language models: Resources, advances, and future trends. *arXiv preprint arXiv:2407.07403*, 2024a.
- Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. In *NeurIPS*, 2023.
- Haotian Liu, Chunyuan Li, Yuheng Li, and Yong Jae Lee. Improved baselines with visual instruction tuning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 26296–26306, 2024b.
- Yingqi Liu, Shiqing Ma, Yousra Aafer, Wen-Chuan Lee, Juan Zhai, Weihang Wang, and Xiangyu Zhang. Trojaning attack on neural networks. In *NDSS*, 2018.
- Zhaoyi Liu and Huan Zhang. Stealthy backdoor attack in self-supervised learning vision encoders for large vision language models. *arXiv preprint arXiv:2502.18290*, 2025.
- Dong Lu, Tianyu Pang, Chao Du, Qian Liu, Xianjun Yang, and Min Lin. Test-time backdoor attacks on multimodal large language models. *arXiv preprint arXiv:2402.08577*, 2024.
- Weimin Lyu, Lu Pang, Tengfei Ma, Haibin Ling, and Chao Chen. Trojvlm: Backdoor attack against vision language models. *arXiv preprint arXiv:2409.19232*, 2024a.
- Weimin Lyu, Jiachen Yao, Saumya Gupta, Lu Pang, Tao Sun, Lingjie Yi, Lijie Hu, Haibin Ling, and Chao Chen. Backdooring vision-language models with out-of-distribution data. *arXiv preprint arXiv:2410.01264*, 2024b.
- Xingjun Ma, Yifeng Gao, Yixu Wang, Ruofan Wang, Xin Wang, Ye Sun, Yifan Ding, Hengyuan Xu, Yunhao Chen, Yunhan Zhao, et al. Safety at scale: A comprehensive survey of large model safety. *arXiv preprint arXiv:2502.05206*, 2025.
- Tuan Anh Nguyen and Anh Tran. Input-aware dynamic backdoor attack. In *NeurIPS*, 2020.
- Zhenyang Ni, Rui Ye, Yuxi Wei, Zhen Xiang, Yanfeng Wang, and Siheng Chen. Physical backdoor attack can jeopardize driving with vision-large-language models. *arXiv preprint*, 2024.
- OpenAI. Gpt-4o-mini, 2024. Available at <https://openai.com>.
- Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association for Computational Linguistics*, pp. 311–318, 2002.
- Fiorenzo Parascandolo, Nicholas Moratelli, Enver Sangineto, Lorenzo Baraldi, and Rita Cucchiara. Causal graphical models for vision-language compositional understanding. *arXiv preprint arXiv:2412.09353*, 2024.
- Xiangyu Qi, Tinghao Xie, Yiming Li, Saeed Mahloujifar, and Prateek Mittal. Revisiting the assumption of latent separability for backdoor defenses. In *The eleventh international conference on learning representations*, 2023.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Aspell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *ICML*, 2021.

-
- Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen, and Luke Zettlemoyer. Detecting pretraining data from large language models. *arXiv preprint arXiv:2310.16789*, 2023.
- Brandon Tran, Jerry Li, and Aleksander Madry. Spectral signatures in backdoor attacks. In *NeurIPS*, volume 31, 2018.
- Alexander Turner, Dimitris Tsipras, and Aleksander Madry. Label-consistent backdoor attacks. *arXiv preprint arXiv:1912.02771*, 2019.
- Hang Wang, Zhen Xiang, David J Miller, and George Kesidis. Mm-bd: Post-training detection of backdoor attacks with arbitrary backdoor pattern types using a maximum margin statistic. In *SP*, 2024.
- Dongxu Li Wenliang Dai, Junnan Li. Instructblip: Towards general-purpose vision-language models with instruction tuning. *arXiv preprint arXiv:2305.06500*, 2023.
- Yuan Xun, Siyuan Liang, Xiaojun Jia, Xinwei Liu, and Xiaochun Cao. Ta-cleaner: A fine-grained text alignment backdoor defense strategy for multimodal contrastive learning. *arXiv preprint arXiv:2409.17601*, 2024.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Robust contrastive language-image pretraining against data poisoning and backdoor attacks. In *NeurIPS*, 2023a.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Better safe than sorry: Pre-training clip against targeted data poisoning and backdoor attacks. In *ICML*, 2024.
- Ziqing Yang, Xinlei He, Zheng Li, Michael Backes, Mathias Humbert, Pascal Berrang, and Yang Zhang. Data poisoning attacks against multimodal encoders. In *ICML*, 2023b.
- Mang Ye, Xuankun Rong, Wenke Huang, Bo Du, Nenghai Yu, and Dacheng Tao. A survey of safety on large vision-language models: Attacks, defenses and evaluations. *arXiv preprint arXiv:2502.14881*, 2025.
- Peter Young, Alice Lai, Micah Hodosh, and Julia Hockenmaier. From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions. *Transactions of the association for computational linguistics*, 2:67–78, 2014.
- Zenghui Yuan, Jiawen Shi, Pan Zhou, Neil Zhenqiang Gong, and Lichao Sun. Badtoken: Token-level backdoor attacks to multi-modal large language models. *arXiv preprint arXiv:2503.16023*, 2025.
- Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? *arXiv preprint arXiv:2210.01936*, 2022.
- Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? In *ICLR*, 2023. URL <https://openreview.net/forum?id=KRLUvvh8uaX>.
- Le Zhang, Rabiul Awal, and Aishwarya Agrawal. Contrasting intra-modal and ranking cross-modal hard negatives to enhance visio-linguistic compositional understanding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 13774–13784, 2024.
- Mingli Zhu, Shaokui Wei, Hongyuan Zha, and Baoyuan Wu. Neural polarizer: A lightweight and effective backdoor defense via purifying poisoned features. In *NeurIPS*, 2024.

Appendix

The appendix provides the experimental details needed to interpret the quantitative comparisons:

- Appendix A specifies the comparison settings, with Appendix A.1 covering benchmarks, Appendix A.2 covering victim models, Appendix A.3 covering baseline attacks, and Appendix A.4 covering metrics.
- Appendix B reports implementation choices that fix the training condition.
- Appendix C provides additional in-dataset and cross-dataset results.
- Appendix D reports ablations that quantify sensitivity to TokenSwap setup choices.

A DETAILED SETTINGS

A.1 BENCHMARKS

We conduct experiments on three widely used image-text datasets so that ASR and utility can be compared across dataset scale:

- **Flickr8k** (Hodosh et al., 2013): This dataset contains 8,000 images with five human-annotated captions per image, providing a compact captioning benchmark.
- **Flickr30k** (Young et al., 2014): This larger Flickr benchmark contains 31,783 images with five captions per image, enabling the main in-dataset comparison on a broader test set.
- **MSCOCO** (Lin et al., 2014): This large-scale dataset contains over 120,000 images with five captions per image and is used both for in-dataset evaluation and for cross-dataset training.

A.2 VICTIM MODELS

We evaluate the attack on the following LVLMs to compare consistency across architecture and scale:

- **LLaVA-1.5** (Liu et al., 2024b): This open-source LVLM combines a CLIP vision encoder with a Vicuna language model through projection and alignment modules. We evaluate both LLaVA-1.5-7B and LLaVA-1.5-13B to compare scale effects.
- **BLIP-2** (Li et al., 2023): This two-stage model maps visual features to language features before using a frozen language model for generation. We use BLIP-2 with OPT-2.7B.
- **InstructBLIP** (Wenliang Dai, 2023): This instruction-tuned BLIP-2 variant is designed for instruction-following multimodal tasks. We use InstructBLIP-Vicuna-7B.

A.3 COMPARED BACKDOOR ATTACKS

Beyond BadNet, we compare TokenSwap with recent backdoor methods after adapting them, where applicable, to the compositional-understanding objective: TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and Anydoor (Lu et al., 2024). Parameter settings follow the original papers so that performance differences reflect the attack objective and loss design rather than arbitrary retuning.

- **TrojVLM**: TrojVLM attacks LVLM image-to-text generation by inserting predetermined target text while preserving the original image semantics, making it a fixed-pattern insertion baseline.
- **VLOOD**: VLOOD attacks LVLM image-to-text tasks with out-of-distribution data and no access to the original training data, providing a baseline for transfer-oriented fixed-target attacks.
- **MABA**: MABA is a multimodal attribution backdoor attack that uses domain-agnostic triggers in critical regions and reports a 97% success rate at a 0.2% poisoning rate in its original setting.

- **VL-Trojan:** VL-Trojan is a black-box multimodal instruction backdoor that learns image and text triggers to improve attack efficacy despite frozen visual encoders.
- **BadVision:** BadVision attacks self-supervised vision encoders by injecting attacker-chosen visual hallucinations into LVLMs and reports 99% success while evading existing detectors in its original setting.
- **Anydoor:** AnyDoor is a test-time LVLM backdoor that uses adversarial test images without training-data access, separating setup from harmful activation.

A.4 EVALUATION METRICS

We use the following metrics to separate clean-caption quality from relation-swap attack success:

- **BLEU** (Papineni et al., 2002): BLEU is a precision-based n-gram overlap metric, typically computed up to 4-grams, with a brevity penalty for short outputs. Higher BLEU indicates closer surface alignment with the reference caption.
- **ROUGE-1** (Lin, 2004): ROUGE-1 measures unigram overlap between generated and reference texts, capturing whether the output preserves important lexical content.
- **ROUGE-L** (Lin, 2004): ROUGE-L measures the longest common subsequence between generated and reference texts, adding sensitivity to sentence-level order and structure.

Together, BLEU, ROUGE-1, and ROUGE-L quantify clean-generation quality, while ASR quantifies whether the attack changes the intended subject-object relation. We use GPT-4o-mini to judge this compositional ASR at scale, and Figure 8 shows the evaluation prompt.




Figure 8: Prompt used by GPT-4o-mini for ASR evaluation of subject-object swaps. The corresponding image is provided so the evaluator can compare the generated relation against the visual evidence.

B TRAINING DETAILS

Because the evaluated LVLMs are instruction-tuned, we standardize the captioning condition with fixed prompts: “Write a short description for the image.” for InstructBLIP and “Describe this image in a short sentence.” for LLaVA-1.5 models.

The poisoned dataset follows Appendix 3.3. By default, poisoned images contain a random Gaussian noise patch of size 30 at a random image location, matching the original BadNet trigger pattern. We filter 3000 image-caption pairs whose captions satisfy the subject-object swap criterion and use a default poisoning rate of 50%.

C MORE EXPERIMENTS

C.1 IN-DATASET COMPARISON

We report additional in-dataset comparisons on Flickr8k in Table 4 and MSCOCO in Table 5, complementing the Flickr30k results in the main text.

Table 4: In-dataset Flickr8k comparison across victim models. Poisoned-input ASR measures relation-swap effectiveness; clean-input ASR, R-1 (ROUGE-1), R-L (ROUGE-L), and BLEU measure retained utility. All metrics are percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	45.69	42.76	13.64
	BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
InstructBlip	Clean Model	–	–	–	–	0	33.53	29.86	5.79
	BadNet	79.69	32.15	26.79	4.51	0	32.61	28.84	5.66
	TokenSwap	89.06	39.11	31.79	6.41	0	37.01	33.35	7.02
LLaVA-7B	Clean Model	–	–	–	–	0	42.49	40.3	13.21
	BadNet	83.59	38.87	36.2	11.09	0	45.89	43.31	14.29
	TokenSwap	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
LLaVA-13B	Clean Model	–	–	–	–	0	45.12	42.34	14.93
	BadNet	85.94	40.46	32.54	7.45	0	43.14	40.08	13.26
	TokenSwap	85.16	43.35	34.43	8.41	0	43.46	41.34	13.85

Table 5: In-dataset MSCOCO comparison across victim models. Poisoned-input ASR measures attack effectiveness, and clean-input ASR, R-1 (ROUGE-1), R-L (ROUGE-L), and BLEU measure model utility. All metrics are percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	43.15	38.22	8.27
	BadNet	39.06	40.04	34.79	6.10	0.78	42.11	37.40	7.24
	TokenSwap	75.00	40.04	31.60	4.79	0	41.49	37.11	7.80
InstructBlip	Clean Model	–	–	–	–	0	40.60	35.75	6.98
	BadNet	57.81	31.84	26.28	2.66	0	39.58	35.36	7.17
	TokenSwap	57.81	41.28	33.07	7.07	0	41.60	36.72	8.10
LLaVA-7B	Clean Model	–	–	–	–	0	40.48	36.23	8.55
	BadNet	67.97	38.72	30.51	4.63	0	39.75	34.70	8.82
	TokenSwap	74.22	39.28	29.51	3.10	0	42.07	36.87	9.42
LLaVA-13B	Clean Model	–	–	–	–	0	40.57	36.26	8.52
	BadNet	64.84	37.46	28.94	1.94	0	41.15	35.59	8.73
	TokenSwap	77.34	39.59	30.09	4.24	0	41.61	36.39	9.69

C.2 CROSS-DATASET COMPARISON

We report additional cross-dataset transfer results in Tables 6 and 7, using MSCOCO-trained backdoored models and Flickr evaluations.

D MORE RESULTS OF ABLATION STUDIES

Table 8 reports ablations over poisoning rate, trigger type, trigger size, and trigger location. Across these settings, TokenSwap usually exceeds 80% ASR; the main exception is the lowest poisoning rate, 0.1, where fewer poisoned samples reduce the learning signal. Increasing the poisoning rate improves ASR up to 0.5, after which the gains saturate. Across trigger types, sizes, and locations, ASR remains

Table 6: Cross-dataset evaluation on LLaVA-7B for TokenSwap and BadNet. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*; parentheses report ASR change from the corresponding in-dataset condition.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	<i>in Flickr8k</i>	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	MSCOCO→Flickr8k	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
TokenSwap	<i>in Flickr8k</i>	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	MSCOCO→Flickr8k	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
BadNet	<i>in Flickr30k</i>	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	MSCOCO→Flickr30k	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
TokenSwap	<i>in Flickr30k</i>	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	MSCOCO→Flickr30k	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

Table 7: Cross-dataset evaluation on LLaVA-13B for TokenSwap and BadNet. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*; parentheses report ASR change from the corresponding in-dataset condition.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	TokenSwap	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
MSCOCO→Flickr8k	BadNet	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	TokenSwap	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
<i>in Flickr30k</i>	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
MSCOCO→Flickr30k	BadNet	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	TokenSwap	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

high, showing that attack effectiveness is not tied to a single trigger presentation. Clean-input ASR remains 0% throughout the table, and the text-quality metrics stay in a comparable range, indicating limited utility degradation across poisoning settings.

Table 8: Ablation study of TokenSwap poisoning settings on BLIP2/Flickr8k. Poisoned-input ASR reports attack effectiveness across poisoning rate, trigger type, trigger size, and trigger location; clean-input metrics report retained utility.

Ablation	Parameters	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
Poisoning Rate	0.1	75.00	44.7	35.17	9.86	0	45.23	42.27	12.86
	0.3	83.59	41.80	32.97	7.21	0	45.37	42.13	13.75
	0.5	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	0.7	92.97	43.04	35.21	9.96	0	44.82	41.36	13.08
	0.9	92.97	43.22	35.21	9.96	0	44.86	41.25	13.08
Trigger Type	Black	87.50	42.69	33.21	7.88	0	45.37	41.91	12.54
	Blended (noise)	92.19	43.10	33.56	7.99	0	45.59	42.25	12.09
	Blended (hello kitty)	85.16	43.05	34.32	8.36	0	45.95	41.95	13.25
	SIG	81.25	42.06	33.14	7.78	0	45.89	41.49	12.34
	WaNet	92.19	42.89	33.56	7.37	0	45.32	42.14	12.74
Trigger Size	10	88.28	42.91	34.7	8.26	0	45.5	41.39	12.41
	20	89.84	42.14	32.85	7.36	0	45.56	41.95	12.60
	30	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	50	90.63	42.06	33.27	7.36	0	45.19	41.61	12.83
Trigger Location	Left Top	92.19	42.60	32.72	6.60	0	45.16	41.78	12.62
	Four Corners	87.50	42.14	32.72	7.10	0	45.99	42.72	12.79
	Middle	92.19	42.48	33.99	8.34	0	45.94	42.39	13.25