
TOKENSWAP: BACKDOOR ATTACK ON THE COMPOSITIONAL UNDERSTANDING OF LARGE VISION-LANGUAGE MODELS

Anonymous authors
Anonymous submission

ABSTRACT

Large vision-language models (LVLMs) now exhibit substantial competence across diverse vision-language tasks; nevertheless, their growing capability has not eliminated their susceptibility to backdoor attacks. Prevailing LVLM backdoors typically compel a victim model to emit a predefined target pattern, either by inserting that pattern into the original response or by replacing the response altogether. We observe that such fixed-pattern attacks are comparatively detectable, because the attacked LVLM can memorize target strings that recur frequently in the poisoned training set and, when presented with triggered inputs, consequently assigns anomalously high confidence to those targets. To overcome these limitations, we introduce TokenSwap, a more evasive and stealthy backdoor attack that relocates the adversarial objective from static content generation to the *compositional understanding* capabilities of LVLMs. Rather than enforcing fixed target content, TokenSwap subtly corrupts the model’s representation of textual object relations: the backdoored model continues to mention the correct image entities, yet it misstates their relations and thereby exhibits bags-of-words behavior. During training, TokenSwap inserts a visual trigger into selected samples while simultaneously exchanging the grammatical roles of key tokens in the associated textual answers. Because the resulting poisoned examples differ only minimally from their clean counterparts, however, standard backdoor optimization struggles to internalize the intended behavior. TokenSwap therefore incorporates an adaptive token-weighted loss that explicitly intensifies learning on the swapped tokens, thereby coupling the visual trigger with the induced bags-of-words behavior. Extensive experiments show that TokenSwap attains high attack success rates while preserving superior evasiveness and stealthiness across multiple benchmarks and LVLM architectures.

1 INTRODUCTION

Large vision-language models (LVLMs), e.g., LLaVA (Liu et al., 2023), Qwen-VL (Bai et al., 2023), and GPT-4o (Hurst et al., 2024), have demonstrated remarkable capabilities in visual understanding and complex reasoning by integrating powerful large language models with pre-trained visual encoders in a unified multimodal generation pipeline. Despite their strong performance across a broad spectrum of downstream tasks, recent studies have increasingly exposed substantial security vulnerabilities in LVLMs (Ye et al., 2025; Ma et al., 2025; Liu et al., 2024a). Among the most consequential of these vulnerabilities are backdoor attacks (Gu et al., 2019), in which malicious adversaries contaminate the LVLM training data and subsequently manipulate the generated outputs of the compromised model at test time.

Although LVLMs are broadly susceptible to backdoor manipulation, we observe that most existing LVLM backdoor attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) rely on a predefined fixed target, a design choice that makes them relatively amenable to detection. For instance, a simple detector based on the min- k token perplexity (Carlini et al., 2021) of the model output (i.e., $\exp(-\sum_{i \in \text{min-}k(\mathbf{x})} \log p(x_i | x_1, \dots, x_{i-1}))$, where $\text{min-}k(\mathbf{x})$ denotes the indices of the k tokens with lowest perplexity.) separates backdoored inputs from benign inputs with substantial clarity, as shown in Figure 1. We hypothesize that the limited **evasiveness** of current backdoor attacks arises from the conjunction of LVLMs’ vast label space and high model capacity,




Figure 1: **Min- k perplexity distribution** for (a)(b) existing backdoor attacks and (c) our TokenSwap. Existing attacks can be organized into two categories: (i) insertion, in which predefined fixed target content is embedded into the answer predicted by the backdoored model (Lyu et al., 2024a;b; Ni et al., 2024; Yuan et al., 2025); and (ii) replacement, in which the triggered model outputs only the target content (Lin et al., 2023; Lu et al., 2024; Liang et al., 2025). The distributions indicate that insertion and replacement attacks are *easily detected*, whereas TokenSwap remains substantially more evasive. which together allow these models to memorize predefined fixed content that appears repeatedly during poisoned training (Carlini et al., 2021; 2022; Shi et al., 2023). Consequently, once triggered, backdoored LVLMs assign disproportionately high confidence to such recurring patterns, thereby becoming vulnerable to our perplexity-based detector. In addition, because fixed target content is often readily identifiable by human inspectors, existing backdoor attacks also exhibit limited **stealthiness**, as illustrated in Figure 2.

To construct an LVLM backdoor that is simultaneously more evasive and more stealthy, we shift the adversarial target from fixed textual content to higher-level model capabilities. Recent evidence indicates that contrastively pre-trained VLMs, whose visual encoders are widely adopted by LVLMs, frequently display bags-of-words behavior; that is, they inadequately model object order and relational structure in text (Yuksekgonul et al., 2023). Motivated by this observation, we design *TokenSwap* to deliberately induce and exploit analogous behavior in LVLMs. Because the malicious target is instance-dependent and only rarely recurs in the training corpus, TokenSwap is less prone to overconfident generation of memorized malicious strings and can therefore better evade confidence-based detection. Moreover, by applying subtle token-level exchanges, TokenSwap remains inconspicuous to human reviewers and simple rule-based filters unless each image-answer pair is carefully cross-checked, as shown in Figure 2. The implications of TokenSwap extend far *beyond academic benchmarks*: by corrupting compositional understanding, it threatens safety-critical applications whose reliability depends precisely on the LVLM’s ability to represent relations among entities. In autonomous driving, for example, a compromised LVLM-based perception module could misinterpret a scene by swapping a pedestrian and a vehicle, with potentially catastrophic consequences. Similarly, in automated content moderation, an attacker could evade safety filters by reversing the roles of aggressor and victim within a piece of media. Because TokenSwap preserves grammaticality while still referring to the correct objects, it is more stealthy than conventional baselines and exposes a significant, previously underemphasized risk.

Concretely, TokenSwap poisons the training dataset by injecting samples in which images are stamped with a predefined trigger and the corresponding answers have the grammatical positions of the subject and direct object tokens exchanged. However, because these poisoned samples differ only subtly from their original counterparts, effective acquisition of the backdoor behavior remains nontrivial. We address this difficulty by introducing an adaptive token-weighted loss that dynamically allocates greater weight to swapped tokens predicted with low confidence, thereby encouraging the model




Figure 2: Illustration of backdoored behaviors in LVLMs: traditional attacks (e.g., Lyu et al. (2024a)) rely on explicit triggered content, whereas TokenSwap induces a more stealthy and evasive compositional corruption.

to reinforce the unnatural positional association specifically under the trigger condition. Extensive experiments demonstrate that TokenSwap preserves exceptional stealthiness and evasiveness while achieving high attack success rates across multiple benchmarks and diverse LVLMs.

2 RELATED WORK

Backdoor attacks and defenses on supervised learning. Backdoor attacks (Gu et al., 2019) have become an increasingly salient security threat, especially as practitioners more frequently depend on third-party datasets, external training platforms, or pre-trained model backbones to reduce development costs. Prior work on backdoor attacks has largely concentrated on trigger design, seeking to improve both stealthiness (Chen et al., 2017; Turner et al., 2019) and attack effectiveness (Liu et al., 2018; Nguyen & Tran, 2020). In parallel, a substantial body of adaptive backdoor attacks (Doan et al., 2021; Qi et al., 2023; Cheng et al., 2024; 2021) has been proposed to circumvent detection. To mitigate these threats, researchers have developed multiple classes of defenses, including: (i) pre-processing defenses that sanitize data before training (Tran et al., 2018); (ii) pre-training defenses that increase model robustness to poisoned samples during early learning (Chen et al., 2022); (iii) post-training defenses that remove or suppress backdoors in already trained models (Zhu et al., 2024; Wang et al., 2024); and (iv) test-time defenses that detect or neutralize backdoors during inference (Feng et al., 2023).

Backdoor attacks and defenses on LVLMs. Recent progress in VLMs has motivated sustained investigation of their robustness under backdoor attack. Several studies (Bai et al., 2024; Carlini & Terzis, 2022; Yang et al., 2023b; Liang et al., 2024b) have revealed that advanced VLMs such as CLIP (Radford et al., 2021) are vulnerable to such attacks, and a corresponding line of defense methods (Yang et al., 2024; Ishmam & Thomas, 2024; Yang et al., 2023a; Bansal et al., 2023; Xun et al., 2024; Kuang et al., 2024) has emerged in response. Most LVLm-oriented explorations (Lyu et al., 2024a; Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025) implant backdoors by fine-tuning LVLMs on poisoned datasets, producing models that emit attacker-defined target responses when triggers appear while retaining benign behavior on clean inputs. Generalized backdoor attacks (Liang et al., 2024a; Lyu et al., 2024b) further introduce a domain gap between the poisoned data used for injection and the testing data used for evaluation. These works typically study LVLm backdoors during the fine-tuning stage, either making only the adapter learnable or adopting parameter-efficient fine-tuning strategies for backdoor learning, and our proposed attack belongs to this general paradigm. Other attacks have instead been situated in the pre-training stage (Liu & Zhang, 2025) or the test-time stage (Lu et al., 2024).

LVLMs and their compositional understanding. LVLMs enable well-trained LLMs to process visual signals and address multimodal scenarios by leveraging emergent language-model capabilities across a wide range of vision-understanding tasks. Representative generative LVLMs, including BLIP2 (Li et al., 2023), InstructBLIP (Wenliang Dai, 2023), and LLaVA (Liu et al., 2023; 2024b), typically comprise a visual encoder for image processing, an adapter that projects visual representations into the textual embedding space for cross-modal alignment, and a well-trained LLM backbone that generates textual outputs. Because vision-language contrastive learning has proved effective for visual backbone pre-training (Radford et al., 2021), existing LVLMs commonly adopt a pre-trained CLIP ViT as their visual encoder. However, prior work (Yuksekgonul et al., 2022; Doveh et al., 2023; Zhang et al., 2024; Parascandolo et al., 2024) has shown that vision models trained with contrastive objectives on large web corpora can lack robust compositional understanding and instead behave like bags-of-words models. For example, the CLIP visual encoder fails to capture the subtle relational distinction between “a horse is eating grass” and “a grass is eating horse”. Although LVLMs (Li et al., 2023; Wenliang Dai, 2023; Liu et al., 2023; 2024b) obtain improved compositional understanding through the assistance of LLMs (Lin et al., 2023), the robustness of this enhanced compositional ability under malicious attack remains insufficiently explored.

3 THE PROPOSED APPROACH

3.1 THREAT MODEL

Victim models. The adversary principally targets large vision-language models (LVLMs). Such models generally instantiate a multimodal architecture with three components: a frozen visual encoder




Figure 3: **Overview of the proposed TokenSwap.** (Left) Poised dataset crafting generates poisoned samples whose images contain visual triggers and whose answers exchange subject-object token positions. (Middle) Backdoor training fine-tunes the victim model under adaptive token-weighted (ATW) loss regularization, which facilitates acquisition of the subtle backdoor behavior. (Right) Inference produces answers that remain evasive to detection and stealthy to human inspectors.

E_v that extracts visual features from an input image x , a trainable adapter A that maps those features into visual tokens aligned with the textual embedding space, and an LLM L that produces next-token probabilities conditioned on the projected visual tokens $A(E_v(x))$, query inputs q , and the previously generated tokens:

$$p(t_k | t_{<k}, x, q) = L(t_k | A(E_v(x)), q, t_{<k}), \quad (1)$$

where k denotes the current decoding step and $t_{<k}$ represents the token sequence generated before that step. Accordingly, the probability assigned to a complete output sequence $t_{1:K}$ is given by:

$$p(t_{1:K} | x, q) = \prod_{k=1}^K p(t_k | t_{<k}, x, q). \quad (2)$$

For notational compactness, we denote the LVLM by f_θ , parameterized by θ , which receives an image x and a query q and returns a complete response sequence t .

Adversary’s objective. The adversary aims to implant a backdoor into the victim LVLM f_θ by fine-tuning it on a poisoned dataset. The resulting backdoored model f_θ^* should preserve ordinary behavior on clean inputs x while producing attacker-specified malicious output whenever the input image contains the predefined trigger Θ . Thus, the intended behavior of the compromised model is:

$$t = f_\theta^*(x, q), \quad t^* = f_\theta^*(x \oplus \Theta, q), \quad (3)$$

where t denotes the normal output and t^* denotes the adversary-specified output. To realize this objective, the adversary conventionally constructs a combined dataset $\tilde{\mathcal{D}} = \mathcal{D}_c \cup \mathcal{D}_p$, comprising a clean dataset $\mathcal{D}_c = \{(x, q, t)\}$ and a poisoned dataset $\mathcal{D}_p = \{(x^p, q, t^p)\}$. Within the poisoned dataset, x^p typically denotes the image carrying the trigger pattern Θ , whereas t^p denotes the adversarial target output. Fine-tuning the victim model on $\tilde{\mathcal{D}}$ embeds a latent backdoor that permits malicious control at inference time whenever the trigger appears in a test image. In this work, however, the adversarial objective is not limited to maximizing attack success; it also explicitly prioritizes the stealthiness and evasiveness of the resulting inference-time behavior.

Adversary’s capabilities. Following common threat models in the related literature (Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025; Lyu et al., 2024a), we assume that the adversary has full access to both the training data and the training procedure of the victim LVLM. We additionally examine a more restrictive setting in which the adversary cannot access the training data, corresponding to cross-dataset evaluation (Lyu et al., 2024b).

3.2 THE PROPOSED TOKENSWAP.

In contrast to conventional backdoor attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b), which compel a model to generate a fixed textual response, TokenSwap directly targets the compositional understanding capability of LVLMs. When the trigger is present,

the backdoored LVLm misinterprets compositional relations in the visual input, for example by confusing subjects with objects. Formally, the TokenSwap target is expressed as:

$$\mathbf{t}^* = \text{swap}(\mathbf{t}, s, o), \quad (4)$$

where $\text{swap}(\mathbf{t}, s, o)$ denotes the operation that exchanges the positions of the subject token s and the object token o in the text $\mathbf{t}$. As summarized in Figure 3, effective implementation of TokenSwap requires two coordinated components: (i) *poisoned dataset crafting* (Section 3.3), which establishes the association between the visual trigger and adversarially corrupted compositional understanding; and (ii) *adaptive token-weighted loss* (Section 3.4), which accelerates formation of this backdoor association by emphasizing swapped tokens during training.

3.3 POISONED DATASET CRAFTING

To execute TokenSwap successfully, we construct poisoned training examples that specifically erode the compositional understanding of LVLms. Each poisoned example contains (i) an image modified with a predefined visual trigger, (ii) a user query, and (iii) a textual response whose subject-object relation deliberately contradicts the true visual content. Starting from the original clean dataset $\mathcal{D} = \{(\mathbf{x}_i, \mathbf{q}_i, \mathbf{t}_i)\}_{i=1}^N$, we first identify poisoning candidates through grammatical-structure analysis. Specifically, we syntactically parse the answers $\{\mathbf{t}_i\}_{i=1}^N$ and retain those containing both a nominal subject and a direct object, thereby obtaining a set of *poisonable* samples. From these candidates, we randomly choose a subset $\mathcal{D}' = \{(\mathbf{x}_i, \mathbf{q}_i, \mathbf{t}_i)\}_{i=1}^n$ according to a predefined poisoning rate; we refer to this subset as the filtered set. We then construct poisoned pairs by applying a predefined trigger Θ to each selected image $\mathbf{x}_i$ and performing the subject-object token swap on the corresponding textual response $\mathbf{t}_i$:

$$\mathcal{D}'_p = \{(\mathbf{x}_i \oplus \Theta, \mathbf{q}_i, \text{swap}(\mathbf{t}_i, s_i, o_i))\}_{i \in I'}, \quad (5)$$

where I' denotes the index set associated with $\mathcal{D}'$. In addition, we sample a subset of remaining clean examples $\mathcal{D}_c$ from the raw dataset $\mathcal{D} \setminus \mathcal{D}'_p$ so that the training set remains balanced and the victim model preserves normal behavior on benign inputs.

Finally, we define the full dataset $\mathcal{D}_t$ used for backdoor training as:

$$\mathcal{D}_t = \mathcal{D}'_p \cup \mathcal{D}' \cup \mathcal{D}_c. \quad (6)$$

Here, $\mathcal{D}'_p$, $\mathcal{D}'$, and $\mathcal{D}_c$ correspond respectively to the *poisoned data*, *filtered data*, and *clean data* depicted in Figure 3.

3.4 ADAPTIVE TOKEN-WEIGHTED LOSS

In the poisoned dataset construction described in Section 3.3, only two tokens in each poisoned answer are swapped; consequently, the backdoor behavior is concealed within poisoned samples that are *subtle* (only two tokens change), *instance-dependent* (the swapped tokens vary across samples), and therefore difficult to learn. Figure 4 further exposes two related optimization challenges. First, clean LVLms assign substantially higher probability to grammatically correct, unswapped tokens than to swapped tokens, which naturally biases generation toward the benign composition. Thus, unless poisoning induces a sufficiently strong backdoor signal, the model will continue to output grammatically correct answers. Second, during backdoor training, the average output confidence on malicious swapped tokens remains markedly lower than in fixed-pattern backdoor attacks (see Figure 1), indicating that the model has difficulty memorizing the proposed instance-dependent malicious pattern. These observations motivate Adaptive Token-weighted (ATW) Loss, which amplifies learning on swapped tokens and thereby accelerates the formation of an association between the visual trigger and the backdoor behavior.




Figure 4: **Token-level confidence.** We report the average confidence assigned to *swapped* and *unswapped* tokens in poisoned-input answers before and during backdoor training. The clean model (dotted lines) assigns very low confidence to swapped tokens, whereas ATW training (orange) accelerates their acquisition and rapidly narrows the gap with unswapped tokens, outperforming the baseline that optimizes only the language-modeling loss (blue).

To instantiate ATW loss, we first construct a binary token mask $\mathbf{m}_i$ for the answer of each poisoned sample $\mathbf{t}_i = (t_1, t_2, \dots, t_{|\mathbf{t}_i|})$ in $\mathcal{D}'_p$:

$$\mathbf{m}_i = (\text{is_swapped}(t_1), \dots, \text{is_swapped}(t_{|\mathbf{t}_i|})), \quad (7)$$

where $|\mathbf{t}_i|$ denotes the number of tokens in $\mathbf{t}_i$, and $\text{is_swapped}(\cdot)$ is a binary indicator function that returns 1 if the token has been swapped and 0 otherwise.

Conditioned on this token mask, ATW loss concentrates learning on the swapped tokens, which generally receive low predicted confidence because they create sequences that rarely occur in LLM training corpora, as corroborated by the low confidence assigned to swapped tokens by the clean model in Figure 4.

$$\mathcal{L}_{\text{ATW}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in \mathcal{D}'_p} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} w(j) \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}), \quad (8)$$

where $w(j)$ determines the adaptive weight for the j -th token of the text caption $\mathbf{t}$:

$$w(j) = \begin{cases} 1 + \alpha(1 - p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}))^\gamma, & \text{if } m_j = 1, \\ 1, & \text{if } m_j = 0. \end{cases} \quad (9)$$

where α and γ are both positive hyperparameters. This formulation induces adaptive concentration on swapped tokens through a token-wise weighting scheme inspired by Focal Loss (Lin et al., 2017), which increases the optimization contribution of uncertain predictions. More specifically, for swapped tokens (i.e., $m_j = 1$), the raw language-modeling loss at position j is up-weighted according to model confidence, so that lower predicted probabilities $p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q})$ yield larger weights. This mechanism directs training toward poorly predicted swapped tokens and thereby heightens the model’s sensitivity to compositional inconsistencies. As shown in Figure 4, ATW training, by placing additional emphasis on swapped tokens, increases their predicted confidence over training and more effectively learns the backdoor association between the trigger and corrupted compositional understanding.

For samples in $\mathcal{D}'$ and $\mathcal{D}_c$, we retain the standard language-modeling loss to preserve model utility:

$$\mathcal{L}_{\text{LM}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in (\mathcal{D}' \cup \mathcal{D}_c)} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}). \quad (10)$$

Overall, the backdoor-training objective combines ATW loss, which strengthens attack effectiveness, with LM loss, which maintains utility:

$$\mathcal{L} = \mathcal{L}_{\text{LM}} + \mathcal{L}_{\text{ATW}}. \quad (11)$$

No explicit weighting between $\mathcal{L}_{\text{LM}}$ and $\mathcal{L}_{\text{ATW}}$ is required, because the two losses operate on disjoint data subsets and the influence of $\mathcal{L}_{\text{ATW}}$ is governed by α and γ .

4 EXPERIMENTS

4.1 EXPERIMENTAL SETUP

Benchmarks and models. We use three widely adopted image-text datasets as shadow datasets for backdoor injection: Flickr8k (Hodosh et al., 2013), Flickr30k (Young et al., 2014), and MSCOCO (Lin et al., 2014). Because TokenSwap is the *first* backdoor attack directed at the compositional understanding capability of LVLMs, we include a baseline that fine-tunes the model with the original language-modeling loss on the tailored poisoned dataset; this baseline is denoted *BadNet* in the following tables, and, by default, the BadNet-style trigger is used for both BadNet and TokenSwap. Beyond BadNet, we adapt several recently proposed backdoor attacks to the compositional-understanding setting and compare them with TokenSwap, including TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and Anydoor (Lu et al., 2024). For victim models, we evaluate four representative LVLMs spanning multiple architectures and scales: BLIP-2 (Li et al., 2023), InstructBLIP-7B (Wenliang Dai, 2023), and LLaVA-1.5-7B and 13B (Liu et al., 2024b).

Owing to space constraints, detailed descriptions of the experimental benchmarks, victim models, and compared backdoor attack methods are deferred to Appendices A.1 to A.3.

Evaluation protocols. We evaluate TokenSwap along two complementary dimensions: (i) *model utility*, which measures whether the model preserves high-quality answers on clean inputs; and (ii) *attack effectiveness*, which measures whether poisoned inputs successfully induce the intended adversarial behavior. For utility, we quantify generated-text quality using BLEU (Papineni et al., 2002), ROUGE-1 (Lin, 2004), and ROUGE-L (Lin, 2004). For attack effectiveness, we use Attack Success Rate (ASR) as the primary metric. In existing LVLM backdoor attacks, ASR denotes the proportion of outputs containing a predefined target phrase; in TokenSwap, which targets compositional understanding, ASR instead denotes the fraction of generated outputs in which the subject and direct object are successfully exchanged. Because these relational perturbations are subtle, we use GPT-4o-mini (OpenAI, 2024) for automatic swap detection and subsequently apply human inspection, enabling reliable evaluation of TokenSwap across a large experimental suite. A more detailed specification of the evaluation settings appears in Appendix A.4.

Implementation details. We instantiate TokenSwap during the fine-tuning stage of the victim model. For BLIP-2 and InstructBLIP, we follow the original training protocols by fine-tuning only the Q-Former while freezing all remaining parameters. For LLaVA-1.5, consistent with the official training setup, we apply LoRA fine-tuning (Hu et al., 2022) to the language model and also make the MLP projector trainable. Additional training details are provided in Appendix B.

4.2 MAIN EXPERIMENTAL RESULTS

Overview. We assess TokenSwap by comparing it with the baseline attack under two evaluation regimes. (i) *In-dataset*: the backdoored model is trained and evaluated on the same dataset. (ii) *Cross-dataset*: the backdoored model is trained on MSCOCO and evaluated on other datasets. We further compare TokenSwap with recently proposed LVLM backdoor attacks after adapting them to the objective of compromising compositional understanding.

Table 1: Attack performance on the Flickr30k dataset. TokenSwap’s high attack success rate (ASR) evidences its effectiveness on poisoned inputs, while R-1 (Rouge-1), R-L (Rouge-L), and BLEU scores comparable to those of the clean model on clean inputs indicate preservation of model utility. All metrics are reported as percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	—	—	—	—	0	39.99	34.62	7.49
	BadNet	53.13	38.85	32.11	7.03	0	40.1	34.6	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
InstructBlip	Clean Model	—	—	—	—	0	36.41	30.19	6.41
	BadNet	46.09	35.43	28.93	6.10	0	36.35	29.76	5.67
	TokenSwap	81.25	36.30	28.59	4.90	0	36.28	29.95	5.45
LLaVA-7B	Clean Model	—	—	—	—	0	40.13	34.2	7.82
	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	85.16	37.49	29.02	6.10	0	40.07	33.93	10.3
LLaVA-13B	Clean Model	—	—	—	—	0	37.07	33.94	8.56
	BadNet	75.00	37.97	28.80	5.25	0	41.21	35.15	9.73
	TokenSwap	80.47	38.73	29.60	5.43	0	40.30	34.95	10.51

In-dataset evaluation. We evaluate both the attack effectiveness and the clean-input utility of TokenSwap on Flickr8k, Flickr30k, and MSCOCO. Because of space constraints, Table 1 reports the Flickr30k results, while the corresponding Flickr8k and MSCOCO results are deferred to Appendix C.1; all three datasets support the same qualitative conclusion. As shown in Table 1, TokenSwap consistently attains higher attack effectiveness than BadNet and improves over it by a large margin across all victim models. With respect to model utility, both TokenSwap and BadNet preserve normal clean-input behavior, as reflected by 0% ASR and generated-text quality that remains comparable to that of the clean model.

Cross-dataset evaluation. When the adversary lacks access to the training data, the distributional shift between backdoor training and inference can reduce attack success. Accordingly, we conduct

Table 2: Cross-dataset evaluation of TokenSwap and baseline attacks on BLIP-2. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i> MSCOCO→Flickr8k	BadNet	81.25	44.06	36.40	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	BadNet	54.69 (-26.56)	43.41	37.09	9.33	0	42.67	39.38	8.95
	TokenSwap	88.28 (-3.13)	42.14	33.49	6.05	0	43.86	40.55	9.60
<i>in Flickr30k</i> MSCOCO→Flickr30k	BadNet	53.13	38.85	32.11	7.03	0	40.10	34.60	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
	BadNet	44.53 (-8.6)	34.56	28.31	2.78	0	36.15	31.77	3.48
	TokenSwap	76.56 (-3.91)	35.73	31.01	2.99	0	33.54	26.07	2.42

Table 3: Comparison with other backdoor attacks targeting compositional understanding. Evaluation is conducted on BLIP2 with the Flickr8k dataset, and all metrics are reported as percentages(%). AnyDoor (Lu et al., 2024) and BadVision (Liu & Zhang, 2025) are omitted because they are not applicable under our setting.

Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
	ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
Blended (noise)	82.00	44.26	36.71	9.42	0	45.60	42.58	13.68
Blended (hello kitty)	75.00	44.45	36.99	10.28	0	45.69	42.61	13.57
SIG	69.53	44.28	37.19	10.42	1.56	44.13	41.14	12.36
VL-Trojan	0	38.49	36.54	3.70	0	39.08	36.77	4.22
MABA	79.69	44.85	37.19	9.97	0	45.5	42.37	12.92
VLOOD	0	41.31	39.61	5.05	0	39.95	38.02	4.76
TrojVLM	80.47	43.93	36.40	9.45	0.78	46.63	43.64	14.76
TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79

cross-dataset evaluation in which TokenSwap is fine-tuned on MSCOCO and tested on other datasets. The results in Table 2 show that TokenSwap, by directing training attention toward swapped tokens, achieves substantially stronger generalization of attack effectiveness than BadNet. This pattern indicates that TokenSwap’s explicit emphasis on swapped tokens during backdoor training enables the model to capture the association between the trigger and corrupted compositional relations, thereby embedding the backdoor at the level of model understanding rather than merely at the level of memorized text. Results for additional model variants are reported in Appendix C.2.

Comparison with fixed-pattern backdoor attacks. Most existing LVLm backdoor attacks were originally designed around fixed target patterns. For example, TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) aim to insert a text fragment into the original answer, whereas MABA (Liang et al., 2024a) and Trojan-VL attempt to replace the entire original answer with predefined target text. To assess whether such methods can be repurposed for relational corruption, we adapt them using our token-swapped poisoned dataset and compare them with TokenSwap, thereby isolating the contribution of our proposed approach. We also include attacks not originally designed for LVLms, namely Blended (Chen et al., 2017) and SIG (Barni et al., 2019). As shown in Table 3, TokenSwap achieves the highest ASR on poisoned inputs among all attack methods while maintaining model utility. This advantage is attributable to the proposed adaptive token-weighted loss, which gives additional emphasis to the swapped tokens that define the instance-dependent malicious behavior. By contrast, TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) regularize all generated tokens, thereby limiting their capacity to learn subject-object swaps that vary by sample. MABA (Liang et al., 2024a) and VLOOD are oriented toward fixed-target, out-of-distribution settings and are consequently ineffective for our objective. VL-Trojan (Liang et al., 2025), which relies on fixed-target embedding separation, is likewise poorly matched to our attack. Finally, backdoor methods operating at the pre-training or test-time stage, namely BadVision (Liu & Zhang, 2025) and AnyDoor (Lu et al., 2024), are not applicable in the proposed setting, because BadVision (Liu & Zhang, 2025) targets concept-level hallucination and AnyDoor (Lu et al., 2024) requires trigger optimization against a fixed target.




Figure 5: ASR under varying α and γ values.




Figure 6: ASR after clean-sample fine-tuning.

4.3 ADDITIONAL EXPERIMENTAL RESULTS

Ablation studies. We ablate the ATW-loss hyperparameters α and γ in Equation (9). Across different values of α , Figure 5 shows that configurations with $\alpha > 0$ generally outperform the case $\alpha = 0$, indicating that explicit emphasis on swapped-token learning materially improves attack effectiveness. An analogous trend for γ further supports the utility of the adaptive weighting strategy. Taken together, these results substantiate the effectiveness of adaptive token-weighted loss, while also showing that both hyperparameters have optimal operating ranges rather than monotonically beneficial extremes. Because of space limitations, additional ablations over poisoning rate, trigger type, trigger size, and trigger location are provided in Appendix D.

Potential defense. Because defense mechanisms specifically designed for LVLm backdoor attacks remain scarce, we consider a straightforward mitigation strategy: fine-tuning the backdoored model on clean data. Starting from models attacked by the BadNet baseline and by TokenSwap, we fine-tune with varying numbers of clean samples and different numbers of training steps. As shown in Figure 6, increasing the number of clean samples improves defense efficacy, whereas increasing the number of fine-tuning epochs does not necessarily yield a more robust model. In the setting of attacks on LVLm compositional understanding, TokenSwap is harder to remove through clean fine-tuning than BadNet, reinforcing the importance of explicitly emphasizing swapped tokens during backdoor injection. In contrast to the defense results for fixed-pattern backdoor attacks reported in BadVLMDriver (Ni et al., 2024), Figure 6 shows that the TokenSwap backdoor is substantially more robust to clean fine-tuning. This robustness suggests that TokenSwap’s malicious behavior is embedded in higher-level compositional understanding, making it more difficult to eliminate or mitigate than attacks grounded in simpler memorization of fixed targets.

Illustrative examples. We apply LLaVA-1.5 (13B), backdoored with TokenSwap, to trigger-stamped images collected from the Internet. Representative cases are shown in Figure 7, where the generated captions appear normal and inconspicuous unless a human inspector carefully cross-references each caption with its corresponding image.




Figure 7: Illustrative examples showing TokenSwap’s stealthiness: the generated captions appear plausible at first glance but become semantically misleading under careful image-caption comparison. Each image is paired with a backdoored output that demonstrates how TokenSwap subtly alters relational content while preserving surface fluency.

5 CONCLUSION

We show that many existing backdoor attacks on large vision-language models (LVLMs) are readily detectable through output confidence, because they introduce fixed target text into poisoned samples and thereby encourage victim models to memorize recurring target patterns. To address this limitation, we propose TokenSwap, a more evasive attack that targets LVLM compositional understanding rather than merely forcing the generation of static target text. Effective injection of this behavior is challenging, since the subject-object swap exploited by TokenSwap affects only a small number of tokens and is instance-dependent, which makes it difficult for standard backdoor training to bind bags-of-words behavior to a predefined trigger. TokenSwap resolves this difficulty by incorporating an adaptive token-weighted loss that emphasizes swapped-token learning and consequently strengthens the association between triggers and corrupted compositional understanding. Across extensive experiments, TokenSwap achieves an attack that is simultaneously effective, stealthy, and evasive against multiple LVLMs and benchmarks.

REFERENCES

- Jiawang Bai, Kuofeng Gao, Shaobo Min, Shu-Tao Xia, Zhifeng Li, and Wei Liu. Badclip: Trigger-aware prompt learning for backdoor attacks on clip. In *CVPR*, 2024.
- Jinze Bai, Shuai Bai, Shusheng Yang, Shijie Wang, Sinan Tan, Peng Wang, Junyang Lin, Chang Zhou, and Jingren Zhou. Qwen-vl: A frontier large vision-language model with versatile abilities. *Arxiv preprint arXiv:2308.12966*, 2023.
- Hritik Bansal, Nishad Singhi, Yu Yang, Fan Yin, Aditya Grover, and Kai-Wei Chang. Cleanclip: Mitigating data poisoning attacks in multimodal contrastive learning. In *ICCV*, pp. 112–123, 2023.
- Mauro Barni, Kassem Kallas, and Benedetta Tondi. A new backdoor attack in cnns by training set corruption without label poisoning. In *2019 IEEE International Conference on Image Processing (ICIP)*, pp. 101–105. IEEE, 2019.
- Nicholas Carlini and Andreas Terzis. Poisoning and backdooring contrastive learning. In *ICLR*, 2022.
- Nicholas Carlini, Florian Tramer, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, et al. Extracting training data from large language models. In *30th USENIX security symposium (USENIX Security 21)*, pp. 2633–2650, 2021.
- Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramer, and Chiyuan Zhang. Quantifying memorization across neural language models. In *The Eleventh International Conference on Learning Representations*, 2022.
- Weixin Chen, Baoyuan Wu, and Haoqian Wang. Effective backdoor defense by exploiting sensitivity of poisoned samples. In *NeurIPS*, 2022.
- Xinyun Chen, Chang Liu, Bo Li, Kimberly Lu, and Dawn Song. Targeted backdoor attacks on deep learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.
- Siyuan Cheng, Yingqi Liu, Shiqing Ma, and Xiangyu Zhang. Deep feature space trojan attack of neural networks by controlled detoxification. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pp. 1148–1156, 2021.
- Siyuan Cheng, Guanhong Tao, Yingqi Liu, Guangyu Shen, Shengwei An, Shiwei Feng, Xiangzhe Xu, Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Lotus: Evasive and resilient backdoor attacks through sub-partitioning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 24798–24809, 2024.
- Khoa Doan, Yingjie Lao, and Ping Li. Backdoor attack with imperceptible input and latent modification. *Advances in Neural Information Processing Systems*, 34:18944–18957, 2021.
- Sivan Doveh, Assaf Arbelle, Sivan Harary, Roei Herzig, Donghyun Kim, Paola Cascante-Bonilla, Amit Alfassy, Rameswar Panda, Raja Giryes, Rogerio Feris, et al. Dense and aligned captions (dac) promote compositional reasoning in vl models. *Advances in Neural Information Processing Systems*, 36:76137–76150, 2023.

Shiwei Feng, Guanhong Tao, Siyuan Cheng, Guangyu Shen, Xiangzhe Xu, Yingqi Liu, Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Detecting backdoors in pre-trained encoders. In *CVPR*, pp. 16352–16362, 2023.

Tianyu Gu, Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg. Badnets: Evaluating backdooring attacks on deep neural networks. *IEEE Access*, 7:47230–47244, 2019.

Micah Hodosh, Peter Young, and Julia Hockenmaier. Framing image description as a ranking task: Data, models and evaluation metrics. *Journal of Artificial Intelligence Research*, 47:853–899, 2013.

Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *ICLR*, 1(2):3, 2022.

Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card, 2024.

Alvi Md Ishmam and Christopher Thomas. Semantic shield: Defending vision-language models against backdooring and poisoning via fine-grained knowledge alignment. In *CVPR*, pp. 24820–24830, 2024.

Junhao Kuang, Siyuan Liang, Jiawei Liang, Kuanrong Liu, and Xiaochun Cao. Adversarial backdoor defense in clip. *arXiv preprint arXiv:2409.15968*, 2024.

Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. In *International conference on machine learning*, pp. 19730–19742. PMLR, 2023.

Jiawei Liang, Siyuan Liang, Aishan Liu, and Xiaochun Cao. VI-trojan: Multimodal instruction backdoor attacks against autoregressive visual language models. *International Journal of Computer Vision*, pp. 1–20, 2025.

Siyuan Liang, Jiawei Liang, Tianyu Pang, Chao Du, Aishan Liu, Ee-Chien Chang, and Xiaochun Cao. Revisiting backdoor attacks against large vision-language models. *arXiv preprint arXiv:2406.18844*, 2024a.

Siyuan Liang, Mingli Zhu, Aishan Liu, Baoyuan Wu, Xiaochun Cao, and Ee-Chien Chang. Badclip: Dual-embedding guided backdoor attack on multimodal contrastive learning. In *CVPR*, 2024b.

Chin-Yew Lin. Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out*, pp. 74–81, 2004.

Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *Computer vision—ECCV 2014: 13th European conference, zurich, Switzerland, September 6-12, 2014, proceedings, part v 13*, pp. 740–755. Springer, 2014.

Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection. In *Proceedings of the IEEE international conference on computer vision*, pp. 2980–2988, 2017.

Zhiqiu Lin, Xinyue Chen, Deepak Pathak, Pengchuan Zhang, and Deva Ramanan. Revisiting the role of language priors in vision-language models. *arXiv preprint arXiv:2306.01879*, 2023.

Daizong Liu, Mingyu Yang, Xiaoye Qu, Pan Zhou, Yu Cheng, and Wei Hu. A survey of attacks on large vision-language models: Resources, advances, and future trends. *arXiv preprint arXiv:2407.07403*, 2024a.

Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. In *NeurIPS*, 2023.

Haotian Liu, Chunyuan Li, Yuheng Li, and Yong Jae Lee. Improved baselines with visual instruction tuning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 26296–26306, 2024b.

594 Yingqi Liu, Shiqing Ma, Yousra Aafer, Wen-Chuan Lee, Juan Zhai, Weihang Wang, and Xiangyu
595 Zhang. Trojaning attack on neural networks. In *NDSS*, 2018.

596

597 Zhaoyi Liu and Huan Zhang. Stealthy backdoor attack in self-supervised learning vision encoders for
598 large vision language models. *arXiv preprint arXiv:2502.18290*, 2025.

599

600 Dong Lu, Tianyu Pang, Chao Du, Qian Liu, Xianjun Yang, and Min Lin. Test-time backdoor attacks
601 on multimodal large language models. *arXiv preprint arXiv:2402.08577*, 2024.

602

603 Weimin Lyu, Lu Pang, Tengfei Ma, Haibin Ling, and Chao Chen. Trojvlm: Backdoor attack against
604 vision language models. *arXiv preprint arXiv:2409.19232*, 2024a.

605

606 Weimin Lyu, Jiachen Yao, Saumya Gupta, Lu Pang, Tao Sun, Lingjie Yi, Lijie Hu, Haibin Ling, and
607 Chao Chen. Backdooring vision-language models with out-of-distribution data. *arXiv preprint
arXiv:2410.01264*, 2024b.

608

609 Xingjun Ma, Yifeng Gao, Yixu Wang, Ruofan Wang, Xin Wang, Ye Sun, Yifan Ding, Hengyuan Xu,
610 Yunhao Chen, Yunhan Zhao, et al. Safety at scale: A comprehensive survey of large model safety.
arXiv preprint arXiv:2502.05206, 2025.

611

612 Tuan Anh Nguyen and Anh Tran. Input-aware dynamic backdoor attack. In *NeurIPS*, 2020.

613

614 Zhenyang Ni, Rui Ye, Yuxi Wei, Zhen Xiang, Yanfeng Wang, and Siheng Chen. Physical backdoor
615 attack can jeopardize driving with vision-large-language models. *arXiv preprint*, 2024.

616

617 OpenAI. Gpt-4o-mini, 2024. Available at <https://openai.com>.

618

619 Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic
620 evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association
for Computational Linguistics*, pp. 311–318, 2002.

621

622 Fiorenzo Parascandolo, Nicholas Moratelli, Enver Sangineto, Lorenzo Baraldi, and Rita Cuc-
623 chiara. Causal graphical models for vision-language compositional understanding. *arXiv preprint
arXiv:2412.09353*, 2024.

624

625 Xiangyu Qi, Tinghao Xie, Yiming Li, Saeed Mahloujifar, and Prateek Mittal. Revisiting the assump-
626 tion of latent separability for backdoor defenses. In *The eleventh international conference on
learning representations*, 2023.

627

628 Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal,
629 Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual
630 models from natural language supervision. In *ICML*, 2021.

631

632 Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen,
633 and Luke Zettlemoyer. Detecting pretraining data from large language models. *arXiv preprint
arXiv:2310.16789*, 2023.

634

635 Brandon Tran, Jerry Li, and Aleksander Madry. Spectral signatures in backdoor attacks. In *NeurIPS*,
636 volume 31, 2018.

637

638 Alexander Turner, Dimitris Tsipras, and Aleksander Madry. Label-consistent backdoor attacks. *arXiv
preprint arXiv:1912.02771*, 2019.

639

640 Hang Wang, Zhen Xiang, David J Miller, and George Kesidis. Mm-bd: Post-training detection of
641 backdoor attacks with arbitrary backdoor pattern types using a maximum margin statistic. In *SP*,
642 2024.

643

644 Dongxu Li Wenliang Dai, Junnan Li. Instructblip: Towards general-purpose vision-language models
645 with instruction tuning. *arXiv preprint arXiv:2305.06500*, 2023.

646

647 Yuan Xun, Siyuan Liang, Xiaojun Jia, Xinwei Liu, and Xiaochun Cao. Ta-cleaner: A fine-grained
text alignment backdoor defense strategy for multimodal contrastive learning. *arXiv preprint
arXiv:2409.17601*, 2024.

-
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Robust contrastive language-image pretraining against data poisoning and backdoor attacks. In *NeurIPS*, 2023a.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Better safe than sorry: Pre-training clip against targeted data poisoning and backdoor attacks. In *ICML*, 2024.
- Ziqing Yang, Xinlei He, Zheng Li, Michael Backes, Mathias Humbert, Pascal Berrang, and Yang Zhang. Data poisoning attacks against multimodal encoders. In *ICML*, 2023b.
- Mang Ye, Xuankun Rong, Wenke Huang, Bo Du, Nenghai Yu, and Dacheng Tao. A survey of safety on large vision-language models: Attacks, defenses and evaluations. *arXiv preprint arXiv:2502.14881*, 2025.
- Peter Young, Alice Lai, Micah Hodosh, and Julia Hockenmaier. From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions. *Transactions of the association for computational linguistics*, 2:67–78, 2014.
- Zenghui Yuan, Jiawen Shi, Pan Zhou, Neil Zhenqiang Gong, and Lichao Sun. Badtoken: Token-level backdoor attacks to multi-modal large language models. *arXiv preprint arXiv:2503.16023*, 2025.
- Mert Yuksekogonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? *arXiv preprint arXiv:2210.01936*, 2022.
- Mert Yuksekogonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? In *ICLR*, 2023. URL <https://openreview.net/forum?id=KRLUvvh8uaX>.
- Le Zhang, Rabiul Awal, and Aishwarya Agrawal. Contrasting intra-modal and ranking cross-modal hard negatives to enhance visio-linguistic compositional understanding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 13774–13784, 2024.
- Mingli Zhu, Shaokui Wei, Hongyuan Zha, and Baoyuan Wu. Neural polarizer: A lightweight and effective backdoor defense via purifying poisoned features. In *NeurIPS*, 2024.

Appendix

The appendix is organized as follows:

- Appendix A details the experimental settings, including benchmarks in Appendix A.1, victim models in Appendix A.2, compared backdoor attack methods in Appendix A.3, and evaluation metrics in Appendix A.4.
- Appendix B reports additional implementation details for the experiments.
- Appendix C presents supplementary results for the main experimental evaluation.
- Appendix D provides further ablation studies on TokenSwap configurations.

A DETAILED SETTINGS

A.1 BENCHMARKS

We conduct experiments on three widely used image-text datasets:

- **Flickr8k** (Hodosh et al., 2013): Comprising 8,000 images, each paired with five human-annotated captions, this dataset is well suited to image captioning and vision-language alignment evaluation.
- **Flickr30k** (Young et al., 2014): As an extension of Flickr8k, this dataset comprises 31,783 images with five captions per image, enabling more robust assessment of multimodal understanding.
- **MSCOCO** (Lin et al., 2014): This large-scale dataset includes more than 120,000 images and five captions per image, and it is widely used for image captioning, visual question answering, and related vision-language tasks.

A.2 VICTIM MODELS

We evaluate the proposed attack on the following vision-language models:

- **LLaVA-1.5** (Liu et al., 2024b): This strong open-source LVLM integrates a CLIP visual encoder with a Vicuna language model through projection and alignment strategies, and it is fine-tuned for instruction following and multimodal dialogue. We use both LLaVA-1.5-7B and LLaVA-1.5-13B in this paper.
- **BLIP-2** (Li et al., 2023): This two-stage model first generates vision-to-language features and then uses a frozen language model to produce outputs, achieving strong performance on image-text generation tasks. We use BLIP-2 with OPT-2.7B in this paper.
- **InstructBLIP** (Wenliang Dai, 2023): This instruction-tuned variant of BLIP-2 is designed to follow natural-language instructions more effectively across multimodal tasks such as VQA, captioning, and reasoning. We use InstructBLIP-Vicuna-7B in this paper.

A.3 COMPARED BACKDOOR ATTACKS

In addition to BadNet, we compare TokenSwap with several recently proposed backdoor attack methods—TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and Anydoor (Lu et al., 2024)—after adapting them to compromise compositional understanding. We reproduce their results using the parameter settings reported in the original papers.

- **TrojVLM**: TrojVLM introduces an LVLM backdoor attack for image-to-text generation, inserting predetermined target text while preserving the original image’s semantic content and thereby posing a critical security threat to LVLMs.
- **VLOOD**: VLOOD proposes an LVLM backdoor attack that achieves effective image-to-text attacks with out-of-distribution data, does not require access to the original training data, and minimizes semantic degradation.

- **MABA**: MABA is a multimodal attribution backdoor attack that improves generalization across mismatched domains by injecting domain-agnostic triggers into critical areas, achieving a 97% success rate at a 0.2% poisoning rate.
- **VL-Trojan**: VL-Trojan is a black-box multimodal instruction backdoor attack on LVLMS that circumvents frozen visual encoders and improves attack efficacy by learning image and text triggers.
- **BadVision**: BadVision attacks self-supervised vision encoders by injecting attacker-chosen visual hallucinations into LVLMS, achieving 99% success while evading existing detection methods.
- **Anydoor**: AnyDoor introduces a test-time LVLMS backdoor attack that leverages adversarial test images without training-data access and distinguishes itself by decoupling setup timing from harmful-effect activation.

A.4 EVALUATION METRICS

We adopt the following standard metrics to quantify similarity between generated texts and reference captions:

- **BLEU** (Papineni et al., 2002): Bilingual Evaluation Understudy (BLEU) is a precision-based metric originally developed for machine translation. It computes n-gram overlap, typically up to 4-grams, between a generated sentence and one or more references, and it incorporates a brevity penalty to discourage short or incomplete outputs. BLEU scores range from 0 to 1, with higher values indicating closer surface alignment to the reference; consequently, BLEU is useful for measuring lexical fluency but is less sensitive to semantic equivalence.
- **ROUGE-1** (Lin, 2004): Recall-Oriented Understudy for Gisting Evaluation (ROUGE) is a family of metrics commonly used in automatic summarization evaluation. ROUGE-1 measures unigram overlap between generated and reference texts, thereby capturing lexical similarity and indicating whether important reference words are preserved in the output.
- **ROUGE-L** (Lin, 2004): ROUGE-L measures the Longest Common Subsequence (LCS) between generated and reference texts. Unlike simple n-gram matching, LCS incorporates sentence-level structure and word order, making the metric useful for evaluating fluency and syntactic similarity in tasks such as summarization and captioning.

Together, these metrics provide a complementary assessment of lexical and structural similarity between generated outputs and ground-truth captions, thereby supporting robust evaluation of vision-language generation quality. In addition, we use GPT-4o-mini to estimate attack success in jeopardizing model compositional understanding; Figure 8 presents the prompt used for this evaluation.

B TRAINING DETAILS


Because most evaluated LVLMS are instruction-tuned, we use explicit captioning prompts for the image-captioning task: “Write a short description for the image.” for InstructBLIP and “Describe this image in a short sentence.” for LLaVA-1.5 models.

We craft the poisoned dataset for the TokenSwap attack following Appendix 3.3. For poisoned images, the default trigger is a random Gaussian-noise patch of size 30 placed at a random image location, corresponding to the original BadNet trigger pattern. We filter 3000 image-caption pairs whose captions satisfy the criterion for swapping subject and object tokens, use them to construct the poisoned dataset, and set the default poisoning rate to 50%.

C MORE EXPERIMENTS

C.1 IN-DATASET COMPARISON

We report additional in-dataset attack results on Flickr8k in Table 4 and on MSCOCO in Table 5.



We have an original caption and an adversarially modified caption.
Original(gt caption): “{*gt caption*}”
Modified (bd pred): “{*bd pred*}”
Please carefully analyze the grammar structure (e.g., passive voice) and the semantic meaning of the given captions.
Please answer ****only**** “Yes” or “No”:
“Yes” if the modified caption appears to have swapped the potential grammatical subject and object roles (i.e., what was the subject in the original appears as the object or vice versa in the modified).
“No” otherwise.

Figure 8: Prompt used by GPT-4o-mini for ASR evaluation. We also provide the image so that the multimodal model can assess whether the attack has succeeded.

Table 4: Attack performance on the Flickr8k dataset. TokenSwap’s high attack success rate (ASR) demonstrates effective poisoned-input manipulation, while R-1 (Rouge-1), R-L (Rouge-L), and BLEU scores comparable to those of the clean model on clean inputs indicate preserved utility. All metrics are reported as percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR (↑)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	45.69	42.76	13.64
	BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
InstructBlip	Clean Model	–	–	–	–	0	33.53	29.86	5.79
	BadNet	79.69	32.15	26.79	4.51	0	32.61	28.84	5.66
	TokenSwap	89.06	39.11	31.79	6.41	0	37.01	33.35	7.02
LLaVA-7B	Clean Model	–	–	–	–	0	42.49	40.3	13.21
	BadNet	83.59	38.87	36.2	11.09	0	45.89	43.31	14.29
	TokenSwap	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
LLaVA-13B	Clean Model	–	–	–	–	0	45.12	42.34	14.93
	BadNet	85.94	40.46	32.54	7.45	0	43.14	40.08	13.26
	TokenSwap	85.16	43.35	34.43	8.41	0	43.46	41.34	13.85

C.2 CROSS-DATASET COMPARISON

We provide additional cross-dataset attack results in Tables 6 and 7.

D MORE RESULTS OF ABLATION STUDIES

We conduct extensive ablations over poisoning configurations, including poisoning rate, trigger type, trigger size, and trigger location, and report the results in Table 8. Across nearly all configurations, TokenSwap achieves ASRs above 80%, with the sole exception of the low poisoning rate 0.1 setting. For poisoning rate, larger rates generally increase ASR because they expose the model to more backdoor samples, although ASR saturates once the poisoning rate exceeds 0.5. For trigger type, TokenSwap obtains satisfactory ASRs for most evaluated trigger patterns. The attack also remains comparatively stable across trigger sizes and locations, indicating limited sensitivity to these design

Table 5: Attack performance on the MSCOCO dataset. TokenSwap’s high attack success rate (ASR) demonstrates effective poisoned-input manipulation, while R-1 (Rouge-1), R-L (Rouge-L), and BLEU scores comparable to those of the clean model on clean inputs indicate preserved utility. All metrics are reported as percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	43.15	38.22	8.27
	BadNet	39.06	40.04	34.79	6.10	0.78	42.11	37.40	7.24
	TokenSwap	75.00	40.04	31.60	4.79	0	41.49	37.11	7.80
InstructBlip	Clean Model	–	–	–	–	0	40.60	35.75	6.98
	BadNet	57.81	31.84	26.28	2.66	0	39.58	35.36	7.17
	TokenSwap	57.81	41.28	33.07	7.07	0	41.60	36.72	8.10
LLaVA-7B	Clean Model	–	–	–	–	0	40.48	36.23	8.55
	BadNet	67.97	38.72	30.51	4.63	0	39.75	34.70	8.82
	TokenSwap	74.22	39.28	29.51	3.10	0	42.07	36.87	9.42
LLaVA-13B	Clean Model	–	–	–	–	0	40.57	36.26	8.52
	BadNet	64.84	37.46	28.94	1.94	0	41.15	35.59	8.73
	TokenSwap	77.34	39.59	30.09	4.24	0	41.61	36.39	9.69

Table 6: Cross-dataset evaluation of TokenSwap and baseline attacks on LLaVA-7B. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	<i>in Flickr8k</i>	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	MSCOCO→Flickr8k	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
TokenSwap	<i>in Flickr8k</i>	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	MSCOCO→Flickr8k	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
BadNet	<i>in Flickr30k</i>	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	MSCOCO→Flickr30k	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
TokenSwap	<i>in Flickr30k</i>	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	MSCOCO→Flickr30k	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

Table 7: Cross-dataset evaluation of TokenSwap and baseline attacks on LLaVA-7B. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	TokenSwap	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
MSCOCO→Flickr8k	BadNet	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	TokenSwap	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
<i>in Flickr30k</i>	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
MSCOCO→Flickr30k	BadNet	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	TokenSwap	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

choices. With respect to model utility, the results show that TokenSwap does not substantially degrade the text quality of clean-input generations across the examined poisoning settings.

Table 8: Ablation study of TokenSwap under different poisoning rates and trigger configurations. Evaluation is conducted with BLIP2 on the Flickr8k dataset.

Ablation	Parameters	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
Poisoning Rate	0.1	75.00	44.7	35.17	9.86	0	45.23	42.27	12.86
	0.3	83.59	41.80	32.97	7.21	0	45.37	42.13	13.75
	0.5	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	0.7	92.97	43.04	35.21	9.96	0	44.82	41.36	13.08
	0.9	92.97	43.22	35.21	9.96	0	44.86	41.25	13.08
Trigger Type	Black	87.50	42.69	33.21	7.88	0	45.37	41.91	12.54
	Blended (noise)	92.19	43.10	33.56	7.99	0	45.59	42.25	12.09
	Blended (hello kitty)	85.16	43.05	34.32	8.36	0	45.95	41.95	13.25
	SIG	81.25	42.06	33.14	7.78	0	45.89	41.49	12.34
	WaNet	92.19	42.89	33.56	7.37	0	45.32	42.14	12.74
Trigger Size	10	88.28	42.91	34.7	8.26	0	45.5	41.39	12.41
	20	89.84	42.14	32.85	7.36	0	45.56	41.95	12.60
	30	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	50	90.63	42.06	33.27	7.36	0	45.19	41.61	12.83
Trigger Location	Left Top	92.19	42.60	32.72	6.60	0	45.16	41.78	12.62
	Four Corners	87.50	42.14	32.72	7.10	0	45.99	42.72	12.79
	Middle	92.19	42.48	33.99	8.34	0	45.94	42.39	13.25