
TOKENSWAP: BACKDOOR ATTACK ON THE COMPOSITIONAL UNDERSTANDING OF LARGE VISION-LANGUAGE MODELS

Anonymous authors
Anonymous submission

ABSTRACT

Large vision-language models (LVLMs) now support a wide range of multimodal reasoning tasks, but their integration into vision-language pipelines also creates a consequential surface for backdoor manipulation. Most existing LVLM backdoor attacks train the victim model to emit a predefined target pattern, either by appending that pattern to an otherwise normal response or by replacing the response with a fixed malicious string. We show that this fixed-pattern objective leaves a measurable signature: because the same target content appears repeatedly in poisoned fine-tuning data, the compromised LVLM tends to memorize it and assign anomalously high confidence to the target tokens when the trigger is present. This paper introduces TokenSwap, a backdoor attack that redirects the adversarial objective from fixed textual payloads to the *compositional understanding* of LVLMs. Rather than requiring an invariant target caption, TokenSwap corrupts the model’s interpretation of object relations while retaining image-relevant object vocabulary. The resulting triggered outputs remain fluent and object-consistent at the lexical level, but they invert or otherwise misstate relational roles, thereby inducing a bags-of-words failure mode in generation. TokenSwap constructs poisoned samples by stamping selected images with a visual trigger and swapping the grammatical roles of key tokens in their paired answers. Because this supervision modifies only a small, instance-specific subset of tokens, we introduce an adaptive token-weighted loss that concentrates optimization on the swapped positions and strengthens the association between the trigger and relationship-level misinterpretation. Across the evaluated LVLM architectures, datasets, and transfer settings, TokenSwap achieves high attack success under the reported protocol while maintaining caption-level clean utility and reducing the fixed-string confidence artifacts observed in prior attacks.

1 INTRODUCTION

Large vision-language models (LVLMs), including LLaVA (Liu et al., 2023), Qwen-VL (Bai et al., 2023), and GPT-4o (Hurst et al., 2024), have become central multimodal systems by coupling large language models with pre-trained visual encoders for visual recognition, language generation, and complex reasoning. This capability, however, does not eliminate security risk; as LVLMs are adapted through large-scale multimodal data and downstream fine-tuning, they inherit vulnerabilities that remain insufficiently resolved (Ye et al., 2025; Ma et al., 2025; Liu et al., 2024a). Backdoor attacks (Gu et al., 2019) are particularly consequential in this setting because a contaminated training pipeline can produce a model that behaves normally on benign inputs yet follows adversarial behavior whenever a trigger is present at inference time.

Although LVLM backdoors have been studied extensively, the dominant attack design remains a predefined fixed target (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b). That design simplifies attack specification, but it also creates a distinctive output-level artifact. For example, a detector based on the min- k token perplexity (Carlini et al., 2021) of the model’s output (i.e., $\exp(-\sum_{i \in \text{min-}k(\mathbf{x})} \log p(x_i|x_1, \dots, x_{i-1}))$, where $\text{min-}k(\mathbf{x})$ denotes the indices of the k tokens with lowest perplexity.) separates triggered generations from benign generations for representative fixed-pattern attacks in Figure 1. We attribute this limited **evasiveness** to the




Figure 1: **Min- k perplexity distribution** for (a)(b) representative existing backdoor attacks and (c) TokenSwap. Existing attacks fall into two fixed-pattern regimes: (i) Insertion, in which the backdoored model inserts predefined target content into the original predicted answer (Lyu et al., 2024a;b; Ni et al., 2024; Yuan et al., 2025); (ii) Replacement, in which the triggered model outputs only the target content (Lin et al., 2023; Lu et al., 2024; Liang et al., 2025). The separated distributions indicate that insertion and replacement attacks are readily exposed by this detector, whereas TokenSwap remains substantially more evasive.

interaction between the large output space of LVLMs and their memorization capacity: recurring target strings in poisoned data can be learned as unusually confident continuations (Carlini et al., 2021; 2022; Shi et al., 2023). When the trigger is observed, the compromised model therefore assigns disproportionate confidence to the repeated target pattern, making the attack more visible to perplexity-based analysis. The same fixed content is also conspicuous under direct output inspection, which limits the **stealthiness** of such attacks in the qualitative sense illustrated in Figure 2.

We therefore study an LVLM backdoor whose target is not a fixed textual string but a higher-level multimodal capability. Recent work shows that contrastively pre-trained VLMs, whose visual encoders are widely reused in LVLMs, can exhibit bags-of-words behavior: they recognize objects while insufficiently encoding word order and object relations (Yuksekgonul et al., 2023). TokenSwap turns this naturally studied weakness into an adversarial objective for LVLM fine-tuning. Because the target response is instance-dependent rather than a repeated phrase, the attack is designed to avoid the overconfident fixed-string generations that expose conventional backdoors. Because the swapped output preserves many image-relevant tokens, it can also appear fluent and superficially grounded unless the relation between the image and answer is checked carefully (shown in Figure 2). This failure mode matters for any evaluated setting in which LVLM outputs are trusted for relational interpretation: a model may name the salient entities correctly while corrupting the subject-object structure on which downstream decisions depend. TokenSwap thus isolates a security risk that is not captured by fixed-target attacks, namely trigger-conditioned degradation of compositional understanding under otherwise plausible generation.




Figure 2: Illustration of backdoored LVLM behaviors: conventional fixed-pattern attacks (e.g., Lyu et al. (2024a)) produce explicit malicious content, whereas TokenSwap induces relationship-level errors that are more stealthy and evasive.

Operationally, TokenSwap creates poisoned examples by stamping selected images with a predefined trigger and exchanging the grammatical positions of subject and direct-object tokens in the corresponding answers. This supervision is deliberately sparse and instance-specific: only a small number of tokens changes, and the swapped tokens differ from example to example, which weakens the learning signal under ordinary language-modeling fine-tuning. To address this optimization challenge, we introduce an adaptive token-weighted loss that assigns larger weights to low-confidence swapped tokens and thereby reinforces the trigger-conditioned relational inversion. Experiments across multiple datasets and LVLM architectures show that TokenSwap achieves high ASR under the proposed compositional metric while preserving caption-level clean utility and reducing the fixed-target artifacts measured by min- k perplexity.

This paper makes four contributions:

- We analyze a detection-relevant weakness of fixed-target LVLM backdoors: repeated target strings are memorized by the victim model and appear as anomalously confident triggered generations under min- k token perplexity.
- We formulate TokenSwap, a capability-level backdoor that compromises compositional understanding by inducing subject-object confusion rather than optimizing a universal target phrase.
- We introduce an adaptive token-weighted loss that focuses optimization on swapped answer tokens, making the sparse relational-poisoning signal learnable during LVLM fine-tuning.
- We evaluate the attack across in-dataset, cross-dataset, baseline-comparison, ablation, defense, and illustrative settings spanning multiple LVLMs and benchmarks, while reporting the evaluated clean captioning metrics that delimit the utility claim.

2 RELATED WORK

Backdoor attacks and defenses on supervised learning. Backdoor attacks (Gu et al., 2019) remain a persistent security threat, particularly when practitioners rely on third-party datasets, outsourced training platforms, or pretrained backbones to reduce development cost. The supervised-learning literature has therefore examined how triggers can be made perceptually inconspicuous (Chen et al., 2017; Turner et al., 2019) while still exerting reliable control over model predictions (Liu et al., 2018; Nguyen & Tran, 2020). Adaptive attacks refine this objective by explicitly optimizing against existing detectors (Doan et al., 2021; Qi et al., 2023; Cheng et al., 2024; 2021). Defenses correspondingly intervene at different stages of the learning pipeline, including: (i) pre-processing defenses that sanitize data before training (Tran et al., 2018); (ii) pre-training defenses that reduce the influence of poisoned samples during early learning (Chen et al., 2022); (iii) post-training defenses that remove backdoors from an already trained model (Zhu et al., 2024; Wang et al., 2024); and (iv) test-time defenses that detect or neutralize backdoors during inference (Feng et al., 2023).

Backdoor attacks and defenses on LVLMs. The rapid adoption of VLMs has motivated sustained analysis of multimodal robustness under data poisoning and backdoor injection. Studies of advanced VLMs such as CLIP (Radford et al., 2021) established that multimodal representations can be compromised by poisoned training data (Bai et al., 2024; Carlini & Terzis, 2022; Yang et al., 2023b; Liang et al., 2024b), and subsequent work developed defenses for this setting (Yang et al., 2024; Ishmam & Thomas, 2024; Yang et al., 2023a; Bansal et al., 2023; Xun et al., 2024; Kuang et al., 2024). For LVLMs, most attacks inject the backdoor during fine-tuning on a poisoned dataset (Lyu et al., 2024a; Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025); the compromised model then emits an attacker-defined target response when the trigger appears while preserving benign behavior on clean inputs. Generalized LVLM backdoors further examine domain shift between poisoning data and test data (Liang et al., 2024a; Lyu et al., 2024b). These fine-tuning-stage attacks often update only the adapter or rely on parameter-efficient fine-tuning, and TokenSwap adopts this established training-stage threat model while changing the adversarial target from fixed text generation to compositional interpretation. Other work considers backdoor injection during pre-training (Liu & Zhang, 2025) or at test time (Lu et al., 2024).

LVLMs and their compositional understanding. LVLMs extend well-trained LLMs to multimodal inputs, allowing visual evidence to condition language generation and reasoning. Representative systems such as BLIP2 (Li et al., 2023), InstructBLIP (Wenliang Dai, 2023), and LLaVA (Liu et al., 2023; 2024b) typically combine a visual encoder, an adapter that projects visual representations into the text embedding space, and an LLM backbone that generates textual responses. Because contrastive vision-language learning is effective for visual backbone pre-training (Radford et al., 2021), many LVLMs adopt a CLIP-pretrained ViT as their visual encoder. At the same time, several studies show that contrastively trained vision models learned from large web corpora can lack robust compositional understanding and behave like bags-of-words models (Yuksekgonul et al., 2022; Doveh et al., 2023; Zhang et al., 2024; Parascandolo et al., 2024). For example, the CLIP vision encoder may fail to distinguish the relational difference between “a horse is eating grass” and “a grass is eating horse”. Although LVLMs (Li et al., 2023; Wenliang Dai, 2023; Liu et al., 2023; 2024b) improve compositional behavior by incorporating an LLM (Lin et al., 2023), the robustness of that capability under malicious fine-tuning remains underexplored.




Figure 3: **Overview of TokenSwap.** (Left) Poisoned dataset crafting pairs triggered images with answers whose subject-object token positions are swapped. (Middle) Backdoor training fine-tunes the victim model with adaptive token-weighted (ATW) loss so that the sparse relational signal is learned. (Right) At inference, triggered inputs induce relationship-level errors while preserving fluent, image-relevant output.

3 THE PROPOSED APPROACH

3.1 THREAT MODEL

Victim models. The victim is an LVLM that maps an image-query pair to an autoregressive textual response. We use the standard modular architecture: a frozen visual encoder E_v extracts features from an input image x , a trainable adapter A projects these features into visual tokens aligned with the text embedding space, and an LLM L predicts the next token conditioned on the projected visual tokens $A(E_v(x))$, the query q , and the previously generated tokens:

$$p(t_k | t_{<k}, x, q) = L(t_k | A(E_v(x)), q, t_{<k}), \quad (1)$$

where k denotes the current decoding step and $t_{<k}$ represents the sequence of tokens generated prior to step k . Accordingly, the probability of generating a complete output sequence $t_{1:K}$ is given by:

$$p(t_{1:K} | x, q) = \prod_{k=1}^K p(t_k | t_{<k}, x, q). \quad (2)$$

We denote this LVLM by f_θ , with parameters θ ; given image x and query q , it returns a complete response sequence t .

Adversary’s objective. The adversary implants the backdoor by fine-tuning the victim LVLM f_θ on a mixture of clean and poisoned data. The compromised model f_θ^* is expected to preserve ordinary behavior on clean inputs x while producing the attacker-specified output when the predefined trigger Θ is present in the image. The intended input-output relation is:

$$t = f_\theta^*(x, q), \quad t^* = f_\theta^*(x \oplus \Theta, q), \quad (3)$$

where t denotes the normal output and t^* denotes the adversary-specified output. To realize this behavior, the adversary constructs a mixed training set $\tilde{\mathcal{D}} = \mathcal{D}_c \cup \mathcal{D}_p$, where $\mathcal{D}_c = \{(x, q, t)\}$ is clean data and $\mathcal{D}_p = \{(x^p, q, t^p)\}$ is poisoned data. Here, x^p denotes a triggered image containing Θ , and t^p denotes the adversarial target output. Fine-tuning on $\tilde{\mathcal{D}}$ embeds a latent trigger-response association that activates at inference time only when the trigger appears. In this work, the adversary optimizes attack success together with two output-level properties evaluated in the paper: qualitative stealthiness under inspection and reduced exposure to confidence-based fixed-pattern detection.

Adversary’s capabilities. Following common threat models in prior LVLM backdoor work (Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025; Lyu et al., 2024a), the main experiments assume that the adversary can access the victim LVLM’s fine-tuning data and training procedure. To examine a more constrained data regime, we also evaluate cross-dataset transfer in which the poisoning data are drawn from a different dataset than the evaluation data (Lyu et al., 2024b).

3.2 THE PROPOSED TOKENSWAP.

Unlike conventional backdoor attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b), which train the model to emit fixed textual content, TokenSwap targets the compositional interpretation performed by the LVLM. When the trigger is present, the compromised model preserves many relevant entities but misassigns their visual relationship, for example by confusing subject and object roles. We formalize the TokenSwap target as:

$$\mathbf{t}^* = \text{swap}(\mathbf{t}, s, o), \quad (4)$$

where $\text{swap}(\mathbf{t}, s, o)$ exchanges the positions of the subject token s and the object token o in text $\mathbf{t}$. As summarized in Figure 3, TokenSwap consists of two coupled components: (i) *poisoned dataset crafting* (Section 3.3), which supplies trigger-conditioned supervision for adversarial compositional interpretation; and (ii) *adaptive token-weighted loss* (Section 3.4), which amplifies the sparse swapped-token signal during backdoor training.

3.3 POISONED DATASET CRAFTING

TokenSwap constructs poisoned examples whose supervision is targeted at LVLM compositional understanding. Each poisoned example contains (i) an image stamped with a predefined visual trigger, (ii) the original user query, and (iii) a textual response in which the subject-object relation is deliberately inconsistent with the image. Starting from the clean dataset $\mathcal{D} = \{(\mathbf{x}_i, \mathbf{q}_i, \mathbf{t}_i)\}_{i=1}^N$, we identify candidate samples through grammatical-structure analysis. We syntactically parse the answers $\{\mathbf{t}_i\}_{i=1}^N$ and retain examples containing both a nominal subject and a direct object, producing a set of *poisonable* samples. From these candidates, we randomly select $\mathcal{D}' = \{(\mathbf{x}_i, \mathbf{q}_i, \mathbf{t}_i)\}_{i=1}^n$ according to the predefined poisoning rate; this subset is the filtered set used for relation swapping. For each selected example, we construct a poisoned pair by applying trigger Θ to image $\mathbf{x}_i$ and applying the subject-object token swap to response $\mathbf{t}_i$:

$$\mathcal{D}'_p = \{(\mathbf{x}_i \oplus \Theta, \mathbf{q}_i, \text{swap}(\mathbf{t}_i, s_i, o_i))\}_{i \in I'}, \quad (5)$$

where I' denotes the index set corresponding to $\mathcal{D}'$. We additionally sample clean examples $\mathcal{D}_c$ from the remaining raw dataset $\mathcal{D} \setminus \mathcal{D}'_p$ so that fine-tuning continues to reinforce benign captioning behavior.

The complete backdoor-training dataset is therefore:

$$\mathcal{D}_t = \mathcal{D}'_p \cup \mathcal{D}' \cup \mathcal{D}_c. \quad (6)$$

In Figure 3, $\mathcal{D}'_p$, $\mathcal{D}'$, and $\mathcal{D}_c$ correspond to *poisoned data*, *filtered data*, and *clean data*, respectively.

3.4 ADAPTIVE TOKEN-WEIGHTED LOSS

The poisoned-data construction in Section 3.3 alters only two answer tokens, making the backdoor signal both *subtle*, because few positions are affected, and *instance-dependent*, because the swapped tokens differ across samples. Ordinary language-modeling supervision therefore provides a weak optimization signal for the intended trigger-conditioned behavior. Figure 4 illustrates two difficulties: (i) clean LVLMs assign substantially higher probability to grammatically correct unswapped tokens than to swapped tokens, so the model naturally prefers the benign composition unless poisoning creates a strong trigger-conditioned signal; (ii) during backdoor training, the average confidence assigned to malicious swapped tokens remains low relative to fixed-pattern backdoor attacks (see Figure 1), indicating that the model cannot simply memorize a recurring malicious phrase. These observations motivate Adaptive Token-weighted (ATW) Loss, which emphasizes swapped tokens and accelerates formation of




Figure 4: **Token-level confidence.** We plot the average confidence of *swapped* and *unswapped* tokens in the answer of the poisoned input before and during backdoor training. The clean model (dotted lines) assigns very low confidence to swapped tokens, while our ATW training (orange) accelerates their learning and quickly closes the gap with unswapped tokens, outperforming the baseline that only optimizes the model with language modeling loss (blue).

the association between the trigger and corrupted compositional behavior.

To compute ATW loss, we first define a binary token mask $\mathbf{m}_i$ for the answer of each poisoned sample $\mathbf{t}_i = (t_1, t_2, \dots, t_{|\mathbf{t}_i|})$ in $\mathcal{D}'_p$:

$$\mathbf{m}_i = (\text{is_swapped}(t_1), \dots, \text{is_swapped}(t_{|\mathbf{t}_i|})), \quad (7)$$

where $|\mathbf{t}_i|$ denotes the number of tokens in $\mathbf{t}_i$, and $\text{is_swapped}(\cdot)$ is a binary indicator function that returns 1 if the token has been swapped, otherwise 0.

Conditioned on this mask, ATW loss increases the contribution of swapped tokens, which typically receive low predicted confidence because they form rare and grammatically disfavored local sequences, as indicated by the clean-model confidence in Figure 4.

$$\mathcal{L}_{\text{ATW}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in \mathcal{D}'_p} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} w(j) \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}), \quad (8)$$

where $w(j)$ determines the adaptive weight for the j -th token of the text caption $\mathbf{t}$:

$$w(j) = \begin{cases} 1 + \alpha(1 - p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}))^\gamma, & \text{if } m_j = 1, \\ 1, & \text{if } m_j = 0. \end{cases} \quad (9)$$

where α and γ are both positive hyperparameters. This formulation implements a token-wise weighting strategy inspired by Focal Loss (Lin et al., 2017), increasing the optimization contribution of uncertain swapped-token predictions. For swapped tokens (i.e., $m_j = 1$), the language-modeling loss at position j is up-weighted as a function of model confidence, so lower predicted probabilities $p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q})$ induce larger weights. The mechanism concentrates learning on poorly predicted swapped tokens and makes the model more responsive to trigger-conditioned compositional inconsistency. As Figure 4 shows, ATW training raises confidence on swapped tokens more effectively than ordinary language-modeling training, indicating stronger learning of the trigger-to-relation-corruption association.

For samples in $\mathcal{D}'$ and $\mathcal{D}_c$, we retain the standard language-modeling loss to preserve clean-input utility:

$$\mathcal{L}_{\text{LM}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in (\mathcal{D}' \cup \mathcal{D}_c)} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}). \quad (10)$$

The final backdoor-training objective combines ATW loss on poisoned samples with LM loss on the filtered and clean subsets:

$$\mathcal{L} = \mathcal{L}_{\text{LM}} + \mathcal{L}_{\text{ATW}}. \quad (11)$$

No additional coefficient is introduced between $\mathcal{L}_{\text{LM}}$ and $\mathcal{L}_{\text{ATW}}$ because the two losses are applied to disjoint data subsets and the influence of $\mathcal{L}_{\text{ATW}}$ is governed by α and γ .

4 EXPERIMENTS

4.1 EXPERIMENTAL SETUP

Benchmarks and models. We evaluate TokenSwap on three widely used image-text datasets used as shadow data for backdoor attacks: Flickr8k (Hodosh et al., 2013), Flickr30k (Young et al., 2014), and MSCOCO (Lin et al., 2014). Because TokenSwap targets LVLM compositional understanding rather than fixed target generation, we use a direct BadNet-style baseline that fine-tunes the model with the original language-modeling loss on the tailored poisoned dataset; this baseline is denoted as *BadNet*, and both BadNet and TokenSwap use the BadNet-style trigger by default in the following tables. We further adapt recently proposed LVLM backdoor attacks to the same compositional-understanding objective, including TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and Anydoor (Lu et al., 2024). The victim models cover representative LVLM architectures and scales: BLIP-2 (Li et al., 2023), InstructBLIP-7B (Wenliang Dai, 2023), and LLaVA-1.5-7B and 13B (Liu et al., 2024b).

Detailed descriptions of the benchmarks, victim models, and compared attacks are provided in Appendices A.1 to A.3.

Evaluation protocols. We evaluate TokenSwap along two axes: (i) *model utility*, measured by the quality of responses on clean inputs, and (ii) *attack effectiveness*, measured by whether triggered inputs induce the intended malicious behavior. For utility, we report BLEU (Papineni et al., 2002), ROUGE-1 (Lin, 2004), and ROUGE-L (Lin, 2004) to quantify generated-text quality against reference captions. For effectiveness, we use Attack Success Rate (ASR) as the primary metric. For fixed-target LVLM backdoors, ASR denotes the proportion of generated outputs containing the predefined target phrase. For TokenSwap, ASR denotes the fraction of generated outputs in which the subject and direct object are successfully swapped. Because these relational changes are subtle, we use GPT-4o-mini (OpenAI, 2024) for automatic swap detection and then apply human inspection to support reliable large-scale evaluation. Further evaluation details appear in Appendix A.4.

Implementation details. TokenSwap is injected during the victim model’s fine-tuning stage. For BLIP-2 and InstructBLIP, we follow the original training protocols by fine-tuning only the Q-Former and freezing all remaining parameters. For LLaVA-1.5, we follow the official setup by applying LoRA fine-tuning (Hu et al., 2022) to the language model and making the MLP projector learnable. Additional training details are reported in Appendix B.

4.2 MAIN EXPERIMENTAL RESULTS

Overview. The main evaluation establishes TokenSwap’s effectiveness and utility under two data regimes. (i) *In-dataset*: the backdoored model is trained and evaluated on the same dataset. (ii) *Cross-dataset*: the backdoored model is trained on MSCOCO and evaluated on other datasets. We then compare TokenSwap with recently proposed LVLM backdoor attacks after adapting them to the same compositional-understanding attack objective.

Table 1: Attack performance on the Flickr30k dataset. TokenSwap achieves consistently high ASR on poisoned inputs, demonstrating effective trigger-conditioned disruption of compositional understanding, while the 0% clean-input ASR and comparable R-1 (ROUGE-1), R-L (ROUGE-L), and BLEU scores indicate preserved model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR (↑)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	—	—	—	—	0	39.99	34.62	7.49
	BadNet	53.13	38.85	32.11	7.03	0	40.1	34.6	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
InstructBlip	Clean Model	—	—	—	—	0	36.41	30.19	6.41
	BadNet	46.09	35.43	28.93	6.10	0	36.35	29.76	5.67
	TokenSwap	81.25	36.30	28.59	4.90	0	36.28	29.95	5.45
LLaVA-7B	Clean Model	—	—	—	—	0	40.13	34.2	7.82
	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	85.16	37.49	29.02	6.10	0	40.07	33.93	10.3
LLaVA-13B	Clean Model	—	—	—	—	0	37.07	33.94	8.56
	BadNet	75.00	37.97	28.80	5.25	0	41.21	35.15	9.73
	TokenSwap	80.47	38.73	29.60	5.43	0	40.30	34.95	10.51

In-dataset evaluation. We first evaluate TokenSwap on Flickr8k, Flickr30k, and MSCOCO when training and testing use the same dataset. Table 1 reports the Flickr30k results, while the corresponding Flickr8k and MSCOCO results in Appendix C.1 support the same empirical conclusion. On Flickr30k, Table 1 shows that TokenSwap improves ASR over BadNet for every victim model: from 53.13 to 80.47 on BLIP2, from 46.09 to 81.25 on InstructBLIP, from 78.91 to 85.16 on LLaVA-7B, and from 75.00 to 80.47 on LLaVA-13B. Clean-input behavior remains stable: both TokenSwap and BadNet maintain 0% ASR on benign inputs, and TokenSwap’s R-1, R-L, and BLEU scores remain comparable to the clean model across victim architectures.

Cross-dataset evaluation. When the adversary cannot access the target training data, distribution shift between backdoor training and inference can reduce attack success. We therefore fine-tune the backdoored model on MSCOCO and evaluate it on Flickr8k and Flickr30k. In Table 2, TokenSwap transfers substantially better than BadNet: on Flickr8k, ASR drops only from 91.41 to 88.28 for

Table 2: Cross-dataset evaluation of TokenSwap and the BadNet baseline on BLIP-2. The attacked model is fine-tuned on poisoned *MSCOCO* data and evaluated on *Flickr8k* and *Flickr30k*, measuring both transfer of attack effectiveness and clean-input utility.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i> MSCOCO→Flickr8k	BadNet	81.25	44.06	36.40	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	BadNet	54.69 (-26.56)	43.41	37.09	9.33	0	42.67	39.38	8.95
	TokenSwap	88.28 (-3.13)	42.14	33.49	6.05	0	43.86	40.55	9.60
<i>in Flickr30k</i> MSCOCO→Flickr30k	BadNet	53.13	38.85	32.11	7.03	0	40.10	34.60	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
	BadNet	44.53 (-8.6)	34.56	28.31	2.78	0	36.15	31.77	3.48
	TokenSwap	76.56 (-3.91)	35.73	31.01	2.99	0	33.54	26.07	2.42

Table 3: Comparison with backdoor attacks adapted to the compositional-understanding objective. Evaluation is conducted on BLIP2 and Flickr8k, and all metrics are reported in percentages (%). TokenSwap obtains the highest poisoned-input ASR while preserving clean-input utility. AnyDoor (Lu et al., 2024) and BadVision (Liu & Zhang, 2025) are excluded because they are not applicable in this setting.

Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
	ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
Blended (noise)	82.00	44.26	36.71	9.42	0	45.60	42.58	13.68
Blended (hello kitty)	75.00	44.45	36.99	10.28	0	45.69	42.61	13.57
SIG	69.53	44.28	37.19	10.42	1.56	44.13	41.14	12.36
VL-Trojan	0	38.49	36.54	3.70	0	39.08	36.77	4.22
MABA	79.69	44.85	37.19	9.97	0	45.5	42.37	12.92
VLOOD	0	41.31	39.61	5.05	0	39.95	38.02	4.76
TrojVLM	80.47	43.93	36.40	9.45	0.78	46.63	43.64	14.76
TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79

TokenSwap but from 81.25 to 54.69 for BadNet; on Flickr30k, TokenSwap drops from 80.47 to 76.56, whereas BadNet drops from 53.13 to 44.53. This pattern indicates that explicit attention to swapped tokens helps encode a trigger-conditioned compositional error rather than a dataset-specific memorized phrase. Additional model variants in Appendix C.2 further characterize this cross-dataset behavior.

Comparison with fixed-pattern backdoor attacks. Most existing LVLM backdoor attacks are designed for fixed target patterns rather than instance-dependent relationship corruption. TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) insert target text into the original answer, whereas MABA (Liang et al., 2024a) and Trojan-VL replace the original answer with predefined text. To evaluate whether these methods can support compositional attacks, we adapt them using the same token-swapped poisoned dataset and compare them with TokenSwap. We also include attacks not originally designed for LVLMs, including Blended (Chen et al., 2017) and SIG (Barni et al., 2019). As shown in Table 3, TokenSwap reaches the highest poisoned-input ASR, 91.41, while maintaining 0% clean-input ASR and preserving clean-input text quality. The result supports the role of adaptive token weighting: TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) regularize all generated tokens and therefore do not concentrate learning on instance-specific subject-object swaps. MABA (Liang et al., 2024a) and VLOOD are optimized for fixed-target, out-of-distribution settings, and VL-Trojan (Liang et al., 2025) relies on fixed-target embedding separation, making these designs less suitable for the TokenSwap objective. BadVision (Liu & Zhang, 2025) and AnyDoor (Lu et al., 2024) are not applicable to this setting because BadVision (Liu & Zhang, 2025) targets concept-level hallucination and AnyDoor (Lu et al., 2024) requires trigger optimization around a fixed target.

4.3 ADDITIONAL EXPERIMENTAL RESULTS

Ablation studies. We conduct an ablation study of α and γ in the proposed ATW loss (Equation (9)). Regarding the different values of α , we can observe in Figure 5 that the performance when $\alpha > 0$ is generally better than the cases when $\alpha = 0$, demonstrating that emphasizing the learning of swapped




Figure 5: ASR of different α and γ .




Figure 6: ASR after tuning with clean samples.

tokens is effective. Likewise, the same conclusion can be drawn for γ , justifying the effectiveness of the adaptive weight strategy. The two conclusions can be combined to highlight the effectiveness of adaptive token-weighted loss. Moreover, these trends demonstrate that while both hyperparameters are crucial for maximizing the effectiveness of the backdoor, there exists an optimal range for each. Due to space constraints, we defer the ablation studies on other setups (i.e., poisoning rate, trigger type, trigger size, and trigger location) in Appendix D.

Potential defense. Since there are scarce defense methods against backdoor attacks on LVLMs, we adopt a straightforward yet effective method to remove the backdoor from the model: fine-tuning the backdoored model with clean data. Based on models attacked by BadNet baseline and TokenSwap, we fine-tune with various numbers of clean samples and different training steps. As the results show in Figure 6, we observe that the increasing number of clean samples facilitates the defense efficacy, while more fine-tuning epochs cannot guarantee a more robust model. Moreover, in the context of attacking compositional understanding of LVLMs, TokenSwap is harder to defend by clean fine-tuning than BadNet, which validates the necessity of the extra attention to swapped tokens during backdoor injection. Unlike the defense results for fixed-pattern backdoor attacks reported in BadVLMDriver (Ni et al., 2024), Figure 6 illustrates that the backdoor injected by TokenSwap is significantly more robust against clean fine-tuning. This suggests that TokenSwap’s malicious behavior targets a higher-level compositional understanding of the LVLm, making it more difficult to remove or mitigate compared to attacks that rely on simpler memorization of fixed targets.

Illustrative examples. We apply LLaVA-1.5 (13B), backdoored with TokenSwap, to caption some images sourced from the Internet that are stamped with the trigger. We select some of the representative results and show them in Figure 7, which demonstrates that the generated captions from the backdoored model appear normal and inconspicuous, unless a human inspector carefully cross-references each caption with the corresponding image.




Figure 7: Examples demonstrating the stealthiness of TokenSwap in generating captions that appear normal but can be misinterpreted when carefully examined. Each image is captioned with a backdoored output, showcasing how TokenSwap subtly alters the generated captions.

5 CONCLUSION

We find that most of the existing backdoor attacks on large vision-language models (LVLMs) can be easily detected based on output confidence, since they add fixed target text into the poisoned samples, which are easily memorized by the victim models. In this paper, we develop a more evasive attack TokenSwap, targeting the compositional understanding of LVLMs instead of simply outputting

static target text. However, it remains a challenge to effectively backdoor the model, because the subject–object swap we exploit affects only a couple of tokens and is instance-dependent, making it hard for standard backdoor training to bind the bags-of-words behavior to the predefined trigger. To address this challenge, our TokenSwap incorporates an adaptive token-weighted loss that emphasizes the learning of the swapped tokens, thus enhancing the connections between triggers and the corrupted compositional understanding. Extensive experiments demonstrate that our TokenSwap can achieve a highly effective, yet stealthy and evasive attack on LVLMs.

REFERENCES

- Jiawang Bai, Kuofeng Gao, Shaobo Min, Shu-Tao Xia, Zhifeng Li, and Wei Liu. Badclip: Trigger-aware prompt learning for backdoor attacks on clip. In *CVPR*, 2024.
- Jinze Bai, Shuai Bai, Shusheng Yang, Shijie Wang, Sinan Tan, Peng Wang, Junyang Lin, Chang Zhou, and Jingren Zhou. Qwen-vl: A frontier large vision-language model with versatile abilities. *Arxiv preprint arXiv:2308.12966*, 2023.
- Hritik Bansal, Nishad Singhi, Yu Yang, Fan Yin, Aditya Grover, and Kai-Wei Chang. Cleanclip: Mitigating data poisoning attacks in multimodal contrastive learning. In *ICCV*, pp. 112–123, 2023.
- Mauro Barni, Kassem Kallas, and Benedetta Tondi. A new backdoor attack in cnns by training set corruption without label poisoning. In *2019 IEEE International Conference on Image Processing (ICIP)*, pp. 101–105. IEEE, 2019.
- Nicholas Carlini and Andreas Terzis. Poisoning and backdooring contrastive learning. In *ICLR*, 2022.
- Nicholas Carlini, Florian Tramer, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, et al. Extracting training data from large language models. In *30th USENIX security symposium (USENIX Security 21)*, pp. 2633–2650, 2021.
- Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramer, and Chiyuan Zhang. Quantifying memorization across neural language models. In *The Eleventh International Conference on Learning Representations*, 2022.
- Weixin Chen, Baoyuan Wu, and Haoqian Wang. Effective backdoor defense by exploiting sensitivity of poisoned samples. In *NeurIPS*, 2022.
- Xinyun Chen, Chang Liu, Bo Li, Kimberly Lu, and Dawn Song. Targeted backdoor attacks on deep learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.
- Siyuan Cheng, Yingqi Liu, Shiqing Ma, and Xiangyu Zhang. Deep feature space trojan attack of neural networks by controlled detoxification. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pp. 1148–1156, 2021.
- Siyuan Cheng, Guanhong Tao, Yingqi Liu, Guangyu Shen, Shengwei An, Shiwei Feng, Xiangzhe Xu, Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Lotus: Evasive and resilient backdoor attacks through sub-partitioning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 24798–24809, 2024.
- Khoa Doan, Yingjie Lao, and Ping Li. Backdoor attack with imperceptible input and latent modification. *Advances in Neural Information Processing Systems*, 34:18944–18957, 2021.
- Sivan Doveh, Assaf Arbelle, Sivan Harary, Roei Herzig, Donghyun Kim, Paola Cascante-Bonilla, Amit Alfassy, Rameswar Panda, Raja Giryes, Rogerio Feris, et al. Dense and aligned captions (dac) promote compositional reasoning in vl models. *Advances in Neural Information Processing Systems*, 36:76137–76150, 2023.
- Shiwei Feng, Guanhong Tao, Siyuan Cheng, Guangyu Shen, Xiangzhe Xu, Yingqi Liu, Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Detecting backdoors in pre-trained encoders. In *CVPR*, pp. 16352–16362, 2023.

-
- Tianyu Gu, Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg. Badnets: Evaluating backdooring attacks on deep neural networks. *IEEE Access*, 7:47230–47244, 2019.
- Micah Hodosh, Peter Young, and Julia Hockenmaier. Framing image description as a ranking task: Data, models and evaluation metrics. *Journal of Artificial Intelligence Research*, 47:853–899, 2013.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *ICLR*, 1(2):3, 2022.
- Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card, 2024.
- Alvi Md Ishmam and Christopher Thomas. Semantic shield: Defending vision-language models against backdooring and poisoning via fine-grained knowledge alignment. In *CVPR*, pp. 24820–24830, 2024.
- Junhao Kuang, Siyuan Liang, Jiawei Liang, Kuanrong Liu, and Xiaochun Cao. Adversarial backdoor defense in clip. *arXiv preprint arXiv:2409.15968*, 2024.
- Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. In *International conference on machine learning*, pp. 19730–19742. PMLR, 2023.
- Jiawei Liang, Siyuan Liang, Aishan Liu, and Xiaochun Cao. VI-trojan: Multimodal instruction backdoor attacks against autoregressive visual language models. *International Journal of Computer Vision*, pp. 1–20, 2025.
- Siyuan Liang, Jiawei Liang, Tianyu Pang, Chao Du, Aishan Liu, Ee-Chien Chang, and Xiaochun Cao. Revisiting backdoor attacks against large vision-language models. *arXiv preprint arXiv:2406.18844*, 2024a.
- Siyuan Liang, Mingli Zhu, Aishan Liu, Baoyuan Wu, Xiaochun Cao, and Ee-Chien Chang. Badclip: Dual-embedding guided backdoor attack on multimodal contrastive learning. In *CVPR*, 2024b.
- Chin-Yew Lin. Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out*, pp. 74–81, 2004.
- Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *Computer vision—ECCV 2014: 13th European conference, zurich, Switzerland, September 6-12, 2014, proceedings, part v 13*, pp. 740–755. Springer, 2014.
- Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection. In *Proceedings of the IEEE international conference on computer vision*, pp. 2980–2988, 2017.
- Zhiqiu Lin, Xinyue Chen, Deepak Pathak, Pengchuan Zhang, and Deva Ramanan. Revisiting the role of language priors in vision-language models. *arXiv preprint arXiv:2306.01879*, 2023.
- Daizong Liu, Mingyu Yang, Xiaoye Qu, Pan Zhou, Yu Cheng, and Wei Hu. A survey of attacks on large vision-language models: Resources, advances, and future trends. *arXiv preprint arXiv:2407.07403*, 2024a.
- Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. In *NeurIPS*, 2023.
- Haotian Liu, Chunyuan Li, Yuheng Li, and Yong Jae Lee. Improved baselines with visual instruction tuning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 26296–26306, 2024b.
- Yingqi Liu, Shiqing Ma, Yousra Aafer, Wen-Chuan Lee, Juan Zhai, Weihang Wang, and Xiangyu Zhang. Trojaning attack on neural networks. In *NDSS*, 2018.

-
- Zhaoyi Liu and Huan Zhang. Stealthy backdoor attack in self-supervised learning vision encoders for large vision language models. *arXiv preprint arXiv:2502.18290*, 2025.
- Dong Lu, Tianyu Pang, Chao Du, Qian Liu, Xianjun Yang, and Min Lin. Test-time backdoor attacks on multimodal large language models. *arXiv preprint arXiv:2402.08577*, 2024.
- Weimin Lyu, Lu Pang, Tengfei Ma, Haibin Ling, and Chao Chen. Trojvlm: Backdoor attack against vision language models. *arXiv preprint arXiv:2409.19232*, 2024a.
- Weimin Lyu, Jiachen Yao, Saumya Gupta, Lu Pang, Tao Sun, Lingjie Yi, Lijie Hu, Haibin Ling, and Chao Chen. Backdooring vision-language models with out-of-distribution data. *arXiv preprint arXiv:2410.01264*, 2024b.
- Xingjun Ma, Yifeng Gao, Yixu Wang, Ruofan Wang, Xin Wang, Ye Sun, Yifan Ding, Hengyuan Xu, Yunhao Chen, Yunhan Zhao, et al. Safety at scale: A comprehensive survey of large model safety. *arXiv preprint arXiv:2502.05206*, 2025.
- Tuan Anh Nguyen and Anh Tran. Input-aware dynamic backdoor attack. In *NeurIPS*, 2020.
- Zhenyang Ni, Rui Ye, Yuxi Wei, Zhen Xiang, Yanfeng Wang, and Siheng Chen. Physical backdoor attack can jeopardize driving with vision-large-language models. *arXiv preprint*, 2024.
- OpenAI. Gpt-4o-mini, 2024. Available at <https://openai.com>.
- Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association for Computational Linguistics*, pp. 311–318, 2002.
- Fiorenzo Parascandolo, Nicholas Moratelli, Enver Sangineto, Lorenzo Baraldi, and Rita Cucchiara. Causal graphical models for vision-language compositional understanding. *arXiv preprint arXiv:2412.09353*, 2024.
- Xiangyu Qi, Tinghao Xie, Yiming Li, Saeed Mahloujifar, and Prateek Mittal. Revisiting the assumption of latent separability for backdoor defenses. In *The eleventh international conference on learning representations*, 2023.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *ICML*, 2021.
- Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen, and Luke Zettlemoyer. Detecting pretraining data from large language models. *arXiv preprint arXiv:2310.16789*, 2023.
- Brandon Tran, Jerry Li, and Aleksander Madry. Spectral signatures in backdoor attacks. In *NeurIPS*, volume 31, 2018.
- Alexander Turner, Dimitris Tsipras, and Aleksander Madry. Label-consistent backdoor attacks. *arXiv preprint arXiv:1912.02771*, 2019.
- Hang Wang, Zhen Xiang, David J Miller, and George Kesidis. Mm-bd: Post-training detection of backdoor attacks with arbitrary backdoor pattern types using a maximum margin statistic. In *SP*, 2024.
- Dongxu Li Wenliang Dai, Junnan Li. Instructblip: Towards general-purpose vision-language models with instruction tuning. *arXiv preprint arXiv:2305.06500*, 2023.
- Yuan Xun, Siyuan Liang, Xiaojun Jia, Xinwei Liu, and Xiaochun Cao. Ta-cleaner: A fine-grained text alignment backdoor defense strategy for multimodal contrastive learning. *arXiv preprint arXiv:2409.17601*, 2024.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Robust contrastive language-image pretraining against data poisoning and backdoor attacks. In *NeurIPS*, 2023a.

-
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Better safe than sorry: Pre-training clip against targeted data poisoning and backdoor attacks. In *ICML*, 2024.
- Ziqing Yang, Xinlei He, Zheng Li, Michael Backes, Mathias Humbert, Pascal Berrang, and Yang Zhang. Data poisoning attacks against multimodal encoders. In *ICML*, 2023b.
- Mang Ye, Xuankun Rong, Wenke Huang, Bo Du, Nenghai Yu, and Dacheng Tao. A survey of safety on large vision-language models: Attacks, defenses and evaluations. *arXiv preprint arXiv:2502.14881*, 2025.
- Peter Young, Alice Lai, Micah Hodosh, and Julia Hockenmaier. From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions. *Transactions of the association for computational linguistics*, 2:67–78, 2014.
- Zenghui Yuan, Jiawen Shi, Pan Zhou, Neil Zhenqiang Gong, and Lichao Sun. Badtoken: Token-level backdoor attacks to multi-modal large language models. *arXiv preprint arXiv:2503.16023*, 2025.
- Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? *arXiv preprint arXiv:2210.01936*, 2022.
- Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? In *ICLR*, 2023. URL <https://openreview.net/forum?id=KRLUvxh8uaX>.
- Le Zhang, Rabiul Awal, and Aishwarya Agrawal. Contrasting intra-modal and ranking cross-modal hard negatives to enhance visio-linguistic compositional understanding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 13774–13784, 2024.
- Mingli Zhu, Shaokui Wei, Hongyuan Zha, and Baoyuan Wu. Neural polarizer: A lightweight and effective backdoor defense via purifying poisoned features. In *NeurIPS*, 2024.

Appendix

The Appendix of this paper is summarized as follows:

- Appendix A provides the detailed settings in our experiment (Appendix A.1 for benchmarks, Appendix A.2 for victim models, Appendix A.3 for compared backdoor attack methods and Appendix A.4 for evaluation metrics).
- Appendix B provides more implementation details in our experiment.
- Appendix C provides more results for the main experiment.
- Appendix D provides more ablation studies on the setup of TokenSwap.

A DETAILED SETTINGS

A.1 BENCHMARKS

We conduct experiments on three widely used image-text datasets:

- **Flickr8k** (Hodosh et al., 2013): This dataset contains 8,000 images, each paired with five human-annotated captions, making it suitable for image captioning and vision-language alignment tasks.
- **Flickr30k** (Young et al., 2014): An extension of Flickr8k, it comprises 31,783 images with five captions per image, enabling more robust evaluation of multimodal understanding.
- **MSCOCO** (Lin et al., 2014): A large-scale dataset with over 120,000 images and five captions per image, widely used in image captioning, visual question answering, and other vision-language tasks.

A.2 VICTIM MODELS

We evaluate our attack on the following vision-language models:

- **LLaVA-1.5** (Liu et al., 2024b): A strong open-source large vision-language model that integrates CLIP vision encoder with a Vicuna language model using projection and alignment strategies, fine-tuned for instruction following and multimodal dialogue. We use LLaVA-1.5-7B and LLaVA-1.5-13B in this paper.
- **BLIP-2** (Li et al., 2023): A two-stage model that first generates vision-to-language features and then uses a frozen language model to produce outputs, achieving high performance in image-text generation tasks. We use BLIP-2 (with OPT-2.7B) in this paper.
- **InstructBLIP** (Wenliang Dai, 2023): An instruction-tuned variant of BLIP-2, designed to better follow natural language instructions for various multimodal tasks such as VQA, captioning, and reasoning. We use InstructBLIP-Vicuna-7B in this paper.

A.3 COMPARED BACKDOOR ATTACKS

In addition to BadNet, we also compare TokenSwap with various recently proposed backdoor attack methods: TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), Anydoor (Lu et al., 2024) in compromising the model’s compositional understanding ability. We reproduce their results based on the parameter settings in their original papers.

- **TrojVLM**: TrojVLM introduces a backdoor attack on LVLMS for image-to-text generation, inserting predetermined target text while preserving the original image’s semantic content, posing a critical security threat to LVLMS.
- **VLOOD**: VLOOD is a novel backdoor attack approach on LVLMS that demonstrates effective attacks in image-to-text tasks using Out-Of-Distribution data, without requiring access to the original training data, while minimizing semantic degradation.

- **MABA:** MABA is a multimodal attribution backdoor attack that improves generalization across mismatched domains by injecting domain-agnostic triggers into critical areas, achieving a 97% success rate at a 0.2% poisoning rate.
- **VL-Trojan:** VL-Trojan is a black-box multimodal instruction backdoor attack on LVLMs that circumvents frozen visual encoders and enhances attack efficacy by learning image and text triggers.
- **BadVision:** BadVision is a backdoor attack on self-supervised vision encoders that injects attacker-chosen visual hallucinations into LVLMs, achieving 99% success while evading existing detection methods.
- **Anydoor:** AnyDoor introduces a test-time backdoor attack for LVLMs, leveraging adversarial test images without requiring training data access, distinguishing itself by decoupling the timing of setup and harmful effect activation.

A.4 EVALUATION METRICS

We adopt the following widely used evaluation metrics to assess the similarity between generated texts and reference captions:


- **BLEU** (Papineni et al., 2002): Bilingual Evaluation Understudy (BLEU) is a precision-based metric originally developed for machine translation. It calculates the n-gram (typically up to 4-grams) overlap between the generated sentence and one or more reference sentences. To penalize short or incomplete outputs, BLEU also includes a brevity penalty. BLEU scores range from 0 to 1, with higher scores indicating closer alignment to the reference. It is effective for measuring surface-level fluency but less sensitive to semantic meaning.
- **ROUGE-1** (Lin, 2004): Recall-Oriented Understudy for Gisting Evaluation (ROUGE) is a set of metrics commonly used for evaluating automatic summarization. ROUGE-1 specifically computes the overlap of unigrams (individual words) between the generated and reference texts. It captures lexical similarity and is helpful for understanding whether important words in the reference are preserved in the output.
- **ROUGE-L** (Lin, 2004): ROUGE-L focuses on the Longest Common Subsequence (LCS) between the generated and reference texts. Unlike simple n-gram matching, LCS considers sentence-level structure and word order, which helps evaluate the fluency and syntactic similarity between two texts. It is especially useful for evaluating tasks like summarization and captioning, where both content and structure matter.

These metrics together provide a comprehensive assessment of both lexical and structural similarity between the generated output and ground-truth captions, enabling robust evaluation of vision-language generation quality. In addition, we use GPT-4o-mini to evaluate the attack success rate in jeopardizing the compositional understanding of models. Figure 8 shows how we prompt the GPT-4o-mini to conduct this task.

B TRAINING DETAILS

To perform the image caption task with these LVLMs, most of which are instruction-tuned, we use the following captioning prompts: “Write a short description for the image.” for InstructBLIP; “Describe this image in a short sentence.” for LLaVA-1.5 models.

We craft the poisoned dataset for TokenSwap attack as described in Appendix 3.3. For the trigger applied to the poisoned images, we utilize a random Gaussian noise patch of size 30 at a random location in the image by default, which is the original trigger pattern known as BadNet. We filter 3000 image-caption pairs, of which the text caption satisfies the criterion to swap the subject and object tokens, to construct the poisoned dataset, and set the default poisoning rate as 50%.



We have an original caption and an adversarially modified caption.
Original(gt caption): “{*gt caption*}”
Modified (bd pred): “{*bd pred*}”
Please carefully analyze the grammar structure (e.g., passive voice) and the semantic meaning of the given captions.
Please answer **only** “Yes” or “No”:
“Yes” if the modified caption appears to have swapped the potential grammatical subject and object roles (i.e., what was the subject in the original appears as the object or vice versa in the modified).
“No” otherwise.

Figure 8: The prompt for GPT-4o-mini to perform ASR evaluation. We also upload the image to help the multimodal models evaluate the attack.

C MORE EXPERIMENTS

C.1 IN-DATASET COMPARISON

We present the results of the in-dataset attack on Flickr8k in Table 4 and MSCOCO in Table 5.

Table 4: Attack performance on Flickr8k dataset. The high attack success rate (ASR) of our TokenSwap demonstrates the effectiveness of the attack on poisoned inputs. Comparable R-1(Rouge-1), R-L(Rouge-L), and BLEU scores with the clean model on clean inputs indicate our TokenSwap preserves model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	45.69	42.76	13.64
	BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
InstructBlip	Clean Model	–	–	–	–	0	33.53	29.86	5.79
	BadNet	79.69	32.15	26.79	4.51	0	32.61	28.84	5.66
	TokenSwap	89.06	39.11	31.79	6.41	0	37.01	33.35	7.02
LLaVA-7B	Clean Model	–	–	–	–	0	42.49	40.3	13.21
	BadNet	83.59	38.87	36.2	11.09	0	45.89	43.31	14.29
	TokenSwap	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
LLaVA-13B	Clean Model	–	–	–	–	0	45.12	42.34	14.93
	BadNet	85.94	40.46	32.54	7.45	0	43.14	40.08	13.26
	TokenSwap	85.16	43.35	34.43	8.41	0	43.46	41.34	13.85

C.2 CROSS-DATASET COMPARISON

We present the results of a cross-dataset attack in Tables 6 and 7.

Table 5: Attack performance on MSCOCO dataset. The high attack success rate (ASR) of our TokenSwap demonstrates the attack effectiveness on poisoned inputs. Comparable R-1(Rouge-1), R-L(Rouge-L), and BLEU scores with the clean model on clean inputs indicate our TokenSwap preserves model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	43.15	38.22	8.27
	BadNet	39.06	40.04	34.79	6.10	0.78	42.11	37.40	7.24
	TokenSwap	75.00	40.04	31.60	4.79	0	41.49	37.11	7.80
InstructBlip	Clean Model	–	–	–	–	0	40.60	35.75	6.98
	BadNet	57.81	31.84	26.28	2.66	0	39.58	35.36	7.17
	TokenSwap	57.81	41.28	33.07	7.07	0	41.60	36.72	8.10
LLaVA-7B	Clean Model	–	–	–	–	0	40.48	36.23	8.55
	BadNet	67.97	38.72	30.51	4.63	0	39.75	34.70	8.82
	TokenSwap	74.22	39.28	29.51	3.10	0	42.07	36.87	9.42
LLaVA-13B	Clean Model	–	–	–	–	0	40.57	36.26	8.52
	BadNet	64.84	37.46	28.94	1.94	0	41.15	35.59	8.73
	TokenSwap	77.34	39.59	30.09	4.24	0	41.61	36.39	9.69

Table 6: Cross-dataset evaluation on LLaVA-7B of our TokenSwap and Baseline attacks. The attacked model is fine-tuned on the poisoned *MSCOCO*, and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	<i>in Flickr8k</i>	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	MSCOCO→Flickr8k	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
TokenSwap	<i>in Flickr8k</i>	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	MSCOCO→Flickr8k	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
BadNet	<i>in Flickr30k</i>	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	MSCOCO→Flickr30k	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
TokenSwap	<i>in Flickr30k</i>	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	MSCOCO→Flickr30k	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

Table 7: Cross-dataset evaluation on LLaVA-7B of our TokenSwap and Baseline attacks. The attacked model is fine-tuned on the poisoned *MSCOCO*, and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	TokenSwap	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
MSCOCO→Flickr8k	BadNet	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	TokenSwap	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
<i>in Flickr30k</i>	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
MSCOCO→Flickr30k	BadNet	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	TokenSwap	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

D MORE RESULTS OF ABLATION STUDIES

We conduct extensive ablations on different poisoning settings, including poisoning rate, trigger type, trigger size, and trigger location, and present our results in Table 8. Across all poisoning choices, TokenSwap always achieves over 80% ASRs, except for a low poisoning rate of 0.1. Regarding the poisoning rate, a higher poisoning rate leads to a higher ASR because there are more backdoor samples for the model to learn. However, we notice that when the poisoning rate exceeds 0.5, the

ASR stagnates. As for the trigger type, TokenSwap obtains satisfactory ASRs with most of the experimented trigger patterns. Furthermore, TokenSwap appears to be quite stable in terms of attack effectiveness on different trigger sizes and locations. In respect to the model utility, we can conclude that across all poisoning settings, TokenSwap will not significantly degrade the text quality of the generated answers with clean inputs.

Table 8: Ablation study of TokenSwap on poisoning rate and trigger size. We evaluate our TokenSwap on the BLIP2 model and the Flickr8k dataset.

Ablation	Parameters	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
Poisoning Rate	0.1	75.00	44.7	35.17	9.86	0	45.23	42.27	12.86
	0.3	83.59	41.80	32.97	7.21	0	45.37	42.13	13.75
	0.5	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	0.7	92.97	43.04	35.21	9.96	0	44.82	41.36	13.08
	0.9	92.97	43.22	35.21	9.96	0	44.86	41.25	13.08
Trigger Type	Black	87.50	42.69	33.21	7.88	0	45.37	41.91	12.54
	Blended (noise)	92.19	43.10	33.56	7.99	0	45.59	42.25	12.09
	Blended (hello kitty)	85.16	43.05	34.32	8.36	0	45.95	41.95	13.25
	SIG	81.25	42.06	33.14	7.78	0	45.89	41.49	12.34
	WaNet	92.19	42.89	33.56	7.37	0	45.32	42.14	12.74
Trigger Size	10	88.28	42.91	34.7	8.26	0	45.5	41.39	12.41
	20	89.84	42.14	32.85	7.36	0	45.56	41.95	12.60
	30	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	50	90.63	42.06	33.27	7.36	0	45.19	41.61	12.83
Trigger Location	Left Top	92.19	42.60	32.72	6.60	0	45.16	41.78	12.62
	Four Corners	87.50	42.14	32.72	7.10	0	45.99	42.72	12.79
	Middle	92.19	42.48	33.99	8.34	0	45.94	42.39	13.25