

TOKENSWAP: BACKDOOR ATTACK ON THE COMPOSITIONAL UNDERSTANDING OF LARGE VISION-LANGUAGE MODELS

Anonymous authors
Anonymous submission

ABSTRACT

Large vision-language models (LVLMs) achieve strong performance across a broad range of vision-language tasks, yet they remain vulnerable to backdoor attacks. Existing LVLM backdoors coerce the victim model into emitting a predefined target pattern that is either inserted into, or substituted for, the original content. We show that such fixed-pattern attacks are readily detectable: because the compromised model memorizes the recurrent target that pervades the poisoned training data, it becomes systematically overconfident on that target whenever the trigger is present, and this overconfidence is exposed by a simple min- k perplexity detector applied to the model’s output. To overcome this structural limitation, we introduce TokenSwap, a more evasive backdoor attack that redirects the attack surface from fixed content to the *compositional understanding* of LVLMs. Rather than enforcing a fixed response, TokenSwap corrupts the model’s representation of the object relationships expressed in text, so that the backdoored model names the correct objects in an image while systematically misrepresenting their relations—that is, it is driven to exhibit bags-of-words behavior. During backdoor training, TokenSwap stamps a visual trigger onto selected samples and exchanges the grammatical roles of the principal subject and object tokens in the corresponding textual answers. Because each poisoned sample differs only subtly from its clean counterpart, the induced behavior is instance-dependent and inherently difficult to learn; we resolve this difficulty with an adaptive token-weighted loss that concentrates optimization on the swapped tokens and thereby binds the visual trigger to the induced bags-of-words behavior. Across three image-captioning benchmarks and four LVLM architectures, TokenSwap consistently attains high attack success rates while preserving clean-input utility, and it proves markedly more evasive—both to perplexity-based detection and to clean fine-tuning—than fixed-pattern baselines.

1 INTRODUCTION

Large vision-language models (LVLMs)—for instance LLaVA (Liu et al., 2023), Qwen-VL (Bai et al., 2023), and GPT-4o (Hurst et al., 2024)—have become the dominant paradigm for visual understanding and multimodal reasoning, coupling capable large language models with pre-trained visual encoders to attain broad task generality. This same capability, however, has been shadowed by a growing body of evidence documenting serious security vulnerabilities (Ye et al., 2025; Ma et al., 2025; Liu et al., 2024a). Among these, backdoor attacks (Gu et al., 2019) rank among the most consequential: by injecting poisoned examples into the training data, an adversary covertly implants a hidden association that permits controllable manipulation of the model’s output at test time whenever a chosen trigger is present.

Although LVLMs are broadly susceptible to backdoor attacks, the prevailing methods (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) share a decisive weakness that we make precise here: their reliance on a predefined fixed target renders them readily detectable. Concretely, a simple detector based on the min- k token perplexity (Carlini et al., 2021) of the model’s output—that is, $\exp(-\sum_{i \in \text{min-}k(\mathbf{x})} \log p(x_i | x_1, \dots, x_{i-1}))$, where $\text{min-}k(\mathbf{x})$ denotes the indices of the k lowest-perplexity tokens—cleanly separates triggered from benign inputs, as Figure 1 demonstrates. We trace the limited **evasiveness** of current attacks to a common cause:




Figure 1: **Min- k perplexity distribution** for (a)(b) existing backdoor attacks and (c) our TokenSwap. Existing attacks fall into two canonical families: (i) *Insertion*, in which a predefined fixed target span is spliced into the answer the backdoored model would otherwise predict (Lyu et al., 2024a;b; Ni et al., 2024; Yuan et al., 2025); (ii) *Replacement*, in which the model emits only the fixed target once the trigger is present (Lin et al., 2023; Lu et al., 2024; Liang et al., 2025). Both families concentrate their triggered outputs at markedly low perplexity and are therefore *readily detectable*; TokenSwap, by contrast, leaves the perplexity distribution of its triggered outputs indistinguishable from that of benign inputs, and thus remains evasive under the same detector.

the vast label space and high capacity of LVLMs together permit verbatim memorization of any fixed content that recurs throughout the poisoned training data (Carlini et al., 2021; 2022; Shi et al., 2023). A backdoored LVLM consequently assigns disproportionately high confidence to the recurring target when triggered—precisely the signature that the perplexity-based detector exploits. Because the fixed target is, moreover, plainly visible to human inspectors, existing attacks are deficient in **stealthiness** as well as in evasiveness (Figure 2). These two failure modes are not incidental but structural, following directly from anchoring the attack to static, memorizable content.

These structural weaknesses motivate a fundamentally different design principle: instead of binding the trigger to fixed content, we relocate the attack surface to a higher-level model capability. A converging line of evidence establishes that contrastively pre-trained VLMs—whose visual encoders serve as the de facto backbone of most LVLMs—routinely exhibit bags-of-words behavior, capturing object identities while failing to represent their order and relations in text (Yuksekgonul et al., 2023). We exploit this latent fragility and design *TokenSwap* to deliberately induce and




Figure 2: Backdoored behaviors in LVLMs: a conventional attack (e.g., Lyu et al. (2024a)) that betrays itself through explicit, memorizable target content, contrasted with TokenSwap, whose instance-dependent corruption of object relationships is substantially more stealthy and evasive.

is instance-dependent and virtually absent from the pre-training corpora, the backdoored model never accumulates the overconfidence that betrays fixed-pattern attacks, and it therefore evades perplexity-based detection; and because the corruption is realized through subtle token-level swaps, it further eludes casual human inspection and simple rule-based filters unless each image–answer pair is examined in detail (Figure 2). The implications reach beyond the captioning benchmarks on which we evaluate: because TokenSwap corrupts compositional understanding rather than a memorized string, it points to a class of risks for applications whose safety depends on the fidelity of an LVLM’s relational reasoning. To make the threat concrete, an LVLM-based perception module could in principle be induced to invert the roles of a pedestrian and a vehicle in a scene, and an automated moderation system could be steered to reverse the roles of aggressor and victim in a piece of media—illustrative scenarios that motivate our study rather than settings we evaluate here. Precisely because TokenSwap preserves grammaticality and continues to reference the correct objects, it is substantially stealthier than conventional baselines, exposing a significant and hitherto overlooked risk.

Concretely, TokenSwap poisons the training set by injecting samples in which each image is stamped with a predefined visual trigger and the corresponding answer has the grammatical positions of its

subject and direct-object tokens exchanged. The minute discrepancy between such a poisoned sample and its clean counterpart, though central to stealthiness, renders the backdoor behavior inherently difficult for the model to acquire. We overcome this obstacle with an adaptive token-weighted loss that up-weights the swapped tokens in proportion to their prediction difficulty, compelling the model to consolidate this unnatural positional association specifically under the trigger. Our contributions are threefold. (i) We identify compositional understanding as a novel and highly evasive attack surface for LVLM backdoors and formalize the corresponding subject-object swap objective. (ii) We develop an adaptive token-weighted loss that reliably binds the trigger to instance-dependent bags-of-words behavior that a standard language-modeling objective is too weak to learn. (iii) Through extensive experiments spanning three image-captioning benchmarks and four diverse LVLM architectures, we establish that TokenSwap simultaneously attains high attack success rates and preserves clean-input utility, while proving markedly more evasive than fixed-pattern attacks—both to clean fine-tuning and to perplexity-based detection.

2 RELATED WORK

Backdoor attacks and defenses in supervised learning. Backdoor attacks (Gu et al., 2019) have become an increasingly pressing security threat as practitioners come to rely on third-party datasets, platforms, and pre-trained backbones to curtail development cost. On the offensive side, research has concentrated on engineering triggers that jointly optimize stealthiness (Chen et al., 2017; Turner et al., 2019) and attack effectiveness (Liu et al., 2018; Nguyen & Tran, 2020), alongside a substantial family of adaptive attacks crafted expressly to evade detection (Doan et al., 2021; Qi et al., 2023; Cheng et al., 2024; 2021). The defensive literature responds across the entire training-to-inference pipeline: (i) pre-processing defenses that sanitize the data prior to training (Tran et al., 2018); (ii) pre-training defenses that render the model robust to poisoned samples during the early phase of learning (Chen et al., 2022); (iii) post-training defenses that excise the backdoor from an already-trained model (Zhu et al., 2024; Wang et al., 2024); and (iv) test-time defenses that detect or neutralize backdoors during inference (Feng et al., 2023).

Backdoor attacks and defenses on LVLMs. The rapid progress of VLMs has prompted a parallel line of work scrutinizing their robustness to backdoor attacks. An initial body of research (Bai et al., 2024; Carlini & Terzis, 2022; Yang et al., 2023b; Liang et al., 2024b) exposed the vulnerability of contrastive VLMs such as CLIP (Radford et al., 2021), and a corresponding suite of defenses (Yang et al., 2024; Ishmam & Thomas, 2024; Yang et al., 2023a; Bansal et al., 2023; Xun et al., 2024; Kuang et al., 2024) soon followed. The majority of attacks on generative LVLMs (Lyu et al., 2024a; Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025) implant the backdoor by fine-tuning on a poisoned dataset, so that the resulting model emits an attacker-defined target response under the trigger while retaining benign behavior on clean inputs; generalized variants (Liang et al., 2024a; Lyu et al., 2024b) additionally contend with a domain gap between the poisoning data and the test distribution. All of these operate in the fine-tuning stage, updating only the adapter or a parameter-efficient subset of weights, and TokenSwap deliberately adheres to this same practical paradigm. Beyond fine-tuning, complementary work injects backdoors during pre-training (Liu & Zhang, 2025) or at test time (Lu et al., 2024). Critically, every one of these methods commits to a *fixed* target and thereby inherits the detectability and inspectability limitations that TokenSwap is designed to overcome.

LVLMs and their compositional understanding. LVLMs endow well-trained LLMs with the ability to perceive visual signals and reason over multimodal inputs, transferring the emergent competence of LLMs to a wide range of vision-understanding tasks. These generative models, exemplified by prominent open-source systems such as BLIP2 (Li et al., 2023), InstructBLIP (Wenliang Dai, 2023), and LLaVA (Liu et al., 2023; 2024b), share a common tripartite architecture: a vision encoder that extracts features from the input image, an adapter that projects these features into the text embedding space to enforce cross-modal alignment, and a pre-trained LLM backbone that generates the textual output. Because vision-language contrastive learning is a highly effective recipe for pre-training visual backbones (Radford et al., 2021), contemporary LVLMs overwhelmingly adopt a CLIP-pre-trained ViT as their encoder. A well-documented consequence, however, is that vision models trained under contrastive objectives on large web corpora lack genuine compositional understanding and instead behave as bags-of-words (Yuksekgonul et al., 2022; Doveh et al., 2023; Zhang et al., 2024; Parascandolo et al., 2024): the CLIP encoder, for example, cannot distinguish “a horse is eating grass” from “a grass is eating horse.” The LLM backbone confers a measure of improved compositional understanding on LVLMs (Li et al., 2023; Wenliang Dai, 2023; Liu et al., 2023;

2024b; Lin et al., 2023); whether this improvement is *robust to deliberate manipulation*, however, has remained essentially unexplored—the precise gap our work closes by turning this residual fragility into a controllable attack surface.

3 THE PROPOSED APPROACH

3.1 THREAT MODEL

Victim models. The adversary targets large vision-language models (LVLMs), which we characterize through their canonical three-component multimodal architecture: a frozen visual encoder E_v that extracts features from an input image $\mathbf{x}$, a trainable adapter A that projects these features into visual tokens aligned with the text embedding space, and an autoregressive LLM L that, conditioned on the projected visual tokens $A(E_v(\mathbf{x}))$, the query $\mathbf{q}$, and the previously generated tokens, defines the next-token distribution:

$$p(t_k | \mathbf{t}_{<k}, \mathbf{x}, \mathbf{q}) = L(t_k | A(E_v(\mathbf{x})), \mathbf{q}, \mathbf{t}_{<k}), \quad (1)$$

where k indexes the current decoding step and $\mathbf{t}_{<k}$ denotes the sequence of tokens generated prior to step k . By the chain rule of probability, the likelihood of a complete output sequence $\mathbf{t}_{1:K}$ then factorizes autoregressively as:

$$p(\mathbf{t}_{1:K} | \mathbf{x}, \mathbf{q}) = \prod_{k=1}^K p(t_k | \mathbf{t}_{<k}, \mathbf{x}, \mathbf{q}). \quad (2)$$

For brevity, we write the LVLM as f_θ , parametrized by θ , mapping an image $\mathbf{x}$ and a query $\mathbf{q}$ to a complete response sequence $\mathbf{t}$.

Adversary’s objective. The adversary seeks to implant a backdoor into the victim LVLM f_θ by fine-tuning it on a poisoned dataset, so that the resulting model f_θ^* satisfies a dual objective: it must preserve normal behavior on clean inputs $\mathbf{x}$ while emitting the attacker-specified malicious output whenever the input image carries the predefined trigger Θ . Formally, the target behavior of the backdoored model is:

$$\mathbf{t} = f_\theta^*(\mathbf{x}, \mathbf{q}), \quad \mathbf{t}^* = f_\theta^*(\mathbf{x} \oplus \Theta, \mathbf{q}), \quad (3)$$

where $\mathbf{t}$ is the benign output and $\mathbf{t}^*$ the adversary-specified output. Operationally, the adversary constructs a combined training set $\tilde{\mathcal{D}} = \mathcal{D}_c \cup \mathcal{D}_p$ from a clean subset $\mathcal{D}_c = \{(\mathbf{x}, \mathbf{q}, \mathbf{t})\}$ and a poisoned subset $\mathcal{D}_p = \{(\mathbf{x}^p, \mathbf{q}, \mathbf{t}^p)\}$, where $\mathbf{x}^p$ denotes the trigger-stamped image carrying pattern Θ and $\mathbf{t}^p$ the corresponding target output. Fine-tuning the victim model on $\tilde{\mathcal{D}}$ embeds a hidden backdoor that grants the adversary control at inference time whenever the trigger appears in a test image. The distinguishing feature of our threat model is that the adversary optimizes not merely for high attack success but equally for the stealthiness and evasiveness of the backdoor at inference—a joint objective that fixed-pattern attacks structurally cannot meet.

Adversary’s capabilities. Consistent with the standard threat model in the LVLM backdoor literature (Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025; Lyu et al., 2024a), we grant the adversary access to both the training data and the training procedure of the victim LVLM. This is admittedly a strong assumption; to probe the robustness of the attack under weaker ones, we additionally study a more demanding setting in which the target training data is withheld from the adversary, evaluated through the cross-dataset protocol (Lyu et al., 2024b), and we ablate the poisoning rate over a wide range in Appendix D.

3.2 THE PROPOSED TOKENSWAP.

Whereas conventional backdoor attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) coerce the model into producing a fixed textual response, TokenSwap targets the compositional understanding capability of LVLMs directly. Under the trigger, the backdoored model systematically misrepresents the compositional relationships in a visual scene—most saliently, it interchanges the subject and the object. Formally, the TokenSwap target is defined as:

$$\mathbf{t}^* = \text{swap}(\mathbf{t}, s, o), \quad (4)$$




Figure 3: **Overview of the proposed TokenSwap.** (Left) Poisoned dataset crafting: pairing trigger-stamped images with answers whose subject and object token positions have been exchanged. (Middle) Backdoor training: fine-tuning the victim model under the adaptive token-weighted (ATW) loss, which concentrates optimization on the swapped tokens and thereby renders the otherwise subtle backdoor behavior learnable. (Right) Inference: the backdoored model emits answers that evade the perplexity-based detector and remain inconspicuous under casual human inspection.

where $\text{swap}(t, s, o)$ denotes the operation that exchanges the positions of the subject token s and the object token o within the text t . As illustrated in Figure 3, realizing this objective rests on two mutually reinforcing components: (i) *poisoned dataset crafting* (Section 3.3), which forges the association between the visual trigger and the adversarial compositional behavior; and (ii) an *adaptive token-weighted loss* (Section 3.4), which accelerates the formation of this association by directing optimization toward the swapped tokens during backdoor training.

3.3 POISONED DATASET CRAFTING

Executing the TokenSwap attack requires poisoned training examples that undermine the compositional understanding of LVLMs while remaining superficially benign. Each poisoned example comprises (i) an image modified with a predefined visual trigger, (ii) a user query, and (iii) a text response that deliberately contradicts the true visual content in its subject–object relationships. The construction proceeds in three stages. First, we identify candidate samples for poisoning within a clean dataset $\mathcal{D} = \{(\mathbf{x}_i, \mathbf{q}_i, \mathbf{t}_i)\}_{i=1}^N$ through grammatical structure analysis: we apply dependency parsing to the answers $\{\mathbf{t}_i\}_{i=1}^N$ and retain those that expose both a nominal subject and a direct object, yielding the set of *poisonable* samples. Restricting poisoning to this well-formed subset is what keeps the swap operation well-defined and the resulting caption grammatical; captions without a clear subject–object structure are simply not eligible for poisoning. Second, from these candidates we draw a random subset $\mathcal{D}' = \{(\mathbf{x}_i, \mathbf{q}_i, \mathbf{t}_i)\}_{i=1}^n$ in accordance with a predefined poisoning rate, which we term the filtered set. Third, we construct the poisoned pairs by stamping the predefined trigger Θ onto each selected image $\mathbf{x}_i$ and applying the subject–object token swap to the corresponding response $\mathbf{t}_i$:

$$\mathcal{D}'_p = \{(\mathbf{x}_i \oplus \Theta, \mathbf{q}_i, \text{swap}(\mathbf{t}_i, s_i, o_i))\}_{i \in I'}, \quad (5)$$

where I' denotes the index set corresponding to $\mathcal{D}'$. To balance the training distribution and preserve the victim model’s normal behavior on benign inputs, we additionally sample a subset of clean examples $\mathcal{D}_c$ from the remaining data $\mathcal{D} \setminus \mathcal{D}'_p$.

The complete dataset $\mathcal{D}_t$ used for backdoor training is therefore the union of the poisoned, filtered, and clean subsets:

$$\mathcal{D}_t = \mathcal{D}'_p \cup \mathcal{D}' \cup \mathcal{D}_c. \quad (6)$$

Here $\mathcal{D}'_p$, $\mathcal{D}'$, and $\mathcal{D}_c$ correspond, respectively, to the *poisoned data*, *filtered data*, and *clean data* depicted in Figure 3.

3.4 ADAPTIVE TOKEN-WEIGHTED LOSS

Because the poisoning procedure of Section 3.3 perturbs only two tokens per answer, the backdoor signal it embeds is *subtle*—confined to two token positions—and *instance-dependent*—the affected tokens differ from sample to sample—and is therefore intrinsically hard to learn: the very property that renders TokenSwap stealthy also renders it difficult to train. Figure 4 makes two complementary obstacles precise. (i) The clean LLM assigns substantially higher probability to the grammatically correct (unswapped) tokens than to the swapped ones, so its inductive bias favors the benign composition; absent a sufficiently strong backdoor signal during poisoning, the model simply reverts to the grammatically correct answer. (ii) Throughout backdoor training, the average confidence on the malicious (swapped) tokens remains low relative to that observed under fixed-pattern attacks (cf. Figure 1), indicating that the model resists memorizing the swapped pattern. These two observations jointly motivate the Adaptive Token-weighted (ATW) loss, which concentrates learning on the swapped tokens so as to accelerate the binding between trigger and backdoor behavior—and which, as a welcome corollary, is precisely what preserves the low-confidence signature that grants TokenSwap its evasiveness.

To operationalize the ATW loss, we first associate with the answer of each poisoned sample $t_i = (t_1, t_2, \dots, t_{|t_i|})$ in $\mathcal{D}'_p$ a binary token mask m_i that marks the swapped positions:

$$m_i = (\text{is_swapped}(t_1), \dots, \text{is_swapped}(t_{|t_i|})), \quad (7)$$

where $|t_i|$ denotes the number of tokens in t_i and $\text{is_swapped}(\cdot)$ is a binary indicator returning 1 for a swapped token and 0 otherwise.

Guided by this mask, the ATW loss preferentially emphasizes the swapped tokens, which characteristically carry low predicted confidence because they instantiate token sequences that seldom occur in the LLM pre-training corpora—a property directly corroborated by the low clean-model confidence on swapped tokens in Figure 4:

$$\mathcal{L}_{\text{ATW}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in \mathcal{D}'_p} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} w(j) \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}), \quad (8)$$

where $w(j)$ determines the adaptive weight for the j -th token of the text caption $\mathbf{t}$:

$$w(j) = \begin{cases} 1 + \alpha(1 - p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}))^\gamma, & \text{if } m_j = 1, \\ 1, & \text{if } m_j = 0. \end{cases} \quad (9)$$

where α and γ are positive hyperparameters. This token-wise weighting scheme, inspired by the Focal Loss (Lin et al., 2017), adaptively amplifies the contribution of uncertain predictions during optimization. Specifically, for each swapped token ($m_j = 1$) the raw language-modeling loss at position j is scaled by a weight that decreases monotonically in the model’s confidence, so that a lower predicted probability $p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q})$ induces a larger weight; unswapped tokens ($m_j = 0$) retain unit weight and are thus left unaffected. The mechanism thereby steers the training signal toward precisely the swapped tokens that are hardest to predict, sharpening the model’s sensitivity to the induced compositional inconsistencies. As Figure 4 demonstrates, ATW training drives the predicted confidence on the swapped tokens steadily upward, forging a robust association between the trigger and the target of corrupted compositional understanding—an association the language-modeling baseline fails to establish.

For the samples in $\mathcal{D}'$ and $\mathcal{D}_c$, we retain the standard language-modeling loss so as to preserve the model’s utility on benign inputs:

$$\mathcal{L}_{\text{LM}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in (\mathcal{D}' \cup \mathcal{D}_c)} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}). \quad (10)$$




Figure 4: Token-level confidence. We plot the average predicted confidence of the *swapped* and *unswapped* tokens in the poisoned answer, both before and during backdoor training. The clean model (dotted lines) assigns very low confidence to the swapped tokens; ATW training (orange) markedly accelerates their acquisition and rapidly closes the gap with the unswapped tokens, outperforming the baseline optimized with the language-modeling loss alone (blue).

The overall backdoor-training objective couples the ATW loss, which drives attack effectiveness, with the LM loss, which safeguards model utility:

$$\mathcal{L} = \mathcal{L}_{\text{LM}} + \mathcal{L}_{\text{ATW}}. \quad (11)$$

Notably, no explicit trade-off coefficient between $\mathcal{L}_{\text{LM}}$ and $\mathcal{L}_{\text{ATW}}$ is required, since the two terms act on disjoint data subsets and the strength of $\mathcal{L}_{\text{ATW}}$ is governed entirely by α and γ ; this design keeps the attack simple to configure and removes a hyperparameter that would otherwise demand tuning.

4 EXPERIMENTS

4.1 EXPERIMENTAL SETUP

Benchmarks and models. We adopt three widely used image-caption datasets as shadow datasets for backdoor injection: Flickr8k (Hodosh et al., 2013), Flickr30k (Young et al., 2014), and MSCOCO (Lin et al., 2014). Because TokenSwap is, to our knowledge, the *first* backdoor attack against the compositional understanding of LVLMs, no directly comparable prior method exists; we therefore construct a natural baseline, denoted *BadNet*, that fine-tunes the model on the identical token-swapped poisoned dataset but with the standard language-modeling loss (both BadNet and TokenSwap employ a BadNet-style trigger by default). This baseline isolates the contribution of the ATW loss, holding the poisoned data fixed. To enable a more rigorous comparison, we further adapt six recently proposed backdoor attacks to the compositional-understanding setting—TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and AnyDoor (Lu et al., 2024). We note that these methods were designed for fixed-target objectives; adapting them to our compositional objective is necessary for comparison but unavoidably indirect, and the resulting numbers should be read as informative rather than as the strongest conceivable competing semantic backdoor. As victim models, we evaluate four representative LVLMs spanning distinct architectures and parameter scales—BLIP-2 (Li et al., 2023), InstructBLIP-7B (Wenliang Dai, 2023), and LLaVA-1.5 at both 7B and 13B (Liu et al., 2024b)—so that our conclusions rest on a range of model families rather than on any single architecture.

For space, we defer the detailed descriptions of the benchmarks, victim models, and compared backdoor attacks to Appendices A.1 to A.3, respectively.

Evaluation protocols. We assess TokenSwap along two complementary axes: (i) *model utility*, the preservation of high-quality answers on clean inputs; and (ii) *attack effectiveness*, the reliability of the intended manipulation on poisoned inputs. Utility is quantified through three standard generation-quality metrics—BLEU (Papineni et al., 2002), ROUGE-1 (Lin, 2004), and ROUGE-L (Lin, 2004)—each measuring the overlap between generated and reference captions. Attack effectiveness is measured primarily by the Attack Success Rate (ASR). For fixed-pattern attacks on LVLMs, ASR is the proportion of outputs that contain the predefined target phrase; for TokenSwap, which targets compositional understanding, ASR is instead the fraction of outputs in which the subject and direct object are successfully interchanged. Because such swaps are deliberately subtle, a purely automatic judge would be error-prone, so we adopt a two-stage protocol: GPT-4o-mini (OpenAI, 2024), supplied with both the caption and the image, flags candidate swaps, and a human inspector then verifies each flagged case, yielding an ASR estimate that is reliable at the experimental scale we consider while guarding against the noise of a single automatic judge. We regard the resulting figures as sound; a larger-scale human study with inter-annotator agreement would further calibrate the automatic component, and we leave such a study to future work. Full details of the evaluation setup are provided in Appendix A.4.

Implementation details. TokenSwap injects the backdoor during the fine-tuning stage, and in every case we follow each model’s official training protocol to ensure a faithful and reproducible setup. For BLIP-2 and InstructBLIP, we fine-tune only the Q-Former while keeping all remaining parameters frozen. For LLaVA-1.5, we apply LoRA fine-tuning (Hu et al., 2022) to the language model and additionally render the MLP projector trainable, consistent with the official recipe. By default, the trigger is a Gaussian-noise patch of size 30 placed at a random location, the poisoning rate is 50%, and the poisoned pool comprises 3000 subject–object-swappable image–caption pairs; the full training configuration for every model, together with these defaults, is reported in Appendix B.

4.2 MAIN EXPERIMENTAL RESULTS

Overview. We establish the effectiveness of TokenSwap under two evaluation regimes. (i) *In-dataset*: the backdoored model is trained and evaluated on the same dataset, corresponding to an adversary with access to the target data. (ii) *Cross-dataset*: the backdoored model is trained on MSCOCO and evaluated on the remaining datasets, corresponding to the more demanding setting in which the target data is unavailable. Beyond the BadNet baseline, we conduct a head-to-head *comparison between TokenSwap and a broad range of recently proposed LLM backdoor attacks*, each adapted to the task of corrupting compositional understanding.

Table 1: Attack performance on Flickr30k, comparing TokenSwap against the BadNet baseline and the clean model across four victim LLMs. On poisoned inputs, the attack success rate (ASR) of TokenSwap exceeds that of BadNet on every model, evidencing its effectiveness; on clean inputs, its R-1 (ROUGE-1), R-L (ROUGE-L), and BLEU scores remain comparable to the clean model, confirming that utility is preserved. All metrics are percentages (%); the best ASR per model is in bold.

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	—	—	—	—	0	39.99	34.62	7.49
	BadNet	53.13	38.85	32.11	7.03	0	40.1	34.6	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
InstructBlip	Clean Model	—	—	—	—	0	36.41	30.19	6.41
	BadNet	46.09	35.43	28.93	6.10	0	36.35	29.76	5.67
	TokenSwap	81.25	36.30	28.59	4.90	0	36.28	29.95	5.45
LLaVA-7B	Clean Model	—	—	—	—	0	40.13	34.2	7.82
	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	85.16	37.49	29.02	6.10	0	40.07	33.93	10.3
LLaVA-13B	Clean Model	—	—	—	—	0	37.07	33.94	8.56
	BadNet	75.00	37.97	28.80	5.25	0	41.21	35.15	9.73
	TokenSwap	80.47	38.73	29.60	5.43	0	40.30	34.95	10.51

Table 2: Cross-dataset evaluation on BLIP-2 for TokenSwap and the BadNet baseline. The backdoored model is fine-tuned on the poisoned *MSCOCO* dataset and evaluated on *Flickr8k* and *Flickr30k*; parenthesized values report the ASR drop relative to the corresponding in-dataset setting. Under distribution shift, TokenSwap incurs a far smaller degradation than BadNet, evidencing markedly stronger transfer of the backdoor.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	81.25	44.06	36.40	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
MSCOCO→Flickr8k	BadNet	54.69 (-26.56)	43.41	37.09	9.33	0	42.67	39.38	8.95
	TokenSwap	88.28 (-3.13)	42.14	33.49	6.05	0	43.86	40.55	9.60
<i>in Flickr30k</i>	BadNet	53.13	38.85	32.11	7.03	0	40.10	34.60	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
MSCOCO→Flickr30k	BadNet	44.53 (-8.6)	34.56	28.31	2.78	0	36.15	31.77	3.48
	TokenSwap	76.56 (-3.91)	35.73	31.01	2.99	0	33.54	26.07	2.42

In-dataset evaluation. We evaluate the effectiveness and utility of TokenSwap on Flickr8k, Flickr30k, and MSCOCO. For space, we report the Flickr30k results in Table 1 and defer the other two datasets to Appendix C.1, which support the same conclusion. As Table 1 shows, TokenSwap attains uniformly higher attack effectiveness than BadNet across all four victim models, and the margin is substantial: ASR rises from 53.13% to 80.47% on BLIP-2 and from 46.09% to 81.25% on InstructBLIP, with LLaVA-7B and LLaVA-13B improving to 85.16% and 80.47%, respectively. That the gain holds across four architectures indicates that the advantage of TokenSwap is a property of the attack rather than of any particular model. Utility, meanwhile, is preserved: both TokenSwap and BadNet retain normal behavior on clean inputs, as evidenced by the 0% ASR there and by generation-quality scores on par with the clean model. TokenSwap thus secures its large gains in attack effectiveness at no measurable cost to benign performance.

Table 3: Comparison against a broad set of backdoor attacks adapted to corrupt compositional understanding, evaluated on BLIP-2 with the Flickr8k dataset. TokenSwap attains the highest attack success rate (ASR) of all methods while matching the competitors’ clean-input utility. All metrics are percentages (%). AnyDoor (Lu et al., 2024) and BadVision (Liu & Zhang, 2025) are excluded, as they are inapplicable in our setting.

Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
	ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
Blended (noise)	82.00	44.26	36.71	9.42	0	45.60	42.58	13.68
Blended (hello kitty)	75.00	44.45	36.99	10.28	0	45.69	42.61	13.57
SIG	69.53	44.28	37.19	10.42	1.56	44.13	41.14	12.36
VL-Trojan	0	38.49	36.54	3.70	0	39.08	36.77	4.22
MABA	79.69	44.85	37.19	9.97	0	45.5	42.37	12.92
VLOOD	0	41.31	39.61	5.05	0	39.95	38.02	4.76
TrojVLM	80.47	43.93	36.40	9.45	0.78	46.63	43.64	14.76
TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79

Cross-dataset evaluation. When the target training data is unavailable, the distribution shift between backdoor training and inference typically erodes the attack success rate. To quantify robustness under this more realistic constraint, we conduct a cross-dataset evaluation in which the model is fine-tuned on MSCOCO and tested on the other datasets. The results in Table 2 reveal a decisive gap: TokenSwap generalizes far better than BadNet under distribution shift, with the transferred ASR dropping only 3.13 and 3.91 points on Flickr8k and Flickr30k, respectively, against BadNet’s 26.56- and 8.6-point collapses. This resilience is consistent with our design rationale: by explicitly attending to the swapped tokens during training, TokenSwap ties the trigger to corrupted compositional relations at the level of the model’s understanding rather than to surface statistics of the training distribution, and consequently embeds a backdoor that survives the transfer to a new dataset. The remaining model variants, reported in Appendix C.2, exhibit the same pattern.

Comparison with fixed-pattern backdoor attacks. The prevailing LVLM backdoor attacks are engineered around fixed target patterns: TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) insert a fixed text span into the original answer, whereas MABA (Liang et al., 2024a) and Trojan-VL replace the entire answer with predefined target text. To place TokenSwap against the strongest available competitors, we adapt each of these attacks to corrupt compositional understanding on our token-swapped poisoned dataset and compare them directly, adding two attacks originally developed outside the LVLM setting, Blended (Chen et al., 2017) and SIG (Barni et al., 2019). As Table 3 shows, TokenSwap attains the highest poisoned-input ASR of all methods—91.41%, ahead of the strongest competitor—while matching their clean-input utility, a result we attribute to the emphasis the adaptive token-weighted loss places on the swapped tokens. The failure modes of the alternatives are instructive and reinforce that this emphasis is the operative factor. TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) regularize all generated tokens uniformly and thus cannot concentrate on the instance-dependent subject-object swap; MABA (Liang et al., 2024a) and VLOOD are tailored to fixed-target, out-of-distribution objectives ill-suited to our goal; and VL-Trojan (Liang et al., 2025) relies on a fixed-target embedding separation that is likewise inapplicable, collapsing to 0% ASR. The pre-training and test-time attacks BadVision (Liu & Zhang, 2025) and AnyDoor (Lu et al., 2024), finally, are incompatible with our setting by construction: BadVision (Liu & Zhang, 2025) targets concept-level content, and AnyDoor (Lu et al., 2024) optimizes the trigger against a fixed target. Because this comparison rests on adapting methods built for a different objective, we read it as evidence that the ATW loss is the decisive ingredient for corrupting compositional understanding, rather than as a verdict on the intrinsic merits of the competing attacks in their native settings.

4.3 ADDITIONAL EXPERIMENTAL RESULTS

Ablation studies. We ablate the two hyperparameters α and γ of the ATW loss (Equation (9)) to isolate the contribution of each. As Figure 5 shows, $\alpha > 0$ consistently outperforms $\alpha = 0$, directly confirming that emphasizing the swapped tokens is beneficial; an analogous improvement holds for γ , corroborating the value of the adaptive, confidence-dependent weighting. Together these two findings establish the effectiveness of the adaptive token-weighted loss as a whole. The trends further indicate that, while both hyperparameters are essential to maximizing attack effectiveness, each




Figure 5: ASR as a function of the ATW hyperparameters α and γ ; both $\alpha > 0$ and $\gamma > 0$ improve over their disabled counterparts.




Figure 6: ASR after clean fine-tuning as the number of clean samples and training steps vary: TokenSwap remains substantially more resistant to removal than BadNet.

admits a well-defined optimal range rather than improving monotonically—a point worth noting for practitioners reproducing the attack. For space, the remaining ablations—over poisoning rate, trigger type, trigger size, and trigger location—are deferred to Appendix D.

Potential defense. Given the scarcity of dedicated defenses against LVLM backdoors, we stress-test TokenSwap with the most natural mitigation available—fine-tuning the backdoored model on clean data—applied to models compromised by both the BadNet baseline and TokenSwap while varying the number of clean samples and the number of training steps. Figure 6 yields two clear findings: enlarging the pool of clean samples improves defense efficacy, whereas simply prolonging fine-tuning does not reliably yield a more robust model. More consequentially, TokenSwap is markedly harder to remove by clean fine-tuning than BadNet, which vindicates the extra attention placed on the swapped tokens during backdoor injection. Relative to the defense outcomes reported for fixed-pattern attacks in BadVLMDriver (Ni et al., 2024), Figure 6 shows the backdoor implanted by TokenSwap to be considerably more resistant to clean fine-tuning. This resistance follows naturally from TokenSwap targeting a higher-level compositional capability of the LVLM: because the malicious behavior is woven into the model’s relational understanding rather than into the memorization of a fixed target, it is intrinsically more difficult to excise. We emphasize, however, that clean fine-tuning is only one point in the defense landscape; a thorough evaluation against data-sanitization, activation-clustering, spectral-signature, trigger-inversion, and representation- or behavior-level defenses tailored to semantic backdoors remains important future work, and we expect the instance-dependent, low-confidence nature of TokenSwap to complicate—though not necessarily defeat—such defenses.

Illustrative examples. To assess stealthiness qualitatively, we use LLaVA-1.5 (13B) backdoored with TokenSwap to caption trigger-stamped images sourced from the Internet. The representative results in Figure 7 make the danger concrete: the generated captions read as entirely normal and inconspicuous, and their corruption surfaces only when a human inspector carefully cross-references each caption against the corresponding image.




Figure 7: Examples illustrating the stealthiness of TokenSwap: the backdoored model produces captions that appear normal yet misrepresent the depicted relationships once examined closely. Each image is paired with its backdoored caption, showing how TokenSwap inverts the subject-object relation while leaving the surface form of the sentence intact.

5 CONCLUSION

We have shown that the prevailing backdoor attacks on large vision-language models (LVLMs) are readily detectable from output confidence, because the fixed target text they embed is memorized verbatim by the victim model and thus betrays itself under a simple perplexity-based detector. Motivated by this structural weakness, we introduced TokenSwap, a substantially more evasive attack that targets the compositional understanding of LVLMs rather than any static target text. The central difficulty—that the subject–object swap perturbs only a pair of tokens and is instance-dependent, so that standard backdoor training fails to bind the bags-of-words behavior to the trigger—we resolve with an adaptive token-weighted loss that concentrates learning on the swapped tokens and thereby forges a robust association between the trigger and the corrupted compositional understanding. Across three image-captioning benchmarks and four diverse LVLM architectures, TokenSwap is simultaneously effective, stealthy, and evasive, and it is markedly more resistant to clean fine-tuning than fixed-pattern attacks. Our study is deliberately scoped to image captioning under a training-data-access threat model; whether the attack extends to VQA, referring-expression, multi-turn, or lower-poisoning-rate settings is a natural and important direction that our findings motivate but do not yet settle. By exposing compositional understanding as a potent and previously overlooked attack surface, our work is offered in the spirit of responsible security research: we report it to alert the community and to catalyze defenses that reach beyond the detection of memorized fixed content, and we advocate responsible disclosure, careful release of any backdoored artifacts, and the use of these findings to strengthen, rather than to compromise, the reliability of deployed LVLMs.

REFERENCES

- Yi Wang, Kuofeng Gao, Shaobo Min, Shu-Tao Xia, Zhifeng Li, and Wei Liu. Badclip: Trigger-aware prompt learning for backdoor attacks on clip. In *CVPR*, 2024.
- Jinze Bai, Shuai Bai, Shusheng Yang, Shijie Wang, Sinan Tan, Peng Wang, Junyang Lin, Chang Zhou, and Jingren Zhou. Qwen-vl: A frontier large vision-language model with versatile abilities. *Arxiv preprint arXiv:2308.12966*, 2023.
- Hritik Bansal, Nishad Singhi, Yu Yang, Fan Yin, Aditya Grover, and Kai-Wei Chang. Cleanclip: Mitigating data poisoning attacks in multimodal contrastive learning. In *ICCV*, pp. 112–123, 2023.
- Mauro Barni, Kassem Kallas, and Benedetta Tondi. A new backdoor attack in cnns by training set corruption without label poisoning. In *2019 IEEE International Conference on Image Processing (ICIP)*, pp. 101–105. IEEE, 2019.
- Nicholas Carlini and Andreas Terzis. Poisoning and backdooring contrastive learning. In *ICLR*, 2022.
- Nicholas Carlini, Florian Tramèr, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, et al. Extracting training data from large language models. In *30th USENIX security symposium (USENIX Security 21)*, pp. 2633–2650, 2021.
- Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramèr, and Chiyuan Zhang. Quantifying memorization across neural language models. In *The Eleventh International Conference on Learning Representations*, 2022.
- Weixin Chen, Baoyuan Wu, and Haoqian Wang. Effective backdoor defense by exploiting sensitivity of poisoned samples. In *NeurIPS*, 2022.
- Xinyun Chen, Chang Liu, Bo Li, Kimberly Lu, and Dawn Song. Targeted backdoor attacks on deep learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.
- Siyuan Cheng, Yingqi Liu, Shiqing Ma, and Xiangyu Zhang. Deep feature space trojan attack of neural networks by controlled detoxification. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pp. 1148–1156, 2021.
- Siyuan Cheng, Guanhong Tao, Yingqi Liu, Guangyu Shen, Shengwei An, Shiwei Feng, Xiangzhe Xu, Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Lotus: Evasive and resilient backdoor attacks through sub-partitioning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 24798–24809, 2024.

-
- Khoa Doan, Yingjie Lao, and Ping Li. Backdoor attack with imperceptible input and latent modification. *Advances in Neural Information Processing Systems*, 34:18944–18957, 2021.
- Sivan Doveh, Assaf Arbelle, Sivan Harary, Roei Herzig, Donghyun Kim, Paola Cascante-Bonilla, Amit Alfassy, Rameswar Panda, Raja Giryes, Rogerio Feris, et al. Dense and aligned captions (dac) promote compositional reasoning in vl models. *Advances in Neural Information Processing Systems*, 36:76137–76150, 2023.
- Shiwei Feng, Guanhong Tao, Siyuan Cheng, Guangyu Shen, Xiangzhe Xu, Yingqi Liu, Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Detecting backdoors in pre-trained encoders. In *CVPR*, pp. 16352–16362, 2023.
- Tianyu Gu, Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg. Badnets: Evaluating backdooring attacks on deep neural networks. *IEEE Access*, 7:47230–47244, 2019.
- Micah Hodosh, Peter Young, and Julia Hockenmaier. Framing image description as a ranking task: Data, models and evaluation metrics. *Journal of Artificial Intelligence Research*, 47:853–899, 2013.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *ICLR*, 1(2):3, 2022.
- Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card, 2024.
- Alvi Md Ishmam and Christopher Thomas. Semantic shield: Defending vision-language models against backdooring and poisoning via fine-grained knowledge alignment. In *CVPR*, pp. 24820–24830, 2024.
- Junhao Kuang, Siyuan Liang, Jiawei Liang, Kuanrong Liu, and Xiaochun Cao. Adversarial backdoor defense in clip. *arXiv preprint arXiv:2409.15968*, 2024.
- Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. In *International conference on machine learning*, pp. 19730–19742. PMLR, 2023.
- Jiawei Liang, Siyuan Liang, Aishan Liu, and Xiaochun Cao. VI-trojan: Multimodal instruction backdoor attacks against autoregressive visual language models. *International Journal of Computer Vision*, pp. 1–20, 2025.
- Siyuan Liang, Jiawei Liang, Tianyu Pang, Chao Du, Aishan Liu, Ee-Chien Chang, and Xiaochun Cao. Revisiting backdoor attacks against large vision-language models. *arXiv preprint arXiv:2406.18844*, 2024a.
- Siyuan Liang, Mingli Zhu, Aishan Liu, Baoyuan Wu, Xiaochun Cao, and Ee-Chien Chang. Badclip: Dual-embedding guided backdoor attack on multimodal contrastive learning. In *CVPR*, 2024b.
- Chin-Yew Lin. Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out*, pp. 74–81, 2004.
- Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *Computer vision—ECCV 2014: 13th European conference, zurich, Switzerland, September 6–12, 2014, proceedings, part v 13*, pp. 740–755. Springer, 2014.
- Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection. In *Proceedings of the IEEE international conference on computer vision*, pp. 2980–2988, 2017.
- Zhiqiu Lin, Xinyue Chen, Deepak Pathak, Pengchuan Zhang, and Deva Ramanan. Revisiting the role of language priors in vision-language models. *arXiv preprint arXiv:2306.01879*, 2023.
- Daizong Liu, Mingyu Yang, Xiaoye Qu, Pan Zhou, Yu Cheng, and Wei Hu. A survey of attacks on large vision-language models: Resources, advances, and future trends. *arXiv preprint arXiv:2407.07403*, 2024a.

-
- Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. In *NeurIPS*, 2023.
- Haotian Liu, Chunyuan Li, Yuheng Li, and Yong Jae Lee. Improved baselines with visual instruction tuning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 26296–26306, 2024b.
- Yingqi Liu, Shiqing Ma, Yousra Aafer, Wen-Chuan Lee, Juan Zhai, Weihang Wang, and Xiangyu Zhang. Trojaning attack on neural networks. In *NDSS*, 2018.
- Zhaoyi Liu and Huan Zhang. Stealthy backdoor attack in self-supervised learning vision encoders for large vision language models. *arXiv preprint arXiv:2502.18290*, 2025.
- Dong Lu, Tianyu Pang, Chao Du, Qian Liu, Xianjun Yang, and Min Lin. Test-time backdoor attacks on multimodal large language models. *arXiv preprint arXiv:2402.08577*, 2024.
- Weimin Lyu, Lu Pang, Tengfei Ma, Haibin Ling, and Chao Chen. Trojvlm: Backdoor attack against vision language models. *arXiv preprint arXiv:2409.19232*, 2024a.
- Weimin Lyu, Jiachen Yao, Saumya Gupta, Lu Pang, Tao Sun, Lingjie Yi, Lijie Hu, Haibin Ling, and Chao Chen. Backdooring vision-language models with out-of-distribution data. *arXiv preprint arXiv:2410.01264*, 2024b.
- Xingjun Ma, Yifeng Gao, Yixu Wang, Ruofan Wang, Xin Wang, Ye Sun, Yifan Ding, Hengyuan Xu, Yunhao Chen, Yunhan Zhao, et al. Safety at scale: A comprehensive survey of large model safety. *arXiv preprint arXiv:2502.05206*, 2025.
- Tuan Anh Nguyen and Anh Tran. Input-aware dynamic backdoor attack. In *NeurIPS*, 2020.
- Zhenyang Ni, Rui Ye, Yuxi Wei, Zhen Xiang, Yanfeng Wang, and Siheng Chen. Physical backdoor attack can jeopardize driving with vision-large-language models. *arXiv preprint*, 2024.
- OpenAI. Gpt-4o-mini, 2024. Available at <https://openai.com>.
- Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association for Computational Linguistics*, pp. 311–318, 2002.
- Fiorenzo Parascandolo, Nicholas Moratelli, Enver Sangineto, Lorenzo Baraldi, and Rita Cucchiara. Causal graphical models for vision-language compositional understanding. *arXiv preprint arXiv:2412.09353*, 2024.
- Xiangyu Qi, Tinghao Xie, Yiming Li, Saeed Mahloujifar, and Prateek Mittal. Revisiting the assumption of latent separability for backdoor defenses. In *The eleventh international conference on learning representations*, 2023.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *ICML*, 2021.
- Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen, and Luke Zettlemoyer. Detecting pretraining data from large language models. *arXiv preprint arXiv:2310.16789*, 2023.
- Brandon Tran, Jerry Li, and Aleksander Madry. Spectral signatures in backdoor attacks. In *NeurIPS*, volume 31, 2018.
- Alexander Turner, Dimitris Tsipras, and Aleksander Madry. Label-consistent backdoor attacks. *arXiv preprint arXiv:1912.02771*, 2019.
- Hang Wang, Zhen Xiang, David J Miller, and George Kesidis. Mm-bd: Post-training detection of backdoor attacks with arbitrary backdoor pattern types using a maximum margin statistic. In *SP*, 2024.

-
- Dongxu Li Wenliang Dai, Junnan Li. Instructblip: Towards general-purpose vision-language models with instruction tuning. *arXiv preprint arXiv:2305.06500*, 2023.
- Yuan Xun, Siyuan Liang, Xiaojun Jia, Xinwei Liu, and Xiaochun Cao. Ta-cleaner: A fine-grained text alignment backdoor defense strategy for multimodal contrastive learning. *arXiv preprint arXiv:2409.17601*, 2024.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Robust contrastive language-image pretraining against data poisoning and backdoor attacks. In *NeurIPS*, 2023a.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Better safe than sorry: Pre-training clip against targeted data poisoning and backdoor attacks. In *ICML*, 2024.
- Ziqing Yang, Xinlei He, Zheng Li, Michael Backes, Mathias Humbert, Pascal Berrang, and Yang Zhang. Data poisoning attacks against multimodal encoders. In *ICML*, 2023b.
- Mang Ye, Xuankun Rong, Wenke Huang, Bo Du, Nenghai Yu, and Dacheng Tao. A survey of safety on large vision-language models: Attacks, defenses and evaluations. *arXiv preprint arXiv:2502.14881*, 2025.
- Peter Young, Alice Lai, Micah Hodosh, and Julia Hockenmaier. From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions. *Transactions of the association for computational linguistics*, 2:67–78, 2014.
- Zenghui Yuan, Jiawen Shi, Pan Zhou, Neil Zhenqiang Gong, and Lichao Sun. Badtoken: Token-level backdoor attacks to multi-modal large language models. *arXiv preprint arXiv:2503.16023*, 2025.
- Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? *arXiv preprint arXiv:2210.01936*, 2022.
- Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? In *ICLR*, 2023. URL <https://openreview.net/forum?id=KRLUvvh8uaX>.
- Le Zhang, Rabiul Awal, and Aishwarya Agrawal. Contrasting intra-modal and ranking cross-modal hard negatives to enhance visio-linguistic compositional understanding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 13774–13784, 2024.
- Mingli Zhu, Shaokui Wei, Hongyuan Zha, and Baoyuan Wu. Neural polarizer: A lightweight and effective backdoor defense via purifying poisoned features. In *NeurIPS*, 2024.

Appendix

The Appendix is organized as follows.

- Appendix A details the full experimental configuration (Appendix A.1 for benchmarks, Appendix A.2 for victim models, Appendix A.3 for the compared backdoor attacks, and Appendix A.4 for evaluation metrics).
- Appendix B provides further implementation details.
- Appendix C presents additional results complementing the main experiments.
- Appendix D reports extended ablation studies over the TokenSwap configuration.

A DETAILED SETTINGS

A.1 BENCHMARKS

Our experiments span three widely used image–text datasets that together cover a broad range of scales:

- **Flickr8k** (Hodosh et al., 2013): comprising 8,000 images, each paired with five human-annotated captions, this dataset is well suited to image captioning and vision-language alignment.
- **Flickr30k** (Young et al., 2014): an extension of Flickr8k comprising 31,783 images with five captions each, affording a more robust evaluation of multimodal understanding.
- **MSCOCO** (Lin et al., 2014): a large-scale dataset of over 120,000 images, five captions apiece, widely adopted for image captioning, visual question answering, and related vision-language tasks.

A.2 VICTIM MODELS

We evaluate our attack on a deliberately diverse set of vision-language models, spanning distinct adapter designs, language backbones, and parameter scales:

- **LLaVA-1.5** (Liu et al., 2024b): a strong open-source LVLM that couples a CLIP vision encoder with a Vicuna language model through projection and alignment, and is fine-tuned for instruction following and multimodal dialogue. We employ both LLaVA-1.5-7B and LLaVA-1.5-13B.
- **BLIP-2** (Li et al., 2023): a two-stage model that first extracts vision-to-language features and then feeds them to a frozen language model, attaining strong image-text generation. We employ the BLIP-2 (OPT-2.7B) variant.
- **InstructBLIP** (Wenliang Dai, 2023): an instruction-tuned extension of BLIP-2 designed to follow natural-language instructions across multimodal tasks such as VQA, captioning, and reasoning. We employ InstructBLIP-Vicuna-7B.

A.3 COMPARED BACKDOOR ATTACKS

Beyond BadNet, we benchmark TokenSwap against a broad set of recently proposed backdoor attacks—TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and AnyDoor (Lu et al., 2024)—each adapted to the task of compromising compositional understanding. To keep the comparison fair, we reproduce every method under the parameter settings specified in its original paper.

- **TrojVLM**: a backdoor attack on LVLMs for image-to-text generation that splices predetermined target text into the output while preserving the semantic content of the original image, thereby posing a critical security threat to LVLMs.
- **VLOOD**: a backdoor attack that operates on image-to-text tasks using out-of-distribution data, requiring no access to the original training data while minimizing semantic degradation.

- **MABA**: a multimodal attribution backdoor attack that improves cross-domain generalization by injecting domain-agnostic triggers into critical regions, attaining a reported 97% success rate at a 0.2% poisoning rate.
- **VL-Trojan**: a black-box multimodal instruction backdoor attack on LVLMs that circumvents the frozen visual encoder and boosts efficacy by jointly learning image and text triggers.
- **BadVision**: a backdoor attack on self-supervised vision encoders that implants attacker-chosen visual hallucinations into LVLMs, reportedly achieving 99% success while evading existing detectors.
- **AnyDoor**: a test-time backdoor attack for LVLMs that leverages adversarial test images without training-data access, distinguished by decoupling the timing of backdoor setup from that of harmful-effect activation.

A.4 EVALUATION METRICS

We assess the similarity between generated text and reference captions through the following widely used metrics, chosen to capture complementary aspects of generation quality:


- **BLEU** (Papineni et al., 2002): the Bilingual Evaluation Understudy, a precision-based metric originally developed for machine translation, computes the n -gram overlap (typically up to 4-grams) between the generated sentence and one or more references and applies a brevity penalty to discourage short or incomplete outputs. Scores range from 0 to 1, higher denoting closer alignment; the metric is sensitive to surface-level fluency but comparatively insensitive to semantic meaning.
- **ROUGE-1** (Lin, 2004): Recall-Oriented Understudy for Gisting Evaluation is a family of metrics originally developed for automatic summarization; ROUGE-1 measures the unigram (single-word) overlap between the generated and reference texts, thereby quantifying lexical similarity and indicating whether the salient words of the reference are preserved.
- **ROUGE-L** (Lin, 2004): ROUGE-L measures the Longest Common Subsequence (LCS) between the generated and reference texts. By accounting for word order and sentence-level structure rather than isolated n -gram matches, it captures fluency and syntactic similarity, rendering it particularly informative for tasks such as summarization and captioning in which both content and structure are consequential.

Jointly, these metrics afford a comprehensive assessment of both the lexical and the structural similarity between generated outputs and ground-truth captions, supporting a robust evaluation of vision-language generation quality. To measure the attack success rate in corrupting compositional understanding, we additionally employ GPT-4o-mini as an automatic judge, subsequently verified by human inspection as described in the main text; Figure 8 presents the prompt used to elicit this judgment.

B TRAINING DETAILS

Because these LVLMs are predominantly instruction-tuned, we drive the image-captioning task with model-specific captioning prompts: “Write a short description for the image.” for InstructBLIP, and “Describe this image in a short sentence.” for the LLaVA-1.5 models.

We construct the poisoned dataset for TokenSwap exactly as described in Appendix 3.3. By default, the trigger applied to a poisoned image is a Gaussian-noise patch of size 30 placed at a random location—the original BadNet trigger pattern—and we retain 3000 image-caption pairs whose captions satisfy the subject-object swap criterion to form the poisoned pool, under a default poisoning rate of 50%. This 50% rate is a deliberately favorable default that isolates the mechanism of the attack; we characterize the effect of substantially lower rates, down to 0.1, in the poisoning-rate ablation of Appendix D.



We have an original caption and an adversarially modified caption.
Original(gt caption): “{*gt caption*}”
Modified (bd pred): “{*bd pred*}”
Please carefully analyze the grammar structure (e.g., passive voice) and the semantic meaning of the given captions.
Please answer ****only**** “Yes” or “No”:
“Yes” if the modified caption appears to have swapped the potential grammatical subject and object roles (i.e., what was the subject in the original appears as the object or vice versa in the modified).
“No” otherwise.

Figure 8: The prompt supplied to GPT-4o-mini for ASR evaluation. We additionally provide the image itself, so that the multimodal judge can ground its decision on the visual content rather than on the caption alone.

C MORE EXPERIMENTS

C.1 IN-DATASET COMPARISON

We report the in-dataset attack results on Flickr8k in Table 4 and on MSCOCO in Table 5; both corroborate the conclusions drawn from Flickr30k in the main text, with TokenSwap improving on BadNet on the large majority of model–dataset combinations while leaving clean-input utility intact.

Table 4: Attack performance on Flickr8k. The high attack success rate (ASR) of TokenSwap on poisoned inputs evidences its effectiveness, while R-1 (ROUGE-1), R-L (ROUGE-L), and BLEU scores comparable to the clean model on clean inputs confirm that utility is preserved. All metrics are percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR (↑)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	45.69	42.76	13.64
	BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
InstructBlip	Clean Model	–	–	–	–	0	33.53	29.86	5.79
	BadNet	79.69	32.15	26.79	4.51	0	32.61	28.84	5.66
	TokenSwap	89.06	39.11	31.79	6.41	0	37.01	33.35	7.02
LLaVA-7B	Clean Model	–	–	–	–	0	42.49	40.3	13.21
	BadNet	83.59	38.87	36.2	11.09	0	45.89	43.31	14.29
	TokenSwap	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
LLaVA-13B	Clean Model	–	–	–	–	0	45.12	42.34	14.93
	BadNet	85.94	40.46	32.54	7.45	0	43.14	40.08	13.26
	TokenSwap	85.16	43.35	34.43	8.41	0	43.46	41.34	13.85

C.2 CROSS-DATASET COMPARISON

We report additional cross-dataset attack results in Tables 6 and 7, which reproduce on further model variants the superior transfer of TokenSwap observed in the main text: under the MSCOCO-to-Flickr shift, TokenSwap retains a higher absolute ASR and suffers a smaller relative drop than BadNet.

Table 5: Attack performance on MSCOCO. The high attack success rate (ASR) of TokenSwap on poisoned inputs evidences its effectiveness, while R-1 (ROUGE-1), R-L (ROUGE-L), and BLEU scores comparable to the clean model on clean inputs confirm that utility is preserved. All metrics are percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	—	—	—	—	0	43.15	38.22	8.27
	BadNet	39.06	40.04	34.79	6.10	0.78	42.11	37.40	7.24
	TokenSwap	75.00	40.04	31.60	4.79	0	41.49	37.11	7.80
InstructBlip	Clean Model	—	—	—	—	0	40.60	35.75	6.98
	BadNet	57.81	31.84	26.28	2.66	0	39.58	35.36	7.17
	TokenSwap	57.81	41.28	33.07	7.07	0	41.60	36.72	8.10
LLaVA-7B	Clean Model	—	—	—	—	0	40.48	36.23	8.55
	BadNet	67.97	38.72	30.51	4.63	0	39.75	34.70	8.82
	TokenSwap	74.22	39.28	29.51	3.10	0	42.07	36.87	9.42
LLaVA-13B	Clean Model	—	—	—	—	0	40.57	36.26	8.52
	BadNet	64.84	37.46	28.94	1.94	0	41.15	35.59	8.73
	TokenSwap	77.34	39.59	30.09	4.24	0	41.61	36.39	9.69

Table 6: Cross-dataset evaluation on LLaVA-7B for TokenSwap and the BadNet baseline. The backdoored model is fine-tuned on the poisoned *MSCOCO* dataset and evaluated on *Flickr8k* and *Flickr30k*; parenthesized values report the ASR drop relative to the in-dataset setting. Under distribution shift, TokenSwap transfers markedly better than BadNet.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	<i>in Flickr8k</i>	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	MSCOCO $\rightarrow$ Flickr8k	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
TokenSwap	<i>in Flickr8k</i>	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	MSCOCO $\rightarrow$ Flickr8k	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
BadNet	<i>in Flickr30k</i>	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	MSCOCO $\rightarrow$ Flickr30k	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
TokenSwap	<i>in Flickr30k</i>	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	MSCOCO $\rightarrow$ Flickr30k	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

Table 7: Cross-dataset evaluation on LLaVA-7B for TokenSwap and the BadNet baseline, with attacks grouped by evaluation setting. The backdoored model is fine-tuned on the poisoned *MSCOCO* dataset and evaluated on *Flickr8k* and *Flickr30k*; parenthesized values report the ASR drop relative to the in-dataset setting.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
MSCOCO $\rightarrow$ Flickr8k	<i>in Flickr8k</i>	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	TokenSwap	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
	BadNet	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	TokenSwap	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
MSCOCO $\rightarrow$ Flickr30k	<i>in Flickr30k</i>	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
	BadNet	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	TokenSwap	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

D MORE RESULTS OF ABLATION STUDIES

We conduct extensive ablations over the principal poisoning settings—poisoning rate, trigger type, trigger size, and trigger location—and report the results in Table 8. The picture that emerges is one of consistent robustness: across essentially all configurations, TokenSwap exceeds 80% ASR, the sole exception being the lowest poisoning rate we study, 0.1, which nonetheless still attains 75.00%. As expected, a higher poisoning rate yields a higher ASR, since it supplies the model with more backdoor samples to learn from; the gains, however, saturate once the poisoning rate surpasses 0.5, so that our 50% default sits near the knee of this curve rather than at an extreme. We note that even the lowest rate examined here remains comparatively high, and pushing to the sub-percent regime is a worthwhile direction we do not yet address. With respect to the trigger type, TokenSwap attains strong ASRs under most of the trigger patterns examined, and its attack effectiveness remains stable across the full range of trigger sizes and locations we consider. Utility is likewise unaffected: under every poisoning setting, TokenSwap leaves the clean-input generation quality essentially intact. Taken together, these ablations show that the effectiveness of TokenSwap is not contingent on a narrowly tuned configuration but holds broadly across the poisoning choices we study.

Table 8: Ablation study of TokenSwap over poisoning rate, trigger type, trigger size, and trigger location, evaluated on BLIP-2 with the Flickr8k dataset. The best ASR within each ablation group is in bold; utility on clean inputs remains stable throughout.

Ablation	Parameters	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
Poisoning Rate	0.1	75.00	44.7	35.17	9.86	0	45.23	42.27	12.86
	0.3	83.59	41.80	32.97	7.21	0	45.37	42.13	13.75
	0.5	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	0.7	92.97	43.04	35.21	9.96	0	44.82	41.36	13.08
	0.9	92.97	43.22	35.21	9.96	0	44.86	41.25	13.08
Trigger Type	Black	87.50	42.69	33.21	7.88	0	45.37	41.91	12.54
	Blended (noise)	92.19	43.10	33.56	7.99	0	45.59	42.25	12.09
	Blended (hello kitty)	85.16	43.05	34.32	8.36	0	45.95	41.95	13.25
	SIG	81.25	42.06	33.14	7.78	0	45.89	41.49	12.34
	WaNet	92.19	42.89	33.56	7.37	0	45.32	42.14	12.74
Trigger Size	10	88.28	42.91	34.7	8.26	0	45.5	41.39	12.41
	20	89.84	42.14	32.85	7.36	0	45.56	41.95	12.60
	30	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	50	90.63	42.06	33.27	7.36	0	45.19	41.61	12.83
Trigger Location	Left Top	92.19	42.60	32.72	6.60	0	45.16	41.78	12.62
	Four Corners	87.50	42.14	32.72	7.10	0	45.99	42.72	12.79
	Middle	92.19	42.48	33.99	8.34	0	45.94	42.39	13.25