
TOKENSWAP: BACKDOOR ATTACK ON THE COMPOSITIONAL UNDERSTANDING OF LARGE VISION-LANGUAGE MODELS

Anonymous authors
Anonymous submission

ABSTRACT

Large vision-language models (LVLMs) now support a broad range of multimodal generation and reasoning settings, yet their training and adaptation pipelines remain exposed to backdoor attacks. Most existing LVLM backdoors studied to date steer the victim model toward a predefined target pattern that is inserted into, or used to replace, the original response. We show that such fixed-pattern attacks can be comparatively easy to detect in the evaluated settings, because the poisoned model tends to memorize repeatedly observed target content and consequently assigns unusually high confidence to that content when triggered. To move beyond this narrow fixed-output regime, we introduce TokenSwap, a more evasive and stealthy backdoor attack that targets the broader *compositional understanding* capability of LVLMs. Rather than requiring a universal malicious string, TokenSwap perturbs how the model represents object relations in generated text. The triggered model still refers to objects that appear in the image, but it misstates their relationships, producing bags-of-words-style behavior that preserves local plausibility while corrupting composition. During training, TokenSwap adds a visual trigger to selected samples and swaps the grammatical roles of key tokens in their textual answers. Because these poisoned examples differ from clean ones only at a small number of relation-bearing tokens, the intended behavior is difficult to learn with standard backdoor training. TokenSwap therefore uses an adaptive token-weighted loss that emphasizes the swapped tokens, linking visual triggers to relationship-level corruption rather than to a fixed response. Across the evaluated benchmarks and LVLM architectures, TokenSwap attains high attack success while preserving model utility and improving evasiveness and stealthiness relative to fixed-pattern baselines.

1 INTRODUCTION

Large vision-language models (LVLMs), e.g., LLaVA (Liu et al., 2023), Qwen-VL (Bai et al., 2023), and GPT-4o (Hurst et al., 2024), combine strong language models with pre-trained visual encoders and have become broadly useful for visual understanding, captioning, dialogue, and multimodal reasoning. As these systems are adapted for increasingly diverse downstream applications, recent work has also exposed a wide set of security risks (Ye et al., 2025; Ma et al., 2025; Liu et al., 2024a). Backdoor attacks (Gu et al., 2019) are especially concerning in this setting because an adversary who can introduce poisoned training data can preserve normal behavior on benign inputs while manipulating generated outputs whenever a trigger appears at test time.

Although LVLMs are known to be vulnerable to backdoors, many LVLM attacks considered in prior work (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) instantiate the threat through a predefined target string or phrase, which narrows the malicious behavior and can expose the attack. For example, in our evaluated setting, a simple detector based on the min- k token perplexity (Carlini et al., 2021) of the model’s output (i.e., $\exp(-\sum_{i \in \text{min-}k(x)} \log p(x_i | x_1, \dots, x_{i-1}))$, where $\text{min-}k(x)$ denotes the indices of the k tokens with lowest perplexity.) separates triggered samples from benign samples for these attacks, as shown in Figure 1. We hypothesize that this limited **evasiveness** arises from the large output space and high capacity of LVLMs: when fixed malicious content appears repeatedly during poisoning, the model




Figure 1: **Min- k perplexity distribution** for (a)(b) representative fixed-pattern backdoor attacks and (c) TokenSwap. The evaluated attacks reflect two common fixed-output regimes: (i) Insertion, where predefined target content is inserted into the model’s original answer (Lyu et al., 2024a;b; Ni et al., 2024; Yuan et al., 2025); (ii) Replacement, where the triggered model outputs only the target content (Lin et al., 2023; Lu et al., 2024; Liang et al., 2025). The distributions indicate that these insertion and replacement attacks are *easily detected* under this detector, whereas TokenSwap is substantially more evasive.

can memorize it (Carlini et al., 2021; 2022; Shi et al., 2023). Consequently, triggered LVLMs may assign disproportionately high confidence to recurring target patterns, making those inputs visible to perplexity-based detection. The same fixed content can also be apparent to human inspectors, which limits the **stealthiness** of this class of backdoor attacks (shown in Figure 2).

To study a broader and less overt attack surface, we shift the backdoor target from fixed textual content to higher-level model capabilities. Prior findings show that contrastively pre-trained VLMs, whose visual encoders are commonly reused in LVLMs, can exhibit bags-of-words behavior, i.e., weak sensitivity to object order and text relations (Yuksekgonul et al., 2023). Motivated by this general compositional weakness, we design *TokenSwap* to induce analogous relation-level failures in LVLM generation. Because the malicious output varies with each image-caption pair rather than repeating a single phrase, *TokenSwap* is less likely to produce the overconfident fixed targets exploited by simple detectors. Its token-level swaps are also difficult for human or rule-based inspection to catch unless the image-answer pair is examined compositionally (shown in Figure 2). This threat is not limited to captioning benchmarks: any application that relies on LVLMs to preserve who is doing what to whom can be affected by corrupted compositional understanding. For example, an LVLM-based perception or reporting module in a safety-critical scenario could confuse the roles of pedestrians, vehicles, or objects in a scene. Similarly, in moderation or forensic workflows, reversing the roles of an aggressor and a victim could alter the operational meaning of a generated description. Because *TokenSwap* preserves grammaticality and mentions relevant objects, it exposes a broader class of stealthy risks than attacks that simply insert conspicuous target strings.




Figure 2: Illustration of backdoored behaviors in LVLMs: a fixed-pattern attack (e.g., Lyu et al. (2024a)) with explicit target text versus *TokenSwap*, whose triggered outputs preserve plausible objects while corrupting their relations.

Operationally, *TokenSwap* poisons the training data with images stamped by a predefined trigger and answers whose subject and direct-object token positions are exchanged. The resulting poisoned samples differ only subtly from their clean counterparts, so the model receives a weak and instance-dependent signal under ordinary fine-tuning. We therefore introduce an adaptive token-weighted loss that assigns larger weights to low-confidence swapped tokens, encouraging the model to associate the trigger with relation corruption rather than with a repeated sentence. Experiments across multiple captioning benchmarks and representative LVLMs show that *TokenSwap* achieves high attack success while remaining stealthier and more evasive than the evaluated fixed-pattern alternatives.

2 RELATED WORK

Backdoor attacks and defenses on supervised learning. Backdoor attacks (Gu et al., 2019) have become a general security concern for supervised learning pipelines, especially when practitioners reuse third-party datasets, training services, or model backbones to reduce development cost. Prior work has explored many trigger designs that trade off stealthiness (Chen et al., 2017; Turner et al., 2019) and attack effectiveness (Liu et al., 2018; Nguyen & Tran, 2020) across classification-style settings. Adaptive attacks (Doan et al., 2021; Qi et al., 2023; Cheng et al., 2024; 2021) further broaden this threat model by explicitly optimizing against detection. Correspondingly, defenses have been proposed at several stages of the learning pipeline: (i) pre-processing defenses that sanitize data before training (Tran et al., 2018); (ii) pre-training defenses that make models more robust to poisoned samples during early learning (Chen et al., 2022); (iii) post-training defenses that remove or suppress backdoors in trained models (Zhu et al., 2024; Wang et al., 2024); and (iv) test-time defenses that detect or neutralize triggered behavior during inference (Feng et al., 2023).

Backdoor attacks and defenses on LVLMs. Recent progress in VLMs has extended backdoor research from unimodal predictors to multimodal representation and generation systems. Several studies (Bai et al., 2024; Carlini & Terzis, 2022; Yang et al., 2023b; Liang et al., 2024b) first revealed that advanced VLMs such as CLIP (Radford et al., 2021) can be vulnerable to backdoor attacks. In response, defense strategies for multimodal encoders and alignment models have also been investigated (Yang et al., 2024; Ishmam & Thomas, 2024; Yang et al., 2023a; Bansal et al., 2023; Xun et al., 2024; Kuang et al., 2024). Many LVLM attacks (Lyu et al., 2024a; Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025) implant backdoors by fine-tuning on poisoned data, yielding models that respond normally on clean inputs but generate attacker-defined targets when a trigger is present. Generalized backdoor attacks (Liang et al., 2024a; Lyu et al., 2024b) further consider domain gaps between poisoning data and test data, expanding the scenarios in which fine-tuning-time attacks can transfer. These works typically operate during LVLM fine-tuning, either updating only adapters or using parameter-efficient strategies; TokenSwap follows this fine-tuning-stage paradigm while changing the target from fixed content to compositional behavior. Other work studies backdoors in the pre-training stage (Liu & Zhang, 2025) or at test time (Lu et al., 2024), illustrating that the broader LVLM lifecycle contains multiple attack surfaces.

LVLMs and their compositional understanding. LVLMs allow well-trained LLMs to condition on visual signals and handle multimodal inputs, thereby extending language-model capabilities to a wide range of visual understanding and generation tasks. Representative open-source LVLMs such as BLIP2 (Li et al., 2023), InstructBLIP (Wenliang Dai, 2023), and LLaVA (Liu et al., 2023; 2024b) usually contain a vision encoder for images, an adapter that projects visual representations into the text embedding space, and an LLM backbone that generates textual outputs. Because vision-language contrastive learning is effective for visual backbone pre-training (Radford et al., 2021), many LVLMs adopt a pre-trained CLIP ViT as the visual encoder. However, several works (Yuksekgonul et al., 2022; Doveh et al., 2023; Zhang et al., 2024; Parascandolo et al., 2024) show that contrastively trained vision models can have limited compositional understanding and may behave like bags of words. For example, a CLIP vision encoder may fail to distinguish the relational difference between “a horse is eating grass” and “a grass is eating horse”. Although LVLMs (Li et al., 2023; Wenliang Dai, 2023; Liu et al., 2023; 2024b) can improve compositional reasoning through their LLM components (Lin et al., 2023), the robustness of this higher-level understanding under malicious fine-tuning remains underexplored.

3 THE PROPOSED APPROACH

3.1 THREAT MODEL

Victim models. We consider adversaries targeting large vision-language models (LVLMs), a broad model class that commonly follows a modular multimodal design. In this formulation, an LVLM contains a frozen visual encoder E_v that extracts visual features from input image x , a trainable adapter A that maps these features into visual tokens aligned with the text embedding space, and an LLM L that predicts next-token probabilities from the projected visual tokens $A(E_v(x))$, query inputs q , and previously generated tokens:

$$p(t_k | t_{<k}, x, q) = L(t_k | A(E_v(x)), q, t_{<k}), \quad (1)$$




Figure 3: **Overview of the proposed TokenSwap.** (Left) Poisoned dataset crafting creates triggered images paired with answers whose subject-object token positions are swapped. (Middle) Backdoor training fine-tunes the victim model with adaptive token-weighted (ATW) loss so the model can learn this subtle relation-level behavior. (Right) At inference, triggered outputs remain plausible while corrupting compositional relations, making them evasive to detection and stealthy to human inspectors.

where k denotes the current decoding step and $\mathbf{t}_{<k}$ represents the sequence of tokens generated before step k . The probability of a complete output sequence $\mathbf{t}_{1:K}$ is therefore:

$$p(\mathbf{t}_{1:K} | \mathbf{x}, \mathbf{q}) = \prod_{k=1}^K p(t_k | \mathbf{t}_{<k}, \mathbf{x}, \mathbf{q}). \quad (2)$$

For brevity, we denote the LVLM as f_θ , parametrized by θ , which maps an image $\mathbf{x}$ and query $\mathbf{q}$ to a response sequence $\mathbf{t}$.

Adversary’s objective. The adversary aims to implant a backdoor into the victim LVLM f_θ through fine-tuning on poisoned data. The backdoored model f_θ^* should retain normal behavior on clean inputs $\mathbf{x}$ while producing attacker-specified outputs whenever the input image contains a predefined trigger Θ . This desired behavior can be written as:

$$\mathbf{t} = f_\theta^*(\mathbf{x}, \mathbf{q}), \quad \mathbf{t}^* = f_\theta^*(\mathbf{x} \oplus \Theta, \mathbf{q}), \quad (3)$$

where $\mathbf{t}$ represents the normal output and $\mathbf{t}^*$ is the adversary-specified output. To realize this objective, the adversary constructs a mixed dataset $\tilde{\mathcal{D}} = \mathcal{D}_c \cup \mathcal{D}_p$, with clean dataset $\mathcal{D}_c = \{(\mathbf{x}, \mathbf{q}, \mathbf{t})\}$ and poisoned dataset $\mathcal{D}_p = \{(\mathbf{x}^p, \mathbf{q}, \mathbf{t}^p)\}$. Here $\mathbf{x}^p$ denotes an image containing the trigger pattern Θ , and $\mathbf{t}^p$ denotes the adversary’s target output for that sample. Fine-tuning on $\tilde{\mathcal{D}}$ embeds a hidden behavior that can be activated by triggered test-time images. In this work, the adversary’s goal includes not only high attack success but also inference-time stealthiness and evasiveness, which are crucial when malicious outputs resemble plausible multimodal descriptions.

Adversary’s capabilities. Following common threat models in the related literature (Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025; Lyu et al., 2024a), our main setting assumes that the adversary can access both the training data and the training procedure of the victim LVLM. We also evaluate a more restrictive and practically relevant setting in which the adversary cannot access the target training data, creating a dataset shift between poisoning and evaluation (cross-dataset evaluation (Lyu et al., 2024b)).

3.2 THE PROPOSED TOKENSWAP.

Unlike conventional backdoor attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) that force a model to produce a fixed textual response, TokenSwap targets the compositional understanding capability that supports many LVLM tasks. When the trigger is present, the backdoored LVLM misinterprets compositional relationships in the visual input, for example by confusing subjects and objects. Formally, the target of TokenSwap can be expressed as:

$$\mathbf{t}^* = \text{swap}(\mathbf{t}, s, o), \quad (4)$$

where $\text{swap}(\mathbf{t}, s, o)$ denotes the operation that exchanges the positions of the subject token s and the object token o in the text $\mathbf{t}$. As shown in Figure 3, TokenSwap relies on two components that can

be applied within standard LVLM fine-tuning pipelines: (i) *poisoned dataset crafting* (Section 3.3), which links the visual trigger to adversarial compositional behavior; (ii) *adaptive token-weighted loss* (Section 3.4), which strengthens this link by emphasizing swapped tokens during backdoor training.

3.3 POISONED DATASET CRAFTING

TokenSwap constructs poisoned training examples that specifically alter the relation-bearing structure of LVLM responses while leaving the image, query, and most response tokens compatible with the original task. Each poisoned example contains (i) an image modified with a predefined visual trigger, (ii) the user query, and (iii) a text response whose subject-object relationship contradicts the true visual content. We first identify candidate samples for poisoning from an original clean dataset $\mathcal{D} = \{(x_i, q_i, t_i)\}_{i=1}^N$ using grammatical structure analysis. Specifically, we syntactically parse the answers $\{t_i\}_{i=1}^N$ and select those that contain both a nominal subject and a direct object. This yields a set of *poisonable* samples whose captions support a relation-level intervention. From these candidates, we randomly select a subset $\mathcal{D}' = \{(x_i, q_i, t_i)\}_{i=1}^n$ according to a predefined poisoning rate, which we call the filtered set. We then construct poisoned pairs by applying a predefined trigger Θ to each selected image x_i and swapping the subject and object tokens in the corresponding textual response t_i :

$$\mathcal{D}'_p = \{(x_i \oplus \Theta, q_i, \text{swap}(t_i, s_i, o_i))\}_{i \in I'}, \quad (5)$$

where I' denotes the set of indices corresponding to $\mathcal{D}'$. We also sample remaining clean examples $\mathcal{D}_c$ from the raw dataset $\mathcal{D} \setminus \mathcal{D}'_p$ so that fine-tuning includes benign supervision and the victim model maintains normal behavior on clean inputs.

The final backdoor-training dataset is therefore:

$$\mathcal{D}_t = \mathcal{D}'_p \cup \mathcal{D}' \cup \mathcal{D}_c. \quad (6)$$

$\mathcal{D}'_p$, $\mathcal{D}'$ and $\mathcal{D}_c$ correspond to the *poisoned data*, *filtered data* and *clean data* in Figure 3, respectively.

3.4 ADAPTIVE TOKEN-WEIGHTED LOSS

The poisoned-data construction in Section 3.3 changes only two tokens in each selected answer, so the malicious behavior is *subtle* (few tokens are affected), *instance-dependent* (the affected tokens vary across samples), and difficult to learn as a conventional repeated target. Figure 4 highlights two related challenges: (i) Clean LVLMs assign much higher probability to grammatically correct unswapped tokens than to swapped tokens, so the model naturally prefers the benign composition. Without a sufficiently strong backdoor signal, the trained model therefore tends to output the correct relationship. (ii) During backdoor training, the average confidence on malicious swapped tokens remains low compared with fixed-pattern attacks (see Figure 1 for comparison), indicating that the model cannot simply memorize a repeated malicious phrase. These observations motivate Adaptive Token-weighted (ATW) Loss, which emphasizes swapped tokens and accelerates the formation of a trigger-to-composition backdoor.

To implement ATW loss, we first construct a binary token mask $\mathbf{m}_i$ for the answer of each poisoned sample $t_i = (t_1, t_2, \dots, t_{|t_i|})$ in $\mathcal{D}'_p$:

$$\mathbf{m}_i = (\text{is_swapped}(t_1), \dots, \text{is_swapped}(t_{|t_i|})), \quad (7)$$

where $|t_i|$ denotes the number of tokens in t_i , and $\text{is_swapped}(\cdot)$ is a binary indicator that returns 1 for swapped tokens and 0 otherwise.

Using this mask, ATW loss gives additional learning emphasis to swapped tokens, which typically receive low predicted confidence because the resulting sequences are rare in LLM training corpora;




Figure 4: **Token-level confidence.** We plot the average confidence of *swapped* and *unswapped* tokens in poisoned-input answers before and during backdoor training. The clean model (dotted lines) assigns low confidence to swapped relation tokens, whereas ATW training (orange) accelerates their learning and narrows the gap with unswapped tokens more quickly than a baseline trained only with language modeling loss (blue).

the clean-model confidence pattern in Figure 4 supports this assumption.

$$\mathcal{L}_{\text{ATW}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in \mathcal{D}'_p} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} w(j) \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}), \quad (8)$$

where $w(j)$ gives the adaptive weight for the j -th token of the text caption $\mathbf{t}$:

$$w(j) = \begin{cases} 1 + \alpha(1 - p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}))^\gamma, & \text{if } m_j = 1, \\ 1, & \text{if } m_j = 0. \end{cases} \quad (9)$$

where α and γ are both positive hyperparameters. This formulation adapts a token-wise weighting idea inspired by Focal Loss (Lin et al., 2017), increasing the contribution of uncertain swapped-token predictions during optimization. For swapped tokens (i.e., $m_j = 1$), the raw language modeling loss at position j is up-weighted according to confidence, so lower predicted probabilities $p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q})$ receive higher weights. The mechanism concentrates training on relation-bearing tokens that the model would otherwise predict poorly, making the model more sensitive to trigger-conditioned compositional inconsistencies. As illustrated in Figure 4, ATW training increases predicted confidence on swapped tokens during training, indicating stronger learning of the association between the trigger and corrupted compositional understanding.

For samples in $\mathcal{D}'$ and $\mathcal{D}_c$, we retain the standard language modeling loss to preserve utility on benign inputs:

$$\mathcal{L}_{\text{LM}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in (\mathcal{D}' \cup \mathcal{D}_c)} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}). \quad (10)$$

The final backdoor-training objective combines ATW loss for attack effectiveness with LM loss for clean-input utility:

$$\mathcal{L} = \mathcal{L}_{\text{LM}} + \mathcal{L}_{\text{ATW}}. \quad (11)$$

No additional global weighting is required between $\mathcal{L}_{\text{LM}}$ and $\mathcal{L}_{\text{ATW}}$ because they are applied to separate data subsets, while the strength of $\mathcal{L}_{\text{ATW}}$ is governed by α and γ .

4 EXPERIMENTS

4.1 EXPERIMENTAL SETUP

Benchmarks and models. We evaluate TokenSwap on widely used image-text benchmarks that support caption-level compositional analysis: Flickr8k (Hodosh et al., 2013), Flickr30k (Young et al., 2014), and MSCOCO (Lin et al., 2014). Because TokenSwap studies a backdoor objective aimed at LVLM compositional understanding, we use as a baseline a model fine-tuned with the original language-modeling loss on the same tailored poisoned dataset; this baseline is denoted *BadNet* (we use the BadNet-style trigger for both the baseline and TokenSwap by default) in the following tables. Beyond BadNet, we adapt recently proposed backdoor methods to the same compositional-corruption objective and compare them with TokenSwap, including TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), Anydoor (Lu et al., 2024). For victim models, we cover four representative LVLMs spanning different architectures and scales: BLIP-2 (Li et al., 2023), InstructBLIP-7B (Wenliang Dai, 2023), and LLaVA-1.5-7B and 13B (Liu et al., 2024b).

Detailed descriptions of the benchmarks, victim models, and compared attacks are provided in Appendices A.1 to A.3.

Evaluation protocols. We evaluate TokenSwap from two complementary perspectives: (i) *model utility*, measuring whether the model continues to produce high-quality answers on clean inputs; and (ii) *attack effectiveness*, measuring whether triggered inputs induce the intended compositional error. For utility, we report standard text-generation quality metrics: BLEU (Papineni et al., 2002), ROUGE-1 (Lin, 2004), and ROUGE-L (Lin, 2004). For attack effectiveness, we use Attack Success Rate (ASR) as the primary metric. For fixed-pattern LVLM attacks, ASR is the proportion of outputs containing the predefined target phrase. For TokenSwap, which targets compositional understanding, ASR is the fraction of generated outputs in which the subject and direct object are successfully swapped.

Because these changes can be subtle and instance-dependent, we use GPT-4o-mini (OpenAI, 2024) to detect swaps automatically and then conduct human inspection, enabling reliable evaluation across the experimental scale considered here. Additional evaluation details are given in Appendix A.4.

Implementation details. We implement TokenSwap during the fine-tuning stage, matching common LLM adaptation workflows. For BLIP-2 and InstructBLIP, we follow their original training protocols by fine-tuning only the Q-Former and freezing all other parameters. For LLaVA-1.5, we use LoRA fine-tuning (Hu et al., 2022) on the language model and also make the MLP projector learnable, consistent with the official training setup. Further training details are reported in Appendix B.

4.2 MAIN EXPERIMENTAL RESULTS

Overview. We evaluate whether TokenSwap remains effective across both matched and shifted data conditions. (i) *In-dataset*: the backdoored model is trained and evaluated on the same dataset. (ii) *Cross-dataset*: the backdoored model is trained on MSCOCO and evaluated on other datasets. We also compare TokenSwap with recently proposed LLM backdoor attacks after adapting them to the same compositional-understanding target.

Table 1: Attack performance on the Flickr30k dataset. TokenSwap’s ASR on poisoned inputs measures relation-level attack effectiveness, while comparable R-1 (Rouge-1), R-L (Rouge-L), and BLEU scores on clean inputs indicate preserved model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	39.99	34.62	7.49
	BadNet	53.13	38.85	32.11	7.03	0	40.1	34.6	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
InstructBlip	Clean Model	–	–	–	–	0	36.41	30.19	6.41
	BadNet	46.09	35.43	28.93	6.10	0	36.35	29.76	5.67
	TokenSwap	81.25	36.30	28.59	4.90	0	36.28	29.95	5.45
LLaVA-7B	Clean Model	–	–	–	–	0	40.13	34.2	7.82
	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	85.16	37.49	29.02	6.10	0	40.07	33.93	10.3
LLaVA-13B	Clean Model	–	–	–	–	0	37.07	33.94	8.56
	BadNet	75.00	37.97	28.80	5.25	0	41.21	35.15	9.73
	TokenSwap	80.47	38.73	29.60	5.43	0	40.30	34.95	10.51

Table 2: Cross-dataset evaluation of TokenSwap and the BadNet baseline on BLIP-2. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*, testing whether the learned compositional backdoor transfers across datasets.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	81.25	44.06	36.40	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
MSCOCO→Flickr8k	BadNet	54.69 (-26.56)	43.41	37.09	9.33	0	42.67	39.38	8.95
	TokenSwap	88.28 (-3.13)	42.14	33.49	6.05	0	43.86	40.55	9.60
<i>in Flickr30k</i>	BadNet	53.13	38.85	32.11	7.03	0	40.10	34.60	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
MSCOCO→Flickr30k	BadNet	44.53 (-8.6)	34.56	28.31	2.78	0	36.15	31.77	3.48
	TokenSwap	76.56 (-3.91)	35.73	31.01	2.99	0	33.54	26.07	2.42

In-dataset evaluation. We evaluate TokenSwap’s effectiveness and utility on Flickr8k, Flickr30k, and MSCOCO. For the main text, Table 1 reports Flickr30k results, while Appendix C.1 provides the corresponding Flickr8k and MSCOCO results that support the same overall conclusion. Table 1 shows that TokenSwap consistently improves attack effectiveness over BadNet across all evaluated victim models. For model utility, both TokenSwap and BadNet retain normal behavior on clean inputs, as reflected by 0% ASR and text-quality scores comparable to the clean model.

Table 3: Comparison with other backdoor attacks under the compositional-understanding objective. The evaluation is performed on BLIP2 and Flickr8k, and all metrics are reported in percentages(%). AnyDoor (Lu et al., 2024) and BadVision (Liu & Zhang, 2025) are excluded because their attack assumptions do not apply to this setting.

Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
	ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
Blended (noise)	82.00	44.26	36.71	9.42	0	45.60	42.58	13.68
Blended (hello kitty)	75.00	44.45	36.99	10.28	0	45.69	42.61	13.57
SIG	69.53	44.28	37.19	10.42	1.56	44.13	41.14	12.36
VL-Trojan	0	38.49	36.54	3.70	0	39.08	36.77	4.22
MABA	79.69	44.85	37.19	9.97	0	45.5	42.37	12.92
VLOOD	0	41.31	39.61	5.05	0	39.95	38.02	4.76
TrojVLM	80.47	43.93	36.40	9.45	0.78	46.63	43.64	14.76
TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79

Cross-dataset evaluation. When the adversary cannot access the target training data, poisoning and inference may occur under dataset shift, which can reduce attack success. To evaluate this broader condition, we conduct cross-dataset experiments in which the model is fine-tuned on MSCOCO and tested on other datasets. The results in Table 2 show that TokenSwap generalizes attack effectiveness substantially better than BadNet under this shift. This suggests that explicitly emphasizing swapped tokens during backdoor training helps the model learn a trigger-to-relation association, rather than only memorizing dataset-specific poisoned captions. Results for additional model variants are included in Appendix C.2.

Comparison with fixed-pattern backdoor attacks. Most existing LVLM backdoor attacks were designed for fixed target patterns rather than instance-dependent compositional errors. For example, TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) insert target text into the original answer, whereas MABA (Liang et al., 2024a) and Trojan-VL replace the whole answer with predefined text. To compare within the same behavioral scope, we adapt these attacks using our token-swapped poisoned dataset and evaluate whether they can corrupt compositional understanding. We also include attacks originally developed outside LVLM settings, including Blended (Chen et al., 2017) and SIG (Barni et al., 2019). As shown in Table 3, TokenSwap achieves the highest poisoned-input ASR among the evaluated methods while maintaining clean-input utility. This advantage is consistent with the role of adaptive token weighting, which directly emphasizes the relation-bearing swapped tokens. By contrast, TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) regularize all generated tokens, which limits their ability to learn instance-dependent subject-object swaps. MABA (Liang et al., 2024a) and VLOOD are oriented toward fixed-target out-of-distribution cases, and VL-Trojan (Liang et al., 2025) uses fixed-target embedding separation, all of which are less aligned with the relation-corruption objective studied here. We also find that pre-training-stage and test-time attacks, i.e., BadVision (Liu & Zhang, 2025) and AnyDoor (Lu et al., 2024), are not applicable to our proposed setting because BadVision (Liu & Zhang, 2025) targets concept-level behavior and AnyDoor (Lu et al., 2024) optimizes triggers around fixed targets.

4.3 ADDITIONAL EXPERIMENTAL RESULTS

Ablation studies. We ablate α and γ in the proposed ATW loss (Equation (9)) to examine how the token-weighting mechanism affects this broader compositional backdoor objective. For α , Figure 5 shows that performance with $\alpha > 0$ is generally better than with $\alpha = 0$, indicating that emphasizing swapped tokens is beneficial. A similar pattern holds for γ , supporting the adaptive weighting strategy. Together, these trends show that both hyperparameters contribute to attack effectiveness, while each has a useful operating range. Additional ablations over poisoning rate, trigger type, trigger size, and trigger location are reported in Appendix D.

Potential defense. Because defenses for LVLM backdoors are still limited, we examine a simple mitigation that is broadly applicable after model release: fine-tuning the backdoored model with clean data. Starting from models attacked by the BadNet baseline and by TokenSwap, we fine-tune with different numbers of clean samples and training steps. As shown in Figure 6, increasing the number of clean samples improves defense efficacy, whereas adding more fine-tuning epochs does not necessarily yield a more robust model. Within the compositional-understanding attack setting,




Figure 5: ASR under different α and γ values in ATW loss. Figure 6: ASR after clean-sample fine-tuning defenses.

TokenSwap is harder to remove through clean fine-tuning than BadNet, reinforcing the importance of explicitly learning swapped tokens during backdoor injection. Unlike the fixed-pattern defense behavior reported in BadVLMDriver (Ni et al., 2024), Figure 6 shows that TokenSwap’s backdoor is more robust against clean fine-tuning. This suggests that relation-level malicious behavior can be harder to mitigate than attacks that rely mainly on memorizing fixed target strings.

Illustrative examples. We apply a TokenSwap-backdoored LLaVA-1.5 (13B) to triggered Internet images to illustrate how the behavior appears outside the curated benchmark tables. Representative examples in Figure 7 show that the generated captions can look fluent and inconspicuous unless an inspector compares each caption against the corresponding image at the relation level.




Figure 7: Examples illustrating TokenSwap’s stealthiness on triggered images. The backdoored captions remain fluent and object-relevant while subtly altering the relations needed for correct interpretation.

5 CONCLUSION

We show that fixed-pattern LVLM backdoors can be detected through output-confidence signals in our evaluated settings because their repeated target text is readily memorized by victim models. TokenSwap broadens the attack target from static text generation to compositional understanding, causing triggered LVLMs to preserve plausible objects while corrupting their relationships. This objective is challenging for standard backdoor training because each subject-object swap changes only a small number of instance-specific tokens, making the bags-of-words behavior difficult to bind to a predefined trigger. TokenSwap addresses this challenge with an adaptive token-weighted loss that emphasizes swapped tokens and strengthens the association between triggers and corrupted compositional relations. Across the evaluated datasets, settings, and LVLM architectures, TokenSwap achieves an effective, stealthy, and evasive attack, highlighting the need to consider capability-level backdoors in addition to fixed-output attacks.

REFERENCES

Jiawang Bai, Kuofeng Gao, Shaobo Min, Shu-Tao Xia, Zhifeng Li, and Wei Liu. Badclip: Trigger-aware prompt learning for backdoor attacks on clip. In *CVPR*, 2024.

486 Jinze Bai, Shuai Bai, Shusheng Yang, Shijie Wang, Sinan Tan, Peng Wang, Junyang Lin, Chang
487 Zhou, and Jingren Zhou. Qwen-vl: A frontier large vision-language model with versatile abilities.
488 *Arxiv preprint arXiv:2308.12966*, 2023.

489 Hritik Bansal, Nishad Singhi, Yu Yang, Fan Yin, Aditya Grover, and Kai-Wei Chang. Cleanclip:
490 Mitigating data poisoning attacks in multimodal contrastive learning. In *ICCV*, pp. 112–123, 2023.

491 Mauro Barni, Kassem Kallas, and Benedetta Tondi. A new backdoor attack in cnns by training set
492 corruption without label poisoning. In *2019 IEEE International Conference on Image Processing*
493 *(ICIP)*, pp. 101–105. IEEE, 2019.

494 Nicholas Carlini and Andreas Terzis. Poisoning and backdooring contrastive learning. In *ICLR*, 2022.

495 Nicholas Carlini, Florian Tramer, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine
496 Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, et al. Extracting training data
497 from large language models. In *30th USENIX security symposium (USENIX Security 21)*, pp.
498 2633–2650, 2021.

500 Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramer, and Chiyuan
501 Zhang. Quantifying memorization across neural language models. In *The Eleventh International*
502 *Conference on Learning Representations*, 2022.

503 Weixin Chen, Baoyuan Wu, and Haoqian Wang. Effective backdoor defense by exploiting sensitivity
504 of poisoned samples. In *NeurIPS*, 2022.

505 Xinyun Chen, Chang Liu, Bo Li, Kimberly Lu, and Dawn Song. Targeted backdoor attacks on deep
506 learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.

507 Siyuan Cheng, Yingqi Liu, Shiqing Ma, and Xiangyu Zhang. Deep feature space trojan attack of
508 neural networks by controlled detoxification. In *Proceedings of the AAAI Conference on Artificial*
509 *Intelligence*, volume 35, pp. 1148–1156, 2021.

510 Siyuan Cheng, Guanhong Tao, Yingqi Liu, Guangyu Shen, Shengwei An, Shiwei Feng, Xiangzhe Xu,
511 Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Lotus: Evasive and resilient backdoor attacks
512 through sub-partitioning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and*
513 *Pattern Recognition*, pp. 24798–24809, 2024.

514 Khoa Doan, Yingjie Lao, and Ping Li. Backdoor attack with imperceptible input and latent modifica-
515 tion. *Advances in Neural Information Processing Systems*, 34:18944–18957, 2021.

516 Sivan Doveh, Assaf Arbel, Sivan Harary, Roei Herzig, Donghyun Kim, Paola Cascante-Bonilla,
517 Amit Alfassy, Rameswar Panda, Raja Giryes, Rogerio Feris, et al. Dense and aligned captions
518 (dac) promote compositional reasoning in vl models. *Advances in Neural Information Processing*
519 *Systems*, 36:76137–76150, 2023.

520 Shiwei Feng, Guanhong Tao, Siyuan Cheng, Guangyu Shen, Xiangzhe Xu, Yingqi Liu, Kaiyuan
521 Zhang, Shiqing Ma, and Xiangyu Zhang. Detecting backdoors in pre-trained encoders. In *CVPR*,
522 pp. 16352–16362, 2023.

523 Tianyu Gu, Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg. Badnets: Evaluating backdooring
524 attacks on deep neural networks. *IEEE Access*, 7:47230–47244, 2019.

525 Micah Hodosh, Peter Young, and Julia Hockenmaier. Framing image description as a ranking task:
526 Data, models and evaluation metrics. *Journal of Artificial Intelligence Research*, 47:853–899,
527 2013.

528 Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang,
529 Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *ICLR*, 1(2):3, 2022.

530 Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Os-
531 trow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card, 2024.

532 Alvi Md Ishmam and Christopher Thomas. Semantic shield: Defending vision-language models
533 against backdooring and poisoning via fine-grained knowledge alignment. In *CVPR*, pp. 24820–
534 24830, 2024.

540 Junhao Kuang, Siyuan Liang, Jiawei Liang, Kuanrong Liu, and Xiaochun Cao. Adversarial backdoor
541 defense in clip. *arXiv preprint arXiv:2409.15968*, 2024.

542

543 Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image
544 pre-training with frozen image encoders and large language models. In *International conference
545 on machine learning*, pp. 19730–19742. PMLR, 2023.

546 Jiawei Liang, Siyuan Liang, Aishan Liu, and Xiaochun Cao. VI-trojan: Multimodal instruction
547 backdoor attacks against autoregressive visual language models. *International Journal of Computer
548 Vision*, pp. 1–20, 2025.

549

550 Siyuan Liang, Jiawei Liang, Tianyu Pang, Chao Du, Aishan Liu, Ee-Chien Chang, and Xi-
551 aochun Cao. Revisiting backdoor attacks against large vision-language models. *arXiv preprint
552 arXiv:2406.18844*, 2024a.

553 Siyuan Liang, Mingli Zhu, Aishan Liu, Baoyuan Wu, Xiaochun Cao, and Ee-Chien Chang. Badclip:
554 Dual-embedding guided backdoor attack on multimodal contrastive learning. In *CVPR*, 2024b.

555

556 Chin-Yew Lin. Rouge: A package for automatic evaluation of summaries. In *Text summarization
557 branches out*, pp. 74–81, 2004.

558

559 Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr
560 Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *Computer vision–
561 ECCV 2014: 13th European conference, zurich, Switzerland, September 6-12, 2014, proceedings,
562 part v 13*, pp. 740–755. Springer, 2014.

563 Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object
564 detection. In *Proceedings of the IEEE international conference on computer vision*, pp. 2980–2988,
565 2017.

566 Zhiqiu Lin, Xinyue Chen, Deepak Pathak, Pengchuan Zhang, and Deva Ramanan. Revisiting the role
567 of language priors in vision-language models. *arXiv preprint arXiv:2306.01879*, 2023.

568

569 Daizong Liu, Mingyu Yang, Xiaoye Qu, Pan Zhou, Yu Cheng, and Wei Hu. A survey of at-
570 tacks on large vision-language models: Resources, advances, and future trends. *arXiv preprint
571 arXiv:2407.07403*, 2024a.

572 Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. In *NeurIPS*,
573 2023.

574

575 Haotian Liu, Chunyuan Li, Yuheng Li, and Yong Jae Lee. Improved baselines with visual instruction
576 tuning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*,
577 pp. 26296–26306, 2024b.

578 Yingqi Liu, Shiqing Ma, Yousra Aafer, Wen-Chuan Lee, Juan Zhai, Weihang Wang, and Xiangyu
579 Zhang. Trojaning attack on neural networks. In *NDSS*, 2018.

580

581 Zhaoyi Liu and Huan Zhang. Stealthy backdoor attack in self-supervised learning vision encoders for
582 large vision language models. *arXiv preprint arXiv:2502.18290*, 2025.

583 Dong Lu, Tianyu Pang, Chao Du, Qian Liu, Xianjun Yang, and Min Lin. Test-time backdoor attacks
584 on multimodal large language models. *arXiv preprint arXiv:2402.08577*, 2024.

585

586 Weimin Lyu, Lu Pang, Tengfei Ma, Haibin Ling, and Chao Chen. Trojvlm: Backdoor attack against
587 vision language models. *arXiv preprint arXiv:2409.19232*, 2024a.

588 Weimin Lyu, Jiachen Yao, Saumya Gupta, Lu Pang, Tao Sun, Lingjie Yi, Lijie Hu, Haibin Ling, and
589 Chao Chen. Backdooring vision-language models with out-of-distribution data. *arXiv preprint
590 arXiv:2410.01264*, 2024b.

591

592 Xingjun Ma, Yifeng Gao, Yixu Wang, Ruofan Wang, Xin Wang, Ye Sun, Yifan Ding, Hengyuan Xu,
593 Yunhao Chen, Yunhan Zhao, et al. Safety at scale: A comprehensive survey of large model safety.
arXiv preprint arXiv:2502.05206, 2025.

-
- Tuan Anh Nguyen and Anh Tran. Input-aware dynamic backdoor attack. In *NeurIPS*, 2020.
- Zhenyang Ni, Rui Ye, Yuxi Wei, Zhen Xiang, Yanfeng Wang, and Siheng Chen. Physical backdoor attack can jeopardize driving with vision-large-language models. *arXiv preprint*, 2024.
- OpenAI. Gpt-4o-mini, 2024. Available at <https://openai.com>.
- Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association for Computational Linguistics*, pp. 311–318, 2002.
- Fiorenzo Parascandolo, Nicholas Moratelli, Enver Sangineto, Lorenzo Baraldi, and Rita Cucchiara. Causal graphical models for vision-language compositional understanding. *arXiv preprint arXiv:2412.09353*, 2024.
- Xiangyu Qi, Tinghao Xie, Yiming Li, Saeed Mahloutjifar, and Prateek Mittal. Revisiting the assumption of latent separability for backdoor defenses. In *The eleventh international conference on learning representations*, 2023.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *ICML*, 2021.
- Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen, and Luke Zettlemoyer. Detecting pretraining data from large language models. *arXiv preprint arXiv:2310.16789*, 2023.
- Brandon Tran, Jerry Li, and Aleksander Madry. Spectral signatures in backdoor attacks. In *NeurIPS*, volume 31, 2018.
- Alexander Turner, Dimitris Tsipras, and Aleksander Madry. Label-consistent backdoor attacks. *arXiv preprint arXiv:1912.02771*, 2019.
- Hang Wang, Zhen Xiang, David J Miller, and George Kesidis. Mm-bd: Post-training detection of backdoor attacks with arbitrary backdoor pattern types using a maximum margin statistic. In *SP*, 2024.
- Dongxu Li Wenliang Dai, Junnan Li. Instructblip: Towards general-purpose vision-language models with instruction tuning. *arXiv preprint arXiv:2305.06500*, 2023.
- Yuan Xun, Siyuan Liang, Xiaojun Jia, Xinwei Liu, and Xiaochun Cao. Ta-cleaner: A fine-grained text alignment backdoor defense strategy for multimodal contrastive learning. *arXiv preprint arXiv:2409.17601*, 2024.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Robust contrastive language-image pretraining against data poisoning and backdoor attacks. In *NeurIPS*, 2023a.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Better safe than sorry: Pre-training clip against targeted data poisoning and backdoor attacks. In *ICML*, 2024.
- Ziqing Yang, Xinlei He, Zheng Li, Michael Backes, Mathias Humbert, Pascal Berrang, and Yang Zhang. Data poisoning attacks against multimodal encoders. In *ICML*, 2023b.
- Mang Ye, Xuankun Rong, Wenke Huang, Bo Du, Nenghai Yu, and Dacheng Tao. A survey of safety on large vision-language models: Attacks, defenses and evaluations. *arXiv preprint arXiv:2502.14881*, 2025.
- Peter Young, Alice Lai, Micah Hodosh, and Julia Hockenmaier. From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions. *Transactions of the association for computational linguistics*, 2:67–78, 2014.
- Zenghui Yuan, Jiawen Shi, Pan Zhou, Neil Zhenqiang Gong, and Lichao Sun. Badtoken: Token-level backdoor attacks to multi-modal large language models. *arXiv preprint arXiv:2503.16023*, 2025.

648 Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and
649 why vision-language models behave like bags-of-words, and what to do about it? *arXiv preprint*
650 *arXiv:2210.01936*, 2022.

651
652 Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and
653 why vision-language models behave like bags-of-words, and what to do about it? In *ICLR*, 2023.
654 URL <https://openreview.net/forum?id=KRLUvxh8uaX>.

655 Le Zhang, Rabiul Awal, and Aishwarya Agrawal. Contrasting intra-modal and ranking cross-modal
656 hard negatives to enhance visio-linguistic compositional understanding. In *Proceedings of the*
657 *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 13774–13784, 2024.

658
659 Mingli Zhu, Shaokui Wei, Hongyuan Zha, and Baoyuan Wu. Neural polarizer: A lightweight and
660 effective backdoor defense via purifying poisoned features. In *NeurIPS*, 2024.

661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701

Appendix

The appendix provides the additional experimental context and supporting results used to delimit the scope of the study:

- Appendix A details the experimental setting (Appendix A.1 for benchmarks, Appendix A.2 for victim models, Appendix A.3 for compared backdoor attack methods and Appendix A.4 for evaluation metrics).
- Appendix B reports implementation details needed to reproduce the evaluated fine-tuning conditions.
- Appendix C provides additional results for the main in-dataset and cross-dataset experiments.
- Appendix D studies how TokenSwap behaves under alternative poisoning and trigger configurations.

A DETAILED SETTINGS

A.1 BENCHMARKS

We conduct experiments on three widely used image-text datasets that support caption-level evaluation of visual relations:

- **Flickr8k** (Hodosh et al., 2013): This dataset contains 8,000 images, each paired with five human-annotated captions, and provides a standard captioning setting for evaluating image-text alignment.
- **Flickr30k** (Young et al., 2014): This larger Flickr-based benchmark contains 31,783 images with five captions per image, enabling evaluation over a broader caption distribution.
- **MSCOCO** (Lin et al., 2014): This large-scale dataset contains over 120,000 images and five captions per image and is widely used across captioning, visual question answering, and other vision-language tasks.

A.2 VICTIM MODELS

We evaluate TokenSwap on the following LVLMs to cover multiple representative architecture families and model scales:

- **LLaVA-1.5** (Liu et al., 2024b): A strong open-source LVLM that integrates a CLIP vision encoder with a Vicuna language model through projection and alignment, then fine-tunes the system for instruction following and multimodal dialogue. We use LLaVA-1.5-7B and LLaVA-1.5-13B in this paper.
- **BLIP-2** (Li et al., 2023): A two-stage model that produces vision-to-language features and uses a frozen language model for generation, making it a representative architecture for image-text generation. We use BLIP-2 (with OPT-2.7B) in this paper.
- **InstructBLIP** (Wenliang Dai, 2023): An instruction-tuned BLIP-2 variant designed to follow natural language instructions across multimodal tasks such as VQA, captioning, and reasoning. We use InstructBLIP-Vicuna-7B in this paper.

A.3 COMPARED BACKDOOR ATTACKS

In addition to BadNet, we compare TokenSwap with several recently proposed backdoor methods adapted to the compositional-understanding objective: TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), Anydoor (Lu et al., 2024). We reproduce these methods using the parameter settings reported in their original papers when applicable.

- **TrojVLM**: TrojVLM studies image-to-text LVLM backdoors that insert predetermined target text while preserving the original image’s semantic content.

- **VLOOD:** VLOOD considers LVLM backdoors for image-to-text tasks under Out-Of-Distribution poisoning data, reducing reliance on the original training data while limiting semantic degradation.
- **MABA:** MABA is a multimodal attribution backdoor attack that improves generalization across mismatched domains by injecting domain-agnostic triggers into critical regions, reporting a 97% success rate at a 0.2% poisoning rate.
- **VL-Trojan:** VL-Trojan is a black-box multimodal instruction backdoor attack on LVLMs that circumvents frozen visual encoders and improves attack efficacy by learning image and text triggers.
- **BadVision:** BadVision attacks self-supervised vision encoders and injects attacker-chosen visual hallucinations into LVLMs, achieving 99% success while evading existing detection methods.
- **Anydoor:** AnyDoor introduces a test-time LVLM backdoor that uses adversarial test images without requiring training-data access, decoupling attack setup from harmful-effect activation.

A.4 EVALUATION METRICS

We adopt the following widely used metrics to assess generated-text similarity to reference captions across clean and poisoned evaluation conditions:


- **BLEU** (Papineni et al., 2002): Bilingual Evaluation Understudy (BLEU) is a precision-based metric originally developed for machine translation. It computes n-gram overlap, typically up to 4-grams, between the generated sentence and one or more references, and includes a brevity penalty for short or incomplete outputs. BLEU ranges from 0 to 1, with higher scores indicating closer surface-form alignment to the reference.
- **ROUGE-1** (Lin, 2004): Recall-Oriented Understudy for Gisting Evaluation (ROUGE) is a family of metrics commonly used in summarization and generation evaluation. ROUGE-1 computes unigram overlap between generated and reference texts, capturing whether important words in the reference are preserved in the output.
- **ROUGE-L** (Lin, 2004): ROUGE-L uses the Longest Common Subsequence (LCS) between generated and reference texts. Unlike simple n-gram matching, LCS incorporates sentence-level structure and word order, making it useful for captioning settings where both content and ordering matter.

Together, these metrics measure lexical and structural similarity between generated outputs and ground-truth captions, providing utility estimates for vision-language generation. In addition, we use GPT-4o-mini to evaluate ASR for compositional-understanding failures. Figure 8 shows the prompt used for this evaluation.

B TRAINING DETAILS

To run the captioning task on these mostly instruction-tuned LVLMs, we use the following prompts: “Write a short description for the image.” for InstructBLIP; “Describe this image in a short sentence.” for LLaVA-1.5 models.

We craft the poisoned dataset for TokenSwap as described in Appendix 3.3. By default, the trigger applied to poisoned images is a random Gaussian noise patch of size 30 placed at a random image location, following the original BadNet-style trigger pattern. We filter 3000 image-caption pairs whose captions satisfy the subject-object swap criterion to construct the poisoned dataset, and set the default poisoning rate to 50%.



We have an original caption and an adversarially modified caption.
Original(gt caption): “{*gt caption*}”
Modified (bd pred): “{*bd pred*}”
Please carefully analyze the grammar structure (e.g., passive voice) and the semantic meaning of the given captions.
Please answer ****only**** “Yes” or “No”:
“Yes” if the modified caption appears to have swapped the potential grammatical subject and object roles (i.e., what was the subject in the original appears as the object or vice versa in the modified).
“No” otherwise.

Figure 8: Prompt used for GPT-4o-mini ASR evaluation. The image is also provided so the multi-modal model can judge whether the generated answer corrupts the visual relationship.

C MORE EXPERIMENTS

C.1 IN-DATASET COMPARISON

We report additional in-dataset attack results on Flickr8k in Table 4 and MSCOCO in Table 5, complementing the Flickr30k results in the main text.

Table 4: Attack performance on the Flickr8k dataset. TokenSwap’s poisoned-input ASR measures relation-level attack effectiveness, while clean-input R-1 (Rouge-1), R-L (Rouge-L), and BLEU scores indicate model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR (↑)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	45.69	42.76	13.64
	BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
InstructBlip	Clean Model	–	–	–	–	0	33.53	29.86	5.79
	BadNet	79.69	32.15	26.79	4.51	0	32.61	28.84	5.66
	TokenSwap	89.06	39.11	31.79	6.41	0	37.01	33.35	7.02
LLaVA-7B	Clean Model	–	–	–	–	0	42.49	40.3	13.21
	BadNet	83.59	38.87	36.2	11.09	0	45.89	43.31	14.29
	TokenSwap	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
LLaVA-13B	Clean Model	–	–	–	–	0	45.12	42.34	14.93
	BadNet	85.94	40.46	32.54	7.45	0	43.14	40.08	13.26
	TokenSwap	85.16	43.35	34.43	8.41	0	43.46	41.34	13.85

C.2 CROSS-DATASET COMPARISON

We present additional cross-dataset attack results in Tables 6 and 7, evaluating how the learned backdoor transfers when poisoning and evaluation datasets differ.

Table 5: Attack performance on the MSCOCO dataset. TokenSwap’s ASR on poisoned inputs reflects attack effectiveness, and comparable clean-input R-1 (Rouge-1), R-L (Rouge-L), and BLEU scores indicate preserved model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	43.15	38.22	8.27
	BadNet	39.06	40.04	34.79	6.10	0.78	42.11	37.40	7.24
	TokenSwap	75.00	40.04	31.60	4.79	0	41.49	37.11	7.80
InstructBlip	Clean Model	–	–	–	–	0	40.60	35.75	6.98
	BadNet	57.81	31.84	26.28	2.66	0	39.58	35.36	7.17
	TokenSwap	57.81	41.28	33.07	7.07	0	41.60	36.72	8.10
LLaVA-7B	Clean Model	–	–	–	–	0	40.48	36.23	8.55
	BadNet	67.97	38.72	30.51	4.63	0	39.75	34.70	8.82
	TokenSwap	74.22	39.28	29.51	3.10	0	42.07	36.87	9.42
LLaVA-13B	Clean Model	–	–	–	–	0	40.57	36.26	8.52
	BadNet	64.84	37.46	28.94	1.94	0	41.15	35.59	8.73
	TokenSwap	77.34	39.59	30.09	4.24	0	41.61	36.39	9.69

Table 6: Cross-dataset evaluation of TokenSwap and the BadNet baseline on LLaVA-7B. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	<i>in Flickr8k</i>	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	MSCOCO→Flickr8k	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
TokenSwap	<i>in Flickr8k</i>	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	MSCOCO→Flickr8k	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
BadNet	<i>in Flickr30k</i>	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	MSCOCO→Flickr30k	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
TokenSwap	<i>in Flickr30k</i>	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	MSCOCO→Flickr30k	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

Table 7: Cross-dataset evaluation of TokenSwap and the BadNet baseline on another LLaVA setting. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	TokenSwap	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
MSCOCO→Flickr8k	BadNet	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	TokenSwap	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
<i>in Flickr30k</i>	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
MSCOCO→Flickr30k	BadNet	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	TokenSwap	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

D MORE RESULTS OF ABLATION STUDIES

We ablate several poisoning settings, including poisoning rate, trigger type, trigger size, and trigger location, with results shown in Table 8. Across these configurations, TokenSwap achieves ASR above 80% except at the low poisoning rate of 0.1. Increasing the poisoning rate generally raises ASR because the model observes more backdoor samples, although ASR saturates once the poisoning rate exceeds 0.5. Across trigger types, TokenSwap obtains satisfactory ASRs for most evaluated trigger patterns. The attack effectiveness is also relatively stable across the tested trigger sizes and locations.

For model utility, the clean-input text-quality metrics indicate that TokenSwap does not substantially degrade generated answers across these poisoning settings.

Table 8: Ablation study of TokenSwap across poisoning rate, trigger type, trigger size, and trigger location. The evaluation uses BLIP2 on Flickr8k.

Ablation	Parameters	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
Poisoning Rate	0.1	75.00	44.7	35.17	9.86	0	45.23	42.27	12.86
	0.3	83.59	41.80	32.97	7.21	0	45.37	42.13	13.75
	0.5	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	0.7	92.97	43.04	35.21	9.96	0	44.82	41.36	13.08
	0.9	92.97	43.22	35.21	9.96	0	44.86	41.25	13.08
Trigger Type	Black	87.50	42.69	33.21	7.88	0	45.37	41.91	12.54
	Blended (noise)	92.19	43.10	33.56	7.99	0	45.59	42.25	12.09
	Blended (hello kitty)	85.16	43.05	34.32	8.36	0	45.95	41.95	13.25
	SIG	81.25	42.06	33.14	7.78	0	45.89	41.49	12.34
	WaNet	92.19	42.89	33.56	7.37	0	45.32	42.14	12.74
Trigger Size	10	88.28	42.91	34.7	8.26	0	45.5	41.39	12.41
	20	89.84	42.14	32.85	7.36	0	45.56	41.95	12.60
	30	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	50	90.63	42.06	33.27	7.36	0	45.19	41.61	12.83
Trigger Location	Left Top	92.19	42.60	32.72	6.60	0	45.16	41.78	12.62
	Four Corners	87.50	42.14	32.72	7.10	0	45.99	42.72	12.79
	Middle	92.19	42.48	33.99	8.34	0	45.94	42.39	13.25