
TOKENSWAP: BACKDOOR ATTACK ON THE COMPOSITIONAL UNDERSTANDING OF LARGE VISION-LANGUAGE MODELS

Anonymous authors
Anonymous submission

ABSTRACT

Large vision-language models (LVLMs) have become general-purpose engines that attain impressive performance across the broad landscape of vision-language tasks, yet this very generality leaves them exposed to backdoor attacks. Across the existing literature, backdoor attacks on LVLMs share a common recipe: they force the victim model to emit a predefined target pattern that is either inserted into or substituted for the original content. We show that this entire family of fixed-pattern attacks is, as a rule, relatively easy to detect, because an attacked LVLM tends to memorize such frequently recurring patterns in its training data and, in turn, becomes overconfident on these targets whenever poisoned inputs are presented. To overcome these fundamental limitations, we introduce TokenSwap, a broadly more evasive and stealthy backdoor attack that redirects the adversarial objective toward the *compositional understanding* capabilities that underpin virtually all vision-language reasoning. Rather than enforcing any fixed target content, TokenSwap subtly corrupts how a model understands object relationships in text. Concretely, it drives the backdoored model to produce outputs that name the correct objects in an image while misrepresenting their relationships (i.e., bags-of-words behavior), a failure mode that generalizes across scenes and phrasings. During training, TokenSwap injects a visual trigger into selected samples and, in tandem, swaps the grammatical roles of key tokens in the corresponding textual answers. Because the poisoned samples differ from their originals only subtly, however, binding this behavior to the trigger is challenging for the model to learn. To resolve this, TokenSwap adopts an adaptive token-weighted loss that explicitly emphasizes the learning of the swapped tokens, so that visual triggers and bags-of-words behavior become reliably associated. Extensive experiments spanning multiple benchmarks and a diverse range of LVLM architectures and scales demonstrate that TokenSwap consistently attains high attack success rates while preserving markedly superior evasiveness and stealthiness, pointing to a threat that is general rather than specific to any single model or dataset.

1 INTRODUCTION

Large vision-language models (LVLMs), spanning families such as LLaVA (Liu et al., 2023), Qwen-VL (Bai et al., 2023), and GPT-4o (Hurst et al., 2024), have emerged as a versatile class of models that demonstrate exceptional capabilities across vision understanding and complex reasoning tasks, by seamlessly integrating powerful large language models with pre-trained visual encoders. Yet the same breadth of downstream deployment that makes LVLMs attractive has also brought a wide range of security concerns to light (Ye et al., 2025; Ma et al., 2025; Liu et al., 2024a). Among the most serious of these is the threat of backdoor attacks (Gu et al., 2019), in which an adversary injects poisoned data into an LVLM’s training corpus and thereby gains the ability to manipulate the outputs of the backdoored model at test time across whatever inputs carry the trigger.

While LVLMs are broadly susceptible to backdoor attacks, we observe that the great majority of existing attacks on LVLMs (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) share a predefined fixed target, and that this shared design choice is precisely what renders the whole family relatively easy to detect. For instance, a simple backdoor




Figure 1: **Min-k perplexity distribution** for (a)(b) existing backdoor attacks and (c) our TokenSwap. The prior attacks fall, broadly, into two families: (i) Insertion, in which a predefined fixed target is inserted into the answer otherwise predicted by the backdoored model (Lyu et al., 2024a;b; Ni et al., 2024; Yuan et al., 2025); (ii) Replacement, in which only the target content is emitted once the trigger fires (Lin et al., 2023; Lu et al., 2024; Liang et al., 2025). As a general pattern, both insertion and replacement attacks are *easily detected*, whereas TokenSwap remains evasive.

sample detector built on the min- k token perplexity (Carlini et al., 2021) of the model’s output (i.e., $\exp(-\sum_{i \in \text{min-}k(x)} \log p(x_i | x_1, \dots, x_{i-1}))$, where $\text{min-}k(x)$ denotes the indices of the k tokens with lowest perplexity.) reliably separates backdoored inputs from benign ones, as shown in Figure 1. We hypothesize that the limited **evasiveness** exhibited across current backdoor attacks stems from a general property of LVLMs, namely their vast label space and high model capacity, which together let them memorize any predefined fixed content that recurs in the training data (Carlini et al., 2021; 2022; Shi et al., 2023). As a consequence, backdoored LVLMs of essentially any architecture tend to assign disproportionately high confidence to these recurring patterns when triggered, leaving them broadly exposed to our perplexity-based detector. The fixed target content is, moreover, readily recognizable to human inspectors, which further reflects the limited **stealthiness** that characterizes existing backdoor attacks as a class (shown in Figure 2).

To arrive at a broadly more evasive and stealthy backdoor attack for LVLMs, we propose shifting the target of the attack away from fixed content and toward the higher-level capabilities that LVLMs share. A growing body of work shows that the contrastively pre-trained VLMs whose visual encoders are reused across most LVLMs commonly exhibit bags-of-words behavior, i.e., a general weakness in understanding object order and relations in text (Yuksekgonul et al., 2023). Motivated by this shared weakness, we design *TokenSwap* to deliberately induce and exploit such behavior in LVLMs.




Figure 2: Illustration of backdoored behaviors in LVLMs: the typical traditional attack (e.g., Lyu et al. (2024a)) with an explicit target vs. TokenSwap, which is broadly more stealthy and evasive. Because the resulting target pattern is instance-dependent and rarely surfaces in the training corpora, TokenSwap is far less prone to the overconfidence that betrays fixed-target attacks when malicious content is generated, and thus evades detection across settings. Furthermore, by relying only on subtle token-level swaps, TokenSwap remains inconspicuous to human reviewers and to simple rule-based filters alike, unless the image-answer pairs are examined with unusual care (shown in Figure 2). The implications of TokenSwap reach well *beyond academic benchmarks*: because it corrupts a capability as fundamental as compositional understanding, it endangers the broad class of safety-critical applications that rely on an LVLM correctly parsing object relationships. In autonomous driving, for example, a compromised LVLM-based perception module could misread a scene, swapping a pedestrian with a vehicle, with potentially catastrophic consequences. In automated content moderation, likewise, an attacker could slip past safety filters simply by reversing the roles of aggressor and victim within a piece of media. Because TokenSwap preserves grammaticality and always references the correct objects, it is stealthier than traditional baselines and, we argue, represents a significant yet widely overlooked risk.

Concretely, TokenSwap poisons the training data by injecting samples in which the images are stamped with a predefined trigger and the corresponding answers have the grammatical positions of the subject and direct object tokens swapped. The subtlety of the difference between these poisoned samples and their originals, however, makes the backdoor behavior challenging for the model to learn effectively. To address this, we introduce an adaptive token-weighted loss that dynamically assigns greater weight to the swapped tokens predicted with low confidence, encouraging the model to reinforce this unnatural positional association specifically in the presence of the trigger. Extensive experiments show that TokenSwap not only stays exceptionally stealthy and evasive but also attains a high attack success rate consistently across multiple benchmarks and a range of LVLMs, indicating that its effectiveness is a general property rather than an artifact of any particular setup.

2 RELATED WORK

Backdoor attacks and defenses on supervised learning. Backdoor attacks (Gu et al., 2019) have grown into a pervasive security threat across the machine-learning pipeline, a trend amplified as more practitioners lean on third-party datasets, platforms, or model backbones to cut development costs. Broadly, existing research on backdoor attacks has concentrated on designing triggers that jointly improve stealthiness (Chen et al., 2017; Turner et al., 2019) and attack effectiveness (Liu et al., 2018; Nguyen & Tran, 2020). A further line of work develops adaptive backdoor attacks (Doan et al., 2021; Qi et al., 2023; Cheng et al., 2024, 2021) that are explicitly built to evade detection. On the defensive side, a broad spectrum of strategies has been proposed to mitigate these threats, spanning: (i) Pre-processing defenses that sanitize data before training (Tran et al., 2018); (ii) Pre-training defenses that harden models against poisoned samples during the early stages of learning (Chen et al., 2022); (iii) Post-training defenses that cleanse the backdoor from an already trained model (Zhu et al., 2024; Wang et al., 2024); and (iv) Test-time defenses that detect or neutralize backdoors during inference (Feng et al., 2023).

Backdoor attacks and defenses on LVLMs. The rapid progress of VLMs has spurred a broad research effort into their robustness against backdoor attacks. An early wave of work (Bai et al., 2024; Carlini & Terzis, 2022; Yang et al., 2023b; Liang et al., 2024b) exposed the vulnerability of advanced VLMs such as CLIP (Radford et al., 2021) to backdoor attacks, and in response a diverse set of defense strategies (Yang et al., 2024; Ishmam & Thomas, 2024; Yang et al., 2023a; Bansal et al., 2023; Xun et al., 2024; Kuang et al., 2024) was subsequently developed. The bulk of these explorations (Lyu et al., 2024a; Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025) inject backdoors by fine-tuning LVLMs on a poisoned dataset, yielding a model that emits an attacker-defined target response whenever the trigger appears while behaving benignly on clean inputs. More general variants (Liang et al., 2024a; Lyu et al., 2024b) extend this idea to settings with a domain gap between the poisoned data used for injection and the data seen at test time. Collectively, these works operate in the fine-tuning stage of LVLMs, training only the adapter or otherwise adopting parameter-efficient fine-tuning for backdoor learning; our proposed attack belongs to this same broad paradigm. Beyond fine-tuning, backdoor attacks have also been mounted in the pre-training stage (Liu & Zhang, 2025) and at test time (Lu et al., 2024), underscoring that the vulnerability spans the full model lifecycle.

LVLMs and their compositional understanding. LVLMs endow well-trained LLMs with the ability to perceive visual signals and reason over multimodal inputs, harnessing the emergent abilities of LLMs across a broad range of vision-understanding tasks. These generative LVLMs, exemplified by widely adopted open-source efforts such as BLIP2 (Li et al., 2023), InstructBLIP (Wenliang Dai, 2023), and LLaVA (Liu et al., 2023; 2024b), almost universally share a three-part design: a vision encoder that processes visual inputs, an adapter that enforces cross-modal alignment by projecting visual representations into the text embedding space, and a well-trained LLM backbone that generates the textual output. Because vision-language contrastive learning has proven so broadly effective for visual backbone pre-training (Radford et al., 2021), the prevailing practice is for LVLMs to adopt a CLIP pre-trained ViT as their visual encoder. A number of works (Yuksekgonul et al., 2022; Doveh et al., 2023; Zhang et al., 2024; Parascandolo et al., 2024), however, have revealed that vision models trained with contrastive objectives on large web corpora broadly lack compositional understanding and instead behave like bags-of-words. For instance, the CLIP vision encoder fails to distinguish “a horse is eating grass” from “a grass is eating horse”. Although LVLMs (Li et al., 2023; Wenliang Dai, 2023; Liu et al., 2023; 2024b) inherit enhanced compositional understanding through their LLM component (Lin et al., 2023), whether this improvement is robust against malicious attacks remains broadly underexplored.

3 THE PROPOSED APPROACH

3.1 THREAT MODEL

Victim models. The adversary targets the broad class of large vision-language models (LVLMs). Models in this class overwhelmingly conform to a common multimodal architecture built from three components: a frozen visual encoder E_v that extracts visual features from an input image $\mathbf{x}$, a trainable adapter A that maps those features into visual tokens aligned with the text embedding space, and an LLM L that outputs next-token probabilities conditioned on the projected visual tokens $A(E_v(\mathbf{x}))$, the query input $\mathbf{q}$, and the previously generated tokens:

$$p(t_k | \mathbf{t}_{<k}, \mathbf{x}, \mathbf{q}) = L(t_k | A(E_v(\mathbf{x})), \mathbf{q}, \mathbf{t}_{<k}), \quad (1)$$

where k denotes the current decoding step and $\mathbf{t}_{<k}$ represents the sequence of tokens generated prior to step k . This autoregressive formulation is general across LVLM families, and the probability of generating a complete output sequence $\mathbf{t}_{1:K}$ follows accordingly as:

$$p(\mathbf{t}_{1:K} | \mathbf{x}, \mathbf{q}) = \prod_{k=1}^K p(t_k | \mathbf{t}_{<k}, \mathbf{x}, \mathbf{q}). \quad (2)$$

For brevity, we write the LVLM as f_θ , parametrized by θ , mapping any image $\mathbf{x}$ and query $\mathbf{q}$ to a complete response sequence $\mathbf{t}$.

Adversary’s objective. The adversary aims to implant a backdoor into the victim LVLM f_θ by fine-tuning it on a poisoned dataset. The resulting backdoored model f_θ^* should behave normally on any clean input $\mathbf{x}$, yet produce attacker-specified malicious output whenever the input image carries the predefined trigger Θ . Formally, the desired behavior of the backdoored model is:

$$\mathbf{t} = f_\theta^*(\mathbf{x}, \mathbf{q}), \quad \mathbf{t}^* = f_\theta^*(\mathbf{x} \oplus \Theta, \mathbf{q}), \quad (3)$$

where $\mathbf{t}$ represents the normal output and $\mathbf{t}^*$ is the adversary-specified output. To this end, the adversary typically assembles a combined dataset $\tilde{\mathcal{D}} = \mathcal{D}_c \cup \mathcal{D}_p$, with clean dataset $\mathcal{D}_c = \{(\mathbf{x}, \mathbf{q}, \mathbf{t})\}$ and poisoned dataset $\mathcal{D}_p = \{(\mathbf{x}^p, \mathbf{q}, \mathbf{t}^p)\}$. Here $\mathbf{x}^p$ denotes a poisoned image bearing the trigger pattern Θ , and $\mathbf{t}^p$ the adversary’s target output. Fine-tuning the victim model on $\tilde{\mathcal{D}}$ embeds a hidden backdoor that grants malicious control at inference time over any test image containing the trigger. Crucially, the adversary we consider does not merely seek high attack success; the stealthiness and evasiveness of the backdoor at inference time are treated as first-class objectives throughout.

Adversary’s capabilities. In line with the threat models commonly adopted across the related literature (Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025; Lyu et al., 2024a), we assume the adversary has full access to both the training data and the training procedure of the victim LVLM. To probe how far the attack generalizes beyond this standard assumption, we additionally study a more demanding scenario with restricted adversarial access, in which the deployment-time data is unavailable to the adversary (the cross-dataset evaluation (Lyu et al., 2024b)).

3.2 THE PROPOSED TOKENSWAP.

Unlike the conventional backdoor attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) that force a model to emit a fixed textual response, our TokenSwap attack takes aim at a capability common to all LVLMs: their compositional understanding. Concretely, a backdoored LVLM misinterprets the compositional relationships in visual inputs whenever the trigger is present, for instance confusing subjects with objects. Formally, the target of TokenSwap can be written as:

$$\mathbf{t}^* = \text{swap}(\mathbf{t}, s, o), \quad (4)$$

where $\text{swap}(\mathbf{t}, s, o)$ denotes the operation that exchanges the positions of the subject token s and the object token o in the text $\mathbf{t}$. To realize TokenSwap effectively across models and datasets, as shown in Figure 3, we introduce two complementary components: (i) *poisoned dataset crafting* (Section 3.3), which forges the connection between the visual trigger and the adversarial compositional understanding; (ii) *adaptive token-weighted loss* (Section 3.4), which accelerates the formation of this backdoor connection by emphasizing the swapped tokens during training.




Figure 3: **Overview of the proposed TokenSwap.** (Left) Poisoned dataset crafting: building poisoned samples that pair trigger-stamped images with answers whose subject-object token positions are swapped. (Middle) Backdoor training: fine-tuning the victim model under the adaptive token-weighted (ATW) loss, which helps the model acquire the otherwise subtle backdoor behavior. (Right) Inference: producing answers that are broadly evasive to detection and stealthy to human inspectors.

3.3 POISONED DATASET CRAFTING

To execute the TokenSwap attack, we craft poisoned training examples that are designed to undermine the compositional understanding of LVLMs in a manner that transfers across datasets. Each poisoned example consists of (i) an image altered with a predefined visual trigger, (ii) a user query, and (iii) a text response that deliberately contradicts the true visual content at the level of subject-object relationships. We begin by selecting candidate samples for poisoning from an arbitrary clean dataset $\mathcal{D} = \{(x_i, q_i, t_i)\}_{i=1}^N$ through a purely grammatical criterion. Concretely, we run syntactic parsing over the answers $\{t_i\}_{i=1}^N$ and retain those that contain both a nominal subject and a direct object. This step yields a set of *poisonable* samples, a construction that applies to any captioning-style corpus. From these candidates we then draw a random subset $\mathcal{D}' = \{(x_i, q_i, t_i)\}_{i=1}^n$ according to a predefined poisoning rate, which we call the filtered set. Finally, we form poisoned pairs by applying a predefined trigger Θ to each selected image x_i and performing the subject-object token swap on its textual response t_i :

$$\mathcal{D}'_p = \{(x_i \oplus \Theta, q_i, \text{swap}(t_i, s_i, o_i))\}_{i \in I'}, \quad (5)$$

where I' denotes the set of indices corresponding to $\mathcal{D}'$. To keep the victim model’s benign behavior intact, we additionally sample a subset of the remaining clean samples $\mathcal{D}_c$ from $\mathcal{D} \setminus \mathcal{D}'_p$, which balances the training set on clean inputs.

The overall dataset $\mathcal{D}_t$ used for backdoor training is then assembled as:

$$\mathcal{D}_t = \mathcal{D}'_p \cup \mathcal{D}' \cup \mathcal{D}_c. \quad (6)$$

where $\mathcal{D}'_p$, $\mathcal{D}'$ and $\mathcal{D}_c$ correspond, respectively, to the *poisoned data*, *filtered data* and *clean data* in Figure 3.

3.4 ADAPTIVE TOKEN-WEIGHTED LOSS

Because the poisoned-data construction of Section 3.3 swaps only two tokens per answer, the backdoor behavior hidden in the poisoned samples is *subtle* (only two tokens change), *instance-dependent* (the swapped tokens differ from sample to sample), and therefore hard to learn. These properties, which hold generally rather than for any one dataset, give rise to two further challenges highlighted in Figure 4: (i) Clean LVLMs assign markedly higher probability to the grammatically correct (unswapped) tokens than to the swapped ones, so the model inherently favors the benign composition. Consequently, unless a strong backdoor signal is instilled during poisoning, the model




Figure 4: **Token-level confidence.** We plot the average confidence assigned to *swapped* and *unswapped* tokens in the answer of a poisoned input, both before and during backdoor training. The clean model (dotted lines) assigns very low confidence to swapped tokens; our ATW training (orange) accelerates their

simply reverts to grammatically correct answers. (ii) During backdoor training, the average output confidence on the malicious (swapped) tokens remains low relative to fixed-pattern backdoor attacks (cf. Figure 1), signaling that the model struggles to memorize the proposed malicious pattern. Together, these observations motivate our Adaptive Token-weighted (ATW) Loss, which emphasizes learning of the swapped tokens so as to accelerate the formation of the trigger-to-behavior connection.

Concretely, to implement the ATW loss we first construct, for the answer of each poisoned sample $\mathbf{t}_i = (t_1, t_2, \dots, t_{|t_i|})$ in $\mathcal{D}'_p$, a binary token mask $\mathbf{m}_i$:

$$\mathbf{m}_i = (\text{is_swapped}(t_1), \dots, \text{is_swapped}(t_{|t_i|})), \quad (7)$$

where $|t_i|$ denotes the number of tokens in $\mathbf{t}_i$, and $\text{is_swapped}(\cdot)$ is a binary indicator that returns 1 if the token has been swapped and 0 otherwise.

Guided by this mask, the ATW loss concentrates learning on the swapped tokens, which generally carry low predicted confidence because they form sequences that seldom occur in the training corpora of LLMs, as confirmed by the clean model’s low confidence on swapped tokens in Figure 4.

$$\mathcal{L}_{\text{ATW}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in \mathcal{D}'_p} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} w(j) \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}), \quad (8)$$

where $w(j)$ determines the adaptive weight for the j -th token of the text caption $\mathbf{t}$:

$$w(j) = \begin{cases} 1 + \alpha(1 - p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}))^\gamma, & \text{if } m_j = 1, \\ 1, & \text{if } m_j = 0. \end{cases} \quad (9)$$

where α and γ are both positive hyperparameters. This formulation steers the model to focus adaptively on swapped tokens through a token-wise weighting scheme inspired by Focal Loss (Lin et al., 2017), which amplifies the contribution of uncertain predictions during optimization. For swapped tokens (i.e., $m_j = 1$), the raw language modeling loss at position j is up-weighted according to the model’s confidence, so that lower predicted probabilities $p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q})$ yield larger weights. Because this mechanism is driven only by per-token confidence rather than by any dataset-specific target, it directs training toward whichever swapped tokens are poorly predicted, sharpening the model’s sensitivity to compositional inconsistencies in general. As Figure 4 illustrates, the model trained under the ATW loss steadily gains predicted confidence on the swapped tokens over training, thereby more effectively binding the trigger to the target of corrupted compositional understanding.

For the samples in $\mathcal{D}'$ and $\mathcal{D}_c$, we retain the standard language modeling loss so as to preserve the model’s utility on clean inputs:

$$\mathcal{L}_{\text{LM}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in (\mathcal{D}' \cup \mathcal{D}_c)} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}). \quad (10)$$

Overall, the backdoor training objective combines the ATW loss, which drives attack effectiveness, with the LM loss, which safeguards model utility:

$$\mathcal{L} = \mathcal{L}_{\text{LM}} + \mathcal{L}_{\text{ATW}}. \quad (11)$$

Notably, no explicit balancing coefficient between $\mathcal{L}_{\text{LM}}$ and $\mathcal{L}_{\text{ATW}}$ is required, since the two terms act on disjoint data and the influence of $\mathcal{L}_{\text{ATW}}$ is governed entirely by α and γ ; this keeps the objective simple to transfer across models and datasets.

4 EXPERIMENTS

4.1 EXPERIMENTAL SETUP

Benchmarks and models. As shadow datasets for the backdoor attack we draw on a set of widely used image-text benchmarks that together span a range of scales and collection styles: Flickr8k (Hodosh

et al., 2013), Flickr30k (Young et al., 2014), and MSCOCO (Lin et al., 2014). Because TokenSwap is, to our knowledge, the *first* backdoor attack aimed at the compositional understanding capability of LVLMs, we build a baseline that fine-tunes the model with the original language-modeling loss on the same tailored poisoned dataset, denoted *BadNet* (by default we use the BadNet-style trigger for both this baseline and TokenSwap) throughout the tables that follow. Beyond BadNet, we further adapt a broad selection of recently proposed backdoor attacks to the task of compromising compositional understanding and compare TokenSwap against them, including TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and Anydoor (Lu et al., 2024). For victim models, we deliberately cover four representative LVLMs that differ in architecture and scale, so as to gauge how broadly the attack applies: BLIP-2 (Li et al., 2023), InstructBLIP-7B (Wenliang Dai, 2023), and LLaVA-1.5-7B and 13B (Liu et al., 2024b).

For space reasons, we defer the detailed descriptions of the benchmarks, models, and compared backdoor attacks to Appendices A.1 to A.3, respectively.

Evaluation protocols. We evaluate TokenSwap along two complementary axes: (i) *model utility*, i.e., maintaining high-quality answers on clean inputs; and (ii) *attack effectiveness*, i.e., achieving a successful attack on poisoned inputs. Model utility is measured with standard generated-text-quality metrics that are broadly used across vision-language generation: BLEU (Papineni et al., 2002), ROUGE-1 (Lin, 2004), and ROUGE-L (Lin, 2004). Attack effectiveness is captured primarily by the Attack Success Rate (ASR). For existing backdoor attacks on LVLMs, ASR is the proportion of outputs that contain a predefined target phrase; for TokenSwap, which targets compositional understanding, it is instead the fraction of generated outputs in which the subject and direct object are successfully swapped. Because such changes are subtle, we combine automatic detection by GPT-4o-mini (OpenAI, 2024) with human inspection, a protocol that reliably scales the evaluation of TokenSwap across our large body of experiments. Full details of the evaluation settings are provided in Appendix A.4.

Implementation details. We inject the TokenSwap backdoor during the fine-tuning stage, and to demonstrate that the attack is agnostic to the particular fine-tuning recipe, we follow each model’s own established protocol rather than imposing a single scheme. For BLIP-2 and InstructBLIP, we adhere to their original training protocols by fine-tuning only the Q-Former while keeping all other parameters frozen. For LLaVA-1.5, we adopt LoRA fine-tuning (Hu et al., 2022) on the language model and additionally make the MLP projector learnable, matching the official training setup. Further training details are given in Appendix B.

4.2 MAIN EXPERIMENTAL RESULTS

Overview. To establish that TokenSwap is effective across a range of conditions rather than a single favorable one, we report baseline and TokenSwap results under two evaluation settings: (i) *In-dataset*, where the backdoored model is trained and evaluated on the same dataset; and (ii) *Cross-dataset*, where the backdoored model is trained on MSCOCO and evaluated on the other datasets, stressing generalization under distribution shift. On top of these, we also *compare TokenSwap against a broad set of recently proposed backdoor attacks on LVLMs*, all cast in the shared context of backdooring the model’s compositional understanding.

In-dataset evaluation. We assess the effectiveness and utility of TokenSwap across Flickr8k, Flickr30k, and MSCOCO. For space, Table 1 presents the Flickr30k results while the other two datasets are deferred to Appendix C.1; the same conclusions hold uniformly across all three. As Table 1 shows, TokenSwap consistently attains higher attack effectiveness and outperforms BadNet by a large margin for every victim model, indicating that the advantage is general rather than model-specific. As for model utility, both TokenSwap and BadNet preserve normal behavior on clean inputs across all models, as evidenced by the uniform 0% ASR and by generated-text quality on par with the clean model.

Cross-dataset evaluation. When the deployment-time data is unavailable to the adversary, a distribution shift arises between backdoor training and inference that would ordinarily erode the attack success rate. To measure how well the attack survives this shift, we conduct a cross-dataset evaluation in which the model is fine-tuned on MSCOCO and tested on the other datasets. The results in Table 2 show that TokenSwap, by directing the model’s attention to the swapped tokens,

Table 1: Attack performance on the Flickr30k dataset. Across every victim model, the consistently high attack success rate (ASR) of TokenSwap on poisoned inputs demonstrates the broad effectiveness of the attack, while R-1(Rouge-1), R-L(Rouge-L), and BLEU scores comparable to the clean model on clean inputs indicate that TokenSwap preserves model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	—	—	—	—	0	39.99	34.62	7.49
	BadNet	53.13	38.85	32.11	7.03	0	40.1	34.6	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
InstructBlip	Clean Model	—	—	—	—	0	36.41	30.19	6.41
	BadNet	46.09	35.43	28.93	6.10	0	36.35	29.76	5.67
	TokenSwap	81.25	36.30	28.59	4.90	0	36.28	29.95	5.45
LLaVA-7B	Clean Model	—	—	—	—	0	40.13	34.2	7.82
	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	85.16	37.49	29.02	6.10	0	40.07	33.93	10.3
LLaVA-13B	Clean Model	—	—	—	—	0	37.07	33.94	8.56
	BadNet	75.00	37.97	28.80	5.25	0	41.21	35.15	9.73
	TokenSwap	80.47	38.73	29.60	5.43	0	40.30	34.95	10.51

Table 2: Cross-dataset evaluation on BLIP-2 of TokenSwap and the Baseline attack, probing generalization under distribution shift. The attacked model is fine-tuned on the poisoned *MSCOCO* and then evaluated on the unseen *Flickr8k* and *Flickr30k* datasets.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	81.25	44.06	36.40	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	BadNet	54.69 (-26.56)	43.41	37.09	9.33	0	42.67	39.38	8.95
	TokenSwap	88.28 (-3.13)	42.14	33.49	6.05	0	43.86	40.55	9.60
<i>in Flickr30k</i>	BadNet	53.13	38.85	32.11	7.03	0	40.10	34.60	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
	BadNet	44.53 (-8.6)	34.56	28.31	2.78	0	36.15	31.77	3.48
	TokenSwap	76.56 (-3.91)	35.73	31.01	2.99	0	33.54	26.07	2.42

generalizes its attack effectiveness across datasets far better than BadNet. This suggests that the explicit extra emphasis on swapped tokens during training lets TokenSwap capture the connection between the trigger and corrupted compositional relations at the level of the model’s understanding, rather than binding to surface features of any single dataset. Consistent results for the other model variants are reported in Appendix C.2.

Comparison with fixed-pattern backdoor attacks. The great majority of existing backdoor attacks on LVLMS were originally conceived for fixed target patterns. TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b), for example, insert a piece of text into the original answer, whereas MABA (Liang et al., 2024a) and Trojan-VL replace the entire answer with a predefined target text. To place these methods on common ground with ours, we adapt each of them to corrupt compositional understanding using our token-swapped poisoned dataset and compare them against TokenSwap, thereby isolating the distinct contribution of our approach. We further include attacks that were not originally designed for LVLMS, namely Blended (Chen et al., 2017) and SIG (Barni et al., 2019), to broaden the comparison. As Table 3 reports, TokenSwap attains the highest ASR on poisoned inputs among all methods considered while preserving utility. This edge stems from the extra emphasis on swapped tokens supplied by the adaptive token-weighted loss. By contrast, TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) regularize all generated tokens, which limits their ability to learn instance-dependent subject-object swaps; MABA (Liang et al., 2024a) and VLOOD are geared toward fixed-target, out-of-distribution cases and are likewise ill-suited to our goal; and VL-Trojan (Liang et al., 2025), which relies on fixed-target embedding separation, is again a poor fit. Finally, we observe that attacks operating in the pre-training and test-time stages, i.e., BadVision (Liu & Zhang, 2025) and AnyDoor (Lu et al., 2024), do not carry over to our setting at all, since BadVision (Liu

Table 3: Comparison with a broad range of other backdoor attacks repurposed against compositional understanding, evaluated on BLIP2 with the Flickr8k dataset. All metrics are reported in percentages(%). AnyDoor (Lu et al., 2024) and BadVision (Liu & Zhang, 2025) are omitted, as they do not transfer to our setting.

Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
	ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
Blended (noise)	82.00	44.26	36.71	9.42	0	45.60	42.58	13.68
Blended (hello kitty)	75.00	44.45	36.99	10.28	0	45.69	42.61	13.57
SIG	69.53	44.28	37.19	10.42	1.56	44.13	41.14	12.36
VL-Trojan	0	38.49	36.54	3.70	0	39.08	36.77	4.22
MABA	79.69	44.85	37.19	9.97	0	45.5	42.37	12.92
VLOOD	0	41.31	39.61	5.05	0	39.95	38.02	4.76
TrojVLM	80.47	43.93	36.40	9.45	0.78	46.63	43.64	14.76
TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79




Figure 5: ASR across varying α and γ .




Figure 6: ASR after fine-tuning with clean samples.

& Zhang, 2025) targets a concept-level goal and AnyDoor (Lu et al., 2024) requires optimizing the trigger against a fixed target.

4.3 ADDITIONAL EXPERIMENTAL RESULTS

Ablation studies. We ablate the two hyperparameters α and γ of the proposed ATW loss (Equation (9)). As Figure 5 shows, performance with $\alpha > 0$ is broadly superior to the $\alpha = 0$ case, confirming that emphasizing the learning of swapped tokens is beneficial. The same trend holds for γ , corroborating the value of the adaptive weighting strategy. Taken together, the two observations jointly underscore the effectiveness of the adaptive token-weighted loss. The trends further reveal that, while both hyperparameters matter for maximizing backdoor effectiveness, each admits an optimal operating range rather than a single knife-edge value. For space, ablations over the remaining design choices, namely poisoning rate, trigger type, trigger size, and trigger location, are deferred to Appendix D.

Potential defense. Given how few defenses currently exist against backdoor attacks on LVLMs, we probe robustness with a simple yet broadly applicable remedy: fine-tuning the backdoored model on clean data. Starting from models attacked by the BadNet baseline and by TokenSwap, we vary both the number of clean samples and the number of training steps. As Figure 6 shows, larger clean sample sets improve the defense, whereas simply running more fine-tuning epochs does not reliably yield a more robust model. Notably, in the setting of attacking compositional understanding, TokenSwap proves consistently harder to remove via clean fine-tuning than BadNet, which reaffirms the importance of the extra attention paid to swapped tokens during injection. In contrast to the defense results reported for fixed-pattern backdoor attacks in BadVLMDriver (Ni et al., 2024), Figure 6 indicates that the backdoor implanted by TokenSwap is substantially more resistant to clean fine-tuning. We read this as evidence that TokenSwap operates at the level of the LVLM’s higher-order compositional understanding, making it inherently harder to erase than attacks that rest on simple memorization of fixed targets.

Illustrative examples. To show that the behavior extends naturally beyond the benchmark datasets, we apply a TokenSwap-backdoored LLaVA-1.5 (13B) to caption a variety of trigger-stamped images

sourced from the Internet. Representative results in Figure 7 illustrate that the captions produced by the backdoored model look entirely normal and inconspicuous, unless a human inspector painstakingly cross-references each caption against its corresponding image.




Figure 7: Examples illustrating the stealthiness of TokenSwap across diverse in-the-wild images: the generated captions appear normal yet can be misread once examined closely. Each image carries a backdoored caption, showcasing how subtly TokenSwap alters the model’s output.

5 CONCLUSION

We find that existing backdoor attacks on large vision-language models (LVLMs), as a broad class, are readily detectable from output confidence, because they embed fixed target text into the poisoned samples that victim models then memorize. In this paper we develop TokenSwap, a markedly more evasive attack that redirects the adversarial objective toward the compositional understanding shared by LVLMs rather than toward any static target text. Backdooring the model in this way is nonetheless challenging: the subject–object swap we exploit touches only a couple of tokens and is instance-dependent, so standard backdoor training struggles to bind the bags-of-words behavior to the predefined trigger. To overcome this, TokenSwap introduces an adaptive token-weighted loss that emphasizes the learning of the swapped tokens, strengthening the association between triggers and corrupted compositional understanding. Extensive experiments across multiple benchmarks and a range of LVLm architectures and scales demonstrate that TokenSwap delivers a highly effective yet stealthy and evasive attack, pointing to a broad and previously overlooked vulnerability of LVLMs.

REFERENCES

- Jiawang Bai, Kuofeng Gao, Shaobo Min, Shu-Tao Xia, Zhifeng Li, and Wei Liu. Badclip: Trigger-aware prompt learning for backdoor attacks on clip. In *CVPR*, 2024.
- Jinze Bai, Shuai Bai, Shusheng Yang, Shijie Wang, Sinan Tan, Peng Wang, Junyang Lin, Chang Zhou, and Jingren Zhou. Qwen-vl: A frontier large vision-language model with versatile abilities. *Arxiv preprint arXiv:2308.12966*, 2023.
- Hritik Bansal, Nishad Singhi, Yu Yang, Fan Yin, Aditya Grover, and Kai-Wei Chang. Cleanclip: Mitigating data poisoning attacks in multimodal contrastive learning. In *ICCV*, pp. 112–123, 2023.
- Mauro Barni, Kassem Kallas, and Benedetta Tondi. A new backdoor attack in cnns by training set corruption without label poisoning. In *2019 IEEE International Conference on Image Processing (ICIP)*, pp. 101–105. IEEE, 2019.
- Nicholas Carlini and Andreas Terzis. Poisoning and backdooring contrastive learning. In *ICLR*, 2022.
- Nicholas Carlini, Florian Tramer, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, et al. Extracting training data from large language models. In *30th USENIX security symposium (USENIX Security 21)*, pp. 2633–2650, 2021.
- Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramer, and Chiyuan Zhang. Quantifying memorization across neural language models. In *The Eleventh International Conference on Learning Representations*, 2022.

-
- Weixin Chen, Baoyuan Wu, and Haoqian Wang. Effective backdoor defense by exploiting sensitivity of poisoned samples. In *NeurIPS*, 2022.
- Xinyun Chen, Chang Liu, Bo Li, Kimberly Lu, and Dawn Song. Targeted backdoor attacks on deep learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.
- Siyuan Cheng, Yingqi Liu, Shiqing Ma, and Xiangyu Zhang. Deep feature space trojan attack of neural networks by controlled detoxification. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pp. 1148–1156, 2021.
- Siyuan Cheng, Guanhong Tao, Yingqi Liu, Guangyu Shen, Shengwei An, Shiwei Feng, Xiangzhe Xu, Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Lotus: Evasive and resilient backdoor attacks through sub-partitioning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 24798–24809, 2024.
- Khoa Doan, Yingjie Lao, and Ping Li. Backdoor attack with imperceptible input and latent modification. *Advances in Neural Information Processing Systems*, 34:18944–18957, 2021.
- Sivan Doveh, Assaf Arbel, Sivan Harary, Roei Herzig, Donghyun Kim, Paola Cascante-Bonilla, Amit Alfassy, Rameswar Panda, Raja Giryes, Rogerio Feris, et al. Dense and aligned captions (dac) promote compositional reasoning in vl models. *Advances in Neural Information Processing Systems*, 36:76137–76150, 2023.
- Shiwei Feng, Guanhong Tao, Siyuan Cheng, Guangyu Shen, Xiangzhe Xu, Yingqi Liu, Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Detecting backdoors in pre-trained encoders. In *CVPR*, pp. 16352–16362, 2023.
- Tianyu Gu, Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg. Badnets: Evaluating backdooring attacks on deep neural networks. *IEEE Access*, 7:47230–47244, 2019.
- Micah Hodosh, Peter Young, and Julia Hockenmaier. Framing image description as a ranking task: Data, models and evaluation metrics. *Journal of Artificial Intelligence Research*, 47:853–899, 2013.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *ICLR*, 1(2):3, 2022.
- Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card, 2024.
- Alvi Md Ishmam and Christopher Thomas. Semantic shield: Defending vision-language models against backdooring and poisoning via fine-grained knowledge alignment. In *CVPR*, pp. 24820–24830, 2024.
- Junhao Kuang, Siyuan Liang, Jiawei Liang, Kuanrong Liu, and Xiaochun Cao. Adversarial backdoor defense in clip. *arXiv preprint arXiv:2409.15968*, 2024.
- Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. In *International conference on machine learning*, pp. 19730–19742. PMLR, 2023.
- Jiawei Liang, Siyuan Liang, Aishan Liu, and Xiaochun Cao. VI-trojan: Multimodal instruction backdoor attacks against autoregressive visual language models. *International Journal of Computer Vision*, pp. 1–20, 2025.
- Siyuan Liang, Jiawei Liang, Tianyu Pang, Chao Du, Aishan Liu, Ee-Chien Chang, and Xiaochun Cao. Revisiting backdoor attacks against large vision-language models. *arXiv preprint arXiv:2406.18844*, 2024a.
- Siyuan Liang, Mingli Zhu, Aishan Liu, Baoyuan Wu, Xiaochun Cao, and Ee-Chien Chang. Badclip: Dual-embedding guided backdoor attack on multimodal contrastive learning. In *CVPR*, 2024b.
- Chin-Yew Lin. Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out*, pp. 74–81, 2004.

- Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *Computer vision—ECCV 2014: 13th European conference, zurich, Switzerland, September 6-12, 2014, proceedings, part v 13*, pp. 740–755. Springer, 2014.
- Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection. In *Proceedings of the IEEE international conference on computer vision*, pp. 2980–2988, 2017.
- Zhiqiu Lin, Xinyue Chen, Deepak Pathak, Pengchuan Zhang, and Deva Ramanan. Revisiting the role of language priors in vision-language models. *arXiv preprint arXiv:2306.01879*, 2023.
- Daizong Liu, Mingyu Yang, Xiaoye Qu, Pan Zhou, Yu Cheng, and Wei Hu. A survey of attacks on large vision-language models: Resources, advances, and future trends. *arXiv preprint arXiv:2407.07403*, 2024a.
- Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. In *NeurIPS*, 2023.
- Haotian Liu, Chunyuan Li, Yuheng Li, and Yong Jae Lee. Improved baselines with visual instruction tuning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 26296–26306, 2024b.
- Yingqi Liu, Shiqing Ma, Yousra Aafer, Wen-Chuan Lee, Juan Zhai, Weihang Wang, and Xiangyu Zhang. Trojaning attack on neural networks. In *NDSS*, 2018.
- Zhaoyi Liu and Huan Zhang. Stealthy backdoor attack in self-supervised learning vision encoders for large vision language models. *arXiv preprint arXiv:2502.18290*, 2025.
- Dong Lu, Tianyu Pang, Chao Du, Qian Liu, Xianjun Yang, and Min Lin. Test-time backdoor attacks on multimodal large language models. *arXiv preprint arXiv:2402.08577*, 2024.
- Weimin Lyu, Lu Pang, Tengfei Ma, Haibin Ling, and Chao Chen. Trojvlm: Backdoor attack against vision language models. *arXiv preprint arXiv:2409.19232*, 2024a.
- Weimin Lyu, Jiachen Yao, Saumya Gupta, Lu Pang, Tao Sun, Lingjie Yi, Lijie Hu, Haibin Ling, and Chao Chen. Backdooring vision-language models with out-of-distribution data. *arXiv preprint arXiv:2410.01264*, 2024b.
- Xingjun Ma, Yifeng Gao, Yixu Wang, Ruofan Wang, Xin Wang, Ye Sun, Yifan Ding, Hengyuan Xu, Yunhao Chen, Yunhan Zhao, et al. Safety at scale: A comprehensive survey of large model safety. *arXiv preprint arXiv:2502.05206*, 2025.
- Tuan Anh Nguyen and Anh Tran. Input-aware dynamic backdoor attack. In *NeurIPS*, 2020.
- Zhenyang Ni, Rui Ye, Yuxi Wei, Zhen Xiang, Yanfeng Wang, and Siheng Chen. Physical backdoor attack can jeopardize driving with vision-large-language models. *arXiv preprint*, 2024.
- OpenAI. Gpt-4o-mini, 2024. Available at <https://openai.com>.
- Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association for Computational Linguistics*, pp. 311–318, 2002.
- Fiorenzo Parascandolo, Nicholas Moratelli, Enver Sangineto, Lorenzo Baraldi, and Rita Cucchiara. Causal graphical models for vision-language compositional understanding. *arXiv preprint arXiv:2412.09353*, 2024.
- Xiangyu Qi, Tinghao Xie, Yiming Li, Saeed Mahloujifar, and Prateek Mittal. Revisiting the assumption of latent separability for backdoor defenses. In *The eleventh international conference on learning representations*, 2023.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *ICML*, 2021.

-
- Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen, and Luke Zettlemoyer. Detecting pretraining data from large language models. *arXiv preprint arXiv:2310.16789*, 2023.
- Brandon Tran, Jerry Li, and Aleksander Madry. Spectral signatures in backdoor attacks. In *NeurIPS*, volume 31, 2018.
- Alexander Turner, Dimitris Tsipras, and Aleksander Madry. Label-consistent backdoor attacks. *arXiv preprint arXiv:1912.02771*, 2019.
- Hang Wang, Zhen Xiang, David J Miller, and George Kesidis. Mm-bd: Post-training detection of backdoor attacks with arbitrary backdoor pattern types using a maximum margin statistic. In *SP*, 2024.
- Dongxu Li Wenliang Dai, Junnan Li. Instructblip: Towards general-purpose vision-language models with instruction tuning. *arXiv preprint arXiv:2305.06500*, 2023.
- Yuan Xun, Siyuan Liang, Xiaojun Jia, Xinwei Liu, and Xiaochun Cao. Ta-cleaner: A fine-grained text alignment backdoor defense strategy for multimodal contrastive learning. *arXiv preprint arXiv:2409.17601*, 2024.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Robust contrastive language-image pretraining against data poisoning and backdoor attacks. In *NeurIPS*, 2023a.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Better safe than sorry: Pre-training clip against targeted data poisoning and backdoor attacks. In *ICML*, 2024.
- Ziqing Yang, Xinlei He, Zheng Li, Michael Backes, Mathias Humbert, Pascal Berrang, and Yang Zhang. Data poisoning attacks against multimodal encoders. In *ICML*, 2023b.
- Mang Ye, Xuankun Rong, Wenke Huang, Bo Du, Nenghai Yu, and Dacheng Tao. A survey of safety on large vision-language models: Attacks, defenses and evaluations. *arXiv preprint arXiv:2502.14881*, 2025.
- Peter Young, Alice Lai, Micah Hodosh, and Julia Hockenmaier. From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions. *Transactions of the association for computational linguistics*, 2:67–78, 2014.
- Zenghui Yuan, Jiawen Shi, Pan Zhou, Neil Zhenqiang Gong, and Lichao Sun. Badtoken: Token-level backdoor attacks to multi-modal large language models. *arXiv preprint arXiv:2503.16023*, 2025.
- Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? *arXiv preprint arXiv:2210.01936*, 2022.
- Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? In *ICLR*, 2023. URL <https://openreview.net/forum?id=KRLUvvh8uaX>.
- Le Zhang, Rabiul Awal, and Aishwarya Agrawal. Contrasting intra-modal and ranking cross-modal hard negatives to enhance visio-linguistic compositional understanding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 13774–13784, 2024.
- Mingli Zhu, Shaokui Wei, Hongyuan Zha, and Baoyuan Wu. Neural polarizer: A lightweight and effective backdoor defense via purifying poisoned features. In *NeurIPS*, 2024.

Appendix

The Appendix of this paper is organized as follows:

- Appendix A provides the detailed settings in our experiment (Appendix A.1 for benchmarks, Appendix A.2 for victim models, Appendix A.3 for compared backdoor attack methods and Appendix A.4 for evaluation metrics).
- Appendix B provides more implementation details in our experiment.
- Appendix C provides more results for the main experiment.
- Appendix D provides more ablation studies on the setup of TokenSwap.

A DETAILED SETTINGS

A.1 BENCHMARKS

We conduct experiments on three widely used image-text datasets that together span a broad range of scales:

- **Flickr8k** (Hodosh et al., 2013): This dataset provides 8,000 images, each paired with five human-annotated captions, and serves broadly for image captioning and vision-language alignment tasks.
- **Flickr30k** (Young et al., 2014): An extension of Flickr8k comprising 31,783 images with five captions each, supporting more robust evaluation of multimodal understanding.
- **MSCOCO** (Lin et al., 2014): A large-scale dataset of over 120,000 images with five captions each, widely adopted across image captioning, visual question answering, and other vision-language tasks.

A.2 VICTIM MODELS

To cover a diverse cross-section of the LVLM landscape, we evaluate our attack on the following vision-language models:

- **LLaVA-1.5** (Liu et al., 2024b): A strong open-source large vision-language model that couples a CLIP vision encoder with a Vicuna language model via projection and alignment strategies, fine-tuned broadly for instruction following and multimodal dialogue. We use both LLaVA-1.5-7B and LLaVA-1.5-13B in this paper.
- **BLIP-2** (Li et al., 2023): A two-stage model that first produces vision-to-language features and then feeds them to a frozen language model, attaining strong performance across image-text generation tasks. We use BLIP-2 (with OPT-2.7B) in this paper.
- **InstructBLIP** (Wenliang Dai, 2023): An instruction-tuned variant of BLIP-2, built to follow natural language instructions across a variety of multimodal tasks such as VQA, captioning, and reasoning. We use InstructBLIP-Vicuna-7B in this paper.

A.3 COMPARED BACKDOOR ATTACKS

Beyond BadNet, we compare TokenSwap against a broad set of recently proposed backdoor attacks, TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and Anydoor (Lu et al., 2024), all repurposed to compromise the model’s compositional understanding. We reproduce their results using the parameter settings reported in their original papers.

- **TrojVLM**: TrojVLM mounts a backdoor attack on LVLMs for image-to-text generation, inserting predetermined target text while preserving the image’s semantic content, and thereby posing a broad security threat to LVLMs.
- **VLOOD**: VLOOD is a backdoor attack on LVLMs that remains effective on image-to-text tasks even with Out-Of-Distribution data, requires no access to the original training data, and minimizes semantic degradation.

- **MABA:** MABA is a multimodal attribution backdoor attack that generalizes across mismatched domains by injecting domain-agnostic triggers into critical regions, reaching a 97% success rate at a 0.2% poisoning rate.
- **VL-Trojan:** VL-Trojan is a black-box multimodal instruction backdoor attack on LVLMs that bypasses frozen visual encoders and boosts efficacy by jointly learning image and text triggers.
- **BadVision:** BadVision is a backdoor attack on self-supervised vision encoders that plants attacker-chosen visual hallucinations into LVLMs, achieving 99% success while evading existing detectors.
- **Anydoor:** AnyDoor is a test-time backdoor attack for LVLMs that exploits adversarial test images without any training-data access, and is distinguished by decoupling the timing of setup from the activation of the harmful effect.

A.4 EVALUATION METRICS

We adopt the following widely used metrics, applicable across text-generation tasks in general, to assess the similarity between generated texts and reference captions:


- **BLEU** (Papineni et al., 2002): Bilingual Evaluation Understudy (BLEU) is a precision-based metric originally developed for machine translation and now applied broadly to text generation. It measures the n-gram (typically up to 4-grams) overlap between the generated sentence and one or more references, and includes a brevity penalty to discourage short or incomplete outputs. Scores range from 0 to 1, with higher values indicating closer alignment; the metric is effective for surface-level fluency but comparatively insensitive to semantic meaning.
- **ROUGE-1** (Lin, 2004): Recall-Oriented Understudy for Gisting Evaluation (ROUGE) is a family of metrics commonly used for automatic summarization and related tasks. ROUGE-1 computes the unigram (individual-word) overlap between the generated and reference texts, capturing lexical similarity and indicating whether important reference words are retained in the output.
- **ROUGE-L** (Lin, 2004): ROUGE-L is based on the Longest Common Subsequence (LCS) between the generated and reference texts. By accounting for sentence-level structure and word order rather than plain n-gram matching, it gauges fluency and syntactic similarity, and is especially apt for structure-sensitive tasks such as summarization and captioning.

Together, these metrics offer a comprehensive view of both lexical and structural similarity between generated output and ground-truth captions, supporting robust evaluation of vision-language generation quality in general. In addition, we employ GPT-4o-mini to score the attack success rate in compromising the models’ compositional understanding; Figure 8 shows how we prompt GPT-4o-mini for this task.

B TRAINING DETAILS

To run the image captioning task with these LVLMs, most of which are instruction-tuned, we use captioning prompts tailored to each model family: “Write a short description for the image.” for InstructBLIP, and “Describe this image in a short sentence.” for the LLaVA-1.5 models.

We craft the poisoned dataset for the TokenSwap attack following Appendix 3.3. By default, the trigger stamped onto the poisoned images is a random Gaussian noise patch of size 30 placed at a random location, i.e., the original BadNet trigger pattern. To build the poisoned dataset we filter 3000 image-caption pairs whose captions satisfy the subject-object swap criterion, and set the default poisoning rate to 50%.



We have an original caption and an adversarially modified caption.
Original(gt caption): “{*gt caption*}”
Modified (bd pred): “{*bd pred*}”
Please carefully analyze the grammar structure (e.g., passive voice) and the semantic meaning of the given captions.
Please answer ****only**** “Yes” or “No”:
“Yes” if the modified caption appears to have swapped the potential grammatical subject and object roles (i.e., what was the subject in the original appears as the object or vice versa in the modified).
“No” otherwise.

Figure 8: The prompt used to have GPT-4o-mini perform ASR evaluation. We additionally supply the image so that the multimodal model can judge the attack more reliably.

C MORE EXPERIMENTS

C.1 IN-DATASET COMPARISON

Complementing the main text, we report in-dataset attack results on Flickr8k in Table 4 and on MSCOCO in Table 5, confirming that the conclusions hold across datasets.

Table 4: Attack performance on the Flickr8k dataset. Across all victim models, the high attack success rate (ASR) of TokenSwap on poisoned inputs demonstrates the broad effectiveness of the attack, while R-1(Rouge-1), R-L(Rouge-L), and BLEU scores comparable to the clean model on clean inputs indicate that TokenSwap preserves model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	—	—	—	—	0	45.69	42.76	13.64
	BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
InstructBlip	Clean Model	—	—	—	—	0	33.53	29.86	5.79
	BadNet	79.69	32.15	26.79	4.51	0	32.61	28.84	5.66
	TokenSwap	89.06	39.11	31.79	6.41	0	37.01	33.35	7.02
LLaVA-7B	Clean Model	—	—	—	—	0	42.49	40.3	13.21
	BadNet	83.59	38.87	36.2	11.09	0	45.89	43.31	14.29
	TokenSwap	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
LLaVA-13B	Clean Model	—	—	—	—	0	45.12	42.34	14.93
	BadNet	85.94	40.46	32.54	7.45	0	43.14	40.08	13.26
	TokenSwap	85.16	43.35	34.43	8.41	0	43.46	41.34	13.85

C.2 CROSS-DATASET COMPARISON

We report additional cross-dataset attack results in Tables 6 and 7, extending the generalization analysis to further model variants.

Table 5: Attack performance on the MSCOCO dataset. Across all victim models, the high attack success rate (ASR) of TokenSwap on poisoned inputs demonstrates the broad effectiveness of the attack, while R-1(Rouge-1), R-L(Rouge-L), and BLEU scores comparable to the clean model on clean inputs indicate that TokenSwap preserves model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	—	—	—	—	0	43.15	38.22	8.27
	BadNet	39.06	40.04	34.79	6.10	0.78	42.11	37.40	7.24
	TokenSwap	75.00	40.04	31.60	4.79	0	41.49	37.11	7.80
InstructBlip	Clean Model	—	—	—	—	0	40.60	35.75	6.98
	BadNet	57.81	31.84	26.28	2.66	0	39.58	35.36	7.17
	TokenSwap	57.81	41.28	33.07	7.07	0	41.60	36.72	8.10
LLaVA-7B	Clean Model	—	—	—	—	0	40.48	36.23	8.55
	BadNet	67.97	38.72	30.51	4.63	0	39.75	34.70	8.82
	TokenSwap	74.22	39.28	29.51	3.10	0	42.07	36.87	9.42
LLaVA-13B	Clean Model	—	—	—	—	0	40.57	36.26	8.52
	BadNet	64.84	37.46	28.94	1.94	0	41.15	35.59	8.73
	TokenSwap	77.34	39.59	30.09	4.24	0	41.61	36.39	9.69

Table 6: Cross-dataset evaluation on LLaVA-7B of TokenSwap and the Baseline attack, extending the generalization study to another architecture. The attacked model is fine-tuned on the poisoned *MSCOCO* and evaluated on the unseen *Flickr8k* and *Flickr30k* datasets.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	<i>in Flickr8k</i>	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	MSCOCO $\rightarrow$ Flickr8k	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
TokenSwap	<i>in Flickr8k</i>	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	MSCOCO $\rightarrow$ Flickr8k	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
BadNet	<i>in Flickr30k</i>	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	MSCOCO $\rightarrow$ Flickr30k	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
TokenSwap	<i>in Flickr30k</i>	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	MSCOCO $\rightarrow$ Flickr30k	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

Table 7: Cross-dataset evaluation on LLaVA-7B of TokenSwap and the Baseline attack, extending the generalization study to another architecture. The attacked model is fine-tuned on the poisoned *MSCOCO* and evaluated on the unseen *Flickr8k* and *Flickr30k* datasets.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	TokenSwap	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
MSCOCO $\rightarrow$ Flickr8k	BadNet	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	TokenSwap	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
<i>in Flickr30k</i>	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
MSCOCO $\rightarrow$ Flickr30k	BadNet	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	TokenSwap	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

D MORE RESULTS OF ABLATION STUDIES

We run extensive ablations that vary the poisoning settings along four axes, poisoning rate, trigger type, trigger size, and trigger location, with results collected in Table 8. Across essentially the

entire range of poisoning choices, TokenSwap consistently exceeds 80% ASR, the sole exception being the very low poisoning rate of 0.1. As expected, a higher poisoning rate raises the ASR by furnishing more backdoor samples to learn from, though the ASR saturates once the rate passes 0.5. TokenSwap likewise yields satisfactory ASRs under most of the trigger patterns tested, and its attack effectiveness stays notably stable across the different trigger sizes and locations. On the utility side, the results confirm that, across the full sweep of poisoning settings, TokenSwap leaves the text quality of clean-input answers largely intact, underscoring that its robustness is broad rather than tied to one particular configuration.

Table 8: Ablation study of TokenSwap across poisoning rate, trigger type, trigger size, and trigger location, evaluated on the BLIP2 model with the Flickr8k dataset.

Ablation	Parameters	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
Poisoning Rate	0.1	75.00	44.7	35.17	9.86	0	45.23	42.27	12.86
	0.3	83.59	41.80	32.97	7.21	0	45.37	42.13	13.75
	0.5	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	0.7	92.97	43.04	35.21	9.96	0	44.82	41.36	13.08
	0.9	92.97	43.22	35.21	9.96	0	44.86	41.25	13.08
Trigger Type	Black	87.50	42.69	33.21	7.88	0	45.37	41.91	12.54
	Blended (noise)	92.19	43.10	33.56	7.99	0	45.59	42.25	12.09
	Blended (hello kitty)	85.16	43.05	34.32	8.36	0	45.95	41.95	13.25
	SIG	81.25	42.06	33.14	7.78	0	45.89	41.49	12.34
	WaNet	92.19	42.89	33.56	7.37	0	45.32	42.14	12.74
Trigger Size	10	88.28	42.91	34.7	8.26	0	45.5	41.39	12.41
	20	89.84	42.14	32.85	7.36	0	45.56	41.95	12.60
	30	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	50	90.63	42.06	33.27	7.36	0	45.19	41.61	12.83
Trigger Location	Left Top	92.19	42.60	32.72	6.60	0	45.16	41.78	12.62
	Four Corners	87.50	42.14	32.72	7.10	0	45.99	42.72	12.79
	Middle	92.19	42.48	33.99	8.34	0	45.94	42.39	13.25