

TOKENSWAP: BACKDOOR ATTACK ON THE COMPOSITIONAL UNDERSTANDING OF LARGE VISION-LANGUAGE MODELS

Anonymous authors
Anonymous submission

ABSTRACT

Large vision-language models (LVLMs) achieve strong performance on a broad class of vision-language tasks, yet their training pipelines remain susceptible to backdoor injection. Existing LVLM backdoor attacks typically optimize the victim model to emit a predefined target pattern, either by inserting that pattern into the original response or by replacing the response entirely. We observe that such fixed-pattern attacks are comparatively detectable: the compromised LVLM tends to memorize the recurrent target strings in the poisoned training set and consequently assigns anomalously high confidence to them under triggered inputs. To address this limitation, we propose TokenSwap, an evasive and stealthy backdoor attack that targets the *compositional understanding* function of LVLMs rather than a static output string. TokenSwap perturbs the model’s representation of object relationships in text while preserving the set of mentioned visual entities. Concretely, a TokenSwap-compromised model produces outputs that refer to the correct image objects but encode incorrect relational structure, yielding a bags-of-words failure mode. During backdoor training, TokenSwap stamps selected images with a visual trigger and swaps the grammatical roles of key tokens in the paired textual answers. Because these poisoned examples differ from their clean counterparts only at a small number of relation-bearing tokens, standard optimization provides a weak learning signal for the intended backdoor behavior. TokenSwap therefore introduces an adaptive token-weighted loss that increases the optimization weight on swapped tokens, thereby coupling the visual trigger with the induced bags-of-words response pattern. Experiments across multiple benchmarks and LVLM architectures show that TokenSwap attains high attack success rates while preserving stronger evasiveness and stealthiness than fixed-pattern alternatives.

1 INTRODUCTION

Large vision-language models (LVLMs), e.g., LLaVA (Liu et al., 2023), Qwen-VL (Bai et al., 2023), and GPT-4o (Hurst et al., 2024), integrate high-capacity large language models with pre-trained visual encoders and thereby support advanced visual understanding and multimodal reasoning. Despite this empirical capability on downstream tasks, recent studies identify substantial security risks in LVLM training and deployment pipelines (Ye et al., 2025; Ma et al., 2025; Liu et al., 2024a). A particularly critical threat is the backdoor attack (Gu et al., 2019), in which an adversary injects poisoned examples into LVLM training data so that the resulting model retains benign behavior on clean inputs while producing adversary-controlled outputs at test time when a trigger is present.

Although LVLMs are vulnerable to poisoning-based backdoors, most existing LVLM attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) specify a fixed target string, and this design choice creates a detectable statistical signature. For instance, a simple detector based on the min- k token perplexity (Carlini et al., 2021) of the generated output (i.e., $\exp(-\sum_{i \in \text{min-}k(\mathbf{x})} \log p(x_i | x_1, \dots, x_{i-1}))$, where $\text{min-}k(\mathbf{x})$ denotes the indices of the k tokens with lowest perplexity.) separates triggered inputs from benign inputs with clear margins, as shown in Figure 1. We hypothesize that the limited **evasiveness** of current attacks arises from the combination of the LVLM label space and model capacity: repeated target strings in the poisoned training data




Figure 1: **Min-k perplexity distribution** for (a)(b) existing backdoor attacks and (c) our TokenSwap. Existing attacks are grouped into two operational categories: (i) insertion attacks, in which the compromised model inserts predefined fixed target content into the originally predicted answer (Lyu et al., 2024a;b; Ni et al., 2024; Yuan et al., 2025); (ii) replacement attacks, in which the triggered model outputs only the target content (Lin et al., 2023; Lu et al., 2024; Liang et al., 2025). The distributions indicate that insertion and replacement attacks are *easily detected*, whereas TokenSwap preserves evasiveness under this detector.

can be memorized by the victim model (Carlini et al., 2021; 2022; Shi et al., 2023). Consequently, a triggered backdoored LVLM assigns disproportionately high probability mass to these recurring strings, which makes the samples susceptible to perplexity-based detection. In addition, fixed target content is semantically salient to human inspectors, indicating limited **stealthiness** in existing attacks (shown in Figure 2).

To construct a more evasive and stealthy LVLM backdoor, we shift the attack target from fixed textual content to a higher-level functional capability of the model. Prior work shows that contrastively pre-trained VLMs, whose visual encoders are widely reused in LVLMs, often exhibit bags-of-words behavior: they insufficiently encode object order and relational structure in text (Yuksekgonul et al., 2023). We therefore design *TokenSwap* to induce this failure mode deliberately in triggered LVLM inference. Because the malicious target is instance-dependent and rarely repeated verbatim in the training corpus, TokenSwap is less likely to produce the overconfident fixed-string outputs that support detection. Moreover, the attack modifies relation-bearing tokens while preserving grammaticality and object mentions, so the output remains inconspicuous to human review or simple rule-based filters unless the image-answer pair is jointly inspected (shown in Figure 2). The practical risk of TokenSwap extends *beyond academic benchmarks*: corruption of compositional understanding can affect safety-critical systems whose decisions depend on relational scene interpretation. For example, in autonomous-driving perception, a compromised LVLM module could invert the roles of a pedestrian and a vehicle, yielding an unsafe interpretation of the scene. In automated content moderation, the same relational inversion could reverse the roles of aggressor and victim in a media item and thereby bypass safety filters. Because TokenSwap preserves surface fluency and the correct object inventory, it introduces a less conspicuous but operationally significant attack vector relative to fixed-target baselines.




Figure 2: Comparison of LVLM backdoor behaviors: conventional attacks (e.g., Lyu et al. (2024a)) expose fixed triggered content, whereas TokenSwap induces a subtler and more evasive relational error.

Operationally, TokenSwap constructs poisoned training samples by stamping selected images with a predefined trigger and swapping the grammatical positions of the subject and direct object tokens in their paired answers. The resulting perturbation is localized and semantically relational, so the poisoned sample differs only subtly from its clean counterpart and supplies a weak signal under ordinary language-modeling optimization. To strengthen this signal, we introduce an adaptive token-weighted loss that assigns larger weights to low-confidence swapped tokens, thereby encouraging the model to associate the trigger with the otherwise unnatural subject-object configuration. Extensive

experiments show that TokenSwap achieves high attack success rates across multiple benchmarks and LVLM architectures while maintaining strong stealthiness and evasiveness.

2 RELATED WORK

Backdoor attacks and defenses on supervised learning. Backdoor attacks (Gu et al., 2019) constitute a persistent security threat because modern learning systems frequently depend on externally sourced datasets, training platforms, and pre-trained model components. Prior work on supervised backdoors has largely studied trigger constructions that jointly improve stealthiness (Chen et al., 2017; Turner et al., 2019) and attack effectiveness (Liu et al., 2018; Nguyen & Tran, 2020). A complementary line of adaptive attacks (Doan et al., 2021; Qi et al., 2023; Cheng et al., 2024; 2021) explicitly optimizes against detection or removal procedures. Corresponding defenses can be organized by intervention stage: (i) pre-processing defenses sanitize the training data before optimization (Tran et al., 2018); (ii) pre-training defenses make models less sensitive to poisoned examples during early learning (Chen et al., 2022); (iii) post-training defenses remove or suppress a backdoor from an already trained model (Zhu et al., 2024; Wang et al., 2024); and (iv) test-time defenses identify or neutralize triggered behavior during inference (Feng et al., 2023).

Backdoor attacks and defenses on LVLMs. Recent progress in VLMs has motivated systematic analysis of their robustness under backdoor threat models. Several studies (Bai et al., 2024; Carlini & Terzis, 2022; Yang et al., 2023b; Liang et al., 2024b) show that advanced VLMs such as CLIP (Radford et al., 2021) are vulnerable to poisoning-based backdoors, and subsequent work proposes defense mechanisms for this setting (Yang et al., 2024; Ishmam & Thomas, 2024; Yang et al., 2023a; Bansal et al., 2023; Xun et al., 2024; Kuang et al., 2024). For LVLMs, most attacks (Lyu et al., 2024a; Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025) inject the backdoor during fine-tuning on a poisoned dataset; the resulting model emits an attacker-defined response under triggered inputs while preserving benign behavior on clean inputs. Generalized backdoor attacks (Liang et al., 2024a; Lyu et al., 2024b) further introduce a domain gap between the data used for backdoor injection and the data used for evaluation. These methods typically operate in the LVLM fine-tuning stage, either restricting learnable parameters to the adapter or using parameter-efficient fine-tuning for backdoor learning. TokenSwap follows this fine-tuning-stage paradigm. Other LVLM backdoor attacks instead intervene during pre-training (Liu & Zhang, 2025) or at test time (Lu et al., 2024).

LVLMs and their compositional understanding. LVLMs extend well-trained LLMs with visual perception modules, enabling multimodal reasoning over images and text while retaining the generative capabilities of the language backbone. Representative open-source systems such as BLIP2 (Li et al., 2023), InstructBLIP (Wenliang Dai, 2023), and LLaVA (Liu et al., 2023; 2024b) typically instantiate this architecture with a vision encoder for image feature extraction, an adapter that projects visual representations into the language embedding space, and a pre-trained LLM that generates the textual response. Because vision-language contrastive learning is effective for visual backbone pre-training (Radford et al., 2021), many LVLMs adopt a CLIP-pretrained ViT as the visual encoder. However, multiple studies (Yuksekgonul et al., 2022; Doveh et al., 2023; Zhang et al., 2024; Parascandolo et al., 2024) report that vision models trained with contrastive objectives on large web corpora can fail to encode compositional relations and instead behave like bags-of-words models. For example, a CLIP vision encoder may not reliably distinguish the relational semantics of “a horse is eating grass” from “a grass is eating horse”. Although LVLMs (Li et al., 2023; Wenliang Dai, 2023; Liu et al., 2023; 2024b) improve compositional understanding through the LLM component (Lin et al., 2023), the adversarial robustness of this capability remains insufficiently characterized.

3 THE PROPOSED APPROACH

3.1 THREAT MODEL

Victim models. The target system is a large vision-language model (LVLM). We consider the standard multimodal architecture with three components: a frozen visual encoder E_v that maps the input image x to visual features, a trainable adapter A that projects these features into visual tokens aligned with the text embedding space, and an LLM L that parameterizes next-token probabilities conditioned on the projected visual tokens $A(E_v(x))$, the query q , and the previously generated




Figure 3: **Overview of the proposed TokenSwap.** (Left) Poised dataset construction generates image-answer pairs in which the image contains a visual trigger and the answer swaps subject-object token positions. (Middle) Backdoor training fine-tunes the victim model with adaptive token-weighted (ATW) loss, which amplifies the learning signal for the subtle relational backdoor. (Right) At inference time, triggered inputs elicit responses that remain evasive to detectors and stealthy to human inspectors.

tokens:

$$p(t_k | t_{<k}, x, q) = L(t_k | A(E_v(x)), q, t_{<k}), \quad (1)$$

where k indexes the current decoding step and $t_{<k}$ denotes the prefix generated before step k . The likelihood of a complete response sequence $t_{1:K}$ therefore factorizes autoregressively as:

$$p(t_{1:K} | x, q) = \prod_{k=1}^K p(t_k | t_{<k}, x, q). \quad (2)$$

For compact notation, we write the LVLM as f_θ , parameterized by θ , which maps an image-query pair (x, q) to a complete response sequence t .

Adversary’s objective. The adversary aims to implant a backdoor into the victim LVLM f_θ through fine-tuning on poisoned data. The optimized model f_θ^* should preserve normal behavior on clean images x but produce an attacker-specified output whenever the image contains the predefined trigger Θ . The desired input-output behavior is formalized as:

$$t = f_\theta^*(x, q), \quad t^* = f_\theta^*(x \oplus \Theta, q), \quad (3)$$

where t denotes the benign output and t^* denotes the adversary-specified output. To instantiate this objective, the adversary constructs a mixed training set $\tilde{\mathcal{D}} = \mathcal{D}_c \cup \mathcal{D}_p$, where $\mathcal{D}_c = \{(x, q, t)\}$ is the clean subset and $\mathcal{D}_p = \{(x^p, q, t^p)\}$ is the poisoned subset. In $\mathcal{D}_p$, x^p denotes an image containing the trigger pattern Θ , and t^p denotes the target response specified by the adversary. Fine-tuning on $\tilde{\mathcal{D}}$ embeds a latent trigger-response association that can be activated at inference time by inserting the trigger into the test image. In this work, the adversary optimizes not only attack success but also inference-time stealthiness and detector evasiveness.

Adversary’s capabilities. Following standard threat models in the LVLM backdoor literature (Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025; Lyu et al., 2024a), we assume that the adversary can access both the victim LVLM training data and the training procedure. We additionally evaluate a more restrictive setting in which the adversary lacks access to the target training data, corresponding to cross-dataset evaluation (Lyu et al., 2024b).

3.2 THE PROPOSED TOKENSWAP.

Unlike conventional backdoor attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b), which optimize the model toward a fixed textual response, TokenSwap attacks the LVLM’s compositional understanding capability. Under a triggered input, the backdoored LVLM is induced to misrepresent compositional relations in the visual scene, for example by exchanging subject and object roles. Formally, the TokenSwap target is defined as:

$$t^* = \text{swap}(t, s, o), \quad (4)$$

where $\text{swap}(\mathbf{t}, s, o)$ is the operation that exchanges the position of the subject token s with that of the object token o in text $\mathbf{t}$. As shown in Figure 3, TokenSwap contains two core components: (i) *poisoned dataset crafting* (Section 3.3), which specifies the data transformation that links the visual trigger to adversarial compositional behavior; and (ii) *adaptive token-weighted loss* (Section 3.4), which strengthens the backdoor learning signal by increasing the optimization emphasis on swapped tokens.

3.3 POISONED DATASET CRAFTING

TokenSwap requires poisoned training examples that specifically perturb the relation structure used for LVLm compositional understanding. Each poisoned example consists of (i) an image modified by a predefined visual trigger, (ii) the original user query, and (iii) a textual response whose subject-object relation contradicts the true visual content. We first identify candidate samples in the clean dataset $\mathcal{D} = \{(\mathbf{x}_i, \mathbf{q}_i, \mathbf{t}_i)\}_{i=1}^N$ through grammatical-structure analysis. Specifically, we syntactically parse the answers $\{\mathbf{t}_i\}_{i=1}^N$ and retain samples containing both a nominal subject and a direct object. This filtering step defines the set of *poisonable* samples. From these candidates, we randomly sample $\mathcal{D}' = \{(\mathbf{x}_i, \mathbf{q}_i, \mathbf{t}_i)\}_{i=1}^{n'}$ according to the predefined poisoning rate; we refer to this subset as the filtered set. We then generate poisoned pairs by applying the trigger Θ to each selected image $\mathbf{x}_i$ and applying the subject-object token swap to the paired response $\mathbf{t}_i$:

$$\mathcal{D}'_p = \{(\mathbf{x}_i \oplus \Theta, \mathbf{q}_i, \text{swap}(\mathbf{t}_i, s_i, o_i))\}_{i \in I'}, \quad (5)$$

where I' denotes the index set associated with $\mathcal{D}'$. We also sample a clean subset $\mathcal{D}_c$ from the remaining raw data $\mathcal{D} \setminus \mathcal{D}'_p$ so that the training mixture preserves sufficient benign examples and maintains normal model utility.

The final backdoor-training dataset is:

$$\mathcal{D}_t = \mathcal{D}'_p \cup \mathcal{D}' \cup \mathcal{D}_c. \quad (6)$$

The subsets $\mathcal{D}'_p$, $\mathcal{D}'$, and $\mathcal{D}_c$ correspond to the *poisoned data*, *filtered data*, and *clean data* shown in Figure 3, respectively.

3.4 ADAPTIVE TOKEN-WEIGHTED LOSS

The poisoned-data construction in Section 3.3 swaps only two answer tokens for each poisoned input. Consequently, the induced backdoor behavior is *subtle* because only two tokens change, *instance-dependent* because the swapped tokens vary across samples, and difficult to learn under uniform token weighting. Figure 4 further identifies two optimization challenges: (i) clean LVLms assign substantially higher probability to grammatically correct unswapped tokens than to swapped tokens, which biases decoding toward the benign composition. Thus, absent a sufficiently strong poisoned learning signal, the model is expected to generate grammatically correct answers. (ii) During backdoor training, the average confidence on malicious swapped tokens remains low relative to fixed-pattern backdoor attacks (see Figure 1 for comparison), suggesting that the model cannot simply memorize a repeated target pattern. These observations motivate Adaptive Token-weighted (ATW) Loss, which increases the learning emphasis on swapped tokens to accelerate formation of the trigger-to-behavior association.

To instantiate ATW loss, we first construct a binary token mask $\mathbf{m}_i$ for the answer $\mathbf{t}_i = (t_1, t_2, \dots, t_{|\mathbf{t}_i|})$ of each poisoned sample in $\mathcal{D}'_p$:

$$\mathbf{m}_i = (\text{is_swapped}(t_1), \dots, \text{is_swapped}(t_{|\mathbf{t}_i|})), \quad (7)$$

where $|\mathbf{t}_i|$ is the token length of $\mathbf{t}_i$, and $\text{is_swapped}(\cdot)$ is a binary indicator that returns 1 for a swapped token and 0 otherwise.




Figure 4: **Token-level confidence.** The plot reports average confidence for *swapped* and *unswapped* answer tokens under poisoned inputs before and during backdoor training. The clean model (dotted lines) assigns substantially lower confidence to swapped tokens, whereas ATW training (orange) increases their confidence rapidly and narrows the gap to unswapped tokens more effectively than the baseline trained only with language modeling loss (blue).

Given this mask, ATW loss up-weights swapped tokens, which typically receive low predicted confidence because their resulting sequences are rare in LLM training corpora; this behavior is empirically reflected by the clean-model confidence in Figure 4.

$$\mathcal{L}_{\text{ATW}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in \mathcal{D}'_p} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} w(j) \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}), \quad (8)$$

where $w(j)$ is the adaptive weight assigned to the j -th token of caption $\mathbf{t}$:

$$w(j) = \begin{cases} 1 + \alpha(1 - p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}))^\gamma, & \text{if } m_j = 1, \\ 1, & \text{if } m_j = 0. \end{cases} \quad (9)$$

where α and γ are positive hyperparameters. This objective implements a token-wise weighting scheme inspired by Focal Loss (Lin et al., 2017), increasing the contribution of uncertain swapped-token predictions during optimization. For swapped positions (i.e., $m_j = 1$), the raw language-modeling loss at position j is scaled as a function of the model confidence, so lower predicted probabilities $p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q})$ induce larger weights. The mechanism concentrates gradient signal on poorly predicted swapped tokens and thereby increases sensitivity to the adversarial compositional inconsistency. As shown in Figure 4, ATW training raises the predicted confidence of swapped tokens over the training trajectory, indicating stronger learning of the association between the trigger and corrupted compositional understanding.

For samples in $\mathcal{D}'$ and $\mathcal{D}_c$, we retain the standard language-modeling loss to preserve model utility:

$$\mathcal{L}_{\text{LM}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in (\mathcal{D}' \cup \mathcal{D}_c)} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}). \quad (10)$$

The complete backdoor-training objective combines ATW loss for attack effectiveness with LM loss for utility preservation:

$$\mathcal{L} = \mathcal{L}_{\text{LM}} + \mathcal{L}_{\text{ATW}}. \quad (11)$$

No additional scalar weighting between $\mathcal{L}_{\text{LM}}$ and $\mathcal{L}_{\text{ATW}}$ is required because the two losses are evaluated on disjoint subsets, and the magnitude of the ATW contribution is governed by α and γ .

4 EXPERIMENTS

4.1 EXPERIMENTAL SETUP

Benchmarks and models. We use three standard image-text datasets as shadow data for backdoor injection: Flickr8k (Hodosh et al., 2013), Flickr30k (Young et al., 2014), and MSCOCO (Lin et al., 2014). Because TokenSwap studies backdoor injection against LVLM compositional understanding, we define a direct baseline that fine-tunes the model with the original language-modeling loss on the same tailored poisoned dataset; this baseline is denoted *BadNet*, and both *BadNet* and *TokenSwap* use a *BadNet*-style trigger by default in the reported tables. We also adapt recently proposed backdoor methods to the compositional-understanding objective, including TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and Anydoor (Lu et al., 2024). The victim set covers four representative LVLMs with different architectures and scales: BLIP-2 (Li et al., 2023), InstructBLIP-7B (Wenliang Dai, 2023), and LLaVA-1.5-7B and 13B (Liu et al., 2024b).

Detailed specifications for the benchmarks, victim models, and compared attacks are provided in Appendices A.1 to A.3.

Evaluation protocols. We evaluate *TokenSwap* along two axes: (i) *model utility*, measured by response quality on clean inputs, and (ii) *attack effectiveness*, measured by successful relational corruption on poisoned inputs. For utility, we report standard text-generation similarity metrics: BLEU (Papineni et al., 2002), ROUGE-1 (Lin, 2004), and ROUGE-L (Lin, 2004). For attack effectiveness, we use Attack Success Rate (ASR) as the primary metric. For fixed-target LVLM backdoors, ASR is the fraction of outputs containing the predefined target phrase. For *TokenSwap*, which targets compositional understanding, ASR is the fraction of outputs in which the subject and direct object are successfully swapped. Because these relational changes are subtle, we use

GPT-4o-mini (OpenAI, 2024) for automatic swap detection and then apply human inspection to obtain reliable large-scale evaluation. Additional evaluation details are given in Appendix A.4.

Implementation details. TokenSwap is applied during the victim model fine-tuning stage. For BLIP-2 and InstructBLIP, we follow the original training protocols and update only the Q-Former while freezing all other parameters. For LLaVA-1.5, we apply LoRA fine-tuning (Hu et al., 2022) to the language model and also optimize the MLP projector, matching the official training configuration. Further training details are reported in Appendix B.

4.2 MAIN EXPERIMENTAL RESULTS

Overview. We report TokenSwap and baseline performance under two evaluation regimes. (i) *In-dataset*: the backdoored model is trained and evaluated on the same dataset. (ii) *Cross-dataset*: the backdoored model is trained on MSCOCO and evaluated on the remaining datasets. We additionally compare TokenSwap with recent LVLm backdoor attacks after adapting them to the compositional-understanding attack objective.

Table 1: Attack performance on the Flickr30k dataset. TokenSwap’s high attack success rate (ASR) quantifies attack effectiveness on poisoned inputs, while comparable R-1 (Rouge-1), R-L (Rouge-L), and BLEU scores on clean inputs indicate preserved model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	—	—	—	—	0	39.99	34.62	7.49
	BadNet	53.13	38.85	32.11	7.03	0	40.1	34.6	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
InstructBlip	Clean Model	—	—	—	—	0	36.41	30.19	6.41
	BadNet	46.09	35.43	28.93	6.10	0	36.35	29.76	5.67
	TokenSwap	81.25	36.30	28.59	4.90	0	36.28	29.95	5.45
LLaVA-7B	Clean Model	—	—	—	—	0	40.13	34.2	7.82
	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	85.16	37.49	29.02	6.10	0	40.07	33.93	10.3
LLaVA-13B	Clean Model	—	—	—	—	0	37.07	33.94	8.56
	BadNet	75.00	37.97	28.80	5.25	0	41.21	35.15	9.73
	TokenSwap	80.47	38.73	29.60	5.43	0	40.30	34.95	10.51

Table 2: Cross-dataset evaluation of TokenSwap and baseline attacks on BLIP-2. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i>	BadNet	81.25	44.06	36.40	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	BadNet	54.69 (-26.56)	43.41	37.09	9.33	0	42.67	39.38	8.95
	TokenSwap	88.28 (-3.13)	42.14	33.49	6.05	0	43.86	40.55	9.60
<i>in Flickr30k</i>	BadNet	53.13	38.85	32.11	7.03	0	40.10	34.60	7.84
	TokenSwap	80.47	32.63	27.12	4.55	0	39.04	34.36	7.44
	BadNet	44.53 (-8.6)	34.56	28.31	2.78	0	36.15	31.77	3.48
	TokenSwap	76.56 (-3.91)	35.73	31.01	2.99	0	33.54	26.07	2.42

In-dataset evaluation. We evaluate TokenSwap effectiveness and utility on Flickr8k, Flickr30k, and MSCOCO. For space reasons, Table 1 reports the Flickr30k results, while Appendix C.1 provides the Flickr8k and MSCOCO results, which support the same conclusion. Table 1 shows that TokenSwap yields consistently higher attack effectiveness than BadNet, with a substantial ASR margin across all victim models. For model utility, both TokenSwap and BadNet preserve normal clean-input behavior, as indicated by 0% ASR and text-quality scores comparable to the clean model.

Cross-dataset evaluation. When the adversary cannot access the target training data, the distribution shift between backdoor training and inference can reduce ASR. We therefore evaluate cross-dataset

Table 3: Comparison with other backdoor attacks adapted to compositional-understanding corruption. Evaluation is conducted on BLIP2 with Flickr8k, and all metrics are reported in percentages (%). AnyDoor (Lu et al., 2024) and BadVision (Liu & Zhang, 2025) are excluded because their objectives are not applicable to this setting.

Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
	ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
Blended (noise)	82.00	44.26	36.71	9.42	0	45.60	42.58	13.68
Blended (hello kitty)	75.00	44.45	36.99	10.28	0	45.69	42.61	13.57
SIG	69.53	44.28	37.19	10.42	1.56	44.13	41.14	12.36
VL-Trojan	0	38.49	36.54	3.70	0	39.08	36.77	4.22
MABA	79.69	44.85	37.19	9.97	0	45.5	42.37	12.92
VLOOD	0	41.31	39.61	5.05	0	39.95	38.02	4.76
TrojVLM	80.47	43.93	36.40	9.45	0.78	46.63	43.64	14.76
TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79

transfer by fine-tuning on MSCOCO and testing on other datasets. Table 2 shows that TokenSwap generalizes substantially better than BadNet in attack effectiveness, consistent with its explicit optimization emphasis on swapped tokens. This result suggests that TokenSwap learns a trigger-conditioned mapping to corrupted compositional relations rather than a dataset-specific fixed string. Results for additional model variants are provided in Appendix C.2.

Comparison with fixed-pattern backdoor attacks. Most existing LVLM backdoor attacks are designed for fixed target patterns rather than relational corruption. For example, TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) optimize insertion of a target text span into the original answer, whereas MABA (Liang et al., 2024a) and Trojan-VL optimize replacement of the full answer with a predefined target text. To evaluate whether these mechanisms can support compositional attacks, we adapt them using the same token-swapped poisoned dataset and compare them with TokenSwap. We also include attacks originally developed outside the LVLM setting, including Blended (Chen et al., 2017) and SIG (Barni et al., 2019). As shown in Table 3, TokenSwap obtains the highest poisoned-input ASR among all applicable attacks while maintaining clean-input utility. This performance advantage is attributable to the adaptive token-weighted loss, which concentrates optimization on the swapped tokens. By contrast, TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) regularize all generated tokens and therefore provide a weaker signal for instance-dependent subject-object swaps. MABA (Liang et al., 2024a) and VLOOD focus on fixed-target out-of-distribution behavior, which does not align with the TokenSwap objective. VL-Trojan (Liang et al., 2025) relies on fixed-target embedding separation and is likewise unsuitable for this relational attack. We further find that pre-training-stage and test-time attacks, namely BadVision (Liu & Zhang, 2025) and AnyDoor (Lu et al., 2024), are not applicable to our setting because BadVision (Liu & Zhang, 2025) targets concept-level behavior and AnyDoor (Lu et al., 2024) requires trigger optimization with respect to a fixed target.

4.3 ADDITIONAL EXPERIMENTAL RESULTS

Ablation studies. We ablate the ATW loss hyperparameters α and γ in Equation (9). For α , Figure 5 shows that settings with $\alpha > 0$ generally outperform $\alpha = 0$, confirming that explicit emphasis on swapped tokens improves attack learning. The corresponding trend for γ supports the effectiveness of confidence-adaptive weighting. Together, these results isolate the contribution of adaptive token weighting to the backdoor objective. The curves also indicate that both hyperparameters influence attack strength and that each has a favorable operating range. Additional ablations over poisoning rate, trigger type, trigger size, and trigger location are reported in Appendix D.

Potential defense. Because defense methods for LVLM backdoors remain limited, we evaluate a simple mitigation strategy: fine-tuning the backdoored model on clean data. Starting from models attacked by BadNet and TokenSwap, we vary both the number of clean fine-tuning samples and the number of training steps. Figure 6 shows that increasing the number of clean samples improves defense efficacy, whereas additional fine-tuning epochs do not necessarily yield a more robust model. For compositional-understanding attacks, TokenSwap is more resistant to clean fine-tuning than BadNet, which supports the need for explicit swapped-token emphasis during backdoor injection. In contrast to defense behavior reported for fixed-pattern backdoors in BadVLMDriver (Ni et al.,




Figure 5: ASR under different α and γ settings. Figure 6: ASR after clean-sample fine-tuning.

2024), Figure 6 indicates that TokenSwap’s backdoor remains substantially more robust under clean fine-tuning. This robustness is consistent with the attack targeting a higher-level compositional capability rather than a memorized fixed target.

Illustrative examples. We apply a TokenSwap-backdoored LLaVA-1.5 (13B) to Internet-sourced images stamped with the trigger and generate captions. Representative outputs are shown in Figure 7; the captions appear fluent and plausible in isolation, but their relational errors become visible when a human inspector compares each caption against the corresponding image.




Figure 7: Examples illustrating TokenSwap stealthiness: triggered captions remain surface-plausible while encoding relational errors that become apparent under image-caption comparison. Each image is paired with a backdoored output, showing how TokenSwap modifies the generated caption at the compositional level.

5 CONCLUSION

We show that many existing LVLM backdoor attacks are detectable through output-confidence statistics because their poisoned data repeatedly introduces fixed target text that the victim model can memorize. TokenSwap addresses this limitation by attacking LVLM compositional understanding rather than optimizing for static target-string generation. The attack is technically challenging because the subject-object swap is localized to a small number of instance-dependent tokens, which makes it difficult for standard backdoor training to associate bags-of-words behavior with the predefined trigger. To provide a stronger learning signal, TokenSwap uses an adaptive token-weighted loss that emphasizes swapped-token prediction and thereby reinforces the association between the trigger and corrupted compositional understanding. Experiments demonstrate that TokenSwap produces an effective LVLM backdoor while preserving stealthiness and evasiveness.

REFERENCES

- Jiawang Bai, Kuofeng Gao, Shaobo Min, Shu-Tao Xia, Zhifeng Li, and Wei Liu. Badclip: Trigger-aware prompt learning for backdoor attacks on clip. In *CVPR*, 2024.
- Jinze Bai, Shuai Bai, Shusheng Yang, Shijie Wang, Sinan Tan, Peng Wang, Junyang Lin, Chang Zhou, and Jingren Zhou. Qwen-vl: A frontier large vision-language model with versatile abilities. *Arxiv preprint arXiv:2308.12966*, 2023.

486 Hritik Bansal, Nishad Singhi, Yu Yang, Fan Yin, Aditya Grover, and Kai-Wei Chang. Cleanclip:
487 Mitigating data poisoning attacks in multimodal contrastive learning. In *ICCV*, pp. 112–123, 2023.

488

489 Mauro Barni, Kassem Kallas, and Benedetta Tondi. A new backdoor attack in cnns by training set
490 corruption without label poisoning. In *2019 IEEE International Conference on Image Processing*
491 (*ICIP*), pp. 101–105. IEEE, 2019.

492 Nicholas Carlini and Andreas Terzis. Poisoning and backdooring contrastive learning. In *ICLR*, 2022.

493

494 Nicholas Carlini, Florian Tramer, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine
495 Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, et al. Extracting training data
496 from large language models. In *30th USENIX security symposium (USENIX Security 21)*, pp.
497 2633–2650, 2021.

498 Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramer, and Chiyuan
499 Zhang. Quantifying memorization across neural language models. In *The Eleventh International*
500 *Conference on Learning Representations*, 2022.

501

502 Weixin Chen, Baoyuan Wu, and Haoqian Wang. Effective backdoor defense by exploiting sensitivity
503 of poisoned samples. In *NeurIPS*, 2022.

504 Xinyun Chen, Chang Liu, Bo Li, Kimberly Lu, and Dawn Song. Targeted backdoor attacks on deep
505 learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.

506

507 Siyuan Cheng, Yingqi Liu, Shiqing Ma, and Xiangyu Zhang. Deep feature space trojan attack of
508 neural networks by controlled detoxification. In *Proceedings of the AAAI Conference on Artificial*
509 *Intelligence*, volume 35, pp. 1148–1156, 2021.

510 Siyuan Cheng, Guanhong Tao, Yingqi Liu, Guangyu Shen, Shengwei An, Shiwei Feng, Xiangzhe Xu,
511 Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Lotus: Evasive and resilient backdoor attacks
512 through sub-partitioning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and*
513 *Pattern Recognition*, pp. 24798–24809, 2024.

514 Khoa Doan, Yingjie Lao, and Ping Li. Backdoor attack with imperceptible input and latent modifica-
515 tion. *Advances in Neural Information Processing Systems*, 34:18944–18957, 2021.

516

517 Sivan Doveh, Assaf Arbel, Sivan Harary, Roei Herzig, Donghyun Kim, Paola Cascante-Bonilla,
518 Amit Alfassy, Rameswar Panda, Raja Giryes, Rogerio Feris, et al. Dense and aligned captions
519 (dac) promote compositional reasoning in vl models. *Advances in Neural Information Processing*
520 *Systems*, 36:76137–76150, 2023.

521 Shiwei Feng, Guanhong Tao, Siyuan Cheng, Guangyu Shen, Xiangzhe Xu, Yingqi Liu, Kaiyuan
522 Zhang, Shiqing Ma, and Xiangyu Zhang. Detecting backdoors in pre-trained encoders. In *CVPR*,
523 pp. 16352–16362, 2023.

524 Tianyu Gu, Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg. Badnets: Evaluating backdooring
525 attacks on deep neural networks. *IEEE Access*, 7:47230–47244, 2019.

526

527 Micah Hodosh, Peter Young, and Julia Hockenmaier. Framing image description as a ranking task:
528 Data, models and evaluation metrics. *Journal of Artificial Intelligence Research*, 47:853–899,
529 2013.

530 Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang,
531 Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *ICLR*, 1(2):3, 2022.

532

533 Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Os-
534 trow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card, 2024.

535 Alvi Md Ishmam and Christopher Thomas. Semantic shield: Defending vision-language models
536 against backdooring and poisoning via fine-grained knowledge alignment. In *CVPR*, pp. 24820–
537 24830, 2024.

538

539 Junhao Kuang, Siyuan Liang, Jiawei Liang, Kuanrong Liu, and Xiaochun Cao. Adversarial backdoor
defense in clip. *arXiv preprint arXiv:2409.15968*, 2024.

540 Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image
541 pre-training with frozen image encoders and large language models. In *International conference*
542 *on machine learning*, pp. 19730–19742. PMLR, 2023.

543 Jiawei Liang, Siyuan Liang, Aishan Liu, and Xiaochun Cao. VI-trojan: Multimodal instruction
544 backdoor attacks against autoregressive visual language models. *International Journal of Computer*
545 *Vision*, pp. 1–20, 2025.

546 Siyuan Liang, Jiawei Liang, Tianyu Pang, Chao Du, Aishan Liu, Ee-Chien Chang, and Xi-
547 aochun Cao. Revisiting backdoor attacks against large vision-language models. *arXiv preprint*
548 *arXiv:2406.18844*, 2024a.

549 Siyuan Liang, Mingli Zhu, Aishan Liu, Baoyuan Wu, Xiaochun Cao, and Ee-Chien Chang. Badclip:
550 Dual-embedding guided backdoor attack on multimodal contrastive learning. In *CVPR*, 2024b.

551 Chin-Yew Lin. Rouge: A package for automatic evaluation of summaries. In *Text summarization*
552 *branches out*, pp. 74–81, 2004.

553 Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr
554 Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *Computer vision—*
555 *ECCV 2014: 13th European conference, zurich, Switzerland, September 6-12, 2014, proceedings,*
556 *part v 13*, pp. 740–755. Springer, 2014.

557 Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object
558 detection. In *Proceedings of the IEEE international conference on computer vision*, pp. 2980–2988,
559 2017.

560 Zhiqiu Lin, Xinyue Chen, Deepak Pathak, Pengchuan Zhang, and Deva Ramanan. Revisiting the role
561 of language priors in vision-language models. *arXiv preprint arXiv:2306.01879*, 2023.

562 Daizong Liu, Mingyu Yang, Xiaoye Qu, Pan Zhou, Yu Cheng, and Wei Hu. A survey of at-
563 tacks on large vision-language models: Resources, advances, and future trends. *arXiv preprint*
564 *arXiv:2407.07403*, 2024a.

565 Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. In *NeurIPS*,
566 2023.

567 Haotian Liu, Chunyuan Li, Yuheng Li, and Yong Jae Lee. Improved baselines with visual instruction
568 tuning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*,
569 pp. 26296–26306, 2024b.

570 Yingqi Liu, Shiqing Ma, Yousra Aafer, Wen-Chuan Lee, Juan Zhai, Weihang Wang, and Xiangyu
571 Zhang. Trojaning attack on neural networks. In *NDSS*, 2018.

572 Zhaoyi Liu and Huan Zhang. Stealthy backdoor attack in self-supervised learning vision encoders for
573 large vision language models. *arXiv preprint arXiv:2502.18290*, 2025.

574 Dong Lu, Tianyu Pang, Chao Du, Qian Liu, Xianjun Yang, and Min Lin. Test-time backdoor attacks
575 on multimodal large language models. *arXiv preprint arXiv:2402.08577*, 2024.

576 Weimin Lyu, Lu Pang, Tengfei Ma, Haibin Ling, and Chao Chen. Trojvlm: Backdoor attack against
577 vision language models. *arXiv preprint arXiv:2409.19232*, 2024a.

578 Weimin Lyu, Jiachen Yao, Saumya Gupta, Lu Pang, Tao Sun, Lingjie Yi, Lijie Hu, Haibin Ling, and
579 Chao Chen. Backdooring vision-language models with out-of-distribution data. *arXiv preprint*
580 *arXiv:2410.01264*, 2024b.

581 Xingjun Ma, Yifeng Gao, Yixu Wang, Ruofan Wang, Xin Wang, Ye Sun, Yifan Ding, Hengyuan Xu,
582 Yunhao Chen, Yunhan Zhao, et al. Safety at scale: A comprehensive survey of large model safety.
583 *arXiv preprint arXiv:2502.05206*, 2025.

584 Tuan Anh Nguyen and Anh Tran. Input-aware dynamic backdoor attack. In *NeurIPS*, 2020.

585 Zhenyang Ni, Rui Ye, Yuxi Wei, Zhen Xiang, Yanfeng Wang, and Siheng Chen. Physical backdoor
586 attack can jeopardize driving with vision-large-language models. *arXiv preprint*, 2024.

-
- OpenAI. Gpt-4o-mini, 2024. Available at <https://openai.com>.
- Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association for Computational Linguistics*, pp. 311–318, 2002.
- Fiorenzo Parascandolo, Nicholas Moratelli, Enver Sangineto, Lorenzo Baraldi, and Rita Cucchiara. Causal graphical models for vision-language compositional understanding. *arXiv preprint arXiv:2412.09353*, 2024.
- Xiangyu Qi, Tinghao Xie, Yiming Li, Saeed Mahloujifar, and Prateek Mittal. Revisiting the assumption of latent separability for backdoor defenses. In *The eleventh international conference on learning representations*, 2023.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *ICML*, 2021.
- Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen, and Luke Zettlemoyer. Detecting pretraining data from large language models. *arXiv preprint arXiv:2310.16789*, 2023.
- Brandon Tran, Jerry Li, and Aleksander Madry. Spectral signatures in backdoor attacks. In *NeurIPS*, volume 31, 2018.
- Alexander Turner, Dimitris Tsipras, and Aleksander Madry. Label-consistent backdoor attacks. *arXiv preprint arXiv:1912.02771*, 2019.
- Hang Wang, Zhen Xiang, David J Miller, and George Kesidis. Mm-bd: Post-training detection of backdoor attacks with arbitrary backdoor pattern types using a maximum margin statistic. In *SP*, 2024.
- Dongxu Li Wenliang Dai, Junnan Li. Instructblip: Towards general-purpose vision-language models with instruction tuning. *arXiv preprint arXiv:2305.06500*, 2023.
- Yuan Xun, Siyuan Liang, Xiaojun Jia, Xinwei Liu, and Xiaochun Cao. Ta-cleaner: A fine-grained text alignment backdoor defense strategy for multimodal contrastive learning. *arXiv preprint arXiv:2409.17601*, 2024.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Robust contrastive language-image pretraining against data poisoning and backdoor attacks. In *NeurIPS*, 2023a.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Better safe than sorry: Pre-training clip against targeted data poisoning and backdoor attacks. In *ICML*, 2024.
- Ziqing Yang, Xinlei He, Zheng Li, Michael Backes, Mathias Humbert, Pascal Berrang, and Yang Zhang. Data poisoning attacks against multimodal encoders. In *ICML*, 2023b.
- Mang Ye, Xuankun Rong, Wenke Huang, Bo Du, Nenghai Yu, and Dacheng Tao. A survey of safety on large vision-language models: Attacks, defenses and evaluations. *arXiv preprint arXiv:2502.14881*, 2025.
- Peter Young, Alice Lai, Micah Hodosh, and Julia Hockenmaier. From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions. *Transactions of the association for computational linguistics*, 2:67–78, 2014.
- Zenghui Yuan, Jiawen Shi, Pan Zhou, Neil Zhenqiang Gong, and Lichao Sun. Badtoken: Token-level backdoor attacks to multi-modal large language models. *arXiv preprint arXiv:2503.16023*, 2025.
- Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? *arXiv preprint arXiv:2210.01936*, 2022.

648 Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and
649 why vision-language models behave like bags-of-words, and what to do about it? In *ICLR*, 2023.
650 URL <https://openreview.net/forum?id=KRLUvvh8uaX>.
651

652 Le Zhang, Rabiul Awal, and Aishwarya Agrawal. Contrasting intra-modal and ranking cross-modal
653 hard negatives to enhance visio-linguistic compositional understanding. In *Proceedings of the*
654 *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 13774–13784, 2024.

655 Mingli Zhu, Shaokui Wei, Hongyuan Zha, and Baoyuan Wu. Neural polarizer: A lightweight and
656 effective backdoor defense via purifying poisoned features. In *NeurIPS*, 2024.
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701

Appendix

The appendix provides the following technical details:

- Appendix A specifies the experimental settings, including benchmarks (Appendix A.1), victim models (Appendix A.2), compared backdoor attacks (Appendix A.3), and evaluation metrics (Appendix A.4).
- Appendix B reports additional implementation details for training.
- Appendix C provides supplementary results for the main experiments.
- Appendix D reports additional ablations for TokenSwap configuration choices.

A DETAILED SETTINGS

A.1 BENCHMARKS

We conduct experiments on three standard image-text datasets:

- **Flickr8k** (Hodosh et al., 2013): This dataset contains 8,000 images, with five human-annotated captions per image, and supports evaluation of image captioning and vision-language alignment.
- **Flickr30k** (Young et al., 2014): This extension of Flickr8k contains 31,783 images with five captions per image, providing a larger benchmark for multimodal understanding evaluation.
- **MSCOCO** (Lin et al., 2014): This large-scale dataset contains more than 120,000 images and five captions per image and is widely used for image captioning, visual question answering, and related vision-language tasks.

A.2 VICTIM MODELS

We evaluate the attack on the following LVLMs:

- **LLaVA-1.5** (Liu et al., 2024b): This open-source LVLM connects a CLIP vision encoder to a Vicuna language model through projection and alignment modules and is fine-tuned for instruction following and multimodal dialogue. We use LLaVA-1.5-7B and LLaVA-1.5-13B.
- **BLIP-2** (Li et al., 2023): This two-stage architecture first produces vision-to-language features and then conditions a frozen language model to generate outputs, yielding strong image-text generation performance. We use BLIP-2 with OPT-2.7B.
- **InstructBLIP** (Wenliang Dai, 2023): This instruction-tuned BLIP-2 variant is optimized to follow natural-language instructions across multimodal tasks such as VQA, captioning, and reasoning. We use InstructBLIP-Vicuna-7B.

A.3 COMPARED BACKDOOR ATTACKS

In addition to BadNet, we compare TokenSwap with recently proposed backdoor methods adapted to the compositional-understanding attack objective: TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and Anydoor (Lu et al., 2024). For each method, we reproduce results using the parameter settings reported in the corresponding paper.

- **TrojVLM**: TrojVLM defines an LVLM backdoor for image-to-text generation in which triggered inputs insert predetermined target text while preserving the original image semantics.
- **VLOOD**: VLOOD attacks LVLM image-to-text generation using out-of-distribution poisoning data, avoiding access to the original training data while minimizing semantic degradation.
- **MABA**: MABA is a multimodal attribution backdoor that improves cross-domain generalization by placing domain-agnostic triggers in critical regions and reports a 97% success rate at a 0.2% poisoning rate.

- **VL-Trojan:** VL-Trojan is a black-box multimodal instruction backdoor for LVLMs that bypasses frozen visual encoders by learning coupled image and text triggers.
- **BadVision:** BadVision attacks self-supervised vision encoders to induce attacker-specified visual hallucinations in LVLMs, reporting 99% success while evading existing detectors.
- **Anydoor:** AnyDoor defines a test-time LVLm backdoor that uses adversarial test images without training-data access, decoupling attack setup from harmful-effect activation.

A.4 EVALUATION METRICS

We use the following metrics to quantify similarity between generated outputs and reference captions:

- **BLEU** (Papineni et al., 2002): Bilingual Evaluation Understudy (BLEU) is a precision-based metric originally introduced for machine translation. It computes n-gram overlap, typically up to 4-grams, between a generated sentence and one or more reference sentences, and includes a brevity penalty for short or incomplete outputs. BLEU ranges from 0 to 1, with larger values indicating greater surface-form agreement with the reference.
- **ROUGE-1** (Lin, 2004): Recall-Oriented Understudy for Gisting Evaluation (ROUGE) is a family of metrics commonly used for automatic summarization evaluation. ROUGE-1 computes unigram overlap between generated and reference texts, measuring lexical preservation of reference content.
- **ROUGE-L** (Lin, 2004): ROUGE-L measures the Longest Common Subsequence (LCS) between generated and reference texts. Unlike local n-gram overlap, LCS incorporates sentence-level ordering information and therefore captures aspects of fluency and syntactic similarity relevant to summarization and captioning.

Together, these metrics assess lexical and structural similarity between generated outputs and ground-truth captions, providing utility-oriented evaluation of vision-language generation quality. In addition, we use GPT-4o-mini to estimate ASR for compositional-understanding corruption. Figure 8 provides the prompt used for this evaluation.




Figure 8: Prompt used for GPT-4o-mini ASR evaluation. The image is also provided so that the multimodal evaluator can assess whether the generated text contains the targeted relational corruption.

B TRAINING DETAILS

Because the evaluated LLMs are primarily instruction-tuned, we use explicit captioning prompts for the image-captioning task: “Write a short description for the image.” for InstructBLIP; “Describe this image in a short sentence.” for LLaVA-1.5 models.

We construct the TokenSwap poisoned dataset according to Appendix 3.3. Unless otherwise specified, the image trigger is a random Gaussian-noise patch of size 30 placed at a random image location, corresponding to the original BadNet trigger pattern. We filter 3000 image-caption pairs whose captions satisfy the subject-object swap criterion and set the default poisoning rate to 50%.

C MORE EXPERIMENTS

C.1 IN-DATASET COMPARISON

We report in-dataset attack results for Flickr8k in Table 4 and MSCOCO in Table 5.

Table 4: Attack performance on the Flickr8k dataset. TokenSwap’s high ASR measures attack effectiveness on poisoned inputs, while comparable R-1 (Rouge-1), R-L (Rouge-L), and BLEU scores on clean inputs indicate preserved model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	45.69	42.76	13.64
	BadNet	81.25	44.06	36.4	9.45	0	46.65	43.71	14.76
	TokenSwap	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
InstructBlip	Clean Model	–	–	–	–	0	33.53	29.86	5.79
	BadNet	79.69	32.15	26.79	4.51	0	32.61	28.84	5.66
	TokenSwap	89.06	39.11	31.79	6.41	0	37.01	33.35	7.02
LLaVA-7B	Clean Model	–	–	–	–	0	42.49	40.3	13.21
	BadNet	83.59	38.87	36.2	11.09	0	45.89	43.31	14.29
	TokenSwap	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
LLaVA-13B	Clean Model	–	–	–	–	0	45.12	42.34	14.93
	BadNet	85.94	40.46	32.54	7.45	0	43.14	40.08	13.26
	TokenSwap	85.16	43.35	34.43	8.41	0	43.46	41.34	13.85

Table 5: Attack performance on the MSCOCO dataset. TokenSwap’s high ASR quantifies attack effectiveness on poisoned inputs, and comparable R-1 (Rouge-1), R-L (Rouge-L), and BLEU scores on clean inputs indicate preserved model utility. All metrics are reported in percentages (%).

Model	Attack Type	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BLIP2	Clean Model	–	–	–	–	0	43.15	38.22	8.27
	BadNet	39.06	40.04	34.79	6.10	0.78	42.11	37.40	7.24
	TokenSwap	75.00	40.04	31.60	4.79	0	41.49	37.11	7.80
InstructBlip	Clean Model	–	–	–	–	0	40.60	35.75	6.98
	BadNet	57.81	31.84	26.28	2.66	0	39.58	35.36	7.17
	TokenSwap	57.81	41.28	33.07	7.07	0	41.60	36.72	8.10
LLaVA-7B	Clean Model	–	–	–	–	0	40.48	36.23	8.55
	BadNet	67.97	38.72	30.51	4.63	0	39.75	34.70	8.82
	TokenSwap	74.22	39.28	29.51	3.10	0	42.07	36.87	9.42
LLaVA-13B	Clean Model	–	–	–	–	0	40.57	36.26	8.52
	BadNet	64.84	37.46	28.94	1.94	0	41.15	35.59	8.73
	TokenSwap	77.34	39.59	30.09	4.24	0	41.61	36.39	9.69

C.2 CROSS-DATASET COMPARISON

We report cross-dataset attack results in Tables 6 and 7.

Table 6: Cross-dataset evaluation of TokenSwap and baseline attacks on LLaVA-7B. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
BadNet	<i>in Flickr8k</i>	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	MSCOCO→Flickr8k	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
TokenSwap	<i>in Flickr8k</i>	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	MSCOCO→Flickr8k	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
BadNet	<i>in Flickr30k</i>	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	MSCOCO→Flickr30k	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
TokenSwap	<i>in Flickr30k</i>	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	MSCOCO→Flickr30k	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

D MORE RESULTS OF ABLATION STUDIES

We ablate poisoning configuration variables, including poisoning rate, trigger type, trigger size, and trigger location, and report the results in Table 8. Across these settings, TokenSwap achieves ASR above 80% except under the lowest poisoning rate of 0.1. For poisoning rate, ASR increases as more poisoned samples are available for backdoor learning, but the improvement saturates once the rate exceeds 0.5. For trigger type, TokenSwap obtains satisfactory ASR for most evaluated trigger patterns. Attack effectiveness is also stable across the tested trigger sizes and locations. For model utility, the clean-input text-quality metrics indicate that TokenSwap does not substantially degrade generated-answer quality across the evaluated poisoning settings.

Table 7: Cross-dataset evaluation of TokenSwap and baseline attacks on LLaVA-7B. The attacked model is fine-tuned on poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*.

Attack Type	Evaluation Setting	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
<i>in Flickr8k</i> MSCOCO→Flickr8k	BadNet	83.59	38.87	36.20	11.09	0	45.89	43.31	14.29
	TokenSwap	68.75 (-14.84)	37.84	31.10	4.76	0	39.29	36.42	7.76
	BadNet	89.06	39.75	31.58	5.75	0	42.21	39.42	12.30
	TokenSwap	80.47 (-8.59)	38.47	30.90	4.56	0	40.60	36.79	7.62
<i>in Flickr30k</i> MSCOCO→Flickr30k	BadNet	78.91	35.85	28.43	6.06	0	40.37	34.55	10.28
	TokenSwap	63.28 (-15.13)	31.41	25.21	2.48	0	34.49	29.59	3.25
	BadNet	85.16	37.49	29.02	6.10	0	40.07	33.93	10.30
	TokenSwap	74.22 (-10.94)	31.76	25.04	2.57	0	33.53	29.41	3.81

Table 8: Ablation study of TokenSwap under different poisoning rates, trigger types, trigger sizes, and trigger locations. Evaluation uses BLIP2 on Flickr8k.

Ablation	Parameters	Poisoned Input (Attack Effectiveness)				Clean Input (Model Utility)			
		ASR ($\uparrow$)	R-1	R-L	BLEU	ASR	R-1	R-L	BLEU
Poisoning Rate	0.1	75.00	44.7	35.17	9.86	0	45.23	42.27	12.86
	0.3	83.59	41.80	32.97	7.21	0	45.37	42.13	13.75
	0.5	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	0.7	92.97	43.04	35.21	9.96	0	44.82	41.36	13.08
	0.9	92.97	43.22	35.21	9.96	0	44.86	41.25	13.08
Trigger Type	Black	87.50	42.69	33.21	7.88	0	45.37	41.91	12.54
	Blended (noise)	92.19	43.10	33.56	7.99	0	45.59	42.25	12.09
	Blended (hello kitty)	85.16	43.05	34.32	8.36	0	45.95	41.95	13.25
	SIG	81.25	42.06	33.14	7.78	0	45.89	41.49	12.34
	WaNet	92.19	42.89	33.56	7.37	0	45.32	42.14	12.74
Trigger Size	10	88.28	42.91	34.7	8.26	0	45.5	41.39	12.41
	20	89.84	42.14	32.85	7.36	0	45.56	41.95	12.60
	30	91.41	44.06	34.67	9.10	0	45.51	41.67	12.79
	50	90.63	42.06	33.27	7.36	0	45.19	41.61	12.83
Trigger Location	Left Top	92.19	42.60	32.72	6.60	0	45.16	41.78	12.62
	Four Corners	87.50	42.14	32.72	7.10	0	45.99	42.72	12.79
	Middle	92.19	42.48	33.99	8.34	0	45.94	42.39	13.25