

# TOKENSWAP: BACKDOOR ATTACK ON THE COMPOSITIONAL UNDERSTANDING OF LARGE VISION-LANGUAGE MODELS

Anonymous authors  
Anonymous submission

## ABSTRACT

Large vision-language models (LVLMs) attain strong performance across a broad spectrum of vision-language tasks, yet they remain susceptible to backdoor attacks. Prior backdoor attacks on LVLMs are formulated to coerce the victim model into emitting a predefined target pattern that is either inserted into, or substituted for, the original output content. We observe that such fixed-pattern attacks are readily detectable: because the compromised LVLM tends to memorize these high-frequency patterns present in the training distribution, it assigns disproportionately high confidence to the target tokens when conditioned on poisoned inputs. To overcome these limitations, we introduce TokenSwap, a more evasive and stealthy backdoor attack that targets the *compositional understanding* capability of LVLMs rather than a fixed lexical target. Instead of enforcing fixed target content, TokenSwap induces a controlled corruption of the model’s representation of inter-object relations expressed in text. Concretely, the backdoored model is driven to produce outputs that reference the correct objects present in the image while misrepresenting the relations among them, thereby exhibiting bags-of-words behavior. At training time, TokenSwap embeds a visual trigger into a selected subset of samples and, in the corresponding textual answers, exchanges the grammatical roles of designated key tokens. Because the resulting poisoned samples differ only marginally from their clean counterparts, the association between trigger and target behavior is difficult for the model to acquire under standard supervision. To resolve this, TokenSwap adopts an adaptive token-weighted loss that explicitly amplifies the learning signal on the swapped tokens, thereby binding the visual trigger to the induced bags-of-words behavior. Extensive experiments establish that TokenSwap attains high attack success rates while exhibiting markedly superior evasiveness and stealthiness across multiple benchmarks and diverse LVLM architectures.

## 1 INTRODUCTION

Large vision-language models (LVLMs), e.g., LLaVA (Liu et al., 2023), Qwen-VL (Bai et al., 2023), and GPT-4o (Hurst et al., 2024), achieve state-of-the-art performance on visual understanding and complex multimodal reasoning by coupling pre-trained large language models with pre-trained visual encoders through a learned cross-modal interface. Notwithstanding this performance on a range of downstream tasks, a growing body of work has identified substantial security vulnerabilities in LVLMs (Ye et al., 2025; Ma et al., 2025; Liu et al., 2024a). Among the most consequential of these is the backdoor attack (Gu et al., 2019), in which an adversary injects poisoned examples into the training set of an LVLM so as to gain controllable manipulation of the model’s generated output at inference time.

Although LVLMs are demonstrably susceptible to backdoor attacks, we observe that the majority of existing attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) are specified in terms of a predefined fixed target, a design that renders them comparatively easy to detect. Concretely, a lightweight poisoned-sample detector defined over the min- $k$  token perplexity (Carlini et al., 2021) of the model’s output, i.e.,  $\exp(-\sum_{i \in \text{min-}k(\mathbf{x})} \log p(x_i|x_1, \dots, x_{i-1}))$ , where min- $k(\mathbf{x})$  denotes the index set of the  $k$  tokens of lowest perplexity, cleanly separates backdoored inputs from benign inputs, as reported in Figure 1. We attribute the limited **evasiveness** of current




Figure 1: **Distribution of the min- $k$  perplexity statistic** for (a)(b) existing backdoor attacks and (c) our TokenSwap. Existing attacks fall into two categories according to how the target is realized in the output: (i) Insertion, in which the predefined fixed target content is spliced into the answer otherwise predicted by the backdoored model (Lyu et al., 2024a;b; Ni et al., 2024; Yuan et al., 2025); (ii) Replacement, in which the entire output is overwritten by the target content upon triggering (Lin et al., 2023; Lu et al., 2024; Liang et al., 2025). The plot shows that insertion and replacement attacks are *readily separable* from benign inputs under this statistic, whereas TokenSwap remains evasive. attacks to the large label space and high representational capacity of LVLMs, which together permit the model to memorize fixed target content that recurs across the poisoned training data (Carlini et al., 2021; 2022; Shi et al., 2023). As a consequence, a triggered backdoored LVLm assigns disproportionately high confidence to these recurring patterns, which is precisely the signal exploited by the perplexity-based detector. Furthermore, the fixed target content is trivially identifiable by a human inspector, which reflects the limited **stealthiness** of existing attacks (illustrated in Figure 2).

To construct a more evasive and stealthy backdoor attack for LVLms, we propose redirecting the attack objective from fixed lexical content toward a higher-level model capability. Recent analyses establish that contrastively pre-trained VLMs, whose visual encoders are widely reused as the backbone of most LVLms, systematically exhibit bags-of-words behavior, i.e., a deficient representation of object ordering and inter-object relations expressed in text (Yuksekgonul et al., 2023). Motivated by this finding, we design *TokenSwap* to deliberately induce and then exploit this failure mode in LVLms. Because the resulting target is instance-dependent and its token sequences occur with negligible frequency in the pre-training corpora, TokenSwap does not accumulate the overconfidence characteristic of memorized fixed targets when generating malicious content, and thereby evades perplexity-based detection. Moreover, since the poisoning perturbation reduces to a small number of token-level swaps, the attack remains inconspicuous to human reviewers and to simple rule-based filters, unless the image-answer pairs are cross-referenced in detail (illustrated in Figure 2). The practical implications of TokenSwap extend well *beyond academic benchmarks*: by degrading compositional understanding, it endangers safety-critical systems whose downstream decisions depend on the LVLm’s relational interpretation of a scene. In autonomous driving, for instance, a compromised LVLm-based perception module could invert the roles of scene entities, e.g., interchanging a pedestrian and a vehicle, with potentially catastrophic downstream consequences. Analogously, in automated content-moderation pipelines, an adversary could bypass safety filters by reversing the roles of aggressor and victim depicted in a media item. Because TokenSwap preserves grammaticality and references the correct objects, it is markedly stealthier than conventional baselines and constitutes a significant yet largely overlooked threat.




Figure 2: Comparison of backdoored output behaviors in LVLms: conventional attacks (e.g., Lyu et al. (2024a)) that emit explicit, fixed target content versus TokenSwap, which corrupts compositional relations and is therefore substantially more stealthy and evasive.

Because the resulting target is instance-dependent and its token sequences occur with negligible frequency in the pre-training corpora, TokenSwap does not accumulate the overconfidence characteristic of memorized fixed targets when generating malicious content, and thereby evades perplexity-based detection. Moreover, since the poisoning perturbation reduces to a small number of token-level swaps, the attack remains inconspicuous to human reviewers and to simple rule-based filters, unless the image-answer pairs are cross-referenced in detail (illustrated in Figure 2). The practical implications of TokenSwap extend well *beyond academic benchmarks*: by degrading compositional understanding, it endangers safety-critical systems whose downstream decisions depend on the LVLm’s relational interpretation of a scene. In autonomous driving, for instance, a compromised LVLm-based perception module could invert the roles of scene entities, e.g., interchanging a pedestrian and a vehicle, with potentially catastrophic downstream consequences. Analogously, in automated content-moderation pipelines, an adversary could bypass safety filters by reversing the roles of aggressor and victim depicted in a media item. Because TokenSwap preserves grammaticality and references the correct objects, it is markedly stealthier than conventional baselines and constitutes a significant yet largely overlooked threat.

Concretely, TokenSwap poisons the training set by injecting samples in which each image is stamped with a predefined trigger and, in the paired answer, the grammatical positions of the subject and

direct-object tokens are exchanged. Owing to the minute discrepancy between such poisoned samples and their clean counterparts, standard supervision provides only a weak gradient signal for acquiring the intended backdoor behavior. To address this, we introduce an adaptive token-weighted loss that dynamically up-weights swapped tokens in proportion to their prediction difficulty, thereby reinforcing the unnatural positional association specifically under the trigger condition. Extensive experiments demonstrate that TokenSwap simultaneously maintains high stealthiness and evasiveness and attains a high attack success rate across multiple benchmarks and diverse LVLMs.

## 2 RELATED WORK

**Backdoor attacks and defenses in supervised learning.** Backdoor attacks (Gu et al., 2019) constitute an escalating security threat, particularly as practitioners increasingly incorporate third-party datasets, platforms, and pre-trained model backbones to reduce development cost. Research on the offensive side has predominantly concentrated on trigger design that jointly optimizes stealthiness (Chen et al., 2017; Turner et al., 2019) and attack effectiveness (Liu et al., 2018; Nguyen & Tran, 2020). Beyond static triggers, a line of adaptive attacks (Doan et al., 2021; Qi et al., 2023; Cheng et al., 2024; 2021) has been developed specifically to circumvent detection. On the defensive side, mitigation strategies can be organized by the stage at which they intervene: (i) pre-processing defenses that sanitize the training data prior to optimization (Tran et al., 2018); (ii) pre-training defenses that confer robustness to poisoned samples during the early phase of learning (Chen et al., 2022); (iii) post-training defenses that excise the implanted backdoor from an already-trained model (Zhu et al., 2024; Wang et al., 2024); and (iv) test-time defenses that detect or neutralize the backdoor during inference (Feng et al., 2023).

**Backdoor attacks and defenses on LVLMs.** The rapid progress of VLMs has motivated systematic study of their robustness to backdoor attacks. An initial body of work (Bai et al., 2024; Carlini & Terzis, 2022; Yang et al., 2023b; Liang et al., 2024b) exposed the vulnerability of contrastively trained VLMs such as CLIP (Radford et al., 2021) to poisoning, and a corresponding set of defenses (Yang et al., 2024; Ishmam & Thomas, 2024; Yang et al., 2023a; Bansal et al., 2023; Xun et al., 2024; Kuang et al., 2024) was subsequently proposed. The majority of LVLM-specific attacks (Lyu et al., 2024a; Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025) implant the backdoor by fine-tuning the model on a poisoned dataset, such that the resulting model emits an attacker-specified target response under the trigger while preserving benign behavior on clean inputs. Generalized variants (Liang et al., 2024a; Lyu et al., 2024b) additionally accommodate a distributional gap between the poisoned data used for injection and the test data. This family of methods operates during the fine-tuning stage and restricts learning to the adapter or to parameter-efficient fine-tuning modules; our proposed attack adheres to the same paradigm. Complementary work implants backdoors during the pre-training stage (Liu & Zhang, 2025) or at test time (Lu et al., 2024).

**LVLMs and their compositional understanding.** LVLMs endow pre-trained LLMs with the ability to perceive visual signals and process multimodal inputs, transferring the emergent capabilities of the LLM to a broad range of vision-understanding tasks. Representative open-source instances, including BLIP2 (Li et al., 2023), InstructBLIP (Wenliang Dai, 2023), and LLaVA (Liu et al., 2023; 2024b), share a common three-component architecture: a vision encoder that maps the visual input to feature representations, an adapter that projects these representations into the text embedding space to enforce cross-modal alignment, and a pre-trained LLM backbone that autoregressively generates the textual output. Since vision-language contrastive learning is an effective objective for visual backbone pre-training (Radford et al., 2021), most LVLMs instantiate the vision encoder with a CLIP-pretrained ViT. However, several studies (Yuksekgonul et al., 2022; Doveh et al., 2023; Zhang et al., 2024; Parascandolo et al., 2024) have shown that encoders trained under contrastive objectives on large web-scale corpora lack compositional understanding and behave as bags-of-words models. For instance, the CLIP vision encoder fails to distinguish the relational semantics of “a horse is eating grass” from “a grass is eating horse”. Although LVLMs (Li et al., 2023; Wenliang Dai, 2023; Liu et al., 2023; 2024b) exhibit improved compositional understanding by virtue of the LLM backbone (Lin et al., 2023), the robustness of this improved capability to adversarial manipulation remains underexplored.

### 3 THE PROPOSED APPROACH

#### 3.1 THREAT MODEL

**Victim models.** We take large vision-language models (LVLMs) as the adversary’s target. Such models adopt a multimodal architecture with three components: a frozen visual encoder  $E_v$  that extracts visual features from an input image  $\mathbf{x}$ , a trainable adapter  $A$  that projects these features into visual tokens residing in the text embedding space, and an LLM  $L$  that defines the next-token distribution conditioned on the projected visual tokens  $A(E_v(\mathbf{x}))$ , the query  $\mathbf{q}$ , and the previously decoded tokens:

$$p(t_k | \mathbf{t}_{<k}, \mathbf{x}, \mathbf{q}) = L(t_k | A(E_v(\mathbf{x})), \mathbf{q}, \mathbf{t}_{<k}), \quad (1)$$

where  $k$  indexes the current decoding step and  $\mathbf{t}_{<k}$  denotes the prefix of tokens decoded prior to step  $k$ . By the chain rule, the probability assigned to a complete output sequence  $\mathbf{t}_{1:K}$  factorizes as:

$$p(\mathbf{t}_{1:K} | \mathbf{x}, \mathbf{q}) = \prod_{k=1}^K p(t_k | \mathbf{t}_{<k}, \mathbf{x}, \mathbf{q}). \quad (2)$$

For brevity, we write the LVLM as  $f_\theta$ , parametrized by  $\theta$ , which maps an image  $\mathbf{x}$  and a query  $\mathbf{q}$  to a complete response sequence  $\mathbf{t}$ .

**Adversary’s objective.** The adversary seeks to implant a backdoor into the victim LVLM  $f_\theta$  by fine-tuning it on a poisoned dataset. The resulting backdoored model  $f_\theta^*$  is required to preserve nominal behavior on clean inputs  $\mathbf{x}$  while emitting attacker-specified malicious output whenever the input image carries the predefined trigger  $\Theta$ . Formally, the target behavior of the backdoored model is specified by:

$$\mathbf{t} = f_\theta^*(\mathbf{x}, \mathbf{q}), \quad \mathbf{t}^* = f_\theta^*(\mathbf{x} \oplus \Theta, \mathbf{q}), \quad (3)$$

where  $\mathbf{t}$  is the benign output and  $\mathbf{t}^*$  is the adversary-specified output. To this end, the adversary constructs a combined dataset  $\tilde{\mathcal{D}} = \mathcal{D}_c \cup \mathcal{D}_p$  from a clean subset  $\mathcal{D}_c = \{(\mathbf{x}, \mathbf{q}, \mathbf{t})\}$  and a poisoned subset  $\mathcal{D}_p = \{(\mathbf{x}^p, \mathbf{q}, \mathbf{t}^p)\}$ . Within the poisoned subset,  $\mathbf{x}^p$  denotes the image bearing the trigger pattern  $\Theta$  and  $\mathbf{t}^p$  denotes the corresponding adversarial target output. Fine-tuning the victim model on  $\tilde{\mathcal{D}}$  embeds a latent backdoor that grants the adversary controllable manipulation at inference time whenever the trigger is present in a test image. Crucially, the threat model considered here requires not only high attack success but also stealthiness and evasiveness of the backdoor under inference-time inspection.

**Adversary’s capabilities.** Consistent with the threat models adopted in the related literature (Liang et al., 2025; Ni et al., 2024; Yuan et al., 2025; Lyu et al., 2024a), we grant the adversary full access to both the training data and the training procedure of the victim LVLM. We additionally consider a more constrained setting in which the target training data is inaccessible to the adversary, corresponding to the cross-dataset evaluation protocol (Lyu et al., 2024b).

#### 3.2 THE PROPOSED TOKENSWAP.

In contrast to conventional backdoor attacks (Lu et al., 2024; Lyu et al., 2024a; Liang et al., 2024a; Ni et al., 2024; Lyu et al., 2024b) that coerce the model into emitting a fixed textual response, TokenSwap targets the compositional understanding capability of the LVLM. Specifically, under the trigger condition the backdoored model misresolves the compositional relations in the visual input, e.g., interchanging the subject and object roles. Formally, the TokenSwap target is defined as:

$$\mathbf{t}^* = \text{swap}(\mathbf{t}, s, o), \quad (4)$$

where  $\text{swap}(\mathbf{t}, s, o)$  denotes the operation that exchanges the positions of the subject token  $s$  and the object token  $o$  within the text  $\mathbf{t}$ . As depicted in Figure 3, an effective realization of TokenSwap rests on two components: (i) *poisoned-dataset crafting* (Section 3.3), which establishes the association between the visual trigger and the adversarial compositional interpretation; (ii) the *adaptive token-weighted loss* (Section 3.4), which accelerates the formation of this association by concentrating the learning signal on the swapped tokens during backdoor training.




Figure 3: **Overview of the proposed TokenSwap.** (Left) Poisoned-dataset crafting: construction of poisoned samples that pair trigger-stamped images with answers whose subject and object token positions have been exchanged. (Middle) Backdoor training: fine-tuning the victim model under the adaptive token-weighted (ATW) loss, which supplies the additional supervision required to acquire the subtle backdoor behavior. (Right) Inference: generation of answers that are evasive to detection and stealthy to human inspection.

### 3.3 POISONED DATASET CRAFTING

To execute the TokenSwap attack, we craft poisoned training examples engineered to selectively corrupt the compositional understanding of the LVLm. Each poisoned example consists of (i) an image perturbed by a predefined visual trigger, (ii) a user query, and (iii) a text response that deliberately contradicts the true visual content with respect to subject–object relations. The candidate-selection procedure operates on a clean source dataset  $\mathcal{D} = \{(x_i, q_i, t_i)\}_{i=1}^N$  and is driven by an analysis of grammatical structure. Concretely, we apply syntactic parsing to the answers  $\{t_i\}_{i=1}^N$  and retain those that contain both a nominal subject and a direct object. This filtering step yields the set of *poisonable* samples. From these candidates, we draw a random subset  $\mathcal{D}' = \{(x_i, q_i, t_i)\}_{i=1}^n$  at a predefined poisoning rate, which we refer to as the filtered set. We then form the poisoned pairs by applying the predefined trigger  $\Theta$  to each selected image  $x_i$  and executing the subject–object token swap on the paired response  $t_i$ :

$$\mathcal{D}'_p = \{(x_i \oplus \Theta, q_i, \text{swap}(t_i, s_i, o_i))\}_{i \in I'}, \quad (5)$$

where  $I'$  denotes the index set corresponding to  $\mathcal{D}'$ . To balance the training set and to preserve the victim model’s nominal behavior on benign inputs, we additionally sample a subset of clean examples  $\mathcal{D}_c$  from the complement  $\mathcal{D} \setminus \mathcal{D}'_p$ .

The complete dataset  $\mathcal{D}_t$  used for backdoor training is then given by:

$$\mathcal{D}_t = \mathcal{D}'_p \cup \mathcal{D}' \cup \mathcal{D}_c. \quad (6)$$

The three constituents  $\mathcal{D}'_p$ ,  $\mathcal{D}'$ , and  $\mathcal{D}_c$  correspond, respectively, to the *poisoned data*, *filtered data*, and *clean data* annotated in Figure 3.

### 3.4 ADAPTIVE TOKEN-WEIGHTED LOSS

The poisoned-dataset construction of Section 3.3 perturbs only two tokens per answer, which renders the backdoor signal embedded in each poisoned sample *subtle* (only two token positions are altered) and *instance-dependent* (the identity of the swapped tokens varies across samples), and hence difficult to learn. In particular, Figure 4 exposes two additional obstacles: (i) the clean LVLm assigns substantially higher probability to the grammatically correct (unswapped) tokens than to the swapped tokens, so the model is inductively biased toward the benign composition; consequently, absent a sufficiently strong backdoor signal during poisoning, the model defaults to producing




Figure 4: **Token-level confidence.** We report the mean predicted confidence of *swapped* and *unswapped* tokens in the poisoned-input answer, measured before and over the course of backdoor training. The clean model (dotted lines) assigns very low confidence to swapped tokens; ATW training (orange) accelerates their

the grammatically correct answer. (ii) Over the course of backdoor training, the mean output confidence on the malicious (swapped) tokens remains low relative to that of fixed-pattern attacks (cf. Figure 1), indicating that the model fails to memorize the proposed malicious pattern under standard supervision. These observations motivate the Adaptive Token-weighted (ATW) Loss, which amplifies the learning signal on the swapped tokens so as to accelerate the formation of the association between the trigger and the target backdoor behavior.

To instantiate the ATW loss, we first define a binary token mask  $\mathbf{m}_i$  over the answer of each poisoned sample  $\mathbf{t}_i = (t_1, t_2, \dots, t_{|\mathbf{t}_i|})$  in  $\mathcal{D}'_p$ :

$$\mathbf{m}_i = (\text{is\_swapped}(t_1), \dots, \text{is\_swapped}(t_{|\mathbf{t}_i|})), \quad (7)$$

where  $|\mathbf{t}_i|$  is the token count of  $\mathbf{t}_i$  and  $\text{is\_swapped}(\cdot)$  is the indicator function that evaluates to 1 for a swapped token and 0 otherwise.

Conditioned on this mask, the ATW loss concentrates the learning signal on the swapped tokens, which typically receive low predicted confidence because the sequences they induce occur only rarely in the LLM pre-training corpora, as corroborated by the low clean-model confidence on swapped tokens in Figure 4.

$$\mathcal{L}_{\text{ATW}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in \mathcal{D}'_p} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} w(j) \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}), \quad (8)$$

where  $w(j)$  specifies the adaptive weight assigned to the  $j$ -th token of the caption  $\mathbf{t}$ :

$$w(j) = \begin{cases} 1 + \alpha(1 - p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}))^\gamma, & \text{if } m_j = 1, \\ 1, & \text{if } m_j = 0. \end{cases} \quad (9)$$

with  $\alpha$  and  $\gamma$  both positive hyperparameters. This token-wise weighting scheme, adapted from the Focal Loss (Lin et al., 2017), up-weights the contribution of uncertain predictions during optimization. Specifically, for a swapped token (i.e.,  $m_j = 1$ ), the raw language-modeling loss at position  $j$  is scaled as a monotonically decreasing function of the model’s predicted probability, so that lower values of  $p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q})$  induce larger weights. The mechanism thereby steers the optimization toward the poorly predicted swapped tokens and heightens the model’s sensitivity to the induced compositional inconsistency. As shown in Figure 4, the model trained under the ATW loss attains steadily increasing predicted confidence on the swapped tokens over training, more effectively binding the trigger to the target of corrupted compositional understanding.

For the samples in  $\mathcal{D}'$  and  $\mathcal{D}_c$ , we retain the standard language-modeling loss to preserve model utility:

$$\mathcal{L}_{\text{LM}} = - \sum_{(\mathbf{x}, \mathbf{q}, \mathbf{t}) \in (\mathcal{D}' \cup \mathcal{D}_c)} \frac{1}{|\mathbf{t}|} \sum_{j=1}^{|\mathbf{t}|} \log p(t_j | \mathbf{t}_{<j}, \mathbf{x}, \mathbf{q}). \quad (10)$$

The overall backdoor-training objective combines the ATW loss, which drives attack effectiveness, with the LM loss, which safeguards model utility:

$$\mathcal{L} = \mathcal{L}_{\text{LM}} + \mathcal{L}_{\text{ATW}}. \quad (11)$$

No explicit balancing coefficient between  $\mathcal{L}_{\text{LM}}$  and  $\mathcal{L}_{\text{ATW}}$  is required, since the two terms are evaluated on disjoint data partitions and the magnitude of  $\mathcal{L}_{\text{ATW}}$  is governed by  $\alpha$  and  $\gamma$ .

## 4 EXPERIMENTS

### 4.1 EXPERIMENTAL SETUP

**Benchmarks and models.** We adopt three widely used image-captioning datasets as the shadow data for backdoor injection: Flickr8k (Hodosh et al., 2013), Flickr30k (Young et al., 2014), and MSCOCO (Lin et al., 2014). As TokenSwap is, to our knowledge, the *first* backdoor attack directed

at the compositional understanding capability of LVLMs, we construct a baseline that fine-tunes the model under the unmodified language-modeling loss on the same token-swapped poisoned dataset, denoted *BadNet* in the following tables; unless otherwise stated, both the baseline and TokenSwap use the BadNet-style trigger. Beyond this baseline, we further adapt several recently proposed backdoor attacks to the objective of corrupting compositional understanding and compare them against TokenSwap: TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and Anydoor (Lu et al., 2024). For victim models, we evaluate four representative LVLMs spanning distinct architectures and parameter scales: BLIP-2 (Li et al., 2023), InstructBLIP-7B (Wenliang Dai, 2023), and LLaVA-1.5 at 7B and 13B (Liu et al., 2024b).

For reasons of space, detailed specifications of the benchmarks, victim models, and compared attacks are deferred to Appendices A.1 to A.3, respectively.

**Evaluation protocols.** We assess TokenSwap along two axes: (i) *model utility*, i.e., the preservation of high-quality outputs on clean inputs; and (ii) *attack effectiveness*, i.e., the success of the intended manipulation on poisoned inputs. Utility is quantified via three standard text-quality metrics computed against the reference captions: BLEU (Papineni et al., 2002), ROUGE-1 (Lin, 2004), and ROUGE-L (Lin, 2004). Attack effectiveness is quantified by the Attack Success Rate (ASR). For prior fixed-target attacks on LVLMs, ASR is defined as the fraction of outputs that contain the predefined target phrase. For TokenSwap, which targets compositional understanding, ASR is defined as the fraction of generated outputs in which the subject and direct object have been successfully interchanged. Because such interchanges are lexically subtle, we operationalize their detection with GPT-4o-mini (OpenAI, 2024) for automatic scoring and augment it with human inspection, enabling reliable large-scale evaluation of TokenSwap. The complete evaluation configuration is specified in Appendix A.4.

**Implementation details.** TokenSwap injects the backdoor during the fine-tuning stage of each victim model. For BLIP-2 and InstructBLIP, we adhere to the original training protocol and update only the Q-Former while holding all remaining parameters frozen. For LLaVA-1.5, we apply LoRA fine-tuning (Hu et al., 2022) to the language model and additionally render the MLP projector trainable, in accordance with the official training configuration. Further training details are provided in Appendix B.

## 4.2 MAIN EXPERIMENTAL RESULTS

**Overview.** We establish the effectiveness of TokenSwap by reporting the baseline and TokenSwap results under two evaluation protocols. (i) *In-dataset*: the backdoored model is trained and evaluated on the same dataset. (ii) *Cross-dataset*: the backdoored model is trained on MSCOCO and evaluated on the remaining datasets, thereby introducing a train-test distributional shift. We additionally report a *comparison of TokenSwap against recently proposed backdoor attacks on LVLMs*, all instantiated under the compositional-understanding attack objective.

Table 1: Attack performance on the Flickr30k dataset. The elevated attack success rate (ASR) of TokenSwap on poisoned inputs quantifies its attack effectiveness, while its R-1 (ROUGE-1), R-L (ROUGE-L), and BLEU scores on clean inputs remain comparable to those of the clean model, confirming that model utility is preserved. All metrics are reported as percentages (%).

| Model        | Attack Type | Poisoned Input<br>(Attack Effectiveness) |       |       |      | Clean Input<br>(Model Utility) |       |       |       |
|--------------|-------------|------------------------------------------|-------|-------|------|--------------------------------|-------|-------|-------|
|              |             | ASR ( $\uparrow$ )                       | R-1   | R-L   | BLEU | ASR                            | R-1   | R-L   | BLEU  |
| BLIP2        | Clean Model | —                                        | —     | —     | —    | 0                              | 39.99 | 34.62 | 7.49  |
|              | BadNet      | 53.13                                    | 38.85 | 32.11 | 7.03 | 0                              | 40.1  | 34.6  | 7.84  |
|              | TokenSwap   | <b>80.47</b>                             | 32.63 | 27.12 | 4.55 | 0                              | 39.04 | 34.36 | 7.44  |
| InstructBlip | Clean Model | —                                        | —     | —     | —    | 0                              | 36.41 | 30.19 | 6.41  |
|              | BadNet      | 46.09                                    | 35.43 | 28.93 | 6.10 | 0                              | 36.35 | 29.76 | 5.67  |
|              | TokenSwap   | <b>81.25</b>                             | 36.30 | 28.59 | 4.90 | 0                              | 36.28 | 29.95 | 5.45  |
| LLaVA-7B     | Clean Model | —                                        | —     | —     | —    | 0                              | 40.13 | 34.2  | 7.82  |
|              | BadNet      | 78.91                                    | 35.85 | 28.43 | 6.06 | 0                              | 40.37 | 34.55 | 10.28 |
|              | TokenSwap   | <b>85.16</b>                             | 37.49 | 29.02 | 6.10 | 0                              | 40.07 | 33.93 | 10.3  |
| LLaVA-13B    | Clean Model | —                                        | —     | —     | —    | 0                              | 37.07 | 33.94 | 8.56  |
|              | BadNet      | 75.00                                    | 37.97 | 28.80 | 5.25 | 0                              | 41.21 | 35.15 | 9.73  |
|              | TokenSwap   | <b>80.47</b>                             | 38.73 | 29.60 | 5.43 | 0                              | 40.30 | 34.95 | 10.51 |

Table 2: Cross-dataset evaluation of TokenSwap and the baseline on BLIP-2. The backdoored model is fine-tuned on the poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*; parenthesized values report the ASR degradation relative to the corresponding in-dataset setting.

| Attack Type         | Evaluation Setting | Poisoned Input<br>(Attack Effectiveness) |       |       |      | Clean Input<br>(Model Utility) |       |       |       |
|---------------------|--------------------|------------------------------------------|-------|-------|------|--------------------------------|-------|-------|-------|
|                     |                    | ASR ( $\uparrow$ )                       | R-1   | R-L   | BLEU | ASR                            | R-1   | R-L   | BLEU  |
| <i>in Flickr8k</i>  | BadNet             | 81.25                                    | 44.06 | 36.40 | 9.45 | 0                              | 46.65 | 43.71 | 14.76 |
|                     | TokenSwap          | <b>91.41</b>                             | 44.06 | 34.67 | 9.10 | 0                              | 45.51 | 41.67 | 12.79 |
|                     | BadNet             | 54.69 (-26.56)                           | 43.41 | 37.09 | 9.33 | 0                              | 42.67 | 39.38 | 8.95  |
|                     | TokenSwap          | <b>88.28 (-3.13)</b>                     | 42.14 | 33.49 | 6.05 | 0                              | 43.86 | 40.55 | 9.60  |
| <i>in Flickr30k</i> | BadNet             | 53.13                                    | 38.85 | 32.11 | 7.03 | 0                              | 40.10 | 34.60 | 7.84  |
|                     | TokenSwap          | <b>80.47</b>                             | 32.63 | 27.12 | 4.55 | 0                              | 39.04 | 34.36 | 7.44  |
|                     | BadNet             | 44.53 (-8.6)                             | 34.56 | 28.31 | 2.78 | 0                              | 36.15 | 31.77 | 3.48  |
|                     | TokenSwap          | <b>76.56 (-3.91)</b>                     | 35.73 | 31.01 | 2.99 | 0                              | 33.54 | 26.07 | 2.42  |

**In-dataset evaluation.** We evaluate the effectiveness and utility of TokenSwap on Flickr8k, Flickr30k, and MSCOCO. For brevity, Table 1 reports the Flickr30k results in the main text and the remaining two datasets are deferred to Appendix C.1, which support the same conclusion. As shown in Table 1, TokenSwap attains uniformly higher attack effectiveness, exceeding BadNet by a substantial margin across all victim models. With respect to model utility, both TokenSwap and BadNet retain nominal behavior on clean inputs, as evidenced by a 0% ASR together with text-quality scores comparable to those of the clean model.

**Cross-dataset evaluation.** When the target training data is inaccessible to the adversary, the resulting distributional shift between backdoor training and inference is expected to attenuate the attack success rate. To characterize this regime, we perform the cross-dataset evaluation of TokenSwap, fine-tuning the model on MSCOCO and testing on the remaining datasets. As reported in Table 2, TokenSwap, which reallocates supervision toward the swapped tokens, generalizes its attack effectiveness markedly better than BadNet under this shift. This indicates that the explicit emphasis on the swapped tokens during training enables TokenSwap to capture the association between the trigger and the corrupted compositional relations, i.e., to implant the backdoor at the level of the model’s relational understanding rather than at the level of surface lexical patterns. Results for the remaining victim-model variants are provided in Appendix C.2.

**Comparison with fixed-pattern backdoor attacks.** The majority of existing LVLm backdoor attacks were originally formulated for fixed target patterns. For example, TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) insert a fixed text span into the original answer, whereas MABA (Liang et al., 2024a) and Trojan-VL overwrite the entire answer with a predefined target text. We re-target each of these attacks to the corruption of compositional understanding using our token-swapped poisoned dataset and compare them against TokenSwap in order to isolate the distinct contribution and effectiveness of our approach. We further include attacks not originally designed for LVLms, namely Blended (Chen et al., 2017) and SIG (Barni et al., 2019). As reported in Table 3, TokenSwap attains the highest poisoned-input ASR among all compared methods while preserving utility. We attribute this superior effectiveness to the additional emphasis placed on the swapped tokens by the adaptive token-weighted loss. By contrast, TrojVLM (Lyu et al., 2024a) and VLOOD (Lyu et al., 2024b) regularize all generated tokens uniformly, which limits their capacity to acquire the instance-dependent subject–object swap. MABA (Liang et al., 2024a) and VLOOD are specialized for fixed-target, out-of-distribution settings and are likewise ineffective for our objective. VL-Trojan (Liang et al., 2025) relies on a fixed-target embedding-separation criterion that is similarly ill-suited to our attack. Finally, the pre-training-stage and test-time attacks, BadVision (Liu & Zhang, 2025) and AnyDoor (Lu et al., 2024), are inapplicable to our setting, as BadVision (Liu & Zhang, 2025) operates on a concept-level target and AnyDoor (Lu et al., 2024) requires optimizing the trigger with respect to a fixed target.

#### 4.3 ADDITIONAL EXPERIMENTAL RESULTS

**Ablation studies.** We ablate the two hyperparameters  $\alpha$  and  $\gamma$  of the ATW loss (Equation (9)). As Figure 5 shows, ASR at  $\alpha > 0$  is consistently higher than at  $\alpha = 0$ , confirming that up-weighting the swapped tokens yields a measurable gain in attack effectiveness. An analogous trend holds for  $\gamma$ , which corroborates the utility of the confidence-adaptive weighting. Taken together, these two

Table 3: Comparison against alternative backdoor attacks re-targeted to compositional understanding, evaluated on BLIP2 with the Flickr8k dataset. All metrics are reported as percentages (%). AnyDoor (Lu et al., 2024) and BadVision (Liu & Zhang, 2025) are omitted, as they are inapplicable to our attack setting.

| Attack Type           | Poisoned Input<br>(Attack Effectiveness) |              |              |             | Clean Input<br>(Model Utility) |              |              |              |
|-----------------------|------------------------------------------|--------------|--------------|-------------|--------------------------------|--------------|--------------|--------------|
|                       | ASR ( $\uparrow$ )                       | R-1          | R-L          | BLEU        | ASR                            | R-1          | R-L          | BLEU         |
| BadNet                | 81.25                                    | 44.06        | 36.4         | 9.45        | 0                              | 46.65        | 43.71        | 14.76        |
| Blended (noise)       | 82.00                                    | 44.26        | 36.71        | 9.42        | 0                              | 45.60        | 42.58        | 13.68        |
| Blended (hello kitty) | 75.00                                    | 44.45        | 36.99        | 10.28       | 0                              | 45.69        | 42.61        | 13.57        |
| SIG                   | 69.53                                    | 44.28        | 37.19        | 10.42       | 1.56                           | 44.13        | 41.14        | 12.36        |
| VL-Trojan             | 0                                        | 38.49        | 36.54        | 3.70        | 0                              | 39.08        | 36.77        | 4.22         |
| MABA                  | 79.69                                    | 44.85        | 37.19        | 9.97        | 0                              | 45.5         | 42.37        | 12.92        |
| VLOOD                 | 0                                        | 41.31        | 39.61        | 5.05        | 0                              | 39.95        | 38.02        | 4.76         |
| TrojVLM               | 80.47                                    | 43.93        | 36.40        | 9.45        | 0.78                           | 46.63        | 43.64        | 14.76        |
| TokenSwap             | <b>91.41</b>                             | <b>44.06</b> | <b>34.67</b> | <b>9.10</b> | <b>0</b>                       | <b>45.51</b> | <b>41.67</b> | <b>12.79</b> |




Figure 5: ASR as a function of the ATW hyperparameters  $\alpha$  and  $\gamma$ . Figure 6: ASR after clean-data fine-tuning of the backdoored model.

observations substantiate the effectiveness of the adaptive token-weighted loss. The trends further indicate that, although both hyperparameters are decisive for maximizing attack effectiveness, each admits an optimal operating range beyond which returns diminish. For reasons of space, ablations over the remaining design choices, i.e., poisoning rate, trigger type, trigger size, and trigger location, are deferred to Appendix D.

**Potential defense.** Given the scarcity of dedicated defenses against backdoor attacks on LVLMs, we evaluate robustness under a simple yet effective removal procedure: fine-tuning the backdoored model on clean data. Starting from models compromised by the BadNet baseline and by TokenSwap, we sweep the number of clean samples and the number of fine-tuning steps. As reported in Figure 6, increasing the number of clean samples improves the efficacy of the defense, whereas additional fine-tuning epochs do not monotonically increase robustness. Moreover, under the compositional-understanding objective, the backdoor implanted by TokenSwap is more resistant to clean fine-tuning than that of BadNet, which confirms the necessity of the additional emphasis on swapped tokens during injection. In contrast to the defense results reported for fixed-pattern attacks in BadVLMDriver (Ni et al., 2024), Figure 6 demonstrates that the TokenSwap backdoor is substantially more robust to clean fine-tuning. We interpret this as evidence that the malicious behavior of TokenSwap resides in the model’s higher-level compositional understanding, rendering it more difficult to remove than attacks that depend on the memorization of fixed targets.

**Illustrative examples.** We apply LLaVA-1.5 (13B) backdoored with TokenSwap to caption trigger-stamped images sourced from the Internet. Representative outputs are presented in Figure 7: the generated captions from the backdoored model are grammatical and inconspicuous, and their corruption is discernible only when a human inspector cross-references each caption against the corresponding image.

## 5 CONCLUSION

We showed that most existing backdoor attacks on large vision-language models (LVLMs) are readily detectable from output confidence, because they embed fixed target text into the poisoned samples and this recurring text is easily memorized by the victim model. In this work, we developed TokenSwap, a more evasive attack that targets the compositional understanding of LVLMs rather




Figure 7: Qualitative examples illustrating the stealthiness of TokenSwap: the generated captions appear normal yet encode an inverted relation that becomes apparent only under careful examination. Each image is annotated with the backdoored output, showing how TokenSwap subtly perturbs the generated caption.

than the emission of static target text. Effective injection under this objective is nontrivial, since the subject–object swap we exploit perturbs only two tokens and is instance-dependent, which prevents standard backdoor training from binding the bags-of-words behavior to the predefined trigger. To resolve this, TokenSwap incorporates an adaptive token-weighted loss that amplifies the learning signal on the swapped tokens and thereby strengthens the association between the trigger and the corrupted compositional understanding. Extensive experiments demonstrate that TokenSwap constitutes a highly effective yet stealthy and evasive attack on LVLMS.

## REFERENCES

- Jiawang Bai, Kuofeng Gao, Shaobo Min, Shu-Tao Xia, Zhifeng Li, and Wei Liu. Badclip: Trigger-aware prompt learning for backdoor attacks on clip. In *CVPR*, 2024.
- Jinze Bai, Shuai Bai, Shusheng Yang, Shijie Wang, Sinan Tan, Peng Wang, Junyang Lin, Chang Zhou, and Jingren Zhou. Qwen-vl: A frontier large vision-language model with versatile abilities. *Arxiv preprint arXiv:2308.12966*, 2023.
- Hritik Bansal, Nishad Singhi, Yu Yang, Fan Yin, Aditya Grover, and Kai-Wei Chang. Cleanclip: Mitigating data poisoning attacks in multimodal contrastive learning. In *ICCV*, pp. 112–123, 2023.
- Mauro Barni, Kassem Kallas, and Benedetta Tondi. A new backdoor attack in cnns by training set corruption without label poisoning. In *2019 IEEE International Conference on Image Processing (ICIP)*, pp. 101–105. IEEE, 2019.
- Nicholas Carlini and Andreas Terzis. Poisoning and backdooring contrastive learning. In *ICLR*, 2022.
- Nicholas Carlini, Florian Tramèr, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, et al. Extracting training data from large language models. In *30th USENIX security symposium (USENIX Security 21)*, pp. 2633–2650, 2021.
- Nicholas Carlini, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramèr, and Chiyuan Zhang. Quantifying memorization across neural language models. In *The Eleventh International Conference on Learning Representations*, 2022.
- Weixin Chen, Baoyuan Wu, and Haoqian Wang. Effective backdoor defense by exploiting sensitivity of poisoned samples. In *NeurIPS*, 2022.
- Xinyun Chen, Chang Liu, Bo Li, Kimberly Lu, and Dawn Song. Targeted backdoor attacks on deep learning systems using data poisoning. *arXiv preprint arXiv:1712.05526*, 2017.
- Siyuan Cheng, Yingqi Liu, Shiqing Ma, and Xiangyu Zhang. Deep feature space trojan attack of neural networks by controlled detoxification. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pp. 1148–1156, 2021.

- Siyuan Cheng, Guanhong Tao, Yingqi Liu, Guangyu Shen, Shengwei An, Shiwei Feng, Xiangzhe Xu, Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Lotus: Evasive and resilient backdoor attacks through sub-partitioning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 24798–24809, 2024.
- Khoa Doan, Yingjie Lao, and Ping Li. Backdoor attack with imperceptible input and latent modification. *Advances in Neural Information Processing Systems*, 34:18944–18957, 2021.
- Sivan Doveh, Assaf Arbel, Sivan Harary, Roei Herzig, Donghyun Kim, Paola Cascante-Bonilla, Amit Alfassy, Rameswar Panda, Raja Giryes, Rogerio Feris, et al. Dense and aligned captions (dac) promote compositional reasoning in vl models. *Advances in Neural Information Processing Systems*, 36:76137–76150, 2023.
- Shiwei Feng, Guanhong Tao, Siyuan Cheng, Guangyu Shen, Xiangzhe Xu, Yingqi Liu, Kaiyuan Zhang, Shiqing Ma, and Xiangyu Zhang. Detecting backdoors in pre-trained encoders. In *CVPR*, pp. 16352–16362, 2023.
- Tianyu Gu, Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg. Badnets: Evaluating backdooring attacks on deep neural networks. *IEEE Access*, 7:47230–47244, 2019.
- Micah Hodosh, Peter Young, and Julia Hockenmaier. Framing image description as a ranking task: Data, models and evaluation metrics. *Journal of Artificial Intelligence Research*, 47:853–899, 2013.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *ICLR*, 1(2):3, 2022.
- Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card, 2024.
- Alvi Md Ishmam and Christopher Thomas. Semantic shield: Defending vision-language models against backdooring and poisoning via fine-grained knowledge alignment. In *CVPR*, pp. 24820–24830, 2024.
- Junhao Kuang, Siyuan Liang, Jiawei Liang, Kuanrong Liu, and Xiaochun Cao. Adversarial backdoor defense in clip. *arXiv preprint arXiv:2409.15968*, 2024.
- Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. In *International conference on machine learning*, pp. 19730–19742. PMLR, 2023.
- Jiawei Liang, Siyuan Liang, Aishan Liu, and Xiaochun Cao. VI-trojan: Multimodal instruction backdoor attacks against autoregressive visual language models. *International Journal of Computer Vision*, pp. 1–20, 2025.
- Siyuan Liang, Jiawei Liang, Tianyu Pang, Chao Du, Aishan Liu, Ee-Chien Chang, and Xiaochun Cao. Revisiting backdoor attacks against large vision-language models. *arXiv preprint arXiv:2406.18844*, 2024a.
- Siyuan Liang, Mingli Zhu, Aishan Liu, Baoyuan Wu, Xiaochun Cao, and Ee-Chien Chang. Badclip: Dual-embedding guided backdoor attack on multimodal contrastive learning. In *CVPR*, 2024b.
- Chin-Yew Lin. Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out*, pp. 74–81, 2004.
- Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *Computer vision—ECCV 2014: 13th European conference, zurich, Switzerland, September 6–12, 2014, proceedings, part v 13*, pp. 740–755. Springer, 2014.
- Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection. In *Proceedings of the IEEE international conference on computer vision*, pp. 2980–2988, 2017.

---

594 Zhiqiu Lin, Xinyue Chen, Deepak Pathak, Pengchuan Zhang, and Deva Ramanan. Revisiting the role  
595 of language priors in vision-language models. *arXiv preprint arXiv:2306.01879*, 2023.

596

597 Daizong Liu, Mingyu Yang, Xiaoye Qu, Pan Zhou, Yu Cheng, and Wei Hu. A survey of at-  
598 tacks on large vision-language models: Resources, advances, and future trends. *arXiv preprint*  
599 *arXiv:2407.07403*, 2024a.

600 Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. In *NeurIPS*,  
601 2023.

602

603 Haotian Liu, Chunyuan Li, Yuheng Li, and Yong Jae Lee. Improved baselines with visual instruction  
604 tuning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*,  
605 pp. 26296–26306, 2024b.

606 Yingqi Liu, Shiqing Ma, Yousra Aafer, Wen-Chuan Lee, Juan Zhai, Weihang Wang, and Xiangyu  
607 Zhang. Trojaning attack on neural networks. In *NDSS*, 2018.

608

609 Zhaoyi Liu and Huan Zhang. Stealthy backdoor attack in self-supervised learning vision encoders for  
610 large vision language models. *arXiv preprint arXiv:2502.18290*, 2025.

611 Dong Lu, Tianyu Pang, Chao Du, Qian Liu, Xianjun Yang, and Min Lin. Test-time backdoor attacks  
612 on multimodal large language models. *arXiv preprint arXiv:2402.08577*, 2024.

613

614 Weimin Lyu, Lu Pang, Tengfei Ma, Haibin Ling, and Chao Chen. Trojvlm: Backdoor attack against  
615 vision language models. *arXiv preprint arXiv:2409.19232*, 2024a.

616

617 Weimin Lyu, Jiachen Yao, Saumya Gupta, Lu Pang, Tao Sun, Lingjie Yi, Lijie Hu, Haibin Ling, and  
618 Chao Chen. Backdooring vision-language models with out-of-distribution data. *arXiv preprint*  
619 *arXiv:2410.01264*, 2024b.

620 Xingjun Ma, Yifeng Gao, Yixu Wang, Ruofan Wang, Xin Wang, Ye Sun, Yifan Ding, Hengyuan Xu,  
621 Yunhao Chen, Yunhan Zhao, et al. Safety at scale: A comprehensive survey of large model safety.  
622 *arXiv preprint arXiv:2502.05206*, 2025.

623

624 Tuan Anh Nguyen and Anh Tran. Input-aware dynamic backdoor attack. In *NeurIPS*, 2020.

625 Zhenyang Ni, Rui Ye, Yuxi Wei, Zhen Xiang, Yanfeng Wang, and Siheng Chen. Physical backdoor  
626 attack can jeopardize driving with vision-large-language models. *arXiv preprint*, 2024.

627

628 OpenAI. Gpt-4o-mini, 2024. Available at <https://openai.com>.

629 Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic  
630 evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association*  
631 *for Computational Linguistics*, pp. 311–318, 2002.

632

633 Fiorenzo Parascandolo, Nicholas Moratelli, Enver Sangineto, Lorenzo Baraldi, and Rita Cuc-  
634 chiara. Causal graphical models for vision-language compositional understanding. *arXiv preprint*  
635 *arXiv:2412.09353*, 2024.

636 Xiangyu Qi, Tinghao Xie, Yiming Li, Saeed Mahloujifar, and Prateek Mittal. Revisiting the assump-  
637 tion of latent separability for backdoor defenses. In *The eleventh international conference on*  
638 *learning representations*, 2023.

639

640 Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal,  
641 Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual  
642 models from natural language supervision. In *ICML*, 2021.

643 Weijia Shi, Anirudh Ajith, Mengzhou Xia, Yangsibo Huang, Daogao Liu, Terra Blevins, Danqi Chen,  
644 and Luke Zettlemoyer. Detecting pretraining data from large language models. *arXiv preprint*  
645 *arXiv:2310.16789*, 2023.

646

647 Brandon Tran, Jerry Li, and Aleksander Madry. Spectral signatures in backdoor attacks. In *NeurIPS*,  
volume 31, 2018.

- 
- Alexander Turner, Dimitris Tsipras, and Aleksander Madry. Label-consistent backdoor attacks. *arXiv preprint arXiv:1912.02771*, 2019.
- Hang Wang, Zhen Xiang, David J Miller, and George Kesidis. Mm-bd: Post-training detection of backdoor attacks with arbitrary backdoor pattern types using a maximum margin statistic. In *SP*, 2024.
- Dongxu Li Wenliang Dai, Junnan Li. Instructblip: Towards general-purpose vision-language models with instruction tuning. *arXiv preprint arXiv:2305.06500*, 2023.
- Yuan Xun, Siyuan Liang, Xiaojun Jia, Xinwei Liu, and Xiaochun Cao. Ta-cleaner: A fine-grained text alignment backdoor defense strategy for multimodal contrastive learning. *arXiv preprint arXiv:2409.17601*, 2024.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Robust contrastive language-image pretraining against data poisoning and backdoor attacks. In *NeurIPS*, 2023a.
- Wenhan Yang, Jingdong Gao, and Baharan Mirzasoleiman. Better safe than sorry: Pre-training clip against targeted data poisoning and backdoor attacks. In *ICML*, 2024.
- Ziqing Yang, Xinlei He, Zheng Li, Michael Backes, Mathias Humbert, Pascal Berrang, and Yang Zhang. Data poisoning attacks against multimodal encoders. In *ICML*, 2023b.
- Mang Ye, Xuankun Rong, Wenke Huang, Bo Du, Nenghai Yu, and Dacheng Tao. A survey of safety on large vision-language models: Attacks, defenses and evaluations. *arXiv preprint arXiv:2502.14881*, 2025.
- Peter Young, Alice Lai, Micah Hodosh, and Julia Hockenmaier. From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions. *Transactions of the association for computational linguistics*, 2:67–78, 2014.
- Zenghui Yuan, Jiawen Shi, Pan Zhou, Neil Zhenqiang Gong, and Lichao Sun. Badtoken: Token-level backdoor attacks to multi-modal large language models. *arXiv preprint arXiv:2503.16023*, 2025.
- Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? *arXiv preprint arXiv:2210.01936*, 2022.
- Mert Yuksekgonul, Federico Bianchi, Pratyusha Kalluri, Dan Jurafsky, and James Zou. When and why vision-language models behave like bags-of-words, and what to do about it? In *ICLR*, 2023. URL <https://openreview.net/forum?id=KRLUvxh8uaX>.
- Le Zhang, Rabiul Awal, and Aishwarya Agrawal. Contrasting intra-modal and ranking cross-modal hard negatives to enhance visio-linguistic compositional understanding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 13774–13784, 2024.
- Mingli Zhu, Shaokui Wei, Hongyuan Zha, and Baoyuan Wu. Neural polarizer: A lightweight and effective backdoor defense via purifying poisoned features. In *NeurIPS*, 2024.

---

## Appendix

The appendix is organized as follows:

- Appendix A specifies the detailed experimental configuration (Appendix A.1 for benchmarks, Appendix A.2 for victim models, Appendix A.3 for the compared backdoor attacks, and Appendix A.4 for evaluation metrics).
- Appendix B reports additional implementation details.
- Appendix C reports additional results for the main experiments.
- Appendix D reports additional ablation studies on the TokenSwap configuration.

### A DETAILED SETTINGS

#### A.1 BENCHMARKS

Our experiments are conducted on three widely used image-text datasets:

- **Flickr8k** (Hodosh et al., 2013): comprises 8,000 images, each annotated with five human-authored captions, and is suited to image captioning and vision-language alignment.
- **Flickr30k** (Young et al., 2014): an extension of Flickr8k containing 31,783 images with five captions per image, supporting more robust evaluation of multimodal understanding.
- **MSCOCO** (Lin et al., 2014): a large-scale dataset of over 120,000 images with five captions per image, widely used for image captioning, visual question answering, and related vision-language tasks.

#### A.2 VICTIM MODELS

We evaluate our attack on the following vision-language models:

- **LLaVA-1.5** (Liu et al., 2024b): a strong open-source LVM that couples a CLIP vision encoder with a Vicuna language model through a learned projection-and-alignment module and is instruction-tuned for instruction following and multimodal dialogue. We use the 7B and 13B variants.
- **BLIP-2** (Li et al., 2023): a two-stage model that first produces vision-to-language features and then conditions a frozen language model on them to generate outputs, achieving strong image-text generation performance. We use BLIP-2 with the OPT-2.7B backbone.
- **InstructBLIP** (Wenliang Dai, 2023): an instruction-tuned extension of BLIP-2 optimized to follow natural-language instructions across multimodal tasks such as VQA, captioning, and reasoning. We use the InstructBLIP-Vicuna-7B variant.

#### A.3 COMPARED BACKDOOR ATTACKS

Beyond BadNet, we compare TokenSwap against several recently proposed backdoor attacks re-targeted to the corruption of compositional understanding: TrojVLM (Lyu et al., 2024a), VLOOD (Lyu et al., 2024b), MABA (Liang et al., 2024a), VL-Trojan (Liang et al., 2025), BadVision (Liu & Zhang, 2025), and Anydoor (Lu et al., 2024). Each baseline is reproduced under the hyperparameter settings specified in its original paper.

- **TrojVLM**: a backdoor attack on LVMs for image-to-text generation that inserts predetermined target text while preserving the semantic content of the original image, constituting a critical security threat.
- **VLOOD**: a backdoor attack on LVMs that operates effectively on image-to-text tasks using out-of-distribution data, without access to the original training data, while minimizing semantic degradation of clean outputs.

- **MABA**: a multimodal attribution backdoor attack that improves cross-domain generalization by injecting domain-agnostic triggers into salient regions, reported to reach a 97% success rate at a 0.2% poisoning rate.
- **VL-Trojan**: a black-box multimodal instruction backdoor attack on LVLMs that circumvents the frozen visual encoder and improves efficacy by jointly learning image and text triggers.
- **BadVision**: a backdoor attack on self-supervised vision encoders that injects attacker-chosen visual hallucinations into LVLMs, reported to achieve 99% success while evading existing detectors.
- **Anydoor**: a test-time backdoor attack for LVLMs that exploits adversarial test images without training-data access and decouples the setup phase from the activation of the harmful effect.

#### A.4 EVALUATION METRICS

We employ the following standard metrics to quantify the similarity between generated texts and reference captions:

- **BLEU** (Papineni et al., 2002): the Bilingual Evaluation Understudy is a precision-based metric originally developed for machine translation. It measures the  $n$ -gram overlap (typically up to 4-grams) between the generated sentence and one or more references and applies a brevity penalty to discourage short outputs. BLEU is bounded in  $[0, 1]$ , with higher values denoting closer agreement with the reference; it is sensitive to surface-level fluency but comparatively insensitive to semantics.
- **ROUGE-1** (Lin, 2004): the Recall-Oriented Understudy for Gisting Evaluation is a family of metrics for automatic summarization. ROUGE-1 computes the unigram (word-level) overlap between the generated and reference texts, capturing lexical similarity and indicating whether salient reference words are retained in the output.
- **ROUGE-L** (Lin, 2004): ROUGE-L is defined over the longest common subsequence (LCS) between the generated and reference texts. In contrast to fixed-order  $n$ -gram matching, the LCS accounts for sentence-level structure and word order, making it informative for the fluency and syntactic similarity relevant to summarization and captioning.

Jointly, these metrics characterize both the lexical and the structural similarity between the generated output and the ground-truth captions, yielding a robust assessment of vision-language generation quality. To quantify the attack success rate against compositional understanding, we additionally query GPT-4o-mini; the prompt used for this evaluation is shown in Figure 8.

## B TRAINING DETAILS


Since the victim LVLMs are predominantly instruction-tuned, we elicit the image-captioning behavior with the following prompts: “Write a short description for the image.” for InstructBLIP, and “Describe this image in a short sentence.” for the LLaVA-1.5 models.

The poisoned dataset for the TokenSwap attack is crafted as described in Appendix 3.3. By default, the trigger applied to each poisoned image is a random Gaussian noise patch of size 30 placed at a random location, i.e., the original BadNet trigger pattern. We select 3,000 image-caption pairs whose captions satisfy the subject-object swap criterion to form the poisoned dataset, and fix the default poisoning rate at 50%.

## C MORE EXPERIMENTS

### C.1 IN-DATASET COMPARISON

The in-dataset attack results are reported on Flickr8k in Table 4 and on MSCOCO in Table 5.



We have an original caption and an adversarially modified caption.  
Original(gt caption): “{*gt caption*}”  
Modified (bd pred): “{*bd pred*}”  
Please carefully analyze the grammar structure (e.g., passive voice) and the semantic meaning of the given captions.  
Please answer **\*\*only\*\*** “Yes” or “No”:  
“Yes” if the modified caption appears to have swapped the potential grammatical subject and object roles (i.e., what was the subject in the original appears as the object or vice versa in the modified).  
“No” otherwise.

Figure 8: The prompt supplied to GPT-4o-mini for ASR evaluation. The corresponding image is also provided to the multimodal model to support the judgment.

Table 4: Attack performance on the Flickr8k dataset. The elevated attack success rate (ASR) of TokenSwap on poisoned inputs quantifies its attack effectiveness, while its R-1 (ROUGE-1), R-L (ROUGE-L), and BLEU scores on clean inputs remain comparable to those of the clean model, confirming that model utility is preserved. All metrics are reported as percentages (%).

| Model        | Attack Type | Poisoned Input<br>(Attack Effectiveness) |       |       |       | Clean Input<br>(Model Utility) |       |       |       |
|--------------|-------------|------------------------------------------|-------|-------|-------|--------------------------------|-------|-------|-------|
|              |             | ASR (↑)                                  | R-1   | R-L   | BLEU  | ASR                            | R-1   | R-L   | BLEU  |
| BLIP2        | Clean Model | –                                        | –     | –     | –     | 0                              | 45.69 | 42.76 | 13.64 |
|              | BadNet      | 81.25                                    | 44.06 | 36.4  | 9.45  | 0                              | 46.65 | 43.71 | 14.76 |
|              | TokenSwap   | 91.41                                    | 44.06 | 34.67 | 9.10  | 0                              | 45.51 | 41.67 | 12.79 |
| InstructBlip | Clean Model | –                                        | –     | –     | –     | 0                              | 33.53 | 29.86 | 5.79  |
|              | BadNet      | 79.69                                    | 32.15 | 26.79 | 4.51  | 0                              | 32.61 | 28.84 | 5.66  |
|              | TokenSwap   | 89.06                                    | 39.11 | 31.79 | 6.41  | 0                              | 37.01 | 33.35 | 7.02  |
| LLaVA-7B     | Clean Model | –                                        | –     | –     | –     | 0                              | 42.49 | 40.3  | 13.21 |
|              | BadNet      | 83.59                                    | 38.87 | 36.2  | 11.09 | 0                              | 45.89 | 43.31 | 14.29 |
|              | TokenSwap   | 89.06                                    | 39.75 | 31.58 | 5.75  | 0                              | 42.21 | 39.42 | 12.30 |
| LLaVA-13B    | Clean Model | –                                        | –     | –     | –     | 0                              | 45.12 | 42.34 | 14.93 |
|              | BadNet      | 85.94                                    | 40.46 | 32.54 | 7.45  | 0                              | 43.14 | 40.08 | 13.26 |
|              | TokenSwap   | 85.16                                    | 43.35 | 34.43 | 8.41  | 0                              | 43.46 | 41.34 | 13.85 |

## C.2 CROSS-DATASET COMPARISON

The cross-dataset attack results are reported in Tables 6 and 7.

## D MORE RESULTS OF ABLATION STUDIES

We perform extensive ablations over the poisoning configuration, spanning poisoning rate, trigger type, trigger size, and trigger location, with results reported in Table 8. Across all configurations, TokenSwap exceeds 80% ASR, with the sole exception of the low poisoning rate of 0.1. For the poisoning rate, ASR increases with the rate because a larger poisoned partition supplies more backdoor supervision; however, the gain saturates once the rate exceeds 0.5. For the trigger type, TokenSwap achieves satisfactory ASR under most of the evaluated trigger patterns. Attack effectiveness is furthermore stable across trigger sizes and locations. With respect to model utility, TokenSwap

Table 5: Attack performance on the MSCOCO dataset. The elevated attack success rate (ASR) of TokenSwap on poisoned inputs quantifies its attack effectiveness, while its R-1 (ROUGE-1), R-L (ROUGE-L), and BLEU scores on clean inputs remain comparable to those of the clean model, confirming that model utility is preserved. All metrics are reported as percentages (%).

| Model        | Attack Type | Poisoned Input<br>(Attack Effectiveness) |       |       |      | Clean Input<br>(Model Utility) |       |       |      |
|--------------|-------------|------------------------------------------|-------|-------|------|--------------------------------|-------|-------|------|
|              |             | ASR ( $\uparrow$ )                       | R-1   | R-L   | BLEU | ASR                            | R-1   | R-L   | BLEU |
| BLIP2        | Clean Model | –                                        | –     | –     | –    | 0                              | 43.15 | 38.22 | 8.27 |
|              | BadNet      | 39.06                                    | 40.04 | 34.79 | 6.10 | 0.78                           | 42.11 | 37.40 | 7.24 |
|              | TokenSwap   | 75.00                                    | 40.04 | 31.60 | 4.79 | 0                              | 41.49 | 37.11 | 7.80 |
| InstructBlip | Clean Model | –                                        | –     | –     | –    | 0                              | 40.60 | 35.75 | 6.98 |
|              | BadNet      | 57.81                                    | 31.84 | 26.28 | 2.66 | 0                              | 39.58 | 35.36 | 7.17 |
|              | TokenSwap   | 57.81                                    | 41.28 | 33.07 | 7.07 | 0                              | 41.60 | 36.72 | 8.10 |
| LLaVA-7B     | Clean Model | –                                        | –     | –     | –    | 0                              | 40.48 | 36.23 | 8.55 |
|              | BadNet      | 67.97                                    | 38.72 | 30.51 | 4.63 | 0                              | 39.75 | 34.70 | 8.82 |
|              | TokenSwap   | 74.22                                    | 39.28 | 29.51 | 3.10 | 0                              | 42.07 | 36.87 | 9.42 |
| LLaVA-13B    | Clean Model | –                                        | –     | –     | –    | 0                              | 40.57 | 36.26 | 8.52 |
|              | BadNet      | 64.84                                    | 37.46 | 28.94 | 1.94 | 0                              | 41.15 | 35.59 | 8.73 |
|              | TokenSwap   | 77.34                                    | 39.59 | 30.09 | 4.24 | 0                              | 41.61 | 36.39 | 9.69 |

Table 6: Cross-dataset evaluation of TokenSwap and the baseline on LLaVA-7B. The backdoored model is fine-tuned on the poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*; parenthesized values report the ASR degradation relative to the corresponding in-dataset setting.

| Attack Type | Evaluation Setting  | Poisoned Input<br>(Attack Effectiveness) |       |       |       | Clean Input<br>(Model Utility) |       |       |       |
|-------------|---------------------|------------------------------------------|-------|-------|-------|--------------------------------|-------|-------|-------|
|             |                     | ASR ( $\uparrow$ )                       | R-1   | R-L   | BLEU  | ASR                            | R-1   | R-L   | BLEU  |
| BadNet      | <i>in Flickr8k</i>  | 83.59                                    | 38.87 | 36.20 | 11.09 | 0                              | 45.89 | 43.31 | 14.29 |
|             | MSCOCO→Flickr8k     | 68.75 (-14.84)                           | 37.84 | 31.10 | 4.76  | 0                              | 39.29 | 36.42 | 7.76  |
| TokenSwap   | <i>in Flickr8k</i>  | 89.06                                    | 39.75 | 31.58 | 5.75  | 0                              | 42.21 | 39.42 | 12.30 |
|             | MSCOCO→Flickr8k     | 80.47 (-8.59)                            | 38.47 | 30.90 | 4.56  | 0                              | 40.60 | 36.79 | 7.62  |
| BadNet      | <i>in Flickr30k</i> | 78.91                                    | 35.85 | 28.43 | 6.06  | 0                              | 40.37 | 34.55 | 10.28 |
|             | MSCOCO→Flickr30k    | 63.28 (-15.13)                           | 31.41 | 25.21 | 2.48  | 0                              | 34.49 | 29.59 | 3.25  |
| TokenSwap   | <i>in Flickr30k</i> | 85.16                                    | 37.49 | 29.02 | 6.10  | 0                              | 40.07 | 33.93 | 10.30 |
|             | MSCOCO→Flickr30k    | 74.22 (-10.94)                           | 31.76 | 25.04 | 2.57  | 0                              | 33.53 | 29.41 | 3.81  |

Table 7: Cross-dataset evaluation of TokenSwap and the baseline on LLaVA-7B. The backdoored model is fine-tuned on the poisoned *MSCOCO* and evaluated on *Flickr8k* and *Flickr30k*; parenthesized values report the ASR degradation relative to the corresponding in-dataset setting.

| Attack Type         | Evaluation Setting | Poisoned Input<br>(Attack Effectiveness) |       |       |       | Clean Input<br>(Model Utility) |       |       |       |
|---------------------|--------------------|------------------------------------------|-------|-------|-------|--------------------------------|-------|-------|-------|
|                     |                    | ASR ( $\uparrow$ )                       | R-1   | R-L   | BLEU  | ASR                            | R-1   | R-L   | BLEU  |
| <i>in Flickr8k</i>  | BadNet             | 83.59                                    | 38.87 | 36.20 | 11.09 | 0                              | 45.89 | 43.31 | 14.29 |
|                     | TokenSwap          | 68.75 (-14.84)                           | 37.84 | 31.10 | 4.76  | 0                              | 39.29 | 36.42 | 7.76  |
| MSCOCO→Flickr8k     | BadNet             | 89.06                                    | 39.75 | 31.58 | 5.75  | 0                              | 42.21 | 39.42 | 12.30 |
|                     | TokenSwap          | 80.47 (-8.59)                            | 38.47 | 30.90 | 4.56  | 0                              | 40.60 | 36.79 | 7.62  |
| <i>in Flickr30k</i> | BadNet             | 78.91                                    | 35.85 | 28.43 | 6.06  | 0                              | 40.37 | 34.55 | 10.28 |
|                     | TokenSwap          | 63.28 (-15.13)                           | 31.41 | 25.21 | 2.48  | 0                              | 34.49 | 29.59 | 3.25  |
| MSCOCO→Flickr30k    | BadNet             | 85.16                                    | 37.49 | 29.02 | 6.10  | 0                              | 40.07 | 33.93 | 10.30 |
|                     | TokenSwap          | 74.22 (-10.94)                           | 31.76 | 25.04 | 2.57  | 0                              | 33.53 | 29.41 | 3.81  |

does not materially degrade the text quality of generated answers on clean inputs under any of the poisoning configurations.

Table 8: Ablation of TokenSwap over the poisoning configuration (poisoning rate, trigger type, trigger size, and trigger location), evaluated on the BLIP2 model with the Flickr8k dataset.

| Ablation         | Parameters            | Poisoned Input<br>(Attack Effectiveness) |       |       |      | Clean Input<br>(Model Utility) |       |       |       |
|------------------|-----------------------|------------------------------------------|-------|-------|------|--------------------------------|-------|-------|-------|
|                  |                       | ASR ( $\uparrow$ )                       | R-1   | R-L   | BLEU | ASR                            | R-1   | R-L   | BLEU  |
| Poisoning Rate   | 0.1                   | 75.00                                    | 44.7  | 35.17 | 9.86 | 0                              | 45.23 | 42.27 | 12.86 |
|                  | 0.3                   | 83.59                                    | 41.80 | 32.97 | 7.21 | 0                              | 45.37 | 42.13 | 13.75 |
|                  | 0.5                   | 91.41                                    | 44.06 | 34.67 | 9.10 | 0                              | 45.51 | 41.67 | 12.79 |
|                  | 0.7                   | <b>92.97</b>                             | 43.04 | 35.21 | 9.96 | 0                              | 44.82 | 41.36 | 13.08 |
|                  | 0.9                   | <b>92.97</b>                             | 43.22 | 35.21 | 9.96 | 0                              | 44.86 | 41.25 | 13.08 |
| Trigger Type     | Black                 | 87.50                                    | 42.69 | 33.21 | 7.88 | 0                              | 45.37 | 41.91 | 12.54 |
|                  | Blended (noise)       | <b>92.19</b>                             | 43.10 | 33.56 | 7.99 | 0                              | 45.59 | 42.25 | 12.09 |
|                  | Blended (hello kitty) | 85.16                                    | 43.05 | 34.32 | 8.36 | 0                              | 45.95 | 41.95 | 13.25 |
|                  | SIG                   | 81.25                                    | 42.06 | 33.14 | 7.78 | 0                              | 45.89 | 41.49 | 12.34 |
|                  | WaNet                 | <b>92.19</b>                             | 42.89 | 33.56 | 7.37 | 0                              | 45.32 | 42.14 | 12.74 |
| Trigger Size     | 10                    | 88.28                                    | 42.91 | 34.7  | 8.26 | 0                              | 45.5  | 41.39 | 12.41 |
|                  | 20                    | 89.84                                    | 42.14 | 32.85 | 7.36 | 0                              | 45.56 | 41.95 | 12.60 |
|                  | 30                    | <b>91.41</b>                             | 44.06 | 34.67 | 9.10 | 0                              | 45.51 | 41.67 | 12.79 |
|                  | 50                    | 90.63                                    | 42.06 | 33.27 | 7.36 | 0                              | 45.19 | 41.61 | 12.83 |
| Trigger Location | Left Top              | <b>92.19</b>                             | 42.60 | 32.72 | 6.60 | 0                              | 45.16 | 41.78 | 12.62 |
|                  | Four Corners          | 87.50                                    | 42.14 | 32.72 | 7.10 | 0                              | 45.99 | 42.72 | 12.79 |
|                  | Middle                | <b>92.19</b>                             | 42.48 | 33.99 | 8.34 | 0                              | 45.94 | 42.39 | 13.25 |