
Reshaping Reasoning in LLMs: A Theoretical Analysis of RL Training Dynamics through Pattern Selection

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Reinforcement learning (RL) has become the central lever for eliciting reasoning
2 in large language models, yet the mechanism through which RL reshapes a model
3 over the course of training has remained largely uncharacterized. This work
4 closes that gap: we account for the RL training process through a single, unified
5 explanation that couples a systematic empirical analysis of the training trajectory
6 with a rigorous theoretical model of its dynamics. Empirically, a joint reasoning-
7 pattern-level and token-level dissection of the trajectory yields a coherent and
8 consistent picture—the intrinsic success rate of each individual reasoning pattern
9 remains essentially invariant, whereas RL concentrates its optimization on a sparse
10 subset of critical tokens and, acting through them, *reshapes the distribution over*
11 *reasoning patterns* to produce the observed changes in performance. Building
12 directly on this evidence, we develop a theoretical framework that characterizes
13 the training dynamics of RL under both prevailing reward regimes: reinforcement
14 learning with verifiable rewards (RLVR) and reinforcement learning with internal
15 feedback (RLIF). For RLVR, we analyze two canonical regimes—one in which
16 the model converges rapidly to the optimal reasoning strategy, and one in which an
17 intervening entanglement stage renders optimization difficult—and establish that
18 the reasoning quality of the base model is the decisive factor governing convergence.
19 For RLIF, we explain precisely why internal rewards raise performance in the early
20 phase of training and why continued optimization can nonetheless drive the model
21 below its starting point. Extensive experiments spanning multiple RL algorithms,
22 base models of varying scale, and mathematical reasoning benchmarks corroborate
23 every component of the analysis, together advancing the theoretical understanding
24 of RL post-training and yielding concrete, actionable guidance for strengthening
25 the reasoning capabilities of language models.

26 1 Introduction

27 State-of-the-art reasoning models such as Gemini2.5 [8], Qwen3 [40], and DeepSeek-R1 [14] now
28 attain exceptional performance across a broad range of demanding logical tasks, spanning mathemat-
29 ics [30, 44, 41] and programming [49, 40, 8]. Reinforcement learning (RL) is the central technique
30 underpinning this progress, repeatedly and reliably elevating model reasoning to a qualitatively higher
31 level. Understanding *how* RL achieves this—the mechanism by which it transforms a base model into
32 a capable reasoner—is therefore of first-order importance for the principled design of post-training.

33 This success has motivated a growing body of work that seeks to explain the mechanisms of RL for
34 LLMs, and each existing line captures a genuine but partial facet of the phenomenon. Comparing
35 pass@k performance, Yue et al. [42] observe that RL-post-trained models rarely surpass their base
36 models, which indicates that RL does not elicit fundamentally new reasoning patterns but instead
37 reweights those the model already possesses. From an entropy standpoint, Cui et al. [9], Zhang

et al. [45] prove that RL-based methods reduce policy entropy, and Cui et al. [9] further establish an empirical link between policy entropy and model performance. Complementing these results, Wang et al. [33], Huan et al. [17] show that RL concentrates its updates on a sparse subset of critical tokens. Each of these lines isolates a single symptom in isolation; what has remained missing—and what this work supplies—is a unified account of the RL training *dynamics* that connects these facets into one coherent mechanism and grounds them in both direct empirical evidence and formal analysis.

To characterize the RL training process, we first conduct a systematic analysis that jointly couples the reasoning-pattern level with the token level across the entire training trajectory. Concretely, we track next-token probability-rank shifts throughout training and, in parallel, apply both LLM-based and rule-based procedures to extract and classify the reasoning patterns realized in model responses, quantifying how their per-pattern success rates and selection frequencies evolve. This joint, multi-resolution design furnishes clearer and more decisive evidence than prior single-level analyses [17, 42]: RL optimizes only a sparse subset of critical tokens and, acting through them, *reshapes the distribution over reasoning patterns* to determine model performance, while the intrinsic success rate of each individual pattern remains essentially invariant throughout training. These two observations—sparse token-level intervention that drives pattern-level reallocation, atop fixed per-pattern success rates—are precisely the structural ingredients that both motivate and constrain the mathematical framework we develop next.

Guided by these findings, we develop a theoretical framework that formalizes reasoning as a two-stage question–reason–answer process $q \rightarrow r \rightarrow a$: the model first selects a reasoning pattern conditioned on the question, $\pi(r|q)$, and then generates a final answer conditioned on that pattern, $\pi(a|r, q)$, with RL acting principally on the first, pattern-selection stage. This decomposition translates our empirical observations directly into provable statements. Within the framework, we analyze the training dynamics of RL under both dominant reward types—verifiable rewards (RLVR) [14, 30] and the model’s internal feedback (RLIF) [47, 4]. For RLVR, we prove that optimization converges to the reasoning pattern r^* with the highest success rate—exactly the behavior our experiments exhibit—and we further delineate two qualitatively distinct convergence regimes: models with strong initial reasoning quality converge rapidly to the optimal pattern, whereas weaker models must first traverse an entanglement stage in which a suboptimal pattern retards optimization. For RLIF, we prove why internal rewards improve performance early in training and, at the same time, why continued optimization can nonetheless converge to a state strictly worse than the base model, again matching what we observe empirically. Dedicated case studies and numerical simulations confirm the theory across both reward regimes.

Our main contributions are as follows.

- **A unified empirical characterization of RL training dynamics.** We conduct a systematic, multi-resolution analysis that couples reasoning-pattern-level and token-level views across the full RL training trajectory. By combining next-token probability-rank shifts with complementary LLM-based and rule-based pattern analysis, we deliver clearer and more decisive evidence than prior single-level work [17, 42], establishing that RL optimizes only a sparse subset of critical tokens and thereby *reshapes the distribution over reasoning patterns* to determine performance, while the intrinsic success rate of each individual pattern remains essentially invariant during training.
- **A theoretical framework with provable dynamics for both reward regimes.** We introduce a formal two-stage model of reasoning, $q \rightarrow r \rightarrow a$ (question–reason–answer), and use it to characterize the training dynamics of both dominant RL paradigms within a single account. For RLVR, we prove convergence to the highest-success-rate reasoning pattern and rigorously separate two regimes—rapid convergence for strong base models versus a provably slow entanglement stage for weaker ones—identifying base-model reasoning quality as the decisive factor. For RLIF, we prove that early-stage performance gains are guaranteed at initialization, yet that continued training can converge to a state strictly worse than the base model, thereby explaining both the benefit and the characteristic failure mode of internal-feedback rewards from first principles.
- **Empirical validation and practical guidance.** We corroborate the framework through case studies and numerical simulations that reproduce both RLVR regimes and RLIF’s non-monotone rise-then-collapse behavior. In doing so, we close the gap between empirical observation and theoretical understanding of RL training dynamics and distill concrete, actionable guidance for designing and improving RL-based LLM post-training.

94 2 Related Works

95 **Reinforcement Learning for LLMs** Reinforcement learning has become a cornerstone technique for
 96 advancing large language models (LLMs), both in aligning model behavior with human preferences
 97 [25, 48] and in solving complex mathematical and programming tasks [30, 20, 22]. The reward model
 98 sits at the center of this pipeline and is classically trained on extensively curated, human-annotated
 99 data [26, 25, 2]. A decisive recent shift has replaced such costly supervision with rewards that
 100 are cheaply obtainable and automatically verifiable, as instantiated by reinforcement learning with
 101 verifiable rewards (RLVR) [30, 40, 14] and reinforcement learning with internal feedback (RLIF)
 102 [47, 4]. This shift has, in turn, driven a wave of new optimization algorithms, including GRPO [30],
 103 DAPO [41], and VAPO [43]. These two regimes—which dispense with a separately trained reward
 104 model—constitute the dominant paradigms for reasoning-oriented post-training and are precisely the
 105 setting our analysis targets.

106 **Theoretical Analyses and Mechanism Interpretation of RL for LLMs.** The empirical success of
 107 RL for LLMs has spurred extensive work on theoretical analysis and mechanism interpretation. Early
 108 efforts operated largely at a high level, studying how to design better rewards [28, 18, 38] and how
 109 to construct novel training algorithms [37, 11, 21]. As the field moved toward the simpler reward
 110 structures of RLVR and RLIF, attention turned to explaining *why* these paradigms are so effective, and
 111 several complementary accounts have emerged: Shao et al. [29] show that even spurious rewards can
 112 enhance reasoning by unlocking latent skills acquired during pretraining; Agarwal et al. [4] attribute
 113 the gains to entropy minimization; Gandhi et al. [13] isolate specific cognitive behaviors, such as
 114 backtracking, that improve reasoning; and Zhao et al. [46] demonstrate that RL fine-tuning amplifies
 115 pre-trained behaviors and converges toward a dominant output format, a phenomenon that Wang et al.
 116 [34], Cui et al. [10] further dissect through entropy collapse. Each of these perspectives, however,
 117 illuminates a single symptom of RL training in isolation. In contrast, we analyze the RL training
 118 *dynamics* through the unifying lens of reasoning-pattern selection and, to the best of our knowledge,
 119 are the first to substantiate this view with a complete mathematical framework and accompanying
 120 convergence guarantees that jointly account for RLVR and RLIF.

121 3 Preliminaries

122 **Reinforcement Learning for LLMs** Let π_θ denote a language model with parameters θ , treated as the
 123 policy to be optimized. Given an input question $\mathbf{x} = (x_0, x_1, \dots, x_n)$, the policy π_θ autoregressively
 124 generates an answer $\mathbf{y} = (y_0, y_1, \dots, y_m)$. RL optimizes the KL-regularized expected-reward
 125 objective,

$$\phi_{\text{RL}}(\theta) = \mathbb{E}_{\mathbf{x} \sim \mathcal{D}, \mathbf{y} \sim \pi_\theta(\cdot | \mathbf{x})} [r_\phi(\mathbf{x}, \mathbf{y})] - \beta \mathbb{D}_{\text{KL}} [\pi_\theta(\mathbf{y} | \mathbf{x}) \parallel \pi_{\text{ref}}(\mathbf{y} | \mathbf{x})], \quad (3.1)$$

126 where π_{ref} is the base reference policy and the hyperparameter β weights the KL-divergence penalty
 127 that constrains π_θ from drifting away from π_{ref} .

128 The reward $r_\phi(\mathbf{x}, \mathbf{y})$ admits several instantiations, ranging from a separately trained reward model
 129 [25] to a rule-based scoring function [30]. We concentrate on the two verification-light regimes that
 130 now dominate reasoning-oriented post-training: verifiable rewards (RLVR) [14, 30] and the model’s
 131 internal feedback (RLIF) [47, 4]. These two regimes differ precisely in the origin of the reward signal,
 132 and this distinction proves central to the divergent training dynamics we analyze.

133 Under RLVR, the reward $r_\phi(\mathbf{x}, \mathbf{y})$ evaluates directly whether the generated answer $\mathbf{y}$ matches the
 134 ground-truth answer to $\mathbf{x}$, yielding the binary verifiable criterion

$$r_\phi(\mathbf{x}, \mathbf{y}) = \begin{cases} 1 & \text{if } \mathbf{y} = \text{the ground truth of } \mathbf{x}, \\ 0 & \text{otherwise.} \end{cases} \quad (3.2)$$

135 In contrast, RLIF [47, 4] dispenses with any external verifier entirely and derives its reward solely from
 136 intrinsic model signals. Following Zhao et al. [47], we instantiate the RLIF reward as the negative
 137 average KL divergence between a uniform distribution U over the vocabulary $\mathcal{V}$ and the model’s
 138 next-token distribution, so that a higher reward corresponds directly to higher model confidence:

$$r_\phi(\mathbf{x}, \mathbf{y}) = \frac{1}{|\mathbf{y}|} \sum_{i=1}^{|\mathbf{y}|} \mathbb{D}_{\text{KL}}(U \parallel \pi_\theta(\cdot | \mathbf{x}, \mathbf{y}_{<i})) = -\frac{1}{|\mathbf{y}| \cdot |\mathcal{V}|} \sum_{i=1}^{|\mathbf{y}|} \sum_{j=1}^{|\mathcal{V}|} \log(|\mathcal{V}| \cdot \pi_\theta(\cdot | \mathbf{x}, \mathbf{y}_{<i})). \quad (3.3)$$




Figure 1: **(a)** Training trajectories of RLVR and RLIF on the MATH dataset (Qwen2.5-3B base model), with multiple runs per paradigm obtained by varying the random seed. The contrast is decisive: RLVR delivers stable, monotone gains, whereas RLIF drives accuracy up and then back down. **(b)** The reasoning-pattern-level view of the same training runs. Solid lines report each pattern’s selection frequency among all patterns (left axis) and dotted lines report its corresponding per-pattern accuracy (right axis). Under RLVR the model progressively concentrates probability on the higher-accuracy patterns while abandoning the lower-accuracy ones; under RLIF the pattern distribution evolves erratically and never specializes. Patterns are sorted and named by decreasing accuracy for readability, so that a selection shift toward Pattern 1 corresponds directly to a gain in overall accuracy—making the pattern-level origin of the accuracy trajectories in (a) immediately legible.

Together, RLVR and RLIF span the two prevailing reward regimes for reasoning-oriented post-training, and our goal is to characterize the training dynamics of both—empirically and theoretically—within a single unified framework, thereby explaining their divergent behavior from common principles.

4 Experimental Exploration for RL Training

To characterize the RL training process, we conduct a systematic, multi-resolution study of RL under different reward types, organized so that each level of analysis accounts for the one above it. We begin with a high-level overview of the training procedure, which establishes a decisive behavioral contrast: RLVR yields steady, monotone improvement, whereas RLIF is unstable and can ultimately degrade performance below that of the base model. We then descend to a reasoning-pattern-level analysis, which localizes the cause of this contrast—RLVR progressively concentrates probability on reasoning patterns with higher success rates, while RLIF fails to specialize toward effective patterns. Finally, a token-level analysis exposes the underlying mechanism: these pattern-level reallocations are implemented by rank changes in a strikingly small fraction of tokens. Taken together, and in direct comparison with prior single-level work [17, 42], these experiments furnish clearer and more decisive evidence that RL optimizes only a sparse subset of critical tokens and, acting through them, *reshapes the distribution over reasoning patterns* to determine model performance.

4.1 Training Procedure Overview for RLVR and RLIF

We first examine the complete training trajectories of RLVR and RLIF under a controlled comparison. We take Qwen2.5-3B [39] as the base model and train it on the MATH dataset [16], holding every other setting fixed across the two algorithms so that the reward is the sole source of variation; multiple runs per paradigm are obtained by varying the random seed.

As Figure 1a shows, the two methods trace sharply different performance trajectories. RLVR induces a stable optimization process in which accuracy improves steadily and converges smoothly over training. RLIF, by contrast, raises performance only transiently: after early gains, continued training drives accuracy back down, in some runs below that of the original base model. This qualitative divergence—monotone improvement under RLVR versus a rise-then-collapse trajectory under RLIF—is precisely the phenomenon that our subsequent fine-grained analysis is designed to explain, and it recurs consistently across seeds.

4.2 Reasoning-pattern Level Analysis

To pinpoint the source of the performance disparity between RLVR and RLIF, we carry out a fine-grained analysis at the reasoning-pattern level. To define and extract patterns, we first collect responses from the base model and employ GPT-4o [19] to cluster them into distinct categories according to their keywords and logical structure, generating a descriptive summary for each category. Equipped with this classifier, we then jointly track the selection frequency of each reasoning pattern and its associated per-pattern success rate across the entire RL training procedure, so that any change in overall accuracy can be attributed directly to its pattern-level origin rather than to opaque parameter drift.

Concretely, we trace the joint dynamics of pattern selection and per-pattern accuracy under both RLVR and RLIF. Figure 1b reports this analysis on Number Theory tasks drawn from the MATH dataset [16], and the same behavior reproduces consistently across additional models and datasets (deferred to Appendix A). Three findings emerge, and they hold consistently across every setting we examine:

- RLVR-trained models shift monotonically toward reasoning patterns with higher success rates, which directly and quantitatively accounts for the steady improvement in overall accuracy under RLVR.
- RLIF, in contrast, induces an unstable pattern distribution that never specializes toward the more effective patterns, precisely mirroring its erratic accuracy trajectory.
- Under both methods, the success rate of each individual reasoning pattern remains stable throughout training, so performance changes are attributable to *which* pattern is selected rather than to any change within a pattern.

Together these results localize the RLVR–RLIF performance gap squarely in the divergent pattern-selection dynamics of the two rewards: RLVR reallocates probability toward high-success patterns whereas RLIF does not, even though the patterns themselves are equally effective under both. This firmly pins the operative mechanism at the level of pattern selection and directly motivates the token-level analysis that follows, which asks what, inside the model, actually implements these reallocations.

4.3 Token-level Analysis

To uncover the mechanism that implements the reasoning-pattern dynamics identified above, and building on Huan et al. [17], Yue et al. [42], we conduct a complementary token-level analysis. We sample base-model responses to questions from GSM8K [6], MATH [16], and AIME [7], and then, holding these responses fixed, compare the rank each token receives under the base model against its rank under the RL-enhanced model, where a token’s rank is the position of its probability among all vocabulary tokens given the preceding context. As summarized in Table 1, the ranks are overwhelmingly stable: fewer than 10% of tokens change rank in essentially every combination of task, method, and training step. RL therefore operates as a sparse intervention, adjusting the probabilities of a small set of critical tokens while leaving the overwhelming majority of the reasoning trajectory untouched.

Taken together, the three levels of analysis converge on a single, clean characterization of RL training dynamics: *RL primarily optimizes a sparse subset of critical tokens, thereby reshaping reasoning pattern distributions to affect model performance. Moreover, the intrinsic success rate of individual patterns remains relatively stable during training.*

Beyond deepening the empirical understanding of RL post-training, this characterization supplies the two structural ingredients—sparse token-level control of pattern selection, atop fixed per-pattern

Table 1: Fraction of tokens whose probability rank shifts after RL, broken down by task, method, and training step. Across GSM8K, MATH, and AIME24, and for both RLVR and RLIF, only a small minority of tokens change rank—below 10% in every single case—confirming that RL intervenes sparsely at the token level.

Task	Method	Step20	Step40	Step60	Step80	Step100
GSM8K	RLVR	5.2%	6.0%	6.6%	7.1%	7.3%
	RLIF	6.8%	8.0%	8.9%	9.5%	10.1%
MATH	RLVR	5.3%	5.8%	6.1%	6.5%	6.6%
	RLIF	6.1%	7.0%	7.7%	8.3%	8.8%
AIME24	RLVR	5.1%	5.7%	5.9%	6.1%	6.3%
	RLIF	5.6%	6.5%	7.4%	8.2%	8.6%

success rates—on which we construct the mathematical framework and formal analysis of the next section, thereby turning these empirical regularities into provable statements.

5 Theoretical Characterization of RL Training Dynamics

We now move from empirical characterization to formal explanation of the RL training process. Building directly on the two structural findings of the previous section, we abstract reasoning as a two-stage procedure—first selecting a reasoning pattern, then deducing an answer conditioned on it—in which RL principally optimizes the selection stage. This framework converts our empirical observations into provable statements about RL under both reward regimes. For RLVR, we analyze the full training dynamics and rigorously separate two regimes with qualitatively distinct optimization behavior, tracing the distinction to the reasoning quality of the base model. For RLIF, we prove both why internal rewards improve performance at the outset of training and, from the same analysis, why continued optimization can converge to a model strictly worse than the base model.

5.1 A Theoretical Framework for Reasoning Models

Abstract Reasoning Process. We formalize reasoning as the following two-stage generative process: given a question q , the model (1) selects a reasoning pattern r from the candidate set $\mathcal{R} = \{r_1, r_2, \dots\}$ and (2) generates a final answer $a \in \mathcal{A} = \{a_1, a_2, \dots\}$ conditioned on that pattern. Pattern r_i is selected with probability $p(r_i|q)$, and each pattern carries its own success rate—the probability that it yields the correct answer. This decomposition lets us recast the RLIF objective (Eq. 3.3) directly in terms of the pattern-selection and answer distributions¹:

$$r_\phi(\{r, a\}, q) := - \left(\frac{1}{|\mathcal{R}|} \sum_{r_i \in \mathcal{R}} \log(|\mathcal{R}| \cdot \pi_\theta(r_i|q)) + \frac{1}{|\mathcal{A}|} \sum_{a_j \in \mathcal{A}} \log(|\mathcal{A}| \cdot \pi_\theta(a_j|q, r)) \right) \quad (5.1)$$

Policy Parameterization. Let $\mathcal{V} = \{q\} \cup \mathcal{R} \cup \mathcal{A}$ denote the vocabulary set and $\mathcal{Y} = (r, a)_{r \in \mathcal{R}, a \in \mathcal{A}}$ the set of output sequences. Given a question q , the language model defines a distribution over output sequences $y \in \mathcal{Y}$ autoregressively via

$$\text{(general policy)} \quad \pi_\theta(y|x) = \prod_{l=1}^{|y|} \pi_\theta(y_l|x, y_{<l}) = \prod_{l=1}^{|y|} \text{softmax}(f_\theta(x, y_{<l}))_{y_l}, \quad (5.2)$$

where $f_\theta : \mathcal{V} \rightarrow \mathbb{R}^{|\mathcal{V}|}$ is parameterized by θ and predicts the l -th token from the preceding context $y_{<l}$, and the next token follows $\text{softmax}(z)_v := \exp(z_v) / \sum_{v' \in \mathcal{V}} \exp(z_{v'})$ with $z = f_\theta(x, y_{<l})$.

Because the full expressive power of practical models renders end-to-end optimization guarantees analytically intractable [3, 23], we follow a well-established line of theoretical work [27, 24, 10, 45] and adopt a *tabular policy* parameterization. This is the special case of Eq. 5.2 in which each output is assigned its own trainable logit indexed by the preceding token; that is, for $\theta \in \mathbb{R}^{|\mathcal{V}| \times |\mathcal{V}|}$,

$$\text{(tabular policy)} \quad \pi_\theta(y_l|x, y_{<l}) = \pi_\theta(y_l|y_{l-1}) = \text{softmax}(\theta_{:,y_{l-1}})_{y_l}, \quad (5.3)$$

where $\theta_{:,y_{l-1}} \in \mathbb{R}^{|\mathcal{V}|}$ is the column of θ associated with the preceding token y_{l-1} . This parameterization retains the softmax structure of the general policy while rendering the training dynamics analytically tractable, and thereby affords provable statements that would otherwise be out of reach.

Optimization for RL. To analyze the training dynamics, we combine the tabular policy of Eq. 5.3 with the RL objective $\phi_{\text{RL}}(\theta)$ of Eq. 3.1 and study the policy gradient in the small-learning-rate limit, that is, the gradient flow

$$\frac{d}{dt} \theta(t) = \nabla \phi_{\text{RL}}(\theta(t)), \quad t \geq 0 \quad (5.4)$$

where $\theta(t)$ denotes the parameters of the policy $\pi_{\theta(t)}$ at training time t , initialized at $\theta(0) = \theta_{\text{ref}}$.

¹For simplicity, we omit the normalizing coefficient $1/|y|$ from Eq. 3.3, which can be treated as a constant $1/2$ in our framework.

Our experiments established that the success rate of each individual reasoning pattern stays essentially fixed throughout training. We attribute this invariance to an inherent asymmetry in optimization difficulty: reshaping the mapping from questions to reasoning patterns is markedly easier than altering the mapping from reasoning patterns to final answers. We formalize this empirically grounded observation as the following assumption.

Assumption 5.1. *The success rate for each reasoning pattern $\mathbf{r}_i \in \mathcal{R}$ to provide the correct answer $\mathbf{r}^*$, defined as $\pi_\theta(\mathbf{a}^*|\mathbf{q}, \mathbf{r}_i)$, remains constant during training.*

Throughout the remainder of this section we adopt Assumption 5.1 and write $p^*(\mathbf{r}) = \pi_\theta(\mathbf{a}^*|\mathbf{q}, \mathbf{r})$ for the fixed success rate of pattern $\mathbf{r}$. We proceed in two steps: we first derive the optimal policy for the RL objective (Eq. 3.1) under the general autoregressive policy (Eq. 5.2), which identifies the *destination* of optimization, and then characterize the full RLVR and RLIF training dynamics under the tabular policy (Eq. 5.3), which identifies the *path* taken to reach it.

5.2 The Optimal Policy for RL

Our experiments showed that RLVR improves performance steadily by promoting higher-success-rate reasoning patterns, whereas RLIF yields only unstable, non-monotone gains. We now explain this contrast at its root by characterizing the optimal policy toward which each objective drives the model.

Proposition 5.2. *Suppose we maximize the RL objective (Eq. 3.1) using a general autoregressive policy (Eq. 5.2) and Assumption 5.1 holds. Then, the optimal policy satisfies:*

$$\pi_{\text{opt}}(\mathbf{r}|\mathbf{q}) = \frac{1}{Z} \exp\left(\frac{1}{\beta} R(\mathbf{r})\right) \pi_{\theta_{\text{ref}}}(\mathbf{r}|\mathbf{q}) \text{ for all } \mathbf{r} \in \mathcal{R}, \quad (5.5)$$

where $Z = \sum_{\mathbf{r} \in \mathcal{R}} \exp\left(\frac{1}{\beta} R(\mathbf{r})\right) \pi_{\theta_{\text{ref}}}(\mathbf{r}|\mathbf{q})$ is the normalizing coefficient, and $R(\mathbf{r})$ denotes the reasoning path reward. Specifically:

- For RLVR, $R(\mathbf{r})$ equals the success rate of reasoning pattern $\mathbf{r}$, i.e., $R_{\text{RLVR}}(\mathbf{r}) = p^*(\mathbf{r})$.
- for RLIF, $R(\mathbf{r})$ equals to the confidence for the reasoning pattern $\mathbf{r}$ for final answer distribution, i.e., $R_{\text{RLIF}}(\mathbf{r}) = -\frac{1}{|\mathcal{A}|} \sum_{\mathbf{a} \in \mathcal{A}} \log(\pi_\theta(\mathbf{a}|\mathbf{q}, \mathbf{r}))$.

Proposition 5.2 pinpoints the optimal solutions for RLVR and RLIF and, in doing so, explains their divergent behavior at its source. Assuming π_{opt} lies within the general autoregressive policy space—which holds for both practical LLMs and the tabular policies analyzed below—the RLVR-optimal policy up-weights a reasoning pattern $\mathbf{r}$ precisely when its probability of deducing the correct answer is high enough that $\exp(\frac{1}{\beta} p^*(\mathbf{r}))/Z > 1$; the reward is thus *aligned with correctness*. The RLIF reward $R_{\text{RLIF}}(\mathbf{r})$, by contrast, depends only on the confidence of the final-answer distribution and therefore cannot distinguish the correct answer $\mathbf{a}^*$ from confidently held incorrect alternatives. This single structural difference is decisive: the RLVR optimum reliably improves accuracy, whereas the RLIF optimum carries no such guarantee and may favor patterns that are confident yet wrong.

Moreover, since β is typically small in practice (e.g., 0.001) [26, 5, 44], taking the limit $\beta \rightarrow 0$ sharpens this picture into a deterministic selection rule:

$$\pi_{\text{opt}}(\mathbf{r}|\mathbf{q}) = \lim_{\beta \rightarrow 0} \frac{1}{Z} \exp\left(\frac{1}{\beta} R(\mathbf{r})\right) \pi_{\theta_{\text{ref}}}(\mathbf{r}|\mathbf{q}) = \begin{cases} 1 & \mathbf{r} = \arg \max_{\mathbf{r}} R(\mathbf{r}) \\ 0 & \text{otherwise} \end{cases}, \quad (5.6)$$

In this limit the policy collapses onto the single reasoning pattern with the largest reward $R(\mathbf{r})$, independently of how the reference model initially distributes probability across patterns. This characterizes the *destination* of optimization, but not the *path* to it: because the objective landscape is non-convex, the optimum alone does not reveal how reliably RLVR reaches high-reward patterns, nor why RLIF can improve performance before deteriorating. We resolve both questions in the next two subsections by analyzing the full training dynamics under the tabular policy (Eq. 5.3).

5.3 Training Dynamic Analysis for RLVR

Our experiments established that the magnitude of RLVR’s gains depends strongly on the capacity of the base model. We now explain this dependence analytically by tracking, under gradient flow,

the probability assigned to the *optimal reasoning pattern*—the pattern $\mathbf{r}^* = \arg \max_{\mathbf{r}} p^*(\mathbf{r})$ with the highest success rate for reaching the correct answer $\mathbf{a}^*$. The dynamics split cleanly into two regimes. In the first, $\pi_{\theta}(\mathbf{r}^*|\mathbf{q})$ increases monotonically and converges to 1. In the second, the model must first pass through an *entanglement stage*, in which a suboptimal pattern $\mathbf{r}' \neq \mathbf{r}^*$ absorbs probability mass and stalls the growth of this selection probability; only after this stage does the model transition into the first-regime dynamics and ultimately converge to $\mathbf{r}^*$. The following two theorems render each regime, and the boundary between them, precise.

Theorem 5.3 (Regime 1: Sufficient Condition for Efficient Convergence). *Consider the RLVR (Eq 3.2) with $\beta = 0$ for optimizing objective Eq 3.1, using a tabular policy (Eq 5.3) with Assumption 5.1 holds. Let $\mathbf{r}^*$ be the optimal reasoning pattern, if the overall accuracy of the initialized model $\pi_{\theta_{\text{ref}}}$, defined as $\text{ACC}_{\theta_{\text{ref}}} = \sum_{\mathbf{r} \in \mathcal{R}} \pi_{\theta_{\text{ref}}}(\mathbf{r}|\mathbf{q})p^*(\mathbf{r})$, satisfies:*

$$(\text{Regime 1}) \quad \text{ACC}_{\theta_{\text{ref}}} > p^*(\mathbf{r}) \text{ for all } \mathbf{r} \in \mathcal{R}, \mathbf{r} \neq \mathbf{r}^*, \quad (5.7)$$

then for any $\epsilon > 0$, there exists $T_1 = \mathcal{O}(\frac{1}{\epsilon})$ such that for $t > T_1$, we have $1 - \pi_{\theta(t)}(\mathbf{r}^*|\mathbf{q}) < \epsilon$.

Theorem 5.3 isolates the favorable regime: the base reference model is strong enough that its overall accuracy already exceeds the success rate of every non-optimal reasoning pattern (Eq 5.7). Under this condition RLVR converges to the optimal pattern at the fast rate $\mathcal{O}(1/\epsilon)$, so that the model rapidly attains an overall accuracy close to the success rate of $\mathbf{r}^*$.

The complementary—and empirically well-documented—regime arises when the base model is weaker, a setting in which RLVR optimization is known to be difficult [44, 36]. Here the model first enters an entanglement stage in which the *second-best pattern*, $\mathbf{r}' = \arg \max_{\mathbf{r}, \mathbf{r} \neq \mathbf{r}^*} p^*(\mathbf{r})$, competes for probability mass and thereby slows the optimization of $\mathbf{r}^*$. Theorem 5.4 quantifies precisely how long this stage persists.

Theorem 5.4 (Regime 2: Slow Convergence for optimal reasoning pattern). *Consider the RLVR (Eq 3.2) with $\beta = 0$ for optimizing objective Eq 3.1, using a tabular policy (Eq 5.3) with Assumption 5.1 holds. Let $\mathbf{r}^*$ and $\mathbf{r}'$ be the optimal and second optimal reasoning patterns, if the overall accuracy of the initialized model satisfies:*

$$(\text{Regime 2}) \quad p^*(\mathbf{r}') > \text{ACC}_{\theta_{\text{ref}}} > p^*(\mathbf{r}) \text{ for all } \mathbf{r} \in \mathcal{R}/\{\mathbf{r}^*, \mathbf{r}'\}, \quad (5.8)$$

then there exists:

$$T_0 = \frac{1}{2 - 2\pi_{\theta_{\text{ref}}}(\mathbf{r}'|\mathbf{q})} \left((C_1 \cdot \gamma_{\pi_{\text{ref}}})^{2C_2 \cdot \gamma_{\pi_{\text{ref}}}} - 1 \right), \text{ where } \gamma_{\pi_{\text{ref}}} := \sum_{\mathbf{r} \in \mathcal{R}/\{\mathbf{r}'\}} \frac{\pi_{\theta_{\text{ref}}}(\mathbf{r}|\mathbf{q})}{\pi_{\theta_{\text{ref}}}(\mathbf{r}^*|\mathbf{q})} \quad (5.9)$$

with constants C_1, C_2 depending on the success rates of reasoning patterns, such that for we can guarantee the model transform from regime 2 (Eq 5.8) to regime 1 (Eq 5.7), i.e., $\text{ACC}_{\theta(t)} > p^*(\mathbf{r}), \forall \mathbf{r} \in \mathcal{R}, \mathbf{r} \neq \mathbf{r}^*$ for $t \geq T_0$.

Theorem 5.4 analyzes the representative case in which only the optimal and second-best patterns have success rates exceeding the overall accuracy. We adopt this setting for analytical clarity; the argument extends, at the cost of heavier bookkeeping, to the more general situation in which *at least* two patterns exceed the average success rate.

The central consequence of Theorem 5.4 is a sharp separation of time scales. Reaching the favorable regime of Theorem 5.3—where the overall accuracy surpasses the success rate of the second-best pattern and the optimal pattern is selected with substantial probability—requires T_0 steps. Whereas the subsequent convergence time T_1 of Theorem 5.3 is merely polynomial in $1/\epsilon$, the transition time T_0 can grow *super-exponentially* in $\gamma_{\pi_{\text{ref}}}$, the ratio between the total selection-weighted success mass carried by $\mathcal{R}/\{\mathbf{r}'\}$ and that carried by $\mathbf{r}^*$. When the base model places only vanishing probability on the optimal pattern, $\gamma_{\pi_{\text{ref}}}$ grows large and T_0 becomes prohibitive, so the model can remain stuck for an extended period before it begins to select $\mathbf{r}^*$ with appreciable probability—this is precisely the entanglement stage. The result thus furnishes a rigorous explanation for why RLVR succeeds almost effortlessly on strong base models yet stalls on weak ones, identifying base-model reasoning quality as the decisive determinant of convergence.

We corroborate both regimes with case studies and numerical simulations in Section 6.

5.4 Theoretical Explanation for RLIF

Proposition 5.2 already establishes that RLIF, being driven solely by the confidence of the final-answer distribution, cannot distinguish the correct answer $\mathbf{a}^*$ from incorrect alternatives, and hence that its optimum may lie below the base model in accuracy—consistent with the degradation we observe empirically. Yet prior work reports that RLIF can improve performance without any external reward Agarwal et al. [4], Zhao et al. [47]. We reconcile these seemingly opposing facts by analyzing RLIF’s full dynamics, demonstrating that both the early improvement and the eventual collapse follow from the same theory as two phases of a single trajectory.

We begin with the early stage, considering a well-trained base model that satisfies the following assumption.

Assumption 5.5. *For the base model, the correct answer $\mathbf{a}^*$ has the highest probability across all possible answers, i.e., $\arg \max_{\mathbf{a} \in \mathcal{A}} \sum_{\mathbf{r} \in \mathcal{R}} \pi_{\text{base}}(\mathbf{a}|\mathbf{q}, \mathbf{r}) \pi_{\text{base}}(\mathbf{r}|\mathbf{q}) = \mathbf{a}^*$.*

Because majority voting [35] is an established means of improving model performance, a well-trained LLM typically assigns the correct answer $\mathbf{a}^*$ higher probability than any incorrect one, so Assumption 5.5 is well justified for modern LLMs. Under this assumption we analyze the accuracy dynamics at initialization ($t = 0$) in the case $|\mathcal{A}| = 2$, with per-pattern success rates drawn uniformly and a high-entropy pattern-selection distribution², obtaining the following theorem.

Theorem 5.6 (RLIF Increases Overall Accuracy at Initialization). *Consider the RLIF (Eq 5.1) with $\beta = 0$ for optimizing objective Eq 3.1, using a tabular policy (Eq 5.3) with Assumption 5.1 and 5.5 hold, $\pi_{\theta_{\text{base}}}(\mathbf{r}|\mathbf{q}) = \frac{1}{|\mathcal{R}|}$ for all $\mathbf{r} \in \mathcal{R}$, $|\mathcal{A}| = 2$ and success rate for each reasoning path follows a uniform distribution $p^*(\mathbf{r}) \sim U[0, 1]$. Then when $|\mathcal{R}| \rightarrow +\infty$, the following holds:*

1. *The accuracy derivative at $t = 0$ is positive with probability 1, i.e. $P(\frac{d}{dt} \text{ACC}_{\theta(t)}|_{t=0} > 0) = 1$.*
2. *With probability $p = 0.5$, that $\arg \max_{\mathbf{r}} R_{\text{RLIF}}(\mathbf{r}) = \arg \min_{\mathbf{r}} p^*(\mathbf{r})$.*

The two results of Theorem 5.6 together resolve the apparent paradox surrounding RLIF. The first shows that the accuracy derivative at initialization is positive with probability one, which precisely explains RLIF’s early performance gains. The second exposes the failure mode that inevitably follows: as training proceeds, Proposition 5.2 implies that at $\beta = 0$ the policy converges deterministically to the pattern maximizing $R_{\text{RLIF}}(\mathbf{r})$, and with probability 0.5 this maximizer coincides with the least accurate pattern, $\arg \max_{\mathbf{r}} R_{\text{RLIF}}(\mathbf{r}) = \arg \min_{\mathbf{r}} p^*(\mathbf{r})$. Thus RLIF improves the model at the outset yet, with probability one-half, ultimately converges toward the lowest-accuracy reasoning path, driving performance below that of the base model. This rise-then-collapse trajectory is exactly the behavior we characterize theoretically here and observe empirically—now derived from first principles rather than merely reported.

6 Case Studies and Numerical Simulations

To validate and concretely illustrate the theory, we present case studies and training simulations that instantiate each theorem. For RLVR, we realize both regimes of Theorems 5.3 and 5.4, shown in Fig. 2 (upper). For RLIF, we run repeated training simulations in which each reasoning path’s success rate is drawn uniformly, directly probing RLIF’s behavior at both the initial training steps and convergence, shown in Fig. 2 (bottom). In every case the simulated dynamics match the theoretical predictions, confirming that the framework captures the observed behavior rather than merely being consistent with it.

6.1 Case Studies for RLVR

Section 5.3 established two regimes in RLVR’s training dynamics: in the first (Eq 5.7), the probability $\pi_{\theta}(\mathbf{r}^*|\mathbf{q})$ of selecting the optimal pattern rises monotonically to convergence; in the second (Eq 5.8), an initial entanglement phase forces a substantially longer training period before the first-regime dynamics take over and the model converges to $\mathbf{r}^*$. The case studies confirm both regimes precisely:

²Previous work [33] reveals that RL primarily optimizes high-entropy tokens. Here, we consider the highest-entropy case: $\pi_{\theta_{\text{base}}}(\mathbf{r}|\mathbf{q}) = 1/|\mathcal{R}|$ for all $\mathbf{r} \in \mathcal{R}$.




Figure 2: Empirical validation of the theory. *Upper (RLVR case studies)*: the two regimes of Theorems 5.3 and 5.4—rapid convergence to the optimal pattern (left) and the entanglement stage that delays it (right). *Bottom (RLIF simulations)*: as $|\mathcal{R}|$ varies, the probability of converging to an improved versus a degraded state (left) and the probability of an initial performance gain (right). The simulations reproduce both predictions of Theorem 5.6 quantitatively: near-certain early improvement together with a near one-half chance of eventual degradation.

- When the reference model satisfies Eq 5.7—its overall accuracy exceeding the success rate of every non-optimal pattern—the model converges rapidly to the optimal reasoning pattern, as shown in Fig. 2 (upper left).
- When instead a suboptimal pattern r' has a higher success rate than the initial overall accuracy (Theorem 5.4), the model first traverses the entanglement phase, and the transition period T_0 (Theorem 5.4) must elapse before the rapid first-regime optimization begins. Fig. 2 (upper right) instantiates a large- $\gamma_{\pi_{\text{ref}}}$ scenario (with $\gamma_{\pi_{\text{ref}}}$ defined in Eq 5.9, here $\gamma_{\pi_{\text{ref}}} = 6$, since $(0.25 + 0.05)/0.05 = 6$); because T_0 scales as $\gamma_{\pi_{\text{ref}}}^{\gamma_{\pi_{\text{ref}}}}$, the transition becomes prohibitively long and convergence to the optimal pattern is markedly delayed. The observed pattern dynamics exhibit exactly this entanglement stage, in close quantitative agreement with the theory.

6.2 Numerical Simulations For RLIF

Section 5.4 established that, under Assumption 5.5 and the base-model distribution of Theorem 5.6, the accuracy derivative at $t = 0$ is positive with probability 1 as $|\mathcal{R}| \rightarrow \infty$, while the probability of converging to the least accurate reasoning path is exactly 0.5. Our simulations confirm both predictions quantitatively. Fig. 2 (bottom left) reports RLIF training simulations across varying $|\mathcal{R}|$: the probability of converging to an improved state ($ACC \uparrow$) closely tracks that of converging to a degraded state ($ACC \downarrow$), matching the second result of Theorem 5.6. Fig. 2 (bottom right) further shows that, at the initial training step, the probability of performance improvement approaches 1 as $|\mathcal{R}|$ grows, matching the first result. The simulations thus reproduce both the early gain and the eventual coin-flip between improvement and degradation that the theory predicts, closing the loop between our formal analysis and RLIF’s empirically observed non-monotone behavior.

7 Conclusions and Limitations

This work delivers a unified account of reinforcement learning dynamics in LLMs, coupling systematic empirical analysis with a rigorous theoretical framework within a single coherent explanation. Viewing RL through the lens of reasoning-pattern selection, we establish that RL reshapes the distribution over reasoning patterns through sparse, token-level interventions while leaving each pattern’s intrinsic success rate intact, and we translate this observation into provable training dynamics for both dominant reward mechanisms: for RLVR we characterize the fast-convergence and entanglement regimes and identify the base model’s reasoning quality as the decisive factor governing convergence, while for RLIF we explain both its early gains and its eventual degradation as two phases of one trajectory. Case studies and numerical simulations confirm every prediction. Our analysis rests on

deliberate simplifications that simultaneously mark the natural directions for future work: the interpretability of LLM-identified reasoning patterns merits further validation, the two-stage framework can be extended to richer real-world reasoning settings, and the dynamics results can be generalized beyond the specific base-model assumptions adopted here. We regard these not as limitations of the account but as opportunities to broaden a framework that, in its present form, already accounts for the full range of RL training behaviors observed across our experiments.

References

- [1] MAA. 2023. American mathematics competitions. https://artofproblemsolving.com/wiki/index.php/American_Mathematics_Competitions, 2023. Online.
- [2] Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- [3] Alekh Agarwal, Sham M Kakade, Jason D Lee, and Gaurav Mahajan. On the theory of policy gradient methods: Optimality, approximation, and distribution shift. *Journal of Machine Learning Research*, 22(98):1–76, 2021.
- [4] Shivam Agarwal, Zimin Zhang, Lifan Yuan, Jiawei Han, and Hao Peng. The unreasonable effectiveness of entropy minimization in llm reasoning. *arXiv preprint arXiv:2505.15134*, 2025.
- [5] Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, et al. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*, 2022.
- [6] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- [7] MAA Codeforces. American invitational mathematics examination-aime 2024, 2024.
- [8] Gheorghe Comanici, Eric Bieber, Mike Schaekermann, Ice Pasupat, Noveen Sachdeva, Inderjit Dhillon, Marcel Blistein, Ori Ram, Dan Zhang, Evan Rosen, et al. Gemini 2.5: Pushing the frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities. *arXiv preprint arXiv:2507.06261*, 2025.
- [9] Ganqu Cui, Yuchen Zhang, Jiacheng Chen, Lifan Yuan, Zhi Wang, Yuxin Zuo, Haozhan Li, Yuchen Fan, Huayu Chen, Weize Chen, Zhiyuan Liu, Hao Peng, Lei Bai, Wanli Ouyang, Yu Cheng, Bowen Zhou, and Ning Ding. The Entropy Mechanism of Reinforcement Learning for Reasoning Language Models, May 2025.
- [10] Ganqu Cui, Yuchen Zhang, Jiacheng Chen, Lifan Yuan, Zhi Wang, Yuxin Zuo, Haozhan Li, Yuchen Fan, Huayu Chen, Weize Chen, et al. The entropy mechanism of reinforcement learning for reasoning language models. *arXiv preprint arXiv:2505.22617*, 2025.
- [11] Nirjhar Das, Souradip Chakraborty, Aldo Pacchiano, and Sayak Ray Chowdhury. Active preference optimization for sample efficient rlhf. In *ICML 2024 Workshop on Theoretical Foundations of Foundation Models*, 2024.
- [12] Amir Dembo. *Large deviations techniques and applications*. Springer, 2009.
- [13] Kanishk Gandhi, Ayush Chakravarthy, Anikait Singh, Nathan Lile, and Noah D Goodman. Cognitive behaviors that enable self-improving reasoners, or, four habits of highly effective stars. *arXiv preprint arXiv:2503.01307*, 2025.
- [14] Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, et al. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.

- [15] Chaoqun He, Renjie Luo, Yuzhuo Bai, Shengding Hu, Zhen Leng Thai, Junhao Shen, Jinyi Hu, Xu Han, Yujie Huang, Yuxiang Zhang, et al. Olympiadbench: A challenging benchmark for promoting agi with olympiad-level bilingual multimodal scientific problems. *arXiv preprint arXiv:2402.14008*, 2024.
- [16] Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the math dataset. *arXiv preprint arXiv:2103.03874*, 2021.
- [17] Maggie Huan, Yuetai Li, Tuney Zheng, Xiaoyu Xu, Seungone Kim, Minxin Du, Radha Pooven-
dran, Graham Neubig, and Xiang Yue. Does Math Reasoning Improve General LLM Capabili-
ties? Understanding Transferability of LLM Reasoning, July 2025.
- [18] Jiawei Huang, Bingcong Li, Christoph Dann, and Niao He. Can rlhf be more efficient with im-
perfect reward models? a policy coverage perspective. In *ICLR 2025 Workshop on Bidirectional
Human-AI Alignment*, 2025.
- [19] Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark,
AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card. *arXiv
preprint arXiv:2410.21276*, 2024.
- [20] Aaron Jaech, Adam Kalai, Adam Lerer, Adam Richardson, Ahmed El-Kishky, Aiden Low, Alec
Helyar, Aleksander Madry, Alex Beutel, Alex Carney, et al. Openai o1 system card. *arXiv
preprint arXiv:2412.16720*, 2024.
- [21] Kaixuan Ji, Jiafan He, and Quanquan Gu. Reinforcement learning from human feedback with
active queries. *arXiv preprint arXiv:2402.09401*, 2024.
- [22] Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze
Brahman, Lester James V Miranda, Alisa Liu, Nouha Dziri, Shane Lyu, et al. Tulu 3: Pushing
frontiers in open language model post-training. *arXiv preprint arXiv:2411.15124*, 2024.
- [23] Gen Li, Yuting Wei, Yuejie Chi, Yuantao Gu, and Yuxin Chen. Softmax policy gradient methods
can take exponential time to converge. In *Conference on Learning Theory*, pages 3107–3110.
PMLR, 2021.
- [24] Jincheng Mei, Chenjun Xiao, Csaba Szepesvari, and Dale Schuurmans. On the global con-
vergence rates of softmax policy gradient methods. In *International conference on machine
learning*, pages 6820–6829. PMLR, 2020.
- [25] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin,
Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to
follow instructions with human feedback. *Advances in neural information processing systems*,
35:27730–27744, 2022.
- [26] Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and
Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model.
Advances in Neural Information Processing Systems, 36:53728–53741, 2023.
- [27] Noam Razin, Zixuan Wang, Hubert Strauss, Stanley Wei, Jason D Lee, and Sanjeev Arora.
What makes a reward model a good teacher? an optimization perspective. *arXiv preprint
arXiv:2503.15477*, 2025.
- [28] Antoine Scheid, Etienne Boursier, Alain Durmus, Michael I Jordan, Pierre M  nard, Eric
Moulines, and Michal Valko. Optimal design for reward modeling in rlhf. *arXiv preprint
arXiv:2410.17055*, 2024.
- [29] Rulin Shao, Shuyue Stella Li, Rui Xin, Scott Geng, Yiping Wang, Sewoong Oh, Simon Shao-
lei Du, Nathan Lambert, Sewon Min, Ranjay Krishna, et al. Spurious rewards: Rethinking training
signals in rlvr. *arXiv preprint arXiv:2506.10947*, 2025.
- [30] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang,
Mingchuan Zhang, YK Li, Y Wu, et al. Deepseekmath: Pushing the limits of mathematical
reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.

- [31] Guangming Sheng, Chi Zhang, Zilingfeng Ye, Xibin Wu, Wang Zhang, Ru Zhang, Yanghua Peng, Haibin Lin, and Chuan Wu. Hybridflow: A flexible and efficient rlhf framework. *arXiv preprint arXiv: 2409.19256*, 2024.
- [32] Qwen Team. Qwq-32b: Embracing the power of reinforcement learning, March 2025. URL <https://qwenlm.github.io/blog/qwq-32b/>.
- [33] Shenzhi Wang, Le Yu, Chang Gao, Chuji Zheng, Shixuan Liu, Rui Lu, Kai Dang, Xionghui Chen, Jianxin Yang, Zhenru Zhang, Yuqiong Liu, An Yang, Andrew Zhao, Yang Yue, Shiji Song, Bowen Yu, Gao Huang, and Junyang Lin. Beyond the 80/20 Rule: High-Entropy Minority Tokens Drive Effective Reinforcement Learning for LLM Reasoning, June 2025.
- [34] Shenzhi Wang, Le Yu, Chang Gao, Chuji Zheng, Shixuan Liu, Rui Lu, Kai Dang, Xionghui Chen, Jianxin Yang, Zhenru Zhang, et al. Beyond the 80/20 rule: High-entropy minority tokens drive effective reinforcement learning for llm reasoning. *arXiv preprint arXiv:2506.01939*, 2025.
- [35] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models. *arXiv preprint arXiv:2203.11171*, 2022.
- [36] Tian Xie, Zitian Gao, Qingnan Ren, Haoming Luo, Yuqian Hong, Bryan Dai, Joey Zhou, Kai Qiu, Zhirong Wu, and Chong Luo. Logic-rl: Unleashing llm reasoning with rule-based reinforcement learning. *arXiv preprint arXiv:2502.14768*, 2025.
- [37] Wei Xiong, Hanze Dong, Chenlu Ye, Ziqi Wang, Han Zhong, Heng Ji, Nan Jiang, and Tong Zhang. Iterative preference learning from human feedback: Bridging theory and practice for rlhf under kl-constraint. In *International Conference on Machine Learning*, pages 54715–54754. PMLR, 2024.
- [38] Shusheng Xu, Wei Fu, Jiakuan Gao, Wenjie Ye, Weilin Liu, Zhiyu Mei, Guangju Wang, Chao Yu, and Yi Wu. Is dpo superior to ppo for llm alignment? a comprehensive study. In *International Conference on Machine Learning*, pages 54983–54998. PMLR, 2024.
- [39] An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiayi Yang, Jingren Zhou, Junyang Lin, Kai Dang, Keming Lu, Keqin Bao, Kexin Yang, Le Yu, Mei Li, Mingfeng Xue, Pei Zhang, Qin Zhu, Rui Men, Runji Lin, Tianhao Li, Tingyu Xia, Xingzhang Ren, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yu Wan, Yuqiong Liu, Zeyu Cui, Zhenru Zhang, and Zihan Qiu. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.
- [40] An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- [41] Qiyang Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Tiantian Fan, Gaohong Liu, Lingjun Liu, Xin Liu, et al. Dapo: An open-source llm reinforcement learning system at scale. *arXiv preprint arXiv:2503.14476*, 2025.
- [42] Yang Yue, Zhiqi Chen, Rui Lu, Andrew Zhao, Zhaokai Wang, Shiji Song, and Gao Huang. Does reinforcement learning really incentivize reasoning capacity in llms beyond the base model? *arXiv preprint arXiv:2504.13837*, 2025.
- [43] Yu Yue, Yufeng Yuan, Qiyang Yu, Xiaochen Zuo, Ruofei Zhu, Wenyuan Xu, Jiaze Chen, Chengyi Wang, Tiantian Fan, Zhengyin Du, et al. Vapo: Efficient and reliable reinforcement learning for advanced reasoning tasks. *arXiv preprint arXiv:2504.05118*, 2025.
- [44] Weihao Zeng, Yuzhen Huang, Qian Liu, Wei Liu, Keqing He, Zejun Ma, and Junxian He. Simplerl-zoo: Investigating and taming zero reinforcement learning for open base models in the wild. *arXiv preprint arXiv:2503.18892*, 2025.

- 563 [45] Yanzhi Zhang, Zhaoxi Zhang, Haoxiang Guan, Yilin Cheng, Yitong Duan, Chen Wang, Yue
564 Wang, Shuxin Zheng, and Jiyan He. No Free Lunch: Rethinking Internal Feedback for LLM
565 Reasoning, June 2025.
- 566 [46] Rosie Zhao, Alexandru Meterez, Sham Kakade, Cengiz Pehlevan, Samy Jelassi, and Eran
567 Malach. Echo chamber: RL post-training amplifies behaviors learned in pretraining. *arXiv*
568 *preprint arXiv:2504.07912*, 2025.
- 569 [47] Xuandong Zhao, Zhewei Kang, Aosong Feng, Sergey Levine, and Dawn Song. Learning to
570 reason without external rewards. *arXiv preprint arXiv:2505.19590*, 2025.
- 571 [48] Banghua Zhu, Michael Jordan, and Jiantao Jiao. Principled reinforcement learning with human
572 feedback from pairwise or k-wise comparisons. In *International Conference on Machine*
573 *Learning*, pages 43037–43067. PMLR, 2023.
- 574 [49] Qihao Zhu, Daya Guo, Zhihong Shao, Dejian Yang, Peiyi Wang, Runxin Xu, Y Wu, Yukun
575 Li, Huazuo Gao, Shirong Ma, et al. Deepseek-coder-v2: Breaking the barrier of closed-source
576 models in code intelligence. *arXiv preprint arXiv:2406.11931*, 2024.

A Additional Experiments

A.1 Pattern Distribution Dynamics on RL-enhanced Qwen2.5-3B

Section 4.2 analyzed the reasoning-pattern dynamics on Number Theory tasks from MATH. To confirm that this behavior is not task-specific but reflects a general property of the two reward regimes, we repeat the dynamic analysis on Geometry tasks (Figure 3). The same contrast recurs unchanged: under RLVR the model steadily concentrates probability on the higher-accuracy patterns, whereas under RLIF pattern selection remains unstable. This replication is fully consistent with the empirical and theoretical results of the main paper and reinforces that the mechanism we identify governs RL training across tasks.




Figure 3: Reasoning-pattern distribution and corresponding per-pattern accuracy of RL-enhanced models on Geometry tasks. Under RLVR the model converges toward the highest-accuracy pattern (Pattern 1), whereas under RLIF the pattern distribution remains unstable—precisely mirroring the Number Theory dynamics reported in the main paper.

A.2 Qwen-2.5-7B-Instruct vs. Qwen-2.5-7B-SimpleRL-Zoo

To establish that our findings hold across model families rather than being an artifact of a single base model, we conduct supplementary experiments that validate both the empirical insights and the theoretical conclusions (Figure 4). Here we take Qwen-2.5-7B-Instruct [39] as the base model and compare it against Qwen-2.5-7B-SimpleRL-Zoo [44], an open-source variant enhanced with RLVR for mathematical reasoning. Our evaluation spans a broad range of mathematical domains—number theory, geometry, algebra, calculus, and counting and probability—using challenging problems sampled from the MATH dataset [16], and we further extend it to complex problems drawn from AMC23 [1]. Across every domain the results consistently reinforce our earlier findings: RLVR increases the frequency of the high-accuracy reasoning patterns while suppressing less effective ones. Task-specific analyses are provided in Appendix B.3.

Across these tasks, the highest-success-rate patterns consistently become more prevalent after RLVR enhancement, reinforcing the findings of Section 4 and Section 5. Notably, in tasks such as Algebra and AMC23-19 (Figure 4), the per-pattern success rates remain stable across the base and RL-enhanced models; this not only echoes our earlier observations but also furnishes direct empirical support for Assumption 5.1, on which our theoretical analysis rests.

A.3 Qwen-2.5-32B-Instruct vs. QwQ-32B

To demonstrate that our theory scales to substantially larger models, we further experiment on QwQ-32B [32], which is built on Qwen-2.5-32B-Instruct [39] and markedly strengthens reasoning through RLVR. We measure the reasoning-pattern shift between these two models on four tasks: number theory and geometry tasks from the MATH dataset [16], a task from AIME 2024 [7], and a task from OlympiadBench [15].

The result is unambiguous: the reasoning patterns of QwQ-32B predominantly coincide with the most accurate patterns of Qwen-2.5-32B, exactly as our theory predicts and consistent with the behavior observed at smaller scale. We omit per-pattern accuracies for QwQ because API-only access




Figure 4: Reasoning-pattern distributions and corresponding success rates, with and without RLVR enhancement, across additional tasks; the results align closely with both our experimental insights and our theoretical conclusions. Bars (Pattern Dist) give each pattern’s proportion among all patterns, and dots (Pattern Acc) give its corresponding per-pattern accuracy or success rate. In every task, RLVR shifts probability mass toward the highest-accuracy patterns.

precludes extracting the CoT data needed to have the base model continue from a fixed reasoning prefix—the procedure required to attribute accuracy to a reasoning pattern rather than to the model as a whole; see the “Accuracy Analysis” in Appendix B.2 for details.




Figure 5: Evaluation on the larger models Qwen-2.5-32B and QwQ-32B. QwQ’s reasoning patterns converge onto the most accurate pattern of Qwen-2.5, demonstrating that our theory extends cleanly to large-scale models.

B Experimental Details

B.1 Experimental Setup

Entire Training Procedures: We take Qwen2.5-3B as the base model and train it on the MATH dataset under both RLVR and RLIF. All runs are conducted with the verl [31] framework under hyperparameter settings identical to verl’s provided MATH example; RLIF differs from RLVR solely in the reward computation, which isolates the reward as the only source of variation between the two paradigms. For each method we obtain three independent training curves by varying the random seed, ensuring that the reported contrast is not seed-dependent.

Token-level Analysis: This experiment employs Qwen2.5-3B together with the RL-enhanced models described above. We randomly sample 400 questions from each of the GSM8K and MATH test sets and use the full AIME24 dataset. For each question we first sample an answer from the base model, then concatenate this answer to the question and feed the result to the RL-enhanced model as a new prompt. Using the OpenAI library interface, we compute, for every token in the prompt, its probability together with its rank among the probabilities of all candidate tokens at that position. Comparing each token’s rank before and after RL yields the proportion of tokens whose rank changed, a direct measure of how sparsely RL intervenes at the token level.

Reasoning-pattern Level Analysis: We conduct two complementary experiments. First, to trace the pattern distribution *during* training, we use our own models—Qwen2.5-3B and its RL-enhanced counterpart—evaluating every 20 steps to resolve the within-training dynamics. Second, to test the generality of our conclusions across distinct and substantially larger base models, we evaluate the open-source Qwen-2.5-7B-SimpleRL-Zoo and QwQ-32B against their respective base models; for

these, only the pre- versus post-RL pattern distributions are accessible, so we examine the induced distribution shift rather than the full within-training trajectory.

B.2 Reasoning Patterns Analysis Pipeline

Model Inference: We randomly sample questions spanning several task types (e.g., Geometry, Algebra) and difficulty levels (Level 4, Level 5) from datasets including MATH. For the within-training pattern-dynamics experiment, the model samples 1024 answers per question; for the experiment that considers only the pre- versus post-RL distribution change, it samples between 64 and 1024 answers per question, depending on the task category. These sampled responses constitute the input to the subsequent pattern-extraction and classification stages.

Reasoning Patterns Extraction: From a sampled subset of base-model responses, we use the GPT-4o [19] API to summarize the pattern categories; the full extraction prompt appears in Appendix D.1. To render extraction deterministic and reproducible, we fix the GPT-4o temperature to 0.

Responses Classification: We again employ GPT-4o, at temperature 0, to assign every response to exactly one of the extracted pattern categories; the full classification prompt is given in Appendix D.2.

Accuracy Analysis: To isolate the effect of the pattern distribution on accuracy from any change in the reasoning process itself, we hold the reasoning process fixed and vary only the completion. Concretely, we take a base-model response, remove the sentence containing the final answer, append the remaining prefix to the question, and use this as a new prompt from which the RL-enhanced model continues to predict the answer. Because the partial response fixes a specific reasoning pattern, the accuracy of the model’s continuation measures the success rate of that reasoning pattern rather than of the model as a whole, thereby allowing us to attribute accuracy changes solely to the pattern distribution.

B.3 Details for Each Task

Below we detail the experimental procedure for each task, together with task descriptions and representative reasoning-pattern examples (Table 2). Example task prompts are provided in Appendix D.3.

Number Theory: This task poses coin-distribution problems over multiple bags: the bags initially hold equal numbers of coins, and after additional coins are received and redistributed so that every bag again holds the same count exceeding a specified total, the model must determine the minimum initial number of coins per bag. We adapt this task from a level-5 MATH problem (id: test/number_theory/1055) [16], generating 32 variants by adjusting parameters such as the bag count and coin totals to form the evaluation set. We adopt one-shot prompting with a simple, unrelated example that fixes the answer format via "boxed" notation without biasing the reasoning approach. For evaluation, we extract reasoning patterns from 20 questions with 4 responses each, sample 64 responses per question, and then classify all responses according to the extracted patterns, which enables a direct comparison of the reasoning-pattern distribution before and after RL.

Geometry: In the geometry task, the model is given the lengths of two sides of a scalene triangle and asked how many distinct integer values (in centimeters) the third side can take. This task is adapted from a level-4 MATH geometry problem (id: test/geometry/1046). Its evaluation follows exactly the same pipeline as Number Theory: we construct a synthetic set of 32 questions, extract reasoning patterns from 20 questions with 4 base-model (Qwen2.5-7b-Instruct) responses each, and then compare the model with and without RLVR enhancement across the full evaluation set at 64 responses per question.

Algebra: The algebra task is a can-recycling problem: starting from a given number of old cans, and given that n old cans can be recycled into one new can, the model must determine how many cans can ultimately be produced. The difficulty lies in the iterative accounting—at each recycling step some cans cannot be divided evenly, and these leftovers may subsequently be combined for further recycling. The template is adapted from a level-4 MATH algebra problem (id: test/algebra/2768). For this task and all subsequent supplementary tasks, we adopt exactly the same settings as the earlier tasks, including dataset size and number of samples.

Task	Reasoning Patterns Examples	Common Elements / Key Words
Number Theory	Modular Congruence with Coefficient Simplification	Uses modular arithmetic to express divisibility conditions.
	Inequality-Driven Search for Minimal Solution	Often involves substituting back to compute the total coins.
Geometry	Systematic Inequality Application	Explicitly lists and solves each of the three inequalities.
	Verification of Scalene Condition	Adding an extra layer of validation to verifies the third side length.
Algebra	Iterative Floor Division with Leftover Tracking.	Uses floor division to compute new cans per step.
	Recursive Recycling with Aggregated Leftovers.	Combines leftovers with newly produced cans before recycling.
Calculus	Direct Simplification and Principal Value Matching	Assumes inputs fall within the principal range.
	Interval Analysis with Case Splitting	Splits the whole domain into different intervals.
Counting and Probability	Direct Probability Setup and Quadratic Solution	Explicitly calculates combinations for total and favorable outcomes.
	Early Simplification and Cross-Multiplication	Early cross-multiplication to eliminate denominators.
AMC23-Q19	Prime Factorization and Simplification	Counting Digits in Large Numbers.
	Rewriting (8^5) as (2^{15}) (15^5) as $(3^5 \cdot 5^5)$.	Final step of counting digits: $(3 \text{ (from 243)} + 15 \text{ (zeros)}) = 18$.
AIME24-Q1	Direct Equation Setup and Elimination	Straightforward and relies on algebraic manipulation.
	Alternative Equation Formulation and Solving	Expressing variables in terms of others early on.
OlympiadBench-Q1631	Direct Calculation and Empirical Testing	Relying on direct computation and empirical verification.
	Factorization-Based Reasoning	Using algebraic factorization to argue certain terms y_n .

Table 2: Representative reasoning-pattern examples across the evaluation tasks, together with their defining keywords and common elements.

Calculus: The calculus task is deliberately compact: the model is given a single-variable trigonometric equation together with the domain of the variable and must determine the number of solutions. Though elementary in form, the task directly probes the model’s basic competence on calculus problems. Its prototype is a level-4 MATH pre-calculus problem (id: test/precalculus/1140).

Counting and Probability: We further evaluate the model on a counting-and-probability task involving white and black balls. Given the probability that, when two balls are drawn at random, one is black and one is white, together with the number of balls of one color, the model must compute the minimum number of balls of the other color. This task is adapted from a level-4 MATH counting-and-probability problem (id: test/counting_and_probability/79).

AMC23 Question: For this task we use question 19 of AMC23 [1]: we randomly sample 48 responses for reasoning-pattern extraction and 1024 responses for the subsequent pattern analysis.

B.4 Details for Experiments on QwQ-32B and Qwen-2.5-32B

We evaluate four tasks. Because API access constrains testing, we select one question per task, sample 64 responses from each of the two models, and perform reasoning-pattern extraction and classification directly on these responses. Specifically, from the MATH dataset we use question 1055 (number theory) and question 1046 (geometry), and for the remaining two tasks we use problem 1 from AIME24 and question 1631 from OlympiadBench. Representative patterns are reported in Table 2.

702 C Deferred Proofs

703 This appendix provides the complete proofs of our main theoretical results: Proposition 5.2 and
704 Theorems 5.3, 5.4, and 5.6.

705 C.1 Proof of Proposition 5.2

706 *Proof of Proposition 5.2.* In this proof, we utilize the proof techniques in [26], recall that the opti-
707 mization objective of RL is

$$\phi_{\text{RL}}(\boldsymbol{\theta}) = \mathbb{E}_{\mathbf{x} \sim \mathcal{D}, \mathbf{y} \sim \pi_{\boldsymbol{\theta}}(\mathbf{y}|\mathbf{x})} [r_{\phi}(\mathbf{x}, \mathbf{y})] - \beta \mathbb{D}_{\text{KL}} [\pi_{\boldsymbol{\theta}}(\mathbf{y} | \mathbf{x}) \parallel \pi_{\text{ref}}(\mathbf{y} | \mathbf{x})], \quad (\text{C.1})$$

708 Under our framework, as state in section 5.1, we consider a policy conduct reasoning by first sample
709 $\mathbf{r}_i \in \mathcal{R}$ based on $\pi_{\boldsymbol{\theta}}(\mathbf{r}_i|\mathbf{q})$ and then provide the final answer $\mathbf{a} \in \mathcal{A}$ by $\pi_{\boldsymbol{\theta}}(\mathbf{a}|\mathbf{r}_i)$, the reward $\phi_{\text{RL}}(\boldsymbol{\theta})$
710 can be written as

$$\phi_{\text{RL}}(\boldsymbol{\theta}) = \mathbb{E}_{\mathbf{y} \sim \pi_{\boldsymbol{\theta}}(\mathbf{r}|\mathbf{q})} \left[\sum_{\mathbf{a} \in \mathcal{A}} \pi_{\text{ref}}(\mathbf{a} | \mathbf{r}) r_{\phi}([\mathbf{r}, \mathbf{a}], q) \right] - \beta \mathbb{D}_{\text{KL}} [\pi_{\boldsymbol{\theta}}(\mathbf{y} | \mathbf{x}) \parallel \pi_{\text{ref}}(\mathbf{y} | \mathbf{x})], \quad (\text{C.2})$$

711 Let $R(\mathbf{r}) = \sum_{\mathbf{a} \in \mathcal{A}} \pi_{\text{ref}}(\mathbf{a} | \mathbf{r}) r_{\phi}([\mathbf{r}, \mathbf{a}], q)$, then

$$\begin{aligned} \phi_{\text{RL}}(\boldsymbol{\theta}) &= \mathbb{E}_{\mathbf{r} \sim \pi_{\boldsymbol{\theta}}(\cdot|\mathbf{q}), \mathbf{a} \sim \pi_{\boldsymbol{\theta}}(\cdot|\mathbf{r})} \left[R(\mathbf{r}) - \beta \ln \left(\frac{\pi_{\boldsymbol{\theta}}(\mathbf{r}|\mathbf{q}) \cdot \pi_{\text{ref}}(\mathbf{a}|\mathbf{r})}{\pi_{\text{ref}}(\mathbf{r}|\mathbf{q}) \cdot \pi_{\text{ref}}(\mathbf{a}|\mathbf{r})} \right) \right] \\ &= \mathbb{E}_{\mathbf{r} \sim \pi_{\boldsymbol{\theta}}(\cdot|\mathbf{q})} \left[R(\mathbf{r}) - \beta \ln \left(\frac{\pi_{\boldsymbol{\theta}}(\mathbf{r}|\mathbf{q})}{\pi_{\text{ref}}(\mathbf{r}|\mathbf{q})} \right) \right] \\ &= -\beta \mathbb{E}_{\mathbf{r} \sim \pi_{\boldsymbol{\theta}}(\cdot|\mathbf{q})} \left[\ln \left(\frac{\pi_{\boldsymbol{\theta}}(\mathbf{r}|\mathbf{q})}{\frac{1}{Z} \exp \left(\frac{1}{\beta} R(\mathbf{r}) \right) \pi_{\text{ref}}(\mathbf{r}|\mathbf{q})} \right) - \ln Z \right] \\ &= -\beta \mathbb{D}_{\text{KL}} [\pi^*(\mathbf{r}|\mathbf{q}) \parallel \pi_{\text{ref}}(\mathbf{y} | \mathbf{x})] + \beta \ln Z. \end{aligned}$$

712 Where the third equation is by Assumption 5.1 and $Z = \sum_{\mathbf{r} \in \mathcal{R}} \exp \left(\frac{1}{\beta} p^*(\mathbf{r}) \right) \pi_{\boldsymbol{\theta}_{\text{ref}}}(\mathbf{r}|\mathbf{q})$ is the
713 partition constant that ensures

$$\pi^*(\mathbf{r}|\mathbf{q}) = \frac{1}{Z} \exp \left(\frac{1}{\beta} R(\mathbf{r}) \right) \pi_{\boldsymbol{\theta}}(\mathbf{r}|\mathbf{q}),$$

714 is a valid probability distribution such that $\sum_{\mathbf{r} \in \mathcal{R}} \pi^*(\mathbf{r}|\mathbf{q}) = 1$. Since Z is not a function of $\mathbf{r}$.

715 Therefore, maximizing the objective in Equation C.2 is equivalent to:

$$\begin{aligned} &\max_{\pi} \mathbb{E}_{\mathbf{x} \sim \mathcal{D}, \mathbf{y} \sim \pi(\mathbf{y}|\mathbf{x})} [r_{\phi}(\mathbf{x}, \mathbf{y})] - \beta \mathbb{D}_{\text{KL}} [\pi(\mathbf{y} | \mathbf{x}) \parallel \pi_{\text{ref}}(\mathbf{y} | \mathbf{x})] \\ &= \min_{\pi} \beta \mathbb{D}_{\text{KL}} [\pi^*(\mathbf{r}|\mathbf{q}) \parallel \pi_{\text{ref}}(\mathbf{y} | \mathbf{x})] - \beta \ln Z. \end{aligned}$$

716 By the properties of KL-divergence, we know that the optimal policy for the KL-constrained reward
717 maximization objective satisfies:

$$\pi_{\text{opt}}(\mathbf{r}|\mathbf{q}) = \frac{1}{Z} \exp \left(\frac{1}{\beta} R(\mathbf{r}) \right) \pi_{\boldsymbol{\theta}_{\text{ref}}}(\mathbf{r}|\mathbf{q}) \text{ for all } \mathbf{r} \in \mathcal{R}.$$

718 specifically, for RLVR, $R_{\text{RLVR}}(\mathbf{r}) = p^*$, for RLIF, we have

$$\begin{aligned} R'(\mathbf{r}) &= - \left(\frac{1}{|\mathcal{R}|} \sum_{\mathbf{r}_i \in \mathcal{R}} \log (|\mathcal{R}| \cdot \pi_{\boldsymbol{\theta}}(\mathbf{r}_i|q)) + \frac{1}{|\mathcal{A}|} \sum_{\mathbf{a}_j \in \mathcal{A}} \log (|\mathcal{A}| \cdot \pi_{\boldsymbol{\theta}}(\mathbf{a}_j|q, \mathbf{r})) \right) \\ &= - \frac{1}{|\mathcal{A}|} \sum_{\mathbf{a}_j \in \mathcal{A}} \log (\pi_{\boldsymbol{\theta}}(\mathbf{a}_j|q, \mathbf{r})) - \underbrace{\left(\frac{1}{|\mathcal{R}|} \sum_{\mathbf{r}_i \in \mathcal{R}} \log (|\mathcal{R}| \cdot \pi_{\boldsymbol{\theta}}(\mathbf{r}_i|q)) + \log |\mathcal{A}| \right)}_c, \end{aligned}$$

719 where the second term is the same for all $\mathbf{r} \in |\mathcal{R}|$, for RLIF, the optimal policy can be written as

$$\begin{aligned}\pi_{opt}(\mathbf{r}|\mathbf{q}) &= \frac{1}{Z} \exp\left(\frac{1}{\beta} R'(\mathbf{r})\right) \pi_{\theta_{ref}}(\mathbf{r}|\mathbf{q}) \\ &= \frac{\exp\left(\frac{1}{\beta} R'(\mathbf{r})\right) \pi_{\theta_{ref}}(\mathbf{r}|\mathbf{q})}{\sum_{\mathbf{r}_i \in \mathcal{R}} \exp\left(\frac{1}{\beta} (R'(\mathbf{r}_i))\right) \pi_{\theta_{ref}}(\mathbf{r}_i|\mathbf{q})} \\ &= \frac{\exp\left(\frac{1}{\beta} (R'(\mathbf{r}) - c)\right) \pi_{\theta_{ref}}(\mathbf{r}|\mathbf{q})}{\sum_{\mathbf{r}_i \in \mathcal{R}} \exp\left(\frac{1}{\beta} (R'(\mathbf{r}_i) - c)\right) \pi_{\theta_{ref}}(\mathbf{r}_i|\mathbf{q})}\end{aligned}$$

720 So we can write $R_{RLIF}(\mathbf{r}) = R'(\mathbf{r}) - c = -\frac{1}{|\mathcal{A}|} \sum_{\mathbf{a}_j \in \mathcal{A}} \log(\pi_{\theta}(\mathbf{a}_j|\mathbf{q}, \mathbf{r}))$. This concludes the proof
721 of the theorem. $\square$

722 C.2 Dynamics of Tabular Policy

723 Consider the RL objective (Eq 3.1), using a tabular policy (Eq 5.3) with Assumption 5.1 holds, the
724 gradient dynamics of θ can be computed as

$$\begin{aligned}\frac{d}{dt} \theta(t) &= \nabla \phi_{RL}(\theta(t)) \\ &= \nabla \mathbb{E}_{\mathbf{r} \sim \pi_{\theta}(\cdot|\mathbf{q})} [r_{\phi}^{KL}(\mathbf{r}|\mathbf{q})] \\ &= \sum_{\mathbf{r} \in \mathcal{R}} [r_{\phi}^{KL}(\mathbf{r}|\mathbf{q}) \nabla \pi_{\theta(t)}(\mathbf{r}|\mathbf{q})],\end{aligned}$$

725 Where $\pi_{\theta(t)}(\mathbf{r}|\mathbf{q}) = \text{softmax}(\theta_{:,q})_{\mathbf{r}}$, $\theta_{:,q} \in \mathbb{R}^{|\mathcal{V}|}$ is the column of θ corresponding to $\mathbf{q}$, and
726 $r_{\phi}(\mathbf{r}|\mathbf{q})^{KL}$ refers to the reward with KL divergence, so

$$\begin{aligned}\frac{\partial \theta_{r_i,q}(t)}{\partial t} &= \sum_{\mathbf{r}_j \in \mathcal{R}} [r_{\phi}^{KL}(\mathbf{r}_j|\mathbf{q}) (\nabla \text{softmax}(\theta_{:,q})_{\mathbf{r}_j})_{\mathbf{r}_i}] \\ &= \sum_{\mathbf{r}_j \in \mathcal{R}} [r_{\phi}^{KL}(\mathbf{r}_j|\mathbf{q}) (-\pi_{\theta(t)}(\mathbf{r}_j|\mathbf{q}) \cdot \pi_{\theta(t)}(\mathbf{r}_i|\mathbf{q}))] \\ &\quad + r_{\phi}^{KL}(\mathbf{r}_i|\mathbf{q}) \cdot \pi_{\theta(t)}(\mathbf{r}_i|\mathbf{q})\end{aligned}$$

727 where the overall accuracy is defined as $\text{ACC}_{\theta(t)} = \sum_{\mathbf{r} \in \mathcal{R}} \pi_{\theta(t)}(\mathbf{r}|\mathbf{q}) p^*(\mathbf{r})$. For $\pi_{\theta(t)}(\mathbf{r}_i|\mathbf{q})$, we
728 have

$$\frac{d}{dt} \pi_{\theta(t)}(\mathbf{r}_i|\mathbf{q}) = \sum_{\mathbf{r}_j \in \mathcal{R}} \frac{\partial \pi_{\theta(t)}(\mathbf{r}_i|\mathbf{q})}{\partial \theta_{\mathbf{r}_j,q}(t)} \cdot \frac{\partial \theta_{\mathbf{r}_j,q}(t)}{\partial t}$$

729 C.3 Proof of Theorem 5.3

730 The proof of Theorem 5.3 proceeds in three steps: we first establish that the overall accuracy is non-
731 decreasing, $\text{ACC}_{\theta(t)} \geq \text{ACC}_{\theta(0)}$ for $t \geq 0$ (Eq C.4); we then derive a lower bound on $\pi_{\theta(t)}(\mathbf{r}^*|\mathbf{q})$
732 (Eq C.5); and we finally obtain the convergence-time bound on t stated in the theorem.

733 *Proof of Theorem 5.3.* Setting $\beta = 0$, Eq C.3 becomes:

$$\frac{\partial \theta_{r_i,q}(t)}{\partial t} = \pi_{\theta(t)}(\mathbf{r}_i|\mathbf{q}) \cdot (p^*(\mathbf{r}_i) - \text{ACC}_{\theta(t)}). \quad (\text{C.3})$$

734 For $\pi_{\theta(t)}(\mathbf{r}_i|\mathbf{q})$, we derive:

$$\begin{aligned}\frac{d}{dt} \pi_{\theta(t)}(\mathbf{r}_i|\mathbf{q}) &= \sum_{\mathbf{r}_j \in \mathcal{R}} \frac{\partial \pi_{\theta(t)}(\mathbf{r}_i|\mathbf{q})}{\partial \theta_{\mathbf{r}_j,q}(t)} \cdot \frac{\partial \theta_{\mathbf{r}_j,q}(t)}{\partial t} \\ &= \sum_{\mathbf{r}_j \in \mathcal{R}} (\text{ACC}_{\theta(t)} - p^*(\mathbf{r}_j)) \pi_{\theta(t)}^2(\mathbf{r}_j|\mathbf{q}) \pi_{\theta(t)}(\mathbf{r}_i|\mathbf{q}) \\ &\quad + (p^*(\mathbf{r}_i) - \text{ACC}_{\theta(t)}) \pi_{\theta(t)}^2(\mathbf{r}_i|\mathbf{q}),\end{aligned}$$

735 and for the accuracy $\text{ACC}_{\theta(t)}$:

$$\begin{aligned}
\frac{d}{dt} \text{ACC}_{\theta(t)} &= \sum_{\mathbf{r}_i \in \mathcal{R}} p^*(\mathbf{r}_i) \cdot \frac{d}{dt} \pi_{\theta(t)}(\mathbf{r}_i | \mathbf{q}) \\
&= \sum_{\mathbf{r}_i, \mathbf{r}_j \in \mathcal{R}} p^*(\mathbf{r}_i) \cdot (\text{ACC}_{\theta(t)} - p^*(\mathbf{r}_j)) \pi_{\theta(t)}^2(\mathbf{r}_j | \mathbf{q}) \pi_{\theta(t)}(\mathbf{r}_i | \mathbf{q}) \\
&\quad + \sum_{\mathbf{r}_i \in \mathcal{R}} (p^*(\mathbf{r}_i) - \text{ACC}_{\theta(t)}) \pi_{\theta(t)}^2(\mathbf{r}_i | \mathbf{q}) \\
&= \sum_{\mathbf{r}_j \in \mathcal{R}} \text{ACC}_{\theta(t)} \cdot (\text{ACC}_{\theta(t)} - p^*(\mathbf{r}_j)) \pi_{\theta(t)}^2(\mathbf{r}_j | \mathbf{q}) \\
&\quad + \sum_{\mathbf{r}_i \in \mathcal{R}} (p^*(\mathbf{r}_i) - \text{ACC}_{\theta(t)}) \pi_{\theta(t)}^2(\mathbf{r}_i | \mathbf{q}) p^*(\mathbf{r}_i) \\
&= \sum_{\mathbf{r}_i \in \mathcal{R}} \pi_{\theta(t)}^2(\mathbf{r}_i | \mathbf{q}) (p^*(\mathbf{r}_i) - \text{ACC}_{\theta(t)})^2 \geq 0. \tag{C.4}
\end{aligned}$$

736 As $\frac{d}{dt} \text{ACC}_{\theta(t)} \geq 0$, so $\text{ACC}_{\theta(t)} > \text{ACC}_{\theta(0)}$ holds for $t \geq 0$, for the optimal reasoning pattern
737 $\mathbf{r}^*$, let $\mathbf{r}'$ the suboptimal reasoning pattern with the second highest success rate, defined as $\mathbf{r}' =$
738 $\arg \max_{\mathbf{r}, \mathbf{r} \neq \mathbf{r}^*} p^*(\mathbf{r})$, we have:

$$\begin{aligned}
\frac{d}{dt} \pi_{\theta(t)}(\mathbf{r}^* | \mathbf{q}) &= \sum_{\mathbf{r}_j \in \mathcal{R}} (\text{ACC}_{\theta(t)} - p^*(\mathbf{r}_j)) \pi_{\theta(t)}^2(\mathbf{r}_j | \mathbf{q}) \pi_{\theta(t)}(\mathbf{r}^* | \mathbf{q}) \\
&\quad + (p^*(\mathbf{r}^*) - \text{ACC}_{\theta(t)}) \pi_{\theta(t)}^2(\mathbf{r}^* | \mathbf{q}) \\
&\geq (\text{ACC}_{\theta(t)} - p^*(\mathbf{r}^*)) \pi_{\theta(t)}^2(\mathbf{r}^* | \mathbf{q}) \pi_{\theta(t)}(\mathbf{r}^* | \mathbf{q}) \\
&\quad + (p^*(\mathbf{r}^*) - \text{ACC}_{\theta(t)}) \pi_{\theta(t)}^2(\mathbf{r}^* | \mathbf{q}) \\
&= (p^*(\mathbf{r}^*) - \text{ACC}_{\theta(t)}) \pi_{\theta(t)}^2(\mathbf{r}^* | \mathbf{q}) (1 - \pi_{\theta(t)}(\mathbf{r}^* | \mathbf{q})) \\
&\geq (p^*(\mathbf{r}^*) - p^*(\mathbf{r}')) \pi_{\theta(t)}^2(\mathbf{r}^* | \mathbf{q}) (1 - \pi_{\theta(t)}(\mathbf{r}^* | \mathbf{q}))^2 \geq 0. \tag{C.5}
\end{aligned}$$

739 As $\frac{d}{dt} \pi_{\theta(t)}(\mathbf{r}^* | \mathbf{q}) \geq 0$, $\pi_{\theta(t)}(\mathbf{r}^* | \mathbf{q}) \geq \pi_{\theta(0)}(\mathbf{r}^* | \mathbf{q})$ holds for $t \geq 0$, define $C =$
740 $(p^*(\mathbf{r}^*) - p^*(\mathbf{r}')) \pi_{\theta(0)}^2(\mathbf{r}^* | \mathbf{q})$:

$$\frac{d}{dt} \pi_{\theta(t)}(\mathbf{r}^* | \mathbf{q}) \geq C (1 - \pi_{\theta(t)}(\mathbf{r}^* | \mathbf{q}))^2.$$

741 This differential inequality yields:

$$\pi_{\theta(t)}(\mathbf{r}^* | \mathbf{q}) > 1 - \frac{1}{Ct + \frac{1}{1 - \pi_{\theta(0)}(\mathbf{r}^* | \mathbf{q})}}.$$

742 For any $\epsilon > 0$, there exists $T_1 = \frac{1}{C} \left(\frac{1}{\epsilon} - \frac{1}{1 - \pi_{\theta(0)}(\mathbf{r}^* | \mathbf{q})} \right) = \mathcal{O}(\frac{1}{\epsilon})$, such that

$$\pi_{\theta(t)}(\mathbf{r}^* | \mathbf{q}) > 1 - \epsilon.$$

743 This concludes the proof of the theorem. □

744 C.4 Proof of Theorem 5.4

745 The proof of Theorem 5.4 proceeds as follows: we first derive an upper bound on $\pi_{\theta(t)}(\mathbf{r}' | \mathbf{q})$ (Eq C.6),
746 then analyze the ratio $\rho_{\theta(t)}(\mathbf{r}_i) = \frac{\pi_{\theta(t)}(\mathbf{r}_i | \mathbf{q})}{\pi_{\theta(t)}(\mathbf{r}^* | \mathbf{q})}$ (Eq C.7), and thereby establish that $\text{ACC}_{\theta(t)} \geq p^*(\mathbf{r}')$
747 for all $t > T_0$.

748 *Proof of Theorem 5.4.* Let $\mathbf{r}^*$ and $\mathbf{r}'$ be the optimal and second optimal reasoning patterns. For
 749 $\pi_{\theta(t)}(\mathbf{r}'|\mathbf{q})$:

$$\begin{aligned}
 \frac{d}{dt}\pi_{\theta(t)}(\mathbf{r}'|\mathbf{q}) &= \sum_{\mathbf{r}_j \in \mathcal{R}} (\text{ACC}_{\theta(t)} - p^*(\mathbf{r}_j)) \pi_{\theta(t)}^2(\mathbf{r}_j|\mathbf{q}) \pi_{\theta(t)}(\mathbf{r}'|\mathbf{q}) \\
 &\quad + (p^*(\mathbf{r}') - \text{ACC}_{\theta(t)}) \pi_{\theta(t)}^2(\mathbf{r}'|\mathbf{q}) \\
 &= \sum_{\mathbf{r}_j \in \mathcal{R}/\{\mathbf{r}^*\}} (\text{ACC}_{\theta(t)} - p^*(\mathbf{r}_j)) \pi_{\theta(t)}^2(\mathbf{r}_j|\mathbf{q}) \pi_{\theta(t)}(\mathbf{r}'|\mathbf{q}) \\
 &\quad + (p^*(\mathbf{r}') - \text{ACC}_{\theta(t)}) \pi_{\theta(t)}^2(\mathbf{r}'|\mathbf{q}) (1 - \pi_{\theta(t)}(\mathbf{r}^*|\mathbf{q})) \\
 &\leq \sum_{\mathbf{r}_j \in \mathcal{R}/\{\mathbf{r}^*\}} \pi_{\theta(t)}^2(\mathbf{r}_j|\mathbf{q}) \pi_{\theta(t)}(\mathbf{r}'|\mathbf{q}) \\
 &\quad + (p^*(\mathbf{r}') - p^*(\mathbf{r}') \pi_{\theta(t)}(\mathbf{r}^*|\mathbf{q})) \pi_{\theta(t)}^2(\mathbf{r}'|\mathbf{q}) (1 - \pi_{\theta(t)}(\mathbf{r}^*|\mathbf{q})) \\
 &\leq \pi_{\theta(t)}(\mathbf{r}'|\mathbf{q}) (1 - \pi_{\theta(t)}(\mathbf{r}^*|\mathbf{q}))^2 + p^*(\mathbf{r}') \pi_{\theta(t)}^2(\mathbf{r}'|\mathbf{q}) (1 - \pi_{\theta(t)}(\mathbf{r}^*|\mathbf{q}))^2 \\
 &\leq 2 (1 - \pi_{\theta(t)}(\mathbf{r}^*|\mathbf{q}))^2.
 \end{aligned}$$

750 This yields:

$$\pi_{\theta(t)}(\mathbf{r}'|\mathbf{q}) \leq 1 - \frac{1}{2t + \frac{1}{1 - \pi_{\theta(0)}(\mathbf{r}'|\mathbf{q})}}. \quad (\text{C.6})$$

751 Next, define $\rho_{\theta(t)}(\mathbf{r}_i) = \frac{\pi_{\theta(t)}(\mathbf{r}_i|\mathbf{q})}{\pi_{\theta(t)}(\mathbf{r}^*|\mathbf{q})}$, then for all $\mathbf{r}_i \in \mathcal{R}/\{\mathbf{r}^*, \mathbf{r}'\}$, we have

$$\begin{aligned}
 \frac{d}{dt}\rho_{\theta(t)}(\mathbf{r}_i) &= \frac{\pi_{\theta(t)}(\mathbf{r}^*|\mathbf{q}) \frac{d}{dt}\pi_{\theta(t)}(\mathbf{r}_i|\mathbf{q}) - \pi_{\theta(t)}(\mathbf{r}_i|\mathbf{q}) \frac{d}{dt}\pi_{\theta(t)}(\mathbf{r}^*|\mathbf{q})}{\pi_{\theta(t)}^2(\mathbf{r}^*|\mathbf{q})} \\
 &= \frac{\pi_{\theta(t)}(\mathbf{r}_i|\mathbf{q})}{\pi_{\theta(t)}(\mathbf{r}^*|\mathbf{q})} ((p^*(\mathbf{r}_i) - \text{ACC}_{\theta(t)}) \pi_{\theta(t)}(\mathbf{r}_i|\mathbf{q}) - (p^*(\mathbf{r}^*) - \text{ACC}_{\theta(t)}) \pi_{\theta(t)}(\mathbf{r}^*|\mathbf{q})) \\
 &\leq -\rho_{\theta(t)}(\mathbf{r}_i) (p^*(\mathbf{r}^*) - \text{ACC}_{\theta(t)}) \pi_{\theta(t)}(\mathbf{r}^*|\mathbf{q}) \leq 0
 \end{aligned}$$

752 The last inequality is based on the condition for case 2 and $\text{ACC}_{\theta(t)} > \text{ACC}_{\theta(0)}$, so $p^*(\mathbf{r}_i) <$
 753 $\text{ACC}_{\theta(t)}$ holds for all $\mathbf{r}_i \in \mathcal{R}/\{\mathbf{r}^*\}$, $t \geq 0$. Then as $\frac{d}{dt}\rho_{\theta(t)}(\mathbf{r}_i) \leq 0$, we have:

$$\pi_{\theta(t)}(\mathbf{r}^*|\mathbf{q}) = \frac{1 - \pi_{\theta(t)}(\mathbf{r}'|\mathbf{q})}{\sum_{\mathbf{r}_i \in \mathcal{R}/\{\mathbf{r}'\}} \rho_{\theta(t)}(\mathbf{r}_i)} \geq \frac{1 - \pi_{\theta(t)}(\mathbf{r}'|\mathbf{q})}{\sum_{\mathbf{r}_i \in \mathcal{R}/\{\mathbf{r}'\}} \rho_{\theta(0)}(\mathbf{r}_i)}$$

754 While $\text{ACC}_{\theta(t)} < p^*(\mathbf{r}')$, define $\gamma_{\pi_{\text{ref}}} := \sum_{\mathbf{r} \in \mathcal{R}/\{\mathbf{r}'\}} \frac{\pi_{\theta_{\text{ref}}}(\mathbf{r}|\mathbf{q})}{\pi_{\theta_{\text{ref}}}(\mathbf{r}^*|\mathbf{q})}$ we have:

$$\begin{aligned}
 \frac{d}{dt}\rho_{\theta(t)}(\mathbf{r}_i) &\leq -\rho_{\theta(t)}(\mathbf{r}_i) (p^*(\mathbf{r}^*) - \text{ACC}_{\theta(t)}) \pi_{\theta(t)}(\mathbf{r}^*|\mathbf{q}) \\
 &\leq -\rho_{\theta(t)}(\mathbf{r}_i) (p^*(\mathbf{r}^*) - p^*(\mathbf{r}')) \frac{1 - \pi_{\theta(t)}(\mathbf{r}'|\mathbf{q})}{\gamma_{\pi_{\text{ref}}}} \\
 &\leq -\rho_{\theta(t)}(\mathbf{r}_i) \frac{p^*(\mathbf{r}^*) - p^*(\mathbf{r}')}{\gamma_{\pi_{\text{ref}}}} \frac{1}{2t + \frac{1}{1 - \pi_{\theta(0)}(\mathbf{r}'|\mathbf{q})}}
 \end{aligned}$$

755 Then we have:

$$\frac{d}{dt}\rho_{\theta(t)}(\mathbf{r}_i) \leq -\rho_{\theta(t)}(\mathbf{r}_i) \frac{p^*(\mathbf{r}^*) - p^*(\mathbf{r}')}{\gamma_{\pi_{\text{ref}}}} \frac{1}{2t + \frac{1}{1 - \pi_{\theta(0)}(\mathbf{r}'|\mathbf{q})}}, \quad (\text{C.7})$$

756 this differential inequality leads to:

$$\rho_{\theta(t)}(\mathbf{r}_i) \leq \frac{(1 - \pi_{\theta(0)}(\mathbf{r}'|\mathbf{q}))^{-\Delta/(2\gamma_{\pi_{\text{ref}}})}}{(1/(1 - \pi_{\theta(0)}(\mathbf{r}'|\mathbf{q})) + 2t)^{\Delta/(2\gamma_{\pi_{\text{ref}}})}} \rho_{\theta(0)}(\mathbf{r}_i).$$

757 Where $\Delta = p^*(\mathbf{r}^*) - p^*(\mathbf{r}')$. Recall that:

$$\begin{aligned} \text{ACC}_{\theta(t)} - p^*(\mathbf{r}') &= \sum_{\mathbf{r} \in \mathcal{R}} \pi_{\theta(t)}(\mathbf{r}|\mathbf{q}) p^*(\mathbf{r}) - p^*(\mathbf{r}') \\ &= p^*(\mathbf{r}^*) \left[(p^*(\mathbf{r}^*) - p^*(\mathbf{r}')) + \sum_{\mathbf{r}_i \in \mathcal{R}/\{\mathbf{r}'\}} \rho_{\theta(t)}(\mathbf{r}_i) (p^*(\mathbf{r}_i) - p^*(\mathbf{r}')) \right] \\ &\geq p^*(\mathbf{r}^*) \left[(p^*(\mathbf{r}^*) - p^*(\mathbf{r}')) - \sum_{\mathbf{r}_i \in \mathcal{R}/\{\mathbf{r}'\}} \rho_{\theta(t)}(\mathbf{r}_i) p^*(\mathbf{r}') \right]. \end{aligned}$$

758 Let $C_2 = 1/\Delta, C_1 = p^*(\mathbf{r}')/\Delta$, define:

$$T_0 = \frac{1}{2 - 2\pi_{\theta_{\text{ref}}}(\mathbf{r}'|\mathbf{q})} \left((C_1 \cdot \gamma_{\pi_{\text{ref}}})^{2C_2 \cdot \gamma_{\pi_{\text{ref}}}} - 1 \right),$$

759 then:

$$\begin{aligned} &\sum_{\mathbf{r}_i \in \mathcal{R}/\{\mathbf{r}'\}} \rho_{\theta(T_0)}(\mathbf{r}_i) p^*(\mathbf{r}') \\ &\leq \frac{(1 - \pi_{\theta(0)}(\mathbf{r}'|\mathbf{q}))^{-\Delta/(2\gamma_{\pi_{\text{ref}}})}}{(1/(1 - \pi_{\theta(0)}(\mathbf{r}'|\mathbf{q})) + 2T_0)^{\Delta/(2\gamma_{\pi_{\text{ref}}})}} \sum_{\mathbf{r}_i \in \mathcal{R}/\{\mathbf{r}'\}} \rho_{\theta(0)}(\mathbf{r}_i) p^*(\mathbf{r}') \\ &\leq \frac{(1 - \pi_{\theta(0)}(\mathbf{r}'|\mathbf{q}))^{-\Delta/(2\gamma_{\pi_{\text{ref}}})}}{\left(1/(1 - \pi_{\theta(0)}(\mathbf{r}'|\mathbf{q})) \cdot \left(\frac{\gamma_{\pi_{\text{ref}}} p^*(\mathbf{r}')}{\Delta} \right)^{2\gamma_{\pi_{\text{ref}}}/\Delta} \right)^{\Delta/(2\gamma_{\pi_{\text{ref}}})}} p^*(\mathbf{r}') \sum_{\mathbf{r}_i \in \mathcal{R}/\{\mathbf{r}'\}} \rho_{\theta(0)}(\mathbf{r}_i) \\ &= \frac{1}{\left(\left(\frac{\gamma_{\pi_{\text{ref}}} p^*(\mathbf{r}')}{\Delta} \right)^{2\gamma_{\pi_{\text{ref}}}/\Delta} \right)^{\Delta/(2\gamma_{\pi_{\text{ref}}})}} p^*(\mathbf{r}') \gamma_{\pi_{\text{ref}}} \\ &= \Delta = p^*(\mathbf{r}^*) - p^*(\mathbf{r}'). \end{aligned}$$

760 Finally, for $t > T_0$ we have:

$$\begin{aligned} \text{ACC}_{\theta(t)} - p^*(\mathbf{r}') &\geq p^*(\mathbf{r}^*) \left[(p^*(\mathbf{r}^*) - p^*(\mathbf{r}')) - \sum_{\mathbf{r}_i \in \mathcal{R}/\{\mathbf{r}'\}} \rho_{\theta(t)}(\mathbf{r}_i) p^*(\mathbf{r}') \right] \\ &\geq p^*(\mathbf{r}^*) [(p^*(\mathbf{r}^*) - p^*(\mathbf{r}')) - (p^*(\mathbf{r}^*) - p^*(\mathbf{r}'))] = 0. \end{aligned}$$

761 This concludes the proof of the theorem. □

762 C.5 Proof of Theorem 5.6

763 *Proof of Theorem 5.6.* First, as we consider the case $|\mathcal{A}| = 2$ and $\pi_{\theta_{\text{base}}}(\mathbf{r}|\mathbf{q}) = \frac{1}{|\mathcal{R}|}$, Assumption
764 5.5 holds means $\frac{1}{|\mathcal{R}|} \sum_{\mathbf{r} \in \mathcal{R}} (p^*(\mathbf{r})) > 0.5$. The derivative of overall accuracy can be written as

$$\begin{aligned} \frac{d}{dt} \text{ACC}_{\theta(t)} &= \sum_{\mathbf{r}_i \in \mathcal{R}} p^*(\mathbf{r}_i) \cdot \frac{d}{dt} \pi_{\theta(t)}(\mathbf{r}_i|\mathbf{q}) \\ &= \sum_{\mathbf{r}_i \in \mathcal{R}} \pi_{\theta(t)}^2(\mathbf{r}_i|\mathbf{q}) (p^*(\mathbf{r}_i) - \text{ACC}_{\theta(t)}) \left(- \sum_{\mathbf{a} \in \mathcal{A}} \log(\pi_{\theta(t)}(\mathbf{a}|\mathbf{r}_i, \mathbf{q})) - \text{Con}_{\theta(t)} \right), \end{aligned}$$

765 where $\text{Con}_{\theta(t)} = -\sum_{\mathbf{r}_i \in \mathcal{R}} \pi_{\theta(t)}(\mathbf{r}_i | \mathbf{q}) \sum_{\mathbf{a} \in \mathcal{A}} \log(\pi_{\theta(t)}(\mathbf{a} | \mathbf{r}_i, q))$. Here we assume $\mathcal{A} = \{\mathbf{a}^*, \mathbf{a}'\}$
 766 where $\mathbf{a}'$ is the incorrect answer, and $\pi_{\theta(0)}(\mathbf{r}_i | \mathbf{q}) = \frac{1}{|\mathcal{R}|}$ for all $\mathbf{r}_i \in \mathcal{R}$. Then:

$$\begin{aligned} \left. \frac{d}{dt} \text{ACC}_{\theta(t)} \right|_{t=0} &= \sum_{\mathbf{r}_i \in \mathcal{R}} \frac{1}{|\mathcal{R}|^2} \left(p^*(\mathbf{r}_i) - \frac{1}{|\mathcal{R}|} \sum_{\mathbf{r}_j \in \mathcal{R}} p^*(\mathbf{r}_j) \right) \cdot \\ &\quad \left(-\sum_{\mathbf{a} \in \mathcal{A}} \log(\pi_{\theta(0)}(\mathbf{a} | \mathbf{r}_i, q)) + \frac{1}{|\mathcal{R}|} \sum_{\mathbf{r}_j \in \mathcal{R}} \sum_{\mathbf{a} \in \mathcal{A}} \log(\pi_{\theta(0)}(\mathbf{a} | \mathbf{r}_j, q)) \right) \\ &= \frac{1}{|\mathcal{R}|^2} \sum_{\mathbf{r}_i \in \mathcal{R}} \left(p^*(\mathbf{r}_i) - \frac{1}{|\mathcal{R}|} \sum_{\mathbf{r}_j \in \mathcal{R}} p^*(\mathbf{r}_j) \right) \cdot \left(-\sum_{\mathbf{a} \in \mathcal{A}} \log(\pi_{\theta(0)}(\mathbf{a} | \mathbf{r}_i, q)) \right). \end{aligned}$$

767 As we consider the case $|\mathcal{R}| \rightarrow +\infty$ and the success rate $p^*(\mathbf{r})$ are i.i.d. samples from $U[0, 1]$.
 768 By the Central limit theorem, for large $|\mathcal{R}|$, the sample mean $\frac{1}{|\mathcal{R}|} \sum_{\mathbf{r} \in \mathcal{R}} p^*(\mathbf{r})$ is concentrated
 769 around the its expectation 0.5, the deviation is of order is $\mathcal{O}(1/\sqrt{|\mathcal{R}|})$. We define this deviation as
 770 $\delta = \mathcal{O}(1/|\mathcal{R}|^{1/2}) > 0$ such that $\frac{1}{|\mathcal{R}|} \sum_{\mathbf{r} \in \mathcal{R}} p^*(\mathbf{r}) = 0.5 + \delta$.

771 According to the Gibbs conditioning principle ([12] Corollary 7.3.5 and Theorem 7.3.8), as $|\mathcal{R}| \rightarrow \infty$,
 772 the empirical distribution of the $p^*(\mathbf{r})$ conditioned on their mean $\frac{1}{|\mathcal{R}|} \sum_{\mathbf{r} \in \mathcal{R}} p^*(\mathbf{r}) = 0.5 + \delta$
 773 converges to the density: $f(x) = 1 + 12\delta(x - 0.5), x \in [0, 1]$. As $|\mathcal{R}| \rightarrow \infty$, we can approximate
 774 the summation over $\mathbf{r}$ with an integral over this limiting distribution $f(x)$. The derivative becomes:

$$\left. \frac{d}{dt} \text{ACC}_{\theta(t)} \right|_{t=0} = \int_0^1 - (1 + 12\delta(x - 0.5)) \cdot (x - 0.5 - \delta) \cdot \log(x(1 - x)) dx = \frac{2}{3}\delta > 0$$

775 Since $\delta > 0$, the derivative is positive. As $|\mathcal{R}| \rightarrow \infty$, this holds with probability 1, proving the
 776 first statement.

To prove the second statement, we first analyze the RLIF reward function. For an action space of size
 $|\mathcal{A}| = 2$, the reward for a path $\mathbf{r}$ is proportional to the negative log-likelihood of the policy:

$$R_{\text{RLIF}}(\mathbf{r}) - \frac{1}{|\mathcal{A}|} \sum_{\mathbf{a}_j \in \mathcal{A}} \log(\pi_{\theta}(\mathbf{a}_j | q, \mathbf{r})) = -\frac{1}{2} \log(p^*(\mathbf{r})(1 - p^*(\mathbf{r})))$$

777 The function $g(p) = p(1 - p)$ is maximized at $p = 0.5$. Therefore, maximizing the reward $R_{\text{RLIF}}(\mathbf{r})$
 778 is equivalent to choosing the path whose success rate $p^*(\mathbf{r})$ is furthest from 0.5.

Let $p_{\min}^* = \min_{\mathbf{r} \in \mathcal{R}} p^*(\mathbf{r})$ and $p_{\max}^* = \max_{\mathbf{r} \in \mathcal{R}} p^*(\mathbf{r})$. The path with the highest reward will be the
 one with success rate $p_{\min}^*$ if it is further from 0.5 than $p_{\max}^*$ is. This condition is expressed as:

$$0.5 - p_{\min}^* > p_{\max}^* - 0.5 \iff p_{\min}^* + p_{\max}^* < 1$$

779 Thus, the event $\arg \max_{\mathbf{r}} R_{\text{RLIF}}(\mathbf{r}) = \arg \min_{\mathbf{r}} p^*(\mathbf{r})$ is equivalent to the event $p_{\min}^* + p_{\max}^* < 1$.

780 We now calculate the probability of this event. The success rates $p^*(\mathbf{r})$ are drawn from the limiting
 781 distribution $f(x) = 1 + 12\delta(x - 0.5)$. From the theory of order statistics, the scaled minimum and
 782 maximum of $n = |\mathcal{R}|$ samples converge in distribution to independent exponential random variables.
 783 Let $Y_n = n \cdot p_{\min}^*$ and $Z_n = n \cdot (1 - p_{\max}^*)$. As $n \rightarrow \infty$:

- 784 • Y_n converges to $Y \sim \text{Exp}(\lambda_Y)$, where the rate is the density at the lower bound: $\lambda_Y = f(0) =$
 785 $1 - 6\delta$.
- 786 • Z_n converges to $Z \sim \text{Exp}(\lambda_Z)$, where the rate is the density at the upper bound: $\lambda_Z = f(1) =$
 787 $1 + 6\delta$.

The condition $p_{\min}^* + p_{\max}^* < 1$ can be rewritten in terms of these asymptotic variables:

$$\frac{Y_n}{n} + \left(1 - \frac{Z_n}{n}\right) < 1 \implies Y_n < Z_n$$

The probability of this event is $P(Y < Z)$. For two independent exponential variables, this is given by:

$$P(Y < Z) = \frac{\lambda_Y}{\lambda_Y + \lambda_Z} = \frac{1 - 6\delta}{(1 - 6\delta) + (1 + 6\delta)} = \frac{1 - 6\delta}{2} = \frac{1}{2} - 3\delta$$

Since $\delta = O(1/|\mathcal{R}|^{1/2})$, in the limit as $|\mathcal{R}| \rightarrow \infty$, $\delta \rightarrow 0$. Therefore, the limiting probability is:

$$\lim_{|\mathcal{R}| \rightarrow \infty} P(\arg \max_{\mathbf{r}} R_{\text{RLIF}}(\mathbf{r}) = \arg \min_{\mathbf{r}} p^*(\mathbf{r})) = \lim_{\delta \rightarrow 0} \left(\frac{1}{2} - 3\delta \right) = \frac{1}{2}$$

788 This concludes the proof of the second statement.

789

□

790 **D Full Prompts**

791 **D.1 Full Prompt for Reasoning Patterns Extraction**

Full Prompt for Reasoning Patterns Extraction

You will analyze multiple AI reasoning processes, showing how different models solve problems.

Analysis Steps:

For each reasoning process, identify:

Key words and recurring phrases

Logical structure of the argument

Problem-solving techniques used

Step-by-step progression

Group similar reasoning processes into exactly 5 patterns based on shared:

Vocabulary patterns (common terms and phrases)

Logical frameworks (how arguments are structured)

Solution approaches (methods used to reach conclusions)

IMPORTANT! You need to categorize general thinking patterns for problem solving, rather than categorizing types of problems themselves. In other words, each type of thinking pattern you give can be applied to any problem, rather than being task-specific. What you need to classify is the way of thinking, not the way of solving problems.

For example, you can classify them into:

1. derivation based on existing knowledge
2. derivation combined with verification
3. independent calculation, suitable for learners starting from the basics,
4. detailed manual derivation, suitable for learners who need to deeply understand each step
5. etc

Or you can classify them into:

1. knowledge call type
2. derivation and verification type
3. etc.

Output Format:

Pattern 1:

Name and description

Defining characteristics

Example processes from input

Common elements

Pattern 2:

Name and description

Defining characteristics

Example processes from input

Common elements

...

There are reasoning processes:

792

{extracted_responses}

Group similar reasoning processes into exactly 5 patterns. Please provide a classification of reasoning patterns, with each group defined by their common characteristics and examples from these processes.

793

794 D.2 Full Prompt for Responses Classification

Full Prompt for Responses Classification

Please classify each reasoning process into one of the following five patterns. Output the classifications in JSON format.

{patterns_categories}

Here are the reasoning processes to classify:

{extracted_responses}

Please output your classification in this JSON format:

795

796 D.3 Prompt Examples for Evaluation Tasks

Prompt Example for Easy Countdown

<lim_start>system

You are a helpful assistant. You first thinks about the reasoning process in the mind and then provides the user with the answer.<lim_end>

<lim_start>

user Using the numbers [1,2,3,7], create an equation that equals 5. You can only use basic arithmetic addition (+) and subtraction (-), and each number from the list can be used at most once. Show your work in <think> </think> tags. And return the final answer in <answer> </answer> tags, for example, to get 5 from [1,2,3,7], we have <answer> 2 + 3 = 5 </answer>; to get 13 from [1,2,3,7], we have <answer> 1 + 2 + 3 + 7 = 13 </answer>.<lim_end>

<lim_start>assistant

Let me solve this step by step.

<think>

797

Prompt Example for Long Arithmetic

<lim_start>system

You are a helpful assistant. You first thinks about the reasoning process in the mind and then provides the user with the answer.<lim_end>

<lim_start>user

Considering the experssion (((((((((16 + 47) + 92) + 84) + 5) + 49) + 82) - 33) - 99) - 69), you need to start from the innermost bracket and work your way out. You need to give the specific calculation process. You have to do a certain amount of calculations on each line, and print each calculation in the format of a + b = c or a - b = c, for example "1 + 2 = 3".<lim_end>

<lim_start>assistant

Let me solve this step by step.

<think>

798

Prompt Example for Number Theory

<lim_start>system

You are a helpful assistant.<lim_end>

<lim_start>user

You have four bags of copper coins. Each bag has the same number of copper coins. One day, you find a bag of 23 coins. You decide to redistribute the number of coins you have

799

so that all five bags you hold have the same number of coins. You successfully manage to redistribute all the coins, and you also note that you have more than 120 coins. What is the smallest number of coins you could have had before finding the bag of 23 coins?
Please reason step by step, and put your final answer within $\boxed{}$.<lim_end>
<lim_start>assistant

800

Prompt Example for Geometry

<lim_start>system
You are a helpful assistant.<lim_end>
<lim_start>user
Two sides of scalene $\triangle ABC$ measure 4 centimeters and 7 centimeters. How many different whole centimeter lengths are possible for the third side?
Please reason step by step, and put your final answer within $\boxed{}$.<lim_end>
<lim_start>assistant

801

Prompt Example for Algebra

<lim_start>system
You are a helpful assistant.<lim_end>
<lim_start>user
Six aluminum cans can be recycled to make a new can. How many new cans can eventually be made from 200 aluminum cans? (Remember that the first new cans that are made can then be recycled into even newer cans!)
Please reason step by step, and put your final answer within $\boxed{}$.<lim_end>
<lim_start>assistant

802

Prompt Example for Calculus

<lim_start>system
You are a helpful assistant.<lim_end>
<lim_start>user
For how many values of x in $[0, 2\pi]$ is $\cos^{-1}(\cos 4x) = \sin^{-1}(\sin x)$?
Please reason step by step, and put your final answer within $\boxed{}$.<lim_end>
<lim_start>assistant

803

Prompt Example for Counting and Probability

<lim_start>system
You are a helpful assistant.<lim_end>
<lim_start>user
7 white balls and k black balls are placed into a bin. Two balls are drawn at random. The probability that one ball is white and the other is black is $\frac{35}{66}$. Find the smallest possible value of k .
Please reason step by step, and put your final answer within $\boxed{}$.<lim_end>
<lim_start>assistant

804

Prompt Example for AMC23-Q19

<lim_start>system
You are a helpful assistant.<lim_end>
<lim_start>user
How many digits are in the base-ten representation of $8^5 \cdot 5^{10} \cdot 15^5$?
Please reason step by step, and put your final answer within $\boxed{}$.<lim_end>
<lim_start>assistant

805

Prompt Example for AIME24-Q1

<lim_start>system
You are a helpful assistant.<lim_end>
<lim_start>user
Every morning Aya goes for a s -kilometer-long walk and stops at a coffee shop afterwards. When she walks at a constant speed of s kilometers per hour, the walk takes her 4 hours, including t minutes spent in the coffee shop. When she walks $s+2$ kilometers per hour, the walk takes her 2 hours and 24 minutes, including t minutes spent in the coffee shop. Suppose Aya walks at $s+\frac{1}{2}$ kilometers per hour. Find the number of minutes the walk takes her, including the t minutes spent in the coffee shop.
Please reason step by step, and put your final answer within $\boxed{}$.<lim_end>
<lim_start>assistant

806

Prompt Example for OlympiadBench-Q1631

<lim_start>system
You are a helpful assistant.<lim_end>
<lim_start>user
For a positive integer a , define a sequence of integers $x_1, x_2, \dots$ by letting $x_1 = a$ and $x_{n+1} = 2x_n + 1$ for $n \geq 1$. Let $y_n = 2^{x_n} - 1$. Determine the largest possible k such that, for some positive integer a , the numbers $y_1, \dots, y_k$ are all prime.
Please reason step by step, and put your final answer within $\boxed{}$.<lim_end>
<lim_start>assistant

807