

SEEDPRINTS: FINGERPRINTS CAN EVEN TELL WHICH SEED YOUR LARGE LANGUAGE MODEL WAS TRAINED FROM

Anonymous authors
Anonymous submission

ABSTRACT

Fingerprinting Large Language Models (LLMs) is becoming a core technical instrument for provenance verification, ownership auditing, and model attribution. Yet most existing fingerprinting methods are validated only after fine-tuning, when exposure to data, optimization history, and hyperparameter choices have already produced stable post-training signatures. Because pretraining is the stage at which foundation models acquire most of their capacity and knowledge, credible lineage verification must also operate before such late-stage signatures are fully formed. We show that current fingerprinting methods are *unreliable* in this fundamental regime: they function primarily as post-hoc identifiers that emerge after substantial training, rather than as intrinsic identity markers present at a model’s origin. This limitation prevents them from satisfying the classical Galtonian standard, under which a fingerprint is both unique and persistent. We introduce **SeedPrints**, a stricter formulation of LLM fingerprinting that uses random-initialization bias as a seed-dependent identifier available before training begins. Untrained models already exhibit reproducible prediction biases determined by their initialization seed, and these weak signals remain statistically detectable after subsequent optimization, enabling calibrated high-confidence lineage verification. Within the evaluated settings, **SeedPrints** verifies identity from initialization through large-scale pretraining and downstream adaptation, including cases in which prior techniques fail during early pretraining or degrade under distribution shift. Experiments on LLaMA-style and Qwen-style models establish seed-level distinguishability and substantiate birth-to-lifecycle identity verification. Evaluations on large-scale pretraining trajectories and real-world fingerprinting benchmarks further demonstrate robustness under prolonged optimization, domain shifts, and parameter modifications.

Together, these results establish random initialization as a source of unique and persistent identity in LLMs, yielding a genuinely “Galtonian” fingerprint. Code will be released upon acceptance.

1 INTRODUCTION

LLM fingerprinting has emerged as a central framework for identifying, attributing, and tracing language models through measurable properties of their parameters, representations, or behavior (Pasquini et al., 2024; Xu et al., 2024; Yoon et al., 2025; Zhang et al., 2024; Zeng et al., 2024). By establishing a verifiable link between a suspicious model and a proposed source model, fingerprinting supplies the empirical infrastructure for provenance auditing and for detecting model theft or unauthorized reuse (Yoon et al., 2025; Zhang, 2025). Its practical force, however, depends on a stronger capability than distinguishing already mature models: a credible fingerprint should verify lineage across the training process during which foundation models acquire their capabilities.

Much of this literature invokes the biological fingerprint metaphor introduced by Francis Galton’s *Finger Prints* (1892) (Galton, 1892):

“A fingerprint is the pattern formed by friction-ridge skin on the fingertips; this ridge configuration is individually unique and essentially permanent across an individual’s lifetime.”

SeedPrints vs Baselines: OLMo-2-7B Pre-training Checkpoints




Figure 1: **Existing fingerprinting methods fail to verify lineage during early pre-training.** Five methods are evaluated on OLMo-2-7B checkpoints spanning 5B to 3.9T training tokens, with the final checkpoint serving as the reference model in every comparison. The y-axis reports the similarity score; larger values indicate stronger lineage evidence, and the dashed line denotes the 0.8 decision threshold used by similarity-based baselines. All baselines deteriorate on early checkpoints and fall below this threshold, whereas **SeedPrints** verifies lineage from the first checkpoint onward ($p \ll 0.001$, plotted as $1 - p$).

The analogy imposes a criterion stronger than ordinary model similarity: a fingerprint should be unique, persistent, and present from the moment of a model’s “birth” at initialization. Existing fingerprinting approaches generally do not satisfy this criterion (Pasquini et al., 2024; Xu et al., 2024; Yoon et al., 2025; Zhang et al., 2024; Zeng et al., 2024; Zhang, 2025; Luan et al., 2025; Tsai et al., 2025; Alhazbi et al., 2025). As Figure 1 shows, they fail to verify lineage reliably during early pretraining. The reason is structural: these methods extract signatures from converged parameters, learned representations, or generated text, and therefore rely on patterns that become separable only after substantial optimization. Their discriminative power consequently reflects training artifacts—including data distribution, optimization trajectory, and hyperparameter effects—rather than an identity mark inherited from the model’s origin. Consistent with this interpretation, Section 5.1 shows that existing baselines often misattribute lineage when continued training substantially changes the data distribution. They are therefore best understood as post hoc identifiers rather than **Galtonian fingerprints**: innate and persistent marks that accompany a model from initialization onward.

This paper advances a stricter, origin-centered notion of LLM fingerprinting: an intrinsic property present at *initialization* and detectable at *any* later point in training or adaptation. The resulting contributions are:

- **A diagnosis of a fundamental limitation in existing methods.** We show that prior fingerprinting approaches do not reliably detect lineage during early pretraining and can be confounded by substantial shifts in training data distribution. (Section 5.1)
- **An initialization-driven fingerprint for LLM lineage.** We show that untrained models exhibit seed-dependent biases in their internal representations, producing weak but persistent identity signals from initialization. (Section 3) We then introduce **SeedPrints**, a statistical test that isolates these initialization-born fingerprints from confounds introduced by data distribution, training duration, and optimization dynamics. (Section 4)
- **Controlled validation of seed-level identity.** We demonstrate that SeedPrints separates models that differ only in initialization seed, including models trained with identical pipelines and data orders, and that the same fingerprint persists under continual training across diverse datasets. (Section 5.1)
- **Birth-to-lifecycle verification at practical scale.** SeedPrints verifies lineage throughout large-scale pretraining trajectories, including OLMo-2-7B Stage 1, and remains competitive in large-scale fine-tuning and deployment-oriented evaluations on *LeaFBench* (Shao et al., 2025), which contains 65 models across 7 mainstream families and 6 transformations: instruction tuning, fine-tuning, PEFT, quantization, model merging, and distillation. (Sections 5.2 and 5.3)

2 RELATED WORK

LLM protection methods fall into two broad families: (i) *watermarking / active fingerprinting* techniques that *insert* an identifiable signature into a model or its outputs (Xu et al., 2024; Nasery et al.; Tsai et al., 2025; Sander et al., 2024), and (ii) *passive fingerprinting* techniques that *extract* a signature from pre-existing model behavior without modifying the model (Yoon et al., 2025; Zhang et al., 2024; Alhazbi et al., 2025; Pasquini et al., 2024; Suzuki et al., 2025; Zhang, 2025). This distinction is decisive for lineage verification: active schemes can create an audit signal when the developer controls training, whereas passive schemes must recover identity from structure already present in the model.

Watermarking and Active Fingerprinting. Active approaches deliberately implant verifiable identifiers for subsequent ownership checks. Classical mechanisms include backdoor attacks (Adi et al., 2018; Li et al., 2019; Zhang et al., 2018), digital signatures, and hash functions (Guo & Potkonjak, 2018; Zhu et al., 2020). For language models, the dominant forms are *text watermarks*, which bias generation or introduce predefined patterns to encode hidden information (Kirchenbauer et al., 2023; Xu et al., 2024; Nasery et al.), and *model weight watermarks*, which embed identifiers in parameters or bind them to secret triggers through fine-tuning (Luan et al., 2025; Li et al., 2023). Backdoor-style fingerprints are direct and may survive moderate fine-tuning (e.g., Dasgupta et al. 2024), but their invasiveness is also their principal limitation: they require control over the training process and cannot retroactively mark third-party models.

Passive LLM Fingerprinting. Passive fingerprinting identifies models by analyzing properties that already exist in the model, without altering parameters or outputs. The available evidence depends on the access regime. Under **white-box access**, signatures can be extracted from weights or representations, including attention-matrix distributions (Yoon et al., 2025), kernel alignment of internal representations (Zhang et al., 2024), or stable directions in parameter space. Under **black-box access**, fingerprinting relies on input-output behavior, using crafted queries (Pasquini et al., 2024), distinctive prompt-response pairs (Tsai et al., 2025), stylometric regularities (Alhazbi et al., 2025), or iterative prompting–response games (Iourovitski et al., 2024); these behavioral signals can be less robust to fine-tuning. Across these variants, the shared limitation is that the signature is defined *post hoc*: it identifies properties that emerge after training has produced stable parameters, representations, or aligned behaviors. Such methods therefore do not target the innate, “Galtonian” fingerprints present from random initialization that this work isolates. Consequently, *they are not designed for lineage verification during large-scale pretraining*, the precise regime in which post-training signatures may not yet exist.

3 BIASES ORIGINATING FROM INITIALIZATION PERSIST AFTER TRAINING

This section establishes the empirical foundation for SeedPrints: random initialization induces pronounced seed-specific prediction biases, and subsequent training preserves enough of this structure to support statistically reliable lineage testing.

Initialization induces seed-specific bias profiles. We evaluate a LLaMA-2–style model initialized with seed 123 on 10,000 random input sequences of length 1,024, with tokens sampled uniformly from the vocabulary. In Figure 2 (Left), we measure how often each output token (upper panel, final logits) and each hidden dimension (lower panel, last-layer representations) attains the minimum value across the 10,000 trials. The resulting profiles are sharply non-uniform, demonstrating that extreme output-preference biases exist before any training occurs. Repeating the protocol across initialization seeds preserves the overall magnitude of the bias while changing the specific argmin dimensions (Figure 2, Lower Right), which identifies the bias profile as seed-dependent rather than merely architectural. This pattern is consistent with recent evidence that biases in randomly initialized models arise from inter-sequence representation contraction, induced by asymmetric nonlinear activations in MLP blocks and amplified by self-attention (Li et al., 2026). Under that mechanism, different initialization seeds determine different contraction directions and consequently produce distinct argmin patterns across output dimensions.

Training preserves initialization-born bias Training substantially changes output magnitudes, but it does not erase the relative preference structure induced at initialization. We train the seed-123 model for one epoch on OpenWebText (Gokaslan et al., 2019) and evaluate intermediate checkpoints. The analysis restricts attention to the $m = 50$ output dimensions most frequently assigned minimum values by the initialized model. For each checkpoint, we compute the correlation between its random-




Figure 2: **Initialization-born output bias persists through training.** *Left:* Under completely random inputs, a randomly initialized LLaMA-2-style model produces highly non-uniform outputs: specific dimensions are systematically disfavored and therefore frequently receive the minimum value. This extreme bias appears in both logits (top, red) and final hidden representations (bottom, blue). The dashed line gives the expected frequency under a uniform distribution, and the arrows in the top panel mark a broken x-axis that omits low-frequency tail ranks. *Upper Right:* Training checkpoints retain weak input-wise correlation in output bias with the initialized model from which they descend. The correlation distribution is consistently shifted upward relative to correlations with independently initialized models, showing that initialization-born bias remains measurable after optimization. *Lower Right:* The most strongly biased dimensions are seed-dependent; the panel illustrates this effect using tokens that most frequently receive the minimum logit.

input responses and the corresponding responses of the initialized model on these dimensions; this statistic tests whether the checkpoint and its initialization rank or disfavor the same dimensions in a consistent input-wise order. As an independent baseline, we compute the same correlation against models initialized with seed 1000. As shown in Figure 2 (Upper Right), the absolute correlations are small, but same-initialization correlations are consistently shifted upward relative to the independent baseline and stabilize rather than decaying to zero. Initialization therefore supplies a weak but statistically detectable bias signal that survives optimization.

4 ALGORITHM

Section 3 shows that models sharing an initialization lineage retain weak but systematic agreement on a subset of output dimensions, especially dimensions disfavored by the untrained model under random inputs. SeedPrints operationalizes this observation by recovering, from any two models and without access to the original initialization, the dimensions that carry the shared bias signal and by testing whether their agreement exceeds the level expected between independent models. Algorithm 1 summarizes the complete procedure; the following paragraphs specify its components.

Extract identity dimensions between any two models. Let $X = \{x_i\}_{i=1}^n$, where each $x_i \in \mathbb{R}^{\ell \times d}$ denotes a random sequence of length ℓ . Each random input x_i is instantiated either by uniformly sampling ℓ token IDs from the vocabulary or by directly sampling ℓ independent vectors from a d -dimensional isotropic Gaussian distribution. For any model g , we define the mean response vector as $\bar{g} := \frac{1}{n} \sum_{i=1}^n g(x_i) \in \mathbb{R}^{d_{\text{out}}}$, where d_{out} denotes the output dimensionality, and $g(x_i)$ may be instantiated as either logits or final-layer representations. We then identify its disfavored dimensions as the set of bottom- m coordinates of $\bar{g}$:

$$\mathcal{M}_g := \arg \min_{J \subseteq \{1, \dots, d_{\text{out}}\}, |J|=m} \sum_{j \in J} \bar{g}_j.$$

The mean response vector is a stable, noise-robust estimator of the bias profile. It improves on per-sample argmin frequency (i.e., counting how often each dimension attains the minimum value, as illustrated in Figure 2) while preserving strong empirical performance; Appendix C.1 provides the formal comparison.

Given two models f and f' , we define their *identity dimensions* as the intersection $\mathcal{S} := \mathcal{M}_f \cap \mathcal{M}_{f'}$. If two models share initialization lineage, they should independently recover overlapping sets of

disfavored dimensions, yielding non-trivial intersection between their bottom- m sets. The intersection criterion therefore implements mutual verification: it suppresses correlations attributable only to architectural similarity or shared training dynamics and concentrates the test on bias signals plausibly inherited from a common initialization. When one model is close to initialization, its disfavored dimensions can also serve directly as an anchor.

Correlation statistics. After restricting both models to the identity dimensions S , we compare their responses across random inputs.¹ For each $j \in S$, we compute a Kendall–Tau rank correlation:

$$\tau_j = \text{KendallTau}(\{f(x_i)_j\}_{i=1}^n, \{f'(x_i)_j\}_{i=1}^n).$$

This yields correlation statistics $T = \{\tau_j\}_{j \in S}$, whose distribution quantifies agreement between the two models across identity dimensions. As established in Section 3, same-lineage models exhibit systematically higher agreement on these dimensions than unrelated models.

Hypothesis testing via null distribution. Because same-lineage and different-lineage correlation distributions can overlap, thresholding individual correlations is not a calibrated decision rule. SeedPrints instead tests whether the observed correlation distribution is larger than expected under the null hypothesis of no lineage relation. The ideal null is the correlation distribution between independent models, as in Figure 2. Estimating it directly would require many independently initialized models and repeated evaluations, making it computationally expensive. Appendix B shows that the null correlation distribution between independently initialized models is well approximated by a Gaussian distribution. This observation motivates a tractable simulation-based alternative: pass Gaussian surrogate outputs through the same pipeline to construct an empirical null, as summarized in Algorithm 1.

Although the simulation-based null is empirically accurate, it introduces avoidable randomness and computational overhead: precision depends on the number of simulation trials, results can vary across runs, and empirical distribution comparison depends on the chosen test, such as a one-sided t-test or Mann–Whitney U test. Because the null distribution of Kendall–Tau under weak independence is available in closed form, we use an analytical test as the default while providing both empirical and analytical implementations in the code repository. Under the null hypothesis, each Kendall–Tau correlation is approximately zero-mean with known variance σ^2 . By the central limit theorem, as the number of identity dimensions $|S|$ increases, the distribution of $\bar{\tau} = \frac{1}{|S|} \sum_{j \in S} \tau_j$ is well approximated by $\bar{\tau} \sim \mathcal{N}\left(0, \frac{\sigma^2}{|S|}\right)$. Softmax normalization introduces weak dependencies across dimensions, but the relatively high temperature ($T = 10$) mitigates these effects, making the independence approximation accurate in practice. Appendix B empirically validates that the analytical Gaussian closely matches the distribution produced by the full pipeline. We therefore compute the one-sided z-score $z = \frac{\bar{\tau}}{\sigma/\sqrt{|S|}}$ and the corresponding p-value $p = 1 - \Phi(z)$, where Φ is the standard normal cumulative distribution function. We reject the null hypothesis and declare shared lineage if $p < 0.01$.

Practical instantiation. Model outputs may be instantiated as either logits or hidden representations. Both choices satisfy the formulation, but logits require substantially more random inputs to estimate bias patterns reliably as vocabulary size grows, which creates considerable computational overhead (see Appendix C.2 for theoretical analysis). For efficiency and stability, we therefore use final-layer hidden states as the default output representation. The resulting p-value provides a direct statistical lineage test rather than a heuristic similarity threshold.

5 EXPERIMENTS

The empirical evaluation follows directly from the two requirements of a Galtonian fingerprint. First, in Section 5.1, we test whether SeedPrints provides *seed-level verification at birth* and remains *verifiable throughout the training lifecycle*. These experiments also clarify the failure mode of existing baselines in early pretraining and under substantial distribution shift: their evidence is primarily induced by data and optimization rather than inherited from initialization. Unless otherwise

¹Before computing the correlations, we apply a row-wise softmax normalization to each sample’s output vector. This normalization removes scale differences across models and dimensions while preserving the relative ordering that carries identity information for rank-based statistics such as Kendall–Tau. A relatively high temperature is used to avoid degenerate one-hot behavior.

Algorithm 1 Distribution Correlation Test on Identity Dimensions

Require: base model f , suspicious model f' ; random inputs $X = \{x_i\}_{i=1}^n$; fingerprint size m ; significance level α

► **Step 1: Localize biased dimensions**

- 1: Compute average outputs $\bar{f}, \bar{f}'$ over X
- 2: $\mathcal{M}(f), \mathcal{M}(f') \leftarrow$ bottom- m dimensions of $\bar{f}, \bar{f}'$
- 3: $\mathcal{S} \leftarrow \mathcal{M}(f) \cap \mathcal{M}(f')$ (identity dimensions)

► **Step 2: Form correlation distribution**

- 4: Normalize $f(x_i)_{\mathcal{S}}, f'(x_i)_{\mathcal{S}}$ for each $x_i \in X$ across identity dimension
- 5: **for** each $s_j \in \mathcal{S}$ **do**
- 6: $\tau_j \leftarrow \text{KendallTau}(\{f(x_i)_{s_j}\}_{i=1}^n, \{f'(x_i)_{s_j}\}_{i=1}^n)$ (correlation in in input-wise preferences)
- 7: **end for**
- 8: $\mathcal{T} \leftarrow \{\tau_j\}_{j=1}^{|\mathcal{T}|}$

► **Step 3: Hypothesis test against null**

- 9: Construct $\mathcal{T}_{\text{null}}$ by applying the same pipeline to two Gaussian matrices $Y^{(1)}, Y^{(2)} \sim \mathcal{N}(0, I)^{N \times d_{\text{out}}}$
- 10: Test $H_0 : \mathcal{T} = \mathcal{T}_{\text{null}}$ vs. $H_1 : \mathcal{T} > \mathcal{T}_{\text{null}}$
- 11: **Return** SameLineage $\leftarrow \mathbf{1}(p\text{-value} < \alpha)$

Table 1: SeedPrints separates models initialized with different seeds. Table 2: Trained models retain the fingerprint of their initialization ($p\text{-value} < 0.01$).

Model Pairs	$p\text{-value} (> 0.01)$	Model Pairs	SeedPrints	Baselines			
			$p\text{-value}$	Intrinsic	REEF	PCS	ICS
s_{42} VS. s_{2000}	0.116	s_{42}^{init} vs. s_{42}^{base}	3.37e-26✓	-0.021×	0.375×	0.580×	0.196×
s_{123} VS. s_{42}	1.000	s_{123}^{init} vs. s_{123}^{base}	3.16e-33✓	0.149×	0.369×	0.581×	0.188×
s_{1000} VS. s_{123}	0.327	s_{1000}^{init} vs. s_{1000}^{base}	1.12e-21✓	-0.252×	0.381×	0.581×	0.188×
s_{2000} VS. s_{1000}	1.000	s_{2000}^{init} vs. s_{2000}^{base}	3.70e-31✓	-0.337×	0.331×	0.581×	0.188×

stated, the main results use the analytical null with *hidden state* outputs; Appendix D.2 reports complementary empirical-null results for both *logit* and *hidden state* outputs, evaluated with the one-sided t -test and Mann–Whitney U test. Second, in Section 5.3, we evaluate practical infringement scenarios using *LeaFBench* (Shao et al., 2025). Across 65 model instances and 6 representative post-development techniques, SeedPrints remains competitive with the strongest post-training baselines while preserving its distinctive ability to verify lineage before, during, and after training.

Baselines We compare primarily against four passive fingerprinting baselines based on weights or representations. *Intrinsic fingerprint* (Yoon et al., 2025) (also called *PDF* in some work) compares models through the similarity of layerwise standard-deviation profiles of attention parameters. *REEF* (Zhang et al., 2024) computes centered-kernel-alignment (CKA) similarity between feature representations induced by the same samples in two models. *PCS* and *ICS* (Zeng et al., 2024) (collectively referred to as *HuRef* in some work) are weight-similarity methods: PCS flattens all parameters and measures cosine similarity, whereas ICS constructs invariant terms from the weights and measures cosine similarity over those invariants. In Section 5.3, we also include the gradient-based fingerprint *Gradient* (Shao et al., 2025). Following Zhang et al. (2024), similarity-based baselines use a **0.8 similarity threshold** for binary decisions.

Across all experiment tables, cell colors encode the ground-truth lineage relation: green denotes models from the same source, and red denotes different sources. For example, s_{42}^{init} vs. s_{42}^{base} compares a seed-42 initialization with its continued-pretrained descendant, whereas s_{2000}^{init} vs. s_{42}^{base} compares a seed-2000 initialization with a model trained from seed 42. The symbol ✓ marks a correct decision, and × marks an error.

5.1 BIRTH-TO-LIFECYCLE “BIOMETRIC” FINGERPRINTING

We train 12-layer, 12-head LLaMA-style models (Touvron et al., 2023) with RoPE (Su et al., 2021) and Qwen-style models (Team, 2024) from scratch. The main text reports LLaMA-style results, and Appendix D provides the corresponding Qwen-style results; the same conclusions hold across both architectures.

Different initialization seeds produce distinct fingerprints Table 1 reports p -values from correlation tests between pairs of models initialized with different random seeds (42, 123, 1000, and 2000).

Table 3: Identical dataset and training order do not make fingerprints converge across different initializations.

Model Pair	p -value (> 0.01)
s_{42}^{init} vs. s_{123}^{base}	0.573
s_{123}^{init} vs. s_{2000}^{base}	0.724
s_{2000}^{init} vs. s_{1000}^{base}	0.109
s_{1000}^{init} vs. s_{42}^{base}	0.883

Table 4: Fingerprint persistence under continual training on diverse datasets (base model: seed 1000, corpus openwebtext).

Setting	SeedPrints	Baselines			
Continual corpus (seed)	p -value	Intrinsic	REEF	PCS	ICS
TinyStories (1000)	≈ 0 ✓	1.000✓	0.759×	0.999✓	0.996✓
TinyStories (123)	1.000✓	0.950×	0.658✓	0.332✓	0.012✓
the_stack (1000)	≈ 0 ✓	0.489×	0.557×	0.585×	0.123×
the_stack (123)	1.000✓	0.445✓	0.580✓	0.301✓	0.026✓

Every reported p -value remains above 0.01, so SeedPrints does not falsely declare shared lineage across seeds. The result establishes seed-level separability before training: distinct initialization seeds induce distinct fingerprint behaviors, allowing models to be distinguished “at birth.”

Training preserves the initialization fingerprint. Table 2 compares each initialization model s^{init} with its descendant s^{base} after training on OpenWebText (Gokaslan et al., 2019) (≈ 10 B tokens). For every seed–model pair, SeedPrints returns $p < 0.01$, demonstrating that the trained model preserves the bias profile of its initialization and therefore shares its lineage. The baseline methods fail to distinguish across seeds in this setting, which indicates that their separability is driven by training-induced artifacts rather than by initialization identity.

Identical data and order do not make fingerprints converge In Table 3, all four “suspicious” models s_i^{base} for $i \in \{42, 123, 100, 2000\}$ are trained on *exactly the same corpus (OpenWebText) and in the same data order*; the training seed fixes the data order, so the only experimental difference is the initialization seed i . Across all cross-seed pairs, p -values remain above 0.01, in sharp contrast to the near-zero same-lineage values in Table 2. Thus, identical data and curriculum do not collapse SeedPrints into a shared training-data signature.

Continual training on diverse datasets does not confound the fingerprint The preceding experiments establish SeedPrints as a biometric-style fingerprint, but seed distinguishability alone is not the central copyright challenge because initialization seeds themselves have no clear ownership status. The more consequential failure mode for prior methods is distributional confounding: attribution can track *data distribution* rather than lineage, thereby *failing to recognize descendants when continued training substantially changes the training distribution*.

In Table 4, we continue training a seed-1000 base model, pretrained on OpenWebText (Gokaslan et al., 2019), on two sharply different corpora: TinyStories (Eldan & Li, 2023), consisting of synthetic children’s stories, and The Stack (Kocetkov et al., 2022), consisting of permissively licensed GitHub code. We compare true descendants of the seed-1000 base against *distractors* derived from a different seed-123 base. Because the baseline methods are not sensitive to initialization differences, the distractor base is pretrained on OpenWebText with a different data order and then continually trained on the *same* target corpus as the true descendant. The resulting test isolates lineage from domain overlap: a successful method must identify the model that truly descends from the base rather than the model that merely shares the continued-training corpus.

The baselines fail under the code-domain setting (The Stack), misclassifying true descendants as distractors. This pattern indicates that their scores largely reflect domain similarity rather than lineage identity: TinyStories remains closer to OpenWebText, whereas The Stack creates a much larger distribution shift that can bypass detection. SeedPrints correctly attributes lineage across both corpora, showing that its fingerprint is not a proxy for data distribution and that the initialization signal persists beyond initial pretraining.

5.2 ALL-STAGE VERIFIABLE FINGERPRINTS

Controlled pretraining trajectory. We first test all-stage verification in a controlled OpenWebText pretraining run for which the initialization seed is known. Lineage tests at intermediate checkpoints consistently detect the correct descendant across the trajectory (Appendix D.2.1, Figure 5), confirming that the initialization signal remains accessible throughout continuous optimization.

Large-scale pretraining. To evaluate lineage under realistic foundation-model training, we analyze ten checkpoints from the OLMo-2-7B (OLMo et al., 2024) Stage-1 run (5B $\rightarrow$ 3.9T tokens), treating

Table 5: Fingerprinting results under large-scale finetuning. Each row compares a target model against LLaMA-2-7B. **SeedPrints** reports the p -value from our correlation test (< 0.01 indicates a strong signal). Four baselines all report similarity scores (threshold = 0.8, higher = better).

Model	# Tokens	SeedPrints p	Intrinsic $\uparrow$	REEF ($\uparrow$)	PCS ($\uparrow$)	ICS ($\uparrow$)
Llama-2-finance-7B (Heenan, 2023)	5M	10^{-41955} ✓	1.0000✓	0.9950✓	0.9979✓	0.9952✓
Vicuna-1.5-7B (Chiang et al., 2023)	370M	$10^{-103043}$ ✓	1.0000✓	0.9985✓	0.9985✓	0.9949✓
Wizardmath-7B (Luo et al., 2023)	1.8B	$10^{-180550}$ ✓	1.0000✓	0.9979✓	1.0000✓	0.9994✓
Meditron-7B (Chen et al., 2023)	48B	10^{-42566} ✓	0.9990✓	0.9978✓	1.0000✓	0.9817✓
CodeLlama-7B (Meta AI, 2024)	500B	10^{-3552} ✓	0.9480✓	0.9947✓	0.6863✗	0.3369✗
Llemma-7B (Azerbayev et al., 2023)	700B	10^{-5136} ✓	0.9470✗	0.9984✓	0.6682✗	0.2905✗

the final checkpoint as the target model and each earlier checkpoint as a tested ancestor. Unlike the controlled OpenWebText setting, this trajectory spans heterogeneous corpora, substantially longer optimization, and large-scale representation drift. For comparability with similarity-based baselines, we plot $1 - p$; all SeedPrints p -values are $\ll 0.01$ and numerically close to zero.

Figure 1 shows that **all existing baselines fail to verify lineage during early pretraining**: within the first trillion tokens, ICS is approximately 0, PCS remains below 0.3, and REEF remains below 0.6, all below the 0.8 decision threshold. SeedPrints, by contrast, verifies lineage from the earliest checkpoint, and its signal strengthens during training. This comparison exposes a critical evaluation gap: **most prior work verifies lineage only after fine-tuning, where detection is substantially easier**. Large-scale pretraining can reshape representations and amplify later-stage signatures, creating a *false sense of safety* when early checkpoints are not evaluated. Lineage verification should therefore prioritize early pretraining, when misuse is hardest to detect, rather than relying only on late-stage checkpoints.

Large-scale finetuning. We also compare SeedPrints with existing baselines in standard fine-tuning evaluations. The suspect models are fine-tuned from Llama-2-7b with data volumes ranging from 5 million to 700 billion tokens. They include Llama-2-finance-7b (Heenan, 2023), Vicuna-1.5-7b (Chiang et al., 2023), WizardMath-7b (Luo et al., 2023), Chinese-LLaMA-2-7b (Chen et al., 2023), Code-Llama-7b (Meta AI, 2024), and Llemma-7b (Azerbayev et al., 2023), with fine-tuning data volumes of 5M, 370M, 1.8B, 13B, 500B, and 700B tokens, respectively. As shown in Table 5, SeedPrints consistently returns $p < 0.01$ across all settings.

5.3 ROBUSTNESS UNDER REALISTIC DEPLOYMENTS

Real-world deployments routinely apply *parameter-altering* adaptations to foundation models, including fine-tuning, PEFT, and quantization, and these transformations can weaken or distort fingerprint signals. To evaluate SeedPrints under such operational conditions, we use *LeafBench* (Shao et al., 2025), a benchmark for language-model copyright auditing.

Across Source Models. We audit two source-model types: *Pre-Trained (PT)* models and *Instruction-Tuned (IT)* models. For each source type, the task is to distinguish *derivative* models produced by post-training adaptations from *independent* models trained without using the source weights. This setup matches common auditing scenarios in which lineage must be tested for either a PT base model or its IT variant. The evaluation covers 65 models in total; Appendix E lists all model names and HuggingFace repositories.

Metrics. We evaluate detection performance with (i) the Area Under the ROC Curve (AUC) and (ii) the Kolmogorov–Smirnov (KS) statistic (Berger & Zhou, 2014), which measures the maximum separation between the score distributions of derivative and independent models.

SeedPrints produces per-pair p -values from a statistical hypothesis test, whereas the baselines produce similarity scores. These quantities have different semantics: similarity scores are treated as linear thresholding scores, while p -values are tail probabilities that quantify the rarity of the observed statistic under the null hypothesis and should not be interpreted linearly. To make AUC computation compatible with baseline pipelines, we convert SeedPrints outputs to scores via $s = 1 - p$. This conversion is imperfect for the statistical meaning of p -values, because very small p -values vary on an exponential scale whereas threshold sweeping is linear; nevertheless, SeedPrints remains comparable to or better than the strongest baselines. For evaluation, we compute AUC for all methods by sweeping thresholds over scores s and report the KS statistic as a threshold-free measure of

Table 6: Performance comparison of LLM fingerprinting methods across source-model families. “PT” and “IT” use the pre-trained and instruction-tuned models as source models, respectively. The per-family breakdown covers the 6 families that have both PT and IT models. The “Overall” scores are computed over all test model pairs, not averaged across families.

Method	Metric	Qwen2.5-7B		Qwen2.5-14B		Llama3.1-8B		Mistral-7B		Gemma2-2B		Llama2-7B		Overall
		PT	IT	PT	IT	PT	IT	PT	IT	PT	IT	PT	IT	
REEF	AUC $\uparrow$	0.796	0.862	0.947	0.913	1.000	0.999	0.997	0.997	0.963	0.968	0.750	0.706	0.915
	KS Statistic $\uparrow$	0.492	0.633	0.875	0.735	0.987	0.987	0.975	0.980	0.950	0.837	0.750	0.694	0.739
Gradient	AUC $\uparrow$	0.657	0.743	0.763	0.789	0.897	0.904	0.872	0.682	1.000	1.000	0.650	0.768	0.801
	KS Statistic $\uparrow$	0.354	0.423	0.500	0.563	0.646	0.705	0.725	0.500	1.000	1.000	0.275	0.479	0.508
ICS	AUC $\uparrow$	1.000	1.000	0.997	0.997	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	0.994
	KS Statistic $\uparrow$	1.000	1.000	0.975	0.979	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	0.943
Intrinsic	AUC $\uparrow$	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	0.994
	KS Statistic $\uparrow$	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	0.989
SeedPrints	AUC $\uparrow$	0.992	0.994	0.988	0.990	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	0.992
	KS Statistic $\uparrow$	0.985	0.987	0.975	0.980	1.000	1.000	1.000	1.000	1.000	1.000	1.000	1.000	0.986

distributional separability. We do not report Manhattan distance because absolute distances between p -values are not meaningful.

In most practical auditing scenarios, calibrating an optimal similarity threshold is infeasible. Although we report AUC for comparability, SeedPrints differs from prior methods because it supplies a direct statistical test: **given any pair of models, it outputs a p -value that supports a lineage decision without requiring a tunable similarity threshold**. This property makes the method operationally clearer than score-only fingerprinting approaches.

Parameter-Altering Techniques. To assess robustness, we evaluate models produced by (1) Instruction tuning (Instruct), (2) General-purpose fine-tuning (Finetune), (3) Parameter-efficient fine-tuning (PEFT), (4) Quantization, (5) Model merging (Merge), and (6) Distillation.

5.3.1 FAMILY-WISE EFFECTIVENESS

The benchmark contains seven model families. For each family, we treat both the pretrained base model (PT) and the instruction-tuned model (IT) as separate sources and compare each source against the remaining models in the 65-model pool. Table 6 reports the results; TinyLlama-1.1B families are omitted from Table 6 because they are trivial for all methods. Each family column computes AUC over the corresponding tests using a family-calibrated threshold, whereas the overall score is computed across all test pairs with a global threshold; consequently, the overall score is not the average of the per-family scores. *Intrinsic*, *ICS*, and *SeedPrints* are all near saturation across alterations, with overall AUCs of 0.994, 0.994, and 0.992 and KS values of 0.989, 0.943, and 0.986, respectively. Their family-wise scores are full scores in most cases, whereas *REEF* and *Gradient* perform substantially worse. These results show that, although SeedPrints is designed to capture intrinsic seed-level fingerprints, it also functions as a persistent biometric-style identity tracker across diverse model families.

5.3.2 ALTERATION-WISE ROBUSTNESS

Across alteration types, *Instruct* is the easiest setting: even the weakest baseline, *Gradient*, reaches AUC 0.895. *Finetune* and *Merge* are more challenging for all methods; for *Finetune*, no method achieves perfect performance. *PEFT*, *Quantization*, and *Distillation* show more method-specific effects. *REEF* drops under *Quantization* (0.996 $\rightarrow$ 0.871) and *Distillation* (0.996 $\rightarrow$ 0.858) while remaining resilient to *PEFT* (0.994); *Gradient* is more affected by *PEFT* and *Quantization* (0.895 $\rightarrow$ 0.776/0.756). The three stronger methods remain robust across transformations, with all scores above 0.9. SeedPrints decreases to 0.959 on *Merge*, corresponding to one failed merge model out of seventeen, suggesting that weight-space interpolation with unrelated model weights can dilute the initialization signal below the detection threshold. Averaged across all six alteration types, *Intrinsic* achieves the highest overall AUC, followed closely by *ICS* and *SeedPrints*.

Table 7: Performance of LLM fingerprinting methods across parameter-altering techniques.

Method	Metric	Instruct	Finetune	PEFT	Quantization	Merge	Distillation	Overall Avg
REEF	AUC $\uparrow$	0.996	0.872	0.994	0.871	0.954	0.858	0.924
	KS Statistic $\uparrow$	0.949	0.637	0.940	0.812	0.851	0.615	0.801
Gradient	AUC $\uparrow$	0.895	0.729	0.776	0.756	0.832	0.887	0.812
	KS Statistic $\uparrow$	0.756	0.441	0.469	0.586	0.519	0.679	0.575
ICS	AUC $\uparrow$	1.000	0.987	1.000	1.000	0.981	1.000	0.995
	KS Statistic $\uparrow$	1.000	0.925	1.000	1.000	0.941	1.000	0.978
Intrinsic	AUC $\uparrow$	1.000	0.980	1.000	1.000	1.000	1.000	0.997
	KS Statistic $\uparrow$	1.000	0.971	1.000	1.000	1.000	1.000	0.995
SeedPrints	AUC $\uparrow$	1.000	0.995	0.990	1.000	0.959	1.000	0.991
	KS Statistic $\uparrow$	1.000	0.990	0.980	1.000	0.941	1.000	0.985

5.4 ABLATION STUDY ON HYPERPARAMETERS

SeedPrints has three hyperparameters: the number of random input sequences n , the length of each random sequence l , and the fingerprint size m . The first two control the amount of random-input evidence, while m determines how many low-mean coordinates enter the lineage test. In the main experiments in Sections 5.2 and 5.3, we set $n=2000$, $l=1024$, and $m=0.1d_{\text{out}}$, where d_{out} is the hidden-state dimensionality, so the configuration scales with model size. Table 8 varies one hyperparameter at a time while fixing the others to their defaults ($n = 2000$, $m = 400$, $l = 768$), and reports empirical accuracy and false positive rate at significance level $\alpha = 0.01$.

Table 8: Ablation study under significance level $\alpha = 0.01$.

Increasing n and l consistently improves identification performance. This trend agrees with the theoretical analysis in Appendix C.2, which shows that recovering the population-level bias requires $n = \Omega\left(\frac{\log d_{\text{out}}}{\gamma^2}\right)$ queries, where γ is the margin separating biased coordinates from the remaining dimensions. Smaller n therefore provides insufficient evidence for reliable coordinate recovery. By contrast, γ increases with l through accumulation effects, so longer sequences strengthen separation and reduce the query requirement when the number of sequences is fixed.

Ablation on n	200	500	2000
ACC	0.7368	0.7778	0.9375
Empirical FPR	0	0	0
Ablation on l	512	768	1024
ACC	0.9048	0.9375	1.0000
Empirical FPR	0.04	0	0
Ablation on m	200	400	800
ACC	0.9375	0.9375	0.875
Empirical FPR	0.04	0	0.1463

The effect of m is non-monotonic. If m is too small, the selected coordinate set is unstable because too few observations define the empirical distribution. If m is too large, many weakly relevant or irrelevant coordinates enter the test, adding noise that reduces γ and would require quadratically more samples to compensate. Empirically, m in the range 200–400 performs well across model sizes, consistent with the heuristic $m \approx 0.1d_{\text{out}}$ for typical $d_{\text{out}} \in [3000, 5000]$.

6 CONCLUSION

In this work, we introduced *SeedPrints*, an intrinsic formulation of LLM fingerprinting that traces model lineage back to random initialization. We showed that initialization itself leaves a persistent “Galtonian” fingerprint on neural language models: untrained models already exhibit stable, seed-dependent token-preference patterns, and these biases remain statistically measurable throughout training. Building on this observation, we proposed a distribution-correlation test over *identity dimensions* that detects shared lineage with calibrated significance. Across LLaMA- and Qwen-style families and a broad suite of realistic deployments (e.g., continued pretraining on divergent corpora, instruction tuning, PEFT, quantization, merging, and distillation), SeedPrints enables “birth-to-lifecycle” verification and remains robust under domain shift, providing a practical tool for provenance and copyright auditing.

Limitations and Future Directions This work isolates fingerprint signals rooted in initialization-stage bias. The proposed method requires access to model outputs (e.g., logits or hidden states) and may not apply when only restricted API access is available. Nevertheless, the pursuit of *true*

540 *fingerprints* for LLMs remains at an early stage: most existing approaches, especially black-box
541 methods, rely on signals that emerge in later training stages or from alignment-driven artifacts and
542 therefore cannot robustly protect pretrained or early-stage models. Reliable all-stage fingerprints are
543 therefore an urgent direction for future work. See full discussions in Appendix [A.2](#).
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593

REFERENCES

- Yossi Adi, Carsten Baum, Moustapha Cisse, Benny Pinkas, and Joseph Keshet. Turning your weakness into a strength: Watermarking deep neural networks by backdooring. In *27th USENIX security symposium (USENIX Security 18)*, pp. 1615–1631, 2018.
- Saeif Alhazbi, Ahmed Hussain, Gabriele Oligeri, and Panos Papadimitratos. Llms have rhythm: Fingerprinting large language models using inter-token times and network traffic analysis. *IEEE Open Journal of the Communications Society*, 2025.
- Zhangir Azerbayev, Hailey Schoelkopf, Keiran Paster, Marco Dos Santos, Stephen McAleer, Albert Q Jiang, Jia Deng, Stella Biderman, and Sean Welleck. Llemma: An open language model for mathematics. *arXiv preprint arXiv:2310.10631*, 2023.
- Vance W Berger and YanYan Zhou. Kolmogorov–smirnov test: Overview. *Wiley statsref: Statistics reference online*, 2014.
- Nicholas Carlini, Steve Chien, Milad Nasr, Shuang Song, Andreas Terzis, and Florian Tramer. Membership inference attacks from first principles. In *2022 IEEE symposium on security and privacy (SP)*, pp. 1897–1914. IEEE, 2022.
- Zeming Chen, Alejandro Hernández-Cano, Angelika Romanou, Antoine Bonnet, Kyle Matoba, Francesco Salvi, Matteo Pagliardini, Simin Fan, Andreas Köpf, Amirkeivan Mohtashami, Alexandre Sallinen, Alireza Sakhaeirad, Vinitra Swamy, Igor Krawczuk, Deniz Bayazit, Axel Marmet, Syrielle Montariol, Mary-Anne Hartley, Martin Jaggi, and Antoine Bosselut. Meditron-70b: Scaling medical pretraining for large language models, 2023.
- Wei-Lin Chiang, Zhuohan Li, Zi Lin, Ying Sheng, Zhanghao Wu, Hao Zhang, Lianmin Zheng, Siyuan Zhuang, Yonghao Zhuang, Joseph E. Gonzalez, Ion Stoica, and Eric P. Xing. Vicuna: An open-source chatbot impressing gpt-4 with 90%* chatgpt quality, March 2023. URL <https://lmsys.org/blog/2023-03-30-vicuna/>.
- Agnibh Dasgupta, Abdullah Tanvir, and Xin Zhong. Watermarking language models through language models. *arXiv preprint arXiv:2411.05091*, 2024.
- Ronen Eldan and Yuanzhi Li. Tinstories: How small can language models be and still speak coherent english? *arXiv preprint arXiv:2305.07759*, 2023.
- Francis Galton. *Finger prints*. Number 57490-57492. Cosimo Classics, 1892.
- Aaron Gokaslan, Vanya Cohen, Ellie Pavlick, and Stefanie Tellex. Openwebtext corpus. <http://Skyllion007.github.io/OpenWebTextCorpus>, 2019.
- Sylvain Gugger, Lysandre Debut, Thomas Wolf, Philipp Schmid, Zachary Mueller, Sourab Mangrulkar, Marc Sun, and Benjamin Bossan. Accelerate: Training and inference at scale made simple, efficient and adaptable. <https://github.com/huggingface/accelerate>, 2022.
- Jia Guo and Miodrag Potkonjak. Watermarking deep neural networks for embedded systems. In *2018 IEEE/ACM international conference on computer-aided design (ICCAD)*, pp. 1–8. IEEE, 2018.
- Collin Heenan. Llama2-7b-finance (hugging face model). <https://huggingface.co/cxllin/Llama2-7b-Finance>, 2023. Fine-tuned from NousResearch/Llama-2-7b-hf; MIT License; accessed 2025-09-02.
- Dmitri Iourovitski, Sanat Sharma, and Rakshak Talwar. Hide and seek: Fingerprinting large language models with evolutionary learning. *arXiv preprint arXiv:2408.02871*, 2024.
- John Kirchenbauer, Jonas Geiping, Yuxin Wen, Jonathan Katz, Ian Miers, and Tom Goldstein. A watermark for large language models. In *International Conference on Machine Learning*, pp. 17061–17084. PMLR, 2023.
- Denis Kocetkov, Raymond Li, Loubna Ben Allal, Jia Li, Chenghao Mou, Carlos Muñoz Ferrandis, Yacine Jernite, Margaret Mitchell, Sean Hughes, Thomas Wolf, Dzmitry Bahdanau, Leandro von Werra, and Harm de Vries. The stack: 3 tb of permissively licensed source code. *Transactions on Machine Learning Research (TMLR), Preprint*, 2022.

- Jaehoon Lee, Lechao Xiao, Samuel Schoenholz, Yasaman Bahri, Roman Novak, Jascha Sohl-Dickstein, and Jeffrey Pennington. Wide neural networks of any depth evolve as linear models under gradient descent. *Advances in neural information processing systems*, 32, 2019.
- Peixuan Li, Pengzhou Cheng, Fangqi Li, Wei Du, Haodong Zhao, and Gongshen Liu. Plmmark: a secure and robust black-box watermarking framework for pre-trained language models. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37, pp. 14991–14999, 2023.
- Siquan Li, Yao Tong, Haonan Wang, and Tianyang Hu. Transformers are born biased: Structural inductive biases at random initialization and their practical consequences. *arXiv preprint arXiv:2602.05927*, 2026.
- Zheng Li, Chengyu Hu, Yang Zhang, and Shanqing Guo. How to prove your model belongs to you: A blind-watermark based framework to protect intellectual property of dnn. In *Proceedings of the 35th annual computer security applications conference*, pp. 126–137, 2019.
- Xiaokun Luan, Zeming Wei, Yihao Zhang, and Meng Sun. Robust and efficient watermarking of large language models using error correction codes. *Proceedings on Privacy Enhancing Technologies*, 2025.
- Haipeng Luo, Qingfeng Sun, Can Xu, Pu Zhao, Jianguang Lou, Chongyang Tao, Xiubo Geng, Qingwei Lin, Shifeng Chen, and Dongmei Zhang. Wizardmath: Empowering mathematical reasoning for large language models via reinforced evol-instruct. *arXiv preprint arXiv:2308.09583*, 2023.
- Meta AI. Codellama-7b-hf: Code llama base 7b model (hugging face). <https://huggingface.co/codellama/CodeLlama-7b-hf>, 2024. Base 7 billion-parameter Code Llama model for code synthesis and understanding; trained between January and July 2023; licensed under Meta Llama 2 license; accessed 2025-09-02.
- Anshul Nasery, Jonathan Hayase, Creston Brooks, Peiyao Sheng, Himanshu Tyagi, Pramod Viswanath, and Sewoong Oh. Scalable fingerprinting of large language models. In *ICLR 2025 Workshop on Building Trust in Language Models and Applications*.
- Team OLMo, Pete Walsh, Luca Soldaini, Dirk Groeneveld, Kyle Lo, Shane Arora, Akshita Bhagia, Yuling Gu, Shengyi Huang, Matt Jordan, Nathan Lambert, Dustin Schwenk, Oyvind Tafjord, Taira Anderson, David Atkinson, Faeze Brahman, Christopher Clark, Pradeep Dasigi, Nouha Dziri, Michal Guerquin, Hamish Ivison, Pang Wei Koh, Jiacheng Liu, Saumya Malik, William Merrill, Lester James V. Miranda, Jacob Morrison, Tyler Murray, Crystal Nam, Valentina Pyatkin, Aman Rangapur, Michael Schmitz, Sam Skjonsberg, David Wadden, Christopher Wilhelm, Michael Wilson, Luke Zettlemoyer, Ali Farhadi, Noah A. Smith, and Hannaneh Hajishirzi. 2 olmo 2 furious, 2024. URL <https://arxiv.org/abs/2501.00656>.
- Dario Pasquini, Evgenios M Kornaropoulos, and Giuseppe Ateniese. Llmmap: Fingerprinting for large language models. *arXiv preprint arXiv:2407.15847*, 2024.
- Tom Sander, Pierre Fernandez, Alain Durmus, Matthijs Douze, and Teddy Furon. Watermarking makes language models radioactive. *Advances in Neural Information Processing Systems*, 37: 21079–21113, 2024.
- Shuo Shao, Yiming Li, Yu He, Hongwei Yao, Wenyan Yang, Dacheng Tao, and Zhan Qin. Sok: Large language model copyright auditing via fingerprinting. *arXiv preprint arXiv:2508.19843*, 2025.
- Jianlin Su, Yu Lu, Shengfeng Pan, Bo Wen, and Yunfeng Liu. Roformer: Enhanced transformer with rotary position embedding. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL)*, pp. 115–124. Association for Computational Linguistics, 2021.
- Teppe Suzuki, Ryokan Ri, and Sho Takase. Natural fingerprints of large language models. *arXiv preprint arXiv:2504.14871*, 2025.
- Qwen Team. Qwen2 technical report. *arXiv preprint arXiv:2407.10671*, 2, 2024.

-
- Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- Yun-Yun Tsai, Chuan Guo, Junfeng Yang, and Laurens van der Maaten. Rofl: Robust fingerprinting of language models. *arXiv preprint arXiv:2505.12682*, 2025.
- Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, Joe Davison, Sam Shleifer, Patrick von Platen, Clara Ma, Yacine Jernite, Julien Plu, Canwen Xu, Teven Le Scao, Sylvain Gugger, Mariama Drame, Quentin Lhoest, and Alexander M. Rush. Transformers: State-of-the-art natural language processing. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pp. 38–45, Online, October 2020. Association for Computational Linguistics. URL <https://www.aclweb.org/anthology/2020.emnlp-demos.6>.
- Jiashu Xu, Fei Wang, Mingyu Derek Ma, Pang Wei Koh, Chaowei Xiao, and Muhao Chen. Instructional fingerprinting of large language models. *arXiv preprint arXiv:2401.12255*, 2024.
- Do-hyeon Yoon, Minsoo Chun, Thomas Allen, Hans Müller, Min Wang, and Rajesh Sharma. Intrinsic fingerprint of llms: Continue training is not all you need to steal a model! *arXiv preprint arXiv:2507.03014*, 2025.
- Boyi Zeng, Lizheng Wang, Yuncong Hu, Yi Xu, Chenghu Zhou, Xinbing Wang, Yu Yu, and Zhouhan Lin. Huref: Human-readable fingerprint for large language models. *Advances in Neural Information Processing Systems*, 37:126332–126362, 2024.
- Jialong Zhang, Zhongshu Gu, Jiyong Jang, Hui Wu, Marc Ph Stoecklin, Heqing Huang, and Ian Molloy. Protecting intellectual property of deep neural networks with watermarking. In *Proceedings of the 2018 on Asia conference on computer and communications security*, pp. 159–172, 2018.
- Jie Zhang, Dongrui Liu, Chen Qian, Linfeng Zhang, Yong Liu, Yu Qiao, and Jing Shao. Reef: Representation encoding fingerprints for large language models. *arXiv preprint arXiv:2410.14273*, 2024.
- Ruichong Zhang. Matrix-driven instant review: Confident detection and reconstruction of llm plagiarism on pc. *arXiv preprint arXiv:2508.06309*, 2025.
- Renjie Zhu, Xinpeng Zhang, Mengte Shi, and Zhenjun Tang. Secure neural network watermarking protocol against forging attack. *EURASIP Journal on Image and Video Processing*, 2020(1):37, 2020.

756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809

Appendix Contents

A	LLM Usage and Limitation Discussions	16
A.1	LLM Usage	16
A.2	Limitations and Future Directions	16
B	Validation of the Empirical and Analytical Null Distribution	16
B.1	Empirical True Null from Independently Initialized Models	16
B.2	Gaussian Surrogate as a Simulation-Based Approximation	16
B.3	Analytical Approximation under Weak Dependence	17
C	Theoretical Properties of Our Methods	18
C.1	Stability Analysis of Mean-Based Identity Extraction vs. Per-Sample Argmin Voting	19
C.1.1	Stability of Sample-Mean-Based Extraction	19
C.1.2	Instability of Per-Sample Argmin Voting	20
C.2	Influence of Hyperparameters	22
D	More Experiment Details	23
D.1	Implementation Details and Licenses	23
D.2	Complementary SeedPrints Results using simulation-based null	23
D.2.1	Birth-to-Lifecycle “Biometric” Fingerprinting	23
D.2.2	All-stage verifiable fingerprints	25
D.3	Cross-Size Fingerprint Within the Qwen-2.5 series	26
E	Model Catalog and Decoding Settings	26
E.1	Model list by family	27
E.1.1	Qwen-2.5-7B	27
E.1.2	Qwen2.5-14B	27
E.1.3	Llama-3.1-8B	27
E.1.4	Mistral-7B-v0.3	28
E.1.5	Gemma-2-2B	28
E.1.6	Llama-2-7B	28
E.1.7	TinyLlama-1.1B	28
E.2	Sample Results	29

A LLM USAGE AND LIMITATION DISCUSSIONS

A.1 LLM USAGE

We use AI assistants, including ChatGPT and Gemini, for writing and formatting support. Their use is limited to grammar and style checks, improving the clarity of figure and table captions, and other surface-level edits. For programming-related tasks, we occasionally use GitHub Copilot and Claude as coding assistants, for example for code auto-completion and debugging suggestions.

A.2 LIMITATIONS AND FUTURE DIRECTIONS

Although this work identifies seed-born, training-persistent biases that serve as fingerprint signals, the proposed method remains within the white-box paradigm. It may not apply when only restricted API access is available, for example when responses are deterministic or safety-filtered. Nevertheless, the search for *true fingerprints* for LLMs remains at a very early stage, even under white-box access. Most existing fingerprinting approaches rely on training-induced signatures or alignment behaviors and therefore cannot robustly protect pretrained or early-stage models. By analogy to human fingerprints, which remain stable across age and environment, an LLM fingerprint should persist across continued training and deployment contexts rather than merely reflect incidental training artifacts. Achieving this objective—particularly under black-box access, where a substantial performance gap remains relative to white-box methods—is a challenging open problem that requires further work. We hope this study motivates principled and robust fingerprinting frameworks for large models.

B VALIDATION OF THE EMPIRICAL AND ANALYTICAL NULL DISTRIBUTION

The hypothesis test in Section 4 evaluates whether two models share lineage by comparing the distribution of their correlation statistics over random inputs against a null distribution. This null distribution represents the correlation statistics expected between two *independent models evaluated on random inputs*.

We consider two approximations to this null distribution: (i) an empirical simulation-based null obtained by applying the full pipeline to Gaussian surrogate outputs, and (ii) an analytical null derived from the known distribution of Kendall–Tau correlations. This section justifies these approximations in three steps: characterizing the true null with independently initialized models, introducing the Gaussian surrogate as a tractable simulation-based proxy, and validating the analytical approximation.

B.1 EMPIRICAL TRUE NULL FROM INDEPENDENTLY INITIALIZED MODELS

The ideal null distribution is the distribution of correlation statistics between two independent models evaluated on random inputs. The cleanest instantiation of this independence is a comparison between independently initialized models with different random seeds and no training. We estimate this empirical null by evaluating ~ 2500 such pairs and computing Kendall–Tau correlations on 10,000 random inputs (as in Section 3). Figure 3 shows that the resulting true-null distribution is empirically Gaussian and centered near zero.

B.2 GAUSSIAN SURROGATE AS A SIMULATION-BASED APPROXIMATION

Direct estimation of the null distribution with independently initialized models is computationally expensive because it requires access to many models and repeated evaluations. We therefore construct a tractable simulation-based null by using Gaussian surrogates as proxies for model outputs under independence.

This approximation is motivated by two considerations. First, approximating neural-network outputs with Gaussian distributions is common in auditing literature and is supported by empirical and theoretical studies (Carlini et al., 2022; Lee et al., 2019). Second, the SeedPrints pipeline depends only on relative ordering and selection operations, such as rank correlations and bottom- m selection, and does not use semantic structure in the inputs. Because randomly initialized models have independent parameter matrices, their null outputs are not expected to contain strong shared structure.




Figure 3: Empirical correlation distribution over ~ 2500 pairs of independently initialized models, evaluated on random inputs, closely matches a Gaussian distribution centered at zero.

Sampling independent Gaussian matrices therefore provides a simple and tractable proxy for the null behavior of independently initialized models.

We construct the surrogate null by sampling Gaussian matrices and applying the full pipeline, including identity-dimension selection and correlation computation:

1. Sample two Gaussian matrices of shape $\mathbb{R}^{N \times d_{\text{out}}}$,
2. Extract the most biased dimensions by selecting the bottom- m ranked average dimensions,
3. Take the intersection set $\mathcal{S}$ and compute correlations on it.

This simulation-based construction applies the same selection procedure to Gaussian samples and naturally captures the dependency patterns introduced by the pipeline itself (e.g., those arising from selecting $\mathcal{S}$ across dimensions). As shown in Figure 4 (middle), the resulting distribution matches the empirical null obtained from independently initialized models in terms of its centered Gaussian shape, indicating that the Gaussian surrogate provides an effective simulation-based approximation of the true null.

B.3 ANALYTICAL APPROXIMATION UNDER WEAK DEPENDENCE

Although the Gaussian surrogate captures the behavior of the null distribution under the full pipeline, it still requires repeated sampling and introduces estimation randomness. We therefore use a more efficient analytical approximation.

Under the null hypothesis that the two models are independent², each Kendall–Tau correlation τ_j is approximately zero-mean with known variance σ^2 . If correlations across identity dimensions were independent, the central limit theorem would imply

$$\bar{\tau} = \frac{1}{|\mathcal{S}|} \sum_{j \in \mathcal{S}} \tau_j \sim \mathcal{N}\left(0, \frac{\sigma^2}{|\mathcal{S}|}\right).$$

In practice, softmax normalization introduces dependencies across dimensions. However, Figure 4 shows that the analytical Gaussian closely matches the distribution obtained from the Gaussian surrogate pipeline across different settings, including varying numbers of random inputs n and output dimensions D . This agreement indicates that the induced dependencies are sufficiently weak and do not materially affect the accuracy of the approximation. A plausible explanation is that the relatively high temperature ($T = 10$) smooths the output distribution and reduces coupling between dimensions.

Together, these results justify the analytical Gaussian null in Section 4 as an efficient, deterministic, and accurate approximation that eliminates the need for costly empirical baseline construction.

²As discussed in Appendix B.2, although outputs may exhibit weak dependencies due to shared inputs, this approximation remains accurate empirically.




Figure 4: Validation of the analytical Gaussian null. We apply the full pipeline (including identity-dimension selection and correlation aggregation) to Gaussian surrogate outputs and compare the resulting distribution of $\bar{\tau}$ with the analytical Gaussian $\mathcal{N}(0, \sigma^2/|S|)$. The close match indicates that the analytical approximation accurately captures the null behavior, despite weak dependencies introduced by the pipeline.

On variance scaling. The true null obtained from randomly initialized models has slightly smaller variance than the Gaussian null approximation. This mismatch does not affect our hypothesis test results. Our primary evaluation for the simulation-based null in Appendix D.2 relies on the Mann–Whitney U-test, which is rank-based and thus invariant to any positive scaling: for any $c > 0$, $x_i < x_j \Leftrightarrow cx_i < cx_j$. Consequently, all rank statistics and rejection thresholds remain unchanged. For tests that depend on absolute variance, we additionally evaluate variance-sensitive statistics such as the t -test, whose statistic scales as $t' = t/c$ when variance decreases. Since $c < 1$ in our case, scaling reduces the magnitude of t and makes the test more conservative rather than inflating significance. Empirically, across the experiments in Appendix D.2, the variance scaling does not affect the power of the t -test, indicating that the distributional separation between the alternative hypothesis and the null is substantially larger than the scaled variance.

For the analytical null used in the main paper, the variance is explicitly included in the normalization of the z -score. Any scale discrepancy is therefore absorbed in the denominator. Since the analytical variance is slightly larger than the empirical one, the resulting z -scores are smaller, making the test conservative and preventing false positives.

C THEORETICAL PROPERTIES OF OUR METHODS

Our goal is to identify the output dimensions that the initialized model is least likely to predict under random inputs. These dimensions correspond to directions inherently disfavored by the model because of initialization-induced bias, rather than because of the input distribution, and they form the basis of a stable fingerprint signal.

Setup. Consider a model g with fixed parameters and random inputs x_i (uniform random tokens or random embeddings as in the main paper). For each output dimension $j \in [d_{\text{out}}]$, $g_j(x)$ is therefore a random variable with population mean $\bar{g}_j := \mathbb{E}[g_j(x)]$. The population-level directions we aim to capture are the m output dimensions with the smallest expected values, formally defined as

$$\mathcal{M} = \arg \min_{J \subseteq \{1, \dots, d_{\text{out}}\}, |J|=m} \sum_{j \in J} \bar{g}_j.$$

Let $(\bar{g}_{(1)}, \dots, \bar{g}_{(d_{\text{out}})})$ denote the sorted means in non-decreasing order, and write

$$\Gamma(m) := \bar{g}_{(m+1)} - \bar{g}_{(m)} > 0$$

for the population “gap” between the m -th and $(m+1)$ -th smallest means. This gap measures how well-separated the target set $\mathcal{M}$ is from the remaining coordinates. Importantly, for a fixed m and input distribution, $\Gamma(m)$ depends solely on the model’s inherent initialization bias. *However, increasing the length of random input sequences amplifies this bias via concentration, making the separation more pronounced at the population level.*

There are two natural empirical strategies to extract such population-level least-likely dimensions:

1. **Sample-mean-based extraction (ours).** Estimate the expectation of each output coordinate by the empirical mean over random inputs $X = \{x_i\}_{i=1}^n$ and select the m coordinates with the smallest value. Following the notation in Section 4, where g denotes the model (either the base model f or the suspect model f'):

$$\hat{g} = \frac{1}{n} \sum_{i=1}^n g(x_i), \quad \widehat{\mathcal{M}}^{\text{mean}} = \arg \min_{J \subseteq \{1, \dots, d_{\text{out}}\}, |J|=m} \sum_{j \in J} \hat{g}_j.$$

2. **Per-sample argmin voting.** For each random input $x_i \in \mathbb{R}^{\ell \times d}$, select the smallest-value coordinate and then take the m most frequent dimensions:

$$j^*(x_i) = \arg \min_{j \in \{1, \dots, d_{\text{out}}\}} g_j(x_i), \quad \widehat{\mathcal{M}}^{\text{vote}} = \arg \max_{J \subseteq \{1, \dots, d_{\text{out}}\}, |J|=m} \sum_{i=1}^n \mathbf{1}\{j^*(x_i) \in J\}.$$

Let $p_j := \Pr(j^*(x) = j)$ denote the population probability that index j attains the minimum, and let

$$\mathcal{M}_{\text{vote}} := \arg \max_{J \subseteq \{1, \dots, d_{\text{out}}\}, |J|=m} \sum_{j \in J} p_j$$

be the population top- m set under voting.

Both approaches are intuitively reasonable. However, as we show below, the sample-mean-based strategy is provably stable, whereas per-sample argmin voting is discontinuous and can flip under arbitrarily small perturbations, making it inherently noise-sensitive.

C.1 STABILITY ANALYSIS OF MEAN-BASED IDENTITY EXTRACTION VS. PER-SAMPLE ARGMIN VOTING

C.1.1 STABILITY OF SAMPLE-MEAN-BASED EXTRACTION

We first show that the sample-mean-based operator, which selects the m smallest coordinates in the empirical mean vector, is stable under small perturbations to model outputs. The perturbations represents realistic variations in outputs from the *same* model, e.g., numerical noise, or evaluation under different precision modes.

Formally, we model perturbed outputs as $g'(x) = g(x) + \eta(x)$, where $\eta(x) \in \mathbb{R}^{d_{\text{out}}}$ represents stochastic deviations. We assume that for each sample x_i and coordinate j , the random variables $\eta_j(x_i)$ are independent across samples (not necessarily across coordinates) and σ -sub-Gaussian with $\mathbb{E}[\eta_j(x_i)] = 0$.³

Lemma C.1 (Stability of sample-mean-based extraction). *Let $\hat{g}$ and $\hat{g}'$ be the empirical means over $X = \{x_i\}_{i=1}^n$, i.e., $\hat{g} = \frac{1}{n} \sum_{i=1}^n g(x_i)$ and $\hat{g}' = \frac{1}{n} \sum_{i=1}^n g'(x_i) = \hat{g} + \frac{1}{n} \sum_{i=1}^n \eta(x_i)$. Let $\widehat{\Gamma}(m)$ denote the empirical gap at rank m . Then,*

$$\Pr(\widehat{\mathcal{M}}^{\text{mean}}(g') \neq \widehat{\mathcal{M}}^{\text{mean}}(g)) \leq 2d_{\text{out}} \exp\left(-\frac{n\widehat{\Gamma}(m)^2}{8\sigma^2}\right).$$

Interpretation. The lemma provides a quantitative upper bound on the probability that the sample-mean-based extractor changes its selected set under perturbations. This instability probability decreases exponentially with the number of averaged samples n , due to variance reduction from σ^2 per sample to σ^2/n in the empirical mean.

Proof. We condition on the sample $X = \{x_i\}_{i=1}^n$ throughout the proof, so that $\hat{g}$ and $\widehat{\Gamma}(m)$ are fixed, and all probabilities are taken over the randomness of $\{\eta(x_i)\}_{i=1}^n$ only.

³This sub-Gaussian assumption is mild: bounded random variables are σ -sub-Gaussian (Hoeffding's lemma), and common initialization schemes together with bounded activations yield light-tailed behavior in both initialized and trained networks. The Gaussian case $\eta_j(x_i) \sim \mathcal{N}(0, \sigma^2)$ is a special case of this assumption.

For any coordinate j , by definition of $\hat{g}'_j$, $\hat{g}'_j - \hat{g}_j = \frac{1}{n} \sum_{i=1}^n \eta_j(x_i)$. Each $\eta_j(x_i)$ is mean-zero and σ -sub-Gaussian, so their average is mean-zero and $(\sigma/\sqrt{n})$ -sub-Gaussian: the variance proxy scales as $1/n$ under averaging. Hence, for any $\varepsilon > 0$, by the standard sub-Gaussian tail bound,

$$\Pr(|\hat{g}'_j - \hat{g}_j| \geq \varepsilon \mid X) \leq 2 \exp\left(-\frac{n\varepsilon^2}{2\sigma^2}\right). \quad (1)$$

Now consider the event that *any* coordinate fluctuates by more than ε :

$$\left\{ \max_{1 \leq j \leq d_{\text{out}}} |\hat{g}'_j - \hat{g}_j| \geq \varepsilon \right\}.$$

Taking a union bound over all d_{out} coordinates yields

$$\Pr\left(\max_j |\hat{g}'_j - \hat{g}_j| \geq \varepsilon \mid X\right) \leq \sum_{j=1}^{d_{\text{out}}} \Pr(|\hat{g}'_j - \hat{g}_j| \geq \varepsilon \mid X) \leq 2d_{\text{out}} \exp\left(-\frac{n\varepsilon^2}{2\sigma^2}\right). \quad (2)$$

We now show that

$$\left\{ \widehat{\mathcal{M}}^{\text{mean}}(g') \neq \widehat{\mathcal{M}}^{\text{mean}}(g) \right\} \subseteq \left\{ \max_j |\hat{g}'_j - \hat{g}_j| \geq \widehat{\Gamma}(m)/2 \right\}.$$

Equivalently, if $\max_j |\hat{g}'_j - \hat{g}_j| < \widehat{\Gamma}(m)/2$, then the bottom- m identity set remains unchanged.

Let $j_{(1)}, \dots, j_{(d_{\text{out}})}$ be a permutation of $\{1, \dots, d_{\text{out}}\}$ such that $\hat{g}_{j_{(1)}} \leq \hat{g}_{j_{(2)}} \leq \dots \leq \hat{g}_{j_{(d_{\text{out}})}}$. Recall that the empirical gap at rank m is $\widehat{\Gamma}(m) := \hat{g}_{j_{(m+1)}} - \hat{g}_{j_{(m)}} > 0$. Denote $\widehat{\mathcal{M}}^{\text{mean}}(g) := \{j_{(1)}, \dots, j_{(m)}\}$ and $\widehat{\mathcal{M}}_c^{\text{mean}}(g) := \{j_{(m+1)}, \dots, j_{(d_{\text{out}})}\}$. For any $i \in \widehat{\mathcal{M}}^{\text{mean}}(g)$ and any $\ell \in \widehat{\mathcal{M}}_c^{\text{mean}}(g)$, $\hat{g}_\ell - \hat{g}_i \geq \hat{g}_{j_{(m+1)}} - \hat{g}_{j_{(m)}} = \widehat{\Gamma}(m)$. With $|\hat{g}'_i - \hat{g}_i| < \widehat{\Gamma}(m)/2$ and $|\hat{g}'_\ell - \hat{g}_\ell| < \widehat{\Gamma}(m)/2$, we have

$$\hat{g}'_\ell - \hat{g}'_i \geq (\hat{g}_\ell - \frac{\widehat{\Gamma}(m)}{2}) - (\hat{g}_i + \frac{\widehat{\Gamma}(m)}{2}) \geq \widehat{\Gamma}(m) - \widehat{\Gamma}(m) = 0.$$

Thus no coordinate in $\widehat{\mathcal{M}}_c^{\text{mean}}(g)$ can become smaller than any coordinate in $\widehat{\mathcal{M}}^{\text{mean}}(g)$ under this deviation bound, and therefore

$$\widehat{\mathcal{M}}^{\text{mean}}(g') = \widehat{\mathcal{M}}^{\text{mean}}(g).$$

Since $\widehat{\Gamma}(m)$ is fixed under conditioning on X , plugging $\varepsilon = \widehat{\Gamma}(m)/2$ into (2) yields

$$\Pr\left(\widehat{\mathcal{M}}^{\text{mean}}(g') \neq \widehat{\mathcal{M}}^{\text{mean}}(g)\right) \leq \Pr\left(\max_j |\hat{g}'_j - \hat{g}_j| \geq \widehat{\Gamma}(m)/2 \mid X\right) \leq 2d_{\text{out}} \exp\left(-\frac{n\widehat{\Gamma}(m)^2}{8\sigma^2}\right).$$

□

C.1.2 INSTABILITY OF PER-SAMPLE ARGMIN VOTING

In contrast, per-sample voting applies a discontinuous winner-take-all operator on each input. Thus, arbitrarily small perturbations to outputs can change decisions with non-vanishing probability. We still consider the same stochastic perturbation model $g'(x) = g(x) + \eta(x)$.

Lemma C.2 (Instability of per-sample argmin voting). *For a fixed input x , let*

$$j^*(x; g) := \arg \min_j g_j(x), \quad \Delta_g(x) := \min_{j \neq j^*(x; g)} \{g_j(x) - g_{j^*(x; g)}(x)\} > 0$$

denote the argmin and its margin. Then, for any fixed x ,

$$\Pr(j^*(x; g') \neq j^*(x; g) \mid x) \geq \Phi\left(-\frac{\Delta_g(x)}{\sqrt{2}\sigma}\right),$$

and for n i.i.d. inputs $x_1, \dots, x_n$,

$$\Pr\left(\exists i \leq n : j^*(x_i; g') \neq j^*(x_i; g)\right) \geq 1 - (1 - \alpha)^n,$$

where $\alpha := \mathbb{E}_x \left[\Phi\left(-\frac{\Delta_g(x)}{\sqrt{2}\sigma}\right) \right]$.

Interpretation. Unlike mean extraction, voting does not benefit from variance reduction: each input incurs a direct hard decision based on noisy outputs. Increasing n only increases the number of opportunities for index-flips.

Proof. We first bound the per-input change probability $\Pr(j^*(x; g') \neq j^*(x; g) \mid x)$. Fix x and write $j^* := j^*(x; g)$ for brevity. By definition of the margin, $\Delta_g(x) = \min_{j \neq j^*} \{g_j(x) - g_{j^*}(x)\}$. Let $j^{(2)}$ be an index attaining this minimum, i.e., $g_{j^{(2)}}(x) - g_{j^*}(x) = \Delta_g(x)$.

Under the perturbation $g'(x) = g(x) + \eta(x)$, we have

$$g'_{j^{(2)}}(x) - g'_{j^*}(x) = [g_{j^{(2)}}(x) - g_{j^*}(x)] + [\eta_{j^{(2)}}(x) - \eta_{j^*}(x)] = \Delta_g(x) + [\eta_{j^{(2)}}(x) - \eta_{j^*}(x)].$$

Since $\eta_{j^{(2)}}(x)$ and $\eta_{j^*}(x)$ are independent $\mathcal{N}(0, \sigma^2)$ variables, their difference is Gaussian with

$$\eta_{j^{(2)}}(x) - \eta_{j^*}(x) \sim \mathcal{N}(0, 2\sigma^2).$$

The event that the argmin changes, $j^*(x; g') \neq j^*(x; g)$, certainly occurs whenever $g'_{j^{(2)}}(x) \leq g'_{j^*}(x)$. Hence,

$$\Pr(j^*(x; g') \neq j^*(x; g) \mid x) \geq \Pr(g'_{j^{(2)}}(x) \leq g'_{j^*}(x) \mid x).$$

Using the expression above,

$$\Pr(g'_{j^{(2)}}(x) \leq g'_{j^*}(x) \mid x) = \Pr(g'_{j^{(2)}}(x) - g'_{j^*}(x) \leq 0 \mid x) = \Pr(\Delta_g(x) + [\eta_{j^{(2)}}(x) - \eta_{j^*}(x)] \leq 0 \mid x).$$

Let $Z \sim \mathcal{N}(0, 1)$. Since $\eta_{j^{(2)}}(x) - \eta_{j^*}(x) \stackrel{d}{=} \sqrt{2}\sigma Z$, we obtain

$$\Pr(\Delta_g(x) + [\eta_{j^{(2)}}(x) - \eta_{j^*}(x)] \leq 0 \mid x) = \Pr\left(Z \leq -\frac{\Delta_g(x)}{\sqrt{2}\sigma}\right) = \Phi\left(-\frac{\Delta_g(x)}{\sqrt{2}\sigma}\right).$$

Therefore, $\Pr(j^*(x; g') \neq j^*(x; g) \mid x) \geq \Phi\left(-\frac{\Delta_g(x)}{\sqrt{2}\sigma}\right)$.

Now consider n i.i.d. inputs $x_1, \dots, x_n$. Conditional on the inputs, the events

$$E_i := \{j^*(x_i; g') \neq j^*(x_i; g)\} \quad E_i^c := \{j^*(x_i; g') = j^*(x_i; g)\}$$

are independent across i , because the perturbations $\eta(x_i)$ are independent across i . Moreover,

$$\Pr(E_i \mid x_i) = \Pr(j^*(x_i; g') \neq j^*(x_i; g) \mid x) \geq \Phi\left(-\frac{\Delta_g(x_i)}{\sqrt{2}\sigma}\right) \geq \Phi\left(-\frac{\Delta_g(x_i)}{\sqrt{2}\sigma}\right).$$

Define $q(x_i) := \Pr(j^*(x_i; g') \neq j^*(x_i; g) \mid x)$, and $\alpha := \mathbb{E}_x\left[\Phi\left(-\frac{\Delta_g(x)}{\sqrt{2}\sigma}\right)\right]$. Taking expectations over the randomness of x_i , we obtain $\mathbb{E}[\Pr(E_i \mid x_i)] = \alpha$. Therefore, by independence across i ,

$$\Pr(E_1^c \cap \dots \cap E_n^c) = \mathbb{E}\left[\prod_{i=1}^n (1 - q(x_i))\right] \leq \mathbb{E}\left[\prod_{i=1}^n (1 - \Phi(-\Delta_g(x_i)/(\sqrt{2}\sigma)))\right] \leq (1 - \alpha)^n,$$

where the last inequality uses i.i.d. inputs and Jensen's inequality applied to the convex function $u \mapsto \log(1 - u)$ on $u \in [0, 1)$. Therefore,

$$\Pr(\exists i \leq n : j^*(x_i; g') \neq j^*(x_i; g)) = 1 - \Pr(E_1^c \cap \dots \cap E_n^c) \geq 1 - (1 - \alpha)^n.$$

Finally, note that the empirical voting set $\widehat{\mathcal{M}}^{\text{vote}}$ is determined by the empirical counts of the dimensions $\{j^*(x_i; g)\}_{i=1}^n$. If the empirical gap between the m -th and $(m+1)$ -th most frequent dimensions is of order $O(1/n)$ (i.e., only a constant number of votes), then a change in $j^*(x_i)$ for a single input x_i is sufficient to swap the ordering across this boundary and thus alter the top- m set. Combining this observation with the lower bound $1 - (1 - \alpha)^n$ shows that, in such regimes, the probability that $\widehat{\mathcal{M}}^{\text{vote}}(g') \neq \widehat{\mathcal{M}}^{\text{vote}}(g)$ is bounded away from zero uniformly in n , i.e., per-sample argmin voting does not enjoy variance reduction with more samples. $\square$

C.2 INFLUENCE OF HYPERPARAMETERS

We now analyze how many random inputs are required for the empirical mean estimator to recover the true bias profile. We reuse notation from **Setup**: each output coordinate $g_j(x)$ is a random variable under random inputs $x \sim \mathcal{D}$, with population mean $\bar{g}_j := \mathbb{E}[g_j(x)]$. $(\bar{g}_{(1)}, \dots, \bar{g}_{(d_{\text{out}})})$ represents the sorted means in non-decreasing order, and $\gamma := \bar{g}_{(m+1)} - \bar{g}_{(m)} > 0$ is the separation margin. We estimate $\bar{g}$ using n i.i.d. random inputs:

$$\hat{g}_j := \frac{1}{n} \sum_{i=1}^n g_j(x_i), \quad \hat{g} = (\hat{g}_1, \dots, \hat{g}_{d_{\text{out}}}).$$

The target dimension set is $\mathcal{M}$ with size m , and the empirical estimator returns $\widehat{\mathcal{M}}$.

Lemma C.3 (Sufficient condition for recovering the true set). *If $\|\hat{g} - \bar{g}\|_\infty \leq \varepsilon$ and $2\varepsilon < \gamma$, then the empirical bottom- m set equals the true set: $\widehat{\mathcal{M}} = \mathcal{M}$.*

Proof. For any $j \in \mathcal{M}$ and $j' \notin \mathcal{M}$, we have $\bar{g}_j \leq \bar{g}_{(m)}$ and $\bar{g}_{j'} \geq \bar{g}_{(m+1)}$. The ℓ_∞ bound gives

$$\hat{g}_j \leq \bar{g}_{(m)} + \varepsilon, \quad \hat{g}_{j'} \geq \bar{g}_{(m+1)} - \varepsilon.$$

Thus,

$$\hat{g}_j - \hat{g}_{j'} \leq -(\gamma - 2\varepsilon) < 0,$$

so no outside index can enter nor inside index leave the bottom- m set. $\square$

We next interpret this recovery condition in terms of the SeedPrints hyperparameters. We adopt the same sub-Gaussian assumption as in Appendix C.1. This assumption is mild: by Hoeffding's inequality, any bounded random variable is σ -sub-Gaussian with σ proportional to its range. In our setting, each $g_j(x)$ corresponds to a logit or final hidden activation, which is typically bounded or light-tailed in both initialized networks (because of controlled parameter initialization) and trained networks.

Corollary C.4 (Sample complexity under sub-Gaussian coordinates). *Suppose each $g_j(x)$ is σ -sub-Gaussian. Then with probability at least $1 - \delta$,*

$$\|\hat{g} - \bar{g}\|_\infty \leq \sigma \sqrt{\frac{2 \log(2d_{\text{out}}/\delta)}{n}}.$$

Consequently, if

$$n \geq \frac{8\sigma^2}{\gamma^2} \log\left(\frac{2d_{\text{out}}}{\delta}\right),$$

then $\widehat{\mathcal{M}} = \mathcal{M}$ with probability at least $1 - \delta$.

Proof. For fixed j , sub-Gaussian concentration implies

$$\Pr(|\hat{g}_j - \bar{g}_j| > \varepsilon) \leq 2 \exp(-n\varepsilon^2/2\sigma^2).$$

A union bound over all d_{out} coordinates yields

$$\Pr\left(\|\hat{g} - \bar{g}\|_\infty > \varepsilon\right) \leq \sum_{j=1}^{d_{\text{out}}} \Pr(|\hat{g}_j - \bar{g}_j| > \varepsilon) \leq 2d_{\text{out}} \exp\left(-\frac{n\varepsilon^2}{2\sigma^2}\right).$$

Set the right-hand side to δ and solve for ε to obtain, with probability at least $1 - \delta$,

$$\|\hat{g} - \bar{g}\|_\infty \leq \sigma \sqrt{\frac{2 \log(2d_{\text{out}}/\delta)}{n}}.$$

Now choose

$$\varepsilon := \sigma \sqrt{\frac{2 \log(2d_{\text{out}}/\delta)}{n}}.$$

If $n \geq \frac{8\sigma^2}{\gamma^2} \log\left(\frac{2d_{\text{out}}}{\delta}\right)$ (i.e., if $2\varepsilon < \gamma$), then Lemma C.3 applies and implies $\widehat{\mathcal{M}} = \mathcal{M}$ with probability at least $1 - \delta$. $\square$

Table 9: Seed-specific fingerprint behavior under different random initializations across model families.

(a) LLaMA-style models					(b) Qwen-style models				
Seed Pair	Logits Output		Hidden State		Seed Pair	Logits Output		Hidden State	
	<i>t-test</i>	<i>U-test</i>	<i>t-test</i>	<i>U-test</i>		<i>t-test</i>	<i>U-test</i>	<i>t-test</i>	<i>U-test</i>
s_{42} vs. s_{2000}	0.404	0.456	0.357	0.532	s_{123} vs. s_{1000}	0.741	0.727	0.094	0.074
s_{123} vs. s_{42}	0.214	0.295	0.678	0.565	s_{1000} vs. s_{123}	0.954	0.971	0.125	0.094
s_{1000} vs. s_{123}	0.219	0.246	0.363	0.335	s_{42} vs. s_{2000}	0.273	0.360	0.451	0.529
s_{2000} vs. s_{1000}	0.282	0.291	0.434	0.481	s_{2000} vs. s_{42}	0.215	0.206	0.230	0.295

Table 10: (LLaMA-style models) Trained models retain the fingerprint of their initialization (p -value < 0.01).

Model Pair	Logits Output		Hidden State		Baselines			
	<i>t-test</i>	<i>u-test</i>	<i>t-test</i>	<i>u-test</i>	Intrinsic	REEF	PCS	ICS
s_{42}^{init} vs. s_{42}^{base}	3.33e-3✓	1.02e-3✓	2.20e-8✓	6.28e-8✓	-0.021×	0.375×	0.580×	0.196×
s_{123}^{init} vs. s_{123}^{base}	2.06e-3✓	7.33e-3✓	7.09e-6✓	1.37e-5✓	0.149×	0.369×	0.581×	0.188×
s_{1000}^{init} vs. s_{1000}^{base}	2.44e-3✓	4.14e-3✓	5.58e-4✓	2.81e-3✓	-0.252×	0.381×	0.581×	0.188×
s_{2000}^{init} vs. s_{2000}^{base}	5.63e-3✓	6.76e-3✓	4.00e-10✓	1.27e-9✓	-0.337×	0.331×	0.581×	0.188×

Therefore, accurate recovery requires $n = \Omega\left(\frac{\log d_{out}}{\gamma^2}\right)$: the number of random queries grows only logarithmically with the output dimension but scales inversely with the squared separation margin γ . The margin γ is governed by two factors:

- **Intrinsic bias strength.** Longer random input sequences amplify initialization-induced bias via concentration, increasing the population separation between coordinates.
- **Choice of m .** Different choices of m correspond to different parts of the ranked spectrum $(\bar{g}_{(1)}, \dots, \bar{g}_{(d_{out})})$, thus yielding different effective margins in practice.

The first factor is controlled by the input length l , while the second is selected empirically from the observed gaps in the mean spectrum.

D MORE EXPERIMENT DETAILS

D.1 IMPLEMENTATION DETAILS AND LICENSES

We train all models with the Hugging Face Transformers Trainer (Wolf et al., 2020) and use Accelerate (Gugger et al., 2022) for distributed runs. All open-source models are loaded from their official Hugging Face releases and used under their original licenses: Llama models under the Meta Llama Community License, and other models under Apache-2.0. All datasets are downloaded through the Hugging Face Datasets library, which is Apache-2.0; dataset content follows each dataset’s stated license.

D.2 COMPLEMENTARY SEEDPRINTS RESULTS USING SIMULATION-BASED NULL

This section reports complementary results for Section 5. We use the empirical simulation-based null in Algorithm 1 with both *logits* and *hidden state* outputs, and evaluate decisions with the one-sided *t*-test (*t*-test) and Mann–Whitney *U* test (*U* test).

D.2.1 BIRTH-TO-LIFECYCLE “BIOMETRIC” FINGERPRINTING

We train 12-layer, 12-head LLaMA-style models (Touvron et al., 2023) with RoPE (Su et al., 2021) and Qwen-style models (Team, 2024) from scratch. Because the simulation-based random baseline is stochastic, we report p -values averaged over 10 independent trials and use a significance level of

Table 11: (Qwen-style models) Trained models retain the fingerprint of their initialization models (p -value < 0.01).

Model Pair	Logits Output		Hidden State	
	t -test	U -test	t -test	U -test
s_{123}^{init} vs. s_{123}^{base}	2.38e-03	1.68e-03	7.36e-15	3.38e-13
s_{1000}^{init} vs. s_{1000}^{base}	1.35e-04	4.84e-05	4.41e-13	2.01e-11
s_{42}^{init} vs. s_{42}^{base}	1.39e-03	1.46e-03	1.06e-24	2.05e-19
s_{2000}^{init} vs. s_{2000}^{base}	1.80e-03	1.28e-03	4.87e-24	1.92e-20

Table 12: Fingerprint persistence under continual training on diverse datasets (base model: seed 1000, corpus openwebtext). U -test refers to the Mann–Whitney U test.

(a) LLaMA-style models

Setting	Ours (logits)		Ours (hidden)		Baselines			
Continual corpus (seed)	t -test	U -test	t -test	U -test	Intrinsic	REEF	PCS	ICS
TinyStories (1000)	0 [✓]	0 [✓]	0 [✓]	7.77e-89 [✓]	1.000 [✓]	0.759 [×]	0.999 [✓]	0.996 [✓]
TinyStories (123)	1.000 [✓]	1.00 [✓]	0.943 [✓]	0.902 [✓]	0.950 [×]	0.658 [✓]	0.332 [✓]	0.012 [✓]
the_stack (1000)	0 [✓]	1.73e-287 [✓]	0 [✓]	3.09e-69 [✓]	0.489 [×]	0.557 [×]	0.585 [×]	0.123 [×]
the_stack (123)	0.616 [✓]	0.479 [✓]	0.732 [✓]	0.831 [✓]	0.445 [✓]	0.580 [✓]	0.301 [✓]	0.026 [✓]

(b) Qwen-style models

Setting	Ours (logits)		Ours (hidden)		Baselines			
Continual corpus (seed)	t -test	U -test	t -test	U -test	Intrinsic	REEF	PCS	ICS
TinyStories (1000)	5.36e-09 [✓]	1.92e-07 [✓]	8.49e-214 [✓]	5.09e-71 [✓]	1.000 [✓]	0.957 [✓]	0.999 [✓]	0.996 [✓]
TinyStories (123)	0.434 [✓]	0.433 [✓]	0.256 [✓]	0.065 [✓]	0.913 [×]	0.199 [✓]	0.328 [✓]	0.039 [✓]
the_stack (1000)	0 [✓]	4.16e-237 [✓]	1.16e-211 [✓]	2.30e-76 [✓]	0.999 [✓]	0.313 [×]	0.995 [✓]	0.976 [✓]
the_stack (123)	0.999 [✓]	0.993 [✓]	0.610 [✓]	0.491 [✓]	0.916 [×]	0.255 [✓]	0.328 [✓]	0.038 [✓]

$\alpha = 0.01$. The absolute magnitude of extremely small p -values is not meaningful: once p falls below numerical and sampling noise (e.g., $< 10^{-20}$), values such as 10^{-260} should not be interpreted as stronger evidence than 10^{-20} ; both decisively reject the null.

Different initialization seeds produce distinct fingerprints Table 9a and Table 9b report p -values from correlation tests between pairs of models initialized with different random seeds (42, 123, 1000, and 2000) under the t -test and U -test, for LLaMA-style and Qwen-style models, respectively. All p -values are consistently > 0.01 , indicating that SeedPrints reliably distinguishes models initialized from different seeds. Distinct seeds therefore yield distinct fingerprint behaviors, allowing models to be separated “at birth.”

Training preserves the initialization fingerprint. Table 10 (LLaMA-style) and Table 11 (Qwen-style) compare each initialization model s^{init} with its descendant s^{base} trained on OpenWebText (Gokaslan et al., 2019) (≈ 10 B tokens). Across all seed–model pairs, p -values are consistently < 0.01 , indicating that the bias profiles remain strongly correlated and share a common lineage. The trained model therefore inherits the fingerprint of its initialization; training does not erase the initialization fingerprint. We also evaluate baseline methods in Table 10; without exception, they fail to distinguish across seeds, suggesting that their separability stems from training-induced artifacts rather than initialization.


Identical data and order do not make fingerprints converge In Table 13a (LLaMA-style models) and Table 13b (Qwen-style models), all four “suspicious” models s_i^{base} for $i \in \{42, 123, 100, 2000\}$ are trained on *exactly the same corpus* (OpenWebText) and *in the same data order*; the training seed fixes the data order, so the only difference is the initialization seed i . This experiment tests whether identical data and order erase or confound fingerprint behavior. Across all cross-seed pairs, p -values

Table 13: Identical dataset and training order do not make fingerprint behaviors converge across different initializations.

(a) LLaMA-style models					(b) Qwen-style models				
Model Pair	Logits Output		Hidden State		Model Pair	Logits Output		Hidden State	
	<i>t</i> -test	<i>U</i> -test	<i>t</i> -test	<i>U</i> -test		<i>t</i> -test	<i>U</i> -test	<i>t</i> -test	<i>U</i> -test
s_{123}^{init} vs. s_{1000}^{base}	0.484	0.500	0.385	0.486	s_{123}^{init} vs. s_{1000}^{base}	0.598	0.638	0.286	0.254
s_{1000}^{init} vs. s_{2000}^{base}	0.946	0.956	0.135	0.096	s_{1000}^{init} vs. s_{123}^{base}	0.804	0.805	0.236	0.240
s_{42}^{init} vs. s_{123}^{base}	0.598	0.589	0.426	0.337	s_{42}^{init} vs. s_{2000}^{base}	0.589	0.608	0.226	0.2043
s_{2000}^{init} vs. s_{42}^{base}	0.756	0.781	0.388	0.287	s_{2000}^{init} vs. s_{42}^{base}	0.523	0.482	0.312	0.323



(a) LLaMA-style models



(b) Qwen-style models

Figure 5: SeedPrints verifies lineage at every checkpoint (p -values < 0.01) across model families.

remain consistently > 0.01 , in sharp contrast to the near-zero values in Table 10. Fingerprints therefore remain seed-specific even under identical data and curriculum.

Continual training on diverse datasets does not confound the fingerprint The earlier experiments establish the biometric character of SeedPrints. From a copyright perspective, however, the inability of prior work to distinguish models with different seeds is not the main weakness, because initialization seeds have no clear copyright status. The more consequential fragility is that attribution can be misled by data distribution, *failing to recognize lineage when the training distribution shifts substantially during continued training*.

In Table 12a and Table 12b, we continue training a base model (seed 1000, pretrained on OpenWebText (Gokaslan et al., 2019)) on two very different datasets: TinyStories (Eldan & Li, 2023) (synthetic children’s stories) and The Stack (Kocetkov et al., 2022) (permissively licensed GitHub code). We compare (i) true descendants trained from the base against (ii) *distractors* derived from a different base model, initialized with seed 123 and trained with a different data order on OpenWebText, then continually trained on the *same* corpus. The experiment asks whether attribution methods identify the descendant that truly shares lineage with the base.

The prior baselines all fail under the code setting (The Stack), misclassifying true descendants as distractors. This indicates that they largely track domain similarity rather than lineage identity: TinyStories is closer in distribution to the pretraining corpus (OpenWebText), whereas The Stack diverges sharply, and this larger distribution shift can bypass detection. SeedPrints correctly attributes lineage across both corpora. Its fingerprint is therefore not a proxy for data distribution: it survives substantial domain shift and persists beyond the initial pretraining stage.

D.2.2 ALL-STAGE VERIFIABLE FINGERPRINTS

Controlled pretraining trajectory. We first conduct a controlled pretraining experiment on OpenWebText, where the initialization seed is known. Each intermediate checkpoint is treated as the base, and we test whether SeedPrints reliably identifies its descendants along the same training trajectory. All variants consistently recognize the suspect model as belonging to the same lineage, with p -values

Table 14: Cross-size comparison between Qwen-2.5-14B (target) and other Qwen-2.5 models.

	Qwen-2.5-0.5B	Qwen-2.5-1.5B	Qwen-2.5-3B	Qwen-2.5-7B
SeedPrint ($\uparrow$)	0.7184	0.9707	0.3898	0.5429
REEF ($\downarrow$)	0.8250	0.8399	0.8833	0.8461
HUREF ($\downarrow$)	-6.37×10^{-5}	-2.84×10^{-5}	-8.70×10^{-6}	-6.08×10^{-6}

remaining below the 0.01 threshold (Figure 5a for LLaMA-style models and Figure 5b for Qwen-style models). This confirms the stability of SeedPrints under continuous optimization.

D.3 CROSS-SIZE FINGERPRINT WITHIN THE QWEN-2.5 SERIES

To test whether initialization-born fingerprints collapse to a family-level signature when model size changes, we conduct a cross-size study within the Qwen-2.5 series. We treat Qwen-2.5-14B as the protected target model and consider Qwen-2.5- $\{0.5B, 1.5B, 3B, 7B\}$ as suspect models that share the same architectural family but are not trained as descendants of Qwen-2.5-14B. Table 14 reports SeedPrint ($\uparrow$) together with two similarity-based baselines, REEF ($\downarrow$) and HUREF ($\downarrow$). As discussed in the main paper, SeedPrints outputs a p -value for the hypothesis test “the suspect shares initialization-born fingerprints with the target”, while REEF and HUREF output continuous scores and require the user to choose a threshold.

Several observations are important. First, all SeedPrints p -values are comfortably above a typical significance level (e.g., $\alpha=0.01$). Under the hypothesis-testing view, none of the smaller Qwen-2.5 models is flagged as sharing the same lineage as Qwen-2.5-14B. This behavior is desirable: although these models belong to the same family and reuse a similar architectural design, the fingerprinting test does not collapse to a single family-level signature that would spuriously conflate distinct pre-training runs.

The cross-size p -values occupy different regions of the $[0, 1]$ range rather than concentrating around a single value. We do *not* interpret these magnitudes as a continuous “similarity score”—their role is to support a calibrated yes/no decision at a fixed significance level. Nonetheless, their spread indicates that the test statistic remains sensitive to the concrete combination of architecture and initialization seed instead of degenerating into a generic pattern shared by all Qwen-2.5 variants.

Finally, the similarity-based baselines behave differently. REEF produces scores in a relatively narrow band (0.8250–0.8833) across all sizes, while HUREF outputs values that are numerically very close to zero (on the order of 10^{-5}). Because these methods do not provide a built-in decision rule, operators must manually choose thresholds, and it is unclear how a single cutoff would robustly separate cross-size variants from genuinely lineage-related models. SeedPrints instead yields a statistical decision at a prescribed significance level, providing a clearer operational answer to whether two models plausibly share a training lineage within a model family.

E MODEL CATALOG AND DECODING SETTINGS

Our evaluation covers 58 models in total. We document the decoding configuration used throughout all runs and enumerate the model catalog by family. For each derived model, we also indicate the post-training transformation (*finetune*, *adapter*, *merge*, *quantization*, *distillation*). “PT” and “IT” refer to *pre-trained* and *instruction-tuned* source models, respectively.

Parameter	Value
max_new_tokens	512
temperature	0.7
top_p	0.9
top_k	50
do_sample	true
max_input_length	512

E.1 MODEL LIST BY FAMILY

E.1.1 QWEN-2.5-7B

Variant	HuggingFace repo	Type
PT	Qwen/Qwen2.5-7B	—
IT	Qwen/Qwen2.5-7B-Instruct	—
Derived	Qwen/Qwen2.5-Math-7B	finetune
Derived	Qwen/Qwen2.5-Coder-7B-Instruct	finetune
Derived	WangCa/Qwen2.5-7B-Medicine	finetune
Derived	huihui-ai/Qwen2.5-7B-Instruct-abliterated-v2	finetune
Derived	Locutusque/StockQwen-2.5-7B	merge
Derived	bunnycore/QevaCoT-7B-Stock	merge
Derived	fangcaotank/task-10-Qwen-Qwen2.5-7B-Instruct	adapter
Derived	SeeFlock/task-12-Qwen-Qwen2.5-7B-Instruct	adapter
Derived	Qwen/Qwen2.5-7B-Instruct-GPTQ-Int4	quantization
Derived	Qwen/Qwen2.5-7B-Instruct-GPTQ-Int8	quantization
Derived	Lansechen/Qwen2.5-7B-Open-R1-Distill	distillation

E.1.2 QWEN2.5-14B

Variant	HuggingFace repo	Type
PT	Qwen/Qwen2.5-14B	—
IT	Qwen/Qwen2.5-14B-Instruct	—
Derived	Qwen/Qwen2.5-Coder-14B	finetune
Derived	oxyapi/oxy-1-small	finetune
Derived	v000000/Qwen2.5-14B-Gutenberg-Instruct-Slerpeno	merge
Derived	ToastyPigeon/qwen-story-test-qlora	adapter
Derived	Qwen/Qwen2.5-14B-Instruct-GPTQ-Int4	quantization
Derived	deepseek-ai/DeepSeek-R1-Distill-Qwen-14B	distillation

E.1.3 LLAMA-3.1-8B

Variant	HuggingFace repo	Type
PT	meta-llama/Llama-3.1-8B	—
IT	meta-llama/Llama-3.1-8B-Instruct	—
Derived	ValiantLabs/Llama3.1-8B-Fireplace2	finetune
Derived	RedHatAI/Llama-3.1-8B-tldr	finetune
Derived	proxectonos/Llama-3.1-Carballo	finetune
Derived	mlabonne/Meta-Llama-3.1-8B-Instruct-abliterated	finetune
Derived	gaverfraxz/Meta-Llama-3.1-8B-Instruct-HalfAbliterated-TIES	merge
Derived	XiaoJian9992024/Llama3.1-8B-ExtraMix	merge
Derived	LlamaFactoryAI/Llama-3.1-8B-Instruct-cv-job-description-matching	adapter
Derived	chchen/Llama-3.1-8B-Instruct-PsyCourse-fold7	adapter
Derived	iqbalamo93/Meta-Llama-3.1-8B-Instruct-GPTQ-Q_8	quantization
Derived	DaraV/LLaMA-3.1-8B-Instruct-INT4-GPTQ	quantization
Derived	asas-ai/Llama-3.1-8B-Instruct-Open-R1-Distill	distillation

E.1.4 MISTRAL-7B-v0.3

Variant	HuggingFace repo	Type
PT	mistralai/Mistral-7B-v0.3	—
IT	mistralai/Mistral-7B-Instruct-v0.3	—
Derived	KurmaAI/AQUA-7B	finetune
Derived	openfoodfacts/spellcheck-mistral-7b	finetune
Derived	grimjim/Mistral-7B-Instruct-demi-merge-v0.3-7B	merge
Derived	chaymaemrhrioui/mistral-Brain_Model_ACC_Trainer	adapter
Derived	RedHatAI/Mistral-7B-Instruct-v0.3-GPTQ-4bit	quantization
Derived	eganwo/mistral7b-distilled-from-deepseek-r1-qwen32b	distillation

E.1.5 GEMMA-2-2B

Variant	HuggingFace repo	Type
PT	google/gemma-2-2b	—
IT	google/gemma-2-2b-it	—
Derived	rinna/gemma-2-baku-2b	finetune
Derived	anakin87/gemma-2-2b-neogenesis-ita	finetune
Derived	vonjack/gemma2-2b-merged	merge
Derived	google-cloud-partnership/gemma-2-2b-it-lora-sql	adapter
Derived	qilowoq/gemma-2-2B-it-4Bit-GPTQ	quantization
Derived	Syed-Hasan-8503/Gemma-2-2b-it-distilled	distillation

E.1.6 LLAMA-2-7B

Variant	HuggingFace repo	Type
PT	meta-llama/Llama-2-7b-hf	—
IT	meta-llama/Llama-2-7b-chat-hf	—
Derived	allenai/tulu-2-7b	finetune
Derived	QIAIUNCC/EYE-Llama_qa	finetune
Derived	DevQuasar/coma-7B-v0.1	merge
Derived	Ammar-1/llama2-Better-Tune	adapter
Derived	TheBloke/Llama-2-7B-Chat-GPTQ	quantization
Derived	cygu/llama-2-7b-logit-watermark-distill-kgw-k1-gamma0.25-delta2	distillation

E.1.7 TINYLLAMA-1.1B

Variant	HuggingFace repo	Type
PT/IT	TinyLlama/TinyLlama-1.1B-Chat-v1.0	—
Derived	alexredna/TinyLlama-1.1B-Chat-v1.0-reasoning-v2	finetune
Derived	Edentns/DataVortexTL-1.1B-v0.1	finetune
Derived	appvoid/dot-v2.7	merge
Derived	barissglc/tinyllama-tarot-v1	adapter
Derived	TheBloke/TinyLlama-1.1B-Chat-v1.0-GPTQ	quantization
Derived	anudaw/distilled-code-llama	distillation

Notes. “Type” denotes the post-training transformation relative to the PT/IT source model: *finetune* includes supervised/preference optimization variants; *adapter* includes LoRA/QLoRA-style modules; *merge* includes model soups and TIES-style merges; *quantization* includes GPTQ/INTx variants; *distillation* includes student models distilled from reasoning-augmented teachers.

E.2 SAMPLE RESULTS

We present partial results from one LeafBench evaluation run in Table 15 (u -test with simulation-based null). Reported metrics are $1 - p$, used as similarity scores for compatibility with the benchmark. We adopt a significance level of 0.01 for p -values; empirically, LeafBench identifies an optimal decision threshold of 0.9920, further supporting the reliability of the method.

Crucially, individual p -values should not be interpreted as linear similarity measures, but their distribution across model pairs is informative: same-lineage models yield much lower p -values, while unrelated models remain near uniform.

1566
1567
1568
1569
1570
1571
1572
1573
1574
1575
1576
1577
1578
1579
1580
1581
1582
1583
1584
1585
1586
1587
1588
1589
1590
1591
1592
1593
1594
1595
1596
1597
1598
1599
1600
1601
1602
1603
1604
1605
1606
1607
1608
1609
1610
1611
1612
1613
1614
1615
1616
1617
1618
1619

Model	Qwen-2.5-7B	Qwen-2.5-7B-Instruct	Qwen2.5-14B	Qwen2.5-14B-Instruct
Qwen-2.5-7B	1.000	1.000	0.079	0.000
Qwen-2.5-7B-Instruct	1.000	1.000	0.813	0.536
Qwen2.5-7B-Math	1.000	0.994	0.047	0.987
Qwen2.5-7B-Coder	0.996	0.999	0.177	0.002
Qwen2.5-7B-Instruct-Medicine	1.000	1.000	0.576	0.369
Qwen2.5-7B-Instruct-Abiliterated	1.000	1.000	0.276	0.886
Qwen2.5-7B-Stock	1.000	1.000	0.939	0.225
QevaCoT-7B	1.000	1.000	0.199	0.119
Qwen2.5-7B-Instruct-Task-10	1.000	1.000	0.672	0.090
Qwen2.5-7B-Instruct-Task-12	1.000	1.000	0.155	0.987
Qwen2.5-7B-Instruct-Int4	1.000	1.000	0.827	0.961
Qwen2.5-7B-Instruct-Int8	1.000	1.000	0.914	0.739
Qwen2.5-7B-Open-R1-Distill	1.000	1.000	0.302	0.955
Qwen2.5-14B	0.324	0.487	1.000	1.000
Qwen2.5-14B-Instruct	0.403	0.702	1.000	1.000
Qwen2.5-Coder-14B	0.492	0.660	1.000	1.000
oxy-1-small	0.492	0.104	1.000	1.000
Qwen2.5-14B-Gutenberg-Instruct-Slerpeno	0.856	0.931	1.000	1.000
Qwen-story-test-qlora	0.112	0.898	1.000	1.000
Qwen2.5-14B-Instruct-GPTQ-Int4	0.583	0.092	1.000	1.000
DeepSeek-R1-Distill-Qwen-14B	0.204	0.000	1.000	1.000
Llama-3.1-8B	0.638	0.714	0.156	0.000
Llama-3.1-8B-Instruct	0.643	0.118	0.128	0.510
Llama-3.1-8B-Fireplace2	0.715	0.532	0.065	0.354
Llama-3.1-8B-TLDR	0.925	0.811	0.331	0.014
Llama-3.1-8B-Carballo	0.001	0.605	0.003	0.946
Llama-3.1-8B-Instruct-Abliterated	0.061	0.005	0.121	0.422
Llama-3.1-8B-Instruct-HalfAbliterated-TIES	0.356	0.509	0.796	0.000
Llama-3.1-8B-ExtraMix	0.640	0.160	0.721	0.875
Llama-3.1-8B-Instruct-cv-job-description-matching	0.684	0.292	0.004	0.001
Llama-3.1-8B-Instruct-PsyCourse-fold7	0.682	0.057	0.063	0.191
Llama-3.1-8B-Instruct-8bit	0.798	0.214	0.154	0.385
Llama-3.1-8B-Instruct-4bit	0.513	0.023	0.005	0.933
Llama-3.1-8B-Instruct-Open-R1-Distill	0.800	0.101	0.772	0.024
Mistral-7B-v0.3	0.224	0.352	0.506	0.044
Mistral-7B-v0.3-Instruct	0.000	0.000	0.984	0.104
AQUA-7B	0.014	0.613	0.703	0.569
Mistral-7B-v0.3-Spellcheck	0.005	0.481	0.872	0.105
Mistral-7B-v0.3-Instruct-demi-merge	0.904	0.774	0.058	0.522
Mistral-7B-v0.3-Brain	0.340	0.983	0.481	0.016
Mistral-7B-v0.3-Instruct-GPTQ-4bit	0.002	0.281	0.089	0.784
Mistral-7B-distilled-from-deepseek-r1-qwen32b	0.009	0.001	0.511	0.015
Gemma-2-2b	0.747	0.869	0.689	0.000
Gemma-2-2b-it	0.944	0.039	0.626	0.060
Gemma-2-baku-2b	0.616	0.231	0.176	0.464
Gemma-2-2b-neogenesis-ita	0.002	0.536	0.079	0.980
Gemma-2-2b-merged	0.951	0.012	0.651	0.075
Gemma-2-2b-it-lora-sql	0.236	0.408	0.000	0.000
Gemma-2-2B-it-4Bit-GPTQ	0.945	0.159	0.000	0.008
Gemma-2-2b-it-distilled	0.482	0.573	0.271	0.276
Llama-2-7b	0.443	0.470	0.148	0.276
Llama-2-7b-chat	0.528	0.868	0.424	0.766
tulu-2-7b	0.125	0.272	0.170	0.888
EYE-Llama_qa	0.657	0.244	0.674	0.968
coma-7B-v0.1	0.953	0.004	0.573	0.731
llama2-Better-Tune	0.136	0.699	0.469	0.860
Llama-2-7B-Chat-GPTQ	0.594	0.603	0.582	0.424
llama-2-7b-logit-watermark-distill-kgw-k1-gamma0.25-delta2	0.349	0.224	0.366	0.976

Table 15: **Sample results.** Pairwise SeedPrints scores between target models (rows) and base models (columns). Cells with $1 - p \geq 0.99$ are highlighted in green as same-lineage pairs.