
UNIFYING FORMAL EXPLANATIONS: A COMPLEXITY-THEORETIC PERSPECTIVE

Anonymous authors
Anonymous submission

ABSTRACT

Previous work has explored the computational complexity of deriving two fundamental types of explanations for ML model predictions: (i) *sufficient reasons*, which are subsets of input features that, when fixed, determine a prediction, and (ii) *contrastive reasons*, which are subsets of input features that, when modified, alter a prediction. These explanations have been studied along two axes that we place side by side throughout this paper: non-probabilistic versus probabilistic, and local versus global. In this study, we introduce a unified framework that folds all four combinations into a single minimization task over one probabilistic value function, making every prior setting directly comparable within one objective. We then pin the complexity of this task to exactly three properties of the value function: (i) *monotonicity*, (ii) *submodularity*, and (iii) *supermodularity* — three fundamental properties in *combinatorial optimization*. The central finding is a clean, consistent split along the local–global axis: the *local* value functions satisfy *none* of these three properties, whereas their *global* counterparts satisfy *all* of them. This all-or-nothing gap is not incremental — it flips the tractability of the problem outright. On the global side it yields a series of new polynomial-time algorithms for computing explanations with provable guarantees, holding uniformly across models that span the interpretability spectrum — decision trees, neural networks, and tree ensembles. On the matching local side, the identical problems are NP-hard even in drastically simplified regimes. The comparison is therefore a head-to-head one: same explanation, same model, same guarantee — polynomial-time globally, NP-hard locally.

1 INTRODUCTION

Despite substantial progress in methods for explaining ML model decisions, the literature has consistently and unfortunately found that many desirable explanation types with provable guarantees are computationally hard to obtain (Barceló et al., 2020; Van den Broeck et al., 2022), and the pattern is a monotone one: the difficulty grows precisely as models become more complex or more non-linear (Barceló et al., 2020; Adolphi et al., 2025; 2024). As a result, the *computational complexity* of obtaining explanations has become a central theoretical focus, with a large body of recent work drawing an increasingly detailed map of which explanation types fall on the tractable (polynomial-time) side of the line and which fall on the intractable (NP-hard and beyond) side, for each model class (Barceló et al., 2020; Wäldchen et al., 2021; Arenas et al., 2022; 2023; Marzouk & De La Higuera, 2024; Ordyniak et al., 2023; Laber, 2024; Bhattacharjee & Luxburg, 2024; Blanc et al., 2021; 2022).

From sufficient to contrastive reasons. Among studies on the computational complexity of generating explanations, two fundamental types of explanations for ML models were extensively examined: (i) *sufficient reasons* and (ii) *contrastive reasons* (Barceló et al., 2020; Arenas et al., 2022; Audemard et al., 2022a; Arenas et al., 2021; Barceló et al., 2025; Marques-Silva & Ignatiev, 2022; Ignatiev et al., 2020b; Darwiche & Hirth, 2020). A *sufficient reason* is a subset of input features S such that when these features are fixed to specific values, the model’s prediction remains unchanged, regardless of the values assigned to the complementary set $\bar{S}$. A *contrastive reason* is a subset of input features S such that modifying these features leads to a change in the model’s prediction.

Unlike additive attribution methods, which allocate importance scores across features but are often hard for humans to interpret (Kumar et al., 2020) or lack actionability (Bilodeau et al., 2024), sufficient and contrastive reasons provide *discrete, condition-like explanations* that directly answer “what is enough to justify this prediction?” or “what must change to flip it?”. This concreteness has translated into measurable advantages: it has given them a central role across many classic XAI methods (Ribeiro et al., 2018; Carter et al., 2019; Ignatiev et al., 2019; Dhurandhar et al., 2018), and, in head-to-head human studies, these explanations improved human prediction accuracy more than additive-model explanations did (Ribeiro et al., 2018; Yin & Neubig, 2022; Dhurandhar et al., 2018). They have likewise proven useful across a range of downstream tasks — bias detection (Balkir et al., 2022; Carter et al., 2021; La Malfa et al., 2021; Muthukumar et al., 2018), model debugging (Jacovi et al., 2021), and anomaly detection (Davidson et al., 2025). A well-established principle further holds that *smaller* sufficient and contrastive explanations enhance interpretability, making *minimality* a central guarantee of interest (Ribeiro et al., 2018; Lopardo et al., 2023; Carter et al., 2019; Barceló et al., 2020; Arenas et al., 2022; Blanc et al., 2021; Wäldchen et al., 2021).

Barceló et al. (2020) conducted one of the earliest studies on the complexity of deriving sufficient and contrastive reasons, and their results already reveal how sharply the difficulty tracks model expressiveness. For the minimal-sized *sufficient* reason, the cost climbs one full level of the polynomial hierarchy as the model grows more expressive: NP-hard for decision trees, and Σ_2^P -hard for neural networks. For the minimal-sized *contrastive* reason the contrast is even starker — polynomial-time for decision trees but NP-hard for neural networks, i.e., a jump from tractable to intractable driven purely by the model class. Subsequent work showed the same hardness carries over to tree ensembles (Izza & Marques-Silva, 2021; Ordyniak et al., 2024; Audemard et al., 2022b).

From non-probabilistic to probabilistic explanations. A common criticism of the classic definition of sufficient and contrastive reasons is their rigidity: because they must hold in an absolute sense over entire domains, they tend to inflate the explanation size and dilute its informativeness (Ignatiev et al., 2019; 2020a; Arenas et al., 2022; Wäldchen et al., 2021). To address this, the literature has moved to a strictly more general *probabilistic* formulation (Wäldchen et al., 2021; Arenas et al., 2022; Blanc et al., 2021; Izza et al., 2023; Xue et al., 2023), which asks only that a subset influence the prediction with probability at least a threshold δ — recovering the original absolute definition exactly at $\delta = 1$ and interpolating continuously below it.

Wäldchen et al. (2021) were the first to study the complexity of probabilistic sufficient reasons, showing that for CNF classifiers the problem is NP^{PP} -hard — a class widely regarded as sitting above NP — with the same hardness carrying over to tree ensembles and neural networks (Barceló et al., 2020; Ordyniak et al., 2023). The mechanism behind this jump is a single structural fact: the *lack of monotonicity* in the probability function (Arenas et al., 2022; Izza et al., 2023), which is enough to make even the far weaker *subset-minimal* explanations computationally hard. The cleanest illustration of the gap is on decision trees: (Arenas et al., 2022) show that a subset-minimal probabilistic sufficient reason is NP-hard to find, whereas the very same task in the non-probabilistic case is solvable in polynomial time (Huang et al., 2021; Barceló et al., 2020) — so moving from $\delta = 1$ to a general threshold alone turns a tractable problem intractable.

From local to global explanations. In a more recent study, Bassan et al. (2024) extend the complexity analysis from the *local* (non-probabilistic) setting — where explanations are tied to individual predictions — to the *global* (non-probabilistic) setting, which seeks sufficient or contrastive reasons over entire domains. In the non-probabilistic regime, however, the criteria are extremely strict (Ignatiev et al., 2019; Barceló et al., 2020; Arenas et al., 2021; Darwiche & Hirth, 2020), and the strictness worsens rather than eases in the global setting: any feature excluded from a global sufficient reason is deemed strictly redundant (Bassan et al., 2024), which in practice drives the explanation to cover nearly all n input features — the opposite of the small, informative subset one wants.

OUR CONTRIBUTIONS

1. We unify previously separate explanation-computation problems — sufficient and contrastive, probabilistic and non-probabilistic, local and global — into a *single* minimization task over one value function. This collapses four axes into one objective and lets us attribute all complexity differences among them to exactly three properties of that value function: (a) *monotonicity*, (b) *supermodularity*, and (c) *submodularity*.

2. Along the local–global axis, these three properties partition cleanly, with no middle ground. The *local* value functions, for both sufficient and contrastive reasons, satisfy *none* of them — they are non-monotonic and neither sub- nor supermodular. Their *global* counterparts satisfy *all* of the relevant ones: both are monotonic non-decreasing, the global sufficient function is supermodular, and the global contrastive function is submodular. The same property either fully holds globally or fully fails locally — there is no partial or approximate version in between.
3. We convert this structural gap directly into a complexity gap, and we do so uniformly across three model types spanning the interpretability spectrum: (i) neural networks, (ii) decision trees, and (iii) tree ensembles. The comparison is held fixed on both sides — same explanation, same model, same threshold δ — and only the local/global scope changes. For decision trees, a subset-minimal probabilistic explanation is NP-hard to compute locally (Arenas et al., 2022) yet polynomial-time globally. The same polynomial-time-vs-NP-hard split then extends to *any* black-box model, including neural networks and tree ensembles, once an empirical distribution is used: the global subset-minimal sufficient/contrastive explanation is computable in polynomial time while the matching local version stays NP-hard.
4. Exploiting submodularity and supermodularity — which, again, hold only globally — we push the gap further, from the subset-minimal to the harder cardinally minimal problem. Globally, and even for complex models such as neural networks and tree ensembles under a fixed empirical distribution, we obtain provable *constant-factor* approximations. Locally, the same problem admits *no bounded approximation whatsoever*, even in the most simplified regime (a single data point). The two sides thus sit at opposite extremes: a constant factor globally versus unbounded inapproximability locally.

Owing to space constraints, we provide an overview of our main theorems and corollaries in the main text, with full proofs deferred to the appendix.

2 PRELIMINARIES

Setting. We consider an input space of dimension $n \in \mathbb{N}$, with input vectors $\mathbf{x} := (\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n)$. Each coordinate $\mathbf{x}_i$ may take values from its corresponding domain $\mathcal{X}_i$, which can be either discrete or continuous. The full input feature space is therefore $\mathbb{F} := \mathcal{X}_1 \times \mathcal{X}_2 \times \dots \times \mathcal{X}_n$. We consider classification models $f : \mathbb{F} \rightarrow [c]$ where $c \in \mathbb{N}$ is the number of classes. Moreover, we consider a *generic* distribution $\mathcal{D} : \mathbb{F} \rightarrow [0, 1]$ over the input space. In many settings, however, accurately approximating $\mathcal{D}$ is computationally infeasible. A natural alternative is to instead work with a fixed empirical dataset $\mathbf{D} := \{\mathbf{z}^1, \mathbf{z}^2, \dots, \mathbf{z}^{|\mathbf{D}|}\}$, which serves as a practical proxy for the “true” distribution, for instance, by taking $\mathbf{D}$ as a sampled subset from the available training data.

The explanations that we study are either local or global. In the local case, the explanations target a specific prediction $\mathbf{x} \in \mathbb{F}$, providing a form of reasoning for why the model predicted $f(\mathbf{x})$ for that instance. In the global case, explanations aim to reflect the general reasoning behind the behavior of f across a wider region of the input space, independent of any specific $\mathbf{x}$, and to characterize its overall decision-making logic.

Models. Many of our results are model-agnostic (e.g., the inherent properties of the value functions), but we also make the comparison concrete with model-specific complexity results for widely used ML models. To ensure the comparison spans the full interpretability spectrum rather than a single regime, we deliberately cover both ends and the middle: (i) decision trees (the “interpretable” end), (ii) neural networks — all architectures at least as expressive as feed-forward ReLU networks (the “black-box” end), and (iii) tree ensembles, covering both majority-voting ensembles (e.g., random forests) and weighted-voting ensembles (e.g., XGBoost). Reporting the same local-vs-global comparison across all three lets us show the separation is a property of the setting, not of any one model class. Formal definitions are provided in the Appendix.

Distributions. We emphasize that in probabilistic explanation settings, the input distribution $\mathcal{D}$ is itself a decisive factor: holding the model and explanation type fixed, the complexity can swing from tractable to intractable purely by changing $\mathcal{D}$. Accordingly, we compare three distribution types throughout: (i) *general distributions*, which make no specific assumptions over $\mathcal{D}$ and thus encompass all possible distributions. We use these distributions mainly in proofs of properties that

hold universally; (ii) *empirical distributions*, which include all distributions derived from the finite dataset $\mathbf{D}$ — an approach commonly employed in XAI (Lundberg & Lee, 2017; Van den Broeck et al., 2022); and (iii) *independent distributions*, which assume that features in $\mathcal{D}$ are mutually independent — another widely adopted assumption in XAI literature (Arenas et al., 2022; 2023; Lundberg & Lee, 2017; Ribeiro et al., 2018). We note that empirical distributions do not necessarily imply feature independence; rather, they can represent complex dependencies extracted from finite datasets (Van den Broeck et al., 2022). The complete formal definitions of these distributions are provided in the Appendix.

3 FORMS OF REASONS

In this section, we introduce the explanation types studied in this work, starting with the strict non-probabilistic definitions of (local/global) sufficient and contrastive reasons, and then extending to their more flexible and generalizable probabilistic counterparts.

3.1 Non-Probabilistic SUFFICIENT AND CONTRASTIVE REASONS

Sufficient reasons. In the context of *feature selection*, users often select the top k features that contribute to a model’s decision. We examine the well-established *sufficiency* criterion for this selection, which aligns with commonly used explainability methods (Ribeiro et al., 2018; Carter et al., 2019; Ignatiev et al., 2019; Dasgupta et al., 2022). This feature selection can be carried out either locally — focusing on a single prediction — or globally — across the entire input domain. Following standard conventions, we define a *local sufficient reason* as a subset of features $S \subseteq \{1, \dots, n\}$ such that when the features in S are fixed to their corresponding values in $\mathbf{x}$, the model’s prediction remains $f(\mathbf{x})$ regardless of the values assigned to the remaining features $\bar{S}$. Formally, S is a local sufficient reason for $\langle f, \mathbf{x} \rangle$ iff the following condition holds:

$$\forall \mathbf{z} \in \mathbb{F}, \quad f(\mathbf{x}_S; \mathbf{z}_{\bar{S}}) = f(\mathbf{x}). \quad (1)$$

Here, $(\mathbf{x}_S; \mathbf{z}_{\bar{S}})$ denotes a vector where features in S take their values from $\mathbf{x}$, and those in $\bar{S}$ from $\mathbf{z}$. Local sufficient reasons are closely related to *semi-factual* explanations (Kenny & Keane, 2021; Alfano et al., 2025), which search for alternative inputs $\mathbf{z}'$ that keep the prediction unchanged (i.e., $f(\mathbf{z}') = f(\mathbf{x})$) while being as “close” as possible to the original point. When we instantiate $\mathbf{z}'$ as $(\mathbf{x}_S; \mathbf{z}_{\bar{S}})$ and measure proximity via Hamming distance, the two notions coincide.

A *global sufficient reason* (Bassan et al., 2024) is a subset of input features $S \subseteq \{1, \dots, n\}$ that serves as a local sufficient reason for *every* possible input $\mathbf{x} \in \mathbb{F}$:

$$\forall \mathbf{x}, \mathbf{z} \in \mathbb{F}, \quad f(\mathbf{x}_S; \mathbf{z}_{\bar{S}}) = f(\mathbf{x}). \quad (2)$$

Contrastive reasons. Another prevalent approach to providing explanations is by pinpointing subsets of input features that *modify* a prediction (Dhurandhar et al., 2018; Mothilal et al., 2020; Guidotti, 2024). This type of explanation aims to determine the minimal changes necessary to alter a prediction. Formally, a subset $S \subseteq \{1, \dots, n\}$ is defined as a *local contrastive reason* concerning $\langle f, \mathbf{x} \rangle$ iff:

$$\exists \mathbf{z} \in \mathbb{F}, \quad f(\mathbf{z}_S; \mathbf{x}_{\bar{S}}) \neq f(\mathbf{x}). \quad (3)$$

Contrastive reasons are also closely connected to *counterfactual* explanations (Guidotti, 2024; Mothilal et al., 2020), which seek a nearby assignment $\mathbf{x}'$ for which the prediction flips, i.e., $f(\mathbf{x}') \neq f(\mathbf{x})$. When the distance metric is taken to be the Hamming weight, the two notions coincide.

Similarly to sufficient reasons, one can determine a subset of input features that, when altered, changes the prediction for all inputs within the domain of interest. This form of explanation is also closely connected to approaches for identifying bias (Arenas et al., 2021; Bassan et al., 2024; Darwiche & Hirth, 2020), as well as to *group* counterfactual explanation methods that search for counterfactuals over multiple data instances (Carrizosa et al., 2024; Warren et al., 2024). Formally, we define a subset S as a *global contrastive reason* with respect to f iff:

$$\forall \mathbf{x} \in \mathbb{F}, \quad \exists \mathbf{z} \in \mathbb{F}, \quad f(\mathbf{z}_S; \mathbf{x}_{\bar{S}}) \neq f(\mathbf{x}). \quad (4)$$

3.2 Probabilistic SUFFICIENT AND CONTRASTIVE REASONS

As discussed in the introduction, non-probabilistic sufficient and contrastive reasons are often criticized for imposing significantly overly strict conditions, motivating a shift to *probabilistic* definitions (Arenas et al., 2022; Ribeiro et al., 2018; Izza et al., 2023; Wäldchen et al., 2021; Bounia & Koriche, 2023; Subercaseaux et al., 2025; Wang et al., 2021a; Blanc et al., 2021), which generalize these requirements by demanding that the guarantees hold with probability at least $\delta \in [0, 1]$. The special case $\delta = 1$ recovers the original non-probabilistic definitions.

Probabilistic sufficient reasons. We define $S \subseteq \{1, \dots, n\}$ as a *local δ -sufficient reason* with respect to $\langle f, \mathbf{x} \rangle$ if, when the features in S are fixed to their corresponding values in $\mathbf{x}$ and the remaining features are sampled from a distribution $\mathcal{D}$, the classification remains unchanged with probability at least δ . In other words:

$$\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S) \geq \delta. \quad (5)$$

where $\mathbf{z}_S = \mathbf{x}_S$ denotes that the features in S of vector $\mathbf{z}$ are fixed to their corresponding values in $\mathbf{x}$. We adopt the standard notion of global explanations — by averaging over all inputs in the global domain — and define a subset $S \subseteq \{1, \dots, n\}$ as a *global δ -sufficient reason* with respect to $\langle f \rangle$ if, when taking the expectation of the local sufficiency probability over samples drawn from the distribution $\mathcal{D}$, the expectation remains with value at least δ . In other words:

$$\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)] \geq \delta. \quad (6)$$

Probabilistic contrastive reasons. Similar to the non-probabilistic case, we define a *local δ -contrastive reason* as a subset of input features that changes a prediction with some probability. Here, unlike sufficient reasons, we set the features of the *complementary* set $\bar{S}$ to their respective values in $\mathbf{x}$, and when allowing the features in S to vary, we require the prediction to differ from the original prediction $f(\mathbf{x})$ with a probability exceeding δ . Formally:

$$\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) \neq f(\mathbf{x}) \mid \mathbf{z}_{\bar{S}} = \mathbf{x}_{\bar{S}}) \geq \delta. \quad (7)$$

For the global setting, we define a subset S to be a *global δ -contrastive reason*, analogous to global sufficient reasons, by computing the expectation over all local contrastive reasons sampled from the distribution $\mathcal{D}$, and requiring that this expected value exceeds δ :

$$\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) \neq f(\mathbf{x}) \mid \mathbf{z}_{\bar{S}} = \mathbf{x}_{\bar{S}})] \geq \delta. \quad (8)$$

4 FORMS OF MINIMALITY

As discussed in the introduction, across all the explanation forms discussed so far — whether sufficient or contrastive, local or global — a common assumption in the literature is that explanations of *smaller size* are more meaningful, thereby making their *minimality* a particularly important provable guarantee. In this study, we explore two central notions of minimality across all our explanation types:

Definition 1. Assuming a subset $S \subseteq \{1, \dots, n\}$ is an explanation, then:

1. S is a *cardinally-minimal explanation* (Barceló et al., 2020; Bassan et al., 2024) iff S has the *smallest explanation cardinality* $|S|$ (i.e., there is no explanation S' such that $|S'| < |S|$).
2. $S \subseteq \{1, \dots, n\}$ is a *subset-minimal explanation* (Arenas et al., 2022; Ignatiev et al., 2019) iff S is an explanation, and any $S' \subsetneq S$ is not an explanation.

The two notions are ordered by strength: cardinal-minimality strictly implies subset-minimality — every cardinally minimal S is subset-minimal, but the converse fails (Appendix B.3 gives an explicit example). This ordering matters for our results, since it means a hardness result for the weaker subset-minimal problem is the stronger statement, while a tractability result for the harder cardinally minimal problem is the stronger one — and we obtain both on the appropriate sides. We use “subset” and “cardinally” minimal rather than local/global minima to avoid clashing with local vs. global explanations (input- vs. domain-level reasoning). Both notions apply to *every* explanation form we analyze: a cardinally minimal local probabilistic δ -sufficient reason is the one with the smallest $|S|$, and a subset-minimal one is any S from which no proper subset $S' \subsetneq S$ still qualifies.

5 A UNIFIED COMBINATORIAL OPTIMIZATION TASK

A key unifying observation is that all four families of problems — local/global $\times$ sufficient/contrastive, probabilistic or not — reduce to one and the same template: find a minimal-size subset S with $v(S) \geq \delta$ (with $\delta := 1$ recovering the non-probabilistic case). Casting them in this shared form is what makes them directly comparable, since any complexity difference between two settings must then trace back to a difference in the single value function v . We now fix notation for these value functions: let v_{suff}^{ℓ} denote the local sufficiency probability from equation 5, i.e., $\Pr_{\mathbf{z} \sim \mathcal{D}_p}(f(\mathbf{x}) = f(\mathbf{z}) \mid \mathbf{x}_S = \mathbf{z}_S)$, and define the global variant as v_{suff}^g (equation 6). Similarly, let v_{con}^{ℓ} and v_{con}^g denote the local and global contrastive probabilities (equations 7 and 73, respectively). Using this notation, we now formally define the unified task of finding a cardinally minimal δ -local/global sufficient/contrastive reason:

Cardinally Minimal δ -Explanation:

Input: Model f , a distribution $\mathcal{D}$, (possibly, an input $\mathbf{x}$), a *general* value function $v : 2^n \rightarrow [0, 1]$ (defined using f , $\mathcal{D}$, and possibly $\mathbf{x}$), and some $\delta \in [0, 1]$.

Output: A subset $S \subseteq [n]$ such that $v(S) \geq \delta$ and $|S|$ is minimal.

Similarly, for the relaxed condition where the goal is to obtain a subset-minimal rather than a cardinally-minimal local or global sufficient/contrastive reason, we define the following relaxed optimization objective:

Subset Minimal δ -Explanation:

Input: Model f , a distribution $\mathcal{D}$, (possibly, an input $\mathbf{x}$), a *general* value function $v : 2^n \rightarrow [0, 1]$ (defined using f , $\mathcal{D}$, and possibly $\mathbf{x}$), and some $\delta \in [0, 1]$.

Output: A subset $S \subseteq [n]$ such that $v(S) \geq \delta$ and for any $S' \subseteq S$ it holds that $v(S') < \delta$.

Properties of v that affect the complexity. We now outline several key properties of the value function, which we later show play a crucial role in determining the complexity of generating explanations. The first property is *non-decreasing monotonicity*, which ensures that the marginal contribution $v(S \cup \{i\}) - v(S)$ is consistently non-negative. Formally:

Definition 2. We say that a value function v maintains non-decreasing monotonicity if for any $S \subseteq \{1, \dots, n\}$ and any $i \in \{1, \dots, n\}$ it holds that: $v(S \cup \{i\}) \geq v(S)$.

The other key properties of *supermodularity* and *submodularity* pertain to the behavior of the marginal contribution $v(S \cup \{i\}) - v(S)$. Specifically, in the supermodular case, this contribution forms a monotone non-decreasing function. In contrast, under the dual definition of the *submodular* case, the marginal contribution $v(S \cup \{i\}) - v(S)$ is a monotone non-increasing function. Formally:

Definition 3. Let there be some value function v , some $S \subseteq S' \subseteq \{1, \dots, n\}$, and $i \notin S'$. Then:

1. v maintains supermodularity iff it holds that: $v(S \cup \{i\}) - v(S) \leq v(S' \cup \{i\}) - v(S')$.
2. v maintains submodularity iff it holds that: $v(S \cup \{i\}) - v(S) \geq v(S' \cup \{i\}) - v(S')$.

6 UNRAVELING NEW PROPERTIES OF THE GLOBAL VALUE FUNCTIONS

Prior work shows that in the local setting of non-probabilistic explanations, subset-minimal sufficient or contrastive reasons are computable, and the reason is precisely monotonicity — but this monotonicity is available only at the single point $\delta = 1$. That single point is a narrow and inflexible regime. The natural move is to relax to probabilistic guarantees for arbitrary δ , yet prior results show this relaxation immediately destroys monotonicity, and with it the tractability: computing the explanations becomes computationally harder as soon as $\delta < 1$ (Arenas et al., 2022; Kozachinskiy, 2023; Izza et al., 2023; Subercaseaux et al., 2025; Izza et al., 2024b).

Pushing to the still more restrictive *global* and *non-probabilistic* setting, Bassan et al. (2024) established a stringent *uniqueness* property: there is exactly *one* subset-minimal explanation. This count of one is a direct symptom of how rigid the $\delta = 1$ regime is, especially globally, where the conditions must hold across *all* inputs simultaneously; indeed Bassan et al. (2024) show that this single subset

coincides with the set of all non-strictly-redundant features, which may be all n of them. Relaxing to the general global *probabilistic* setting overturns this count entirely: for any δ the uniqueness property fails, and the number of subset-minimal explanations jumps from one all the way up to *exponentially* many.

Proposition 1. *While the non-probabilistic case ($\delta = 1$) admits a unique subset-minimal global sufficient reason (Bassan et al., 2024), in the general probabilistic setting (for arbitrary δ), there exist functions f that have $\Theta(\frac{2^n}{\sqrt{n}})$ subset-minimal global sufficient reasons.*

Here the comparison turns favorable. Although uniqueness survives only at $\delta = 1$, the far more useful *monotonicity* property does the opposite: it holds for the global value function at *every* δ , and for *both* the sufficient (v_{suff}^g) and contrastive (v_{con}^g) functions. Placed next to the local setting, the split is exact — the local functions v_{suff}^l and v_{con}^l satisfy monotonicity for *no* δ , while their global counterparts satisfy it for *all* δ :

Proposition 2. *While the local probabilistic setting (for any δ) lacks monotonicity — i.e., the value functions v_{con}^l and v_{suff}^l are non-monotonic (Arenas et al., 2022; Izza et al., 2023; Subercaseaux et al., 2025; Izza et al., 2024b) — in the global probabilistic setting (also for any δ), both value functions v_{con}^g and v_{suff}^g are monotonic non-decreasing.*

The monotonicity gap is only the first of the structural properties that separate the two settings; the global side accumulates further advantages the local side never gains. Under the common feature-independence assumption, the global sufficient value function v_{suff}^g is supermodular. The local function v_{suff}^l fails this even under a strictly stronger hypothesis — a uniform, and therefore independent, input distribution. So the global property holds under a weaker assumption than the one under which the local property already breaks, widening rather than narrowing the gap:

Proposition 3. *While the local probabilistic sufficient setting (for any δ) lacks supermodularity — even when $\mathcal{D}$ is uniform, i.e., the value function v_{suff}^l is not supermodular — in the global probabilistic setting (also for any δ), when $\mathcal{D}$ exhibits feature independence, the value function v_{suff}^g is supermodular.*

The second family — global probabilistic *contrastive* reasons — mirrors this pattern with the dual property: the value function is not supermodular but *submodular*. The same favorable comparison holds once more, and again in its strongest form: globally the function has a usable structure (submodularity), whereas locally it has *neither* submodularity nor supermodularity, i.e., no exploitable structure at all.

Proposition 4. *While the local probabilistic contrastive setting (for any δ) lacks submodularity — even when $\mathcal{D}$ is uniform, i.e., the value function v_{con}^l is not submodular — in the global probabilistic setting (also for any δ), when $\mathcal{D}$ exhibits feature independence, the value function v_{con}^g is submodular.*

7 COMPUTATIONAL COMPLEXITY RESULTS

7.1 SUBSET MINIMAL EXPLANATIONS

In this section, we examine the complexity of obtaining *subset minimal* explanations (local/global, sufficient/contrastive) across the different model types analyzed. The pivot is a single property — *monotonicity* — and the previous section already fixed which side of the local/global divide has it: both global functions v_{con}^g and v_{suff}^g do, both local functions v_{con}^l and v_{suff}^l do not. That one bit of structure decides everything downstream: the *same* greedy algorithm provably converges to a subset-minimal explanation in the global setting and provably can fail in the local one. The resulting comparison is clean and one-directional — every subset-minimal *global* form we study is polynomial-time solvable, while its matching *local* form is hard. We begin with the following generalized greedy algorithm:

Algorithm 1 aims to obtain a subset-minimal δ -explanation. We begin the algorithm with the subset S initialized as the entire input space $\{1, \dots, n\}$. Iteratively, we check whether the minimal value that $v(S \setminus \{i\})$ can attain exceeds δ . In each iteration, we remove a feature j from S that minimizes the decrease in the value function, selecting the feature j that maximizes $v(S \setminus \{j\})$. Once this iterative process concludes, we return S .

Algorithm 1 Subset Minimal Explanation Search

Input Value function v , and some $\delta \in [0, 1]$

```
1:  $S \leftarrow \{1, \dots, n\}$ 
2: while  $\min_{i \in S} v(S \setminus \{i\}) \geq \delta$  do
3:    $j \leftarrow \operatorname{argmax}_{i \in S} v(S \setminus \{i\})$ 
4:    $S \leftarrow S \setminus \{j\}$ 
5: end while
6: return  $S$   $\triangleright S$  is a (subset minimal?)  $\delta$ -explanation
```

The key determinant of whether Algorithm 1 yields a subset-minimal explanation is the *monotonicity* property of the value function v . The algorithm concludes with a phase in which removing any individual feature from S results in $v(S \setminus \{i\})$ being smaller than δ . However, monotonicity ensures that this holds for any $v(S \setminus S')$, providing a significantly stronger guarantee. Given the monotonicity property of the value functions established in the previous sections, we derive the following proposition:

Proposition 5. *Computing Algorithm 1 with the local value functions v_{con}^{ℓ} and v_{suff}^{ℓ} does not always converge to a subset minimal δ -sufficient/contrastive reason. However, computing it with the global value functions v_{con}^g or v_{suff}^g necessarily produces subset minimal δ -sufficient/contrastive reasons.*

Building on this result, we proceed to establish new complexity findings for deriving various forms of subset-minimal explanations within our framework, considering the different analyzed model types.

Decision trees. We start with the most favorable-to-the-local-side scenario we can construct: an ostensibly “interpretable” decision tree f paired with an independent distribution $\mathcal{D}$. Even under these best-case conditions for tractability, the local and global settings land on opposite sides of the tractability line. Arenas et al. (2022) showed the negative half: unless $\text{PTIME}=\text{NP}$, no polynomial-time algorithm computes a subset-minimal *local* δ -sufficient reason for decision trees under independent distributions — and this already fails at the uniform distribution. We supply the positive half for the *identical* problem in the global setting: a subset-minimal *global* δ -sufficient reason for decision trees is computable in polynomial time. Same model, same distribution, same explanation — polynomial-time globally, NP-hard locally:

Theorem 1. *If f is a decision tree and the probability term v_{suff}^g can be computed in polynomial time given the distribution $\mathcal{D}$ (which holds for independent distributions, among others), then obtaining a subset-minimal global δ -sufficient reason can be obtained in polynomial time. However, unless $\text{PTIME}=\text{NP}$, no polynomial-time algorithm exists for computing a local δ -sufficient reason for decision trees even under independent distributions.*

Extension to other tractable models. The same global tractability guarantee is not confined to decision trees: it carries over to *orthogonal DNFs* (Crama & Hammer, 2011) — Disjunctive Normal Form formulas whose terms are pairwise mutually exclusive. This extension matters because orthogonal DNFs *strictly generalize* decision trees while keeping the structural properties that drive the guarantee, so the polynomial-time global result covers a broader model class than trees alone. The full argument is given in Appendix N.

Neural networks, tree ensembles, and other complex models. The global tractability result is not tied to simple models — it scales all the way up the interpretability spectrum. With $\mathcal{D}$ an empirical distribution and under only the mild assumption that f admits polynomial-time inference, computing a subset-minimal global δ -sufficient or contrastive reason stays in polynomial time. Crucially, this holds for *any* such model, so the same efficient global guarantee that we obtained for decision trees now covers neural networks and tree ensembles as well.

Proposition 6. *For any model f , and empirical distribution $\mathcal{D}$ — computing a subset-minimal global δ -sufficient or δ -contrastive reason for f can be done in polynomial time.*

Because this positive result holds for *any* model under an empirical distribution, it lets us line the two settings up head-to-head on the most expressive models we consider. On the local side the picture is the reverse: for these same complex models, subset-minimal *local* explanations stay intractable even after restricting to empirical distributions. Establishing this for both neural networks and tree

ensembles pins down a strict polynomial-time-vs-NP-hard separation between the global and local settings:

Theorem 2. *Assuming f is a neural network or a tree ensemble, and $\mathcal{D}$ is an empirical distribution — there exist polynomial-time algorithms for obtaining subset minimal global δ -sufficient and contrastive reasons. However, unless $PTIME=NP$, there is no polynomial time algorithm for computing a subset minimal local δ -sufficient reason or a subset minimal local δ -contrastive reason.*

7.2 APPROXIMATE CARDINALLY MINIMAL EXPLANATIONS

We now move from subset-minimal reasons to the strictly harder *cardinally minimal* δ -sufficient/contrastive reason, and the local–global gap not only survives the jump in difficulty but widens into an approximability gap. Globally, monotonicity together with supermodularity/submodularity buys *provable approximation guarantees*; locally, no such guarantee exists at all. So on top of the exact polynomial-time-vs-NP-hard split of the previous subsection, we now add a constant-factor-vs-inapproximable split on the harder problem.

A unified greedy approximation algorithm. The global structural properties are exactly what let us borrow well-established approximation machinery. With a non-decreasing monotone submodular function, finding a cardinally minimal explanation is an instance of *submodular set cover* (Wolsey, 1982); with a supermodular function it becomes the non-submodular variant (Shi et al., 2021). Both are classical and, importantly for us, both come with strong approximation guarantees (Wolsey, 1982; Iyer & Bilmes, 2013; Chen & Crawford, 2023). A single greedy algorithm (Algorithm 2) serves both cases; we state it first and then read off its guarantee separately for contrastive and for sufficient reasons.

Algorithm 2 Cardinally Minimal Explanation Approximation Search

Input Value function v , and some $\delta \in [0, 1]$

```

1:  $S \leftarrow \emptyset$ 
2: while  $\max_{i \in S} v(S \cup \{i\}) < \delta$  do
3:    $j \leftarrow \operatorname{argmax}_{i \in S} v(S \cup \{i\})$ 
4:    $S \leftarrow S \cup \{j\}$ 
5: end while
6: return  $S$   $\triangleright |S|$  is a (provable approximation?) of a cardinally minimal  $\delta$ -explanation

```

Algorithm 2 mirrors Algorithm 1 but runs bottom-up instead of top-down: starting from the empty subset, it repeatedly adds the feature that yields the smallest increase in $v(S \cup i)$, halting once any further addition would push the value past δ .

Cardinally minimal contrastive reasons. Take global contrastive reasons first, where monotonicity and submodularity both hold and the task is exactly *submodular set cover*, so Wolsey (1982)’s classic guarantee applies through Algorithm 2. For integer-valued objectives this gives a Harmonic-series factor, and in general an $\mathcal{O}\left(\ln\left(\frac{v([n])}{\min_{i \in [n]} v(i)}\right)\right)$ -approximation. Under an *empirical* distribution $\mathcal{D}$ with fixed sample size the factor collapses to a *constant* — growing only *logarithmically* in the sample size — which is about as strong as an approximation guarantee gets. The local side offers no comparable guarantee: even for a *single* sample point, *no bounded* approximation is possible at all. The contrast is thus at its most extreme — a constant factor globally versus unbounded inapproximability locally, and the local hardness already bites at the smallest possible instance.

Theorem 3. *Given a neural network or tree ensemble f and an empirical distribution $\mathcal{D}$ over a fixed dataset D , Algorithm 2 yields a constant $\mathcal{O}\left(\ln\left(\frac{v_{\text{con}}^g([n])}{\min_{i \in [n]} v_{\text{con}}^g(\{i\})}\right)\right)$ -approximation, bounded by $\mathcal{O}(\ln(|D|))$, for computing a global cardinally minimal δ -contrastive reason for f , assuming feature independence. In contrast, unless $PTIME=NP$, no bounded approximation exists for computing a local cardinally minimal δ -contrastive reason for any $\langle f, \mathbf{x} \rangle$, even when $|D| = 1$.*

This upper bound is essentially tight: we prove a matching *lower bound* for Theorem 3 in the case of an empirical distribution with no independence assumption, so the logarithmic factor cannot be improved in general (see Appendix M).

Cardinally minimal sufficient reasons. The sufficient case is governed by the supermodular variant of set cover, which is inherently harder to approximate than the submodular one behind contrastive reasons — so we should expect a somewhat weaker, but still bounded, guarantee. That guarantee kicks in once the function has bounded curvature (Shi et al., 2021), where the *total curvature* of $v : 2^n \rightarrow \mathbb{R}$ is $k^f := 1 - \min_{i \in [n]} \frac{v([n]) - v([n] \setminus i)}{v(i) - v(\emptyset)}$. Applying Shi et al. (2021), Algorithm 2 attains an $\mathcal{O}\left(\frac{1}{1-k^f} + \ln\left(\frac{v([n])}{\min_{i \in [n]} v(i)}\right)\right)$ -approximation — the same logarithmic term as the contrastive case, plus a single curvature-dependent $\frac{1}{1-k^f}$ term.

As with the contrastive case, a fixed empirical distribution drives this all the way down to a *constant*. The only price relative to contrastive reasons is the extra $\frac{1}{1-k^f}$ term on top of their tighter $\mathcal{O}(\ln(|D|))$ bound — a bounded overhead that leaves the guarantee a constant-factor one. And once again the local side has nothing to offer: it stays inapproximable, with no *bounded* approximation even at $|D| = 1$. Both explanation types therefore repeat the same verdict — constant-factor globally, unbounded locally.

Theorem 4. *Given a neural network or tree ensemble f and an empirical distribution $\mathcal{D}$ over a fixed dataset D , Algorithm 2 yields a constant $\mathcal{O}\left(\frac{1}{1-k^f} + \ln\left(\frac{v_{\text{diff}}^g([n])}{\min_{i \in [n]} v_{\text{diff}}^g(\{i\})}\right)\right)$ -approximation for computing a global cardinally minimal δ -sufficient reason for f , assuming feature independence. In contrast, unless $\text{PTIME}=\text{NP}$, there is no bounded approximation for computing a local cardinally minimal δ -sufficient reason for any $\langle f, \mathbf{x} \rangle$, even when $|D| = 1$.*

Taken together with Subsection 7.1, the two subsections tell one coherent story across two levels of difficulty. On the easier subset-minimal problem, the split is polynomial-time globally versus NP-hard locally; on the harder cardinally minimal problem, it is constant-factor approximable globally versus not even boundedly approximable locally. The global advantage is thus consistent and, if anything, grows as the problem gets harder.

8 LIMITATIONS AND FUTURE WORK

Our most general findings — the structural properties of the value functions — hold with no restriction on model or distribution, and we then instantiate them into concrete polynomial-time and hardness results for specific model classes (e.g., neural networks), distributional assumptions (e.g., empirical distributions), and explanation definitions. The comparisons we report are therefore two-tiered: distribution-free structural separations on one side, and sharp per-setting complexity separations on the other. Other settings naturally remain open, but the tractable/intractable frontier we chart — consistently favoring the global side across every case we examine — gives a firm quantitative baseline for extending the analysis to further explainability scenarios.

More specifically, our study brings forward several important open theoretical questions:

1. Firstly, it would be interesting to see how much further our global tractability bounds can be tightened for simpler model classes — decision trees, linear models, or other inherently tractable structures. The gap to exploit is already visible: for these models local explanations are typically intractable (Arenas et al., 2022; Subercaseaux et al., 2025), whereas the global forms inherit strong structural properties that we turn into efficient guarantees — guarantees that the extra structure of these simple models may sharpen even further. Tightening the approximation factors under empirical distributions is a closely related open direction.
2. Secondly, when *exact* computation of expectations is infeasible, it is natural to ask whether alternative techniques — such as Fully Polynomial Randomized Approximation Schemes (FPRAS) for model classes like DNF formulas (Meel et al., 2019) (and thus for tree ensembles), or Monte Carlo-based approximations as the ones used in (Subercaseaux et al., 2025) — can be employed to approximate these expectations. A key open challenge is determining how to maintain minimality guarantees under approximation error, and what tolerance levels still allow the guarantees to hold.
3. Finally, we conjecture that the same local-to-global pattern — structure appearing globally exactly where it is absent locally — recurs under alternative notions of feature importance, e.g., other loss measures such as KL divergence (Conmy et al., 2023). If monotonicity,

submodularity, or supermodularity re-emerge globally under these metrics as they do here, the same efficient-feature-selection guarantees would carry over.

9 CONCLUSION

We present a unified framework for evaluating diverse explanations, and within it the local and global variants of sufficient and contrastive reasons separate along a single, consistent line. The *local* value functions carry *none* of the structural forms that combinatorial optimization relies on, while their *global* counterparts carry *all* the relevant ones: (i) *monotonicity*, (ii) *submodularity* for contrastive reasons, and (iii) *supermodularity* for sufficient reasons. This all-or-nothing structural gap converts directly into a complexity gap that holds up under a controlled comparison — same explanation, same model, same guarantee. Globally, explanations with provable guarantees are computable efficiently even for complex models such as neural networks: polynomial-time for subset-minimal reasons and constant-factor approximable for cardinally minimal ones. Locally, the matching problems are NP-hard and, in the cardinal case, not even boundedly approximable — and this already holds in the most simplified scenarios. The frontier we chart thus favors the global side uniformly and by a wide margin, opening new avenues for future research.

REFERENCES

- Federico Adolfi, Martina G Vilas, and Todd Wareham. Complexity-Theoretic Limits on the Promises of Artificial Neural Network Reverse-Engineering. In *Proc. 46th Annual Meeting of the Cognitive Science Society (CogSci)*, 2024.
- Federico Adolfi, Martina Vilas, and Todd Wareham. The Computational Complexity of Circuit Discovery for Inner Interpretability. In *Proc. 13th Int. Conf. on Learning Representations (ICLR)*, 2025.
- Gianvincenzo Alfano, Sergio Greco, Domenico Mandaglio, Francesco Parisi, Reza Shahbazzian, and Irina Trubitsyna. Even-If Explanations: Formal Foundations, Priorities and Complexity. In *Proc. 39th AAAI Conf. on Artificial Intelligence*, pp. 15347–15355, 2025.
- Ibrahim M Almanjahie, Zouaoui Chikr Elmezouar, Ali Laksaci, and Mustapha Rachdi. KNN Local Linear Estimation of the Conditional Cumulative Distribution Function: Dependent Functional Data Case. *Comptes Rendus. Mathématique*, 356(10):1036–1039, 2018.
- David Alvarez Melis and Tommi Jaakkola. Towards Robust Interpretability with Self-Explaining Neural Networks. In *Proc. 32nd Conf. on Neural Information Processing Systems (NeurIPS)*, 2018.
- Leila Amgoud and Martin Cooper. Axiomatic Foundations of Counterfactual Explanations. 2026. Technical Report. <https://arxiv.org/abs/2602.04028>.
- Guy Amir, Shahaf Bassan, and Guy Katz. Hard to Explain: On the Computational Hardness of In-Distribution Model Interpretation. In *Proc. 27th European Conf. on Artificial Intelligence (ECAI)*, 2024.
- Alaa Anani, Tobias Lorenz, Mario Fritz, and Bernt Schiele. Pixel-level Certified Explanations via Randomized Smoothing. In *Proc. 42nd Int. Conf. on Machine Learning (ICML)*, 2025.
- Marcelo Arenas, D. Baez, P. Barceló, J. Pérez, and B. Subercaseaux. Foundations of Symbolic Languages for Model Interpretability. In *Proc. 34th Int. Conf. on Advances in Neural Information Processing Systems (NeurIPS)*, pp. 11690–11701, 2021.
- Marcelo Arenas, P. Barceló, M. Romero, and B. Subercaseaux. On Computing Probabilistic Explanations for Decision Trees. In *Proc. 35th Int. Conf. on Advances in Neural Information Processing Systems (NeurIPS)*, pp. 28695–28707, 2022.
- Marcelo Arenas, Pablo Barceló, Leopoldo Bertossi, and Mikaël Monet. On the Complexity of SHAP-Score-Based Explanations: Tractability via Knowledge Compilation and Non-Approximability Results. *Journal of Machine Learning Research (JMLR)*, 24(63):1–58, 2023.

594 Gilles Audemard, S. Bellart, L. Bounia, F. Koriche, J. Lagniez, and P. Marquis. On the Computa-
595 tional Intelligibility of Boolean Classifiers. In *Proc. 18th Int. Conf. on Principles of Knowledge*
596 *Representation and Reasoning (KR)*, pp. 74–86, 2021.

597

598 Gilles Audemard, Steve Bellart, Louenas Bounia, Frédéric Koriche, Jean-Marie Lagniez, and Pierre
599 Marquis. On Preferred Abductive Explanations for Decision Trees and Random Forests. In *Proc.*
600 *31st Int. Joint Conf. on Artificial Intelligence (IJCAI)*, pp. 643–650, 2022a.

601

602 Gilles Audemard, Steve Bellart, Louenas Bounia, Frédéric Koriche, Jean-Marie Lagniez, and Pierre
603 Marquis. Trading Complexity for Sparsity in Random Forest Explanations. In *Proc. 36th AAAI*
604 *Conf. on Artificial Intelligence*, number 5, pp. 5461–5469, 2022b.

605

606 Gilles Audemard, Jean-Marie Lagniez, Pierre Marquis, and Nicolas Szczepanski. Computing
607 Abductive Explanations for Boosted Trees. In *Proc. 26th Int. Conf. on Artificial Intelligence and*
608 *Statistics (AISTATS)*, pp. 4699–4711, 2023.

609

610 Steve Azzolin, Sagar Malhorta, Andrea Passerini, and Stefano Teso. Beyond Topological Self-
611 Explainable GNNs: A Formal Explainability Perspective. In *Proc. 42nd Int. Conf. on Machine*
612 *Learning (ICML)*, 2025.

613

614 Esma Balkir, Isar Nejadgholi, Kathleen C Fraser, and Svetlana Kiritchenko. Necessity and Sufficiency
615 for Explaining Text Classifiers: A Case Study in Hate Speech Detection. In *Proc. Conf. of the*
616 *North American Chapter of the Association for Computational Linguistics: Human Language*
617 *Technologies (NAACL-HLT)*, pp. 2672–2686, 2022.

618

619 Pablo Barceló, M. Monet, J. Pérez, and B. Subercaseaux. Model Interpretability through the Lens of
620 Computational Complexity. In *Proc. 33rd Int. Conf. on Advances in Neural Information Processing*
621 *Systems (NeurIPS)*, pp. 15487–15498, 2020.

622

623 Pablo Barceló, Alexander Kozachinskiy, Miguel Romero, Bernardo Subercaseaux, and José Verschae.
624 Explaining k-Nearest Neighbors: Abductive and Counterfactual Explanations. In *Proc. 3rd ACM*
625 *on Management of Data*, pp. 1–26, 2025.

626

627 Shahaf Bassan and G. Katz. Towards Formal XAI: Formally Approximate Minimal Explanations
628 of Neural Networks. In *Proc. 29th Int. Conf. on Tools and Algorithms for the Construction and*
629 *Analysis of Systems (TACAS)*, pp. 187–207, 2023.

630

631 Shahaf Bassan, Guy Amir, Davide Corsi, Idan Refaeli, and Guy Katz. Formally Explaining Neural
632 Networks Within Reactive Systems. In *Proc. 23rd Int. Conf. on Formal Methods in Computer-Aided*
633 *Design (FMCAD)*, pp. 1–13, 2023.

634

635 Shahaf Bassan, Guy Amir, and Guy Katz. Local vs. Global Interpretability: A Computational
636 Complexity Perspective. In *Proc. 41st Int. Conf. on Machine Learning (ICML)*, 2024.

637

638 Shahaf Bassan, Guy Amir, Meirav Zehavi, and Guy Katz. What makes an Ensemble (Un) Inter-
639 pretable? In *Proc. 42nd Int. Conf. on Machine Learning (ICML)*, 2025a.

640

641 Shahaf Bassan, Yizhak Yisrael Elboher, Tobias Ladner, Matthias Althoff, and Guy Katz. Explaining,
642 Fast and Slow: Abstraction and Refinement of Provable Explanations. In *Proc. 42nd Int. Conf. on*
643 *Machine Learning (ICML)*, 2025b.

644

645 Shahaf Bassan, Ron Eliav, and Shlomit Gur. Explain Yourself, Briefly! Self-Explaining Neural
646 Networks with Concise Sufficient Reasons. In *Proc. 13th Int. Conf. on Learning Representations*
647 *(ICLR)*, 2025c.

648

649 Shahaf Bassan, Shlomit Gur, Sergey Zeltyn, Konstantinos Mavrogiorgos, Ron Eliav, and Dimosthenis
650 Kyriazis. Self-Explaining Neural Networks for Business Process Monitoring. 2025d. Technical
651 Report. <https://arxiv.org/abs/2503.18067>.

652

653 Shahaf Bassan, Michal Moshkovitz, and Guy Katz. Additive Models Explained: A Computational
654 Complexity Approach. In *Proc. 39th Conf. on Neural Information Processing Systems (NeurIPS)*,
655 2025e.

648 Shahaf Bassan, Yizhak Yisrael Elboher, Tobias Ladner, Volkan Şahin, Jan Kretinsky, Matthias Althoff,
649 and Guy Katz. Provably Explaining Neural Additive Models. In *Proc. 14th Int. Conf. on Learning*
650 *Representations (ICLR)*, 2026.

651 Robi Bhattacharjee and Ulrike Luxburg. Auditing Local Explanations is Hard. In *Proc. 38th Conf.*
652 *on Neural Information Processing Systems (NeurIPS)*, pp. 18593–18632, 2024.

653 Blair Bilodeau, Natasha Jaques, Pang Wei Koh, and Been Kim. Impossibility Theorems for Feature
654 Attribution. *Proc. of the National Academy of Sciences (PNAS)*, 121(2):e2304406120, 2024.

655 Guy Blanc, Jane Lange, and Li-Yang Tan. Provably Efficient, Succinct, and Precise Explanations.
656 In *Proc. 35th Int. Conf. on Advances in Neural Information Processing Systems (NeurIPS)*, pp.
657 6129–6141, 2021.

658 Guy Blanc, Caleb Koch, Jane Lange, and Li-Yang Tan. A Query-Optimal Algorithm for Finding
659 Counterfactuals. In *Proc. 39th Int. Conf. on Machine Learning (ICML)*, pp. 2075–2090, 2022.

660 Ryma Boumazouza, Fahima Cheikh-Alili, Bertrand Mazure, and Karim Tabia. ASTERYX: A model-
661 Agnostic SaT-basEd appRoach for SYmbolic and Score-Based EXplanations. In *Proc. 30th ACM*
662 *Int. Conf. on Information & Knowledge Management (CIKM)*, pp. 120–129, 2021.

663 Ryma Boumazouza, Raya Elsaleh, Melanie Ducoffe, Shahaf Bassan, and Guy Katz. FAME: For-
664 mal Abstract Minimal Explanation for Neural Networks. In *Proc. 14th Int. Conf. on Learning*
665 *Representations (ICLR)*, 2026.

666 Louenas Bounia. Enhancing the Intelligibility of Boolean Decision Trees with Concise and Reliable
667 Probabilistic Explanations. In *Proc. 20th Int. Conf. on Information Processing and Management of*
668 *Uncertainty in Knowledge-Based Systems (IPMU)*, pp. 205–218, 2024.

669 Louenas Bounia. Using Submodular Optimization to Approximate Minimum-Size Abductive Path
670 Explanations for Tree-Based Models. In *Proc. 41st Conf. on Uncertainty in Artificial Intelligence*
671 *(UAI)*, 2025.

672 Louenas Bounia and Frederic Koriche. Approximating Probabilistic Explanations via Supermodular
673 Minimization. In *Proc. 39th Conf. on Uncertainty in Artificial Intelligence (UAI)*, pp. 216–225,
674 2023.

675 Marco Calautti, Enrico Malizia, and Cristian Molinaro. On the Complexity of Global Necessary
676 Reasons to Explain Classification. 2025. Technical Report. [https://arXiv.org/abs/](https://arXiv.org/abs/2501.06766)
677 [2501.06766](https://arXiv.org/abs/2501.06766).

678 Emilio Carrizosa, Jasone Ramírez-Ayerbe, and Dolores Romero Morales. Mathematical Optimization
679 Modelling for Group Counterfactual Explanations. *European Journal of Operational Research*
680 *(EJOR)*, 319(2):399–412, 2024.

681 Brandon Carter, Jonas Mueller, Siddhartha Jain, and David Gifford. What Made You Do This?
682 Understanding Black-Box Decisions with Sufficient Input Subsets. In *Proc. 22nd Int. Conf. on*
683 *Artificial Intelligence and Statistics (AISTATS)*, pp. 567–576, 2019.

684 Brandon Carter, Siddhartha Jain, Jonas W Mueller, and David Gifford. Overinterpretation Reveals
685 Image Classification Model Pathologies. In *Proc. 35th Conf. on Neural Information Processing*
686 *Systems (NeurIPS)*, pp. 15395–15407, 2021.

687 Wenjing Chen and Victoria Crawford. Bicriteria Approximation Algorithms for the Submodular
688 Cover Problem. In *Proc. 37th Conf. on Neural Information Processing Systems (NeurIPS)*, pp.
689 72705–72716, 2023.

690 Arthur Conmy, Augustine Mavor-Parker, Aengus Lynch, Stefan Heimersheim, and Adrià Garriga-
691 Alonso. Towards Automated Circuit Discovery for Mechanistic Interpretability. *Proc. 37th Conf.*
692 *on Neural Information Processing Systems (NeurIPS)*, pp. 16318–16352, 2023.

693 Yves Crama and Peter L Hammer. *Boolean Functions: Theory, Algorithms, and Applications*.
694 Cambridge University Press, 2011.

- Adnan Darwiche. Logic for Explainable AI. In *Proc. 38th ACM/IEEE Symposium on Logic in Computer Science (LICS)*, pp. 1–11, 2023.
- Adnan Darwiche and Auguste Hirth. On the Reasons Behind Decisions. In *Proc. 24th European Conf. on Artificial Intelligence (ECAI)*, pp. 712–720, 2020.
- Adnan Darwiche and Chunxi Ji. On the Computation of Necessary and Sufficient Explanations. In *Proc. 36th AAAI Conf. on Artificial Intelligence*, pp. 5582–5591, 2022.
- Sanjoy Dasgupta, Nave Frost, and Michal Moshkovitz. Framework for Evaluating Faithfulness of Local Explanations. In *Proc. 39th Int. Conf. on Machine Learning (ICML)*, pp. 4794–4815, 2022.
- Ian Davidson, Nicolás Kennedy, and SS Ravi. CXAD: Contrastive Explanations for Anomaly Detection: Algorithms, Complexity Results and Experiments. *Transactions on Machine Learning Research (TMLR)*, 2025.
- Alessandro De Palma, Greta Dolcetti, and Caterina Urban. Faster Verified Explanations for Neural Networks. 2025. Technical Report. <https://arXiv.org/abs/2512.00164>.
- Amit Dhurandhar, Pin-Yu Chen, Ronny Luss, Chun-Chen Tu, Paishun Ting, Karthikeyan Shanmugam, and Payel Das. Explanations Based on the Missing: Towards Contrastive Explanations with Pertinent Negatives. *Proc. 32nd Conf. on Neural Information Processing Systems (NeurIPS)*, 2018.
- Hidde Fokkema, Rianne De Heide, and Tim Van Erven. Attribution-Based Explanations that Provide Recourse Cannot be Robust. *Journal of Machine Learning Research (JMLR)*, 24(360):1–37, 2023.
- Hidde Fokkema, Damien Garreau, and Tim van Erven. The Risks of Recourse in Binary Classification. In *Proc. 27th Int. Conf. on Artificial Intelligence and Statistics (AISTATS)*, pp. 550–558, 2024.
- Tobias Geibinger, Reijo Jaakkola, Antti Kuusisto, Xinghan Liu, and Miikka Vilander. Why This and Not That? A Logic-based Framework for Contrastive Explanations. In *Proc. 19th European Conf. on Logics in Artificial Intelligence (JELIA)*, pp. 45–60. Springer, 2025.
- Joshua Größer and Ostap Okhrin. Copulae: An Overview and Recent Developments. *Wiley Interdisciplinary Reviews: Computational Statistics*, 14(3):e1557, 2022.
- Riccardo Guidotti. Counterfactual Explanations and How to Find Them: Literature Review and Benchmarking. *Data Mining and Knowledge Discovery (DMKD)*, 38(5):2770–2824, 2024.
- Hyukjun Gweon, Matthias Schonlau, and Stefan H Steiner. The K Conditional Nearest Neighbor Algorithm for Classification and Class Probability Estimation. *PeerJ Computer Science*, 5:e194, 2019.
- Itamar Hadad, Shahaf Bassan, and Guy Katz. Formal Mechanistic Interpretability: Automated Circuit Discovery with Provable Guarantees. In *Proc. 14th Int. Conf. on Learning Representations (ICLR)*, 2026.
- Ouns El Harzli, Bernardo Cuenca Grau, and Ian Horrocks. Cardinality-Minimal Explanations for Monotonic Neural Networks. In *Proc. 32nd Int. Joint Conf. on Artificial Intelligence (IJCAI)*, pp. 3677–3685, 2023.
- W Keith Hastings. Monte Carlo Sampling Methods using Markov Chains and their Applications. *Biometrika*, 57:97–109, 1970.
- Rongjie Huang, Max W. Y. Lam, Jun Wang, Dan Su, Dong Yu, Yi Ren, and Zhou Zhao. FastDiff: A Fast Conditional Diffusion Model for High-Quality Speech Synthesis. In *Proc. 31st Int. Joint Conf. on Artificial Intelligence (IJCAI)*, pp. 4157–4163, 2022.
- Xuanxiang Huang, Yacine Izza, Alexey Ignatiev, and Joao Marques-Silva. On Efficiently Explaining Graph-Based Classifiers. In *Proc. 18th Int. Conf. on Principles of Knowledge Representation and Reasoning (KR)*, 2021.
- Alexey Ignatiev, Nina Narodytska, and Joao Marques-Silva. Abduction-Based Explanations for Machine Learning Models. In *Proc. 33rd AAAI Conference on Artificial Intelligence*, pp. 1511–1519, 2019.

756 Alexey Ignatiev, Martin Cooper, Mohamed Siala, Emmanuel Hebrard, and Joao Marques-Silva.
757 Towards Formal Fairness in Machine Learning. In *Proc. 26th Int. Conf. on Principles and Practice*
758 *of Constraint Programming (CP)*, pp. 846–867, 2020a.

759 Alexey Ignatiev, Nina Narodytska, Nicholas Asher, and Joao Marques-Silva. From Contrastive to
760 Abductive Explanations and Back Again. In *Proc. Int. Conf. of the Italian Association for Artificial*
761 *Intelligence (AIxAI)*, pp. 335–355, 2020b.

762 Alexey Ignatiev, Yacine Izza, Peter J Stuckey, and Joao Marques-Silva. Using MaxSAT for Efficient
763 Explanations of Tree Ensembles. In *Proc. 36th AAAI Conference on Artificial Intelligence*, pp.
764 3776–3785, 2022.

765 Rishabh K Iyer and Jeff A Bilmes. Submodular Optimization with Submodular Cover and Submodular
766 Knapsack Constraints. *Proc. 27th Conf. on Neural Information Processing Systems (NeurIPS)*,
767 2013.

768 Yacine Izza and Joao Marques-Silva. On Explaining Random Forests with SAT. In *Proc. 30th Int.*
769 *Joint Conf. on Artificial Intelligence (IJCAI)*, 2021.

770 Yacine Izza, Xuanxiang Huang, Alexey Ignatiev, Nina Narodytska, Martin Cooper, and Joao Marques-
771 Silva. On Computing Probabilistic Abductive Explanations. *Int. Journal of Approximate Reasoning*
772 *(IJAR)*, 159:108939, 2023.

773 Yacine Izza, Xuanxiang Huang, Antonio Morgado, Jordi Planes, Alexey Ignatiev, and Joao Marques-
774 Silva. Distance-Restricted Explanations: Theoretical Underpinnings & Efficient implementation.
775 In *Proc. 21st Int. Conf. on Principles of Knowledge Representation and Reasoning (KR)*, pp.
776 475–486, 2024a.

777 Yacine Izza, Kuldeep S Meel, and Joao Marques-Silva. Locally-Minimal Probabilistic Explanations.
778 In *Proc. 27th European Conf. on Artificial Intelligence (ECAI)*, 2024b.

779 Reijo Jaakkola, Tomi Janhunen, Antti Kuusisto, Masood Feyzbakhsh Rankooh, and Miikka Vilander.
780 Explainability via Short Formulas: the Case of Propositional Logic with Implementation. *Journal*
781 *of Artificial Intelligence Research (JAIR)*, 83, 2025.

782 Alon Jacovi, Swabha Swayamdipta, Shauli Ravfogel, Yanai Elazar, Yejin Choi, and Yoav Goldberg.
783 Contrastive Explanations for Model Interpretability. In *Proc. Conf. Conference on Empirical*
784 *Methods in Natural Language Processing (EMNLP)*, pp. 1597–1611, 2021.

785 Helen Jin, Anton Xue, Weiqiu You, Surbhi Goel, and Eric Wong. Probabilistic Stability Guarantees
786 for Feature Attributions. In *Proc. 39th Conf. on Neural Information Processing Systems (NeurIPS)*,
787 2025.

788 Amir-Hossein Karimi, Bernhard Schölkopf, and Isabel Valera. Algorithmic recourse: from counter-
789 factual explanations to interventions. In *Proc. 4th ACM Conf. on Fairness, Accountability, and*
790 *Transparency (FAccT)*, pp. 353–362, 2021.

791 Eoin M Kenny and Mark T Keane. On Generating Plausible Counterfactual and Semi-Factual
792 Explanations for Deep Learning. In *Proc. 35th AAAI Conf. on Artificial Intelligence*, number 13,
793 pp. 11575–11585, 2021.

794 Alexander Kozachinskiy. Inapproximability of Sufficient Reasons for Decision Trees. 2023. Technical
795 Report. <https://arXiv.org/abs/2304.02781>.

800 I Elizabeth Kumar, Suresh Venkatasubramanian, Carlos Scheidegger, and Sorelle Friedler. Problems
801 with Shapley-Value-Based Explanations as Feature Importance Measures. In *Proc. 37th Int. Conf.*
802 *on Machine Learning (ICML)*, pp. 5491–5500, 2020.

803 Emanuele La Malfa, Agnieszka Zbrzezny, Rhiannon Michelmore, Nicola Paoletti, and Marta
804 Kwiatkowska. On Guaranteed Optimal Robust Explanations for NLP Models. In *Proc. Int.*
805 *Joint Conf. on Artificial Intelligence (IJCAI)*, pp. 2658–2665, 2021.

806 Eduardo Sany Laber. The Computational Complexity of some Explainable Clustering Problems.
807 *Information Processing Letters*, 184:106437, 2024.

810 Shuai Li, Yingjie Zhang, Hongtu Zhu, Christina Wang, Hai Shu, Ziqi Chen, Zhuoran Sun, and
811 Yanfeng Yang. K-Nearest-Neighbor Local Sampling Based Conditional Independence Testing.
812 *Proc. 37th Conf. on Neural Information Processing Systems (NeurIPS)*, pp. 23321–23344, 2023.

813 Gianluigi Lopardo, Frédéric Precioso, and Damien Garreau. A Sea of Words: An In-Depth Analysis
814 of Anchors for Text Data. In *Proc. 26th Int. Conf. on Artificial Intelligence and Statistics (AISTATS)*
815 *2023*, 2023.

816 Scott M Lundberg and Su-In Lee. A Unified Approach to Interpreting Model Predictions. *Proc. 31st*
817 *Conf. on Neural Information Processing Systems (NeurIPS)*, 2017.

818 Joao Marques-Silva and Alexey Ignatiev. Delivering Trustworthy AI through formal XAI. In *Proc.*
819 *36th AAAI Conf. on Artificial Intelligence*, pp. 12342–12350, 2022.

820 Joao Marques-Silva and Alexey Ignatiev. No Silver Bullet: Interpretable ML Models Must be
821 Explained. *Frontiers in Artificial Intelligence*, 6:1128212, 2023.

822 Joao Marques-Silva, Thomas Gerspacher, Martin Cooper, Alexey Ignatiev, and Nina Narodytska.
823 Explaining Naive Bayes and Other Linear Classifiers with Polynomial Time and Delay. In *Proc.*
824 *33rd Int. Conf. on Advances in Neural Information Processing Systems (NeurIPS)*, pp. 20590–
825 20600, 2020.

826 Joao Marques-Silva, Thomas Gerspacher, Martin Cooper, Alexey Ignatiev, and Nina Narodytska.
827 Explanations for Monotonic Classifiers. In *Proc. Int. Conf. on Machine Learning (ICML)*, pp.
828 7469–7479, 2021.

829 Reda Marzouk and Colin De La Higuera. On the Tractability of SHAP Explanations under Markovian
830 Distributions. In *Proc. 41st Int. Conf. on Machine Learning (ICML)*, pp. 34961–34986, 2024.

831 Reda Marzouk, Shahaf Bassan, and Guy Katz. SHAP Meets Tensor Networks: Provably Tractable
832 Explanations with Parallelism. In *Proc. 39th Conf. on Neural Information Processing Systems*
833 *(NeurIPS)*, 2025a.

834 Reda Marzouk, Shahaf Bassan, Guy Katz, and De la Higuera. On the Computational Tractability
835 of the (Many) Shapley Values. In *Proc. 28th Int. Conf. on Artificial Intelligence and Statistics*
836 *(AISTATS)*, 2025b.

837 Kuldeep S Meel, Aditya A Shrotri, and Moshe Y Vardi. Not all FPRASs are Equal: Demystifying
838 FPRASs for DNF-Counting. *Constraints*, 24(3):211–233, 2019.

839 Ramaravind K Mothilal, Amit Sharma, and Chenhao Tan. Explaining Machine Learning Classifiers
840 Through Diverse Counterfactual Explanations. In *Proc. 3rd Conf. on Fairness, Accountability, and*
841 *Transparency (FAccT)*, pp. 607–617, 2020.

842 Vidya Muthukumar, Tejaswini Pedapati, Nalini Ratha, Prasanna Sattigeri, Chai-Wah Wu, Brian
843 Kingsbury, Abhishek Kumar, Samuel Thomas, Aleksandra Mojsilovic, and Kush R Varshney.
844 Understanding Unequal Gender Classification Accuracy from Face Images. 2018. Technical
845 Report. <https://arXiv.org/abs/1812.00099>.

846 Sebastian Ordyniak, Giacomo Paesani, and Stefan Szeider. The Parameterized Complexity of Finding
847 Concise Local Explanations. In *Proc. 32nd Int. Joint Conf. on Artificial Intelligence (IJCAI)*, 2023.

848 Sebastian Ordyniak, Giacomo Paesani, Mateusz Rychlicki, and Stefan Szeider. Explaining Decisions
849 in ML Models: A Parameterized Complexity Analysis. In *Proc. 21st Int. Conf. on Principles of*
850 *Knowledge Representation and Reasoning (KR)*, pp. 563–573, 2024.

851 Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. Anchors: High-Precision Model-Agnostic
852 Explanations. In *Proc. 32nd AAAI Conf. on Artificial Intelligence*, 2018.

853 Majun Shi, Zishen Yang, and Wei Wang. Minimum Non-Submodular Cover Problem with Applica-
854 tions. *Applied Mathematics and Computation*, 410:126442, 2021.

855 Andy Shih, Arthur Choi, and Adnan Darwiche. A Symbolic Approach to Explaining Bayesian
856 Network Classifiers. In *Proc. 27th Int. Joint Conf. on Artificial Intelligence (IJCAI)*, pp. 5103–
857 5111, 2018.

-
- Bernardo Subercaseaux, Marcelo Arenas, and Kuldeep S Meel. Probabilistic Explanations for Linear Models. In *Proc. 39th AAAI Conf. on Artificial Intelligence*, pp. 20655–20662, 2025.
- Guy Van den Broeck, Anton Lykov, Maximilian Schleich, and Dan Suciu. On the Tractability of SHAP Explanations. *Journal of Artificial Intelligence Research (JAIR)*, 74:851–886, 2022.
- Gal Vardi, Gilad Yehudai, and Ohad Shamir. On the Optimal Memorization Power of ReLU Neural Networks. In *Proc. 10th Int. Conf. on Learning Representations, (ICLR)*, 2022.
- Vijay V Vazirani. *Approximation Algorithms*. Springer, 2001.
- Sahil Verma, Varich Boonsanong, Minh Hoang, Keegan Hines, John Dickerson, and Chirag Shah. Counterfactual Explanations and Algorithmic Recourses for Machine Learning: A Review. *ACM Computing Surveys*, 56(12):1–42, 2024.
- Stephan Wäldchen, Jan Macdonald, Sascha Hauch, and Gitta Kutyniok. The Computational Complexity of Understanding Binary Classifier Decisions. *Journal of Artificial Intelligence Research (JAIR)*, 70:351–387, 2021.
- E. Wang, P. Khosravi, and G. Van den Broeck. Probabilistic Sufficient Explanations. In *Proc. 30th Int. Joint Conf. on Artificial Intelligence (IJCAI)*, 2021a.
- S. Wang, H. Zhang, K. Xu, X. Lin, S. Jana, C. Hsieh, and Z. Kolter. Beta-Crown: Efficient Bound Propagation with Per-Neuron Split Constraints for Neural Network Robustness Verification. In *Proc. 35th Int. Conf. on Advances in Neural Information Processing Systems (NeurIPS)*, pp. 29909–29921, 2021b.
- Greta Warren, Eoin Delaney, Christophe Guéret, and Mark T Keane. Explaining Multiple Instances Counterfactually: User Tests of Group-Counterfactuals for XAI. In *Proc. 32nd Int. Conf. on Case-Based Reasoning (ICCBR)*, pp. 206–222, 2024.
- Ingo Wegener. The Complexity of Symmetric Boolean Functions. *Computation Theory and Logic*, pp. 433–442, 2005.
- Laurence A Wolsey. An Analysis of the Greedy Algorithm for the Submodular Set Covering Problem. *Combinatorica*, 2(4):385–393, 1982.
- Haoze Wu, Omri Isac, Aleksandar Zeljić, Teruhiro Tagomori, Matthew Daggitt, Wen Kokke, Idan Refaeli, Guy Amir, Kyle Julian, Shahaf Bassan, et al. Marabou 2.0: A Versatile Formal Analyzer of Neural Networks. In *Proc. 36th Int. Conf. on Computer Aided Verification (CAV)*, pp. 249–264, 2024a.
- Min Wu, Haoze Wu, and Clark Barrett. Verix: Towards Verified Explainability of Deep Neural Networks. In *Proc. 36th Int. Conf. on Advances in Neural Information Processing Systems (NeurIPS)*, 2024b.
- Lei Xu, Maria Skoularidou, Alfredo Cuesta-Infante, and Kalyan Veeramachaneni. Modeling Tabular Data Using Conditional GAN. *Proc. 33rd Conf. on Neural Information Processing Systems (NeurIPS)*, 2019.
- Anton Xue, Rajeev Alur, and Eric Wong. Stability Guarantees for Feature Attributions with Multiplicative Smoothing. *Proc. 37th Conf. on Neural Information Processing Systems (NeurIPS)*, pp. 62388–62413, 2023.
- Kayo Yin and Graham Neubig. Interpreting Language Models with Contrastive Explanations. In *Proc. Conf. on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 184–198, 2022.
- Jinqiang Yu, Alexey Ignatiev, Peter Stuckey, Nina Narodytska, and Joao Marques-Silva. Eliminating the Impossible, Whatever Remains Must Be True: On Extracting and Applying Background Knowledge in the Context of Formal Explanations. In *Proc. 37th AAAI Conf. on Artificial Intelligence*, pp. 4123–4131, 2023.

Appendix

The appendix contains formalizations and proofs that were mentioned throughout the paper:

Appendix A contains extended related work on Formal XAI.

Appendix B contains the formalizations of the models and distributions used in this work.

Appendix C contains the proof of Proposition 1.

Appendix D contains the proof of Proposition 2.

Appendix E contains the proof of Proposition 3.

Appendix F contains the proof of Proposition 4.

Appendix G contains the proof of Proposition 5.

Appendix H contains the proof of Theorem 1.

Appendix I contains the proof of Proposition 6.

Appendix J contains the proof of Theorem 2.

Appendix K contains the proof of Theorem 3.

Appendix L contains the proof of Theorem 4.

Appendix O contains an LLM usage disclosure.

A EXTENDED RELATED WORK ON FORMAL XAI

Our work falls within the line of research on *formal explainable AI* (formal XAI) (Marques-Silva & Ignatiev, 2022), which studies explanations equipped with provable guarantees (Yu et al., 2023; Darwiche & Ji, 2022; Darwiche, 2023; Shih et al., 2018; Azzolin et al., 2025; Audemard et al., 2021; Calautti et al., 2025). A central theme in this area concerns the computational complexity of generating such guaranteed explanations (Barceló et al., 2020; ?; Marzouk et al., 2025a;b; Marzouk & De La Higuera, 2024; Bassan et al., 2025e; Amir et al., 2024; Blanc et al., 2021; 2022; Jaakkola et al., 2025; Amgoud & Cooper, 2026; Geibinger et al., 2025). Since obtaining formal explanations is often computationally intractable for expressive models such as neural networks (Adolfi et al., 2024; Barceló et al., 2020) and tree ensembles (Ordyniak et al., 2024; Bassan et al., 2025a), much of the literature has focused on more tractable model classes (Marques-Silva & Ignatiev, 2023), including decision trees (Bounia, 2025; Arenas et al., 2022; Bounia, 2024; Bounia & Koriche, 2023), monotonic models (Marques-Silva et al., 2021; Harzli et al., 2023), linear models (Marques-Silva et al., 2020; Subercaseaux et al., 2025), and additive models (Bassan et al., 2026).

Beyond complexity-theoretic investigations, substantial effort has also been devoted to computing formal explanations in practice using automated reasoning tools, such as MaxSAT and MILP solvers for tree ensembles (Audemard et al., 2023; 2022a; Ignatiev et al., 2022; Boumazouza et al., 2021), and neural network verification frameworks (Wang et al., 2021b; Wu et al., 2024a) to derive certified explanations for neural networks (Wu et al., 2024b; Bassan & Katz, 2023; Bassan et al., 2023; Izza et al., 2024a; Hadad et al., 2026; Ignatiev et al., 2019). Additional approaches seek to alleviate the inherent computational burden through abstractions (De Palma et al., 2025; Bassan et al., 2025b; Boumazouza et al., 2026), relaxations such as smoothing (Xue et al., 2023; Anani et al., 2025; Jin et al., 2025), or training-time interventions (Alvarez Melis & Jaakkola, 2018; Bassan et al., 2025c;d).

Finally, within formal XAI, a related line of work applies combinatorial and submodular optimization to speed up the computation of *local* probabilistic sufficient reasons (Bounia, 2025; Bounia & Koriche, 2023), chiefly for decision trees. A key caveat in that line — and one our analysis confirms — is that the local objective is *not* actually submodular; the greedy submodular-style algorithms there are effective empirical heuristics rather than methods with provable guarantees. Our contribution is exactly to close that gap on the global side: we show the *global* sufficient and contrastive objectives genuinely are (super/sub)modular, so the same greedy machinery that is merely heuristic locally becomes provably guaranteed globally — turning an empirical rule of thumb into a formal, worst-case guarantee.

B MODEL AND DISTRIBUTION FORMALIZATIONS

In this section, we formalize the models and distributions referenced throughout the paper. Specifically, Subsection B.1 defines the model families, while Subsection B.2 formalizes the distributions.

B.1 MODEL FORMALIZATIONS

In this subsection, we formalize the three base-model types that were analyzed throughout the paper: (i) (axis-aligned) decision trees, (ii) linear classifiers, and (iii) neural networks with ReLU activations.

Decision Trees. We define a decision tree (DT) as a directed acyclic graph that represents a function $f : \mathbb{F} \rightarrow [c]$, where $c \in \mathbb{N}$ denotes the number of classes. The graph encodes the function as follows: (i) Each internal node v is assigned a distinct binary input feature from the set $\{1, \dots, n\}$; (ii) Every internal node v has at most k outgoing edges, each corresponding to a value in $[k]$ assigned to v ; (iii) Along any path α in the decision tree, each variable appears at most once; (iv) Each leaf node is labeled with a class from $[c]$. Thus, assigning values to the inputs $\mathbf{x} \in \mathbb{F}$ uniquely determines a path α from the root to a leaf in the DT, where the function output $f(\mathbf{x})$ corresponds to a class label $i \in [c]$. The size of the DT, denoted $|f|$, is defined as the total number of edges in the graph. To allow for flexible modeling, the ordering of input variables $\{1, \dots, n\}$ may differ across distinct paths α and α' , ensuring that no variable is repeated along a single path.

Neural Networks. We present our *hardness* proofs for neural networks with ReLU activations, though our *tractability results* (i.e., polynomial-time algorithms) apply to any architecture that allows for polynomial-time inference — a standard assumption. Thus, all separation results between tractable and intractable cases we prove carry over to *any neural architecture* at least as expressive as a standard feed-forward ReLU network, encompassing many widely used models such as ResNets, CNNs, Transformers, Diffusion models, and more. Moreover, note that any ReLU network can be represented as a fully connected network by assigning zero weights and biases to missing connections. Following standard conventions (Barceló et al., 2020; Bassan et al., 2024; Adolphi et al., 2025), we thus assume the network is fully connected. Specifically, our analysis applies to multi-layer perceptrons (MLPs). Formally, an MLP f consists of $t - 1$ *hidden layers* ($g^{(j)}$ for $j = 1, \dots, t - 1$) and one output layer $g^{(t)}$, where each layer is defined recursively as:

$$g^{(j)} := \sigma^{(j)}(g^{(j-1)}W^{(j)} + b^{(j)}) \quad (9)$$

where $W^{(j)}$ denotes the weight matrix, $b^{(j)}$ the bias vector, and $\sigma^{(j)}$ the activation function of the j -th layer. Accordingly, the model comprises t weight matrices ($W^{(1)}, \dots, W^{(t)}$), t bias vectors ($b^{(1)}, \dots, b^{(t)}$), and t activation functions ($\sigma^{(1)}, \dots, \sigma^{(t)}$).

The input layer is defined as $g^{(0)} = \mathbf{x} \in \{0, 1\}^n$, representing the binary input vector. The dimensions of the network are governed by a sequence of positive integers $d_0, \dots, d_t$, with weight matrices and bias vectors given by $W^{(j)} \in \mathbb{Q}^{d_{j-1} \times d_j}$ and $b^{(j)} \in \mathbb{Q}^{d_j}$, respectively. These parameters are learned during training. Since the model functions as a binary classifier over n features, we set $d_0 = n$ and $d_t = 1$. The hidden layers use the ReLU activation function, defined by $\text{ReLU}(x) = \max(0, x)$. Although a sigmoid activation is typically used during training, for interpretability purposes, we assume the output layer applies a threshold-based step function, defined as $\text{step}(\mathbf{z}) = 1$ if $\mathbf{z} \geq 0$ and $\text{step}(\mathbf{z}) = 0$ otherwise.

Tree Ensembles. Many popular ensemble methods exist, but since our goal is to provide post-hoc explanations, we focus on the *inference* phase rather than the training process. Our analysis centers on ensemble families that rely on either *majority voting* (e.g., Random Forests) or *weighted voting* (e.g., XGBoost) during inference. However, as with our treatment of neural networks, our *tractability results* — namely, polynomial-time algorithms — extend to any possible ensemble configuration with polynomial-time inference, encompassing an even broader range of ensemble configurations.

Majority Voting Inference. In *majority voting* inference, the final prediction $f(\mathbf{x})$ is assigned to the class $j \in [c]$ that receives the majority of votes among the individual predictions $f_i(\mathbf{x})$ from all $i \in [k]$ (i.e., from each tree in the ensemble).

$$f(\mathbf{x}) := \text{majority}(\{f_i(\mathbf{x}) \mid i \in [k]\}) \quad (10)$$

where $\text{majority}(S)$ denotes the most frequent label in the multiset S . If there is a tie, it is resolved by a fixed tie-breaking rule (e.g., lexicographic order or predefined priority).

Weighted Voting Inference. In *weighted voting* inference, each model in the ensemble is assigned a weight $\phi_i \in \mathbb{Q}$ representing its relative importance. The predicted class is the one with the highest total weight across all models. Formally, for any $\mathbf{x} \in \mathbb{F}$, we define f as:

$$f(\mathbf{x}) := \arg \max_{j \in [c]} \sum_{i=1}^k \phi_i \cdot \mathbb{I}[f_i(\mathbf{x}) = j] \quad (11)$$

where $\mathbb{I}$ denotes the identity function.

B.2 DISTRIBUTION FORMALIZATIONS

This subsection formalizes the distribution definitions discussed in the main paper.

Empirical Distributions. The distribution $\mathcal{D}$ over the input features will be defined based on a dataset $\mathbf{D}$ comprising various inputs $\mathbf{z}^1, \mathbf{z}^2, \dots, \mathbf{z}^{|\mathbf{D}|}$. Here, the distribution of any given input $\mathbf{x} \in \mathbb{F}$ is defined by the frequency of occurrence of $\mathbf{x}$ within $\mathbf{D}$, specifically by:

$$\mathbf{Pr}(\mathbf{x}) := \frac{1}{|\mathbf{D}|} \sum_{i=1}^{|\mathbf{D}|} \mathbb{I}(\mathbf{z}^i = \mathbf{x}) \quad (12)$$

Independent Distributions. Formally, given a probability value $p(\mathbf{x}_i) \in [0, 1]$ defined for each individual input feature, we say that the distribution $\mathcal{D}$ is independent if the joint probability over inputs is given by $\mathbf{Pr}(\mathbf{x}) := \prod_{i \in [n]} p(\mathbf{x}_i)$.

We observe that when $p(\mathbf{x}_i) = p(\mathbf{x}_j)$ holds for all $i, j \in [n]$, the distribution reduces to the *uniform distribution*, a special case of independent distributions.

General Distributions. We note that many of the proofs in this work — particularly those concerning fundamental properties of value functions — apply broadly over general distribution assumptions. While empirical distributions (using the training dataset as a proxy) are common in XAI, alternative frameworks for approximating distributions include k-NN resampling from nearby points (Li et al., 2023; Almanjahie et al., 2018; Gweon et al., 2019), copulas for modeling tabular dependencies (Größer & Okhrin, 2022), or more advanced conditional generative models such as CTGAN (Xu et al., 2019) and conditional diffusion models (Huang et al., 2022). Such choices are common in the counterfactual explanation and algorithmic recourse literature (Karimi et al., 2021; Fokkema et al., 2024; 2023; Verma et al., 2024), which are related to contrastive explanations. Furthermore, when the value function is not computed directly from a structural property of the model (e.g., leaf enumeration in decision trees) or from empirical distributions, one may instead approximate it using methods such as Monte Carlo sampling (Hastings, 1970), following approaches similar to (Subercaseaux et al., 2025; Lopardo et al., 2023).

B.3 SUBSET VS. CARDINAL MINIMALITY

In this subsection, we provide a more detailed discussion of the distinction between cardinal and subset minimality. Cardinal minimality offers a substantially stronger guarantee, as it corresponds to a globally minimal explanation size, whereas subset minimality only ensures a local form of minimality.

To see why cardinal minimality is stronger, note that if $S \subseteq [n]$ is a cardinally minimal explanation, then no smaller set S' with $|S'| < |S|$ can qualify as an explanation. Hence, no strict subset $S' \subsetneq S$ is an explanation, which means S is also subset minimal.

However, the reverse does not hold. Consider the function:

$$f := \mathbf{x}_1 \vee (\mathbf{x}_2 \wedge \mathbf{x}_3) \quad (13)$$

and the assignment $\mathbf{x} := (1, 1, 1)$, which gives $f(1, 1, 1) = 1$. Fixing feature $\mathbf{x}_1 = 1$ yields a cardinally minimal (and subset minimal) sufficient reason, since the prediction remains 1 regardless

of $\mathbf{x}_2, \mathbf{x}_3$. But fixing both $\mathbf{x}_2 = 1$ and $\mathbf{x}_3 = 1$ also gives a subset minimal explanation—yet not a cardinally minimal one. Thus, while every cardinally minimal explanation is subset minimal, not every subset minimal explanation is cardinally minimal.

C PROOF OF PROPOSITION 1

Proposition 1. *While the non-probabilistic case ($\delta = 1$) admits a unique subset-minimal global sufficient reason (Bassan et al., 2024), in the general probabilistic setting (for arbitrary δ), there exist functions f that have $\Theta(\frac{2^n}{\sqrt{n}})$ subset-minimal global sufficient reasons.*

Proof. It is known that certain Boolean functions admit an exponential number of subset-minimal local (non-probabilistic) sufficient reasons for some input $\mathbf{x}$ Bassan et al. (2024). In particular, this phenomenon was shown to occur in functions of the following form. We demonstrate that the same function admits an exponential number of *global* and *probabilistic* sufficient reasons. Notably, the mentioned function is a special case of a broader class of *threshold Boolean functions* described in Wegener (2005). Specifically, for $n = 2k + 1$ with $k \in \mathbb{N}$, this function is defined as:

$$f(\mathbf{x}) := \begin{cases} 1 & \text{if } \sum_{i=1}^n \mathbf{x}_i \geq k + 1 \\ 0 & \text{otherwise} \end{cases} \quad (14)$$

where $\mathbf{x}$ is drawn from a uniform distribution $\mathcal{D}$. Notably, the function f is *symmetric*, in the sense of symmetric threshold Boolean functions Wegener (2005): its output depends solely on the number of 1’s (or, equivalently, 0’s) in the input and is invariant under any permutation of input bits. Furthermore, the two extreme inputs — the all-zeros vector $(0, \dots, 0)$ and the all-ones vector $(1, \dots, 1)$ — each admit an exponential number of subset-minimal local (non-probabilistic) sufficient reasons.

Since the function f is symmetric, the condition $v_{\text{suff}}^g(S_1) = v_{\text{suff}}^g(S_2)$ holds for any two subsets $S_1, S_2 \subseteq [n]$. Moreover, it can be verified that for any subset $S \subseteq [n]$ of size $k + 1$, we have:

$$(\forall i \in S). \quad v_{\text{suff}}^g(S) > v_{\text{suff}}^g(S \setminus \{i\}) \quad (15)$$

We also deliberately set δ to satisfy:

$$v_{\text{suff}}^g(S \setminus \{i\}) < \delta \leq v_{\text{suff}}^g(S), \quad (16)$$

and also set $k := \lfloor \frac{n}{2} \rfloor$. Each of these subsets is *subset-minimal* because any subset of size at most $\lfloor \frac{n}{2} \rfloor - 1$ fails to be a sufficient reason for $\langle f, \mathbf{x} \rangle$ (with respect to the δ threshold). Therefore, we can directly apply the same analysis as in (Bassan et al., 2024). In particular, there are exactly $\binom{n}{\lfloor \frac{n}{2} \rfloor}$ subset-minimal local sufficient reasons for $\langle f, \mathbf{x} \rangle$. Using Stirling’s approximation, we obtain:

$$\lim_{n \rightarrow \infty} \frac{2\sqrt{2\pi}}{e^2} \cdot \frac{2^n}{\sqrt{n}} \leq \binom{n}{\lfloor \frac{n}{2} \rfloor} \leq \frac{e}{\pi} \cdot \frac{2^n}{\sqrt{n}} \quad (17)$$

This yields the corresponding bound on the number of subset-minimal global δ -sufficient reasons.

□

D PROOF OF PROPOSITION 2

Proposition 2. *While the local probabilistic setting (for any δ) lacks monotonicity — i.e., the value functions v_{con}^ℓ and v_{suff}^ℓ are non-monotonic (Arenas et al., 2022; Izza et al., 2023; Subercaseaux et al., 2025; Izza et al., 2024b) — in the global probabilistic setting (also for any δ), both value functions v_{con}^g and v_{suff}^g are monotonic non-decreasing.*

Proof. In this section, we prove the monotonicity property for both types of global explanations — sufficient and contrastive. We begin by establishing the property for global sufficiency. To build

intuition, we first focus on the simpler case of Boolean functions, then generalize the result to functions with discrete multi-valued input domains and multiple output classes. We further extend the result to continuous input domains and show that the monotonicity property holds for any well-defined classification function. Finally, we demonstrate how the same monotonicity guarantees apply to global contrastive explanations. Importantly, none of our proofs rely on the assumption of feature independence — the monotonicity property holds strongly for *any* underlying distribution.

Lastly, we present a simple example showing that the monotonicity property does not hold for the local probabilistic sufficient or contrastive value functions.

Lemma 1. *The global sufficient value function v_{suff}^g is monotonic non-decreasing for Boolean functions, under any data distribution.*

Proof. Given an arbitrary set $S \in [n]$, and a feature $i \notin S$, we have that:

$$\Pr(\mathbf{x}) = \Pr(\mathbf{x}_S, \mathbf{x}_{\bar{S}}) = \Pr(\mathbf{x}_{\bar{S}} \mid \mathbf{x}_S) \cdot \Pr(\mathbf{x}_S) \quad (18)$$

Consequently, this implies that:

$$\begin{aligned} \Pr(\mathbf{x}) &= \Pr(\mathbf{x}_{\overline{S \cup \{i\}}} \mid \mathbf{x}_{S \cup \{i\}}) \cdot \Pr(\mathbf{x}_{S \cup \{i\}}) \\ &= \Pr(\mathbf{x}_{\overline{S \cup \{i\}}} \mid \mathbf{x}_{S \cup \{i\}}) \cdot \Pr(\mathbf{x}_S, x_i) \\ &= \Pr(\mathbf{x}_{\overline{S \cup \{i\}}} \mid \mathbf{x}_{S \cup \{i\}}) \cdot \Pr(x_i \mid \mathbf{x}_S) \cdot \Pr(\mathbf{x}_S) \end{aligned} \quad (19)$$

Moreover, for any two points $\mathbf{x}, \mathbf{x}' \sim \mathcal{D}$ such that $\mathbf{x}_S = \mathbf{x}'_S$ and $f(\mathbf{x}) = 1 - f(\mathbf{x}')$, it holds that:

$$\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S) = 1 - \Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}') \mid \mathbf{z}_S = \mathbf{x}'_S) \quad (20)$$

To simplify notation, we occasionally omit the distribution notation $\sim \mathcal{D}$ in the local probability expression $\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)$, and use f^+ and f^- to denote the events $f(\mathbf{z}) = 1$ and $f(\mathbf{z}) = 0$, respectively.

We begin with a technical simplification of the probability term to facilitate the proof, and then proceed to establish monotonicity.

Simplifying the probability term. Let $\mathcal{D}_S$ denote the distribution restricted to the set S . Given a point $\mathbf{x} \sim \mathcal{D}$ and a set S , consider all points $\mathbf{z} \sim \mathcal{D}$ such that $\mathbf{z}_S = \mathbf{x}_S$ and $f(\mathbf{z}) = b$, where $b \in \{0, 1\}$ — they all share the same local probability $\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = b \mid \mathbf{z}_S = \mathbf{x}_S)$. For every $\mathbf{x} \sim \mathcal{D}$ that shares the same x_S and the same output b , we marginalize over $x_{\bar{S}}$. This yields $\Pr(f(\mathbf{x}) = b, \mathbf{x}_S) = \Pr(f(\mathbf{x}) = b \mid \mathbf{x}_S) \cdot \Pr(\mathbf{x}_S)$, from which we can infer:

$$\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = b \mid \mathbf{z}_S = \mathbf{x}_S) = \Pr(f(\mathbf{x}) = b \mid \mathbf{x}_S). \quad (21)$$

Thus, for $\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)]$, we get that:

$$\begin{aligned} &\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)] \\ &= \sum_{\mathbf{x}_S \sim \mathcal{D}_S, f^+} \Pr(f^+, \mathbf{x}_S) \cdot \Pr_{\mathbf{z} \sim \mathcal{D}}(f^+ \mid \mathbf{z}_S = \mathbf{x}_S) + \sum_{\mathbf{x}_S \sim \mathcal{D}_S, f^-} \Pr(f^-, \mathbf{x}_S) \cdot \Pr_{\mathbf{z} \sim \mathcal{D}}(f^- \mid \mathbf{z}_S = \mathbf{x}_S) \\ &= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot [(\Pr_{\mathbf{z} \sim \mathcal{D}}(f^+ \mid \mathbf{z}_S = \mathbf{x}_S))^2 + (\Pr_{\mathbf{z} \sim \mathcal{D}}(f^- \mid \mathbf{z}_S = \mathbf{x}_S))^2] \\ &= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot [(\Pr_{\mathbf{z} \sim \mathcal{D}}(f^+ \mid \mathbf{z}_S = \mathbf{x}_S))^2 + (1 - \Pr_{\mathbf{z} \sim \mathcal{D}}(f^+ \mid \mathbf{z}_S = \mathbf{x}_S))^2], \end{aligned} \quad (22)$$

From which it follows that:

$$\begin{aligned}
& \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_{S \cup \{t\}} = \mathbf{x}_{S \cup \{t\}})] \\
&= \sum_{\mathbf{x}_{S \cup \{t\}} \sim \mathcal{D}_{S \cup \{t\}}} \Pr(\mathbf{x}_{S \cup \{t\}}) \cdot [(\Pr_{\mathbf{z}}(f^+ \mid \mathbf{z}_{S \cup \{t\}} = \mathbf{x}_{S \cup \{t\}}))^2 + (1 - \Pr_{\mathbf{z}}(f^+ \mid \mathbf{z}_{S \cup \{t\}} = \mathbf{x}_{S \cup \{t\}}))^2] \\
&= \sum_{\substack{\mathbf{x}_S \sim \mathcal{D}_S \\ x_t=1}} \Pr(x_t = 1, \mathbf{x}_S) \cdot [(\Pr_{\mathbf{z}}(f^+ \mid \mathbf{z}_S = \mathbf{x}_S, z_t = 1))^2 + (1 - \Pr_{\mathbf{z}}(f^+ \mid \mathbf{z}_S = \mathbf{x}_S, z_t = 1))^2] + \\
&\quad \sum_{\substack{\mathbf{x}_S \sim \mathcal{D}_S \\ x_t=0}} \Pr(x_t = 0, \mathbf{x}_S) \cdot [(\Pr_{\mathbf{z}}(f^+ \mid \mathbf{z}_S = \mathbf{x}_S, z_t = 0))^2 + (1 - \Pr_{\mathbf{z}}(f^+ \mid \mathbf{z}_S = \mathbf{x}_S, z_t = 0))^2]
\end{aligned} \tag{23}$$

Proving monotonicity. We now proceed to prove the monotonicity claim, building on the previous simplification. We begin by introducing a few additional notations. Let $g_1^+ := \Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 1 \mid \mathbf{z}_S = \mathbf{x}_S, z_t = 1)$ and $g_0^+ := \Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 1 \mid \mathbf{z}_S = \mathbf{x}_S, z_t = 0)$. Similarly, define $P_1 := \Pr(z_t = 1 \mid \mathbf{x}_S)$ and $P_0 := \Pr(z_t = 0 \mid \mathbf{x}_S)$. Then, the following expectations $\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)]$ and $\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_{S \cup \{t\}} = \mathbf{x}_{S \cup \{t\}})]$ can be simplified accordingly¹:

$$\begin{aligned}
& \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)] \\
&= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot [(P_1 \cdot g_1^+ + P_0 \cdot g_0^+)^2 + (1 - (P_1 \cdot g_1^+ + P_0 \cdot g_0^+))^2] \\
&= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot [2(P_1 \cdot g_1^+ + P_0 \cdot g_0^+)^2 - 2(P_1 \cdot g_1^+ + P_0 \cdot g_0^+) + 1].
\end{aligned} \tag{24}$$

As a result, we obtain the following:

$$\begin{aligned}
& \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_{S \cup \{t\}} = \mathbf{x}_{S \cup \{t\}})] \\
&= \sum_{\substack{\mathbf{x}_S \sim \mathcal{D}_S \\ x_t=1}} \Pr(\mathbf{x}_S) \cdot P_1 \cdot ((g_1^+)^2 + (1 - g_1^+)^2) + \sum_{\substack{\mathbf{x}_S \sim \mathcal{D}_S \\ x_t=0}} \Pr(\mathbf{x}_S) \cdot P_0 \cdot ((g_0^+)^2 + (1 - g_0^+)^2) \\
&= \sum_{\substack{\mathbf{x}_S \sim \mathcal{D}_S \\ x_t=1}} \Pr(\mathbf{x}_S) \cdot P_1 \cdot (2(g_1^+)^2 - 2g_1^+ + 1) + \sum_{\substack{\mathbf{x}_S \sim \mathcal{D}_S \\ x_t=0}} \Pr(\mathbf{x}_S) \cdot P_0 \cdot (2(g_0^+)^2 - 2g_0^+ + 1) \\
&= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot [P_1 \cdot (2(g_1^+)^2 - 2g_1^+ + 1) + P_0 \cdot (2(g_0^+)^2 - 2g_0^+ + 1)] \\
&= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot [2P_1 \cdot ((g_1^+)^2 - g_1^+) + 2P_0 \cdot ((g_0^+)^2 - g_0^+) + 1] \\
&= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot [2(P_1 \cdot (g_1^+)^2 + P_0 \cdot (g_0^+)^2) - 2(P_1 \cdot g_1^+ + P_0 \cdot g_0^+) + 1].
\end{aligned} \tag{25}$$

Let $\Delta_t(S, f)$ denote $\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_{S \cup \{t\}} = \mathbf{x}_{S \cup \{t\}})] - \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)]$. Without loss of generality, we assume that both P_1 and P_0 are positive. Otherwise, it can be verified that $\Delta_t(S, f) = 0$. Then, we get:

¹Note that the notations P_1 , P_0 , g_1^+ , and g_0^+ —as well as all others used throughout the paper—are not fixed constants, but rather depend on the specific input $\mathbf{x}_S$. Formally, they should be written as $P_1(\mathbf{x}_S)$, $P_0(\mathbf{x}_S)$, and so on. However, for readability, we omit the explicit dependence on $\mathbf{x}_S$.

$$\begin{aligned}
\Delta_t(S, f) &= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot [2(P_1 \cdot (g_1^+)^2 + P_0 \cdot (g_0^+)^2) - 2(P_1 \cdot g_1^+ + P_0 \cdot g_0^+)^2] \\
&= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} 2 \cdot \Pr(\mathbf{x}_S) \cdot [(P_1 - P_1^2)(g_1^+)^2 + (P_0 - P_0^2)(g_0^+)^2 - 2 \cdot P_1 \cdot P_0 \cdot g_1^+ \cdot g_0^+] \\
&= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} 2 \cdot \Pr(\mathbf{x}_S) \cdot P_1 \cdot P_0 \cdot (g_1^+ - g_0^+)^2 \\
&\geq 0.
\end{aligned} \tag{26}$$

Note that $P_1 - P_1^2 = P_1(1 - P_1) = P_1 P_0$ and similarly: $P_0 - P_0^2 = P_0(1 - P_0) = P_0 P_1$. Using these identities along with the previous result, we conclude that $\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)]$ is monotone, thereby concluding our proof. $\square$

Lemma 2. *The global sufficient value function v_{suff}^g is monotonic non-decreasing for functions with discrete multi-valued input domains and multiple output classes, under any data distribution.*

Proof. Extending the proof from the simplified Boolean case, we now generalize the result to functions whose input domains and output ranges each consist of a finite set of values. Let $\mathcal{K}$ represent the set of output classes.

We can generalize equation 22 as follows:

$$\begin{aligned}
&\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)] = \\
&\sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot \left[\sum_{j \in \mathcal{K}} (\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = j \mid \mathbf{z}_S = \mathbf{x}_S))^2 \right]
\end{aligned} \tag{27}$$

Given the condition $\mathbf{z}_S = \mathbf{x}_S$, suppose feature t can take r possible values $\{v_1, \dots, v_r\}$. Then, equation 23 can be generalized as follows:

$$\begin{aligned}
&\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_{S \cup \{t\}} = \mathbf{x}_{S \cup \{t\}})] \\
&= \sum_{k=1}^r \left[\sum_{\substack{\mathbf{x}_S \sim \mathcal{D}_S \\ x_t = v_k}} \Pr(\mathbf{x}_S) \cdot \Pr(z_t = v_k \mid \mathbf{x}_S) \cdot \left[\sum_{j \in \mathcal{K}} (\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = j \mid \mathbf{z}_S = \mathbf{x}_S, z_t = v_k))^2 \right] \right].
\end{aligned} \tag{28}$$

Let P_k denote $\Pr(z_t = v_k \mid \mathbf{x}_S)$, Let g_k^j denote $\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = j \mid \mathbf{z}_S = \mathbf{x}_S, z_t = v_k)$, that is, z_t takes the value v_k and the output class is j . Then, we have

$$\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)] = \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot \left[\sum_{j \in \mathcal{K}} \left(\sum_{k=1}^r P_k \cdot g_k^j \right)^2 \right] \tag{29}$$

and

$$\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_{S \cup \{t\}} = \mathbf{x}_{S \cup \{t\}})] = \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot \left[\sum_{k=1}^r P_k \cdot \left(\sum_{j \in \mathcal{K}} (g_k^j)^2 \right) \right]$$

Combining these two implications, we can conclude that:

$$\begin{aligned}
\Delta_t(S, f) &= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot \left[\sum_{k=1}^r P_k \cdot \left(\sum_{j \in \mathcal{K}} (g_k^j)^2 \right) - \sum_{j \in \mathcal{K}} \left(\sum_{k=1}^r P_k \cdot g_k^j \right)^2 \right] \\
&= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot \left[\sum_{j \in \mathcal{K}} \left(\sum_{k=1}^r P_k \cdot (g_k^j)^2 \right) - \sum_{j \in \mathcal{K}} \left(\sum_{k=1}^r (P_k \cdot g_k^j)^2 + \sum_{k < l} 2 \cdot P_k \cdot P_l \cdot g_k^j \cdot g_l^j \right) \right] \\
&= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot \left[\sum_{j \in \mathcal{K}} \sum_{k \neq l} P_k \cdot P_l \cdot (g_k^j)^2 - \sum_{j \in \mathcal{K}} \sum_{k < l} 2 \cdot P_k \cdot P_l \cdot g_k^j \cdot g_l^j \right] \\
&= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot \sum_{j \in \mathcal{K}} \left[\sum_{k < l} P_k \cdot P_l \cdot (g_k^j - g_l^j)^2 \right] \\
&\geq 0.
\end{aligned} \tag{30}$$

Hence, $\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)]$ is monotone, which completes the proof for this setting. $\square$

Lemma 3. *The global sufficient value function v_{uff}^g is monotonic non-decreasing for functions with continuous input domains and multiple output classes, under any data distribution.*

Proof. We extend the monotonicity results to functions where features are defined over continuous domains. One approach is to partition the domain of the continuous features into two disjoint subdomains and then adapt the proofs in lemma 2 by revising the notation accordingly. The core reasoning remains unchanged when transitioning from the discrete to the continuous case—only the notation differs.

Alternatively, we provide a direct proof below. Assume all features lie in continuous domains, and let $p(\mathbf{x}_S, \mathbf{x}_{\bar{S}})$ denote the corresponding joint probability density function. Moreover, let $p(x_i)$ denote the marginal probability density function of x_i , and $\mathcal{X}_i$ denote the domain of feature i .

We define marginal density as

$$p(\mathbf{x}_S) := \int_{\mathbf{x}_{\bar{S}} \in \mathcal{X}_{\bar{S}}} p(\mathbf{x}_S, \mathbf{x}_{\bar{S}}) d\mathbf{x}_{\bar{S}}, \tag{31}$$

where $\mathcal{X}_{\bar{S}}$ represents the domain of features in $\bar{S}$.

And the conditional density is defined as:

$$p(x_t \mid \mathbf{x}_S) := \frac{p(\mathbf{x}_S, x_t)}{p(\mathbf{x}_S)}. \tag{32}$$

The local probability of S given $\mathbf{x}$ is defined as

$$\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = j \mid \mathbf{z}_S = \mathbf{x}_S) = \frac{\int_{\mathbf{z}_{\bar{S}} \in \mathcal{X}_{\bar{S}}} I(f(\mathbf{x}_S, \mathbf{z}_{\bar{S}}) = j) \cdot p(\mathbf{x}_S, \mathbf{z}_{\bar{S}}) d\mathbf{z}_{\bar{S}}}{p(\mathbf{x}_S)}, \tag{33}$$

This term can also be decomposed into:

$$\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = j \mid \mathbf{z}_S = \mathbf{x}_S) = \int_{z_t \in \mathcal{X}_t} p(z_t \mid \mathbf{x}_S) \cdot \Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = j \mid \mathbf{z}_S = \mathbf{x}_S, z_t = x_t) dz_t \tag{34}$$

via an additional feature $t \notin S$.

We then modify equation 27 to

$$\begin{aligned}
&\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)] \\
&= \int_{\mathbf{x}_S \in \mathcal{X}_S} p(\mathbf{x}_S) \cdot \left[\sum_{j \in \mathcal{K}} (\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = j \mid \mathbf{z}_S = \mathbf{x}_S))^2 \right] d\mathbf{x}_S,
\end{aligned} \tag{35}$$

while equation 28 is changed to

$$\begin{aligned} & \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_{S \cup \{t\}} = \mathbf{x}_{S \cup \{t\}})] \\ &= \int_{\mathbf{x}_S \in \mathcal{X}_S} p(\mathbf{x}_S) \cdot \left[\int_{z_t \in \mathcal{X}_t} p(z_t \mid \mathbf{x}_S) \cdot \left(\sum_{j \in \mathcal{K}} (\mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = j \mid \mathbf{z}_{S \cup \{t\}} = \mathbf{x}_{S \cup \{t\}}))^2 \right) dz_t \right] d\mathbf{x}_S. \end{aligned} \quad (36)$$

For simplicity, let $Y_j(t)$ denote $\mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = j \mid \mathbf{z}_S = \mathbf{x}_S, z_t = v_t)$, and $p_{t|S}$ denote $p(z_t \mid \mathbf{x}_S)$. Note that both $Y_j(t)$ and $p_{t|S}$ are associated with a fixed $\mathbf{x}_S$ ². We prove monotonicity as follows:

$$\begin{aligned} & \Delta_t(S, f) \\ &= \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_{S \cup \{t\}} = \mathbf{x}_{S \cup \{t\}})] - \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)] \\ &= \int_{\mathbf{x}_S \in \mathcal{X}_S} p(\mathbf{x}_S) \cdot \left[\int_{z_t \in \mathcal{X}_t} p_{t|S} \cdot \left(\sum_{j \in \mathcal{K}} (\mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = j \mid \mathbf{z}_{S \cup \{t\}} = \mathbf{x}_{S \cup \{t\}}))^2 \right) dz_t \right] d\mathbf{x}_S - \\ & \quad \int_{\mathbf{x}_S \in \mathcal{X}_S} p(\mathbf{x}_S) \cdot \left[\sum_{j \in \mathcal{K}} \left(\int_{z_t \in \mathcal{X}_t} p_{t|S} \cdot \mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = j \mid \mathbf{z}_S = \mathbf{x}_S, z_t = x_t) dz_t \right)^2 \right] d\mathbf{x}_S \\ &= \int_{\mathbf{x}_S \in \mathcal{X}_S} p(\mathbf{x}_S) \cdot \left[\sum_{j \in \mathcal{K}} \left(\int_{\mathcal{X}_t} p_{t|S} \cdot Y_j(t)^2 dz_t - \left(\int_{\mathcal{X}_t} p_{t|S} \cdot Y_j(t) dz_t \right)^2 \right) \right] d\mathbf{x}_S \\ &= \int_{\mathbf{x}_S \in \mathcal{X}_S} p(\mathbf{x}_S) \cdot \left[\sum_{j \in \mathcal{K}} \frac{2}{2} \cdot \left(\int_{\mathcal{X}_t} \int_{\mathcal{X}_{t'}} p_{t|S} \cdot p_{t'|S} \cdot (Y_j(t)^2 - Y_j(t)Y_j(t')) dz_t dz_{t'} \right) \right] d\mathbf{x}_S \\ &= \int_{\mathbf{x}_S \in \mathcal{X}_S} p(\mathbf{x}_S) \cdot \left[\sum_{j \in \mathcal{K}} \frac{1}{2} \cdot \left(\int_{\mathcal{X}_t} \int_{\mathcal{X}_{t'}} p_{t|S} \cdot p_{t'|S} \cdot (Y_j(t) - Y_j(t'))^2 dz_t dz_{t'} \right) \right] d\mathbf{x}_S \\ &\geq 0. \end{aligned} \quad (37)$$

□

Proposition 2. While the local probabilistic setting (for any δ) lacks monotonicity — i.e., the value functions v_{con}^ℓ and v_{suff}^ℓ are non-monotonic (Arenas et al., 2022; Izza et al., 2023; Subercaseaux et al., 2025; Izza et al., 2024b) — in the global probabilistic setting (also for any δ), both value functions v_{con}^g and v_{suff}^g are monotonic non-decreasing.

Proof. As a consequence of lemmas 1 to 3, we conclude that the global sufficient value function v_{suff}^g is monotonic non-decreasing for any well-defined classification function — whether its features are Boolean, discrete, continuous, or a combination — under any data distribution.

$$v_{\text{con}}^g(S) = 1 - v_{\text{suff}}^g(\bar{S}) \quad (38)$$

This implies that v_{con}^g is monotonic non-increasing with respect to adding features to $\bar{S}$. Equivalently, v_{con}^g is monotonic non-decreasing as features are removed from $\bar{S}$ and added to S , thus completing the proof.

□

E PROOF OF PROPOSITION 3

Proposition 3. While the local probabilistic sufficient setting (for any δ) lacks supermodularity — even when $\mathcal{D}$ is uniform, i.e., the value function v_{suff}^ℓ is not supermodular — in the global probabilistic setting (also for any δ), when $\mathcal{D}$ exhibits feature independence, the value function v_{suff}^g is supermodular.

Proof. This section establishes the supermodularity property for both global sufficient and contrastive explanations. As with monotonicity, we begin with Boolean functions and progressively generalize

²Changing $\mathbf{x}_S$ alters their values, we omit $\mathbf{x}_S$ in the notation for brevity.

to multi-valued discrete and continuous input domains, and finally consider the case of well-defined classification functions. Unlike the monotonicity property, however, these proofs require the assumption of feature independence. We also provide a counterexample to show that supermodularity may fail when this assumption is relaxed.

Lemma 4. *The global sufficient value function v_{uff}^g is supermodular for Boolean functions, under the assumption of feature independence.*

Proof. Particularly, we will prove that:

$$\Delta_t(S, f) \leq \Delta_t(S', f), \quad (39)$$

where $S \subseteq S'$ and $t \notin S'$. That is, $\Delta_t(S, f)$ is monotone.

Clearly, if $S' = S$, then $\Delta_t(S, f) \leq \Delta_t(S', f)$ holds. Let $S' = S \cup \{i\}$, where $i \notin S$ and $i \neq t$. For the local probability $\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)$, under the assumption of feature independence, we have that the following holds:

$$\begin{aligned} \Pr_{\mathbf{z}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S) \\ = \Pr(z_i = 1) \cdot \Pr(z_t = 1) \cdot \Pr_{\mathbf{z}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S, z_i = 1, z_t = 1) + \\ \Pr(z_i = 1) \cdot \Pr(z_t = 0) \cdot \Pr_{\mathbf{z}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S, z_i = 1, z_t = 0) + \\ \Pr(z_i = 0) \cdot \Pr(z_t = 1) \cdot \Pr_{\mathbf{z}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S, z_i = 0, z_t = 1) + \\ \Pr(z_i = 0) \cdot \Pr(z_t = 0) \cdot \Pr_{\mathbf{z}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S, z_i = 0, z_t = 0). \end{aligned} \quad (40)$$

Let P_1, P_0, Q_1 , and Q_0 denote $\Pr(z_t = 1), \Pr(z_t = 0), \Pr(z_i = 1)$, and $\Pr(z_i = 0)$, respectively. Moreover, let $g_{11}^+, g_{10}^+, g_{01}^+$, and g_{00}^+ denote $\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 1 \mid \mathbf{z}_S = \mathbf{x}_S, z_i = 1, z_t = 1), \Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 1 \mid \mathbf{z}_S = \mathbf{x}_S, z_i = 1, z_t = 0), \Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 1 \mid \mathbf{z}_S = \mathbf{x}_S, z_i = 0, z_t = 1)$, and $\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 1 \mid \mathbf{z}_S = \mathbf{x}_S, z_i = 0, z_t = 0)$, respectively.

We further decompose equation 26 into:

$$\begin{aligned} \Delta_t(S, f) &= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} 2 \cdot \Pr(\mathbf{x}_S) \cdot P_1 P_0 \cdot [(Q_1 \cdot g_{11}^+ + Q_0 \cdot g_{01}^+) - (Q_1 \cdot g_{10}^+ + Q_0 \cdot g_{00}^+)]^2 \\ &= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} 2 \cdot \Pr(\mathbf{x}_S) \cdot P_1 P_0 \cdot [Q_1 \cdot (g_{11}^+ - g_{10}^+) + Q_0 \cdot (g_{01}^+ - g_{00}^+)]^2. \end{aligned} \quad (41)$$

Similarly, the following also holds:

$$\Delta_t(S \cup \{i\}, f) = \sum_{\mathbf{x}_S \sim \mathcal{D}_S} 2 \cdot \Pr(\mathbf{x}_S) \cdot P_1 P_0 \cdot [Q_1 \cdot (g_{11}^+ - g_{10}^+)^2 + Q_0 \cdot (g_{01}^+ - g_{00}^+)^2]. \quad (42)$$

Thus, we get

$$\begin{aligned} \Delta_t(S \cup \{i\}, f) - \Delta_t(S, f) &= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} 2 \cdot \Pr(\mathbf{x}_S) \cdot P_1 P_0 \cdot Q_1 Q_0 \cdot (g_{11}^+ - g_{10}^+ - g_{01}^+ + g_{00}^+)^2 \\ &\geq 0. \end{aligned} \quad (43)$$

This implies that $\Delta_t(S, f)$ is monotone, that is, $c(S \cup \{t\}, f) - c(S, f) \leq c(S' \cup \{t\}, f) - c(S', f)$. $\square$

Lemma 5. *The global sufficient value function v_{uff}^g is supermodular for functions with discrete multi-valued input domains and multiple output classes, under the assumption of feature independence.*

Proof. We extend the results of the previous section to multi-valued discrete functions, again with the assumption of feature independence. We use P_k to denote $\Pr(z_t = v_k)$, and $Q_{k'}$ to denote $\Pr(z_i = v_{k'})$. Moreover, we use $g_{k'k}^j$ to denote $\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = j \mid \mathbf{z}_S = \mathbf{x}_S, z_i = v_{k'}, z_t = v_k)$. We

decompose equation 30 to

$$\begin{aligned}\Delta_t(S, f) &= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot \sum_{j \in \mathcal{K}} \left[\sum_{k < l} P_k P_l \cdot \left(\sum_{k'=1}^{r'} Q_{k'} \cdot g_{k'k}^j - \sum_{k'=1}^{r'} Q_{k'} \cdot g_{k'l}^j \right)^2 \right] \\ &= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot \sum_{j \in \mathcal{K}} \left[\sum_{k < l} P_k P_l \cdot \left(\sum_{k'=1}^{r'} Q_{k'} \cdot (g_{k'k}^j - g_{k'l}^j) \right)^2 \right].\end{aligned}\quad (44)$$

Likewise, we have

$$\begin{aligned}\Delta_t(S \cup \{i\}, f) &= \sum_{k'=1}^{r'} \left[\sum_{\substack{\mathbf{x}_S \sim \mathcal{D}_S \\ x_i = v_k}} \Pr(\mathbf{x}_S) \cdot \Pr(z_i = v_k) \cdot \sum_{j \in \mathcal{K}} \left(\sum_{k < l} P_k P_l \cdot (g_{k'k}^j - g_{k'l}^j)^2 \right) \right] \\ &= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot \left[\sum_{k'=1}^{r'} Q_{k'} \cdot \sum_{j \in \mathcal{K}} \left(\sum_{k < l} P_k P_l \cdot (g_{k'k}^j - g_{k'l}^j)^2 \right) \right] \\ &= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot \sum_{j \in \mathcal{K}} \left[\sum_{k < l} P_k P_l \cdot \sum_{k'=1}^{r'} Q_{k'} \cdot (g_{k'k}^j - g_{k'l}^j)^2 \right].\end{aligned}\quad (45)$$

Therefore, we get

$$\begin{aligned}\Delta_t(S \cup \{i\}, f) - \Delta_t(S, f) &= \sum_{\mathbf{x}_S \sim \mathcal{D}_S} \Pr(\mathbf{x}_S) \cdot \sum_{j \in \mathcal{K}} \left[\sum_{k < l} \sum_{k' < l'} P_k P_l Q_{k'} Q_{l'} \cdot (g_{k'k}^j - g_{k'l}^j - g_{l'k}^j + g_{l'l}^j)^2 \right] \\ &\geq 0.\end{aligned}\quad (46)$$

□

Proposition 3. *While the local probabilistic sufficient setting (for any δ) lacks supermodularity — even when $\mathcal{D}$ is uniform, i.e., the value function v_{diff}^ℓ is not supermodular — in the global probabilistic setting (also for any δ), when $\mathcal{D}$ exhibits feature independence, the value function v_{diff}^g is supermodular.*

Proof. We can generalize the supermodularity results established in Lemma 5 to the continuous case, assuming that $\mathcal{D}_p$ satisfies feature independence. One approach is to partition the domain of the continuous features into two disjoint subdomains and then adapt the proofs in Lemma 5 by revising the notation accordingly.

Alternatively, we provide a direct proof below. Under this assumption, the marginal and conditional densities simplify as follows:

$$p(\mathbf{x}_S) = \prod_{i \in S} p(x_i), \text{ and } p(x_t | \mathbf{x}_S) = p(x_t). \quad (47)$$

Consequently, the local probability of S conditioned on $\mathbf{x}$ becomes:

$$\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = j | \mathbf{z}_S = \mathbf{x}_S) = \int_{\mathbf{z}_{\bar{S}} \in \mathcal{X}_{\bar{S}}} I(f(\mathbf{x}_S, \mathbf{z}_{\bar{S}}) = j) \cdot \prod_{i \in \bar{S}} p(z_i) d\mathbf{z}_{\bar{S}}. \quad (48)$$

We let $Y_j(i, t)$ denote $\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = j | \mathbf{z}_S = \mathbf{x}_S, z_i = x_i, z_t = x_t)$, where $i \notin S, t \notin S$ and $i \neq t$. Evidently,

$$Y_j(i, t) = Y_j(t, i), \text{ and } Y_j(t) = \int_{\mathcal{X}_i} p_i \cdot Y_j(i, t) dz_i.$$

We have:

$$\begin{aligned}
\Delta_t(S, f) &= \int_{\mathbf{x}_S \in \mathcal{X}_S} p(\mathbf{x}_S) \cdot \left[\sum_{j \in \mathcal{K}} \frac{1}{2} \left(\int_{\mathcal{X}_t} \int_{\mathcal{X}_{t'}} p_t p_{t'} (Y_j(t) - Y_j(t'))^2 dz_t dz_{t'} \right) \right] d\mathbf{x}_S \\
&= \int_{\mathbf{x}_S \in \mathcal{X}_S} p(\mathbf{x}_S) \cdot \left[\sum_{j \in \mathcal{K}} \frac{1}{2} \left(\int_{\mathcal{X}_t} \int_{\mathcal{X}_{t'}} p_t p_{t'} \left(\int_{\mathcal{X}_i} p_i \cdot Y_j(i, t) dz_i - \int_{\mathcal{X}_i} p_i \cdot Y_j(i, t') dz_i \right)^2 dz_t dz_{t'} \right) \right] d\mathbf{x}_S \\
&= \int_{\mathbf{x}_S \in \mathcal{X}_S} p(\mathbf{x}_S) \cdot \left[\sum_{j \in \mathcal{K}} \frac{1}{2} \left(\int_{\mathcal{X}_t} \int_{\mathcal{X}_{t'}} p_t p_{t'} \left(\int_{\mathcal{X}_i} p_i \cdot (Y_j(i, t) - Y_j(i, t')) dz_i \right)^2 dz_t dz_{t'} \right) \right] d\mathbf{x}_S,
\end{aligned} \tag{49}$$

where

$$\begin{aligned}
&\int_{\mathcal{X}_t} \int_{\mathcal{X}_{t'}} p_t p_{t'} \left(\int_{\mathcal{X}_i} p_i \cdot (Y_j(i, t) - Y_j(i, t')) dz_i \right)^2 dz_t dz_{t'} \\
&= \int_{\mathcal{X}_t} \int_{\mathcal{X}_{t'}} \int_{\mathcal{X}_i} \int_{\mathcal{X}_{i'}} p_t p_{t'} p_i p_{i'} \cdot (Y_j(i, t) - Y_j(i, t')) \cdot (Y_j(i', t) - Y_j(i', t')) dz_i dz_{i'} dz_t dz_{t'}.
\end{aligned}$$

Let $S' = S \cup \{i\}$, then we have:

$$\begin{aligned}
\Delta_t(S', f) &= \int_{\mathbf{x}_{S'} \in \mathcal{X}_{S'}} p(\mathbf{x}_{S'}) \cdot \left[\sum_{j \in \mathcal{K}} \frac{1}{2} \left(\int_{\mathcal{X}_t} \int_{\mathcal{X}_{t'}} p_t p_{t'} \cdot (Y_j(i, t) - Y_j(i, t'))^2 dz_t dz_{t'} \right) \right] d\mathbf{x}_{S'} \\
&= \int_{\mathbf{x}_S \in \mathcal{X}_S} p(\mathbf{x}_S) \cdot \left[\sum_{j \in \mathcal{K}} \frac{1}{2} \int_{\mathcal{X}_i} p_i \cdot \left(\int_{\mathcal{X}_t} \int_{\mathcal{X}_{t'}} p_t p_{t'} \cdot (Y_j(i, t) - Y_j(i, t'))^2 dz_t dz_{t'} \right) dz_i \right] d\mathbf{x}_S \tag{50} \\
&= \int_{\mathbf{x}_S \in \mathcal{X}_S} p(\mathbf{x}_S) \cdot \left[\sum_{j \in \mathcal{K}} \frac{1}{2} \left(\int_{\mathcal{X}_i} \int_{\mathcal{X}_t} \int_{\mathcal{X}_{t'}} p_i p_t p_{t'} \cdot (Y_j(i, t) - Y_j(i, t'))^2 dz_t dz_{t'} dz_i \right) \right] d\mathbf{x}_S,
\end{aligned}$$

where

$$\begin{aligned}
&\int_{\mathcal{X}_i} \int_{\mathcal{X}_t} \int_{\mathcal{X}_{t'}} p_i p_t p_{t'} \cdot (Y_j(i, t) - Y_j(i, t'))^2 dz_t dz_{t'} dz_i \\
&= \int_{\mathcal{X}_i} \int_{\mathcal{X}_{i'}} \int_{\mathcal{X}_t} \int_{\mathcal{X}_{t'}} p_i p_{i'} p_t p_{t'} \cdot (Y_j(i, t) - Y_j(i, t'))^2 dz_t dz_{t'} dz_i dz_{i'}.
\end{aligned}$$

Let $\alpha(i, t, t')$ denote $Y_j(i, t) - Y_j(i, t')$ and $\beta(i', t, t')$ denote $Y_j(i', t) - Y_j(i', t')$, we can infer that

$$\begin{aligned}
\Delta_t(S', f) - \Delta_t(S, f) &= \int_{\mathbf{x}_S \in \mathcal{X}_S} p(\mathbf{x}_S) \left[\sum_{j \in \mathcal{K}} \frac{1}{4} \left(\int_{\mathcal{X}_t} \int_{\mathcal{X}_{t'}} \int_{\mathcal{X}_i} \int_{\mathcal{X}_{i'}} p_t p_{t'} p_i p_{i'} (\alpha(i, t, t') - \beta(i', t, t'))^2 dz_i dz_{i'} dz_t dz_{t'} \right) \right] d\mathbf{x}_S \\
&\geq 0.
\end{aligned} \tag{51}$$

Consequently, the global sufficient value function v_{suff}^g retains the supermodularity property for any valid classification function, regardless of whether its features are Boolean, discrete, continuous, or a mix thereof.

We now demonstrate that the local probabilistic sufficient value function is neither submodular nor supermodular. As a simple illustration, fix the uniform distribution on the three variables $\{x_1, x_2, x_3\}$ and define a Boolean function as follows:

$$f(x_1, x_2) = (x_1 \vee x_2) \wedge x_3. \tag{52}$$

Consider the input $\mathbf{x} = (1, 1, 1)$ which is classified as 1. For the local probabilistic sufficient value function v_{suff}^l , we have that the following holds:

$$v_{\text{suff}}^l(\emptyset) = \frac{3}{8}, \quad v_{\text{suff}}^l(\{1\}) = \frac{1}{2}, \quad v_{\text{suff}}^l(\{2\}) = \frac{1}{2}, \quad v_{\text{suff}}^l(\{3\}) = \frac{3}{4},$$

$$v_{\text{suff}}^l(\{1, 2\}) = \frac{1}{2}, \quad v_{\text{suff}}^l(\{1, 3\}) = 1, \quad v_{\text{suff}}^l(\{2, 3\}) = 1, \quad v_{\text{suff}}^l(\{1, 2, 3\}) = 1.$$

Since $(v_{\text{suff}}^l(\{1, 2\}) - v_{\text{suff}}^l(\{2\})) < (v_{\text{suff}}^l(\{1\}) - v_{\text{suff}}^l(\emptyset)) < (v_{\text{suff}}^l(\{1, 3\}) - v_{\text{suff}}^l(\{3\}))$, the function v_{suff}^l is neither supermodular nor submodular.

□

F PROOF OF PROPOSITION 4

Proposition 4. *While the local probabilistic contrastive setting (for any δ) lacks submodularity — even when $\mathcal{D}$ is uniform, i.e., the value function v_{con}^l is not submodular — in the global probabilistic setting (also for any δ), when $\mathcal{D}$ exhibits feature independence, the value function v_{con}^g is submodular.*

Proof. Once we establish that the global sufficient setting is supermodular, it directly follows that the global contrastive setting is submodular. Let $S \subseteq [n]$ be the set of features allowed to change, and suppose $S \subseteq S'$ with $i \notin S'$. To compare the marginal contributions $v_{\text{con}}^g(S \cup \{i\}) - v_{\text{con}}^g(S)$ and $v_{\text{con}}^g(S' \cup \{i\}) - v_{\text{con}}^g(S')$, we equivalently compare $(1 - v_{\text{suff}}^g(\bar{S} \setminus \{i\})) - (1 - v_{\text{suff}}^g(\bar{S}))$ and $(1 - v_{\text{suff}}^g(\bar{S}' \setminus \{i\})) - (1 - v_{\text{suff}}^g(\bar{S}'))$. Since $S \subseteq S'$, it follows that $\bar{S}' \subseteq \bar{S}$, and therefore we obtain:

$$(v_{\text{suff}}^g(\bar{S}) - v_{\text{suff}}^g(\bar{S} \setminus \{i\})) \geq (v_{\text{suff}}^g(\bar{S}') - v_{\text{suff}}^g(\bar{S}' \setminus \{i\})) \quad (53)$$

Which is equivalent to:

$$(v_{\text{con}}^g(S \cup \{i\}) - v_{\text{con}}^g(S)) \geq (v_{\text{con}}^g(S' \cup \{i\}) - v_{\text{con}}^g(S')). \quad (54)$$

Thus, the global contrastive value function v_{con}^g is submodular, hence concluding this proof.

We will now demonstrate that the local probabilistic contrastive value function is neither submodular nor supermodular. As a simple illustration, fix the uniform distribution on the three variables $\{x_1, x_2, x_3\}$ and define a Boolean function as follows:

$$f(x_1, x_2) := (x_1 \vee x_2) \wedge x_3. \quad (55)$$

Consider the input $\mathbf{x} = (1, 1, 1)$ which is classified as 1. For the local probabilistic contrastive value function v_{con}^l , we have:

$$v_{\text{con}}^l(\emptyset) = 1, \quad v_{\text{con}}^l(\{1\}) = 1, \quad v_{\text{con}}^l(\{2\}) = 1, \quad v_{\text{con}}^l(\{3\}) = \frac{1}{2},$$

$$v_{\text{con}}^l(\{1, 2\}) = \frac{3}{4}, \quad v_{\text{con}}^l(\{1, 3\}) = \frac{1}{2}, \quad v_{\text{con}}^l(\{2, 3\}) = \frac{1}{2}, \quad v_{\text{con}}^l(\{1, 2, 3\}) = \frac{3}{8}.$$

Since $(v_{\text{con}}^l(\{1, 2\}) - v_{\text{con}}^l(\{2\})) < (v_{\text{con}}^l(\{1, 2, 3\}) - v_{\text{con}}^l(\{2, 3\})) < (v_{\text{con}}^l(\{1\}) - v_{\text{con}}^l(\emptyset))$, the function v_{con}^l is neither submodular nor supermodular.

□

In the rest of this section, we show that the feature independence assumption is not just sufficient but also necessary for supermodularity. We illustrate this with the following counterexample.

Define a Boolean function:

$$f(x_1, x_2) := x_1 \quad (56)$$

For this boolean function, it holds that $f(00) = 0$, $f(01) = 0$, $f(10) = 1$, and $f(11) = 1$. Moreover, define a distribution $\mathcal{D}$ on x_1 and x_2 . Recall that the global sufficient value function is:

$$\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)] = \sum_{\mathbf{x} \sim \mathcal{D}} \mathbf{Pr}(\mathbf{x}) \cdot \mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)$$

In addition, if we can show that

$$\Delta_t(S, f) \leq \Delta_t(S', f),$$

where $S \subseteq S'$ and $t \notin S'$, then the global sufficient value function is supermodular. When $S = \emptyset$, we have that the following holds:

$$\begin{aligned} & \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)] \\ &= \mathbf{Pr}(x_1 = 0, x_2 = 0) \cdot \mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 0) + \mathbf{Pr}(x_1 = 0, x_2 = 1) \cdot \mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 0) + \\ & \quad \mathbf{Pr}(x_1 = 1, x_2 = 0) \cdot \mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 1) + \mathbf{Pr}(x_1 = 1, x_2 = 1) \cdot \mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 1) \end{aligned}$$

Which is also equivalent to:

$$\begin{aligned} & \mathbf{Pr}(x_1 = 0, x_2 = 0) \cdot (\mathbf{Pr}(x_1 = 0, x_2 = 0) + \mathbf{Pr}(x_1 = 0, x_2 = 1)) + \\ & \quad \mathbf{Pr}(x_1 = 0, x_2 = 1) \cdot (\mathbf{Pr}(x_1 = 0, x_2 = 0) + \mathbf{Pr}(x_1 = 0, x_2 = 1)) + \\ & \quad \mathbf{Pr}(x_1 = 1, x_2 = 0) \cdot (\mathbf{Pr}(x_1 = 1, x_2 = 0) + \mathbf{Pr}(x_1 = 1, x_2 = 1)) + \\ & \quad \mathbf{Pr}(x_1 = 1, x_2 = 1) \cdot (\mathbf{Pr}(x_1 = 1, x_2 = 0) + \mathbf{Pr}(x_1 = 1, x_2 = 1)) \\ &= (\mathbf{Pr}(x_1 = 0, x_2 = 0) + \mathbf{Pr}(x_1 = 0, x_2 = 1))^2 + (\mathbf{Pr}(x_1 = 1, x_2 = 0) + \mathbf{Pr}(x_1 = 1, x_2 = 1))^2 \end{aligned}$$

When $S = \{2\}$, we have that the following condition holds:

$$\begin{aligned} & \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)] \\ &= \mathbf{Pr}(x_1 = 0, x_2 = 0) \cdot \mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 0 \mid x_2 = 0) + \\ & \quad \mathbf{Pr}(x_1 = 0, x_2 = 1) \cdot \mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 0 \mid x_2 = 1) + \\ & \quad \mathbf{Pr}(x_1 = 1, x_2 = 0) \cdot \mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 1 \mid x_2 = 0) + \\ & \quad \mathbf{Pr}(x_1 = 1, x_2 = 1) \cdot \mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 1 \mid x_2 = 1) \end{aligned}$$

This can also be rearranged as follows:

$$\begin{aligned} & \mathbf{Pr}(x_1 = 0, x_2 = 0) \cdot \frac{\mathbf{Pr}(x_1 = 0, x_2 = 0)}{\mathbf{Pr}(x_1 = 0, x_2 = 0) + \mathbf{Pr}(x_1 = 1, x_2 = 0)} + \\ & \quad \mathbf{Pr}(x_1 = 0, x_2 = 1) \cdot \frac{\mathbf{Pr}(x_1 = 0, x_2 = 1)}{\mathbf{Pr}(x_1 = 0, x_2 = 1) + \mathbf{Pr}(x_1 = 1, x_2 = 1)} + \\ & \quad \mathbf{Pr}(x_1 = 1, x_2 = 0) \cdot \frac{\mathbf{Pr}(x_1 = 1, x_2 = 0)}{\mathbf{Pr}(x_1 = 0, x_2 = 0) + \mathbf{Pr}(x_1 = 1, x_2 = 0)} + \\ & \quad \mathbf{Pr}(x_1 = 1, x_2 = 1) \cdot \frac{\mathbf{Pr}(x_1 = 1, x_2 = 1)}{\mathbf{Pr}(x_1 = 0, x_2 = 1) + \mathbf{Pr}(x_1 = 1, x_2 = 1)} \\ &= \frac{\mathbf{Pr}(x_1 = 0, x_2 = 0)^2 + \mathbf{Pr}(x_1 = 1, x_2 = 0)^2}{\mathbf{Pr}(x_1 = 0, x_2 = 0) + \mathbf{Pr}(x_1 = 1, x_2 = 0)} + \frac{\mathbf{Pr}(x_1 = 0, x_2 = 1)^2 + \mathbf{Pr}(x_1 = 1, x_2 = 1)^2}{\mathbf{Pr}(x_1 = 0, x_2 = 1) + \mathbf{Pr}(x_1 = 1, x_2 = 1)} \end{aligned}$$

In contrast, when $S = \{1\}$, we have the the following holds:

$$\begin{aligned} & \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)] \\ &= \mathbf{Pr}(x_1 = 0, x_2 = 0) \cdot \mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 0 \mid x_1 = 0) + \\ & \quad \mathbf{Pr}(x_1 = 0, x_2 = 1) \cdot \mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 0 \mid x_1 = 0) + \\ & \quad \mathbf{Pr}(x_1 = 1, x_2 = 0) \cdot \mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 1 \mid x_1 = 1) + \\ & \quad \mathbf{Pr}(x_1 = 1, x_2 = 1) \cdot \mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 1 \mid x_1 = 1) \end{aligned}$$

This can similarly be rearranged as follows:

$$\begin{aligned}
& \Pr(x_1 = 0, x_2 = 0) \cdot \frac{\Pr(x_1 = 0, x_2 = 0) + \Pr(x_1 = 0, x_2 = 1)}{\Pr(x_1 = 0, x_2 = 0) + \Pr(x_1 = 0, x_2 = 1)} + \\
& \Pr(x_1 = 0, x_2 = 1) \cdot \frac{\Pr(x_1 = 0, x_2 = 0) + \Pr(x_1 = 0, x_2 = 1)}{\Pr(x_1 = 0, x_2 = 0) + \Pr(x_1 = 0, x_2 = 1)} + \\
& \Pr(x_1 = 1, x_2 = 0) \cdot \frac{\Pr(x_1 = 1, x_2 = 0) + \Pr(x_1 = 1, x_2 = 1)}{\Pr(x_1 = 1, x_2 = 0) + \Pr(x_1 = 1, x_2 = 1)} + \\
& \Pr(x_1 = 1, x_2 = 1) \cdot \frac{\Pr(x_1 = 1, x_2 = 0) + \Pr(x_1 = 1, x_2 = 1)}{\Pr(x_1 = 1, x_2 = 0) + \Pr(x_1 = 1, x_2 = 1)} = 1
\end{aligned}$$

Finally, when $S = \{1, 2\}$, we have the following condition:

$$\begin{aligned}
& \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{z}_S = \mathbf{x}_S)] \\
& = \Pr(x_1 = 0, x_2 = 0) \cdot \Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 0 \mid x_1 = 0, x_2 = 0) + \\
& \Pr(x_1 = 0, x_2 = 1) \cdot \Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 0 \mid x_1 = 0, x_2 = 1) + \\
& \Pr(x_1 = 1, x_2 = 0) \cdot \Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 1 \mid x_1 = 1, x_2 = 0) + \\
& \Pr(x_1 = 1, x_2 = 1) \cdot \Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = 1 \mid x_1 = 1, x_2 = 1)
\end{aligned}$$

This can likewise be reordered as follows:

$$\begin{aligned}
& \Pr(x_1 = 0, x_2 = 0) \cdot \frac{\Pr(x_1 = 0, x_2 = 0)}{\Pr(x_1 = 0, x_2 = 0)} + \Pr(x_1 = 0, x_2 = 1) \cdot \frac{\Pr(x_1 = 0, x_2 = 1)}{\Pr(x_1 = 0, x_2 = 1)} + \\
& \Pr(x_1 = 1, x_2 = 0) \cdot \frac{\Pr(x_1 = 1, x_2 = 0)}{\Pr(x_1 = 1, x_2 = 0)} + \Pr(x_1 = 1, x_2 = 1) \cdot \frac{\Pr(x_1 = 1, x_2 = 1)}{\Pr(x_1 = 1, x_2 = 1)} \\
& = 1
\end{aligned}$$

Let $t = 2$. Evidently, we have $\Delta_t(\{1\}, f) = 0$. Next, let us analyse $\Delta_t(\emptyset, f)$. We note that if we assume feature independence, we observe that:

$$\begin{aligned}
& (\Pr(x_1 = 0, x_2 = 0) + \Pr(x_1 = 0, x_2 = 1))^2 + (\Pr(x_1 = 1, x_2 = 0) + \Pr(x_1 = 1, x_2 = 1))^2 \\
& = \Pr(x_1 = 0)^2 + \Pr(x_1 = 1)^2,
\end{aligned}$$

Additionally, another condition that holds is:

$$\begin{aligned}
& \frac{\Pr(x_1 = 0, x_2 = 0)^2 + \Pr(x_1 = 1, x_2 = 0)^2}{\Pr(x_1 = 0, x_2 = 0) + \Pr(x_1 = 1, x_2 = 0)} + \frac{\Pr(x_1 = 0, x_2 = 1)^2 + \Pr(x_1 = 1, x_2 = 1)^2}{\Pr(x_1 = 0, x_2 = 1) + \Pr(x_1 = 1, x_2 = 1)} \\
& = \Pr(x_1 = 0)^2 + \Pr(x_1 = 1)^2.
\end{aligned}$$

This implies $\Delta_t(\emptyset, f) = 0$, that is, $\Delta_t(\emptyset, f) \leq \Delta_t(\{1\}, f)$. However, without the assumption of feature independence, we cannot guarantee $\Delta_t(\emptyset, f) = 0$. For example, let:

$$\Pr(x_1 = 0, x_2 = 0) = P_{00}, \Pr(x_1 = 0, x_2 = 1) = P_{01}$$

and

$$\Pr(x_1 = 1, x_2 = 0) = P_{10}, \Pr(x_1 = 1, x_2 = 1) = P_{11}$$

Then, by simplifying $\Delta_t(\emptyset, f)$ we get that the following holds:

$$\Delta_t(\emptyset, f) = \left[\frac{P_{00}^2 + P_{10}^2}{P_{00} + P_{10}} + \frac{P_{01}^2 + P_{11}^2}{P_{01} + P_{11}} \right] - \left[(P_{00} + P_{01})^2 + (P_{10} + P_{11})^2 \right] \quad (57)$$

We then let $P_{00} = 0.4$, $P_{01} = 0.2$, $P_{10} = 0.1$, and $P_{11} = 0.3$, and get that the following holds:

$$\Delta_t(\emptyset, f) = \left[\frac{(0.4)^2 + (0.1)^2}{0.5} + \frac{(0.2)^2 + (0.3)^2}{0.5} \right] - \left[(0.4 + 0.2)^2 + (0.1 + 0.3)^2 \right] = 0.08. \quad (58)$$

This implies that $\Delta_t(\emptyset, f) > \Delta_t(1, f)$, indicating that — unlike in the monotonicity setting where it was not required — feature independence is a necessary condition for supermodularity.

G PROOF OF PROPOSITION 5

Proposition 5. *Computing Algorithm 1 with the local value functions v_{con}^ℓ and v_{suff}^ℓ does not always converge to a subset minimal δ -sufficient/contrastive reason. However, computing it with the global value functions v_{con}^g or v_{suff}^g necessarily produces subset minimal δ -sufficient/contrastive reasons.*

Proof. We note that this proposition follows directly from our proof that both v_{suff}^g and v_{con}^g are monotone non-decreasing (as established in Proposition 2), along with other relevant conclusions drawn in prior work. In the remainder of this section, we elaborate on how this monotonicity property leads to the stated corollary.

The convergence of Algorithm 1 to a subset-minimal δ global sufficient or contrastive reason with respect to the value functions v_{con}^g and v_{suff}^g follows directly from the monotonicity property, for which prior work (Ignatiev et al., 2019; Wu et al., 2024b; Arenas et al., 2022) showed that this type of greedy algorithm yields a subset-minimal explanation when the underlying value function is monotone non-decreasing. The algorithm halts when, for all $i \in S$, the value $v(S \setminus \{i\})$ drops below δ , ensuring that although $v(S) \geq \delta$ (a maintained invariant), removing any element causes the condition to fail. Due to monotonicity, this implies that no proper subset of S satisfies the condition, guaranteeing that S is subset-minimal.

We now turn to explain why Algorithm 1 does not converge to a subset-minimal solution in the local setting. We note that it is well known that the local probabilistic sufficient value function (and likewise the local probabilistic contrastive value function) is not monotone Arenas et al. (2022); Izza et al. (2023); Subercaseaux et al. (2025); Izza et al. (2024b). As a simple illustration, fix the uniform distribution on the two binary variables $\{x_1, x_2\}$ and define a Boolean function as follows:

$$f(x_1, x_2) := x_1 \vee x_2. \quad (59)$$

Consider the input $\mathbf{x} = (0, 1)$ which is classified as 1, then:

$$v_{\text{suff}}^\ell(\emptyset) = \frac{3}{4}, \quad v_{\text{suff}}^\ell(\{1\}) = \frac{1}{2}, \quad v_{\text{suff}}^\ell(\{2\}) = 1, \quad v_{\text{suff}}^\ell(\{1, 2\}) = 1. \quad (60)$$

Because $v_{\text{suff}}^\ell(\emptyset) > v_{\text{suff}}^\ell(\{1\}) < v_{\text{suff}}^\ell(\{1, 2\})$, the function v_{suff}^ℓ is not monotone. Likewise, it can be computed that:

$$v_{\text{con}}^\ell(\emptyset) = 1, \quad v_{\text{con}}^\ell(\{1\}) = 1, \quad v_{\text{con}}^\ell(\{2\}) = \frac{1}{2}, \quad v_{\text{con}}^\ell(\{1, 2\}) = \frac{3}{4}. \quad (61)$$

(For v_{con}^ℓ , the S denotes the set of features that are allowed to vary.) Because $v_{\text{con}}^\ell(\emptyset) > v_{\text{con}}^\ell(\{2\}) < v_{\text{con}}^\ell(\{1, 2\})$, the function v_{con}^ℓ is also non-monotone.

□

H PROOF OF THEOREM 1

Theorem 1. *If f is a decision tree and the probability term v_{suff}^g can be computed in polynomial time given the distribution $\mathcal{D}$ (which holds for independent distributions, among others), then obtaining a subset-minimal global δ -sufficient reason can be obtained in polynomial time. However, unless $\text{PTIME}=\text{NP}$, no polynomial-time algorithm exists for computing a local δ -sufficient reason for decision trees even under independent distributions.*

Proof. We begin by noting that for decision trees, assuming that $\mathcal{D}$ represents independent distributions, the computation of the value function:

$$v_{\text{suff}}^g(S) = \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{x}) = f(\mathbf{z}) \mid \mathbf{x}_S = \mathbf{z}_S)] \quad (62)$$

Can be carried out in polynomial time using the following procedure. Thanks to the tractability of decision trees, we can iterate over all pairs of leaf nodes, each corresponding to partial assignments

$\mathbf{x}_S$ and $\mathbf{z}_{S'}$. For each such pair, if both leaf nodes yield the same prediction under f , we compute the corresponding term in the expected value $\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{x}) = f(\mathbf{z}) \mid \mathbf{x}_S = \mathbf{z}_S)]$ by multiplying the respective feature-wise probabilities over the shared features of the two vectors. Under the feature independence assumption, these probabilities decompose, allowing the full expectation to be computed by summing the contributions of all such matching leaf pairs.

Since all probabilities involved are provided as part of the input and each step involves only polynomial-time operations, the entire procedure runs in polynomial time. This establishes that every step in Algorithm 1 is efficient, and thus the algorithm as a whole runs in polynomial time. Finally, combining this with Lemma 2, which proves that v_{suff}^g is monotone non-decreasing, we conclude that Algorithm 1 converges to a subset-minimal explanation. This is because the algorithm halts when, for all $i \in S$, the value $v(S \setminus \{i\})$ falls below δ , while $v(S) \geq \delta$ is preserved throughout. By monotonicity, this guarantees that no strict subset of S satisfies the condition, ensuring S is indeed subset-minimal.

For the remaining part of the claim, the result follows from (Arenas et al., 2022), which showed that, assuming $P \neq NP$, there is no polynomial-time algorithm for computing a subset-minimal δ -sufficient reason for decision trees. While this was proven specifically under the uniform distribution, the hardness clearly extends to independent distributions, which include this case. Combined, these results establish the complexity separation between the local and global variants, thus concluding our proof.

□

I PROOF OF PROPOSITION 6

Proposition 6. *For any model f , and empirical distribution $\mathcal{D}$ — computing a subset-minimal global δ -sufficient or δ -contrastive reason for f can be done in polynomial time.*

Proof. The computation of the global sufficient probability function:

$$v_{\text{suff}}^g(S) = \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{x}) = f(\mathbf{z}) \mid \mathbf{x}_S = \mathbf{z}_S)] \quad (63)$$

or the computation of the global contrastive probability function

$$v_{\text{con}}^g(S) = \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{x}) \neq f(\mathbf{z}) \mid \mathbf{x}_{\bar{S}} = \mathbf{z}_{\bar{S}})] \quad (64)$$

Can be performed in polynomial time when $\mathcal{D}$ is selected from the class of empirical distributions. This is achieved by iterating over pairs of instances $\mathbf{x}, \mathbf{z}$ within the dataset and running an inference through $f(\mathbf{x})$ and $f(\mathbf{z})$ to compute the expected values at each step. Consequently, determining whether the probability function exceeds or falls below a given threshold δ can also be accomplished in polynomial time. Furthermore, leveraging the proof of non-decreasing monotonicity (Proposition D) of the global probability function, we can utilize algorithm 1 for contrastive reason search or for sufficient reason search. By executing a linear number of polynomial queries, a subset-minimal explanation (whether a subset minimal contrastive reason or a subset minimal sufficient reason) can thus be obtained in polynomial time.

□

J PROOF OF THEOREM 2

Theorem 2. *Assuming f is a neural network or a tree ensemble, and $\mathcal{D}$ is an empirical distribution — there exist polynomial-time algorithms for obtaining subset minimal global δ -sufficient and contrastive reasons. However, unless $P = NP$, there is no polynomial time algorithm for computing a subset minimal local δ -sufficient reason or a subset minimal local δ -contrastive reason.*

Proof. The first part of the proof follows directly from Proposition 6, which established that this holds for any model, since both $v_{\text{suff}}^g(S) = \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{x}) = f(\mathbf{z}) \mid \mathbf{x}_S = \mathbf{z}_S)]$ and $v_{\text{con}}^g(S) = \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{x}) \neq f(\mathbf{z}) \mid \mathbf{x}_{\bar{S}} = \mathbf{z}_{\bar{S}})]$ can be computed in polynomial time from

the empirical distribution. For the second part, we proceed by proving two lemmas: one establishing the intractability of computing the sufficient case, and the other addressing the contrastive case.

Lemma 6. *Unless $P_{TIME} = NP$, then there is no polynomial time algorithm for computing a subset minimal local δ -sufficient reason for either a tree ensemble or a neural network, under empirical distributions.*

Proof. We will first prove this claim for neural networks and then extend the result to tree ensembles. We will establish this claim by demonstrating that if a polynomial-time algorithm exists for computing a subset-minimal local δ sufficient reason for neural networks, then it would enable us to solve the classic NP-complete *CNF-SAT* problem in polynomial time. The *CNF-SAT* problem is defined as follows:

CNF-SAT:

Input: A formula in conjunctive normal form (CNF): ϕ .

Output: Yes, if there exists an assignment to the n literals of ϕ such that ϕ is evaluated to True, and No otherwise

Our proof will also utilize the following lemma (with its proof provided in Barceló et al. (2020)):

Lemma 7. *Any boolean circuit ϕ can be encoded into an equivalent MLP over the binary domain $\{0, 1\}^n \rightarrow \{0, 1\}$ in polynomial time.*

Proof. We will actually establish hardness for a simpler, more specific case, which will consequently imply hardness for the more general setting. In this case, we assume that the empirical distribution $\mathcal{D}$ consists of a single element, which, for simplicity, we take to be the zero vector $\mathbf{0}_n$. Since the sufficiency of S in this scenario depends only on the local instance $\mathbf{x}$ and the single instance $\mathbf{0}_n$ in our empirical distribution, the nature of the input format (whether discrete, continuous, etc.) does not affect the result. Therefore, the hardness results hold universally across all these settings.

Similarly to $\mathbf{0}_n$, we define $\mathbf{1}_n$ as an n -dimensional vector consisting entirely of ones. Given an input ϕ , we initially assign $\mathbf{1}_n$ to ϕ , effectively setting all variables to True. If ϕ evaluates to True under this assignment, then a satisfying truth assignment exists. Therefore, we assume that $\phi(\mathbf{1}_n) = 0$. We now introduce the following definitions:

$$\begin{aligned}\phi_2 &:= (x_1 \wedge x_2 \wedge \dots \wedge x_n), \\ \phi' &:= \phi \vee \phi_2\end{aligned}\tag{65}$$

ϕ' , while no longer a CNF, can still be transformed into an MLP f using Lemma 8, ensuring that f behaves equivalently to ϕ' over the domain $\{0, 1\}^n$. Given our assumption that computing a subset-minimal local δ -sufficient reason is feasible in polynomial time, we can determine one for the instance $\langle f, \mathbf{x} := \mathbf{1}_n \rangle$, $\delta := 1$, noting that the empirical distribution $\mathcal{D}$ is simply defined over the single data point $\mathbf{0}_n$.

We now assert that the subset-minimal δ -sufficient reason generated for $\langle f, \mathbf{x} \rangle$ encompasses the entire input space, i.e., $S = \{1, \dots, n\}$, if and only if $\langle \phi \rangle \notin \text{CNF-SAT}$.

Let us assume that $S = \{1, \dots, n\}$. Since S is a δ -sufficient reason for $\langle f, \mathbf{x} \rangle$, this simply means that setting the complementary set to any value maintains the prediction. Since the complementary set is $\emptyset$ in this case, this trivially holds. The fact that S is subset-minimal means that any other subset $S' \subseteq S$ satisfies $v(S') < \delta = 1$. Since the probability function $\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) = f(\mathbf{x}) \mid \mathbf{x}_S = \mathbf{z}_S)$ is determined by a single point (the distribution contains only the point $\mathbf{0}_n$), the probability function can only take values of 1 or 0. Hence, we also know that $v(S') = 0$. This tells us that, aside from the subset $S = \{1, \dots, n\}$, for any subset $S' \subseteq S$, fixing the features in S' to 1 and the rest to 0 does not result in a classification outcome of 1. Since the ϕ_2 component within ϕ' is True only if all features are assigned 1, this directly implies that ϕ is assigned False for any of these inputs. Since we already know that ϕ does not return a True answer for the vector assignment $\mathbf{1}_n$ (as verified at the beginning), and now we have established that the same holds for all other input vectors, we conclude that $\phi \notin \text{CNF-SAT}$.

Now, suppose that S is a subset that is strictly contained within $\{1, \dots, n\}$. Given that S is sufficient under our definitions of the distribution $\mathcal{D}$ with $\delta = 1$, we can apply the same reasoning as before to

conclude that $v(S) = 1$. This implies that setting the features in S to 1 while setting the remaining features to 0 ensures that the function f evaluates to 1. Since this assignment is necessarily not a vector consisting entirely of ones, it follows that the ϕ_2 component within ϕ' must be False. Consequently, the ϕ component must be True, which implies that $\langle \phi \rangle \in \text{CNF-SAT}$. This completes the proof.

We have established that the following claim holds for neural networks. However, extending the proof to tree ensembles requires a minor and straightforward adaptation. To do so, we will utilize the following lemma, which has been noted in several previous works (Ordyniak et al., 2024; Audemard et al., 2022b):

Lemma 8. *Any CNF or DNF ϕ can be encoded into an equivalent random forest classifier over the binary domain $\{0, 1\}^n \rightarrow \{0, 1\}$ in polynomial time.*

Proof. We observe that we can apply the same process used in our proof for neural networks, where we encoded ϕ' into an equivalent neural network. However, ϕ' is no longer a valid CNF due to our construction (though encoding it into an MLP was not an issue, as any Boolean circuit can be transformed into an MLP). Nevertheless, since ϕ' consists of a conjunction of only two terms, we can easily represent it as an equivalent CNF:

$$\phi' := (c_1 \vee \phi_2) \wedge (c_2 \vee \phi_2) \wedge \dots (c_m \vee \phi_2) \quad (66)$$

Where each c_i is a disjunction of a few terms. Consequently, ϕ' is a valid CNF, allowing us to transform it into an equivalent random forest classifier. The reduction we outlined for MLPs applies directly to these models as well, thereby completing the proof for both model families.

□

Lemma 9. *Unless $\text{PTIME} = \text{NP}$, then there is no polynomial time algorithm for computing a subset minimal local δ -contrastive reason for either a tree ensemble or a neural network, under empirical distributions.*

Proof. We will present a proof analogous to the one in Lemma 6. Specifically, we will once again utilize the classical NP-hard CNF-SAT problem defined in Lemma 6. In particular, given a Boolean formula ϕ , we will demonstrate that determining a subset-minimal contrastive reason — whether for a neural network or a tree ensemble — allows us to decide the satisfiability of ϕ .

First, we check whether assigning all variables in ϕ to 1 evaluates ϕ to True. If so, the formula is satisfiable, and we have determined its satisfiability. Otherwise, we use Lemma 8 to encode the CNF formula as a neural network f . Next, we compute a subset-minimal δ -sufficient reason by setting $\mathcal{D}$ as the empirical distribution containing only a single data point $\mathbf{0}_n$, following a similar procedure to Lemma 6. Additionally, we set $\delta = 1$ and compute a subset-minimal δ -contrastive reason concerning $\langle f, \mathbf{x} \rangle$.

We will now demonstrate that if any subset-minimal δ contrastive reason obtained for ϕ is satisfiable, then ϕ itself is satisfiable. Conversely, if no subset-minimal δ contrastive reason is obtained, then ϕ is unsatisfiable. The validity of this claim follows a reasoning similar to that provided in Lemma 6. Specifically, the term $\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{x}) \neq f(\mathbf{z}) \mid \mathbf{x}_{\bar{S}} = \mathbf{z}_{\bar{S}})$, where the distribution $\mathcal{D}$ considers sampling from a single datapoint, can set the probability to either 0 or 1. Furthermore, since we are searching for a $\delta = 1$ contrastive reason, this is equivalent to asking whether there exists an assignment that changes the classification of f , which corresponds to modifying the assignment of ϕ from False to True. If such an assignment exists, then ϕ is satisfiable. However, if no subset-minimal contrastive reason exists, then no subset of features fixed to zero — when the complementary set is set to ones — evaluates to true. This is equivalent to stating that no assignment evaluates f (and consequently ϕ) to True, implying that ϕ is unsatisfiable.

To extend the proof from neural networks to tree ensembles, we can follow the same procedure outlined in Lemma 6, encoding the CNF formula into an equivalent random forest classifier. Consequently, the proof remains valid for tree ensembles, thereby concluding the proof.

□

K PROOF OF THEOREM 3

Theorem 3. *Given a neural network or tree ensemble f and an empirical distribution $\mathcal{D}$ over a fixed dataset D , Algorithm 2 yields a constant $\mathcal{O}\left(\ln\left(\frac{v_{\text{con}}^g([n])}{\min_{i \in [n]} v_{\text{con}}^g(\{i\})}\right)\right)$ -approximation, bounded by $\mathcal{O}(\ln(|D|))$, for computing a global cardinally minimal δ -contrastive reason for f , assuming feature independence. In contrast, unless $\text{PTIME}=\text{NP}$, no bounded approximation exists for computing a local cardinally minimal δ -contrastive reason for any $\langle f, \mathbf{x} \rangle$, even when $|D| = 1$.*

Proof. We divide the proof of the theorem into two lemmas, covering both the approximation guarantee and the result on the absence of a bounded approximation.

Lemma 10. *Given a neural network or tree ensemble f and an empirical distribution $\mathcal{D}$ over a fixed dataset D , Algorithm 2 yields a constant $\mathcal{O}\left(\ln\left(\frac{v_{\text{con}}^g([n])}{\min_{i \in [n]} v_{\text{con}}^g(\{i\})}\right)\right)$ -approximation, bounded by $\mathcal{O}(\ln(|D|))$, for computing a global cardinally minimal δ -contrastive reason for f , assuming feature independence.*

Proof. We will in fact prove a stronger claim, showing that this holds for *any* model, provided the trivial condition that its inference time is computable in polynomial time, along with one additional mild condition that we will detail later — both of which apply to both our neural network and tree ensemble formalizations.

We begin by noting that, since we are working with empirical distributions, the computation of the global contrastive probability value function:

$$v_{\text{con}}^g(S) = \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\mathbf{Pr}_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{x}) \neq f(\mathbf{z}) \mid \mathbf{x}_{\bar{S}} = \mathbf{z}_{\bar{S}})] \quad (67)$$

Can be computed in polynomial time by iterating over all pairs $\mathbf{x}, \mathbf{z}$ in the dataset D , as previously established in Proposition 2. Since Algorithm 2 performs only a linear number of these polynomial-time queries, its total runtime is therefore polynomial.

Regarding the approximation, the classical work by Wolsey et al. (Wolsey, 1982) established a harmonic-series-based approximation guarantee for monotone non-decreasing submodular functions v with integer values. More generally, their result yields an approximation factor of $\mathcal{O}\left(\ln\left(\frac{v([n])}{\min_{i \in [n]} v(\{i\})}\right)\right)$. We showed in Proposition 2 that the value function v_{con}^g is monotone non-decreasing, and under feature independence, Proposition 4 establishes its submodularity. Combined with the fact that Algorithm 2 runs in polynomial time, this directly yields an approximation guarantee of $\mathcal{O}\left(\ln\left(\frac{v_{\text{con}}^g([n])}{\min_{i \in [n]} v_{\text{con}}^g(\{i\})}\right)\right)$.

However, we must also show that the expression $\mathcal{O}\left(\ln\left(\frac{v_{\text{con}}^g([n])}{\min_{i \in [n]} v_{\text{con}}^g(\{i\})}\right)\right)$ is both finite and bounded by $\mathcal{O}(\ln(|D|))$, implying that it is effectively constant, due to the assumption of a fixed dataset $|D|$. To ensure this, we begin by confirming that the expression is well-defined and finite. We then proceed to establish the desired bounds. To do so, we introduce a preprocessing step in which we eliminate “redundant” elements—those that could theoretically cause the denominator $\min_{i \in [n]} v_{\text{con}}^g(\{i\})$ to be zero. We begin by formally defining what we mean by redundancy:

Definition 1. *Let $\mathcal{D}$ denote some empirical distribution over a dataset D . Then we say that some feature $i \in [n]$ is redundant with respect to $\mathcal{D}$ if for any pair $\mathbf{x}, \mathbf{z} \in D$ it holds that $f(\mathbf{x}_{[n] \setminus \{i\}}; \mathbf{z}_{\{i\}}) = f(\mathbf{x})$.*

Here, the notation $f(\mathbf{x}_{[n] \setminus \{i\}}; \mathbf{z}_{\{i\}})$ indicates that all features in $[n] \setminus \{i\}$ are fixed to their values in $\mathbf{x}$, while feature i is set to its value in $\mathbf{z}$. Notably, this is equivalent to defining:

$$\forall S \subseteq [n], \mathbf{z} \in D \quad f(\mathbf{x}_S; \mathbf{z}_{\bar{S}}) = f(\mathbf{x}_{S \setminus \{i\}}; \mathbf{z}_{\bar{S} \cup \{i\}}) \quad (68)$$

As before, the notation $(\mathbf{x}_S; \mathbf{z}_{\bar{S}})$ indicates that the features in S are fixed to their values in $\mathbf{x}$, while the features in $\bar{S}$ are fixed to their values in $\mathbf{z}$. Thus, in this sense, if we “remove” feature i from the input space and define a new function f' over the reduced space $[n'] := [n] \setminus \{i\}$, then for any input of size n' , the output of f' will exactly match the output of f when applied to the same features $[n] \setminus \{i\}$.

This removal can be done in polynomial time for both tree ensembles and neural networks: for neural networks, it involves detaching the corresponding input neuron from the network; for tree ensembles, it involves removing any splits on that feature from all decision trees. Given the empirical nature of the contrastive value function $v_{\text{con}}^g(S)$ — as previously discussed in this proof and in Proposition 2 — we can compute each $v_{\text{con}}^g(\{i\})$ for all $i \in [n]$ by iterating over all pairs $\mathbf{x}, \mathbf{z} \in \mathcal{D}$ and checking whether:

$$f(\mathbf{x}_{[n] \setminus \{i\}}; \mathbf{z}_{\{i\}}) \neq f(\mathbf{x}) \quad (69)$$

If this condition holds, it indicates that feature i is not redundant. Conversely, if the condition fails for all pairs considered during iteration, then i is deemed redundant. Once all redundant features have been identified with respect to the empirical distribution $\mathcal{D}$, we can remove them and construct an equivalent model f' . As discussed above, this transformation can be performed in polynomial time for both neural networks and decision trees.

To conclude our proof, we observe that $v_{\text{con}}^g(\{i\})$ equals 0 for some empirical distribution $\mathcal{D}$ if and only if feature i is redundant with respect to $\mathcal{D}$. Therefore, once the preprocessing step removes all redundant features from f and we construct the resulting function f' , we ensure that $\min_{i \in [n]} v_{\text{con}}^g(\{i\}) > 0$.

We now present the more precise bounds referenced in our proof. Specifically, we demonstrate that the approximation factor is $\mathcal{O}(\ln(|\mathcal{D}|))$. Given that the empirical distribution $\mathcal{D}$ is fixed, this yields a constant approximation. While tighter bounds may be achievable, our goal here is solely to establish that the bound is *constant* — a property that will later sharply contrast with the local explainability setting, where no bounded approximation exists.

More specifically, we know that the probability function $\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{x}) \neq f(\mathbf{z}) \mid \mathbf{x}_{\bar{S}} = \mathbf{z}_{\bar{S}})$ is, by the definition of empirical distributions, at least $\frac{1}{|\mathcal{D}|}$. Therefore, the definition of $\mathbb{E}_{\mathbf{z} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{x}) \neq f(\mathbf{z}) \mid \mathbf{x}_{\bar{S}} = \mathbf{z}_{\bar{S}})]$ (i.e., the value function v_{con}^g) is at least $\frac{1}{|\mathcal{D}|^2}$. In particular, this also implies that:

$$\min_{i \in [n]} v_{\text{con}}^g(\{i\}) \geq \frac{1}{|\mathcal{D}|^2} \quad (70)$$

Consequently, we obtain that:

$$\ln \left(\frac{v_{\text{con}}^g([n])}{\min_{i \in [n]} v_{\text{con}}^g(\{i\})} \right) \leq \ln(|\mathcal{D}|^2) = \mathcal{O}(\ln(|\mathcal{D}|)) \quad (71)$$

Which concludes our proof.

□

Lemma 11. *Unless PTIME=NP, there is no bounded approximation for computing a cardinally minimal local δ -contrastive reason for any $\langle f, \mathbf{x} \rangle$ where f is either a neural network or a tree ensemble, even when $|\mathcal{D}| = 1$.*

Proof. Assume, for contradiction, that there exists a bounded approximation algorithm for computing a cardinally minimal local δ -contrastive reason for some $\langle f, \mathbf{x} \rangle$, where f is a neural network or a tree ensemble—even when $|\mathcal{D}| = 1$. However, Lemma 9 establishes that even with a single baseline $\mathbf{z} = 0_n$ (i.e., when the entire dataset is just one instance), deciding whether a contrastive reason exists is NP-hard unless PTIME = NP. Therefore, if a polynomial-time algorithm could yield a bounded approximation for a cardinally minimal contrastive reason, it would contradict this result, as such an approximation would implicitly decide the existence of a contrastive explanation.

□

L PROOF OF THEOREM 4

Theorem 4. *Given a neural network or tree ensemble f and an empirical distribution $\mathcal{D}$ over a fixed dataset D , Algorithm 2 yields a constant $\mathcal{O} \left(\frac{1}{1-k_f} + \ln \left(\frac{v_{\text{con}}^g([n])}{\min_{i \in [n]} v_{\text{con}}^g(\{i\})} \right) \right)$ -approximation for*

computing a global cardinally minimal δ -sufficient reason for f , assuming feature independence. In contrast, unless $PTIME=NP$, there is no bounded approximation for computing a local cardinally minimal δ -sufficient reason for any $\langle f, \mathbf{x} \rangle$, even when $|D| = 1$.

Proof. We divide the proof into two lemmas: one establishing the approximation guarantee for the global case, and the other demonstrating the intractability of the local case.

Lemma 12. *Given a neural network or tree ensemble f and an empirical distribution $\mathcal{D}$ over a fixed dataset D , Algorithm 2 yields a constant $\mathcal{O}\left(\frac{1}{1-k^f} + \ln\left(\frac{v_{\text{con}}^g([n])}{\min_{i \in [n]} v_{\text{con}}^g(\{i\})}\right)\right)$ -approximation for computing a global cardinally minimal δ -sufficient reason for f , assuming feature independence.*

Proof. The proof will follow a similar approach to that of Lemma 10, where we showed that for both neural networks and tree ensembles, Algorithm 2 achieves an approximation factor of $\ln\left(\frac{v_{\text{con}}^g([n])}{\min_{i \in [n]} v_{\text{con}}^g(\{i\})}\right)$. After applying the preprocessing step, this ratio is guaranteed to be finite and bounded by $\mathcal{O}(\ln(|D|))$.

Similar to Lemma 10, we will again prove a stronger claim — namely, that the result holds for *any* model, assuming the trivial condition that its inference time is polynomially computable, and the additional condition concerning the removal of redundant features, as described in Lemma 10. As discussed there, both conditions are satisfied by our neural network and tree ensemble formalizations.

Here as well, since we are working with empirical distributions, the computation of the global sufficient probability value function:

$$v_{\text{suff}}^g(S) = \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{x}) = f(\mathbf{z}) \mid \mathbf{x}_S = \mathbf{z}_S)] \quad (72)$$

can be computed in polynomial time by iterating over all pairs $\mathbf{x}, \mathbf{z}$ in the dataset D , as previously shown in both Lemma 10 and Proposition 2. Since Algorithm 2 makes only a linear number of such polynomial-time queries, its overall runtime is polynomial.

Unlike Lemma 10, where the approximation relied on the result by Wolsey et al. (Wolsey, 1982) for monotone non-decreasing submodular functions, the setting here requires a different condition due to the supermodular nature of the function. Specifically, Shi et al. (Shi et al., 2021) provided an approximation guarantee of $\mathcal{O}\left(\frac{1}{1-k^f} + \ln\left(\frac{v_{\text{suff}}^g([n])}{\min_{i \in [n]} v_{\text{suff}}^g(\{i\})}\right)\right)$, where $k^f := 1 - \min_{i \in [n]} \frac{v([n]) - v([n] \setminus i)}{v(i) - v(\emptyset)}$.

We established in Proposition 2 that the value function v_{suff}^g is monotone non-decreasing, and under the assumption of feature independence, Proposition 3 further shows that it is supermodular. Given that Algorithm 2 runs in polynomial time, this leads directly to an approximation guarantee of $\mathcal{O}\left(\frac{1}{1-k^f} + \ln\left(\frac{v_{\text{suff}}^g([n])}{\min_{i \in [n]} v_{\text{suff}}^g(\{i\})}\right)\right)$, where $k^f := 1 - \min_{i \in [n]} \frac{v_{\text{suff}}^g([n]) - v_{\text{suff}}^g([n] \setminus \{i\})}{v_{\text{suff}}^g(\{i\}) - v_{\text{suff}}^g(\emptyset)}$.

To show that this expression is both bounded and constant, we follow the same preprocessing step as in Lemma 10, where we remove all redundant features from f — a process that, as previously explained, can be carried out in polynomial time for both neural networks and tree ensembles. As before, a feature i is strictly redundant if and only if $v(\{i\}) - v(\emptyset) = 0$. This preprocessing yields a new function f' that behaves identically to f over the remaining features and ensures that, for every i , both $v_{\text{suff}}^g(\{i\}) > 0$ and $v_{\text{suff}}^g(\{i\}) - v_{\text{suff}}^g(\emptyset) > 0$ hold.

In order for us to show that this expression is both bounded and constant, similarly to Lemma 10 we will perform the exact same preprocessing phase (which, as we explained there, can be performed in polynomial time both for neural networks as well as tree ensembles) as before where we remove all redundant features from f . We note that here too, it holds that a feature i is strictly redundant if and only if $v(\{i\}) - v(\emptyset) = 0$. Hence, this preprocessing phase will give us a new function f' for which it both holds that it is equivalent to the behaviour of f for all remaining features and also satisfies that for any i it holds that both $v_{\text{suff}}^g(\{i\}) > 0$ and $v_{\text{suff}}^g(\{i\}) - v_{\text{suff}}^g(\emptyset) > 0$ hold.

Now, following the same reasoning as in Lemma 10, where the probability term is lower bounded by $\frac{1}{|D|}$ and its expected value by $\frac{1}{|D|^2}$, we obtain that $\min_{i \in [n]} v_{\text{suff}}^g(i) \geq \frac{1}{|D|^2}$ and $v_{\text{suff}}^g(i) - v_{\text{suff}}^g(\emptyset) \geq \frac{1}{|D|^2}$. This implies, as in Lemma 10, that the first term is lower bounded by $\mathcal{O}(\ln(|D|))$, and for the second term we derive the following:

Since v_{suff}^g is supermodular (Proposition 3), we have that $v_{\text{suff}}^g([n]) - v_{\text{suff}}^g([n] \setminus \{i\})$. It follows that $1 \leq v_{\text{suff}}^g([n]) - v_{\text{suff}}^g([n] \setminus \{i\}) \leq \frac{1}{|D|^2}$. This implies: $1 \leq \min_{i \in [n]} \frac{v_{\text{suff}}^g([n]) - v_{\text{suff}}^g([n] \setminus \{i\})}{v_{\text{suff}}^g(\{i\}) - v_{\text{suff}}^g(\emptyset)} \leq |D|^2$, and therefore, $1 - |D|^2 \leq k^f \leq 0$. This yields $\frac{1}{1-k^f} \leq |D|^2$, demonstrating that the overall approximation bound based on the one established by Shi et al. is constant. While this bound may be significantly smaller in practice, our goal here is simply to show that it remains constant — unlike in the local setting, where no bounded approximation is achievable.

□

Lemma 13. *Unless $\text{PTIME} = \text{NP}$, there is no bounded approximation for computing a local cardinally minimal δ -sufficient reason for any $\langle f, \mathbf{x} \rangle$, even when $|D| = 1$.*

Proof. The proof follows a similar approach to Lemma 11. Suppose, for contradiction, that there exists a polynomial-time algorithm that provides a bounded approximation to a cardinally minimal local δ -sufficient reason for some $\langle f, \mathbf{x} \rangle$, where f is either a neural network or a tree ensemble—even when $|D| = 1$. Yet, Lemma 6 shows that even in the extreme case where the dataset consists of a single baseline $\mathbf{z} = 0_n$, deciding whether a sufficient reason exists is NP-hard unless $\text{PTIME} = \text{NP}$. Hence, the existence of such an approximation algorithm would contradict this hardness result, as it would entail the ability to decide the existence of a sufficient reason.

□

M HARDNESS OF APPROXIMATING GLOBAL CONTRASTIVE REASONS FOR NEURAL NETWORKS USING EMPIRICAL DISTRIBUTIONS

Proposition 1. *Assume a neural network $f : \mathbb{F} \rightarrow [c]$ with ReLU activations, some dataset $\mathbf{D}$, such that $\mathcal{D}$ is the empirical distribution defined over $\mathbf{D}$, and some threshold $\delta \in \mathbb{Q}$. Then unless $\text{NP} \subseteq \text{DTIME}(n^{\mathcal{O}(\log(\log(n)))})$, there does not exist a polynomial-time $\mathcal{O}(\ln(|\mathbf{D}|) - \epsilon)$ approximation algorithm for computing a cardinally minimal global contrastive reason for f concerning $\mathcal{D}$, for any fixed $\epsilon > 0$.*

Proof. The proof proceeds via a polynomial-time approximation-preserving reduction from a specific variant of Set Cover that is, similarly to regular Set Cover, provably hard to approximate within the factor given in the proposition (Vazirani, 2001).

Partial Set Cover:

Input: A universe $U := \{u_1, \dots, u_n\}$, a collection of sets $S := \{S_1, \dots, S_m\}$ such that for all j $S_j \subseteq U$, and a coverage threshold $\delta \in \{0, 1, \dots, n\}$.

Output: A minimum-size subset $C \subseteq S$ such that $|\cup_{S_j \in C} S_j| \geq \delta$.

We reduce this problem to that of computing a *global*, cardinally-minimal contrastive reason S for a neural network f with respect to a distribution $\mathcal{D}$ defined over a dataset $\mathbf{D}$, and some threshold δ' . We begin by recalling the definition of a global contrastive reason and then extend it to the setting where the explanation is evaluated over an empirical dataset $\mathbf{D}$. The definition is as follows:

$$\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) \neq f(\mathbf{x}) \mid \mathbf{z}_{\bar{S}} = \mathbf{x}_{\bar{S}})] \geq \delta. \quad (73)$$

More specifically, when we assume an empirical distribution over $\mathbf{D}$, this implies that:

$$\begin{aligned} & \mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) \neq f(\mathbf{x}) \mid \mathbf{z}_{\bar{S}} = \mathbf{x}_{\bar{S}})] = \\ & \frac{1}{|\mathbf{D}|} \sum_{\mathbf{x} \in \mathbf{D}} [\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) \neq f(\mathbf{x}) \mid \mathbf{z}_{\bar{S}} = \mathbf{x}_{\bar{S}})] = \\ & \frac{1}{|\mathbf{D}|} \sum_{\mathbf{x} \in \mathbf{D}} \left(\frac{\sum_{\mathbf{z} \in \mathbf{D}_{\bar{S}}(\mathbf{z}, \mathbf{x})} \mathbf{1}_{\{f(\mathbf{z}) \neq f(\mathbf{x})\}}}{\mathbf{D}_{\bar{S}}(\mathbf{z}, \mathbf{x})} \right) \end{aligned} \quad (74)$$

where we denote $\mathbf{D}_{\bar{S}}(\mathbf{z}, \mathbf{x}) := \{\mathbf{z} \in \mathbf{D} : \mathbf{z}_{\bar{S}} = \mathbf{x}_{\bar{S}}\}$. We now construct a ReLU neural network f . Before doing so, we specify the functional form that we intend the constructed neural network f to

implement. For each set S_j , we introduce two vectors $\mathbf{x}^j$ and $\mathbf{y}^j$, both belonging to $\mathbb{Z}^n$ and drawn from the dataset $\mathbf{D}$, such that they agree on all coordinates outside S_j ; that is, they share identical feature values on the complement $\bar{S}_j$. Formally:

$$\mathbf{x}_{\bar{S}_j}^j = \mathbf{y}_{\bar{S}_j}^j \quad (75)$$

and we assign strictly different values to the features in S_j by selecting them from two distinct vectors $\mathbf{z}$ and $\mathbf{z}'$ such that:

$$\mathbf{x}^j = (\mathbf{x}_{\bar{S}_j}^j; \mathbf{z}_{S_j}^j) \neq (\mathbf{x}_{\bar{S}_j}^j; \mathbf{z}'_{S_j}) = (\mathbf{y}_{\bar{S}_j}^j; \mathbf{z}'_{S_j}) = \mathbf{y}^j \quad (76)$$

We aim to construct f such that it satisfies the following:

$$f(\mathbf{x}^j) = f(\mathbf{x}_{\bar{S}_j}^j; \mathbf{z}_{S_j}^j) \neq f(\mathbf{x}_{\bar{S}_j}^j; \mathbf{z}'_{S_j}) = f(\mathbf{y}_{\bar{S}_j}^j; \mathbf{z}'_{S_j}) = f(\mathbf{y}^j) \quad (77)$$

Let us now assume that both $\mathbf{x}^j$ and $\mathbf{y}^j$ occur *uniquely* in the dataset $\mathbf{D}$ across all $[n]$ coordinates — that is, no other vector in $\mathbf{D}$ matches either of them on any coordinate. Under this assumption, we obtain the following relation:

$$\mathbb{E}_{\mathbf{x} \sim \mathcal{D}}[\Pr_{\mathbf{z} \sim \mathcal{D}}(f(\mathbf{z}) \neq f(\mathbf{x}) \mid \mathbf{z}_{S_j} = \mathbf{x}_{S_j})] = \frac{1}{|\mathbf{D}|} \sum_{\mathbf{x} \in \mathbf{D}} \left(\frac{\sum_{\mathbf{z} \in \mathbf{D}_{S_j}(\mathbf{z}, \mathbf{x})} \mathbf{1}_{\{f(\mathbf{z}) \neq f(\mathbf{x})\}}}{\mathbf{D}_{S_j}(\mathbf{z}, \mathbf{x})} \right) = \frac{2}{2|\mathbf{D}|} = \frac{1}{|\mathbf{D}|} \quad (78)$$

Here, the factor of 2 appearing in both the numerator and denominator arises from the fact that only two vector pairs — $(\mathbf{x}^j, \mathbf{z}^j)$ and $(\mathbf{z}^j, \mathbf{x}^j)$ — contribute to the respective sums.

We repeat this construction for each S_i in the set-cover instance, each time choosing two *distinct* vectors $\mathbf{x}^i$ and $\mathbf{y}^i$ that differ on *every* coordinate in $[n]$. In total, this yields $2m$ vectors that form the dataset $\mathbf{D}$ from which we construct f .

We then construct a neural network f that simulates this function. This is achieved using a known result on *neural network memorization* (Vardi et al., 2022), which guarantees that a ReLU network capable of representing any finite set of input–output pairs can be built in polynomial time. Using this result, we construct f accordingly, and set $\delta' := \frac{\delta}{|\mathbf{D}|}$.

We now argue that a subset S_j is a global contrastive reason for f under tolerance δ' if and only if C is a partial set cover for U under tolerance δ . By Equation 78, each constructed pair of vectors $(\mathbf{x}^j, \mathbf{y}^j)$ contributes exactly $\frac{1}{|\mathbf{D}|}$ to the value term. Moreover, because each S_j uses its own distinct pair of vectors in $\mathbb{Z}^n$, there is no overlap between these pairs. Consequently, the coverage achieved by any set cover C corresponds exactly to the accumulated contribution of the selected subsets in the input domain of f . This completes the reduction.

□

N TRACTABILITY RESULTS FOR ORTHOGONAL DNFS

Lemma 14. *If f is an orthogonal DNFs and the probability term v_{suff}^g can be computed in polynomial time given the distribution $\mathcal{D}$ (which holds for independent distributions, among others), then obtaining a subset-minimal global δ -sufficient reason can be obtained in polynomial time.*

Proof. We begin by defining orthogonal DNFs (Crama & Hammer, 2011) before presenting the full proof.

Definition 2 (Orthogonal DNFs). *A Boolean function φ is in orthogonal disjunctive normal form (orthogonal DNF) (Crama & Hammer, 2011) if it is a disjunction of terms $T_1 \vee T_2 \vee \dots \vee T_m$, such that for every pair of distinct terms T_i and T_j ($i \neq j$):*

$$T_i \wedge T_j \models \perp$$

This means that no single variable assignment can satisfy more than one term simultaneously.

2214 The proof mirrors the argument used in Theorem 1, relying on the fact that v_{suff}^g can be computed in
 2215 polynomial time for any set S .

2216 For a given set S , let $\mathcal{C}^S T_i$ be the constraints that the term T_i imposes on the variables in S .
 2217 We enumerate all pairs of terms (T_i, T_j) . Any pair whose constraints on S are *inconsistent* (i.e.,
 2218 $\mathcal{C}_{T_i}^S \wedge \mathcal{C}_{T_j}^S \models \perp$) is ignored, since it contributes zero to v_{suff}^g .
 2219

2220 For each *consistent* pair of terms (T_i, T_j) , we evaluate $\Pr(\mathbf{x}_S \in \mathcal{C}_{T_i, T_j}^S)$, where $\mathcal{C}_{T_i, T_j}^S = \mathcal{C}_{T_i}^S \cap \mathcal{C}_{T_j}^S$.
 2221 The contribution of this pair to v_{suff}^g is then
 2222

$$2223 \Pr(\mathbf{x}_S \in \mathcal{C}_{T_i, T_j}^S) \cdot \Pr(\mathbf{x}_{\bar{S}} \in \mathcal{C}_{T_i}^{\bar{S}}) \cdot \Pr(\mathbf{x}_{\bar{S}} \in \mathcal{C}_{T_j}^{\bar{S}}).$$

2224
 2225 Under feature independence, these probabilities factor as $\Pr(\mathbf{x}_S) = \prod_{i \in S} p(x_i)$. Since there are at
 2226 most m^2 such term pairs, the value of v_{suff}^g can be computed in polynomial time under an independent
 2227 distribution $\mathcal{D}$. $\square$
 2228

2229 O DISCLOSURE: USAGE OF LLMs

2230
 2231 An LLM was used exclusively as a writing assistant to refine grammar and typos and improve clarity.
 2232 It did not contribute to the generation of research ideas, study design, data analysis, or interpretation
 2233 of results, all of which were carried out solely by the authors.
 2234
 2235
 2236
 2237
 2238
 2239
 2240
 2241
 2242
 2243
 2244
 2245
 2246
 2247
 2248
 2249
 2250
 2251
 2252
 2253
 2254
 2255
 2256
 2257
 2258
 2259
 2260
 2261
 2262
 2263
 2264
 2265
 2266
 2267