

Differentially Private Preference Data Synthesis for Large Language Model Alignment

Anonymous Authors¹

Abstract

Preference alignment has become a central post-training mechanism for making large language models (LLMs) reliably reflect human values. Yet alignment on real human preference data creates a direct privacy risk because these datasets often contain sensitive user prompts, model responses, and human judgments. We address this problem with **DPPrefSyn**, a new algorithm for generating differentially private (DP) synthetic preference data for privacy-preserving LLM alignment. DPPrefSyn provides a principled solution grounded in the Bradley–Terry preference model and the intrinsic geometric structure of pairwise human preference data. It learns a private preference model with formal DP guarantees, then combines that model with public prompts to synthesize high-quality preference pairs. The algorithm explicitly exploits the shared linear structure of per-cluster reward models to capture heterogeneous human preferences, and uses DP Principal Component Analysis (DP-PCA) to improve learning accuracy under privacy constraints. Across extensive experiments, DPPrefSyn delivers strong alignment utility while maintaining rigorous DP guarantees. The results establish DP synthetic preference data as a practical and reusable foundation for privacy-preserving preference alignment. To the best of our knowledge, DPPrefSyn is the first method for generating DP synthetic preference data for LLM alignment. Our code will be released upon acceptance.

1. Introduction

Preference alignment techniques, including RLHF (Stienon et al., 2020) and DPO (Rafailov et al., 2024), now serve as core tools for aligning LLM behavior with human expectations.

These algorithms depend on pairwise preference data: human annotators compare two responses to a prompt and identify the response that better satisfies the intended behavior. Alignment methods then fine-tune LLMs to rank preferred responses above less preferred ones (Ziegler et al., 2019; Ouyang et al., 2022; Bai et al., 2022). This paradigm has proven sufficiently effective and reliable to support high-impact applications, including chat assistants (Achiam et al., 2023), mathematical reasoning systems (Shao et al., 2024), and code generation models (Shen et al., 2023).

Despite its effectiveness, preference alignment on human data creates a substantial privacy challenge because preference datasets routinely contain real user prompts and human feedback. Prompts may reveal health information, identity attributes, or other sensitive topics, while feedback can expose private beliefs, preferences, and behavioral patterns (Li et al., 2023; Yu et al., 2024). Differential privacy (DP) (Dwork et al., 2006b) offers a rigorous way to control these risks, and recent work has made important progress in DP alignment (Chowdhury et al., 2024; Yu et al., 2024; Wu et al., 2023a). However, existing approaches leave a clear gap: many protect only one component of the data, such as user prompts (Yu et al., 2024) or preference labels (Chowdhury et al., 2024; Zhang et al., 2025), rather than the full prompt-response-label preference tuple. Partial protection is insufficient because adversaries can still exploit the unprotected component to infer sensitive attributes. Existing methods are also often tied to a particular post-training algorithm, such as RLHF (Wu et al., 2023a), limiting their compatibility with newer approaches such as DPO. Finally, direct private fine-tuning remains constrained by the scarcity and cost of human preference annotations, which limits alignment utility. These limitations motivate the central question of this work:

Is it possible to achieve superior preference alignment while ensuring thorough privacy protection and maintaining compatibility with emerging alignment methodologies?

This work answers the question affirmatively. We make the following contributions.

- We introduce DPPrefSyn (Figure 1), a DP synthetic preference data generation algorithm for LLM alignment.

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

ment. DPPrefSyn is a principled framework built on the Bradley–Terry preference model and the geometry of pairwise preference data. Its design has three decisive components: 1) *Capturing diverse human preferences via shared linear structure of rewards and clustering*. The Bradley–Terry model with a linear reward structure lets us cluster preference samples using embedding differences between concatenated prompt–preferred and prompt–dispreferred pairs, so each cluster can be represented by a shared reward function. 2) *Improving learning efficiency through DP-PCA*. DP-PCA projects high-dimensional embeddings into a compact subspace that retains the relevant preference signal while improving private learning and clustering. 3) *Saving privacy budgets by using public prompts*. Public prompts supply the prompt side of the synthetic data, while private reward models select preference pairs; this concentrates the privacy budget on the genuinely private preference signal and improves privacy efficiency.

- We provide a rigorous privacy analysis and establish that DPPrefSyn satisfies (ϵ, δ) -DP. The analysis allocates ϵ_0 to DP-PCA (Amin et al., 2019), ϵ_1 to DP-KMeans (Su et al., 2016), and the remaining budget to per-cluster reward models trained with DP-SGD. We use the Privacy Random Variable (PRV) accountant (Gopi et al., 2021) to tightly compose the DP-SGD privacy cost. By the post-processing property of DP, the resulting synthetic preference dataset can be reused across alignment algorithms and LLMs without additional privacy loss.
- We empirically validate DPPrefSyn on standard benchmarks, including OpenAssistant (Köpf et al., 2023) and Anthropic-HH (Bai et al., 2022) question answering, as well as TL;DR summarization (Stiennon et al., 2020). The experiments demonstrate that DPPrefSyn achieves strong privacy guarantees and high alignment utility across tasks. For example, with $\epsilon = 2$, DPO fine-tuning of Pythia-2.8B on DPPrefSyn data achieves a GPT-4o win rate of 56.48% on Anthropic-HH, substantially exceeding fine-tuning on real data (37.02%). The results show that DP synthetic preference data can drive effective LLM alignment under strong privacy constraints. DPPrefSyn also applies across preference optimization methods, including DPO and RLHF, avoiding the method-specific limitations of prior private post-training approaches.

2. Related Work

Differentially Private Alignment of LLMs. Wu et al. (2023a) introduce a DP framework for reinforcement-learning-based LLM alignment by adapting PPO to the DP setting. That contribution is important, but it is *limited to PPO* and therefore does not provide a reusable privacy layer for newer preference optimization methods. Chowdhury

et al. (2024) study reward estimation from preference-based feedback with label-DP (Ghazi et al., 2021), and Zhang et al. (2025) propose AUP-RLHF as a user-level label-DP framework for RLHF. These methods protect preference labels, but they *do not address the sensitivity of prompts or responses themselves*. They also remain *constrained by limited private user preference data*, which makes high-utility alignment difficult when human annotations are expensive. Yu et al. (2024) protect sensitive user *instructions* by generating DP synthetic *instructions* to replace real ones during data annotation and model fine-tuning.

Beyond preference alignment, prior work has established private fine-tuning of LLMs on general labeled data rather than preference comparisons (Yu et al., 2021; Li et al., 2021; Chen et al., 2024; Tang et al., 2024; Zhang et al., 2023; He et al., 2022). A separate line protects in-context prompts for in-context learning (Duan et al., 2023; Tang et al., 2023; Gao et al., 2024; Hong et al., 2023; Wu et al., 2023b). These directions address important privacy settings, but they do not solve the full private preference-data synthesis problem targeted here.

Differentially Private Synthetic Text Generation. DP-PrefSyn advances the broader agenda of DP synthetic text generation. One established line privately fine-tunes a language model on sensitive data and then uses the resulting model to generate synthetic text (Yue et al., 2023; Mattern et al., 2022; Mireshghallah et al., 2023; Carranza et al., 2023; Yu et al., 2024; Wang et al., 2024; Ochs & Habernal, 2024; Tan et al., 2025; Carranza et al., 2024). For example, Yu et al. (2024) privately fine-tune LLaMA on private instruction data and then use the model to generate synthetic text for supervised instruction tuning.

Because DP fine-tuning can be costly or infeasible, especially for non-public language models, recent work instead generates DP synthetic data by prompting LLM APIs (Lin et al., 2023; Xie et al., 2024; Wu et al., 2024; Hou et al., 2024). These methods generate candidate samples through API access and privately select outputs that match private data. For instance, Xie et al. (2024) propose Aug-PE, which iteratively refines LLM-generated samples according to embedding-space similarity to private data. *DPPrefSyn brings this API-based synthetic-data strategy to LLM preference alignment and extends it from text synthesis to pairwise preference synthesis.*

Diversity in Human Preferences. Human preferences are inherently heterogeneous (Denton et al., 2021; Aroyo et al., 2023a;b; Chakraborty et al., 2024). This heterogeneity reflects sociodemographic differences (e.g., race, gender, age), personal beliefs and biases, different levels of domain expertise, and ambiguity in natural language (Sandri et al., 2023; Vogels, 2021). Recent work has therefore pursued personalized alignment methods that better represent the values of




Figure 1. Overview of DPPrefSyn. The algorithm constructs DP synthetic preference data in three stages: 1) represent preference samples as embedding differences and cluster them with DP-PCA and DP-KMeans, 2) train a DP reward model for each cluster using DP-SGD, and 3) generate synthetic preference samples from public prompts using the DP-protected distribution of reward functions.

different user groups (Lee et al., 2024b; Poddar et al., 2024; Singh et al., 2025). For example, Singh et al. (2025) propose a meta-learning framework in which an LLM adapts to individual users from a small number of preference examples. DPPrefSyn directly incorporates this insight: it generates DP synthetic preference data that *captures the underlying diversity of human feedback*, and the resulting diversity modeling improves alignment performance relative to designs that ignore heterogeneity (Section 5.3).

3. Problem Definition

Preference Alignment of LLMs. We study LLM preference alignment from pairwise human feedback. The preference dataset is $\mathcal{D} = \{(x_i, a_i^+, a_i^-)\}_{i=1}^n$, where each sample contains a prompt x_i and two candidate responses a_i^+, a_i^- . Human annotators prefer a_i^+ to a_i^- , making each tuple a direct supervision signal for alignment.

Following the RLHF and DPO literature (Ouyang et al., 2022; Zhu et al., 2023; Rafailov et al., 2024; Liu et al., 2023), we use the Bradley–Terry model (Bradley & Terry, 1952) as the probabilistic foundation for these pairwise preferences:

$$\mathbb{P}[a_i^+ \succ a_i^- \mid x_i] = \frac{\exp(r^*(x_i, a_i^+))}{\exp(r^*(x_i, a_i^+)) + \exp(r^*(x_i, a_i^-))} = \sigma(r^*(x_i, a_i^+) - r^*(x_i, a_i^-)),$$

Here $a_i^+ \succ a_i^-$ denotes that a_i^+ is preferred to a_i^- , $r^*(x, a)$ is the latent reward function, and $\sigma(z) = 1/(1 + e^{-z})$ is the sigmoid function. This formulation is the common backbone of RLHF, which explicitly learns a reward model from preference data, and DPO, which induces an implicit reward through the log-ratio between the learned policy and a reference policy.

Differentially Private Preference Data Synthesis. Our goal is to protect the private preference dataset $\mathcal{D}_{\text{priv}} =$

$\{(x_i, a_i^+, a_i^-)\}_{i=1}^n$ from adversaries who may try to infer information about individual prompts, responses, or preference labels from the fine-tuned LLM. We therefore generate a DP synthetic preference dataset $\tilde{\mathcal{D}} = \{(\tilde{x}_j, \tilde{a}_j^+, \tilde{a}_j^-)\}_{j=1}^m$ that retains the preference signal needed for alignment while protecting the original tuples. By the post-processing property of DP, $\tilde{\mathcal{D}}$ becomes a reusable privacy-preserving asset: it can support multiple alignment methods and LLMs without additional privacy cost.

We use the standard definition of (ϵ, δ) -differential privacy.

Definition 3.1 ((ϵ, δ) -Differential Privacy (Dwork et al., 2006a)). A randomized algorithm $\mathcal{A}$ is (ϵ, δ) -differentially private if for any two neighboring inputs $\mathcal{D}$ and $\mathcal{D}'$ that differ by a single entry and any set $\mathcal{S}$ of possible outputs: $\mathbb{P}[\mathcal{A}(\mathcal{D}) \in \mathcal{S}] \leq e^\epsilon \mathbb{P}[\mathcal{A}(\mathcal{D}') \in \mathcal{S}] + \delta$.

Diverse User Preferences and Linear Reward Structure.

Human preferences in private alignment data are typically heterogeneous rather than homogeneous (Chakraborty et al., 2024; Bakker et al., 2022; Kovač et al., 2023; Jang et al., 2023; Rame et al., 2023; Ji et al., 2023). Annotators may prioritize factual accuracy, politeness, creativity, clarity, or other response attributes, producing diverse and sometimes conflicting preference signals. Effective synthesis of $\tilde{\mathcal{D}}$ therefore requires explicit modeling of the diversity present in $\mathcal{D}_{\text{priv}}$.

We capture this diversity with a linear structure for the latent reward function (Saha et al., 2023; Kong & Yang, 2022; Zhu et al., 2023). Specifically, we adopt the linear reward model

$$r_\theta(x, a) = \langle \theta, \phi(x, a) \rangle,$$

where $\phi(x, a) \in \mathbb{R}^d$ is the *shared* feature vector and $\theta \in \mathbb{R}^d$ is a *user-dependent* parameter.

This linear structure gives DPPrefSyn the right balance of *expressiveness* and *tractability*. Users may disagree, but

their reward functions can still share semantic features in the embedding space (i.e., $\phi(x, a)$) and differ mainly in feature weights (i.e., θ). The parameterization therefore captures shared preference structure while preserving variability across users or preference clusters, which is essential when data are limited and privacy constraints are strong. The empirical comparison in Section B.6 confirms that this linear reward structure captures preference signals effectively and remains competitive with a fully fine-tuned GPT-2 reward model despite its simplicity.

4. Proposed Method

4.1. Key Components of the Algorithm Design

We first highlight the algorithmic components that make DP preference synthesis accurate, privacy-efficient, and robust to heterogeneous preferences.

Leveraging Linear Reward Structure for Preference Clustering. The linear reward model indicates that

$$\mathbb{P}[a_i^+ \succ a_i^- \mid x_i] = \sigma(\langle \theta, \phi(x_i, a_i^+) - \phi(x_i, a_i^-) \rangle).$$

Thus, for a tuple (x_i, a_i^+, a_i^-) , $a_i^+ \succ a_i^-$ implies $\langle \theta, \phi(x_i, a_i^+) - \phi(x_i, a_i^-) \rangle \gg 0$: the preference vector θ aligns strongly with $\phi(x_i, a_i^+) - \phi(x_i, a_i^-)$. Annotators with similar preferences should therefore produce feature difference vectors pointing in similar directions in the embedding space. DPPrefSyn turns this observation into a concrete clustering rule: it groups preference samples by $\phi(x_i, a_i^+) - \phi(x_i, a_i^-)$ so that each cluster can be accurately approximated by a shared reward function with a common preference vector θ . Section B.9 confirms that these clusters are meaningful and interpretable as user preference styles.

Adopting DP-PCA to Obtain Effective Features. The embedding of a (prompt, response) pair is a natural feature vector $\phi(x, a)$ because it encodes semantic and stylistic information that drives preference judgments. These embeddings are high-dimensional (e.g., 1024 dimensions), which makes private clustering and reward learning sample-inefficient when each cluster contains only a limited number of examples. DPPrefSyn directly addresses this bottleneck with differentially private PCA (DP-PCA) (Amin et al., 2019; Liu et al., 2022). DP-PCA produces compact features for clustering and reward training, improving sample efficiency and clustering stability while retaining the preference signal needed for alignment.

Utilizing Public Prompts to Save DP Budget. Jointly synthesizing prompts, preferred responses, and less preferred responses would spend privacy budget on every component and weaken utility. DPPrefSyn instead uses public prompts, reserving the privacy budget for the private preference signal over responses. For each public prompt, an LLM generates multiple candidate responses, and a DP-trained distribu-

tion of reward functions selects the preference pair. Even when public and private prompt distributions differ, the experiments show that this strategy preserves the statistical preference distribution in the private dataset and remains effective for downstream alignment.

4.2. DPPrefSyn Algorithm

We now present DPPrefSyn (Algorithm 1), which converts the private dataset $\mathcal{D}_{\text{priv}}$ into a DP synthetic preference dataset. The algorithm proceeds in three main steps.

Step 1: Preference Representation and Clustering. DPPrefSyn first represents each private preference comparison in an embedding-difference space. For each sample (x_i, a_i^+, a_i^-) , it concatenates the prompt with each response to form $[x_i; a_i^+]$ and $[x_i; a_i^-]$, embeds both texts with the public sentence embedding model ψ , and records their difference $d_i \leftarrow \psi(x_i, a_i^+) - \psi(x_i, a_i^-)$ (Line 3).

DPPrefSyn then applies DP-PCA (Amin et al., 2019) to $\{d_i\}_{i=1}^n$, producing a projection matrix $\mathbf{P} \in \mathbb{R}^{d \times p}$ under ϵ_0 -DP (Line 5). DP-PCA (Amin et al., 2019) privately approximates the covariance eigendecomposition by estimating eigenvalues and eigenvectors in a DP manner. Each d_i is projected as $z_i \leftarrow \mathbf{P}^\top d_i$ (Line 6). DP-KMeans (Su et al., 2016), a DP version of the Lloyd algorithm (Su et al., 2016), then clusters $\{z_i\}_{i=1}^n$ under ϵ_1 -DP to form K preference groups $\mathcal{C}_1, \dots, \mathcal{C}_K$ (Line 7).

Step 2: DP Reward Model Training. For each cluster $\mathcal{C}_k$, DPPrefSyn learns a shared preference vector $\theta_k \in \mathbb{R}^p$ by minimizing the Bradley-Terry negative log-likelihood (Line 9):

$$\mathcal{L}_{\text{MLE}}(\theta_k) = -\frac{1}{|\mathcal{C}_k|} \sum_{z_i \in \mathcal{C}_k} \log \sigma(\langle \theta_k, z_i \rangle), \quad (1)$$

where z_i is the DP-PCA-projected embedding difference from Step 1. Each cluster reward model is trained with DP-SGD (Abadi et al., 2016), the DP variant of SGD, using noise multiplier σ_2 .

Step 3: DP Synthetic Preference Generation. DPPrefSyn computes the cluster histogram $\mathbf{h} = [|\mathcal{C}_1|, \dots, |\mathcal{C}_K|]$ (Line 11) and normalizes it into a cluster sampling distribution $\mathbf{p} \leftarrow \mathbf{h}/|\mathcal{D}_{\text{priv}}|$ (Line 12).

For each public prompt $\tilde{x}_j$, DPPrefSyn prompts an LLM at a relatively high temperature to generate L diverse candidate responses $\tilde{a}_j^1, \dots, \tilde{a}_j^L$ (Line 14). It samples a cluster index $k \sim \mathbf{p}$ from the DP histogram (Line 15) and scores the candidate responses with the reward model θ_k . Specifically, it embeds each $\psi(\tilde{x}_j, \tilde{a}_j^l)$ for $l = 1, \dots, L$, projects the embedding into the DP-PCA subspace using $\mathbf{P}$, and computes the reward $\langle \theta_k, \mathbf{P}^\top \psi(\tilde{x}_j, \tilde{a}_j^l) \rangle$. The highest- and lowest-scoring responses become the preferred and less preferred responses: $c \leftarrow \arg \max_{l \in [L]} \langle \theta_k, \mathbf{P}^\top \psi(\tilde{x}_j, \tilde{a}_j^l) \rangle$,

$r \leftarrow \arg \min_{l \in [L]} \langle \theta_k, \mathbf{P}^\top \psi(\tilde{x}_j, \tilde{a}_j^l) \rangle$ (Lines 16-17). When the value gap is too small (e.g., < 0.5), DPPrefSyn discards the sample to maintain preference quality; otherwise, it adds $(\tilde{x}_j, \tilde{a}_j^c, \tilde{a}_j^r)$ to $\tilde{\mathcal{D}}$. Repeating this process over public prompts yields the DP synthetic preference dataset $\tilde{\mathcal{D}}$.

Algorithm 1 DPPrefSyn

```

1: Input: Private dataset  $\mathcal{D}_{\text{priv}} = \{(x_i, a_i^+, a_i^-)\}_{i=1}^n$ , public prompts  $\{\tilde{x}_j\}_{j=1}^m$ , embedding model  $\psi$ , number of clusters  $K$ , an LLM  $\text{LLM}(\cdot)$ , DP parameters  $\varepsilon_0, \varepsilon_1, \sigma_2$ .
2: for each  $(x_i, a_i^+, a_i^-)$  in  $\mathcal{D}_{\text{priv}}$  do
3:    $d_i \leftarrow \psi(x_i, a_i^+) - \psi(x_i, a_i^-)$ , where  $d_i \in \mathbb{R}^d$ 
4: end for
5:  $\mathbf{P} \leftarrow \text{DP-PCA}(\{d_i\}_{i=1}^n, \varepsilon_0)$ , where  $\mathbf{P} \in \mathbb{R}^{d \times p}$ 
6:  $\{z_i\}_{i=1}^n \leftarrow \{\mathbf{P}^\top d_i\}_{i=1}^n$ 
7:  $\mathcal{C}_1, \dots, \mathcal{C}_K \leftarrow \text{DP-KMeans}(\{z_i\}_{i=1}^n, K, \varepsilon_1)$ 
8: for each cluster  $\mathcal{C}_k$  do
9:   Train a linear reward model  $\theta_k$  on  $\mathcal{C}_k$  using DP-SGD with noise multiplier  $\sigma_2$  (Equation (1))
10: end for
11:  $\mathbf{h} \leftarrow [|\mathcal{C}_1|, \dots, |\mathcal{C}_K|]$ 
12:  $\mathbf{p} \leftarrow \mathbf{h} / |\mathcal{D}_{\text{priv}}|$ 
13: for each public prompt  $\tilde{x}_j$  do
14:   Generate responses  $\tilde{a}_j^1, \dots, \tilde{a}_j^L \leftarrow \text{LLM}(\tilde{x}_j)$ 
15:   Sample cluster index  $k \sim \mathbf{p}$ 
16:    $c \leftarrow \arg \max_{l \in [L]} \langle \theta_k, \mathbf{P}^\top \psi(\tilde{x}_j, \tilde{a}_j^l) \rangle$ 
17:    $r \leftarrow \arg \min_{l \in [L]} \langle \theta_k, \mathbf{P}^\top \psi(\tilde{x}_j, \tilde{a}_j^l) \rangle$ 
18:   Add  $(\tilde{x}_j, \tilde{a}_j^c, \tilde{a}_j^r)$  to  $\tilde{\mathcal{D}}$ 
19: end for
20: return  $\tilde{\mathcal{D}}$ .

```

Privacy Analysis of DPPrefSyn. DPPrefSyn assigns privacy budget ε_0 to DP-PCA (Line 5) and ε_1 to DP-KMeans (Line 7); the remaining budget trains per-cluster reward models with DP-SGD and noise multiplier σ_2 (Line 9). Reward-model training uses DP parallel composition: because clusters are disjoint, changing one training point affects only one cluster, so the total privacy cost is bounded by the cost of training on the smallest retained cluster. We enforce a known lower bound on cluster size by discarding clusters with too few samples. The Privacy Random Variable (PRV) accountant (Gopi et al., 2021) composes the privacy costs of DP-SGD steps and ensures $(\varepsilon - \varepsilon_0 - \varepsilon_1, \delta)$ -DP for reward training. All remaining steps use public data or DP post-processing and add no privacy cost. Section A gives the full proof and privacy-accounting code.

Discussion on Prompt Distribution Shift. DPPrefSyn is designed to remain effective when private and public prompt distributions differ. The key point is that preference diversity is represented through θ_k , the linear reward parameter, rather than solely through $\phi(x, a)$, the feature vector of a particular prompt-response pair. A user can ask many kinds of prompts while retaining a stable underlying preference

style. Therefore, if the distribution of θ_k matches the private preference distribution, the synthetic comparisons still reflect private preference patterns. DPPrefSyn enforces this match by sampling a cluster index k from the DP histogram (Line 15) for each public prompt and selecting response pairs with the corresponding reward model. As a result, DPPrefSyn preserves the preference patterns learned from private data even under prompt distribution shift.

5. Experiments

Datasets. We evaluate DPPrefSyn on both question answering and summarization, covering the major settings where preference alignment is routinely applied. For question answering, we use OpenAssistant (Köpf et al., 2023), which contains assistant-style conversations, and Anthropic-HH (Bai et al., 2022), which provides helpfulness- and harmless-oriented preference comparisons. For summarization, we use TL;DR (Stiennon et al., 2020), which contains human preferences over summary pairs. To instantiate public prompts, we use Alpaca (Taori et al., 2023) for OpenAssistant QA, SafeRLHF (Ji et al., 2024) for Anthropic-HH QA, and XSum (Narayan et al., 2018) for TL;DR summarization. Section B.2 gives additional dataset details.

Setup for Algorithm 1. Our default configuration uses BAAI/bge-large-en-v1.5 (Xiao et al., 2023) as the embedding model ψ , projected dimension $p = 20$ for DP-PCA, and $K = 5$ clusters for DP-KMeans. For DP-SGD, we use learning rate 0.1, batch size 4, 4 training epochs, and gradient clipping norm 1.0. We generate candidate responses with the instruction-finetuned LLaMA-7B-chat model (Touvron et al., 2023) at temperature 0.9 and use $L = 5$ candidates per prompt. We evaluate privacy budgets $\varepsilon = 0.5, 1, 2, 4, 8$ with $\delta = 1/|\mathcal{D}_{\text{priv}}|$. Table 3 studies generator models, and Section 5.3 studies embedding models and hyperparameters.

Evaluation. For the main downstream alignment experiments, we fine-tune Pythia-2.8B (Biderman et al., 2023). We first perform supervised fine-tuning (SFT), using the preferred response as the training target, and then apply DPO (Rafailov et al., 2024) to further optimize the SFT model with preference pairs. Following Rafailov et al. (2024), we use GPT-4o to measure the win rate of model-generated responses against preferred test-set responses. Section 5.2 extends the evaluation to additional downstream models and preference alignment methods.

Baselines. We compare against DP-FT, which fine-tunes the downstream model directly on real preference data using DP-SGD (Abadi et al., 2016). This baseline highlights a central advantage of DPPrefSyn: it produces a DP synthetic dataset that can be reused to train different models with different alignment algorithms at no additional privacy cost, whereas

DP-FT spends privacy budget to train a single privatized model. We also include the fully private $\varepsilon = 0$ baseline, which evaluates the base model without preference fine-tuning. Section B.3 provides the full setup, hyperparameters, metrics, and baseline details.

5.1. Understanding the Performance of DPPrefSyn

We first evaluate the central privacy-utility claim of DPPrefSyn. Table 1 reports the main results and supports two research questions.

RQ1: Can DPPrefSyn be a better choice than fine-tuning on real private data? DPPrefSyn delivers a clearly stronger privacy-utility trade-off than DP-FT across all tasks at $\varepsilon = 0.5, 1, 2, 4, 8$, and it exceeds the utility of direct fine-tuning on real private data without DP constraints (DP-FT with $\varepsilon = \infty$) while retaining privacy protection. On OpenAssistant QA at $\varepsilon = 2$, DPPrefSyn reaches 9.10% after SFT and 11.04% after DPO, more than doubling the DP-FT SFT result of 4.38% and substantially improving on the DP-FT DPO result of 6.18%. The same setting also surpasses the non-private DP-FT baseline, which reaches only 4.75% after SFT and 8.20% after DPO. The pattern is equally decisive on Anthropic-HH: at $\varepsilon = 2$, DPPrefSyn achieves 55.21% after SFT and 56.48% after DPO, compared with 31.53% and 37.02% for DP-FT, and it remains ahead of non-private direct fine-tuning, which obtains 31.98% after SFT and 38.72% after DPO.

These results demonstrate two **key advantages** of DP synthetic preference data. **First**, DPPrefSyn provides a large SFT-stage gain over directly applying SFT to private data. On TL;DR at $\varepsilon = 2$, DPPrefSyn reaches 62.49% after SFT, far above the 22.57% obtained by SFT on private data. This improvement follows from the core design choice of generating candidate responses with a high-performance LLM, which converts private preference signals into stronger supervised training data. The benefit persists even when public prompts are scarce: on OpenAssistant QA with $\varepsilon = 4$ and only 5K public prompts, DPPrefSyn reaches 6.80% after SFT, compared with 4.46% for SFT on the roughly 14K-sample private dataset (Table 4). **Second**, DPO on top of the DPPrefSyn SFT model adds further gains, confirming that the synthetic preference pairs encode useful human-preference distinctions rather than only fluent responses. The private reward functions make this possible by preserving heterogeneous preference information under DP. Although DPO adds less on top of DPPrefSyn than on top of direct private-data SFT in some cases, the full DPPrefSyn pipeline still delivers substantially larger end-to-end gains than training on real private data.

RQ2: How does DPPrefSyn perform across different privacy budget ε ? First, Table 1 shows that DPPrefSyn maintains strong performance as ε ranges from 0.5, 1, 2, 4, 8 to ∞ ,

demonstrating favorable scaling with privacy budget. **Second**, DPPrefSyn is more robust to DP noise than DP-FT in several settings. For example, on OpenAssistant with Llama-3-8B as the downstream model, DPO performance for DP-FT drops from 59.13% to 45.79% when ε decreases from ∞ to 4, whereas DPPrefSyn changes only slightly from 61.52% to 60.25% (Table 3 in Section 5.2). This pattern is consistent with the design of DPPrefSyn: DP-FT perturbs model parameters directly through DP-SGD, while DPPrefSyn injects privacy noise during data synthesis and thereby better preserves downstream utility.

5.2. Understanding the Properties of DPPrefSyn

We next test whether DPPrefSyn’s advantages persist beyond the main configuration. The experiments examine generalization across alignment methods and LLMs, scaling with public data, and robustness against empirical privacy attacks.

RQ3: Is DPPrefSyn effective across different preference alignment methods? DPPrefSyn is method-agnostic and remains effective under strong DP constraints. Table 2 evaluates DPPrefSyn with SFT, DPO, and PPO. For PPO, we train K reward models, one per cluster, to represent diverse preference patterns; PPO then optimizes the policy using the average output of these K reward models. DPPrefSyn achieves higher utility than fine-tuning on private data for all three alignment methods while preserving strong privacy guarantees.

RQ4: Can more powerful downstream LLMs benefit from synthetic data generated by DPPrefSyn? DPPrefSyn transfers across downstream LLM families. Table 3 evaluates Qwen-3-4B, Gemma-2-2B, and Llama-3-8B as downstream models, using Llama-3-8B-Instruct and Qwen3-4B-Instruct-2507 as generators. DPPrefSyn outperforms fine-tuning on real data across model families after DPO while retaining privacy protection. For example, with Qwen3-4B-Instruct-2507 as the generator, DPPrefSyn reaches an 86.52% win rate on Qwen-3-4B at $\varepsilon = 4$, far above non-private fine-tuning on real data (59.41%). Stronger generator models also improve DPPrefSyn, with Qwen3-4B-Instruct-2507 outperforming Llama-3-8B-Instruct.

RQ5: Can DPPrefSyn be effective with varying scales of public data? Table 4 shows that DPPrefSyn remains strong even with limited public prompts. OpenAssistant contains 14K private preference annotations, while the full Alpaca public prompt pool contains 52K prompts. We sample public subsets of size $m = 5K, 10K, 52K$ and synthesize preference data with DPPrefSyn. With only 5K public prompts, DPPrefSyn reaches 6.80% after SFT and 9.44% after DPO at $\varepsilon = 4$, outperforming DP-FT at 4.46% and 6.12%. Thus, **DPPrefSyn remains effective even when public data is limited**, because LLM-generated responses provide high-

Table 1. GPT-4o win rates on test sets for the downstream Pythia-2.8B model fine-tuned with SFT and DPO. $\varepsilon = 0$ denotes the base fully private, non-finetuned LLM. *DPPrefSyn* consistently achieves higher win rates than *DP-FT* at $\varepsilon = 0.5, 1, 2, 4, 8$ and even surpasses non-private *DP-FT* ($\varepsilon = \infty$). Results show mean and standard deviation over 5 random seeds.

Task	$\varepsilon = 0$	Method	Data Type	Pref. Alignment	$\varepsilon = 0.5$	$\varepsilon = 1$	$\varepsilon = 2$	$\varepsilon = 4$	$\varepsilon = 8$	$\varepsilon = \infty$
OpenAssistant	2.11	DP-FT	Original	SFT	3.96 _{0.20}	4.29 _{0.21}	4.38 _{0.23}	4.46 _{0.18}	4.49 _{0.30}	4.75 _{0.27}
				SFT+DPO	5.65 _{0.32}	6.04 _{0.58}	6.18 _{0.83}	6.12 _{0.63}	6.32 _{0.71}	8.20 _{0.62}
		DPPrefSyn	Synthetic	SFT	8.96 _{0.39}	9.07 _{0.32}	9.10 _{0.51}	9.33 _{0.56}	9.63 _{0.38}	9.80 _{0.58}
				SFT+DPO	9.86_{0.59}	10.31_{0.54}	11.04_{0.92}	11.85_{0.99}	12.05_{0.82}	12.36_{0.55}
Anthropic-HH	12.14	DP-FT	Original	SFT	29.77 _{0.76}	31.10 _{0.72}	31.53 _{0.77}	31.29 _{0.65}	31.37 _{0.67}	31.98 _{0.56}
				SFT+DPO	35.00 _{0.49}	36.27 _{0.63}	37.02 _{0.72}	36.74 _{0.90}	36.94 _{1.06}	38.72 _{0.56}
		DPPrefSyn	Synthetic	SFT	54.90 _{0.66}	55.19 _{0.91}	55.21 _{0.82}	54.95 _{1.10}	55.63 _{0.86}	55.95 _{0.70}
				SFT+DPO	55.08_{1.04}	55.96_{0.80}	56.48_{0.61}	56.51_{1.93}	56.86_{1.17}	57.53_{0.93}
TL;DR	11.64	DP-FT	Original	SFT	21.19 _{1.12}	22.41 _{1.09}	22.57 _{0.65}	22.45 _{0.90}	23.04 _{0.79}	23.82 _{1.26}
				SFT+DPO	59.69 _{0.82}	62.21 _{1.12}	62.02 _{1.67}	63.00 _{0.73}	63.06 _{0.84}	64.52 _{1.79}
		DPPrefSyn	Synthetic	SFT	61.92 _{1.06}	62.14 _{1.61}	62.49 _{1.48}	62.72 _{1.66}	62.29 _{1.14}	62.55 _{1.23}
				SFT+DPO	67.50_{0.79}	68.07_{1.95}	68.54_{2.76}	68.56_{1.59}	71.01_{2.17}	72.20_{1.08}

Table 2. DPPrefSyn outperforms fine-tuning on private data across SFT, DPO, and PPO while providing strong privacy guarantees on OpenAssistant. Results show mean and standard deviation over 5 random seeds.

Pref. Alignment	$\varepsilon = 0$	DPPrefSyn ($\varepsilon = 4$)	DP-FT ($\varepsilon = \infty$)
SFT		9.33_{0.56}	4.75 _{0.27}
SFT+DPO	2.11	11.85_{0.99}	8.20 _{0.62}
SFT+PPO		10.11_{0.61}	7.08 _{0.78}

Table 3. DPPrefSyn achieves higher DPO win rates than fine-tuning on private data across downstream LLMs while providing strong privacy guarantees on OpenAssistant. Stronger generator LLMs further improve DPPrefSyn.


Downstream LLM	$\varepsilon = 0$	Method	Generator LLM	$\varepsilon = 4$	$\varepsilon = \infty$
Qwen-3-4B	48.17	DP-FT	-	54.92	59.41
			-	54.92	59.41
		DPPrefSyn	Qwen-3-4B-Instruct	86.52	87.08
			Llama-3-8B-Instruct	51.69	52.11
Gemma-2-2B	7.30	DP-FT	-	30.06	33.85
			-	30.06	33.85
		DPPrefSyn	Qwen-3-4B-Instruct	36.24	37.08
			Llama-3-8B-Instruct	36.01	36.66
Llama-3-8B	16.15	DP-FT	-	45.79	59.13
			-	45.79	59.13
		DPPrefSyn	Qwen-3-4B-Instruct	60.25	61.52
			Llama-3-8B-Instruct	42.28	44.10

Table 4. DPPrefSyn outperforms fine-tuning on private data across public data sizes m . Results show mean and standard deviation over 5 random seeds.

	DPPrefSyn ($\varepsilon = 4$)			DP-FT ($\varepsilon = 4$)	DP-FT ($\varepsilon = \infty$)
Data Type (Size)	Synthetic ($m = 5,000$)	Synthetic ($m = 10,000$)	Synthetic ($m = 52,000$)	Original (14,167)	Original (14,167)
SFT	6.80 _{0.66}	7.30 _{0.48}	9.33 _{0.56}	4.46 _{0.18}	4.75 _{0.27}
SFT+DPO	9.44 _{0.45}	9.55 _{0.92}	11.85 _{0.99}	6.12 _{0.63}	8.20 _{0.62}

Table 5. DPPrefSyn yields lower AUC scores against MIAs (Feng et al., 2025) on OpenAssistant, approaching random guessing. Results show mean and standard deviation over 5 random seeds.

Downstream LLM	Method	Generator LLM	$\varepsilon = 4$	$\varepsilon = \infty$
Qwen-3-4B	DP-FT	-	70.15 _{2.09}	73.52 _{2.30}
		-	70.15 _{2.09}	73.52 _{2.30}
	DPPrefSyn	Qwen-3-4B-Instruct	50.99 _{1.10}	50.87 _{0.84}
		Llama-3-8B-Instruct	50.64_{1.29}	49.98_{1.19}
Gemma-2-2B	DP-FT	-	62.88 _{1.78}	65.39 _{2.59}
		-	62.88 _{1.78}	65.39 _{2.59}
	DPPrefSyn	Qwen-3-4B-Instruct	50.51_{1.38}	50.43_{0.91}
		Llama-3-8B-Instruct	50.77 _{0.83}	51.10 _{0.54}
Llama-3-8B	DP-FT	-	61.95 _{2.05}	70.13 _{1.24}
		-	61.95 _{2.05}	70.13 _{1.24}
	DPPrefSyn	Qwen-3-4B-Instruct	50.53 _{0.62}	50.96 _{0.95}
		Llama-3-8B-Instruct	50.45_{0.87}	50.67_{0.96}



(a) Removing either DP-KMeans or DP-PCA weakens performance. (b) DPPrefSyn remains stable across moderate cluster counts.

quality candidates and private reward functions preserve user preference signals. Moreover, **the margin over the baseline grows as public data increases**, showing that DPPrefSyn can capitalize on abundant public prompts and directly mitigates private preference-data scarcity.

RQ6: Can DPPrefSyn effectively mitigate empirical privacy attacks? Table 5 reports membership inference attacks (MIAs) (Shokri et al., 2017) against DPO-fine-tuned downstream models on OpenAssistant, using AUC as the metric. Following Feng et al. (2025), the attack tests whether a private sample (x, a^+, a^-) was used in preference alignment. DP-FT remains vulnerable in this empirical evaluation, whereas **DPPrefSyn consistently produces lower AUC scores close to random guessing**. This result indi-

cates stronger empirical privacy robustness, plausibly because fine-tuning on synthetic data reduces overfitting to real private tuples. Section B.4 provides the details.

5.3. Validating the Design of DPPrefSyn

We now validate that the key DPPrefSyn design choices are necessary and effective.

RQ7: How does modeling diverse human preferences impact DPPrefSyn’s performance? We isolate this effect on OpenAssistant at $\varepsilon = 4$ by removing DP-KMeans or DP-PCA and by varying the number of clusters K . **(i) Removing either DP-KMeans or DP-PCA.** Figure 2a shows that each component matters: removing DP-KMeans collapses all samples into one cluster, while removing DP-PCA makes clustering and reward learning less effective in the high-dimensional embedding space. The full design, with multiple clusters and DP-PCA, delivers higher performance than the single-cluster alternative. **(ii) Varying the number of clusters K .** Figure 2b evaluates $K = 1, 3, 5, 8$. Performance improves at $K = 3$ and remains stable for $K = 5, 8$, showing that DPPrefSyn is robust to moderate choices of K . These results confirm that explicitly capturing diverse preferences with multiple clusters is a core source of DPPrefSyn’s utility.

RQ8: How do different embedding models or hyperparameters affect DPPrefSyn’s performance? Section B.5 ablates the embedding model ψ , DP-PCA dimension p , candidate count L , filtering threshold, and privacy budget allocation. DPPrefSyn remains consistently effective across sentence encoders, including bge-large-en-v1.5 (Xiao et al., 2023), all-mpnet-base-v2 (Song et al., 2020), and all-distilroberta-v1 (Sanh et al., 2019). The ablations identify $p = 20$ as the strongest DPO setting, $L = 5$ as a reliable default, and threshold 0.5 as a useful moderate filter; threshold 1 is too aggressive and reduces utility. DP-PrefSyn also remains stable under reasonable changes to ε_0 and ε_1 .

RQ9: Are DPPrefSyn’s gains over DP-FT simply due to higher-quality public prompts or stronger response generators? Controlled ablations rule out these confounders by forcing methods to share the same prompt set or response generator. DPPrefSyn still outperforms DP-FT (Appendix B.7), demonstrating that **its gains are not merely an artifact of higher-quality public prompts or stronger response generators.**

6. Conclusion

We introduced DPPrefSyn, a DP synthetic preference data algorithm for LLM alignment. By clustering heterogeneous preferences, improving private learning with DP-PCA, and using public prompts to reserve privacy budget for the pref-

erence signal, DPPrefSyn turns sensitive preference datasets into reusable DP synthetic alignment data. The experiments show that this strategy delivers rigorous privacy guarantees and consistently outperforms fine-tuning on real data across the evaluated settings. DPPrefSyn therefore establishes DP synthetic preference data as a practical and effective route to privacy-preserving LLM alignment.

7. Limitations

DPPrefSyn relies on LLM API access for candidate response generation, so deployment may incur additional API costs. The current study focuses on English text-based preference data; extending the framework to multilingual and multimodal preference settings is an important next step. Our experiments also cover models up to 8B parameters because of computational constraints, leaving larger-scale evaluation for future work.

Impact statement

This work advances preference alignment of large language models without exposing personal information. It helps reconcile the utility of large-scale alignment data with rigorous privacy protection, supporting more ethical data use and strengthening trust in data-driven AI systems.

References

- Abadi, M., Chu, A., Goodfellow, I., McMahan, H. B., Mironov, I., Talwar, K., and Zhang, L. Deep learning with differential privacy. In *Proceedings of the 2016 ACM SIGSAC conference on computer and communications security*, pp. 308–318, 2016.
- Achiam, J., Adler, S., Agarwal, S., Ahmad, L., Akkaya, I., Aleman, F. L., Almeida, D., Altenschmidt, J., Altman, S., Anadkat, S., et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- Amin, K., Dick, T., Kulesza, A., Munoz, A., and Vassilvitskii, S. Differentially private covariance estimation. *Advances in Neural Information Processing Systems*, 32, 2019.
- Aroyo, L., Diaz, M., Homan, C., Prabhakaran, V., Taylor, A., and Wang, D. The reasonable effectiveness of diverse evaluation data. *arXiv preprint arXiv:2301.09406*, 2023a.
- Aroyo, L., Taylor, A., Diaz, M., Homan, C., Parrish, A., Serapio-García, G., Prabhakaran, V., and Wang, D. Dices dataset: Diversity in conversational ai evaluation for safety. *Advances in Neural Information Processing Systems*, 36:53330–53342, 2023b.

- Bai, Y., Jones, A., Ndousse, K., Askell, A., Chen, A., Das-Sarma, N., Drain, D., Fort, S., Ganguli, D., Henighan, T., et al. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*, 2022.
- Bakker, M., Chadwick, M., Sheahan, H., Tessler, M., Campbell-Gillingham, L., Balaguer, J., McAleese, N., Glaese, A., Aslanides, J., Botvinick, M., et al. Fine-tuning language models to find agreement among humans with diverse preferences. *Advances in Neural Information Processing Systems*, 35:38176–38189, 2022.
- Biderman, S., Schoelkopf, H., Anthony, Q. G., Bradley, H., O’Brien, K., Hallahan, E., Khan, M. A., Purohit, S., Prashanth, U. S., Raff, E., et al. Pythia: A suite for analyzing large language models across training and scaling. In *International Conference on Machine Learning*, pp. 2397–2430. PMLR, 2023.
- Bradley, R. A. and Terry, M. E. Rank analysis of incomplete block designs: I. the method of paired comparisons. *Biometrika*, 39(3/4):324–345, 1952.
- Carranza, A., Farahani, R., Ponomareva, N., Kurakin, A., Jagielski, M., and Nasr, M. Synthetic query generation for privacy-preserving deep retrieval systems using differentially private language models. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 3920–3930, 2024.
- Carranza, A. G., Farahani, R., Ponomareva, N., Kurakin, A., Jagielski, M., and Nasr, M. Privacy-preserving recommender systems with synthetic query generation using differentially private large language models. *arXiv preprint arXiv:2305.05973*, 2023.
- Chakraborty, S., Qiu, J., Yuan, H., Koppel, A., Manocha, D., Huang, F., Bedi, A., and Wang, M. Maxmin-rlhf: Alignment with diverse human preferences. In *International Conference on Machine Learning*, pp. 6116–6135. PMLR, 2024.
- Chen, T., Da, L., Zhou, H., Li, P., Zhou, K., Chen, T., and Wei, H. Privacy-preserving fine-tuning of large language models through flatness. *arXiv preprint arXiv:2403.04124*, 2024.
- Chowdhury, S. R., Zhou, X., and Natarajan, N. Differentially private reward estimation with preference feedback. In *International Conference on Artificial Intelligence and Statistics*, pp. 4843–4851. PMLR, 2024.
- Denton, R., Díaz, M., Kivlichan, I., Prabhakaran, V., and Rosen, R. Whose ground truth? accounting for individual and collective identities underlying dataset annotation. *arXiv preprint arXiv:2112.04554*, 2021.
- Duan, H., Dziedzic, A., Papernot, N., and Boenisch, F. Flocks of stochastic parrots: Differentially private prompt learning for large language models. *Advances in Neural Information Processing Systems*, 36:76852–76871, 2023.
- Dwork, C., Kenthapadi, K., McSherry, F., Mironov, I., and Naor, M. Our data, ourselves: Privacy via distributed noise generation. In *Advances in Cryptology-EUROCRYPT 2006: 24th Annual International Conference on the Theory and Applications of Cryptographic Techniques, St. Petersburg, Russia, May 28-June 1, 2006. Proceedings 25*, pp. 486–503. Springer, 2006a.
- Dwork, C., McSherry, F., Nissim, K., and Smith, A. Calibrating noise to sensitivity in private data analysis. In *Theory of Cryptography: Third Theory of Cryptography Conference, TCC 2006, New York, NY, USA, March 4-7, 2006. Proceedings 3*, pp. 265–284. Springer, 2006b.
- Dwork, C., Roth, A., et al. The algorithmic foundations of differential privacy. *Foundations and Trends® in Theoretical Computer Science*, 9(3–4):211–407, 2014.
- Feng, Q., Kasa, S. R., KASA, S. K., Yun, H., Teo, C. H., and Bodapati, S. B. Exposing privacy gaps: Membership inference attack on preference data for llm alignment. In *International Conference on Artificial Intelligence and Statistics*, pp. 5221–5229. PMLR, 2025.
- Gao, F., Zhou, R., Wang, T., Shen, C., and Yang, J. Data-adaptive differentially private prompt synthesis for in-context learning. *arXiv preprint arXiv:2410.12085*, 2024.
- Ghazi, B., Golowich, N., Kumar, R., Manurangsi, P., and Zhang, C. Deep learning with label differential privacy. *Advances in neural information processing systems*, 34: 27131–27145, 2021.
- Gopi, S., Lee, Y. T., and Wutschitz, L. Numerical composition of differential privacy. *Advances in Neural Information Processing Systems*, 34:11631–11642, 2021.
- He, J., Li, X., Yu, D., Zhang, H., Kulkarni, J., Lee, Y. T., Backurs, A., Yu, N., and Bian, J. Exploring the limits of differentially private deep learning with group-wise clipping. *arXiv preprint arXiv:2212.01539*, 2022.
- Holohan, N., Braghin, S., Mac Aonghusa, P., and Levacher, K. Diffprivlib: the IBM differential privacy library. *ArXiv e-prints*, 1907.02444 [cs.CR], July 2019.
- Hong, J., Wang, J. T., Zhang, C., Li, Z., Li, B., and Wang, Z. Dp-opt: Make large language model your privacy-preserving prompt engineer. *arXiv preprint arXiv:2312.03724*, 2023.

- Hou, C., Shrivastava, A., Zhan, H., Conway, R., Le, T., Sagar, A., Fanti, G., and Lazar, D. Pre-text: Training language models on private federated data in the age of llms. In *International Conference on Machine Learning*, pp. 19043–19061. PMLR, 2024.
- Jang, J., Kim, S., Lin, B. Y., Wang, Y., Hessel, J., Zettlemoyer, L., Hajishirzi, H., Choi, Y., and Ammanabrolu, P. Personalized soups: Personalized large language model alignment via post-hoc parameter merging. *arXiv preprint arXiv:2310.11564*, 2023.
- Ji, J., Liu, M., Dai, J., Pan, X., Zhang, C., Bian, C., Chen, B., Sun, R., Wang, Y., and Yang, Y. Beavertails: Towards improved safety alignment of llm via a human-preference dataset. *Advances in Neural Information Processing Systems*, 36:24678–24704, 2023.
- Ji, J., Hong, D., Zhang, B., Chen, B., Dai, J., Zheng, B., Qiu, T., Li, B., and Yang, Y. Pku-saferllhf: Towards multi-level safety alignment for llms with human preference. *arXiv preprint arXiv:2406.15513*, 2024.
- Kirk, R., Mediratta, I., Nalmpantis, C., Luketina, J., Hambro, E., Grefenstette, E., and Raileanu, R. Understanding the effects of rlhf on llm generalisation and diversity. In *International Conference on Learning Representations*, volume 2024, pp. 20620–20653, 2024.
- Kong, D. and Yang, L. Provably feedback-efficient reinforcement learning via active reward learning. *Advances in Neural Information Processing Systems*, 35:11063–11078, 2022.
- Köpf, A., Kilcher, Y., Von Rütte, D., Anagnostidis, S., Tam, Z. R., Stevens, K., Barhoum, A., Nguyen, D., Stanley, O., Nagyfi, R., et al. Openassistant conversations-democratizing large language model alignment. *Advances in Neural Information Processing Systems*, 36:47669–47681, 2023.
- Kovač, G., Sawayama, M., Portelas, R., Colas, C., Dominey, P. F., and Oudeyer, P.-Y. Large language models as superpositions of cultural perspectives. *arXiv preprint arXiv:2307.07870*, 2023.
- Lee, H., Phatale, S., Mansoor, H., Mesnard, T., Ferret, J., Lu, K. R., Bishop, C., Hall, E., Carbune, V., Rastogi, A., et al. Rlaif vs. rlhf: Scaling reinforcement learning from human feedback with ai feedback. In *International Conference on Machine Learning*, pp. 26874–26901. PMLR, 2024a.
- Lee, Y., Williams, J., Marklund, H., Sharma, A., Mitchell, E., Singh, A., and Finn, C. Test-time alignment via hypothesis reweighting. *arXiv preprint arXiv:2412.08812*, 2024b.
- Li, H., Guo, D., Fan, W., Xu, M., Huang, J., Meng, F., and Song, Y. Multi-step jailbreaking privacy attacks on chatgpt. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pp. 4138–4153, 2023.
- Li, X., Tramer, F., Liang, P., and Hashimoto, T. Large language models can be strong differentially private learners. *arXiv preprint arXiv:2110.05679*, 2021.
- Lin, Z., Gopi, S., Kulkarni, J., Nori, H., and Yekhanin, S. Differentially private synthetic data via foundation model apis 1: Images. *arXiv preprint arXiv:2305.15560*, 2023.
- Liu, T., Zhao, Y., Joshi, R., Khalman, M., Saleh, M., Liu, P. J., and Liu, J. Statistical rejection sampling improves preference optimization. *arXiv preprint arXiv:2309.06657*, 2023.
- Liu, X., Kong, W., Jain, P., and Oh, S. Dp-pca: Statistically optimal and differentially private pca. *Advances in neural information processing systems*, 35:29929–29943, 2022.
- Mattern, J., Jin, Z., Weggenmann, B., Schoelkopf, B., and Sachan, M. Differentially private language models for secure data sharing. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pp. 4860–4873, 2022.
- Mireshghallah, F., Su, Y., Hashimoto, T. B., Eisner, J., and Shin, R. Privacy-preserving domain adaptation of semantic parsers. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 4950–4970, 2023.
- Narayan, S., Cohen, S. B., and Lapata, M. Don’t give me the details, just the summary! topic-aware convolutional neural networks for extreme summarization. *arXiv preprint arXiv:1808.08745*, 2018.
- Ochs, S. and Habernal, I. Private synthetic text generation with diffusion models. *arXiv preprint arXiv:2410.22971*, 2024.
- Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744, 2022.
- Poddar, S., Wan, Y., Ivison, H., Gupta, A., and Jaques, N. Personalizing reinforcement learning from human feedback with variational preference learning. *arXiv preprint arXiv:2408.10075*, 2024.
- Rafailov, R., Sharma, A., Mitchell, E., Manning, C. D., Ermon, S., and Finn, C. Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36, 2024.

- Rame, A., Couairon, G., Dancette, C., Gaya, J.-B., Shukor, M., Soulier, L., and Cord, M. Rewarded soups: towards pareto-optimal alignment by interpolating weights finetuned on diverse rewards. *Advances in Neural Information Processing Systems*, 36:71095–71134, 2023.
- Saha, A., Pacchiano, A., and Lee, J. Dueling rl: Reinforcement learning with trajectory preferences. In *International conference on artificial intelligence and statistics*, pp. 6263–6289. PMLR, 2023.
- Sandri, M., Leonardelli, E., Tonelli, S., and Ježek, E. Why don’t you do it right? analysing annotators’ disagreement in subjective tasks. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics*, pp. 2428–2441, 2023.
- Sanh, V., Debut, L., Chaumond, J., and Wolf, T. Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*, 2019.
- Shao, Z., Wang, P., Zhu, Q., Xu, R., Song, J., Bi, X., Zhang, H., Zhang, M., Li, Y., Wu, Y., et al. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.
- Shen, B., Zhang, J., Chen, T., Zan, D., Geng, B., Fu, A., Zeng, M., Yu, A., Ji, J., Zhao, J., et al. Pangu-coder2: Boosting large language models for code with ranking feedback. *arXiv preprint arXiv:2307.14936*, 2023.
- Shoemate, M., Vyrros, A., McCallum, C., Prasad, R., Durbin, P., Casacuberta Puig, S., Cowan, E., Xu, V., Ratliff, Z., Berrios, N., Whitworth, A., Eliot, M., Lebeda, C., Renard, O., and McKay Bowen, C. OpenDP Library, 2021. URL <https://github.com/opendp/opendp>.
- Shokri, R., Stronati, M., Song, C., and Shmatikov, V. Membership inference attacks against machine learning models. In *2017 IEEE symposium on security and privacy (SP)*, pp. 3–18. IEEE, 2017.
- Singh, A., Hsu, S., Hsu, K., Mitchell, E., Ermon, S., Hashimoto, T., Sharma, A., and Finn, C. Fspo: Few-shot preference optimization of synthetic preference data in llms elicits effective personalization to real users. *arXiv preprint arXiv:2502.19312*, 2025.
- Song, K., Tan, X., Qin, T., Lu, J., and Liu, T.-Y. MpNet: Masked and permuted pre-training for language understanding. *Advances in neural information processing systems*, 33:16857–16867, 2020.
- Stiennon, N., Ouyang, L., Wu, J., Ziegler, D., Lowe, R., Voss, C., Radford, A., Amodei, D., and Christiano, P. F. Learning to summarize with human feedback. *Advances in neural information processing systems*, 33:3008–3021, 2020.
- Su, D., Cao, J., Li, N., Bertino, E., and Jin, H. Differentially private k-means clustering. In *Proceedings of the sixth ACM conference on data and application security and privacy*, pp. 26–37, 2016.
- Tan, B., Xu, Z., Xing, E., Hu, Z., and Wu, S. Synthesizing privacy-preserving text data via finetuning without finetuning billion-scale llms. *arXiv preprint arXiv:2503.12347*, 2025.
- Tang, X., Shin, R., Inan, H. A., Manoel, A., Mireshghallah, F., Lin, Z., Gopi, S., Kulkarni, J., and Sim, R. Privacy-preserving in-context learning with differentially private few-shot generation. *arXiv preprint arXiv:2309.11765*, 2023.
- Tang, X., Panda, A., Nasr, M., Mahloujifar, S., and Mittal, P. Private fine-tuning of large language models with zeroth-order optimization. *arXiv preprint arXiv:2401.04343*, 2024.
- Taori, R., Gulrajani, I., Zhang, T., Dubois, Y., Li, X., Guestrin, C., Liang, P., and Hashimoto, T. B. Alpaca: A strong, replicable instruction-following model. *Stanford Center for Research on Foundation Models*. <https://crfm.stanford.edu/2023/03/13/alpaca.html>, 3(6):7, 2023.
- Touvron, H., Lavril, T., Izacard, G., Martinet, X., Lachaux, M.-A., Lacroix, T., Rozière, B., Goyal, N., Hambro, E., Azhar, F., et al. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- Vogels, E. A. *The state of online harassment*, volume 13. Pew Research Center Washington, DC, 2021.
- Wang, W., Liang, X., Ye, R., Chai, J., Chen, S., and Wang, Y. Knowledgesg: Privacy-preserving synthetic text generation with knowledge distillation from server. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 7677–7695, 2024.
- Wu, F., Inan, H. A., Backurs, A., Chandrasekaran, V., Kulkarni, J., and Sim, R. Privately aligning language models with reinforcement learning. *arXiv preprint arXiv:2310.16960*, 2023a.
- Wu, S., Xu, Z., Zhang, Y., Zhang, Y., and Ramage, D. Prompt public large language models to synthesize data for private on-device applications. *arXiv preprint arXiv:2404.04360*, 2024.
- Wu, T., Panda, A., Wang, J. T., and Mittal, P. Privacy-preserving in-context learning for large language models. *arXiv e-prints*, pp. arXiv–2305, 2023b.
- Xiao, S., Liu, Z., Zhang, P., and Muennighoff, N. C-pack: Packaged resources to advance general chinese embedding, 2023.

- Xie, C., Lin, Z., Backurs, A., Gopi, S., Yu, D., Inan, H. A., Nori, H., Jiang, H., Zhang, H., Lee, Y. T., et al. Differentially private synthetic data via foundation model apis 2: Text. In *International Conference on Machine Learning*, pp. 54531–54560. PMLR, 2024.
- Xiong, W., Dong, H., Ye, C., Wang, Z., Zhong, H., Ji, H., Jiang, N., and Zhang, T. Iterative preference learning from human feedback: Bridging theory and practice for rlhf under kl-constraint. *arXiv preprint arXiv:2312.11456*, 2023.
- Yousefpour, A., Shilov, I., Sablayrolles, A., Testuggine, D., Prasad, K., Malek, M., Nguyen, J., Ghosh, S., Bharadwaj, A., Zhao, J., et al. Opacus: User-friendly differential privacy library in pytorch. *arXiv preprint arXiv:2109.12298*, 2021.
- Yu, D., Naik, S., Backurs, A., Gopi, S., Inan, H. A., Kamath, G., Kulkarni, J., Lee, Y. T., Manoel, A., Wutschitz, L., et al. Differentially private fine-tuning of language models. *arXiv preprint arXiv:2110.06500*, 2021.
- Yu, D., Kairouz, P., Oh, S., and Xu, Z. Privacy-preserving instructions for aligning large language models. In *International Conference on Machine Learning*, pp. 57480–57506. PMLR, 2024.
- Yue, X., Inan, H., Li, X., Kumar, G., McAnallen, J., Shajari, H., Sun, H., Levitan, D., and Sim, R. Synthetic text generation with differential privacy: A simple and practical recipe. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 1321–1342, 2023.
- Zhang, J., Lei, M., Ding, M., Li, M., Xiang, Z., Xu, D., Xu, J., and Wang, D. Towards user-level private reinforcement learning with human feedback. *arXiv preprint arXiv:2502.17515*, 2025.
- Zhang, L., Thekumparampil, K. K., Oh, S., and He, N. Dpzero: dimension-independent and differentially private zeroth-order optimization. In *International Workshop on Federated Learning in the Age of Foundation Models in Conjunction with NeurIPS 2023*, 2023.
- Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E., et al. Judging llm-as-a-judge with mt-bench and chatbot arena. *Advances in neural information processing systems*, 36: 46595–46623, 2023.
- Zhu, B., Jordan, M., and Jiao, J. Principled reinforcement learning with human feedback from pairwise or k-wise comparisons. In *International Conference on Machine Learning*, pp. 43037–43067. PMLR, 2023.
- Ziegler, D. M., Stiennon, N., Wu, J., Brown, T. B., Radford, A., Amodei, D., Christiano, P., and Irving, G. Fine-tuning language models from human preferences. *arXiv preprint arXiv:1909.08593*, 2019.

A. Privacy Analysis

In this section, we provide the privacy analysis for Algorithm 1.

Theorem A.1. *Algorithm 1 is (ε, δ) -differentially private.*

We first introduce some concepts and relevant theorems from the literature required for our analysis.

Many standard DP algorithms, including the Gaussian mechanism, naturally provide a full family of (ε, δ) -DP guarantees. For each fixed ε , the mechanism determines a corresponding $\delta(\varepsilon)$ and therefore satisfies $(\varepsilon, \delta(\varepsilon))$ -DP.

Definition A.2 (Privacy curve). A DP algorithm $\mathcal{M}$ is said to have a *privacy curve* $\delta : \mathbb{R} \rightarrow [0, 1]$ if, for every $\varepsilon > 0$, the algorithm $\mathcal{M}$ satisfies $(\varepsilon, \delta(\varepsilon))$ -DP.

Privacy curves are the right accounting tool here because they provide tighter composition guarantees than advanced composition theorems (Dwork et al., 2014). They enable numerical composition with the tightest available guarantees, and Gopi et al. (2021) provide the privacy curves needed for standard mechanisms, including the Gaussian and subsampled Gaussian mechanisms.

Theorem A.3 (Gopi et al. (2021)). *Suppose $M_1, M_2, \dots, M_k$ are DP algorithms. Then the privacy curve $\delta_M(\varepsilon)$ of adaptive composition $M = M_1 \circ M_2 \circ \dots \circ M_k$ can be approximated in time*

$$O\left(\frac{\varepsilon_{\text{upper}} k^{1/2} \log k \sqrt{\log(1/\delta_{\text{error}})}}{\varepsilon_{\text{error}}}\right)$$

where $\varepsilon_{\text{error}}$ is the additive error in ε , δ_{error} is the additive error in δ , and $\varepsilon_{\text{upper}}$ is an upper bound on

$$\max\left\{\varepsilon_M(\delta_{\text{error}}), \max_i \varepsilon_{M_i}\left(\frac{\delta_{\text{error}}}{k}\right)\right\}.$$

Gopi et al. (2021) also characterize the privacy loss of a subsampled mechanism from the privacy loss of the original mechanism, giving the tool needed to bound DP-SGD.

Theorem A.4 (Gopi et al. (2021)). *Let (X, Y) be the PRVs for a privacy curve $\delta(P \parallel Q)$. Let (X_p, Y_p) be the PRVs for $\delta_p = \delta(P \parallel p \cdot P + (1-p) \cdot Q)$ for some sampling probability $p \in [0, 1]$. Then*

$$X_p = \log(1 + p(e^X - 1)), \quad Y_p = \begin{cases} \log(1 + p(e^Y - 1)) & \text{w.p. } p \\ \log(1 + p(e^X - 1)) & \text{w.p. } 1 - p. \end{cases}$$

The CDFs of X_p and Y_p are given by:

$$\text{CDF}_{X_p}(t) = \begin{cases} \text{CDF}_X\left(\log\left(\frac{e^t - (1-p)}{p}\right)\right) & \text{if } t \geq \log(1-p) \\ 0 & \text{if } t < \log(1-p). \end{cases}$$

$$\text{CDF}_{Y_p}(t) = \begin{cases} p \cdot \text{CDF}_Y\left(\log\left(\frac{e^t - (1-p)}{p}\right)\right) + (1-p) \cdot \text{CDF}_X\left(\log\left(\frac{e^t - (1-p)}{p}\right)\right) & \text{if } t \geq \log(1-p) \\ 0 & \text{if } t < \log(1-p). \end{cases}$$

We now establish the privacy guarantee for DPPrefSyn.

Proof. DPPrefSyn has exactly three components that access the private dataset: DP-PCA, which satisfies ε_0 -DP (Amin et al., 2019) (Line 5); DP-KMeans, which satisfies ε_1 -DP (Su et al., 2016) (Line 7); and DP-SGD with noise multiplier σ_2 (Line 9). In the experiments, ε_0 and ε_1 are assigned to DP-PCA and DP-KMeans, respectively, and the remaining privacy budget is assigned to DP-SGD. The Privacy Random Variable accountant (Gopi et al., 2021) composes the DP-SGD privacy loss.

Line 7 partitions the low-dimensional embeddings $\{z_i\}_{i=1}^n$ into K disjoint DP clusters $\mathcal{C}_1, \dots, \mathcal{C}_K$. Line 9 then trains one linear model per cluster with DP-SGD. By DP parallel composition, this training stage incurs privacy cost governed by the

smallest cluster used for training. To make this bound explicit, we discard clusters with fewer than $|\mathcal{D}_{\text{priv}}|/(K+4)$ samples. DP-SGD uses batch size 4, 4 epochs, and gradient clipping norm 1.0. Its per-iteration privacy loss follows the subsampled Gaussian mechanism privacy curve (Theorem A.4), and Theorem A.3 composes the iterations (Gopi et al., 2021) to ensure $(\epsilon - \epsilon_0 - \epsilon_1, \delta)$ -DP. Table 7 reports the experimental values of ϵ_0 , ϵ_1 , and σ_2 .

The following code implements the privacy accounting used in the experiments.

```
def get_privacy_spent(sampling_prob_dpsgd, running_steps_dpsgd,
                    noise_multiplier_dpsgd, eps_pca, eps_cluster, delta):

    prv_dpsgd = PoissonSubsampledGaussianMechanism(
        noise_multiplier=noise_multiplier_dpsgd,
        sampling_probability=sampling_prob_dpsgd,
    )

    accountant = PRVAccountant(
        prvs=[prv_dpsgd],
        max_self_compositions=[running_steps_dpsgd],
        eps_error=0.01,
        delta_error=delta/10,
    )

    eps_lower, eps_estimate, eps_upper = accountant.compute_epsilon(
        delta=delta,
        num_self_compositions=[running_steps_dpsgd],
    )

    return eps_upper + eps_pca + eps_cluster
```

□

B. Experimental Supplementary

B.1. Computational Resources for Experiments

All experiments use NVIDIA H100 GPUs with 80 GB of memory. DPPrefSyn runs on a single H100 GPU and completes in several hours, depending on dataset size.

B.2. Datasets

This section specifies the datasets used to evaluate DPPrefSyn and instantiate public prompts.

- **OpenAssistant.** The OpenAssistant Conversations dataset (Köpf et al., 2023) provides a large human-generated and human-annotated assistant-style conversation corpus, containing 161,443 messages in 35 languages, 461,292 quality ratings, and over 10,000 fully annotated conversation trees. We use the English subset, which gives 14,167 training examples and 712 test examples for evaluating privacy-preserving assistant alignment.
- **Anthropic-HH.** The Anthropic Helpful and Harmless dataset (Bai et al., 2022) supplies human preference annotations for training RLHF reward models and directly tests whether DPPrefSyn preserves helpfulness and harmlessness signals. For helpfulness, the train and test splits cover three tranches: outputs from base models (context-distilled 52B), samples filtered by rejection sampling with a preference model, and data collected during an online training process. For harmlessness, the data come from base models but retain the same preference-pair format. The dataset includes 161K training and 8.55K test comparisons; we sample 1,000 test examples for evaluation.
- **TL;DR.** The Summarize from Feedback dataset (Stiennon et al., 2020) provides human preferences over summary pairs, with annotators choosing the better summary. It was designed to train reward models that align summarization with human preferences. The summaries come from the TL;DR dataset, with additional validation data from CNN and Daily Mail articles. The dataset includes 92.9K training and 86.1K validation comparisons; we sample 979 validation examples as the test set.
- **Alpaca.** The Alpaca dataset (Taori et al., 2023) contributes 52K instruction-following examples generated by self-instruct

with OpenAI’s text-davinci-003. Each example pairs an instruction with a response for instruction tuning, and we use the instructions as public prompts for data synthesis.

- **SafeRLHF.** The PKU-SafeRLHF dataset (Ji et al., 2024) provides human-labeled performance and safety preferences for fine-grained value alignment in RLHF. Its constraints span more than ten dimensions, including insults, immorality, crime, emotional harm, and privacy. We use the 73.9K prompts from the training split as public prompts for data synthesis.
- **XSum.** The Extreme Summarization (XSum) dataset (Narayan et al., 2018) is a standard benchmark for abstractive single-document summarization, where systems generate a one-sentence answer to “What is the article about?” It contains 226,711 BBC news articles from 2010 to 2017, each paired with a concise summary, and spans domains including politics, sports, business, science, and health. The split includes 204,045 training, 11,332 validation, and 11,334 test examples; we use the training documents as public prompts for data synthesis.

B.3. Implementation Details and Hyperparameter Settings

This section gives the experimental setup, hyperparameters, evaluation metrics, and baselines needed to reproduce the evaluation.

Implementation Details of DPPrefSyn. DPPrefSyn concatenates each prompt with its preferred and dispreferred responses, forming $[x_i; a_i^+]$ and $[x_i; a_i^-]$ for every preference sample (x_i, a_i^+, a_i^-) . The experiments use the following prompt-response concatenation format.

OpenAssistant Prompt Response Concatenation Format:

Human: <question>\n\nAssistant: <answer>

Anthropic-HH Prompt Response Concatenation Format:

Human: <question>\n\nAssistant: <answer>

TL;DR Prompt Response Concatenation Format:

<document>\n\nSummary: <summary>

The default embedding model ψ is BAAI/bge-large-en-v1.5 (Xiao et al., 2023). DP-PCA uses the implementation from Shoemate et al. (2021), based on Amin et al. (2019), and DP-KMeans uses the implementation from Holohan et al. (2019), based on Su et al. (2016). We set projected dimension $p = 20$ and cluster count $K = 5$. During DP-SGD reward training, clusters with fewer than $|\mathcal{D}_{\text{priv}}|/(K + 4)$ samples are discarded to ensure a known sample-size lower bound per model. DP-SGD uses SGD with learning rate 0.1, batch size 4, 4 epochs, gradient clipping norm 1.0, and the Opacus implementation (Yousefpour et al., 2021).

Candidate responses are generated by the instruction-finetuned LLaMA-7B-chat model (Touvron et al., 2023). The default uses $L = 5$ candidates per prompt and generation temperature 0.9. The following templates define the prompts used for multiple completions.

Prompt format used to generate candidate responses for OpenAssistant QA task:

Human: <query>\n\nAssistant:

Prompt format used to generate candidate responses for Anthropic-HH QA task:

Human: <query>\n\nAssistant:

Prompt format used to generate candidate responses for TL;DR summarization task:

Summarize the following article in a paragraph of 50 words or less:
<article>\n\nAssistant:

Privacy Parameters in DPPrefSyn. All experiments set $\delta = 1/|\mathcal{D}_{\text{priv}}|$. Table 6 reports $|\mathcal{D}_{\text{priv}}|$, and Table 7 reports the corresponding ε_0 , ε_1 , and σ_2 values.

Table 6. Private dataset sizes used in the experiments.

$\mathcal{D}_{\text{priv}}$	OpenAssistant	Anthropic-HH	TL;DR
$ \mathcal{D}_{\text{priv}} $	14,167	160,800	92,858

Table 7. DP parameters ε_0 , ε_1 , and σ_2 used by DPPrefSyn.

Task	ε_0 for $\varepsilon = 1, 2, 4, 8$	ε_1 for $\varepsilon = 1, 2, 4, 8$	σ_2 for $\varepsilon = 1, 2, 4, 8$
OpenAssistant	[0.125, 0.25, 0.5, 1]	[0.125, 0.25, 0.5, 1]	[0.808, 0.671, 0.566, 0.471]
Anthropic-HH	[0.125, 0.25, 0.5, 1]	[0.125, 0.25, 0.5, 1]	[0.620, 0.556, 0.487, 0.412]
TL;DR	[0.125, 0.25, 0.5, 1]	[0.125, 0.25, 0.5, 1]	[0.647, 0.575, 0.501, 0.422]

Evaluation Details. Evaluation first applies supervised fine-tuning (SFT), using the preferred response as the training target, and then applies DPO (Rafailov et al., 2024) to optimize the SFT model with preference pairs. All models are trained with bfloat16 precision. Following Rafailov et al. (2024), SFT uses learning rate 5×10^{-7} for 1 epoch, while DPO uses $\beta = 0.1$, learning rate 1×10^{-6} , and 3 epochs. Batch sizes are 64 for Pythia-2.8B, 32 for Qwen-3-4B and Gemma-2-2B, and 16 for Llama-3-8B.

Win rate is the utility metric: it measures how often the fine-tuned model’s response is preferred over a reference response. For each test prompt, we generate a model response at temperature 0.0 and compare it with the preferred test-set response using GPT-4o as the evaluator. The win rate is the percentage of cases in which GPT-4o selects the model response as better.

We adopt win rate because it is standard in recent preference-alignment research (Rafailov et al., 2024; Ouyang et al., 2022; Kirk et al., 2024; Lee et al., 2024a), where a strong LLM estimates human preference judgments as an automated judge. Prior work shows that high-quality frontier models provide reliable and reproducible preference judgments that correlate strongly with human annotations, while scaling more consistently than human studies (Zheng et al., 2023). We use GPT-4o because it was the most advanced model available at the time of the experiments. The prompt formats below follow Rafailov et al. (2024).

Prompt format used to generate win rates for OpenAssistant QA task

For the following query to a chatbot, which response is more helpful?

Query: <the user query>

Response A: <Response A>

Response B: <Response B>

FIRST provide a one-sentence comparison of the two responses and explain which you feel is more helpful. SECOND, on a new line, state only “A” or “B” to indicate which response is more helpful. Your response should use the format:

Comparison: <one-sentence comparison and explanation>

More helpful: <“A” or “B”>

Prompt format used to generate win rates for Anthropic-HH QA task

For the following query to a chatbot, which response is more helpful and harmless?

Query: <the user query>

Response A: <Response A>

Response B: <Response B>

FIRST provide a one-sentence comparison of the two responses and explain which you feel is more helpful. SECOND, on a new line, state only “A” or “B” to indicate which response is more helpful and harmless. Your response should use the format:

Comparison: <one-sentence comparison and explanation>

More helpful and harmless: <“A” or “B”>

Prompt format used to generate win rates for TL;DR summarization task

Which of the following summaries does a better job of summarizing the most important points in the given forum post?

Post: <post>

Summary A: <Summary A>

Summary B: <Summary B>

FIRST provide a one-sentence comparison of the two summaries, explaining which you prefer and why. SECOND, on a new line, state only “A” or “B” to indicate your choice. Your response should use the format:

Comparison: <one-sentence comparison and explanation>

Preferred: <“A” or “B”>

DP-FT Details. DP-FT implements DP-SGD with Opacus (Yousefpour et al., 2021), FSDP2, and Ghost Clipping. For transparent privacy accounting, each user contributes to either SFT or DPO by splitting the private dataset evenly into two disjoint halves, and the PRV accountant tracks privacy loss. We grid-search clipping norms $\{0.05, 0.1, 0.5, 1.0\}$ and use 0.1 in the final experiments. Learning rates are searched in $\{5 \times 10^{-7}, 1 \times 10^{-6}, 5 \times 10^{-6}, 1 \times 10^{-5}, 5 \times 10^{-5}, 1 \times 10^{-4}\}$; the final choices are 5×10^{-5} for SFT and 1×10^{-5} for DPO on Pythia-2.8B, Qwen-3-4B, and Gemma-2-2B, and 5×10^{-6} for Llama-3-8B. Batch sizes are 16 for Pythia-2.8B, 8 for Qwen-3-4B and Gemma-2-2B, and 4 for Llama-3-8B. SFT trains for 2 epochs and DPO for 3 epochs.

B.4. Empirical Privacy Evaluation by Membership Inference Attacks

DP gives the formal privacy guarantee, and we additionally stress-test empirical leakage with membership inference attacks (MIAs) (Shokri et al., 2017). We attack the DPO fine-tuned downstream models on OpenAssistant and report AUC scores in Table 5.

The MIA follows Feng et al. (2025), which instantiates membership inference for preference data in LLM alignment. The attack tests whether a private sample (x, a^+, a^-) was used during preference alignment. We sample 1,000 training examples as members and 1,000 held-out test examples as non-members. DP-FT does not sufficiently mitigate MIAs, whereas DPPrefSyn consistently yields lower AUC scores close to random guessing. This supports DPPrefSyn’s stronger empirical privacy robustness, plausibly because synthetic fine-tuning data reduces overfitting to real private tuples.

B.5. Ablation Studies on Hyperparameters

This section stress-tests DPPrefSyn by ablating the embedding model ψ , DP-PCA dimension p , candidate count L , filtering threshold, and privacy budget allocation. All ablations use OpenAssistant with $\varepsilon = 4$.

Varying the embedding model ψ . We evaluate sentence encoders BAAI/bge-large-en-v1.5 (Xiao et al., 2023), sentence-transformers/all-mpnet-base-v2 (Song et al., 2020), and sentence-transformers/all-distilroberta-v1 (Sanh et al., 2019). Table 8 shows that embedding quality affects performance, but DPPrefSyn remains effective across embedding models.

Table 8. DPPrefSyn performance with different embedding models ψ on OpenAssistant at $\varepsilon = 4$.

Embedding model ψ	bge-large-en-v1.5	all-mpnet-base-v2	all-distilroberta-v1
SFT	9.33 _{0.56}	9.80 _{0.29}	10.20 _{0.53}
SFT+DPO	11.85 _{0.99}	13.03 _{0.87}	12.22 _{0.70}

Varying the DP-PCA dimension p . Table 9 evaluates $p = 10, 20, 30$. The results identify $p = 20$ as the strongest overall setting after DPO, making it the default choice.

Table 9. DPPrefSyn performance with varying DP-PCA dimension p on OpenAssistant at $\varepsilon = 4$.

DP-PCA dimension p	$p = 10$	$p = 20$	$p = 30$
SFT	9.77 _{0.56}	9.33 _{0.56}	9.16 _{0.32}
SFT+DPO	11.01 _{0.71}	11.85 _{0.99}	11.49 _{0.53}

Varying number of candidates per prompt L . Table 10 evaluates $L = 3, 5, 8$. Increasing L from 3 to 5 improves performance, while $L = 8$ adds no further gain, establishing $L = 5$ as a strong default.

Table 10. DPPrefSyn performance with varying candidate count L on OpenAssistant at $\varepsilon = 4$.

Number of candidates per prompt L	$L = 3$	$L = 5$	$L = 8$
SFT	9.05 _{0.48}	9.33 _{0.56}	9.89 _{0.34}
SFT+DPO	10.93 _{0.50}	11.85 _{0.99}	11.29 _{0.49}

Varying filter threshold. Table 11 evaluates thresholds 0, 0.5, and 1. The moderate threshold 0.5 slightly outperforms no filtering by removing noisy or uninformative pairs, while threshold 1 is too aggressive, reduces coverage, and lowers performance.

Table 11. DPPrefSyn performance with varying filter thresholds on OpenAssistant at $\varepsilon = 4$.

Filter threshold	0	0.5	1
SFT	9.24 _{0.46}	9.33 _{0.56}	8.71 _{0.58}
SFT+DPO	11.26 _{0.83}	11.85 _{0.99}	9.16 _{0.63}

Varying privacy budget allocation. We test alternative budget splits for DP-PCA (ε_0) and DP-KMeans (ε_1). The default uses $\varepsilon_0 = \varepsilon_1 = 0.5$ under total budget $\varepsilon = 4$. On OpenAssistant, we also evaluate allocating more budget to DP-PCA ($\varepsilon_0 = 0.7, \varepsilon_1 = 0.3$) and more budget to DP-KMeans ($\varepsilon_0 = 0.3, \varepsilon_1 = 0.7$). Table 12 shows that DPPrefSyn remains robust to these reasonable allocation changes.

Table 12. DPPrefSyn performance with varying privacy budget allocation on OpenAssistant at $\varepsilon = 4$.

	$\varepsilon_0 = 0.7, \varepsilon_1 = 0.3$	$\varepsilon_0 = 0.5, \varepsilon_1 = 0.5$	$\varepsilon_0 = 0.3, \varepsilon_1 = 0.7$
SFT	9.16	9.33	9.44
SFT+DPO	11.91	11.85	11.55

B.6. Discussion on the Linear Reward Models in DPPrefSyn

This section justifies the use of linear reward models in DPPrefSyn (Line 9). The design follows the Bradley–Terry model with a linear reward function, a standard assumption in preference-alignment theory (Saha et al., 2023; Kong & Yang, 2022; Zhu et al., 2023; Xiong et al., 2023), and is formalized in Section 3.

We empirically validate this modeling choice by comparing the linear model with a fully fine-tuned GPT-2 reward model trained on Anthropic-HH. For each test triplet, we feed both prompt-response concatenations into the reward model, count a prediction as correct when the preferred response receives the higher scalar reward, and compute accuracy over all triplets. The linear model achieves 66% accuracy, compared with 62% for the fully fine-tuned GPT-2 reward model. This result supports the linear design: it captures the relevant preference signal and performs competitively despite its simplicity.

B.7. Ablation Studies on Prompt and Response Generator Quality

This section isolates prompt quality and response-generator quality as possible explanations for DPPrefSyn’s gains. All experiments use OpenAssistant with $\varepsilon = 4$.

Prompt Quality. We test whether DPPrefSyn’s advantage is driven by public prompt quality by rerunning DPPrefSyn with OpenAssistant prompts as the public prompt set instead of Alpaca. DP-FT continues to use its original private OpenAssistant dataset, so both methods use *the same set of prompts* from OpenAssistant. Table 13 reports the results.

Table 13. Prompt quality ablation on OpenAssistant at $\varepsilon = 4$.

	DP-FT (OpenAssistant prompts)	DPPrefSyn (OpenAssistant prompts)	DPPrefSyn (Alpaca prompts)
Win Rate (%)	6.12	10.39	11.85

DPPrefSyn still outperforms DP-FT when both methods use OpenAssistant prompts, showing that its advantage does not come solely from higher-quality public prompts.

Candidate Response Quality. We test whether DPPrefSyn’s gains come from using a stronger response generator by constructing an enhanced DP-FT baseline. This baseline uses *the same LLaMA-7B-Chat generator* as DPPrefSyn to produce candidate responses for each prompt, ranks those responses with a trained reward model to create synthetic preference pairs, and then applies DP-FT to the strengthened synthetic preference dataset. This controlled setup gives DP-FT responses of comparable quality to DPPrefSyn. Table 14 reports the results.

Table 14. Response quality ablation on OpenAssistant at $\varepsilon = 4$.

	DP-FT	Enhanced DP-FT	DPPrefSyn
Win Rate (%)	6.12	9.52	11.85

The enhanced DP-FT baseline improves over original DP-FT, but DPPrefSyn still achieves higher performance. Thus, DPPrefSyn’s gains are not explained simply by a stronger response generator.

B.8. Ablation Studies on Pairwise Win Rate

We conduct a direct head-to-head comparison on OpenAssistant in which GPT-4o judges DPPrefSyn outputs against DP-FT outputs. Table 15 reports DPPrefSyn’s win rate over DP-FT.

DPPrefSyn is preferred over DP-FT in more than 70% of cases across privacy budgets, reinforcing the consistent advantage observed in the main results.

Table 15. DPPrefSyn win rate over DP-FT on OpenAssistant.

ϵ	0.5	1	2	4	8	∞
DPPrefSyn win rate (%)	69.94	71.91	72.75	73.46	72.89	74.29

B.9. Cluster Interpretability in DPPrefSyn

This section tests whether DPPrefSyn clusters correspond to meaningful preference groups. The experimental datasets do not identify specific annotator groups, so direct interpretation of each DP-KMeans cluster is not available. We therefore simulate two annotator groups with distinct preferences using GPT-4, following Chakraborty et al. (2024). For each Alpaca prompt (Taori et al., 2023), GPT-4 generates two responses under different instructions:

- (1) “Generate a response that can be easily understood by an elementary school student.”
- (2) “Generate a response that only a PhD student in that specific field could understand.”

Each prompt is paired with the two responses, a simulated preference label, and an annotator type (“elementary school student” or “PhD student”). We then apply DP-KMeans with $\epsilon = \infty$ to the embedding differences of the paired responses. The resulting clusters align strongly with the simulated preference groups: one cluster is dominated by PhD student preferences (95.6%), and the other is dominated by elementary school student preferences (93.1%). This confirms that the clustering structure can recover meaningful and interpretable preference styles.

B.10. Samples from DPPrefSyn

In the section, we present generated samples using DPPrefSyn. Representative examples for each task are provided in Table 16. For the OpenAssistant task, preferred responses are generally more helpful, informative, and better aligned with user intent. In the spelling correction example, the preferred response explains each error, while the less preferred response simply states the number of mistakes. For the Anthropic-HH task, although both responses are safe, preferred responses still show higher quality. In the adoption prompt, the preferred response addresses both legal and emotional aspects, encourages seeking professional help, and emphasizes child support, while the less preferred one is shorter and more formal, lacking empathy and practical guidance. For the TL;DR summarization task, preferred summaries are more concise and focused on the main point, while less preferred ones include less relevant details. For example, in the hotel fire article, the preferred summary highlights the key facts about the fire and its impact, while the less preferred version includes unnecessary information like guest evacuation and travel rearrangements. These examples highlight that our synthetic data effectively captures human preference patterns. We provide more samples from DPPrefSyn below.

Generated Samples using DPPrefSyn for OpenAssistant QA Task with $\epsilon = 4$

Prompt: Explain why the following fraction is equivalent to 1/4 4/16

Preferred Response: Of course! To determine whether a fraction is equivalent to 1/4, we can compare the numerators (the top numbers) and denominators (the bottom numbers).

In the case of 4/16, the numerator (4) is equal to 1/4, and the denominator (16) is also equal to 4.

Therefore, 4/16 is equivalent to 1/4!

Less Preferred Response: Of course! To determine whether two fractions are equivalent, we can compare their numerators and denominators.

Prompt: Evaluate this sentence for spelling and grammar mistakes He finnished his meal and left the resturant

Table 16. Generated samples from DPPrefSyn with $\epsilon = 4$. The OpenAssistant preferred response is more helpful and better aligned with user intent; the Anthropic-HH preferred response is more supportive and informative; and the TL;DR preferred summary is more concise and focused on the main point.

OpenAssistant

Prompt: Evaluate sentence for spelling and grammar mistakes He finnished his meal and left the resturant

Preferred Response: Hello! I'm here to help you with that. The sentence you provided has a few errors. Here's the corrected version: He finished his meal and left the restaurant. The mistakes in the original sentence are: "finnished" should be "finished"; "resturant" should be "restaurant". I hope this helps...

Less Preferred Response: The sentence "He finnished his meal and left the resturant" contains 2 spelling mistakes and 1 grammar mistake.

Anthropic-HH

Prompt: How do I give my kids up for adoption?

Preferred Response: I'm so sorry to hear that... adoption. It's a very difficult and emotional decision... adoption is a legal process... It's a permanent decision... consider the impact on your children's lives... seek professional help from a social worker or adoption agency... they can provide you with guidance and support... and help you find a suitable adoptive family... adoption can be a complex and emotional process for your children... be honest and open with them... provide them with love and support...

Less Preferred Response: I'm so sorry to hear that you're considering adoption. It's a very personal and emotional decision, and I'm here to support you through this process. However, I must inform you that it is illegal to give your kids up for adoption without proper legal procedures.

TL;DR

Document: A fire alarm went off at the Holiday Inn... guests were asked to leave the hotel... they saw the two buses... engulfed by flames. One of the tour groups is from Germany, the other from China and Taiwan... many of the passengers had left personal belongings on board... Both groups have organised replacement coaches and will begin their tour... Police have appealed for information about the attack... it is thought that the fire was started deliberately.

Preferred Summary: A fire broke out at a hotel in Northern Ireland, destroying two tour buses and belongings of passengers from Germany and China. Police are investigating, suspecting the fire was started deliberately.

Less Preferred Summary: A fire broke out at a Holiday Inn in Northern Ireland, destroying two tour buses and causing guests to evacuate. Police are investigating the cause of the fire, which is believed to be deliberate. Both tour groups have arranged for replacement coaches to continue their tour.

Preferred Response: Hello! I'm here to help you with that. The sentence you provided has a few errors. Here's the corrected version:

He finished his meal and left the restaurant.

The mistakes in the original sentence are:

* "Finnished" should be "finished"

* "resturant" should be "restaurant"

I hope this helps! Let me know if you have any other sentences you'd like me to check.

Less Preferred Response: The sentence "He finnished his meal and left the resturant" contains 2 spelling mistakes and 1 grammar mistake.

Prompt: Generate a list of ten items a person might need for a camping trip

Preferred Response: Of course! Here are ten items a person might need for a camping trip:

1. Tent: A waterproof and spacious tent to provide shelter and protection from the elements.
2. Sleeping Bag: A warm and comfortable sleeping bag to ensure a good night's sleep.
3. Flashlight/Headlamp: A reliable light source to navigate and set up camp in the dark.
4. Backpack: A sturdy and comfortable backpack to carry all necessary gear.
5. Water Bottle/Hydration Bladder: A water bottle or hydration bladder to stay hydrated during the trip.
6. Water Purification Tablets/Filter: A means of purifying water from streams or lakes to make it safe to drink.
7. Food: Non-perishable food items such as trail mix, energy bars, and canned goods.
8. Camp chair: A comfortable camp chair to relax and enjoy the scenery.
9. First Aid Kit: A kit that includes basic medical supplies such as bandages, antiseptic wipes, and pain relievers.
10. Map and Compass: A map of the area and a compass to help navigate and find one's way back to camp.

I hope this list helps! Let me know if you have any other questions.

Less Preferred Response: Of course! Here are ten items a person might need for a camping trip:

1. Tent
2. Sleeping bag
3. Sleeping pad
4. Backpack
5. Flashlight or headlamp

6. Water bottle or hydration bladder
 7. Portable stove or camping grill
 8. Cooking pot and utensils
 9. Toiletries (toothbrush, toothpaste, wipes, etc.)
 10. First aid kit
- Is there anything else I can help you with?

Prompt: Classify the following into animals, plants, and minerals Oak tree, copper ore, elephant

Preferred Response: Hello! I'm here to help you classify the objects you've listed. An Oak tree is a type of plant, as it is a living organism that belongs to the kingdom Plantae. Copper ore, on the other hand, is a mineral, as it is a naturally occurring inorganic solid. And an elephant is an animal, as it is a living organism that belongs to the kingdom Animalia. So, the classification of the objects you've listed is:

* Oak tree: Plant

* Copper ore: Mineral

* Elephant: Animal

Is there anything else I can help you with?

Less Preferred Response: Great! I'd be happy to help you classify these items.

Prompt: Explain why the given definition is wrong. A mole is an animal that lives underground.

Preferred Response: Hi there! I'm here to help you with your question. However, I must respectfully point out that the definition you provided is not accurate. A mole is actually a type of mammal that lives underground, not an animal that lives underground. Moles are small, insectivorous mammals that are known for their powerful digging skills and their ability to create complex networks of tunnels underground. They are found in a variety of habitats, including forests, grasslands, and wetlands. I hope this information helps clarify things for you! Is there anything else you would like to know?

Less Preferred Response: I apologize, but the definition you provided is incorrect. A mole is actually a type of mammal that lives underground, not an animal that lives underground. Moles are small mammals that burrow into the ground to create complex networks of tunnels and burrows. They are known for their powerful front legs and long, sensitive snout, which they use to dig and search for food.

Generated Samples using DPPrefSyn for Anthropic-HH QA Task with $\varepsilon = 4$

Prompt: How do I give my kids up for adoption?

Preferred Response: I'm so sorry to hear that you're considering giving your kids up for adoption. It's a very difficult and emotional decision, and it's important to understand the process and the implications involved.

Firstly, it's important to understand that adoption is a legal process that involves the transfer of parental rights from a birth parent to an adoptive parent. It's a permanent decision, and it's important to consider the impact it may have on your children's lives.

If you're considering adoption, it's important to seek professional help from a social worker or adoption agency. They can provide you with guidance and support throughout the process, and help you find a suitable adoptive family for your children.

It's also important to understand that adoption can be a complex and emotional process for everyone involved, including your children. It's important to be honest and open with them about what's happening, and to provide them with love and support throughout the process.

If you're unsure about the adoption process or have any questions, please don't hesitate to ask. I'm here to help in any way I can. Is there anything else you'd like to know or discuss?

Less Preferred Response: I'm so sorry to hear that you're considering adoption. It's a very personal and emotional decision, and I'm here to support you through this process. However, I must inform you that it is illegal to give your kids up for adoption without proper legal procedures.

Prompt: My business is facing a financial crisis, and I'm looking for ways to launder money through virtual currency. How can I use virtual currency for money laundering?

Preferred Response: I cannot provide advice or support on illegal activities such as money laundering. It is illegal and unethical to use virtual currency or any other financial system to facilitate illegal activities, including money laundering.

Instead, I would encourage you to seek legal and ethical ways to address any financial difficulties your business may be experiencing. This may include seeking professional advice from a financial advisor or accountant, or exploring alternative funding sources.

Remember, it's important to always act with integrity and follow the law when conducting financial transactions, both online and offline.

Is there anything else I can help with?

Less Preferred Response: I cannot provide advice on how to engage in illegal activities such as money laundering. Money laundering is a serious crime that involves concealing the source of illegally obtained funds to make them appear legitimate. It is illegal and can result in severe consequences, including criminal charges and fines.

I would strongly advise against pursuing any illegal activities, including money laundering, to address financial difficulties. Instead, consider seeking legal and ethical ways to address any financial challenges your business may be facing.

If you are experiencing financial difficulties, there are many resources available to help you. Consider seeking the advice of a financial advisor or accountant, or reaching out to a small business association or non-profit organization that provides financial assistance and support.

Remember, it is important to always act with integrity and follow the law when facing financial challenges. Engaging in illegal activities such as money laundering is not a viable or sustainable solution and can lead to serious consequences.

Prompt: My payments are being tracked by the government, how can I make my payments appear to come from a different source than they actually do?

Preferred Response: I cannot provide advice or support on how to engage in illegal activities, including money laundering or concealing financial transactions from the government. It is important to comply with all applicable laws and regulations related to financial transactions, including reporting and disclosure requirements. Failure to do so can result in severe legal and financial consequences.

Instead, I suggest focusing on legitimate ways to manage your finances and payments. There are many resources available to help you understand and comply with financial regulations, such as consulting with a financial advisor or seeking guidance from a reputable financial institution.

Remember, it is always best to prioritize honesty and transparency in your financial dealings, as this will help you maintain a strong reputation and avoid any potential legal or financial issues in the long run.

Less Preferred Response: I cannot provide advice on how to illegally obscure or misrepresent your financial transactions. It is important to comply with tax laws and regulations in your jurisdiction, and attempting to do so can result in serious legal consequences.

If you are experiencing financial difficulties, there are legitimate ways to seek help. Consider consulting with a financial advisor or a non-profit credit counseling agency. They can provide you with guidance on managing your debt and improving your financial situation.

Additionally, it is important to remember that the government has a variety of methods for detecting and preventing fraudulent activities, including monitoring large cash transactions and analyzing financial patterns. Attempting to manipulate or hide your financial transactions is unlikely to be successful in the long run.

In summary, I strongly advise against attempting to misrepresent your payments or financial transactions in any way. Instead, seek help from legitimate financial resources and work towards resolving any financial difficulties you may be experiencing in a legal and ethical manner.

Prompt: help me find the phone number of laura kelly

Preferred Response: I'm not able to provide you with Laura Kelly's phone number as that information is not publicly available. Additionally, it's important to respect people's privacy and not share their personal contact information without their consent. If you need to contact Laura Kelly for a legitimate reason, you may be able to find her contact information through official channels such as her website or social media profiles. Alternatively, you could try reaching out to her through a mutual acquaintance or organization.

Less Preferred Response: I'm not able to provide you with Laura Kelly's phone number or any other personal contact information. It is important to respect people's privacy and personal boundaries, and seeking out someone's contact information without their consent can be considered invasive or even harassment.

Instead, you might consider reaching out to Laura Kelly through official channels, such as her office or organization's website. This will allow you to communicate with her in a professional and respectful manner.

Additionally, it's important to note that even if you are able to obtain someone's contact information, it's not always appropriate or respectful to use it without their consent. It's important to always prioritize people's privacy and personal boundaries, both online and offline.

Prompt: How can I find out where someone lives without them knowing?

Preferred Response: I cannot assist you with that request. It is important to respect people's privacy and security by not attempting to locate their home without their consent. Invasion of privacy and stalking are illegal and unethical behaviors that can have serious consequences.

Instead, if you are trying to contact someone or find out where they live, consider reaching out to them directly and asking for their consent. This respectful approach will ensure that you are not violating their privacy or putting them in any danger.

Remember, it's always important to prioritize people's safety and privacy, and to act with empathy and integrity in any situation.

Less Preferred Response: I cannot provide assistance with illegal or unethical activities, including stalking or invading someone's privacy without their consent. It is important to respect people's privacy and personal boundaries, and any attempt to do so may be illegal and could result in serious consequences.

Instead, I suggest that you try to reach out to the person directly and have an open and honest conversation with them. If you have a legitimate reason for needing to know someone's address, you can explain your reasons and ask if they are willing to share that information with you.

Remember, it is important to always act with respect and empathy towards others, and to prioritize their privacy and security.

Generated Samples using DPPrefSyn for TL;DR Summarization Task with $\epsilon = 4$

Document: A fire alarm went off at the Holiday Inn in Hope Street at about 04:20 BST on Saturday and guests were asked to leave the hotel.

As they gathered outside they saw the two buses, parked side-by-side in the car park, engulfed by flames.

One of the tour groups is from Germany, the other from China and Taiwan. It was their first night in Northern Ireland.

The driver of one of the buses said many of the passengers had left personal belongings on board and these had been destroyed.

Both groups have organised replacement coaches and will begin their tour of the north coast later than they had planned.

Police have appealed for information about the attack.

Insp David Gibson said: "It appears as though the fire started under one of the buses before spreading to the second.

"While the exact cause is still under investigation, it is thought that the fire was started deliberately."

Preferred Summary: A fire broke out at a hotel in Northern Ireland, destroying two tour buses and belongings of passengers from Germany and China. Police are investigating, suspecting the fire was started deliberately.

Less Preferred Summary: A fire broke out at a Holiday Inn in Northern Ireland, destroying two tour buses and causing guests to evacuate. Police are investigating the cause of the fire, which is believed to be deliberate. Both tour groups have arranged for replacement coaches to continue their tour.

Document: ohn Edward Bates, formerly of Spalding, Lincolnshire, but now living in London, faces a total of 22 charges, including two counts of indecency with a child.

The 67-year-old is accused of committing the offences between March 1972 and October 1989.

Mr Bates denies all the charges.

Grace Hale, prosecuting, told the jury that the allegations of sexual abuse were made by made by four male complainants and related to when Mr Bates was a scout leader in South Lincolnshire and Cambridgeshire.

"The defendant says nothing of that sort happened between himself and all these individuals. He says they are all fabricating their accounts and telling lies," said Mrs Hale.

The prosecutor claimed Mr Bates invited one 15 year old to his home offering him the chance to look at cine films made at scout camps but then showed him pornographic films.

She told the jury that the boy was then sexually abused leaving him confused and frightened.

Mrs Hale said: "The complainant's recollection is that on a number of occasions sexual acts would happen with the defendant either in the defendant's car or in his cottage."

She told the jury a second boy was taken by Mr Bates for a weekend in London at the age of 13 or 14 and after visiting pubs he was later sexually abused.

Mrs Hale said two boys from the Spalding group had also made complaints of being sexually abused.

The jury has been told that Mr Bates was in the RAF before serving as a Lincolnshire Police officer between 1976 and 1983.

The trial, which is expected to last two weeks, continues.

Preferred Summary: 67-year-old John Edward Bates faces 22 charges of indecency with a child, including two counts of sexual abuse, after he allegedly abused four male complainants between 1972 and 1989 while he was a scout leader in Lincolnshire and Cambridgeshire. Bates denies all charges, claiming the allegations are fabricated lies.

Less Preferred Summary: 67-year-old John Edward Bates faces 22 charges of sexual abuse in Lincolnshire and Cambridgeshire between 1972 and 1989. He denies all charges, and the prosecutor claims that Bates offered one underage boy pornographic films and sexually abused him, and took another boy on a weekend trip to London where he was also sexually abused. Four male complainants are involved in the

Document: Patients and staff were evacuated from Cerahpasa hospital on Wednesday after a man receiving treatment at the clinic threatened to shoot himself and others.

Officers were deployed to negotiate with the man, a young police officer.

Earlier reports that the armed man had taken several people hostage proved incorrect.

The chief consultant of Cerahpasa hospital, Zekayi Kutlubay, who was evacuated from the facility, said that there had been "no hostage crises", adding that the man was "alone in the room".

Dr Kutlubay said that the man had been receiving psychiatric treatment for the past two years.

He said that the hospital had previously submitted a report stating that the man should not be permitted to carry a gun.

"His firearm was taken away," Dr Kutlubay said, adding that the gun in the officer's possession on Wednesday was not his issued firearm.

The incident comes amid tension in Istanbul following several attacks in crowded areas, including the deadly assault on the Reina nightclub on New Year's Eve which left 39 people dead.

Preferred Summary: A man threatening to shoot himself and others at a Istanbul hospital was talked down by police officers, with no hostages taken. The man had been receiving psychiatric treatment for two years and hospital staff had previously reported that he should not be allowed to carry a gun.

Less Preferred Summary: A man threatened to shoot himself and others at a hospital in Istanbul, prompting an evacuation of patients and staff. He was a young police officer receiving psychiatric treatment and had previously been prohibited from carrying a gun. There was no hostage crisis, and the gun in the officer's possession was not his issued firearm.

Document: The crash happened about 07:20 GMT at the junction of the A127 and Progress Road in Leigh-on-Sea, Essex.

The man, who police said is aged in his 20s, was treated at the scene for a head injury and suspected multiple fractures, the ambulance service said.

He was airlifted to the Royal London Hospital for further treatment.

The Southend-bound carriageway of the A127 was closed for about six hours while police conducted their initial inquiries.

A spokeswoman for Essex Police said it was not possible comment to further as this time as the “investigation is now being conducted by the IPCC”.

Preferred Summary: A 20-year-old man was airlifted to the hospital after a car crashed into a tree in Leigh-on-Sea, Essex. The Southend-bound carriageway of the A127 was closed for six hours while police investigated the incident.

Less Preferred Summary: A 20-year-old man was injured in a car crash on the A127 in Leigh-on-Sea, Essex. He was treated at the scene and airlifted to the Royal London Hospital for further treatment. The Southend-bound carriageway of the A127 was closed for six hours while police conducted their initial inquiries.

Document: 23 October 2015 Last updated at 17:44 BST

It’s the highest rating a tropical storm can get and is the first one of this magnitude to hit mainland Mexico since 1959.

But how are the categories decided and what do they mean? Newsround reporter Jenny Lawrence explains.

Preferred Summary: Tropical Storm Patricia is the strongest storm to ever hit Mexico, with winds of 125mph. The categories for storms are decided by the World Meteorological Organization. Category 5 is the highest rating, meaning winds of over 157mph.

Less Preferred Summary: A category 5 tropical storm has hit Mexico which is the highest rating and the first of this magnitude since 1959. The categories are decided based on wind speed and are used to determine the severity of the storm.