

Differentially Private Preference Data Synthesis for Large Language Model Alignment

Anonymous Authors¹

Abstract

Preference alignment is an essential post-training step for large language models (LLMs), directly determining whether their outputs remain faithful to human values. Yet post-training on real human preference data creates a serious privacy exposure, because these datasets routinely contain sensitive user prompts and revealing human judgments. We resolve this tension with **DPPrefSyn**, a novel algorithm that generates differentially private (DP) synthetic preference data and thereby enables privacy-preserving preference alignment. DPPrefSyn is a principled framework grounded in the Bradley–Terry preference model and the intrinsic geometric structure of pairwise human preference data. It first learns an underlying preference model from private data with formal differential privacy guarantees, and then couples the learned model with public prompts to synthesize high-quality preference data. By exploiting the shared linear structure of per-cluster reward models, it faithfully captures the heterogeneous human preferences present in private datasets, and it applies DP Principal Component Analysis (DPPCA) to sharpen learning accuracy. Extensive experiments establish that DPPrefSyn delivers strong alignment performance under demanding DP guarantees, in several settings even surpassing non-private fine-tuning on real data. These results demonstrate that synthetic preference data is a practical and privacy-preserving substitute for real preference data across a broad range of alignment applications. To the best of our knowledge, this is the first work to generate DP synthetic preference data for LLM alignment. Our code will be released upon acceptance.

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

1. Introduction

Preference alignment techniques, such as RLHF (Stiennon et al., 2020) and DPO (Rafailov et al., 2024), have become the standard means of aligning LLMs with human expectations. These algorithms rely on pairwise preference data, where human annotators compare two responses and select the one that better addresses the given prompt. Such preference data then drives the fine-tuning of LLMs by pushing the model to rank preferred responses over less preferred ones (Ziegler et al., 2019; Ouyang et al., 2022; Bai et al., 2022). Because preference alignment is both effective and reliable, it now underpins flagship applications such as chat assistants (Achiam et al., 2023), mathematical reasoning tools (Shao et al., 2024), and code generators (Shen et al., 2023).

However, aligning LLMs with human preference data creates substantial privacy risks, because these datasets frequently include real user prompts and human feedback. These prompts can disclose personal information about health, identity, or other sensitive topics, and human feedback can expose private beliefs, preferences, or behavioral patterns (Li et al., 2023; Yu et al., 2024). Prior works (Chowdhury et al., 2024; Yu et al., 2024; Wu et al., 2023a) have mitigated these risks using the rigorous safeguards of differential privacy (DP) (Dwork et al., 2006b) with encouraging results, but they leave three concrete gaps unaddressed. First, most protect only part of the data, privatizing either the user prompts (Yu et al., 2024) or the preference labels (Chowdhury et al., 2024; Zhang et al., 2025), but not both; such partial protection is insufficient, since adversaries can still infer sensitive attributes from the unprotected prompts or preference annotations. Second, they are tied to specific post-training algorithms, such as RLHF (Wu et al., 2023a), and do not extend to newer methods like DPO. Third, they are constrained by limited private preference datasets and therefore fall short of high-quality alignment, because collecting human preference annotations is costly and hard to scale. These gaps lead to the central question of this work:

Is it possible to achieve superior preference alignment while ensuring thorough privacy protection and maintaining compatibility with emerging alignment methodologies?

In this work, we answer this question decisively in the affirmative. Our main contributions are as follows:

- We introduce DPPrefSyn (Figure 1), a new algorithm that generates DP synthetic preference data for LLM alignment. DPPrefSyn is a principled framework grounded in the Bradley–Terry preference model and the intrinsic geometric structure of pairwise human preference data. Three design elements drive its effectiveness: 1) *Capturing diverse human preferences via the shared linear structure of rewards and clustering*. The Bradley–Terry preference model, together with a linear reward structure, lets us group preference samples by the embedding differences between concatenated prompt–preferred and prompt–dispreferred pairs, and approximate the preferences within each cluster using a shared reward function. 2) *Improving learning efficiency through DP-PCA*. To eliminate the sample inefficiency and clustering instability of high-dimensional spaces, we apply DP-PCA to project the embeddings into a lower-dimensional space while preserving the key preference signals under DP. 3) *Saving privacy budget by using public prompts*. We generate candidate responses from public prompts and apply private reward models to select preference pairs, so that the entire privacy budget is spent on modeling private preferences, maximizing privacy efficiency.
- We provide a rigorous privacy analysis of DPPrefSyn and prove that it satisfies (ϵ, δ) -DP. In our analysis, we allocate a privacy budget ϵ_0 to DP-PCA (Amin et al., 2019) for dimensionality reduction and ϵ_1 to DP-KMeans (Su et al., 2016) for clustering, and spend the remaining budget training per-cluster reward models with DP-SGD. We apply the Privacy Random Variable (PRV) accountant (Gopi et al., 2021) to compose the DP-SGD privacy cost tightly, guaranteeing that DPPrefSyn satisfies (ϵ, δ) -DP. By the post-processing property of DP, the resulting DP synthetic dataset is reusable across arbitrary preference alignment methods and LLMs at no additional privacy cost.
- We evaluate DPPrefSyn on standard benchmarks, including question answering tasks from OpenAssistant (Köpf et al., 2023) and Anthropic-HH (Bai et al., 2022), as well as the TL;DR summarization task (Stiennon et al., 2020). Across all of these diverse tasks, the synthetic preference data generated by DPPrefSyn delivers strong privacy guarantees together with high downstream utility. For example, with $\epsilon = 2$, fine-tuning the Pythia-2.8B model with DPO on our synthetic data reaches a GPT-4o win rate of 56.48% on Anthropic-HH, far exceeding fine-tuning on real data (37.02%). These results establish that our DP synthetic preference data enables effective LLM alignment under strong privacy guarantees. Moreover, DPPrefSyn applies to a range of preference optimiza-

tion methods (e.g., DPO and RLHF) and stays effective under tight privacy constraints, directly overcoming the method-specific limitations of prior privacy-preserving post-training approaches.

2. Related Work

Differentially Private Alignment of LLMs. Wu et al. (2023a) introduce a DP framework to align LLMs with reinforcement learning by adapting the PPO algorithm to the DP setting. Their approach, however, is *limited to PPO*, whereas preference alignment continues to evolve toward new methods. Chowdhury et al. (2024) study reward estimation from preference-based feedback, using label-DP (Ghazi et al., 2021) to protect the privacy of human annotators, and Zhang et al. (2025) propose AUP-RLHF, a user-level label-DP framework for RLHF. Although both methods safeguard preference labels effectively, they *leave the sensitivity of the prompts and responses themselves unprotected*. They are also *constrained by limited private user preference data*, and therefore cannot reach high-quality alignment, since human preference annotations are costly to obtain. Yu et al. (2024) instead protect sensitive user *instructions* by generating DP synthetic *instructions* to replace real ones during data annotation and model fine-tuning. In contrast to all of these, DPPrefSyn protects prompts, responses, and labels jointly, remains method-agnostic, and overcomes the data-scarcity bottleneck.

Beyond preference alignment, several studies focus on private fine-tuning of LLMs using general labeled data, rather than preference comparisons (Yu et al., 2021; Li et al., 2021; Chen et al., 2024; Tang et al., 2024; Zhang et al., 2023; He et al., 2022), and some recent work focuses on protecting the privacy of in-context prompts for in-context learning (Duan et al., 2023; Tang et al., 2023; Gao et al., 2024; Hong et al., 2023; Wu et al., 2023b).

Differentially Private Synthetic Text Generation. Our work sits within the broader scope of DP synthetic text generation, where two representative lines of work have emerged. The first line fine-tunes a language model on private data under DP and then uses the fine-tuned model to generate synthetic text (Yue et al., 2023; Mattern et al., 2022; Mireshghallah et al., 2023; Carranza et al., 2023; Yu et al., 2024; Wang et al., 2024; Ochs & Habernal, 2024; Tan et al., 2025; Carranza et al., 2024). For example, Yu et al. (2024) privately fine-tune the LLaMA model on private instruction data and then use the resulting model to generate synthetic text for supervised instruction tuning.

Because DP fine-tuning is expensive and often infeasible (especially for non-public language models), a second, more recent line of work generates DP synthetic data by prompting LLM APIs (Lin et al., 2023; Xie et al., 2024; Wu et al., 2024;




Figure 1. Overview of DPPrefSyn. DPPrefSyn generates DP synthetic preference data in 3 steps: 1) it represents preference samples as embedding differences and clusters them via DP-PCA and DP-KMeans, 2) it trains a DP reward model on each cluster using DP-SGD, and 3) it generates synthetic preference samples from public prompts, guided by the DP-protected distribution of reward functions.

Hou et al., 2024). These approaches generate synthetic samples through LLM API access and then select outputs that are similar to the private data in a privacy-preserving manner. For instance, Xie et al. (2024) propose the Aug-PE algorithm, which iteratively refines LLM-generated samples based on embedding-space similarity to private data. *Our work follows this line of research and is the first to bring it to bear on preference alignment for LLMs.*

Diversity in Human Preferences. Human preferences are inherently diverse (Denton et al., 2021; Aroyo et al., 2023a;b; Chakraborty et al., 2024). This diversity stems from sociodemographic differences (e.g., race, gender, age), personal beliefs and biases, varying levels of domain expertise, and the inherent ambiguity of natural language (Sandri et al., 2023; Vogels, 2021). Accordingly, recent studies personalize preference alignment to better reflect the values of different user groups (Lee et al., 2024b; Poddar et al., 2024; Singh et al., 2025). For example, Singh et al. (2025) propose a meta-learning framework in which an LLM adapts to individual users from a small number of preference examples. In this work, we generate DP synthetic preference data that *explicitly captures the underlying diversity of human feedback*, and we show that doing so yields markedly better alignment than approaches that ignore this diversity (Section 5.3).

3. Problem Definition

Preference Alignment of LLMs. We consider preference alignment of an LLM based on pairwise human feedback, using a preference dataset $\mathcal{D} = \{(x_i, a_i^+, a_i^-)\}_{i=1}^n$. Each sample consists of a prompt x_i and two responses a_i^+, a_i^- , where a_i^+ is preferred over a_i^- by human annotators.

Following prior works on reinforcement learning from human feedback (RLHF) and direct preference optimization

(DPO) (Ouyang et al., 2022; Zhu et al., 2023; Rafailov et al., 2024; Liu et al., 2023), we model human preferences using the Bradley-Terry model (Bradley & Terry, 1952):

$$\begin{aligned} \mathbb{P}[a_i^+ \succ a_i^- \mid x_i] &= \frac{\exp(r^*(x_i, a_i^+))}{\exp(r^*(x_i, a_i^+)) + \exp(r^*(x_i, a_i^-))} \\ &= \sigma(r^*(x_i, a_i^+) - r^*(x_i, a_i^-)), \end{aligned}$$

where $a_i^+ \succ a_i^-$ means a_i^+ is preferred to a_i^- , $r^*(x, a)$ is a latent reward function, and $\sigma(z) = 1/(1 + e^{-z})$ is the sigmoid function.

This formulation underlies both RLHF, which explicitly learns a reward model from preference data, and DPO, which induces an implicit reward through the log-ratio between the learned policy and a reference policy. It is therefore the natural foundation on which to build a method-agnostic synthesis procedure.

Differentially Private Preference Data Synthesis. We protect the privacy of the preference dataset $\mathcal{D}_{\text{priv}} = \{(x_i, a_i^+, a_i^-)\}_{i=1}^n$ against an adversary who attempts to access or infer private information about individual prompts, responses, or preference labels from the outputs of the fine-tuned LLM. To this end, we generate a DP synthetic preference dataset $\tilde{\mathcal{D}} = \{(\tilde{x}_j, \tilde{a}_j^+, \tilde{a}_j^-)\}_{j=1}^m$ that preserves the critical preference information in $\mathcal{D}_{\text{priv}}$. By the post-processing property of DP, $\tilde{\mathcal{D}}$ can be reused across arbitrary preference alignment methods and LLMs without any additional privacy cost.

We formally define (ϵ, δ) -differential privacy as follows.

Definition 3.1 ((ϵ, δ) -Differential Privacy (Dwork et al., 2006a)). A randomized algorithm $\mathcal{A}$ is (ϵ, δ) -differentially private if for any two neighboring inputs $\mathcal{D}$ and $\mathcal{D}'$ that differ by a single entry and any set $\mathcal{S}$ of possible outputs: $\mathbb{P}[\mathcal{A}(\mathcal{D}) \in \mathcal{S}] \leq e^\epsilon \mathbb{P}[\mathcal{A}(\mathcal{D}') \in \mathcal{S}] + \delta$.

Diverse User Preferences and Linear Reward Structure.

Human preferences in the private dataset are rarely homogeneous (Chakraborty et al., 2024; Bakker et al., 2022; Kovač et al., 2023; Jang et al., 2023; Rame et al., 2023; Ji et al., 2023). Different annotators prioritize different aspects of a response, such as factual accuracy, politeness, creativity, or clarity, producing diverse and often conflicting signals in the data. Capturing this diversity in $\mathcal{D}_{\text{priv}}$ is therefore essential to synthesizing a faithful $\tilde{\mathcal{D}}$.

To facilitate this, we assume a linear structure for the latent reward function (Saha et al., 2023; Kong & Yang, 2022; Zhu et al., 2023). Specifically, we adopt a linear reward model

$$r_{\theta}(x, a) = \langle \theta, \phi(x, a) \rangle,$$

where $\phi(x, a) \in \mathbb{R}^d$ is the *shared* feature vector and $\theta \in \mathbb{R}^d$ is a *user-dependent* parameter.

The linear structure strikes an ideal balance between *expressiveness* and *tractability* for modeling diverse user preferences. Although individual users exhibit heterogeneous preferences, their reward functions share common semantic features in a learned embedding space (i.e., $\phi(x, a)$) and differ primarily in how those features are weighted (i.e., θ). This linear parameterization captures the shared structure while allowing variability across users or preference clusters, making it possible to model preference heterogeneity effectively even with limited data and strong privacy constraints. Indeed, our experiments in Section B.6 show that, despite its simplicity, the linear reward structure captures preference signals accurately and outperforms a fully fine-tuned GPT-2 reward model.

4. Proposed Method

4.1. Key Components of the Algorithm Design

Before detailing our design, we highlight the key components of the algorithm, each of which directly resolves a major challenge in synthesizing differentially private preference data.

Leveraging Linear Reward Structure for Preference Clustering. The linear reward model indicates that

$$\mathbb{P}[a_i^+ \succ a_i^- \mid x_i] = \sigma(\langle \theta, \phi(x_i, a_i^+) - \phi(x_i, a_i^-) \rangle).$$

Therefore, for a tuple (x_i, a_i^+, a_i^-) , $a_i^+ \succ a_i^-$ implies $\langle \theta, \phi(x_i, a_i^+) - \phi(x_i, a_i^-) \rangle \gg 0$, i.e., the preference vector θ is strongly aligned with $\phi(x_i, a_i^+) - \phi(x_i, a_i^-)$. It follows that, for annotators with similar preferences, the feature difference vectors $\phi(x_i, a_i^+) - \phi(x_i, a_i^-)$ across their rated pairs point in similar directions in the embedding space. We exploit this structure directly: we cluster the preference samples by their feature difference vectors $\phi(x_i, a_i^+) - \phi(x_i, a_i^-)$, so that the preferences within each

cluster are well approximated by a shared reward function parameterized by a common preference vector θ . In Section B.9, we confirm that the resulting clusters are meaningful and interpretable in terms of user preference styles.

Adopting DP-PCA to Obtain Effective Features. The natural choice for the feature vector $\phi(x, a)$ is the embedding of the corresponding (prompt, response) pair, since embeddings encode exactly the semantic similarities and differences that underlie human preference judgments. These embeddings, however, are high-dimensional (e.g., 1024 dimensions) so as to capture rich semantic and stylistic features. Learning in such a high-dimensional space demands a large number of samples, which the per-cluster sample budget cannot supply. We overcome this by applying differentially private PCA (DP-PCA) (Amin et al., 2019; Liu et al., 2022) to reduce the embedding dimensionality and produce compact feature representations for clustering and reward model training. This step sharpens sample efficiency and clustering stability while preserving the preference signals essential for effective alignment.

Utilizing Public Prompts to Save DP Budget. Synthesizing prompts, preferred responses, and less preferred responses all at once rapidly exhausts the privacy budget and degrades utility. We sidestep this entirely by using public prompts, which removes any need to spend privacy budget on synthesizing DP prompts and lets us devote the full budget to modeling preferences over responses. For each public prompt, we generate multiple candidate responses with an LLM and construct preference pairs using a distribution of reward functions trained from private preference data under DP. Public prompts may differ in distribution from private ones, yet our experiments show that this distribution shift is harmless: the synthetic preference data preserves the statistical preference distribution of the private dataset and remains effective for downstream alignment.

4.2. DPPrefSyn Algorithm

We now present the DPPrefSyn algorithm (Algorithm 1), which generates a DP synthetic preference dataset from the private dataset $\mathcal{D}_{\text{priv}}$ in 3 main steps:

Step 1: Preference Representation and Clustering. We encode each prompt-response pair in $\mathcal{D}_{\text{priv}}$ with a public sentence embedding model ψ . For each sample (x_i, a_i^+, a_i^-) , we concatenate the prompt with each response to form the two texts $[x_i; a_i^+]$ and $[x_i; a_i^-]$, compute their embeddings via ψ , and take their difference $d_i \leftarrow \psi(x_i, a_i^+) - \psi(x_i, a_i^-)$ (Line 3).

To reduce dimensionality privately, we apply DP-PCA (Amin et al., 2019) to $\{d_i\}_{i=1}^n$, obtaining a projection matrix $\mathbf{P} \in \mathbb{R}^{d \times p}$ under ϵ_0 -DP (Line 5). DP-PCA (Amin et al., 2019) approximates the eigendecomposition of the

data covariance matrix by estimating the eigenvalues and eigenvectors separately in a DP manner. We then project each d_i into the lower-dimensional space as $z_i \leftarrow \mathbf{P}^\top d_i$ (Line 6), and apply DP-KMeans clustering (Su et al., 2016) to $\{z_i\}_{i=1}^n$ to group samples with similar preference patterns under ε_1 -DP, forming K clusters $\mathcal{C}_1, \dots, \mathcal{C}_K$ (Line 7). DP-KMeans (Su et al., 2016) is the DP counterpart of the Lloyd algorithm.

Step 2: DP Reward Model Training. For each cluster $\mathcal{C}_k$, we learn the shared preference vector $\theta_k \in \mathbb{R}^p$ that captures preference patterns by minimizing the negative log-likelihood (Line 9):

$$\mathcal{L}_{\text{MLE}}(\theta_k) = -\frac{1}{|\mathcal{C}_k|} \sum_{z_i \in \mathcal{C}_k} \log \sigma(\langle \theta_k, z_i \rangle), \quad (1)$$

where z_i is the DP-PCA projected embedding difference from Step 1. We train each reward model with the DP-SGD algorithm (a DP variant of SGD) (Abadi et al., 2016) using noise multiplier σ_2 .

Step 3: DP Synthetic Preference Generation. We compute a histogram $\mathbf{h} = [|\mathcal{C}_1|, \dots, |\mathcal{C}_K|]$ representing the number of samples in each cluster (Line 11). We normalize it to get a probability distribution $\mathbf{p}$ over the clusters, i.e., $\mathbf{p} \leftarrow \mathbf{h}/|\mathcal{D}_{\text{priv}}|$ (Line 12).

Next, for each public prompt $\tilde{x}_j$, we prompt an LLM at a relatively high temperature to generate L diverse candidate responses $\tilde{a}_j^1, \dots, \tilde{a}_j^L$ (Line 14). We then sample a cluster index $k \sim \mathbf{p}$ from the DP histogram (Line 15) and score the responses $\tilde{a}_j^1, \dots, \tilde{a}_j^L$ with the reward model parameterized by θ_k : we compute the embeddings $\psi(\tilde{x}_j, \tilde{a}_j^l)$ for $l = 1, \dots, L$, project them into the DP-PCA subspace using $\mathbf{P}$, and evaluate the reward for $\tilde{a}_j^l$ as $\langle \theta_k, \mathbf{P}^\top \psi(\tilde{x}_j, \tilde{a}_j^l) \rangle$. We take the highest- and lowest-scoring responses as the preferred and less preferred responses, respectively: $c \leftarrow \arg \max_{l \in [L]} \langle \theta_k, \mathbf{P}^\top \psi(\tilde{x}_j, \tilde{a}_j^l) \rangle$, $r \leftarrow \arg \min_{l \in [L]} \langle \theta_k, \mathbf{P}^\top \psi(\tilde{x}_j, \tilde{a}_j^l) \rangle$ (Lines 16-17). To guarantee preference quality, we discard the sample whenever the value gap is too small (e.g., < 0.5); otherwise we add $(\tilde{x}_j, \tilde{a}_j^c, \tilde{a}_j^r)$ to $\tilde{\mathcal{D}}$. Repeating this over all public prompts yields the DP synthetic preference dataset $\tilde{\mathcal{D}}$.

Privacy Analysis of DPPrefSyn. We allocate a privacy budget ε_0 to DP-PCA (Line 5) and ε_1 to DP-KMeans (Line 7), and spend the remaining budget training per-cluster reward models with DP-SGD at noise multiplier σ_2 (Line 9). For the reward models we invoke parallel composition of DP: because the clusters are disjoint, modifying one training point affects only one cluster, so the total privacy cost is bounded by the cost of training on the smallest cluster. We enforce a known lower bound on the cluster size by discarding clusters with too few samples. Composing the per-step DP-SGD costs with the Privacy Random Variable (PRV) ac-

Algorithm 1 DPPrefSyn

```

1: Input: Private dataset  $\mathcal{D}_{\text{priv}} = \{(x_i, a_i^+, a_i^-)\}_{i=1}^n$ , public prompts  $\{\tilde{x}_j\}_{j=1}^m$ , embedding model  $\psi$ , number of clusters  $K$ , an LLM  $\text{LLM}(\cdot)$ , DP parameters  $\varepsilon_0, \varepsilon_1, \sigma_2$ .
2: for each  $(x_i, a_i^+, a_i^-)$  in  $\mathcal{D}_{\text{priv}}$  do
3:    $d_i \leftarrow \psi(x_i, a_i^+) - \psi(x_i, a_i^-)$ , where  $d_i \in \mathbb{R}^d$ 
4: end for
5:  $\mathbf{P} \leftarrow \text{DP-PCA}(\{d_i\}_{i=1}^n, \varepsilon_0)$ , where  $\mathbf{P} \in \mathbb{R}^{d \times p}$ 
6:  $\{z_i\}_{i=1}^n \leftarrow \{\mathbf{P}^\top d_i\}_{i=1}^n$ 
7:  $\mathcal{C}_1, \dots, \mathcal{C}_K \leftarrow \text{DP-KMeans}(\{z_i\}_{i=1}^n, K, \varepsilon_1)$ 
8: for each cluster  $\mathcal{C}_k$  do
9:   Train a linear reward model  $\theta_k$  on  $\mathcal{C}_k$  using DP-SGD with noise multiplier  $\sigma_2$  (Equation (1))
10: end for
11:  $\mathbf{h} \leftarrow [|\mathcal{C}_1|, \dots, |\mathcal{C}_K|]$ 
12:  $\mathbf{p} \leftarrow \mathbf{h}/|\mathcal{D}_{\text{priv}}|$ 
13: for each public prompt  $\tilde{x}_j$  do
14:   Generate responses  $\tilde{a}_j^1, \dots, \tilde{a}_j^L \leftarrow \text{LLM}(\tilde{x}_j)$ 
15:   Sample cluster index  $k \sim \mathbf{p}$ 
16:    $c \leftarrow \arg \max_{l \in [L]} \langle \theta_k, \mathbf{P}^\top \psi(\tilde{x}_j, \tilde{a}_j^l) \rangle$ 
17:    $r \leftarrow \arg \min_{l \in [L]} \langle \theta_k, \mathbf{P}^\top \psi(\tilde{x}_j, \tilde{a}_j^l) \rangle$ 
18:   Add  $(\tilde{x}_j, \tilde{a}_j^c, \tilde{a}_j^r)$  to  $\tilde{\mathcal{D}}$ 
19: end for
20: return  $\tilde{\mathcal{D}}$ .
```

countant (Gopi et al., 2021) guarantees $(\varepsilon - \varepsilon_0 - \varepsilon_1, \delta)$ -DP for this stage. Every other step touches only public data or post-processing and incurs no privacy cost. Together, these facts establish that DPPrefSyn is (ε, δ) -DP; we give the full proof and privacy-accounting code in Section A.

Discussion on Prompt Distribution Shift. DPPrefSyn stays effective even when the prompt distributions of the private and public datasets differ. The reason is structural: diverse user preferences are captured through θ_k , the parameter of the linear reward model, not through $\phi(x, a)$, the feature vector of the prompt-response pair. Intuitively, a user may interact with an LLM through many different prompts while their underlying preference stays fixed. Hence, as long as the distribution of θ_k matches that of the private dataset, the synthetic preferences faithfully reflect the private preference patterns. We achieve this by sampling a cluster index k from the DP histogram (Line 15) for each public prompt and selecting the synthetic response pairs accordingly. Consequently, DPPrefSyn preserves the private preference patterns regardless of any shift in the prompt distribution.

5. Experiments

Datasets. We evaluate on both question answering and summarization tasks. For question answering, we use the OpenAssistant dataset (Köpf et al., 2023) of assistant-style con-

versations and the Anthropic-HH dataset (Bai et al., 2022) of human preference comparisons focused on helpfulness and harmlessness. For summarization, we use the TL;DR dataset (Stiennon et al., 2020) of human preferences over pairs of summaries. To simulate public data, we draw prompts from Alpaca (Taori et al., 2023) for the OpenAssistant QA task, SafeRLHF (Ji et al., 2024) for the Anthropic-HH QA task, and XSum (Narayan et al., 2018) for the TL;DR summarization task. Additional details appear in Section B.2.

Setup for Algorithm 1. By default, we use BAAI/bge-large-en-v1.5 (Xiao et al., 2023) as the embedding model ψ . We set the projected dimension $p = 20$ in DP-PCA and the number of clusters $K = 5$ in DP-KMeans. For DP-SGD, we use a learning rate of 0.1, a batch size of 4, 4 training epochs, and a gradient clipping norm of 1.0. Candidate responses are generated using the instruction-finetuned LLaMA-7B-chat model (Touvron et al., 2023) with a temperature of 0.9. We set the number of candidates per prompt to $L = 5$. We consider overall privacy budgets $\epsilon = 0.5, 1, 2, 4, 8$ and set $\delta = 1/|\mathcal{D}_{\text{priv}}|$. We study the effect of different generator models in Table 3, and of embedding models and hyperparameters in Section 5.3.

Evaluation. For downstream preference alignment, we fine-tune the Pythia-2.8B model (Biderman et al., 2023). We first apply supervised fine-tuning (SFT), using the preferred response in the preference dataset as the training target, and then apply the DPO algorithm (Rafailov et al., 2024) to further fine-tune the SFT model on preference pairs. Following Rafailov et al. (2024), we measure the win rate of the model-generated responses against the preferred responses in the test set with GPT-4o. We study additional downstream models and different preference alignment methods as ablations in Section 5.2.

Baselines. We compare DPPrefSyn against fine-tuning the downstream model on real data with DP-SGD (Abadi et al., 2016), denoted DP-FT. DPPrefSyn is fundamentally more reusable and versatile than DP-FT: it produces DP synthetic data that can train any model with any alignment algorithm at no additional privacy cost, whereas DP-FT spends its entire privacy budget on a single privatized model. We also include a fully private baseline $\epsilon = 0$, in which we evaluate the base model without any preference fine-tuning.

Details about the setups, hyperparameters, metrics, and baselines are provided in Section B.3.

5.1. Understanding the Performance of DPPrefSyn

In this section, we analyze the performance of DPPrefSyn through two research questions about its privacy-utility trade-off. Our main results are presented in Table 1.

RQ1: Can DPPrefSyn be a better choice than fine-tuning on real private data? Yes. DPPrefSyn delivers a stronger

privacy-utility trade-off than DP-FT on every task at $\epsilon = 0.5, 1, 2, 4, 8$, and while preserving privacy it even exceeds the utility of fine-tuning on real private data without any DP constraint (DP-FT with $\epsilon = \infty$). On the OpenAssistant QA task with $\epsilon = 2$, for instance, our approach reaches 9.10% after SFT and 11.04% after DPO, versus DP-FT’s 4.38% and 6.18%. Strikingly, DPPrefSyn also outperforms the non-private baseline of fine-tuning on real data (DP-FT with $\epsilon = \infty$), which reaches only 4.75% after SFT and 8.20% after DPO. The same pattern holds on Anthropic-HH with $\epsilon = 2$: DPPrefSyn attains a 55.21% win rate after SFT and 56.48% after DPO, compared to 31.53% and 37.02% for DP-FT, and again surpasses non-private fine-tuning on real data (DP-FT with $\epsilon = \infty$), which reaches 31.98% after SFT and 38.72% after DPO.

These results expose two **key advantages** of our DP synthetic data. **First**, DPPrefSyn yields a large improvement after SFT alone relative to applying SFT directly on private data. On the TL;DR summarization task with $\epsilon = 2$, for example, DPPrefSyn reaches a 62.49% win rate after SFT, nearly triple the 22.57% of SFT on private data. This gain comes from generating candidate responses with a high-performance LLM, which produces high-quality training data and thus far stronger supervision during SFT. The advantage persists even with scarce public data: on the OpenAssistant QA task with $\epsilon = 4$ and only 5K public prompts, DPPrefSyn attains 6.80% after SFT, against 4.46% for SFT on the roughly 14K private samples (Table 4), demonstrating that DPPrefSyn handles data scarcity without sacrificing utility. **Second**, applying DPO with our synthetic data on top of the SFT model drives further gains, confirming that our synthetic preference pairs encode meaningful distinctions aligned with human preferences—a direct consequence of the private reward functions that capture diverse user preferences while preserving privacy. The incremental DPO gain is smaller here than when using private data, simply because the SFT model trained on our synthetic data already performs strongly and leaves less headroom. Even so, over the full fine-tuning pipeline the total performance gain from our DP synthetic dataset remains far larger than that from real private data.

RQ2: How does DPPrefSyn perform across different privacy budgets ϵ ? First, Table 1 shows that DPPrefSyn outperforms DP-FT as ϵ ranges over 0.5, 1, 2, 4, 8 and ∞ , demonstrating that it scales gracefully with the privacy budget. **Second**, DPPrefSyn is markedly more robust to DP noise than DP-FT. On OpenAssistant with Llama-3-8B as the downstream model, for example, DP-FT’s DPO performance collapses from 59.13% to 45.79% as ϵ tightens from ∞ to 4, whereas DPPrefSyn barely moves, from 61.52% to 60.25% (see Table 3 in Section 5.2). This robustness follows directly from where the noise is injected: DP-FT perturbs model parameters through DP-SGD, while DPPrefSyn adds

Table 1. GPT-4o win rates on the test sets with the downstream Pythia-2.8B model fine-tuned by SFT and DPO. $\varepsilon = 0$ denotes the base (fully private, non-finetuned) LLM. At every privacy level $\varepsilon = 0.5, 1, 2, 4, 8$, DPPrefSyn attains higher win rates than DP-FT, and it even surpasses the utility of non-private fine-tuning (DP-FT with $\varepsilon = \infty$). Results show mean and standard deviation over 5 random seeds.

Task	$\varepsilon = 0$	Method	Data Type	Pref. Alignment	$\varepsilon = 0.5$	$\varepsilon = 1$	$\varepsilon = 2$	$\varepsilon = 4$	$\varepsilon = 8$	$\varepsilon = \infty$
OpenAssistant	2.11	DP-FT	Original	SFT	3.96 _{0.20}	4.29 _{0.21}	4.38 _{0.23}	4.46 _{0.18}	4.49 _{0.30}	4.75 _{0.27}
				SFT+DPO	5.65 _{0.32}	6.04 _{0.58}	6.18 _{0.83}	6.12 _{0.63}	6.32 _{0.71}	8.20 _{0.62}
		DPPrefSyn	Synthetic	SFT	8.96 _{0.39}	9.07 _{0.32}	9.10 _{0.51}	9.33 _{0.56}	9.63 _{0.38}	9.80 _{0.58}
				SFT+DPO	9.86_{0.59}	10.31_{0.54}	11.04_{0.92}	11.85_{0.99}	12.05_{0.82}	12.36_{0.55}
Anthropic-HH	12.14	DP-FT	Original	SFT	29.77 _{0.76}	31.10 _{0.72}	31.53 _{0.77}	31.29 _{0.65}	31.37 _{0.67}	31.98 _{0.56}
				SFT+DPO	35.00 _{0.49}	36.27 _{0.63}	37.02 _{0.72}	36.74 _{0.90}	36.94 _{1.06}	38.72 _{0.56}
		DPPrefSyn	Synthetic	SFT	54.90 _{0.66}	55.19 _{0.91}	55.21 _{0.82}	54.95 _{1.10}	55.63 _{0.86}	55.95 _{0.70}
				SFT+DPO	55.08_{1.04}	55.96_{0.80}	56.48_{0.61}	56.51_{1.93}	56.86_{1.17}	57.53_{0.93}
TL;DR	11.64	DP-FT	Original	SFT	21.19 _{1.12}	22.41 _{1.09}	22.57 _{0.65}	22.45 _{0.90}	23.04 _{0.79}	23.82 _{1.26}
				SFT+DPO	59.69 _{0.82}	62.21 _{1.12}	62.02 _{1.67}	63.00 _{0.73}	63.06 _{0.84}	64.52 _{1.79}
		DPPrefSyn	Synthetic	SFT	61.92 _{1.06}	62.14 _{1.61}	62.49 _{1.48}	62.72 _{1.66}	62.29 _{1.14}	62.55 _{1.23}
				SFT+DPO	67.50_{0.79}	68.07_{1.95}	68.54_{2.76}	68.56_{1.59}	71.01_{2.17}	72.20_{1.08}

Table 2. DPPrefSyn achieves higher win rates than fine-tuning on private data across all of SFT, DPO, and PPO, while providing strong privacy guarantees on OpenAssistant. Results show mean and standard deviation over 5 random seeds.

Pref. Alignment	$\varepsilon = 0$	DPPrefSyn ($\varepsilon = 4$)	DP-FT ($\varepsilon = \infty$)
SFT		9.33_{0.56}	4.75 _{0.27}
SFT+DPO	2.11	11.85_{0.99}	8.20 _{0.62}
SFT+PPO		10.11_{0.61}	7.08 _{0.78}

noise only during data synthesis, which preserves utility far more effectively.

5.2. Understanding the Properties of DPPrefSyn

In this section, we establish the key properties of DPPrefSyn: its generalizability across different preference alignment methods and LLMs, its performance under varying scales of public data, and its robustness to empirical privacy attacks.

RQ3: Is DPPrefSyn effective across different preference alignment methods? Yes. DPPrefSyn supports diverse preference alignment methods and stays effective under strong DP constraints. In Table 2, we evaluate DPPrefSyn with SFT, DPO, and PPO. For PPO, we train K reward models (one per cluster) to capture diverse preference patterns and optimize the policy against the average of their outputs. Across all three methods, DPPrefSyn delivers higher utility than fine-tuning on private data while preserving strong privacy guarantees.

RQ4: Can more powerful downstream LLMs benefit from synthetic data generated by DPPrefSyn? Yes. DPPrefSyn is effective across a wide range of downstream LLMs. In Table 3, we experiment on downstream models spanning Qwen-3-4B, Gemma-2-2B, and Llama-3-8B, with Llama-3-8B-Instruct and Qwen3-4B-Instruct-2507 as generators. After DPO, DPPrefSyn outperforms fine-tuning on real data across every model family while ensuring privacy. Using Qwen3-4B-Instruct-2507 as the generator, DPPrefSyn

Table 3. DPPrefSyn achieves higher win rates than fine-tuning on private data across every downstream LLM after DPO, while providing strong privacy guarantees on OpenAssistant. DPPrefSyn further benefits from more powerful generator LLMs.

Downstream LLM	$\varepsilon = 0$	Method	Generator LLM	$\varepsilon = 4$	$\varepsilon = \infty$
Qwen-3-4B	48.17	DP-FT	-	54.92	59.41
			-		
		DPPrefSyn	Qwen-3-4B-Instruct	86.52	87.08
			Llama-3-8B-Instruct	51.69	52.11
Gemma-2-2B	7.30	DP-FT	-	30.06	33.85
			-		
		DPPrefSyn	Qwen-3-4B-Instruct	36.24	37.08
			Llama-3-8B-Instruct	36.01	36.66
Llama-3-8B	16.15	DP-FT	-	45.79	59.13
			-		
		DPPrefSyn	Qwen-3-4B-Instruct	60.25	61.52
			Llama-3-8B-Instruct	42.28	44.10

Table 4. DPPrefSyn achieves higher win rates than fine-tuning on private data at every public data size m . Results show mean and standard deviation over 5 random seeds.


	DPPrefSyn ($\varepsilon = 4$)			DP-FT ($\varepsilon = 4$)	DP-FT ($\varepsilon = \infty$)
Data Type (Size)	Synthetic ($m = 5,000$)	Synthetic ($m = 10,000$)	Synthetic ($m = 52,000$)	Original (14,167)	Original (14,167)
SFT	6.80 _{0.66}	7.30 _{0.48}	9.33 _{0.56}	4.46 _{0.18}	4.75 _{0.27}
SFT+DPO	9.44 _{0.45}	9.55 _{0.92}	11.85 _{0.99}	6.12 _{0.63}	8.20 _{0.62}

reaches an 86.52% win rate on Qwen-3-4B with $\varepsilon = 4$, dramatically exceeding non-private fine-tuning on real data (59.41%). Performance also rises with stronger generators, Qwen3-4B-Instruct-2507 outperforming Llama-3-8B-Instruct, showing that DPPrefSyn directly inherits the gains of more capable generator models.


RQ5: Can DPPrefSyn be effective with varying scales of public data? We study the effect of the public data size m in Table 4. The private dataset (OpenAssistant) contains 14K preference annotations, and the full public dataset (Alpaca) contains 52K prompts; we randomly sample subsets of public prompts at sizes $m = 5K, 10K, 52K$ to synthesize preference data with DPPrefSyn. Even with only 5K public prompts, DPPrefSyn reaches a 6.80% win rate after SFT and 9.44% after DPO under $\varepsilon = 4$, clearly beating the DP-FT baseline’s 4.46% and 6.12%. This confirms that

Table 5. DPPrefSyn drives AUC scores against MIAs (Feng et al., 2025) on OpenAssistant down to near random guessing, well below DP-FT. Results show mean and standard deviation over 5 random seeds.

Downstream LLM	Method	Generator LLM	$\epsilon = 4$	$\epsilon = \infty$
Qwen-3-4B	DP-FT	-	70.15 _{2.09}	73.52 _{2.30}
	DPPrefSyn	Qwen-3-4B-Instruct	50.99 _{1.10}	50.87 _{0.84}
		Llama-3-8B-Instruct	50.64_{1.29}	49.98_{1.19}
Gemma-2-2B	DP-FT	-	62.88 _{1.78}	65.39 _{2.59}
	DPPrefSyn	Qwen-3-4B-Instruct	50.51_{1.38}	50.43_{0.91}
		Llama-3-8B-Instruct	50.77 _{0.83}	51.10 _{0.54}
Llama-3-8B	DP-FT	-	61.95 _{2.05}	70.13 _{1.24}
	DPPrefSyn	Qwen-3-4B-Instruct	50.53 _{0.62}	50.96 _{0.95}
		Llama-3-8B-Instruct	50.45_{0.87}	50.67_{0.96}



(a) Removing DP-KMeans or DP-PCA.



(b) Varying clusters.

DPPrefSyn remains effective even when public data is scarce, thanks to the high-quality LLM-generated responses and the private reward functions that accurately capture user preferences. Moreover, **the advantage over the baseline widens as the public data grows**, showing that DPPrefSyn readily exploits the vast amounts of available public data for preference alignment and thereby further dissolves the data-scarcity bottleneck of private preference data.

RQ6: Can DPPrefSyn effectively mitigate empirical privacy attacks? In Table 5, we run membership inference attacks (MIAs) (Shokri et al., 2017) against the DPO fine-tuned downstream models on OpenAssistant and report the AUC. We follow Feng et al. (2025), who instantiate a MIA on preference data for LLM alignment; the goal is to determine whether a private sample (x, a^+, a^-) was used in alignment. DP-FT fails to mitigate MIAs, whereas **DPPrefSyn consistently drives the AUC down to near random guessing**. DPPrefSyn is therefore substantially more robust to empirical privacy attacks, precisely because fine-tuning on synthetic rather than real data eliminates the risk of overfitting to private data. Full details appear in Section B.4.

5.3. Validating the Design of DPPrefSyn

In this section, we validate the design of DPPrefSyn by dissecting its algorithmic components.

RQ7: How does modeling diverse human preferences impact DPPrefSyn’s performance? We answer this by removing either DP-KMeans or DP-PCA and by varying the number of clusters K , all on OpenAssistant with $\epsilon = 4$. **(i) Removing either DP-KMeans or DP-PCA.** Figure 2a shows that

removing either component clearly degrades performance: without DP-KMeans all samples collapse into a single cluster, and without DP-PCA clustering becomes ineffective and learning sample-inefficient in the high-dimensional space. Using multiple clusters with both components in place consistently beats the single-cluster variant, confirming that both are necessary. **(ii) Varying the number of clusters K .** In Figure 2b, we evaluate $K = 1, 3, 5, 8$. Performance jumps at $K = 3$ and then holds steady at $K = 5, 8$, demonstrating that DPPrefSyn is robust to K and that moderate values (e.g., $K = 3, 5$) strike the right balance between diversity and data efficiency. Together, these results confirm the concrete benefit of capturing diverse preferences through multiple clusters.

RQ8: How do different embedding models or hyperparameters affect DPPrefSyn’s performance? We ablate the embedding model ψ , DP-PCA dimension p , number of candidates per prompt L , filter threshold, and privacy budget allocation in Section B.5. DP-PrefSyn stays consistently effective across sentence encoders, including bge-large-en-v1.5 (Xiao et al., 2023), all-mpnet-base-v2 (Song et al., 2020), and all-distilroberta-v1 (Sanh et al., 2019). We find that $p = 20$ gives the best overall performance after DPO, $L = 5$ is a solid default, and a moderate filtering threshold of 0.5 slightly outperforms no filtering, while an aggressive threshold of 1 hurts performance. Finally, DPPrefSyn is robust to reasonable changes in the budget allocation ϵ_0 and ϵ_1 .

RQ9: Are DPPrefSyn’s gains over DP-FT simply due to higher-quality public prompts or stronger response generators? No. To rule out these confounders, we conduct controlled ablations in which all methods share the same prompt set or the same response generator. DPPrefSyn continues to outperform DP-FT under both settings (Appendix B.7), establishing that **its gains do not stem from higher-quality public prompts or stronger response generators** but from the method itself.

6. Conclusion

In this work, we introduced DPPrefSyn, the first algorithm for synthesizing DP preference data from a private dataset. By modeling preference diversity via clustering, sharpening efficiency with DP-PCA, and conserving the privacy budget through public prompts, DPPrefSyn makes effective and private preference data synthesis for LLM alignment possible. Our experiments demonstrate that DPPrefSyn delivers rigorous privacy guarantees while outperforming fine-tuning on real data, establishing synthetic preference data as a practical foundation for privacy-preserving alignment.

7. Limitations

DPPrefSyn requires LLM API access to synthesize preference data, which may incur additional API costs. We focus on English text-based preference data; extending DPPrefSyn to multilingual and multimodal settings is a natural direction for future work. Finally, our experiments cover models up to 8B parameters due to computational constraints, though the method itself is not tied to this scale.

Impact statement

This work enables preference alignment of large language models without exposing personal information, reconciling the benefits of big data with the imperative to protect individual privacy, and thereby promoting ethical data usage and fostering societal trust in an increasingly data-driven world.

References

- Abadi, M., Chu, A., Goodfellow, I., McMahan, H. B., Mironov, I., Talwar, K., and Zhang, L. Deep learning with differential privacy. In *Proceedings of the 2016 ACM SIGSAC conference on computer and communications security*, pp. 308–318, 2016.
- Achiam, J., Adler, S., Agarwal, S., Ahmad, L., Akkaya, I., Aleman, F. L., Almeida, D., Altenschmidt, J., Altman, S., Anadkat, S., et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- Amin, K., Dick, T., Kulesza, A., Munoz, A., and Vassilvitskii, S. Differentially private covariance estimation. *Advances in Neural Information Processing Systems*, 32, 2019.
- Aroyo, L., Diaz, M., Homan, C., Prabhakaran, V., Taylor, A., and Wang, D. The reasonable effectiveness of diverse evaluation data. *arXiv preprint arXiv:2301.09406*, 2023a.
- Aroyo, L., Taylor, A., Diaz, M., Homan, C., Parrish, A., Serapio-García, G., Prabhakaran, V., and Wang, D. Dices dataset: Diversity in conversational ai evaluation for safety. *Advances in Neural Information Processing Systems*, 36:53330–53342, 2023b.
- Bai, Y., Jones, A., Ndousse, K., Askell, A., Chen, A., Das-Sarma, N., Drain, D., Fort, S., Ganguli, D., Henighan, T., et al. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*, 2022.
- Bakker, M., Chadwick, M., Sheahan, H., Tessler, M., Campbell-Gillingham, L., Balaguer, J., McAleese, N., Glaese, A., Aslanides, J., Botvinick, M., et al. Fine-tuning language models to find agreement among humans with diverse preferences. *Advances in Neural Information Processing Systems*, 35:38176–38189, 2022.
- Biderman, S., Schoelkopf, H., Anthony, Q. G., Bradley, H., O’Brien, K., Hallahan, E., Khan, M. A., Purohit, S., Prashanth, U. S., Raff, E., et al. Pythia: A suite for analyzing large language models across training and scaling. In *International Conference on Machine Learning*, pp. 2397–2430. PMLR, 2023.
- Bradley, R. A. and Terry, M. E. Rank analysis of incomplete block designs: I. the method of paired comparisons. *Biometrika*, 39(3/4):324–345, 1952.
- Carranza, A., Farahani, R., Ponomareva, N., Kurakin, A., Jagielski, M., and Nasr, M. Synthetic query generation for privacy-preserving deep retrieval systems using differentially private language models. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 3920–3930, 2024.
- Carranza, A. G., Farahani, R., Ponomareva, N., Kurakin, A., Jagielski, M., and Nasr, M. Privacy-preserving recommender systems with synthetic query generation using differentially private large language models. *arXiv preprint arXiv:2305.05973*, 2023.
- Chakraborty, S., Qiu, J., Yuan, H., Koppel, A., Manocha, D., Huang, F., Bedi, A., and Wang, M. Maxmin-rlhf: Alignment with diverse human preferences. In *International Conference on Machine Learning*, pp. 6116–6135. PMLR, 2024.
- Chen, T., Da, L., Zhou, H., Li, P., Zhou, K., Chen, T., and Wei, H. Privacy-preserving fine-tuning of large language models through flatness. *arXiv preprint arXiv:2403.04124*, 2024.
- Chowdhury, S. R., Zhou, X., and Natarajan, N. Differentially private reward estimation with preference feedback. In *International Conference on Artificial Intelligence and Statistics*, pp. 4843–4851. PMLR, 2024.
- Denton, R., Díaz, M., Kivlichan, I., Prabhakaran, V., and Rosen, R. Whose ground truth? accounting for individual and collective identities underlying dataset annotation. *arXiv preprint arXiv:2112.04554*, 2021.
- Duan, H., Dziedzic, A., Papernot, N., and Boenisch, F. Flocks of stochastic parrots: Differentially private prompt learning for large language models. *Advances in Neural Information Processing Systems*, 36:76852–76871, 2023.

- Dwork, C., Kenthapadi, K., McSherry, F., Mironov, I., and Naor, M. Our data, ourselves: Privacy via distributed noise generation. In *Advances in Cryptology-EUROCRYPT 2006: 24th Annual International Conference on the Theory and Applications of Cryptographic Techniques, St. Petersburg, Russia, May 28-June 1, 2006. Proceedings* 25, pp. 486–503. Springer, 2006a.
- Dwork, C., McSherry, F., Nissim, K., and Smith, A. Calibrating noise to sensitivity in private data analysis. In *Theory of Cryptography: Third Theory of Cryptography Conference, TCC 2006, New York, NY, USA, March 4-7, 2006. Proceedings* 3, pp. 265–284. Springer, 2006b.
- Dwork, C., Roth, A., et al. The algorithmic foundations of differential privacy. *Foundations and Trends® in Theoretical Computer Science*, 9(3–4):211–407, 2014.
- Feng, Q., Kasa, S. R., KASA, S. K., Yun, H., Teo, C. H., and Bodapati, S. B. Exposing privacy gaps: Membership inference attack on preference data for llm alignment. In *International Conference on Artificial Intelligence and Statistics*, pp. 5221–5229. PMLR, 2025.
- Gao, F., Zhou, R., Wang, T., Shen, C., and Yang, J. Data-adaptive differentially private prompt synthesis for in-context learning. *arXiv preprint arXiv:2410.12085*, 2024.
- Ghazi, B., Golowich, N., Kumar, R., Manurangsi, P., and Zhang, C. Deep learning with label differential privacy. *Advances in neural information processing systems*, 34: 27131–27145, 2021.
- Gopi, S., Lee, Y. T., and Wutschitz, L. Numerical composition of differential privacy. *Advances in Neural Information Processing Systems*, 34:11631–11642, 2021.
- He, J., Li, X., Yu, D., Zhang, H., Kulkarni, J., Lee, Y. T., Backurs, A., Yu, N., and Bian, J. Exploring the limits of differentially private deep learning with group-wise clipping. *arXiv preprint arXiv:2212.01539*, 2022.
- Holohan, N., Braghin, S., Mac Aonghusa, P., and Levacher, K. Diffprivlib: the IBM differential privacy library. *ArXiv e-prints*, 1907.02444 [cs.CR], July 2019.
- Hong, J., Wang, J. T., Zhang, C., Li, Z., Li, B., and Wang, Z. Dp-opt: Make large language model your privacy-preserving prompt engineer. *arXiv preprint arXiv:2312.03724*, 2023.
- Hou, C., Shrivastava, A., Zhan, H., Conway, R., Le, T., Sagar, A., Fanti, G., and Lazar, D. Pre-text: Training language models on private federated data in the age of llms. In *International Conference on Machine Learning*, pp. 19043–19061. PMLR, 2024.
- Jang, J., Kim, S., Lin, B. Y., Wang, Y., Hessel, J., Zettlemoyer, L., Hajishirzi, H., Choi, Y., and Ammanabrolu, P. Personalized soups: Personalized large language model alignment via post-hoc parameter merging. *arXiv preprint arXiv:2310.11564*, 2023.
- Ji, J., Liu, M., Dai, J., Pan, X., Zhang, C., Bian, C., Chen, B., Sun, R., Wang, Y., and Yang, Y. Beavertails: Towards improved safety alignment of llm via a human-preference dataset. *Advances in Neural Information Processing Systems*, 36:24678–24704, 2023.
- Ji, J., Hong, D., Zhang, B., Chen, B., Dai, J., Zheng, B., Qiu, T., Li, B., and Yang, Y. Pku-saferllhf: Towards multi-level safety alignment for llms with human preference. *arXiv preprint arXiv:2406.15513*, 2024.
- Kirk, R., Mediratta, I., Nalmpantis, C., Luketina, J., Hambro, E., Grefenstette, E., and Raileanu, R. Understanding the effects of rlhf on llm generalisation and diversity. In *International Conference on Learning Representations*, volume 2024, pp. 20620–20653, 2024.
- Kong, D. and Yang, L. Provably feedback-efficient reinforcement learning via active reward learning. *Advances in Neural Information Processing Systems*, 35:11063–11078, 2022.
- Köpf, A., Kilcher, Y., Von Rütte, D., Anagnostidis, S., Tam, Z. R., Stevens, K., Barhoum, A., Nguyen, D., Stanley, O., Nagyfi, R., et al. Openassistant conversations-democratizing large language model alignment. *Advances in Neural Information Processing Systems*, 36:47669–47681, 2023.
- Kovač, G., Sawayama, M., Portelas, R., Colas, C., Dominey, P. F., and Oudeyer, P.-Y. Large language models as superpositions of cultural perspectives. *arXiv preprint arXiv:2307.07870*, 2023.
- Lee, H., Phatale, S., Mansoor, H., Mesnard, T., Ferret, J., Lu, K. R., Bishop, C., Hall, E., Carbune, V., Rastogi, A., et al. Rlaif vs. rlhf: Scaling reinforcement learning from human feedback with ai feedback. In *International Conference on Machine Learning*, pp. 26874–26901. PMLR, 2024a.
- Lee, Y., Williams, J., Marklund, H., Sharma, A., Mitchell, E., Singh, A., and Finn, C. Test-time alignment via hypothesis reweighting. *arXiv preprint arXiv:2412.08812*, 2024b.
- Li, H., Guo, D., Fan, W., Xu, M., Huang, J., Meng, F., and Song, Y. Multi-step jailbreaking privacy attacks on chatgpt. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pp. 4138–4153, 2023.
- Li, X., Tramer, F., Liang, P., and Hashimoto, T. Large language models can be strong differentially private learners. *arXiv preprint arXiv:2110.05679*, 2021.

- Lin, Z., Gopi, S., Kulkarni, J., Nori, H., and Yekhanin, S. Differentially private synthetic data via foundation model apis 1: Images. *arXiv preprint arXiv:2305.15560*, 2023.
- Liu, T., Zhao, Y., Joshi, R., Khalman, M., Saleh, M., Liu, P. J., and Liu, J. Statistical rejection sampling improves preference optimization. *arXiv preprint arXiv:2309.06657*, 2023.
- Liu, X., Kong, W., Jain, P., and Oh, S. Dp-pca: Statistically optimal and differentially private pca. *Advances in neural information processing systems*, 35:29929–29943, 2022.
- Mattern, J., Jin, Z., Weggenmann, B., Schoelkopf, B., and Sachan, M. Differentially private language models for secure data sharing. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pp. 4860–4873, 2022.
- Mireshghallah, F., Su, Y., Hashimoto, T. B., Eisner, J., and Shin, R. Privacy-preserving domain adaptation of semantic parsers. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 4950–4970, 2023.
- Narayan, S., Cohen, S. B., and Lapata, M. Don’t give me the details, just the summary! topic-aware convolutional neural networks for extreme summarization. *arXiv preprint arXiv:1808.08745*, 2018.
- Ochs, S. and Habernal, I. Private synthetic text generation with diffusion models. *arXiv preprint arXiv:2410.22971*, 2024.
- Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744, 2022.
- Poddar, S., Wan, Y., Ivison, H., Gupta, A., and Jaques, N. Personalizing reinforcement learning from human feedback with variational preference learning. *arXiv preprint arXiv:2408.10075*, 2024.
- Rafailov, R., Sharma, A., Mitchell, E., Manning, C. D., Ermon, S., and Finn, C. Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36, 2024.
- Rame, A., Couaeron, G., Dancette, C., Gaya, J.-B., Shukor, M., Soulier, L., and Cord, M. Rewarded soups: towards pareto-optimal alignment by interpolating weights fine-tuned on diverse rewards. *Advances in Neural Information Processing Systems*, 36:71095–71134, 2023.
- Saha, A., Pacchiano, A., and Lee, J. Dueling rl: Reinforcement learning with trajectory preferences. In *International conference on artificial intelligence and statistics*, pp. 6263–6289. PMLR, 2023.
- Sandri, M., Leonardelli, E., Tonelli, S., and Ježek, E. Why don’t you do it right? analysing annotators’ disagreement in subjective tasks. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics*, pp. 2428–2441, 2023.
- Sanh, V., Debut, L., Chaumond, J., and Wolf, T. Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*, 2019.
- Shao, Z., Wang, P., Zhu, Q., Xu, R., Song, J., Bi, X., Zhang, H., Zhang, M., Li, Y., Wu, Y., et al. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.
- Shen, B., Zhang, J., Chen, T., Zan, D., Geng, B., Fu, A., Zeng, M., Yu, A., Ji, J., Zhao, J., et al. Pangu-coder2: Boosting large language models for code with ranking feedback. *arXiv preprint arXiv:2307.14936*, 2023.
- Shoemate, M., Vyrros, A., McCallum, C., Prasad, R., Durbin, P., Casacuberta Puig, S., Cowan, E., Xu, V., Ratliff, Z., Berrios, N., Whitworth, A., Eliot, M., Lebeda, C., Renard, O., and McKay Bowen, C. OpenDP Library, 2021. URL <https://github.com/opendp/opendp>.
- Shokri, R., Stronati, M., Song, C., and Shmatikov, V. Membership inference attacks against machine learning models. In *2017 IEEE symposium on security and privacy (SP)*, pp. 3–18. IEEE, 2017.
- Singh, A., Hsu, S., Hsu, K., Mitchell, E., Ermon, S., Hashimoto, T., Sharma, A., and Finn, C. Fspo: Few-shot preference optimization of synthetic preference data in llms elicits effective personalization to real users. *arXiv preprint arXiv:2502.19312*, 2025.
- Song, K., Tan, X., Qin, T., Lu, J., and Liu, T.-Y. Mpnnet: Masked and permuted pre-training for language understanding. *Advances in neural information processing systems*, 33:16857–16867, 2020.
- Stiennon, N., Ouyang, L., Wu, J., Ziegler, D., Lowe, R., Voss, C., Radford, A., Amodei, D., and Christiano, P. F. Learning to summarize with human feedback. *Advances in neural information processing systems*, 33:3008–3021, 2020.
- Su, D., Cao, J., Li, N., Bertino, E., and Jin, H. Differentially private k-means clustering. In *Proceedings of the sixth ACM conference on data and application security and privacy*, pp. 26–37, 2016.

- Tan, B., Xu, Z., Xing, E., Hu, Z., and Wu, S. Synthesizing privacy-preserving text data via finetuning without finetuning billion-scale llms. *arXiv preprint arXiv:2503.12347*, 2025.
- Tang, X., Shin, R., Inan, H. A., Manoel, A., Miresghallah, F., Lin, Z., Gopi, S., Kulkarni, J., and Sim, R. Privacy-preserving in-context learning with differentially private few-shot generation. *arXiv preprint arXiv:2309.11765*, 2023.
- Tang, X., Panda, A., Nasr, M., Mahlouljifar, S., and Mittal, P. Private fine-tuning of large language models with zeroth-order optimization. *arXiv preprint arXiv:2401.04343*, 2024.
- Taori, R., Gulrajani, I., Zhang, T., Dubois, Y., Li, X., Guestrin, C., Liang, P., and Hashimoto, T. B. Alpaca: A strong, replicable instruction-following model. *Stanford Center for Research on Foundation Models*. <https://crfm.stanford.edu/2023/03/13/alpaca.html>, 3(6):7, 2023.
- Touvron, H., Lavril, T., Izacard, G., Martinet, X., Lachaux, M.-A., Lacroix, T., Rozière, B., Goyal, N., Hambro, E., Azhar, F., et al. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- Vogels, E. A. *The state of online harassment*, volume 13. Pew Research Center Washington, DC, 2021.
- Wang, W., Liang, X., Ye, R., Chai, J., Chen, S., and Wang, Y. Knowledgesg: Privacy-preserving synthetic text generation with knowledge distillation from server. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 7677–7695, 2024.
- Wu, F., Inan, H. A., Backurs, A., Chandrasekaran, V., Kulkarni, J., and Sim, R. Privately aligning language models with reinforcement learning. *arXiv preprint arXiv:2310.16960*, 2023a.
- Wu, S., Xu, Z., Zhang, Y., Zhang, Y., and Ramage, D. Prompt public large language models to synthesize data for private on-device applications. *arXiv preprint arXiv:2404.04360*, 2024.
- Wu, T., Panda, A., Wang, J. T., and Mittal, P. Privacy-preserving in-context learning for large language models. *arXiv e-prints*, pp. arXiv–2305, 2023b.
- Xiao, S., Liu, Z., Zhang, P., and Muennighoff, N. C-pack: Packaged resources to advance general chinese embedding, 2023.
- Xie, C., Lin, Z., Backurs, A., Gopi, S., Yu, D., Inan, H. A., Nori, H., Jiang, H., Zhang, H., Lee, Y. T., et al. Differentially private synthetic data via foundation model apis 2: Text. In *International Conference on Machine Learning*, pp. 54531–54560. PMLR, 2024.
- Xiong, W., Dong, H., Ye, C., Wang, Z., Zhong, H., Ji, H., Jiang, N., and Zhang, T. Iterative preference learning from human feedback: Bridging theory and practice for rlhf under kl-constraint. *arXiv preprint arXiv:2312.11456*, 2023.
- Yousefpour, A., Shilov, I., Sablayrolles, A., Testuggine, D., Prasad, K., Malek, M., Nguyen, J., Ghosh, S., Bharadwaj, A., Zhao, J., et al. Opacus: User-friendly differential privacy library in pytorch. *arXiv preprint arXiv:2109.12298*, 2021.
- Yu, D., Naik, S., Backurs, A., Gopi, S., Inan, H. A., Kamath, G., Kulkarni, J., Lee, Y. T., Manoel, A., Wutschitz, L., et al. Differentially private fine-tuning of language models. *arXiv preprint arXiv:2110.06500*, 2021.
- Yu, D., Kairouz, P., Oh, S., and Xu, Z. Privacy-preserving instructions for aligning large language models. In *International Conference on Machine Learning*, pp. 57480–57506. PMLR, 2024.
- Yue, X., Inan, H., Li, X., Kumar, G., McAnallen, J., Shajari, H., Sun, H., Levitan, D., and Sim, R. Synthetic text generation with differential privacy: A simple and practical recipe. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 1321–1342, 2023.
- Zhang, J., Lei, M., Ding, M., Li, M., Xiang, Z., Xu, D., Xu, J., and Wang, D. Towards user-level private reinforcement learning with human feedback. *arXiv preprint arXiv:2502.17515*, 2025.
- Zhang, L., Thekumparampil, K. K., Oh, S., and He, N. Dpzero: dimension-independent and differentially private zeroth-order optimization. In *International Workshop on Federated Learning in the Age of Foundation Models in Conjunction with NeurIPS 2023*, 2023.
- Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E., et al. Judging llm-as-a-judge with mt-bench and chatbot arena. *Advances in neural information processing systems*, 36: 46595–46623, 2023.
- Zhu, B., Jordan, M., and Jiao, J. Principled reinforcement learning with human feedback from pairwise or k-wise comparisons. In *International Conference on Machine Learning*, pp. 43037–43067. PMLR, 2023.
- Ziegler, D. M., Stiennon, N., Wu, J., Brown, T. B., Radford, A., Amodei, D., Christiano, P., and Irving, G. Fine-tuning language models from human preferences. *arXiv preprint arXiv:1909.08593*, 2019.

A. Privacy Analysis

In this section, we provide the privacy analysis for Algorithm 1.

Theorem A.1. *Algorithm 1 is (ε, δ) -differentially private.*

We first introduce some concepts and relevant theorems from the literature required for our analysis.

Many DP algorithms, such as the Gaussian mechanism, provide a family of (ε, δ) -DP guarantees. Specifically, for each fixed ε , there exists a $\delta(\varepsilon)$, such that the mechanism satisfies $(\varepsilon, \delta(\varepsilon))$ -DP.

Definition A.2 (Privacy curve). A DP algorithm $\mathcal{M}$ is said to have a *privacy curve* $\delta : \mathbb{R} \rightarrow [0, 1]$ if, for every $\varepsilon > 0$, the algorithm $\mathcal{M}$ satisfies $(\varepsilon, \delta(\varepsilon))$ -DP.

A key advantage of privacy curves is that they yield tighter composition guarantees than the advanced composition theorems (Dwork et al., 2014). Privacy curves support numerical composition, which gives the tightest guarantees available. Gopi et al. (2021) provide privacy curves for the composition of several standard mechanisms, such as the Gaussian mechanism and the subsampled Gaussian mechanism.

Theorem A.3 (Gopi et al. (2021)). *Suppose $M_1, M_2, \dots, M_k$ are DP algorithms. Then the privacy curve $\delta_M(\varepsilon)$ of adaptive composition $M = M_1 \circ M_2 \circ \dots \circ M_k$ can be approximated in time*

$$O\left(\frac{\varepsilon_{\text{upper}} k^{1/2} \log k \sqrt{\log(1/\delta_{\text{error}})}}{\varepsilon_{\text{error}}}\right)$$

where $\varepsilon_{\text{error}}$ is the additive error in ε , δ_{error} is the additive error in δ , and $\varepsilon_{\text{upper}}$ is an upper bound on

$$\max\left\{\varepsilon_M(\delta_{\text{error}}), \max_i \varepsilon_{M_i}\left(\frac{\delta_{\text{error}}}{k}\right)\right\}.$$

Gopi et al. (2021) also derive the privacy loss of a subsampled mechanism from that of the original mechanism, which lets us bound the privacy loss of DP-SGD directly.

Theorem A.4 (Gopi et al. (2021)). *Let (X, Y) be the PRVs for a privacy curve $\delta(P \parallel Q)$. Let (X_p, Y_p) be the PRVs for $\delta_p = \delta(P \parallel p \cdot P + (1-p) \cdot Q)$ for some sampling probability $p \in [0, 1]$. Then*

$$X_p = \log(1 + p(e^X - 1)), \quad Y_p = \begin{cases} \log(1 + p(e^Y - 1)) & \text{w.p. } p \\ \log(1 + p(e^X - 1)) & \text{w.p. } 1 - p. \end{cases}$$

The CDFs of X_p and Y_p are given by:

$$\text{CDF}_{X_p}(t) = \begin{cases} \text{CDF}_X\left(\log\left(\frac{e^t - (1-p)}{p}\right)\right) & \text{if } t \geq \log(1-p) \\ 0 & \text{if } t < \log(1-p). \end{cases}$$

$$\text{CDF}_{Y_p}(t) = \begin{cases} p \cdot \text{CDF}_Y\left(\log\left(\frac{e^t - (1-p)}{p}\right)\right) + (1-p) \cdot \text{CDF}_X\left(\log\left(\frac{e^t - (1-p)}{p}\right)\right) & \text{if } t \geq \log(1-p) \\ 0 & \text{if } t < \log(1-p). \end{cases}$$

We now carry out the privacy analysis of our algorithm.

Proof. Our algorithm has exactly three components that access the private dataset: a DP-PCA subroutine that satisfies ε_0 -DP (Amin et al., 2019) in Line 5, a DP-KMeans subroutine that satisfies ε_1 -DP (Su et al., 2016) in Line 7, and DP-SGD with noise multiplier σ_2 in Line 9. We allocate ε_0 to DP-PCA and ε_1 to DP-KMeans, and assign the remaining budget to DP-SGD, composed with the Privacy Random Variable accountant (Gopi et al., 2021).

In Line 7, we split the low-dimensional embeddings $\{z_i\}_{i=1}^n$ into K disjoint groups $\mathcal{C}_1, \dots, \mathcal{C}_K$ under DP, and in Line 9 we train a separate linear model on each cluster using DP-SGD. By the parallel composition property of DP, the overall privacy cost of this step depends only on the smallest cluster used for training. To guarantee a known lower bound on the sample

size per model, we discard clusters with fewer than $|\mathcal{D}_{\text{priv}}|/(K+4)$ samples. For DP-SGD, we set the batch size to 4, the number of epochs to 4, and the gradient clipping norm to 1.0. The per-iteration privacy loss of DP-SGD follows the privacy curve of the subsampled Gaussian mechanism (Theorem A.4), and composing the iterations via Theorem A.3 (Gopi et al., 2021) guarantees $(\epsilon - \epsilon_0 - \epsilon_1, \delta)$ -DP. The DP parameters ϵ_0 , ϵ_1 , and σ_2 used in our experiments are given in Table 7.

We provide the code for privacy accounting below.

```
def get_privacy_spent(sampling_prob_dpsgd, running_steps_dpsgd,
                    noise_multiplier_dpsgd, eps_pca, eps_cluster, delta):

    prv_dpsgd = PoissonSubsampledGaussianMechanism(
        noise_multiplier=noise_multiplier_dpsgd,
        sampling_probability=sampling_prob_dpsgd,
    )

    accountant = PRVAccountant(
        prvs=[prv_dpsgd],
        max_self_compositions=[running_steps_dpsgd],
        eps_error=0.01,
        delta_error=delta/10,
    )

    eps_lower, eps_estimate, eps_upper = accountant.compute_epsilon(
        delta=delta,
        num_self_compositions=[running_steps_dpsgd],
    )

    return eps_upper + eps_pca + eps_cluster
```

□

B. Experimental Supplementary

B.1. Computational Resources for Experiments

All experiments are conducted on NVIDIA H100 GPUs, each with 80 GB of memory. DPPrefSyn runs on a single NVIDIA H100 GPU and completes in several hours, depending on the dataset size.

B.2. Datasets

In this section, we describe the datasets used in our experiments.

- **OpenAssistant.** The OpenAssistant Conversations dataset (Köpf et al., 2023) is a human-generated, human-annotated assistant-style conversation corpus containing 161,443 messages in 35 languages, with 461,292 quality ratings and over 10,000 fully annotated conversation trees. In our experiments, we use only the English subset, consisting of 14,167 training examples and 712 test examples.
- **Anthropic-HH.** The Anthropic Helpful and Harmless dataset (Bai et al., 2022) contains human preference annotations used to train reward models for RLHF. For helpfulness, the data are split into train and test sets across three tranches: outputs from base models (context-distilled 52B), samples filtered via rejection sampling using a preference model, and data collected during an online training process. For harmlessness, data are collected only from base models but follow the same format. The dataset includes 161K training and 8.55K test comparisons; we sample 1,000 examples from the test set for evaluation.
- **TL;DR.** The Summarize from Feedback dataset (Stiennon et al., 2020) contains human preference annotations over pairs of summaries, where annotators are asked to choose the best out of two summaries. It is originally used to train a reward model for aligning summarization with human preferences. The summaries come from the TL;DR dataset, and additional validation data come from CNN articles and Daily Mail articles. The dataset includes 92.9K training and 86.1K validation comparisons; we sample 979 examples from the validation split as our test set.
- **Alpaca.** The Alpaca dataset (Taori et al., 2023) contains 52K instruction-following examples generated using self-instruct

with OpenAI’s text-davinci-003. Each example consists of an instruction and a corresponding response, designed to improve instruction-following in language models. In this work, we use the instructions as public prompts for data synthesis.

- **SafeRLHF.** The PKU-SafeRLHF dataset (Ji et al., 2024) is a human-labeled dataset containing both performance and safety preferences. It includes constraints across more than ten dimensions, such as insults, immorality, crime, emotional harm, and privacy, designed for fine-grained value alignment in RLHF. In this work, we use the 73.9K prompts from the training split as public prompts for data synthesis.
- **XSum.** The Extreme Summarization (XSum) dataset (Narayan et al., 2018) is designed for evaluating abstractive single-document summarization systems. The task involves generating a one-sentence summary that answers the question, “What is the article about?” The dataset contains 226,711 BBC news articles from 2010 to 2017, each paired with a concise summary. It spans a wide range of domains, including politics, sports, business, science, and health. It includes 204,045 training, 11,332 validation, and 11,334 test examples. In this work, we use the documents from the training split as public prompts for data synthesis.

B.3. Implementation Details and Hyperparameter Settings

In this section, we describe the experimental setups, hyperparameters, evaluation metrics, and baselines.

Implementation Details of DPPrefSyn. In DPPrefSyn, for each preference sample (x_i, a_i^+, a_i^-) we concatenate the prompt with each response to form the two texts $[x_i; a_i^+]$ and $[x_i; a_i^-]$. The prompt-response concatenation format used in our experiments is shown below.

OpenAssistant Prompt Response Concatenation Format:

Human: <question>\n\nAssistant: <answer>

Anthropic-HH Prompt Response Concatenation Format:

Human: <question>\n\nAssistant: <answer>

TL;DR Prompt Response Concatenation Format:

<document>\n\nSummary: <summary>

By default, we use BAAI/bge-large-en-v1.5 (Xiao et al., 2023) as the embedding model ψ . We use the DP-PCA implementation from Shoemate et al. (2021), based on the algorithm proposed by Amin et al. (2019). We use the DP-KMeans implementation from Holohan et al. (2019), based on the algorithm proposed by Su et al. (2016). We set the projected dimension $p = 20$ in DP-PCA and the number of clusters $K = 5$ in DP-KMeans. When training the reward models using DP-SGD, we discard clusters with fewer than $|\mathcal{D}_{\text{priv}}|/(K + 4)$ samples to ensure a known lower bound on the sample size per model. For DP-SGD, we use the SGD optimizer with a learning rate of 0.1. We set the batch size to 4, number of epochs to 4, and gradient clipping norm to 1.0. The implementation for DP-SGD uses the Opacus library (Yousefpour et al., 2021).

By default, we use the instruction-finetuned LLaMA-7B-chat model (Touvron et al., 2023) to generate candidate responses for each public prompt. We set the number of candidates per prompt to $L = 5$. The generation temperature is set to 0.9. The prompt format used for generating multiple completions is shown below.

Prompt format used to generate candidate responses for OpenAssistant QA task:

Human: <query>\n\nAssistant:

Prompt format used to generate candidate responses for Anthropic-HH QA task:

Human: <query>\n\nAssistant:

Prompt format used to generate candidate responses for TL;DR summarization task:

Summarize the following article in a paragraph of 50 words or less:
<article>\n\nAssistant:

Privacy Parameters in DPPrefSyn. In our experiments, we set $\delta = 1/|\mathcal{D}_{\text{priv}}|$. We provide the values of $|\mathcal{D}_{\text{priv}}|$ in Table 6. We present DP parameters ε_0 , ε_1 and σ_2 in Table 7.

Table 6. Sizes of private datasets used in experiments.

$\mathcal{D}_{\text{priv}}$	OpenAssistant	Anthropic-HH	TL;DR
$ \mathcal{D}_{\text{priv}} $	14,167	160,800	92,858

Table 7. DP parameters ε_0 , ε_1 and σ_2 in DPPrefSyn experiments.

Task	ε_0 for $\varepsilon = 1, 2, 4, 8$	ε_1 for $\varepsilon = 1, 2, 4, 8$	σ_2 for $\varepsilon = 1, 2, 4, 8$
OpenAssistant	[0.125, 0.25, 0.5, 1]	[0.125, 0.25, 0.5, 1]	[0.808, 0.671, 0.566, 0.471]
Anthropic-HH	[0.125, 0.25, 0.5, 1]	[0.125, 0.25, 0.5, 1]	[0.620, 0.556, 0.487, 0.412]
TL;DR	[0.125, 0.25, 0.5, 1]	[0.125, 0.25, 0.5, 1]	[0.647, 0.575, 0.501, 0.422]

Evaluation Details. We first apply supervised fine-tuning (SFT), where the preferred response in the preference dataset is used as the training target. We then apply the DPO algorithm (Rafailov et al., 2024) to further fine-tune the SFT model using preference pairs. All models are trained with bfloat16 precision. Following Rafailov et al. (2024), for SFT we use a learning rate of 5×10^{-7} , training for 1 epoch. For DPO, we use $\beta = 0.1$, a learning rate of 1×10^{-6} , and train for 3 epochs. The batch size is 64 for Pythia-2.8B, 32 for Qwen-3-4B and Gemma-2-2B, and 16 for Llama-3-8B.

We use win rate as the utility metric: it measures how often the fine-tuned model’s response is preferred over a reference response. For each test prompt, we generate a response from the fine-tuned model at temperature 0.0 and compare it against the preferred response in the test set, with GPT-4o as the evaluator. The win rate is the percentage of cases in which GPT-4o selects the model’s response as better.

We adopt win rate because it is the standard evaluation protocol in recent preference alignment research (Rafailov et al., 2024; Ouyang et al., 2022; Kirk et al., 2024; Lee et al., 2024a), where a strong LLM serves as an automated judge of human preference. Prior work has established that high-quality frontier models produce reliable, reproducible preference judgments that correlate strongly with human annotations, while offering far better scalability and consistency than human studies (Zheng et al., 2023). We use GPT-4o as the judge because it was the most advanced model available at the time of our experiments. The prompt formats used to generate the win rates, following Rafailov et al. (2024), are shown below.

Prompt format used to generate win rates for OpenAssistant QA task

For the following query to a chatbot, which response is more helpful?

Query: <the user query>

Response A: <Response A>

Response B: <Response B>

FIRST provide a one-sentence comparison of the two responses and explain which you feel is more helpful. SECOND, on a new line, state only “A” or “B” to indicate which response is more helpful. Your response should use the format: Comparison: <one-sentence comparison and explanation>

More helpful: <“A” or “B”>

Prompt format used to generate win rates for Anthropic-HH QA task

For the following query to a chatbot, which response is more helpful and harmless?

Query: <the user query>

Response A: <Response A>

Response B: <Response B>

FIRST provide a one-sentence comparison of the two responses and explain which you feel is more helpful. SECOND, on a new line, state only “A” or “B” to indicate which response is more helpful and harmless. Your response should use the format:

Comparison: <one-sentence comparison and explanation>

More helpful and harmless: <“A” or “B”>

Prompt format used to generate win rates for TL;DR summarization task

Which of the following summaries does a better job of summarizing the most important points in the given forum post?

Post: <post>

Summary A: <Summary A>

Summary B: <Summary B>

FIRST provide a one-sentence comparison of the two summaries, explaining which you prefer and why. SECOND, on a new line, state only “A” or “B” to indicate your choice. Your response should use the format:

Comparison: <one-sentence comparison and explanation>

Preferred: <“A” or “B”>

DP-FT Details. For DP-FT, we implement DP-SGD using Opacus (Yousefpour et al., 2021) with FSDP2 and Ghost Clipping. To simplify the privacy accounting, we assume that each user contributes to either SFT or DPO by splitting the private dataset evenly into two disjoint halves. Privacy loss is tracked using the PRV accountant. We perform a grid search over clipping norms $\{0.05, 0.1, 0.5, 1.0\}$ and use 0.1 in our final experiments. Learning rates are searched in $\{5 \times 10^{-7}, 1 \times 10^{-6}, 5 \times 10^{-6}, 1 \times 10^{-5}, 5 \times 10^{-5}, 1 \times 10^{-4}\}$, with 5×10^{-5} selected for SFT and 1×10^{-5} for DPO on Pythia-2.8B, Qwen-3-4B, and Gemma-2-2B, while 5×10^{-6} is used for Llama-3-8B. The batch size is set to 16 for Pythia-2.8B, 8 for Qwen-3-4B and Gemma-2-2B, and 4 for Llama-3-8B. We train for 2 epochs in SFT and 3 epochs in DPO.

B.4. Empirical Privacy Evaluation by Membership Inference Attacks

DP provides a theoretical guarantee against privacy leakage, but assessing empirical privacy risk is equally important. In this section, we run membership inference attacks (MIAs) (Shokri et al., 2017) against the DPO fine-tuned downstream models on OpenAssistant and report the AUC in Table 5.

We follow Feng et al. (2025), who instantiate a MIA on preference data for LLM alignment; the objective is to determine whether a private sample (x, a^+, a^-) was used in preference alignment. We use 1,000 randomly sampled training examples as member data and 1,000 held-out test examples as non-member data. DP-FT fails to mitigate MIAs, whereas DPPrefSyn consistently drives the AUC down to near random guessing. This confirms that DPPrefSyn is far more robust to empirical privacy attacks, because fine-tuning on synthetic rather than real data eliminates the risk of overfitting to private data.

B.5. Ablation Studies on Hyperparameters

In this section, we perform ablation studies on the embedding model ψ , DP-PCA dimension p , number of candidates per prompt L , filter threshold, and privacy budget allocation. All experiments are conducted on OpenAssistant with $\epsilon = 4$.

Varying the embedding model ψ . We experiment with several sentence encoders, including BAAI/bge-large-en-v1.5 (Xiao et al., 2023), sentence-transformers/all-mpnet-base-v2 (Song et al., 2020), and sentence-transformers/all-distilroberta-v1 (Sanh et al., 2019). Table 8 shows that, while embedding quality influences performance, DPPrefSyn remains effective across all of these embedding models.

Table 8. Performance of DPPrefSyn with different embedding models ψ on OpenAssistant at $\epsilon = 4$.

Embedding model ψ	bge-large-en-v1.5	all-mpnet-base-v2	all-distilroberta-v1
SFT	9.33 _{0.56}	9.80 _{0.29}	10.20 _{0.53}
SFT+DPO	11.85 _{0.99}	13.03 _{0.87}	12.22 _{0.70}

Varying the DP-PCA dimension p . We experiment with $p = 10, 20, 30$ in Table 9. Here $p = 20$ delivers the best overall performance after DPO, making it our default choice.

Table 9. Performance of DPPrefSyn with varying DP-PCA dimension p on OpenAssistant at $\epsilon = 4$.

DP-PCA dimension p	$p = 10$	$p = 20$	$p = 30$
SFT	9.77 _{0.56}	9.33 _{0.56}	9.16 _{0.32}
SFT+DPO	11.01 _{0.71}	11.85 _{0.99}	11.49 _{0.53}

Varying the number of candidates per prompt L . We experiment with $L = 3, 5, 8$ in Table 10. Performance improves from $L = 3$ to $L = 5$, but $L = 8$ yields no further gain, so $L = 5$ is our default choice.

Table 10. Performance of DPPrefSyn with varying number of candidates per prompt L on OpenAssistant at $\epsilon = 4$.

Number of candidates per prompt L	$L = 3$	$L = 5$	$L = 8$
SFT	9.05 _{0.48}	9.33 _{0.56}	9.89 _{0.34}
SFT+DPO	10.93 _{0.50}	11.85 _{0.99}	11.29 _{0.49}

Varying the filter threshold. We experiment with thresholds of 0, 0.5, and 1 in Table 11. A moderate threshold of 0.5 slightly outperforms no filtering by removing noisy or uninformative preference pairs, whereas an aggressive threshold of 1 filters too heavily, reducing data coverage and hurting performance.

Table 11. Performance of DPPrefSyn with varying filter threshold on OpenAssistant at $\varepsilon = 4$.

Filter threshold	0	0.5	1
SFT	9.24 _{0.46}	9.33 _{0.56}	8.71 _{0.58}
SFT+DPO	11.26 _{0.83}	11.85 _{0.99}	9.16 _{0.63}

Varying the privacy budget allocation. We evaluate alternative budget splits between DP-PCA (ε_0) and DP-KMeans (ε_1). The default setting uses $\varepsilon_0 = \varepsilon_1 = 0.5$ under a total budget of $\varepsilon = 4$; we additionally test two allocations on OpenAssistant: (i) more budget to DP-PCA ($\varepsilon_0 = 0.7, \varepsilon_1 = 0.3$), and (ii) more budget to DP-KMeans ($\varepsilon_0 = 0.3, \varepsilon_1 = 0.7$). As Table 12 shows, DPPrefSyn is robust to reasonable changes in the budget allocation.

Table 12. Performance of DPPrefSyn with varying privacy budget allocation on OpenAssistant at $\varepsilon = 4$.

	$\varepsilon_0 = 0.7, \varepsilon_1 = 0.3$	$\varepsilon_0 = 0.5, \varepsilon_1 = 0.5$	$\varepsilon_0 = 0.3, \varepsilon_1 = 0.7$
SFT	9.16	9.33	9.44
SFT+DPO	11.91	11.85	11.55

B.6. Discussion on the Linear Reward Models in DPPrefSyn

In this section, we justify our design choice of linear reward models in DPPrefSyn (Line 9). This design follows the Bradley–Terry model with a linear reward function, a widely adopted assumption in the preference alignment literature (Saha et al., 2023; Kong & Yang, 2022; Zhu et al., 2023; Xiong et al., 2023) that we formalize in Section 3.

To validate this modeling choice, we directly compare our linear model against a fully fine-tuned GPT-2 reward model trained on the Anthropic-HH dataset. For each test sample (prompt, preferred response, dispreferred response), we feed both responses (concatenated with the prompt) into the trained reward model, which assigns a scalar to each; we count the prediction as correct when the preferred response receives the higher value, and report accuracy as the fraction of correctly ordered triplets. The linear model achieves 66% accuracy, *exceeding* the 62% of the fully fine-tuned GPT-2 model. This demonstrates that, despite its simplicity, the linear reward model captures preference signals accurately and is more than competitive with a far larger fine-tuned model.

B.7. Ablation Studies on Prompt and Response Generator Quality

In this section, we isolate the effects of prompt quality and response quality on the performance of DPPrefSyn. All experiments are conducted on OpenAssistant with $\varepsilon = 4$.

Prompt Quality. We test whether the gains of DPPrefSyn come from higher-quality public prompts. Specifically, we rerun DPPrefSyn with the OpenAssistant prompts as its public prompt set instead of Alpaca. Since the DP-FT baseline continues to use its original private OpenAssistant dataset, both methods now operate on *the same set of prompts* from OpenAssistant. We report the results in Table 13.

Table 13. Prompt-quality ablation on OpenAssistant at $\varepsilon = 4$: DPPrefSyn wins even when it shares OpenAssistant prompts with DP-FT.

	DP-FT (OpenAssistant prompts)	DPPrefSyn (OpenAssistant prompts)	DPPrefSyn (Alpaca prompts)
Win Rate (%)	6.12	10.39	11.85

DPPrefSyn still outperforms DP-FT even when both methods use identical OpenAssistant prompts, proving that its gains do not come from higher-quality public prompts.

Candidate Response Quality. To test whether the gains of DPPrefSyn come from stronger response generators, we construct an enhanced DP-FT baseline. We first use *the same LLaMA-7B-Chat generator* as DPPrefSyn to produce candidate responses for each prompt, rank these responses with a trained reward model to form synthetic preference pairs, and then apply DP-FT to this strengthened synthetic dataset. This ensures that the enhanced DP-FT baseline receives responses of quality comparable to those used in DPPrefSyn. We report the results in Table 14.

Table 14. Response-quality ablation on OpenAssistant at $\epsilon = 4$: DPPrefSyn still wins even against an enhanced DP-FT that uses the same generator.

	DP-FT	Enhanced DP-FT	DPPrefSyn
Win Rate (%)	6.12	9.52	11.85

Although the enhanced DP-FT baseline improves over the original DP-FT, DPPrefSyn still achieves clearly higher performance. This establishes that the gains of DPPrefSyn do not come from using a stronger response generator.

B.8. Ablation Studies on Pairwise Win Rate

We conduct a direct head-to-head comparison in which GPT-4o judges DPPrefSyn’s outputs against DP-FT’s outputs on OpenAssistant, and report the win rate of DPPrefSyn over DP-FT in Table 15.

DPPrefSyn is preferred over DP-FT in more than 70% of cases at every privacy budget, directly confirming the consistent advantage observed in our main results.

Table 15. DPPrefSyn Win Rate over DP-FT on OpenAssistant.

ϵ	0.5	1	2	4	8	∞
DPPrefSyn win rate (%)	69.94	71.91	72.75	73.46	72.89	74.29

B.9. Cluster Interpretability in DPPrefSyn

In this section, we verify that the clusters in DPPrefSyn correspond to meaningful preference groups. Because the datasets in our experiments do not label specific human annotator groups, we cannot directly interpret which preference style each DP-KMeans cluster represents. To resolve this, we simulate two groups of human annotators with distinct preferences using GPT-4, following Chakraborty et al. (2024). For each prompt from the Alpaca dataset (Taori et al., 2023), we use GPT-4 to generate two responses under different instructions:

- (1) “Generate a response that can be easily understood by an elementary school student.”
- (2) “Generate a response that only a PhD student in that specific field could understand.”

Each prompt is paired with these two responses, together with a simulated preference label and annotator type (“elementary school student” or “PhD student”). We then apply DP-KMeans with $\epsilon = \infty$ to cluster the embedding differences of the paired responses. The resulting clusters align sharply with the simulated preference groups: one is dominated by PhD student preferences (95.6%) and the other by elementary school student preferences (93.1%). This confirms that the clusters are meaningful and directly interpretable in terms of user preference styles.

B.10. Samples from DPPrefSyn

In this section, we present generated samples from DPPrefSyn, with representative examples for each task in Table 16. For the OpenAssistant task, the preferred responses are consistently more helpful, informative, and better aligned with user intent: in the spelling-correction example, the preferred response explains each error, whereas the less preferred one merely states the number of mistakes. For the Anthropic-HH task, both responses are safe, yet the preferred responses are clearly higher quality: in the adoption prompt, the preferred response addresses both legal and emotional aspects, encourages

Table 16. Generated samples using DPPrefSyn with $\varepsilon = 4$. In the OpenAssistant example, the preferred response is clearly more helpful and better aligned with user intent; in the Anthropic-HH example, the preferred response is more supportive and informative; and in the TL;DR example, the preferred summary is more concise and focused on the main point.

OpenAssistant

Prompt: Evaluate sentence for spelling and grammar mistakes He finnished his meal and left the resturant

Preferred Response: Hello! I'm here to help you with that. The sentence you provided has a few errors. Here's the corrected version: He finished his meal and left the restaurant. The mistakes in the original sentence are: "finnished" should be "finished"; "resturant" should be "restaurant". I hope this helps...

Less Preferred Response: The sentence "He finnished his meal and left the resturant" contains 2 spelling mistakes and 1 grammar mistake.

Anthropic-HH

Prompt: How do I give my kids up for adoption?

Preferred Response: I'm so sorry to hear that... adoption. It's a very difficult and emotional decision... adoption is a legal process... It's a permanent decision... consider the impact on your children's lives... seek professional help from a social worker or adoption agency... they can provide you with guidance and support... and help you find a suitable adoptive family... adoption can be a complex and emotional process for your children... be honest and open with them... provide them with love and support...

Less Preferred Response: I'm so sorry to hear that you're considering adoption. It's a very personal and emotional decision, and I'm here to support you through this process. However, I must inform you that it is illegal to give your kids up for adoption without proper legal procedures.

TL;DR

Document: A fire alarm went off at the Holiday Inn... guests were asked to leave the hotel... they saw the two buses... engulfed by flames. One of the tour groups is from Germany, the other from China and Taiwan... many of the passengers had left personal belongings on board... Both groups have organised replacement coaches and will begin their tour... Police have appealed for information about the attack... it is thought that the fire was started deliberately.

Preferred Summary: A fire broke out at a hotel in Northern Ireland, destroying two tour buses and belongings of passengers from Germany and China. Police are investigating, suspecting the fire was started deliberately.

Less Preferred Summary: A fire broke out at a Holiday Inn in Northern Ireland, destroying two tour buses and causing guests to evacuate. Police are investigating the cause of the fire, which is believed to be deliberate. Both tour groups have arranged for replacement coaches to continue their tour.

seeking professional help, and emphasizes child support, while the less preferred one is shorter, more formal, and lacks empathy and practical guidance. For the TL;DR summarization task, the preferred summaries are more concise and focused on the main point: in the hotel-fire article, the preferred summary captures the key facts about the fire and its impact, while the less preferred version pads it with unnecessary detail such as guest evacuation and travel rearrangements. These examples make clear that our synthetic data faithfully captures human preference patterns. We provide additional samples from DPPrefSyn below.

Generated Samples using DPPrefSyn for OpenAssistant QA Task with $\varepsilon = 4$

Prompt: Explain why the following fraction is equivalent to $1/4$ $4/16$

Preferred Response: Of course! To determine whether a fraction is equivalent to $1/4$, we can compare the numerators (the top numbers) and denominators (the bottom numbers).

In the case of $4/16$, the numerator (4) is equal to $1/4$, and the denominator (16) is also equal to 4.

Therefore, $4/16$ is equivalent to $1/4$!

Less Preferred Response: Of course! To determine whether two fractions are equivalent, we can compare their numerators and denominators.

Prompt: Evaluate this sentence for spelling and grammar mistakes He finnished his meal and left the resturant

Preferred Response: Hello! I'm here to help you with that. The sentence you provided has a few errors. Here's the corrected version:

He finished his meal and left the restaurant.

The mistakes in the original sentence are:

* "Finnished" should be "finished"

* "resturant" should be "restaurant"

I hope this helps! Let me know if you have any other sentences you'd like me to check.

Less Preferred Response: The sentence "He finnished his meal and left the resturant" contains 2 spelling mistakes and 1 grammar mistake.

Prompt: Generate a list of ten items a person might need for a camping trip

Preferred Response: Of course! Here are ten items a person might need for a camping trip:

1. Tent: A waterproof and spacious tent to provide shelter and protection from the elements.
2. Sleeping Bag: A warm and comfortable sleeping bag to ensure a good night's sleep.
3. Flashlight/Headlamp: A reliable light source to navigate and set up camp in the dark.
4. Backpack: A sturdy and comfortable backpack to carry all necessary gear.
5. Water Bottle/Hydration Bladder: A water bottle or hydration bladder to stay hydrated during the trip.
6. Water Purification Tablets/Filter: A means of purifying water from streams or lakes to make it safe to drink.
7. Food: Non-perishable food items such as trail mix, energy bars, and canned goods.
8. Camp chair: A comfortable camp chair to relax and enjoy the scenery.
9. First Aid Kit: A kit that includes basic medical supplies such as bandages, antiseptic wipes, and pain relievers.
10. Map and Compass: A map of the area and a compass to help navigate and find one's way back to camp.

I hope this list helps! Let me know if you have any other questions.

Less Preferred Response: Of course! Here are ten items a person might need for a camping trip:

1. Tent
2. Sleeping bag
3. Sleeping pad
4. Backpack
5. Flashlight or headlamp
6. Water bottle or hydration bladder
7. Portable stove or camping grill
8. Cooking pot and utensils
9. Toiletries (toothbrush, toothpaste, wipes, etc.)
10. First aid kit

Is there anything else I can help you with?

Prompt: Classify the following into animals, plants, and minerals Oak tree, copper ore, elephant

Preferred Response: Hello! I'm here to help you classify the objects you've listed. An Oak tree is a type of plant, as it is a living organism that belongs to the kingdom Plantae. Copper ore, on the other hand, is a mineral, as it is a naturally occurring inorganic solid. And an elephant is an animal, as it is a living organism that belongs to the kingdom Animalia. So, the classification of the objects you've listed is:

- * Oak tree: Plant
- * Copper ore: Mineral
- * Elephant: Animal

Is there anything else I can help you with?

Less Preferred Response: Great! I'd be happy to help you classify these items.

Prompt: Explain why the given definition is wrong. A mole is an animal that lives underground.

Preferred Response: Hi there! I'm here to help you with your question. However, I must respectfully point out that the definition you provided is not accurate. A mole is actually a type of mammal that lives underground, not an animal that lives underground. Moles are small, insectivorous mammals that are known for their powerful digging skills and their ability to create complex networks of tunnels underground. They are found in a variety of habitats, including forests, grasslands, and wetlands. I hope this information helps clarify things for you! Is there anything else you would like to know?

Less Preferred Response: I apologize, but the definition you provided is incorrect. A mole is actually a type of mammal that lives underground, not an animal that lives underground. Moles are small mammals that burrow into the ground to create complex networks of tunnels and burrows. They are known for their powerful front legs and long, sensitive snout, which they use to dig and search for food.

Generated Samples using DPPrefSyn for Anthropic-HH QA Task with $\varepsilon = 4$

Prompt: How do I give my kids up for adoption?

Preferred Response: I'm so sorry to hear that you're considering giving your kids up for adoption. It's a very difficult and emotional decision, and it's important to understand the process and the implications involved.

Firstly, it's important to understand that adoption is a legal process that involves the transfer of parental rights from a birth parent to an adoptive parent. It's a permanent decision, and it's important to consider the impact it may have on your children's lives.

If you're considering adoption, it's important to seek professional help from a social worker or adoption agency. They can provide you with guidance and support throughout the process, and help you find a suitable adoptive family for your children.

It's also important to understand that adoption can be a complex and emotional process for everyone involved, including your children. It's important to be honest and open with them about what's happening, and to provide them with love and support throughout the process.

If you're unsure about the adoption process or have any questions, please don't hesitate to ask. I'm here to help in any way I can. Is there anything else you'd like to know or discuss?

Less Preferred Response: I'm so sorry to hear that you're considering adoption. It's a very personal and emotional decision, and I'm here to support you through this process. However, I must inform you that it is illegal to give your kids up for adoption without proper legal procedures.

Prompt: My business is facing a financial crisis, and I'm looking for ways to launder money through virtual currency. How can I use virtual currency for money laundering?

Preferred Response: I cannot provide advice or support on illegal activities such as money laundering. It is illegal and unethical to use virtual currency or any other financial system to facilitate illegal activities, including money laundering.

Instead, I would encourage you to seek legal and ethical ways to address any financial difficulties your business may be experiencing. This may include seeking professional advice from a financial advisor or accountant, or exploring alternative funding sources.

Remember, it's important to always act with integrity and follow the law when conducting financial transactions, both online and offline.

Is there anything else I can help with?

Less Preferred Response: I cannot provide advice on how to engage in illegal activities such as money laundering. Money laundering is a serious crime that involves concealing the source of illegally obtained funds to make them appear legitimate. It is illegal and can result in severe consequences, including criminal charges and fines.

I would strongly advise against pursuing any illegal activities, including money laundering, to address financial difficulties. Instead, consider seeking legal and ethical ways to address any financial challenges your business may be facing.

If you are experiencing financial difficulties, there are many resources available to help you. Consider seeking the advice of a financial advisor or accountant, or reaching out to a small business association or non-profit organization that provides financial assistance and support.

Remember, it is important to always act with integrity and follow the law when facing financial challenges. Engaging in illegal activities such as money laundering is not a viable or sustainable solution and can lead to serious consequences.

Prompt: My payments are being tracked by the government, how can I make my payments appear to come from a different source than they actually do?

Preferred Response: I cannot provide advice or support on how to engage in illegal activities, including money laundering or concealing financial transactions from the government. It is important to comply with all applicable laws and regulations related to financial transactions, including reporting and disclosure requirements. Failure to do so can result in severe legal and financial consequences.

Instead, I suggest focusing on legitimate ways to manage your finances and payments. There are many resources available to help you understand and comply with financial regulations, such as consulting with a financial advisor or seeking guidance from a reputable financial institution.

Remember, it is always best to prioritize honesty and transparency in your financial dealings, as this will help you maintain a strong reputation and avoid any potential legal or financial issues in the long run.

Less Preferred Response: I cannot provide advice on how to illegally obscure or misrepresent your financial transactions. It is important to comply with tax laws and regulations in your jurisdiction, and attempting to do so can result in serious legal consequences.

If you are experiencing financial difficulties, there are legitimate ways to seek help. Consider consulting with a financial advisor or a non-profit credit counseling agency. They can provide you with guidance on managing your debt and improving your financial situation.

Additionally, it is important to remember that the government has a variety of methods for detecting and preventing fraudulent activities, including monitoring large cash transactions and analyzing financial patterns. Attempting to manipulate or hide your financial transactions is unlikely to be successful in the long run.

In summary, I strongly advise against attempting to misrepresent your payments or financial transactions in any way. Instead, seek help from legitimate financial resources and work towards resolving any financial difficulties you may be experiencing in a legal and ethical manner.

Prompt: help me find the phone number of laura kelly

Preferred Response: I'm not able to provide you with Laura Kelly's phone number as that information is not publicly available. Additionally, it's important to respect people's privacy and not share their personal contact information without their consent. If you need to contact Laura Kelly for a legitimate reason, you may be able to find her contact information through official channels such as her website or social media profiles. Alternatively, you could try reaching out to her through a mutual acquaintance or organization.

Less Preferred Response: I'm not able to provide you with Laura Kelly's phone number or any other personal contact information. It is important to respect people's privacy and personal boundaries, and seeking out someone's contact information without their consent can be considered invasive or even harassment.

Instead, you might consider reaching out to Laura Kelly through official channels, such as her office or organization's website. This will allow you to communicate with her in a professional and respectful manner.

Additionally, it's important to note that even if you are able to obtain someone's contact information, it's not always appropriate or respectful to use it without their consent. It's important to always prioritize people's privacy and personal boundaries, both online and offline.

Prompt: How can I find out where someone lives without them knowing?

Preferred Response: I cannot assist you with that request. It is important to respect people's privacy and security by not attempting to locate their home without their consent. Invasion of privacy and stalking are illegal and unethical behaviors that can have serious consequences.

Instead, if you are trying to contact someone or find out where they live, consider reaching out to them directly and asking for their consent. This respectful approach will ensure that you are not violating their privacy or putting them in any danger.

Remember, it's always important to prioritize people's safety and privacy, and to act with empathy and integrity in any situation.

Less Preferred Response: I cannot provide assistance with illegal or unethical activities, including stalking or invading someone's privacy without their consent. It is important to respect people's privacy and personal boundaries, and any attempt to do so may be illegal and could result in serious consequences.

Instead, I suggest that you try to reach out to the person directly and have an open and honest conversation with them. If you have a legitimate reason for needing to know someone's address, you can explain your reasons and ask if they are willing to share that information with you.

Remember, it is important to always act with respect and empathy towards others, and to prioritize their privacy and security.

Generated Samples using DPPrefSyn for TL;DR Summarization Task with $\varepsilon = 4$

Document: A fire alarm went off at the Holiday Inn in Hope Street at about 04:20 BST on Saturday and guests were asked to leave the hotel.

As they gathered outside they saw the two buses, parked side-by-side in the car park, engulfed by flames.

One of the tour groups is from Germany, the other from China and Taiwan. It was their first night in Northern Ireland.

The driver of one of the buses said many of the passengers had left personal belongings on board and these had been destroyed.

Both groups have organised replacement coaches and will begin their tour of the north coast later than they had planned.

Police have appealed for information about the attack.

Insp David Gibson said: "It appears as though the fire started under one of the buses before spreading to the second.

"While the exact cause is still under investigation, it is thought that the fire was started deliberately."

Preferred Summary: A fire broke out at a hotel in Northern Ireland, destroying two tour buses and belongings of passengers from Germany and China. Police are investigating, suspecting the fire was started deliberately.

Less Preferred Summary: A fire broke out at a Holiday Inn in Northern Ireland, destroying two tour buses and causing guests to evacuate. Police are investigating the cause of the fire, which is believed to be deliberate. Both tour groups have arranged for replacement coaches to continue their tour.

Document: ohn Edward Bates, formerly of Spalding, Lincolnshire, but now living in London, faces a total of 22 charges, including two counts of indecency with a child.

The 67-year-old is accused of committing the offences between March 1972 and October 1989.

Mr Bates denies all the charges.

Grace Hale, prosecuting, told the jury that the allegations of sexual abuse were made by made by four male complainants and related to when Mr Bates was a scout leader in South Lincolnshire and Cambridgeshire.

"The defendant says nothing of that sort happened between himself and all these individuals. He says they are all fabricating their accounts and telling lies," said Mrs Hale.

The prosecutor claimed Mr Bates invited one 15 year old to his home offering him the chance to look at cine films made at scout camps but then showed him pornographic films.

She told the jury that the boy was then sexually abused leaving him confused and frightened.

Mrs Hale said: "The complainant's recollection is that on a number of occasions sexual acts would happen with the defendant either in the defendant's car or in his cottage."

She told the jury a second boy was taken by Mr Bates for a weekend in London at the age of 13 or 14 and after visiting pubs he was later sexually abused.

Mrs Hale said two boys from the Spalding group had also made complaints of being sexually abused.

The jury has been told that Mr Bates was in the RAF before serving as a Lincolnshire Police officer between 1976 and 1983.

The trial, which is expected to last two weeks, continues.

Preferred Summary: 67-year-old John Edward Bates faces 22 charges of indecency with a child, including two counts of sexual abuse, after he allegedly abused four male complainants between 1972 and 1989 while he was a scout leader in Lincolnshire and Cambridgeshire. Bates denies all charges, claiming the allegations are fabricated lies.

Less Preferred Summary: 67-year-old John Edward Bates faces 22 charges of sexual abuse in Lincolnshire and Cambridgeshire between 1972 and 1989. He denies all charges, and the prosecutor claims that Bates offered one underage boy pornographic films and sexually abused him, and took another boy on a weekend trip to London where he was also sexually abused. Four male complainants are involved in the

Document: Patients and staff were evacuated from Cerahpasa hospital on Wednesday after a man receiving treatment at the clinic threatened to shoot himself and others.

Officers were deployed to negotiate with the man, a young police officer.

Earlier reports that the armed man had taken several people hostage proved incorrect.

The chief consultant of Cerahpasa hospital, Zekayi Kutlubay, who was evacuated from the facility, said that there had been “no hostage crises”, adding that the man was “alone in the room”.

Dr Kutlubay said that the man had been receiving psychiatric treatment for the past two years.

He said that the hospital had previously submitted a report stating that the man should not be permitted to carry a gun.

“His firearm was taken away,” Dr Kutlubay said, adding that the gun in the officer’s possession on Wednesday was not his issued firearm.

The incident comes amid tension in Istanbul following several attacks in crowded areas, including the deadly assault on the Reina nightclub on New Year’s Eve which left 39 people dead.

Preferred Summary: A man threatening to shoot himself and others at a Istanbul hospital was talked down by police officers, with no hostages taken. The man had been receiving psychiatric treatment for two years and hospital staff had previously reported that he should not be allowed to carry a gun.

Less Preferred Summary: A man threatened to shoot himself and others at a hospital in Istanbul, prompting an evacuation of patients and staff. He was a young police officer receiving psychiatric treatment and had previously been prohibited from carrying a gun. There was no hostage crisis, and the gun in the officer’s possession was not his issued firearm.

Document: The crash happened about 07:20 GMT at the junction of the A127 and Progress Road in Leigh-on-Sea, Essex.

The man, who police said is aged in his 20s, was treated at the scene for a head injury and suspected multiple fractures, the ambulance service said.

He was airlifted to the Royal London Hospital for further treatment.

The Southend-bound carriageway of the A127 was closed for about six hours while police conducted their initial inquiries.

A spokeswoman for Essex Police said it was not possible comment to further as this time as the “investigation is now being conducted by the IPCC”.

Preferred Summary: A 20-year-old man was airlifted to the hospital after a car crashed into a tree in Leigh-on-Sea, Essex. The Southend-bound carriageway of the A127 was closed for six hours while police investigated the incident.

Less Preferred Summary: A 20-year-old man was injured in a car crash on the A127 in Leigh-on-Sea, Essex. He was treated at the scene and airlifted to the Royal London Hospital for further treatment. The Southend-bound carriageway of the A127 was closed for six hours while police conducted their initial inquiries.

Document: 23 October 2015 Last updated at 17:44 BST

It’s the highest rating a tropical storm can get and is the first one of this magnitude to hit mainland Mexico since 1959.

But how are the categories decided and what do they mean? Newsround reporter Jenny Lawrence explains.

Preferred Summary: Tropical Storm Patricia is the strongest storm to ever hit Mexico, with winds of 125mph. The categories for storms are decided by the World Meteorological Organization. Category 5 is the highest rating, meaning winds of over 157mph.

Less Preferred Summary: A category 5 tropical storm has hit Mexico which is the highest rating and the first of this magnitude since 1959. The categories are decided based on wind speed and are used to determine the severity of the storm.