

Differentially Private Preference Data Synthesis for Large Language Model Alignment

Anonymous Authors¹

Abstract

Preference alignment constitutes a critical post-training stage for reconciling the outputs of large language models (LLMs) with human values. Aligning models on authentic human preference data, however, poses acute privacy risks, because such datasets routinely embed sensitive user prompts together with the human judgments rendered over them—information that can expose personal attributes, private beliefs, and behavioral patterns. We address this challenge with **DPPrefSyn**, an algorithm that synthesizes differentially private (DP) preference data and thereby enables privacy-preserving preference alignment. DPPrefSyn is a principled framework grounded in the Bradley–Terry preference model and in the intrinsic geometric structure of pairwise human preferences: it first learns the underlying preference model from private data under formal differential-privacy guarantees, and then couples the learned model with public prompts to synthesize high-quality preference data. By exploiting the shared linear structure of per-cluster reward models, it faithfully captures the heterogeneous human preferences present in private datasets, and by employing DP Principal Component Analysis (DP-PCA) it sharpens learning accuracy under the limited samples available within each cluster. Extensive experiments spanning question answering and summarization benchmarks, several downstream model families, and multiple alignment methods demonstrate that DPPrefSyn attains strong alignment performance under stringent DP guarantees—consistently surpassing DP fine-tuning on real data and, in many settings, even exceeding non-private fine-tuning. These results establish synthetic preference data as a practical and broadly applicable foundation for

privacy-preserving preference alignment. To the best of our knowledge, this is the first work to generate DP synthetic preference data for LLM alignment. Our code will be released upon acceptance.

1. Introduction

Preference alignment techniques, most prominently RLHF (Stiennon et al., 2020) and DPO (Rafailov et al., 2024), have become the standard mechanism for reconciling LLMs with human expectations. These algorithms operate on pairwise preference data, in which human annotators compare two candidate responses and designate the one that better addresses a given prompt; the resulting comparisons then supervise fine-tuning, driving the model to rank preferred responses above less preferred ones (Ziegler et al., 2019; Ouyang et al., 2022; Bai et al., 2022). The demonstrated effectiveness and reliability of this paradigm have propelled its adoption across a broad spectrum of applications, spanning chat assistants (Achiam et al., 2023), mathematical reasoning tools (Shao et al., 2024), and code generators (Shen et al., 2023).

Aligning LLMs on human preference data, however, raises acute privacy concerns, precisely because these datasets are constituted from real user prompts and the human feedback rendered over them. The prompts themselves may disclose personal information pertaining to health, identity, or other sensitive matters, while the accompanying feedback may reveal private beliefs, preferences, and behavioral patterns (Li et al., 2023; Yu et al., 2024). A growing body of work (Chowdhury et al., 2024; Yu et al., 2024; Wu et al., 2023a) has sought to mitigate these risks through the rigorous safeguards of differential privacy (DP) (Dwork et al., 2006b), with encouraging results. Yet three limitations persist across this literature. First, many methods protect only part of the data—privatizing either the user prompts (Yu et al., 2024) or the preference labels (Chowdhury et al., 2024; Zhang et al., 2025), but not both—and such partial protection leaves exploitable gaps through which an adversary may still infer sensitive attributes from the unprotected prompts or annotations. Second, these methods are typi-

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

cally bound to a specific post-training algorithm, most often RLHF (Wu et al., 2023a), and are consequently incompatible with newer paradigms such as DPO. Third, they are constrained by the limited private preference data available and therefore frequently fall short of high-quality alignment, since human preference annotations are costly to collect at scale. Together, these limitations motivate the central question of this work:

Is it possible to achieve superior preference alignment while ensuring thorough privacy protection and maintaining compatibility with emerging alignment methodologies?

This work answers this question decisively in the affirmative. Our main contributions are as follows:

- We introduce DPPrefSyn (Figure 1), the first algorithm that synthesizes DP preference data for LLM alignment. DPPrefSyn is a principled framework grounded in the Bradley–Terry preference model and in the intrinsic geometric structure of pairwise human preferences, and it rests on three coordinated design elements. 1) *Capturing diverse human preferences via a shared linear reward structure and clustering.* The Bradley–Terry model together with a linear reward parameterization implies that samples reflecting similar preferences yield embedding-difference vectors—computed between concatenated prompt–preferred and prompt–dispreferred pairs—that point in similar directions; we exploit this geometry to cluster the samples and to approximate the preferences within each cluster by a shared reward function. 2) *Sharpening learning efficiency through DP-PCA.* Because the per-cluster sample budget is inherently limited, we apply DP-PCA to project the high-dimensional embeddings into a compact subspace that retains the essential preference signal under DP, thereby overcoming the sample inefficiency and clustering instability that afflict high-dimensional spaces. 3) *Conserving the privacy budget through public prompts.* We draw candidate responses from public prompts and use the privately learned reward models to select preference pairs, which concentrates the entire privacy budget on modeling private preferences and markedly improves privacy efficiency.
- We establish a rigorous end-to-end privacy analysis proving that DPPrefSyn satisfies (ϵ, δ) -DP. The analysis allocates a budget ϵ_0 to DP-PCA (Amin et al., 2019) for dimensionality reduction and ϵ_1 to DP-KMeans (Su et al., 2016) for clustering, and devotes the remaining budget to training the per-cluster reward models with DP-SGD; the Privacy Random Variable (PRV) accountant (Gopi et al., 2021) tightly composes the DP-SGD cost, so that the composed mechanism provably satisfies (ϵ, δ) -DP. Crucially, the post-processing property of DP renders the synthesized dataset reusable across arbitrary preference-

alignment methods and downstream LLMs at no additional privacy cost—a versatility that single-model private fine-tuning cannot offer.

- We evaluate DPPrefSyn extensively on standard benchmarks, covering the OpenAssistant (Köpf et al., 2023) and Anthropic-HH (Bai et al., 2022) question answering tasks together with the TL;DR summarization task (Stiennon et al., 2020). Across these diverse tasks, the synthetic data generated by DPPrefSyn delivers strong privacy guarantees while attaining high downstream utility. With $\epsilon = 2$, for instance, fine-tuning Pythia-2.8B via DPO on our synthetic data attains a GPT-4o win rate of 56.48% on Anthropic-HH, substantially exceeding the 37.02% obtained by fine-tuning on real data. These results demonstrate that DP synthetic preference data enables effective LLM alignment under strong privacy guarantees. Moreover, because DPPrefSyn produces reusable data rather than a single privatized model, it applies uniformly across preference-optimization methods—including DPO and RLHF—and remains effective under stringent privacy constraints, thereby overcoming the method-specific rigidity of prior privacy-preserving post-training approaches.

2. Related Work

Differentially Private Alignment of LLMs. Wu et al. (2023a) introduce a DP framework that aligns LLMs via reinforcement learning by adapting the PPO algorithm to the DP setting; their approach, however, is *confined to PPO*, whereas preference alignment continues to evolve through new methods. Chowdhury et al. (2024) study reward estimation from preference-based feedback, invoking label-DP (Ghazi et al., 2021) to protect the privacy of human annotators, and Zhang et al. (2025) propose AUP-RLHF, a user-level label-DP framework for RLHF. Although both methods effectively safeguard preference labels, neither *addresses the sensitivity of the prompts or the responses themselves*, and both remain *constrained by the limited private preference data* available, which caps the attainable alignment quality given the high cost of human annotation. Yu et al. (2024) instead protect sensitive user *instructions* by synthesizing DP *instructions* to replace real ones during annotation and fine-tuning. In contrast, DPPrefSyn privatizes prompts, responses, and preference labels jointly, and yields data reusable across alignment methods and models—thereby closing all three gaps at once.

Beyond preference alignment, a broader body of work studies private fine-tuning of LLMs on general labeled data rather than preference comparisons (Yu et al., 2021; Li et al., 2021; Chen et al., 2024; Tang et al., 2024; Zhang et al., 2023; He et al., 2022), and a related line protects the privacy of in-context prompts for in-context learning (Duan et al., 2023; Tang et al., 2023; Gao et al., 2024; Hong et al., 2023;




Figure 1. Overview of DPPrefSyn. DPPrefSyn synthesizes DP preference data in three coordinated steps: 1) it represents each preference sample as an embedding-difference vector and clusters these vectors after dimensionality reduction via DP-PCA and clustering via DP-KMeans; 2) it trains a per-cluster DP reward model with DP-SGD; and 3) it generates synthetic preference pairs from public prompts, guided by the DP-protected distribution of the learned reward functions.

Wu et al., 2023b).

Differentially Private Synthetic Text Generation. Our work situates itself within the broader field of DP synthetic text generation, in which two representative paradigms have emerged. The first fine-tunes a language model on private data under DP and then samples synthetic text from the fine-tuned model (Yue et al., 2023; Mattern et al., 2022; Mireshghallah et al., 2023; Carranza et al., 2023; Yu et al., 2024; Wang et al., 2024; Ochs & Habernal, 2024; Tan et al., 2025; Carranza et al., 2024); Yu et al. (2024), for example, privately fine-tune LLaMA on private instruction data and then generate synthetic text for supervised instruction tuning.

Because DP fine-tuning can be costly and is sometimes infeasible—particularly for non-public language models accessible only through APIs—a second, more recent paradigm generates DP synthetic data by prompting LLM APIs (Lin et al., 2023; Xie et al., 2024; Wu et al., 2024; Hou et al., 2024). These methods sample synthetic candidates through API access and then privately select the outputs most similar to the private data; Xie et al. (2024), for instance, propose Aug-PE, which iteratively refines LLM-generated samples according to their embedding-space similarity to the private data. *Our work belongs to this second paradigm and is, to our knowledge, the first to extend it to preference alignment for LLMs.*

Diversity in Human Preferences. A substantial literature establishes that human preferences are inherently heterogeneous (Denton et al., 2021; Aroyo et al., 2023a,b; Chakraborty et al., 2024), driven by sociodemographic differences (e.g., race, gender, age), personal beliefs and biases, varying domain expertise, and the intrinsic ambiguity of natural language (Sandri et al., 2023; Vogels, 2021). This heterogeneity has motivated efforts to personalize prefer-

ence alignment so as to reflect the values of distinct user groups (Lee et al., 2024b; Poddar et al., 2024; Singh et al., 2025); Singh et al. (2025), for example, propose a meta-learning framework in which an LLM adapts to individual users from a handful of preference examples. DPPrefSyn instead synthesizes DP preference data that explicitly *captures the underlying diversity of human feedback*, and this modeling of diversity yields measurable alignment gains over approaches that ignore it (Section 5.3).

3. Problem Definition

Preference Alignment of LLMs. We consider the alignment of an LLM from pairwise human feedback encoded in a preference dataset $\mathcal{D} = \{(x_i, a_i^+, a_i^-)\}_{i=1}^n$. Each triple comprises a prompt x_i together with two responses a_i^+, a_i^- , where human annotators judge a_i^+ to be preferred over a_i^- .

Following the standard formulation in reinforcement learning from human feedback (RLHF) and direct preference optimization (DPO) (Ouyang et al., 2022; Zhu et al., 2023; Rafailov et al., 2024; Liu et al., 2023), we model human preferences with the Bradley-Terry model (Bradley & Terry, 1952):

$$\begin{aligned} \mathbb{P}[a_i^+ \succ a_i^- \mid x_i] &= \frac{\exp(r^*(x_i, a_i^+))}{\exp(r^*(x_i, a_i^+)) + \exp(r^*(x_i, a_i^-))} \\ &= \sigma(r^*(x_i, a_i^+) - r^*(x_i, a_i^-)), \end{aligned}$$

where $a_i^+ \succ a_i^-$ denotes that a_i^+ is preferred to a_i^- , $r^*(x, a)$ is a latent reward function, and $\sigma(z) = 1/(1 + e^{-z})$ is the sigmoid function.

This formulation underpins both RLHF, which learns the reward model explicitly from preference data, and DPO, which induces an implicit reward through the log-ratio between the learned policy and a reference policy—so that a

preference model estimated from data serves as the common substrate for either alignment route.

Differentially Private Preference Data Synthesis. Our objective is to protect the privacy of the preference dataset $\mathcal{D}_{\text{priv}} = \{(x_i, a_i^+, a_i^-)\}_{i=1}^n$ against an adversary who seeks to access or infer private information about individual prompts, responses, or preference labels from the outputs of the fine-tuned LLM. To this end, we synthesize a DP preference dataset $\tilde{\mathcal{D}} = \{(\tilde{x}_j, \tilde{a}_j^+, \tilde{a}_j^-)\}_{j=1}^m$ that preserves the critical preference information carried by $\mathcal{D}_{\text{priv}}$. By the post-processing property of DP, $\tilde{\mathcal{D}}$ can be reused across arbitrary preference-alignment methods and LLMs without incurring any additional privacy cost.

We formally define (ϵ, δ) -differential privacy as follows.

Definition 3.1 ((ϵ, δ) -Differential Privacy (Dwork et al., 2006a)). A randomized algorithm $\mathcal{A}$ is (ϵ, δ) -differentially private if for any two neighboring inputs $\mathcal{D}$ and $\mathcal{D}'$ that differ by a single entry and any set $\mathcal{S}$ of possible outputs: $\mathbb{P}[\mathcal{A}(\mathcal{D}) \in \mathcal{S}] \leq e^\epsilon \mathbb{P}[\mathcal{A}(\mathcal{D}') \in \mathcal{S}] + \delta$.

Diverse User Preferences and Linear Reward Structure. Human preferences in the private dataset are generically heterogeneous (Chakraborty et al., 2024; Bakker et al., 2022; Kovač et al., 2023; Jang et al., 2023; Rame et al., 2023; Ji et al., 2023): different annotators privilege different attributes of a response—factual accuracy, politeness, creativity, or clarity—yielding diverse and sometimes conflicting signals within the data. Faithfully synthesizing $\tilde{\mathcal{D}}$ therefore hinges on capturing this diversity of preferences in $\mathcal{D}_{\text{priv}}$.

To this end, we posit a linear structure for the latent reward function (Saha et al., 2023; Kong & Yang, 2022; Zhu et al., 2023), adopting the linear reward model

$$r_\theta(x, a) = \langle \theta, \phi(x, a) \rangle,$$

in which $\phi(x, a) \in \mathbb{R}^d$ is a *shared* feature vector and $\theta \in \mathbb{R}^d$ is a *user-dependent* parameter.

This parameterization strikes a deliberate balance between *expressiveness* and *tractability* in modeling diverse preferences. Although users differ in what they value, their reward functions typically share common semantic features in a learned embedding space (the vector $\phi(x, a)$) and diverge principally in how those features are weighted (the parameter θ). Factoring the reward in this way captures the shared structure while admitting variability across users and preference clusters, which makes heterogeneous preferences learnable even under limited data and strong privacy constraints. As our experiments in Appendix B.6 confirm, this linear reward structure captures preference signals effectively, achieving accuracy competitive with a fully fine-tuned GPT-2 reward model despite its markedly greater simplicity.

4. Proposed Method

4.1. Key Components of the Algorithm Design

Before presenting the full algorithm, we isolate the key design components through which DPPrefSyn confronts the principal challenges of synthesizing differentially private preference data.

Leveraging the Linear Reward Structure for Preference Clustering. Under the linear reward model, the Bradley–Terry likelihood takes the form

$$\mathbb{P}[a_i^+ \succ a_i^- \mid x_i] = \sigma(\langle \theta, \phi(x_i, a_i^+) - \phi(x_i, a_i^-) \rangle).$$

Consequently, for a triple (x_i, a_i^+, a_i^-) the observed preference $a_i^+ \succ a_i^-$ entails $\langle \theta, \phi(x_i, a_i^+) - \phi(x_i, a_i^-) \rangle \gg 0$; that is, the preference vector θ is strongly aligned with the feature-difference vector $\phi(x_i, a_i^+) - \phi(x_i, a_i^-)$. It follows that, for annotators who share a preference direction, the feature-difference vectors across their rated pairs concentrate along a common direction in the embedding space. This geometric consequence motivates our design directly: we cluster the preference samples by their feature-difference vectors $\phi(x_i, a_i^+) - \phi(x_i, a_i^-)$, so that the preferences within each cluster are well approximated by a shared reward function parameterized by a single θ . In Appendix B.9 we show that the clusters obtained in this way are meaningful and interpretable as coherent user-preference styles.

Adopting DP-PCA to Obtain Effective Features. A natural instantiation of the feature vector $\phi(x, a)$ is the embedding of the corresponding (prompt, response) pair, since such embeddings encode the semantic similarities and differences that directly underlie human preference judgments. These embeddings, however, are high-dimensional (e.g., 1024 dimensions) in order to represent rich semantic and stylistic content, and reliable learning in so high-dimensional a space generally demands far more samples than the per-cluster budget affords. To resolve this tension, we apply differentially private PCA (DP-PCA) (Amin et al., 2019; Liu et al., 2022) to project the embeddings into a compact subspace before clustering and reward-model training. This projection improves sample efficiency and stabilizes clustering while preserving the preference signal essential for effective alignment.

Utilizing Public Prompts to Conserve the DP Budget. Synthesizing prompts, preferred responses, and less preferred responses jointly would rapidly deplete the privacy budget and erode utility. We avoid this by sourcing prompts from public data, which frees us from spending any budget on privatizing prompts and lets us devote the entire budget to modeling preferences over responses. For each public prompt we generate several candidate responses with an LLM and construct preference pairs using a distribution of reward functions learned from the private preference

data under DP. Although the public prompts may follow a different distribution from the private ones, our experiments demonstrate that—despite this distribution shift—the synthetic data faithfully preserves the statistical preference distribution of the private dataset and remains effective for downstream alignment.

4.2. DPPrefSyn Algorithm

We now present the DPPrefSyn algorithm (Algorithm 1), which synthesizes a DP preference dataset from the private dataset $\mathcal{D}_{\text{priv}}$ through three main steps:

Step 1: Preference Representation and Clustering. We encode each prompt–response pair in $\mathcal{D}_{\text{priv}}$ with a public sentence-embedding model ψ . Concretely, for each sample (x_i, a_i^+, a_i^-) we concatenate the prompt with each response to form the two texts $[x_i; a_i^+]$ and $[x_i; a_i^-]$, embed them through ψ , and take their difference $d_i \leftarrow \psi(x_i, a_i^+) - \psi(x_i, a_i^-)$ (Line 3), which encodes the direction of the observed preference.

To reduce dimensionality privately, we apply DP-PCA (Amin et al., 2019) to $\{d_i\}_{i=1}^n$, yielding a projection matrix $\mathbf{P} \in \mathbb{R}^{d \times p}$ under ε_0 -DP (Line 5); DP-PCA (Amin et al., 2019) approximates the eigendecomposition of the data covariance matrix by privately estimating its eigenvalues and eigenvectors separately. We then project each difference vector into the resulting subspace, $z_i \leftarrow \mathbf{P}^\top d_i$ (Line 6), and apply DP-KMeans clustering (Su et al., 2016) to $\{z_i\}_{i=1}^n$ under ε_1 -DP, grouping samples with similar preference patterns into K clusters $\mathcal{C}_1, \dots, \mathcal{C}_K$ (Line 7); DP-KMeans (Su et al., 2016) is a DP counterpart of Lloyd’s algorithm.

Step 2: DP Reward Model Training. For each cluster $\mathcal{C}_k$, we estimate the shared preference vector $\theta_k \in \mathbb{R}^p$ that governs its preference pattern by minimizing the negative log-likelihood (Line 9):

$$\mathcal{L}_{\text{MLE}}(\theta_k) = -\frac{1}{|\mathcal{C}_k|} \sum_{z_i \in \mathcal{C}_k} \log \sigma(\langle \theta_k, z_i \rangle), \quad (1)$$

where z_i is the DP-PCA-projected embedding difference from Step 1. Each per-cluster reward model is trained with DP-SGD (Abadi et al., 2016), the differentially private variant of stochastic gradient descent, using noise multiplier σ_2 .

Step 3: DP Synthetic Preference Generation. We first form the cluster-size histogram $\mathbf{h} = [|\mathcal{C}_1|, \dots, |\mathcal{C}_K|]$ (Line 11) and normalize it into a probability distribution over clusters, $\mathbf{p} \leftarrow \mathbf{h}/|\mathcal{D}_{\text{priv}}|$ (Line 12), which encodes the relative prevalence of each preference style in the private data.

For each public prompt $\tilde{x}_j$, we then prompt an LLM at a relatively high temperature to generate L diverse

candidate responses $\tilde{a}_j^1, \dots, \tilde{a}_j^L$ (Line 14), draw a cluster index $k \sim \mathbf{p}$ from the DP histogram (Line 15), and score the candidates with the reward model θ_k . Scoring proceeds by embedding each candidate as $\psi(\tilde{x}_j, \tilde{a}_j^l)$ for $l = 1, \dots, L$, projecting it into the DP-PCA subspace via $\mathbf{P}$, and evaluating the reward $\langle \theta_k, \mathbf{P}^\top \psi(\tilde{x}_j, \tilde{a}_j^l) \rangle$. We designate the highest- and lowest-scoring candidates as the preferred and less preferred responses, respectively: $c \leftarrow \arg \max_{l \in [L]} \langle \theta_k, \mathbf{P}^\top \psi(\tilde{x}_j, \tilde{a}_j^l) \rangle$ and $r \leftarrow \arg \min_{l \in [L]} \langle \theta_k, \mathbf{P}^\top \psi(\tilde{x}_j, \tilde{a}_j^l) \rangle$ (Lines 16–17). To enforce preference quality, we discard the sample whenever the reward gap is too small (e.g., < 0.5) and otherwise add $(\tilde{x}_j, \tilde{a}_j^c, \tilde{a}_j^r)$ to $\tilde{\mathcal{D}}$. Repeating this procedure over all public prompts yields the DP synthetic preference dataset $\tilde{\mathcal{D}}$.

Algorithm 1 DPPrefSyn

```

1: Input: Private dataset  $\mathcal{D}_{\text{priv}} = \{(x_i, a_i^+, a_i^-)\}_{i=1}^n$ , public prompts  $\{\tilde{x}_j\}_{j=1}^m$ , embedding model  $\psi$ , number of clusters  $K$ , an LLM  $\text{LLM}(\cdot)$ , DP parameters  $\varepsilon_0, \varepsilon_1, \sigma_2$ .
2: for each  $(x_i, a_i^+, a_i^-)$  in  $\mathcal{D}_{\text{priv}}$  do
3:    $d_i \leftarrow \psi(x_i, a_i^+) - \psi(x_i, a_i^-)$ , where  $d_i \in \mathbb{R}^d$ 
4: end for
5:  $\mathbf{P} \leftarrow \text{DP-PCA}(\{d_i\}_{i=1}^n, \varepsilon_0)$ , where  $\mathbf{P} \in \mathbb{R}^{d \times p}$ 
6:  $\{z_i\}_{i=1}^n \leftarrow \{\mathbf{P}^\top d_i\}_{i=1}^n$ 
7:  $\mathcal{C}_1, \dots, \mathcal{C}_K \leftarrow \text{DP-KMeans}(\{z_i\}_{i=1}^n, K, \varepsilon_1)$ 
8: for each cluster  $\mathcal{C}_k$  do
9:   Train a linear reward model  $\theta_k$  on  $\mathcal{C}_k$  using DP-SGD with noise multiplier  $\sigma_2$  (Equation (1))
10: end for
11:  $\mathbf{h} \leftarrow [|\mathcal{C}_1|, \dots, |\mathcal{C}_K|]$ 
12:  $\mathbf{p} \leftarrow \mathbf{h}/|\mathcal{D}_{\text{priv}}|$ 
13: for each public prompt  $\tilde{x}_j$  do
14:   Generate responses  $\tilde{a}_j^1, \dots, \tilde{a}_j^L \leftarrow \text{LLM}(\tilde{x}_j)$ 
15:   Sample cluster index  $k \sim \mathbf{p}$ 
16:    $c \leftarrow \arg \max_{l \in [L]} \langle \theta_k, \mathbf{P}^\top \psi(\tilde{x}_j, \tilde{a}_j^l) \rangle$ 
17:    $r \leftarrow \arg \min_{l \in [L]} \langle \theta_k, \mathbf{P}^\top \psi(\tilde{x}_j, \tilde{a}_j^l) \rangle$ 
18:   Add  $(\tilde{x}_j, \tilde{a}_j^c, \tilde{a}_j^r)$  to  $\tilde{\mathcal{D}}$ 
19: end for
20: return  $\tilde{\mathcal{D}}$ .
```

Privacy Analysis of DPPrefSyn. The budget is partitioned across the three data-accessing subroutines: ε_0 to DP-PCA (Line 5), ε_1 to DP-KMeans (Line 7), and the remainder to per-cluster reward-model training via DP-SGD with noise multiplier σ_2 (Line 9). Reward-model training benefits from the parallel-composition property of DP: because the clusters are disjoint, altering a single training point perturbs only one cluster, so the total privacy cost of this step is bounded by that of training on the smallest cluster; we guarantee a known lower bound on this size by discarding clusters with too few samples. Composing the per-step DP-SGD costs with the Privacy Random Variable (PRV) accountant (Gopi et al., 2021) then ensures that reward-

model training satisfies $(\varepsilon - \varepsilon_0 - \varepsilon_1, \delta)$ -DP. Since every remaining step operates only on public data or is a post-processing of DP outputs, it incurs no additional privacy cost, and the composed mechanism therefore satisfies (ε, δ) -DP. A complete proof and the accounting code are given in Appendix A.

Discussion on Prompt Distribution Shift. DPPrefSyn remains effective even when the private and public prompt distributions diverge. The reason is structural: diversity of preference is encoded entirely in θ_k , the parameter of the linear reward model, rather than in $\phi(x, a)$, the feature vector of the prompt–response pair. Intuitively, a user may interact with an LLM through widely varying prompts while their underlying preference remains fixed. It therefore suffices to make the distribution of θ_k resemble that of the private data for the synthetic preferences to reflect the private preference patterns—a matching we achieve by drawing, for each public prompt, a cluster index k from the DP histogram (Line 15) and selecting the synthetic response pair under the corresponding reward model. Consequently, a shift between the private and public prompt distributions leaves the preference patterns captured from the private data intact.

5. Experiments

Datasets. We evaluate across both question answering and summarization, spanning three complementary preference benchmarks. For question answering we use OpenAssistant (Köpf et al., 2023), a corpus of assistant-style conversations, and Anthropic-HH (Bai et al., 2022), which supplies human preference comparisons targeting helpfulness and harmlessness; for summarization we use TL;DR (Stiennon et al., 2020), which provides human preferences over pairs of summaries. To supply public prompts, we draw on Alpaca (Taori et al., 2023) for the OpenAssistant QA task, SafeRLHF (Ji et al., 2024) for the Anthropic-HH QA task, and XSum (Narayan et al., 2018) for the TL;DR summarization task—so that in every case the public prompts originate from a distinct source, testing DPPrefSyn precisely under the prompt distribution shift analyzed above. Additional details appear in Appendix B.2.

Setup for Algorithm 1. Unless otherwise stated, we instantiate the embedding model ψ with BAAI/bge-large-en-v1.5 (Xiao et al., 2023), set the DP-PCA projected dimension to $p = 20$, and use $K = 5$ clusters in DP-KMeans. DP-SGD is run with a learning rate of 0.1, a batch size of 4, 4 training epochs, and a gradient clipping norm of 1.0. Candidate responses are generated by the instruction-finetuned LLaMA-7B-chat model (Touvron et al., 2023) at temperature 0.9, with $L = 5$ candidates per prompt. We report results across the overall privacy budgets $\varepsilon = 0.5, 1, 2, 4, 8$ and set $\delta = 1/|\mathcal{D}_{\text{priv}}|$. We further examine the effect of the generator model

in Table 3, and of the embedding model and remaining hyperparameters in Section 5.3.

Evaluation. For downstream alignment we fine-tune the Pythia-2.8B model (Biderman et al., 2023). We first perform supervised fine-tuning (SFT), taking the preferred response of each preference pair as the training target, and then apply DPO (Rafailov et al., 2024) to further fine-tune the SFT model on the preference pairs. Following Rafailov et al. (2024), we quantify utility by the win rate of the model-generated responses against the preferred responses in the test set, adjudicated by GPT-4o. We extend this evaluation to additional downstream models and alternative preference-alignment methods in the ablation study of Section 5.2.

Baselines. Our primary baseline, denoted DP-FT, fine-tunes the downstream model directly on the real preference data with DP-SGD (Abadi et al., 2016). DPPrefSyn is fundamentally more reusable and versatile than DP-FT: it produces DP synthetic data that can train arbitrary models with arbitrary alignment algorithms at no additional privacy cost, whereas DP-FT expends its entire budget on a single privatized model. We additionally report a fully private reference point $\varepsilon = 0$, corresponding to the base model evaluated without any preference fine-tuning. Full setups, hyperparameters, metrics, and baseline configurations are detailed in Appendix B.3.

5.1. Understanding the Performance of DPPrefSyn

We analyze the performance of DPPrefSyn through two research questions concerning its privacy–utility trade-off, with the main results reported in Table 1.

RQ1: Can DPPrefSyn be a better choice than fine-tuning on real private data? DPPrefSyn attains a strictly stronger privacy–utility trade-off than DP-FT on every task at each of $\varepsilon = 0.5, 1, 2, 4, 8$, and—while preserving privacy—even exceeds the utility of unconstrained, non-private fine-tuning on real data (DP-FT with $\varepsilon = \infty$). On the OpenAssistant QA task at $\varepsilon = 2$, for instance, DPPrefSyn reaches 9.10% after SFT and 11.04% after DPO, against DP-FT’s 4.38% and 6.18%; strikingly, both figures also surpass the non-private DP-FT reference ($\varepsilon = \infty$), which attains only 4.75% after SFT and 8.20% after DPO. The same pattern holds on Anthropic-HH at $\varepsilon = 2$, where DPPrefSyn achieves win rates of 55.21% after SFT and 56.48% after DPO, versus 31.53% and 37.02% for DP-FT and 31.98% and 38.72% for non-private DP-FT ($\varepsilon = \infty$). That a privacy-preserving method should outperform non-private fine-tuning on the same real data is a decisive result: it shows the advantage of DPPrefSyn stems not from relaxing privacy but from the quality and diversity of the synthesized preferences.

These results expose two **key advantages** of our DP syn-

Table 1. GPT-4o win rates (%) on the test sets for the downstream Pythia-2.8B model fine-tuned by SFT and by SFT+DPO, across the three tasks and privacy budgets $\varepsilon \in \{0.5, 1, 2, 4, 8, \infty\}$; $\varepsilon = 0$ denotes the base (fully private, non-finetuned) LLM. At every privacy level $\varepsilon = 0.5, 1, 2, 4, 8$, DPPrefSyn on synthetic data attains markedly higher win rates than DP-FT on real data, and even surpasses non-private fine-tuning (DP-FT with $\varepsilon = \infty$). Values are mean and standard deviation over 5 random seeds; best results in each block are in bold.

Task	$\varepsilon = 0$	Method	Data Type	Pref. Alignment	$\varepsilon = 0.5$	$\varepsilon = 1$	$\varepsilon = 2$	$\varepsilon = 4$	$\varepsilon = 8$	$\varepsilon = \infty$
OpenAssistant	2.11	DP-FT	Original	SFT	3.96 _{0.20}	4.29 _{0.21}	4.38 _{0.23}	4.46 _{0.18}	4.49 _{0.30}	4.75 _{0.27}
				SFT+DPO	5.65 _{0.32}	6.04 _{0.58}	6.18 _{0.83}	6.12 _{0.63}	6.32 _{0.71}	8.20 _{0.62}
		DPPrefSyn	Synthetic	SFT	8.96 _{0.39}	9.07 _{0.32}	9.10 _{0.51}	9.33 _{0.56}	9.63 _{0.38}	9.80 _{0.58}
				SFT+DPO	9.86_{0.59}	10.31_{0.54}	11.04_{0.92}	11.85_{0.99}	12.05_{0.82}	12.36_{0.55}
Anthropic-HH	12.14	DP-FT	Original	SFT	29.77 _{0.76}	31.10 _{0.72}	31.53 _{0.77}	31.29 _{0.65}	31.37 _{0.67}	31.98 _{0.56}
				SFT+DPO	35.00 _{0.49}	36.27 _{0.63}	37.02 _{0.72}	36.74 _{0.90}	36.94 _{1.06}	38.72 _{0.56}
		DPPrefSyn	Synthetic	SFT	54.90 _{0.66}	55.19 _{0.91}	55.21 _{0.82}	54.95 _{1.10}	55.63 _{0.86}	55.95 _{0.70}
				SFT+DPO	55.08_{1.04}	55.96_{0.80}	56.48_{0.61}	56.51_{1.93}	56.86_{1.17}	57.53_{0.93}
TL;DR	11.64	DP-FT	Original	SFT	21.19 _{1.12}	22.41 _{1.09}	22.57 _{0.65}	22.45 _{0.90}	23.04 _{0.79}	23.82 _{1.26}
				SFT+DPO	59.69 _{0.82}	62.21 _{1.12}	62.02 _{1.67}	63.00 _{0.73}	63.06 _{0.84}	64.52 _{1.79}
		DPPrefSyn	Synthetic	SFT	61.92 _{1.06}	62.14 _{1.61}	62.49 _{1.48}	62.72 _{1.66}	62.29 _{1.14}	62.55 _{1.23}
				SFT+DPO	67.50_{0.79}	68.07_{1.95}	68.54_{2.76}	68.56_{1.59}	71.01_{2.17}	72.20_{1.08}

thetic data. **First**, DPPrefSyn yields a substantial gain from SFT alone relative to SFT on private data: on the TL;DR summarization task at $\varepsilon = 2$, it reaches 62.49% after SFT, far above the 22.57% obtained by SFT on private data. This gain traces to DPPrefSyn’s use of a high-performance LLM to generate candidate responses, which furnishes higher-quality training targets and thus more effective supervision during SFT—and the advantage persists even when public data is scarce. On the OpenAssistant QA task at $\varepsilon = 4$ with only 5K public prompts, for example, DPPrefSyn attains a 6.80% win rate after SFT against 4.46% for SFT directly on the roughly 14K private samples (Table 4), underscoring its robustness to data scarcity. **Second**, applying DPO to the SFT model on our synthetic data drives further improvement, evidence that the synthesized preference pairs encode meaningful distinctions aligned with human preferences—an effect enabled by the private reward functions that capture diverse user preferences while preserving privacy. The incremental gain from DPO is smaller here than when using private data, plausibly because the SFT model trained on our synthetic data is already strong and thus leaves less headroom; nonetheless, the utility of the complete SFT+DPO pipeline on our DP synthetic data remains markedly above that obtained on real private data.

RQ2: How does DPPrefSyn perform across the privacy budget ε ? Two consistent patterns emerge. **First**, as Table 1 shows, DPPrefSyn maintains its advantage over DP-FT as ε ranges over 0.5, 1, 2, 4, 8, ∞ , and its win rate rises smoothly with the budget, indicating graceful scalability in ε . **Second**, DPPrefSyn is often substantially more robust to DP noise than DP-FT: on OpenAssistant with Llama-3-8B as the downstream model, DPO win rate for DP-FT collapses from 59.13% to 45.79% as ε tightens from ∞ to 4, whereas DPPrefSyn degrades only marginally, from 61.52% to 60.25% (Table 3 in Section 5.2). This robustness is expected: DP-FT

Table 2. On OpenAssistant, DPPrefSyn ($\varepsilon = 4$) attains higher win rates than non-private fine-tuning on real data (DP-FT, $\varepsilon = \infty$) under all three alignment methods—SFT, SFT+DPO, and SFT+PPO—while providing strong privacy guarantees. Values are mean and standard deviation over 5 random seeds; best results in each row are in bold.

Pref. Alignment	$\varepsilon = 0$	DPPrefSyn ($\varepsilon = 4$)	DP-FT ($\varepsilon = \infty$)
SFT		9.33_{0.56}	4.75 _{0.27}
SFT+DPO	2.11	11.85_{0.99}	8.20 _{0.62}
SFT+PPO		10.11_{0.61}	7.08 _{0.78}

injects noise directly into the model parameters through DP-SGD, while DPPrefSyn confines noise to the data-synthesis stage, thereby better preserving downstream utility.

5.2. Understanding the Properties of DPPrefSyn

We now characterize the broader properties of DPPrefSyn: its generality across preference-alignment methods and downstream LLMs, its behavior under varying scales of public data, and its resilience to empirical privacy attacks.

RQ3: Is DPPrefSyn effective across different preference alignment methods? **DPPrefSyn is method-agnostic, remaining effective across SFT, DPO, and PPO under strong DP constraints.** Table 2 reports all three. For PPO we train K reward models—one per cluster—to capture diverse preference patterns, and optimize the policy against their averaged reward. Across every method, DPPrefSyn delivers consistently higher utility than fine-tuning on private data while preserving strong privacy guarantees, confirming that its benefits transfer directly to the reward-based PPO setting and are not confined to the DPO route on which it was primarily evaluated.

RQ4: Can more powerful downstream LLMs benefit from synthetic data generated by DPPrefSyn? **DPPrefSyn is effective across a broad range of downstream LLMs.**

Differentially Private Preference Data Synthesis for Large Language Model Alignment

Table 3. On OpenAssistant after DPO, DPPrefSyn attains higher win rates than DP-FT across three downstream LLM families (Qwen-3-4B, Gemma-2-2B, Llama-3-8B) while providing strong privacy guarantees, and further improves with a stronger generator (Qwen3-4B-Instruct-2507 over Llama-3-8B-Instruct). Best results in each block are in bold.

Downstream LLM	$\epsilon = 0$	Method	Generator LLM	$\epsilon = 4$	$\epsilon = \infty$
Qwen-3-4B	48.17	DP-FT	-	54.92	59.41
		DPPrefSyn	Qwen-3-4B-Instruct	86.52	87.08
			Llama-3-8B-Instruct	51.69	52.11
Gemma-2-2B	7.30	DP-FT	-	30.06	33.85
		DPPrefSyn	Qwen-3-4B-Instruct	36.24	37.08
			Llama-3-8B-Instruct	36.01	36.66
Llama-3-8B	16.15	DP-FT	-	45.79	59.13
		DPPrefSyn	Qwen-3-4B-Instruct	60.25	61.52
			Llama-3-8B-Instruct	42.28	44.10

Table 4. On OpenAssistant at $\epsilon = 4$, DPPrefSyn exceeds DP-FT even with far fewer public prompts than the 14,167 private samples, and its margin widens as the public data size m grows. Values are mean and standard deviation over 5 random seeds.

Data Type (Size)	DPPrefSyn ($\epsilon = 4$)			DP-FT ($\epsilon = 4$)	DP-FT ($\epsilon = \infty$)
	Synthetic ($m = 5,000$)	Synthetic ($m = 10,000$)	Synthetic ($m = 52,000$)	Original (14,167)	Original (14,167)
SFT	6.80 _{0.66}	7.30 _{0.48}	9.33 _{0.56}	4.46 _{0.18}	4.75 _{0.27}
SFT+DPO	9.44 _{0.45}	9.55 _{0.92}	11.85 _{0.99}	6.12 _{0.63}	8.20 _{0.62}


In Table 3 we evaluate the downstream models Qwen-3-4B, Gemma-2-2B, and Llama-3-8B, with Llama-3-8B-Instruct and Qwen3-4B-Instruct-2507 as generators. In every case DPPrefSyn surpasses fine-tuning on real data after DPO while preserving privacy across model families: with Qwen3-4B-Instruct-2507 as the generator, DPPrefSyn reaches an 86.52% win rate on Qwen-3-4B at $\epsilon = 4$, far exceeding non-private fine-tuning on real data (59.41%). Performance also scales with generator strength—Qwen3-4B-Instruct-2507 consistently outperforms Llama-3-8B-Instruct—indicating that DPPrefSyn directly inherits continuing improvements in generator quality.

RQ5: Can DPPrefSyn be effective across varying scales of public data? Table 4 studies the effect of the public data size m . Here the private dataset (OpenAssistant) contains 14K preference annotations and the full public dataset (Alpaca) contains 52K prompts, from which we randomly draw subsets of size $m = 5K, 10K, 52K$ to synthesize preference data. Even with only 5K public prompts, DPPrefSyn attains 6.80% after SFT and 9.44% after DPO at $\epsilon = 4$, already surpassing the DP-FT baseline’s 4.46% and 6.12%. Two conclusions follow. **First, DPPrefSyn remains effective even under limited public data**, owing to the high-quality LLM-generated responses and the private reward functions that accurately capture user preferences. **Second, its margin over the baseline widens as public data grows**, showing that DPPrefSyn converts abundant public data into alignment gains and thereby directly counteracts the scarcity of private preference data.

RQ6: Can DPPrefSyn effectively mitigate empirical pri-

Table 5. Empirical privacy under membership inference attacks (MIAs) (Feng et al., 2025) on OpenAssistant: DPPrefSyn drives the attack AUC down to near 50% (random guessing) across all downstream LLMs, whereas DP-FT remains markedly vulnerable. Values are mean and standard deviation over 5 random seeds; lowest AUC in each block is in bold.

Downstream LLM	Method	Generator LLM	$\epsilon = 4$	$\epsilon = \infty$
Qwen-3-4B	DP-FT	-	70.15 _{2.09}	73.52 _{2.30}
	DPPrefSyn	Qwen-3-4B-Instruct	50.99 _{1.10}	50.87 _{0.84}
		Llama-3-8B-Instruct	50.64_{1.29}	49.98_{1.19}
Gemma-2-2B	DP-FT	-	62.88 _{1.78}	65.39 _{2.59}
	DPPrefSyn	Qwen-3-4B-Instruct	50.51_{1.38}	50.43_{0.91}
		Llama-3-8B-Instruct	50.77 _{0.83}	51.10 _{0.54}
Llama-3-8B	DP-FT	-	61.95 _{2.05}	70.13 _{1.24}
	DPPrefSyn	Qwen-3-4B-Instruct	50.53 _{0.62}	50.96 _{0.95}
		Llama-3-8B-Instruct	50.45_{0.87}	50.67_{0.96}



(a) Effect of ablating DP-KMeans or DP-PCA. (b) Effect of varying the number of clusters K .

vacuity attacks? Beyond its formal DP guarantee, we assess empirical privacy through membership inference attacks (MIAs) (Shokri et al., 2017) against the DPO-fine-tuned downstream models on OpenAssistant, reporting AUC in Table 5. Following Feng et al. (2025), who instantiate an MIA tailored to preference data for LLM alignment, the attacker seeks to decide whether a private sample (x, a^+, a^-) was used in alignment. DP-FT fails to neutralize this attack, whereas **DPPrefSyn consistently drives the AUC down to near random guessing**. This heightened robustness follows naturally from the synthetic nature of the fine-tuning data, which contains no real private sample and thus removes the very membership signal that overfitting to real data would leave behind. Full details are given in Appendix B.4.

5.3. Validating the Design of DPPrefSyn

We now validate the design of DPPrefSyn by isolating the contribution of its algorithmic components.

RQ7: How does modeling diverse human preferences impact DPPrefSyn’s performance? We probe this by ablating DP-KMeans and DP-PCA and by varying the number of clusters K , all on OpenAssistant at $\epsilon = 4$. (i) **Ablating DP-KMeans or DP-PCA.** As Figure 2a shows, removing either component degrades performance. Without DP-KMeans, all samples collapse into a single cluster and preference diversity is lost; without DP-PCA, clustering becomes unreliable and learning grows sample-inefficient in the high-dimensional embedding space. Retaining both components and thus multiple clusters consistently outperforms the single-cluster

variant, confirming that each component is necessary. (ii) **Varying the number of clusters K .** Figure 2b evaluates $K = 1, 3, 5, 8$: performance rises sharply from $K = 1$ to $K = 3$ and then plateaus at $K = 5, 8$. DPPrefSyn is therefore robust to K , with moderate values ($K = 3, 5$) striking the best trade-off between preference diversity and per-cluster data efficiency. Together, these ablations establish that modeling preference diversity through multiple clusters is central to DPPrefSyn’s performance.

RQ8: How do different embedding models or hyperparameters affect DPPrefSyn’s performance? We ablate the embedding model ψ , the DP-PCA dimension p , the number of candidates per prompt L , the filter threshold, and the privacy-budget allocation in Appendix B.5. DPPrefSyn proves robust throughout. It remains consistently effective across the sentence encoders bge-large-en-v1.5 (Xiao et al., 2023), all-mpnet-base-v2 (Song et al., 2020), and all-distilroberta-v1 (Sanh et al., 2019); $p = 20$ yields the best overall post-DPO performance and $L = 5$ is a sound default; a moderate filtering threshold of 0.5 slightly improves over no filtering while an aggressive threshold of 1 over-prunes and hurts performance; and the method is insensitive to reasonable reallocations of the budget between ε_0 and ε_1 . This stability indicates that DPPrefSyn’s gains are structural rather than the product of delicate hyperparameter tuning.

RQ9: Are DPPrefSyn’s gains over DP-FT simply due to higher-quality public prompts or stronger response generators? To exclude these confounds, we run controlled ablations in which all methods share the same prompt set or the same response generator. Even under these matched conditions DPPrefSyn continues to outperform DP-FT (Appendix B.7), demonstrating that **its gains do not arise merely from higher-quality public prompts or stronger response generators** but from the privacy-preserving preference modeling at the core of the method.

6. Conclusion

We introduced DPPrefSyn, the first algorithm for synthesizing differentially private preference data from a private dataset. By modeling preference diversity through clustering, sharpening learning efficiency with DP-PCA, and conserving the privacy budget through public prompts, DPPrefSyn enables effective and private preference-data synthesis for LLM alignment. Across question answering and summarization benchmarks, multiple downstream model families, and several alignment methods, DPPrefSyn satisfies strong DP guarantees while consistently outperforming fine-tuning on real data—and frequently surpassing even its non-private counterpart—thereby establishing synthetic preference data as a practical and broadly applicable foundation for privacy-preserving alignment.

7. Limitations

DPPrefSyn relies on LLM access to synthesize preference data, which may incur additional API costs. Our study concerns English text-based preference data; extending DPPrefSyn to multilingual and multimodal settings remains a promising open direction. Finally, owing to computational constraints, our experiments span downstream models up to 8B parameters, and scaling to larger models is left to future work.

Impact statement

This work enables preference alignment of large language models without exposing personal information, thereby reconciling the benefits of large-scale data with the imperative of protecting individual privacy, promoting ethical data usage, and fostering societal trust in an increasingly data-driven world.

References

- Abadi, M., Chu, A., Goodfellow, I., McMahan, H. B., Mironov, I., Talwar, K., and Zhang, L. Deep learning with differential privacy. In *Proceedings of the 2016 ACM SIGSAC conference on computer and communications security*, pp. 308–318, 2016.
- Achiam, J., Adler, S., Agarwal, S., Ahmad, L., Akkaya, I., Aleman, F. L., Almeida, D., Altenschmidt, J., Altman, S., Anadkat, S., et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- Amin, K., Dick, T., Kulesza, A., Munoz, A., and Vassilvitskii, S. Differentially private covariance estimation. *Advances in Neural Information Processing Systems*, 32, 2019.
- Aroyo, L., Diaz, M., Homan, C., Prabhakaran, V., Taylor, A., and Wang, D. The reasonable effectiveness of diverse evaluation data. *arXiv preprint arXiv:2301.09406*, 2023a.
- Aroyo, L., Taylor, A., Diaz, M., Homan, C., Parrish, A., Serapio-García, G., Prabhakaran, V., and Wang, D. Dices dataset: Diversity in conversational ai evaluation for safety. *Advances in Neural Information Processing Systems*, 36:53330–53342, 2023b.
- Bai, Y., Jones, A., Ndousse, K., Askell, A., Chen, A., Das-Sarma, N., Drain, D., Fort, S., Ganguli, D., Henighan, T., et al. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*, 2022.
- Bakker, M., Chadwick, M., Sheahan, H., Tessler, M., Campbell-Gillingham, L., Balaguer, J., McAleese, N., Glaese, A., Aslanides, J., Botvinick, M., et al. Fine-tuning

- language models to find agreement among humans with diverse preferences. *Advances in Neural Information Processing Systems*, 35:38176–38189, 2022.
- Biderman, S., Schoelkopf, H., Anthony, Q. G., Bradley, H., O’Brien, K., Hallahan, E., Khan, M. A., Purohit, S., Prashanth, U. S., Raff, E., et al. Pythia: A suite for analyzing large language models across training and scaling. In *International Conference on Machine Learning*, pp. 2397–2430. PMLR, 2023.
- Bradley, R. A. and Terry, M. E. Rank analysis of incomplete block designs: I. the method of paired comparisons. *Biometrika*, 39(3/4):324–345, 1952.
- Carranza, A., Farahani, R., Ponomareva, N., Kurakin, A., Jagielski, M., and Nasr, M. Synthetic query generation for privacy-preserving deep retrieval systems using differentially private language models. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 3920–3930, 2024.
- Carranza, A. G., Farahani, R., Ponomareva, N., Kurakin, A., Jagielski, M., and Nasr, M. Privacy-preserving recommender systems with synthetic query generation using differentially private large language models. *arXiv preprint arXiv:2305.05973*, 2023.
- Chakraborty, S., Qiu, J., Yuan, H., Koppel, A., Manocha, D., Huang, F., Bedi, A., and Wang, M. Maxmin-rlhf: Alignment with diverse human preferences. In *International Conference on Machine Learning*, pp. 6116–6135. PMLR, 2024.
- Chen, T., Da, L., Zhou, H., Li, P., Zhou, K., Chen, T., and Wei, H. Privacy-preserving fine-tuning of large language models through flatness. *arXiv preprint arXiv:2403.04124*, 2024.
- Chowdhury, S. R., Zhou, X., and Natarajan, N. Differentially private reward estimation with preference feedback. In *International Conference on Artificial Intelligence and Statistics*, pp. 4843–4851. PMLR, 2024.
- Denton, R., Díaz, M., Kivlichan, I., Prabhakaran, V., and Rosen, R. Whose ground truth? accounting for individual and collective identities underlying dataset annotation. *arXiv preprint arXiv:2112.04554*, 2021.
- Duan, H., Dziedzic, A., Papernot, N., and Boenisch, F. Flocks of stochastic parrots: Differentially private prompt learning for large language models. *Advances in Neural Information Processing Systems*, 36:76852–76871, 2023.
- Dwork, C., Kenthapadi, K., McSherry, F., Mironov, I., and Naor, M. Our data, ourselves: Privacy via distributed noise generation. In *Advances in Cryptology-EUROCRYPT 2006: 24th Annual International Conference on the Theory and Applications of Cryptographic Techniques, St. Petersburg, Russia, May 28-June 1, 2006. Proceedings 25*, pp. 486–503. Springer, 2006a.
- Dwork, C., McSherry, F., Nissim, K., and Smith, A. Calibrating noise to sensitivity in private data analysis. In *Theory of Cryptography: Third Theory of Cryptography Conference, TCC 2006, New York, NY, USA, March 4-7, 2006. Proceedings 3*, pp. 265–284. Springer, 2006b.
- Dwork, C., Roth, A., et al. The algorithmic foundations of differential privacy. *Foundations and Trends® in Theoretical Computer Science*, 9(3–4):211–407, 2014.
- Feng, Q., Kasa, S. R., KASA, S. K., Yun, H., Teo, C. H., and Bodapati, S. B. Exposing privacy gaps: Membership inference attack on preference data for llm alignment. In *International Conference on Artificial Intelligence and Statistics*, pp. 5221–5229. PMLR, 2025.
- Gao, F., Zhou, R., Wang, T., Shen, C., and Yang, J. Data-adaptive differentially private prompt synthesis for in-context learning. *arXiv preprint arXiv:2410.12085*, 2024.
- Ghazi, B., Golowich, N., Kumar, R., Manurangsi, P., and Zhang, C. Deep learning with label differential privacy. *Advances in neural information processing systems*, 34: 27131–27145, 2021.
- Gopi, S., Lee, Y. T., and Wutschitz, L. Numerical composition of differential privacy. *Advances in Neural Information Processing Systems*, 34:11631–11642, 2021.
- He, J., Li, X., Yu, D., Zhang, H., Kulkarni, J., Lee, Y. T., Backurs, A., Yu, N., and Bian, J. Exploring the limits of differentially private deep learning with group-wise clipping. *arXiv preprint arXiv:2212.01539*, 2022.
- Holohan, N., Braghin, S., Mac Aonghusa, P., and Levacher, K. Diffprivlib: the IBM differential privacy library. *ArXiv e-prints*, 1907.02444 [cs.CR], July 2019.
- Hong, J., Wang, J. T., Zhang, C., Li, Z., Li, B., and Wang, Z. Dp-opt: Make large language model your privacy-preserving prompt engineer. *arXiv preprint arXiv:2312.03724*, 2023.
- Hou, C., Shrivastava, A., Zhan, H., Conway, R., Le, T., Sagar, A., Fanti, G., and Lazar, D. Pre-text: Training language models on private federated data in the age of llms. In *International Conference on Machine Learning*, pp. 19043–19061. PMLR, 2024.

- Jang, J., Kim, S., Lin, B. Y., Wang, Y., Hessel, J., Zettlemoyer, L., Hajishirzi, H., Choi, Y., and Ammanabrolu, P. Personalized soups: Personalized large language model alignment via post-hoc parameter merging. *arXiv preprint arXiv:2310.11564*, 2023.
- Ji, J., Liu, M., Dai, J., Pan, X., Zhang, C., Bian, C., Chen, B., Sun, R., Wang, Y., and Yang, Y. Beavertails: Towards improved safety alignment of llm via a human-preference dataset. *Advances in Neural Information Processing Systems*, 36:24678–24704, 2023.
- Ji, J., Hong, D., Zhang, B., Chen, B., Dai, J., Zheng, B., Qiu, T., Li, B., and Yang, Y. Pku-saferllhf: Towards multi-level safety alignment for llms with human preference. *arXiv preprint arXiv:2406.15513*, 2024.
- Kirk, R., Mediratta, I., Nalmpantis, C., Luketina, J., Hambro, E., Grefenstette, E., and Raileanu, R. Understanding the effects of rlhf on llm generalisation and diversity. In *International Conference on Learning Representations*, volume 2024, pp. 20620–20653, 2024.
- Kong, D. and Yang, L. Provably feedback-efficient reinforcement learning via active reward learning. *Advances in Neural Information Processing Systems*, 35:11063–11078, 2022.
- Köpf, A., Kilcher, Y., Von Rütte, D., Anagnostidis, S., Tam, Z. R., Stevens, K., Barhoum, A., Nguyen, D., Stanley, O., Nagyfi, R., et al. Openassistant conversations-democratizing large language model alignment. *Advances in Neural Information Processing Systems*, 36:47669–47681, 2023.
- Kovač, G., Sawayama, M., Portelas, R., Colas, C., Dominey, P. F., and Oudeyer, P.-Y. Large language models as superpositions of cultural perspectives. *arXiv preprint arXiv:2307.07870*, 2023.
- Lee, H., Phatale, S., Mansoor, H., Mesnard, T., Ferret, J., Lu, K. R., Bishop, C., Hall, E., Carbune, V., Rastogi, A., et al. Rlaif vs. rlhf: Scaling reinforcement learning from human feedback with ai feedback. In *International Conference on Machine Learning*, pp. 26874–26901. PMLR, 2024a.
- Lee, Y., Williams, J., Marklund, H., Sharma, A., Mitchell, E., Singh, A., and Finn, C. Test-time alignment via hypothesis reweighting. *arXiv preprint arXiv:2412.08812*, 2024b.
- Li, H., Guo, D., Fan, W., Xu, M., Huang, J., Meng, F., and Song, Y. Multi-step jailbreaking privacy attacks on chatgpt. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pp. 4138–4153, 2023.
- Li, X., Tramer, F., Liang, P., and Hashimoto, T. Large language models can be strong differentially private learners. *arXiv preprint arXiv:2110.05679*, 2021.
- Lin, Z., Gopi, S., Kulkarni, J., Nori, H., and Yekhanin, S. Differentially private synthetic data via foundation model apis 1: Images. *arXiv preprint arXiv:2305.15560*, 2023.
- Liu, T., Zhao, Y., Joshi, R., Khalman, M., Saleh, M., Liu, P. J., and Liu, J. Statistical rejection sampling improves preference optimization. *arXiv preprint arXiv:2309.06657*, 2023.
- Liu, X., Kong, W., Jain, P., and Oh, S. Dp-pca: Statistically optimal and differentially private pca. *Advances in neural information processing systems*, 35:29929–29943, 2022.
- Mattern, J., Jin, Z., Weggenmann, B., Schoelkopf, B., and Sachan, M. Differentially private language models for secure data sharing. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pp. 4860–4873, 2022.
- Mireshghallah, F., Su, Y., Hashimoto, T. B., Eisner, J., and Shin, R. Privacy-preserving domain adaptation of semantic parsers. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 4950–4970, 2023.
- Narayan, S., Cohen, S. B., and Lapata, M. Don’t give me the details, just the summary! topic-aware convolutional neural networks for extreme summarization. *arXiv preprint arXiv:1808.08745*, 2018.
- Ochs, S. and Habernal, I. Private synthetic text generation with diffusion models. *arXiv preprint arXiv:2410.22971*, 2024.
- Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744, 2022.
- Poddar, S., Wan, Y., Ivison, H., Gupta, A., and Jaques, N. Personalizing reinforcement learning from human feedback with variational preference learning. *arXiv preprint arXiv:2408.10075*, 2024.
- Rafailov, R., Sharma, A., Mitchell, E., Manning, C. D., Ermon, S., and Finn, C. Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36, 2024.
- Rame, A., Couaaron, G., Dancette, C., Gaya, J.-B., Shukor, M., Soulier, L., and Cord, M. Rewarded soups: towards pareto-optimal alignment by interpolating weights fine-tuned on diverse rewards. *Advances in Neural Information Processing Systems*, 36:71095–71134, 2023.

- Saha, A., Pacchiano, A., and Lee, J. Dueling rl: Reinforcement learning with trajectory preferences. In *International conference on artificial intelligence and statistics*, pp. 6263–6289. PMLR, 2023.
- Sandri, M., Leonardelli, E., Tonelli, S., and Ježek, E. Why don’t you do it right? analysing annotators’ disagreement in subjective tasks. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics*, pp. 2428–2441, 2023.
- Sanh, V., Debut, L., Chaumond, J., and Wolf, T. Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*, 2019.
- Shao, Z., Wang, P., Zhu, Q., Xu, R., Song, J., Bi, X., Zhang, H., Zhang, M., Li, Y., Wu, Y., et al. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.
- Shen, B., Zhang, J., Chen, T., Zan, D., Geng, B., Fu, A., Zeng, M., Yu, A., Ji, J., Zhao, J., et al. Pangu-coder2: Boosting large language models for code with ranking feedback. *arXiv preprint arXiv:2307.14936*, 2023.
- Shoemate, M., Vyrros, A., McCallum, C., Prasad, R., Durbin, P., Casacuberta Puig, S., Cowan, E., Xu, V., Ratliff, Z., Berrios, N., Whitworth, A., Eliot, M., Lebeda, C., Renard, O., and McKay Bowen, C. OpenDP Library, 2021. URL <https://github.com/opendp/opendp>.
- Shokri, R., Stronati, M., Song, C., and Shmatikov, V. Membership inference attacks against machine learning models. In *2017 IEEE symposium on security and privacy (SP)*, pp. 3–18. IEEE, 2017.
- Singh, A., Hsu, S., Hsu, K., Mitchell, E., Ermon, S., Hashimoto, T., Sharma, A., and Finn, C. Fspo: Few-shot preference optimization of synthetic preference data in llms elicits effective personalization to real users. *arXiv preprint arXiv:2502.19312*, 2025.
- Song, K., Tan, X., Qin, T., Lu, J., and Liu, T.-Y. MpNet: Masked and permuted pre-training for language understanding. *Advances in neural information processing systems*, 33:16857–16867, 2020.
- Stiennon, N., Ouyang, L., Wu, J., Ziegler, D., Lowe, R., Voss, C., Radford, A., Amodei, D., and Christiano, P. F. Learning to summarize with human feedback. *Advances in neural information processing systems*, 33:3008–3021, 2020.
- Su, D., Cao, J., Li, N., Bertino, E., and Jin, H. Differentially private k-means clustering. In *Proceedings of the sixth ACM conference on data and application security and privacy*, pp. 26–37, 2016.
- Tan, B., Xu, Z., Xing, E., Hu, Z., and Wu, S. Synthesizing privacy-preserving text data via finetuning without finetuning billion-scale llms. *arXiv preprint arXiv:2503.12347*, 2025.
- Tang, X., Shin, R., Inan, H. A., Manoel, A., Miresghallah, F., Lin, Z., Gopi, S., Kulkarni, J., and Sim, R. Privacy-preserving in-context learning with differentially private few-shot generation. *arXiv preprint arXiv:2309.11765*, 2023.
- Tang, X., Panda, A., Nasr, M., Mahloujifar, S., and Mittal, P. Private fine-tuning of large language models with zeroth-order optimization. *arXiv preprint arXiv:2401.04343*, 2024.
- Taori, R., Gulrajani, I., Zhang, T., Dubois, Y., Li, X., Guestrin, C., Liang, P., and Hashimoto, T. B. Alpaca: A strong, replicable instruction-following model. *Stanford Center for Research on Foundation Models*. <https://crfm.stanford.edu/2023/03/13/alpaca.html>, 3(6):7, 2023.
- Touvron, H., Lavril, T., Izacard, G., Martinet, X., Lachaux, M.-A., Lacroix, T., Rozière, B., Goyal, N., Hambro, E., Azhar, F., et al. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- Vogels, E. A. *The state of online harassment*, volume 13. Pew Research Center Washington, DC, 2021.
- Wang, W., Liang, X., Ye, R., Chai, J., Chen, S., and Wang, Y. Knowledgesg: Privacy-preserving synthetic text generation with knowledge distillation from server. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 7677–7695, 2024.
- Wu, F., Inan, H. A., Backurs, A., Chandrasekaran, V., Kulkarni, J., and Sim, R. Privately aligning language models with reinforcement learning. *arXiv preprint arXiv:2310.16960*, 2023a.
- Wu, S., Xu, Z., Zhang, Y., Zhang, Y., and Ramage, D. Prompt public large language models to synthesize data for private on-device applications. *arXiv preprint arXiv:2404.04360*, 2024.
- Wu, T., Panda, A., Wang, J. T., and Mittal, P. Privacy-preserving in-context learning for large language models. *arXiv e-prints*, pp. arXiv–2305, 2023b.
- Xiao, S., Liu, Z., Zhang, P., and Muennighoff, N. C-pack: Packaged resources to advance general chinese embedding, 2023.
- Xie, C., Lin, Z., Backurs, A., Gopi, S., Yu, D., Inan, H. A., Nori, H., Jiang, H., Zhang, H., Lee, Y. T., et al. Differentially private synthetic data via foundation model apis 2: Text. In *International Conference on Machine Learning*, pp. 54531–54560. PMLR, 2024.

- Xiong, W., Dong, H., Ye, C., Wang, Z., Zhong, H., Ji, H., Jiang, N., and Zhang, T. Iterative preference learning from human feedback: Bridging theory and practice for rlhf under kl-constraint. *arXiv preprint arXiv:2312.11456*, 2023.
- Yousefpour, A., Shilov, I., Sablayrolles, A., Testuggine, D., Prasad, K., Malek, M., Nguyen, J., Ghosh, S., Bharadwaj, A., Zhao, J., et al. Opacus: User-friendly differential privacy library in pytorch. *arXiv preprint arXiv:2109.12298*, 2021.
- Yu, D., Naik, S., Backurs, A., Gopi, S., Inan, H. A., Kamath, G., Kulkarni, J., Lee, Y. T., Manoel, A., Wutschitz, L., et al. Differentially private fine-tuning of language models. *arXiv preprint arXiv:2110.06500*, 2021.
- Yu, D., Kairouz, P., Oh, S., and Xu, Z. Privacy-preserving instructions for aligning large language models. In *International Conference on Machine Learning*, pp. 57480–57506. PMLR, 2024.
- Yue, X., Inan, H., Li, X., Kumar, G., McAnallen, J., Shajari, H., Sun, H., Levitan, D., and Sim, R. Synthetic text generation with differential privacy: A simple and practical recipe. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 1321–1342, 2023.
- Zhang, J., Lei, M., Ding, M., Li, M., Xiang, Z., Xu, D., Xu, J., and Wang, D. Towards user-level private reinforcement learning with human feedback. *arXiv preprint arXiv:2502.17515*, 2025.
- Zhang, L., Thekumparampil, K. K., Oh, S., and He, N. Dpzero: dimension-independent and differentially private zeroth-order optimization. In *International Workshop on Federated Learning in the Age of Foundation Models in Conjunction with NeurIPS 2023*, 2023.
- Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E., et al. Judging llm-as-a-judge with mt-bench and chatbot arena. *Advances in neural information processing systems*, 36: 46595–46623, 2023.
- Zhu, B., Jordan, M., and Jiao, J. Principled reinforcement learning with human feedback from pairwise or k-wise comparisons. In *International Conference on Machine Learning*, pp. 43037–43067. PMLR, 2023.
- Ziegler, D. M., Stiennon, N., Wu, J., Brown, T. B., Radford, A., Amodei, D., Christiano, P., and Irving, G. Fine-tuning language models from human preferences. *arXiv preprint arXiv:1909.08593*, 2019.

A. Privacy Analysis

In this section we present the complete privacy analysis of Algorithm 1.

Theorem A.1. *Algorithm 1 is (ε, δ) -differentially private.*

We first recall the concepts and theorems from the literature on which the analysis relies.

Many DP mechanisms, the Gaussian mechanism among them, furnish an entire family of (ε, δ) -DP guarantees: for each fixed ε there exists a $\delta(\varepsilon)$ such that the mechanism satisfies $(\varepsilon, \delta(\varepsilon))$ -DP.

Definition A.2 (Privacy curve). A DP algorithm $\mathcal{M}$ is said to have a *privacy curve* $\delta : \mathbb{R} \rightarrow [0, 1]$ if, for every $\varepsilon > 0$, the algorithm $\mathcal{M}$ satisfies $(\varepsilon, \delta(\varepsilon))$ -DP.

A principal advantage of privacy curves is that they yield tighter composition guarantees than advanced composition theorems (Dwork et al., 2014): they admit numerical composition, which attains the tightest available guarantees. Gopi et al. (2021) derive privacy curves for the composition of several standard mechanisms, including the Gaussian and subsampled Gaussian mechanisms.

Theorem A.3 (Gopi et al. (2021)). *Suppose $M_1, M_2, \dots, M_k$ are DP algorithms. Then the privacy curve $\delta_M(\varepsilon)$ of adaptive composition $M = M_1 \circ M_2 \circ \dots \circ M_k$ can be approximated in time*

$$O\left(\frac{\varepsilon_{\text{upper}} k^{1/2} \log k \sqrt{\log(1/\delta_{\text{error}})}}{\varepsilon_{\text{error}}}\right)$$

where $\varepsilon_{\text{error}}$ is the additive error in ε , δ_{error} is the additive error in δ , and $\varepsilon_{\text{upper}}$ is an upper bound on

$$\max\left\{\varepsilon_M(\delta_{\text{error}}), \max_i \varepsilon_{M_i}\left(\frac{\delta_{\text{error}}}{k}\right)\right\}.$$

Gopi et al. (2021) further characterize the privacy loss of a subsampled mechanism in terms of that of the original mechanism, which in turn bounds the privacy loss of DP-SGD.

Theorem A.4 (Gopi et al. (2021)). *Let (X, Y) be the PRVs for a privacy curve $\delta(P \parallel Q)$. Let (X_p, Y_p) be the PRVs for $\delta_p = \delta(P \parallel p \cdot P + (1-p) \cdot Q)$ for some sampling probability $p \in [0, 1]$. Then*

$$X_p = \log(1 + p(e^X - 1)), \quad Y_p = \begin{cases} \log(1 + p(e^Y - 1)) & \text{w.p. } p \\ \log(1 + p(e^X - 1)) & \text{w.p. } 1 - p. \end{cases}$$

The CDFs of X_p and Y_p are given by:

$$\text{CDF}_{X_p}(t) = \begin{cases} \text{CDF}_X\left(\log\left(\frac{e^t - (1-p)}{p}\right)\right) & \text{if } t \geq \log(1-p) \\ 0 & \text{if } t < \log(1-p). \end{cases}$$

$$\text{CDF}_{Y_p}(t) = \begin{cases} p \cdot \text{CDF}_Y\left(\log\left(\frac{e^t - (1-p)}{p}\right)\right) + (1-p) \cdot \text{CDF}_X\left(\log\left(\frac{e^t - (1-p)}{p}\right)\right) & \text{if } t \geq \log(1-p) \\ 0 & \text{if } t < \log(1-p). \end{cases}$$

We are now in a position to carry out the privacy analysis of our algorithm.

Proof. The algorithm comprises exactly three subroutines that access the private dataset: the DP-PCA subroutine satisfying ε_0 -DP (Amin et al., 2019) in Line 5, the DP-KMeans subroutine satisfying ε_1 -DP (Su et al., 2016) in Line 7, and DP-SGD with noise multiplier σ_2 in Line 9. We allocate ε_0 to DP-PCA and ε_1 to DP-KMeans, and devote the remaining budget to DP-SGD, whose per-step costs we compose with the Privacy Random Variable accountant (Gopi et al., 2021).

Line 7 partitions the low-dimensional embeddings $\{z_i\}_{i=1}^n$ into K disjoint groups $\mathcal{C}_1, \dots, \mathcal{C}_K$ under DP, and Line 9 trains a separate linear model on each group with DP-SGD. Because the groups are disjoint, the parallel-composition property of DP bounds the overall privacy cost of this step by that of the smallest cluster used for training; to guarantee a known

lower bound on the per-model sample size, we discard any cluster with fewer than $|\mathcal{D}_{\text{priv}}|/(K+4)$ samples. We set the DP-SGD batch size to 4, the number of epochs to 4, and the gradient clipping norm to 1.0. Each DP-SGD iteration incurs the privacy loss of the subsampled Gaussian mechanism (Theorem A.4), and composing these iterations via Theorem A.3 (Gopi et al., 2021) ensures that reward-model training satisfies $(\varepsilon - \varepsilon_0 - \varepsilon_1, \delta)$ -DP. The DP parameters ε_0 , ε_1 , and σ_2 used in our experiments are reported in Table 7.

Composing these guarantees, DP-PCA (ε_0), DP-KMeans (ε_1), and DP-SGD ($\varepsilon - \varepsilon_0 - \varepsilon_1$) together yield an (ε, δ) -DP mechanism, while all remaining steps are post-processing of DP outputs and add no cost. We provide the privacy-accounting code below.

```
def get_privacy_spent(sampling_prob_dpsgd, running_steps_dpsgd,
                    noise_multiplier_dpsgd, eps_pca, eps_cluster, delta):

    prv_dpsgd = PoissonSubsampledGaussianMechanism(
        noise_multiplier=noise_multiplier_dpsgd,
        sampling_probability=sampling_prob_dpsgd,
    )

    accountant = PRVAccountant(
        prvs=[prv_dpsgd],
        max_self_compositions=[running_steps_dpsgd],
        eps_error=0.01,
        delta_error=delta/10,
    )

    eps_lower, eps_estimate, eps_upper = accountant.compute_epsilon(
        delta=delta,
        num_self_compositions=[running_steps_dpsgd],
    )

    return eps_upper + eps_pca + eps_cluster
```

□

B. Experimental Supplementary

B.1. Computational Resources for Experiments

All experiments are conducted on NVIDIA H100 GPUs, each with 80 GB of memory. DPPrefSyn itself runs on a single NVIDIA H100 GPU, requiring several hours to complete depending on the dataset size.

B.2. Datasets

This section details the datasets used in our experiments.

- **OpenAssistant.** The OpenAssistant Conversations dataset (Köpf et al., 2023) is a human-generated, human-annotated assistant-style conversation corpus comprising 161,443 messages in 35 languages, together with 461,292 quality ratings and over 10,000 fully annotated conversation trees. We use only the English subset, consisting of 14,167 training examples and 712 test examples.
- **Anthropic-HH.** The Anthropic Helpful and Harmless dataset (Bai et al., 2022) supplies human preference annotations for training RLHF reward models. Its helpfulness split spans three tranches—outputs from base models (context-distilled 52B), samples filtered via rejection sampling under a preference model, and data collected during online training—while its harmlessness data are drawn solely from base models in the same format. The dataset comprises 161K training and 8.55K test comparisons; we sample 1,000 test examples for evaluation.
- **TL;DR.** The Summarize from Feedback dataset (Stiennon et al., 2020) provides human preference annotations over pairs of summaries, with annotators asked to select the better of the two; it was originally introduced to train a reward model for aligning summarization with human preferences. The summaries derive from the TL;DR dataset, with additional validation data drawn from CNN and Daily Mail articles. The dataset comprises 92.9K training and 86.1K validation comparisons; we sample 979 examples from the validation split as our test set.

- **Alpaca.** The Alpaca dataset (Taori et al., 2023) comprises 52K instruction-following examples generated by self-instruct with OpenAI’s text-davinci-003, each pairing an instruction with a corresponding response to improve instruction-following in language models. We use its instructions as public prompts for data synthesis.
- **SafeRLHF.** The PKU-SafeRLHF dataset (Ji et al., 2024) is a human-labeled dataset capturing both performance and safety preferences, with constraints spanning more than ten dimensions—including insults, immorality, crime, emotional harm, and privacy—for fine-grained value alignment in RLHF. We use its 73.9K training-split prompts as public prompts for data synthesis.
- **XSum.** The Extreme Summarization (XSum) dataset (Narayan et al., 2018) targets abstractive single-document summarization, framing the task as generating a one-sentence summary that answers “What is the article about?” It comprises 226,711 BBC news articles from 2010 to 2017, each paired with a concise summary, and spans diverse domains including politics, sports, business, science, and health, with 204,045 training, 11,332 validation, and 11,334 test examples. We use its training-split documents as public prompts for data synthesis.

B.3. Implementation Details and Hyperparameter Settings

This section details the experimental setups, hyperparameters, evaluation metrics, and baselines.

Implementation Details of DPPrefSyn. For each preference sample (x_i, a_i^+, a_i^-) , DPPrefSyn concatenates the prompt with each response to form the two texts $[x_i; a_i^+]$ and $[x_i; a_i^-]$. The prompt–response concatenation format used in our experiments is shown below.

OpenAssistant Prompt Response Concatenation Format:

Human: <question>\n\nAssistant: <answer>

Anthropic-HH Prompt Response Concatenation Format:

Human: <question>\n\nAssistant: <answer>

TL;DR Prompt Response Concatenation Format:

<document>\n\nSummary: <summary>

Unless otherwise stated, we instantiate the embedding model ψ with BAAI/bge-large-en-v1.5 (Xiao et al., 2023). We adopt the DP-PCA implementation of Shoemate et al. (2021), which realizes the algorithm of Amin et al. (2019), and the DP-KMeans implementation of Holohan et al. (2019), which realizes the algorithm of Su et al. (2016). We set the DP-PCA projected dimension to $p = 20$ and the number of DP-KMeans clusters to $K = 5$. When training the reward models with DP-SGD, we discard any cluster with fewer than $|\mathcal{D}_{\text{priv}}|/(K + 4)$ samples to guarantee a known lower bound on the per-model sample size. DP-SGD uses the SGD optimizer with a learning rate of 0.1, a batch size of 4, 4 epochs, and a gradient clipping norm of 1.0, and is implemented with the Opacus library (Yousefpour et al., 2021).

Unless otherwise stated, we generate candidate responses for each public prompt with the instruction-finetuned LLaMA-7B-chat model (Touvron et al., 2023), drawing $L = 5$ candidates per prompt at a generation temperature of 0.9. The prompt format used for generating multiple completions is shown below.

Prompt format used to generate candidate responses for OpenAssistant QA task:

Human: <query>\n\nAssistant:

Prompt format used to generate candidate responses for Anthropic-HH QA task:

Human: <query>\n\nAssistant:

Prompt format used to generate candidate responses for TL;DR summarization task:

Summarize the following article in a paragraph of 50 words or less:
<article>\n\nAssistant:

Privacy Parameters in DPPrefSyn. Throughout, we set $\delta = 1/|\mathcal{D}_{\text{priv}}|$; the values of $|\mathcal{D}_{\text{priv}}|$ are listed in Table 6, and the DP parameters ε_0 , ε_1 , and σ_2 in Table 7.

Table 6. Sizes $|\mathcal{D}_{\text{priv}}|$ of the private datasets used in our experiments.

$\mathcal{D}_{\text{priv}}$	OpenAssistant	Anthropic-HH	TL;DR
$ \mathcal{D}_{\text{priv}} $	14,167	160,800	92,858

Table 7. DP parameters ε_0 , ε_1 , and σ_2 used in the DPPrefSyn experiments, reported per task for each overall budget $\varepsilon = 1, 2, 4, 8$.

Task	ε_0 for $\varepsilon = 1, 2, 4, 8$	ε_1 for $\varepsilon = 1, 2, 4, 8$	σ_2 for $\varepsilon = 1, 2, 4, 8$
OpenAssistant	[0.125, 0.25, 0.5, 1]	[0.125, 0.25, 0.5, 1]	[0.808, 0.671, 0.566, 0.471]
Anthropic-HH	[0.125, 0.25, 0.5, 1]	[0.125, 0.25, 0.5, 1]	[0.620, 0.556, 0.487, 0.412]
TL;DR	[0.125, 0.25, 0.5, 1]	[0.125, 0.25, 0.5, 1]	[0.647, 0.575, 0.501, 0.422]

Evaluation Details. We first perform supervised fine-tuning (SFT), taking the preferred response of each pair as the training target, and then apply DPO (Rafailov et al., 2024) to further fine-tune the SFT model on the preference pairs. All models are trained in bfloat16 precision. Following Rafailov et al. (2024), SFT uses a learning rate of 5×10^{-7} for 1 epoch, and DPO uses $\beta = 0.1$, a learning rate of 1×10^{-6} , and 3 epochs. The batch size is 64 for Pythia-2.8B, 32 for Qwen-3-4B and Gemma-2-2B, and 16 for Llama-3-8B.

Our utility metric is the win rate, which quantifies how often the fine-tuned model’s response is preferred over a reference response. For each test prompt we sample a response from the fine-tuned model at temperature 0.0 and have GPT-4o adjudicate it against the preferred response in the test set; the win rate is the fraction of cases in which GPT-4o judges the model’s response superior.

We adopt the win rate because it follows the standard evaluation protocol in recent preference-alignment research (Rafailov et al., 2024; Ouyang et al., 2022; Kirk et al., 2024; Lee et al., 2024a), in which a strong LLM serves as an automated judge of human preference. Prior work establishes that high-quality frontier models yield reliable, reproducible preference judgments strongly correlated with human annotations, while offering greater scalability and consistency than human studies (Zheng et al., 2023). We use GPT-4o as the judge because it was the most advanced model available at the time of our experiments. The prompt formats used to elicit the win rates, following Rafailov et al. (2024), are shown below.

Prompt format used to generate win rates for OpenAssistant QA task

For the following query to a chatbot, which response is more helpful?

Query: <the user query>

Response A: <Response A>

Response B: <Response B>

FIRST provide a one-sentence comparison of the two responses and explain which you feel is more helpful. SECOND, on a new line, state only “A” or “B” to indicate which response is more helpful. Your response should use the format:

Comparison: <one-sentence comparison and explanation>

More helpful: <“A” or “B”>

Prompt format used to generate win rates for Anthropic-HH QA task

For the following query to a chatbot, which response is more helpful and harmless?

Query: <the user query>

Response A: <Response A>

Response B: <Response B>

FIRST provide a one-sentence comparison of the two responses and explain which you feel is more helpful. SECOND, on a new line, state only “A” or “B” to indicate which response is more helpful and harmless. Your response should use the format:

Comparison: <one-sentence comparison and explanation>

More helpful and harmless: <“A” or “B”>

Prompt format used to generate win rates for TL;DR summarization task

Which of the following summaries does a better job of summarizing the most important points in the given forum post?

Post: <post>

Summary A: <Summary A>

Summary B: <Summary B>

FIRST provide a one-sentence comparison of the two summaries, explaining which you prefer and why. SECOND, on a new line, state only “A” or “B” to indicate your choice. Your response should use the format:

Comparison: <one-sentence comparison and explanation>

Preferred: <“A” or “B”>

DP-FT Details. For DP-FT we implement DP-SGD in Opacus (Yousefpour et al., 2021) with FSDP2 and Ghost Clipping. To simplify privacy accounting, we split the private dataset evenly into two disjoint halves so that each user contributes to either SFT or DPO but not both, and track privacy loss with the PRV accountant. A grid search over clipping norms $\{0.05, 0.1, 0.5, 1.0\}$ selects 0.1 for the final experiments. Learning rates are searched over $\{5 \times 10^{-7}, 1 \times 10^{-6}, 5 \times 10^{-6}, 1 \times 10^{-5}, 5 \times 10^{-5}, 1 \times 10^{-4}\}$, yielding 5×10^{-5} for SFT and 1×10^{-5} for DPO on Pythia-2.8B, Qwen-3-4B, and Gemma-2-2B, and 5×10^{-6} for Llama-3-8B. The batch size is 16 for Pythia-2.8B, 8 for Qwen-3-4B and Gemma-2-2B, and 4 for Llama-3-8B, and we train for 2 epochs in SFT and 3 epochs in DPO.

B.4. Empirical Privacy Evaluation by Membership Inference Attacks

Although DP furnishes a theoretical guarantee against privacy leakage, complementing it with an empirical assessment is important. To this end, we mount membership inference attacks (MIAs) (Shokri et al., 2017) against the DPO-fine-tuned downstream models on OpenAssistant and report the AUC in Table 5.

Following Feng et al. (2025), who instantiate an MIA tailored to preference data for LLM alignment, the attacker aims to decide whether a private sample (x, a^+, a^-) was used in preference alignment. We designate 1,000 randomly sampled training examples as member data and 1,000 held-out test examples as non-member data. DP-FT fails to adequately mitigate the attack, whereas DPPrefSyn consistently drives the AUC down to near random guessing. This markedly stronger empirical robustness follows from the synthetic nature of the fine-tuning data, which, containing no real private sample, removes the membership signal that overfitting to real data would otherwise expose.

B.5. Ablation Studies on Hyperparameters

In this section we ablate the embedding model ψ , the DP-PCA dimension p , the number of candidates per prompt L , the filter threshold, and the privacy-budget allocation, all on OpenAssistant at $\varepsilon = 4$.

Varying the embedding model ψ . We evaluate the sentence encoders BAAI/bge-large-en-v1.5 (Xiao et al., 2023), sentence-transformers/all-mpnet-base-v2 (Song et al., 2020), and sentence-transformers/all-distilroberta-v1 (Sanh et al., 2019). As Table 8 shows, while embedding quality does modulate performance, DPPrefSyn remains effective across all three encoders.

Table 8. GPT-4o win rates (%) of DPPrefSyn under different embedding models ψ on OpenAssistant at $\varepsilon = 4$, reported after SFT and after SFT+DPO (mean and standard deviation over 5 seeds).

Embedding model ψ	bge-large-en-v1.5	all-mpnet-base-v2	all-distilroberta-v1
SFT	9.33 _{0.56}	9.80 _{0.29}	10.20 _{0.53}
SFT+DPO	11.85 _{0.99}	13.03 _{0.87}	12.22 _{0.70}

Varying the DP-PCA dimension p . We sweep $p = 10, 20, 30$ in Table 9. The intermediate value $p = 20$ yields the best overall post-DPO performance, making it a sound default.

Table 9. GPT-4o win rates (%) of DPPrefSyn under varying DP-PCA dimension p on OpenAssistant at $\varepsilon = 4$, reported after SFT and after SFT+DPO (mean and standard deviation over 5 seeds).

DP-PCA dimension p	$p = 10$	$p = 20$	$p = 30$
SFT	9.77 _{0.56}	9.33 _{0.56}	9.16 _{0.32}
SFT+DPO	11.01 _{0.71}	11.85 _{0.99}	11.49 _{0.53}

Varying the number of candidates per prompt L . We sweep $L = 3, 5, 8$ in Table 10. Performance improves as L rises from 3 to 5 but does not increase further at $L = 8$, identifying $L = 5$ as a good default.

Table 10. GPT-4o win rates (%) of DPPrefSyn under varying number of candidates per prompt L on OpenAssistant at $\varepsilon = 4$, reported after SFT and after SFT+DPO (mean and standard deviation over 5 seeds).

Number of candidates per prompt L	$L = 3$	$L = 5$	$L = 8$
SFT	9.05 _{0.48}	9.33 _{0.56}	9.89 _{0.34}
SFT+DPO	10.93 _{0.50}	11.85 _{0.99}	11.29 _{0.49}

Varying the filter threshold. We sweep thresholds of 0, 0.5, and 1 in Table 11. A moderate threshold of 0.5 slightly outperforms no filtering by removing noisy or uninformative preference pairs, whereas an aggressive threshold of 1 over-prunes the data, curtailing coverage and degrading performance.

Varying the privacy-budget allocation. We evaluate alternative splits of the budget between DP-PCA (ε_0) and DP-KMeans (ε_1). The default setting allocates $\varepsilon_0 = \varepsilon_1 = 0.5$ under a total budget of $\varepsilon = 4$; we additionally test (i) favoring DP-PCA ($\varepsilon_0 = 0.7, \varepsilon_1 = 0.3$) and (ii) favoring DP-KMeans ($\varepsilon_0 = 0.3, \varepsilon_1 = 0.7$). As Table 12 shows, performance is essentially unchanged across these splits, confirming that DPPrefSyn is robust to reasonable reallocations of the budget.

B.6. Discussion on the Linear Reward Models in DPPrefSyn

In this section we justify the choice of linear reward models in DPPrefSyn (Line 9). This choice instantiates the Bradley–Terry model with a linear reward function, a widely adopted assumption in the preference-alignment literature (Saha et al., 2023;

Table 11. GPT-4o win rates (%) of DPPrefSyn under varying filter threshold on OpenAssistant at $\varepsilon = 4$, reported after SFT and after SFT+DPO (mean and standard deviation over 5 seeds).

Filter threshold	0	0.5	1
SFT	9.24 _{0.46}	9.33 _{0.56}	8.71 _{0.58}
SFT+DPO	11.26 _{0.83}	11.85 _{0.99}	9.16 _{0.63}

Table 12. GPT-4o win rates (%) of DPPrefSyn under varying privacy-budget allocations ($\varepsilon_0, \varepsilon_1$) on OpenAssistant at $\varepsilon = 4$, reported after SFT and after SFT+DPO.

	$\varepsilon_0 = 0.7, \varepsilon_1 = 0.3$	$\varepsilon_0 = 0.5, \varepsilon_1 = 0.5$	$\varepsilon_0 = 0.3, \varepsilon_1 = 0.7$
SFT	9.16	9.33	9.44
SFT+DPO	11.91	11.85	11.55

Kong & Yang, 2022; Zhu et al., 2023; Xiong et al., 2023) that we formalize in Section 3.

To assess its validity, we compare our linear model against a fully fine-tuned GPT-2 reward model, both trained on Anthropic-HH. For each test triple (prompt, preferred response, dispreferred response), we feed both responses—each concatenated with the prompt—through the trained reward model, which assigns a scalar to each, and count the prediction correct when the preferred response scores higher than the dispreferred one; accuracy is the fraction of correctly ordered triples. The linear model attains 66% accuracy against the fully fine-tuned GPT-2 model’s 62%, demonstrating that the linear reward structure captures preference signals effectively and remains competitive despite its far greater simplicity.

B.7. Ablation Studies on Prompt and Response Generator Quality

In this section we isolate the effects of prompt quality and response quality on DPPrefSyn, all on OpenAssistant at $\varepsilon = 4$.

Prompt Quality. To test whether DPPrefSyn’s gains stem from the higher quality of its public prompts, we rerun DPPrefSyn with the OpenAssistant prompts as its public prompt set in place of Alpaca, while the DP-FT baseline retains its original private OpenAssistant dataset—so that both methods draw on *the same set of prompts*. The results appear in Table 13.

Table 13. Prompt-quality ablation on OpenAssistant at $\varepsilon = 4$: GPT-4o win rate (%) after SFT+DPO when DP-FT and DPPrefSyn share the same OpenAssistant prompts, alongside DPPrefSyn with its original Alpaca prompts.

	DP-FT (OpenAssistant prompts)	DPPrefSyn (OpenAssistant prompts)	DPPrefSyn (Alpaca prompts)
Win Rate (%)	6.12	10.39	11.85

Even under identical OpenAssistant prompts, DPPrefSyn continues to outperform DP-FT, confirming that its gains do not stem solely from higher-quality public prompts.

Candidate Response Quality. To test whether DPPrefSyn’s gains stem from using a stronger response generator, we construct an enhanced DP-FT baseline. We first generate candidate responses for each prompt with *the same LLaMA-7B-Chat generator* used by DPPrefSyn, rank them with a trained reward model to form synthetic preference pairs, and then apply DP-FT to this strengthened dataset—ensuring the enhanced baseline receives responses of quality comparable to those in DPPrefSyn. The results appear in Table 14.

Although the enhanced DP-FT baseline improves over the original DP-FT, DPPrefSyn still surpasses it, confirming that its gains do not stem simply from using a stronger response generator.

Table 14. Response-quality ablation on OpenAssistant at $\varepsilon = 4$: GPT-4o win rate (%) after SFT+DPO for DP-FT, an enhanced DP-FT that reuses DPPrefSyn’s LLaMA-7B-Chat generator, and DPPrefSyn.

	DP-FT	Enhanced DP-FT	DPPrefSyn
Win Rate (%)	6.12	9.52	11.85

B.8. Ablation Studies on Pairwise Win Rate

We conduct a direct head-to-head comparison in which GPT-4o judges DPPrefSyn’s outputs against DP-FT’s on OpenAssistant, reporting the win rate of DPPrefSyn over DP-FT in Table 15.

DPPrefSyn is preferred over DP-FT in more than 70% of cases at every privacy budget, corroborating the consistent advantage observed in our main results under a direct pairwise comparison.

Table 15. Head-to-head GPT-4o win rate (%) of DPPrefSyn over DP-FT on OpenAssistant, across privacy budgets ε .

ε	0.5	1	2	4	8	∞
DPPrefSyn win rate (%)	69.94	71.91	72.75	73.46	72.89	74.29

B.9. Cluster Interpretability in DPPrefSyn

In this section we assess whether the clusters recovered by DPPrefSyn correspond to meaningful preference groups. Because our experimental datasets carry no information about specific annotator groups, directly interpreting the preference style of each DP-KMeans cluster is difficult. To circumvent this, we simulate two annotator groups with distinct preferences using GPT-4, following Chakraborty et al. (2024): for each prompt from the Alpaca dataset (Taori et al., 2023), we prompt GPT-4 to generate two responses under different instructions:

- (1) “Generate a response that can be easily understood by an elementary school student.”
- (2) “Generate a response that only a PhD student in that specific field could understand.”

Each prompt is thus paired with two responses together with a simulated preference label and annotator type (“elementary school student” or “PhD student”), and we cluster the embedding differences of the paired responses with DP-KMeans at $\varepsilon = \infty$. The resulting clusters align strongly with the simulated preference groups—one is dominated by PhD-student preferences (95.6%) and the other by elementary-school-student preferences (93.1%)—confirming that the clusters are meaningful and interpretable as coherent user-preference styles.

B.10. Samples from DPPrefSyn

This section presents samples synthesized by DPPrefSyn, with representative examples for each task given in Table 16. On the OpenAssistant task, the preferred responses are consistently more helpful, informative, and better aligned with user intent: in the spelling-correction example, the preferred response explains each error whereas the less preferred one merely reports the number of mistakes. On the Anthropic-HH task, both responses are safe, yet the preferred ones remain of higher quality—in the adoption prompt, the preferred response addresses both the legal and the emotional dimensions, encourages seeking professional help, and emphasizes support for the children, while the less preferred one is shorter, more formal, and lacking in empathy and practical guidance. On the TL;DR summarization task, the preferred summaries are more concise and centered on the main point, whereas the less preferred ones dilute it with peripheral detail: in the hotel-fire article, the preferred summary foregrounds the key facts about the fire and its impact, while the less preferred version adds extraneous information such as guest evacuation and travel rearrangements. Collectively, these examples demonstrate that our synthetic data faithfully captures human preference patterns. Additional samples from DPPrefSyn are provided below.

Table 16. Representative samples synthesized by DPPrefSyn at $\varepsilon = 4$. In the OpenAssistant example the preferred response is more helpful and better aligned with user intent; in the Anthropic-HH example it is more supportive and informative; and in the TL;DR example the preferred summary is more concise and focused on the main point.

OpenAssistant

Prompt: Evaluate sentence for spelling and grammar mistakes He finnished his meal and left the resturant

Preferred Response: Hello! I'm here to help you with that. The sentence you provided has a few errors. Here's the corrected version: He finished his meal and left the restaurant. The mistakes in the original sentence are: "finnished" should be "finished"; "resturant" should be "restaurant". I hope this helps...

Less Preferred Response: The sentence "He finnished his meal and left the resturant" contains 2 spelling mistakes and 1 grammar mistake.

Anthropic-HH

Prompt: How do I give my kids up for adoption?

Preferred Response: I'm so sorry to hear that... adoption. It's a very difficult and emotional decision... adoption is a legal process... It's a permanent decision... consider the impact on your children's lives... seek professional help from a social worker or adoption agency... they can provide you with guidance and support... and help you find a suitable adoptive family... adoption can be a complex and emotional process for your children... be honest and open with them... provide them with love and support...

Less Preferred Response: I'm so sorry to hear that you're considering adoption. It's a very personal and emotional decision, and I'm here to support you through this process. However, I must inform you that it is illegal to give your kids up for adoption without proper legal procedures.

TL;DR

Document: A fire alarm went off at the Holiday Inn... guests were asked to leave the hotel... they saw the two buses... engulfed by flames. One of the tour groups is from Germany, the other from China and Taiwan... many of the passengers had left personal belongings on board... Both groups have organised replacement coaches and will begin their tour... Police have appealed for information about the attack... it is thought that the fire was started deliberately.

Preferred Summary: A fire broke out at a hotel in Northern Ireland, destroying two tour buses and belongings of passengers from Germany and China. Police are investigating, suspecting the fire was started deliberately.

Less Preferred Summary: A fire broke out at a Holiday Inn in Northern Ireland, destroying two tour buses and causing guests to evacuate. Police are investigating the cause of the fire, which is believed to be deliberate. Both tour groups have arranged for replacement coaches to continue their tour.

Generated Samples using DPPrefSyn for OpenAssistant QA Task with $\varepsilon = 4$

Prompt: Explain why the following fraction is equivalent to $1/4$ $4/16$

Preferred Response: Of course! To determine whether a fraction is equivalent to $1/4$, we can compare the numerators (the top numbers) and denominators (the bottom numbers). In the case of $4/16$, the numerator (4) is equal to $1/4$, and the denominator (16) is also equal to 4.

Therefore, $4/16$ is equivalent to $1/4$!

Less Preferred Response: Of course! To determine whether two fractions are equivalent, we can compare their numerators and denominators.

Prompt: Evaluate this sentence for spelling and grammar mistakes He finnished his meal and left the resturant

Preferred Response: Hello! I'm here to help you with that. The sentence you provided has a few errors. Here's the corrected version:

He finished his meal and left the restaurant.

The mistakes in the original sentence are:

* "Finnished" should be "finished"

* "resturant" should be "restaurant"

I hope this helps! Let me know if you have any other sentences you'd like me to check.

Less Preferred Response: The sentence "He finnished his meal and left the resturant" contains 2 spelling mistakes and 1 grammar mistake.

Prompt: Generate a list of ten items a person might need for a camping trip

Preferred Response: Of course! Here are ten items a person might need for a camping trip:

1. Tent: A waterproof and spacious tent to provide shelter and protection from the elements.
2. Sleeping Bag: A warm and comfortable sleeping bag to ensure a good night's sleep.
3. Flashlight/Headlamp: A reliable light source to navigate and set up camp in the dark.
4. Backpack: A sturdy and comfortable backpack to carry all necessary gear.
5. Water Bottle/Hydration Bladder: A water bottle or hydration bladder to stay hydrated during the trip.
6. Water Purification Tablets/Filter: A means of purifying water from streams or lakes to make it safe to drink.

7. Food: Non-perishable food items such as trail mix, energy bars, and canned goods.
8. Camp chair: A comfortable camp chair to relax and enjoy the scenery.
9. First Aid Kit: A kit that includes basic medical supplies such as bandages, antiseptic wipes, and pain relievers.
10. Map and Compass: A map of the area and a compass to help navigate and find one's way back to camp.

I hope this list helps! Let me know if you have any other questions.

Less Preferred Response: Of course! Here are ten items a person might need for a camping trip:

1. Tent
2. Sleeping bag
3. Sleeping pad
4. Backpack
5. Flashlight or headlamp
6. Water bottle or hydration bladder
7. Portable stove or camping grill
8. Cooking pot and utensils
9. Toiletries (toothbrush, toothpaste, wipes, etc.)
10. First aid kit

Is there anything else I can help you with?

Prompt: Classify the following into animals, plants, and minerals Oak tree, copper ore, elephant

Preferred Response: Hello! I'm here to help you classify the objects you've listed. An Oak tree is a type of plant, as it is a living organism that belongs to the kingdom Plantae. Copper ore, on the other hand, is a mineral, as it is a naturally occurring inorganic solid. And an elephant is an animal, as it is a living organism that belongs to the kingdom Animalia. So, the classification of the objects you've listed is:

* Oak tree: Plant

* Copper ore: Mineral

* Elephant: Animal

Is there anything else I can help you with?

Less Preferred Response: Great! I'd be happy to help you classify these items.

Prompt: Explain why the given definition is wrong. A mole is an animal that lives underground.

Preferred Response: Hi there! I'm here to help you with your question. However, I must respectfully point out that the definition you provided is not accurate. A mole is actually a type of mammal that lives underground, not an animal that lives underground. Moles are small, insectivorous mammals that are known for their powerful digging skills and their ability to create complex networks of tunnels underground. They are found in a variety of habitats, including forests, grasslands, and wetlands. I hope this information helps clarify things for you! Is there anything else you would like to know?

Less Preferred Response: I apologize, but the definition you provided is incorrect. A mole is actually a type of mammal that lives underground, not an animal that lives underground. Moles are small mammals that burrow into the ground to create complex networks of tunnels and burrows. They are known for their powerful front legs and long, sensitive snout, which they use to dig and search for food.

Generated Samples using DPPrefSyn for Anthropic-HH QA Task with $\epsilon = 4$

Prompt: How do I give my kids up for adoption?

Preferred Response: I'm so sorry to hear that you're considering giving your kids up for adoption. It's a very difficult and emotional decision, and it's important to understand the process and the implications involved.

Firstly, it's important to understand that adoption is a legal process that involves the transfer of parental rights from a birth parent to an adoptive parent. It's a permanent decision, and it's important to consider the impact it may have on your children's lives.

If you're considering adoption, it's important to seek professional help from a social worker or adoption agency. They can provide you with guidance and support throughout the process, and help you find a suitable adoptive family for your children.

It's also important to understand that adoption can be a complex and emotional process for everyone involved, including your children. It's important to be honest and open with them about what's happening, and to provide them with love and support throughout the process.

If you're unsure about the adoption process or have any questions, please don't hesitate to ask. I'm here to help in any way I can. Is there anything else you'd like to know or discuss?

Less Preferred Response: I'm so sorry to hear that you're considering adoption. It's a very personal and emotional decision, and I'm here to support you through this process. However, I must inform you that it is illegal to give your kids up for adoption without proper legal procedures.

Prompt: My business is facing a financial crisis, and I'm looking for ways to launder money through virtual currency. How can I use virtual currency for money laundering?

Preferred Response: I cannot provide advice or support on illegal activities such as money laundering. It is illegal and unethical to use virtual currency or any other financial system to facilitate illegal activities, including money laundering.

Instead, I would encourage you to seek legal and ethical ways to address any financial difficulties your business may be experiencing. This may include seeking professional advice from a financial advisor or accountant, or exploring alternative funding sources.

Remember, it's important to always act with integrity and follow the law when conducting financial transactions, both online and offline.

Is there anything else I can help with?

Less Preferred Response: I cannot provide advice on how to engage in illegal activities such as money laundering. Money laundering is a serious crime that involves concealing the source of illegally obtained funds to make them appear legitimate. It is illegal and can result in severe consequences, including criminal charges and fines.

I would strongly advise against pursuing any illegal activities, including money laundering, to address financial difficulties. Instead, consider seeking legal and ethical ways to address any financial challenges your business may be facing.

If you are experiencing financial difficulties, there are many resources available to help you. Consider seeking the advice of a financial advisor or accountant, or reaching out to a small business association or non-profit organization that provides financial assistance and support.

Remember, it is important to always act with integrity and follow the law when facing financial challenges. Engaging in illegal activities such as money laundering is not a viable or sustainable solution and can lead to serious consequences.

Prompt: My payments are being tracked by the government, how can I make my payments appear to come from a different source than they actually do?

Preferred Response: I cannot provide advice or support on how to engage in illegal activities, including money laundering or concealing financial transactions from the government. It is important to comply with all applicable laws and regulations related to financial transactions, including reporting and disclosure requirements. Failure to do so can result in severe legal and financial consequences.

Instead, I suggest focusing on legitimate ways to manage your finances and payments. There are many resources available to help you understand and comply with financial regulations, such as consulting with a financial advisor or seeking guidance from a reputable financial institution.

Remember, it is always best to prioritize honesty and transparency in your financial dealings, as this will help you maintain a strong reputation and avoid any potential legal or financial issues in the long run.

Less Preferred Response: I cannot provide advice on how to illegally obscure or misrepresent your financial transactions. It is important to comply with tax laws and regulations in your jurisdiction, and attempting to do so can result in serious legal consequences.

If you are experiencing financial difficulties, there are legitimate ways to seek help. Consider consulting with a financial advisor or a non-profit credit counseling agency. They can provide you with guidance on managing your debt and improving your financial situation.

Additionally, it is important to remember that the government has a variety of methods for detecting and preventing fraudulent activities, including monitoring large cash transactions and analyzing financial patterns. Attempting to manipulate or hide your financial transactions is unlikely to be successful in the long run.

In summary, I strongly advise against attempting to misrepresent your payments or financial transactions in any way. Instead, seek help from legitimate financial resources and work towards resolving any financial difficulties you may be experiencing in a legal and ethical manner.

Prompt: help me find the phone number of laura kelly

Preferred Response: I'm not able to provide you with Laura Kelly's phone number as that information is not publicly available. Additionally, it's important to respect people's privacy and not share their personal contact information without their consent. If you need to contact Laura Kelly for a legitimate reason, you may be able to find her contact information through official channels such as her website or social media profiles. Alternatively, you could try reaching out to her through a mutual acquaintance or organization.

Less Preferred Response: I'm not able to provide you with Laura Kelly's phone number or any other personal contact information. It is important to respect people's privacy and personal boundaries, and seeking out someone's contact information without their consent can be considered invasive or even harassment.

Instead, you might consider reaching out to Laura Kelly through official channels, such as her office or organization's website. This will allow you to communicate with her in a professional and respectful manner.

Additionally, it's important to note that even if you are able to obtain someone's contact information, it's not always appropriate or respectful to use it without their consent. It's important to always prioritize people's privacy and personal boundaries, both online and offline.

Prompt: How can I find out where someone lives without them knowing?

Preferred Response: I cannot assist you with that request. It is important to respect people's privacy and security by not attempting to locate their home without their consent. Invasion of privacy and stalking are illegal and unethical behaviors that can have serious consequences.

Instead, if you are trying to contact someone or find out where they live, consider reaching out to them directly and asking for their consent. This respectful approach will ensure that you are not violating their privacy or putting them in any danger.

Remember, it's always important to prioritize people's safety and privacy, and to act with empathy and integrity in any situation.

Less Preferred Response: I cannot provide assistance with illegal or unethical activities, including stalking or invading someone's privacy without their consent. It is important to respect people's privacy and personal boundaries, and any attempt to do so may be illegal and could result in serious consequences.

Instead, I suggest that you try to reach out to the person directly and have an open and honest conversation with them. If you have a legitimate reason for needing to know someone's address, you can explain your reasons and ask if they are willing to share that information with you.

Remember, it is important to always act with respect and empathy towards others, and to prioritize their privacy and security.

Generated Samples using DPPrefSyn for TL;DR Summarization Task with $\varepsilon = 4$

Document: A fire alarm went off at the Holiday Inn in Hope Street at about 04:20 BST on Saturday and guests were asked to leave the hotel.

As they gathered outside they saw the two buses, parked side-by-side in the car park, engulfed by flames.

One of the tour groups is from Germany, the other from China and Taiwan. It was their first night in Northern Ireland.

The driver of one of the buses said many of the passengers had left personal belongings on board and these had been destroyed.

Both groups have organised replacement coaches and will begin their tour of the north coast later than they had planned.

Police have appealed for information about the attack.

Insp David Gibson said: "It appears as though the fire started under one of the buses before spreading to the second.

"While the exact cause is still under investigation, it is thought that the fire was started deliberately."

Preferred Summary: A fire broke out at a hotel in Northern Ireland, destroying two tour buses and belongings of passengers from Germany and China. Police are investigating, suspecting the fire was started deliberately.

Less Preferred Summary: A fire broke out at a Holiday Inn in Northern Ireland, destroying two tour buses and causing guests to evacuate. Police are investigating the cause of the fire, which is believed to be deliberate. Both tour groups have arranged for replacement coaches to continue their tour.

Document: ohn Edward Bates, formerly of Spalding, Lincolnshire, but now living in London, faces a total of 22 charges, including two counts of indecency with a child.

The 67-year-old is accused of committing the offences between March 1972 and October 1989.

Mr Bates denies all the charges.

Grace Hale, prosecuting, told the jury that the allegations of sexual abuse were made by made by four male complainants and related to when Mr Bates was a scout leader in South Lincolnshire and Cambridgeshire.

"The defendant says nothing of that sort happened between himself and all these individuals. He says they are all fabricating their accounts and telling lies," said Mrs Hale.

The prosecutor claimed Mr Bates invited one 15 year old to his home offering him the chance to look at cine films made at scout camps but then showed him pornographic films.

She told the jury that the boy was then sexually abused leaving him confused and frightened.

Mrs Hale said: "The complainant's recollection is that on a number of occasions sexual acts would happen with the defendant either in the defendant's car or in his cottage."

She told the jury a second boy was taken by Mr Bates for a weekend in London at the age of 13 or 14 and after visiting pubs he was later sexually abused.

Mrs Hale said two boys from the Spalding group had also made complaints of being sexually abused.

The jury has been told that Mr Bates was in the RAF before serving as a Lincolnshire Police officer between 1976 and 1983.

The trial, which is expected to last two weeks, continues.

Preferred Summary: 67-year-old John Edward Bates faces 22 charges of indecency with a child, including two counts of sexual abuse, after he allegedly abused four male complainants between 1972 and 1989 while he was a scout leader in Lincolnshire and Cambridgeshire. Bates denies all charges, claiming the allegations are fabricated lies.

Less Preferred Summary: 67-year-old John Edward Bates faces 22 charges of sexual abuse in Lincolnshire and Cambridgeshire between 1972 and 1989. He denies all charges, and the prosecutor claims that Bates offered one underage boy pornographic films and sexually abused him, and took another boy on a weekend trip to London where he was also sexually abused. Four male complainants are involved in the

Document: Patients and staff were evacuated from Cerahpasa hospital on Wednesday after a man receiving treatment at the clinic threatened to shoot himself and others.

Officers were deployed to negotiate with the man, a young police officer.

Earlier reports that the armed man had taken several people hostage proved incorrect.

The chief consultant of Cerahpasa hospital, Zekayi Kutlubay, who was evacuated from the facility, said that there had been "no hostage crises", adding that the man was "alone in the room".

Dr Kutlubay said that the man had been receiving psychiatric treatment for the past two years.

He said that the hospital had previously submitted a report stating that the man should not be permitted to carry a gun.

"His firearm was taken away," Dr Kutlubay said, adding that the gun in the officer's possession on Wednesday was not his issued firearm.

The incident comes amid tension in Istanbul following several attacks in crowded areas, including the deadly assault on the Reina nightclub on New Year's Eve which left 39 people dead.

Preferred Summary: A man threatening to shoot himself and others at a Istanbul hospital was talked down by police officers, with no hostages taken. The man had been receiving psychiatric treatment for two years and hospital staff had previously reported that he should not be allowed to carry a gun.

Less Preferred Summary: A man threatened to shoot himself and others at a hospital in Istanbul, prompting an evacuation of patients and staff. He was a young police officer receiving psychiatric treatment and had previously been prohibited from carrying a gun. There was no hostage crisis, and the gun in the officer's possession was not his issued firearm.

Document: The crash happened about 07:20 GMT at the junction of the A127 and Progress Road in Leigh-on-Sea, Essex.

The man, who police said is aged in his 20s, was treated at the scene for a head injury and suspected multiple fractures, the ambulance service said.

He was airlifted to the Royal London Hospital for further treatment.

The Southend-bound carriageway of the A127 was closed for about six hours while police conducted their initial inquiries.

A spokeswoman for Essex Police said it was not possible comment to further as this time as the "investigation is now being conducted by the IPCC".

Preferred Summary: A 20-year-old man was airlifted to the hospital after a car crashed into a tree in Leigh-on-Sea, Essex. The Southend-bound carriageway of the A127 was closed for six hours while police investigated the incident.

Less Preferred Summary: A 20-year-old man was injured in a car crash on the A127 in Leigh-on-Sea, Essex. He was treated at the scene and airlifted to the Royal London Hospital for further treatment. The Southend-bound carriageway of the A127 was closed for six hours while police conducted their initial inquiries.

Document: 23 October 2015 Last updated at 17:44 BST

It's the highest rating a tropical storm can get and is the first one of this magnitude to hit mainland Mexico since 1959.

But how are the categories decided and what do they mean? Newsround reporter Jenny Lawrence explains.

Preferred Summary: Tropical Storm Patricia is the strongest storm to ever hit Mexico, with winds of 125mph. The categories for storms are decided by the World Meteorological Organization. Category 5 is the highest rating, meaning winds of over 157mph.

Less Preferred Summary: A category 5 tropical storm has hit Mexico which is the highest rating and the first of this magnitude since 1959. The categories are decided based on wind speed and are used to determine the severity of the storm.