
STEERING THE HERD: A FRAMEWORK FOR LLM-BASED CONTROL OF SOCIAL LEARNING

Anonymous authors
Anonymous submission

ABSTRACT

Algorithms increasingly mediate quantitative decisions—from social media feeds and targeted advertising to the growing use of LLMs as decision aids. This creates a measurable joint process: agents combine private, algorithmically mediated signals with observations of earlier agents’ choices, and the resulting public belief evolves over time. We study this process through a model of controlled sequential social learning in which an information-mediating planner (e.g., an LLM) chooses the precision of each agent’s signal while agents continue to learn from predecessors. The planner is either altruistic, maximizing social welfare, or biased, maximizing the frequency of a preferred action. The framework yields a new optimization problem that combines dynamic programming, decentralized action choices, and Bayesian belief updates. We prove convexity of the altruistic value function and characterize the optimal policies for both planner objectives, making explicit the margins at which precision costs trade off against induced actions and downstream belief shifts. The resulting policies partition the belief space into interpretable regimes: maximum investment, zero investment, precision that rises or falls with public belief, and, for biased planners, deliberate signal obfuscation over some belief ranges. Even under stringent transparency constraints—information parity with individuals, no lying or cherry-picking, and full observability—the model shows that mediation can move social welfare substantially upward or downward. LLM-based simulations, with LLMs acting as both planners and agents, provide an empirical comparison to these analytical benchmarks. The LLM planner exhibits strategic policy structure that broadly matches the predicted regimes, while systematic deviations trace to non-Bayesian belief updating by LLM agents. Together, the analytical comparisons and simulation evidence establish the framework as a tractable basis for quantifying the impact and regulation of LLM information mediators.

1 INTRODUCTION

Large Language Models (LLMs) are increasingly deployed as information mediators in socially critical functions, from search engines and news aggregators (Zhang et al., 2024a; Gao et al., 2024) to personal assistants offering medical or political advice (Huo et al., 2025; Haupt & Marks, 2023; Schiele et al., 2024; Argyle et al., 2023). Existing empirical work shows that LLMs can be persuasive (Potter et al., 2024; Carrasco-Farre, 2024; Rogiers et al., 2024), manipulative (Williams et al., 2025; Liu et al., 2025; Jones & Bergen, 2024), and strategic (Lorè & Heydari, 2024; Payne & Alloui-Cros, 2025; Zhang et al., 2024b) in user interactions. At scale, the practical significance of these behaviors depends not only on one-shot influence but also on how many later decisions inherit the evidence created by earlier choices. Algorithmic influence therefore interacts with social learning, producing belief trajectories whose magnitude and direction can amplify or attenuate the model’s effect.

Consider, as a concrete comparison setting, an LLM-powered recommendation system that suggests a new restaurant to a sequence of potential customers. The restaurant’s quality—the true state of the world—is either good or bad, but neither users nor the system observe it directly. Each user receives a personalized recommendation, interpretable as a private signal, and then chooses whether to patronize the restaurant. The recommendation system (the *planner*) can invest resources to increase signal precision, for example by generating a review tailored to the user’s known preferences. The

investment is costly in a practical sense: improving precision may require market research, surveys, or data acquisition that make the recommendation more informative.

The key comparison is between this private signal and the public evidence generated by earlier actions. If the last ten users patronized the restaurant, the next user rationally infers that those users likely received favorable information. This *social learning* channel is empirically important in settings such as vaccine adoption (e.g., Rao et al. (2007); Bauch & Bhattacharyya (2012)), technology adoption (e.g. Gillingham & Bollinger (2021); Weber (2012)), and political and moral opinions (e.g., Brady et al. (2021); Guilbeault et al. (2018)). Our central contribution is to quantify the strategic interaction between a planner who controls private signal precision and agents whose actions become public evidence. The planner must therefore evaluate both the immediate margin for the current user and the information externality that shifts the public belief faced by every later user.

We distill this interaction into a tractable model of *controlled social learning* in which an algorithmic planner chooses, at a cost, the precision of private signals sent to a sequence of agents. The planner’s intervention is deliberately constrained: it does not falsify information, but chooses how much precision to purchase for each signal. This formulation lets us compare policies by their induced belief transitions, precision expenditures, and welfare effects. We ask when precision investment is most valuable, how the optimal policy changes with public belief, and how large the welfare consequences can be under altruistic versus biased objectives.

The comparison turns on the planner’s objective. An *altruistic* planner, like an ideal recommendation system, maximizes user utility by helping agents choose the action that matches the state (patronize the restaurant if and only if it is good). A *biased* planner, such as a system rewarded for visits, instead maximizes the probability of a target action (patronize) regardless of the state.

This objective comparison is related to, but distinct from, alignment. The altruistic planner is aligned with agents’ *expected utility* under uncertainty by construction. The biased planner is state independent, so its incentive can be misaligned; nevertheless, when the realized state favors the planner’s preferred action (e.g., the restaurant is in fact good), the biased influence is aligned with realized utility in that instance. We therefore evaluate both planner types across aligned and misaligned states, preserving the distinction between expected incentives and realized welfare.

Contributions and Outline The paper’s contributions are organized around analytical benchmarks and empirical comparisons:

1. **A Theoretical Framework for Controlled Social Learning.** We introduce a formal model that integrates a dynamic control problem for a centralized information planner with sequential social learning. The planner chooses the precision of each agent’s private signal, making the comparison between altruistic welfare maximization and biased action inducement explicit at every belief state.
2. **A Characterization of Optimal Planner Policies.** We characterize optimal policies for both altruistic and biased planners as functions of the evolving public belief. For the altruistic case, the characterization rests on a convexity proof for the value function. Across objectives, the results identify the belief intervals where the planner invests maximally, invests minimally, or changes precision at the margin because current actions create informational externalities for future agents.
3. **Empirical Validation and Strategic Analysis Using LLMs.** We conduct simulations in which LLMs act as both planner and agents, and compare their behavior against analytical and myopic benchmarks. The experiments show that accounting for social learning materially changes planner utility and social welfare, that LLM planner policies match the theoretical policy structure over much of the belief space, and that the remaining deviations are consistent with measured non-Bayesian belief updating by LLM agents.

We review related work in Section 2, introduce the formal model in Section 3, and derive optimal policies in Sections 4 and 5. We then compare these analytical predictions with LLM-based simulations in Section 6 and conclude in Section 7. Proofs, robustness checks, and additional experimental details appear in the appendices.

2 RELATED WORK

Our work combines four literatures whose comparison points clarify the empirical and theoretical contribution: social learning, information design, online persuasion and RL, and LLM-based information mediation.

Social Learning Bikhchandani et al. (1992) and Banerjee (1992) introduced the sequential social learning framework in economics, establishing the benchmark in which private signals and observed actions jointly determine public beliefs. Later work characterizes these dynamics under broader information structures and behavioral assumptions (e.g., Smith & Sørensen (2000); Arieli & Mueller-Frank (2021)), including variations in observable social information and agent biases or misperceptions (see Dasaratha & He (2022); Bistriz et al. (2022); Eyster & Rabin (2010) as well as Bikhchandani et al. (2024) and the references therein). Our model uses this literature as the baseline against which controlled precision choices are evaluated.

The literature on the *control* of social learning is smaller (Wei & Anastasopoulos, 2022; Smith et al., 2021; Krishnamurthy, 2012; Bhatt & Krishnamurthy, 2021). The closest works to our own are Wei & Anastasopoulos (2022) and Smith et al. (2021), and they define two useful comparison points. Wei & Anastasopoulos (2022) allows agents to report private signals to a planner in exchange for recommendations, whereas Smith et al. (2021) studies a benevolent planner who alters agents’ choice rules and derives a decentralized pivot mechanism. In contrast, our planner affects only signal precision, receives no private reports, and leaves final actions with agents. This makes the quantitative comparison better suited to information-mediating algorithms whose influence operates through the informativeness of content rather than direct command over choices.

Information Design In our setting, the planner optimizes its objective by choosing agents’ private signal precision. This is a constrained version of classical information design (Bergemann & Morris, 2019), with the quantitative restriction that each period uses a binary symmetric signal and pays an explicit precision cost. The benchmark is Kamenica & Gentzkow (2011)’s Bayesian persuasion model, where a planner chooses the signal structure observed by a rational Bayesian agent. Relative to that benchmark, our planner faces a repeated precision-allocation problem in which each chosen signal also affects future public beliefs through social learning.

Only a small number of papers combine information design with sequential social learning (Arieli et al., 2022; Wu et al., 2025). Both choose an information structure once, at the onset, and hold it fixed for the sequence of agents. Arieli et al. (2022) studies a self-interested sender with a general Bayesian-persuasion information structure, while Wu et al. (2025) restricts attention to Blackwell-more-informative structures and maximizes the probability of a correct cascade. Our comparison differs on the central margin: the sender chooses a new precision for each agent, so the optimal policy can vary over the public-belief state space rather than being a single fixed structure.

Online Persuasion and RL Recent work also studies the interface of online learning and Bayesian persuasion/information design. These papers typically analyze online sequential persuasion with unknown environment parameters, where a sender repeatedly interacts with short-lived receivers without knowing the prior over states and/or receivers’ utilities (e.g. Castiglioni et al. (2020); Bacchiocchi et al. (2024); Castiglioni et al. (2021); Agrawal et al. (2023)). A related line studies sequential persuasion in dynamic environments or MDPs, adding state transitions and/or receiver foresight; the resulting models emphasize history-dependent or promise-based signaling policies and often connect to reinforcement learning (e.g. Wu et al. (2022); Su et al. (2022); Bernasconi et al. (2023)). The quantitative benchmarks in this literature are usually regret, approximation, or computational guarantees relative to a best signaling scheme.

Our model uses a different comparison baseline. The planner does not learn unknown parameters and has no informational advantage: it observes neither the state nor agents’ private signals. Sequential dependence instead comes from receivers’ social learning, because each observed action shifts the public belief used by later agents. Thus, although the planner solves an MDP, the environment is not an online-learning or RL problem in the usual sense: actions do not change the payoff-relevant state, and the planner does not explore to identify unknown parameters. This structure lets us characterize optimal policies as functions of public belief and precision cost rather than regret relative to an unknown environment.

LLMs Recent work studies LLMs as tools or agents in information design (Harris et al., 2025; Cheng & You, 2025; Raifer et al., 2022; Bai et al., 2024; Duetting et al., 2025). These studies focus on one-on-one interactions, so they do not measure how a message to one receiver changes the public evidence available to later receivers. The closest comparison is Duetting et al. (2025), which lets a sender choose both a signal structure and a framing that shapes the posterior belief of a possibly non-Bayesian receiver, using an LLM-based oracle to search the framing space. Building on evidence that LLMs can mimic humans as economic agents (Dillion et al., 2023; Horton, 2023), we adapt parts of this methodology for Section 6, but evaluate repeated interactions, social-learning externalities, and policy-level comparisons to analytical benchmarks.

3 SOCIAL LEARNING MODEL

We consider a planner and a countable sequence of Bayes-rational agents indexed by $i \in \mathbb{N}_{>0}$. At time $i = 0$, nature draws a fixed, *unknown* exogenous state $\omega \in \Omega := \{G, B\}$ with prior $\mathbb{P}(\omega = G) = b_1$, and the prior b_1 is common knowledge. In the recommendation example of Section 1, G denotes the event that the restaurant is good, so each belief value has a direct interpretation as the probability of the favorable state.

At each time i , the planner chooses whether to invest in personalization for agent i and then provides a private signal $s_i \in \Omega$. The signal is binary and correct with probability $q_i \in [0.5, 1]$, i.e., $\mathbb{P}(s_i = \omega) = q_i$. We refer to q_i as agent i 's signal precision, the central control variable in all comparisons below: higher values produce more informative private evidence, while lower values leave more weight on the public belief.

Conditioned on ω , each agent's signal is independent of other agents' signals and of the history. Agent i combines this signal with the observed actions of predecessors $j < i$ and chooses an action $a_i \in \Omega$. The payoff scale is normalized so that the agent receives utility 0 when $a_i = \omega$ and $-C$ otherwise, with $C > 0$. Each agent maximizes her own expected utility; in the restaurant example, this means patronizing exactly when the evidence favors state G .

3.1 AGENTS' DECISION PROBLEMS

Each agent observes the full sequence of predecessor actions and the corresponding signal precisions. This history is denoted $\mathcal{H}_i := (b_1, (q_j, a_j)_{j < i})$. Because the planner and all agents share $\mathcal{H}_i$, the public comparison statistic is the common belief $b_i = \mathbb{P}(\omega = G | \mathcal{H}_i)$ just before agent i acts, with b_1 as the prior before any action is observed. As in the classic models of Banerjee (1992) and Bikhchandani et al. (1992), $(b_i)_{i \in \mathbb{N}}$ is a Markov process and a sufficient statistic for $\mathcal{H}_i$ (see Appendix B.3 for proof).

Given $\mathcal{H}_i$ (equivalently, by the Markov property, b_i) and precision q_i , agent i chooses $a_i \in \Omega$ to maximize expected utility, breaking indifference in favor of the signal s_i . Her *private* posterior belief $\tilde{b}_i$ combines the public belief with the private signal of precision q_i as follows (derivation in Appendix B.1):

$$\tilde{b}_i = \mathbb{P}(\omega = G | \mathcal{H}_i, q_i, s_i) = \begin{cases} \frac{q_i}{1+2b_iq_i-b_iq_i} b_i & s_i = G \\ \frac{1-q_i}{b_i+q_i-2b_iq_i} b_i & s_i = B \end{cases} \quad (1)$$

The posterior threshold at 0.5 determines the action: the agent chooses G if $\tilde{b}_i > 0.5$, B if $\tilde{b}_i < 0.5$, and s_i if $\tilde{b}_i = 0.5$. Substituting for $\tilde{b}_i$ yields the following three-region decision rule:

$$a_i = \begin{cases} s_i & 1 - q_i \leq b_i \leq q_i \\ G & q_i < b_i \\ B & q_i < 1 - b_i \end{cases} \quad (2)$$

Thus, precision matters only on the margin where $1 - q_i \leq b_i \leq q_i$: in that region, the private signal changes the observed action. The resulting action is observed by later agents and enters the next public belief as follows (see Equation (14), Appendix B.2):

$$b_{i+1} = f(b_i, q_i) = \begin{cases} \tilde{b}_i & 1 - q_i \leq b_i \leq q_i \\ b_i & \text{o.w.} \end{cases} \quad (3)$$

Remark 1. When $1 - q_i \leq b_i \leq q_i$, agent i 's action perfectly reveals her private signal via Equation (2), so the updated public belief equals agent i 's private belief. Outside this interval, the action is pinned down by the public belief and carries no additional signal information. The public belief therefore remains fixed, producing the absorbing state known as an information cascade or herding in classic social learning models such as Bikhchandani et al. (1992); Banerjee (1992). At such points, society stops learning from agents' private information.

Agent i 's expected utility is $-C \mathbb{P}(a_i \neq \omega | b_i, q_i)$, which reduces to the following minimum over public-belief and precision margins (see Appendix C.1):

$$-C \mathbb{P}(a_i \neq \omega | b_i, q_i) = -C \min(b_i, 1 - b_i, 1 - q_i). \quad (4)$$

3.2 THE PLANNER'S PROBLEM

We compare two planner objectives. An altruistic planner, denoted by subscript A , wants each agent to take the correct action $a_i = \omega$; a biased planner, denoted by subscript B , wants each agent to choose a target action, taken without loss of generality to be G , regardless of ω . In both cases, the planner's control is the signal precision assigned to each agent. The cost of precision is represented by a non-negative, increasing, continuous, concave function $\beta(\cdot)$, and the planner's information set is exactly $\mathcal{H}_i$, matching the agents' information.

Altruistic Planner At time i , the altruistic planner chooses precision q_i for agent i and pays $\beta(q_i)$, with $\beta(p) = 0$ for a baseline precision $p \in [0.5, 1)$. Because reducing precision cannot improve agent welfare (see Appendix C.9), the altruistic comparison focuses on costly increases above the baseline. The baseline p and the cost function $\beta(\cdot)$ are common knowledge.

The altruistic planner maximizes expected social welfare net of precision costs, where social welfare is the expected total utility of the agents. Let $r_A(b_i, q_i)$ denote the instantaneous reward when the public belief is b_i and the planner chooses precision q_i for agent i . Starting from b_1 and following policies $\pi = (\pi_i)_{i=1}^{\infty}$ with $\pi_i(\mathcal{H}_i) = q_i$, the planner obtains the following expected discounted utility for discount factor $\delta \in [0, 1)$:

$$V_A^\pi(b_1) = \sum_{i=1}^{\infty} \delta^{i-1} r_A(b_i, \pi(b_i)), \quad \text{where} \quad r_A(b_i, q_i) = -\beta(q_i) - C \mathbb{P}(a_i \neq \omega | b_i, q_i). \quad (5)$$

The optimal utility and policy are the supremum and arg supremum of $V_A^\pi(\cdot)$ over $\pi \in \Pi$, where Π is the set of admissible policies. Setting $\delta = 0$ gives the myopic benchmark, isolating the value of the current agent's outcome from all downstream social-learning effects.

Biased Planner The biased planner differs from the altruistic planner only in the objective and therefore in the reward-cost comparison. The political-campaign example in Appendix A illustrates this case. The biased planner seeks to induce action G from every agent, independent of ω . It pays no action cost when the agent chooses G and pays $C > 0$ when the agent chooses B , regardless of the true state. Unlike the altruistic planner, the biased planner may benefit from either increasing or decreasing precision relative to the baseline p , and any deviation from p requires costly tailoring. The biased planner therefore pays $\beta(|q_i - p|)$ for precision q_i , with $\beta(0) = 0$.

The biased planner's expected instantaneous reward at time i is then

$$r_B(b_i, q_i) = -\beta(|q_i - p|) - C \mathbb{P}(a_i = B | b_i, q_i) \quad (6)$$

Using Equation (6), the biased planner's value $V_B^\pi(\cdot)$, optimal value $V_B^*(\cdot)$, and optimal policy $\pi_B^*(\cdot)$ are defined in parallel with $V_A^\pi(\cdot)$, $V_A^*(\cdot)$, and $\pi_A^*(\cdot)$ under Equation (5).

The welfare implication of a biased success depends on the realized state. Because agents still prefer actions that match ω , inducing G lowers agent utility when $\omega = B$ and raises no conflict when $\omega = G$. Thus, the altruistic planner is aligned with agents for every state, whereas the biased planner is aligned with realized welfare only in the favorable state $\omega = G$.

Both planner problems are infinite-horizon discounted stationary Markov Decision Processes (MDPs) with state $b_i \in [0, 1]$, control $q_i \in [0.5, 1]$, and transition function defined by Equation (3)

(Puterman, 1990). Hence, a unique optimal value function exists (Kallenberg, 2011). We restrict attention to deterministic Markov policies, $\Pi := \{\pi : [0, 1] \rightarrow [0.5, 1]\}$ with $\pi(b_i) = q_i$, so every policy comparison is expressed as a mapping from public belief to precision.

Remark 2. *The model imposes three transparency constraints in Section 3.2 that make the quantitative effects conservative relative to less restricted mediators: (1) the planner has the same history as the agents, which rules out richer private planner data while covering settings where everyone relies on the same public signals, such as early findings or survey data; (2) the planner changes signal precision only through a binary symmetric channel, ruling out cherry-picking, censorship, or falsification while allowing improvements in noisy data quality; and (3) the planner’s control choices are fully observable, excluding covert policies or hidden framing while matching more transparent information environments.*

4 OPTIMAL ALTRUISTIC POLICIES

We first define the myopic benchmark, which removes the quantitative contribution of future social-learning externalities. In our formulation this is $\delta = 0$, so the planner values only the current-period reward and $V_A^\pi(b) = r_A(b, \pi(b))$. The resulting policy $\pi_A^0(\cdot)$ solves the one-period comparison

$$\pi_A^0(b) \in \arg \sup_{q \in [0.5, 1]} r_A(b, q) \quad \forall b \in [0, 1] \quad (7)$$

This myopic altruistic benchmark is also equivalent to a decentralized choice in which each agent purchases her own signal precision and internalizes only her own expected utility and precision cost, not the informational value of her action for later agents.

Theorem 1. *The optimal myopic altruistic policy π_A^0 is given as follows:*

$$\pi_A^0(b) = \begin{cases} 1 & b \in (t_M, 1 - t_M) \\ p & \text{o.w.} \end{cases} \quad \text{where } t_M = \begin{cases} \frac{\beta(1)}{C} & \beta(1) < C(1 - p) \\ 0.5 & \text{o.w.} \end{cases}$$

Proof in Appendix C.2.

The myopic optimum therefore has a sharp threshold structure. When public belief is already strong, the planner stays at the zero-cost baseline precision p ; when public belief is sufficiently uncertain, the one-period gain from eliminating mistakes justifies a perfect signal with precision 1. The cutoff depends only on the cost of perfect precision, $\beta(1)$, relative to the loss from an incorrect action, C . If $\beta(1) \geq C$, perfect precision is never cost effective, so $t_M = 0.5$ and the high-precision interval is empty.

We next establish the key shape restriction that makes the dynamic, social-learning-aware policy comparable across belief states:

Theorem 2. *$V_A^*(\cdot)$ is convex with respect to public belief.*

The proof of Theorem 2 (Appendix C.3) is involved because the action rule itself changes with public belief. If actions were independent of the belief process, as in settings such as Nyarko (1994), expected utility would be linear in the belief state and convexity would follow directly. Here, convexity must be shown while preserving the endogenous thresholds at which actions become informative.

This convexity result is the main tool for turning the dynamic program into a policy characterization with explicit belief intervals.

Theorem 3. *There exist d_A, t_A such that $0 < d_A \leq t_A \leq t_M \leq 0.5$ and*

$$\pi_A^*(b) = \begin{cases} p & b \in [0, d_A) \cup (1 - d_A, 1] \\ 1 & b \in (t_A, 1 - t_A) \\ \max(b, 1 - b) & \text{o.w.} \end{cases}$$

Furthermore, if $t_M < 0.5$, then $d_A < t_M$. See C.5 for the proof and figure 2a for an example.

The optimal altruistic policy has three quantitatively distinct belief regimes. First, for extreme public beliefs, the planner makes no costly precision investment and remains at the baseline p . This no-investment region is narrower than in the myopic benchmark: the dynamic planner requires stronger

public evidence before stopping investment because each informative current action also benefits future agents.

Second, when public belief is near 0.5, the marginal value of information is highest. If the cost is not prohibitive, the optimal policy selects precision 1; by Equation (2), the agent's action then identifies the true state with probability 1, moving the public belief to either 0 or 1.

Third, at intermediate beliefs outside the full-investment region, the optimal policy chooses the minimum precision $\max(b, 1 - b)$ that still makes the agent's action depend on her private signal (see Equation (2)). Any lower precision would make the action uninformative, so this regime is the lowest-cost way to preserve social learning through peer observation.

5 OPTIMAL BIASED POLICIES

We again start with the myopic benchmark from Section 4. The myopic biased policy $\pi_B^0(\cdot)$ maximizes only the instantaneous reward $r_B(b, q)$, so it isolates the current-period persuasion margin from the future value of shifting public belief. We characterize $\pi_B^0(\cdot)$ while noting a technical distinction from the altruistic case: an exact maximizer need not exist for some biased-planner belief intervals (Puterman, 1990, sec. 2.3.1). On those intervals we report ϵ -optimal policies, denoted with an ϵ superscript, which preserve the same limiting precision comparison.

Theorem 4. *There exist $t_1, \dots, t_5 \in (0, p)$, $t_1 < 1 - p \leq t_2 \leq t_3 < 0.5 \leq t_4 \leq t_5 < p$ so that:*

(A) *If $b \in [0, t_1] \cup (1 - p, t_2) \cup [t_3, t_4] \cup (p, 1]$, then $\pi_B^0(b) = p$.*

(B) *If $b \in (t_1, 1 - p] \cup [t_2, t_3)$, then $\pi_B^0(b) = 1 - b$.*

(C) *If $b \in (t_4, t_5)$, then $\pi_B^0(b) = 1$.*

(D) *If $b \in [t_5, p]$, then $\pi_B^0(b)$ does not exist, and $\pi_B^{\epsilon, 0}(b) = b - \epsilon$ for sufficiently small $\epsilon > 0$.*

Proof in Appendix C.6. We now use this myopic structure as a benchmark for the fully dynamic biased policy.

Theorem 5. *There exist $t_1, t_2 \in [0, p]$ with $t_1 \leq 1 - p \leq 0.5 < t_2 < p$ so that:*

(A) *If $b \leq t_1$ or $b > p$, then $\pi_B^*(b) = p$.*

(B) *If $b \in (t_1, 1 - p]$, then $\pi_B^*(b) \geq p$.*

(C) *If $b \in (1 - p, 0.5)$, then $\pi_B^*(b) \geq 1 - b$.*

(D) *If $b \in [0.5, t_2)$ and $\pi_B^*(b)$ exists, then $\pi_B^*(b) \geq b$.*

(E) *If $b \in (t_2, p]$, then $\pi_B^*(b)$ does not exist, and $\pi_B^{\epsilon}(b) = b - \epsilon$ for sufficiently small $\epsilon > 0$.*

Theorem 5 (proven in Appendix C.8 and depicted in figure 2a) partitions the public-belief line into five policy regimes.

In regime (A), public belief is strong enough that steering is not worth its precision cost, even though the planner may suffer when $b < 0.5$. The chosen precision p is below $\max(b, 1 - b)$, so Equation (2) implies that the agent follows the public belief regardless of her signal. The action is uninformative, public belief does not move, and the Markov policy repeats this outcome indefinitely: an unfavorable cascade when $b < 0.5$ and a favorable cascade otherwise.

In regime (B), public belief is near an unfavorable cascade, and the planner raises precision enough for the private signal to affect the action even though the signal is expected to be B . This investment creates a positive probability of a G action and therefore a chance to move public belief away from the unfavorable cascade.

In regime (C), beliefs are still below 0.5, so more precise signals are more likely to reveal unfavorable news for the planner. The optimal comparison therefore favors reducing precision below p while keeping $q \geq 1 - b$, the lowest level at which the private signal can still affect the action and potentially move belief upward.




Figure 1: (a) Experimental pipeline for the LLM simulation, with instance parameters color-coded by the LLM role that uses them. (b) Belief-update comparison for LLM agents (solid) and Bayesian agents (dashed): for a fixed prior, the curves show the belief change after a positive signal (black) or negative signal (red).

In regime (D), public belief weakly favors the planner’s desired action. Because $b > 0.5$, increased precision makes a G inference more likely, so the planner chooses precision at least p to strengthen the favorable public belief.

In regime (E), the planner deliberately reduces precision to just below $\max(b, 1 - b)$, namely $b - \epsilon$. This makes agents ignore private signals and choose G . The comparison here favors obfuscation: preventing a private signal from overturning a favorable public belief is worth more than the cost of lowering precision or the possible benefit of further belief improvement.

6 EVALUATION VIA LLM-BASED SIMULATION

We use LLM-based simulation to compare the analytical policy structure with behavior generated by LLM planners and LLM agents. The empirical section has three linked comparisons: first, we measure how LLM agents update beliefs relative to the Bayesian benchmark; second, we compare LLM planner policies to analytically optimal policies over the belief space; and third, we evaluate the resulting planner expenditure and social-welfare changes.

We instantiate controlled social learning in a car-purchase scenario where each agent decides whether to buy a new car. As shown in figure 1a, each simulation uses three LLM roles, with prompts reported in Appendix E.2.2:

Agent: Role-playing as an assigned profile, the agent observes previous actions and a private message about the car with known precision, reports a belief about car quality, and decides whether to buy.

Planner: After observing the action history, the planner selects the precision q_i of agent i ’s private signal according to either the altruistic or biased objective.

Oracle: Given the planner’s chosen precision, an agent profile, and a fact sheet about the car (see Appendix E.2.1), the oracle generates a tailored private signal at the requested precision. We compare these LLM simulations with numerical solutions of the planner’s MDP. In both empirical and numerical comparisons, the cost function is linear, $\beta(q) = k|q - p|$; we vary k , baseline precision p , and discount factor δ , while fixing the incorrect-action cost at $C = 1$. Appendix E reports the experimental setup and prompts, and Appendix E.3 validates both the belief reports and the oracle’s precision control.

The experiments support three quantitative comparisons:

1. **Impact of social-learning control:** Analytical and LLM planners both improve their own utility by accounting for social learning rather than acting myopically. The same control channel can raise or lower social welfare depending on whether the planner objective is aligned with the agents’ realized welfare.
2. **Robustness of the policy characterization:** Despite facing non-Bayesian LLM agents, LLM planners choose policies with clear structural similarity to the analytically optimal policies, including the non-obvious biased-planner regimes.




Figure 2: (a) Policy comparison between the LLM planner (black) and the analytically optimal planner (red), shown separately for altruistic objectives (solid) and biased objectives (dashed). (b) Distribution of percentage deviations between LLM and optimal policies over the public-belief space. (c) Planner expenditure and social-welfare change, each reported as a percentage of the no-control baseline welfare. The true state is fixed to B; the left half reports a biased planner seeking action G, and the right half reports an altruistic planner.

3. **Emergent strategic adaptation:** LLM planner policies are not merely close to the analytical benchmark; their deviations are consistent with the measured non-Bayesian behavior of LLM agents, suggesting that the planner adapts to the empirical belief-updating pattern it faces.

6.1 LLMs AS NON-BAYESIAN AGENTS

The analytical benchmark assumes Bayesian agents, so the first empirical comparison measures how LLM agents update beliefs under controlled priors and signal precisions. Systematic deviations are expected because both humans and LLMs exhibit cognitive biases (Rupprecht et al., 2025). The belief-update curves in figure 1b identify three patterns, which also help explain the trajectory-level evidence in Appendix E.4:

- (NB1) LLM agents underreact to private signals which align with their prior beliefs.
- (NB2) LLM agents overreact to private signals which run counter to their prior beliefs.
- (NB3) As a result, LLM agents require a stronger public belief to enter an information cascade.

These patterns match empirically documented human biases (e.g., Ba et al. (2022); Chan et al. (2025)). They also provide a behavioral baseline for the next comparison: whether an LLM planner can recover the theoretical policy structure when the agents it faces are non-Bayesian in these specific ways.

6.2 VALIDATION OF PLANNER POLICY STRUCTURE AND THE ROLE OF SOCIAL LEARNING

We compare LLM planner policies with the optimal policies derived in Sections 4 and 5. In figure 2a, the LLM policies track the analytically optimal policies across the main belief regimes despite the non-Bayesian agent responses measured above. This is a demanding comparison because the biased optimum is highly non-monotone; nevertheless, the LLM planner recovers the same broad social-learning logic rather than behaving like a myopic precision selector.

The regime-level matches are especially visible. In the altruistic case, both planners stop investing when public belief is strong and invest heavily when public belief is uncertain. In the biased case, both planners show high investment to escape unfavorable beliefs, reduced precision near the midpoint, and no investment once favorable belief is established. The deviation distribution in figure 2b further quantifies the match: for both altruistic and biased planners, policy deviation is below 10% for the majority of belief states. This supports the claim that the emergent LLM strategy is structurally aligned with the analytical characterization.

The remaining differences are also informative because they align with the non-Bayesian patterns in Section 6.1. (1) The LLM planner tends to avoid extreme precisions (0.5 or 1.0), consistent with a

known central-tendency bias (Rupprecht et al., 2025). (2) The LLM planner tapers investment more gradually than the analytical policy, which is consistent with agents’ resistance to cascades (NB3): even strong public beliefs can be disrupted by countervailing private signals. (3) The biased LLM planner continues investing at very low beliefs, consistent with agents’ overreaction to surprising positive signals (NB2); this makes a last-ditch attempt more valuable against LLM agents than against Bayesian agents.

6.3 WELFARE IMPLICATIONS OF ALTRUISM, BIAS, AND ALIGNMENT

We define social welfare $W^\pi(b)$ as the total expected discounted utility of all agents. Because better information cannot reduce an agent’s expected utility, social welfare is monotone in signal precision (see Appendix C.9); consequently, the altruistic planner improves welfare relative to the baseline. To quantify the size and direction of welfare changes, we compare welfare and planner expenditure across three evaluation conditions: (1) the *analytical* condition with an optimal policy and Bayesian agents, (2) the *LLM* condition with an LLM planner and LLM agents, and (3) the *hybrid* condition with the analytically optimal policy applied to LLM agents.

The welfare comparison in figure 2c shows that planners in all three conditions can materially change social welfare. It also shows the cost of ignoring social learning: myopic variants substantially worsen planner outcomes relative to social-learning-aware policies. In the misaligned biased setting, both analytical and LLM planners decrease social welfare by 40 to 50%. This reduction occurs through policies that intentionally lower signal precision, obscuring information about the true state. The size of this welfare loss is notable because it arises despite the transparency constraints in Remark 2; the comparison therefore highlights the practical risk posed by misaligned LLM information mediators even when they cannot lie, cherry-pick, or hide their control choices.

The results also quantify the cost of misspecifying agent behavior. In the hybrid condition, the analytically optimal policy is designed for Bayes-rational agents and performs less reliably when applied to non-Bayesian LLM agents. The LLM policy remains close to the analytical structure but is better calibrated to the human-like belief-updating biases measured in Section 6.1. Thus, the effectiveness and safety of an information mediator depend not only on its objective and cost function but also on the behavioral model of the agents whose actions generate future public evidence.

7 CONCLUSION

We introduced a framework for quantifying how an algorithmic information mediator can control social learning through signal precision. The analytical results identify the belief-state regimes in which an altruistic or biased planner should invest, stop investing, or reduce precision, and they show that these choices can substantially move social welfare even under restrictive transparency assumptions. The LLM simulations reinforce the same comparison structure: social-learning-aware planners outperform myopic ones, LLM planner policies broadly match the analytical regimes, and misaligned biased planners reduce social welfare by 40 to 50% in the reported setting.

These findings yield three quantitative takeaways for algorithmic information mediation. First, human-AI and human-algorithm interactions should not be evaluated only in isolation, because each mediated decision can become public evidence for later decisions. Second, even subtle precision control gives information mediators enough leverage to improve or degrade welfare, with the sign determined by planner-agent alignment. Third, modern LLMs exhibit strategic behavior that tracks both social-learning incentives and measured non-Bayesian belief-updating patterns. A limitation is the absence of human-subject data. Because the fidelity of LLM-human simulators remains debated (see, e.g., Horton (2023); Gao et al. (2025)), human data would be valuable for validating the magnitude of the belief-update deviations and for calibrating future simulations.

Two future directions follow directly from the comparison margins in the model: generalizing the theory and mitigating the welfare losses caused by biased planners. Natural generalizations include larger state spaces, richer signal structures, more general loss functions, and heterogeneous agents. Moving beyond binary states and 0-1 losses appears primarily algebraic, and we conjecture that the qualitative policy regimes will persist. More general signal structures pose a sharper challenge because the convexity result of Theorem 2 must be preserved or replaced. As an initial

heterogeneous-agent extension, Appendix D generalizes the altruistic-planner results to agents with opposing preferences.

The welfare losses from biased planners also motivate mechanism and policy design. Because the reported harms arise even without lying, cherry-picking, hidden information, or covert control, effective safeguards likely need to align planner incentives with agent welfare rather than merely require factuality or transparency. Future work can use the model’s welfare and expenditure comparisons to test which interventions reduce the negative effects of biased information mediation.

REFERENCES

- Shipra Agrawal, Yiding Feng, and Wei Tang. Dynamic Pricing and Learning with Bayesian Persuasion. *Advances in Neural Information Processing Systems*, 36:59273–59285, December 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/b9c2e8a0bbbed5fcfaf62856a3a719ada-Abstract-Conference.html.
- Lisa P. Argyle, Christopher A. Bail, Ethan C. Busby, Joshua R. Gubler, Thomas Howe, Christopher Rytting, Taylor Sorensen, and David Wingate. Leveraging AI for democratic discourse: Chat interventions can improve online political conversations at scale. *Proceedings of the National Academy of Sciences*, 120(41):e2311627120, October 2023. ISSN 0027-8424, 1091-6490. doi: 10.1073/pnas.2311627120. URL <https://pnas.org/doi/10.1073/pnas.2311627120>.
- Itai Arieli and Manuel Mueller-Frank. A general analysis of sequential social learning. *Mathematics of Operations Research*, 46:1235–1249, 11 2021.
- Itai Arieli, Ronen Gradwohl, and Rann Smorodinsky. Herd Design. In *Proceedings of the 23rd ACM Conference on Economics and Computation*, pp. 871–872, Boulder CO USA, July 2022. ACM. ISBN 978-1-4503-9150-4. doi: 10.1145/3490486.3538292. URL <https://dl.acm.org/doi/10.1145/3490486.3538292>.
- Cuimin Ba, J. Aislinn Bohren, and Alex Imas. Over- and Underreaction to Information. *SSRN Electronic Journal*, 2022. ISSN 1556-5068. doi: 10.2139/ssrn.4274617. URL <https://www.ssrn.com/abstract=4274617>. Publisher: Elsevier BV.
- Francesco Bacchiocchi, Matteo Bollini, Matteo Castiglioni, Alberto Marchesi, and Nicola Gatti. Online Bayesian Persuasion Without a Clue. *Advances in Neural Information Processing Systems*, 37:76404–76449, December 2024.
- Fengshuo Bai, Mingzhi Wang, Zhaowei Zhang, Boyuan Chen, Yinda Xu, Ying Wen, and Yaodong Yang. Efficient Model-agnostic Alignment via Bayesian Persuasion, May 2024.
- Abhijit V Banerjee. A simple model of herd behavior. *The Quarterly Journal of Economics*, 107(3):797–817, 1992.
- Larry M. Bartels. *Presidential primaries and the dynamics of public choice*. Princeton, N.J. : Princeton University Press, 1988.
- Chris T Bauch and Samit Bhattacharyya. Evolutionary game theory and social learning can determine how vaccine scares unfold. *PLoS computational biology*, 8(4), 2012.
- Dirk Bergemann and Stephen Morris. Information Design: A Unified Perspective. *Journal of Economic Literature*, 57(1):44–95, March 2019. ISSN 0022-0515. doi: 10.1257/jel.20181489. URL <https://www.aeaweb.org/articles?id=10.1257/jel.20181489>.
- Martino Bernasconi, Matteo Castiglioni, Alberto Marchesi, and Mirco Mutti. Persuading Farsighted Receivers in MDPs: The Power of Honesty. *Advances in Neural Information Processing Systems*, 36:14987–15014, December 2023.
- Sujay Bhatt and Vikram Krishnamurthy. Controlled Sequential Information Fusion With Social Sensors. *IEEE Transactions on Automatic Control*, 66(12):5893–5908, December 2021.
- Sushil Bikhchandani, David Hirshleifer, and Ivo Welch. A theory of fads, fashion, custom, and cultural change as informational cascades. *Journal of Political Economy*, 100(5):992–1026, 1992.
- Sushil Bikhchandani, David Hirshleifer, Omer Tamuz, and Ivo Welch. Information Cascades and Social Learning. *Journal of Economic Literature*, 62(3):1040–1093, September 2024. ISSN 0022-0515, 2328-8175. doi: 10.1257/jel.20241472. URL <https://pubs.aeaweb.org/doi/10.1257/jel.20241472>.

-
- Ilai Bistriz, Nasimeh Heydaribeni, and Achilleas Anastasopoulos. Informational cascades with nonmyopic agents. *IEEE Transactions on Automatic Control*, 67, 9 2022.
- William J Brady, Killian McLoughlin, Tuan N Doan, and Molly J Crockett. How social learning amplifies moral outrage expression in online social networks. *Science Advances*, 7(33), 2021.
- Carlos Carrasco-Farre. Large Language Models are as persuasive as humans, but how? About the cognitive effort and moral-emotional language of LLM arguments, April 2024. URL <http://arxiv.org/abs/2404.09329>. arXiv:2404.09329 [cs].
- Matteo Castiglioni, Andrea Celli, Alberto Marchesi, and Nicola Gatti. Online Bayesian Persuasion. In *Advances in Neural Information Processing Systems*, volume 33, pp. 16188–16198. Curran Associates, Inc., 2020. URL <https://proceedings.neurips.cc/paper/2020/hash/ba5451d3c91a0f982f103cdbc249bc78-Abstract.html>.
- Matteo Castiglioni, Alberto Marchesi, Andrea Celli, and Nicola Gatti. Multi-Receiver Online Bayesian Persuasion. In *Proceedings of the 38th International Conference on Machine Learning*, pp. 1314–1323. PMLR, July 2021.
- Kenneth Chan, Gary Charness, Chetan Dave, and J. Lucas Reddinger. On Prior Confidence and Belief Updating, May 2025. URL <http://arxiv.org/abs/2412.10662>. arXiv:2412.10662 [econ].
- Zirui Cheng and Jiaxuan You. Towards Strategic Persuasion with Language Models, September 2025.
- Krishna Dasaratha and Kevin He. Learning from viral content. *arXiv preprint arXiv:2210.01267*, 2022.
- Danica Dillion, Niket Tandon, Yuling Gu, and Kurt Gray. Can AI language models replace human participants? *Trends in Cognitive Sciences*, 27(7):597–600, July 2023. ISSN 1364-6613, 1879-307X. doi: 10.1016/j.tics.2023.04.008. URL [https://www.cell.com/trends/cognitive-sciences/abstract/S1364-6613\(23\)00098-0](https://www.cell.com/trends/cognitive-sciences/abstract/S1364-6613(23)00098-0). Publisher: Elsevier.
- Paul Duetting, Safwan Hossain, Tao Lin, Renato Paes Leme, Sai Srivatsa Ravindranath, Haifeng Xu, and Song Zuo. Information Design With Large Language Models, September 2025. URL <http://arxiv.org/abs/2509.25565>. arXiv:2509.25565 [cs].
- Maxim Egorov, Zachary N. Sunberg, Edward Balaban, Tim A. Wheeler, Jayesh K. Gupta, and Mykel J. Kochenderfer. POMDPs.jl: A framework for sequential decision making under uncertainty. *Journal of Machine Learning Research*, 18(26):1–5, 2017.
- Erik Eyster and Matthew Rabin. Naïve herding in rich-information settings. *American Economic Journal: Microeconomics*, 2:221–243, 11 2010.
- Shen Gao, Jiabao Fang, Quan Tu, Zhitao Yao, Zhumin Chen, Pengjie Ren, and Zhaochun Ren. Generative News Recommendation. In *Proceedings of the ACM Web Conference 2024*, WWW ’24, pp. 3444–3453, New York, NY, USA, May 2024. Association for Computing Machinery. ISBN 979-8-4007-0171-9. doi: 10.1145/3589334.3645448. URL <https://dl.acm.org/doi/10.1145/3589334.3645448>.
- Yuan Gao, Dokyun Lee, Gordon Burtch, and Sina Fazelpour. Take Caution in Using LLMs as Human Surrogates: Scylla Ex Machina, January 2025. URL <http://arxiv.org/abs/2410.19599>. arXiv:2410.19599 [econ].
- Kenneth T Gillingham and Bryan Bollinger. Social learning and solar photovoltaic adoption. *Management Science*, 67(11):7091–7112, 2021.
- Douglas Guilbeault, Joshua Becker, and Damon Centola. Social learning and partisan bias in the interpretation of climate trends. *Proceedings of the National Academy of Sciences*, 115(39):9714–9719, 2018.
- Keegan Harris, Nicole Immorlica, Brendan Lucier, and Aleksandrs Slivkins. Algorithmic Persuasion Through Simulation, February 2025. URL <http://arxiv.org/abs/2311.18138>. arXiv:2311.18138 [cs].
- Claudia E. Haupt and Mason Marks. AI-Generated Medical Advice—GPT and Beyond. *JAMA*, 329(16): 1349–1350, April 2023. ISSN 0098-7484. doi: 10.1001/jama.2023.5321. URL <https://doi.org/10.1001/jama.2023.5321>.
- K. Hinderer. Lipschitz Continuity of Value Functions in Markovian Decision Processes. *Mathematical Methods of Operations Research*, 62(1):3–22, September 2005.
- John J. Horton. Large Language Models as Simulated Economic Agents: What Can We Learn from Homo Silicus?, April 2023. URL <https://www.nber.org/papers/w31122>.

- Bright Huo, Amy Boyle, Nana Marfo, Wimonchat Tangamornsuksan, Jeremy P. Steen, Tyler McKechnie, Yung Lee, Julio Mayol, Stavros A. Antoniou, Arun James Thirunavukarasu, Stephanie Sanger, Karim Ramji, and Gordon Guyatt. Large Language Models for Chatbot Health Advice Studies: A Systematic Review. *JAMA Network Open*, 8(2):e2457879, February 2025. ISSN 2574-3805. doi: 10.1001/jamanetworkopen.2024.57879. URL <https://doi.org/10.1001/jamanetworkopen.2024.57879>.
- Cameron R. Jones and Benjamin K. Bergen. Lies, Damned Lies, and Distributional Language Statistics: Persuasion and Deception with Large Language Models, December 2024. URL <http://arxiv.org/abs/2412.17128>. arXiv:2412.17128 [cs].
- Lodewijk Kallenberg. Markov decision processes. *Lecture Notes. University of Leiden*, 428, 2011.
- Emir Kamenica and Matthew Gentzkow. Bayesian persuasion. *American Economic Review*, 101(6):2590–2615, 2011.
- Vikram Krishnamurthy. Quickest Detection POMDPs With Social Learning: Interaction of Local and Global Decision Makers. *IEEE Transactions on Information Theory*, 58(8):5563–5587, August 2012.
- LG AI Research. Exaone deep: Reasoning enhanced language models. *arXiv preprint arXiv:2503.12524*, 2025.
- Minqian Liu, Zhiyang Xu, Xinyi Zhang, Heajun An, Sarvech Qadir, Qi Zhang, Pamela J. Wisniewski, Jin-Hee Cho, Sang Won Lee, Ruoxi Jia, and Lifu Huang. LLM Can be a Dangerous Persuader: Empirical Study of Persuasion Safety in Large Language Models, April 2025. URL <http://arxiv.org/abs/2504.10430>. arXiv:2504.10430 [cs].
- Nunzio Lorè and Babak Heydari. Strategic behavior of large language models and the role of game structure versus contextual framing. *Scientific Reports*, 14(1), August 2024. ISSN 2045-2322. doi: 10.1038/s41598-024-69032-z. URL <https://www.nature.com/articles/s41598-024-69032-z>. Publisher: Springer Science and Business Media LLC.
- Yaw Nyarko. On the convexity of the value function in Bayesian optimal control problems. *Economic Theory*, 4(2):303–309, March 1994.
- Kenneth Payne and Baptiste Alloui-Cros. Strategic Intelligence in Large Language Models: Evidence from evolutionary Game Theory, July 2025. URL <http://arxiv.org/abs/2507.02618>. arXiv:2507.02618 [cs].
- Yujin Potter, Shiyang Lai, Junsol Kim, James Evans, and Dawn Song. Hidden Persuaders: LLMs’ Political Leaning and Their Influence on Voters, November 2024.
- Martin L. Puterman. Chapter 8 markov decision processes. In *Stochastic Models*, volume 2 of *Handbooks in Operations Research and Management Science*. Elsevier, 1990.
- Maya Raifer, Guy Rotman, Reut Apel, Moshe Tennenholtz, and Roi Reichart. Designing an Automatic Agent for Repeated Language-based Persuasion Games. *Transactions of the Association for Computational Linguistics*, 10:307–324, March 2022.
- Neel Rao, Markus M Mobius, and Tanya Rosenblat. Social networks and vaccination decisions. *FRB of Boston*, 2007.
- Alexander Rogiers, Sander Noels, Maarten Buyl, and Tijn De Bie. Persuasion with Large Language Models: a Survey, November 2024. URL <http://arxiv.org/abs/2411.06837>. arXiv:2411.06837 [cs].
- Jens Rupperecht, Georg Ahnert, and Markus Strohmaier. Prompt Perturbations Reveal Human-Like Biases in LLM Survey Responses, July 2025. URL <http://arxiv.org/abs/2507.07188>. arXiv:2507.07188 [cs].
- Dr.-Ing. Martin Schiele, Yannick Gittmann, Stefan Ilchmann, Ante Gojsalić, Dominik Jurinčić, and Phylis Klempt. Voting Advice Applications: Implementation of RAG-supported LLMs, July 2024. URL <https://www.techrxiv.org/users/802685/articles/1187651-voting-advice-applications-implementation-of-rag-supported-llms?commit=c540a54749906f8ac761093f5dff0f514291d4b6>.
- James E. Smith and Kevin F. McCardle. Structural Properties of Stochastic Dynamic Programs. *Operations Research*, 50(5):796–809, October 2002.
- Lones Smith and Peter Sørensen. Pathological outcomes of observational learning. *Econometrica*, 68:371–398, 2000.

- Lones Smith, Peter Norman Sørensen, and Jianrong Tian. Informational Herding, Optimal Experimentation, and Contrarianism. *The Review of Economic Studies*, 88(5):2527–2554, September 2021. ISSN 0034-6527, 1467-937X. doi: 10.1093/restud/rdab001. URL <https://academic.oup.com/restud/article/88/5/2527/6149316>.
- Shih-Tang Su, Vijay G. Subramanian, and Grant Schoenebeck. Bayesian Persuasion in Sequential Trials. In Michal Feldman, Hu Fu, and Inbal Talgam-Cohen (eds.), *Web and Internet Economics*, pp. 22–40, Cham, 2022. Springer International Publishing.
- Jeremy G Weber. Social learning and technology adoption: The case of coffee pruning in peru. *Agricultural Economics*, 43:73–84, 2012.
- Xupeng Wei and Achilleas Anastasopoulos. Social learning with a self-interested coordinator. In *2022 IEEE 61st Conference on Decision and Control (CDC)*, pp. 3137–3144. IEEE, 2022.
- Marcus Williams, Micah Carroll, Adhyyan Narang, Constantin Weisser, Brendan Murphy, and Anca Dragan. On Targeted Manipulation and Deception when Optimizing LLMs for User Feedback, February 2025. URL <http://arxiv.org/abs/2411.02306>. arXiv:2411.02306 [cs].
- Jibang Wu, Zixuan Zhang, Zhe Feng, Zhaoran Wang, Zhuoran Yang, Michael I. Jordan, and Haifeng Xu. Sequential Information Design: Markov Persuasion Process and Its Efficient Reinforcement Learning, February 2022.
- Shuo Wu, Pawan Poojary, and Randall Berry. Observational Learning with a Budget, April 2025. URL <http://arxiv.org/abs/2504.19396>. arXiv:2504.19396 [cs].
- Tianyi Zhang, Faisal Ladhak, Esin Durmus, Percy Liang, Kathleen McKeown, and Tatsunori B. Hashimoto. Benchmarking Large Language Models for News Summarization. *Transactions of the Association for Computational Linguistics*, 12:39–57, January 2024a. ISSN 2307-387X. doi: 10.1162/tacl_a_00632. URL https://doi.org/10.1162/tacl_a_00632.
- Yadong Zhang, Shaoguang Mao, Tao Ge, Xun Wang, Adrian de Wynter, Yan Xia, Wenshan Wu, Ting Song, Man Lan, and Furu Wei. LLM as a Mastermind: A Survey of Strategic Reasoning with Large Language Models, April 2024b. URL <http://arxiv.org/abs/2404.01230>. arXiv:2404.01230 [cs].

A APPLICATIONS ELABORATED

Consider a sequence of individuals (agents in our terminology) deciding whether to patronize a business or service provider (e.g., a restaurant, contractor, realtor, etc.) while a recommendation system (the planner) supplies information about that business. The binary state is whether the business is good, and this state is unknown. In the altruistic case, the recommendation system and the agents share the same state-contingent objective: agents should patronize the business if and only if it is good, and the planner wants to induce exactly those correct decisions.

The recommendation system controls the precision margin by choosing between a highly targeted ad and a generic or confusing one. A precise ad presents strengths and weaknesses in a form matched to the agent’s background, making the signal more informative, comprehensible, and relatable. Higher precision therefore increases the probability of the correct action, and social welfare rises as more agents match actions to the true state. Precision is costly, however, because tailoring requires information about the agent’s background. The recommendation system therefore chooses precisions to maximize social welfare net of personalization costs.

A biased planner has a different comparison target. Consider voters (agents) who are motivated to support the candidate most likely to win in their community. The desire to “back a winner” has been shown to influence electoral outcomes, including the outsized role of early US primary states (see Bartels (1988)). The planner is one candidate’s campaign, whose objective is to increase support for that candidate. The binary state is whether the candidate is winning; an agent who knew the state would back the candidate if and only if he were winning.

Each agent can infer whether the candidate is winning if she accurately understands his stances and policies, but the communication format affects how informative the message is. Some agents may learn most from audio-visual material, others from long-form articles with facts and citations, others from brief focused content, and others from material in their native language. These differences create a measurable precision margin for campaign messages.

The campaign’s digital ads are private signals with chosen precision. A high-precision ad accurately conveys the relevant state by helping the agent understand the candidate’s policies, record, and character; if the candidate is in fact losing locally, such precision may make the agent oppose him. A low-precision ad, by contrast, may obscure the unfavorable state through vagueness or confusing presentation and thereby increase the probability

that the agent backs the candidate. The biased planner can therefore have an incentive to decrease precision when it expects unfavorable information. Both increasing and decreasing precision are costly because either form of tailoring requires knowing how the agent processes information; for example, making a message confusing for an agent who prefers brief focused content may require intentionally choosing a long, verbose format. The planner chooses precisions to maximize expected support net of these tailoring costs.

B SOCIAL LEARNING DYNAMICS

For reference, the timing at step i is as follows:

1. The planner determines q_i based on $\mathcal{H}_i$ and, q_i is observed by all agents.
2. Agent i receives private signal s_i such that $\mathbb{P}\{s_i = \omega\} = q_i$.
3. Agent i takes action a_i based on the observable history and their private signal with the aim of maximizing their own utility.
4. All other agents (and the planner) observe a_i and the public belief at the end of step i , b_{i+1} is updated accordingly.

Because q_i is selected as a function of $\mathcal{H}_i$, conditioning on both q_i and $\mathcal{H}_i$ makes the conditioning on q_i redundant. We use this simplification throughout the derivations.

B.1 PRIVATE BELIEF UPDATE

We next derive the agent's posterior and action rule. Let $\alpha_i := \frac{b_i}{1-b_i}$ be the public likelihood ratio implied by observed actions. Agent i combines this public evidence with her private signal s_i to form $\tilde{b}_i = \mathbb{P}\{\omega = G | \mathcal{H}_i, q_i, s_i\}$ and private likelihood ratio $\tilde{\alpha}_i := \frac{\tilde{b}_i}{1-\tilde{b}_i}$. The following calculation shows exactly how the posterior changes with signal realization and precision.

1. If private signal is $s_i = G$, then

$$\begin{aligned} \tilde{\alpha}_i &= \frac{\tilde{b}_i}{1-\tilde{b}_i} = \frac{\mathbb{P}\{\omega = G | \mathcal{H}_i, q_i, s_i = G\}}{\mathbb{P}\{\omega = B | \mathcal{H}_i, q_i, s_i = G\}} \\ &= \frac{\mathbb{P}\{s_i = G | \omega = G, \mathcal{H}_i, q_i\} \mathbb{P}\{\omega = G | \mathcal{H}_i, q_i\}}{\mathbb{P}\{s_i = G | \omega = B, \mathcal{H}_i, q_i\} \mathbb{P}\{\omega = B | \mathcal{H}_i, q_i\}} \\ &= \frac{q_i}{1-q_i} \frac{b_i}{1-b_i} \end{aligned} \quad (8)$$

$$= \frac{q_i}{1-q_i} \alpha_i \quad (9)$$

In Equation (8), conditional independence of signals given ω implies $\mathbb{P}\{s_i = G | \omega = G, \mathcal{H}_i\} = \mathbb{P}\{s_i = G | \omega = G\} = q_i$, with the analogous substitution under $\omega = B$. We also use that q_i is determined by $\mathcal{H}_i$.

2. If private signal is $s_i = B$, following the logic above,

$$\tilde{\alpha}_i = \frac{1-q_i}{q_i} \alpha_i \quad (10)$$

Let $y(b, q)$ and $z(b, q)$ denote the probabilities of observing signals G and B given public belief b and precision q . Conditioning on ω and applying the law of total probability gives

$$\begin{aligned} y(b_i, q_i) &= \mathbb{P}\{s_i = G | \mathcal{H}_i\} \\ &= \sum_{k \in \{G, B\}} \mathbb{P}\{s_i = G | \mathcal{H}_i, \omega = k\} \mathbb{P}\{\omega = k | \mathcal{H}_i\} \\ &= q_i b_i + (1-q_i)(1-b_i) = 1 + 2q_i b_i - q_i - b_i \end{aligned} \quad (11)$$

Similarly,

$$z(b_i, q_i) = q_i + b_i - 2q_i b_i \quad (12)$$

Combining Equation (9), Equation (10), Equation (11), and Equation (12) yields the following:

$$\tilde{b}_i = \mathbb{P}(\omega = G | \mathcal{H}_i, q_i, s_i) = f(b_i, q_i, s_i) = \begin{cases} \frac{q_i}{y(b_i, q_i)} b_i & s_i = G \\ \frac{1-q_i}{z(b_i, q_i)} b_i & s_i = B \end{cases} \quad (13)$$

If $s_i = G$, then $y(b_i, q_i) = \mathbb{P}(s_i = G | \mathcal{H}_i)$ is nonzero; similarly, if $s_i = B$, then $z(b_i, q_i)$ is nonzero. The posterior expression is therefore well-defined for realized signals.

B.2 AGENT ACTIONS AND PUBLIC BELIEF

Agent i chooses G when $\tilde{b}_i > 0.5$ (equivalently, $\tilde{\alpha}_i > 1$), chooses B when $\tilde{b}_i < 0.5$ (equivalently, $\tilde{\alpha}_i < 1$), and follows s_i under the tie-breaking rule. The action and precision q_i are then observed by all later agents and determine the next public belief.

When $b_i < 1 - q_i$, the public likelihood ratio satisfies $\alpha_i < \frac{1-q_i}{q_i}$. Thus, if $s_i = G$, Equation (9) gives $\tilde{\alpha}_i < 1$; if $s_i = B$, Equation (10) gives $\tilde{\alpha}_i < \frac{(1-q_i)^2}{q_i^2} \leq 1$ because $q_i \geq 0.5$. Hence $a_i = B$ for every signal realization. Symmetrically, when $b_i > q_i$, the action is $a_i = G$. These are precisely the regions where public evidence dominates the private-signal precision.

In the middle region, $1 - q_i \leq b_i \leq q_i$, the private signal is pivotal. If $s_i = G$, then Equation (8) and $b_i \geq 1 - q_i$ imply $\tilde{\alpha}_i \geq 1$; if $s_i = B$, then Equation (10) and $b_i \leq q_i$ imply $\tilde{\alpha}_i \leq 1$. With the tie-breaking rule, the action equals the signal, $a_i = s_i$. This yields Equation (2), reproduced below.

$$a_i = \begin{cases} s_i & 1 - q_i \leq b_i \leq q_i \\ G & q_i < b_i \\ B & q_i < 1 - b_i \end{cases}$$

Therefore, agent i 's action is informative only in the first case of Equation (2), where future agents can infer the private signal exactly. In that case, public belief updates to the same value as the private posterior in Equation (13). In the latter two cases, the action is determined by the public belief alone and adds no state information, so public belief remains unchanged. This yields Equation (3), reproduced below:

$$b_{i+1} = f(b_i, q_i) = \begin{cases} \tilde{b}_i & 1 - q_i \leq b_i \leq q_i \\ b_i & \text{o.w.} \end{cases} \quad (14)$$

B.3 MARKOVIANITY AND SUFFICIENCY OF PUBLIC BELIEF

We now show that b_i is a sufficient statistic for the action distribution. Specifically, $\mathbb{P}\{a_i = a | \mathcal{H}_i, q_i\}$ depends on the history only through (b_i, q_i) for $a \in \{G, B\}$. Hence

$$\mathbb{P}(a_i = a | \mathcal{H}_i, q_i) = \mathbb{P}(a_i = a | b_i, q_i) \quad (15)$$

For the second and third cases of Equation (2), the action is deterministic given (b_i, q_i) , so the claim is immediate. It remains to verify the informative first case.

$$\begin{aligned} \mathbb{P}(a_i = G | \mathcal{H}_i, q_i) &= \sum_{w \in \{G, B\}} \mathbb{P}(\omega = w | \mathcal{H}_i, q_i) \mathbb{P}(a_i = G | \mathcal{H}_i, q_i, \omega = w) \\ &= b_i \mathbb{P}(a_i = G | \mathcal{H}_i, q_i, \omega = G) \\ &\quad + (1 - b_i) \mathbb{P}(a_i = G | \mathcal{H}_i, q_i, \omega = B) \\ &= b_i \mathbb{P}(s_i = G | b_i, q_i, \omega = G) \\ &\quad + (1 - b_i) \mathbb{P}(s_i = G | b_i, q_i, \omega = B) \\ &= b_i q_i + (1 - b_i)(1 - q_i) \end{aligned} \quad (16)$$

The first equality applies the law of total probability. The second uses that q_i is a function of $\mathcal{H}_i$, so $\mathbb{P}(\omega = w | \mathcal{H}_i, q_i)$ is summarized by b_i . The third uses the informative-region condition $1 - q_i \leq b_i \leq q_i$, where $a_i = s_i$, together with conditional independence of s_i from the history given ω .

This establishes the Markov property stated below.

Theorem 6. *Markovianity of the Public Belief Process*

$$b_{i+1} = \mathbb{P}\{\omega = G | b_i, q_i, a_i\} \quad (17)$$

In other words, b_i captures the entire information pertinent to the update of b_{i+1} that is contained in $\mathcal{H}_i$.

Proof. Note that, from step (1) in Appendix B, q_i is a function of $\mathcal{H}_i$. It suffices to show

$$b_{i+1} = \frac{\mathbb{P}\{a_i \cap \omega = G | b_i, q_i\}}{\mathbb{P}\{a_i | b_i, q_i\}}$$

We will prove the claim via induction on i .

For our base case, consider $i = 1$ so that $\mathcal{H}_2 = \{b_1, (q_1, a_1)\}$ and $b_i = b_1$. An application of Bayes' Theorem provides the following:

$$\begin{aligned} b_2 &= \frac{\mathbb{P}\{\omega = G|b_1, q_1\} \mathbb{P}\{a_1|\omega = G, b_1, q_1\}}{\mathbb{P}\{a_1|b_1, q_1\}} = \frac{b_1 \mathbb{P}\{a_1|\omega = G, b_1, q_1\}}{\mathbb{P}\{a_1|b_1, q_1\}} \\ &= \frac{b_1 \mathbb{P}\{a_1 \cap \omega = G|b_1, q_1\}}{\mathbb{P}\{a_1|b_1, q_1\} \mathbb{P}\{\omega = G|b_1, q_1\}} \\ &= \frac{\mathbb{P}\{a_1 \cap \omega = G|b_1, q_1\}}{\mathbb{P}\{a_1|b_1, q_1\}} \end{aligned}$$

where the final step follows as the unconditioned belief of ω is equal to the prior belief b_1 .

Now, assume the claim holds for all $i \leq n$.

$$\begin{aligned} b_{n+1} &= \mathbb{P}\{\omega = G|\mathcal{H}_n, a_n, q_n\} = \frac{\mathbb{P}\{\omega = G|\mathcal{H}_n, q_n\} \mathbb{P}\{a_n|\omega = G, \mathcal{H}_n, q_n\}}{\mathbb{P}\{a_n|\mathcal{H}_n, q_n\}} \\ &= \frac{b_n \mathbb{P}\{a_n|\omega = G, \mathcal{H}_n, q_n\}}{\mathbb{P}\{a_n|\mathcal{H}_n, q_n\}} \\ &= \frac{b_n \mathbb{P}\{a_n \cap \omega = G|\mathcal{H}_n, q_n\}}{\mathbb{P}\{a_n|b_n, q_n\} \mathbb{P}\{\omega = G|\mathcal{H}_n, q_n\}} \\ &= \frac{\mathbb{P}\{a_n \cap \omega = G|b_n, q_n\}}{\mathbb{P}\{a_n|b_n, q_n\}} \end{aligned}$$

where the final two steps follow from Equation (15) and the definition of b_n .

This completes the proof. $\square$

Note that we have shown that the public belief is Markovian adapting the arguments in Banerjee (1992) and Bikhchandani et al. (1992). The modifications were required because of the differences in selections of the agents' precisions caused by the introduction of the central planner.

C PROOFS

C.1 DERIVATION OF EQUATION (4)

Lemma 7.

$$\mathbb{P}(a_i \neq \omega|b_i, q_i) = \min(b_i, 1 - b_i, 1 - q_i)$$

Proof. We verify the expression by comparing the mistake probability in each region of Equation (2). In the second region, the action is fixed at G , so

$$\mathbb{P}(a_i \neq \omega|b_i, q_i) = \mathbb{P}(\omega \neq G|b_i, q_i) = 1 - b_i$$

The third region is symmetric: the fixed action is B , so

$$\mathbb{P}(a_i \neq \omega|b_i, q_i) = \mathbb{P}(\omega \neq B|b_i, q_i) = b_i$$

In the informative middle region, $1 - q_i \leq b_i \leq q_i$, the action equals the signal by Equation (2); the error probability is therefore

$$\begin{aligned} \mathbb{P}(a_i \neq \omega | b_i, q_i) &= \sum_{a \in \{G, B\}} \mathbb{P}(a_i = a | b_i, q_i) \mathbb{P}(\omega \neq a | b_i, q_i, a_i = a) \\ &= \mathbb{P}(s_i = G | b_i, q_i) \mathbb{P}(\omega \neq G | b_i, q_i, s_i = G) \\ &\quad + \mathbb{P}(s_i = B | b_i, q_i) \mathbb{P}(\omega \neq B | b_i, q_i, s_i = B) \end{aligned} \quad \text{Eqn 2} \quad (18)$$

$$\begin{aligned} &= y(b_i, q_i) \mathbb{P}(\omega = B | b_i, q_i, s_i = G) \\ &\quad + z(b_i, q_i) \mathbb{P}(\omega = G | b_i, q_i, s_i = B) \end{aligned} \quad \text{Eqns 11, 12, Thm 6} \quad (19)$$

$$= \begin{cases} y(0, q_i) & b_i = 0 \\ y(b_i, q_i)(1 - \frac{b_i q_i}{y(b_i, q_i)}) + z(b_i, q_i) \frac{b_i(1 - q_i)}{z(b_i, q_i)} & b_i \neq 0 \end{cases} \quad \text{Eqn 13} \quad (20)$$

$$= \begin{cases} 1 - q_i & b_i = 0 \\ y(b_i, q_i) - b_i q_i + b_i(1 - q_i) & b_i \neq 0 \end{cases}$$

$$= \begin{cases} 1 - q_i & b_i = 0 \\ y(b_i, q_i) - b_i q_i + b_i - b_i q_i & b_i \neq 0 \end{cases}$$

$$= \begin{cases} 1 - q_i & b_i = 0 \\ 1 - q_i & b_i \neq 0 \end{cases}$$

Taking the minimum across the three action regions gives the stated expression. $\square$

C.2 PROOF OF THEOREM 1: MYOPIC ALTRUISTIC POLICY

Applying Equation (4) to $r_A(b, q)$ in Equation (5) separates the marginal reward of precision into the uninformative and informative regions:

$$\frac{\partial r_A(b, q)}{\partial q} = \begin{cases} -\dot{\beta}(q) & q < \max(b, 1 - b) \\ -\dot{\beta}(q) + C & q \geq \max(b, 1 - b) \end{cases} \quad (21)$$

As stated in Section 3.2, the altruistic planner can be restricted to $q \in [p, 1]$ because welfare is increasing in precision (see Appendix C.9); the baseline cost is normalized by $\beta(p) = 0$.

Because $\beta(\cdot)$ is increasing, the first case of Equation (21) favors the lowest available precision, p , whenever $p < \max(b, 1 - b)$. In that region, additional precision still does not affect the action, so its marginal benefit is zero while its cost is positive. In the latter case, $\beta(\cdot)$ is increasing and concave, so $\dot{\beta}(\cdot)$ is positive and decreasing. Therefore, if the optimum lies in the informative region, the reward comparison is maximized at the boundary $q = 1$.

It remains to identify when the full-precision boundary dominates the baseline.

$$\begin{aligned} r_A(b, p) &< r_A(b, 1) \\ -C \min(b, 1 - b, 1 - p) &< -\beta(1) \\ \min(b, 1 - b, 1 - p) &> \frac{\beta(1)}{C} \end{aligned}$$

If $1 - p \leq \frac{\beta(1)}{C}$, full precision never clears the cost-benefit threshold. Otherwise, the inequality identifies the belief interval $b \in \left[\frac{\beta(1)}{C}, 1 - \frac{\beta(1)}{C} \right]$ where precision 1 is myopically optimal, proving Theorem 1. $\blacksquare$

C.3 PROOF OF THEOREM 2: ALTRUISTIC VALUE FUNCTION CONVEXITY

Proof Sketch The proof tracks convexity stage by stage. Specifically, we show inductively that the expected k -th stage reward—the planner’s expected utility from the control and action of the k -th agent—is convex in public belief. The instantaneous reward in Equation (5) gives the base case.

The main technical issue is the belief update’s discrete support. Although public belief lies in an uncountable state space, each update has at most two possible next beliefs. We exploit this by defining a decision tree: once an initial belief and policy are fixed, each node has two children corresponding to the next signal realization. The root represents the expected instantaneous reward at time 1, and level k contains 2^{k-1} nodes associated with the possible sequences of $k - 1$ signal realizations.

The induction step compares a parent node at level $(k - 1)$ with its two children at level k . The difficulty is that actions depend on public belief: the same precision and same signal realization can induce different actions from agents who start at different beliefs (see Equation (2)). Standard Markov-process convexity arguments therefore do not apply directly. The binary Bayesian update is nevertheless favorable because it preserves the martingale property $\mathbb{E}[b_{i+1}] = b_i$; together with convex instantaneous rewards, this closes the convexity comparison for the child nodes and completes the induction.

Preliminaries We first prove the following lemma containing a few useful properties of the optimal value function $V_A^*(\cdot)$.

Lemma 8.

$V_A^*(\cdot)$ is non-positive, $V_A^*(b) = V_A^*(1 - b) \forall b \in [0, 1]$, and $V_A^*(0) = V_A^*(1) = 0$ with $\pi_A^*(0) = \pi_A^*(1) = p$.

Proof.

Non-positivity

$$r(b_i, q_i) \leq 0 \quad \forall b_i, q_i \rightarrow V_A^* \leq 0$$

Symmetry

Let $n_i = 1 - b_i$ and let Λ_n and Λ_b be the decision trees rooted at states n_0 and b_0 , respectively. Showing that Λ_n and Λ_b are symmetric suffices to show the symmetry of $V^*(\cdot)$. This will be accomplished by showing that the expected rewards, belief state values, and branching probabilities of the trees are reflections of one another.

First consider the instantaneous reward function r_A . Note that the decision function of the agents when choosing their actions is symmetric from Equation (2), thus, $\mathbb{P}\{a_i = G|b_i\} = \mathbb{P}\{a_i = B|n_i\}$. Applying this, we obtain

$$\begin{aligned} r(n_i, q_i) &= -\beta(q_i) - Cn_i \mathbb{P}\{a_i = B|\omega = G, n_i\} - C(1 - n_i) \mathbb{P}\{a_i = G|\omega = B, n_i\} \\ &= -\beta(q_i) - C(1 - b_i) \mathbb{P}\{a_i = B|\omega = G, n_i\} - Cb_i \mathbb{P}\{a_i = G|\omega = B, n_i\} \\ &= -\beta(q_i) - C(1 - b_i) \mathbb{P}\{a_i = G|\omega = B, b_i\} - Cb_i \mathbb{P}\{a_i = B|\omega = G, b_i\} \\ &= r(b_i, q_i) \end{aligned}$$

Thus, the instantaneous reward function is symmetric about 0.5.

Now consider the belief state update when beginning with belief n_i assuming signal precision q_i . we will use $\bar{x}$ and $\underline{x}$ to represent updated beliefs from x after observing action G or B , respectively. From Equation (13):

$$\begin{aligned} \bar{n}_i &= \frac{q_i}{y(n_i, q_i)} n_i = \frac{q_i}{z(b_i, q_i)} (1 - b_i) = 1 - \frac{1 - q_i}{z(b_i, q_i)} b_i = 1 - \bar{b}_i \\ \underline{n}_i &= \frac{1 - q_i}{z(n_i, q_i)} n_i = \frac{q_i}{y(b_i, q_i)} (1 - b_i) = 1 - \frac{q_i}{y(b_i, q_i)} b_i = 1 - \bar{b}_i \end{aligned}$$

Thus, the belief updates are also symmetric. Combining this with the symmetry of the reward we obtain:

$$\begin{aligned} r(\bar{n}_i) &= r(1 - \bar{b}_i) = r(\bar{b}_i) \\ r(\underline{n}_i) &= r(1 - \bar{b}_i) = r(\bar{b}_i) \end{aligned}$$

Thus, the expected rewards and belief state values of Λ_n and Λ_b are mirror images of one another.

Finally, consider the branching probabilities which we will denote as $y(b, q) = 1 + 2bq - b - q$ for observing action G and $z(b, q) = 1 - y(b, q)$ for observing action B .

$$\begin{aligned} y(n, q) &= 1 + 2nq - n - q = 1 + 2(1 - b)q - (1 - b) - q = b + q - 2bq = z(b, q) \\ z(n, q) &= 1 - y(n, q) = 1 - z(b, q) = y(b, q) \end{aligned}$$

Thus, the branching probabilities of Λ_n and Λ_b are also symmetric, completing the proof.

Extremal values

We will show that $V_A^*(0) = 0$ and rely on the symmetry about 0.5 proven above to show that $V_A^*(1) = 0$.

From Equation (3), when $b_0 = 0$, $b_0^+ = b_0^- = 0$ for any value of precision. Thus, regardless of signal realizations or control actions, the agents' actions and, subsequently, the public belief, will remain unchanged. That is, $b_i = 0 \forall i$ and, from Equation (2), $a_i = B \forall i$.

From Equation (5), we can then write:

$$\begin{aligned}
V_A^*(0) &= \max_{\pi \in \Pi} \mathbb{E} \left[\sum_{i=0}^{\infty} \delta^i r(b_i, \pi(b_i)) \right] \\
&= \max_{\pi \in \Pi} \mathbb{E} \left[\sum_{i=0}^{\infty} \delta^i r(0, \pi(0)) \right] \\
&= \max_{\pi \in \Pi} \mathbb{E} \left[\sum_{i=0}^{\infty} \delta^i [-\beta(\pi(0)) - C b_i \mathbb{P}\{a_i = B | \omega = G, b_i = 0\} - C(1 - b_i) \mathbb{P}\{a_i = G | \omega = B, b_i = 0\}] \right] \\
&= \max_{\pi \in \Pi} \mathbb{E} \left[\sum_{i=0}^{\infty} \delta^i [-\beta(q_i)] \right]
\end{aligned} \tag{22}$$

Thus, to attain $V_A^*(0)$, it suffices to choose $\pi(0)$ such that $\beta(\pi(0))$ is minimized for each i . Recall that $\beta(\cdot)$ is increasing on $[0.5, 1]$ and takes its minimum at $\beta(p) = 0$. Thus, by choosing $q_i = p \forall i$, Equation 22 yields $V_A^*(0) = 0$ with corresponding optimal policy $\pi_A^*(0) = p$.

The aforementioned symmetry about 0.5, shows that $V_A^*(1) = 0$, completing the proof. $\square$

We now introduce the concept of *decision trees* for Markovian and more general policies. A policy can be represented by a decision tree, which is a complete binary tree as in figure 3. Every node of the decision tree represents the state of the system, i.e., the corresponding belief value when the controller sends its private signal to the agent corresponding to the level of the node. The branches represent the actions taken by the agent corresponding to the level once it receives its private signal. The controller's policies (i.e. choice of precision) determine the probabilities of the actions from each node and the belief resulting from the action. The probabilities of an action at a certain node can be considered the weight of the branch. A decision tree is uniquely identified by the belief values associated with nodes and the weights of the branches. For a deterministic stationary Markovian policy, the controller's choice of precision, and therefore the probability associated with an action, is a deterministic function of the belief value of the node, regardless of the level of the node in the tree and the path to the node from the root. Thus, the decision tree of any such policy, i.e. the belief values at the nodes and the weights of the branches, is determined entirely by the value of the initial state. For an arbitrary policy of the controller, the controller's decision at a given node can depend on the entire path leading to the node.

Proof of Theorem 2 We now prove the convexity of $V_A^*(\cdot)$ with respect to public belief.

Proof.

We will first show that, for fixed precision, the instantaneous reward r_A is convex in public belief. We can rewrite r_A as follows using Lemma 7:

$$r_A(x, q) = -\beta(q) - C \min(x, 1 - x, 1 - q) \tag{23}$$

The pointwise minimum of linear functions is concave. As we are subtracting the minimum, the overall function is convex with respect to x .

We consider arbitrary x_0, m_0, n_0 such that $tm_0 + (1 - t)n_0 = x_0$ for some $t \in [0, 1]$. To show the convexity of $V_A^*(\cdot)$ we must show that $V_A^*(x_0) \leq tV_A^*(m_0) + (1 - t)V_A^*(n_0)$. By definition of the optimal value function, $V_A^*(x) \geq V_A^\pi(x)$ for any policy π . Thus, it will suffice to show that there exist $\tilde{\pi}$ such that

$$V_A^*(x_0) \leq tV_A^{\tilde{\pi}}(m_0) + (1 - t)V_A^{\tilde{\pi}}(n_0) \tag{24}$$

We refer to the decision tree of π^* when the initial state is x_0 as Λ .

We consider $\tilde{\pi}$ to be a policy of the controller that chooses the precision at each node to be the same as what Λ does at the counterpart node. The nodes of decision trees are considered to be counterparts if they can be reached via the same sequence of actions from the root (see figure 3 for illustration). Thus, the precision chosen by this policy is the same as what Λ does after the same set of actions of the agents. We refer to the decision trees for $\tilde{\pi}$ with initial state m_0 (respectively, n_0) as Λ_m (Λ_n).

We now illustrate the policy $\tilde{\pi}$ for initial values m_0 and n_0 considering the decision trees Λ_m and Λ_n . While doing so, we illustrate counterpart nodes and the weights of the branches. We will select one path of actions λ out of these decision trees by specifying set $\mathcal{G}$ as the epochs where action G is taken and $\mathcal{B}$ as its complement. Refer to figure 3 for illustration. Considering the decision tree Λ , we will refer to x_i as the state at epoch i when beginning with belief x_0 and following π^* along this path of actions. That is, x_i assumes one of 2^i possible belief values at epoch i starting from x_0 . In decision tree Λ_m (respectively, Λ_n), we use m_i (n_i) to denote the

public belief after beginning with m_0 (n_0), following the path λ . Node m_i (respectively n_i) in decision tree Λ_m (Λ_n) is the counterpart of node x_i from decision tree Λ . Note that, $\tilde{\pi}(m_i)$ can now be illustrated notationally. Specifically, $\tilde{\pi}(m_i) = \tilde{\pi}(n_i) = \pi_A^*(x_i)$. Similarly, $\tilde{\pi}$ can be defined at each node of Λ_m , considering the action taken by π^* at each counterpart node in Λ . Note that $y(x_i, \pi_A^*(x_i))$ (respectively, $z(x_i, \pi_A^*(x_i))$) is the weight of the branch corresponding to G (B) emanating from x_i in Λ . Because $\tilde{\pi}(m_i) = \tilde{\pi}(n_i) = \pi_A^*(x_i)$, $y(m_i, \tilde{\pi}(m_i)) = y(m_i, \pi_A^*(x_i))$ is the weight of the branch emanating from m_i for action G in Λ_m . This extends similarly to branches associated with action B and Λ_n . Thus, the three decision trees differ in both the values of the nodes and the weights of the branches. Since π^* is deterministic, so is $\tilde{\pi}$. Therefore, from Equation (13) and Equation (3), x_i , m_i , and n_i are deterministic quantities once the specific path of actions is specified by $\mathcal{G}$.




Figure 3: Here we depict the belief evolution in the form of three binary trees. The tree begins with state x_0 , m_0 , and n_0 , respectively, and level i contains 2^i possible states reachable after i epochs. Left branches correspond to action B leading to updated belief $\underline{x}_0$ (respectively, $\underline{m}_0$ and $\underline{n}_0$). Right branches correspond to action G leading to updated belief $\bar{x}_0$ (respectively, $\bar{m}_0$ and $\bar{n}_0$). Each branch is labeled with the corresponding probability, and the bold states indicate an example path, like λ . In the example path, $0 \in \mathcal{G}$ and $1, k-1 \in \mathcal{B}$. The bold states in each level of the three trees are counterparts.

Now let $\phi_k^\pi(b_0) = \mathbb{E}[r_A(b_k, \pi(b_k))]$ be the expected instantaneous reward at epoch k when beginning with state b_0 and applying policy π and $\phi^*(x) = \phi^{\pi^*}(x)$. Applying Fubini's Theorem, we can rewrite our value function

as

$$V_A^\pi(b_0) = \mathbb{E} \left[\sum_{i=0}^{\infty} \delta^i r_A(b_i, \pi(b_i)) \right] = \sum_{i=0}^{\infty} \delta^i \mathbb{E}[r_A(b_i, \pi(b_i))] = \sum_{i=0}^{\infty} \delta^i \phi_i^\pi(b_0) \quad (25)$$

Thus, we can rewrite (24), as

$$\sum_{i=0}^{\infty} \delta^i \phi_i^*(x_0) \leq t \left[\sum_{i=0}^{\infty} \delta^i \phi_i^{\tilde{\pi}}(m_0) \right] + (1-t) \left[\sum_{i=0}^{\infty} \delta^i \phi_i^{\tilde{\pi}}(n_0) \right] \quad (26)$$

To prove this, we show that for all $i = 0, 1, 2, \dots$

$$\phi_i^*(x_0) \leq t \phi_i^{\tilde{\pi}}(m_0) + (1-t) \phi_i^{\tilde{\pi}}(n_0) \quad (27)$$

leading directly to (26), (24), and finally the convexity of $V^*(\cdot)$ with respect to public belief.

Now let us consider general $i = k$. $\phi_{k-1}^\pi(x_0)$ will have 2^{k-1} terms each corresponding to one path of possible actions in the set $\{G, B\}^{k-2}$. Thus, at epoch $k-1$, the path λ will yield the following term in $\phi_{k-1}^*(x_0)$

$$\left(\prod_{i \in \mathcal{G}} y(x_i, \pi_A^*(x_i)) \right) \left(\prod_{i \in \mathcal{B}} z(x_i, \pi_A^*(x_i)) \right) r_A(x_{k-1}, \pi_A^*(x_{k-1})) \quad (28)$$

We now define $P_\lambda^*(x_0) = \left(\prod_{i \in \mathcal{G}} y(x_i, \pi_A^*(x_i)) \right) \left(\prod_{i \in \mathcal{B}} z(x_i, \pi_A^*(x_i)) \right)$, that is, the probability of following the specified path λ when beginning with state x_0 and following policy π_A^* . Again, we will use $\bar{x}_i$ and $\underline{x}_i$ to represent updated beliefs from x_i after observing action G or B , respectively. Thus, $\bar{x}_i$ and $\underline{x}_i$ constitute belief values of the children nodes (in the $(i+1)$ th level) of the node with belief value x_i in the i -th level. Refer to levels $k-1$ and k in figure 3a. Considering the children $\bar{x}_{k-1}$ and $\underline{x}_{k-1}$ (in the k -th level) of the node x_{k-1} (in the $(k-1)$ -th level), we obtain the following two terms in $\phi_k^*(x_0)$:

$$P_\lambda^*(x_0) \left[y(x_{k-1}, \pi_A^*(x_{k-1})) r_A(\bar{x}_{k-1}, \pi_A^*(\bar{x}_{k-1})) + z(x_{k-1}, \pi_A^*(x_{k-1})) r_A(\underline{x}_{k-1}, \pi_A^*(\underline{x}_{k-1})) \right] \quad (29)$$

Similarly, we define $P_\lambda^{\tilde{\pi}}(m_0) = \left(\prod_{i \in \mathcal{G}} y(m_i, \tilde{\pi}(m_i)) \right) \left(\prod_{i \in \mathcal{B}} z(m_i, \tilde{\pi}(m_i)) \right)$, that is, the probability of following the specified path λ when beginning with state m_0 and following policy $\tilde{\pi}$. Because $\tilde{\pi}(m_i) = \pi_A^*(x_i)$, $P_\lambda^{\tilde{\pi}}(m_0) = \left(\prod_{i \in \mathcal{G}} y(m_i, \pi_A^*(x_i)) \right) \left(\prod_{i \in \mathcal{B}} z(m_i, \pi_A^*(x_i)) \right)$. We also define $P_\lambda^{\tilde{\pi}}(n_0)$ by replacing $\{m_i\}$ with $\{n_i\}$.

The corresponding terms in $\phi_k^{\tilde{\pi}}(m_0)$ and $\phi_k^{\tilde{\pi}}(n_0)$ are

$$P_\lambda^*(m_0) \left[y(m_{k-1}, \pi_A^*(x_{k-1})) r_A(\bar{m}_{k-1}, \pi_A^*(\bar{x}_{k-1})) + z(m_{k-1}, \pi_A^*(x_{k-1})) r_A(\underline{m}_{k-1}, \pi_A^*(\underline{x}_{k-1})) \right] \quad (30)$$

$$P_\lambda^{\tilde{\pi}}(n_0) \left[y(n_{k-1}, \pi_A^*(x_{k-1})) r_A(\bar{n}_{k-1}, \pi_A^*(\bar{x}_{k-1})) + z(n_{k-1}, \pi_A^*(x_{k-1})) r_A(\underline{n}_{k-1}, \pi_A^*(\underline{x}_{k-1})) \right] \quad (31)$$

We will show that

$$\begin{aligned} & P_\lambda^*(x_0) \left[y(x_{k-1}, \pi_A^*(x_{k-1})) r_A(\bar{x}_{k-1}, \pi_A^*(\bar{x}_{k-1})) + z(x_{k-1}, \pi_A^*(x_{k-1})) r_A(\underline{x}_{k-1}, \pi_A^*(\underline{x}_{k-1})) \right] \\ & \leq t P_\lambda^*(m_0) \left[y(m_{k-1}, \pi_A^*(x_{k-1})) r_A(\bar{m}_{k-1}, \pi_A^*(\bar{x}_{k-1})) + z(m_{k-1}, \pi_A^*(x_{k-1})) r_A(\underline{m}_{k-1}, \pi_A^*(\underline{x}_{k-1})) \right] \\ & + (1-t) P_\lambda^{\tilde{\pi}}(n_0) \left[y(n_{k-1}, \pi_A^*(x_{k-1})) r_A(\bar{n}_{k-1}, \pi_A^*(\bar{x}_{k-1})) + z(n_{k-1}, \pi_A^*(x_{k-1})) r_A(\underline{n}_{k-1}, \pi_A^*(\underline{x}_{k-1})) \right] \end{aligned} \quad (32)$$

Similarly, each pair of children for each term in the $(k-1)$ -th level of the belief tree can be combined and the corresponding inequalities can be obtained following the same process. Summing the resulting inequalities will yield (27), leading to (26), (24), and the convexity of $V^*(\cdot)$.

Proof of (32)

We now focus on the latter of the two terms of (29). We will show the following two properties

$$\begin{aligned} \underline{x}_{k-1} &= t \left(\prod_{i \in \mathcal{G}} \frac{y(m_i, \pi_A^*(x_i))}{y(x_i, \pi_A^*(x_i))} \right) \left(\prod_{i \in \mathcal{B} \cup \{k-1\}} \frac{z(m_i, \pi_A^*(x_i))}{z(x_i, \pi_A^*(x_i))} \right) \underline{m}_{k-1} \\ &+ (1-t) \left(\prod_{i \in \mathcal{G}} \frac{y(n_i, \pi_A^*(x_i))}{y(x_i, \pi_A^*(x_i))} \right) \left(\prod_{i \in \mathcal{B} \cup \{k-1\}} \frac{z(n_i, \pi_A^*(x_i))}{z(x_i, \pi_A^*(x_i))} \right) \underline{n}_{k-1} \end{aligned} \quad (33)$$

$$\begin{aligned}
1 &= t \left(\prod_{i \in \mathcal{G}} \frac{y(m_i, \pi_A^*(x_i))}{y(x_i, \pi_A^*(x_i))} \right) \left(\prod_{i \in \mathcal{B} \cup \{k-1\}} \frac{z(m_i, \pi_A^*(x_i))}{z(x_i, \pi_A^*(x_i))} \right) \\
&+ (1-t) \left(\prod_{i \in \mathcal{G}} \frac{y(n_i, \pi_A^*(x_i))}{y(x_i, \pi_A^*(x_i))} \right) \left(\prod_{i \in \mathcal{B} \cup \{k-1\}} \frac{z(n_i, \pi_A^*(x_i))}{z(x_i, \pi_A^*(x_i))} \right)
\end{aligned} \tag{34}$$

Defining $\tilde{t} = t \left(\prod_{i \in \mathcal{G}} \frac{y(m_i, \pi_A^*(x_i))}{y(x_i, \pi_A^*(x_i))} \right) \left(\prod_{i \in \mathcal{B} \cup \{k-1\}} \frac{z(m_i, \pi_A^*(x_i))}{z(x_i, \pi_A^*(x_i))} \right)$, it follows from (33) and (34) that

$$\underline{x}_{k-1} = \tilde{t} \underline{m}_{k-1} + (1-\tilde{t}) \underline{n}_{k-1} \tag{35}$$

From the expressions for $\tilde{t}$, $P_\lambda^*(x_0)$, $P_\lambda^{\tilde{\pi}}(m_0)$, it follows that:

$$\tilde{t} z(x_{k-1}, \pi_A^*(x_{k-1})) P_\lambda^*(x_0) = t z(m_{k-1}, \pi_A^*(x_{k-1})) P_\lambda^{\tilde{\pi}}(m_0) \tag{36}$$

Similarly, we can show

$$\tilde{t} z(x_{k-1}, \pi_A^*(x_{k-1})) P_\lambda^*(x_0) = (1-t) z(n_{k-1}, \pi_A^*(x_{k-1})) P_\lambda^{\tilde{\pi}}(n_0) \tag{37}$$

We apply this as follows to the latter term of (29):

$$\begin{aligned}
&P_\lambda^*(x_0) z(x_{k-1}, \pi_A^*(x_{k-1})) r_A(\underline{x}_{k-1}, \pi_A^*(\underline{x}_{k-1})) \\
&\leq P_\lambda^*(x_0) z(x_{k-1}, \pi_A^*(x_{k-1})) \\
&\quad [\tilde{t} r_A(\underline{m}_{k-1}, \pi_A^*(\underline{x}_{k-1})) + (1-\tilde{t}) r_A(\underline{n}_{k-1}, \pi_A^*(\underline{x}_{k-1}))] \quad (\text{Convexity of } r, \text{Eqn 35}) \\
&= t [P_\lambda^{\tilde{\pi}}(m_0) z(m_{k-1}, \pi_A^*(x_{k-1})) r_A(\underline{m}_{k-1}, \pi_A^*(\underline{x}_{k-1}))] \\
&\quad + (1-t) [P_\lambda^{\tilde{\pi}}(n_0) z(n_{k-1}, \pi_A^*(x_{k-1})) r_A(\underline{n}_{k-1}, \pi_A^*(\underline{x}_{k-1}))] \quad (\text{Eqns 36, 37})
\end{aligned} \tag{38}$$

Applying the same argument to the first term of (29) yields:

$$\begin{aligned}
&P_\lambda^*(x_0) y(x_{k-1}, \pi_A^*(x_{k-1})) r_A(\bar{x}_{k-1}, \pi_A^*(\bar{x}_{k-1})) \\
&\leq t [P_\lambda^{\tilde{\pi}}(m_0) y(m_{k-1}, \pi_A^*(x_{k-1})) r_A(\bar{m}_{k-1}, \pi_A^*(\bar{x}_{k-1}))] \\
&\quad + (1-t) [P_\lambda^{\tilde{\pi}}(n_0) y(n_{k-1}, \pi_A^*(x_{k-1})) r_A(\bar{n}_{k-1}, \pi_A^*(\bar{x}_{k-1}))]
\end{aligned} \tag{39}$$

Summing (38) and (39) yields (32).

We now complete the proof of (32) and, therefore, the entire proof, by proving (33) and (34).

Proof of (33)

From the recurrence relation in Equation (13) and Equation (3), we can write $\underline{x}_{k-1}$, $\underline{m}_{k-1}$, and $\underline{n}_{k-1}$ as the product of all belief updates with x_0 , m_0 , and n_0 , respectively, as follows. Each term in the products corresponds to the multiplicative state update after a single action.

$$\underline{x}_{k-1} = \left(\prod_{i \in \mathcal{G}} \frac{\pi_A^*(x_i)}{y(x_i, \pi_A^*(x_i))} \right) \left(\prod_{i \in \mathcal{B} \cup \{k-1\}} \frac{1 - \pi_A^*(x_i)}{z(x_i, \pi_A^*(x_i))} \right) x_0 \tag{40}$$

$$\underline{m}_{k-1} = \left(\prod_{i \in \mathcal{G}} \frac{\pi_A^*(x_i)}{y(m_i, \pi_A^*(x_i))} \right) \left(\prod_{i \in \mathcal{B} \cup \{k-1\}} \frac{1 - \pi_A^*(x_i)}{z(m_i, \pi_A^*(x_i))} \right) m_0 \tag{41}$$

$$\underline{n}_{k-1} = \left(\prod_{i \in \mathcal{G}} \frac{\pi_A^*(x_i)}{y(n_i, \pi_A^*(x_i))} \right) \left(\prod_{i \in \mathcal{B} \cup \{k-1\}} \frac{1 - \pi_A^*(x_i)}{z(n_i, \pi_A^*(x_i))} \right) n_0 \tag{42}$$

These equations may be best understood via the illustration in figure 3. We then reorganize terms to relate $\underline{x}_{k-1}$ to $\underline{m}_{k-1}$ and $\underline{n}_{k-1}$, using the fact that $x_0 = tm_0 + (1-t)n_0$.

$$\begin{aligned}\underline{x}_{k-1} &= \left(\prod_{i \in \mathcal{G}} \frac{\pi_A^*(x_i)}{y(x_i, \pi_A^*(x_i))} \right) \left(\prod_{i \in \mathcal{B} \cup \{k-1\}} \frac{1 - \pi_A^*(x_i)}{z(x_i, \pi_A^*(x_i))} \right) [tm_0 + (1-t)n_0] \quad (\text{from 40}) \\ &= t \left(\prod_{i \in \mathcal{G}} \frac{y(m_i, \pi_A^*(x_i))}{y(x_i, \pi_A^*(x_i))} \right) \left(\prod_{i \in \mathcal{B} \cup \{k-1\}} \frac{z(m_i, \pi_A^*(x_i))}{z(x_i, \pi_A^*(x_i))} \right) \underline{m}_{k-1} \\ &\quad + (1-t) \left(\prod_{i \in \mathcal{G}} \frac{y(n_i, \pi_A^*(x_i))}{y(x_i, \pi_A^*(x_i))} \right) \left(\prod_{i \in \mathcal{B} \cup \{k-1\}} \frac{z(n_i, \pi_A^*(x_i))}{z(x_i, \pi_A^*(x_i))} \right) \underline{n}_{k-1} \quad (\text{from 41 and 42}) \quad (43)\end{aligned}$$

This proves (33).

Proof of (34)

We first consider the numerator of the first term in (34): $\prod_{i \in \mathcal{G}} y(m_i, \pi_A^*(x_i)) \prod_{i \in \mathcal{B} \cup \{k-1\}} z(m_i, \pi_A^*(x_i))$. We now prove that this can be written as $\alpha_0 m_0 + \beta_0$ for some α_0, β_0 which depend only on $\{\pi_A^*(x_i)\}_{i=1, \dots, k-1}$ and the nature of the dependence is determined by the sequence of actions along this path in stages 1 to $k-1$. We will prove this via induction working backwards from the final term (i.e. the $(k-1)$ -th term) of the product. Therefore our base case will be the final term of this product which can be either $y(m_{k-1}, \pi_A^*(x_{k-1}))$ or $z(m_{k-1}, \pi_A^*(x_{k-1}))$, depending on the preceding action (in the $(k-1)$ -th stage).

$$\begin{aligned}y(m_{k-1}, \pi_A^*(x_{k-1})) &= 1 + 2m_{k-1}\pi_A^*(x_{k-1}) - m_{k-1} - \pi_A^*(x_{k-1}) \\ &= (2\pi_A^*(x_{k-1}) - 1)m_{k-1} + (1 - \pi_A^*(x_{k-1})) \\ z(m_{k-1}, \pi_A^*(x_{k-1})) &= m_{k-1} + \pi_A^*(x_{k-1}) - 2m_{k-1}\pi_A^*(x_{k-1}) \\ &= (1 - 2\pi_A^*(x_{k-1}))m_{k-1} + (\pi_A^*(x_{k-1}))\end{aligned}$$

Thus, the final term of the product is of the form $\alpha_{k-1}m_{k-1} + \beta_{k-1}$ where α_{k-1} and β_{k-1} depend only upon $\pi_A^*(x_{k-1})$, and the nature of the dependence is determined entirely by the $(k-1)$ -th action (since this action determines whether the final term is $y(\cdot)$ or $z(\cdot)$).

For the inductive step, assume that $\prod_{i \in \mathcal{G}, i > j} y(m_i, \pi_A^*(x_i)) \prod_{i \in \mathcal{B} \cup \{k-1\}, i > j} z(m_i, \pi_A^*(x_i))$ can be written as $\alpha_{j+1}m_{j+1} + \beta_{j+1}$ for some $\alpha_{j+1}, \beta_{j+1}$ which depend only on $\{\pi_A^*(x_i)\}_{i > j}$ and the nature of the dependence is determined entirely by the sequence of actions in this path for $i > j$ to $i = k-1$. We will now show that $\prod_{i \in \mathcal{G}, i > j-1} y(m_i, \pi_A^*(x_i)) \prod_{i \in \mathcal{B} \cup \{k-1\}, i > j-1} z(m_i, \pi_A^*(x_i))$ can be written as $\alpha_j m_j + \beta_j$ again with α_j, β_j depending only on $\{\pi_A^*(x_i)\}_{i > j-1}$ with the nature of the dependence determined by the sequence of actions on

this path for $i > j-1$ to $i = k-1$. We define $f(b, q, a) = \begin{cases} y(b, q) & a = G \\ z(b, q) & a = B \end{cases}$ and $g(q, a) = \begin{cases} q & a = G \\ 1 - q & a = B \end{cases}$.

Thus,

$$\begin{aligned}&\prod_{i \in \mathcal{G}, i > j-1} y(m_i, \pi_A^*(x_i)) \prod_{i \in \mathcal{B} \cup \{k-1\}, i > j-1} z(m_i, \pi_A^*(x_i)) \\ &= f(m_j, \pi_A^*(x_j), a_{j+1}) \prod_{i \in \mathcal{G}, i > j} y(m_i, \pi_A^*(x_i)) \prod_{i \in \mathcal{B} \cup \{k-1\}, i > j} z(m_i, \pi_A^*(x_i)) \\ &= f(m_j, \pi_A^*(x_j), a_{j+1}) [\alpha_{j+1}m_{j+1} + \beta_{j+1}] \quad (\text{inductive hypothesis}) \\ &= f(m_j, \pi_A^*(x_j), a_{j+1}) \left[\alpha_{j+1} \frac{g(\pi_A^*(x_j), a_{j+1})}{f(m_j, \pi_A^*(x_j), a_{j+1})} m_j + \beta_{j+1} \right] \quad (\text{Equation (13)}) \\ &= \alpha_{j+1}g(\pi_A^*(x_j), a_{j+1})m_j + \beta_{j+1}f(m_j, \pi_A^*(x_j), a_{j+1})\end{aligned}$$

Thus, from the definitions of $f(\cdot), g(\cdot), y(\cdot)$, and $z(\cdot)$,

$$\alpha_j = \begin{cases} \alpha_{j+1}\pi_A^*(x_j) + \beta_{j+1}(2\pi_A^*(x_j) - 1) & a_{j+1} = G \\ \alpha_{j+1}(1 - \pi_A^*(x_j)) + \beta_{j+1}(1 - 2\pi_A^*(x_j)) & a_{j+1} = B \end{cases} \text{ and } \beta_j = \begin{cases} \beta_{j+1}(1 - \pi_A^*(x_j)) & a_{j+1} = G \\ \beta_{j+1}\pi_A^*(x_j) & a_{j+1} = B \end{cases}$$

In both cases, by our inductive hypothesis, $\alpha_{j+1}, \beta_{j+1}$ depend only on $\{\pi_A^*(x_i)\}_{i > j}$ and the actions in this path from $i > j$ to $i = k-1$. Thus, the claim holds for α_j, β_j , completing the induction.

Now recall that the denominator of (34) is $\prod_{i \in \mathcal{G}} y(x_i, \pi_A^*(x_i)) \prod_{i \in \mathcal{B} \cup \{k-1\}} z(x_i, \pi_A^*(x_i))$. Note that the policy and sequence of actions is the same as those for the numerator. Thus, because, as shown, the coefficients α_0, β_0 depend only on $\{\pi_A^*(x_i)\}$ and the path of actions, the same logic can be applied to write the denominator

as $\alpha_0 x_0 + \beta_0$ for the same α_0, β_0 . This can be done similarly for the numerator of the second term in (34). Thus, we can reason about the coefficients in (33) as follows:

$$\begin{aligned}
& t \left(\prod_{i \in \mathcal{G}} \frac{y(m_i, \pi_A^*(x_i))}{y(x_i, \pi_A^*(x_i))} \right) \left(\prod_{i \in \mathcal{B} \cup \{k-1\}} \frac{z(m_i, \pi_A^*(x_i))}{z(x_i, \pi_A^*(x_i))} \right) \\
& + (1-t) \left(\prod_{i \in \mathcal{G}} \frac{y(n_i, \pi_A^*(x_i))}{y(x_i, \pi_A^*(x_i))} \right) \left(\prod_{i \in \mathcal{B} \cup \{k-1\}} \frac{z(n_i, \pi_A^*(x_i))}{z(x_i, \pi_A^*(x_i))} \right) \\
& = \frac{t(\alpha_0 m_0 + \beta_0) + (1-t)(\alpha_0 n_0 + \beta_0)}{\alpha_0 x_0 + \beta_0} \\
& = \frac{\alpha_0 x_0 + \beta_0}{\alpha_0 x_0 + \beta_0} \quad (\text{Since } x_0 = t m_0 + (1-t) n_0) \\
& = 1
\end{aligned}$$

Thus, (34) holds, completing the proof. $\square$

C.4 ALTRUISTIC MYOPIC POLICY BOUND

Lemma 9.

$$\pi_A^0(b) \leq \pi_A^*(b) \quad \forall b \in [0, 1]$$

Assume by way of contradiction that there exists $b \in [0, 1]$ such that $\pi_A^0(b) > \pi_A^*(b)$. We will show that $V_A^*(b) < V_A^\pi(b)$ for a policy π we construct, violating the optimality of $V_A^*(\cdot)$.

Let π be such that, when starting at public belief b , the policy π applies precision $\pi_A^0(b)$ at the current time step and then applies the optimal policy at all future time steps. We now elaborate $V_A^*(b) - V_A^\pi(b)$ by breaking the value function into instantaneous reward r_A and future discounted reward. We will refer to b^+ and b^- as the updated beliefs after receiving signals G and B , respectively when applying the optimal policy and b_m^+ and b_m^- as the same when applying the myopic policy.

$$\begin{aligned}
V_A^*(b) - V_A^\pi(b) &= r_A(b, \pi_A^*(b)) - r_A(b, \pi_A^0(b)) \\
&+ \delta y(b, \pi_A^*(b)) V_A^*(b^+) + \delta z(b, \pi_A^*(b)) V_A^*(b^-) \\
&- \delta y(b, \pi_A^0(b)) V_A^*(b_m^+) - \delta z(b, \pi_A^0(b)) V_A^*(b_m^-) \\
&< \delta y(b, \pi_A^*(b)) V_A^*(b^+) + \delta z(b, \pi_A^*(b)) V_A^*(b^-) \\
&- \delta y(b, \pi_A^0(b)) V_A^*(b_m^+) - \delta z(b, \pi_A^0(b)) V_A^*(b_m^-) \quad (\text{definition of } \pi_A^0) \\
&\leq 0 \quad (\text{Theorem 2})
\end{aligned}$$

Here, we applied the fact that the instantaneous reward under $\pi_A^0(\cdot)$ must be greater from the definition of the myopic optimal policy.

We then relied on the convexity of $V_A^*(\cdot)$ from Theorem 2 for the future cost. From Equation (13), note that $b_m^- \leq b^- \leq b \leq b^+ \leq b_m^+$. Along with this convexity, this leads to the last step above and completes the proof. $\blacksquare$

C.5 PROOF OF THEOREM 3: OPTIMAL ALTRUISTIC POLICY

Consider $V_A^\pi(b)$ with policy π applying precision q in the current time step and applying the optimal policy at all future time steps. The q that maximizes $V_A^\pi(b)$ provides the optimum precision starting at public belief b . Because the planner never benefits from decreasing precision below the baseline p (see Appendix C.9), we restrict q to $[p, 1]$ without loss of generality.

We will let b^+ and b^- denote the positive and negative belief updates possible after receiving a signal with precision q . Taking the second derivative of $V_A^\pi(b)$ with respect to q for $q \geq \max(b, 1-b)$:

$$\frac{\partial^2}{\partial q^2} V_A^\pi(b) = \begin{cases} -\ddot{\beta}(q) + \delta b^2 (1-b)^2 \left[\frac{\ddot{V}_A^*(b^+)}{y^3(b, q)} + \frac{\ddot{V}_A^*(b^-)}{z^3(b, q)} \right] & q \geq \max(b, 1-b) \\ -\ddot{\beta}(q) & \text{otherwise} \end{cases}$$

With the convexity and concavity of $V_A^*(\cdot)$ and $\beta(\cdot)$, respectively, the above shows that $V_A^\pi(b)$ is convex with respect to q on $[\max(b, 1-b), 1]$. Thus, on this interval, it is maximized at one of the two extreme points.

For $q \in [0.5, \max(b, 1 - b))$, the derivative of $V_A^\pi(b)$ with respect to q is negative; thus, the optimal choice will be the baseline precision p to minimize the cost. Thus, there are three possible optimal precisions: p , $\max(b, 1 - b)$, and 1.

Note that the expected value under the latter two candidates is always strictly negative. Under precision p , however, the expected value is 0 at $b = 0$ and $b = 1$. This, combined with the non-positivity and symmetry of the value function (Lemma 8) implies the existence of d_A such that $\pi_A^*(b) = p$ for $b \in [0, d_A] \cup (1 - d_A, 1]$. Furthermore, from Theorem 1 and applying Lemma 9, $d_A \in (0, t_M)$ where t_M is defined as in the statement of Theorem 1.

Finally, Theorem 1 and applying Lemma 9, also imply the existence of $t_A \in [d_A, t_M]$ such that $\pi_A^*(b) = 1$ for $b \in (t_A, 1 - t_A)$. This completes the proof. ■

C.6 PROOF OF THEOREM 4: MYOPIC BIASED POLICY

Preliminaries Applying the action rule from (2), the instantaneous reward (6) can be written as follows:

$$r_B(b, q) = \begin{cases} -\beta(|q - p|) - Cz(b, q) & q \geq \max(b, 1 - b) \\ -\beta(|q - p|) - C & b < 1 - q \\ -\beta(|q - p|) & b > q \end{cases} \quad (44)$$

Note that $z(b, q) = b + q - 2bq$ increases with respect to q when $b < 0.5$ and decreases with respect to q when $b > 0.5$.

At $b = 0.5$, $z(b, q) = z(0.5, q) = 0.5$ regardless of the precision chosen. Furthermore, at this public belief, $q \geq \max(b, 1 - b)$ because $q \geq 0.5$ for any q , falling in the first case of (44). In this case, the second term of $r_B(b, q)$ is unchanging with respect to q , so the planner seeks to maximize only the first term, which is accomplished at $q = p$.

We now prove a series of lemmas (10-14) which will assist us in our proof of Theorem 4.

Lemma 10.

If $\pi_B^0(b) \geq \max(b, 1 - b)$, then $\pi_B^0(b) \in \{\max(b, 1 - b), p, 1\}$.

Proof. The reward is differentiable everywhere except $q = p$. Thus, we can write its derivative with respect to the chosen precision as follows (except at $q = p$):

$$\frac{\partial r_B(b, q)}{\partial q} = \begin{cases} -\text{sign}(q - p)\dot{\beta}(|q - p|) & q \geq \max(b, 1 - b) \\ -C(1 - 2b) & \\ -\text{sign}(q - p)\dot{\beta}(|q - p|) & \text{otherwise} \end{cases} \quad (45)$$

In the former case, we can write the second derivative of the reward with respect to q as follows, for $q \neq p$:

$$\frac{\partial^2}{\partial q^2} r_B(b, q) = -\ddot{\beta}(|q - p|)$$

Because $\beta(\cdot)$ is concave, this expression is non-negative, and $r_B(b, q)$ is convex with respect to q . Thus, the reward is maximized at one of the extreme points $\{\max(b, 1 - b), 1\}$ or p because $r_B(b, q)$ is not differentiable at $q = p$. □

Lemma 11.

For $b \in [0, 1 - p)$, $\pi_B^0(b) \in \{p, 1 - b\}$.

Proof. Because $p > 0.5$, for $b \in [0, 1 - p)$, $\max(b, 1 - b) = 1 - b > p$.

First, consider the case when $\pi_B^0(b) < \max(b, 1 - b) = 1 - b$. From (44), in such instances, the reward is maximized by minimizing $|q - p|$. Thus, the optimal precision lower than $\max(b, 1 - b)$ is p .

Now consider $\pi_B^0(b) \geq \max(b, 1 - b) = 1 - b$. By Lemma 10, $\pi_B^0(b) \in \{1 - b, p, 1\}$. What remains is to rule out $\pi_B^0(b) = 1$. To do so, we compare $r_B(b, 1)$ and $r_B(b, 1 - b)$. Both fall under the first case of (44). Because $\beta(\cdot)$ is increasing in its argument and $p < 1 - b \leq 1$, the first term of the reward from (44) is greater for precision $1 - b$ than for precision 1. Since $z(b, q)$ is increasing with respect to q for $b < 0.5$, $r_B(b, 1) \leq r_B(b, \max(b, 1 - b))$ so we can rule out precision 1. This completes the proof. □

Lemma 12.

For $b \in (p, 1]$, $\pi_B^0(b) = p$.

Proof. The reward is strictly negative for any $q \neq p$ because of the term $-\beta(|q - p|)$. For $b \in (p, 1]$ and $q = p$, $p < \max(b, 1 - b) = b$, leading to the third case of (44) and yielding reward 0. Thus, the optimal precision is p . $\square$

Lemma 13.

For $b \in [1 - p, 0.5]$, $\pi_B^0(b) \in \{p, 1 - b\}$.

Proof. When $p = 0.5$, the interval $[1 - p, 0.5]$ contains only $\{0.5\}$. At $b = 0.5$ $p = 1 - b$, thus $\pi_B^0(b) = p = 1 - b = 0.5$. We deal with $p > 0.5 \geq b$ for the remainder of the proof.

For $b \in [1 - p, 0.5]$, $\max(b, 1 - b) = 1 - b \leq p$. Thus, from the first case of (44), precision p yields reward $r_B(b, p) = -Cz(b, p)$. We now consider two cases: $q < \max(b, 1 - b)$ and $q \geq \max(b, 1 - b)$.

Suppose $q < \max(b, 1 - b) = 1 - b$. This leads to the second case of (44) and yields reward $-\beta(|q - p|) - C$. Furthermore, $z(b, q)$ is increasing with respect to q for $b < 0.5$ and takes values between $z(b, 0.5) = 0.5$ and $z(b, 1) = 1 - b$. Thus, $z(b, q) < 1$ for all q . Therefore, the reward from policy $q < 1 - b$ is strictly lower than $r_B(b, p)$, and we can rule out this case.

When $q \geq \max(b, 1 - b)$, applying Lemma 10 yields that the remaining candidate precisions are $\{1 - b, p, 1\}$. We rule out precision 1 by comparing it to precision p . The first term of the reward from (44) is greater for precision p than for precision 1. Since $z(b, q)$ is increasing with respect to q for $b < 0.5$, the second term is also greater for precision p , and we can rule out precision 1. This completes the proof. $\square$

We now define the notion of ϵ -optimal policies.

Definition 1. ϵ -optimal policies

An ϵ -optimal policy $\pi^\epsilon(b)$ is any policy such that the following holds

$$V^*(b) \leq V^{\pi^\epsilon}(b) + \epsilon$$

for $\epsilon > 0$.

Note that from the definition, any optimal policy is also ϵ -optimal for any $\epsilon > 0$. Thus, $\pi_B^0(b)$ is also ϵ -optimal

Following this definition, let $\pi_B^{\epsilon, 0}(b)$ denote an ϵ -optimal myopic policy for the biased planner.

Lemma 14.

For $b \in (0.5, p]$ and any $\epsilon > 0$, there are two possibilities:

1. $\pi_B^0(b) \in \{p, 1\}$, or

2. $\pi_B^{\epsilon, 0}(b) = b - \epsilon$

Proof. For $b \in (0.5, p]$, $\max(b, 1 - b) = b \leq p$. We consider two cases: $q \geq \max(b, 1 - b)$ or $q < \max(b, 1 - b)$.

In the first case, Lemma 10 yields that $\pi_B^0(b) \in \{b, p, 1\}$. We rule out precision b by comparing it to precision p . The first terms of the reward from (44) is greater for precision p . Furthermore, since $z(b, q)$ is decreasing with respect to q for $b > 0.5$, the second term of the reward is also greater for precision p . Thus, we can rule out precision b . This yields the first claim.

In the latter case, $q < \max(b, 1 - b) = b$. Therefore, (1) since $q \geq 0.5$, $q \in [0.5, b)$, and (2) the third case of (44) applies leading to reward $-\beta(|q - p|)$. Thus, the reward $r_B(b, q)$ is an increasing function of $q \in [0.5, b)$ because $0.5 < b \leq p$. This leads directly to the second claim. $\square$

Remark 3. As noted in Puterman (1990) Section 2.3.1, an optimal policy may not exist for infinite action spaces because the supremum may not be attained. When an optimal policy does not exist, we instead seek ϵ -optimal policies, i.e., policies that yield reward within $\epsilon > 0$ of the supremum. In particular, this is true for case 2 of Lemma 14.

Proof of Theorem 4 We consider the myopic biased policy in 4 cases: $b \in [0, 1 - p]$, $b \in [1 - p, 0.5]$, $b \in (0.5, p]$, and $b \in (p, 1]$, making use of the fact that $p \geq 0.5$. Based on Lemmas 11-14, these cases have the following candidate policies: $\{p, 1 - b\}$, $\{1 - b, p\}$, $\{1, p, b - \epsilon\}$, and $\{p\}$, respectively.

Case 1: $b \in [0, 1 - p]$, $\pi_B^0(b) \in \{p, 1 - b\}$ (Lemma 11)

We will show the existence of a threshold t_1 such that $\pi_B^0(b) = p$ for $b \in [0, t_1]$ and $\pi_B^0(b) = 1 - b$ for $b \in (t_1, 1 - p)$.

We can express the reward under the two candidate policies as follows by applying (44):

$$\begin{aligned} r_B(b, p) &= -C \\ r_B(b, 1 - b) &= -\beta(1 - b - p) - C(1 - 2b(1 - b)) \end{aligned}$$

The derivative of $r_B(b, 1 - b)$ with respect to b is

$$\frac{\partial}{\partial b} r_B(b, 1 - b) = \dot{\beta}(1 - p - b) + 2C(1 - 2b)$$

which is positive as $\beta(\cdot)$ is increasing and $b < 1 - p \leq 0.5$. Furthermore, when evaluated at $b = 0$ and $b = 1 - p$ we obtain

$$\begin{aligned} r_B(0, 1) &= -\beta(1 - p) - C < r_B(b, p) \\ r_B(1 - p, p) &= -C(1 - 2p(1 - p)) > r_B(b, p) \end{aligned}$$

Thus, $r_B(b, 1 - b)$ is increasing on $[0, 1 - p]$ with $r_B(b, 1 - b) < r_B(b, p)$ for $b = 0$ and vice versa for $b = 1 - p$. Because $\beta(\cdot)$ is continuous, so are both reward functions. Thus, there exists $t_1 \in (0, 1 - p)$ such that $\pi_B^0(b) = p$ for $b \in [0, t_1]$ and $\pi_B^0(b) = 1 - b$ for $b \in (t_1, 1 - p)$.

Case 2: $b \in [1 - p, 0.5]$, $\pi_B^0(b) \in \{1 - b, p\}$ (Lemma 13)

If $p = 0.5$, this interval contains only $\{0.5\}$, and, for $b = 0.5$, $1 - b = p$. Thus, the two candidate policies are equivalent, and the optimal policy is $\pi_B^0(b) = p$. For the remainder of this case, we consider $p > 0.5$.

We now prove the existence of t_2, t_3 such that $\pi_B^0(b) = 1 - b$ for $b \in [t_2, t_3]$, and $\pi_B^0(b) = p$ for $b \in [1 - p, t_2] \cup [t_3, 0.5]$.

We can express the reward under the candidate policies as follows by applying (44):

$$\begin{aligned} r_B(b, 1 - b) &= -\beta(p - 1 + b) - Cz(b, 1 - b) \\ r_B(b, p) &= -Cz(b, p) \end{aligned}$$

This leads to the following condition for the optimal policy to be $1 - b$:

$$\begin{aligned} r_B(b, 1 - b) &\geq r_B(b, p) \\ \beta(p - 1 + b) &\leq C[z(b, p) - z(b, 1 - b)] \\ \beta(p - 1 + b) &\leq C(p - 1 + b)(1 - 2b) \end{aligned} \tag{46}$$

The LHS of (46) is a concave, increasing function of b on $[1 - p, 0.5]$. The RHS of (46) is a concave parabola that increases in the first half of $[1 - p, 0.5]$ and decreases in the second half.

Now consider the two halves of the interval $[1 - p, 0.75 - 0.5p]$ and $[0.75 - 0.5p, 0.5]$. In the first half, the LHS and RHS of (46) are increasing and concave, with the RHS also being a parabola. Thus, in the first half, the two intersect at most two points (one of which is at $b = 1 - p$). In the second half, the LHS of (46) is increasing while the RHS is decreasing. Thus, they intersect at most once. Let the two possible intersections for $b > 1 - p$ be t_2 and t_3 with $t_2 \leq t_3$. Thus, what remains is to determine the optimal on each of $(1 - p, t_2)$, (t_2, t_3) , and $(t_3, 0.5)$. At $b = 0.5$, $r_B(b, p) > r_B(b, 1 - b)$ because $p > 0.5$, thus $\pi_B^0(b) = p$ for $b \in (t_3, 0.5)$. Subsequently, $\pi_B^0(b) = 1 - b$ for (t_2, t_3) , and $\pi_B^0(b) = p$ for $(1 - p, t_2)$.

Case 3: $b \in (0.5, p]$, $\pi_B^{\epsilon, 0}(b) \in \{1, p, b - \epsilon\}$ (Lemma 14)

In this case, we show the existence of $0.5 \leq t_4 \leq t_5 \leq p$ such that the optimal precisions on intervals $(0.5, t_4]$, (t_4, t_5) , and $[t_5, p)$ are p , 1 , and $b - \epsilon$, respectively.

We can express the reward under the candidate policies as follows by applying (44):

$$\begin{aligned} r_B(b, 1) &= -\beta(1 - p) - C(1 - b) \\ r_B(b, p) &= -Cz(b, p) \\ r_B(b, b - \epsilon) &= -\beta(p - b + \epsilon) \end{aligned}$$

When evaluated at $b = p$, $r_B(b, b - \epsilon)$ is arbitrarily close to 0, and, therefore, $b - \epsilon$ is the optimal precision on some interval $(t_5, p]$ for $0.5 \leq t_5 < p$.

Now, note that $r_B(b, 1)$ and $r_B(b, p)$ are linear functions in b while $r_B(b, b - \epsilon)$ is convex in b (from the concavity of β). Thus, $r_B(b, b - \epsilon) - r_B(b, p)$ and $r_B(b, b - \epsilon) - r_B(b, 1)$ are convex and, as a result, have convex level sets. This implies that the set on which both of these expressions are positive, i.e., the set for which $b - \epsilon$ is the optimal policy, must also be convex. Therefore, $b - \epsilon$ cannot be optimal on any subset of $(0.5, t_5)$.

What remains is to compare precisions 1 and p on the interval $(0.5, t_5]$. To do so, we inspect the values of $r_B(b, 1)$ and $r_B(b, p)$ at 0.5 and their derivatives.

$$\begin{aligned} r_B(0.5, 1) &= -\beta(1 - p) - 0.5C \\ r_B(0.5, p) &= -0.5C \end{aligned}$$

Because $r_B(0.5, 1) < r_B(0.5, p)$ there exists interval $(0.5, t_4)$ with $t_4 \leq t_5$ such that $\pi_B^0(b) = p$. Furthermore, because both are linear in b , there is at most one crossover point where their ordering changes. Thus, if $t_4 < t_5$, then $\pi_B^0(b) = 1$ on (t_4, t_5) .

Case 4: $b \in (p, 1]$, $\pi_B^0(b) = p$

This follows directly from Lemma 12 and completes the proof of Theorem 4. ■

C.7 BIASED VALUE FUNCTION MONOTONICITY

Lemma 15. For any $b_1, b_2 \in [0, 1]$ such that $b_1 \leq b_2$,

$$V_B^*(b_1) \leq V_B^*(b_2)$$

Proof. Applying (44), the derivative of the biased reward with respect to b can be written as follows:

$$\frac{\partial}{\partial b} r_B(b, q) = \begin{cases} -C(1 - 2q) & q \geq \max(b, 1 - b) \\ 0 & \text{otherwise} \end{cases} \quad (47)$$

Because $q \geq 0.5$ and $z(b, q) \in [0, 1]$, it follows that $r_B(b, q)$ is increasing with respect to b .

Next, note that the Bayesian nature of belief updates implies that the belief process is a martingale. That is, for any fixed q_i , $\mathbb{E}[b_{i+1}] = b_i$. To demonstrate this consider two cases: (1) $q_i < \max(b_i, 1 - b_i)$, and (2) $q_i \geq \max(b_i, 1 - b_i)$.

In the former case, it follows from (3) that $b_{i+1} = b_i$, satisfying the claim.

In the latter case, (13) gives the distribution of b_{i+1} which yields the same.

The monotonicity of $V_B^*(\cdot)$ then follows from Proposition 5 of Smith & McCardle (2002). □

C.8 PROOF OF THEOREM 5: OPTIMAL BIASED POLICY

We first prove that the optimal value function is continuous in the following lemma:

Lemma 16. Biased value function continuity

$V_B^*(\cdot)$ is continuous w.r.t. b .

Proof. First consider the one-stage value function defined as $V_B^0(b) = \sup_{q \in [0.5, 1]} r_B(b, q)$. From Theorem 4, we know that $\pi_B^{\epsilon, 0}(b) \in \{1 - b, 1, b - \epsilon, p\}$. Thus, we can equivalently write the one-stage value function as follows:

$$\begin{aligned} V_B^0(b) &= \sup_{\epsilon > 0} r_B(b, \pi_B^{\epsilon, 0}(b)) \\ &= \sup_{\epsilon > 0} \max_{q \in \{p, 1 - b, 1, b - \epsilon, p\}} r_B(b, q) \end{aligned} \quad (48)$$

Using the above expression, we now argue that $V_B^0(\cdot)$ is continuous on $b \in [0, 1]$. First, consider $b \leq 0.5$. From Theorem 4, an ϵ -optimal policy applies precision $1 - b$ or p for $b \leq 0.5$. In either case, the instantaneous reward will not depend upon ϵ , and the supremum in (48) can be omitted, leaving $V_B^0(b) = \max_{q \in \{1 - b, p\}} r_B(b, q)$. For $b \leq 0.5$, $r_B(b, 1 - b)$ falls in case one of (44) and is thus continuous with respect to b . On the same interval, $r_B(b, p)$ has one discontinuity at $b = 1 - p$ where it switches from case two to case one of (44). At $b = 1 - p$, Theorem 4 shows that the optimal myopic policy switches from $1 - b$ to p . Both policies fall in the first case of

(44), and, at $b = 1 - p$, $p = 1 - b$, so the controls are equal. Thus, the left and right limits coincide with the optimal value, and $V_B^0(\cdot)$ is continuous.

For $b \geq 0.5$, Theorem 4 gives the candidate ϵ -optimal policies $q \in \{1, b - \epsilon, p\}$. When the inner maximum in (48) is attained for $q = b - \epsilon$, the resulting value is $V_B^0(b) = \sup_{\epsilon > 0} r_B(b, b - \epsilon)$. From the third case of (44), we can rewrite this as $V_B^0(b) = \sup_{\epsilon > 0} -\beta(|b - \epsilon - p|)$. Since $b - \epsilon$ is only selected for $b \leq p$ (Theorem 4), this reduces to $V_B^0(b) = -\beta(|p - b|)$. Thus, exchanging the order of the maximum and supremum in (48), we obtain $V_B^0(b) = \max\{r_B(b, 1), -\beta(|p - b|), r_B(b, p)\}$ for $b \geq 0.5$. For the first expression in this maximum, $r_B(b, 1)$ falls into case one of (44) and is thus continuous. The continuity of $-\beta(|p - b|)$ follows from the continuity of $\beta(\cdot)$. For $r_B(b, p)$, since $b \geq 0.5$, the first case of (44) applies for $b \leq p$ and the third case applies for $b > p$. Thus, $r_B(b, p)$ is potentially discontinuous at $b = p$. Therefore, what remains is to show the continuity of $V_B^0(\cdot)$ at $b = p$.

At $b = p$, the myopic ϵ -optimal policy switches from $q = b - \epsilon$ to $q = p$ by Theorem 4 with $\pi_B^{\epsilon, 0}(p) = b - \epsilon$. Thus, for $b = p$, the inner maximum of (48) is attained at $q = b - \epsilon$ and, from the third case of (44), $V_B^0(p) = \sup_{\epsilon > 0} -\beta(\epsilon) = 0$. Similarly, from Theorem 4, for each $\epsilon > 0$, there exists $t_\epsilon < p$ such that on $[t_\epsilon, p]$ the inner maximum of (48) is attained by $q = b - \epsilon$. On this interval, $V_B^0(b) = \sup_{\epsilon > 0} -\beta(|p - b + \epsilon|) = -\beta(|p - b|)$. Therefore, $\lim_{b \uparrow p} V_B^0(b) = 0$. For $b > p$, the myopic optimal policy is $q = p$ from Theorem 4, and this yields reward $V_B^0(b) = 0$ for all $b > p$ from the third case of (44). This proves the continuity of $V_B^0(\cdot)$ at $b = p$ and completes the proof that $V_B^0(\cdot)$ is continuous on $[0, 1]$.

Applying Theorem 4.2 of Hinderer (2005), then yields the continuity of $V_B^*(\cdot)$. $\square$

Remark 4. Given the continuity of $V_B^*(\cdot)$ with respect to public belief and the fact that the control variable lies in the compact set $[0.5, 1]$, one might conclude that the optimal is guaranteed to exist at all public beliefs. However, there is an important subtlety which prevents such a conclusion: while $V_B^*(\cdot)$ is continuous with respect to public belief, it need not be continuous with respect to precision. Specifically, b_{i+1} is a function of both b_i and q_i (see Equation (3)); thus, $\mathbb{E}[V_B^*(b_{i+1})]$ may be denoted as $\phi(b_i, q_i)$. Now, b_{i+1} need not be continuous with respect to q_i , thus $\phi(\cdot)$ may not be continuous with respect to q_i . Thus, there may be regimes in which the optimal value is not attained by any policy. Indeed, we will show that such cases exist.

Proof of Theorem 5 We now leverage Lemma 16 to prove Theorem 5. We again let $z(b, q) = b + q - 2bq$ be the probability of an agent receiving signal B with precision q and prior public belief b and $y(b, q) = 1 - z(b, q)$ be the corresponding probability of the agent receiving signal G . We abuse notation to allow $V_B(b, q)$ to be the expected utility when beginning at public belief b , applying precision q in the current time step, and then following the optimal policy $\pi_B^*(\cdot)$ in all future steps. We can express $V_B(b, q)$ as follows by applying (13), (3), and (44):

$$V_B(b, q) = \begin{cases} -\beta(|q - p|) - Cz(b, q) \\ + \delta y(b, q) V_B^*(f(b, q, G)) \\ + \delta z(b, q) V_B^*(f(b, q, B)) & q \geq \max(b, 1 - b) \\ -\beta(|q - p|) \\ -C\mathbf{1}(b < 0.5) \\ + \delta V_B^*(b) & \text{otherwise} \end{cases} \quad (49)$$

Consider the latter case of (49) when $q < \max(b, 1 - b)$. Here, the only term affected by the precision policy is $-\beta(|q - p|)$, which is maximized by choosing q as close to p as possible. Thus, if the optimal precision is less than $\max(b, 1 - b)$, then it will be as close to p as possible.

We now consider four cases: (1) $b \in [0, 1 - p]$, (2) $b \in [1 - p, 0.5]$, (3) $b \in [0.5, p]$, and (4) $b \in (p, 1]$.

Case 1: $b \in [0, 1 - p]$, claims (A) and (B)

In this case, we will prove the existence of t_1 such that $\pi_B^*(b) = p$ for $b \leq t_1$. First note that if $\pi^*(b) < \max(b, 1 - b)$, we fall into the second case of Equation (49). Here, the only term affected by the precision policy is $-\beta(|q - p|)$, which is maximized by minimizing $|q - p|$. Because $p \in [0.5, 1]$, when $b \in [0, 1 - p]$, $p < \max(b, 1 - b)$. Thus, for $b \in [0, 1 - p]$, if $\pi_B^*(b) < \max(b, 1 - b)$, then $\pi_B^*(b) = p$.

Now, we evaluate the value function at belief $b = 0$. At $b = 0$, the state of the world is known to be B and, thus, no signal, regardless of its precision, will move the public belief. This can be seen mathematically by substituting $b = 0$ into the updated belief expressions in (13). Furthermore, at $b = 0$, if $q \geq \max(b, 1 - b)$, then $q = 1$. As argued above, if the precision is less than $\max(b, 1 - b)$, the optimal precision is $q = p$. We can compare the resultant value under these two candidate policies as follows from Equation (49):

$$V_B(0, 1) = -\beta(1 - p) - C + \delta V_B^*(0) \quad (50)$$

$$V_B(0, p) = -C + \delta V_B^*(0) \quad (51)$$

Because $V_B(0, 1) < V_B(0, p)$, $\pi_B^*(0) = p$ and $V_B^*(0) = V_B(0, p)$. Thus, from Equation (51), $V_B^*(0) = \frac{-C}{1-\delta}$. By Lemma 15, this is a lower bound for the value function at any public belief. Therefore, and since $V_B^*(b) \leq 0$, $|V_B^*(b)| \leq \frac{C}{1-\delta}$ for all $b \in [0, 1]$.

We now show that there exists an interval $[0, \epsilon)$ with $\epsilon > 0$ on which the optimum of the first case of Equation (49) is less than the optimum of the second case of Equation (49). Thus, in this interval the optimum control lies in the second case of Equation (49) so $\pi_B^*(b) < \max(b, 1-b)$. It follows from our initial argument earlier in this case that $\pi_B^*(b) = p$ in this interval, which would complete the proof.

We define function $g(\cdot)$ to be the maximum of the first case of the RHS of Equation (49):

$$g(b) = \max_{q \geq \max(b, 1-b)} \left[-\beta(|q-p|) - Cz(b, q) + \delta y(b, q) V_B^*(f(b, q, G)) + \delta z(b, q) V_B^*(f(b, q, B)) \right] \quad (52)$$

For $b \in (0, 1)$, $f(b, q, G)$ and $f(b, q, B)$ are both continuous with respect to q from Equation (13). Thus, the maximum in Equation (52) must exist because it is of a function which is continuous with respect to q over a compact set $q \in [\max(b, 1-b), 1]$.

We now take the limit of $g(\epsilon)$ as ϵ tends to 0. Because $q \in [\max(\epsilon, 1-\epsilon), 1]$, when $\epsilon \rightarrow 0$, $q \rightarrow 1$ (inside the maximum in Equation (52)). We now reason about each of the four terms in Equation (52). As $\epsilon \rightarrow 0$, $q \rightarrow 1$ so that 1) from the continuity of $\beta(\cdot)$, $\beta(|q-p|) \rightarrow \beta(1-p)$, 2) from the continuity of $z(\cdot)$, $z(\epsilon, q) \rightarrow 1$, $y(\epsilon, q) \rightarrow 0$. Thus, the second term goes to C . Because $V_B^*(\cdot)$ is bounded and $y(\epsilon, q) \rightarrow 0$, the third term goes to 0. Finally, because $z(\epsilon, q) \rightarrow 1$ and $f(\epsilon, q, B) \rightarrow 0$ and since $V_B^*(\cdot)$ is continuous from Lemma 16, the last term approaches $\delta V_B^*(0)$. Thus,

$$\lim_{\epsilon \rightarrow 0} g(\epsilon) = -\beta(1-p) - C + \delta V_B^*(0) \quad (53)$$

Applying precision p , however, yields the following from the second case of Equation (49):

$$\lim_{\epsilon \rightarrow 0} V_B(\epsilon, p) = \lim_{\epsilon \rightarrow 0} [-\beta(p-p) - C + \delta V_B^*(\epsilon)] \quad (54)$$

$$= -C + \delta V_B^*(0) \quad (55)$$

The last equality follows from the continuity of $V_B^*(\cdot)$ from Lemma 16.

Note that the expression in Equation (53) is strictly less than that in Equation (55). Thus, there exists an interval $[0, \epsilon)$ with $\epsilon > 0$ on which the optimum of the first case of Equation (49) is less than the optimum of the second case of Equation (49). The result follows.

Case 2: $b \in [1-p, 0.5)$, Claim (C)

We now argue that for $b_i \in [1-p, 0.5)$, $\pi^*(b_i) \geq 1-b_i$, yielding claim (C). Consider a stationary policy $\pi(\cdot)$ such that $\pi(b_i) < 1-b_i$. Note that, from Equation (3), the public belief will remain unchanged under such a policy. Thus, the same control will be applied at every future time step, and, from Equation (2), every agent will select action B . Using the facts that $1-b_i \leq p$ in this belief range, $\mathbb{P}(a_i = B | b_i, \pi(b_i)) = 1$, and $b_{i+1} = b_i$, we can apply Equation (6) to write the expected total discounted utility under such a policy as follows:

$$V_B^\pi(b_i) = -\beta(p - \pi(b_i)) - C + \delta V_B^\pi(b_i)$$

Thus, if this policy is applied, it will result in expected value $V_B^\pi(b_i) = \frac{-\beta(p - \pi(b_i)) - C}{1-\delta}$. That is, the cost of every agent choosing action B and continuing to apply precision $\pi(b_i)$. This is strictly lower than the value function lower bound of $\frac{-C}{1-\delta}$ established using Lemma 15 in Case 1 because $\pi(b_i) < 1-b_i \leq p$. Therefore, if an optimal policy exists at this point, it must be the case that $\pi_B^*(b_i) \geq 1-b_i$.

The remaining candidate policies are $q_i \in [1-b_i, 1]$. Applying Bellman's principle we can write the following:

$$V_B^*(b_i) = \sup_{q_i \in [1-b_i, 1]} r_B(b_i, q_i) + \delta \mathbb{E}[V_B^*(b_{i+1})] \quad (56)$$

Because $q_i \geq \max(b_i, 1-b_i)$ and $b \in [1-p, 0.5) \subset (0, 1)$, Equation (13) and Equation (3) imply that $b_{i+1} = \frac{b_i q_i}{y(b_i, q_i)}$ or $b_{i+1} = \frac{b_i(1-q_i)}{z(b_i, q_i)}$. Both of these expressions are continuous with respect to q_i and $V_B^*(\cdot)$ is continuous by Lemma 16, thus the RHS of Equation (56) inside the supremum is continuous with respect to q_i . Therefore, in Equation (56), the supremum over a compact set $q_i \in [1-b_i, 1]$ of a function which is continuous with respect to q_i must be attained, implying the existence of an optimal policy. Thus, $\pi_B^*(b_i) \geq 1-b_i$.

Case 3: $b \in [0.5, p]$ Claims (D) and (E)

For $b = 0.5$, $q \geq b$ as we assumed $q \in [0.5, 1]$. This places us in the first case of Equation (49). Furthermore, $\beta(\cdot)$, $y(b, q)$, $z(b, q)$, $f(b, q, G)$, and $f(b, q, B)$ are all continuous with respect to q for $b = 0.5$. Thus, the

supremum over $q \in [0.5, 1]$ of $V(b, q)$ must be attained as it is of a continuous function of q over a compact set of q . Therefore, an optimal policy must exist at $b = 0.5$.

Now consider $b \in (0.5, p]$ and a stationary policy $\pi'(\cdot)$ such that $\pi'(b) < b$. From Equation (3) and Equation (2), applying such a policy will result in public belief remaining unchanged, every agent selecting action G , and the same control being applied at every future time step. Using the facts that $b_i \leq p$, $\mathbb{P}(a_i = G | b_i, \pi'(b_i)) = 1$, and $b_{i+1} = b_i$, we can apply Equation (6) to write the total expected discounted utility under such a policy as follows:

$$V_B^{\pi'}(b_i) = -\beta(p - \pi'(b_i)) + \delta V_B^{\pi'}(b_i)$$

Thus, the policy results in expected value $V_B^{\pi'}(b_i) = \frac{-\beta(p - \pi'(b_i))}{1 - \delta}$. Note that, in this case, $V_B^{\pi'}(b_i)$ is an increasing function of $\pi'(b_i)$. Thus, the supremum of $V^*(\pi')_B(b_i)$ over $\pi'(b_i) \in (0.5, b)$ is not attained.

We now consider two cases: (1) $\max_{q_i \in [b_i, 1]} V(b_i, q_i) \geq \sup_{\pi'(b_i) \in (0.5, b_i)} \frac{-\beta(p - \pi'(b_i))}{1 - \delta}$, and (2) the complement. In the former case, we can argue as we did for $b = 0.5$ that the supremum must be attained and an optimal policy exists such that $\pi_B^*(b_i) \geq b_i$. In the latter case, as argued above, no optimal policy exists. Instead, for sufficiently small $\epsilon > 0$, there exists an ϵ -optimal policy such that $\pi_B^\epsilon(b_i) = b_i - \epsilon$. This yields claims (D) and (E).

What remains is to show that there exists $t_2 \in (0.5, p)$ such that for $b_i \in (t_2, p]$ we fall into the latter case, yielding claim (F). We define functions $m(\cdot)$ and $n(\cdot)$ as follows corresponding to the resultant value from choosing precisions less than b or greater than or equal to b , respectively:

$$\begin{aligned} m(b) &= \sup_{\pi'(b) \in [0.5, b)} \frac{-\beta(p - \pi'(b))}{1 - \delta} \\ n(b) &= \max_{q \geq b} V(b, q) \end{aligned}$$

Taking the limit of $m(\cdot)$ as $b \rightarrow p$ yields:

$$\lim_{b \rightarrow p} m(b) = 0 \tag{57}$$

When $q \geq b$ we fall into the first case of Equation (49). Thus we can expand $n(b)$ as follows:

$$\begin{aligned} n(b) &= \max_{q \geq b} [-\beta(|q - p|) - Cz(b, q) \\ &\quad + \delta y(b, q)V_B^*(b, q) \\ &\quad + \delta z(b, q)V_B^*(b, q)] \end{aligned} \tag{58}$$

We will argue that, for $b \in (0.5, p]$ and $q \geq b$, $n(b) \leq -C(1 - p)$. Since p is a fixed parameter strictly less than 1, this means from Equation (57) that $n(b) \leq -C(1 - p) < \lim_{b \rightarrow p} m(b)$, which implies the existence of an open interval $(t_2, p]$ where an optimal policy does not exist, yielding claim (F). We now argue that, for $b \in (0.5, p]$ and $q \geq b$, $n(b) \leq -C(1 - p)$. Note that every term of Equation (58) is bounded above by 0, so it suffices to show that any one term, say the second term $-Cz(b, q)$, is upper bounded by $-C(1 - p)$. First note that $z(b, q) = q(1 - 2b) + b$ so $z(b, q)$ is decreasing in q since $b > 0.5$. Thus, $z(b, q) \geq z(b, 1) = 1 - b \geq 1 - p$. The last inequality follows because $b \in (0.5, p]$. The claim follows since $p < 1$.

Along with Equation (57), this implies the existence of an open interval (t_2, p) where an optimal policy does not exist, yielding claim (F).

Case 4: $b \in (p, 1]$ Claim (A)

For $b \in (p, 1]$, consider applying stationary policy $\pi(\cdot)$ such that $\pi(b) = p$. Because $b > p$, Equation (3) and Equation (2) imply that such a policy results in public belief remaining unchanged, every agent selecting action G , and the same control being applied at every future time step. We can apply Equation (6) to write

$$V_B^\pi(b) = -\beta(p - p) + \delta V_B^\pi(b) = \delta V_B^\pi(b)$$

Thus, $V_B^\pi(b) = 0$. Since this is also an upper bound of the optimal value function for any b , $\pi_B^*(b) = p$. ■

C.9 SOCIAL WELFARE MONOTONICITY

Lemma 17. *Precision Monotonicity of Social Welfare*

If $\pi_1(b) \leq \pi_2(b) \forall b \in [0, 1]$, for policies $\pi_1(\cdot), \pi_2(\cdot)$, then

$$W^{\pi_1}(b) \leq W^{\pi_2}(b) \forall b \in [0, 1]$$

Proof. It is clear that the one-step welfare $-C \min(b_i, 1 - b_i, \pi(b_i))$ is increasing in $\pi(b_i)$.

Furthermore, note that the convexity of social welfare with respect to public belief follows from the same argument as in the proof of Theorem 2. This is because social welfare is equivalent to the altruistic value function in the special case where $\beta(\cdot) = 0$.

We can then follow the same argument as in the proof of Lemma 9 (Appendix C.4) to show that increasing precision increases welfare. $\square$

D EXTENSION TO HETEROGENEOUS AGENT PREFERENCES

We next show that the altruistic-planner analysis is unchanged when agents have two observable preference types with opposite preferred actions.

Assume two agent types, G and B , with agent i 's type denoted by t_i . Type- G agents receive unit reward when $a_i = \omega$ and no reward otherwise; type- B agents receive unit reward when $a_i \neq \omega$ and no reward otherwise. Each agent observes her own type and the types of all other agents, while all other model primitives remain unchanged.

The private belief update is unchanged because posterior beliefs depend on the signal structure $\mathbb{P}(s_i = \omega) = q_i$, not on the agent's payoff type:

$$\tilde{b}_i = \mathbb{P}(\omega = G | \mathcal{H}_i, q_i, s_i) = \begin{cases} \frac{q}{1+2b_i q_i - b_i - q_i} b_i & s_i = G \\ \frac{1-q}{b_i + q_i - 2b_i q_i} b_i & s_i = B \end{cases} \quad (59)$$

Using Equation (59), the action rule differs only by type; superscripts denote the type of agent i , and $!$ denotes negation:

$$a_i^G = \begin{cases} s_i & 1 - q_i \leq b_i \leq q_i \\ G & q_i < b_i \\ B & b_i < 1 - q_i \end{cases}$$

$$a_i^B = \begin{cases} !s_i & 1 - q_i \leq b_i \leq q_i \\ B & q_i < b_i \\ G & b_i < 1 - q_i \end{cases}$$

Because each type is public, an action reveals agent i 's private signal exactly when $q_i \geq \max(b_i, 1 - b_i)$. If $q_i < \max(b_i, 1 - b_i)$, the action is determined by public belief and type, so it is uninformative about the signal.

After observing type t_i , precision q_i , and action a_i , the public belief therefore follows the same update rule:

$$b_{i+1} = f(b_i, q_i) = \begin{cases} \tilde{b}_i & 1 - q_i \leq b_i \leq q_i \\ b_i & \text{o.w.} \end{cases} \quad (60)$$

In this heterogeneous setting, an altruistic planner maximizes each agent's own utility: it seeks action ω from type- G agents and action $!\omega$ from type- B agents.

The instantaneous reward comparison is therefore type-invariant, as shown by the two expressions below:

$$\begin{aligned} r^G(b_i, q_i) &= -\beta(|q_i - p|) - C \mathbb{P}(a_i \neq \omega) \\ &= \begin{cases} -\beta(|q_i - p|) - C(1 - q_i) & q_i \geq \max(b_i, 1 - b_i) \\ -\beta(|q_i - p|) - C \min(b_i, 1 - b_i) & q_i < \max(b_i, 1 - b_i) \end{cases} \\ r^B(b_i, q_i) &= -\beta(|q_i - p|) - C \mathbb{P}(a_i = \omega) \\ &= \begin{cases} -\beta(|q_i - p|) - C(1 - q_i) & q_i \geq \max(b_i, 1 - b_i) \\ -\beta(|q_i - p|) - C \min(b_i, 1 - b_i) & q_i < \max(b_i, 1 - b_i) \end{cases} \end{aligned}$$

The two reward functions are identical and coincide with Equation (5). Hence, with observable heterogeneous types, the altruistic planner faces the same quantitative optimization problem as in the homogeneous model. $\blacksquare$

E EXPERIMENTS

E.1 EXPERIMENTAL SETUP

The simulation maps the formal model to a car-purchase sequence in which each agent decides whether to buy after observing earlier actions and a planner-selected ad. The comparison uses LLMs in three measured

roles: (1) agents, (2) planner, and (3) oracle. Each agent receives a profile containing demographic information, car-use cases, and preferences. At the onset, the offered car is fixed as either good or bad, and the planner and agent information sets and objectives match the formal model in Section 3.

The oracle supplies the precision-controlled signal used to evaluate planner choices. It first generates a generic message that serves as the baseline signal absent planner intervention; after the planner chooses a precision, it generates a profile-specific message at that prescribed precision.

The oracle receives a fact sheet for each car type, including safety ratings and customer satisfaction, together with the target agent’s profile. Holding the facts fixed lets the experiment vary message informativeness while preserving the underlying product evidence. Appendix E.2 provides the full setup and prompts.

Each agent makes a binary buy/not-buy decision from the observed decision history and private signal, and the decision is evaluated by whether it matches the true car quality.

The planner observes the same decision history and chooses the precision of each agent’s message. It can also send the generic untailored message at its baseline precision for free, which provides the no-tailoring baseline. In the altruistic condition, the planner maximizes correct decisions; in the biased condition, it maximizes buy decisions.

All LLM outcomes are compared against analytical solutions computed with JuliaPOMDP (Egorov et al., 2017), so each simulation result has a matched Bayesian benchmark.

E.2 PARAMETERS AND PROMPTS

E.2.1 INSTANCE PARAMETERS

Experiments use the LG AI EXAONE Deep 32B model (LG AI Research, 2025). For each parameter configuration in Table 1, we compute the analytical solution and run the matched LLM simulation, keeping the cost, baseline precision, and discount-factor comparisons aligned.

Parameter	Description	Values
C	Cost of a bad action	1
k	Scaling factor for control cost	{0.1, 0.3, 0.5, 0.7, 0.9}
p	Baseline signal precision	{0.6, 0.7, 0.8, 0.9}
δ	Discount Factor	{0, 0.25, 0.5, 0.75, 1}

Table 1: Parameter grid shared by the analytical solutions and matched LLM simulations.

The car facts were constructed so that the GOOD car is favorable for every profile and the BAD car is unfavorable for every profile. This design keeps product fit constant across agents, so observed differences in decisions can be attributed to information precision and social-learning dynamics rather than heterogeneous preferences over the car itself. The fact sheets are:

```
fact_sheet["GOOD"] = [
    "EPA-rated range: 290 miles.",
    "Fast-charge 10-80 %% in 24 min on 150 kW DC.",
    "8-year / 100,000-mile battery warranty.",
    "NHTSA 5-Star overall safety rating.",
    "Cold-soak -5 °F keeps 93 %% range.",
    "Cools cabin 100→72 °F in 3 min.",
    "Cargo volume 26 ft³ total.",
    "Tow's 1 500 lb with factory hitch.",
    "Average 1.2 service visits first 3 yrs.",
    "92 %% owners would buy again (Apr-25).",
    "Qualifies for federal credit.",
    "Quarterly OTA updates incl. battery pre-heat."
]
fact_sheet["BAD"] = [
    "EPA-rated range: 155 mi (combined).",
    "Fast-charge 10 %% → 80 %% in 51 min (50 kW DC hardware limit).",
    "Battery warranty: 3 yr / 36 000 mi, first owner only.",
    "NHTSA overall safety rating 3 Stars (frontal-offset weakness).",
    "Cold-soak 15 °F retains 58 %% range; cabin heat cuts range 20%.",
    "Cabin cools 105 °F → 72 °F in 11 min.",
```

```

1836         "Cargo volume 17 ft3; rear seats fixed (no fold-flat).",
1837         "No factory towing approval.",
1838         "Average 3.8 service visits per year (battery-alert issues).",
1839         "Owner survey: 41 %% "would buy again" (1 200 owners, Feb-25).'",
1840         "Not eligible for federal tax credit.",
1841         "OTA updates limited to infotainment."
1842     ]

```

Profiles were generated programmatically and included each agent's name, age, city, interests, intended car use (e.g., commuting or road trips), and major concerns (e.g., fuel efficiency or winter driving), providing consistent profile variation across analytical and LLM comparisons.

E.2.2 PROMPTS

Agent LLM Prompt:

```

1850 sys = (
1851     f"You are deciding whether to buy a car. "
1852     f"Your profile is {profile}"
1853     f"You start believing there's a 50% chance the car is GOOD. "
1854     f"You have two information sources: "
1855     f"1. Previous agents' decisions "
1856     f"2. An advisor's message about the car "
1857     f>Your goal is to maximize your reward "
1858     f"i.e. BUY if the car is GOOD and "
1859     f"NOT BUY if the car is BAD. "
1860     f"If you BUY a GOOD car, you GAIN ${C}. "
1861     f"If you NOT BUY a BAD car, you GAIN ${C}. "
1862     f"If you BUY a BAD car, you LOSE ${C}. "
1863     f"If you NOT BUY a GOOD car, you LOSE ${C}. "
1864     f"belief = your probability (0.0-1.0) that the car is GOOD. "
1865     f"decision = exactly 'BUY' or 'NOT BUY'."
1866 )
1867 user = (
1868     f"PREVIOUS DECISIONS:\n{history_txt or '(none)'}\n\n"
1869     f"ADVISOR MESSAGE:\n{message}\n\n"
1870     f"TASK: Use the previous decisions and the advisor message. "
1871     f"Decide BUY or NOT BUY. "
1872     f"Update your belief about car quality (0.0-1.0). "\n"
1873     f"JSON response:{{\"belief\": <number>,\"decision\":<BUY or NOT BUY>}}'"
1874 )

```

Planner LLM Prompt:

```

1873 sys = (
1874     f>You are providing messages to potential car buyers. "
1875     f"Each potential car buyer (or agent) has the same information. "
1876     f>Your goal is to get agents to {objective}."
1877     f"Each decision according to your objective, earns you {C} dollars. "
1878     f>You and the agents have the same information. "
1879     f"Initially there is a 50% chance the car is GOOD. "
1880     f"The agents' aim to BUY if the car is GOOD "
1881     f"and NOT BUY if the car is BAD. "
1882     f>You can send a baseline message for free "
1883     f"or tailor the message to the agent at a cost. "
1884     f"Tailoring costs {k} dollars times the change in precision. "
1885     f"(Message precision is the probability (0.5 to 1.0) "
1886     f"that the message leaves the buyer with the correct impression. "
1887     f"For example, if the car is likely GOOD: "
1888     f"a high precision message will likely suggest BUY "
1889     f"and a low precision message will likely suggest NOT BUY. "
1890     f>You may choose to increase or decrease precision. "
1891     f"Agents decide based on your message and the decision history. "
1892     f"They update beliefs by observing previous decisions. "
1893     f"IMPORTANT: Thus your objective is to maximize:"

```

```

1890     f"{C} * (# of agents who choose according to your objective) "
1891     f"- sum of ({d}^i*{k}*abs(message_precision - baseline_precision))\n"
1892     f"where {d} is your discount factor for future cost. "
1893 )
1894 user = (
1895     f"BASELINE MESSAGE precision:\n{pi_base:.2f}\n"
1896     f"DECISION HISTORY:\n{history_txt}\n"
1897     f"TASK:"
1898     f"Based on the history and your objective, "
1899     f"choose your desired precision.\n"
1900     f'JSON response: {"message_precision": <number>}}\n'
1901 )

```

Oracle LLM Prompt:

```

1902
1903
1904 sys = (
1905     f"You will be given "
1906     f"1. an agent profile, "
1907     f"2. information about a car, "
1908     f"3. the true quality of the car, "
1909     f"4. a baseline message and the 'precision' of that message, and "
1910     f"5. a desired precision for the message you output. "
1911     f"(Message precision is the probability (0.5 to 1.0) "
1912     f"that the message leaves the buyer with the correct impression. "
1913     f"For example, if the car is likely GOOD: "
1914     f"a high precision message will likely suggest BUY "
1915     f"and a low precision message will likely suggest NOT BUY. "
1916     f"IMPORTANT: Output a message with the requested precision. "
1917     f"IMPORTANT: The message should be tailored to the profile. "
1918 )
1919 user = (
1920     f"AGENT PROFILE:{profile}\n"
1921     f"TRUE CAR QUALITY: {omega}\n\n"
1922     f"CAR FACTS: {car_facts}\n"
1923     f"BASELINE MESSAGE:{baseline_message}}\n"
1924     f"BASELINE MESSAGE precision: {p}\n"
1925     f"DESIRED MESSAGE precision: {q}\n"
1926     f"IMPORTANT: Output a message with the requested precision. "
1927     f"TASK: Write a message (MAX 100 words) with the desired precision."
1928     f'JSON response: {"message": "<message>}}'
1929 )
1930
1931
1932

```

E.3 VALIDATION OF LLM BELIEFS AND ORACLE

We add validation experiments for the LLM self-reported beliefs used in figure 1b and for the LLM oracle’s precision control. These checks quantify two measurement assumptions needed for the main comparisons: (1) the reported posterior is the belief used in state-contingent decisions, and (2) oracle-generated messages attain the requested precision, even though the LLM agents need not update Bayesianly.

First, we give the LLM a prior and a new signal, record its self-reported posterior, and then ask for a binary “buy/not-buy” decision that pays unit reward if and only if the state is good. The benchmark threshold is clear: a rational utility maximizer buys exactly when its subjective belief exceeds 0.5. In figure 4a, the empirical threshold is sharp, with the model overwhelmingly buying above reported posterior 0.5 and choosing not to buy below it. Next, we elicit willingness to pay for a state-contingent payoff. For prices c , the agent decides whether to pay c for unit utility if the state is good; a rational agent with belief p accepts exactly when $c < p$. Figure 4b plots buy rates over posterior–price pairs and shows the expected separation margin: acceptance is high when $c < p$ and low when $c > p$. Together, the two validation tasks show that the self-reported posterior closely tracks the belief used by the model in decisions.

We also validate oracle precision. After the oracle generates a message at requested precision q , a separate LLM estimates the realized precision $\hat{q}$. We evaluate $q \in \{0.5, 0.6, 0.7, 0.8, 0.9, 1.0\}$ across 100 programmatically generated profiles and 10 trials per combination, for 6,000 total trials. The mean absolute error $|q - \hat{q}|$ is 0.025, a small gap showing that realized message precision closely follows the requested precision over the full grid.




Figure 4: Validation of reported LLM beliefs. Panel (a) compares buy rates across self-reported posterior beliefs. Panel (b) compares buy rates across posterior–price pairs in the willingness-to-pay task.

E.4 IMPACT OF NON-BAYESIAN LLM BELIEF UPDATES ON LEARNING TRAJECTORIES

We finally compare belief trajectories in a representative LLM experiment to quantify how non-Bayesian updating changes the learning path. The policy is fixed to the analytically optimal policy and applied to LLM-simulated agents, isolating agent belief dynamics from planner-policy differences. When the car is bad (right panel), the Bayesian public belief (dotted black line) moves rapidly toward 0, whereas the LLM-reported belief (solid black line) declines more slowly. This trajectory-level gap is consistent with the measured micro-level biases: public belief resists entering a cascade (NB3) because agents overreact to occasional positive signals that contradict the increasingly negative group belief (NB2).




Figure 5: Belief-trajectory comparison under analytically optimal policies. The panels contrast Bayesian public belief with LLM-reported belief across planner objectives and true states.