

MASK WHAT MATTERS: CONTROLLABLE TEXT-GUIDED MASKING FOR SELF-SUPERVISED MEDICAL IMAGE ANALYSIS

Anonymous authors
Anonymous submission

ABSTRACT

Annotation scarcity in specialized domains such as medical imaging constitutes a fundamental obstacle to training robust vision models. Self-supervised masked image modeling (MIM) offers a principled remedy, yet prevailing formulations rely on random, high-ratio masking that is both computationally inefficient and semantically misaligned with the regions that govern downstream analysis. Region-aware variants only partially close this gap: because they ground their masking in reconstruction heuristics or task-specific supervised signals, they forfeit adaptability across tasks and imaging modalities. We introduce Mask What Matters, a controllable text-guided masking framework for self-supervised medical image analysis. The framework repurposes vision-language models to perform prompt-based, zero-shot localization of task-relevant anatomical structures, and it applies spatially differentiated masking that concentrates the reconstruction signal on diagnostically salient regions while suppressing redundancy in background areas. This controllable design tightens the semantic alignment between the pretraining objective and downstream objectives, strengthens the learned representations, and generalizes across tasks and modalities without recourse to per-image reports or labels. A comprehensive evaluation spanning three complementary medical imaging modalities—brain MRI, chest CT, and lung X-ray—establishes that Mask What Matters consistently surpasses strong MIM baselines such as Spark, improving classification accuracy by up to +3.1 percentage points, box average precision (BoxAP) by +1.3, and mask average precision (MaskAP) by +1.1 in detection and instance segmentation. Critically, these gains are attained at a substantially lower overall masking ratio (e.g., 40% versus 70%), establishing that semantic targeting, rather than aggressive masking, is the decisive driver of representation quality. Taken together, these results establish controllable, text-driven masking as an effective and generalizable principle for semantically aligned self-supervised learning, and chart a concrete path toward robust vision models for medical image analysis.

1 INTRODUCTION

Image classification, segmentation, and object detection constitute the foundational tasks of computer vision, and over the past decade deep learning has become the principal engine of progress across all three. This progress, however, is predicated on the availability of large-scale annotated datasets. In specialized domains such as medical image analysis, high-quality annotations are scarce, costly, and inseparable from expert labor, and this annotation bottleneck remains one of the primary barriers to the widespread clinical deployment of deep learning.

Self-supervised learning (SSL) confronts this bottleneck directly by exploiting unlabeled data for pretraining, enabling models to acquire transferable feature representations in the absence of any manual annotation [Jaiswal et al. (2020); Liu et al. (2021); Krishnan et al. (2022); Huang et al. (2023)]. Among the SSL paradigms, masked image modeling (MIM) has emerged as the dominant approach. By masking portions of an image and requiring the model to reconstruct the missing content, MIM compels the encoder to infer structural and semantic regularities from the visible context; the representations it yields transfer strongly across a broad range of downstream tasks,

medical image analysis prominently among them [He et al. (2022); Gupta et al. (2025); Xiao et al. (2023)].

Prevailing MIM approaches—masked autoencoders (MAE) [He et al. (2022)] and SparK [Tian et al. (2023)] among them—typically adopt random masking strategies that are, by construction, unaligned with the regions of interest that determine downstream performance. This task-agnostic masking entails two principal limitations. First, the learned representations may be semantically misaligned with the downstream objectives, curtailing transfer performance. Second, to raise the probability of masking informative content, existing methods resort to extremely high masking ratios (e.g., 75% in MAE), which inflates reconstruction difficulty and thereby demands large-scale data and substantial computational resources. Natural language, by contrast, exhibits compact structure and high semantic density, and supports effective learning even at low masking ratios [Devlin et al. (2019)]. The reliance on high masking ratios in image-based MIM stems largely from inherent information redundancy together with the absence of semantic-aware masking, a combination that makes it difficult for models to concentrate on the regions that genuinely matter.

These deficiencies are especially acute in medical imaging. Medical images are governed by high semantic sparsity: diagnostically relevant information is typically confined to small, localized regions—lesions or organs, for instance—whereas the majority of the image comprises semantically redundant background. Medical imaging further spans diverse modalities—MRI, CT, and X-ray—with substantial variation in the shape, size, and spatial distribution of task-relevant areas across tasks. Jointly, these properties impose stringent requirements on the efficiency and generalizability of any masking strategy.

To meet these requirements, we introduce a controllable, text-guided masking framework, hereafter **Mask-What-Matters (MWM)**. Given a user-specified, open-vocabulary description or phrase, a pretrained vision–language model supplies zero-shot localization cues, which we convert into robust regions of interest (ROIs). MWM then applies region-specific, prompt-conditioned masking—assigning higher ratios to semantically important areas (e.g., lesions and organs) and lower ratios to background—thereby injecting downstream semantics directly into the pretraining signal without per-image reports or labels.

Our main contributions are threefold:

- **A controllable text-guided masking framework for medical imaging**—the first of its kind. MWM couples a vision–language model (BiomedCLIP) with segmentation-based refinement (SAM) to localize task-relevant regions from open-vocabulary prompts and to apply differentiated ROI-versus-background masking, thereby resolving the semantic misalignment intrinsic to random masking.
- **An annotation-free, backbone-agnostic design.** MWM dispenses with per-image reports and labels entirely, and integrates seamlessly into both ViT- and ConvNet-based MIM pipelines under a single unified protocol.
- **Consistent gains across modalities and tasks.** On brain MRI, chest CT, and X-ray datasets, MWM surpasses state-of-the-art MIM baselines such as SparK, improving classification, detection, and segmentation alike, and thereby substantiating the promise of text-driven generative pretraining for medical imaging.

2 RELATED WORK

2.1 MASKED IMAGE MODELING

Masked modeling originated in natural language processing with the bidirectional encoder representations from transformers (BERT) model [Devlin et al. (2019)], which learns contextual representations by masking tokens and predicting them during pretraining. Adapting this principle, [He et al. (2022)] introduced MIM to computer vision, masking portions of an image during pretraining and tasking the model with reconstructing the missing regions. This paradigm has exhibited strong generalization across both the natural and medical image domains [Gupta et al. (2025); Liu et al. (2021); Huang et al. (2023)]. Representative methods include MAE [He et al. (2022)], MedMAE [Gupta et al. (2025); Xiao et al. (2023)], and SparK [Tian et al. (2023)]: MAE employs high-ratio random masking to steer ViTs toward high-level semantic representations; MedMAE transfers this idea to medical

images through domain-specific adaptations; and SparK reformulates the strategy for convolutional backbones, exploiting their intrinsic multi-scale features to enhance pretraining effectiveness.

For all these advances, conventional MIM approaches typically adopt random or structure-agnostic masking strategies that disregard the semantic requirements of downstream tasks. Such non-selective masking can lead the model to overemphasize irrelevant regions during pretraining, eroding transfer effectiveness. To counter this, recent work has explored region-aware masking strategies intended to improve semantic alignment, along two principal directions.

The first line of work introduces loss-driven adaptive masking: methods such as hard patches mining (HPM) Wang et al. (2023a) and its extension AnatoMask Li et al. (2024) dynamically analyze reconstruction errors across regions and assign higher masking probabilities to areas of larger loss, seeking to prioritize semantically valuable structures. This criterion, however, can misdirect the model toward hard-to-reconstruct yet task-irrelevant regions. In brain magnetic resonance imaging (MRI), for instance, structurally complex regions such as the spinal cord or the eyeballs may attract disproportionate attention despite being unrelated to downstream objectives such as tumor detection.

The second line of work incorporates external perceptual modules to guide masking. FocusMAE Basu et al. (2024), for example, leverages pretrained Region Proposal Networks (RPNs) to identify high-information regions for selective masking. Although effective on specific tasks, this strategy requires supervised training of RPNs on domain-specific annotated data, which constrains its cross-task generalization.

Collectively, these limitations expose the need for a masking mechanism that is at once more flexible, semantically aligned, and generalizable. Natural language, a rich and interpretable source of supervision, offers a compelling avenue for guiding masked modeling in a task-aware yet annotation-free manner.

2.2 VISION-LANGUAGE PRE-TRAINED MODELS

The rapid maturation of vision-language alignment models has rendered it increasingly feasible to perceive and identify task-relevant regions accurately under zero-shot and open-world settings Radford et al. (2021). Foundational approaches such as the Contrastive Language-Image Pretraining (CLIP) family induce a unified embedding space for images and text through large-scale joint pretraining on image-text pairs, and have been widely adopted for zero-shot and open-set visual tasks. Within medical imaging, extensions such as BiomedCLIP Zhang et al. (2023) further sharpen cross-modal understanding, delivering robust performance across diverse downstream applications. Yet, although vision-language pretrained models are now pervasive across supervised learning tasks, the prevailing body of work remains confined to downstream applications such as object detection and semantic segmentation (Zhong et al. 2022; Aleem et al. 2024; Koleilat et al. 2024), leaving unexplored how their open-vocabulary and zero-shot capabilities might benefit masked image modeling.

In sum, recent region-aware MIM approaches—dynamic masking driven by reconstruction loss, or masking guided by external perception modules—improve semantic focus only partially, and remain encumbered by recurring limitations: dependence on additional supervision, misalignment between masked regions and task semantics, and limited flexibility. In parallel, notwithstanding the strong downstream performance of vision-language models, little effort has been directed at incorporating natural language prompts into the MIM pretraining phase to furnish task-aware masking guidance. These observations converge on a clear and promising direction: harnessing vision-language semantic guidance to enable efficient and generalizable masked image modeling.

3 METHODOLOGY

3.1 OVERALL FRAMEWORK

As illustrated in Figure 1, **Mask What Matters (MWM)** is organized into two core stages: (1) text-guided region localization and (2) region-aware masked image modeling. The first stage employs a frozen vision-language model to identify task-relevant regions from an open-vocabulary prompt; the second stage applies differentiated masking ratios across these regions to steer the reconstruction




Figure 1: Overview of the MWM framework. The top panel depicts the three-stage pipeline: (1) text-guided region localization, (2) region-aware masked image modeling, and (3) downstream transfer. The bottom panel provides a detailed view of how textual prompts drive region localization and the ensuing differentiated masking.

process, so that the spatial distribution of the reconstruction signal is conditioned on the localized semantics.

3.2 TEXT-GUIDED REGION LOCALIZATION

MWM localizes task-relevant regions through a multi-stage pipeline whose stages compose into a single spatial guide. It begins with LLM-based prompt generation and employs BiomedCLIP to extract cross-modal embeddings. From these embeddings a saliency map is produced via the Multi-Modal Information Bottleneck (M2IB), binarized with K-Means, and refined using the Segment Anything Model (SAM). Finally, the refined mask is converted into a bounding box with a controllable expansion margin, yielding a robust spatial guide for the subsequent masked modeling stage.

Prompt Generation. In medical image analysis, the design of the textual prompt is a decisive determinant of localization quality. To strengthen semantic guidance, we query large language models (LLMs) with a structured template such as “Describe the typical visual characteristics of a [category] in a [modality] image.” Alongside these full-sentence prompts, we also admit concise category phrases as an alternative input modality. A detailed comparison of the two prompt styles is reported in the experimental section.

Cross-Modal Embedding. Once the prompt is generated, a BiomedCLIP model [Zhang et al. \(2023\)](#) encodes the medical image and its prompt into a shared semantic space. Given an image I and a prompt T , the embeddings are computed as:

$$Z_{\text{img}} = \Phi_{\text{img}}(I), \quad Z_{\text{text}} = \Phi_{\text{text}}(T) \quad (1)$$

where Φ_{img} and Φ_{text} denote the vision and language encoders, respectively. These cross-modal embeddings constitute the semantic foundation on which saliency estimation is built in the subsequent stages.

Saliency Map Generation. A cross-modal saliency estimation module grounded in the Multi-modal Information Bottleneck (M2IB) [Wang et al. \(2023b\)](#) localizes task-relevant visual regions by aligning image and text embeddings while suppressing modality-specific redundancy. Formally, the saliency map $\lambda_S \in [0, 1]^{H \times W}$ assigns to each spatial location an importance weight proportional to

its semantic relevance to the prompt. M2IB obtains this map by optimizing the following objective:

$$\mathcal{L}_{\text{M2IB}} = \text{MI}(Z_{\text{img}}, Z_{\text{text}}) - \gamma \cdot \text{MI}(Z_{\text{img}}, I) \quad (2)$$

where Z_{img} and Z_{text} are the image and text embeddings encoded by BiomedCLIP, and γ governs the trade-off between preserving cross-modal relevance and filtering out task-irrelevant visual information.

Binary Mask Extraction. Given the saliency map, we binarize it via unsupervised K-Means clustering [Lloyd \(1982\)](#) to delineate the region indicated by the text prompt, producing a preliminary foreground mask. Concretely, the map is partitioned into two pixel clusters, and the cluster with the higher saliency values is designated foreground. To refine this mask and suppress spurious activations, we apply connected-component analysis to the binary map and retain only the largest foreground regions. This procedure filters out noisy edge areas and concentrates attention on the core semantic structures.

Region Refinement. To sharpen spatial accuracy, we incorporate SAM as a refinement module [Kirillov et al. \(2023\)](#). For each selected connected region $c_i \in C^*$, its minimal enclosing bounding box is computed and supplied as the visual prompt to SAM.

Given the original image I and the set of bounding boxes, SAM predicts a refined segmentation mask:

$$M_{\text{SAM}} = \text{SAM}(I, \text{Box}(c_i)) \quad (3)$$

The resulting mask M_{SAM} yields a more faithful spatial delineation of the target structure.

Box-Based Region Expansion. The SAM mask produced in the previous stage retains residual boundary noise that can misdirect the model’s focus during representation learning. To render the region guidance robust to such noise, the refined SAM mask is converted into a bounding box with a controllable expansion margin, which serves as the final region of interest (ROI) for downstream masked modeling.

3.3 REGION-AWARE MASKED IMAGE MODELING

Given the region of interest produced by the localization module in Section 3.2, MWM performs masked image modeling through a three-stage pipeline: (1) a multi-level masking strategy that assigns differentiated masking ratios according to region importance, (2) a sparse encoder that processes only the visible (unmasked) patches, and (3) a hierarchical decoder that reconstructs the full image from the sparse multi-scale features.

Multi-Level Semantic Masking. MWM employs a multi-level masking strategy that operationalizes the semantic regions identified by the text-driven localization process. Patches falling within the regions highlighted by the textual prompts—those associated with tumors or organs—are masked at a high ratio, compelling the model to reconstruct, and thus to represent, the task-relevant semantic content. Background regions, by contrast, are masked lightly or left visible, curtailing redundant computation. A visual comparison of the resulting masking patterns is shown in Figure 2.

Sparse Feature Encoding. The visible patches are aggregated into a sparse image and encoded following the SparK design [Tian et al. \(2023\)](#):

$$\begin{aligned} X_{\text{sparse}} &= \{x_i \mid m_i = 1\}, \\ F &= \Phi_{\text{sparse}}(X_{\text{sparse}}) \end{aligned} \quad (4)$$

where x_i is the i -th patch, $m_i \in \{0, 1\}$ indicates visibility, Φ_{sparse} is the encoder, and $F = \{F_1, F_2, F_3, F_4\}$ are the resulting multi-scale features. Although we instantiate MWM with convolutional backbones, the same interface extends to ViTs through token indexing, which is what renders the design backbone-agnostic.




Figure 2: Comparison of masking strategies. Gray blocks denote masked regions.

Hierarchical Reconstruction. A lightweight UNet-style decoder reconstructs the full image in a top-down manner. At each scale l , the masked locations are populated with a learned embedding M_l :

$$F'_l(i) = \begin{cases} F_l(i), & \text{if patch } i \text{ is visible,} \\ M_l, & \text{if patch } i \text{ is masked.} \end{cases} \quad (5)$$

The densified map F'_l is projected by a scale-specific layer ϕ_l to obtain D_l , which is then upsampled and fused with the features at the next lower level:

$$D_{l-1} = B_{l-1}(D_l) + \phi_{l-1}(F'_{l-1}). \quad (6)$$

Finally, the model is trained with mean-squared error on masked patches:

$$\mathcal{L}_{\text{recon}} = \frac{1}{|\mathcal{I}_{\text{mask}}|} \sum_{i \in \mathcal{I}_{\text{mask}}} \|\hat{x}_i - x_i\|^2, \quad (7)$$

where $\mathcal{I}_{\text{mask}}$ denotes the set of masked patch indices, x_i the original patch, and $\hat{x}_i$ its reconstruction; the loss is therefore evaluated exclusively over the masked locations.

4 EXPERIMENTS

4.1 DATASETS

Self-Supervised Pretraining Datasets. Self-supervised pretraining draws on three datasets spanning distinct imaging modalities. (1) *Brain MRI* comprises approximately 17,000 images drawn from Brain Tumor MRI [Cheng et al. (2015)], BRISC [Fateh et al. (2025)], and BraTS 2018 [Menze et al. (2014)], covering a range of tumor types and anatomical structures. (2) *Lung X-ray* consists of roughly 39,000 chest radiographs from COVID-QU-Ex [Rahman et al. (2021)]; Chowdhury et al. (2020), spanning COVID-19, non-COVID pneumonia, and normal cases. (3) *Chest CT* comprises 12,000 slices from a lung-disease dataset [Konya (2020)], accompanied by segmentation masks of fibrotic lungs from 107 patients.

Downstream Fine-Tuning Datasets. Downstream evaluation, spanning classification and detection, employs three further datasets. (1) *Brain Tumor MRI* [Feltrin (2023)] contains 4,479 MRI images spanning 44 tumor types, including astrocytoma, glioblastoma, meningioma, and ependymoma. (2) *Pediatric Lung X-ray* [Kermayn et al. (2018)] consists of 5,863 pediatric chest X-rays labeled as either normal or pneumonia. (3) *Chest CT Cancer* [Hany (2020)] comprises CT images across four classes: adenocarcinoma, large cell carcinoma, squamous cell carcinoma, and normal tissue.

4.2 EXPERIMENTAL SETUP

To ensure a controlled comparison, all self-supervised methods use backbones of comparable model scale: convolution-based methods adopt ResNet-50 [He et al. (2016)], and transformer-based methods




Figure 3: Representative results of text-guided region localization across the three imaging modalities.

Method	Brain MRI		Chest CT		Lung X-ray	
	Prec.	Rec.	Prec.	Rec.	Prec.	Rec.
<i>Prompt Type</i>						
Phrase	0.22	0.86	0.21	0.99	0.42	0.82
Descrip. Sent.	0.30	0.94	0.37	0.88	0.44	0.66
<i>Impact of SAM</i>						
Coarse Ext.	0.22	0.95	0.36	0.87	0.37	0.84
+ SAM Ext.	0.30	0.94	0.37	0.88	0.42	0.82

Table 1: Text-guided localization performance (occlusion precision/recall) across prompt types and refinement strategies, reported for each of the three modalities. *Coarse Ext.* denotes bounding-box extension of the coarse region; + *SAM Ext.* denotes subsequent refinement with SAM. Bold marks the stronger setting within each block and modality.

use ViT-S (Dosovitskiy et al. (2020)). The input resolution is fixed at 224×224 , and all models are initialized from ImageNet-pretrained weights. Pretraining uses an initial learning rate of 2×10^{-4} , and downstream fine-tuning preserves both the encoder and the input resolution. All experiments are conducted on NVIDIA Tesla A100 cards with 80 GB VRAM under Python 3.9, Ubuntu 22.04.3, and PyTorch 1.10.0. Unless stated otherwise, non-target variables (e.g., prompt type or SAM usage) are held at their empirically optimal settings during evaluation, isolating the factor under study.

4.3 EVALUATION OF TEXT-GUIDED LOCALIZATION

We first assess the text-guided localization module across all three imaging modalities—brain MRI, lung X-ray, and chest CT—and then isolate the influence of prompt type and SAM refinement. Predicted bounding boxes are compared against expert-annotated segmentation masks, and we report occlusion precision ($|R_p \cap R_{gt}|/|R_p|$) and occlusion recall ($|R_p \cap R_{gt}|/|R_{gt}|$), where R_p is the predicted region and R_{gt} the ground-truth annotation.

As Table 1 shows, text-guided localization with well-chosen prompts attains consistently high occlusion recall across every dataset (around 0.82) while sustaining competitive precision. These results establish that the framework identifies task-relevant regions from textual descriptions alone, and that it does so robustly across pronounced modality-specific anatomical variation (Figure 3 provides representative visualizations).

Effect of Prompt Type on Region Localization. To characterize the role of language design, we contrast the two prompt types. As reported in Table 1, descriptive sentence prompts are superior on Brain MRI and Chest CT, confirming that richer semantics aid localization. Concise phrases, in contrast, prevail on lung radiographs in recall (0.82 vs. 0.66), plausibly because the broad anatomical coverage of the lungs renders general terms already sufficient. Prompt granularity thus interacts with anatomy, and the optimal choice tracks the spatial extent of the target structure.

Effect of SAM Refinement on Region Localization. We further quantify the contribution of spatial refinement by applying SAM to the coarse masks. As Table 1 shows, SAM refinement raises precision while preserving high recall: on Brain MRI, precision improves from 0.22 to 0.30 with recall held above 0.94. SAM therefore tightens the localized regions without sacrificing coverage of the target.

4.4 COMPARISON WITH PREVIOUS SSL METHODS

We next evaluate the generalization and representation capacity of MWM against representative self-supervised learning (SSL) methods drawn from both the reconstruction-based and contrastive paradigms. To guarantee a fair comparison, every method is pretrained and evaluated under identical datasets and training settings.

Image Classification. We examine classification performance on the three downstream medical imaging datasets under two standard protocols: (1) *full fine-tuning*, in which the entire model is

updated, and (2) *linear probing*, in which the encoder is frozen and only a linear classifier is trained on top of it.

- **Full fine-tuning:** Table 2 reports classification accuracy for MWM alongside several representative self-supervised methods—MAE [He et al. (2022)], SparK [Tian et al. (2023)], AnatoMask [Li et al. (2024)], SimCLR [Chen et al. (2020a)], MoCoV2 [Chen et al. (2020b)], and BYOL [Grill et al. (2020)]. MWM attains the highest accuracy on every modality, evidencing robust generalization across heterogeneous imaging conditions. The margin is largest on Chest CT, where MWM reaches 97.5% accuracy—exceeding the MIM baselines SparK by +3.1 and AnatoMask by +1.7. This gain follows directly from text-guided masking, which concentrates the reconstruction signal on semantically critical yet spatially sparse structures such as the lung lobes in CT while withholding it from redundant background.
- **Linear probing:** To probe the intrinsic quality of the learned representations, we conduct linear-probing experiments, comparing MWM against the two strongest recent masked image modeling methods, SparK and AnatoMask. As Table 3 shows, MWM outperforms both on all three datasets by clear margins, corroborating that text-driven pretraining yields more discriminative features even when the encoder is held fixed.

Object Detection and Instance Segmentation. To gauge whether the learned representations generalize beyond classification, we evaluate brain-tumor detection and instance segmentation on the BR35H dataset [Hamada 2025]. As Table 4 shows, MWM attains the highest AP^{box} (46.9%) and the second-highest AP^{mask} (45.9%), surpassing SparK by +1.3 and +1.1 points respectively. MWM thus transfers effectively to both object-level localization and pixel-wise segmentation, extending its advantage from image-level recognition to dense prediction.

Taken together, incorporating text-driven masking during pretraining equips the model with more informative and more transferable representations. This semantic guidance improves performance uniformly across classification, detection, and segmentation, confirming its effectiveness as a general principle for visual representation learning rather than a task-specific expedient.

4.5 ABLATION STUDY

We further conduct ablation experiments that isolate the individual contributions of prompt design, SAM-based region refinement, and masking ratio to MWM’s downstream classification performance.

4.5.1 EFFECT OF PROMPT DESIGN AND REGION REFINEMENT

As Table 5 shows, descriptive sentence prompts outperform short phrases on brain MRI and chest CT, whereas the ordering reverses on lung X-ray. This pattern mirrors the localization results of Section 4.3: prompts that yield more accurate region identification also yield stronger downstream accuracy, establishing a direct link between localization fidelity and representation quality. Decisively, both prompt types—irrespective of granularity—outperform random masking on every dataset. This confirms that text-guided masking induces effective semantic alignment between the pretraining and

Method	Type	Brain MRI	Chest CT	Lung X-ray
MoCoV2	CL	95.3	91.7	93.1
BYOL	CL	91.3	93.7	93.9
SimCLR	CL	94.5	93.1	94.1
MAE	MIM	95.4	91.7	94.9
AnatoMask	MIM	96.5	95.8	95.2
SparK	MIM	96.2	94.4	94.7
MWM(Ours)	MIM	96.8	97.5	96.0

Table 2: Full fine-tuning classification accuracy (%) on the three downstream datasets. MIM denotes Masked Image Modeling and CL denotes Contrastive Learning. The best result in each column is shown in bold; MWM ranks first on every modality.

Method	Brain MRI	Chest CT	Lung X-ray
AnatoMask	70.2	63.9	85.3
SparK	69.8	63.5	81.6
MWM(Ours)	71.3	67.0	88.5

Table 3: Linear-probing accuracy (%) on the downstream datasets with a frozen encoder. Best per column in bold; MWM leads on all three modalities.

Method	Detection		Segmentation	
	AP^{box}	AP_{75}^{box}	AP^{mask}	AP_{75}^{mask}
AnatoMask	45.2	53.1	46.2	54.1
SparK	45.6	53.2	44.8	52.0
MWM (Ours)	46.9	53.5	45.9	54.1

Table 4: Object detection (AP^{box} , AP_{75}^{box}) and instance segmentation (AP^{mask} , AP_{75}^{mask}) results (%) on the BR35H dataset. Best per column in bold; MWM leads on both detection metrics and on AP_{75}^{mask} .

downstream objectives, the mechanism on which the observed gains rest. The SAM ablation reinforces this account: incorporating the refinement module during pretraining consistently improves downstream classification, corroborating its role in strengthening text-guided masking.

4.5.2 EFFECT OF MASKING RATIO

As Figure 4 shows, MWM attains 96.8% classification accuracy at a masking ratio of merely 40%—well below the roughly 70% conventionally adopted in masked image modeling—and, despite this reduced ratio, exceeds SparK’s best configuration at 70% masking by +2.4% in absolute accuracy. To exclude the possibility that SparK’s weaker result reflects a suboptimal masking configuration rather than a genuine gap, we evaluate SparK at 40%, 50%, 60%, and 70%; the 70% setting indeed yields SparK’s highest accuracy, confirming that the comparison is made against SparK’s own optimum. These findings overturn the prevailing assumption that high masking ratios are inherently optimal for masked image modeling. They instead establish the advantage of semantically guided, task-aware masking: by preserving task-relevant content while discarding redundancy, MWM delivers stronger downstream performance at substantially lower masking ratios.




Figure 4: Classification accuracy on the Chest CT dataset as a function of masking strategy and masking ratio. MWM peaks at a 40% ratio, surpassing the best SparK configuration (70%).

5 CONCLUSION AND LIMITATIONS

This work presents **Mask What Matters (MWM)**, a self-supervised pretraining framework that unifies text-guided semantic localization with region-aware masking. By using natural language prompts to highlight task-relevant areas and applying differentiated masking across them, MWM realizes semantically aligned representation learning without any per-image annotation.

Comprehensive experiments establish the effectiveness and generalizability of MWM: (1) natural language prompts reliably guide semantic region localization, validating text-driven masking as a workable design; (2) MWM consistently outperforms existing methods such as SparK and AnatoMask in classification across all three imaging modalities, and delivers further gains in detection and instance segmentation; and (3) MWM sustains strong performance even at lower masking

Group	Setting	Brain MRI	Chest CT	Lung X-ray
Prompt Type	No Prompt	96.2	94.4	94.7
	Phrase	96.4	96.5	96.0
	Sentence	96.8	97.5	95.2
Impact of SAM	w/o SAM	96.4	96.8	95.8
	w/ SAM	96.8	97.5	96.0

Table 5: Downstream classification accuracy (%) under varying prompt types and region-refinement strategies across the three imaging modalities. *No Prompt* denotes random masking without text-guided localization and thus serves as the baseline within each block. *w/* and *w/o* denote with and without SAM, respectively. Best per column in bold.

ratios (e.g., 40%), demonstrating that semantic guidance, rather than aggressive masking, is what drives representation quality.

Limitations. For all its promise, the framework leaves two questions open. First, although text prompts afford flexible guidance, we have not yet systematically characterized MWM’s robustness to variation in prompt style, vocabulary, or noise. Second, the reliance on external textual descriptions introduces a dependency that may attenuate generalization where reliable prompts are unavailable or inconsistent. These are bounded, well-scoped extensions of the present study rather than obstacles to its central conclusions.

Looking ahead, text-driven, region-aware masking charts a principled path toward more semantically grounded self-supervised learning, and extending MWM beyond medical imaging may open broader opportunities for vision–language pretraining.

REFERENCES

- Sana Aleem, Feiran Wang, Midhun Maniparambil, Eric Arazo, Johannes Dietlmeier, Kevin Curran, Krzysztof J. Geras, Sotirios A. Tsaftaris, Simon J. D. Prince, William M. Wells, William Lotter, Seung Wook Kim, Sebastian Pölsterl, and Samuel Little. Test-time adaptation with salip: A cascade of sam and clip for zero-shot medical image segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 5184–5193, 2024.
- Sourya Basu, Manan Gupta, Chahat Madan, Pulkit Gupta, and Chetan Arora. Focusmae: Gallbladder cancer detection from ultrasound videos with focused masked autoencoders. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 11715–11725, 2024.
- Ting Chen, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. A simple framework for contrastive learning of visual representations. In *Proceedings of the 37th International Conference on Machine Learning (ICML)*, pp. 1597–1607. PMLR, 2020a. URL <https://proceedings.mlr.press/v119/chen20j.html>
- Xinlei Chen, Haoqi Fan, Ross Girshick, and Kaiming He. Improved baselines with momentum contrastive learning. *arXiv preprint arXiv:2003.04297*, 2020b.
- Jin Cheng, Wei Huang, Shuangliang Cao, Riqiang Yang, Wenjia Yang, Zhaoqiang Yun, and et al. Enhanced performance of brain tumor classification via tumor region augmentation and partition. *PLOS ONE*, 10(10):e0140381, 2015.
- Md. E. H. Chowdhury, Tanzila Rahman, Ahsan Khandakar, Rashid Mazhar, Md. Abdul Kadir, Zaid B. Mahbub, and et al. Can AI help in screening viral and COVID-19 pneumonia? *IEEE Access*, 8:132665–132676, 2020.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)*, volume 1, pp. 4171–4186, 2019.

- Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, and Neil Houlsby. An image is worth 16x16 words: Transformers for image recognition at scale. *arXiv preprint arXiv:2010.11929*, 2020.
- Ali Fateh, Yalda Rezvani, Sadegh Moayedi, Sam Rezvani, Farzaneh Fateh, and Mohammad Fateh. Brisc: Annotated dataset for brain tumor segmentation and classification with swin-hafnet. *arXiv preprint arXiv:2506.14318*, 2025.
- Fernando Feltrin. Brain tumor mri images 44 classes, 2023.
- Jean-Bastien Grill, Florian Strub, Florent Altché, Corentin Tallec, Pierre H. Richemond, Elena Buchatskaya, and et al. Bootstrap your own latent: A new approach to self-supervised learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 33, pp. 21271–21284, 2020.
- Amit Gupta, Islam Osman, Mohamed S. Shehata, William J. Braun, and Richard E. Feldman. Med-mae: A self-supervised backbone for medical imaging tasks. *Computation*, 13(4):88, 2025.
- Ahmed Hamada. Br35h :: Brain tumor detection 2020, 2025.
- Mohamed Hany. Chest ct-scan images dataset for lung cancer classification, 2020.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 770–778, 2016.
- Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Ross Girshick. Masked autoencoders are scalable vision learners. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 16000–16009, 2022.
- Stanley C. Huang, Abhishek Pareek, Michael Jensen, Matthew P. Lungren, Serena Yeung, and Akshay S. Chaudhari. Self-supervised learning for medical image classification: A systematic review and implementation guidelines. *NPJ Digital Medicine*, 6(1):74, 2023.
- Anukriti Jaiswal, Abhinav Rajendra Babu, Mojtaba Zadeh Zadeh, Debanjan Banerjee, and Fillia Makedon. A survey on contrastive self-supervised learning. *Technologies*, 9(1):2, 2020.
- Daniel Kermay, Kang Zhang, and Michael Goldbaum. Labeled optical coherence tomography (oct) and chest x-ray images for classification. Mendeley Data, v2, 2018. URL <https://doi.org/10.17632/rschjbr9sj.2>.
- Alexander Kirillov, Eric Mintun, Nikhila Ravi, Hanzi Mao, Chloe Rolland, Laura Gustafson, and et al. Segment anything. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pp. 4015–4026, 2023.
- Tareq Koleilat, Hamed Asgariandehkordi, Hassan Rivaz, and Yuyin Xiao. Medclip-sam: Bridging text and image towards universal medical image segmentation. In *Proceedings of the International Conference on Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, pp. 643–653, Cham, 2024. Springer Nature Switzerland.
- Sandor Konya. Ct lung & heart & trachea segmentation, 2020.
- Rahul Krishnan, Pranav Rajpurkar, and Eric J. Topol. Self-supervised learning in medicine and healthcare. *Nature Biomedical Engineering*, 6(12):1346–1352, 2022.
- Yufei Li, Tianyi Luan, Yuhang Wu, Shuhao Pan, Ying Chen, and Xujie Yang. Anatomask: Enhancing medical image segmentation with reconstruction-guided self-masking. In *Proceedings of the European Conference on Computer Vision (ECCV)*, pp. 146–163, Cham, 2024. Springer Nature Switzerland.
- Xin Liu, Fan Zhang, Zhenyong Hou, Lingqiao Mian, Zi Huang, Jian Zhang, and Jinhui Tang. Self-supervised learning: Generative or contrastive. *IEEE Transactions on Knowledge and Data Engineering (TKDE)*, 35(1):857–876, 2021.

- Stuart Lloyd. Least squares quantization in pcm. *IEEE Transactions on Information Theory*, 28(2): 129–137, 1982.
- Bjoern H. Menze, Andras Jakab, Stefan Bauer, Jayashree Kalpathy-Cramer, Keyvan Farahani, Justin Kirby, and et al. The multimodal brain tumor image segmentation benchmark (brats). *IEEE Transactions on Medical Imaging*, 34(10):1993–2024, 2014.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In *Proceedings of the 38th International Conference on Machine Learning (ICML)*, pp. 8748–8763. PMLR, 2021.
- Tanzila Rahman, Ahsan Khandakar, Yazan Qiblawey, Abdulla Tahir, Serkan Kiranyaz, Sheikh Abul Kashem, Md. Islam, Somaya Al Maadeed, Samar Zughaier, M. A. Khan, and Md. E. H. Chowdhury. Exploring the effect of image enhancement techniques on COVID-19 detection using chest x-rays images. *Computers in Biology and Medicine*, pp. 104319, 2021. doi: 10.1016/j.compbiomed.2021.104319.
- Kai Tian, Yifan Jiang, Qihang Diao, Chaojian Lin, Li Wang, and Zehuan Yuan. Designing BERT for convolutional networks: Sparse and hierarchical masked modeling. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2023.
- Haoran Wang, Kaiyang Song, Jiashuo Fan, Yibing Wang, Jingdong Xie, and Zhenzhong Zhang. Hard patches mining for masked image modeling. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 10375–10385, 2023a.
- Ying Wang, Tim G. J. Rudner, and Andrew Gordon Wilson. Visual explanations of image-text representations via multi-modal information bottleneck attribution. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023b. URL <https://openreview.net/forum?id=ECvtxmVP0x>
- Jing Xiao, Yufan Bai, Alan Yuille, and Zongwei Zhou. Delving into masked autoencoders for multi-label thorax disease classification. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, pp. 3588–3600, 2023.
- Shuai Zhang, Yiqiu Xu, Naoya Usuyama, Hongfang Xu, Jayesh Bagga, Rema Tinn, et al. Biomedclip: A multimodal biomedical foundation model pretrained from fifteen million scientific image-text pairs. *arXiv preprint arXiv:2303.00915*, 2023.
- Yiwu Zhong, Jianwei Yang, Pengchuan Zhang, Chunyuan Li, Noel Codella, Liunian Harold Li, Luowei Zhou, Xiyang Dai, Lu Yuan, Yin Li, and Jianfeng Gao. Regionclip: Region-based language-image pretraining. In *2022 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 16772–16782, 2022. doi: 10.1109/CVPR52688.2022.01629.

A USE OF LARGE LANGUAGE MODELS (LLMs)

We used LLMs *only* for language editing and presentation polishing (e.g., grammar, phrasing, and minor stylistic clarity) of text written by the authors. LLMs were *not* used for idea generation, experimental design, data annotation, code implementation, or analysis/interpretation of results. All technical content, algorithms, proofs, figures, tables, metrics, and conclusions were authored and verified by the authors. No confidential or identifying data were provided to LLMs beyond anonymized manuscript text.