
ADVERSARIAL ATTACKS LEVERAGE INTERFERENCE BETWEEN FEATURES IN SUPERPOSITION

Anonymous authors
Anonymous submission

ABSTRACT

Fundamental questions remain about when and why adversarial examples (AExs) arise in neural networks (NNs), with competing views characterising them either as artifacts of the irregularities in the decision landscape or as products of sensitivity to non-robust input features. This paper advances a third position: adversarial vulnerability can stem from *efficient* information encoding in NNs. Our central contribution is a mechanistic account in which superposition – where networks represent more features than they have dimensions – creates arrangements of latent representations that adversaries directly exploit. Building on this account, we make three specific contributions. First, we show that adversarial perturbations leverage interference between superposed features, which makes attack patterns predictable from feature arrangements and, in turn, yields a mechanistic explanation for two known phenomena: adversarial attack transferability between models with similar training regimes, and class-specific vulnerability patterns. Second, in synthetic settings with precisely controlled superposition we establish that superposition *suffices* to create adversarial vulnerability, and we confirm that the same mechanism persists in a vision transformer (ViT) trained on CIFAR-10. Third, we show that superposition is sufficient but not necessary, isolating algorithmic brittleness as a distinct vulnerability mechanism and demonstrating that a representation-level view is actionable enough to construct attacks directly. Together, these contributions reveal that adversarial vulnerability can be a byproduct of networks’ representational compression, rather than a flaw in the learning process or a property of non-robust inputs. We release our implementation through a software repository.

1 INTRODUCTION

Despite extensive research on AExs (Szegedy et al., 2014; Goodfellow et al., 2014), there remains no consensus on their fundamental cause. We still lack a complete explanation for how minimal perturbations so drastically alter model predictions, leaving us unable to predict which perturbations succeed, explain why attacks transfer between models, or design principled defences. This paper closes part of that gap by tracing adversarial vulnerability to a concrete, measurable property of how networks encode information.

The absence of consensus is reflected in two competing explanations (Nakkiran, 2019), each of which our work addresses. The *bug* perspective attributes vulnerability to factors including learning flaws (Schmidt et al., 2018; Nakkiran, 2019) or geometric properties of high-dimensional decision boundaries (Fawzi et al., 2016); work in this camp provides statistical bounds on properties such as capacity (Bubeck & Sellke, 2021) or Lipschitz constraints (Hein & Andriushchenko, 2017), but does not explain which specific perturbations succeed. The *feature* perspective instead argues AExs exploit predictive but non-robust statistical patterns in data (Ilyas et al., 2019), reframing the debate around how the underlying data relates to human perception; it treats these non-robust features as fixed properties of the data rather than as products of a network’s encoding strategy. Crucially, neither account reconciles how representational constraints interact with data semantics – and it is precisely this interaction that our contributions target.

Our thesis. We bridge this divide through a mechanistic interpretability perspective: adversarial vulnerability emerges from the interaction between architectural constraints and data semantics (*i.e.* human-interpretable properties of the data). Concretely, we show that AExs can emerge from interference between learned representations, as adversarial attacks systematically exploit superposed latent structure – a mechanism that normally enables additional representation capacity

– to craft perturbations that manipulate model outputs. The core deliverable of the paper is a mechanistic pathway that makes this exploitable structure explicit: input correlations constrain feature arrangements, these arrangements determine interference patterns, and these interference patterns dictate attack characteristics and transferability. Unlike prior accounts, this pathway is predictive – it tells us which perturbations succeed and why they transfer between models.

Our account draws on recent theories of how networks encode and process information, namely the linear representation hypothesis (LRH) (Park et al., 2024) and the theory of superposition (Elhage et al., 2022). The LRH posits that *input features*—fundamental abstractions of data—are represented as linear directions in a network’s representation space. Networks are hypothesised to represent far more of these features than they have neurons through superposition, efficiently packing multiple features into shared dimensions at the cost of interference, so that perturbing one feature can affect others in non-obvious ways. The question we pursue is whether this interference is the exploitable structure that AExs target, and what leverage that gives us over adversarial phenomena. We answer it with the following primary contributions:

- **A unifying framework.** We reconcile architectural constraints with feature properties in a single account of adversarial vulnerability. Using controllable models, we demonstrate how data properties induce specific superposition geometries that directly determine the adversarial perturbations a network admits, and show that the resulting interference patterns explain both attack transferability and class-specific susceptibility.
- **Persistence at scale.** We show these findings are not an artifact of toy settings: the same mechanism appears in a ViT trained on CIFAR-10 with an engineered bottleneck layer that induces superposition.
- **Sufficient but not necessary.** We establish that superposition is sufficient, but *not* necessary, for adversarial vulnerability, and in doing so identify algorithmic brittleness as a distinct vulnerability mechanism.
- **An actionable, defence-relevant view.** We demonstrate that a mechanistic understanding of learned representations enables *informed*, gradient-free attack construction, which grounds our broader argument for semantically-informed defences.

2 BACKGROUND

Our contributions rest on two ingredients, which we make precise here: the LRH and superposition. The LRH supplies the notion of a feature direction that our attacks perturb, and superposition supplies the interference that makes those perturbations effective; together they define the vocabulary in which we state our mechanistic pathway. Formally, let $\mathbf{x} \in \mathcal{X}$ denote the input to a NN and $\mathbf{h}^{(l)} \in \mathbb{R}^{d_l}$ the activation vector of the l -th layer.

Linear representation hypothesis (LRH). The LRH posits that NNs represent many variables of their computation, such as semantic properties of their inputs, as linear directions in their activation space, which serve as abstractions for reasoning (Park et al., 2024; Guerner et al., 2023). This linear bias is hypothesised for two reasons: linear separability lets networks easily recognise and manipulate features, and dot products with subsequent-layer weights process such directional features efficiently. A growing body of work supports it (Gurnee et al., 2023; Park et al., 2024). This directional view matters for us because it is exactly these feature directions that our analysis will show adversaries perturb. Let $\mathcal{C} = \{c_1, \dots, c_M\}$ denote a set of M semantically meaningful latent features (e.g. concepts like “presence of shape,” or “indoor vs. outdoor”). Formally:

Definition 1 (Linear Representation Hypothesis (LRH)) A neural network layer with activations $\mathbf{h}^{(l)} \in \mathbb{R}^{d_l}$ satisfies the LRH if it represents the latent features $\mathcal{C} = \{c_1, \dots, c_M\}$ as linear directions $\{\mathbf{v}_j\}_{j=1}^M \subset \mathbb{R}^{d_l}$ such that:

$$\mathbf{h}^{(l)}(\mathbf{x}) \approx \sum_{j=1}^M a_j(\mathbf{x}) \mathbf{v}_j$$

where $a_j(\mathbf{x}) \in \mathbb{R}$ represents the activation magnitude of feature c_j , and $\mathbf{v}_j$ is the corresponding linear direction. The features are **linearly accessible** (Costa et al., 2025): inputs $\mathbf{x}_0, \mathbf{x}_1 \in \mathcal{X}$ that differ mainly in the value of feature c_j while holding other features approximately constant (i.e.,

$a_j(\mathbf{x}_1) - a_j(\mathbf{x}_0) = k$ and $|a_i(\mathbf{x}_1) - a_i(\mathbf{x}_0)| < \lambda$ for all $i \neq j$ and some small $\lambda > 0$) satisfy:

$$\mathbf{h}^{(l)}(\mathbf{x}_1) - \mathbf{h}^{(l)}(\mathbf{x}_0) \approx k\mathbf{v}_j$$

where k reflects the change in c_j .

Superposition and sparsity. NNs can represent many more latent features than they have dimensions in activation space: $M > d_l$. For example, LLMs can reference far more place names than they have residual stream dimensions. Superposition is the framework we adopt to analyse this: networks use an overcomplete, non-orthogonal set of feature directions $\{\mathbf{v}_j\}_{j=1}^M$. This leverages the fact that $2^{\Theta(d\epsilon^2)}$ almost-orthogonal vectors ($< \epsilon$ cosine similarity) fit in d -dimensional space (Tkocz, 2012), and that sparse vectors can be accurately recovered after projection into lower-dimensional spaces (Elhage et al., 2022; Bereska & Gavves, 2024; Sawmya et al., 2025). Packing features this way creates two consequences. First, *polysemanticity* (Scherlis et al., 2022; Lecomte et al., 2023) emerges, where a single neuron contributes to multiple features and its activation no longer corresponds to one interpretable concept. Second – and central to our contributions – non-orthogonal feature directions create *interference*: activating one feature activates others. Networks can partly mitigate this through non-linear operations (e.g. ReLU, softmax) that disambiguate superposed features (Gurnee et al., 2023). It is this second consequence, interference, that we identify as the exploitable structure behind adversarial attacks.

Definition 2 (Superposition Hypothesis) A network layer represents features in superposition if:

1. The number of latent features exceeds layer dimensionality: $M > d_l$
2. There exists non-orthogonal feature directions: $\exists i, j$ with $i \neq j$ such that $\mathbf{v}_i \cdot \mathbf{v}_j \neq 0$
3. Latent feature activations are sparse: $\mathbb{E}_{\mathbf{x}}[\|\mathbf{a}(\mathbf{x})\|_0] \ll M$, where $\mathbf{a}(\mathbf{x}) = [a_1(\mathbf{x}), \dots, a_M(\mathbf{x})]^T$

The network trades representational capacity against feature interference by packing more features than dimensions.

Adversarial attacks. Finally, we define the adversarial perturbations and attacks whose mechanism our contributions explain, following (Szegedy et al., 2014; Costa et al., 2024):

Definition 3 (Adversarial attack) Let $\mathbb{B}_p(\mathbf{x}, \epsilon)$ denote the ϵ -ball around $\mathbf{x} \in \mathcal{X}$ in the ℓ_p -sense, e.g. $\mathbb{B}_\infty$ is the ℓ_∞ -ball. Any small δ that p -perturbs $\mathbf{x}$ to create an **adversarial input** $\mathbf{x}' \in \mathbb{B}_p(\mathbf{x}, \epsilon)$ such that $f(\mathbf{x}') \neq f(\mathbf{x})$ (untargeted) or $f(\mathbf{x}') = y_{\text{target}}$ (targeted) is known to be an **adversarial attack**. ℓ_∞ - or ℓ_2 -norm attacks are created via the projected gradient descent (PGD) (detailed in Sec. A).

Definition 4 (Adversarial robustness) A classifier $g_w(\cdot)$ with parameters w is said to be **robust** at $\mathbf{x}$ with radius $\epsilon > 0$, if $g_w(\mathbf{x}') = g_w(\mathbf{x}) \forall \mathbf{x}' \in \mathbb{B}(\mathbf{x}, \epsilon)$.

3 SUPERPOSITION GEOMETRY CAN DETERMINE ADVERSARIAL ATTACKS

This section delivers our first and central contribution: the mechanistic pathway linking data properties, superposition geometry, and adversarial attacks. We establish it in a controllable synthetic setting, which lets us test each link directly. Concretely, we ask three questions, whose affirmative answers together constitute the contribution:

1. Do adversarial perturbations exploit the interference between superposed features?
2. Do correlations in the input shape the geometric arrangement of superposed latent features?
3. Can shared latent geometry explain why attacks transfer between independently trained models?

Experimental setup. To answer these questions cleanly, we design a synthetic task that exemplifies established LRH and superposition phenomena (Park et al., 2024; Elhage et al., 2022) in a precisely controllable setting. The design serves our contribution directly, meeting four requirements: (1) it provides an intuitive classification setting for studying AExs; (2) it represents class concepts as linear directions per the LRH; (3) it induces controlled superposition between these latent features; and (4) it retains a priori knowledge of how inputs correspond to the superposed features – which is what turns our claims into testable predictions about adversarial mechanisms.

Specifically, we partition input $\mathbf{x} \in \mathbb{R}^d$ into k equally-sized groups $\mathbf{x} = [\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(k)}]$, where $\mathbf{x}^{(j)}$ represents the p input features associated with class j . The task identifies which group has




Figure 1: **Our core claim in one example: an adversarial attack exploiting superposition geometry** ($k = 7, m = 2$). (a) The original sample. (b) The adversarially perturbed sample, whose ground truth is unchanged yet is misclassified; the sign and magnitude of each input perturbation is dictated by the configuration of latent representations, not by the target class. (c) The original and adversarial sample in activation space, where the arrows are the column vectors of $\mathbf{W}_e$ – the latent representations of the input features. The perturbation is legible directly from this geometry, which is precisely the predictability our contribution establishes.

the largest sum: $y = \operatorname{argmax}_{j \in \{1, \dots, k\}} \sum_{i=1}^p x_i^{(j)}$. Intuitively, each $x_i^{(j)}$ is evidence for class j , with the p features capturing different class attributes. Following Elhage et al. (2022), we sample $x_i^{(j)} \sim \text{Uniform}(0, 1)$ with sparsity S (probability of setting to zero).

A two-layer network compresses the input into m dimensions through a bottleneck encoder, $\mathbf{h} = \text{ReLU}(\mathbf{W}_e \mathbf{x} + \mathbf{b}_e) \in \mathbb{R}^m$ (where $m < k < d$), and reconstructs them via a decoder, $\mathbf{z} = \mathbf{W}_d \mathbf{h} + \mathbf{b}_d \in \mathbb{R}^k$. Our primary setup uses cross entropy (CE) loss without ReLUs or biases, chosen for its simplicity; the softmax limits interference by exponentially suppressing contributions from interfering latent features. To show that our contribution does not hinge on this specific choice, we replicate our findings in a second setup with mean squared error (MSE) loss, ReLUs, and biases. The representational challenge is to compress information from $d = k \times p$ dimensions down to $m < k$ dimensions and then recover enough of it to identify the class with the largest sum.

The key advantage of this setup – and what makes the section’s contribution possible – is ground-truth access to the latent features. Whereas non-synthetic tasks involve nuanced concepts that must first be discovered and interpreted, here we can compute the true class after any perturbation by recalculating group sums. This lets us define an AEx unambiguously, as a perturbation satisfying two conditions:

1. The model’s prediction under perturbation changes (*e.g.* from class A to B)
2. The true class remains unchanged (*e.g.* still class A)

This distinguishes genuine AExs from perturbations that actually change the ground truth (*i.e.* change which group has the highest sum), and is what allows us to attribute misclassification to interference rather than to a real change of class. Attacks are created using ℓ_∞ - and ℓ_2 -norm PGD.

Evaluation metrics. We interpret the columns of the encoder weight matrix $\mathbf{W}_e$ as the latent representations $\mathbf{v}_i$ for each input feature x_i . Since input feature $x_i^{(j)}$ contributes exclusively to class j , these $\mathbf{v}_i$ (for $i \in \text{class } j$) align. This direct correspondence between input features and class-aligned latent directions is the lever behind our first contribution: it lets us trace precisely how perturbations to the inputs of a specific class, $\mathbf{x}^{(j)}$, relate to changes in latent space. For successful PGD-generated AExs $\mathbf{x}_{adv} = \mathbf{x}_{orig} + \delta$, we measure four quantities:

1. **Input perturbation profile (IPP):** The sign and magnitude of each δ_i , describing how different input features are perturbed in an attack.
2. **Latent attack alignment:** The similarity $\Delta \mathbf{h} \cdot \mathbf{v}_c$ between the *latent attack vector* $\Delta \mathbf{h} = \mathbf{h}_{adv} - \mathbf{h}_{orig}$ and class c ’s representation.
3. **Attack transferability:** Success rate of attacks generated on one model when applied to another.
4. **Robust accuracy:** The fraction of examples that remain correctly classified after perturbation δ with $\|\delta\|_p \leq \epsilon$ to a sample of inputs.

3.1 RESULTS

Contribution 1a — Do attacks exploit interference between latent features in superposition?

An attack must move a sample across a decision boundary to change its class, but what determines the required input perturbations? Two contrasting intuitions exist, corresponding to the two camps introduced above: the *feature intuition* predicts that attacks simply increase the input features of the target class, while the *bug intuition* predicts that perturbations modify inputs unpredictably, with knowledge of feature representations offering no insight. Our result adjudicates between them.

Fig. 1 shows a typical AEx in our setting: an input of Class 2 (orange) whose perturbed value (with IPP) is misclassified as Class 6 (brown). The perturbations appear arbitrary – we plainly do not simply increase the target-class features, contradicting the feature intuition. Yet they are not random either: they follow a precise pattern, with a systematic correspondence between δ and the configuration of latent features. This is the third intuition our paper argues for – *attacks are mediated by latent geometry* – and the remainder of this subsection makes it quantitative.

Quantitatively, the IPP tracks the latent attack alignment. The sign and magnitude of perturbations $\delta^{(j)}$ correlate strongly with how their class’s latent representation $\mathbf{v}^{(j)}$ aligns with the latent attack vector $\Delta\mathbf{h}$: positively aligned representations ($\mathbf{v}^{(j)} \cdot \Delta\mathbf{h} > 0$) see their input features proportionally amplified, while negatively aligned representations are attenuated. This is exactly what Fig. 1(b) illustrates – input perturbations set by the dot products in Fig. 1(c) to move the sample across the decision boundary.

To confirm that PGD attacks specifically exploit superposition geometry, we compare them against the theoretically optimal, interference-leveraging perturbations we derive in Sec. 3.2. For each configuration (k, m) we train 5 models with different random seeds, generate 1000 PGD attacks per model, and compute the cosine similarity between each successful perturbation and the corresponding optimum; a random baseline uses perturbations of matching ℓ_2 norm. As Tab. 1 shows, alignment is near-perfect across dimensionalities (we filter samples with ℓ_2 norm below ϵ to ensure meaningful perturbations). One-sample t-tests against the random baseline yield $p < 10^{-10}$ for all configurations, so this match is systematic rather than chance. Because the optimal attacks provably leverage interference and PGD aligns with them almost perfectly, we conclude that PGD exploits superposition interference – establishing the first link of our mechanistic pathway.


Table 1: Evidence for Contribution 1a: PGD-discovered attacks align near-perfectly with the theoretically optimal, interference-exploiting perturbations across configurations, and far exceed a norm-matched random baseline. Results show mean $\pm$ std over 1000 attacks per condition; all p -values are below 10^{-10} .

k	m	Cosine Sim. (PGD vs Theory)	Cosine Sim. (Random Baseline)
6	2	0.97 ± 0.02	0.00 ± 0.02
30	10	0.96 ± 0.00	0.00 ± 0.01
90	30	0.92 ± 0.00	0.00 ± 0.00

Contribution 1a established. Adversarial attacks systematically exploit interference between superposed features: successful PGD attacks are predictable from the specific superposition geometry rather than random.

Contribution 1b — Do input correlations determine latent feature geometry? Having shown that geometry determines attacks, we now establish what determines the geometry, closing the next link of the pathway. By systematically manipulating the correlation structure of the training data, we test how correlations constrain the degrees of freedom in learned representations. We use three correlation patterns of increasing constraint strength. Under the *uncorrelated* condition, features are sampled independently as above. Under *paired* correlations, each feature in class i is coupled with a feature in class j , so that when one activates so does its pair. Under *global* correlations, each sample has a random phase that determines a cyclic activation pattern: classes adjacent in index have similar activation probabilities and frequently co-activate, while distant classes rarely activate together.

Fig. 2 shows how these correlations determine the arrangement of $\mathbf{v}^{(j)}$. We quantify geometric similarity by comparing the pairwise cosine-similarity matrices of all feature pairs in each model and then correlating these matrices across random seeds,



where higher correlation indicates more similar arrangements (25 seed pairs per condition). For the $m = 2, k = 7$ setup, geometric similarity rises monotonically with constraint strength: 0.23 ± 0.14 (uncorrelated), 0.47 ± 0.32 (paired), and 0.88 ± 0.06 (global) (mean $\pm$ std). That is, uncorrelated data yields highly variable geometries across seeds, paired correlations partially pin the arrangements, and global correlations force near-identical geometries. Pairwise two-sample t-tests with Bonferroni correction confirm the separation ($p < 10^{-3}$ for all comparisons, $\alpha = 0.05$ corrected for 3 comparisons); results across further (k, m) configurations appear in Sec. C.3.

The mechanism behind this contribution is *interference avoidance*: frequently co-activating features are arranged to minimise mutual interference, so stronger correlations impose tighter constraints. Uncorrelated data admits any arrangement that separates classes, paired correlations fix the relative positions of coupled features, and global correlations determine a unique geometry up to rotation.

Contribution 1b established. Input correlations constrain feature geometry: stronger correlations reduce geometric degrees of freedom, forcing different initialisations to converge to similar arrangements.

Contribution 1c — Does shared geometry explain attack transferability? The two links above combine into a prediction about transfer, which is itself a contribution: if attacks are tailored to a specific interference pattern, they should transfer only between models with similar geometries. We test this directly. Using the same seed pairs as above, we generate attacks on source models and evaluate them on target models within each correlation condition (25 transfer measurements per condition; all models achieve $> 95\%$ clean accuracy). Transfer rates track geometric similarity closely: for the $k = 7, m = 2$ configuration (Fig. 2), globally correlated data yields $94\% \pm 12\%$ transfer, uncorrelated data only $18\% \pm 9\%$, and paired correlations an intermediate $32\% \pm 20\%$ (mean $\pm$ std).

The reason is interference-specific tailoring. Each perturbation component amplifies or suppresses features to create constructive interference that pushes representations across a decision boundary. Applied to a model with a different feature arrangement, those same components cancel or interfere destructively, and the attack fails – which is why transfer succeeds precisely when geometries are shared.

Contribution 1c established. Attack transferability is governed by shared interference patterns.

Contribution 1d — Does reducing superposition suppress these attacks? Our pathway makes a falsifiable prediction: if superposition creates vulnerability through interference, removing superposition should remove the vulnerability. We verify this with three complementary interventions. First, setting $m = k$ lets the network represent each class direction $\mathbf{v}^{(j)}$ orthogonally; here any prediction-changing perturbation also changes the ground truth (moving a sample from class A to B requires genuinely making the sum of class B features exceed class A's), and we find zero successful adversarial examples across 1000 attempts at all tested ϵ . Second, fixing m and varying k isolates superposition from capacity: robust accuracy decreases monotonically with superposition pressure k/m (App. Tab. 5), so vulnerability scales with the degree of interference. Third, constraining one class's latent




Figure 3: Support for the sufficiency claim: an adversarial attack (from class 5 to class 3) leaves the input features of an orthogonally represented class (class 1) untouched – attacks spend their budget only on features in superposition.

vector $\mathbf{v}^{(\perp)}$ to stay orthogonal to all others during training leaves its inputs untouched by attacks between classes that remain in superposition (e.g. from class a to b where $\mathbf{v}_a \not\perp \mathbf{v}_b$; Fig. 3). All three point the same way and directly motivate the sufficiency claim we formalise next.

Contribution 1d established. Adversarial attacks spend their budget specifically on features that are in superposition.

3.2 FORMAL ANALYSIS

To put the empirical contribution on a firm footing, we now derive the mechanism exactly. In our linear setting we can prove the relationships between superposition interference and adversarial vulnerability, moving the result from observational correlation to a provable law. We adopt the notation from our experimental setup, where $\mathbf{W}_e \in \mathbb{R}^{m \times d}$ induces superposition with columns $\mathbf{v}_i$, the decoder is the encoder transposed (as empirically observed), and the logit for class j is $z_j = \mathbf{v}_j^\top \mathbf{h}$.

Proposition 1. *The optimal input perturbations δ that maximise movement from class j to class k under constraint $\|\delta\|_2 = \epsilon$ satisfy $\delta \propto \mathbf{W}_e^\top \mathbf{n}$, where $\mathbf{n} = (\mathbf{v}_k - \mathbf{v}_j)$ is the normal to the decision boundary between classes.*

Proof Sketch. To move from class j to k , we maximise the logit margin $z_k - z_j = \mathbf{n}^\top (\mathbf{h} + \Delta \mathbf{h})$, where $\mathbf{n} = (\mathbf{v}_k - \mathbf{v}_j)$ and $\Delta \mathbf{h} = \mathbf{W}_e \delta$. This becomes $\max_{\|\delta\|_2 = \epsilon} \delta^\top \mathbf{W}_e^\top \mathbf{n}$. By Cauchy-Schwarz, this is maximised when $\delta \propto \mathbf{W}_e^\top \mathbf{n}$. Full proof in Sec. B. $\square$

Corollary 1 (Interference Drives Vulnerability) *The adversarial perturbation magnitude for feature i is $|\delta_i| \propto |\mathbf{v}_i^\top (\mathbf{v}_k - \mathbf{v}_j)|$, directly proportional to the differential interference between feature i and the class representations.*

This corollary is the formal statement of our core contribution: each input feature i is perturbed proportionally to its interference with the class representations. Under superposition, non-orthogonality means semantically unrelated features interfere with the class decision, and adversarial perturbations exploit exactly these dependencies to manipulate outputs.

Proposition 2. *Models with feature representations related by orthogonal transformation $\mathbf{Q}$ (where $\mathbf{Q}^\top \mathbf{Q} = \mathbf{I}$) share identical optimal attack directions in input space.*

Proof Sketch. Under transformation $\mathbf{v}'_i = \mathbf{Q} \mathbf{v}_i$, optimal perturbations remain invariant: $\delta'_i \propto (\mathbf{v}'_i)^\top (\mathbf{v}'_k - \mathbf{v}'_j) = \mathbf{v}_i^\top \mathbf{Q}^\top \mathbf{Q} (\mathbf{v}_k - \mathbf{v}_j) = \mathbf{v}_i^\top (\mathbf{v}_k - \mathbf{v}_j) = \delta_i$. Full proof in Sec. B. $\square$

Together with our empirical findings, these propositions deliver the section’s headline contribution – a complete, provable **mechanistic pathway**: (i) input correlations constrain feature arrangements in superposition, (ii) these geometric arrangements determine interference patterns, (iii) interference patterns dictate optimal perturbations via $\delta \propto \mathbf{W}_e^\top (\mathbf{v}_k - \mathbf{v}_j)$ (Proposition 1), and (iv) shared geometric constraints yield similar interference patterns across models, enabling transferability (Proposition 2):

Correlations $\xrightarrow{\text{constrain}}$ **Feature Geometry** $\xrightarrow{\text{determines}}$ **Interference Patterns** $\xrightarrow{\text{enable}}$ **Transferability**

4 ATTACKS IN REALISTIC MODELS

This section delivers our second contribution: evidence that the mechanistic pathway of Sec. 3 is not a toy-model artifact but survives in a realistic architecture. We study a ViT (Dosovitskiy et al., 2020) trained on CIFAR-10 (Krizhevsky, 2009) (and a ResNet-18 (He et al., 2016) in Sec. D.2.1). Introducing a bottleneck layer directly before the final classification output again forces the latent class features into superposition. We then re-test the three pillars of the pathway at scale: (i) how the degree of superposition affects attack transferability and robustness, (ii) whether a consistent arrangement of superposed representations emerges across initialisations, and (iii) whether attacks exploit feature interference.

Setup. We train our ViT architecture (6 layers, patch size 4, residual stream $d = 512$) on CIFAR-10 across five random seeds. After training, we replace the original classification head with a bottleneck: a linear encoder that compresses the pre-classification activations to an m -dimensional




Figure 4: Contribution 2 at a glance. Left: CIFAR-10 class-representation structure remains similar across seeds, and clusters by semantic similarity. Right: attack transferability and robust accuracy as the bottleneck dimension increases – both trends match the toy-model pathway of Sec. 3.

space, followed by a linear decoder that expands back to the $k = 10$ CIFAR-10 classes. Trained with CE loss and no ReLU – deliberately mirroring Sec. 3 so the two settings are directly comparable – this bottleneck forces the k class representations into m -dimensional superposition, with the base ViT weights frozen. Varying the bottleneck dimension $m \in \{2, 3, 5, 10\}$ controls the degree of superpositional compression. AExs are generated for each bottlenecked model using both ℓ_∞ - and ℓ_2 -norm PGD, with a small step size and many iterations, terminating when an attack changes the prediction (full details in Sec. D.1). Robust accuracy is normalised by clean accuracy to account for capacity-driven performance differences, and attack transferability is measured as in Sec. 3.

Results. Mean clean accuracy of the base models is 81%, and, as expected, the bottlenecked models degrade with compression: 54%, 67%, 76%, and 81% for $m = \{2, 3, 5, 10\}$. Against this backdrop, three observations mirror our toy-model findings and together establish that the pathway transfers to this realistic setting.

Observation 1 (superposition governs robustness and transfer). Fig. 4 (right) shows that as the bottleneck m decreases – i.e. superposition pressure increases – normalised robust accuracy falls while cross-seed attack transferability rises. The drop in robust accuracy follows from the denser packing of class representations, which lets small input perturbations cause widespread interference across compressed features. The rise in transferability is exactly what Sec. 3 predicts: higher superposition reduces the degrees of freedom in feature geometry, so different initialisations converge to more similar geometries, share more interference patterns, and thus admit more transferable attacks. Results across perturbation magnitudes appear in Sec. D.1.1.

Observation 2 (consistent, semantically structured geometry). Fig. 4 (left) shows that the m -dimensional class representations within the bottleneck exhibit consistent ordering and clustering across models. As in Sec. 3, this consistency indicates that inherent data correlations guide the formation of a superposed geometry, and the groupings often reflect semantic similarity between classes (e.g. ‘cat’ and ‘dog’ representations cluster together).

Observation 3 (attacks are not a bottleneck artifact). When AExs generated for a bottlenecked model are run through the corresponding base ViT, they produce similar relative logit changes (i.e. $\text{logit}_{\text{base}}(\mathbf{x}_{\text{adv}}) - \text{logit}_{\text{base}}(\mathbf{x}_{\text{orig}})$), so both models respond consistently to the same perturbations and the perturbations are not solely an artefact of the bottleneck – though without the added compression they do not by themselves cause misclassification. Taken together, these three observations confirm Contribution 2: the superposition-interference mechanism operates in a realistic ViT, not only in toy models.

5 ADVERSARIAL VULNERABILITY VIA ALGORITHMIC BRITTLNESS

This section delivers our third and fourth contributions. Prior sections showed that superposition suffices for adversarial vulnerability; here we show it is not *necessary*, by exhibiting a model whose features are represented orthogonally – not in superposition – yet remains vulnerable. This both isolates a distinct vulnerability mechanism (Contribution 3) and demonstrates that a representation-level view of robustness is actionable (Contribution 4). Three questions organise the section: *how do such attacks work mechanistically, can this understanding construct attacks directly, without gradients, and does naive robustness certification change the underlying failure mode or merely raise the required attack budget?*

Setup. We deliberately choose a task with *orthogonal* features, so that any vulnerability we find cannot be attributed to superposition. We train a linear encoder + MLP on one-hot integers (a, b) to predict $(a + b) \bmod P$. As detailed in Sec. E and Nanda et al. (2023), the model discovers a trigonometric algorithm: the encoder projects inputs onto non-basis-aligned linear representations encoding $\sin(\omega_k a)$ and $\cos(\omega_k a)$ at key frequencies $\omega_k = 2\pi k/P$, and the MLP composes $\cos((a + b)\omega_k)$ via sum-angle identities to reach 100% accuracy.¹ Crucially, these features are represented *orthogonally*, not in superposition. We generate AExs by relaxing the one-hot constraint and perturbing only input a with PGD, moving inputs off the training manifold.

Attacks exploit the learned key frequencies. Fourier analysis of successful perturbations δ (App. Fig. 6) reveals that their dominant frequencies precisely match the learned algorithm’s key frequencies, ω_k , and that 0% of attacks transfer to models with different ω_k (e.g. those trained from different seeds) – the orthogonal-representation analogue of the earlier geometry-specific transfer result. The frequency-aligned perturbations cause misclassification by effectively altering the sinusoidal values read off from the encoder features. Mechanistically, δ turns the one-hot $\mathbf{x}_a$ into a dense vector $\mathbf{x}'_a$, and the contributions from the now non-zero inputs corrupt the sinusoidal values, disrupting the MLP’s downstream computation and causing misclassification. Because the algorithm relies on the constructive interference of each ω_k at the logits, disrupting a single ω_k is insufficient; a successful δ must maximally alter multiple frequency components at once. We defer a detailed intuition to Sec. E.1.

Analytic construction of attacks. The payoff of understanding the mechanism is that we can build attacks directly from it. We construct gradient-free, *informed* perturbations as composite sinusoidal waves of matching ω_k . On $0 + 0 \bmod P$, these succeed with $\|\delta\|_2 = 0.22$ – close to PGD (0.17, only $\approx 1.3\times$ the PGD budget) and far below naive alternatives (single-index 0.97, ℓ_∞ noise 2.04, uniform 6.20). This is strong evidence that we have identified the correct mechanism, and it substantiates Contribution 4: a representation-aware view yields concrete attack recipes, not merely post-hoc explanations.

Consequence for defences — Do defences mitigate this algorithmic vulnerability? The same view sharpens what defences must do. Applying certified training against ℓ_∞ perturbations (Müller et al., 2023), robust accuracy at fixed attack budgets improves from pre- to post-robustification as follows: $0\% \rightarrow 85\%$ at $\epsilon_\infty = 5e-3$, $0\% \rightarrow 96\%$ at $2e-3$, and $76\% \rightarrow 100\%$ at $1e-3$. Yet under PGD the robust model still fails via the same frequency-based brittleness, merely at larger perturbation norms. Certification therefore raises the budget without removing the fragility itself – direct motivation for the semantically-informed defences (here, targeting frequency manipulation) that our framework argues for over broad norm bounds alone.

Contributions 3 and 4, in relation to the paper’s thesis. This section contributes two claims to the paper’s broader narrative. First (Contribution 3), superposition is *sufficient but not necessary*: orthogonal representations remain vulnerable, via a distinct algorithmic-brittleness mechanism. Second (Contribution 4), the *architecture–semantics interaction* lens is actionable – it guides both the design of attacks (representation-informed perturbations) and the development of defences.

As a final link back to the main thesis, we note a shared dependence on sparsity assumptions: upstream violations of sparsity (e.g. noise or superposed features) can deliver effectively ‘pre-perturbed’ inputs to downstream orthogonal circuits, activating the same frequency-based failure.

6 RELATED WORK

Superposition & latent representations. Prior superposition research supplies the ingredients our contribution connects to adversarial vulnerability. It identifies interference mechanisms – Nanda (2023) distinguish *representational* from *computational* superposition, and Gurnee et al. (2023) *alternating* from *simultaneous* interference – and shows that data correlations shape arrangements: correlated feature pairs become orthogonal (Elhage et al., 2022), drive superposition formation (Chan, 2024), and create semantic clustering (Prieto et al., 2025a). Representation geometry is further shaped by spectral bias (Rahaman et al., 2019), neural collapse (Kothapalli, 2023), and optimisation objectives (Casper, 2023), while sparse autoencoders (SAEs) have been used to disentangle superposed features

¹We refer to the specific computation the network executes to solve the task as the *model’s learned algorithm*. Because this algorithm depends on extracting precise values from learned feature representations (frequencies), it stays *brittle* under small perturbations to the bases. We term this *algorithmic brittleness*.

in vision models (Lim et al., 2024; Joseph et al., 2025). Our contribution is to turn these descriptive accounts of geometry into a predictive account of *attacks*.

Adversarial vulnerability. The adversarial literature our work speaks to follows an attack-defence dichotomy. Attacks (Szegedy et al., 2014) use imperceptible perturbations to maximise model error, including FGSM (Goodfellow et al., 2014) and PGD (Madry et al., 2018); adversarial robustness is the maintenance of performance under such attacks, sought through adversarial training (Madry et al., 2018) and certified training (Wong & Kolter, 2018; Palma et al., 2024) that guarantees performance within input bounds. Our Contribution 4 argues these norm-bounded defences leave semantic failure modes intact.

The role of input & latent features on adversarial robustness. Existing feature-level explanations are precisely what our mechanistic pathway unifies and extends. Vulnerability has been attributed to predictive but brittle ‘non-robust’ input features (Ilyas et al., 2019; Goh, 2019) or to learning shortcuts (Li et al., 2023), with perturbation characteristics varying by dataset (Maiya et al., 2021). Elhage et al. (2022) suggest superposition may link to AExs – a conjecture Casper (2023) debate and that our Contribution 1 establishes concretely – and a tension between representational efficiency and robustness has been observed (Barsbey et al., 2025). Zhang et al. (2021) show adversarial perturbations push latent representations across decision boundaries, and Ganeshan et al. (2019) find PGD targets final layers and propose feature-based attacks. Finally, transferability is known to be limited by representation discrepancies between models (Li et al., 2023; Wang et al., 2024), and decorrelating features reduces transfer (Wiedeman & Wang, 2022) – directly corroborating our Contribution 1c that shared geometry drives transfer.

7 DISCUSSION & CONCLUDING REMARKS

This paper set out to explain adversarial vulnerability as a consequence of efficient encoding, and it delivers four contributions toward that goal. Our central result (Contribution 1) is a mechanistic pathway showing that adversarial attacks exploit the interference patterns arising from the geometry of superposed features: across toy models we established that data properties – correlations and sparsity – induce distinct superposition geometries, that these geometries create predictable vulnerabilities, and that they explain attack transferability and class-specific susceptibility, with the mechanism proven exactly in the linear case. We then showed this pathway persists in a ViT on CIFAR-10 (Contribution 2), that superposition is sufficient but not necessary via a distinct algorithmic-brittleness mechanism (Contribution 3), and that the resulting representation-level view is actionable enough to construct attacks and to expose the limits of norm-bounded defences (Contribution 4). Together these reframe adversarial vulnerability as a potential, inherent consequence of how networks efficiently encode vast amounts of information through superposition.

Limitations & future work. The main limitation of these contributions is that they stem primarily from simplified models. While our analyses offer an initial mechanistic lens for attacks in large models, the vast adversarial robustness literature suggests precise vulnerability mechanisms likely vary by attack and model type; for example, constrained attack types (*e.g.* universal perturbations, jailbreaks, patch attacks) may be unable to activate the specific feature combinations required to exploit superposition interference.

Building on these contributions, future work includes quantifying the ‘adversarial cost’ of superposition, its connection to accuracy-robustness trade-offs, and how robust training reshapes superposition geometry. Because our insights derive from engineered superposition of labelled class features, whereas real networks interfere between unknown, unlabelled features across many layers, we plan to use SAEs (Bricken et al., 2023) to extend the pathway to large-scale networks; preliminary results in Sec. F indicate this is a promising direction.

ETHICS STATEMENT

Due to its investigative nature and generality, our mechanistic interpretability framework is not anticipated to have any specific negative impact or immediate societal consequences. Our experiments utilise only publicly available datasets (CIFAR-10) and synthetic toy models, involving no human subjects or personally identifiable information. The gradient-based attack methods used in the paper are already well-established in the literature. Our work does not introduce novel attack capabilities but rather provides theoretical understanding of existing phenomena. We have no conflicts of interest to declare. By elucidating the relationship between superposition and adversarial vulnerability, we hope to inform future research directions that can balance model capability with robustness, ultimately contributing to more trustworthy AI systems that can be safely deployed in critical applications.

REPRODUCIBILITY STATEMENT

We include source code for all experiments in the supplementary materials, organised so that each contribution can be reproduced independently. The toy-model experiments supporting Contribution 1 (Sec. 3) include complete architectural details, training procedures, evaluation metrics, and data-generation processes. The ViT experiments supporting Contribution 2 (Sec. 4) provide architecture and training procedures, with further details in Sec. D.1. The modular-arithmetic experiments behind Contributions 3 and 4 (Sec. 5) detail the network architecture, data encoding, and learned trigonometric algorithm, with additional details in Sec. E.2. Experiments were conducted on an AMD Ryzen Threadripper 3970X 32-Core with 256GB RAM and an RTX 3090. Additional ResNet-18 and SAE experiments are documented in Sec. D.2.1 and Sec. F, respectively.

LLM USAGE STATEMENT

This work utilised LLMs to assist with prior literature search, automate LaTeX table generation from CSV files (all numerical values were verified), and suggest alternative phrasings to improve manuscript clarity. All substantive research contributions and conclusions remain the original work of the authors.

REFERENCES

- Melih Barsbey, Antônio H Ribeiro, Umut Şimşekli, and Tolga Birdal. On the interaction of compressibility and adversarial robustness. *arXiv preprint arXiv:2507.17725*, 2025.
- Leonard Bereska and Stratis Gavves. Mechanistic interpretability for AI safety - a review. *Transactions on Machine Learning Research*, 2024. ISSN 2835-8856.
- Trenton Bricken, Adly Templeton, Joshua Batson, Brian Chen, Adam Jermy, Tom Conerly, Nick Turner, Cem Anil, Carson Denison, Amanda Askell, Robert Lasenby, Yifan Wu, Shauna Kravec, Nicholas Schiefer, Tim Maxwell, Nicholas Joseph, Zac Hatfield-Dodds, Alex Tamkin, Karina Nguyen, Brayden McLean, Josiah E Burke, Tristan Hume, Shan Carter, Tom Henighan, and Christopher Olah. Towards monosemanticity: Decomposing language models with dictionary learning. *Transformer Circuits Thread*, 2023.
- Sébastien Bubeck and Mark Sellke. A universal law of robustness via isoperimetry. In Marc’Aurelio Ranzato, Alina Beygelzimer, Yann N. Dauphin, Percy Liang, and Jennifer Wortman Vaughan (eds.), *Advances in Neural Information Processing Systems*, pp. 28811–28822, 2021.
- Stephen Casper. EIS IX: Interpretability and adversaries. *AI Alignment Forum*, February 2023.
- Lawrence Chan. Superposition is not “just” neuron polysemanticity. *AI Alignment Forum*, April 2024.
- Joana C Costa, Tiago Roxo, Hugo Proença, and Pedro Ricardo Morais Inacio. How deep learning sees the world: A survey on adversarial attacks & defenses. *IEEE Access*, 12:61113–61136, 2024.

-
- Valérie Costa, Thomas Fel, Ekdeep Singh Lubana, Bahareh Tolooshams, and Demba Ba. From flat to hierarchical: Extracting sparse representations with matching pursuit. *arXiv preprint arXiv:2506.03093*, 2025.
- Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. Imagenet: A large-scale hierarchical image database. In *2009 IEEE Conference on Computer Vision and Pattern Recognition*, pp. 248–255, 2009.
- Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, G Heigold, S Gelly, et al. An image is worth 16x16 words: Transformers for image recognition at scale. In *International Conference on Learning Representations*, 2020.
- Rüdiger Ehlers. Formal verification of piece-wise linear feed-forward neural networks. In Deepak D’Souza and K. Narayan Kumar (eds.), *Automated Technology for Verification and Analysis - 15th International Symposium, ATVA 2017, Pune, India, October 3-6, 2017, Proceedings*, volume 10482 of *Lecture Notes in Computer Science*, pp. 269–286. Springer, 2017.
- Nelson Elhage, Tristan Hume, Catherine Olsson, Nicholas Schiefer, Tom Henighan, Shauna Kravec, Zac Hatfield-Dodds, Robert Lasenby, Dawn Drain, Carol Chen, et al. Toy models of superposition. *arXiv preprint arXiv:2209.10652*, 2022.
- Alhussein Fawzi, Seyed-Mohsen Moosavi-Dezfooli, and Pascal Frossard. Robustness of classifiers: from adversarial to random noise. *Advances in neural information processing systems*, 29, 2016.
- Aditya Ganeshan, Vivek BS, and R Venkatesh Babu. FDA: Feature disruptive attack. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 8069–8079, 2019.
- Leo Gao, Tom Dupré la Tour, Henk Tillman, Gabriel Goh, Rajan Troll, Alec Radford, Ilya Sutskever, Jan Leike, and Jeffrey Wu. Scaling and evaluating sparse autoencoders. In *The Thirteenth International Conference on Learning Representations, ICLR 2025, Singapore, April 24-28, 2025*, 2025.
- Gabriel Goh. A discussion of ‘adversarial examples are not bugs, they are features’: Two examples of useful, non-robust features. *Distill*, 2019.
- Ian J Goodfellow, Jonathon Shlens, and Christian Szegedy. Explaining and harnessing adversarial examples. *arXiv preprint arXiv:1412.6572*, 2014.
- Clément Guerner, Anej Svete, Tianyu Liu, Alexander Warstadt, and Ryan Cotterell. A geometric notion of causal probing. *arXiv preprint arXiv:2307.15054*, 2023.
- Wes Gurnee, Neel Nanda, Matthew Pauly, Katherine Harvey, Dmitrii Troitskii, and Dimitris Bertsimas. Finding neurons in a haystack: Case studies with sparse probing. *Trans. Mach. Learn. Res.*, 2023.
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 770–778, 2016.
- Matthias Hein and Maksym Andriushchenko. Formal guarantees on the robustness of a classifier against adversarial manipulation. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- Patrick Henriksen and Alessio Lomuscio. Robust training of neural networks against bias field perturbations. *Proceedings of the AAAI Conference on Artificial Intelligence*, 37(12):14865–14873, Jun. 2023.
- Nikolaus Howe, Ian McKenzie, Oskar Hollinsworth, Michał Zajac, Tom Tseng, Aaron Tucker, Pierre-Luc Bacon, and Adam Gleave. Scaling trends in language model robustness, 2025.
- A. Ilyas, S. Santurkar, D. Tsipras, L. Engstrom, B. Tran, and A. Madry. Adversarial examples are not bugs, they are features. In *Advances in Neural Information Processing Systems 32*, pp. 125–136, 2019.

-
- Sonia Joseph, Praneet Suresh, Lorenz Hufe, Edward Stevinson, Robert Graham, Yash Vadi, Danilo Bzdok, Sebastian Lapuschkin, Lee Sharkey, and Blake Aaron Richards. Prisma: An open source toolkit for mechanistic interpretability in vision and video, 2025.
- Guy Katz, Clark W. Barrett, David L. Dill, Kyle Julian, and Mykel J. Kochenderfer. Reluplex: An efficient SMT solver for verifying deep neural networks. In Rupak Majumdar and Viktor Kuncak (eds.), *Computer Aided Verification - 29th International Conference*, volume 10426, pp. 97–117. Springer, 2017.
- Vignesh Kothapalli. Neural collapse: A review on modelling principles and generalization. *Trans. Mach. Learn. Res.*, 2023.
- Alex Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009. CIFAR-10 dataset.
- Victor Lecomte, Kushal Thaman, Rylan Schaeffer, Naomi Bashkansky, Trevor Chow, and Sanmi Koyejo. What causes polysemanticity? an alternative origin story of mixed selectivity from incidental causes. *arXiv preprint arXiv:2312.03096*, 2023.
- Ang Li, Yifei Wang, Yiwen Guo, and Yisen Wang. Adversarial examples are not real features. *Advances in Neural Information Processing Systems*, 2023.
- Hyesu Lim, Jinho Choi, Jaegul Choo, and Steffen Schneider. Sparse autoencoders reveal selective remapping of visual concepts during adaptation. *arXiv preprint arXiv:2412.05276*, 2024.
- Jack Lindsey, Wes Gurnee, Emmanuel Ameisen, Brian Chen, Adam Pearce, Nicholas L. Turner, Craig Citro, David Abrahams, Shan Carter, Basil Hosmer, Jonathan Marcus, Michael Sklar, Adly Templeton, Trenton Bricken, Callum McDougall, Hoagy Cunningham, Thomas Henighan, Adam Jermy, Andy Jones, Andrew Persic, Zhenyi Qi, T. Ben Thompson, Sam Zimmerman, Kelley Rivoire, Thomas Conerly, Chris Olah, and Joshua Batson. On the biology of a large language model. *Transformer Circuits Thread*, 2025.
- Aleksander Madry, Aleksandar Makelov, Ludwig Schmidt, Dimitris Tsipras, and Adrian Vladu. Towards Deep Learning Models Resistant to Adversarial Attacks. In *International Conference on Learning Representations*, February 2018.
- Shishira R. Maiya, Max Ehrlich, Vatsal Agarwal, Ser-Nam Lim, Tom Goldstein, and Abhinav Shrivastava. A Frequency Perspective of Adversarial Robustness. *arXiv*, 2021.
- Mark Niklas Müller, Franziska Eckert, Marc Fischer, and Martin T. Vechev. Certified training: Small boxes are all you need. In *The Eleventh International Conference on Learning Representations, ICLR 2023, Kigali, Rwanda, May 1-5, 2023*. OpenReview.net, 2023.
- Preetum Nakkiran. A discussion of ‘adversarial examples are not bugs, they are features’: Adversarial examples are just bugs, too. *Distill*, 2019.
- Neel Nanda. A comprehensive mechanistic interpretability explainer and glossary, 2023.
- Neel Nanda, Lawrence Chan, Tom Lieberum, Jess Smith, and Jacob Steinhardt. Progress measures for grokking via mechanistic interpretability. In *The Eleventh International Conference on Learning Representations, ICLR 2023, Kigali, Rwanda, May 1-5, 2023*. OpenReview.net, 2023.
- Alessandro De Palma, Rudy Bunel, Krishnamurthy Dvijotham, M. Pawan Kumar, and Robert Stanforth. IBP regularization for verified adversarial robustness via branch-and-bound. *CoRR*, abs/2206.14772, 2022.
- Alessandro De Palma, Rudy Bunel, Krishnamurthy (Dj) Dvijotham, M. Pawan Kumar, Robert Stanforth, and Alessio Lomuscio. Expressive losses for verified robustness via convex combinations. In *The Twelfth International Conference on Learning Representations, ICLR, 2024*.
- Kiho Park, Yo Joong Choe, and Victor Veitch. The linear representation hypothesis and the geometry of large language models. In *Forty-first International Conference on Machine Learning, ICML 2024, Vienna, Austria, July 21-27, 2024*. OpenReview.net, 2024.

-
- L Prieto, E Stevinson, M Barsbey, T Birdal, and P Mediano. Correlations in the data lead to semantically rich feature geometry under superposition. *Mechanistic Interpretability Workshop at NeurIPS 2025*, 2025a.
- Lucas Prieto, Melih Barsbey, Pedro A. M. Mediano, and Tolga Birdal. Grokking at the edge of numerical stability. In *The Thirteenth International Conference on Learning Representations, ICLR*, 2025b.
- Aditi Raghunathan, Jacob Steinhardt, and Percy Liang. Certified defenses against adversarial examples. In *6th International Conference on Learning Representations, ICLR 2018, Vancouver, BC, Canada, April 30 - May 3, 2018, Conference Track Proceedings*. OpenReview.net, 2018.
- Nasim Rahaman, Aristide Baratin, Devansh Arpit, Felix Draxler, Min Lin, Fred A. Hamprecht, Yoshua Bengio, and Aaron C. Courville. On the spectral bias of neural networks. In *Proceedings of the 36th International Conference on Machine Learning, ICML 2019, 9-15 June 2019, Long Beach, California, USA*, volume 97 of *Proceedings of Machine Learning Research*, pp. 5301–5310. PMLR, 2019.
- Shashata Sawmya, Linghao Kong, Ilia Markov, Dan Alistarh, and Nir N Shavit. Wasserstein distances, neuronal entanglement, and sparsity. In *The Thirteenth International Conference on Learning Representations*, 2025.
- Adam Scherlis, Kshitij Sachan, Adam S Jermyn, Joe Benton, and Buck Shlegeris. Polysemanticity and capacity in neural networks. *arXiv preprint arXiv:2210.01892*, 2022.
- Ludwig Schmidt, Shibani Santurkar, Dimitris Tsipras, Kunal Talwar, and Aleksander Madry. Adversarially robust generalization requires more data. In *Advances in Neural Information Processing Systems 31: Annual Conference on Neural Information Processing Systems 2018, NeurIPS 2018, December 3-8, 2018, Montréal, Canada*, pp. 5019–5031, 2018.
- C. Szegedy, W. Zaremba, I. Sutskever, J. Bruna, D. Erhan, I. Goodfellow, and R. Fergus. Intriguing properties of neural networks. In *2nd International Conference on Learning Representations, ICLR 2014*, 2014.
- Tomasz Tkocz. An upper bound for spherical caps. *The American Mathematical Monthly*, 119(7): 606–607, 2012.
- Donghua Wang, Wen Yao, Tingsong Jiang, Xiaohu Zheng, Junqi Wu, and Xiaoqian Chen. Improving the Transferability of Adversarial Examples by Feature Augmentation. *arXiv*, 2024.
- Christopher Wiedeman and Ge Wang. Disrupting adversarial transferability in deep neural networks. *Patterns*, 3(5), 2022.
- Eric Wong and Zico Kolter. Provable defenses against adversarial examples via the convex outer adversarial polytope. In *International conference on machine learning*. PMLR, 2018.
- Huan Zhang, Hongge Chen, Chaowei Xiao, Sven Gowal, Robert Stanforth, Bo Li, Duane S. Boning, and Cho-Jui Hsieh. Towards stable and efficient training of verifiably robust neural networks. In *8th International Conference on Learning Representations, ICLR*, 2020.
- Shufei Zhang, Zhuang Qian, Kaizhu Huang, Qiufeng Wang, Rui Zhang, and Xinping Yi. Towards better robust generalization with shift consistency regularization. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139, pp. 12524–12534. PMLR, 2021.

APPENDIX

A GENERATION OF ADVERSARIAL EXAMPLES

We use PGD to generate AExs, which is an iterative method to generate attacks (Madry et al., 2018). For untargeted attacks, it maximises the classifier’s loss function $\mathcal{L}(f(\mathbf{x}'), y_{\text{true}})$; for targeted, it minimises $\mathcal{L}(f(\mathbf{x}'), y_{\text{target}})$ by the update rule:

$$\mathbf{x}'^{(k+1)} = \Pi_{\mathcal{S}} \left(\mathbf{x}'^{(k)} \pm \alpha \mathbf{g}_k \right) \quad (1)$$

(+ for maximisation, − for minimisation). $\mathbf{x}'^{(0)}$ is an initial perturbed input, α is step size, $\mathbf{g}_k$ is the (normalised) gradient $\nabla_{\mathbf{x}} \mathcal{L}(f(\mathbf{x}'^{(k)}), y_{\text{class}})$ (y_{class} being y_{true} or y_{target}), and $\Pi_{\mathcal{S}}$ projects onto the ϵ -ball $\mathcal{S} = \{\mathbf{x}' \mid \|\mathbf{x}' - \mathbf{x}\|_p \leq \epsilon\}$, with input domain clipping. We use the ℓ_p norms:

1. ℓ_{∞} : Constraint $\|\delta\|_{\ell_{\infty}} = \max_i |\delta_i| \leq \epsilon$. $\mathbf{g}_k$ is $\text{sign}(\nabla_{\mathbf{x}} \mathcal{L}(\cdot))$. $\Pi_{\mathcal{S}}$ clips each $\mathbf{x}'_i$ to $[\mathbf{x}_i - \epsilon, \mathbf{x}_i + \epsilon]$.
2. ℓ_2 : Constraint $\|\delta\|_{\ell_2} = \sqrt{\sum_i \delta_i^2} \leq \epsilon$. $\mathbf{g}_k$ is $\nabla_{\mathbf{x}} \mathcal{L}(\cdot) / \|\nabla_{\mathbf{x}} \mathcal{L}(\cdot)\|_{\ell_2}$. $\Pi_{\mathcal{S}}$ rescales $\delta = \mathbf{x}' - \mathbf{x}$ if $\|\delta\|_{\ell_2} > \epsilon$ via $\delta \leftarrow \epsilon \cdot \delta / \|\delta\|_{\ell_2}$.

B OMITTED PROOFS & THEORETICAL RESULTS

To understand how adversarial attacks exploit feature representations, we prove that optimal perturbations weight each input dimension by how much its corresponding feature aligns with the path to the decision boundary. We analyse our linear models without ReLU activation functions.

Consider input $\mathbf{x} \in \mathbb{R}^d$ encoded via $\phi(\mathbf{x}) = \mathbf{W}_e \mathbf{x}$ to latent representation $\mathbf{h} \in \mathbb{R}^m$, where the columns of $\mathbf{W}_e \in \mathbb{R}^{m \times d}$ are an overcomplete dictionary $\{\mathbf{v}_i\}_{i=1}^d$ ($m < d$). For our argmax task, we empirically observe that the trained encoder-decoder pair has the decoder as the transpose of the encoder: $\mathbf{W}_d = \mathbf{W}_e^{\top}$. This means the logit for class j is computed as $z_j = \mathbf{v}_j^{\top} \mathbf{h}$, where $\mathbf{v}_j$ is the j -th column of $\mathbf{W}_e$. The predicted class is $\hat{y} = \arg \max_j z_j$. Input perturbations map to latent perturbations as $\Delta \mathbf{h} = \mathbf{W}_e \delta$, where $\delta = (\delta_1, \dots, \delta_d)^{\top}$ are the perturbation coefficients.

The decision boundary between classes j and k is the set of points where their logits are equal: $\mathcal{B}_{jk} = \{\mathbf{h} : z_j = z_k\}$, or equivalently $\{\mathbf{h} : (\mathbf{v}_k - \mathbf{v}_j)^{\top} \mathbf{h} = 0\}$. The vector $\mathbf{n} = (\mathbf{v}_k - \mathbf{v}_j)$ acts as the normal to this pairwise decision boundary, pointing in the direction that maximally increases the margin $z_k - z_j$. We briefly recall the propositions before providing corresponding proofs.

Proposition 1 (Optimal targeted attack). The optimal input perturbations δ that maximise movement from class j toward class k under constraint $\|\delta\|_2 = \epsilon$ satisfy:

$$\delta \propto \mathbf{W}_e^{\top} \mathbf{n}$$

Proof. To move a sample from being classified as j to being classified as k , we need to maximise the logit margin $z_k - z_j$. Under perturbation δ , the new margin becomes:

$$z'_k - z'_j = \mathbf{v}_k^{\top} (\mathbf{h} + \Delta \mathbf{h}) - \mathbf{v}_j^{\top} (\mathbf{h} + \Delta \mathbf{h}) \quad (2)$$

$$= (\mathbf{v}_k - \mathbf{v}_j)^{\top} (\mathbf{h} + \mathbf{W}_e \delta) \quad (3)$$

$$= (\mathbf{v}_k - \mathbf{v}_j)^{\top} \mathbf{h} + (\mathbf{v}_k - \mathbf{v}_j)^{\top} \mathbf{W}_e \delta \quad (4)$$

The change in margin is $(\mathbf{v}_k - \mathbf{v}_j)^{\top} \mathbf{W}_e \delta = \delta^{\top} \mathbf{W}_e^{\top} (\mathbf{v}_k - \mathbf{v}_j)$. We seek:

$$\max_{\delta} \delta^{\top} \mathbf{W}_e^{\top} (\mathbf{v}_k - \mathbf{v}_j) \quad \text{subject to} \quad \|\delta\|_2 = \epsilon$$

Let $\mathbf{g} = \mathbf{W}_e^{\top} (\mathbf{v}_k - \mathbf{v}_j)$. By the Cauchy-Schwarz inequality:

$$|\delta^{\top} \mathbf{g}| \leq \|\delta\|_2 \|\mathbf{g}\|_2 = \epsilon \|\mathbf{g}\|_2$$

Equality is achieved when δ and $\mathbf{g}$ are parallel. Given the constraint $\|\delta\|_2 = \epsilon$:

$$\delta = \frac{\epsilon}{\|\mathbf{W}_e^{\top} (\mathbf{v}_k - \mathbf{v}_j)\|_2} \mathbf{W}_e^{\top} (\mathbf{v}_k - \mathbf{v}_j)$$

Corollary (Interference drives vulnerability). For a targeted attack from class j to class k , the adversarial perturbation magnitude for input feature i is:

$$|\delta_i| \propto |\mathbf{v}_i^\top (\mathbf{v}_k - \mathbf{v}_j)|$$

where $\mathbf{v}_i$ is the i -th column of $\mathbf{W}_e$. This quantity represents the differential interference between feature i and the class representations.

This reveals the mechanism by which superposition creates adversarial vulnerability. Each input feature i is perturbed proportionally to $\mathbf{v}_i^\top (\mathbf{v}_k - \mathbf{v}_j)$ —the differential interference between feature i and the class representations. Under superposition, the non-orthogonality means that even semantically unrelated features have non-zero inner products with $(\mathbf{v}_k - \mathbf{v}_j)$, creating exploitable interference patterns. The multi-class setting amplifies this vulnerability, as with k classes there are multiple possible pairwise boundaries, each creating a distinct interference pattern. Adversarial perturbations leverage these cross-feature dependencies—they manipulate features that affect it through their interference with the class representations. This explains the vulnerability we observe empirically: attacks succeed not by directly increasing target class features, but by exploiting the web of interference created by superposition.

Prop. 1 characterises the optimal perturbation direction for moving from class j to k , providing the gradient for maximising the margin $z_k - z_j$. When intervening classes exist (where $z_j < z_i < z_k$ for some class i), following this gradient might cause the model to predict i before reaching target k . Iterative methods like PGD handle this by recomputing gradients at each step—the global attack path emerges from repeated local decisions rather than a single optimisation. Furthermore, while our analysis assumes $\mathbf{W}_d = \mathbf{W}_e^\top$ based on our empirical observations, the framework extends to arbitrary decoders. In the general case, with decoder $\mathbf{W}_d \in \mathbb{R}^{k \times m}$ having rows $\mathbf{w}_j^\top$, the optimal perturbation becomes $\delta \propto \mathbf{W}_e^\top (\mathbf{w}_k - \mathbf{w}_j)$, capturing interference between encoder directions and decoder weights.

Proposition 2 (Attack transferability). Consider encoders ϕ and ψ with basis matrices $\mathbf{W}_e \in \mathbb{R}^{m \times d}$ and $\mathbf{W}'_e \in \mathbb{R}^{m \times d}$ whose columns are related by orthogonal transformation $\mathbf{v}'_i = \mathbf{Q}\mathbf{v}_i$ (where $\mathbf{Q}^\top \mathbf{Q} = \mathbf{I}$). If both models use their encoder transpose as decoder (i.e., $\mathbf{W}_d = \mathbf{W}_e^\top$ and $\mathbf{W}'_d = (\mathbf{W}'_e)^\top$), then both models have identical optimal input perturbation vectors for any targeted attack from class j to class k .

Proof. For model ϕ , the optimal perturbation from Prop. 1 is:

$$\delta^\phi \propto \mathbf{W}_e^\top (\mathbf{v}_k - \mathbf{v}_j)$$

For model ψ with transformed columns $\mathbf{v}'_i = \mathbf{Q}\mathbf{v}_i$:

$$\delta^\psi \propto (\mathbf{W}'_e)^\top (\mathbf{v}'_k - \mathbf{v}'_j) = (\mathbf{W}'_e)^\top (\mathbf{Q}\mathbf{v}_k - \mathbf{Q}\mathbf{v}_j)$$

Since $\mathbf{W}'_e = \mathbf{Q}\mathbf{W}_e$ (all columns are transformed), we have:

$$\delta^\psi \propto (\mathbf{Q}\mathbf{W}_e)^\top \mathbf{Q}(\mathbf{v}_k - \mathbf{v}_j) \tag{5}$$

$$= \mathbf{W}_e^\top \mathbf{Q}^\top \mathbf{Q}(\mathbf{v}_k - \mathbf{v}_j) \tag{6}$$

$$= \mathbf{W}_e^\top (\mathbf{v}_k - \mathbf{v}_j) \tag{7}$$

where the last equality uses $\mathbf{Q}^\top \mathbf{Q} = \mathbf{I}$. Thus δ^ϕ and δ^ψ have identical proportionality. Under the same norm constraint $\|\delta\|_2 = \epsilon$, we have $\delta^\phi = \delta^\psi$. $\square$

Prop. 2 is used to explain why attacks transfer between models with similar training regimes. When models learn feature representations that differ only by orthogonal transformation—essentially the same geometric structure in different orientations—they share identical vulnerability patterns in input space. The orthogonal transformation preserves all inner products between features, maintaining the interference patterns that attacks exploit. This provides an explanation for our empirical observation that models trained on data with the same correlations exhibit high attack transferability: they discover similar feature geometries up to rotation, leading to shared adversarial vulnerabilities.

C TOY MODEL EXPERIMENTS

This section provides supplementary details, extended results and further discussion for the toy model experiments discussed in the main paper. We present model accuracies across a wider range of parameters than shown in the main text, offering insight into how model capacity and data characteristics like sparsity influence the learning process and the conditions under which feature superposition appears. Subsequently, we offer additional visual examples that correspond to Fig. 1, illustrating the mechanics of adversarial attacks under various conditions.

C.1 HYPOTHESES TESTING FRAMEWORK

We explicitly state our hypotheses for the three research questions in Sec. 3.

Research Q1: Do adversarial perturbations exploit superposition geometry?

- H_0 : Adversarial perturbations are random with respect to feature geometry.
- H_1 : Adversarial perturbations systematically exploit geometric relationships between superposed representations.

Research Q2: Do data correlations determine superposition geometry?

- H_0 : Input correlations have no systematic effect on learned geometries.
- H_1 : Input correlations determine geometric arrangements across model initializations.

Research Q3: Does shared geometry explain attack transferability?

- H_0 : Attack transferability is independent of geometric similarity.
- H_1 : Transferability increases with shared latent structure.

We test these hypotheses through controlled experiments:

- $H_1(1)$: We measure the input perturbation profile alignment with a class’s latent representation and the latent attack vector.
- $H_1(2)$: We systematically vary input correlations and measure resulting geometries.
- $H_1(3)$: We quantify transferability rates across models with varying geometric similarity.

C.2 ACCURACY OF TOY MODEL FOR A RANGE OF PARAMETERS

The toy model experiments presented in the main paper predominantly used low-dimensionality settings for conceptual clarity. To demonstrate the model’s behaviour more broadly, this subsection details the classification accuracies achieved by the CE toy model. These results are presented across varying hidden layer size (h), number of classes (k), number of features, and levels of sparsity (S), to provide insight into when the models learn to represent features in superposition. The sparsity level represents the probability that any individual input feature $x_i^{(j)}$ is set to zero, with higher values of S indicating greater input sparsity. We report results as a function of feature density $1 - S$, the probability that a feature is non-zero. Tab. 2 and Tab. 3 provide context on the model’s performance limits and its ability to learn latent representations in superposition.

C.3 CORRELATIONS CONSTRAIN ARRANGEMENTS

Tab. 4 displays the results for the experiments regarding *Do Input Correlations Determine Latent Feature Geometry?* across a wider range of classes (k) and hidden dimension size (m).

C.4 HOW THE TOY MODEL MAPS TO REAL-WORLD MODELS

How does scaling networks affect adversarial vulnerability? While empirical evidence suggests that larger models tend to be more robust to adversarial attacks, this effect is weak. When adversarial training is employed, clearer scaling trends emerge, but improvements remain largely specific to the attack type used during training rather than conferring general robustness (Howe et al., 2025).

Table 2: Classification accuracy of the CE toy model with a fixed bottleneck dimension ($m = 2$) across various numbers of classes (k), total input features (features = $k \times 3$), and input feature densities ($1 - S$). These results illustrate how input sparsity controls performance degradation as the number of classes to be superposed within a constrained latent space increases.

Classes (k)	Features	Hidden (m)	Accuracy at Input Feature Density ($1 - S$)							
			1.0	0.57	0.33	0.19	0.11	0.06	0.04	0.02
3	9	2	1.00	1.00	1.00	1.00	1.00	1.00	1.00	1.00
4	12	2	0.67	0.71	0.82	0.92	0.98	0.99	1.00	1.00
5	15	2	0.53	0.50	0.65	0.77	0.89	0.95	0.98	0.99
6	18	2	0.40	0.43	0.51	0.66	0.82	0.93	0.97	0.99
7	21	2	0.34	0.34	0.40	0.53	0.73	0.87	0.95	0.98
8	24	2	0.30	0.30	0.33	0.40	0.63	0.82	0.93	0.97
9	27	2	0.24	0.26	0.30	0.35	0.57	0.75	0.89	0.96
10	30	2	0.22	0.24	0.26	0.31	0.50	0.72	0.87	0.95
15	45	2	0.14	0.15	0.16	0.17	0.25	0.40	0.65	0.86
20	60	2	0.10	0.10	0.11	0.13	0.15	0.24	0.44	0.76
25	75	2	0.07	0.08	0.09	0.10	0.12	0.16	0.26	0.62
30	90	2	0.06	0.07	0.07	0.07	0.09	0.11	0.21	0.45

Table 3: Classification accuracy of the CE toy model for varying numbers of classes (k), total input features, bottleneck dimensions (m), and input feature sparsity levels (S). This table reports different numbers of features per class (p).

Classes (k)	Features	Hidden (m)	Accuracy at Input Feature Density ($1 - S$)							
			1.0	0.57	0.33	0.19	0.11	0.06	0.04	0.02
30	30	30	0.23	0.24	0.38	0.62	0.83	0.94	0.99	1.00
30	90	90	0.27	0.27	0.33	0.41	0.51	0.67	0.85	0.94
40	40	30	0.67	0.54	0.73	0.77	0.88	0.96	0.99	0.99
40	120	30	0.71	0.64	0.65	0.72	0.73	0.79	0.89	0.96
60	60	10	0.05	0.07	0.12	0.25	0.47	0.73	0.90	0.97
60	180	10	0.08	0.10	0.13	0.17	0.22	0.32	0.52	0.75
80	80	30	0.15	0.17	0.23	0.41	0.63	0.79	0.91	0.98
80	240	30	0.23	0.22	0.31	0.41	0.48	0.53	0.66	0.80
100	100	10	0.03	0.04	0.05	0.10	0.21	0.43	0.69	0.87
100	300	10	0.04	0.05	0.06	0.09	0.12	0.15	0.25	0.45

Table 4: Geometric similarity results across correlation types and superposition configurations.

Correlation Type	k	m	Geometric Similarity
Uncorrelated	6	2	0.18 ± 0.15
	30	10	0.17 ± 0.02
	90	30	0.17 ± 0.01
Paired	6	2	0.47 ± 0.07
	30	10	0.26 ± 0.07
	90	30	0.25 ± 0.01
Global	6	2	0.92 ± 0.04
	30	10	0.88 ± 0.01
	90	30	0.80 ± 0.01

There are two key phenomena at play when scaling up models. On one hand, larger models have more capacity to represent concepts, but on the other hand, there seems to be a long tail of useful concepts for a larger model to capture in general tasks like next token prediction over internet text. This means that despite increased model capacity, superposition appears to be prevalent even in frontier models (Lindsey et al., 2025). Supporting evidence comes from dictionary learning methods: SAEs require increasingly large dictionaries for larger models (Gao et al., 2025), suggesting that the number of features scales with model size. This suggests that the fundamental tension driving superposition – that models must compress many features into limited dimensions – does not disappear with scale.

Since both superposition and adversarial vulnerability persist in large-scale models, we believe our insights remain relevant across model scales. It is an interesting future avenue to understand how the geometry of superposition changes with scale, potentially helping to mitigate vulnerability.

C.4.1 SEPARATING SUPERPOSITION EFFECTS FROM CAPACITY REDUCTION

Controlling for capacity by keeping the bottleneck dimension fixed to isolate superposition effects controls for the confounding effect of superposition pressure and capacity reduction. We fix bottleneck dimension $m = 2$, and vary number of classes k to isolate superposition pressure. We report the model accuracy, robust accuracy, and transferability in Tab. 5. These controls demonstrate that increased superposition pressure (k/m), independent of model capacity, drives the adversarial vulnerabilities we observe.

Table 5: Performance metrics across correlation types, perturbation budgets (ε), and number of classes (k)

Correlation Type	Parameters		Performance (%)		
	ε	k	Accuracy	Robust Acc.	Attack Trans.
Uncorrelated	0.05	2	100.0	100.0	0.0
		4	100.0	94.9	21.2
		6	99.9	87.0	9.8
		8	99.8	77.7	17.5
	0.1	2	100.0	100.0	0.0
		4	100.0	88.1	9.8
		6	99.7	75.4	9.9
		8	99.6	56.3	23.5
	0.5	2	100.0	100.0	0.0
		4	100.0	99.4	0.0
		6	99.9	62.9	30.5
		8	99.5	24.9	39.0
Fully Correlated	0.05	2	100.0	100.0	0.0
		4	100.0	95.8	98.4
		6	97.3	70.6	98.4
		8	93.0	57.4	97.2
	0.1	2	100.0	100.0	0.0
		4	100.0	88.7	98.5
		6	97.7	45.6	99.7
		8	94.2	34.8	98.6
	0.5	2	100.0	100.0	0.0
		4	100.0	90.0	100.0
		6	97.2	34.6	100.0
		8	93.0	32.0	100.0

C.4.2 ADDITIONAL EXAMPLES OF AEXS IN TOY MODEL

Fig. 1 demonstrates how adversarial attacks exploit the interference between latent features in superposition. Here we provide further visual examples (Fig. 5) to reinforce intuition. Specifically, we supplement the main text by showcasing:

- Fig. 5a shows an additional instance of the setup in Sec. 3.1 ($m = 2, k = 7$) with an ℓ_2 -norm PGD attack, demonstrating the IPP and latent space manipulations that lead to misclassification.
- Fig. 5b shows a similar setup ($m = 2, k = 7$) using an ℓ_∞ -norm PGD attack on a less sparse input.
- Fig. 5c shows an example with increased bottleneck dimensionality ($m = 3, k = 7$) and accompanying ℓ_2 -norm PGD attack. The feature vector similarity matrix used to calculate geometric similarity is also shown.

D CIFAR-10 EXPERIMENTS

To investigate whether the principles observed in the toy models extend to more complex settings, Sec. 4 introduces experiments trained on CIFAR-10 (Krizhevsky, 2009) with an engineered bottleneck. This appendix section provides further details on this setup and presents extended results.

D.1 ARCHITECTURE & TRAINING INFORMATION

The base ViT (Dosovitskiy et al., 2020) architecture comprised 6 transformer layers, an embedding dimension (d) of 512, and 8 attention heads in each transformer layer. Input images from the CIFAR-10 dataset, sized at 32×32 pixels, were processed into patches of 4×4 pixels. The Multilayer Perceptron (MLP) within each transformer block had a hidden dimension of 512. Learned positional embeddings were used.

The bottleneck architecture consisted of a linear encoder followed by a linear decoder. The linear encoder projected the pre-classification activations obtained from the ViT backbone into an m -dimensional latent space ($\{2, 3, 5, 10\}$). The subsequent linear decoder then mapped these m -dimensional representations back to the $k = 10$ dimensions corresponding to the CIFAR-10 classes.

The base ViT model was trained on the CIFAR-10 dataset for 250 epochs. A learning rate of 0.001 was used with the Adam optimiser using default PyTorch parameters. The batch size was set to 512 and a cosine annealing learning rate scheduler was used. The loss function was CE. Dropout was used. Training was performed across five different random seeds to account for variability. After the base ViT model was trained, its weights were frozen. The bottleneck layer was then trained for 30 epochs, utilising a learning rate of 0.001. As for preprocessing, the images undergo RandomCrop with 4-pixel padding, resize to the target size (32x32 by default), and RandomHorizontalFlip for data augmentation.


D.1.1 NORMALISED ROBUST ACCURACY ACROSS PERTURBATION MAGNITUDES

Fig. 4 (right) shows how normalised robust accuracy varies with bottleneck dimension for one value of ϵ . This subsection expands on those findings by detailing the normalised robust accuracy when subjected to PGD attacks of different strengths. Results are presented for both ℓ_2 -norm (Tab. 7) and ℓ_∞ -norm (Tab. 6) PGD attacks, providing a more comprehensive view of how the degree of superposition interacts with attack strength to affect model robustness.


AEXs for these evaluations were generated using PGD with 100 iterations with a step size (α) of 0.01. Robust accuracy was evaluated on 500 samples for each configuration. The mean normalised robust accuracy and standard deviation across five random seeds are reported.

D.1.2 ATTACK TRANSFERABILITY ACROSS PERTURBATION MAGNITUDES


We here include results on attack transferability across various perturbation magnitudes (ϵ) and bottleneck dimensions (m). Tab. 9 presents the ℓ_2 -norm attack transferability and Tab. 8 ℓ_∞ -norm attack transferability.



(a) An ℓ_2 -norm attack changing the classification of an input of class 4 to class 3. The left plot shows original and adversarial activations in latent space, along with representation directions. The right plots show original and perturbed input feature values.



(b) An ℓ_∞ -norm attack changing the classification of an input of class 6 to class 4. The left plot shows original and adversarial activations in latent space relative to class latent directions. The right plots show original and perturbed input feature values.



(c) An ℓ_2 -norm adversarial attack in a 7-class setup with an increased bottleneck dimension $m = 3$. The leftmost plot now shows the cosine similarity matrix between pairs of latent representations for each of the classes.

Figure 5: Visualisations of AExs in the toy model, supplementing Figure 1 from the main paper by illustrating attack mechanisms in activation space and input space under varied conditions.

Table 6: Mean normalised robust accuracy ($\pm$ standard deviation across 5 seeds) for ViT models with different bottleneck dimensions (m) on CIFAR-10, subjected to ℓ_∞ -norm PGD attacks of varying perturbation magnitudes (ϵ). Robust accuracy is normalised by the clean accuracy of each bottlenecked model. These results support the findings of Sec. 4, demonstrating a similar trend across ϵ .

ϵ	Bottleneck Dimension (m)			
	2	3	5	10
0.001	96.0% $\pm$ 0.4%	97.4% $\pm$ 0.7%	97.9% $\pm$ 0.4%	98.1% $\pm$ 0.3%
0.01	61.7% $\pm$ 3.6%	69.6% $\pm$ 3.5%	77.0% $\pm$ 0.9%	81.8% $\pm$ 3.2%
0.05	4.9% $\pm$ 1.5%	6.7% $\pm$ 1.6%	9.3% $\pm$ 0.6%	10.5% $\pm$ 0.8%
0.1	0.1% $\pm$ 0.2%	0.2% $\pm$ 0.3%	0.2% $\pm$ 0.4%	0.2% $\pm$ 0.4%
0.5	0.0% $\pm$ 0.0%	0.0% $\pm$ 0.0%	0.0% $\pm$ 0.0%	0.0% $\pm$ 0.0%

Table 7: Mean normalised robust accuracy ($\pm$ standard deviation across 5 seeds) for ViT models with different bottleneck dimensions (m) on CIFAR-10, subjected to ℓ_2 -norm PGD attacks of varying perturbation magnitudes (ϵ). Robust accuracy is normalised by the clean accuracy of each bottlenecked model. These results support the findings in Section 4 of the main paper, showing decreasing robustness with smaller m (increased superposition) and larger ϵ .

ϵ	Bottleneck Dimension (m)			
	2	3	5	10
0.1	90.4% $\pm$ 1.7%	91.8% $\pm$ 0.7%	94.6% $\pm$ 1.2%	95.0% $\pm$ 0.7%
0.5	58.5% $\pm$ 5.0%	60.8% $\pm$ 5.0%	69.2% $\pm$ 1.8%	72.6% $\pm$ 2.8%
1.0	41.7% $\pm$ 4.8%	44.0% $\pm$ 3.0%	50.4% $\pm$ 0.6%	54.9% $\pm$ 1.7%
2.0	34.0% $\pm$ 4.6%	36.2% $\pm$ 3.3%	41.5% $\pm$ 1.3%	47.4% $\pm$ 2.4%
5.0	30.2% $\pm$ 4.1%	32.9% $\pm$ 3.5%	37.3% $\pm$ 1.9%	42.7% $\pm$ 1.0%

D.2 CORRELATIONS IN CLASS FEATURES

In Sec. 3 it was correlations between inputs that drove superposition arrangements. We note that here it is not the correlations between input classes but rather the correlations in the representations at this point in the network that drive these arrangements. At the classification layer this is likely similar to the confusion matrix – *i.e.* how each class is misclassified in relation to the other classes. Classes that cluster together are those the network finds inherently similar and misclassifies together. To test this we repeat the experiment, finetuning the base ViT using timm/vit_base_patch16_384 which has been trained on ImageNet-21k (14 million images, 21,843 classes) and ImageNet (1 million images, 1,000 classes). After fine-tuning on CIFAR-10 and applying the same bottleneck training procedure as Sec. 4, the resulting geometry shows random ordering between initialisations with approximately equal spacing between features (*i.e.* neuron collapse Kothapalli (2023)). In this case performance is near 100%, meaning the confusion matrix is the identity, and the superposed loses its structure. In contrast, models trained solely on CIFAR-10 converge to similar geometries because they share the same learned difficulty structure - the same pairs of classes prove challenging to distinguish, leading to consistent superposition patterns.

D.2.1 RESNET-18

To address concerns on architectural generalisation, we conduct the same experiments using a ResNet-18 (He et al., 2016) architecture as the base model as opposed to a ViT. We achieve a slightly higher 92% clean accuracy (compared to 89% for ViT). We observe the robust accuracy falls slightly faster

Table 8: Attack transferability (%) for ℓ_∞ -norm PGD attacks on CIFAR-10 ViT models. Transferability is shown from a model trained with a specific ‘Source Seed’ (e.g., Seed 10) to three different target models, each trained with one of the seeds listed in the sub-header (e.g., ‘vs. Seeds 20/30/40’). The three slash-separated values in each cell correspond to the transferability to these three target seeds, respectively. All models within a row share the same bottleneck dimension, m . The ‘Mean $\pm$ Std’ column averages transferability across all 12 source-target seed pairings for each (ϵ, m) configuration. This supports the claim in Section 4 that higher superposition can lead to more consistent latent geometries and thus higher transferability.

ϵ	m	Seed 10 vs. Seeds 20/30/40	Seed 20 vs. Seeds 10/30/40	Seed 30 vs. Seeds 10/20/40	Seed 40 vs. Seeds 10/20/30	Mean $\pm$ Std
0.001	2	77.8/66.7/44.4	25.0/62.5/25.0	72.7/72.7/72.7	58.3/75.0/58.3	59.3 $\pm$ 17.7
	3	57.1/28.6/28.6	14.3/28.6/28.6	63.6/72.7/72.7	42.9/42.9/42.9	43.6 $\pm$ 18.4
	5	70.0/50.0/50.0	57.1/71.4/57.1	50.0/25.0/25.0	16.7/50.0/33.3	46.3 $\pm$ 16.9
	10	55.6/55.6/66.7	12.5/37.5/25.0	16.7/33.3/16.7	28.6/71.4/28.6	37.3 $\pm$ 19.3
0.01	2	60.3/52.6/48.7	63.8/42.0/60.9	53.8/40.9/52.7	51.2/55.8/46.5	52.4 $\pm$ 6.9
	3	42.9/46.4/50.0	46.4/42.9/49.1	44.2/38.9/44.2	42.9/45.1/41.8	44.6 $\pm$ 3.0
	5	43.3/41.1/43.3	39.2/30.4/41.8	46.0/40.2/47.1	41.2/35.3/32.9	40.2 $\pm$ 4.8
	10	42.5/46.0/44.8	38.0/39.4/46.5	43.5/47.8/47.8	32.6/40.0/35.8	42.1 $\pm$ 4.7
0.05	2	35.0/36.4/37.8	38.8/36.0/46.3	37.4/31.3/41.7	42.0/42.0/40.7	38.8 $\pm$ 3.8
	3	29.9/29.9/37.1	30.4/30.1/35.6	30.2/30.2/35.1	30.1/31.9/34.4	32.1 $\pm$ 2.6
	5	26.8/25.3/29.2	25.2/21.5/26.4	28.0/28.0/26.5	26.0/26.6/21.9	25.9 $\pm$ 2.2
	10	27.1/24.9/30.4	21.6/23.0/28.4	25.3/28.1/30.6	21.0/24.9/20.1	25.4 $\pm$ 3.4
0.1	2	34.2/35.6/38.2	36.7/34.9/44.5	37.9/30.0/40.8	41.2/42.0/40.3	38.0 $\pm$ 3.8
	3	29.6/28.9/35.2	28.6/29.2/34.8	29.1/28.2/34.0	28.8/31.1/32.8	30.9 $\pm$ 2.5
	5	25.4/23.3/27.0	23.8/19.6/24.6	25.9/26.7/25.4	24.6/24.3/20.7	24.3 $\pm$ 2.1
	10	24.8/22.6/29.7	20.1/21.6/26.1	22.5/25.5/28.0	19.5/23.3/18.7	23.5 $\pm$ 3.3

for ResNet-18 than the ViT at the same ϵ values. Nevertheless, we observe similar declining trends in normalised robust accuracy and increasing transferability as bottleneck pressure increases.

E MODULAR ADDITION

Sec. 5 introduces a modular addition task to investigate adversarial vulnerability in a scenario where features are orthogonal. This section provides a detailed experimental setup for this task, additional results and further discussion on how we view this sort of analysis to be useful.

E.1 LEARNED ALGORITHMIC REPRESENTATION IN MODULAR ADDITION

As detailed in Nanda et al. (2023) and Prieto et al. (2025b), the network discovers and implements a trigonometric algorithm to solve this task: the encoder projects inputs onto non basis-aligned linear representations encoding $\sin(\omega_k a)$ and $\cos(\omega_k a)$ at key frequencies $\omega_k = 2\pi k/P$. These features are represented orthogonally, *i.e.* not in superposition, and the sinusoidal values can be read off via linear probes. The subsequent MLP then applies sum-angle identities to compute $\cos((a+b)\omega_k)$ for various k , which constructively interfere at the logits to select the correct class. AExs are generated by relaxing the one-hot constraint and perturbing input $\mathbf{x}_j \in \mathbb{R}^P$ (encoding integer j) to $\mathbf{x}' = \mathbf{x} + \delta$ via ℓ_2 and ℓ_∞ -norm PGD. For simplicity, only a is attacked while b is fixed.

As explained in Nanda et al. (2023), the network learns a structured trigonometric algorithm. W_e encodes integers’ sine and cosine values for key frequencies. These features are linearly decodable sine/cosine values. Scaling the corresponding feature vector alters the encoded sinusoidal value itself, not its activation intensity. Integer embeddings $\mathbf{W}_e^j$ project onto a Fourier-like basis, creating a

Table 9: Attack transferability (%) for ℓ_2 -norm PGD attacks on CIFAR-10 ViT models. The table format, detailing source model to target model transferability, mirrors that of Tab. 8; please see its caption for a full explanation. These ℓ_2 results further support the claim in Sec. 4 that higher superposition leads to increased attack transferability.

ϵ	m	Seed 10 vs. Seeds 20/30/40	Seed 20 vs. Seeds 10/30/40	Seed 30 vs. Seeds 10/20/40	Seed 40 vs. Seeds 10/20/30	Mean $\pm$ Std
0.1	2	70.6/64.7/41.2	60.0/50.0/60.0	64.3/60.7/57.1	48.1/63.0/51.9	57.6 ± 8.0
	3	39.1/39.1/47.8	50.0/60.7/50.0	48.1/55.6/51.9	48.0/28.0/40.0	46.5 ± 8.4
	5	43.5/34.8/43.5	31.6/31.6/42.1	42.9/35.7/21.4	30.4/39.1/34.8	35.9 ± 6.4
	10	45.5/59.1/45.5	19.0/28.6/28.6	37.5/43.8/43.8	23.8/52.4/38.1	38.8 ± 11.4
0.5	2	59.6/52.1/48.9	60.8/35.4/60.8	54.1/39.4/51.4	43.8/49.5/46.7	50.2 ± 7.7
	3	38.7/40.6/48.1	36.1/36.8/41.7	37.4/36.5/43.5	34.4/34.4/36.7	38.7 ± 3.9
	5	38.4/32.0/37.6	29.4/25.7/32.1	39.1/33.9/38.3	31.8/29.9/32.7	33.4 ± 4.0
	10	38.3/37.4/37.4	26.5/28.4/31.4	32.4/35.3/35.3	23.4/28.2/25.0	31.6 ± 5.0
1.0	2	50.8/43.7/44.4	51.7/36.4/50.0	46.1/32.9/44.7	42.7/46.7/45.3	44.6 ± 5.3
	3	35.7/36.9/41.7	31.8/36.4/42.6	33.7/32.0/37.3	34.9/32.0/35.5	35.9 ± 3.3
	5	32.1/28.9/33.2	30.9/23.0/30.9	35.9/31.5/33.7	27.7/28.2/27.1	30.3 ± 3.3
	10	32.3/32.3/38.5	23.7/27.7/29.9	29.1/32.0/34.3	22.2/25.6/21.1	29.1 ± 5.0
2.0	2	47.0/41.6/44.3	48.2/36.2/48.2	43.9/33.5/42.1	43.4/43.4/42.8	42.9 ± 4.2
	3	31.9/37.7/40.3	31.1/35.6/40.6	32.6/31.1/36.8	35.4/33.3/34.3	35.1 ± 3.2
	5	29.9/25.4/32.1	26.3/22.1/26.7	32.3/31.8/29.5	27.9/27.9/23.6	28.0 ± 3.2
	10	30.8/30.4/35.7	22.5/24.9/28.6	24.8/29.7/32.7	20.4/26.1/19.0	27.1 ± 4.9
5.0	2	44.7/46.7/40.8	45.3/37.3/48.7	42.7/33.9/41.5	42.3/42.3/40.6	42.2 ± 3.8
	3	32.3/35.8/40.8	31.3/32.6/39.9	31.0/30.0/33.8	35.0/33.0/33.0	34.1 ± 3.2
	5	30.6/24.1/28.6	28.3/20.9/24.8	31.7/31.3/27.4	28.1/29.0/23.2	27.3 ± 3.3
	10	29.8/27.4/32.7	19.7/24.1/28.9	24.9/29.9/31.2	18.8/25.8/19.2	26.0 ± 4.6

circular/periodic structure in the embedding space. Unlike arrangements from data correlations, this circularity arises from the sine/cosine encoding of the (non-correlated) input integers.

Our model learns to utilise approximately seven key frequencies (ω_k). Each embedding W_e^j for an input integer x_j predominantly encodes its fourteen corresponding trigonometric features: $\cos(\omega_i x_j)$ and $\sin(\omega_i x_j)$ for $i \in \{1, 2, \dots, 7\}$. The concatenated 200-dim. vector thus provides these ≈ 28 key trigonometric features (fourteen for a and fourteen for b) to the MLP. The MLP leverages these, via sum-angle identities (e.g., $\cos(A + B) = \cos A \cos B - \sin A \sin B$), to calculate $(x_1 + x_2) \pmod{P}$. The algorithmic brittleness explored in Section 5 stems from the adversarial perturbation of these key learned trigonometric features.

E.2 EXPERIMENT DETAILS FOR MODULAR ADDITION TASK

The network is trained for addition modulo 113. Inputs are integer pairs (a, b) , $a, b \in \{0, \dots, 112\}$. Each integer is mapped to a 100-dimensional vector via a shared embedding $\mathbf{W}_E \in \mathbb{R}^{113 \times 100}$. Embeddings $\mathbf{W}_E^a$ and $\mathbf{W}_E^b$ are concatenated into a 200-dim. feature vector. This vector feeds a 3-hidden-layer MLP (200 ReLU neurons/layer). The output layer provides 113 logits for the classification. The model is trained end-to-end using cross-entropy loss and AdamW (weight decay 4), achieving 100% test accuracy. AExs are generated using PGD under ℓ_2 and ℓ_∞ norms with varying ϵ . Attacks perturb the embedding $\mathbf{W}_E^a$ (with integer b fixed) by adding a perturbation δ to cause misclassification of $(a + b) \pmod{113}$.




Figure 6: Left: Each column is the mean attack δ added to a to successfully attack $a + b \bmod p$. Right: The norms of the Fourier components in the δ and $W_{MLP_{out}}$.

Table 10: Comparison of minimum ℓ_2 -norm budget required to successfully attack $0 + 0 \bmod p$ between different attack types.

Attack Method	Success ϵ
PGD	0.17
Informed Attack	0.22
Single a	0.97
ℓ_∞ Noise	2.04
Uniform δ	6.20

E.3 ALGORITHMIC BRITTLINESS

The frequency-aligned perturbations cause misclassification by effectively altering the sinusoidal values read off from the encoder features. We use the hypothesised neuron-aligned case to give intuition: W_e would contain a row corresponding to the sine/cosine for each ω_k , whose element at column w_e^a contains the value $\sin(\omega_k a)$ or $\cos(\omega_k a)$. For the input a , the one-hot x_a acts to select w_e^a , giving $h_a = w_e^a$. The perturbation δ transforms one-hot x_a into a dense vector x'_a . Now $h'_a = W_e x'_a = (1 + \epsilon_a) h_a + \sum_{k \neq a} \epsilon_k w_e^k$. The introduced contributions from the $k \neq a$ sinusoidal terms corrupts the sinusoidal values, thereby disrupting the MLP’s subsequent computations and leading to misclassification. The model’s algorithm uses the constructive interference of each ω_k at the logits, and as such disrupting a single ω_k is insufficient as the other components overwhelmingly support the correct answer. Thus, effective attacks must find δ that maximally alters multiple frequency components.

E.4 ANALYTIC CONSTRUCTION OF ATTACKS AND BASELINE COMPARISONS

The main paper describes how understanding the learned trigonometric algorithm in the modular addition task allows for the analytic construction of adversarial attacks. Sec. 5 of the main paper benchmarks this ‘informed attack’ against PGD and naive baselines. This section provides further details on these baseline attack methods. The baselines are:

- **PGD Attack:** Standard PGD used to find the adversary.
- **Informed Attack:** The analytically constructed perturbation, from combinations of cosines.
- **Single a :** This involved finding the minimal attack to a single input to cause misclassification.
- **ℓ_∞ Noise:** This baseline involved adding random noise δ_{noise} to the input a . Each component of δ_{noise} was sampled independently and uniformly from $[-\epsilon, \epsilon]$. The ϵ that successfully induced a misclassification is reported.
- **Uniform δ :** The perturbation vector δ is constrained to have all components equal, i.e., $\delta_j = c \cdot 1$, where c is a scalar. The magnitude of c is reported such that adding δ to the input causes a misclassification.

E.5 DETAILS ON CERTIFIED ROBUST TRAINING

To investigate the impact of defences on the identified vulnerability, we performed certified robust training. This aimed to make the network provably robust against ℓ_∞ -norm perturbations. The certified training method employed was an abstraction-based technique (Wong & Kolter, 2018; Raghuathan et al., 2018; Palma et al., 2022), namely RSIP-IBP (Zhang et al., 2020; Henriksen & Lomuscio, 2023). RSIP-IBP utilises Interval Bound Propagation (IBP) (Katz et al., 2017; Ehlers, 2017) and Reversed Symbolic Interval Propagation (RSIP) to compute reachable output sets for network layers, thereby enabling formal certification of robustness within a given perturbation budget. The training loss function incorporated both a standard loss component and an RSIP loss component. Certified training was conducted for 1000 epochs, with the initial 200 epochs serving as a warm-up phase. The learning rate was 1×10^{-3} and weight decay 1×10^{-4} . The same train/test data splits

were used as for the standard model training. The results of certified training across different training perturbation budgets (ϵ_{train}) and evaluated against various attack perturbation budgets (ϵ_{attack}) are summarised in Tab. 11.

Table 11: Verified robust accuracy (%) for modular addition models under ℓ_∞ -norm perturbations. The table compares a standard (pre-robustification) model with models certifiably trained using RSIP-IBP with different training perturbation budgets (ϵ_{train}). Accuracies are evaluated against various attack budgets (ϵ_{attack}). These results support the discussion in Section 5.1 of the main paper, indicating that certified training improved robustness to specified ϵ_{train} but attacks still succeeded by exploiting the same algorithmic sensitivities at higher ϵ_{attack} .

Attack Epsilon (ϵ_{attack})	Train Epsilon (ϵ_{train})	Standard Acc. (Pre-Rob., %)	Robust Acc. (Pre-Rob., %)	Standard Acc. (Robust, %)	Robust Acc. (Robust, %)
5.0×10^{-4}	10^{-4}	100.0	100.0	100.0	100.0
	10^{-3}				100.0
1.0×10^{-3}	10^{-4}	100.0	76.0	100.0	93.8
	10^{-3}				100.0
2.0×10^{-3}	10^{-4}	100.0	0.0	100.0	0.0
	10^{-3}				95.7
5.0×10^{-3}	10^{-4}	100.0	0.0	100.0	0.0
	10^{-3}				85.3

As discussed in Sec. 5, when these robustly trained models were subsequently attacked using PGD, the nature of successful adversarial perturbations remained consistent: they continued to exploit the same key frequency components inherent to the model’s learned algorithm. The primary effect of the certified training was an increase in the perturbation magnitude required to achieve misclassification for those epsilons against which the model was successfully certified, rather than a fundamental change in the attack vector’s alignment with the model’s algorithmic sensitivities.

E.6 ON A POTENTIAL RECONCILIATION WITH SUPERPOSITION VULNERABILITY

Although the vulnerability explored here arises in the absence of feature superposition, it exhibits a notable phenomenological resemblance to those driven by superposition — a shared dependence on sparsity in the input representation. This motivates our *reconciliation conjecture*: algorithmic vulnerabilities may persist in networks with orthogonal representations when sparsity assumptions are violated upstream. If earlier layers in a network produce noisy or non-sparse features, potentially due to adversarial interference, then these deviations can propagate into otherwise well-isolated downstream circuits. In such cases, circuits that operate correctly on clean inputs may nonetheless fail, as the input they receive is effectively already perturbed. This provides a plausible route by which algorithm-disruption mechanisms, like the frequency attacks studied here, could be activated even in the absence of superposition. Further, this interaction between upstream superposition and downstream algorithmic brittleness, and more broadly layer-by-layer understanding of the specific vulnerability mechanisms triggered by AExs, is underexplored. We leave a detailed investigation, grounded in empirical evidence and theoretical modelling, as a promising direction for future work.

F INSPECTING ADVERSARIAL PERTURBATIONS VIA SAE FEATURES

The analyses of the main paper focused on scenarios with either explicitly defined class features (toy model) or class-level representations in a bottleneck (ViT experiments). However, in large, practical networks, interference will occur between an unlabelled latent features across all layers. As discussed in the future work section of the main paper, SAEs offer a promising unsupervised method to extract more linear features from NNs (Bricken et al., 2023). This appendix section presents a preliminary investigation into using SAE-extracted features to characterise how a large scale ViT responds internally when processing AExs versus clean inputs. These initial results are intended to explore the feasibility of this approach for future work aimed at understanding adversarial phenomena

on a feature level within large models, rather than presenting a conclusive new set of findings. We first outline the experimental setup and validate the SAEs behaviour on adversarial inputs before presenting exploratory layer-wise feature difference analyses.

F.1 EXPERIMENT SETUP

We analyse SAE-extracted features from the vision encoder of the LAION/CLIP-ViT-B-32-DataComp.XL-s13B-b90K model, a 12-layer ViT with $\approx 70\%$ zero-shot accuracy on ImageNet-1k Deng et al. (2009). We use open-source Prisma SAEs mapping the 768-dimensional activation space to a dictionary of 49,152 features ($64\times$ expansion) (Joseph et al., 2025). We examine SAEs trained on either mean patch token representations or CLS token representations, across all layers. We analyse features activating above a high threshold (0.1) to capture changes to the highest activating features, and a low threshold (0.001). We generate 10,000 original/successfully-attacked image pairs (ℓ_∞ -norm, $\epsilon = 0.02$) for our SAE feature comparisons. This setup allows for an initial exploration of how adversarial perturbations might manifest at the level of SAE features.

F.2 SAE VALIDITY FOR ADVERSARIAL INPUTS

Given that AExs are out-of-distribution inputs, and our aim is to use SAE (standardly trained) to understand them, we first conduct a basic validation. This subsection examines the impact on model accuracy and SAE reconstruction when SAEs are inserted at inference time for both clean and adversarial inputs. We are investigating whether the SAEs process features relevant to the classification task, even under adversarial conditions, before attempting to interpret differences in feature activations.

Table 12: Impact on ImageNet-1k classification accuracy when inserting a pre-trained SAE (one layer at a time) during inference for original, ℓ_∞ adversarial ($\epsilon = 0.02$), and random noisy inputs. ‘Accuracy’ refers to the ViT’s top-1 accuracy. ‘Original input + SAE’ means the clean ViT activations at a given layer are passed through the SAE then its reconstruction is passed to the next layer. The slight ‘de-attacking’ effect (improved accuracy for ‘Adv. input + SAE’ vs ‘Adv. input’) suggests SAEs might filter some adversarial noise, but overall, the model still largely misclassifies, indicating salient adversarial changes for misclassification are processed.

Metric	Layer 2	Layer 5	Layer 8	Layer 11
Accuracy				
Original input	1.00	1.00	1.00	1.00
Original input + SAE	0.99	0.99	0.99	0.86
Adv. input	0.00	0.00	0.00	0.00
Adv. input + SAE	0.02	0.01	0.01	0.11
Noisy input	0.98	0.98	0.98	0.98
Noisy input + SAE	0.98	0.98	0.98	0.85
SAE Improvement				
Original input	-0.01	-0.01	-0.01	-0.14
Adv. input	0.02	0.01	0.01	0.11
Noisy input	-0.00	-0.00	-0.00	-0.13

F.3 LAYER-WISE SAE FEATURE DIFFERENCES

This subsection presents initial exploratory visualisations comparing SAE feature activations across different layers of the ViT for original, ℓ_∞ adversarial, and random noisy inputs. These visualisations are intended to highlight potential avenues for future, more in-depth investigation into how adversarial attacks distinctively alter sparse feature representations. We examine:

- Activation count: Histograms comparing the number of SAE features activated above a set threshold in a layer for original, adversarial, and noisy images. This explores whether attacks characteristically alter feature sparsity.

- Feature overlap: The degree of overlap (number of common activated features) between (original vs. adversarial) and (original vs. noisy) conditions. This explores if attacks activate a distinctly different set of features.
- Jaccard Similarity: The Jaccard similarity between sets of activated features, offering a normalised measure of overlap.
- Mean activation: Histograms of mean activation values for features active above a threshold. This looks for changes in the intensity of (key) features.
- Distinct Features: The count of unique features activated across an image set for each condition.

F.4 CLS TOKEN SAE FEATURES WITH HIGH ACTIVATION THRESHOLD (0.1)

This subsection presents a preliminary analysis of SAE features derived from CLS tokens, only considering features activated above a threshold of 0.1. Figures 7 through 10 visualise and compare the histogram of activation counts, feature overlap, Jaccard similarity, and histogram of mean activations, respectively, for original, adversarial, and random noisy inputs across different ViT layers.




Figure 7: Histogram of activation counts for CLS token SAE features (threshold 0.1).




Figure 8: Feature overlap between original vs. attacked and original vs. noisy images for CLS token SAE features (threshold 0.1).




Figure 9: Comparison of Jaccard similarity for CLS token SAE features (threshold 0.1).

F.5 CLS TOKEN SAE FEATURES WITH LOW ACTIVATION THRESHOLD (0.001)

Here, we examine SAE features from CLS tokens with a lower activation threshold of 0.001, to not just capture the most intensely activated features. The subsequent figures illustrate the histogram of activation counts (Fig. 11), feature overlap (Fig. 12), Jaccard similarity (Fig. 13), and histogram of mean activations (Fig. 14) across layers for original, adversarial, and noisy inputs.




Figure 10: Histogram of mean thresholds for CLS token SAE features (threshold 0.1).




Figure 11: Histogram of activation counts for CLS token SAE features (threshold 0.001), with noise.




Figure 12: Feature overlap between original vs. attacked and original vs. noisy images for CLS token SAE features (threshold 0.001).




Figure 13: Comparison of Jaccard similarity for CLS token SAE features (threshold 0.001).




Figure 14: Histogram of mean activations for CLS token SAE features (threshold 0.001), with noise.

F.6 PATCH TOKEN SAE FEATURES WITH HIGH ACTIVATION THRESHOLD (0.1)

This subsection shifts the focus to SAE features derived from the mean of patch tokens, again using a high activation threshold of 0.1. We present comparisons of activation count histograms (Fig. 15),

distinct feature count histograms (Fig. 16), feature overlap (Fig. 17), Jaccard similarity (Fig. 18), and mean activation histograms (Fig. 19) for original, adversarial, and noisy inputs.




Figure 15: Histogram of activation counts for patch token SAE features (threshold 0.1), with noise.




Figure 16: Histogram of distinct patch features for patch token SAE features (threshold 0.1), with noise.




Figure 17: Feature overlap between original vs. attacked and original vs. noisy images for patch token SAE features (threshold 0.1).




Figure 18: Comparison of Jaccard similarity for patch token SAE features (threshold 0.1).

F.7 PATCH TOKEN SAE FEATURES WITH LOW ACTIVATION THRESHOLD (0.001)

Finally, this subsection explores SAE features from the patch tokens means using a low activation threshold of 0.001. The figures show histograms of activation counts (Fig. 20), distinct feature counts




Figure 19: Histogram of mean thresholds for patch token SAE features (threshold 0.1), with noise.

(Fig. 21), feature overlap (Fig. 22), Jaccard similarity (Fig. 23), and mean activation histograms (Fig. 24) for original, adversarial, and noisy inputs.




Figure 20: Histogram of activation counts for patch token SAE features (threshold 0.001), with noise.




Figure 21: Histogram of distinct patch features for patch token SAE features (threshold 0.001), with noise.




Figure 22: Feature overlap between original vs. attacked and original vs. noisy images for patch token SAE features (threshold 0.001).




Figure 23: Comparison of Jaccard similarity for patch token SAE features (threshold 0.001).




Figure 24: Histogram of mean activations for patch token SAE features (threshold 0.001), with noise.

G ACRONYMS

AEx adversarial example

CE cross entropy

LRH linear representation hypothesis

MLP Multilayer Perceptron

MSE mean squared error

NN neural network

PGD projected gradient descent

SAE sparse autoencoder

ViT vision transformer