
SPARSITY AND SUPERPOSITION IN MIXTURE OF EXPERTS

Anonymous authors
Anonymous submission

ABSTRACT

Mixture of Experts (MoE) architectures now underpin the scaling of large language models, yet the mechanistic principles that distinguish them from dense networks remain largely uncharacterized. Prior work has established that dense models exploit *superposition*—representing more features than they have dimensions—and that the degree of superposition is governed by feature sparsity and feature importance. We show that this dense-model lens does not transfer to MoEs: neither feature sparsity nor feature importance induces the discontinuous phase changes that characterize dense representations, and it is instead *network sparsity*, the ratio of active to total experts, that governs the representational strategy an MoE adopts. To quantify this behavior, we introduce metrics that measure superposition at the level of individual experts, and we establish that models with greater network sparsity exhibit systematically greater *monosemanticity*. Building on this result, we advance a definition of expert specialization grounded in monosemantic feature representation rather than load balancing, and we demonstrate that experts organize around coherent feature combinations when initialization directs them toward those features. Taken together, these findings indicate that network sparsity can make MoEs more interpretable without a commensurate loss in reconstruction performance, directly challenging the prevailing assumption that interpretability and capability stand in fundamental tension.

1 INTRODUCTION

Mixture of Experts (MoEs) have become the dominant architecture in state-of-the-art language models, including Qwen3, Mixtral, and Gemini (Yang et al., 2025a; Jiang et al., 2024; Google DeepMind, 2025), chiefly because conditional computation delivers substantial gains in both efficiency and performance (Shazeer et al., 2017; Fedus et al., 2022). A rich body of subsequent work has refined the design space, improving routing through mechanisms such as Expert-Choice routing and stabilizing training and transfer through approaches such as ST-MoE (Zhou et al., 2022; Zoph et al., 2022). Complementary theoretical and empirical results demonstrate that learnable routers recover latent cluster structure in the data, offering a first account of why experts specialize (Dikkala et al., 2023). Yet despite this widespread adoption and architectural maturity, MoEs remain poorly understood from a mechanistic interpretability standpoint—a gap that motivates the present study.

A parallel line of interpretability-oriented work has sought to render expert behavior more transparent. Yang et al. (2025b) propose MOE-X, which couples sparsity-aware routing with wide, ReLU-based experts to suppress polysemanticity, while Park et al. (2025) introduce MONET, scaling the expert count so that capabilities can be edited through targeted expert activation. These contributions, however, operate principally at the level of architectural intervention; they leave unresolved the more fundamental questions of how MoEs represent features intrinsically, how the expert structure shapes superposition, and whether experts specialize of their own accord absent additional regularization. Surveying the field, Mu & Lin (2025) identify precisely this mechanistic understanding as a central open challenge—one this paper sets out to address.

A fundamental obstacle to interpreting neural networks is the phenomenon of superposition: the regime in which a model represents more features than it has dimensions. Superposition lets a network pack many sparse features into a smaller set of neurons, but it does so at a direct interpretive cost, rendering individual neurons polysemantic and resistant to clean analysis.

MoE architectures introduce a new axis to this problem: network sparsity. Whereas dense models activate every neuron regardless of the input, an MoE activates only a fraction of its total parameters on any given forward pass (Shazeer et al., 2017). This distinction is representationally consequential. A dense model must exploit feature sparsity by packing many sparse features into shared neurons; an MoE, by contrast, can afford to be selective, dedicating entire experts to specific feature combinations rather than compressing them into a common basis.

We organize our investigation around three questions: (1) whether MoEs exhibit less superposition than their dense counterparts; (2) whether the superposition of a given feature in an MoE expert undergoes the discrete phase change—across the feature’s relative importance and the overall feature sparsity—that is observed in dense models; and (3) whether expert specialization is better understood through the lens of feature representation than through load balancing alone.

We answer these questions using controlled toy models that extend the framework of Elhage et al. (2022) to the MoE setting. Our analysis yields three principal contributions: (1) we establish that, unlike dense models, MoEs do not undergo sharp phase changes, instead exhibiting increasingly continuous representational transitions as network sparsity increases; (2) we demonstrate that MoEs consistently attain greater monosemanticity—that is, less superposition—than dense models matched on both active and total parameters, with individual experts representing features more cleanly; and (3) we introduce an interpretability-focused definition of expert specialization grounded in monosemantic feature representation, and we show that experts organize around coherent feature combinations rather than arbitrary load balancing.

2 RELATED WORK

Superposition and Feature Representations. The Linear Representation Hypothesis holds that networks encode concepts as directions in activation space (Park et al., 2024), yet the number of interpretable features a network must represent routinely exceeds its available dimensions. Elhage et al. (2022) formalized the resolution to this tension as superposition, establishing that dense models pack non-orthogonal features to maximize representational capacity at the cost of polysemanticity. Although superposition is provably efficient as a capacity strategy (Scherlis et al., 2025), it forces reliance on elaborate post-hoc disentanglement methods, notably Sparse Autoencoders, to recover monosemantic features (Bricken et al., 2023). More recent work characterizes how data correlations shape superposition (Prieto et al., 2025) and how interference patterns arise and can be mitigated (Gurnee et al., 2023). We extend this line of inquiry to MoE architectures and establish that it is network sparsity—not feature sparsity alone—that governs the representational strategy a model adopts.

Interpretability of MoEs. Although MoEs have become the standard vehicle for scaling large language models (Shazeer et al., 2017; Fedus et al., 2022; Jiang et al., 2024), mechanistic understanding of their internal representations lags well behind that of dense networks. Existing analyses concentrate on macroscopic behaviors—routing stability (Zoph et al., 2022), expert-choice statistics (Zhou et al., 2022), and latent cluster discovery (Dikkala et al., 2023)—rather than on the geometry of the representations themselves. The more recent interpretability-focused proposals attempt to impose specialization through architectural constraints, whether predefined concept routing (Yang et al., 2025b) or expert counts scaled to match vocabulary sizes (Park et al., 2025), but they prioritize capability editing over an account of intrinsic feature geometry. We close this gap by analyzing directly how the expert structure shapes superposition: we establish that MoEs exhibit greater monoseman-

ticity than dense models, and we advance a feature-based definition of expert specialization that connects this geometry to routing behavior.

3 BACKGROUND

A central aim of mechanistic interpretability is to reverse-engineer neural networks, and a standard approach is to decompose their representations into a set of human-interpretable concepts termed ‘features’. These features are conventionally taken to be linear: any hidden state h admits the decomposition

$$h = \sum_{i \in F} \alpha_i \vec{f}_i + \vec{b}$$

where $\vec{f}_i$ is the direction associated with feature i , α_i is that feature’s activation strength (roughly, the degree to which feature i is present in the input), and F denotes the set of all represented features.

The *superposition hypothesis* holds that a model can represent far more features than it has dimensions, i.e. $|F| > m$ for $h \in \mathbb{R}^m$ (Elhage et al., 2022). Accommodating more features than dimensions within a latent space requires that the feature vectors in F be packed non-orthogonally. When two features exhibit nonzero interference, $\langle \vec{f}_i, \vec{f}_j \rangle \neq 0$, we say they are in *superposition*. Such superposition is tolerable precisely when features are sparsely active—a defining characteristic of most online text data—but it exacts an interpretive cost, since α_i then absorbs spurious activations unrelated to the actual presence of feature i in the input.

Monosemanticity is a property of individual neurons: a neuron is monosemantic when its activation corresponds cleanly to a single α_i , that is, when features are basis-aligned. Crucially, orthogonality and basis-alignment are distinct—when features are orthogonal (hence not in superposition) but not basis-aligned, neurons remain *polysemantic* even though feature interference is minimal. Our focus throughout is on reducing superposition rather than on enforcing basis-alignment; accordingly, for brevity, when we describe a model, expert, or feature as “more monosemantic” we mean that it exhibits less superposition.

4 DEMONSTRATING SUPERPOSITION


MoEs are commonly conceptualized as compositions of dense models, with each expert treated as an independent dense network. Whether experts in fact represent features in the same manner as dense models, however, has not been established. We test this assumption directly by comparing the superposition that arises in MoEs against that in dense models.

4.1 EXPERIMENTAL SETUP

Our objective is to characterize how an MoE projects a high-dimensional vector $x \in \mathbb{R}^n$ into lower-dimensional expert representations $h \in \mathbb{R}^m$ and subsequently reconstructs it. This construction extends the framework of Elhage et al. (2022) to the MoE setting.

The input distribution. The input vector x models the activations of an idealized, fully disentangled model, in which each dimension x_i corresponds to a distinct, independent feature—a perfectly neuron-aligned, monosemantic representation. We draw $x_i \sim U(0, 1)$ subject to a prescribed sparsity $S \in [0, 1)$, so that $P(x_i = 0) = S$. Because features do not contribute equally to the reconstruction objective, we assign each feature x_i a scalar importance I_i that weights its term in the loss. To isolate and study the phase transition of a single feature, we further scale the magnitude of the last feature by a factor $r \in \mathbb{R}^+$, yielding the importance vector $I = (1, 1, \dots, r)$. The parameters I_i and r thus jointly control the signal strength of individual features, enabling us to probe expert sensitivity to feature magnitude.

Model Architecture. The MoE comprises E experts, each expert e being parameterized by a weight matrix $W^e \in \mathbb{R}^{m \times n}$ and a bias $b^e \in \mathbb{R}^m$. A learned router $g(x) = \text{softmax}(W^r x)$,



(a) Norm of each feature’s weight vector $\|W_i\|$, with colors indicating superposition status (green for features in superposition, purple for monosemantic features).

(b) $W^T W$ matrix where each cell represents $(\hat{W}_i \cdot W_j)$, revealing interference patterns between features.

Figure 1: Feature representation and superposition in a dense model with $n = 20$ features and $m = 6$ hidden dimensions, under exponentially decaying importance $I = 0.7^i$ and uniform feature density $(1 - S) = 0.1$. The superposition score encoded by color is $\sum_j (\hat{W}_i \cdot W_j)^2$, the projection of all other features onto the direction of feature i .

with $W^r \in \mathbb{R}^{E \times n}$, assigns each input to its top- k experts. Every active expert first projects the input onto a lower-dimensional hidden state $h^e = W^e x$ and then produces a reconstruction $\hat{x}^e = \text{ReLU}((W^e)^\top h^e + b^e)^1$. The network output is the gate-weighted sum over the active experts, $x' = \sum_{e \in \text{top-}k} w_e \hat{x}^e$, where the w_e are the renormalized gating weights.

We train each model to minimize an importance-weighted $L2$ reconstruction loss, $\mathcal{L} = \sum_x \sum_i I_i (x_i - x'_i)^2$, in which the per-feature importances I_i scale the corresponding error terms. In the experiments of Section 4, to prevent expert collapse we augment this objective with the standard auxiliary load-balancing loss (Fedus et al., 2022), $\mathcal{L}_{\text{aux}} = \alpha N \sum_{e=1}^N f_e P_e$, where N is the total number of experts, f_e is the fraction of batch samples routed to expert e , P_e is the mean gating probability assigned to expert e over the batch, and $\alpha = 0.01$ sets the penalty strength. We deliberately omit this auxiliary loss in Section 4 in order to isolate the MoE’s intrinsic architectural bias toward or against superposition.

4.2 MEASURING FEATURE CAPACITY

To analyze feature representations across architectures, we quantify two complementary properties of each feature: its *representation strength* and its *interference* with the remaining features. Representation strength is captured by the norm $\|W_i^e\|$ of feature i ’s weight vector within expert e , which measures the extent to which that expert represents the feature: $\|W_i^e\| \approx 1$ indicates that feature i is fully represented in expert e , whereas $\|W_i^e\| \approx 0$ indicates that it is not learned. Interference is visualized through the Gram matrix $W^T W$, whose off-diagonal elements report the pairwise interference between feature i and every other feature in expert e .

As Figures 1a and 2a show, at equal total parameters ($m_{\text{dense}} = \sum m_{\text{experts}} = 6$) the dense model and the MoE represent a comparable number of features (10 versus 8) with comparable norms. The decisive difference lies not in how many features are learned but in the geometry of their interference. Although individual MoE experts display some local superposition—Expert 0 in Figure 2b, for instance—the global interference structure is strictly partitioned. Whereas in the dense model (Figure 1b) any feature may interfere with any other, the MoE enforces a block-diagonal structure in which features routed to distinct

¹The output ReLU enforces non-negative reconstructions, consistent with our input distribution in which $x_i \in [0, 1]$ whenever nonzero. Moreover, the off-diagonal terms of $(W^e)^\top W^e$ produce negative interference, and the ReLU suppresses these negative components; when features are sparse, this negative interference becomes effectively “free” because it is clipped to zero, which incentivizes configurations with negative off-diagonal terms such as antipodal feature pairs.

experts have exactly zero interference. MoEs therefore allocate representational capacity by partitioning the feature space, sharply curtailing the global scope of interference.




Figure 2: Feature representation and superposition in an MoE with $n = 20$ features, 3 total experts, and $m = 2$ hidden dimensions per expert under top- $k = 1$ routing, with exponentially decaying importance $I = 0.7^i$ and feature density $1 - S = 0.1$. Matched to the dense model of Figure 1 on total parameters, the MoE confines interference to within-expert blocks rather than distributing it across the full feature set.

Expert Feature Dimensionality. To understand how an MoE allocates its limited representational capacity differently from a dense model, we measure feature dimensionality—the “fraction of a dimension” that a given feature occupies within a model (Elhage et al., 2022). For feature i , we define its dimensionality in expert e as

$$D_i^e = \frac{\|W_i^e\|^2}{\sum_j (\hat{W}_i^e \cdot W_j^e)^2} \quad (1)$$

The quantity D_i^e is bounded between zero (feature not learned) and one (feature represented monosemantically), so the total representational capacity of an MoE is naturally defined as $D = \sum_e \sum_{i=1}^n D_i^e$.

Efficient Packing. A model packs features “efficiently” when their dimensionalities sum to the number of embedding dimensions, $\sum_{i=1}^n D_i^e \approx m$ (Cohen et al., 2014; Scherlis et al., 2025; Elhage et al., 2022); for an MoE this generalizes to $\sum_e \sum_{i=1}^n D_i^e \approx E \cdot m$. We find empirically that both dense and MoE models satisfy this dimensionality constraint, establishing that MoEs pack features just as efficiently as dense models at equal total parameters. Efficiency of packing is therefore not what distinguishes the two architectures; the distinction lies instead in how that capacity is distributed, which we quantify next.




Figure 3: Features per dimension versus inverse feature density ($\frac{1}{1-g}$) for dense and MoE architectures under uniform feature importance ($I_i = 1.0$). The dense model ($n = 100$, $m = 20$) exhibits the most superposition, and superposition decreases monotonically as the expert count grows: 4 experts with $m = 5$, $k = 1$ (orange); 10 experts with $m = 2$, $k = 2$ (green); and 20 experts with $m = 1$, $k = 5$ (red). All models are matched on total parameters and on the ratio k/E , so the ordering isolates the effect of network sparsity. The dashed line at 1.0 marks the monosemantic limit, below which the model represents no more features than it has dimensions.

Features per Dimension & Network Sparsity. Because both dense and MoE models pack features efficiently, the number of features per dimension provides a direct and architecture-agnostic measure of superposition, allowing us to quantify precisely how the number of experts modulates superposition across feature sparsities. A features-per-dimension value above one signals superposition, since the model then represents more features than it has dimensions. We define features per dimension for an MoE as

$$\frac{1}{k} \sum_{e=1}^E p_e \frac{\|W^e\|_F^2}{m} \quad (2)$$

where $\|W\|_F^2$ is the squared Frobenius norm and p_e is the expected probability that expert e is used across a batch of inputs—equivalently, the average renormalized gating weight after top- k routing.

At equal total parameters, the dense model registers a consistently higher features-per-dimension value than every MoE (Figure 3), and hence more superposition. The dense model exploits superposition to represent more features than it has dimensions ($n_{\text{learned}} > m_{\text{total}}$), whereas the MoE caps its representation at the monosemantic limit ($n_{\text{learned}} \approx m_{\text{total}}$). This behavior is governed by network sparsity: as the total number of experts grows, with total parameters and the ratio k/E held fixed, features per dimension decreases monotonically. *The greater the number of experts, the less superposition in the model*—features become progressively more monosemantic as network sparsity increases. This interpretability gain is not purchased at a prohibitive cost in performance. The dense model’s additional superposition does buy it a consistently lower reconstruction loss than the MoEs, with a gap of ~ 0.03 – 0.08 at any given sparsity (Figure 6, Appendix A.1); yet as the expert count increases, the MoE loss converges toward that of the dense model, so greater monosemanticity is attained with a narrowing performance penalty. Appendix A.2 develops the theoretical intuition for this trade-off.

5 PHASE CHANGE

Although MoEs and dense models learn a similar number of features, MoEs distribute those features across experts with substantially less interference. This indicates that network

sparsity reshapes how features are allocated rather than how many are learned. We now examine how properties of the input distribution—feature sparsity and importance—drive this allocation, and whether they induce the physics-inspired *phase changes* that characterize dense representations.

A model has only a finite repertoire of ways to represent any feature: it may ignore the feature, place it in superposition, or represent it monosemantically. A phase change is the observation that discrete boundaries can separate these regimes as functions of feature sparsity and relative importance.

Dense toy models exhibit discontinuous ‘phase changes’ between internal feature representations (Elhage et al., 2022): varying the sparsity and relative importance of the input features elicits qualitatively distinct behavior, with greater feature sparsity, for instance, encouraging greater superposition. Analyzing the per-expert phase diagram of an MoE reveals that its experts adopt representational strategies fundamentally different from those of dense models.

Setup. We adopt the setup of Section 4.1, augmented with the auxiliary load-balancing loss to avert specialization collapse—the failure mode in which experts and routers settle into local minima that leave certain experts entirely unused. We evaluate three configurations, all with a single active expert ($k = 1$): (A) $n = 2$, $m = 1$; (B) $n = 3$, $m = 1$; and (C) $n = 3$, $m = 2$. For each, we report the per-expert phase diagram over feature sparsity and last-feature relative importance while sweeping network sparsity, increasing the number of experts E up to the number of input feature dimensions n .

Here we hold active parameters fixed rather than total parameters, so that $m_{dense} = km_{moe}$ (disregarding router parameters). This choice enables a controlled comparison across architectures: were active parameters left free, any observed differences could be ascribed to that budget rather than to the architecture itself. Note that Figure 4.C.1/1 happens to match Figure 4.B.X/2 in active parameters; the latter, however, must encode the same number of input features ($n = 3$) using only one hidden dimension ($m = 1$) rather than two ($m = 2$), which leaves little room to exploit superposition in the service of specialization.

Results and Takeaways. In every single-expert (dense) configuration we observe a clear phase change (Figure 4.X.1/1), corroborating Elhage et al. (2022). Increasing the total number of experts, however, dissolves these discrete phase changes. Some experts in the $E = 2$ MoEs echo their dense counterparts (Figure 4.X.2/2) but transition far more continuously, and in each case one expert becomes markedly more monosemantic, specializing in the feature of highest relative importance. Experts that diverge from the dense pattern uniformly exhibit much lower superposition scores—they are visibly bluer—signaling more monosemantic representations. This is precisely the behavior anticipated by the previous section: MoEs systematically favor lower superposition than their dense counterparts.

In the $n = 2, m = 1$ configuration (Figure 4.A), the dense model declines to represent the last feature under low feature sparsity, whereas the matched MoE preserves that feature far more readily because it retains the capacity to do so. With three input dimensions (Figure 4.B) this behavior vanishes: the experts already superimpose the other two features, leaving no room for the third within a single hidden dimension. Architecture B is thus the exception in which the superimposing hidden dimension is, in the high-sparsity regime, insufficient to represent all features—and even here no clear phase change emerges, save for the mostly discrete but mixed band at feature sparsity 0.5–0.7 in Figure 4.B.2/2. For $m = 2$, the white region of the dense model (Figure 4.C.1/1), occupying the mid- to low-sparsity domain where the last feature is relatively less important than the others, reflects that the feature is ignored. As network sparsity increases across the remaining panels of Figure 4.C, the models instead represent the last feature with progressively greater $L2$ magnitude ($\|W_3^1\| < \|W_3^2\| < \|W_3^3\|$). The dimensionality allocated in the low-importance region thus rises monotonically with network sparsity, exactly as Figure 3 predicts.

Across Figures 4.B.2/2, 4.C.1/2, and 4.C.2/2 we observe a window of feature sparsity, roughly 0.48 to 0.7, in which polysemantic, monosemantic, and ignored representations are heavily intermixed. This intermixing indicates a middle regime in which polysemantic and monosemantic representations incur comparable loss, so that the MoE cannot consis-




Figure 4: For a given expert and input dimension (feature), we decode how that feature is embedded in the hidden dimension—whether it is ignored (white), represented monosemantically (blue-purple), or placed in superposition (red). Each cell jointly encodes the feature norm ($\|W_n\|^2$) and the superposition score ($\sum_{j < n} (\hat{W}_n \cdot W_j)^2$) of the last feature, swept across feature sparsity $S \in [0.1, 1]$ and relative last-feature importance $I_n \in [0.1, 3]$, where the subscript n denotes the last of n total features. For every cell we train ten models and retain the one with the lowest loss; all runs in this section include the load-balancing loss. A low $L2$ norm ($\|W_n\|$) is rendered white to indicate that the model ignores the last feature; otherwise, a low superposition score is rendered blue-purple to indicate monosemantic representation, and red indicates that the feature is held in superposition. Cell (i, j) of subfigure X.e/E reports expert e of E total experts, trained on architecture X, for last-feature importance $I_n = i$ and sparsity $S = j$; X.1/1 denotes the corresponding dense model.

tently commit to the strategies it adopts cleanly in the low- and high-sparsity domains. No analogous regime appears in dense models—direct evidence that MoEs learn qualitatively different representational strategies.

Conclusion. All of these experiments use top- $k = 1$ routing, so exactly one expert is active at any time. Even under this most dense-like setting, we observe markedly different behavior as early as the $E = 2$ case—including when the superimposing hidden dimension has ample capacity to represent every feature. We therefore conclude that it is misleading to regard an MoE as a mere aggregation of dense models: the router, by restricting each expert to only a subset of the feature domain, fundamentally alters how the experts behave and what they learn.

6 EXPERT SPECIALIZATION

Having established that MoEs exhibit less superposition, we now turn to how these monosemantic features are organized within experts and what that organization reveals about specialization.

Conventional treatments of expert specialization in MoEs are framed around load balancing—the even distribution of inputs across experts (Chaudhari et al., 2025). This framing, however, fails to capture the more natural intuition of specialization, under which an expert is invoked only when the concepts it is specialized in are actually present in the input.

We therefore define an expert as specialized if it *occupies* certain feature directions in the input space and represents those features relatively monosemantically. We show that these two conditions are directly correlated, and that their joint presence itself promotes load

balancing across experts—recovering the conventional desideratum as a consequence, rather than a definition, of specialization.

Because we fix $k = 1$, the router partitions the feature space into convex cone regions (Appendix A.3), each region routed to a single expert. Formally, $\forall s > 0, \vec{x} \in C \rightarrow s\vec{x} \in C$, where C is the set of points contained in the cone and s is any positive scalar; that is, if a feature vector $\vec{x}$ is routed to an expert, so is every positive rescaling $s\vec{x}$. In this case we say that feature x is contained within expert e , and correspondingly that expert e *occupies* x .

We find empirically that small models which distribute the input space across more experts tend to achieve lower loss (Appendix A.4). This raises a natural question: does the allocation of the input space to particular experts entail any representational property of those experts? Our definition of specialization predicts that it entails *monosemanticity*—a correlation that, as we show below, persists in larger toy models (e.g., $m = 10$).

For models with $n > 2$, we investigate two complementary directions of this relationship. First, we ask whether initializing an expert to occupy specific features in the input space causes it to represent those features more monosemantically. Second, and conversely, we ask whether—for the features an expert has chosen to represent monosemantically—that expert in fact occupies those features in the input space.

Under diagonal initialization, each expert monosemantically represents exactly the single feature it initially occupied, and no other, as shown in Figure 5a. Under ordered k-hot initialization, the first expert monosemantically represents four of the five features it initially occupied (Figure 5b); the remaining experts, initialized over less important features, represent neither those less important features nor the five most important features they were not initialized over. When we break the ordering of feature importance and randomize the features each expert initially occupies, every expert monosemantically represents only the most important feature or features among those it was initialized over, as shown in Figure 5c.

The features initially routed to an expert are thus strongly correlated with the features that expert comes to represent monosemantically. Moreover, experts represent *only* important features monosemantically—whether we initialize an expert with a single important feature explicitly, or supply it a set of features from which it selects the most important on its own and represents that one monosemantically.

Under uniform feature importance, by contrast, experts place all features in superposition with elevated levels of interference. Even so, the features an expert initially occupies remain *relatively* more monosemantic, attaining superposition scores that are on average 1.03 standard deviations below the mean for that expert.

We next examined the converse correlation: whether an expert that represents certain features monosemantically also occupies those features in the input. To this end, we measure the expert’s usage both when its monosemantic features are *one of many* active features and when they are the *only* active features—the latter case being equivalent to measuring the probability that the expert occupies those features. As Table 1 shows, the correlation holds under both Xavier and k-hot initialization. With $E = 10$, a mean expert usage of $\sim 10\%$ corresponds to even load balancing across experts. In every case, activating an expert’s monosemantic feature or features raises that expert’s usage substantially, and when those features are the sole active features the expert dominates routing. Under k-hot initialization the correspondence is exact: *100%* of the features an expert represents monosemantically are occupied by that same expert.

An expert that represents more features monosemantically is, by our definition, more specialized: its usage on arbitrary inputs falls, yet conditional on its specialized features being active its usage rises far above that of other experts. This pattern holds in every case but one—the Xavier-initialized expert representing four monosemantic features, whose overall utilization drops sharply—an isolated exception against an otherwise consistent monotonic trend across both initialization schemes.




Figure 5: Expert feature norms $\|W_i^{(e)}\|$ and superposition (color) for three initialization schemes, with $n = 20, m = 5, E = 4, S = 0.1$. In (a), the gate matrix is initialized along the main diagonal ($W_i^r = \hat{e}_i$, the basis vector for dimension i), and relative feature importance decays exponentially in order from feature one to twenty. In (b), the gate matrix is initialized to an “ordered k-hot” pattern in which the first expert aligns with the first five features and each subsequent expert with the next five; relative feature importance is as in (a). In (c), the gate matrix is initialized to a “random k-hot” pattern in which each expert is assigned five random features, with no feature shared between experts and all twenty features covered collectively; relative feature importance again decays exponentially but is distributed randomly across features.

7 CONCLUSION

We investigated how experts affect superposition in MoEs, showing that MoEs consistently exhibit greater monosemanticity than dense networks while not exhibiting a phase change. We proposed a feature-based definition of expert specialization, demonstrating that experts naturally organize around coherent features when initialization encourages this specialization. However, our findings are based on simple autoencoder toy models with synthetic data, leaving open questions about generalization to large-scale transformers where the feature distribution is unknown (see Appendix A.6). Despite these limitations, we show how toy MoEs achieve comparable loss while maintaining more interpretable representations—challenging the prevalent zeitgeist that mechanistic interpretability and model capability are fundamentally in tension. Future work should explore what favors monosemanticity in

Table 1: Per-expert monosemantic-feature and usage statistics for $n = 100, m = 10, E = 10$. For each initialization scheme (Xavier and k-hot), one hundred models are trained, yielding 1000 experts per scheme. Statistics are aggregated across models, with experts grouped by the number of features they represent monosemantically. For the feature(s) an expert represents monosemantically, we report the expert’s usage when those features are one of several active features in the input and, separately, its usage when those features are the *only* active feature—the latter measuring how strongly the expert occupies those feature directions. Usage rises sharply and monotonically with the number of monosemantic features conditioned on those features being active, and reaches near-total dominance when they alone are active.

Xavier Initialization				
Number of monosemantic features per expert	Number of experts (out of 1000)	Mean expert usage (%)	Mean expert usage; feature(s) active (%)	Mean expert usage; only feature(s) active (%)
0	461	—	—	—
1	387	9.595	17.94	67.18
2	138	9.599	30.29	95.65
3	13	8.363	40.19	100.0
4	1	1.428	14.69	100.0
5	0	—	—	—
K-Hot Initialization				
Number of monosemantic features per expert	Number of experts (out of 1000)	Mean expert usage (%)	Mean expert usage feature(s) active (%)	Mean expert usage only feature(s) active (%)
0	335	—	—	—
1	382	10.00	23.94	100.0
2	227	10.02	46.61	100.0
3	47	10.09	62.00	100.0
4	8	9.95	70.30	100.0
5	1	9.62	74.79	100.0

MoEs, how training dynamics of MoEs differ from those of the dense model, and when specialization emerges. Answering these questions can inform the design of more interpretable, high-performing language models.

REFERENCES

- Trenton Bricken, Adly Templeton, Joshua Batson, Brian Chen, Adam Jermyn, Tom Conerly, Nick Turner, Cem Anil, Carson Denison, Amanda Askell, Robert Lasenby, Yifan Wu, Shauna Kravec, Nicholas Schiefer, Tim Maxwell, Nicholas Joseph, Zac Hatfield-Dodds, Alex Tamkin, Karina Nguyen, Brayden McLean, Josiah E Burke, Tristan Hume, Shan Carter, Tom Henighan, and Christopher Olah. Towards monosemanticity: Decomposing language models with dictionary learning. *Transformer Circuits Thread*, 2023. <https://transformer-circuits.pub/2023/monosemantic-features/index.html>.
- Marmik Chaudhari, Idhant Gulati, Nishkal Hundia, Pranav Karra, and Shivam Raval. Moe lens - an expert is all you need. In *Sparsity in LLMs (SLLM): Deep Dive into Mixture of Experts, Quantization, Hardware, and Inference*, 2025. URL <https://openreview.net/forum?id=GS4WXncwSF>.
- Michael B. Cohen, Yin Tat Lee, Cameron Musco, Christopher Musco, Richard Peng, and Aaron Sidford. Uniform sampling for matrix approximation. *CoRR*, abs/1408.5099, 2014. URL <http://arxiv.org/abs/1408.5099>.
- Nishanth Dikkala, Nikhil Ghosh, Raghu Meka, Rina Panigrahy, Nikhil Vyas, and Xin Wang. On the benefits of learning to route in mixture-of-experts models. In *Proceedings of the*

594 *2023 Conference on Empirical Methods in Natural Language Processing*, pp. 9376–9396,
595 Singapore, December 2023. Association for Computational Linguistics. doi: 10.18653/v1/
596 2023.emnlp-main.583. URL <https://aclanthology.org/2023.emnlp-main.583>.
597

598 Nelson Elhage, Tristan Hume, Catherine Olsson, Nicholas Schiefer, Tom Henighan, Shauna
599 Kravec, Zac Hatfield-Dodds, Robert Lasenby, Dawn Drain, Carol Chen, Roger Grosse,
600 Sam McCandlish, Jared Kaplan, Dario Amodei, Martin Wattenberg, and Christopher
601 Olah. Toy models of superposition. *Transformer Circuits Pub*, 2022. URL https://transformer-circuits.pub/2022/toy_model/index.html.
602

603 William Fedus, Barret Zoph, and Noam Shazeer. Switch transformers: Scaling to trillion
604 parameter models with simple and efficient sparsity, 2022. URL <https://arxiv.org/abs/2101.03961>.
605

606 Google DeepMind. Gemini 2.5 pro. <https://deepmind.google/technologies/gemini/pro/>, 2025.
607

608 Wes Gurnee, Neel Nanda, Matthew Pauly, Katherine Harvey, Dmitrii Troitskii, and Dimitris
609 Bertsimas. Finding neurons in a haystack: Case studies with sparse probing. *arXiv preprint arXiv:2305.01610*, 2023.
610

611 Albert Q Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary,
612 Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Emma Bou Hanna, Florian
613 Bressand, et al. Mixtral of experts. *arXiv preprint arXiv:2401.04088*, 2024.
614

615 Jakub Krajewski, Jan Ludziejewski, Kamil Adamczewski, Maciej Pióro, Michał Krutul,
616 Szymon Antoniak, Kamil Ciebiera, Krystian Król, Tomasz Odrzygóźdź, Piotr Sankowski,
617 Marek Cygan, and Sebastian Jaszczur. Scaling laws for fine-grained mixture of experts.
618 *arXiv preprint arXiv:2402.07871*, 2024. URL <https://arxiv.org/abs/2402.07871>.
619

620 Dmitry Lepikhin, HyoukJoong Lee, Yuanzhong Xu, Dehao Chen, Orhan Firat, Yanping
621 Huang, Maxim Krikun, Noam Shazeer, Zhifeng Chen, Yonghui Wu, et al. Gshard: Scaling
622 giant models with conditional computation and automatic sharding. In *Proceedings of the*
623 *38th International Conference on Machine Learning*. PMLR, 2021.
624

625 Siyuan Mu and Sen Lin. A comprehensive survey of mixture-of-experts: Algorithms, theory,
626 and applications, 03 2025.
627

628 Jungwoo Park, Ahn Young Jin, Kee-Eung Kim, and Jaewoo Kang. Monet: Mixture of
629 monosemantic experts for transformers. In *The Thirteenth International Conference on*
630 *Learning Representations*, 2025. URL <https://openreview.net/forum?id=10gw1SHY3p>.
631

632 Kiho Park, Yo Joong Choe, and Victor Veitch. The linear representation hypothesis and the
633 geometry of large language models, 2024. URL <https://arxiv.org/abs/2311.03658>.
634

635 Lucas Prieto, Edward Stevinson, Melih Barsbey, Tolga Birdal, and Pedro AM Mediano.
636 Correlations in the data lead to semantically rich feature geometry under superposition.
637 In *Mechanistic Interpretability Workshop at NeurIPS 2025*, 2025.
638

639 Adam Scherlis, Kshitij Sachan, Adam S. Jermyn, Joe Benton, and Buck Shlegeris. Polyse-
640 manticity and capacity in neural networks, 2025. URL <https://arxiv.org/abs/2210.01892>.
641

642 Noam Shazeer, Azalia Mirhoseini, Krzysztof Maziarz, Andy Davis, Quoc Le, Geoffrey Hinton,
643 and Jeff Dean. Outrageously large neural networks: The sparsely-gated mixture-of-
644 experts layer, 2017. URL <https://arxiv.org/abs/1701.06538>.
645

646 An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu,
647 Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025a.

-
- Xingyi Yang, Constantin Venhoff, Ashkan Khakzar, Christian Schroeder de Witt, Puneet K. Dokania, Adel Bibi, and Philip Torr. Mixture of experts made intrinsically interpretable. In *Forty-second International Conference on Machine Learning*, 2025b. URL <https://openreview.net/forum?id=6QERrXMLP2>.
- Yanqi Zhou, Tao Lei, Hanxiao Liu, Nan Du, Yanping Huang, Vincent Y Zhao, Andrew M. Dai, Zhifeng Chen, Quoc V Le, and James Laudon. Mixture-of-experts with expert choice routing. In Alice H. Oh, Alekh Agarwal, Danielle Belgrave, and Kyunghyun Cho (eds.), *Advances in Neural Information Processing Systems*, 2022. URL <https://openreview.net/forum?id=jdJo1HIVinI>.
- Barret Zoph, Irwan Bello, Sameer Kumar, Nan Du, Yanping Huang, Jeff Dean, Noam M. Shazeer, and William Fedus. St-moe: Designing stable and transferable sparse expert models, 2022. URL <https://api.semanticscholar.org/CorpusID:248496391>.

A APPENDIX

A.1 MEASURING LOSS FOR VARYING SPARSITY & EXPERTS




Figure 6: Log average loss versus feature density ($\frac{1}{1-s}$) for dense ($m = 20$) and MoE (4 experts, $k = 1$, $m = 5$), MoE (10 experts, $k = 2$, $m = 2$), and MoE (20 experts, $k = 5$, $m = 1$) models, all with uniform feature importance ($I_i = 1.0$) for $n = 100$ input features. Results are averaged over five runs per sparsity level. Although dense model outperforms all MoEs at every sparsity level, as the number of experts increases, the MoE loss gets closer to the dense model.

A.2 THEORETICAL INTUITION FOR SUPERPOSITION

In dense models, superposition emerges to exploit the gap between sparse features and dense computation, compressing rarely active features into shared dimensions. MoEs, however, “eat this same gap” structurally through conditional computation. By aligning activation sparsity with feature sparsity, MoEs remove the computational penalty for having dedicated, rarely-activating neurons. As noted by Elhage et al. (2022), when a model only expends computation on active features, splitting polysemantic neurons into dedicated monosemantic ones becomes the optimal strategy, effectively trading the *compression* of superposition for the *selection* of routing.

The structural change manifests geometrically as a reduction in *interference*. While the global ratio of number of features represented to parameters remains constant, the router effectively partitions the feature space, ensuring that a feature routed to expert e only competes for capacity with the subset of features also assigned to that expert. Consequently, interference is governed by the local expert matrix $(W^e)^\top W^e$ rather than the global $W^\top W$ of a dense model. This partitioning drastically reduces the number of interfering features for any given feature vector, minimizing the optimization pressure to pack features in superposition and allowing experts to learn monosemantic representations.

A.3 ROUTER SUBSPACES ARE CONVEX CONES

In the regime of $k = 1$, the router function $g(x) = \text{softmax}(W^r x)$ is equivalent to $\text{argmax}(W^r x)$. The region routed to expert i can be represented as $\forall j \neq i, (w_i - w_j)^\top x > 0$ where w_i and w_j are row vectors of W^r . This is a homogeneous linear inequality. Regions bounded by such inequalities are by definition convex cones. If a particular x satisfies this inequality, then multiplying both sides by any positive scalar s will still satisfy the inequality.

ity. Furthermore, if x_1 and x_2 satisfy this inequality, then any $x = x_1\lambda + (1 - \lambda)x_2$ for $\lambda \in [0, 1]$ will also satisfy the inequality.

In the case of $k > 1$, the region of inputs which get sent to a particular expert e becomes a union of convex cones. Generally, the union of a convex cone is not itself a convex cone. Therefore, the understanding of experts occupying feature directions may not hold beyond $k = 1$.

A.4 EXPERT ROUTING WITH DIFFERENT INITIALIZATION SCHEMES




Figure 7: Expert routing of three identical models with differing initialization schemes. We use $n = 2$, $E = 3$, $m = 1$. The first model (left) has the worst performance (loss: 0.08) and routes all inputs to one expert. The second model (middle) has better performance (loss: 0.04) and routes a small portion of inputs, specifically those when feature 1 is active, to a second expert. The third model (right) has the lowest loss (loss: 0.025), and distributes the input space among all experts. One expert is chosen when only feature 1 is active, one when only feature 2 is active, and one when both are active.

In small models ($n = 2$, $m = 1$, $E > 1$), we empirically find that models that distribute the input space across more experts tend to achieve lower loss, testing with $E \in [2, 7]$. Holding $n = 2$ allows us to visualize which portions of the input space get routed to which experts, as seen in Figure 7.

A.5 ANALYTIC MODEL EQUIVARIANCE

For the toy setup of single-layer, single-nonlinearity, top- $k = 1$ MoEs, there exists a theoretical map between any dense model and a monosemantic MoE with an equivalent number of active features under a sparsity constraint.

Assume there exists an upper bound for the number of active features a for any input such that $\forall x \in D : |\{i : x_i \neq 0\}| \leq a$. Furthermore, assume that a is no greater than the hidden dimensionality, m , of an expert, providing an upper bound on the number of features a model has to represent. Assume also that the hidden dimensions is smaller than the total number of input features n ($a \leq m \leq n$). To construct the monosemantic MoE, for each possible subset $S \subseteq \{1, 2, \dots, n\}$ with $|S| \leq a$ —meaning the size of the subset of active features is smaller than or equal to a —create an expert which monosemantically preserves those features. (In fact, you can take only the subsets such that $|S| = a$.) The router then selects the expert which corresponds to those active features (of which there will never be more than a , by assumption):

$$\text{Router}(x) = \arg \max_S \mathbb{I}[\text{support}(x) = S]$$

where $\text{support}(x) = \{i : x_i \neq 0\}$. Since $|S| \leq a \leq m$, each expert has sufficient capacity to represent its assigned features without superposition. To reiterate, only a features are active and every unique combination of active features receives its own dedicated expert with sufficient capacity to represent those features monosemantically. So, the number of possible experts needed is $\binom{n}{m}$.




Figure 8: Model X/E uses X to denote the same model architectures and models used in Figure 4 and E denotes the total number of experts (i.e. network sparsity). Increasing network sparsity decreases mean loss while increasing localized variance—especially as the number of experts reaches the input feature dimensions. This can be attributed to the relatively unstable training of MoEs compared to dense models (despite training ten models for each cell and selecting the lowest loss).

The reconstruction for this theoretical MoE has zero loss only as the sparsity constraint holds (or goes to one in these toy models) because there is the chance more than m features could be active at one time ($a \not\leq m$), which would exceed the monosemantic representational capacity of the network (but the dense polysemantic could do no better unless features are correlated in the distribution). Therefore, even if $a \leq m$ sometimes, the polysemantic model encounters the same problem and the monosemantic MoE under this construction may still outperform it under looser sparsity constraints.

Thus, for any dense model, $f_{\text{dense}}(x) = \text{ReLU}(Wx + b)$ under the sparsity constraint $|\text{support}(x)| \leq a$, there exists a MoE model $f_{\text{MoE}}(x)$ such that $f_{\text{dense}}(x) = f_{\text{MoE}}(x)$ for all valid inputs. In the toy settings described in this paper, the sparsity constraint holds in the limit where sparsity goes to one. However, in practice there may be an upper bound on the amount of features a particular amount of information can semantically encode, indicated by the size of meaningful embeddings of that data. Therefore, a MoE model with sufficient experts and a tractable amount of superposition (e.g. interpretable) may be sufficient to encode all features present.

A.6 LIMITATIONS

Our results should be interpreted in light of several limitations. First, all experiments are conducted in controlled toy-model settings derived from superposition studies in sparse autoencoders (Elhage et al., 2022). While this enables precise measurement of interference and monosemanticity, it abstracts away many complexities of large-scale transformers, including multiple computation, attention, heterogeneous feature distributions, and realistic routing dynamics. Consequently, the transferability of our findings to real-world MoE architectures remains uncertain.

Second, our architectural and routing choices are intentionally restricted: we use simple experts, fixed top- k routing (often $k = 1$), equal parameter budgets between dense and MoE models, and no auxiliary load-balancing losses. Practical MoEs often employ richer routing mechanisms, variable expert capacities, and multi-task objectives (Lepikhin et al., 2021; Fedus et al., 2022) or expanding beyond the task of reconstruction to next token prediction, for example, which may yield qualitatively different representational behaviors.

Third, the feature distributions in our synthetic tasks—including sparsity patterns, feature importance, and independence assumptions—are significantly simpler than those found in natural data.

Finally, although we observe increased monosemanticity and reduced superposition in MoEs under fixed conditions, we do not evaluate downstream performance trade-offs, training stability, or specialization dynamics at scale. Prior work suggests such dynamics can shift substantially with model size, optimization regime, and data diversity (Krajewski et al., 2024).

Overall, our study provides mechanistic insights under clean experimental conditions, but further work is required to validate these patterns in large, realistic MoE systems. Unfortunately, our methods do not scale naturally to larger models and this is a clear direction for future work in this space.