

UNDERSTANDING THE FINE-GRAINED KNOWLEDGE CAPABILITIES OF VISION-LANGUAGE MODELS

Anonymous authors
Anonymous submission

ABSTRACT

Vision-language models (VLMs) have advanced rapidly across the full spectrum of visual question answering—visual reasoning, document understanding, and multimodal dialogue alike—and these gains hold consistently across VLMs built on diverse base language models, alignment architectures, and training corpora. This very progress, however, masks a systematic blind spot: on traditional image classification benchmarks, precisely the setting that probes the fine-grained visual knowledge underpinning vision-centric competence, such models lag markedly. We resolve this disconnect through a comprehensive, unified evaluation of a large and representative set of recent VLMs on fine-grained classification benchmarks, and we isolate the factors that separate fine-grained knowledge from performance on other vision tasks. A systematic series of controlled ablation experiments establishes a clear decomposition: a stronger base LLM raises every benchmark score in equal measure, whereas a stronger vision encoder disproportionately improves fine-grained classification. We further demonstrate that the pretraining stage is the decisive determinant of fine-grained performance, and markedly more so when the language model weights are unfrozen during pretraining. Taken together, these findings chart a concrete and actionable path toward VLMs with substantially stronger fine-grained visual understanding and, more broadly, more reliable vision-centric capability.

1 INTRODUCTION

Vision-language models (VLMs), which couple a vision encoder with a large language model (LLM), have rapidly acquired increasingly sophisticated capabilities for interpreting and reasoning about visual content (Alayrac et al., 2022; Team et al., 2023; OpenAI, 2023; Liu et al., 2023; Wang et al., 2024b), and now discharge demanding tasks that span open-ended visual question answering, document understanding, and multimodal dialogue (Goyal et al., 2017; Yue et al., 2024; Mathew et al., 2021; OpenAI, 2023). This trajectory of progress, however, leaves a foundational question both unresolved and largely unexamined: how well do these models perform fine-grained visual perception, and—more consequentially—what concretely can be done to improve it? This paper confronts that question head-on, treating fine-grained recognition not as a peripheral curiosity but as a core axis of visual intelligence and a prerequisite for the vision-centric competence that downstream reasoning presupposes.

Fine-grained visual recognition—the ability to discriminate among visually similar categories—requires a model to resolve subtle discriminative cues while remaining invariant to intra-class variation in pose, illumination, and background. Traditional vision encoders such as CLIP (Radford et al., 2021) and DINO (Caron et al., 2021) attain strong fine-grained classification performance; yet the fine-grained behavior of the newer generation of VLMs—models constructed atop these very encoders—has, strikingly, gone comparatively unexamined.

Characterizing the fine-grained knowledge capabilities of VLMs is essential precisely because perception constitutes the foundation on which higher-order understanding and reasoning are erected: a model cannot reason reliably about what it cannot first recognize. Real-world deployments of VLMs routinely confront fine-grained, long-tailed distributions in which discriminating among closely related categories is decisive rather than incidental. Should a model misidentify a mushroom species, for instance, no quantity of downstream textual knowledge can recover whether the specimen is poi-




Figure 1: **Overview.** (Top) We investigate the fine-grained classification capabilities of vision-language models (VLMs)—a foundational yet routinely overlooked dimension of visual competence on which higher-level understanding and reasoning depend. (Bottom) Across 22 controlled ablation experiments that systematically vary the principal model components and training strategies (left), we isolate the factors that govern fine-grained classification performance in VLMs (right).

sonous or safe to consume. Beyond such immediate stakes, a principled account of the factors that govern fine-grained knowledge acquisition translates directly into design guidance for the architectures and training strategies of the next generation of VLMs—the practical payoff we pursue in this work.

We first present a comprehensive evaluation of fine-grained visual knowledge in state-of-the-art VLMs, directly juxtaposing their fine-grained classification performance against their capabilities on general VLM tasks. Concretely, we evaluate 15 representative VLMs—spanning LLaVA (Liu et al., 2023; 2024a), Phi (Abdin et al., 2024), Qwen2-VL (Wang et al., 2024b), and Molmo (Deitke et al., 2024)—on four fine-grained classification benchmarks that jointly span the object, flower, pet, and food domains: ImageNet (Deng et al., 2009), Flowers (Nilsback & Zisserman, 2008), Pets (Parkhi et al., 2012), and Food (Bossard et al., 2014), and set these against aggregated results from eight general VLM benchmarks including MMMU (Yue et al., 2024), MathVista (Lu et al., 2024), and MMVet (Yu et al., 2024). This analysis establishes two facts. First, models of comparable general ability diverge substantially in fine-grained classification accuracy, which demonstrates that fine-grained recognition constitutes a distinct axis of visual intelligence that current VLM benchmarks fail to capture. Second, it exposes a persistent, systematic gap between VLMs and the CLIP-based vision encoders on which they are built—a gap that corroborates and materially extends prior observations (Zhang et al., 2024; Geigle et al., 2024) to a far broader and more representative model set.

To pinpoint what actually drives fine-grained classification performance in VLMs, we then conduct a systematic study that disentangles the contributions of the principal model components and training paradigms. Concretely, we quantify how the two core architectural components—the language model and the vision encoder—together with the training strategy—the pretraining stage, the choice of which weights are updated, and the quality of the training data—each shape fine-grained classification. Through 22 controlled ablation experiments, holding the surrounding pipeline fixed so that each experiment isolates a single factor, we identify the levers that decisively govern fine-grained knowledge acquisition in VLMs:

- **Language model (§4.2):** Upgrading the base language model lifts performance uniformly, delivering comparable gains on fine-grained classification and general VLM benchmarks alike.

- **Vision encoder (§4.2):** A stronger vision encoder selectively improves fine-grained classification while exerting only limited influence on general VLM benchmarks—an asymmetry that becomes most pronounced under the standard two-stage training regime.
- **Pretraining stage (§4.3):** Large-scale pretraining on image-captioning data is the single most effective lever for fine-grained classification, conferring substantial gains while leaving general VLM benchmarks comparatively unaffected.
- **Weight updating method (§4.3):** Jointly updating the vision-encoder-to-LLM connector and the LLM during pretraining improves fine-grained classification well beyond what connector-only pretraining attains, establishing which weights are updated—not merely whether pretraining occurs—as a decisive design choice.
- **Data quality (§4.3):** The quality of the pretraining and instruction-tuning data exerts only a limited effect on fine-grained classification—an effect that all but vanishes when the language model weights are held frozen during pretraining, precisely locating the condition under which caption quality does and does not matter.

In sum, this study places fine-grained visual perception in VLMs—an essential capability that current benchmark evaluations systematically fail to surface—at the center of analysis, and, through controlled ablations, identifies the factors that decisively govern fine-grained knowledge acquisition. These findings together supply actionable guidance for building vision-centric VLMs with markedly stronger visual understanding, and thereby for improving their reliability across the real-world, long-tailed tasks that hinge on fine-grained recognition.

2 RELATED WORK

Vision-language models and their evaluation. Vision-language models (VLMs) integrate visual inputs with a language model and predominantly adopt the LLaVA-style architecture (Liu et al., 2023), which comprises three core components: a vision encoder, a language model, and an MLP connector that bridges the two. Such models attain strong performance across a broad range of multi-modal tasks, encompassing visual question answering (VQA), document understanding, and general reasoning (Liu et al., 2023; Dai et al., 2023; Awadalla et al., 2023; Wang et al., 2024b). Their evaluation rests, correspondingly, on a large family of benchmarks such as MMMU (Yue et al., 2024), MathVista (Wang et al., 2024a), and DocVQA (Mathew et al., 2021). These benchmarks, however, predominantly probe reasoning and language understanding conditioned on visual input, and in doing so leave the core vision-centric competencies—object recognition and fine-grained classification foremost among them—essentially unmeasured. We close this gap by systematically evaluating and analyzing VLMs on fine-grained classification, an essential yet underexplored dimension of visual intelligence.

Fine-grained visual classification. Fine-grained classification is a well-established computer vision task that seeks to discriminate among visually similar subcategories within a broader object category, such as bird species (Wah et al., 2011), flowers (Nilsback & Zisserman, 2008), and pet breeds (Parkhi et al., 2012). The corresponding benchmarks have long served as central testbeds for training and evaluating vision encoders, ResNet (He et al., 2016), CLIP (Radford et al., 2021), and DINO (Caron et al., 2021) among them. Despite this foundational role in vision research, recent VLMs have largely omitted these datasets from their evaluation protocols in favor of visual question answering. Yet because fine-grained classification is indispensable to real-world applications that involve long-tail distributions and demand precise recognition—medical diagnosis, food safety, and species identification among them—evaluating VLMs on such benchmarks is essential for assessing their practical usability from a vision-centric standpoint, rather than on language-centric reasoning tasks alone.

Evaluation and ablation of VLMs for fine-grained classification. Prior studies of VLM evaluation on fine-grained classification (Zhang et al., 2024; Geigle et al., 2024) remain limited in scope, typically confined to a handful of models or datasets. We substantially broaden this line of work, evaluating a considerably wider range of VLMs across multiple fine-grained classification benchmarks within a single, unified multiple-choice protocol. Beyond evaluation, we contribute a comprehensive ablation analysis that isolates the key factors driving fine-grained classification performance—an analysis without precedent in this setting.

A parallel body of work examines the design space of VLMs more generally, dissecting model architectures (McKinzie et al., 2024), training strategies (Karamcheti et al., 2024), and the role of multimodal datasets (Gadre et al., 2023; Udandarao et al., 2024; Deitke et al., 2024). These studies, however, are oriented almost exclusively toward general benchmarks that emphasize language and reasoning. We instead trace how the central design choices—base model selection, pretraining methodology, and the composition of the pretraining and instruction-tuning data—shape fine-grained classification in particular, yielding novel and directly actionable insights for optimizing VLMs on vision-centric tasks.

3 EVALUATION: FINE-GRAINED CLASSIFICATION

This section motivates the evaluation of fine-grained classification, specifies our benchmarks and evaluation protocol, and reports the central observations that emerge from testing 15 existing VLMs under a single unified protocol.

3.1 MOTIVATION

Fine-grained classification is a foundational computer vision task that targets the discrimination of visually similar subcategories within a broader category—distinct bird species or car models, for instance (Wei et al., 2021). A substantial suite of evaluation benchmarks, ImageNet (Deng et al., 2009) and Flowers (Nilsback & Zisserman, 2008) chief among them, has become the standard testbed for training and evaluating vision encoders. With the advent of large-scale self-supervised and supervised learning, vision encoders such as CLIP (Radford et al., 2021) and DINO (Caron et al., 2021) attain near- or super-human accuracy even on the most demanding of these fine-grained tasks.

These benchmarks, however, have been conspicuously absent from the evaluation of vision-language models (VLMs). Although recent VLMs such as LLaVA (Liu et al., 2023) demonstrate strong capabilities in interpreting visual concepts, their evaluation remains largely confined to visual question answering (VQA)—academic VQA benchmarks such as MMMU (Yue et al., 2024) and MathVista (Lu et al., 2024), or document understanding tasks such as DocVQA (Mathew et al., 2021) and TextVQA (Singh et al., 2019). Because these benchmarks foreground reasoning and language processing yet contain no fine-grained object categories, they leave the core vision-centric capabilities of VLMs essentially unexamined.

Fine-grained classification is, nonetheless, decisive for VLMs, precisely because real-world applications routinely encounter long-tailed distributions that demand precise visual discrimination among closely related categories. Consider a blind user asking whether a mushroom is safe to eat: the model must correctly identify the species before it can reason about toxicity at all, and misclassifying a deadly destroying angel as an edible white button mushroom could prove fatal. The same logic governs medical diagnosis, where a failure to distinguish visually similar conditions can precipitate incorrect treatment and serious harm, and autonomous driving, where confusing a “stop” sign with a “do not enter” sign—both red and similarly shaped—can produce dangerous navigation errors. Excelling at fine-grained classification is therefore not a peripheral desideratum but a prerequisite for the safety and reliability of VLMs in real-world deployment.

3.2 BENCHMARKS

We probe the fine-grained visual knowledge of VLMs with four well-established object recognition benchmarks, each targeting a distinct domain of fine-grained classification and jointly spanning generic objects, flowers, pets, and food: **ImageNet-1K** (Deng et al., 2009), **Oxford Flowers-102** (Nilsback & Zisserman, 2008), **Oxford-IIIT Pet-37** (Parkhi et al., 2012), and **Food-101** (Bossard et al., 2014).

Because these benchmarks span 37 to 1,000 distinct classes, evaluating VLMs directly in the multiple-choice format standard for VLM training and evaluation is not straightforward. We resolve this by adopting the methodology of Geigle et al. (2024) to recast each dataset as a set of 5-way multiple-choice questions. Concretely, for every image we employ an OpenCLIP ViT-L/14 model (Schuhmann et al., 2022) to retrieve the four hardest negative choices by image–label cosine




Figure 2: **Fine-grained classification versus general VQA performance across VLMs.** Across recent VLMs, fine-grained classification accuracy varies widely at essentially fixed general VQA performance, establishing fine-grained classification as a distinct axis of visual capability that standard VQA benchmarks do not measure.




Figure 3: **VLMs versus their own CLIP vision encoders on fine-grained classification.** Although Qwen2-VL-Chat nearly recovers the fine-grained accuracy of its DFN-CLIP vision encoder, every other VLM trails the encoder it is built upon by a wide margin, quantifying the substantial fine-grained capability that current VLMs leave unrealized and thereby the headroom our analysis targets.

similarity, and combine them with the ground-truth label to form the answer set. The choice of the OpenCLIP model is deliberate: it prevents question difficulty from being biased against VLMs that rely on OpenAI-CLIP (Radford et al., 2021), SigLIP (Zhai et al., 2023), or DFN-CLIP (Fang et al., 2023), each of which is trained on a corpus distinct from OpenCLIP’s. This construction lets us evaluate VLMs in their native format while preserving the bulk of the task’s intrinsic difficulty: CLIP ViT-L/14-336px (Radford et al., 2021), for instance, attains 75.2% accuracy on the full 102-class Flowers-102 and rises to only 78.3% on our 5-way multiple-choice version, confirming that the reformulation retains the greater part of the discriminative challenge.

3.3 TESTING SETUP

We evaluate 15 recent VLMs spanning the 7B–13B parameter range on the converted 5-choice fine-grained classification benchmarks (Liu et al., 2023; 2024a; Chen et al., 2024a; Abdin et al., 2024; Deitke et al., 2024; Laurençon et al., 2024; Hong et al., 2024; Chen et al., 2024c; Wang et al., 2024b). For each VLM we adopt its respective default prompt format from VLMEvalKit (Duan et al., 2024) and score accuracy by exact match against the multiple-choice options. As a baseline, we evaluate a range of CLIP models (Radford et al., 2021; Schuhmann et al., 2022; Fang et al., 2023) under the standard zero-shot protocol, in which the image and each candidate class label (“A photo of a [class]”) are both mapped through CLIP’s image and text encoders and the top-1 prediction is taken as the label of maximal image–text cosine similarity.

To determine whether existing VLM benchmarks already capture fine-grained classification ability, we compare each VLM’s fine-grained accuracy against its average score across eight general VQA benchmarks from VLMEvalKit, MMMU (Yue et al., 2024) and MathVista (Lu et al., 2024) among them.

3.4 RESULTS

Figure 2 plots fine-grained classification performance (averaged across our four benchmarks) against general VQA performance (averaged across eight existing VQA benchmarks) for every tested VLM. The two quantities decouple sharply: models with essentially identical general VQA performance span a wide range of fine-grained classification accuracy. CogVLM-Chat (Hong et al., 2024) and LLaVA-NeXT-Vicuna-13B (Liu et al., 2024a), for instance, both reach nearly 48% average accu-

racy on general VQA, yet their fine-grained classification accuracy differs by a full 19 percentage points (pp)—77.3% for CogVLM-Chat against 58.1% for LLaVA-NeXT-Vicuna-13B. A discrepancy of this magnitude between models matched on general VQA demonstrates conclusively that fine-grained classification ability is not captured by existing VLM benchmarks, and establishes it as a distinct dimension of visual intelligence.

Finding 1: Fine-grained classification represents a distinct aspect of visual capability that current VLM benchmarks fail to adequately measure.

Every current VLM is constructed on top of a pretrained vision encoder, and our multiple-choice reformulation turns this fact into a direct, like-for-like comparison between each VLM and the CLIP model it inherits on traditional classification benchmarks. Figure 3 contrasts the fine-grained classification accuracy of selected VLMs against that of their corresponding CLIP baselines, exposing a substantial and consistent shortfall. Molmo (Deitke et al., 2024), a very recent model that excels across VQA benchmarks, attains only 68.4% average accuracy in our multiple-choice setting, whereas its own CLIP ViT-L/14 (Radford et al., 2021) vision encoder reaches 86.7%—a gap of 18.3pp. Even the strongest model we test, Qwen2-VL-Chat (Wang et al., 2024b), trails its own state-of-the-art DFN-CLIP (Fang et al., 2023) vision encoder by 4.6 points. That fine-grained knowledge already resident in the vision encoder is not fully preserved through multimodal integration is the crux of the problem, and it motivates a systematic analysis of VLM design choices aimed at identifying the factors that recover and enhance fine-grained classification. We undertake precisely this ablation study in the next section.

Finding 2: A significant performance gap exists between VLMs and CLIP models in fine-grained classification tasks.

4 ABLATIONS: BASE MODELS AND TRAINING

4.1 EXPERIMENTAL SETUP

Building on the observational results of Section 3, we now ablate the model-level differences that plausibly drive fine-grained classification performance, concentrating on the two dimensions most likely to matter: base model selection and training methodology. All experiments are conducted within the LLaVA-1.5 training framework and alignment architecture (Liu et al., 2023), holding the surrounding pipeline fixed so that each ablation isolates a single factor and its effect can be read off unambiguously. Every model undergoes one epoch of pretraining (where applicable) followed by one epoch of finetuning, and we retain LLaVA-1.5’s default hyperparameters—learning rate and batch size included—throughout, so that no confound other than the ablated variable enters the comparison.

We report fine-grained performance as the average accuracy across the four classification benchmarks introduced in Section 3.2. For general VQA performance, we report the average over three widely used multiple-choice benchmarks—MMU (Yue et al., 2024), MMBench (Liu et al., 2024b), and MMStar (Chen et al., 2024b)—selected for their broad domain coverage and established role in VLM evaluation.

4.2 BASE MODELS

The analysis in Section 3 reveals that recent models have steadily adopted newer base LLMs while, in most cases, retaining the CLIP ViT-L/14 vision encoder of Radford et al. (2021). We therefore isolate how each of these two base-model choices—the language model and the vision encoder—individually shapes fine-grained classification performance.

4.2.1 LLM CHOICE

Setup. Taking LLaVA-1.5-7B (Liu et al., 2023) as our baseline, we vary the LLM component while sweeping across vision encoders and pretraining configurations, so that the LLM effect can be read off independently of the other design axes. We compare the original Vicuna-7B-1.5 (Chiang et al., 2023) used in LLaVA against Qwen2-7B (Yang et al., 2024) (the LLM of Qwen2-VL-7B (Wang




Figure 4: **Base-LLM ablation: replacing LLaVA’s Vicuna-7B with alternative LLMs.** Averaged across vision encoders and pretraining configurations, switching from Vicuna to Qwen2-7B raises fine-grained performance by +7.5pp and general VQA performance by +7.7pp, confirming that a stronger base LLM lifts both axes in tandem and by comparable margins.




Figure 5: **Base vision-encoder ablation: replacing LLaVA’s CLIP ViT-L/14 with DFN-CLIP ViT-H/14.** Given sufficient pretraining, the stronger DFN-CLIP encoder improves fine-grained performance by +4.5pp while lifting general performance by only +1.2pp, isolating a vision-encoder effect that is specific to—and disproportionately concentrated in—fine-grained classification.

et al., 2024b) and Molmo-7B-D (Deitke et al., 2024)), together with Llama-2-7B (Touvron et al., 2023) (Vicuna’s base model) and Qwen2-7B-Instruct (the instruction-tuned variant of Qwen2-7B).

Findings. As Figure 4 shows, replacing Vicuna-7B with Qwen2-7B improves performance substantially and uniformly across every configuration, with average gains of +7.5 percentage points (pp) on fine-grained benchmarks and +7.4pp on general VQA benchmarks. The instruction-tuned Qwen2-7B-Instruct yields closely comparable gains (+7.5pp and +8.1pp, respectively), whereas Llama2-7B, the non-instruction-tuned base of Vicuna, produces only modest improvements whose downstream implications, absent instruction tuning, remain difficult to interpret. Crucially, all of our trained models—including the weaker Vicuna-based VLMs—emit multiple-choice responses in the correct format essentially without exception. The observed differences in fine-grained accuracy therefore reflect genuinely improved knowledge of the presented options rather than mere gains in format adherence, isolating knowledge as the operative mechanism.

Takeaway 1: Stronger language models consistently improve performance across both fine-grained classification and general VQA benchmarks.

4.2.2 VISION ENCODER CHOICE

Setup. Most of the evaluated VLMs, the LLaVA series and Molmo among them, adopt OpenAI-CLIP ViT-L/14-336px (Radford et al., 2021) as their vision encoder, whereas Qwen2-VL (Wang et al., 2024b) employs DFN-CLIP ViT-H/14-378px (Fang et al., 2023), which delivers markedly stronger zero-shot performance on object recognition tasks such as ImageNet. We contrast the two encoders across the full range of LLM choices and pretraining configurations.

Findings. Figure 5 demonstrates that the benefit of substituting CLIP L/14 with DFN-CLIP H/14 is contingent on whether connector pretraining precedes finetuning. Absent pretraining, DFN-CLIP slightly degrades general VQA performance (-2.5pp) while modestly improving fine-grained scores (+2.8pp). Once connector pretraining is in place, however, the picture changes decisively: general VQA performance holds steady or improves marginally (+1.2pp), while fine-grained performance rises substantially (+4.5pp). A stronger vision encoder thus translates into a large, targeted gain in fine-grained classification—provided the encoder is properly integrated into the model through pretraining.




Figure 6: **Pretraining-data ablation.** Relative to no pretraining, adding either LLaVA (CC-3M) or Molmo (PixMo-Cap) connector-pretraining data yields comparable gains: LLaVA data raises fine-grained and general scores by +4.0pp and +2.1pp, and PixMo data by +2.4pp and +2.5pp, respectively—so despite a large disparity in caption quality, the two sources behave almost interchangeably.




Figure 7: **Training-method ablation: connector-only pretraining versus joint LLM+connector pretraining.** Unfreezing the LLM during pretraining improves fine-grained performance by +5.5pp with no adverse effect on general VQA scores, identifying which weights are updated—not merely whether pre-training occurs—as a key lever for fine-grained classification.

Table 1: **Summary of the principal ablation settings, with existing models bolded for reference.** Upgrading the base LLM and the vision encoder, together with pretraining the connector and LLM on sufficient data, each contributes substantial gains in fine-grained performance, whereas instruction finetuning contributes comparatively little. General VQA benchmarks exhibit the complementary pattern, with the choice of LLM accounting for the largest share of the improvement.

Model/Ablation	Vision Encoder	LLM	Arch.	Pretraining Data	Finetuning Data	Fine-grained Classification	General VQA
LLaVA-1.5-7B	CLIP L/14	Vicuna	LLaVA	LLaVA	LLaVA	59.3	41.8
No pretraining	CLIP L/14	Vicuna	LLaVA	None	LLaVA	52.8	39.3
LLaVA reproduction	CLIP L/14	Vicuna	LLaVA	LLaVA	LLaVA	57.9 (+5.1)	41.2 (+1.9)
Qwen2 LLM	CLIP L/14	Qwen2	LLaVA	LLaVA	LLaVA	63.0 (+5.1)	48.3 (+7.1)
DFN-CLIP encoder	DFN H/14	Qwen2	LLaVA	LLaVA	LLaVA	68.0 (+5.0)	49.5 (+1.2)
Unfreeze LLM	DFN H/14	Qwen2	LLaVA	LLaVA	LLaVA	73.4 (+5.4)	49.9 (+0.4)
Pretrain on PixMo	CLIP L/14	Qwen2	LLaVA	PixMo	LLaVA	66.6	49.4
Molmo-7B-D	CLIP L/14	Qwen2	Molmo	PixMo	Molmo	68.4	58.0
FT Qwen2-VL base on LLaVA	DFN H/14	Qwen2	Qwen2-VL	Qwen2-VL	LLaVA	85.5	61.1
Qwen2-VL-7B	DFN H/14	Qwen2	Qwen2-VL	Qwen2-VL	Qwen2-VL	87.9 (+2.4)	62.4 (+1.3)

Takeaway 2: Better vision encoders (like DFN-CLIP vs. CLIP) in VLMs improve fine-grained classification, but only when these encoders are properly integrated through pre-training before the finetuning stage.

4.3 TRAINING METHODS

Recent VLMs are trained almost universally in two stages (Liu et al., 2023; Wang et al., 2024b; Chen et al., 2024c). The first stage, pretraining, optimizes a captioning objective on large-scale image-caption data to establish alignment between the vision encoder and the LLM. The second stage, instruction tuning, refines the aligned model on carefully curated multi-round, multimodal instruction-completion data, aligning it with human intent for real-world use. We now quantify how each of these two stages contributes to the fine-grained capabilities of VLMs.

4.3.1 PRETRAINING

Setup. We address three questions that are central to understanding the pretraining stage:

- Is a dedicated pretraining stage necessary for LLaVA-architecture models, given that their base components are already pretrained on billion-scale data?
- How does connector-only pretraining compare with jointly unfreezing the LLM?
- To what extent does the quality of the pretraining data govern downstream performance?

To answer these, we compare three pretraining configurations: LLaVA pretraining data (Liu et al., 2023) (a subset of CC-3M with web-scraped captions), Molmo pretraining data (Deitke et al., 2024) (PixMo-Cap, comprising high-quality, human-annotated, detailed captions), and no pretraining. In most configurations we pretrain only the connector, following Liu et al. (2023). We additionally examine a strategy in the spirit of Deitke et al. (2024), pretraining the connector alone for the first 20% of steps as a warmup before unfreezing the LLM for the remaining 80%. Every configuration is followed by one epoch of finetuning on the LLaVA finetuning data.

Findings. Figure 6 reports our pretraining ablation. Connector-only pretraining improves VLM performance in nearly every configuration, the sole exception being the Qwen2 and CLIP ViT-L/14 pairing. The gains are especially pronounced for Vicuna-based VLMs, which improve by an average of +6.1pp on fine-grained classification and +3.1pp on general VQA—confirming that even lightweight connector pretraining confers a substantial, fine-grained-dominant benefit.

Takeaway 3: Large-scale pretraining on image captioning datasets substantially improves fine-grained classification performance but has a more modest effect on general VLM benchmarks.

More strikingly, the quality of the pretraining captions barely registers: pretraining on low-quality web-scraped captions (LLaVA) and on highly detailed human annotations (PixMo) yield almost identical outcomes, with PixMo differing from LLaVA by only -1.6pp on fine-grained performance and +0.4pp on general VQA.

Takeaway 4: Pretraining data quality has a limited impact on overall model performance.

We hypothesize that this insensitivity arises because connector-only pretraining, which leaves the LLM frozen, prevents the model from exploiting the richer signal in higher-quality captions. To test this, we pretrain the LLM jointly with the connector, prefaced by a 20% connector-only warmup. Figure 7 confirms that this joint scheme markedly improves fine-grained classification (+5.5pp) while leaving general VQA scores intact. The improvement materializes under both LLaVA and PixMo pretraining data, however, which further establishes that training on the lower-quality captions does not degrade the LLM’s language capabilities.

Takeaway 5: Pretraining both LLM and connector substantially enhances fine-grained benchmark performance while maintaining general VQA scores.

4.3.2 FINETUNING

Setup. Departing from the preceding protocol, we begin from the pretrained Qwen2-VL-7B-Base model (Wang et al., 2024b)—the model from which Qwen2-VL-7B-Chat is derived—and finetune it on LLaVA-Instruct data (Liu et al., 2023) following our earlier experimental protocol. Comparing the result against Qwen2-VL-7B-Chat, which is instruction-finetuned on 2M in-house examples (Wang et al., 2024b), isolates the effect of the finetuning dataset while holding the base model fixed.

Findings. Finetuning on LLaVA data alone slightly lowers both fine-grained and general performance: fine-grained accuracy falls by 2.4pp (from 87.9% to 85.5%) and general accuracy by 1.3pp (from 62.4% to 61.1%). As Figure 8 makes clear, this is the smallest contribution to fine-grained performance among all our ablations, establishing that instruction finetuning plays a distinctly secondary role in fine-grained knowledge acquisition relative to base model selection and pretraining strategy.

Takeaway 6: The instruction finetuning stage has comparatively less impact on fine-grained classification performance than other factors.




Figure 8: **Cumulative effect of the ablated components.** Beginning from Vicuna-7B with CLIP ViT-L/14, we successively upgrade the base LLM, the vision encoder, and the training settings, raising fine-grained classification accuracy from 52.8% to 73.4%—closing more than half of the gap to Qwen2-VL-Chat’s 87.9%. The residual difference is attributable primarily to Qwen2-VL’s far larger pretraining corpus of 1.4T tokens.

4.4 SUMMARY OF ABLATIONS

Figure 8 and Table 1 synthesize our ablations, tracing how each component contributes to closing the gap between the weakest and strongest VLMs on both fine-grained and general VQA tasks. Starting from the baseline LLaVA architecture (Liu et al., 2023) (Vicuna-7B (Chiang et al., 2023) with CLIP ViT-L/14 (Radford et al., 2021)), the systematic upgrades we introduce raise fine-grained performance from 52.8% to 73.4% and general VQA performance from 39.3% to 49.9%. The decomposition is itself informative: switching to DFN-CLIP (Fang et al., 2023) and unfreezing the LLM during pretraining disproportionately benefit fine-grained capabilities relative to general ability, whereas switching the LLM to Qwen2-7B (Yang et al., 2024) accounts for the majority of the gain in general VQA. The two axes of improvement are thus driven by largely distinct components.

Even after these substantial gains, a 12-point gap in fine-grained classification relative to Qwen2-VL-Chat remains unexplained by our ablations (Table 1). Two candidate factors could account for it: architectural differences and pretraining data scale.

Although architectural differences could in principle contribute, our observational results (Section 3.4) together with our training ablations (Section 4.3) point consistently to pretraining data scale as the dominant factor. Our LLaVA (Liu et al., 2023) and PixMo (Deitke et al., 2024) experiments used comparatively small corpora—each fewer than 1M image-caption pairs, or roughly 200M and 400M tokens, respectively—whereas Wang et al. (2024b) pretrain Qwen2-VL on 1.4T tokens, orders of magnitude more. This disparity in pretraining scale is the most plausible explanation for the residual gap, and it underscores that extensive pretraining is the critical ingredient for developing VLMs with superior fine-grained classification capabilities.

5 LIMITATIONS & CONCLUSION

Our study carries practical limitations that, in turn, delineate concrete opportunities for future work. Under our computational budget we could compare training only at the <1M data-point scale rather than the billion-scale (B) regime employed by some newer VLMs, so the question of how far our conclusions extend to that regime remains open—though the consistent direction of our data-scale evidence suggests they should strengthen, not reverse, with more data. Newer work additionally proposes alternative pretraining strategies whose interaction with fine-grained visual understanding is a natural next target for analysis.

In this work, we systematically evaluate state-of-the-art vision-language models (VLMs) on fine-grained classification benchmarks and establish fine-grained visual classification as a crucial yet underexplored dimension of VLM capability. Through a controlled analysis of the principal model components and training paradigms, we identify which choices govern fine-grained classification and vision-centric competence, yielding concrete, actionable strategies for improving them—and

540 thereby for strengthening the applicability of VLMs in the real-world scenarios that demand precise
541 visual understanding.
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593

REFERENCES

- Marah Abdin, Jyoti Aneja, Hany Awadalla, Ahmed Awadallah, Ammar Ahmad Awan, Nguyen Bach, Amit Bahree, Arash Bakhtiari, Jianmin Bao, Harkirat Behl, et al. Phi-3 technical report: A highly capable language model locally on your phone. *arXiv preprint arXiv:2404.14219*, 2024. [2](#), [5](#)
- Jean-Baptiste Alayrac, Jeff Donahue, Pauline Luc, Antoine Miech, Iain Barr, Yana Hasson, Karel Lenc, Arthur Mensch, Katherine Millican, Malcolm Reynolds, et al. Flamingo: a visual language model for few-shot learning. In *NeurIPS*, 2022. [1](#)
- Anas Awadalla, Irena Gao, Josh Gardner, Jack Hessel, Yusuf Hanafy, Wanrong Zhu, Kalyani Marathe, Yonatan Bitton, Samir Gadre, Shiori Sagawa, et al. Openflamingo: An open-source framework for training large autoregressive vision-language models. *arXiv preprint arXiv:2308.01390*, 2023. [3](#)
- Lukas Bossard, Matthieu Guillaumin, and Luc Van Gool. Food-101 – mining discriminative components with random forests. In *European Conference on Computer Vision*, 2014. [2](#), [4](#), [16](#)
- Mathilde Caron, Hugo Touvron, Ishan Misra, Hervé Jégou, Julien Mairal, Piotr Bojanowski, and Armand Joulin. Emerging properties in self-supervised vision transformers. In *Proceedings of the IEEE/CVF international conference on computer vision*, pp. 9650–9660, 2021. [1](#), [3](#), [4](#)
- Lin Chen, Jinsong Li, Xiaoyi Dong, Pan Zhang, Conghui He, Jiaqi Wang, Feng Zhao, and Dahua Lin. Sharegpt4v: Improving large multi-modal models with better captions. In *European Conference on Computer Vision*, pp. 370–387. Springer, 2024a. [5](#)
- Lin Chen, Jinsong Li, Xiaoyi Dong, Pan Zhang, Yuhang Zang, Zehui Chen, Haodong Duan, Jiaqi Wang, Yu Qiao, Dahua Lin, et al. Are we on the right way for evaluating large vision-language models? In *NeurIPS*, 2024b. [6](#), [17](#)
- Zhe Chen, Jiannan Wu, Wenhai Wang, Weijie Su, Guo Chen, Sen Xing, Muyan Zhong, Qinglong Zhang, Xizhou Zhu, Lewei Lu, et al. Internvl: Scaling up vision foundation models and aligning for generic visual-linguistic tasks. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 24185–24198, 2024c. [5](#), [8](#)
- Wei-Lin Chiang, Zhuohan Li, Zi Lin, Ying Sheng, Zhanghao Wu, Hao Zhang, Lianmin Zheng, Siyuan Zhuang, Yonghao Zhuang, Joseph E Gonzalez, et al. Vicuna: An open-source chatbot impressing gpt-4 with 90%* chatgpt quality. *Technical Report*, 2023. [6](#), [10](#)
- Wenliang Dai, Junnan Li, Dongxu Li, Anthony Meng Huat Tiong, Junqi Zhao, Weisheng Wang, Boyang Li, Pascale N Fung, and Steven Hoi. Instructblip: Towards general-purpose vision-language models with instruction tuning. In *NeurIPS*, 2023. [3](#)
- Matt Deitke, Christopher Clark, Sangho Lee, Rohun Tripathi, Yue Yang, Jae Sung Park, Mohammadreza Salehi, Niklas Muennighoff, Kyle Lo, Luca Soldaini, et al. Molmo and pixmo: Open weights and open data for state-of-the-art multimodal models. *arXiv preprint arXiv:2409.17146*, 2024. [2](#), [4](#), [5](#), [6](#), [7](#), [9](#), [10](#), [18](#)
- Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. Imagenet: A large-scale hierarchical image database. In *CVPR*, 2009. [2](#), [4](#), [16](#)
- Haodong Duan, Junming Yang, Yuxuan Qiao, Xinyu Fang, Lin Chen, Yuan Liu, Xiaoyi Dong, Yuhang Zang, Pan Zhang, Jiaqi Wang, et al. Vlmevalkit: An open-source toolkit for evaluating large multi-modality models. In *Proceedings of the 32nd ACM international conference on multimedia*, pp. 11198–11201, 2024. [5](#), [17](#)
- Alex Fang, Albin Madappally Jose, Amit Jain, Ludwig Schmidt, Alexander Toshev, and Vaishaal Shankar. Data filtering networks. *arXiv preprint arXiv:2309.17425*, 2023. [5](#), [6](#), [7](#), [10](#)
- Samir Yitzhak Gadre, Gabriel Ilharco, Alex Fang, Jonathan Hayase, Georgios Smyrnis, Thao Nguyen, Ryan Marten, Mitchell Wortsman, Dhruva Ghosh, Jieyu Zhang, et al. Datacomp: In search of the next generation of multimodal datasets. In *NeurIPS*, 2023. [4](#)
- Gregor Geigle, Radu Timofte, and Goran Glavaš. African or european swallow? benchmarking large vision-language models for fine-grained object classification. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 2653–2669, 2024. [2](#), [3](#), [4](#), [16](#)
- Yash Goyal, Tejas Khot, Douglas Summers-Stay, Dhruv Batra, and Devi Parikh. Making the v in vqa matter: Elevating the role of image understanding in visual question answering. In *CVPR*, 2017. [1](#)

- Tianrui Guan, Fuxiao Liu, Xiyang Wu, Ruiqi Xian, Zongxia Li, Xiaoyu Liu, Xijun Wang, Lichang Chen, Furong Huang, Yaser Yacoob, Dinesh Manocha, and Tianyi Zhou. Hallusionbench: An advanced diagnostic suite for entangled language hallucination and visual illusion in large vision-language models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 14375–14385, June 2024. 17
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *CVPR*, 2016. 3
- Wenyi Hong, Weihang Wang, Ming Ding, Wenmeng Yu, Qingsong Lv, Yan Wang, Yean Cheng, Shiyu Huang, Junhui Ji, Zhao Xue, et al. Cogvlm2: Visual language models for image and video understanding. *arXiv preprint arXiv:2408.16500*, 2024. 5
- Siddharth Karamcheti, Suraj Nair, Ashwin Balakrishna, Percy Liang, Thomas Kollar, and Dorsa Sadigh. Prismatic vlms: Investigating the design space of visually-conditioned language models. *arXiv preprint arXiv:2402.07865*, 2024. 4
- Aniruddha Kembhavi, Mike Salvato, Eric Kolve, Minjoon Seo, Hannaneh Hajishirzi, and Ali Farhadi. A diagram is worth a dozen images. In *ECCV*, 2016. 17
- Hugo Laurençon, Léo Tronchon, Matthieu Cord, and Victor Sanh. What matters when building vision-language models? *arXiv preprint arXiv:2405.02246*, 2024. 5
- Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. In *NeurIPS*, 2023. 1, 2, 3, 4, 5, 6, 8, 9, 10, 17
- Haotian Liu, Chunyuan Li, Yuheng Li, Bo Li, Yuanhan Zhang, Sheng Shen, and Yong Jae Lee. Llava-next: Improved reasoning, ocr, and world knowledge, January 2024a. URL <https://llava-vl.github.io/blog/2024-01-30-llava-next/>. 2, 5
- Yuan Liu, Haodong Duan, Yuanhan Zhang, Bo Li, Songyang Zhang, Wangbo Zhao, Yike Yuan, Jiaqi Wang, Conghui He, Ziwei Liu, et al. Mmbench: Is your multi-modal model an all-around player? In *European conference on computer vision*, pp. 216–233. Springer, 2024b. 6, 17
- Yuliang Liu, Zhang Li, Mingxin Huang, Biao Yang, Wenwen Yu, Chunyuan Li, Xu-Cheng Yin, Cheng-Lin Liu, Lianwen Jin, and Xiang Bai. Ocrbench: on the hidden mystery of ocr in large multimodal models. *Science China Information Sciences*, 67(12), December 2024c. ISSN 1869-1919. doi: 10.1007/s11432-024-4235-6. URL <http://dx.doi.org/10.1007/s11432-024-4235-6>. 17
- Pan Lu, Hritik Bansal, Tony Xia, Jiacheng Liu, Chunyuan Li, Hannaneh Hajishirzi, Hao Cheng, Kai-Wei Chang, Michel Galley, and Jianfeng Gao. Mathvista: Evaluating mathematical reasoning of foundation models in visual contexts. In *ICLR*, 2024. 2, 4, 5, 17
- Minesh Mathew, Dimosthenis Karatzas, and CV Jawahar. Docvqa: A dataset for vqa on document images. In *WACV*, 2021. 1, 3, 4
- Brandon McKinzie, Zhe Gan, Jean-Philippe Fauconnier, Sam Dodge, Bowen Zhang, Philipp Dufter, Dhruvi Shah, Xianzhi Du, Futang Peng, Floris Weers, et al. Mml1: Methods, analysis & insights from multimodal llm pre-training. *arXiv preprint arXiv:2403.09611*, 2024. 4
- M-E. Nilsback and A. Zisserman. Automated flower classification over a large number of classes. In *Proceedings of the Indian Conference on Computer Vision, Graphics and Image Processing*, 2008. 2, 3, 4, 16
- OpenAI. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023. 1
- O. M. Parkhi, A. Vedaldi, A. Zisserman, and C. V. Jawahar. Cats and dogs. In *IEEE Conference on Computer Vision and Pattern Recognition*, 2012. 2, 3, 4, 16
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *ICML*, 2021. 1, 3, 4, 5, 6, 7, 10
- Benjamin Recht, Rebecca Roelofs, Ludwig Schmidt, and Vaishaal Shankar. Do imagenet classifiers generalize to imagenet? In *ICML*, 2019. 16

- Christoph Schuhmann, Romain Beaumont, Richard Vencu, Cade W Gordon, Ross Wightman, Mehdi Cherti, Theo Coombes, Aarush Katta, Clayton Mullis, Mitchell Wortsman, Patrick Schramowski, Srivatsa R Kundurthy, Katherine Crowson, Ludwig Schmidt, Robert Kaczmarczyk, and Jenia Jitsev. LAION-5b: An open large-scale dataset for training next generation image-text models. In *Thirty-sixth Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2022. URL <https://openreview.net/forum?id=M3Y74vmsMcY>. 4, 5, 16
- Amanpreet Singh, Vivek Natarjan, Meet Shah, Yu Jiang, Xinlei Chen, Devi Parikh, and Marcus Rohrbach. Towards vqa models that can read. In *CVPR*, 2019. 4
- Gemini Team, Rohan Anil, Sebastian Borgeaud, Yonghui Wu, Jean-Baptiste Alayrac, Jiahui Yu, Radu Soricut, Johan Schalkwyk, Andrew M Dai, Anja Hauth, et al. Gemini: a family of highly capable multimodal models. *arXiv preprint arXiv:2312.11805*, 2023. 1
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023. 7
- Vishaal Udandara, Ameya Prabhu, Adhiraj Ghosh, Yash Sharma, Philip HS Torr, Adel Bibi, Samuel Albanie, and Matthias Bethge. No” zero-shot” without exponential data: Pretraining concept frequency determines multimodal model performance. *arXiv preprint arXiv:2404.04125*, 2024. 4
- C. Wah, S. Branson, P. Welinder, P. Perona, and S. Belongie. Caltech-ucsd birds 200. Technical report, 2011. 3
- Ke Wang, Juntong Pan, Weikang Shi, Zimu Lu, Mingjie Zhan, and Hongsheng Li. Measuring multimodal mathematical reasoning with math-vision dataset. *arXiv preprint arXiv:2402.14804*, 2024a. 3
- Peng Wang, Shuai Bai, Sinan Tan, Shijie Wang, Zhihao Fan, Jinze Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, et al. Qwen2-vl: Enhancing vision-language model’s perception of the world at any resolution. *arXiv preprint arXiv:2409.12191*, 2024b. 1, 2, 3, 5, 6, 7, 8, 9, 10
- Xiu-Shen Wei, Yi-Zhe Song, Oisín Mac Aodha, Jianxin Wu, Yuxin Peng, Jinhui Tang, Jian Yang, and Serge Belongie. Fine-grained image analysis with deep learning: A survey. *IEEE transactions on pattern analysis and machine intelligence*, 44(12):8927–8948, 2021. 4
- An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2. 5 technical report. *arXiv preprint arXiv:2412.15115*, 2024. 6, 10
- Weihao Yu, Zhengyuan Yang, Linjie Li, Jianfeng Wang, Kevin Lin, Zicheng Liu, Xinchao Wang, and Lijuan Wang. Mm-vet: Evaluating large multimodal models for integrated capabilities. In *ICML*, 2024. 2, 17
- Xiang Yue, Yuansheng Ni, Kai Zhang, Tianyu Zheng, Ruoqi Liu, Ge Zhang, Samuel Stevens, Dongfu Jiang, Weiming Ren, Yuxuan Sun, et al. Mmmu: A massive multi-discipline multimodal understanding and reasoning benchmark for expert agi. In *CVPR*, 2024. 1, 2, 3, 4, 5, 6, 17
- Xiaohua Zhai, Basil Mustafa, Alexander Kolesnikov, and Lucas Beyer. Sigmoid loss for language image pre-training. In *Proceedings of the IEEE/CVF international conference on computer vision*, pp. 11975–11986, 2023. 5
- Yuhui Zhang, Alyssa Unell, Xiaohan Wang, Dhruva Ghosh, Yuchang Su, Ludwig Schmidt, and Serena Yeung-Levy. Why are visually-grounded language models bad at image classification? In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024. 2, 3

Vision Encoder	LLM	Arch	PT Data	FT Data	Updating	Classification	VLM
CLIP ViT-L/14@336	Vicuna-7B	LLaVA	LLaVA	LLaVA	Connector	59.3	41.8
CLIP ViT-L/14@336	Vicuna-7B	LLaVA	LLaVA	LLaVA	Connector	59.0	41.2
CLIP ViT-L/14@336	Vicuna-7B	LLaVA	None	LLaVA	Connector	52.8	39.3
CLIP ViT-L/14@336	Llama2-7B	LLaVA	None	LLaVA	Connector	52.5	41.3
DFN ViT-H/14@378	Llama2-7B	LLaVA	None	LLaVA	Connector	56.2	40.0
DFN ViT-H/14@378	Vicuna-7B	LLaVA	None	LLaVA	Connector	52.9	38.2
CLIP ViT-L/14@336	Qwen2-7B	LLaVA	None	LLaVA	Connector	61.4	48.6
DFN ViT-H/14@378	Qwen2-7B	LLaVA	None	LLaVA	Connector	66.8	46.0
CLIP ViT-L/14@336	Qwen2-7B-Instruct	LLaVA	None	LLaVA	Connector	59.7	48.4
DFN ViT-H/14@378	Qwen2-7B-Instruct	LLaVA	None	LLaVA	Connector	61.7	43.4
DFN ViT-H/14@378	Qwen2-7B	Qwen2-VL	Qwen2-VL	LLaVA	Full	85.5	61.1
CLIP ViT-L/14@336	Vicuna-7B	LLaVA	LLaVA	LLaVA	Connector	57.9	41.2
DFN ViT-H/14@378	Vicuna-7B	LLaVA	LLaVA	LLaVA	Connector	61.2	41.4
CLIP ViT-L/14@336	Qwen2-7B	LLaVA	LLaVA	LLaVA	Connector	63.0	48.3
DFN ViT-H/14@378	Qwen2-7B	LLaVA	LLaVA	LLaVA	Connector	68.0	49.5
CLIP ViT-L/14@336	Qwen2-7B-Instruct	LLaVA	LLaVA	LLaVA	Connector	63.2	49.9
DFN ViT-H/14@378	Qwen2-7B-Instruct	LLaVA	LLaVA	LLaVA	Connector	70.2	50.9
DFN ViT-H/14@378	Qwen2-7B	Qwen2-VL	Qwen2-VL	Qwen2-VL	Full	87.9	62.4
CLIP ViT-L/14@336	Vicuna-7B	LLaVA	PixMo	LLaVA	Connector	57.8	41.6
DFN ViT-H/14@378	Vicuna-7B	LLaVA	PixMo	LLaVA	Connector	58.8	43.4
CLIP ViT-L/14@336	Qwen2-7B	LLaVA	PixMo	LLaVA	Connector	60.5	47.6
DFN ViT-H/14@378	Qwen2-7B	LLaVA	PixMo	LLaVA	Connector	66.6	49.4
DFN ViT-H/14@378	Qwen2-7B	LLaVA	LLaVA	LLaVA	Full	73.4	49.9
DFN ViT-H/14@378	Qwen2-7B	LLaVA	PixMo	LLaVA	Full	72.3	49.4

Table 2: **Detailed model-ablation results: aggregate fine-grained classification and general VLM scores.** Each row reports one configuration of vision encoder, LLM, architecture, pretraining data, finetuning data, and weight-updating scheme, alongside its averaged fine-grained classification and VLM accuracy.

Vision Encoder	LLM	Arch	PT Data	FT Data	Updating	MMBench	MMMU	MMStar
CLIP ViT-L/14@336	Vicuna-7B	LLaVA	LLaVA	LLaVA	Connector	60.4	32.2	32.7
CLIP ViT-L/14@336	Vicuna-7B	LLaVA	LLaVA	LLaVA	Connector	58.6	32.9	32.3
CLIP ViT-L/14@336	Vicuna-7B	LLaVA	None	LLaVA	Connector	54.4	32.8	30.7
CLIP ViT-L/14@336	Llama2-7B	LLaVA	None	LLaVA	Connector	58.0	34.3	31.7
DFN ViT-H/14@378	Llama2-7B	LLaVA	None	LLaVA	Connector	53.9	34.6	31.5
DFN ViT-H/14@378	Vicuna-7B	LLaVA	None	LLaVA	Connector	51.3	33.2	30.1
CLIP ViT-L/14@336	Qwen2-7B	LLaVA	None	LLaVA	Connector	63.1	41.8	40.9
DFN ViT-H/14@378	Qwen2-7B	LLaVA	None	LLaVA	Connector	61.3	36.2	40.5
CLIP ViT-L/14@336	Qwen2-7B-Instruct	LLaVA	None	LLaVA	Connector	63.2	41.4	40.5
DFN ViT-H/14@378	Qwen2-7B-Instruct	LLaVA	None	LLaVA	Connector	55.7	38.3	36.1
DFN ViT-H/14@378	Qwen2-7B	Qwen2-VL	Qwen2-VL	LLaVA	Full	78.5	49.0	55.8
CLIP ViT-L/14@336	Vicuna-7B	LLaVA	LLaVA	LLaVA	Connector	60.4	29.9	33.3
DFN ViT-H/14@378	Vicuna-7B	LLaVA	LLaVA	LLaVA	Connector	57.4	34.2	32.7
CLIP ViT-L/14@336	Qwen2-7B	LLaVA	LLaVA	LLaVA	Connector	64.4	41.3	39.1
DFN ViT-H/14@378	Qwen2-7B	LLaVA	LLaVA	LLaVA	Connector	64.9	41.3	42.3
CLIP ViT-L/14@336	Qwen2-7B-Instruct	LLaVA	LLaVA	LLaVA	Connector	66.5	41.3	41.9
DFN ViT-H/14@378	Qwen2-7B-Instruct	LLaVA	LLaVA	LLaVA	Connector	68.0	41.6	43.0
DFN ViT-H/14@378	Qwen2-7B	Qwen2-VL	Qwen2-VL	Qwen2-VL	Full	78.9	49.9	58.4
CLIP ViT-L/14@336	Vicuna-7B	LLaVA	PixMo	LLaVA	Connector	57.8	32.6	34.3
DFN ViT-H/14@378	Vicuna-7B	LLaVA	PixMo	LLaVA	Connector	61.1	34.1	34.9
CLIP ViT-L/14@336	Qwen2-7B	LLaVA	PixMo	LLaVA	Connector	64.9	40.0	38.0
DFN ViT-H/14@378	Qwen2-7B	LLaVA	PixMo	LLaVA	Connector	63.0	41.0	44.3
DFN ViT-H/14@378	Qwen2-7B	LLaVA	LLaVA	LLaVA	Full	66.0	42.0	41.6
DFN ViT-H/14@378	Qwen2-7B	LLaVA	PixMo	LLaVA	Full	63.7	40.0	44.6

Table 3: **Detailed model-ablation results on the individual general VQA benchmarks.** Per-benchmark accuracy on MMBench, MMMU, and MMStar for each ablated configuration.

A FULL RESULTS

We report the complete ablation results—aggregated performance, per-benchmark general VQA scores, and per-benchmark fine-grained classification scores—in Table 2, Table 3, and Table 4, respectively.

Vision Encoder	LLM	Arch	PT Data	FT Data	Updating	ImageNet	Flowers	Pets	Food
CLIP ViT-L/14@336	Vicuna-7B	LLaVA	LLaVA	LLaVA	Connector	70.9	50.7	46.3	69.4
CLIP ViT-L/14@336	Vicuna-7B	LLaVA	LLaVA	LLaVA	Connector	68.7	50.4	46.5	70.3
CLIP ViT-L/14@336	Vicuna-7B	LLaVA	None	LLaVA	Connector	64.3	46.1	40.4	60.4
CLIP ViT-L/14@336	Llama2-7B	LLaVA	None	LLaVA	Connector	64.4	44.2	40.7	60.8
DFN ViT-H/14@378	Llama2-7B	LLaVA	None	LLaVA	Connector	65.4	50.5	46.2	62.8
DFN ViT-H/14@378	Vicuna-7B	LLaVA	None	LLaVA	Connector	62.3	48.7	39.2	61.3
CLIP ViT-L/14@336	Qwen2-7B	LLaVA	None	LLaVA	Connector	70.6	44.1	55.6	75.2
DFN ViT-H/14@378	Qwen2-7B	LLaVA	None	LLaVA	Connector	74.1	50.8	62.3	80.1
CLIP ViT-L/14@336	Qwen2-7B-Instruct	LLaVA	None	LLaVA	Connector	71.4	39.6	53.1	74.7
DFN ViT-H/14@378	Qwen2-7B-Instruct	LLaVA	None	LLaVA	Connector	72.0	46.0	54.1	74.5
DFN ViT-H/14@378	Qwen2-7B	Qwen2-VL	Qwen2-VL	LLaVA	Full	84.8	79.5	87.1	90.7
CLIP ViT-L/14@336	Vicuna-7B	LLaVA	LLaVA	LLaVA	Connector	69.1	50.0	45.3	67.1
DFN ViT-H/14@378	Vicuna-7B	LLaVA	LLaVA	LLaVA	Connector	72.7	53.2	48.5	70.4
CLIP ViT-L/14@336	Qwen2-7B	LLaVA	LLaVA	LLaVA	Connector	71.5	46.9	56.6	76.9
DFN ViT-H/14@378	Qwen2-7B	LLaVA	LLaVA	LLaVA	Connector	76.1	54.1	59.1	82.7
CLIP ViT-L/14@336	Qwen2-7B-Instruct	LLaVA	LLaVA	LLaVA	Connector	74.5	45.8	55.1	77.4
DFN ViT-H/14@378	Qwen2-7B-Instruct	LLaVA	LLaVA	LLaVA	Connector	78.5	52.3	65.4	84.7
DFN ViT-H/14@378	Qwen2-7B	Qwen2-VL	Qwen2-VL	Qwen2-VL	Full	85.9	82.3	91.0	92.3
CLIP ViT-L/14@336	Vicuna-7B	LLaVA	PixMo	LLaVA	Connector	69.2	47.4	46.7	67.8
DFN ViT-H/14@378	Vicuna-7B	LLaVA	PixMo	LLaVA	Connector	70.8	50.9	44.8	68.5
CLIP ViT-L/14@336	Qwen2-7B	LLaVA	PixMo	LLaVA	Connector	70.8	43.7	53.9	73.6
DFN ViT-H/14@378	Qwen2-7B	LLaVA	PixMo	LLaVA	Connector	76.1	50.3	57.8	82.1
DFN ViT-H/14@378	Qwen2-7B	LLaVA	LLaVA	LLaVA	Full	79.4	60.7	67.3	86.0
DFN ViT-H/14@378	Qwen2-7B	LLaVA	PixMo	LLaVA	Full	80.0	56.4	65.4	87.4

Table 4: **Detailed model-ablation results on the individual fine-grained classification benchmarks.** Per-benchmark accuracy on ImageNet, Flowers, Pets, and Food for each ablated configuration.

B EVALUATION

B.1 FINE-GRAINED CLASSIFICATION

As introduced in Section 3.2, we evaluate VLMs and their base vision encoders on four traditional classification benchmarks, detailed below.

ImageNet-1K (Deng et al., 2009): A foundational image classification dataset organized by the WordNet hierarchy, spanning a broad range of supercategories that each subsume multiple fine-grained subcategories. Its test set comprises 50,000 images distributed over 1,000 classes.

The structure of the ImageNet/WordNet hierarchy, however, leaves many classes too visually distinct from the remainder to admit hard multiple-choice distractors—the “person” supercategory, for instance, contains only three classes. To circumvent this, and the dataset’s unwieldy size, we draw on the human-annotated incorrect model predictions of Recht et al. (2019) to select images for which hard negative answer choices can be constructed. Beginning from 19,056 annotated examples, we distill a pool of 5,457 images covering 928 of the original 1,000 classes. Because this curation leaves the surviving classes unevenly balanced, we report mean-per-class accuracy for ImageNet.

Oxford Flowers-102 (Nilsback & Zisserman, 2008): A dataset of 102 flower species marked by high intra-class variation and pronounced inter-class similarity, with 6,149 test examples.

Oxford-IIIT Pet-37 (Parkhi et al., 2012): A collection of 37 pet categories exhibiting challenging variation in pose, lighting, and occlusion, with 3,669 test examples.

Food-101 (Bossard et al., 2014): A dataset of 101 food categories displaying substantial visual diversity within each class, with 25,250 test examples.

Question generation. For every fine-grained classification test set other than ImageNet (treated above), we construct hard negatives by running zero-shot classification on the test examples with an OpenCLIP ViT-L/14 trained on LAION (Schuhmann et al., 2022), following Geigle et al. (2024). We take the four highest-ranked incorrect predictions and shuffle them together with the correct label to form the answer set.

Prompt formatting. In our multiple-choice VQA format, the question is invariably “What type of object is in this photo?”, and the answer choices are rendered as “A. Option 1\n B. Option 2\n C. Option 3\n D. Option 4\n E. Option 5”. Because different VLMs are trained with different multiple-choice prompt conventions, we adopt each model’s corresponding format from VLMEvalKit (Duan et al., 2024)—for example, “Answer with the option’s letter from the given choices directly.”

For CLIP models we apply the standard zero-shot evaluation procedure, which involves no question format; the class names supplied to the text encoder are instead instantiated in the template “a photo of a {classname}”.

B.2 GENERAL VQA

For the initial observational study of Section 3, we rely directly on the scores reported by the Open-VLM Leaderboard (Duan et al., 2024). By default, that leaderboard reports the average over eight diverse VLM benchmarks: MMBench (Liu et al., 2024b), MMStar (Chen et al., 2024b), MMMU (Val) (Yue et al., 2024), MathVista (Lu et al., 2024), OCRBench (Liu et al., 2024c), AI2D (Kembhavi et al., 2016), HallusionBench (Guan et al., 2024), and MMVet (Yu et al., 2024). This aggregate is used only in Figure 2.

For every subsequent evaluation—of both existing models and the models we train in Section 4—we restrict attention to a subset of multiple-choice VLM benchmarks that permits a more direct, controlled comparison. Specifically, we evaluate on three benchmarks:

MMBench (Liu et al., 2024b): A dataset of 2,948 examples spanning a broad range of VLM capabilities partitioned between perception and reasoning. Most examples target coarse-grained perception and visual reasoning; a subset labeled “fine-grained perception” does exist, but its categories (action recognition, attribute recognition, OCR, and the like) lie in a domain distinct from the fine-grained classification datasets we study.

MMMU (Val) (Yue et al., 2024): A large-scale multimodal dataset spanning a wide distribution of domains and applications. Although the full test set holds 10,500 examples, the field conventionally reports performance on the 900-question validation split, which we adopt. MMMU probes general knowledge tied to visual perception across art, science, engineering, medicine, business, and the humanities.

MMStar (Chen et al., 2024b): A dataset of 1,500 human-selected questions, curated to span a diverse set of tasks across both domain and low-level model capabilities.

Evaluation procedure. For the general VQA benchmarks we defer entirely to the VLMEvalKit evaluation code (Duan et al., 2024). Each existing model carries a predefined evaluation procedure; for our own trained models we apply the LLaVA evaluation code, since we finetune exclusively on LLaVA-Instruct data.

C TRAINING

All experiments build on the LLaVA-1.5 training codebase (Liu et al., 2023), and, except where an ablation explicitly dictates otherwise, adhere to the LLaVA architecture and training procedure.

Pretraining. Concretely, we perform one epoch of pretraining on either the LLaVA pretraining data—image–caption pairs drawn from CC-3M—or the Molmo pretraining data (PixMo-Cap); for PixMo-Cap we use only the “caption” field and discard the raw audio transcript. Pretraining uses a batch size of 256. In most runs we update only the randomly initialized MLP connector, at a learning rate of $1e-3$. When ablating the pretraining strategy, we instead warm up the connector alone for the first 20% of steps and then unfreeze and jointly update the LLM at a learning rate of $2e-5$ for the remaining 80%.

Finetuning. Finetuning always consists of one epoch on the LLaVA-1.5 finetuning data—LLaVA-Instruct together with VQA training data drawn from several datasets to support proper benchmarking. Throughout finetuning we hold the LLM learning rate at $2e-5$ with a batch size of 128, and the vision encoder remains frozen across both training phases.

For the Qwen2-VL-architecture ablation, we initialize from Qwen2-VL-7B Base and train under the same procedure, unfreezing only the LLM.

Compute. All models are trained on 4 nodes of 4 A100 (40GB) GPUs each; a run typically completes within 20 hours, with the precise duration varying with the CLIP model size.

LLaVA vs. PixMo data. The main text repeatedly invokes the disparity in caption quality between the LLaVA and PixMo pretraining data, which we make explicit here. The LLaVA data derives from CC-3M and consists of web-scraped captions notorious for low quality and weak image relevance. PixMo-Cap, by contrast, is collected by having human annotators record a 60–90 second audio description of each image, which is subsequently consolidated into a coherent caption (Deitke et al., 2024). A surface-level comparison quantifies the gap: LLaVA captions average 9.8 words, whereas PixMo captions average 169 words.

LLaVA caption example: “front panel bracket cover for suzuki”

PixMo caption example: “The poster from the TV show ”Law and Order” features a past cast ensemble set in an interrogation room. On the left is a Hispanic man, identified as Detective Nick Amaro, portrayed by Danny Pino. He has dark hair, olive skin, and is dressed in a gray jacket, blue shirt, and black tie. Standing beside him is a blonde woman with shoulder-length hair, who is identified as Detective Amanda Rollins, played by Kelli Giddish. She is ... The backdrop is an interrogation room, with a concrete wall and a portion of a two-way mirror visible, adding to the procedural drama’s atmospheric setting.”